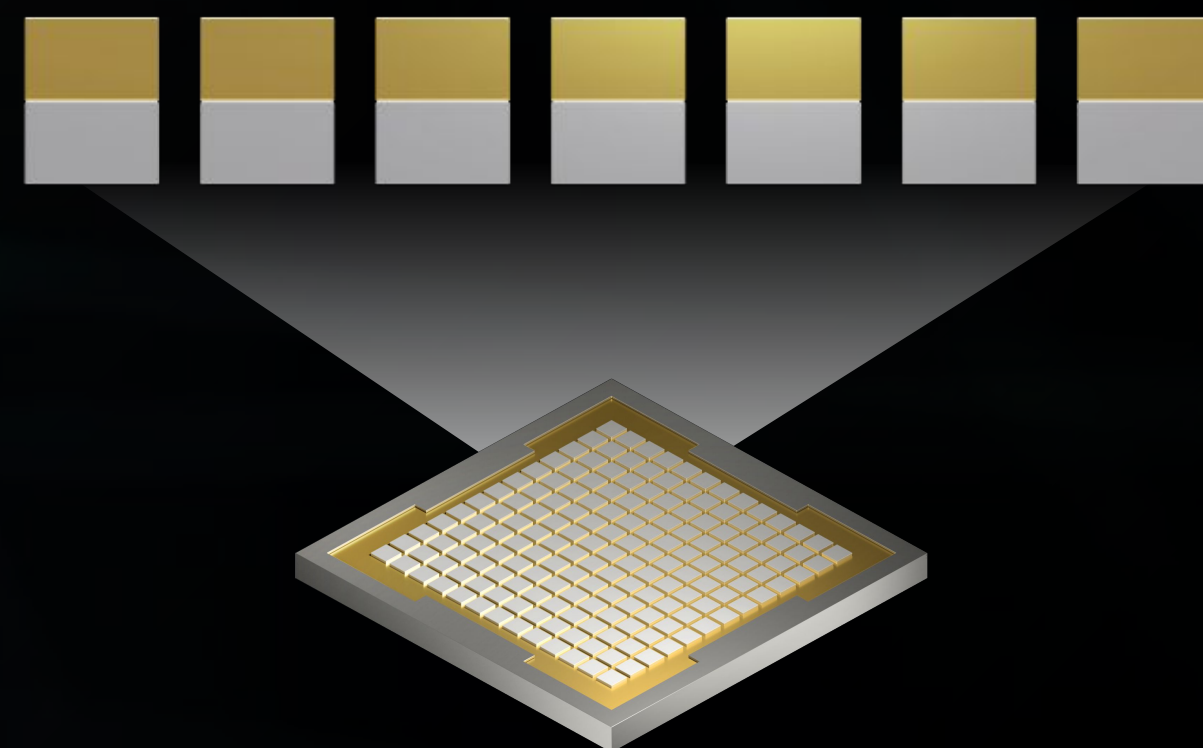


Virtual

MULTI-INSTANCE GPU_s IN CONTAINERS AND KUBERNETES

Kevin Klues
kklues@nvidia.com

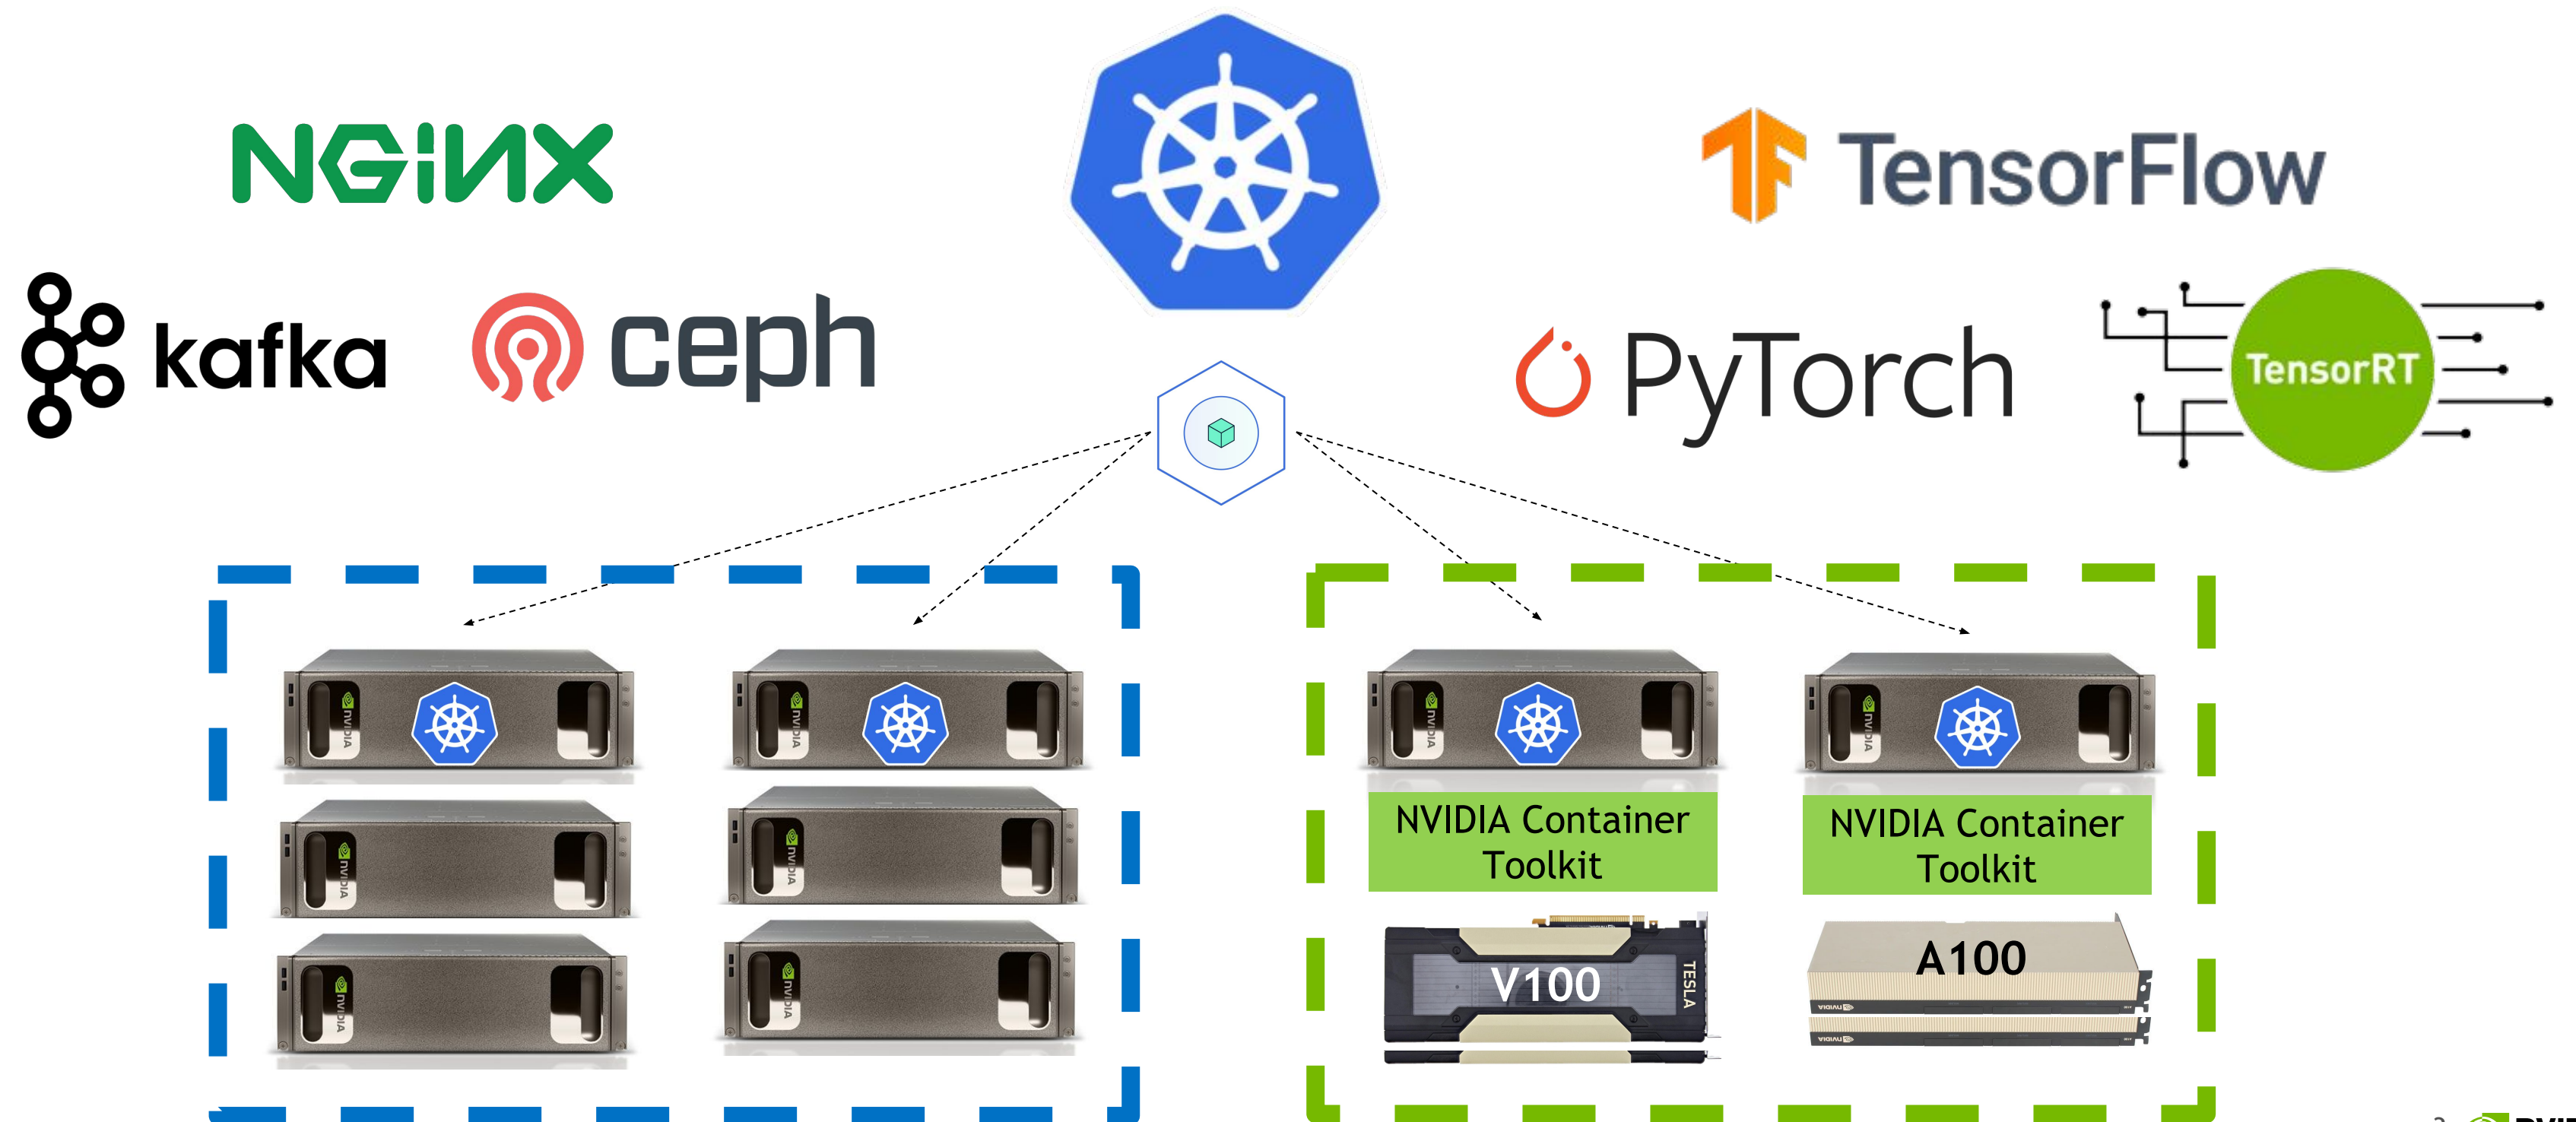


GPUs AND KUBERNETES

Seamlessly scale up training and inference to a cluster of GPU machines

GPUs AND KUBERNETES

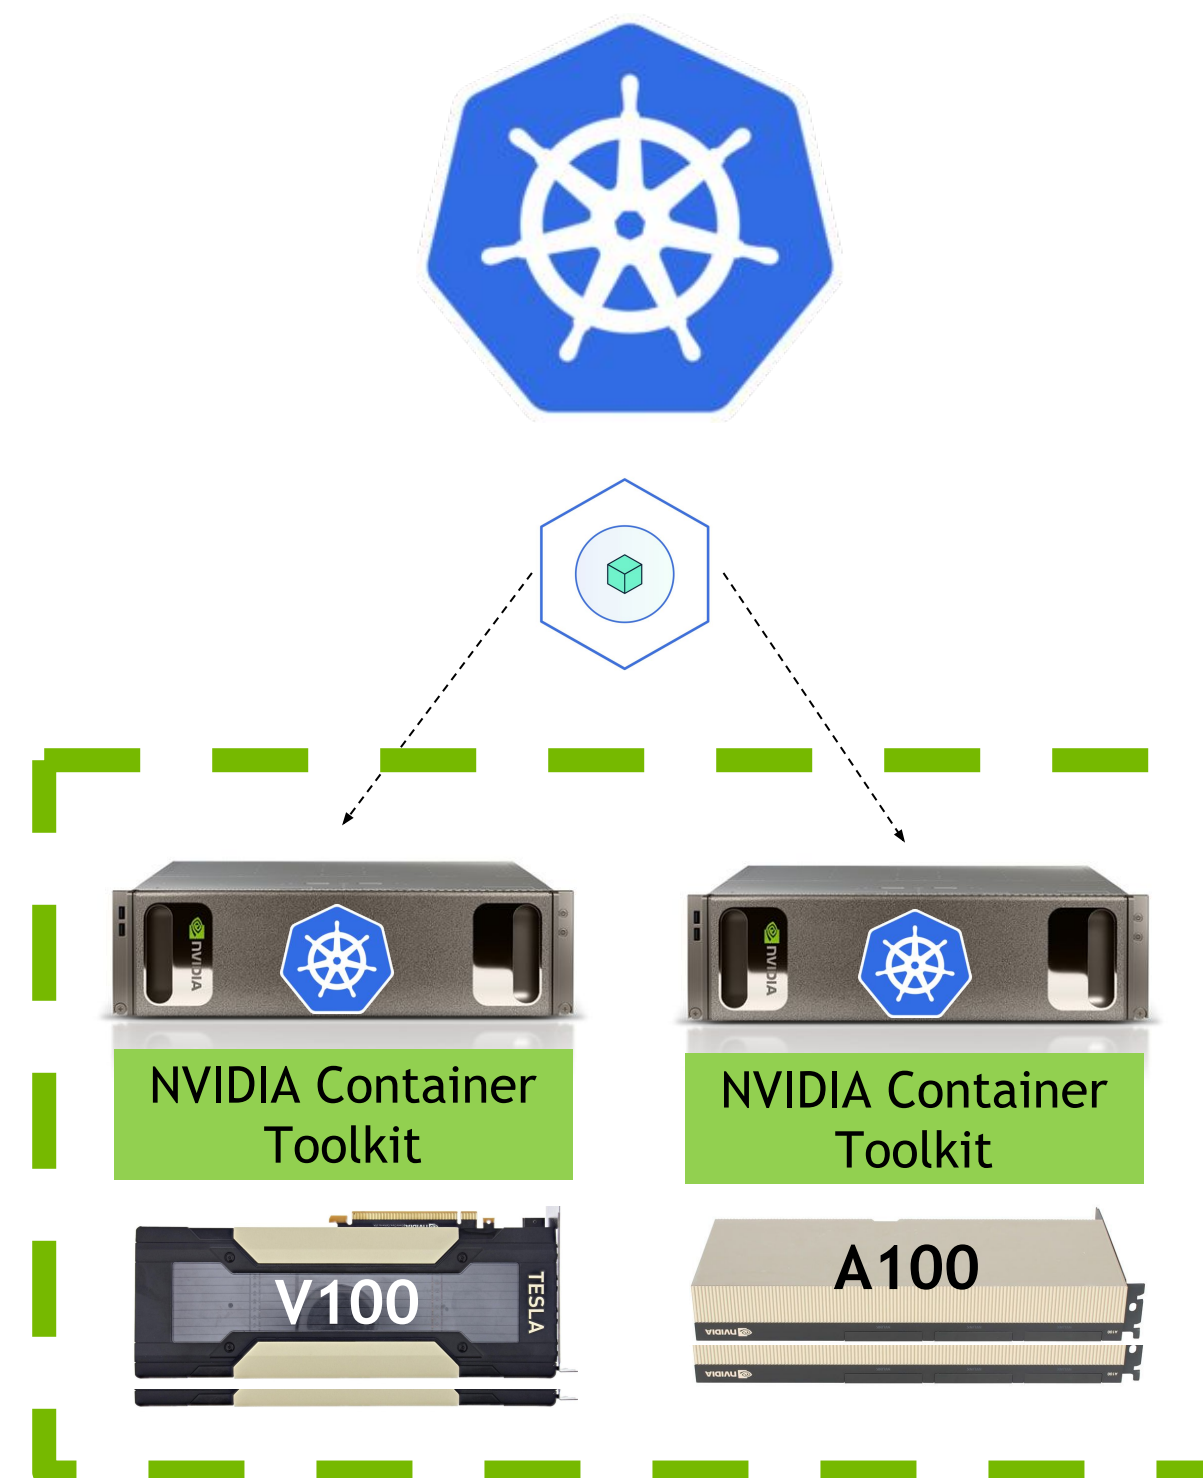
Seamlessly scale up training and inference to a cluster of GPU machines



GPUs AND KUBERNETES

Seamlessly scale up training and inference to a cluster of GPU machines

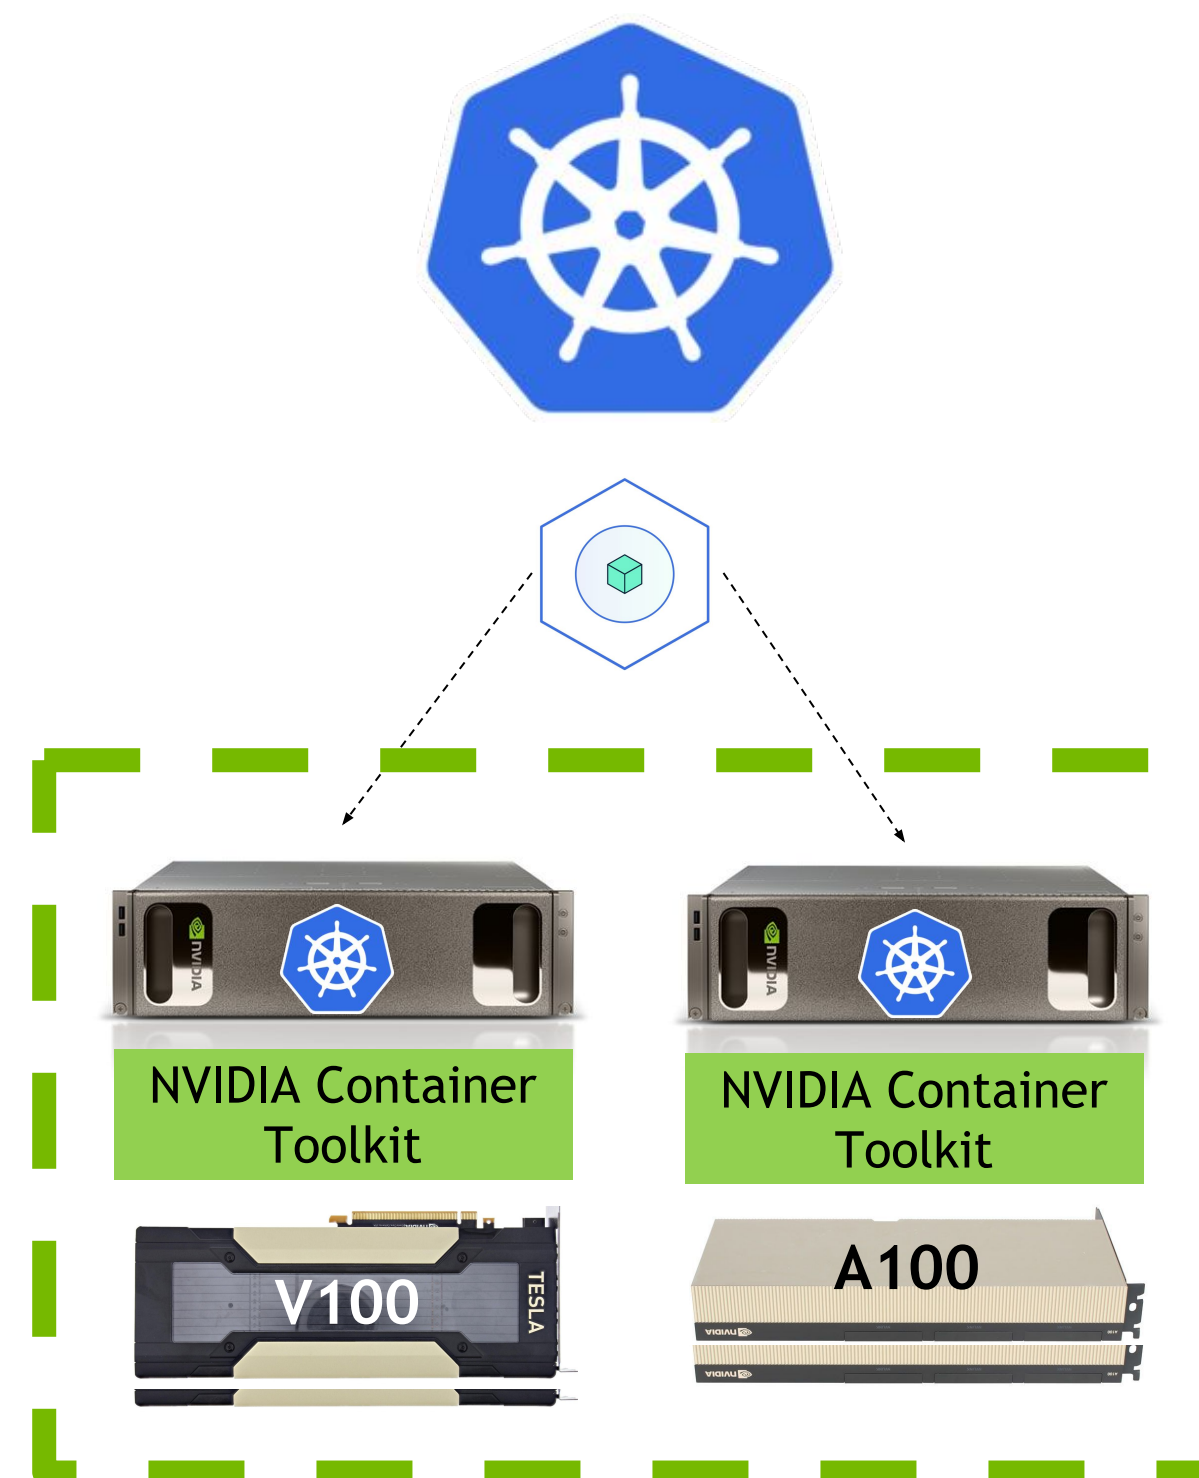
```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/gpu: 4
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
    nvidia.com/cuda.runtime: 11.0
    nvidia.com/cuda.driver: 450.51.06
```



GPUs AND KUBERNETES

Seamlessly scale up training and inference to a cluster of GPU machines

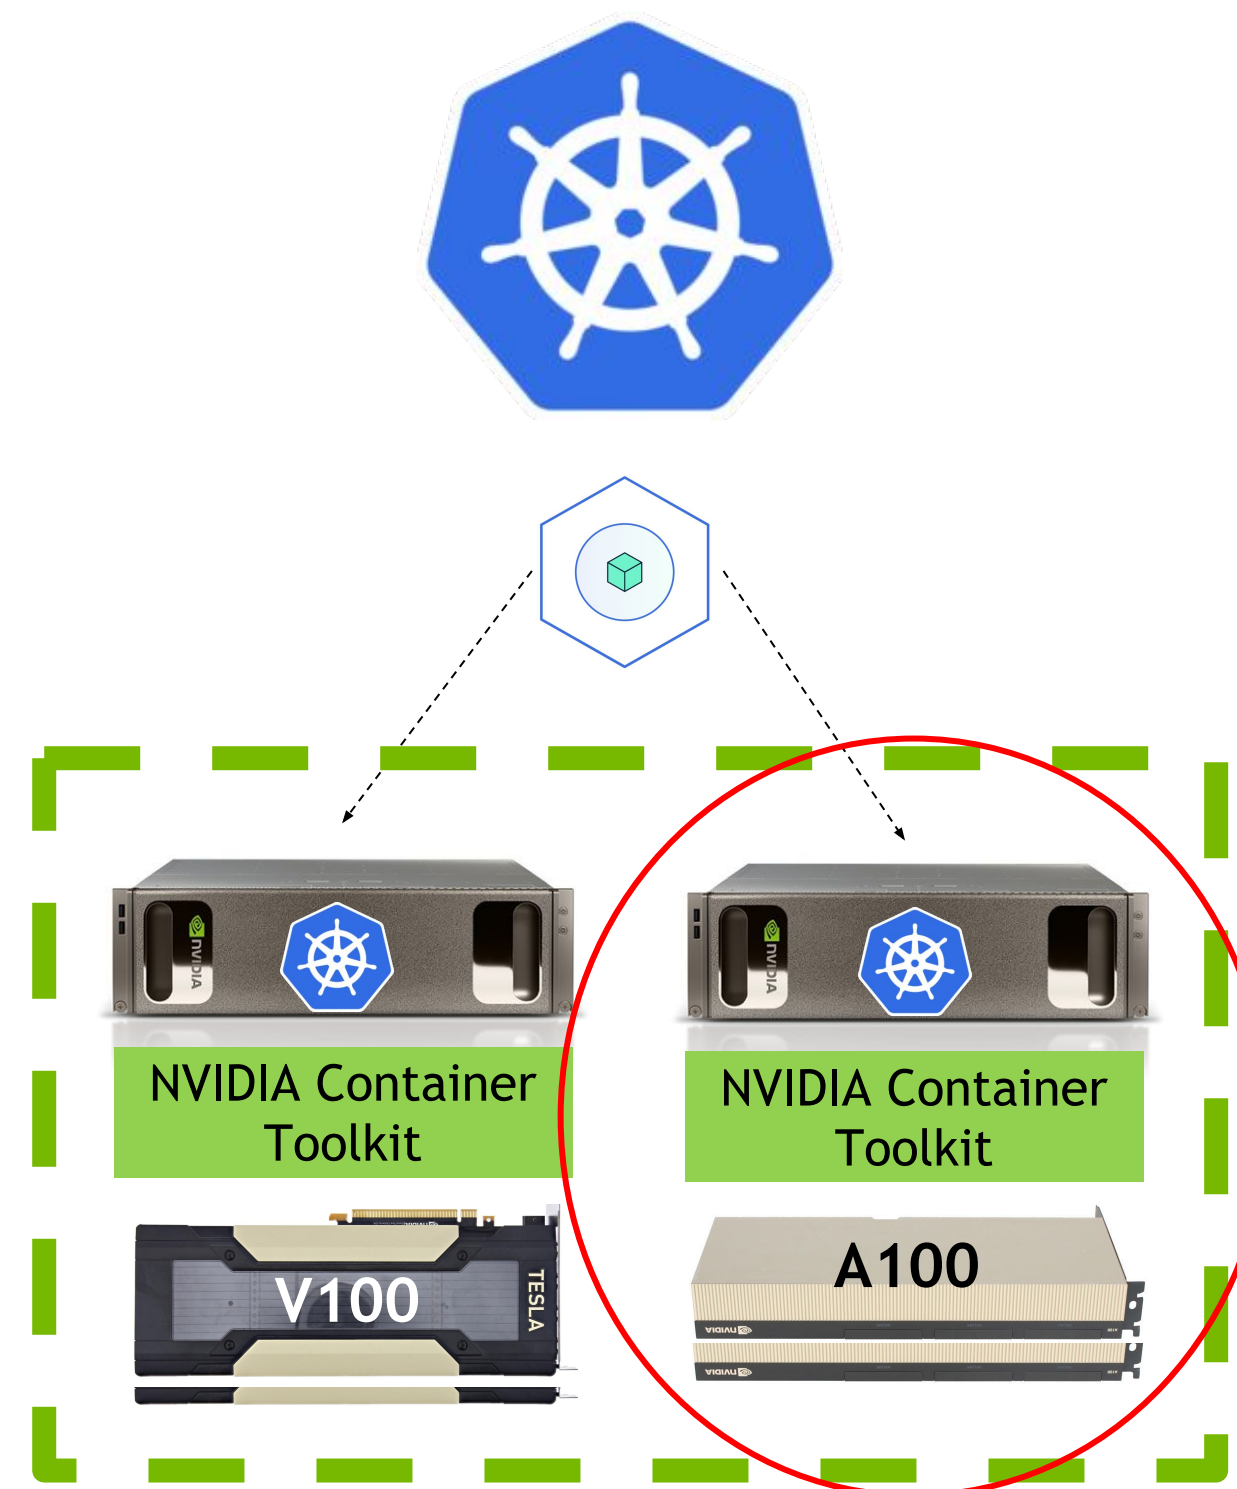
```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/gpu: 4
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
    nvidia.com/cuda.runtime: 11.0
    nvidia.com/cuda.driver: 450.51.06
```



GPUs AND KUBERNETES

Seamlessly scale up training and inference to a cluster of GPU machines

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/gpu: 4
      nodeSelector:
        nvidia.com/gpu.product: A100-PCIE-40GB
        nvidia.com/cuda.runtime: 11.0
        nvidia.com/cuda.driver: 450.51.06
```

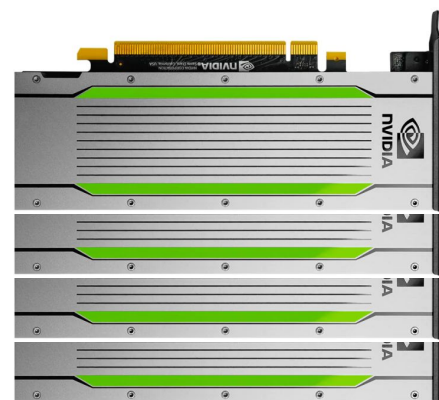


GPUs AND KUBERNETES

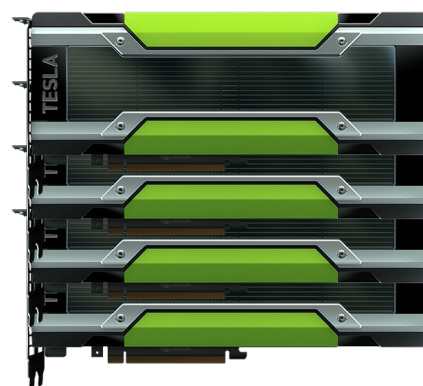
Seamlessly scale up training and inference to a cluster of GPU machines



4 x T4



4 x K80



8 x P100



8 x A100



8 x A100

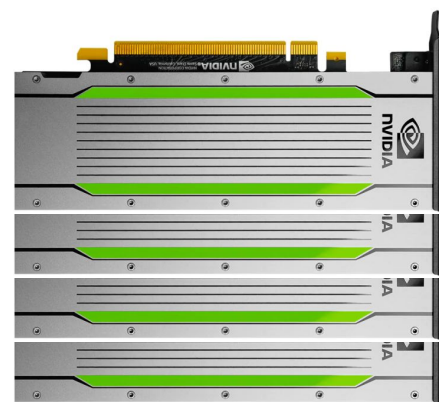


GPUs AND KUBERNETES

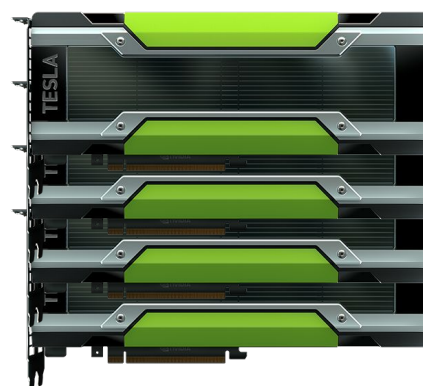
Seamlessly scale up training and inference to a cluster of GPU machines



4 x T4



4 x K80



8 x P100



...

Training



8 x A100



8 x A100



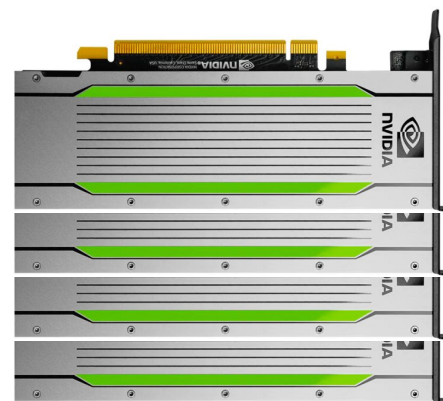
GPUs AND KUBERNETES

Seamlessly scale up training and inference to a cluster of GPU machines

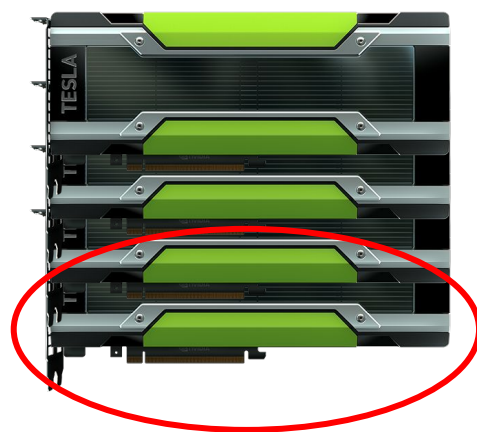
Inference



4 x T4



4 x K80



8 x P100



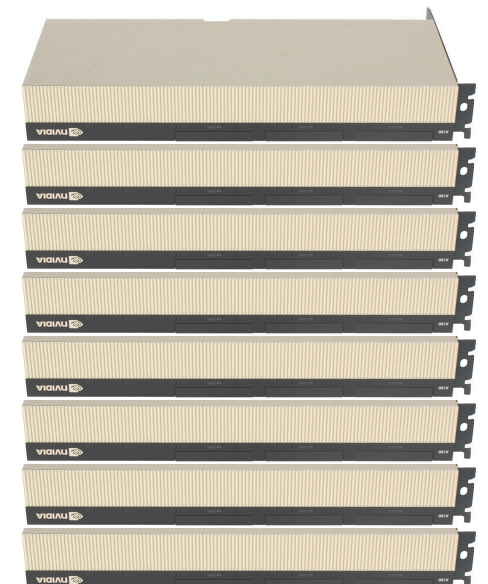
Training



8 x A100



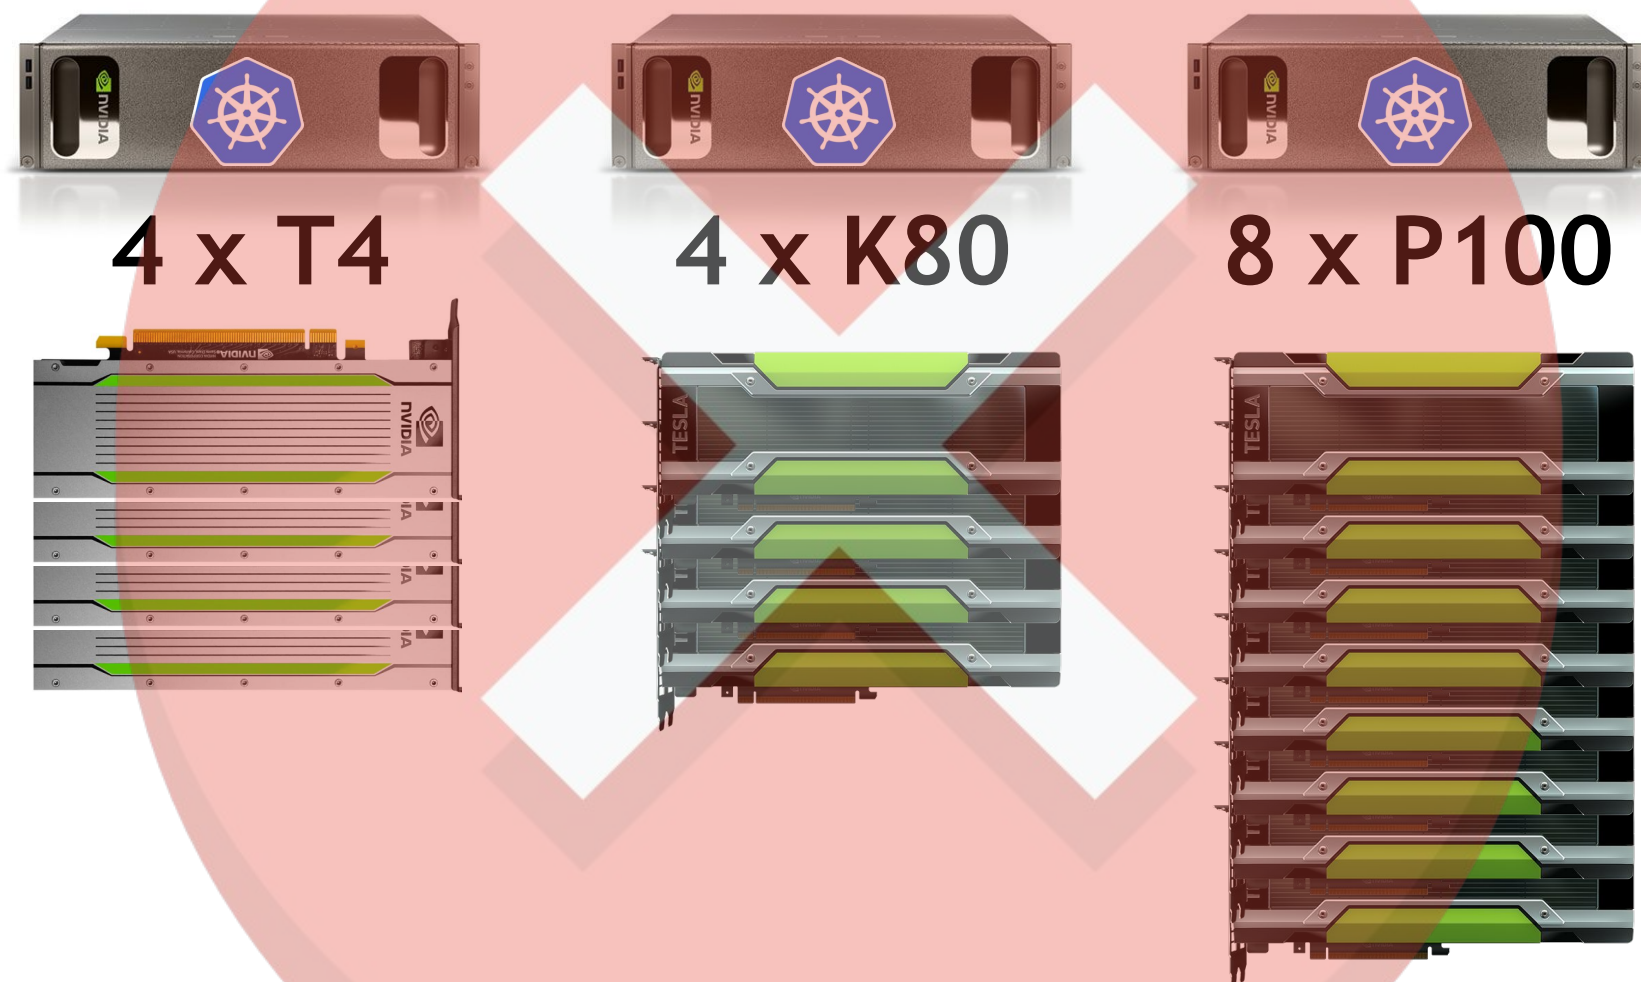
8 x A100



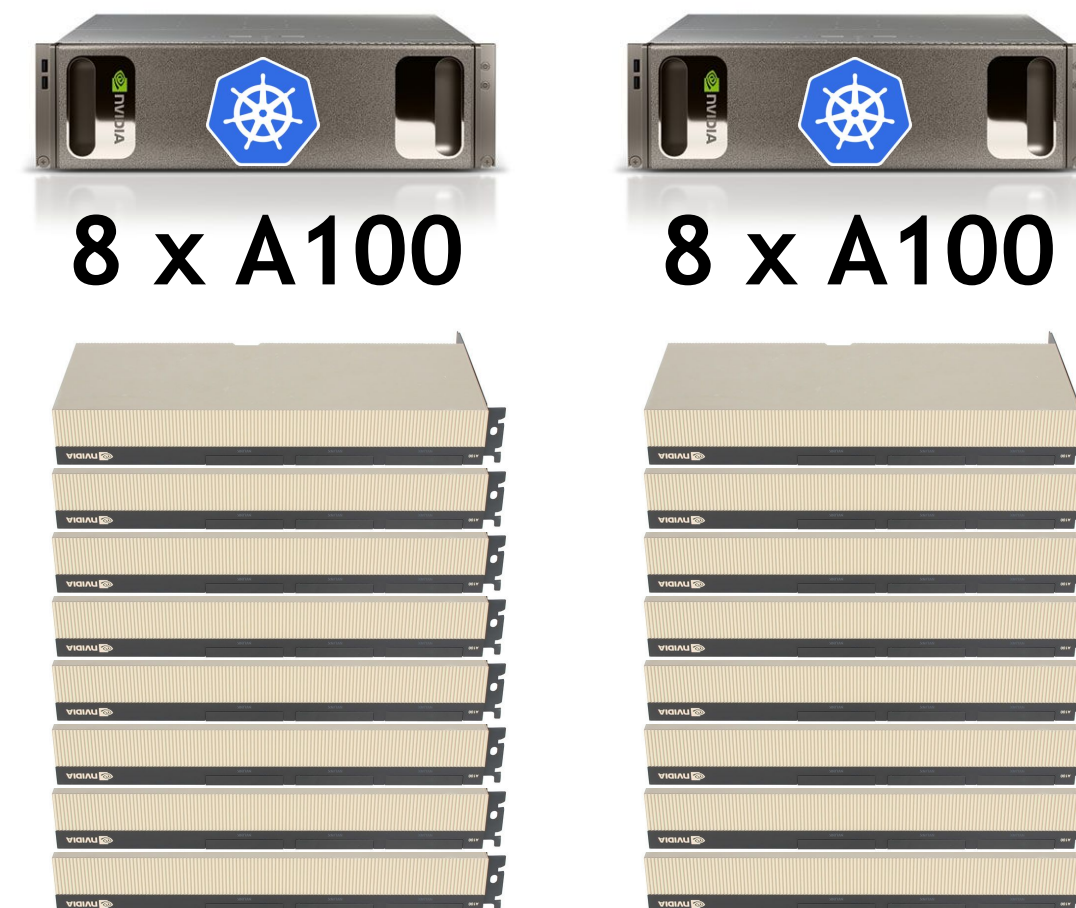
GPUs AND KUBERNETES

Seamlessly scale up training and inference to a cluster of GPU machines

Inference



Training



GPUs AND KUBERNETES

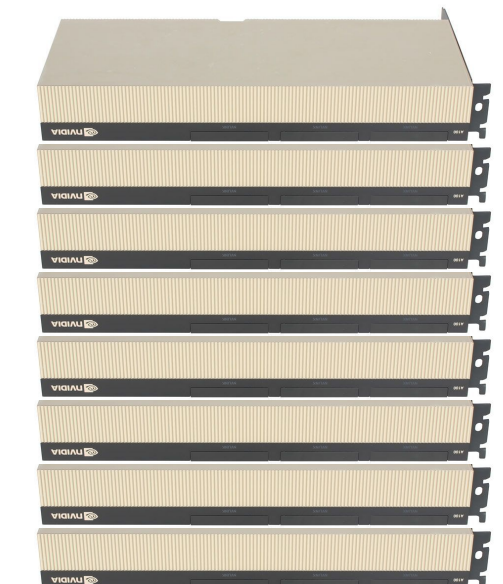
Seamlessly scale up training and inference to a cluster of GPU machines



8 x A100

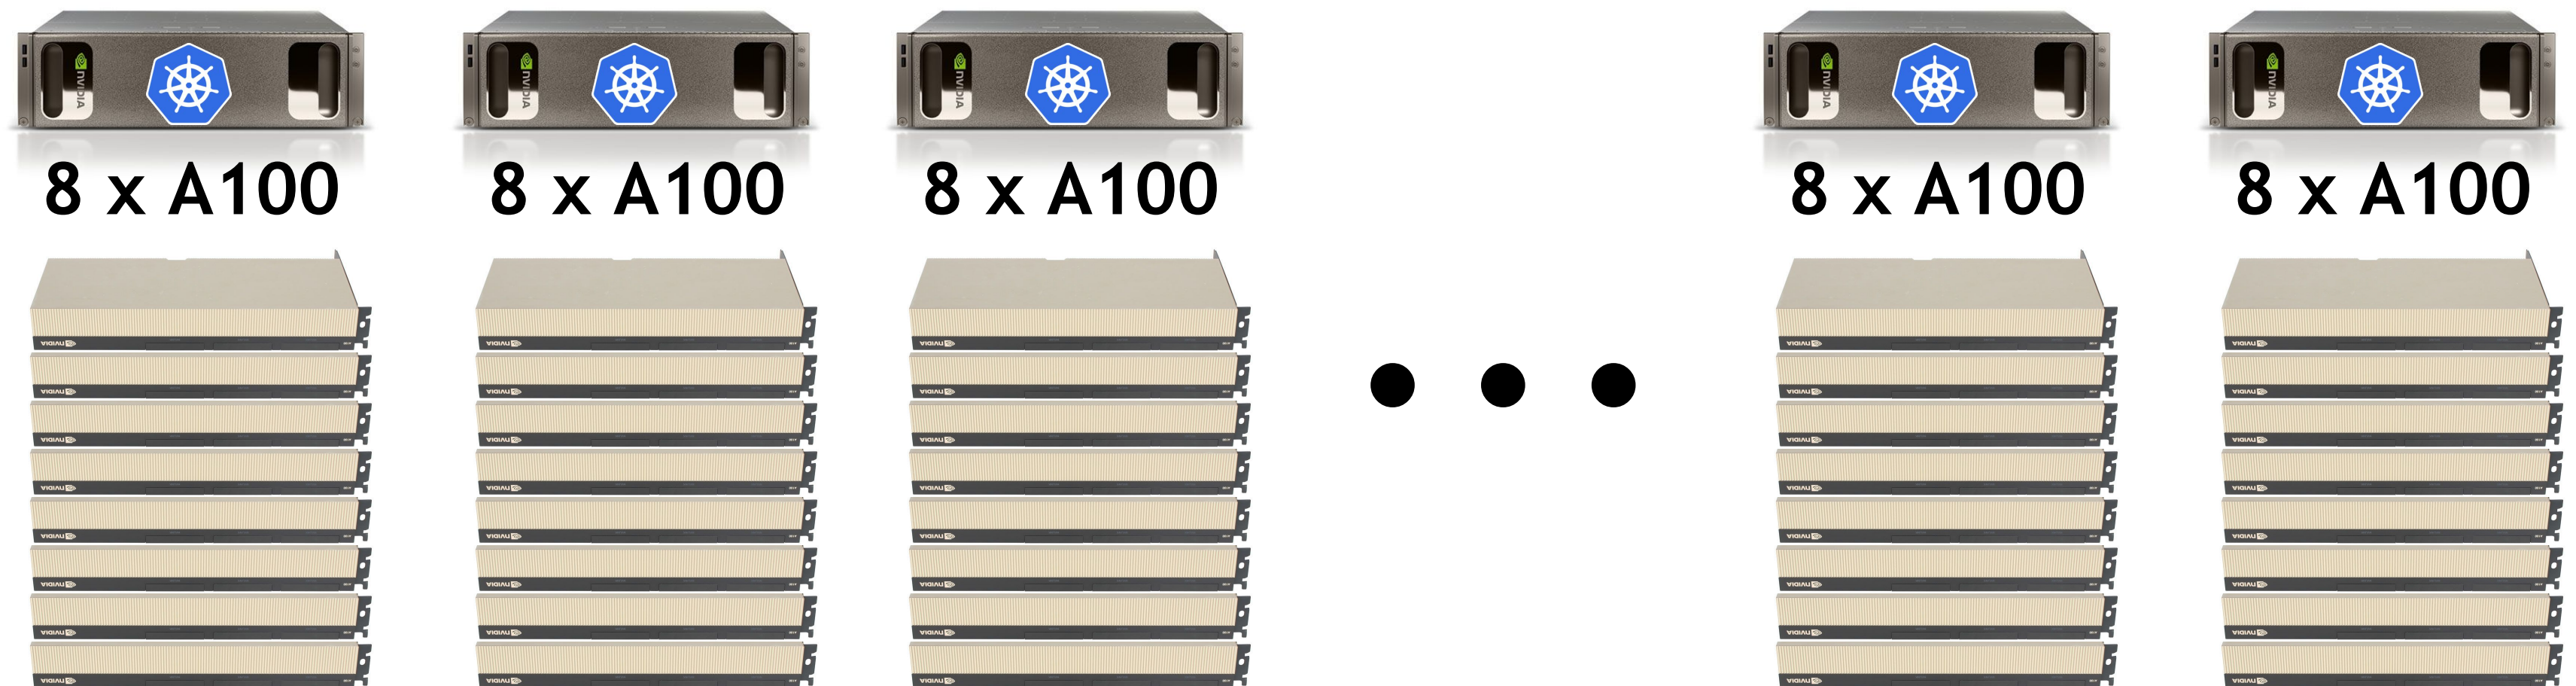


8 x A100



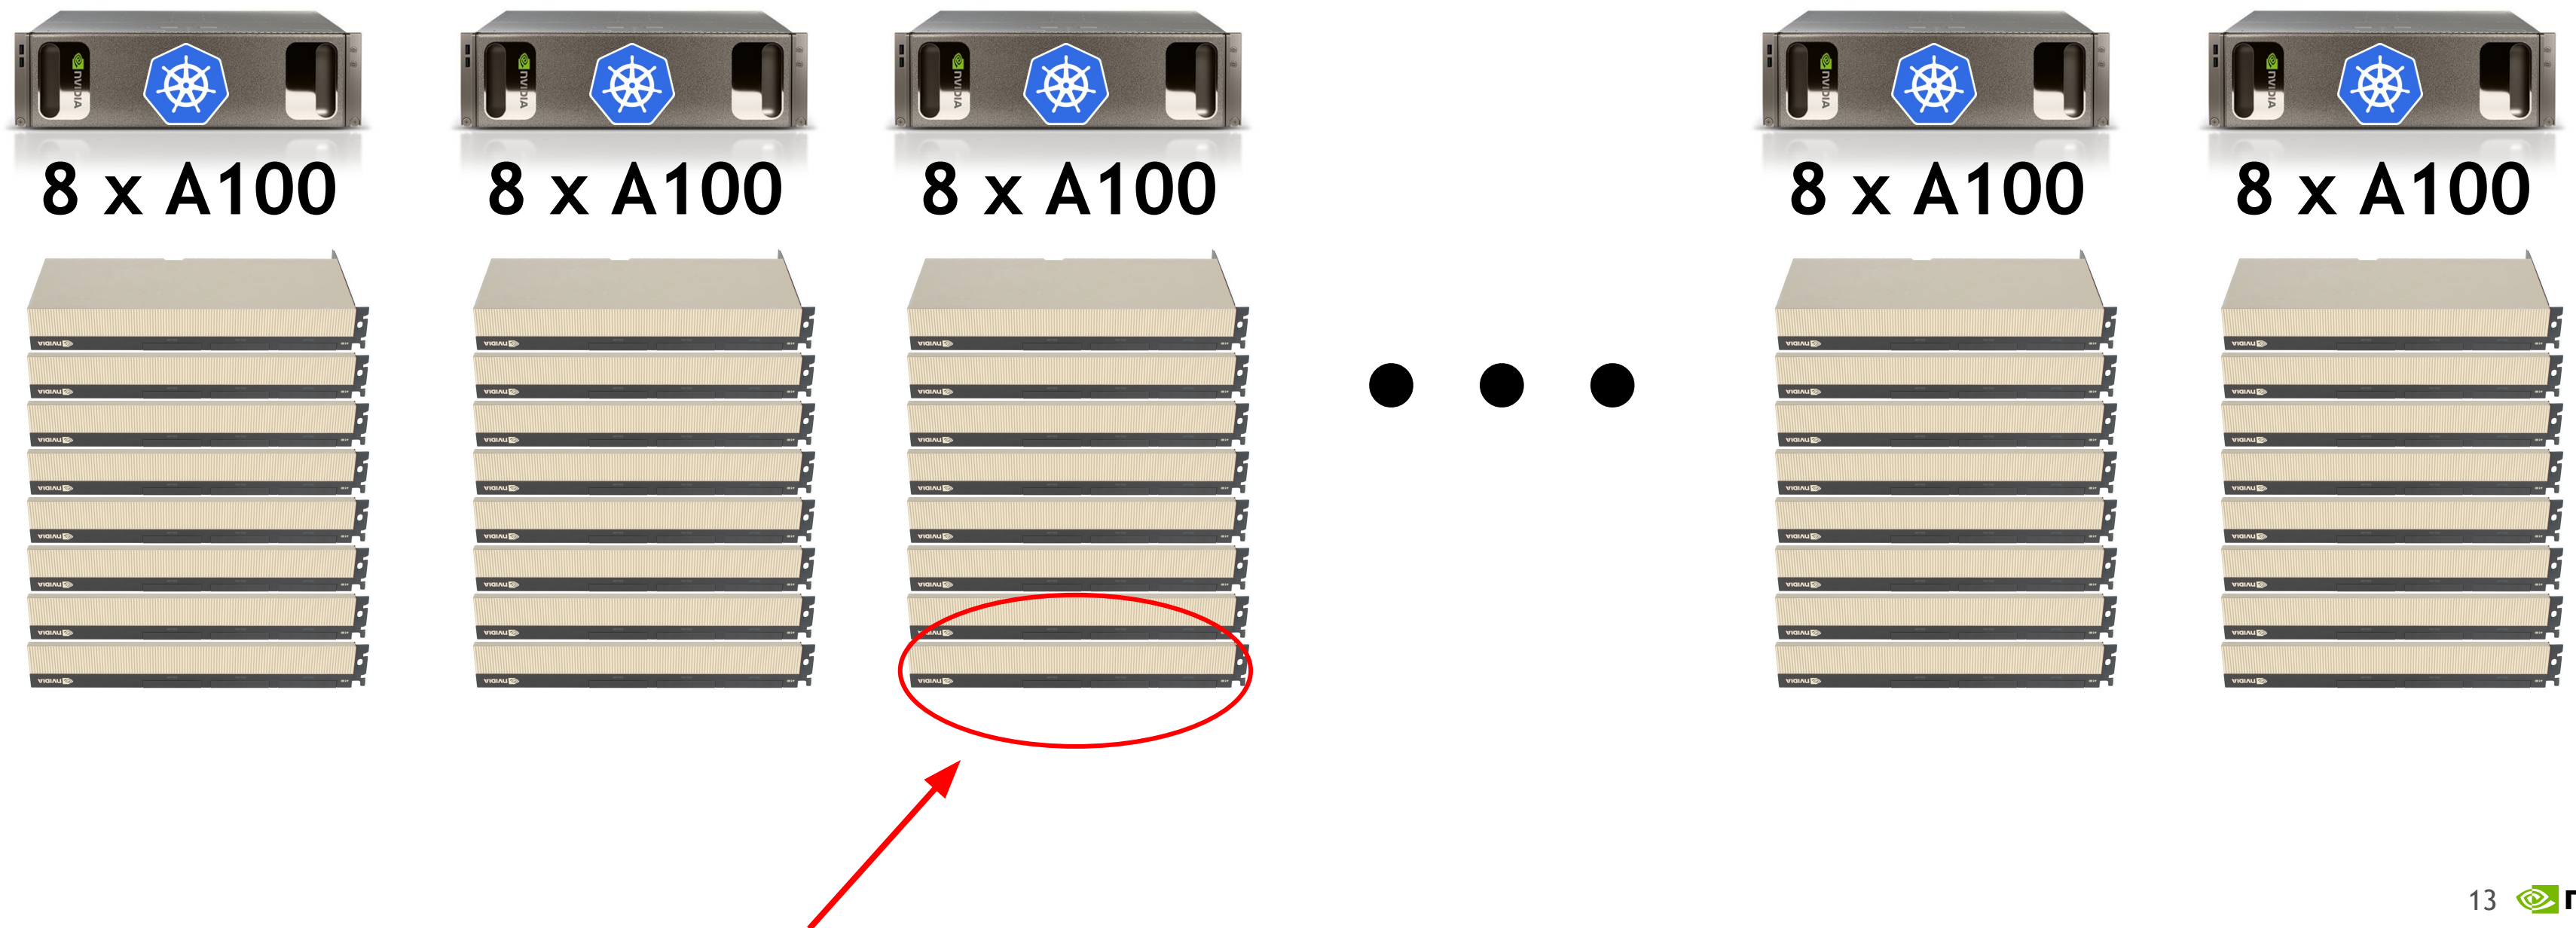
GPUs AND KUBERNETES

Seamlessly scale up training and inference to a cluster of GPU machines



GPUs AND KUBERNETES

Seamlessly scale up training and inference to a cluster of GPU machines

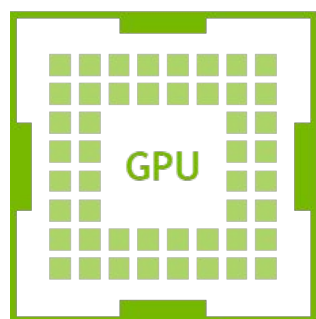


WHAT ARE MULTI-INSTANCE GPU_s?

Slices of a full GPU with dedicated memory and compute resources

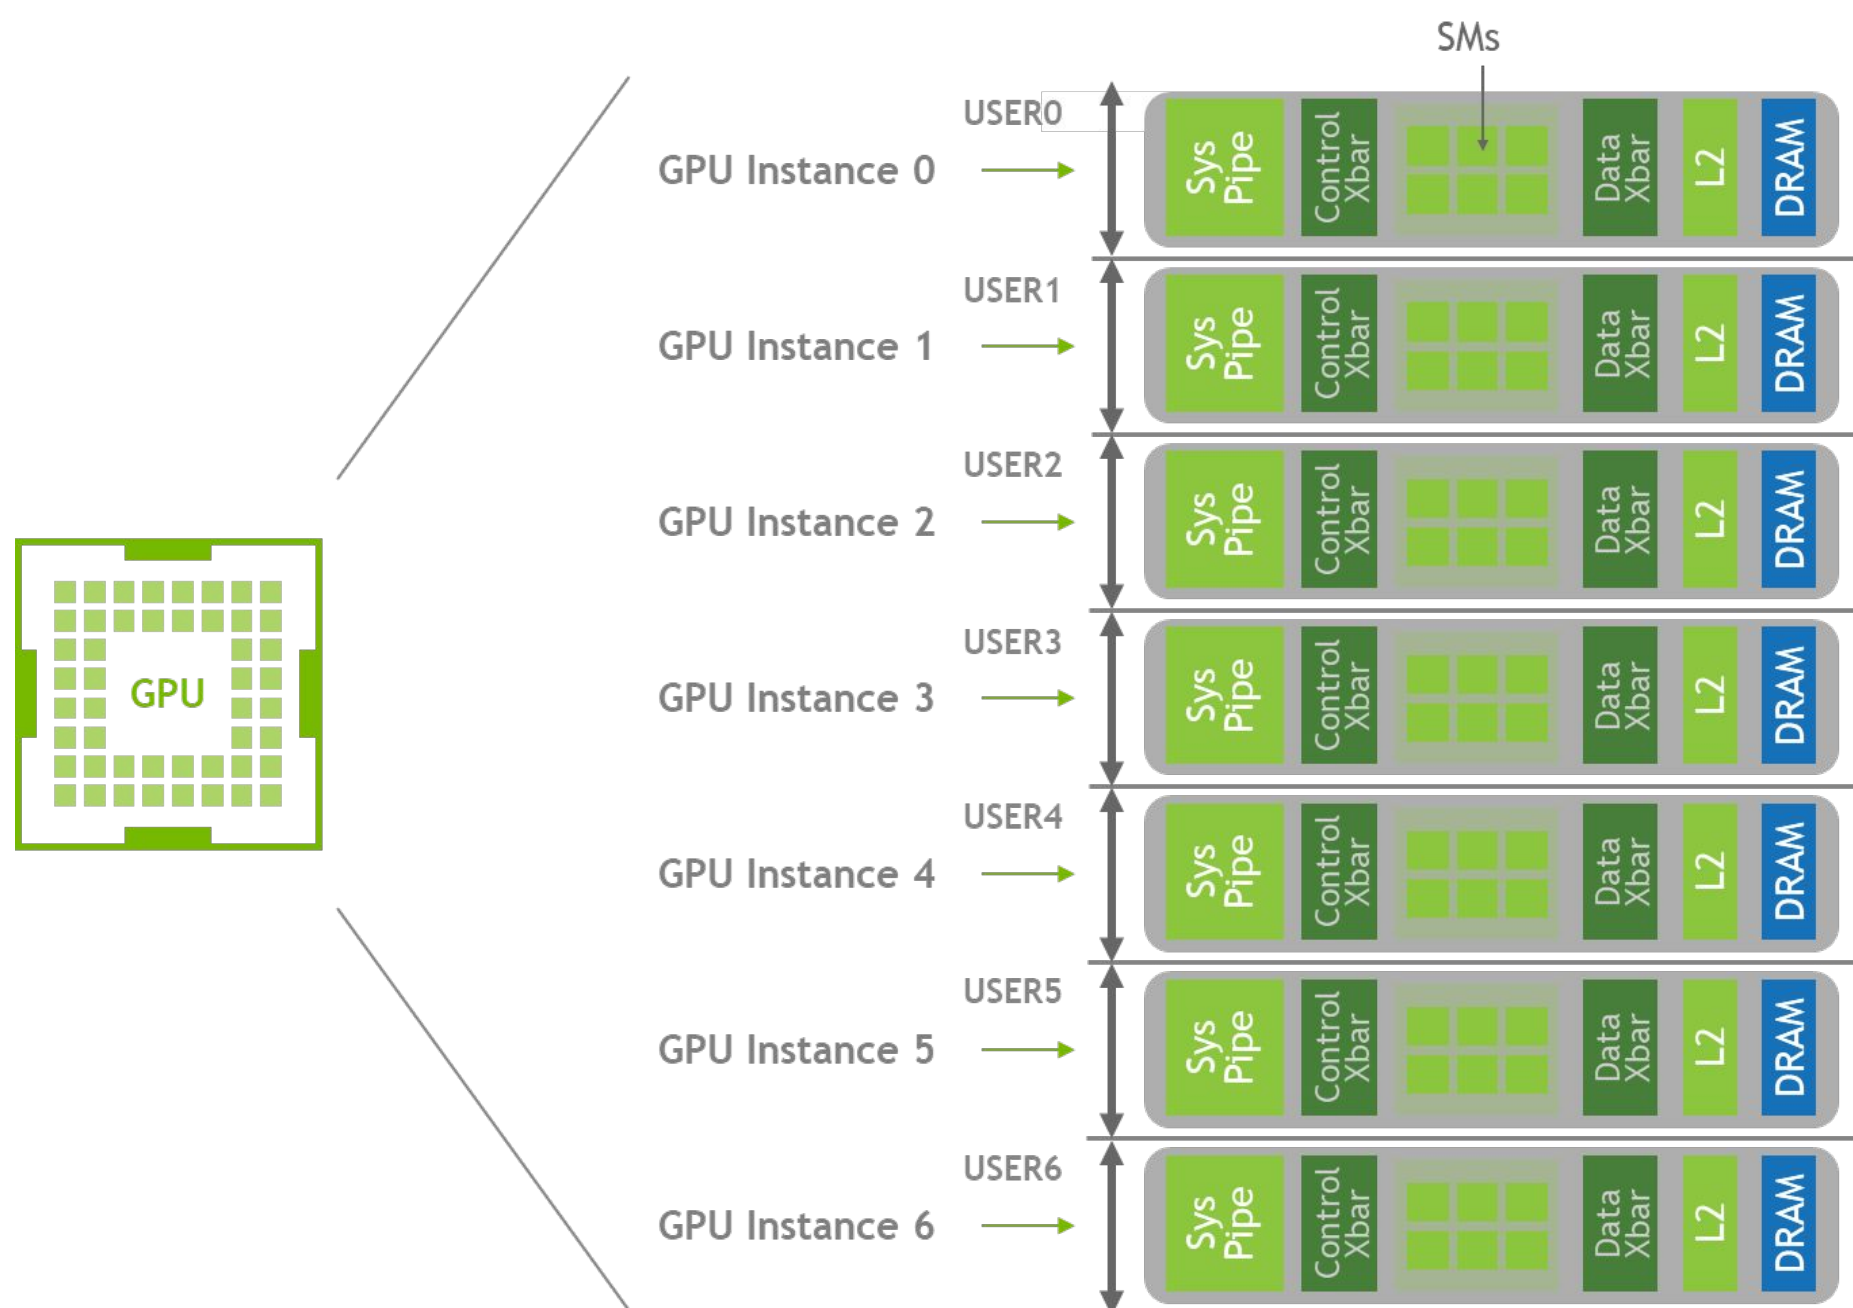
WHAT ARE MULTI-INSTANCE GPU_s?

Slices of a full GPU with dedicated memory and compute resources



WHAT ARE MULTI-INSTANCE GPU_s?

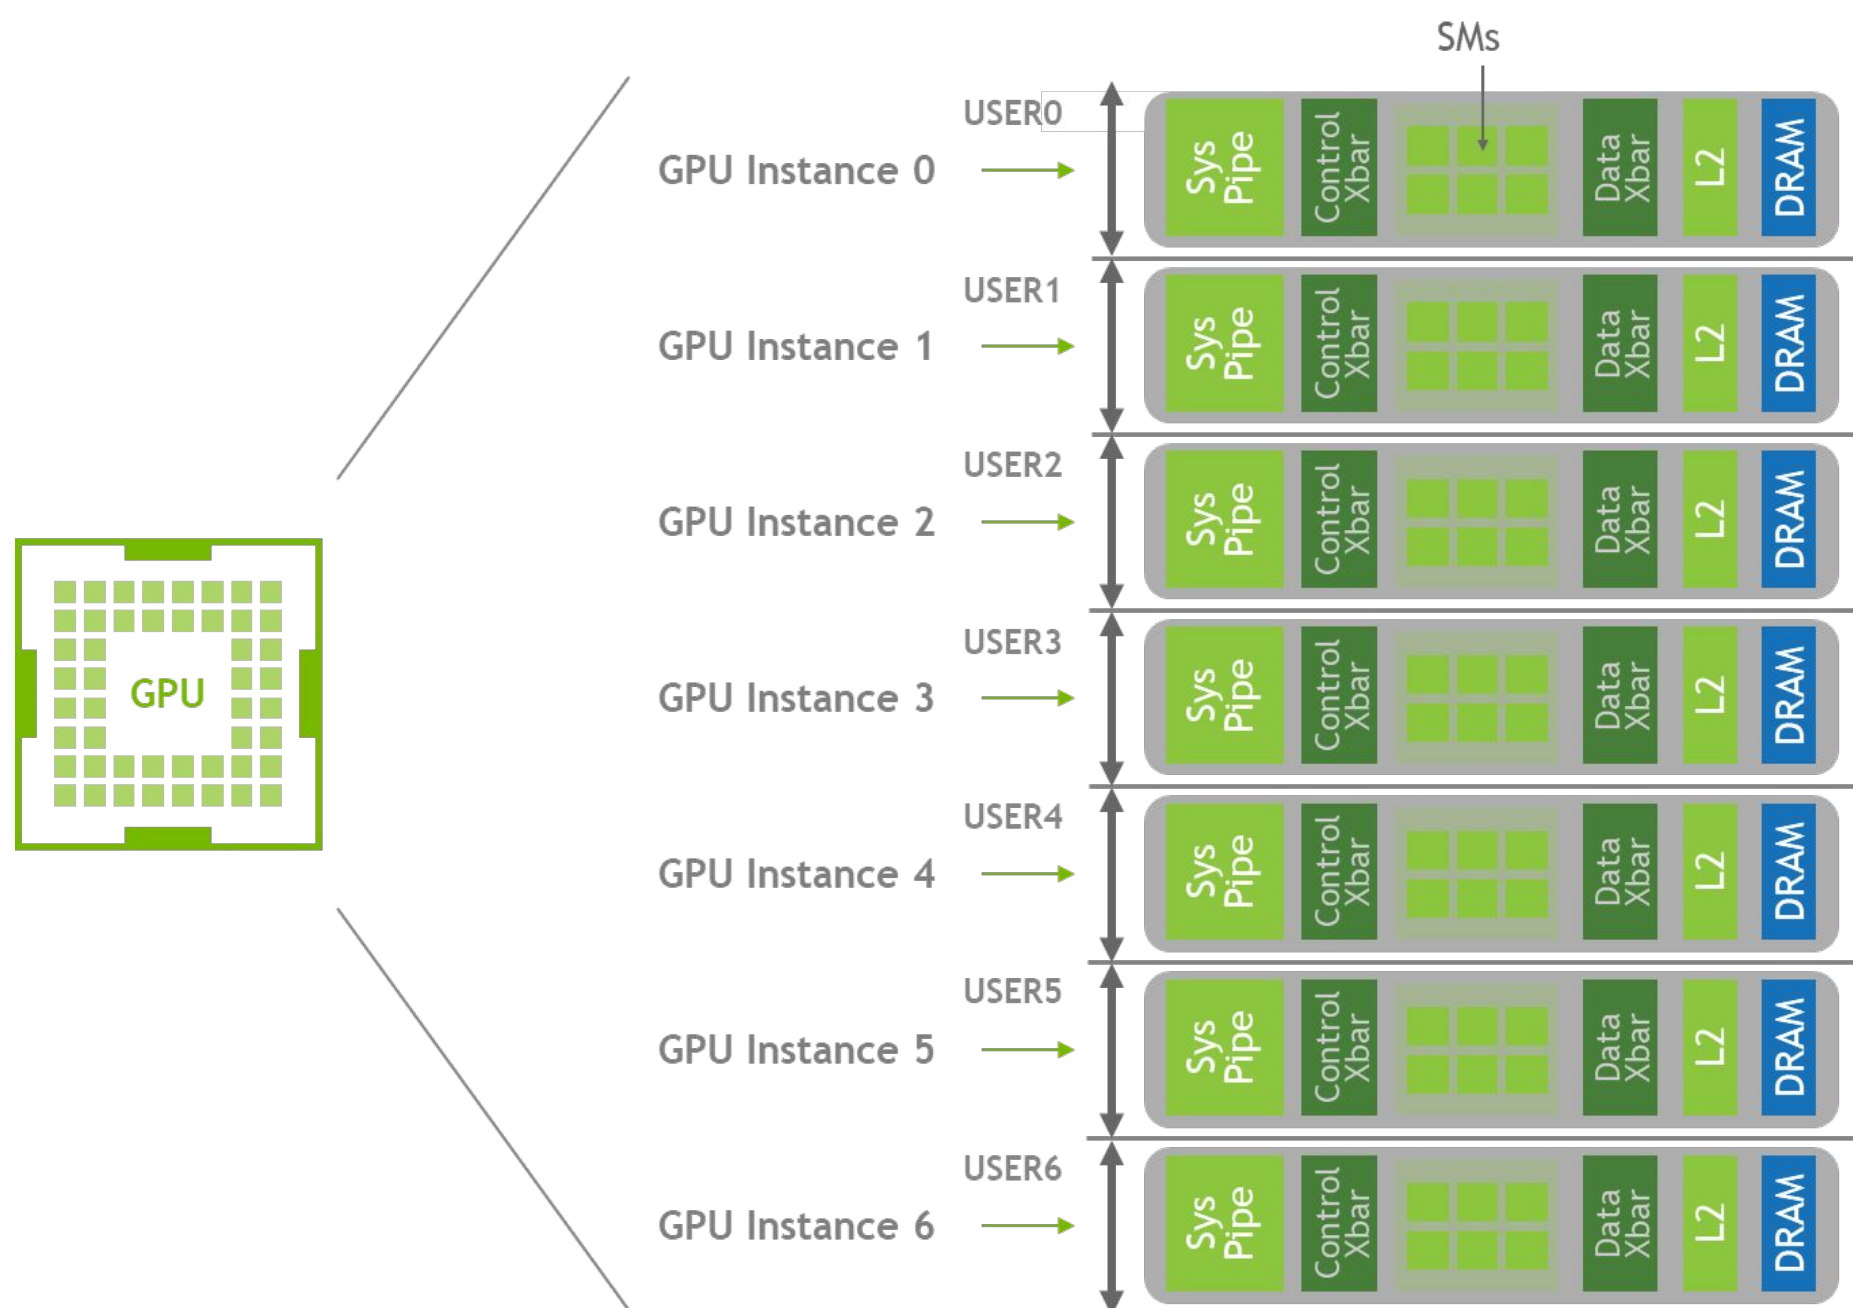
Slices of a full GPU with dedicated memory and compute resources



Create multiple “GPU Instances” on a single GPU:
Dedicated SM, Memory, L2 cache, Bandwidth for
hardware QoS & isolation

WHAT ARE MULTI-INSTANCE GPU_s?

Slices of a full GPU with dedicated memory and compute resources

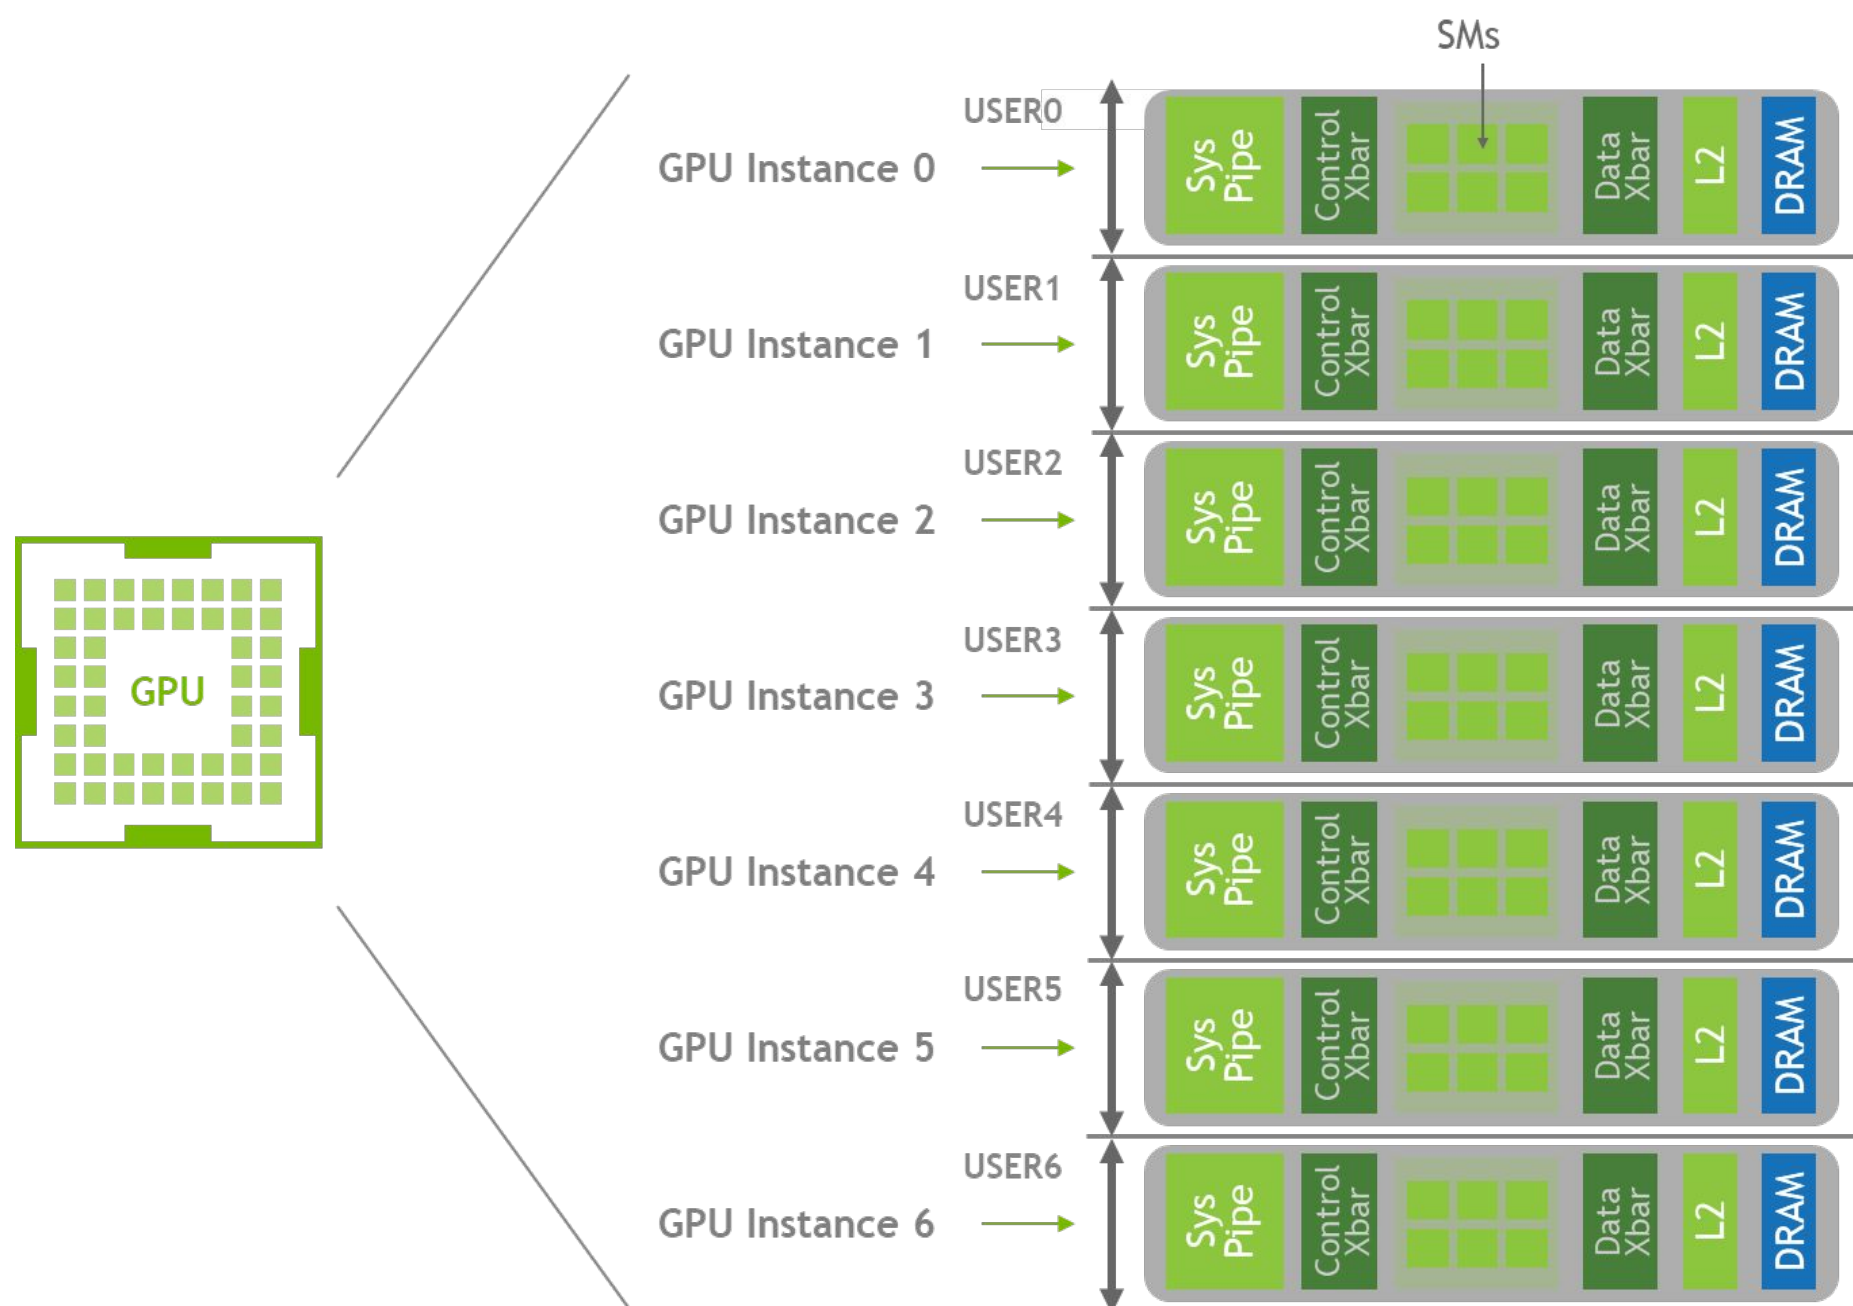


Create multiple “GPU Instances” on a single GPU:
Dedicated SM, Memory, L2 cache, Bandwidth for hardware QoS & isolation

Simultaneous Workload Execution With Guaranteed Quality Of Service:
All GPU instances run in parallel with predictable throughput & latency

WHAT ARE MULTI-INSTANCE GPU_s?

Slices of a full GPU with dedicated memory and compute resources



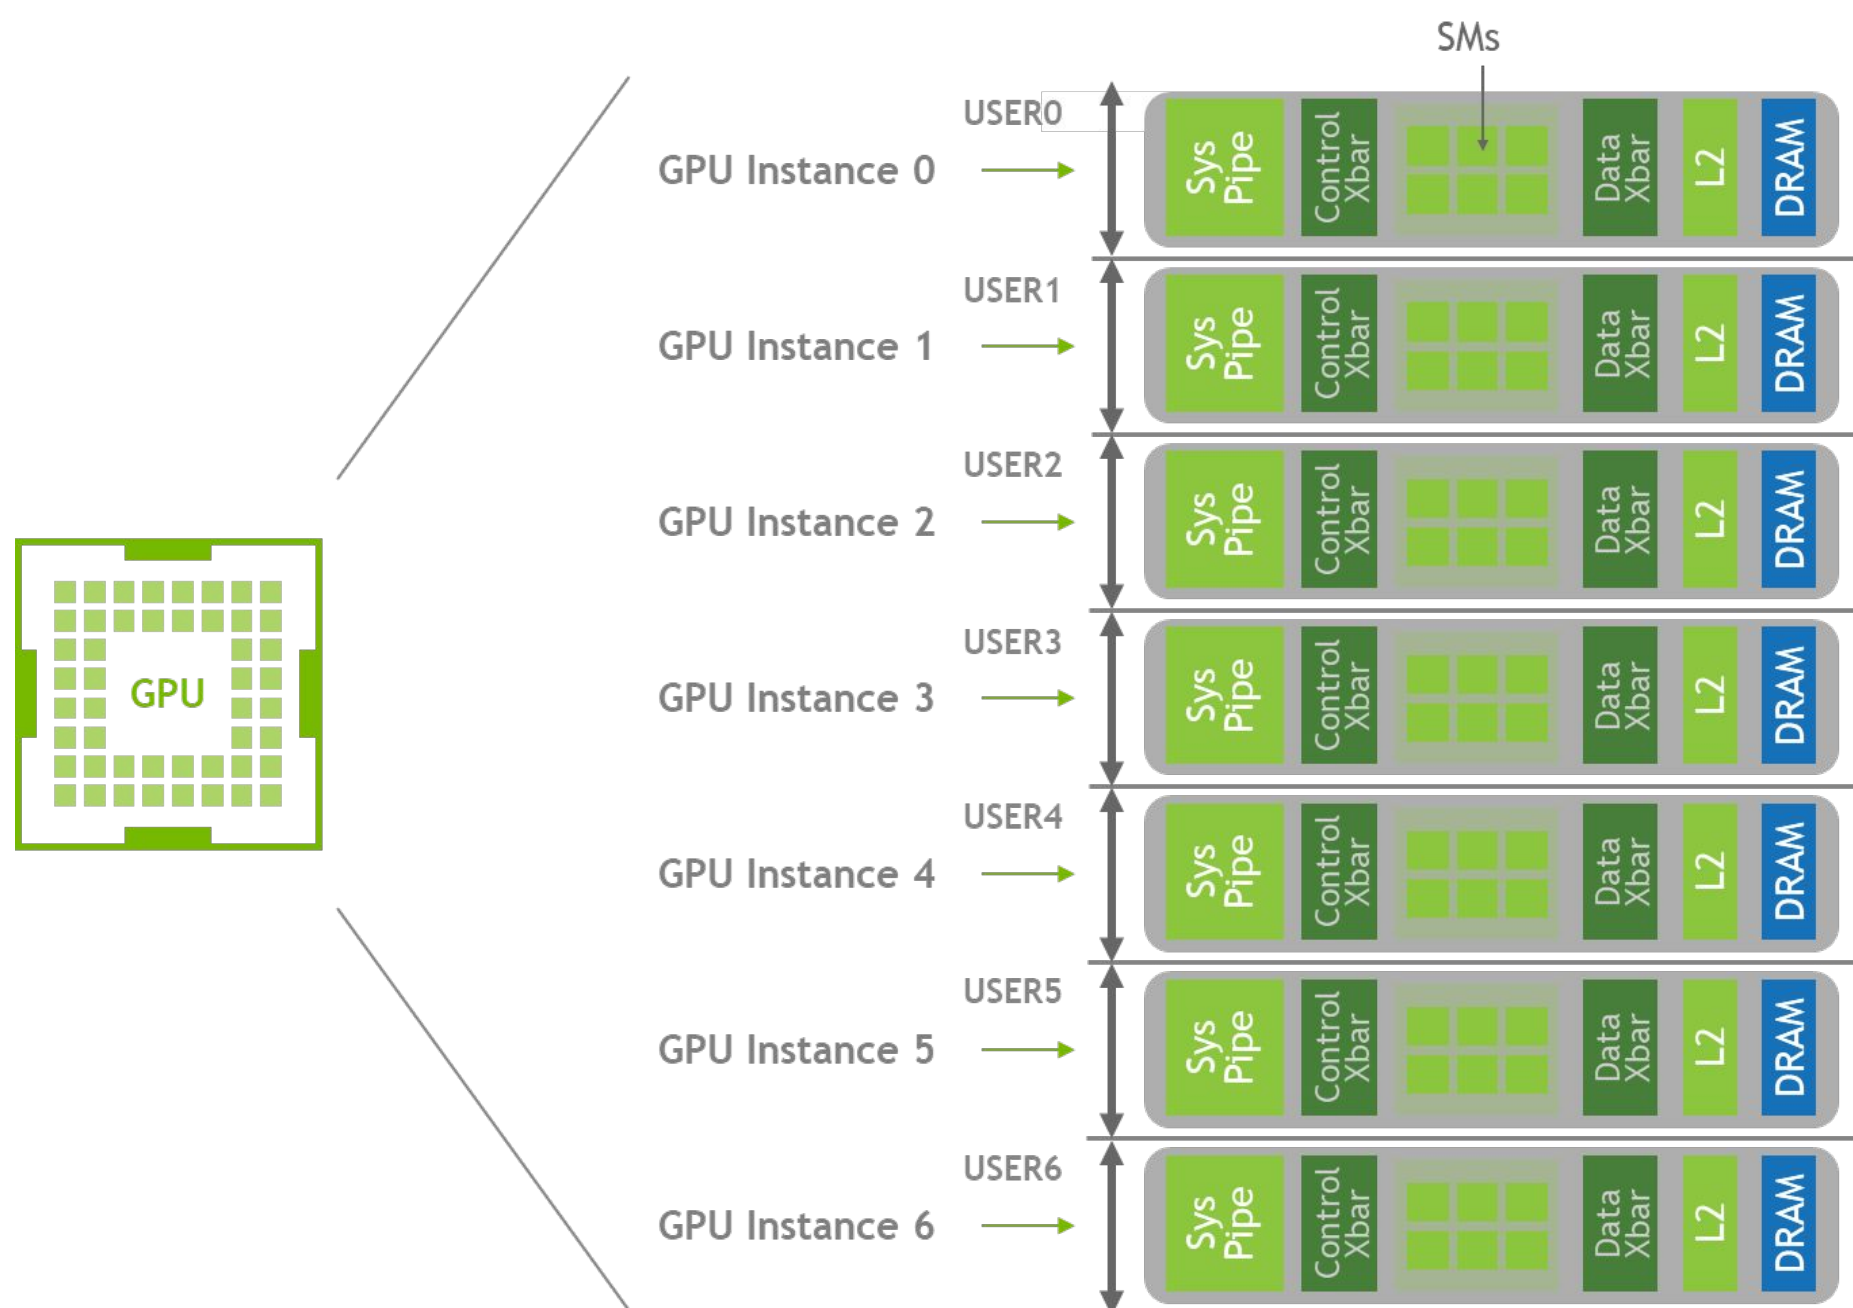
Create multiple “GPU Instances” on a single GPU:
Dedicated SM, Memory, L2 cache, Bandwidth for hardware QoS & isolation

Simultaneous Workload Execution With Guaranteed Quality Of Service:
All GPU instances run in parallel with predictable throughput & latency

Right Sized GPU Allocation:
Different sized GPU instances based on target workloads

WHAT ARE MULTI-INSTANCE GPU_s?

Slices of a full GPU with dedicated memory and compute resources



Create multiple “GPU Instances” on a single GPU:
Dedicated SM, Memory, L2 cache, Bandwidth for hardware QoS & isolation

Simultaneous Workload Execution With Guaranteed Quality Of Service:
All GPU instances run in parallel with predictable throughput & latency

Right Sized GPU Allocation:
Different sized GPU instances based on target workloads

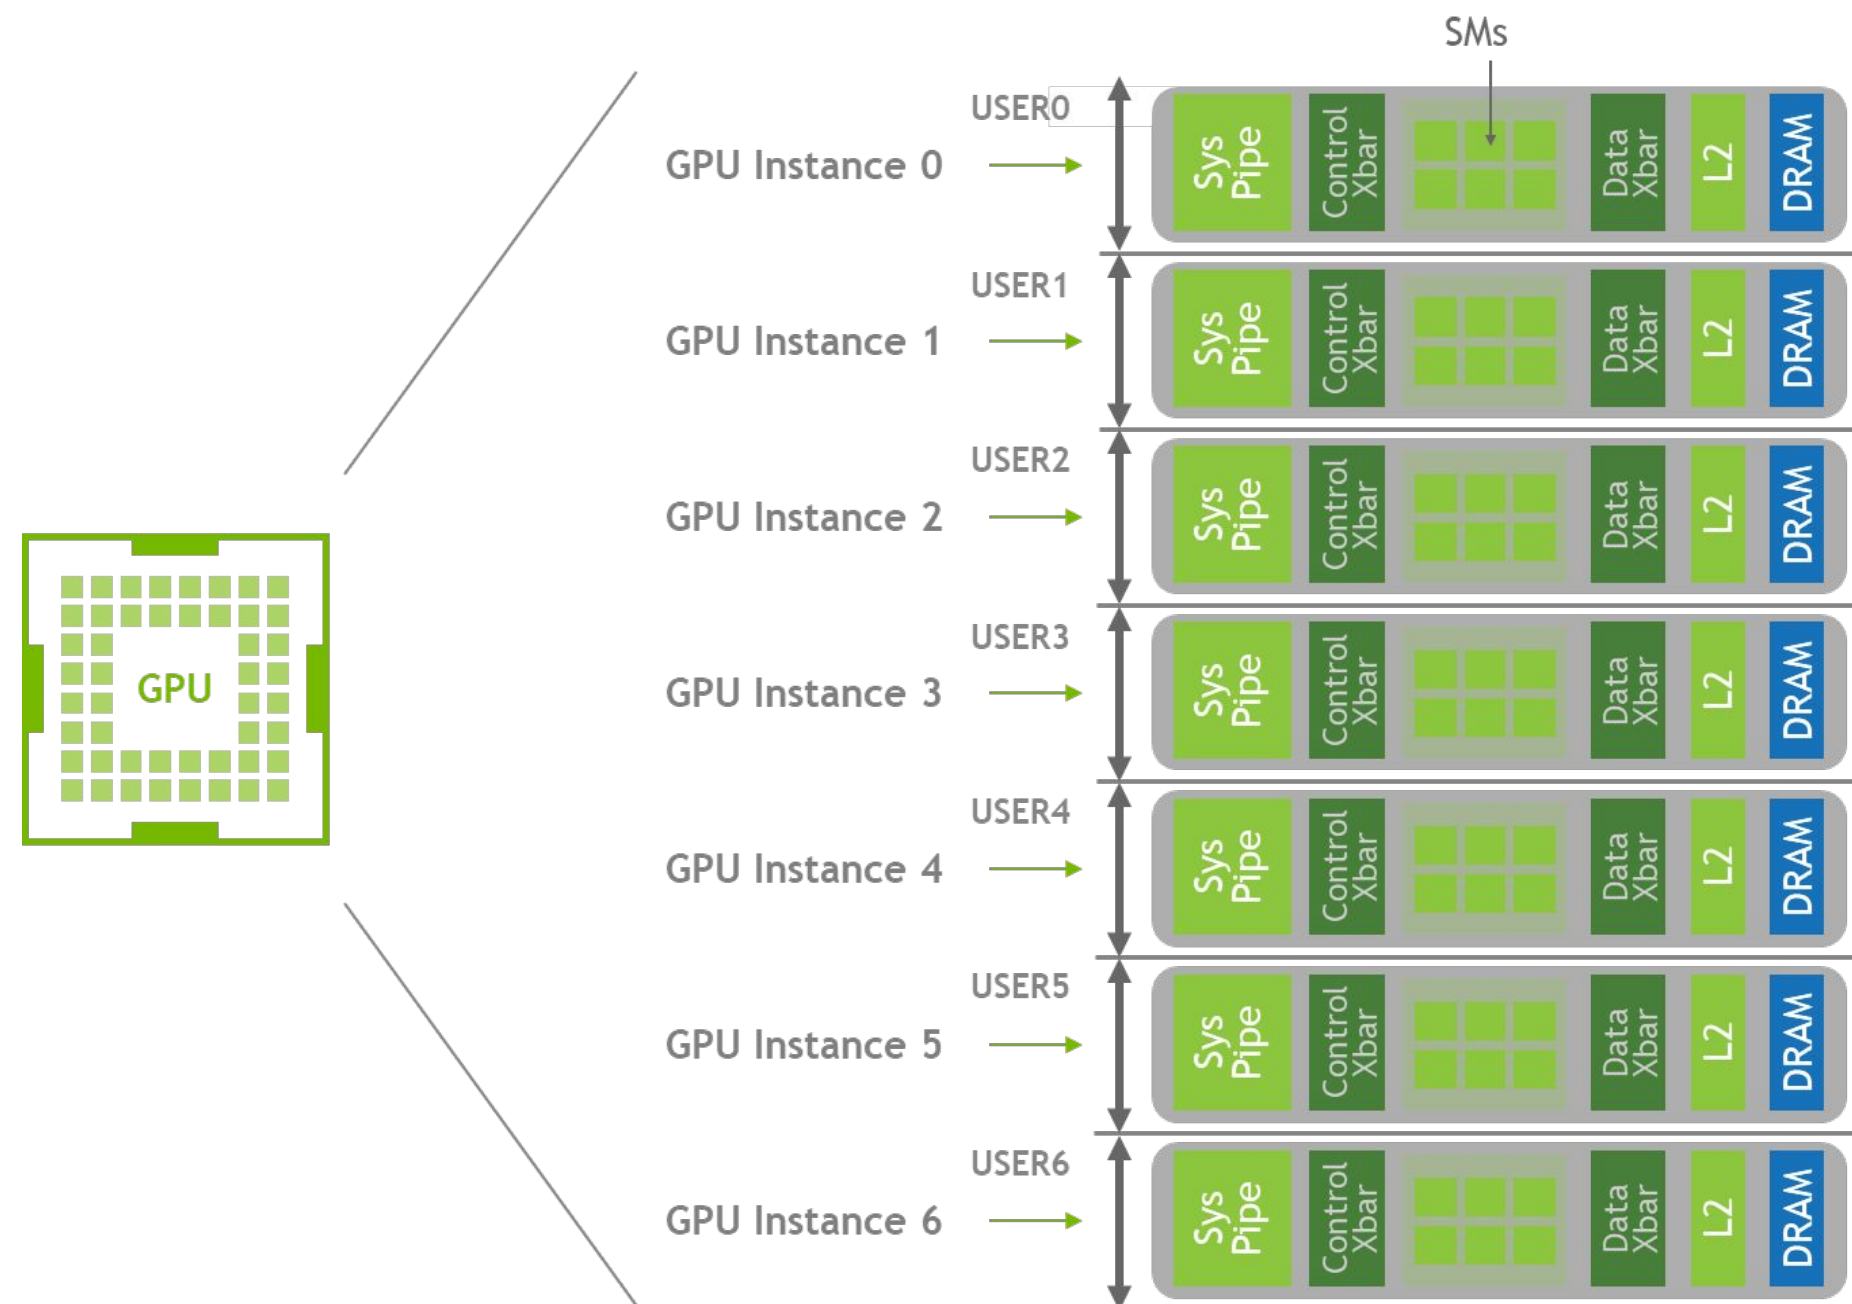
Diverse Deployment Environments:
Supported with Bare metal, Container and Virtualized Env.

OUTLINE

- Multi-Instance GPUs (MIG)
- GPUs and Containers
- GPUs and Kubernetes
- MIG in Containers
- MIG in Kubernetes
- System-Level Interface for MIG
- Best-Practices for Provisioning MIG
- Putting it all together – Demo

MULTI-INSTANCE GPU_s (MIG)

Use Cases



MULTI-INSTANCE GPU_s (MIG)

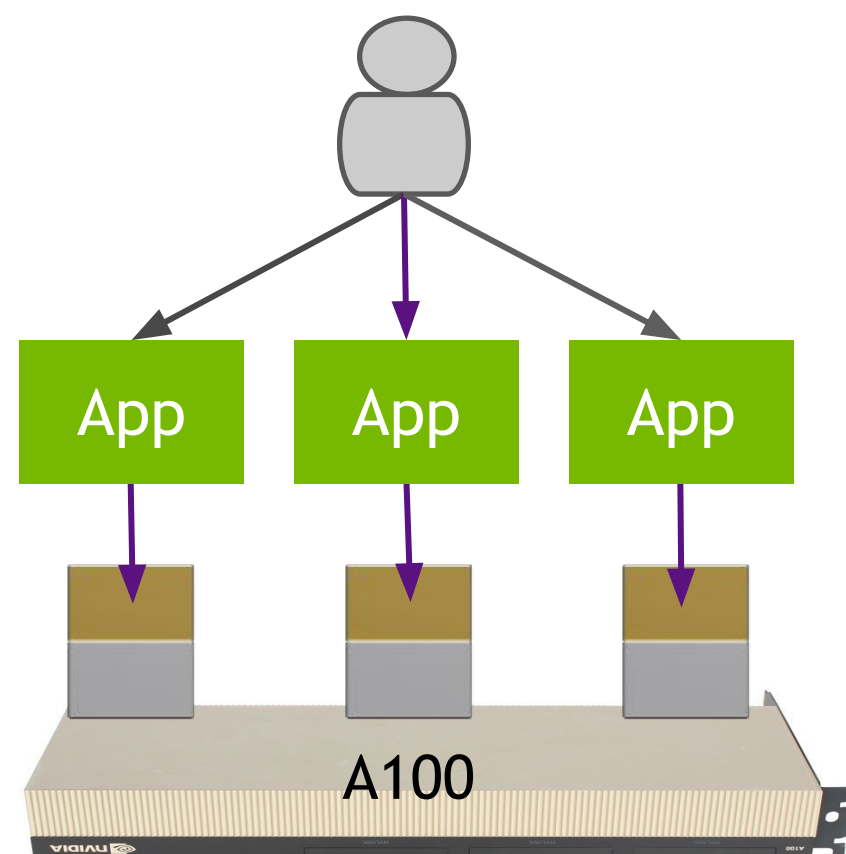
Use Cases



MULTI-INSTANCE GPUs (MIG)

Use Cases

Single User → Multiple Apps

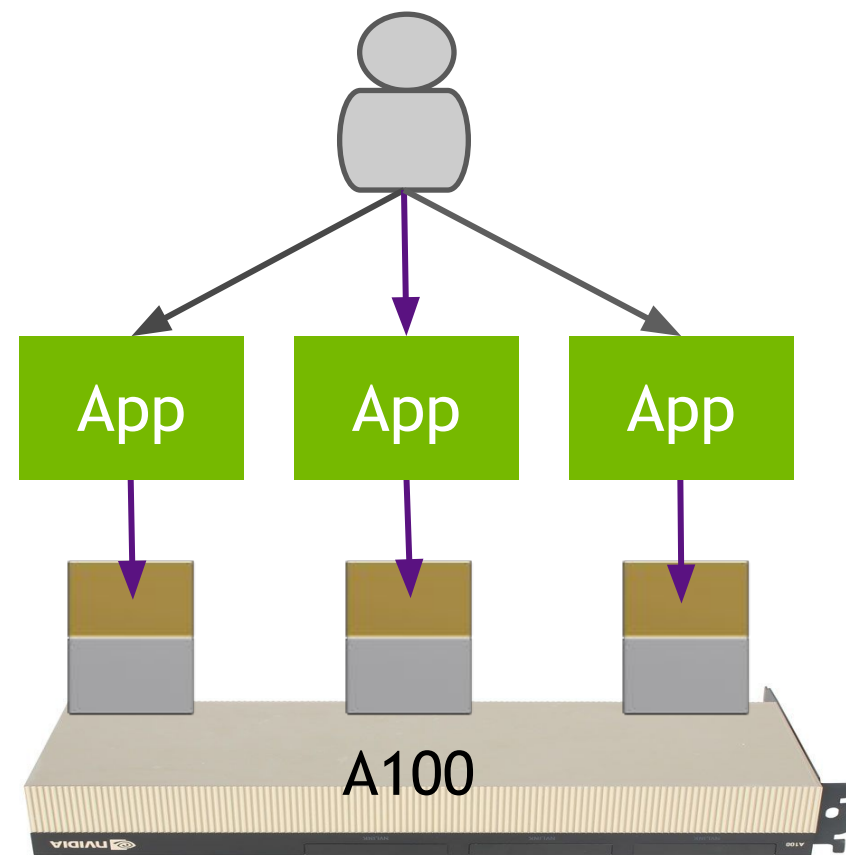


E.g. *Multiple inference jobs*

MULTI-INSTANCE GPUs (MIG)

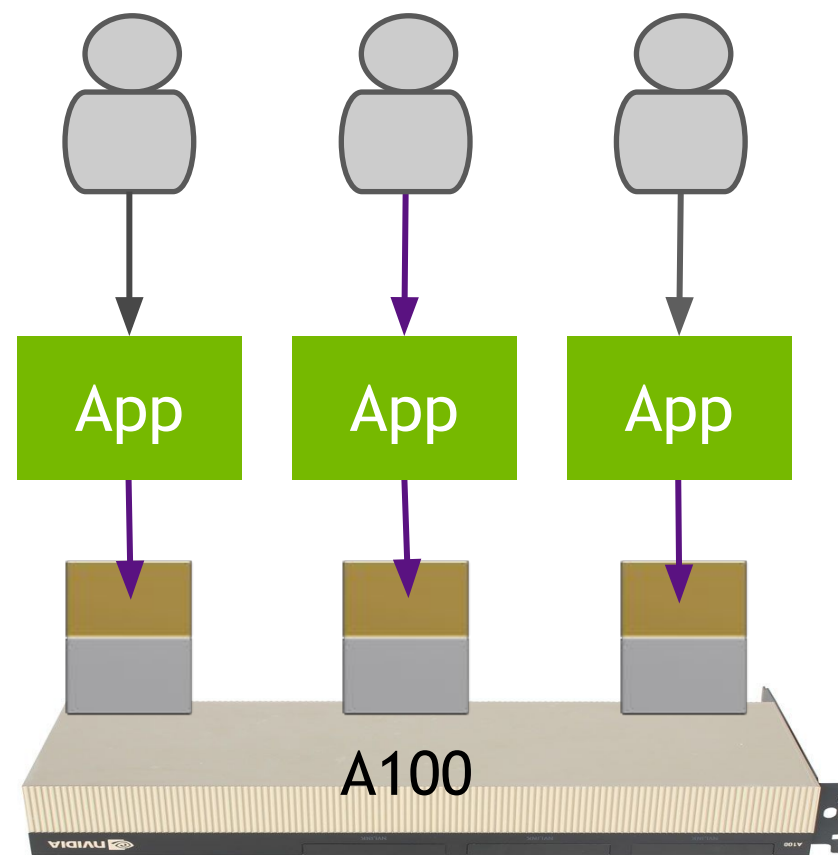
Use Cases

Single User → Multiple Apps



E.g. *Multiple inference jobs*

Single Tenant → Multi-User

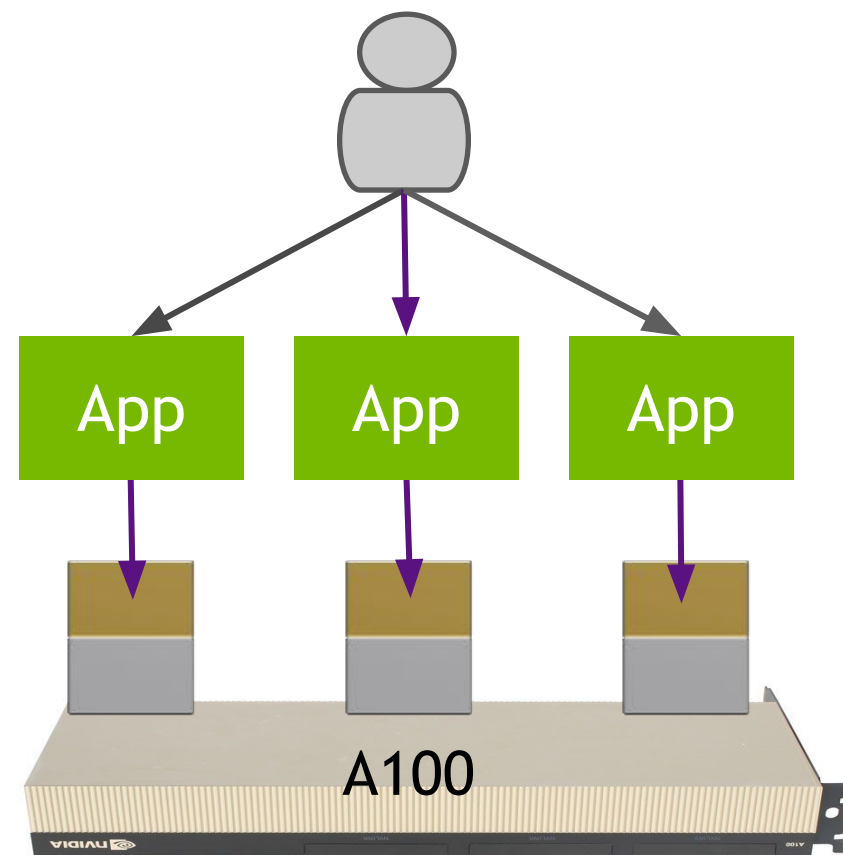


E.g. *Jupyter notebooks for model exploration*

MULTI-INSTANCE GPUs (MIG)

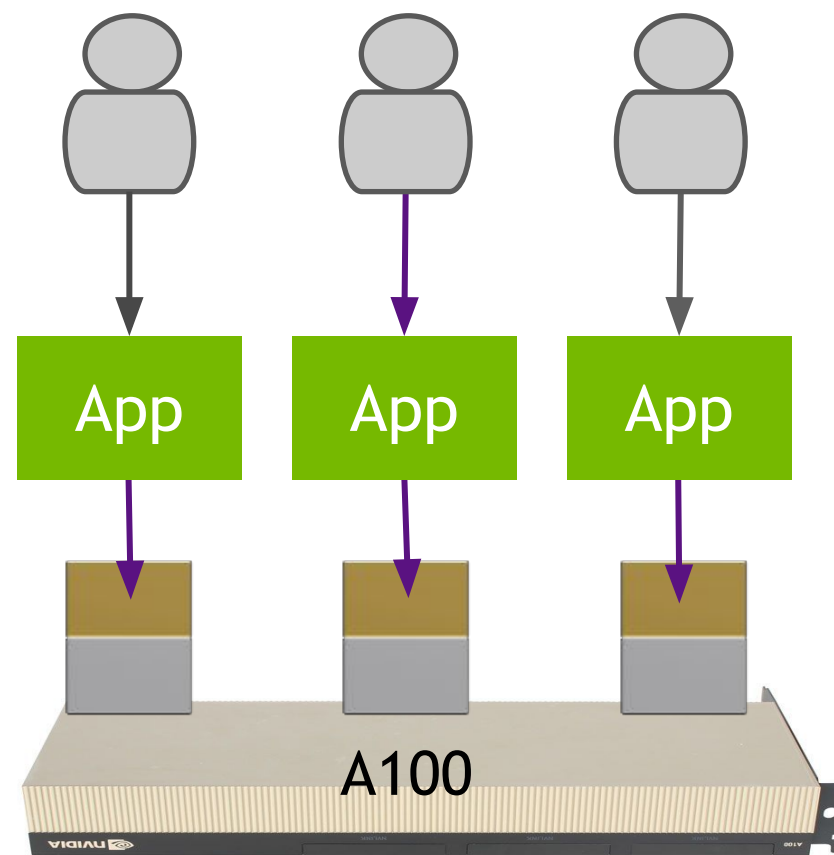
Use Cases

Single User → Multiple Apps



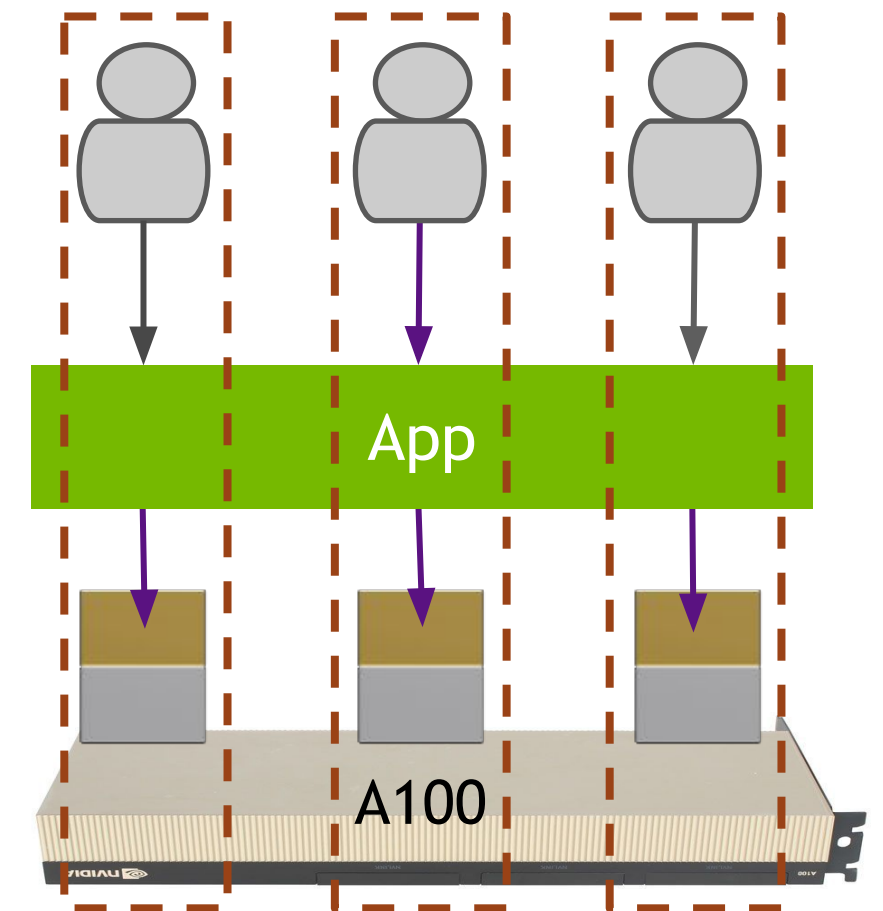
E.g. *Multiple inference jobs*

Single Tenant → Multi-User



E.g. *Jupyter notebooks for model exploration*

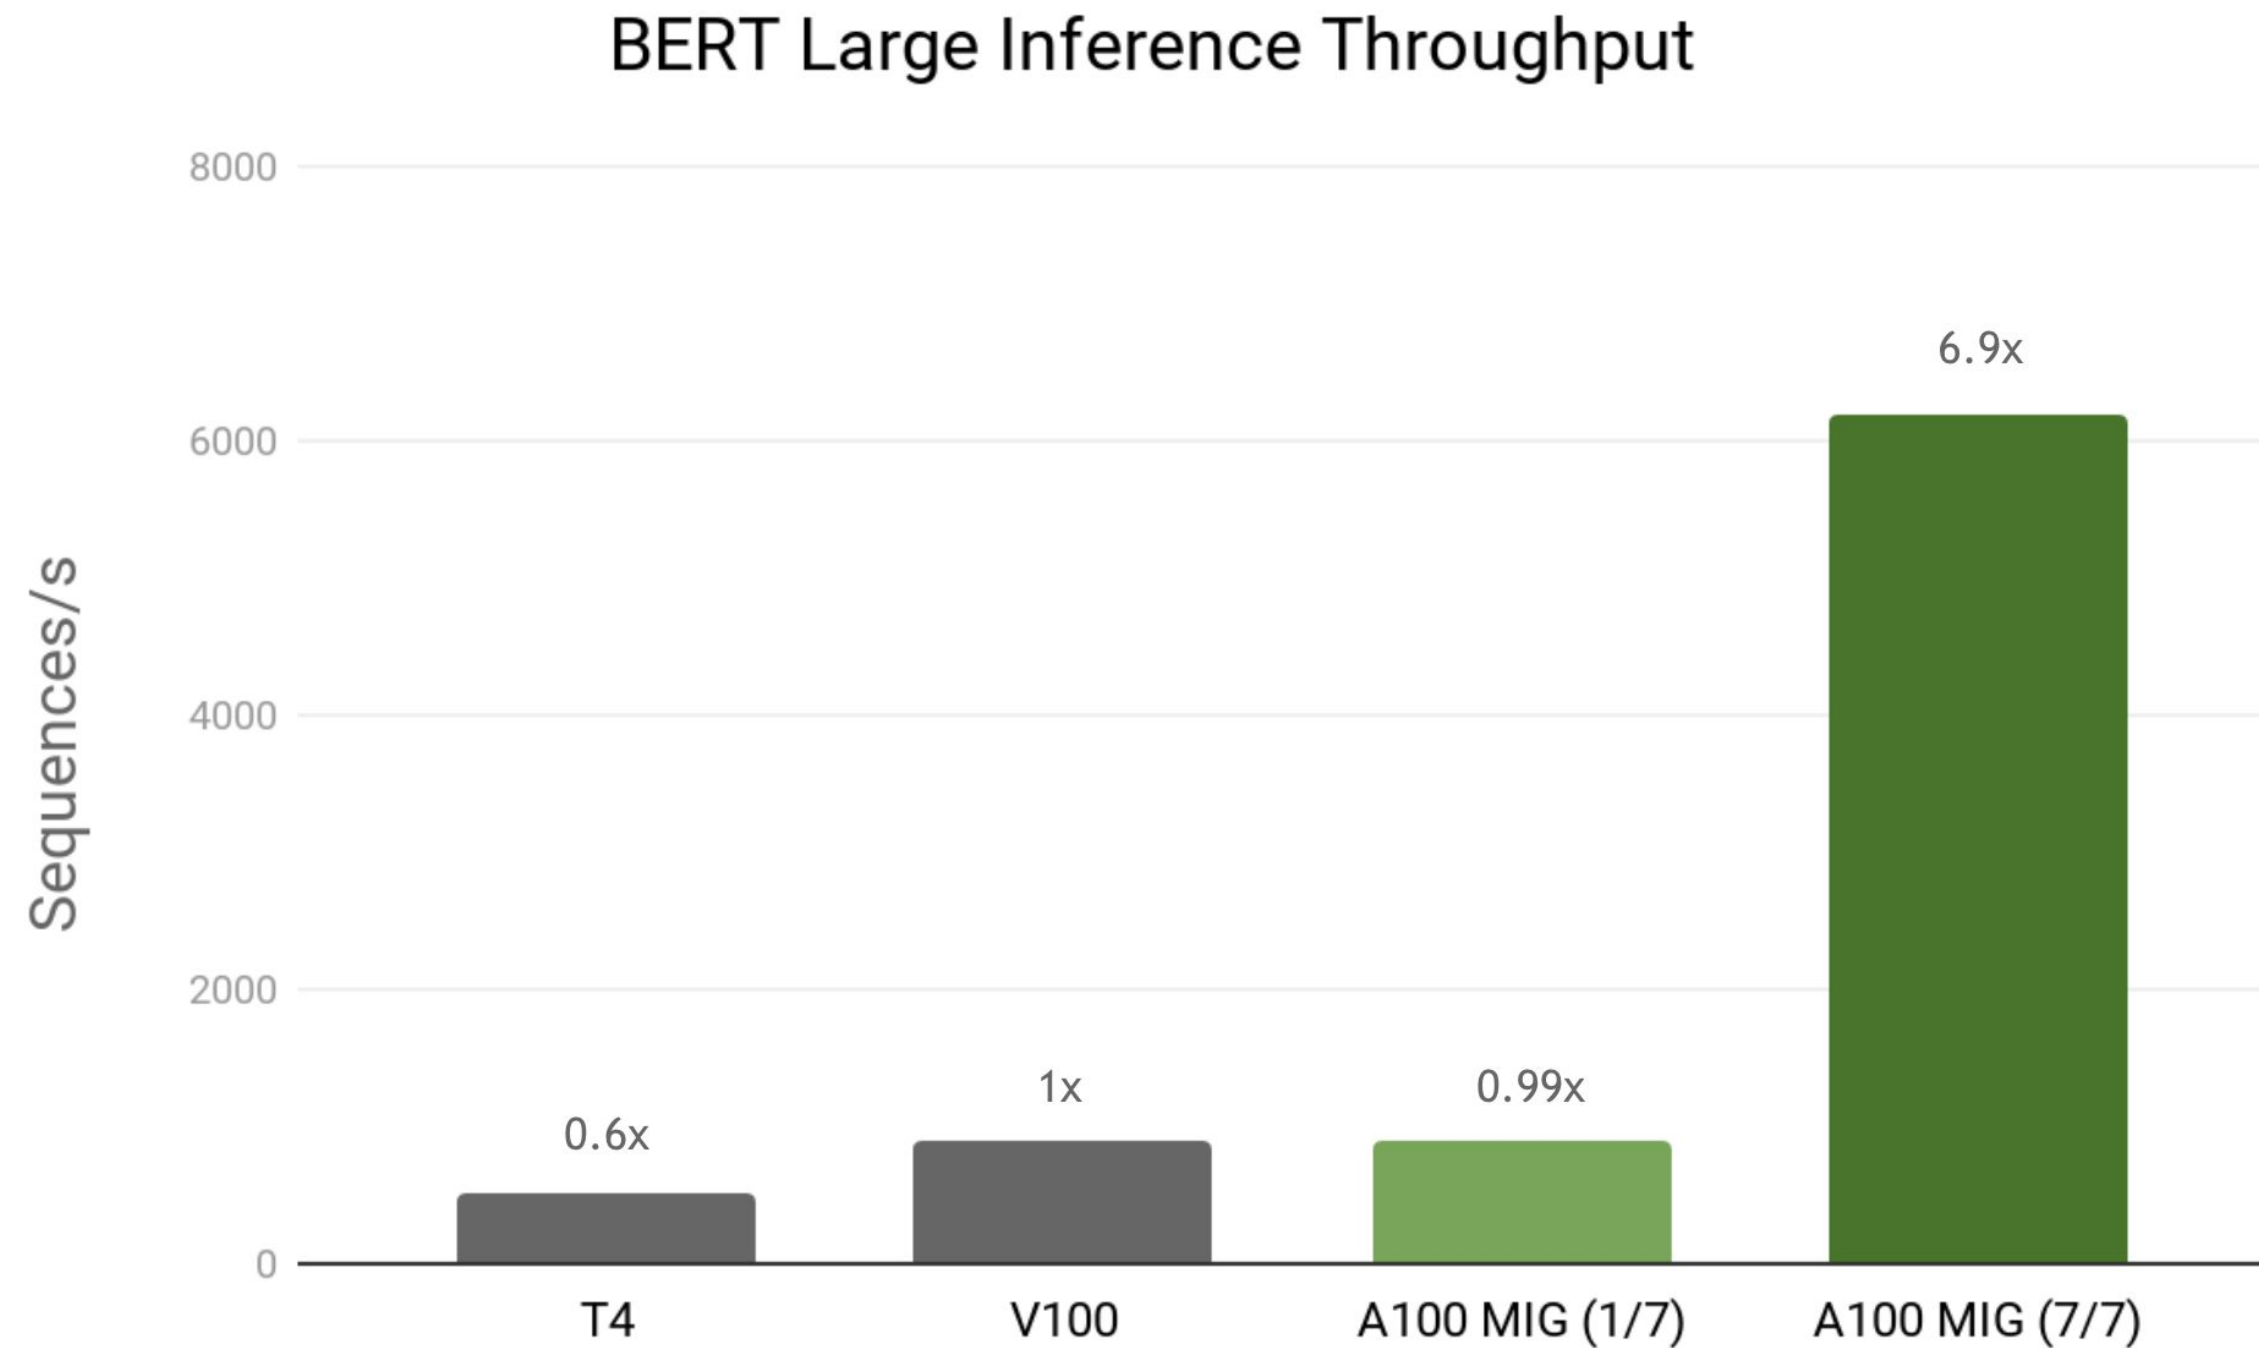
Multi-Tenant → Multi-User



E.g. *Managed Cloud Services*

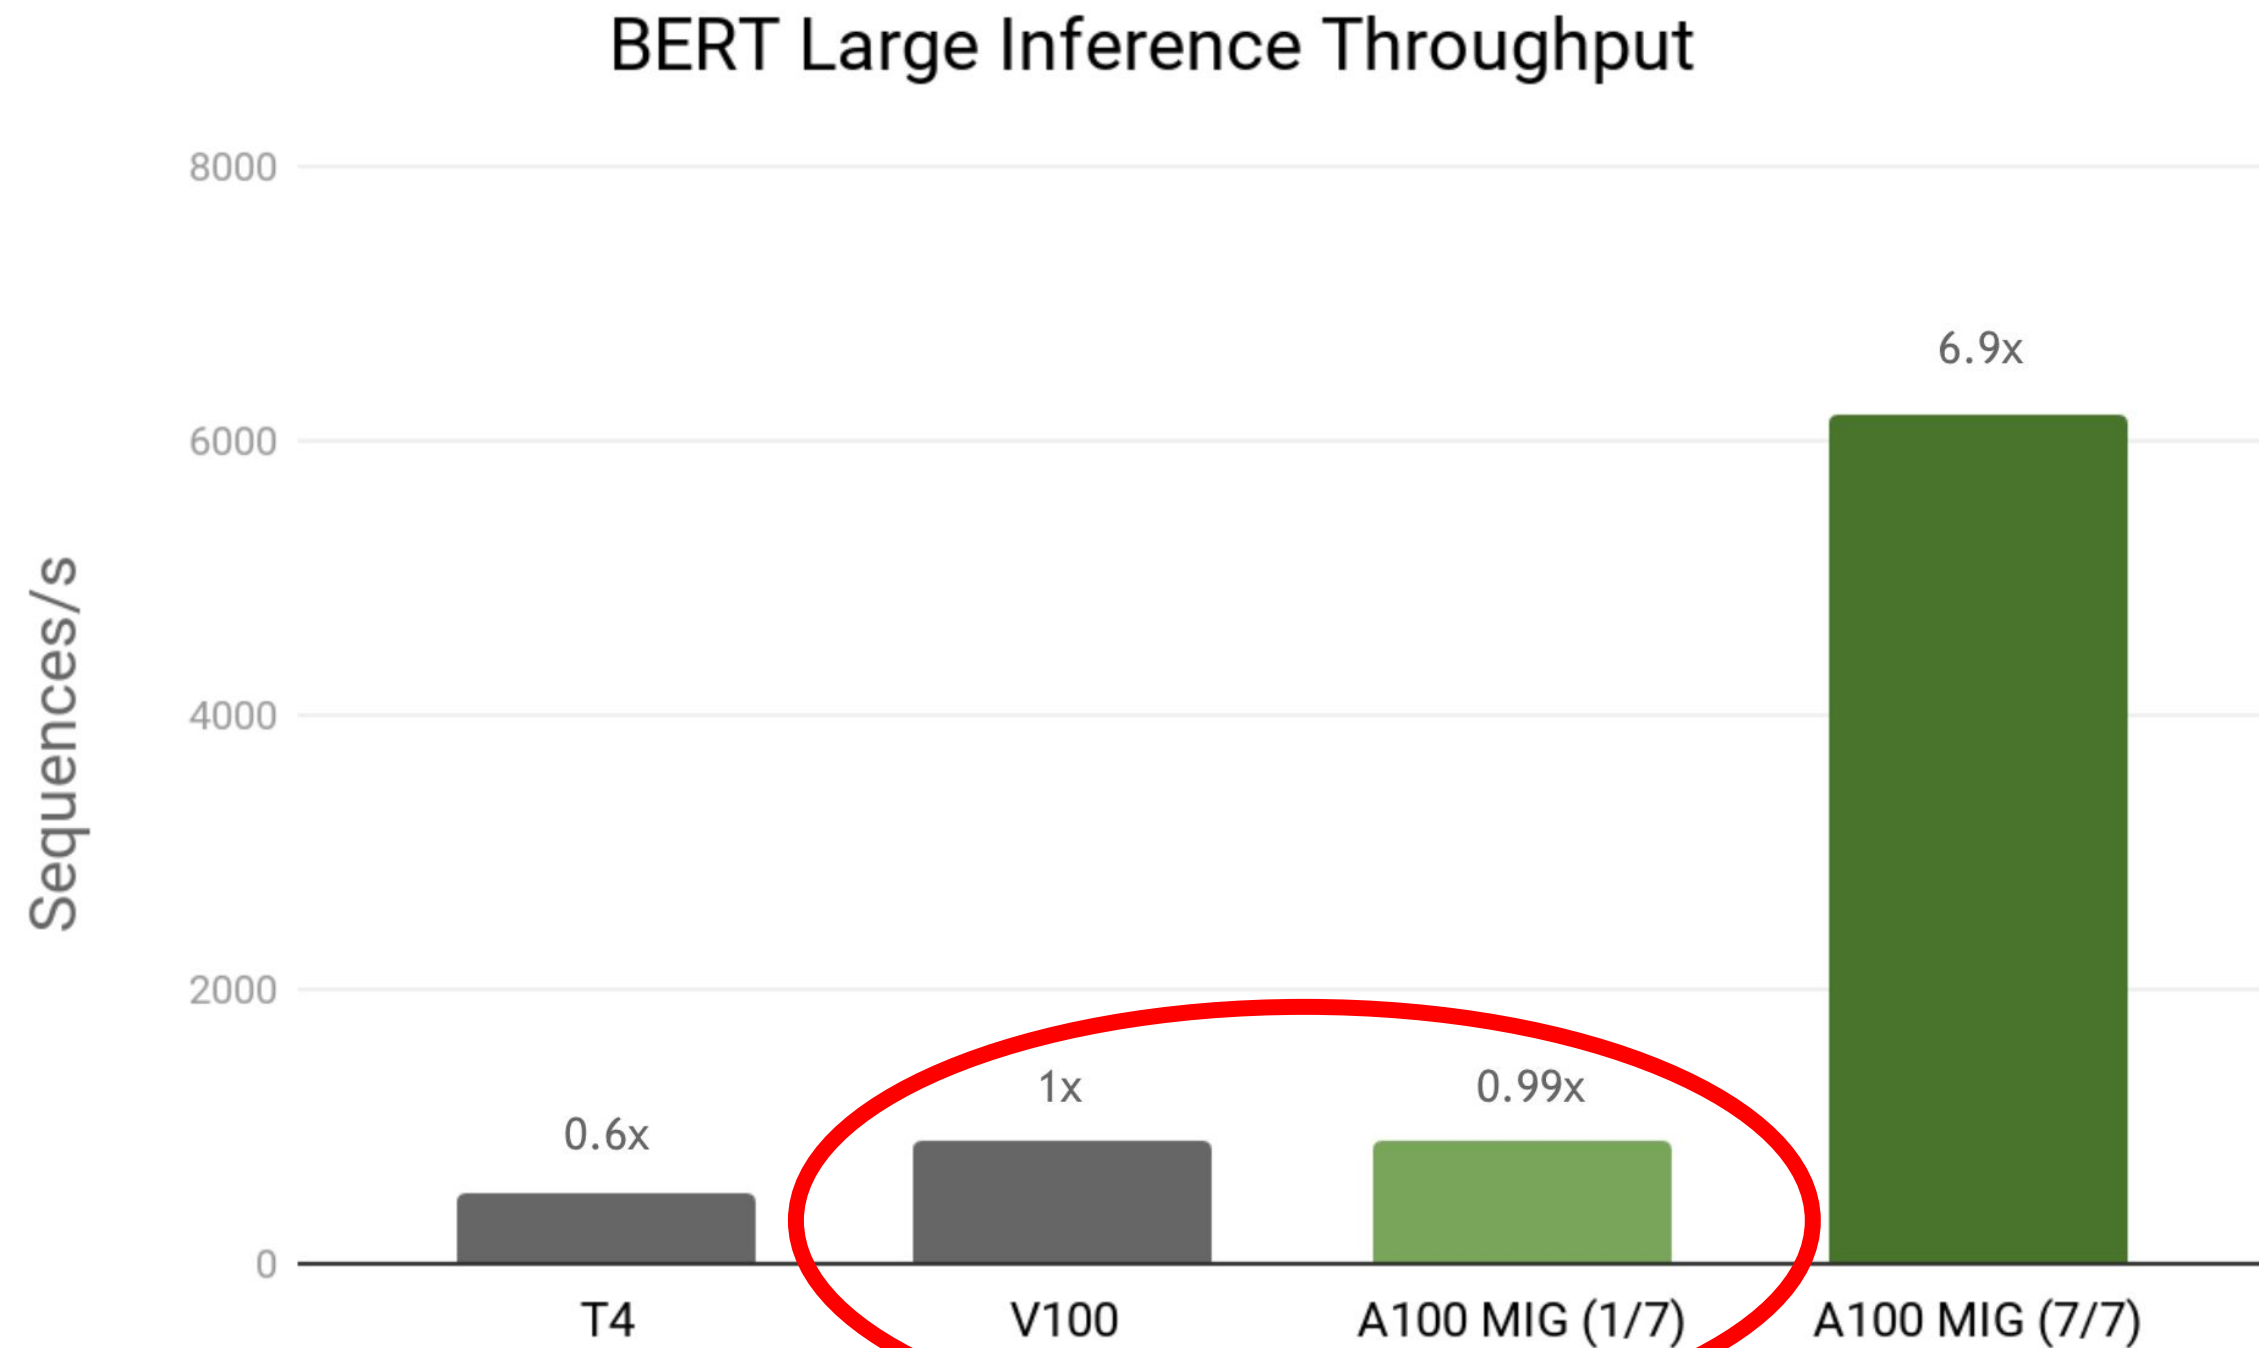
MULTI-INSTANCE GPUs (MIG)

Relative Performance



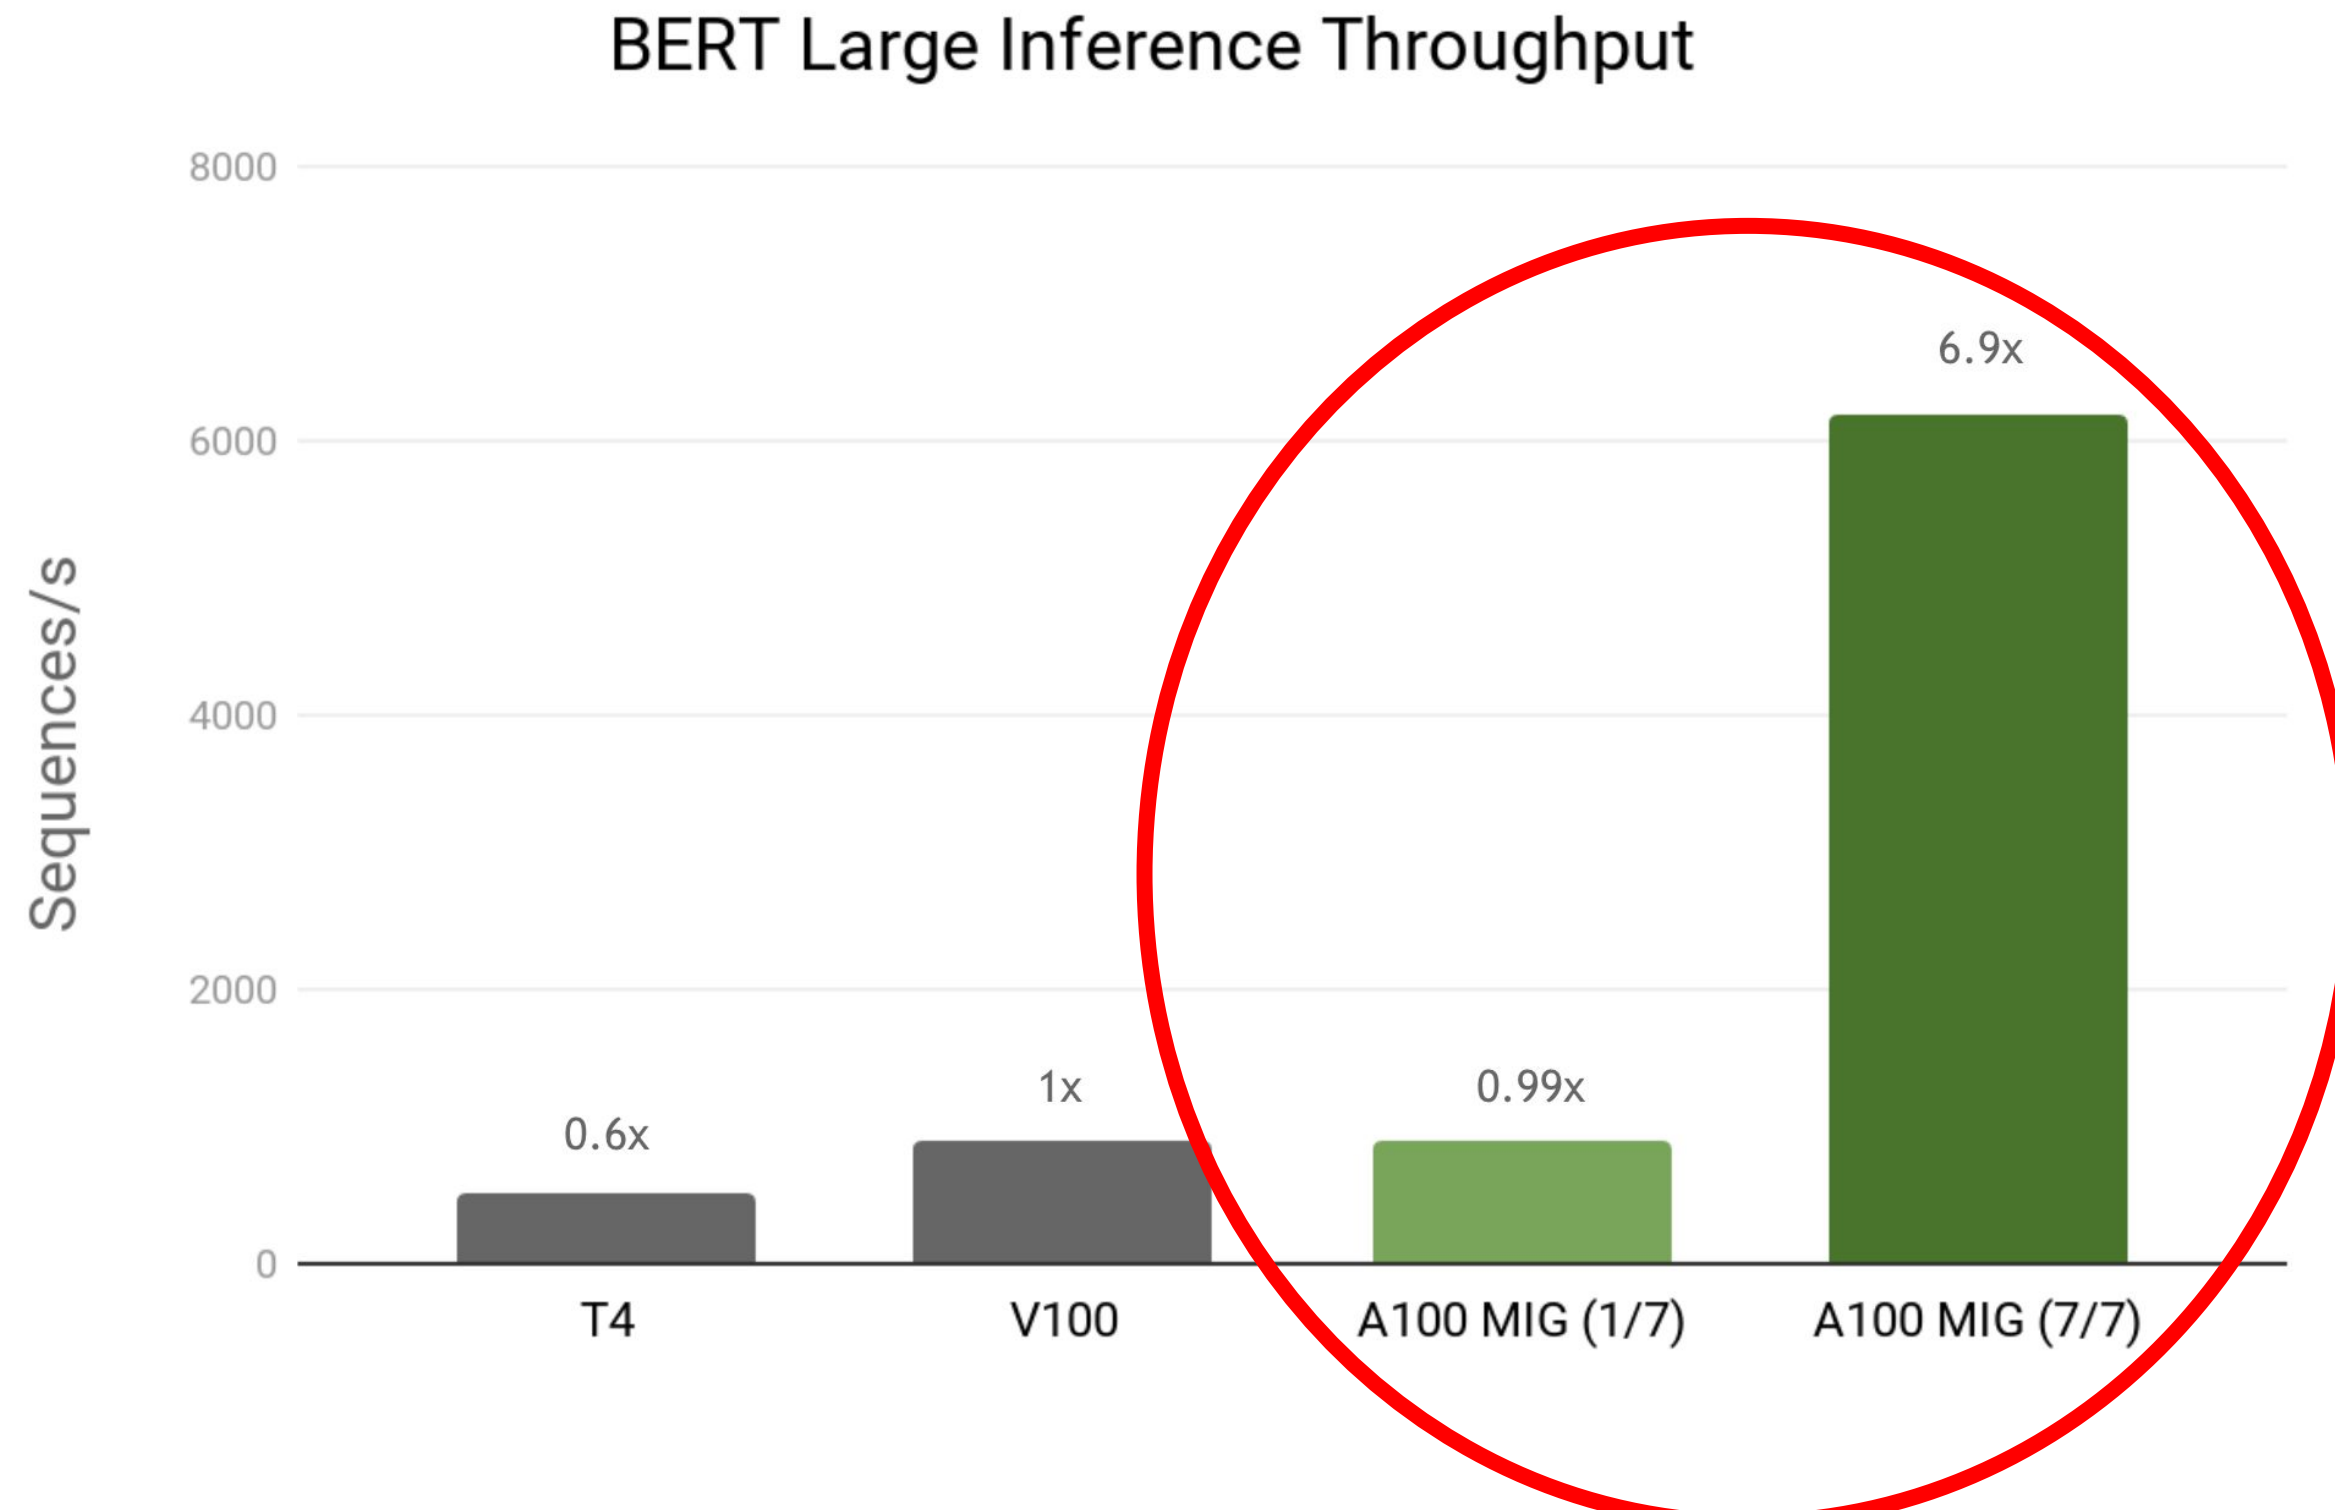
MULTI-INSTANCE GPUs (MIG)

Relative Performance



MULTI-INSTANCE GPUs (MIG)

Relative Performance



MULTI-INSTANCE GPUs (MIG)

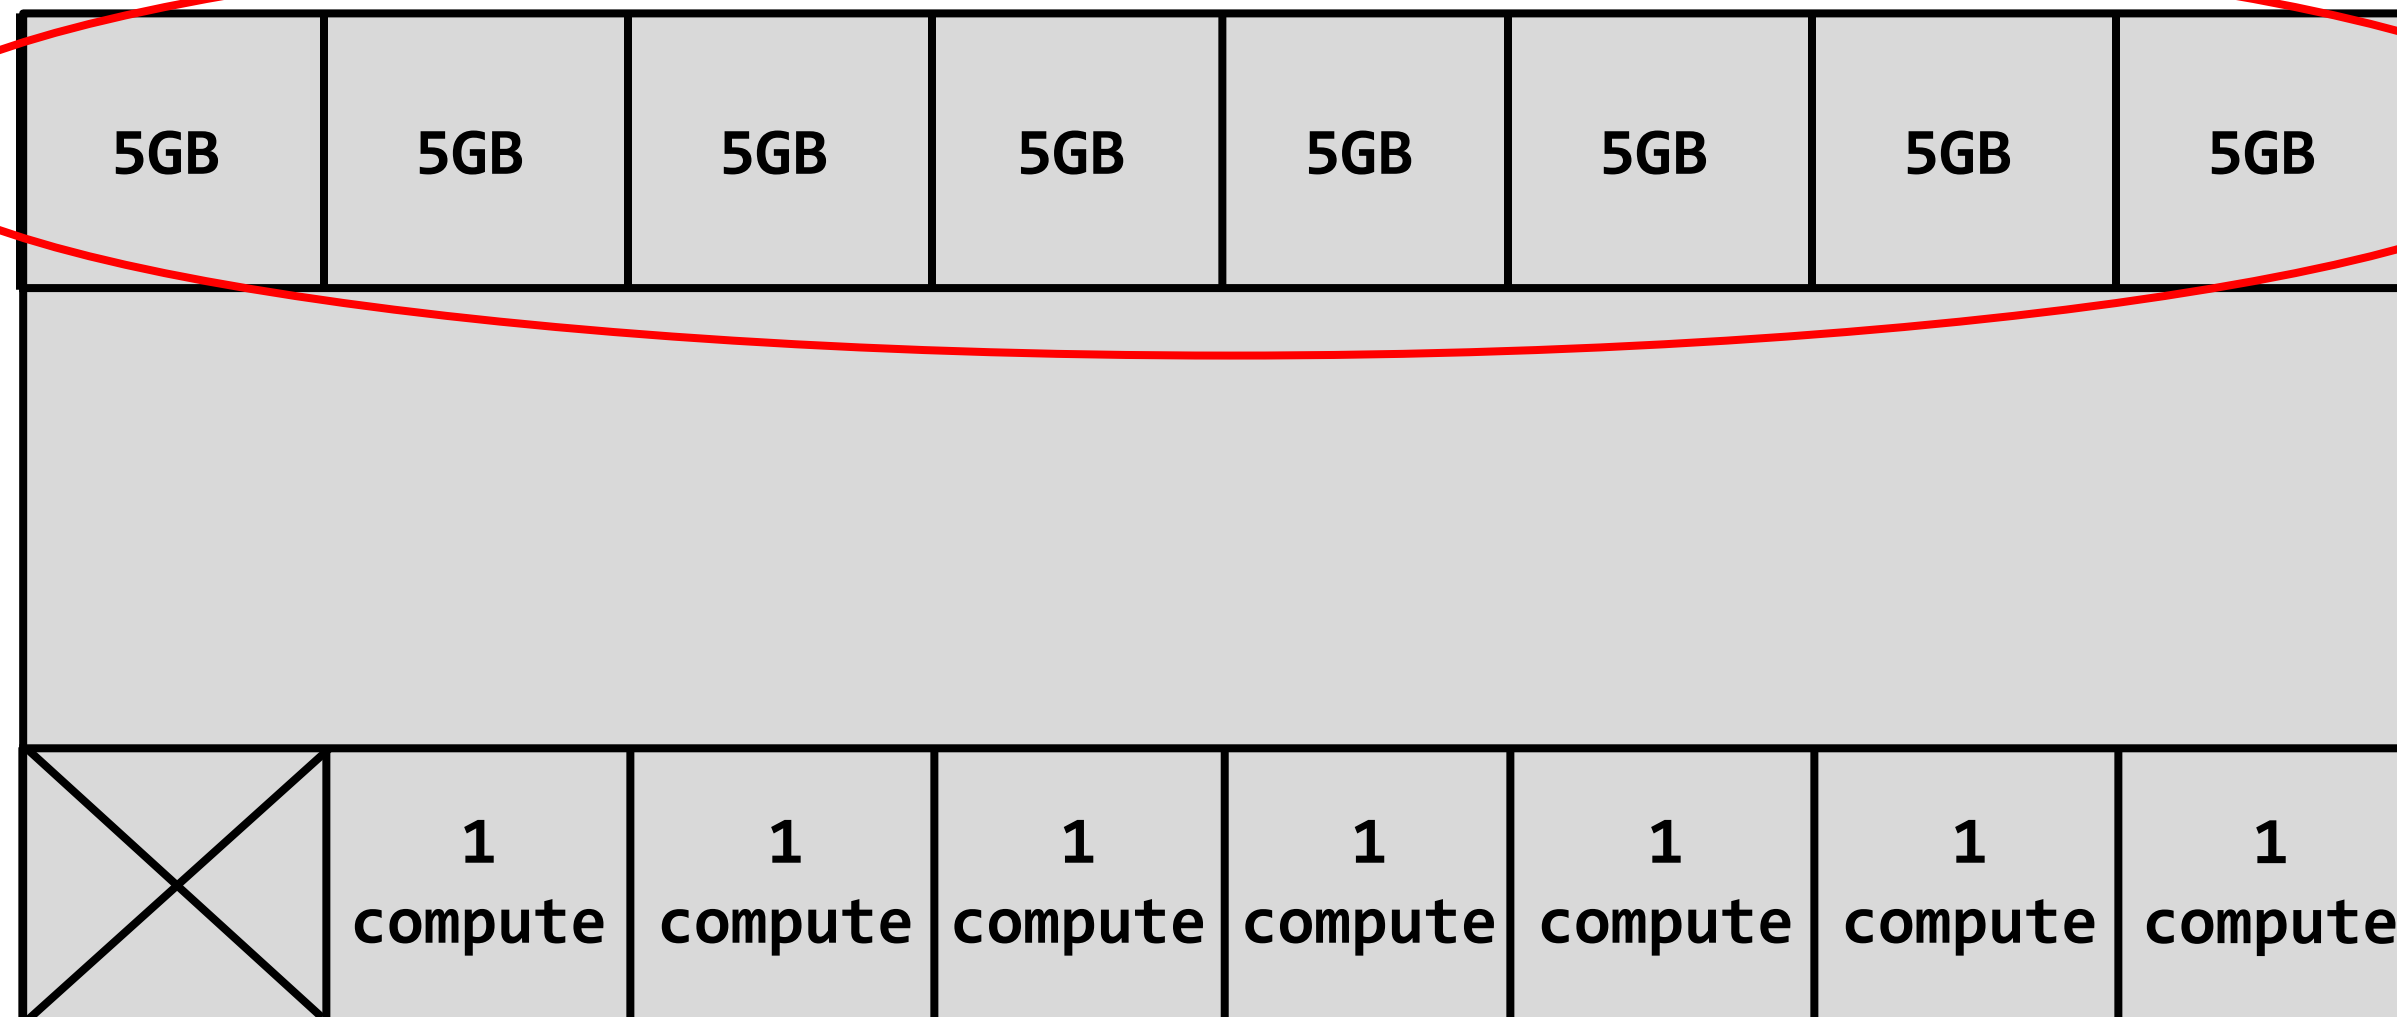
GPU Instances, Compute Instances, and MIG Devices

5GB	5GB	5GB	5GB	5GB	5GB	5GB	5GB
	1 compute	1 compute	1 compute	1 compute	1 compute	1 compute	1 compute

NVIDIA A100 (40GB)

MULTI-INSTANCE GPUs (MIG)

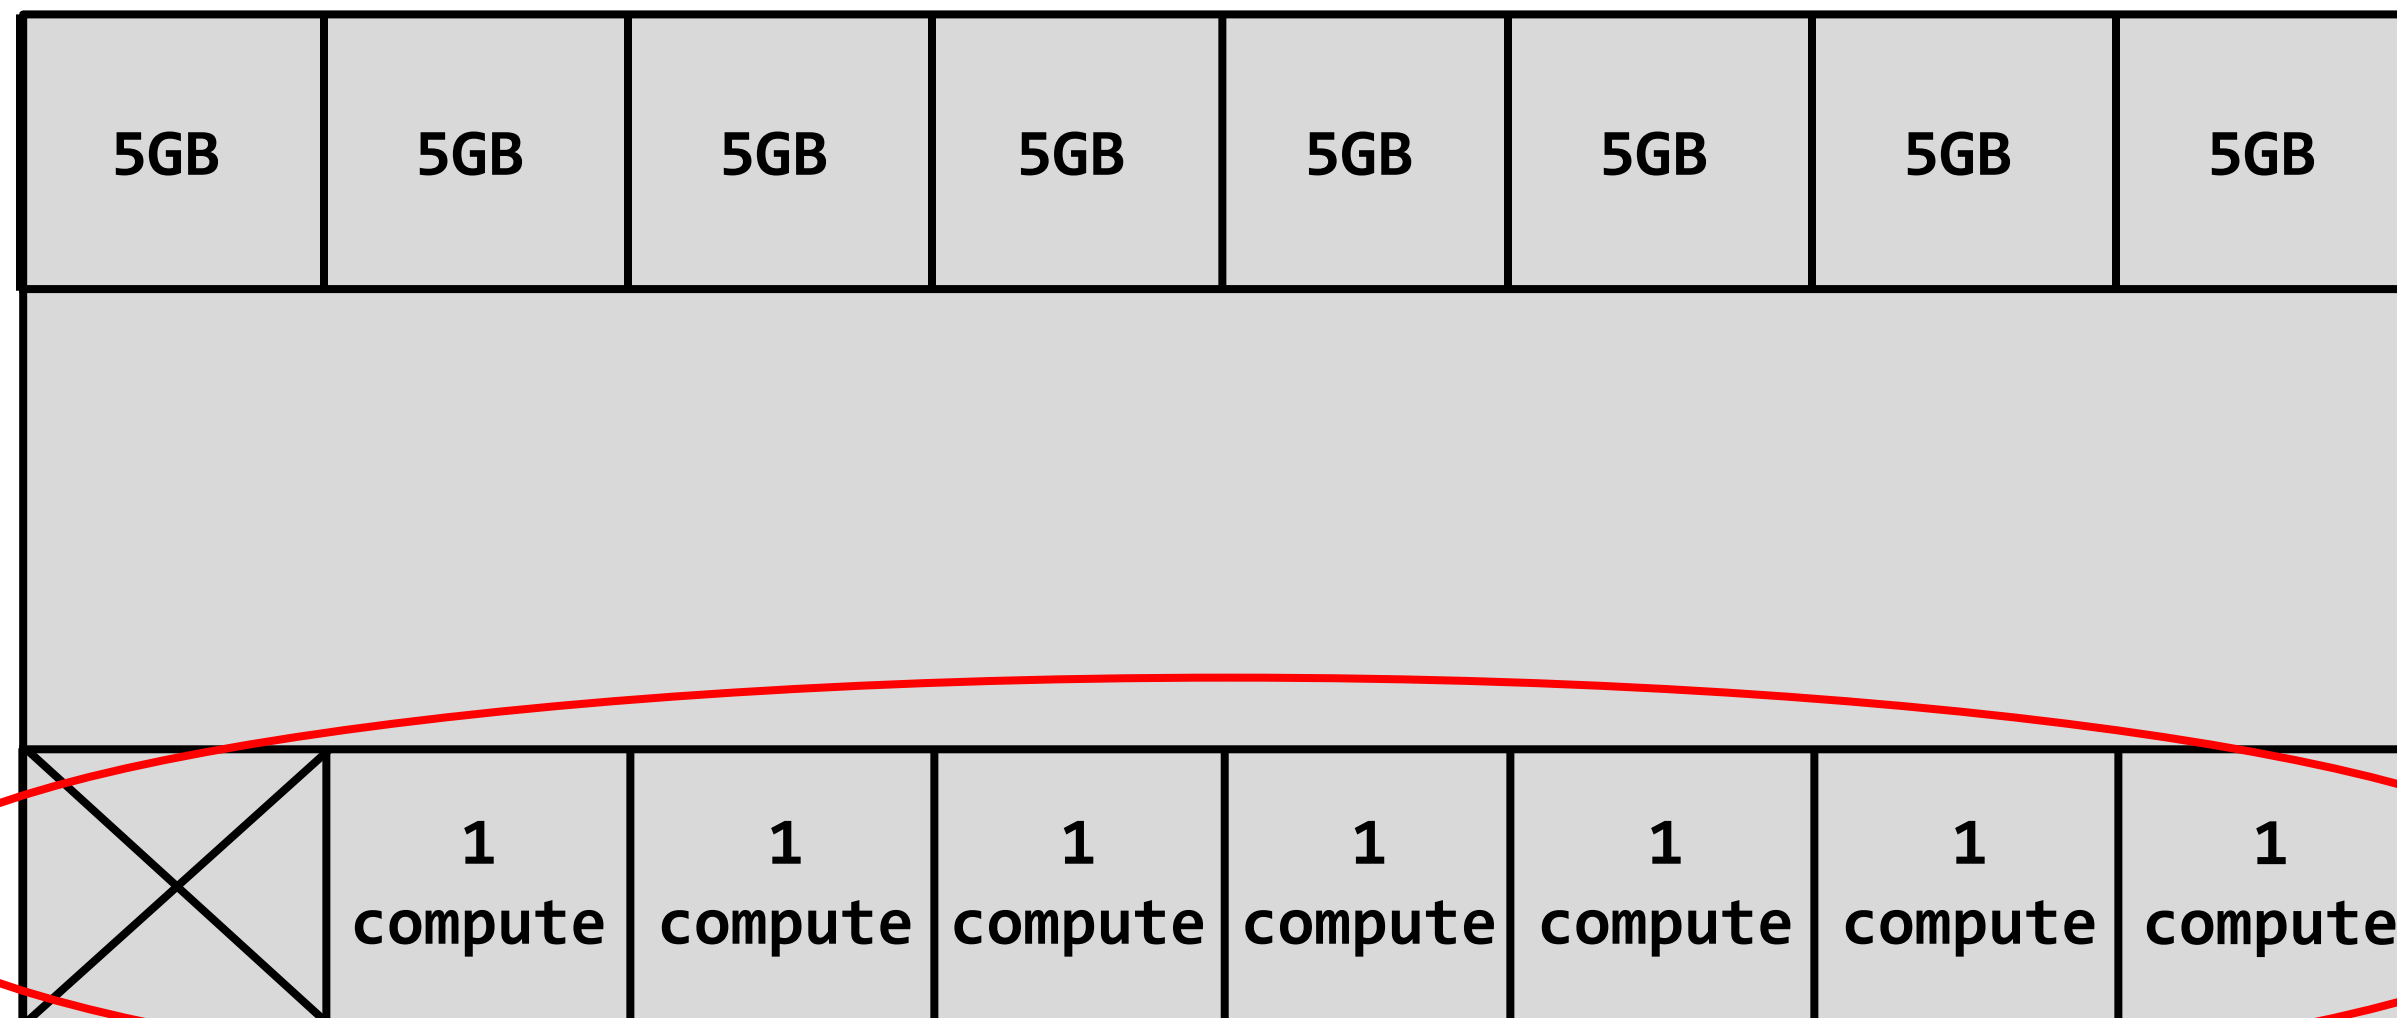
GPU Instances, Compute Instances, and MIG Devices



• NVIDIA A100 (40GB)
8 x 5GB Memory Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices



NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

10GB	10GB	10GB	10GB	10GB	10GB	10GB	10GB
	1 compute	1 compute	1 compute	1 compute	1 compute	1 compute	1 compute

NVIDIA A100 (80GB)

- 8 x 10GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

6GB	6GB	6GB	6GB
1 compute	1 compute	1 compute	1 compute

NVIDIA A30 (24GB)

- 4 x 6GB Memory Slices
- 4 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

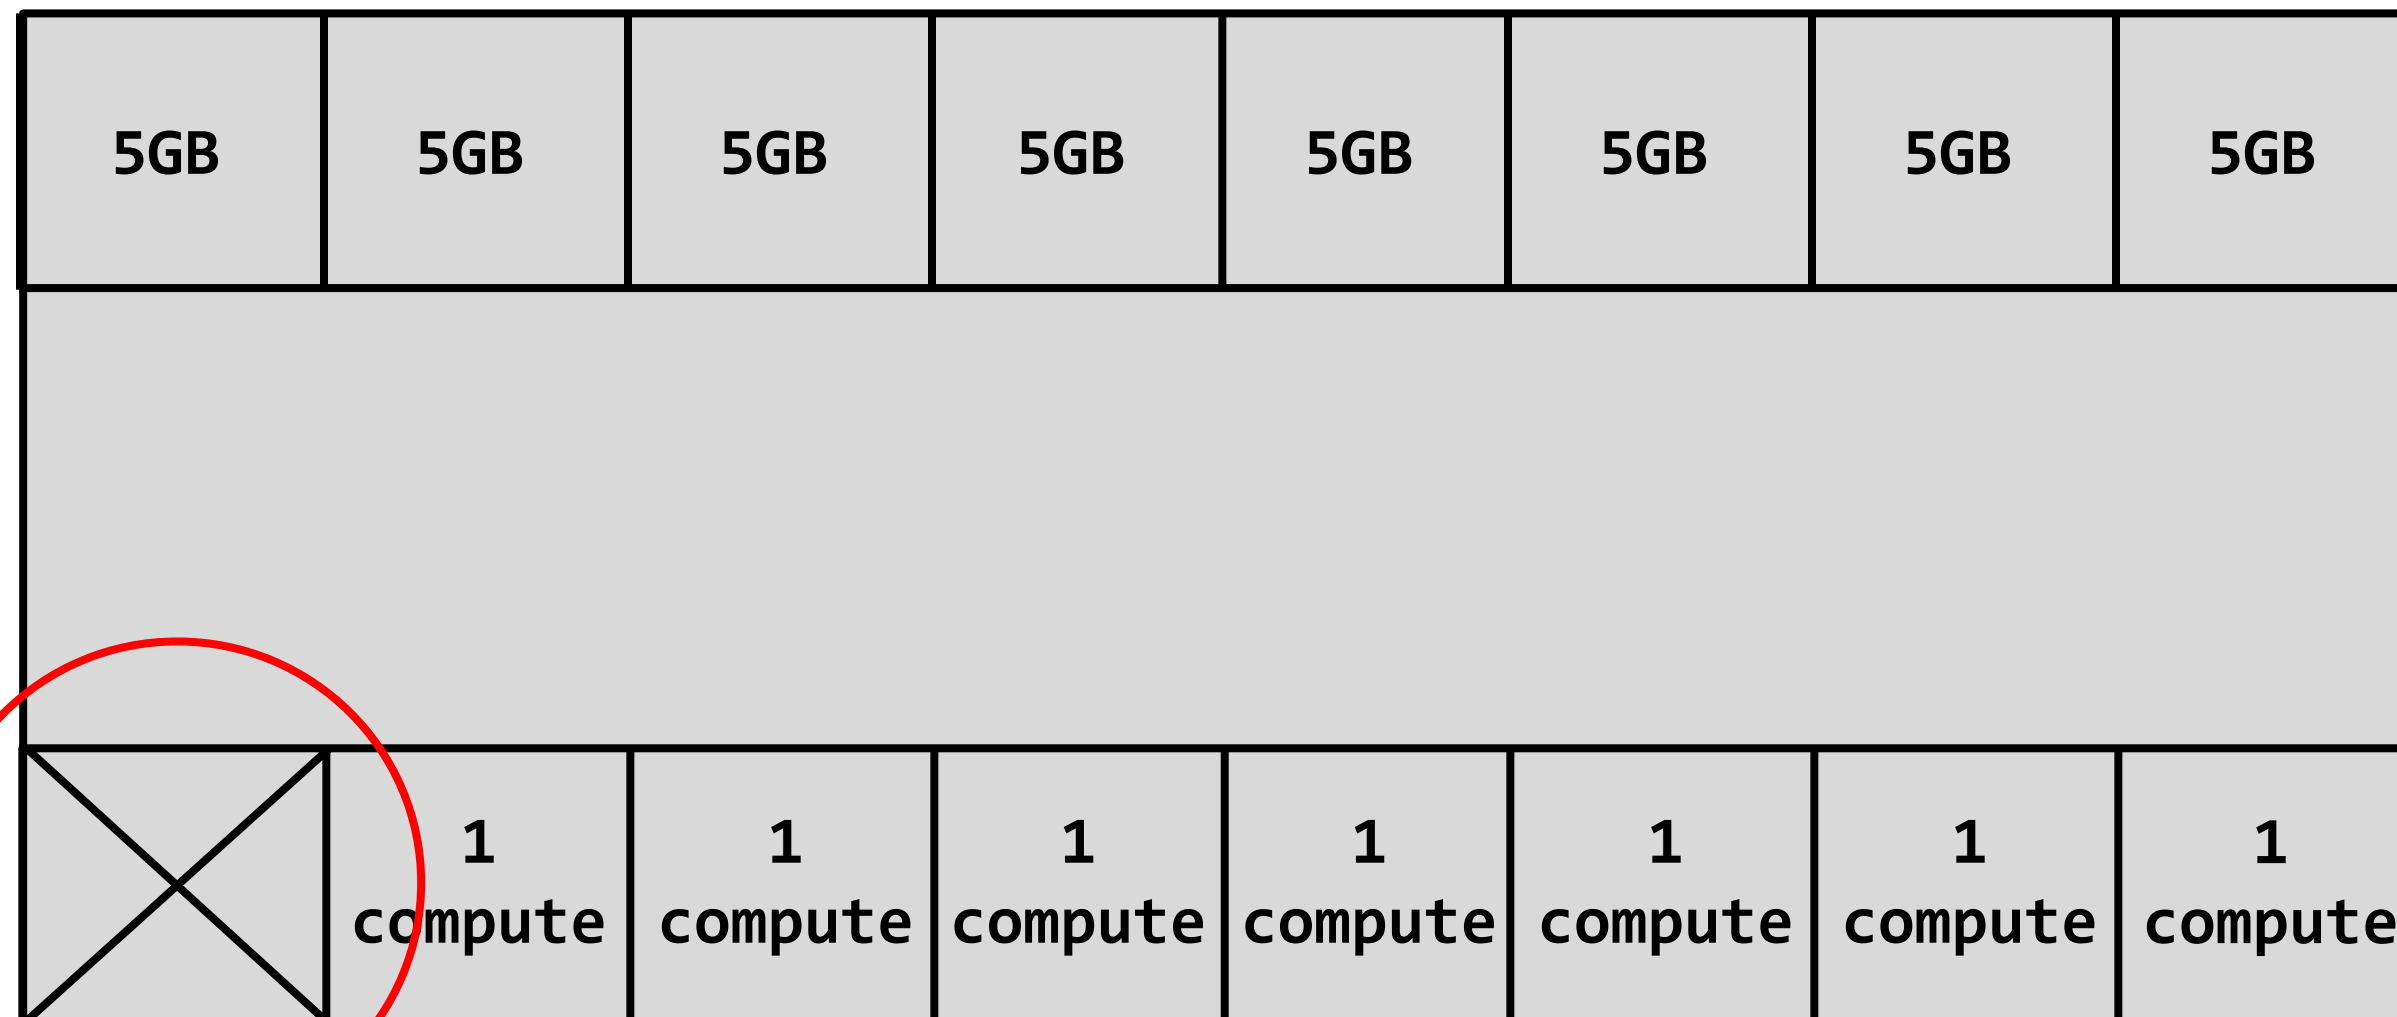
5GB	5GB	5GB	5GB	5GB	5GB	5GB	5GB
	1 compute	1 compute	1 compute	1 compute	1 compute	1 compute	1 compute

NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices



NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

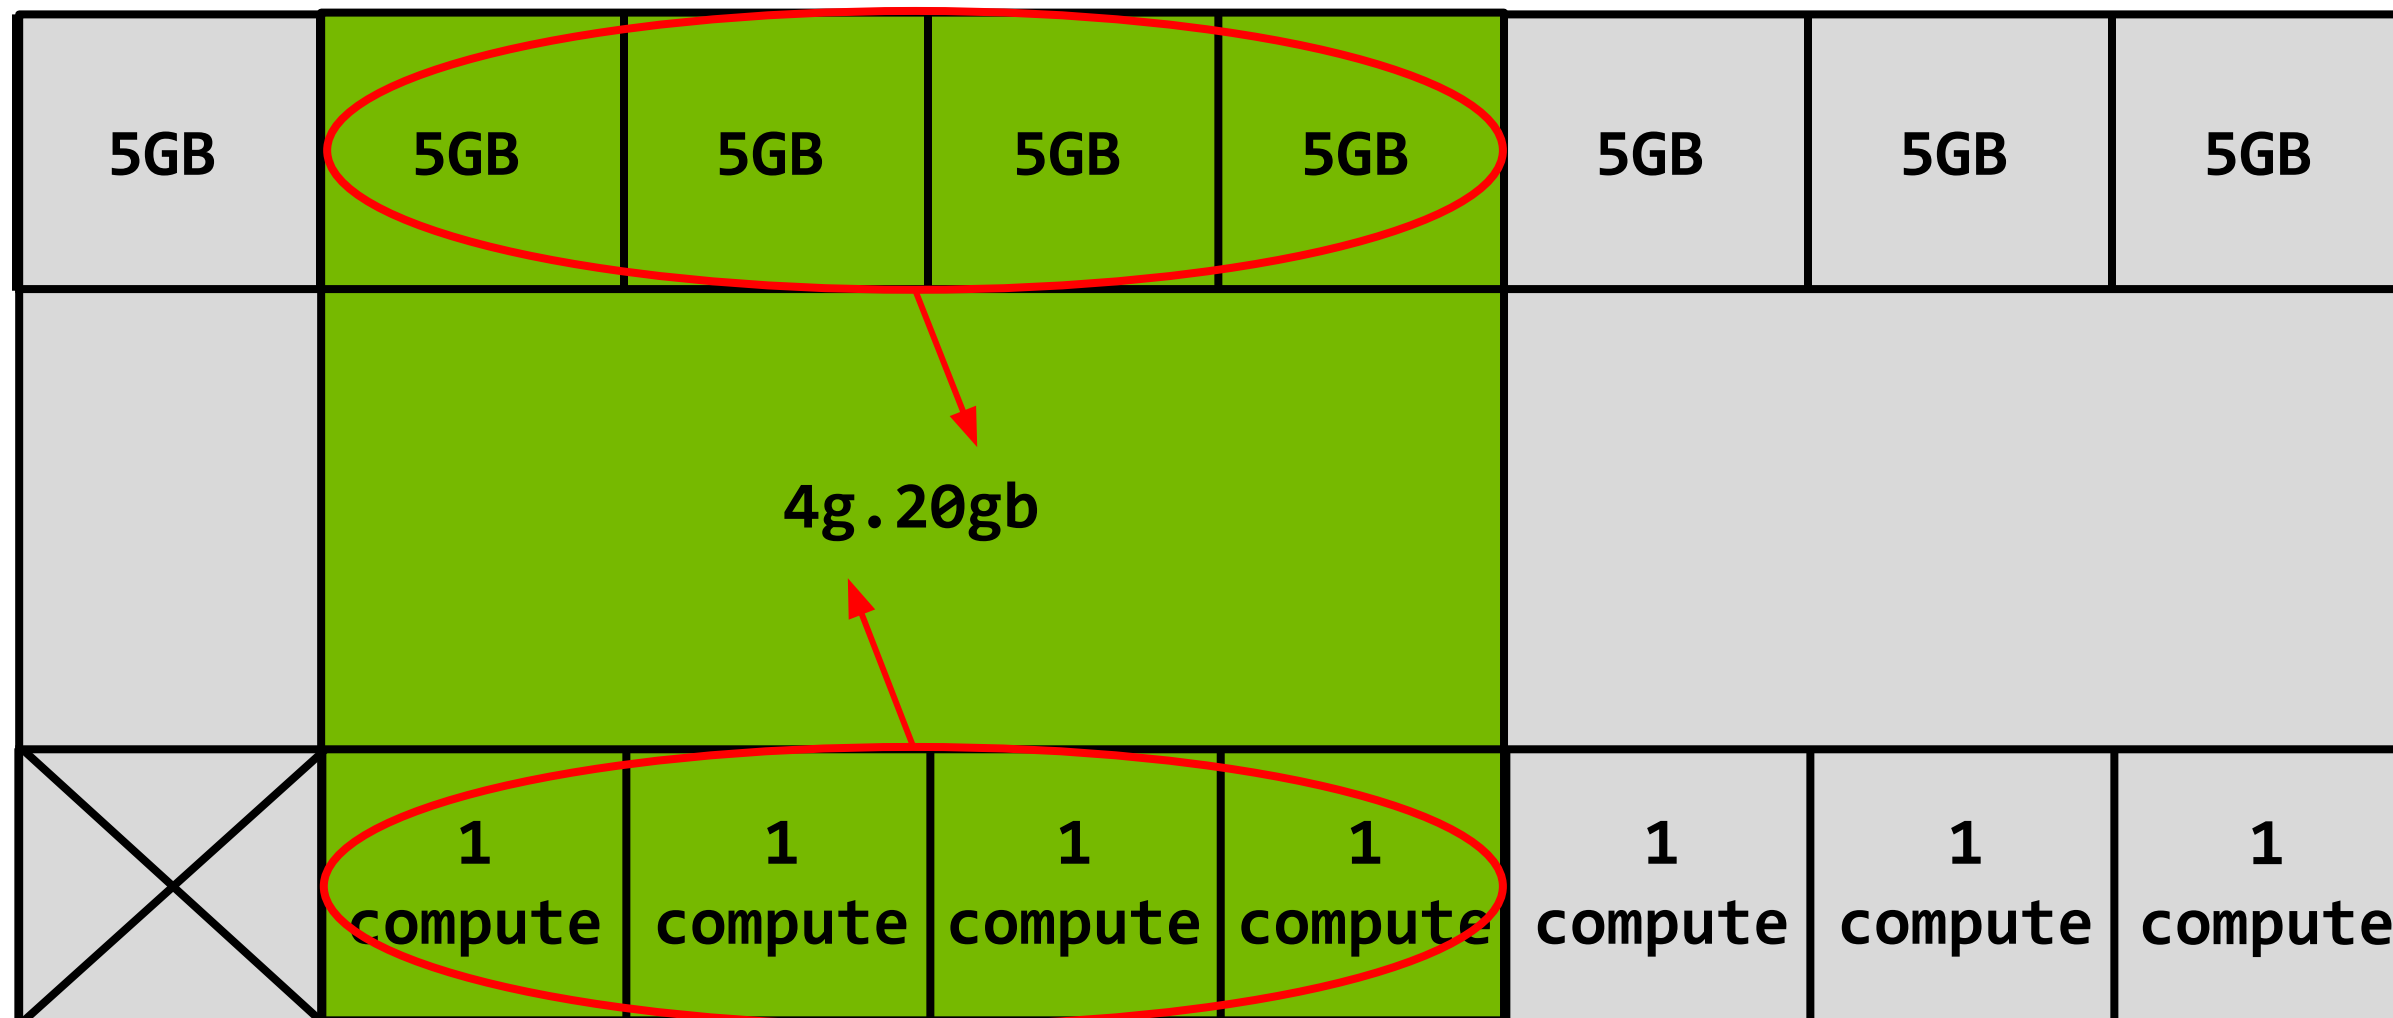
5GB	5GB	5GB	5GB	5GB	5GB	5GB	5GB
	1g.5gb	GPU Instance <ul style="list-style-type: none">• Fixed partition of memory and compute• Fixed amount of “other” GPU Engines (depending on size)					
	1 compute	1 compute	1 compute	1 compute	1 compute	1 compute	1 compute

NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

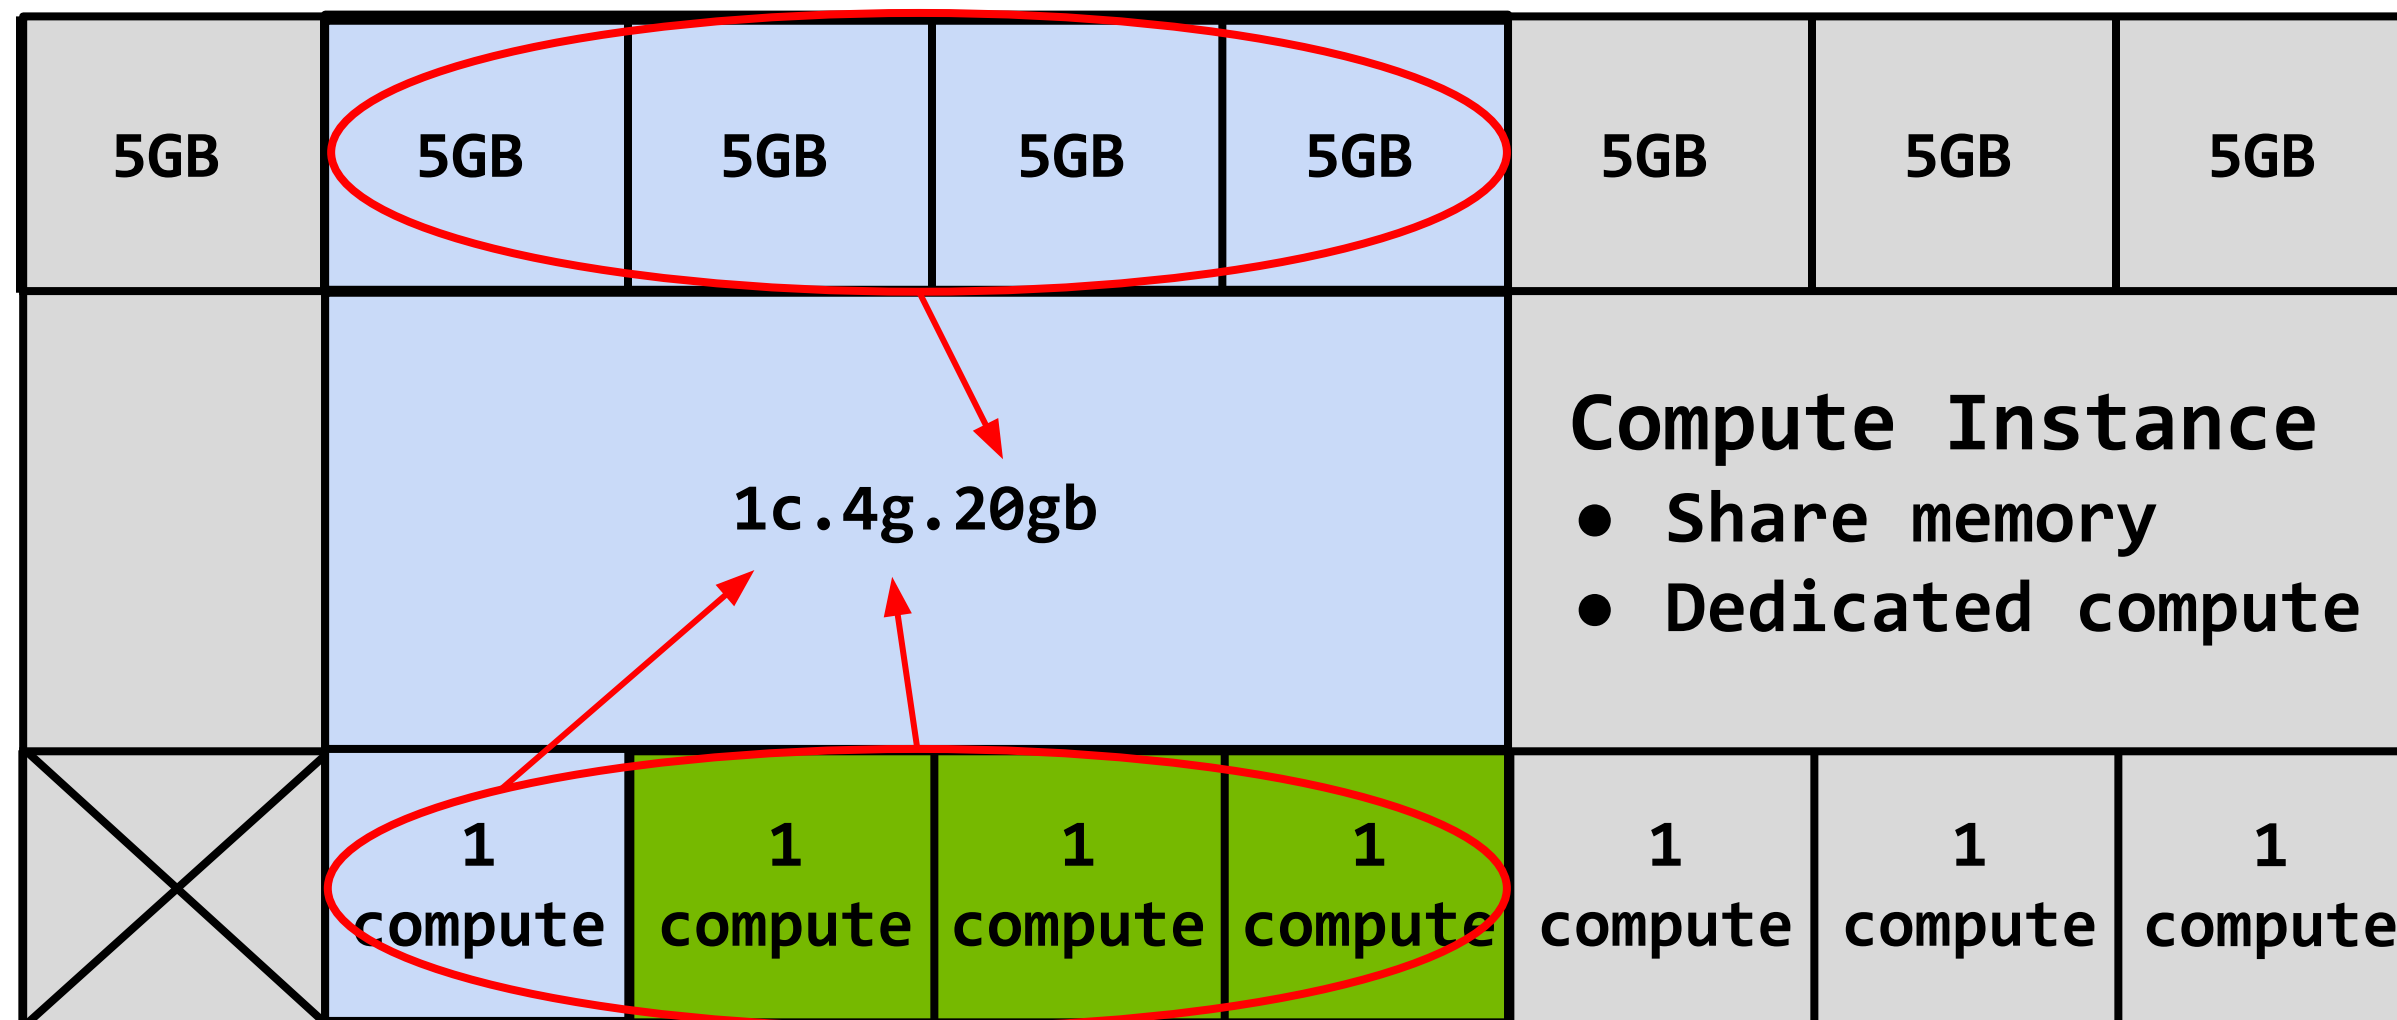


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

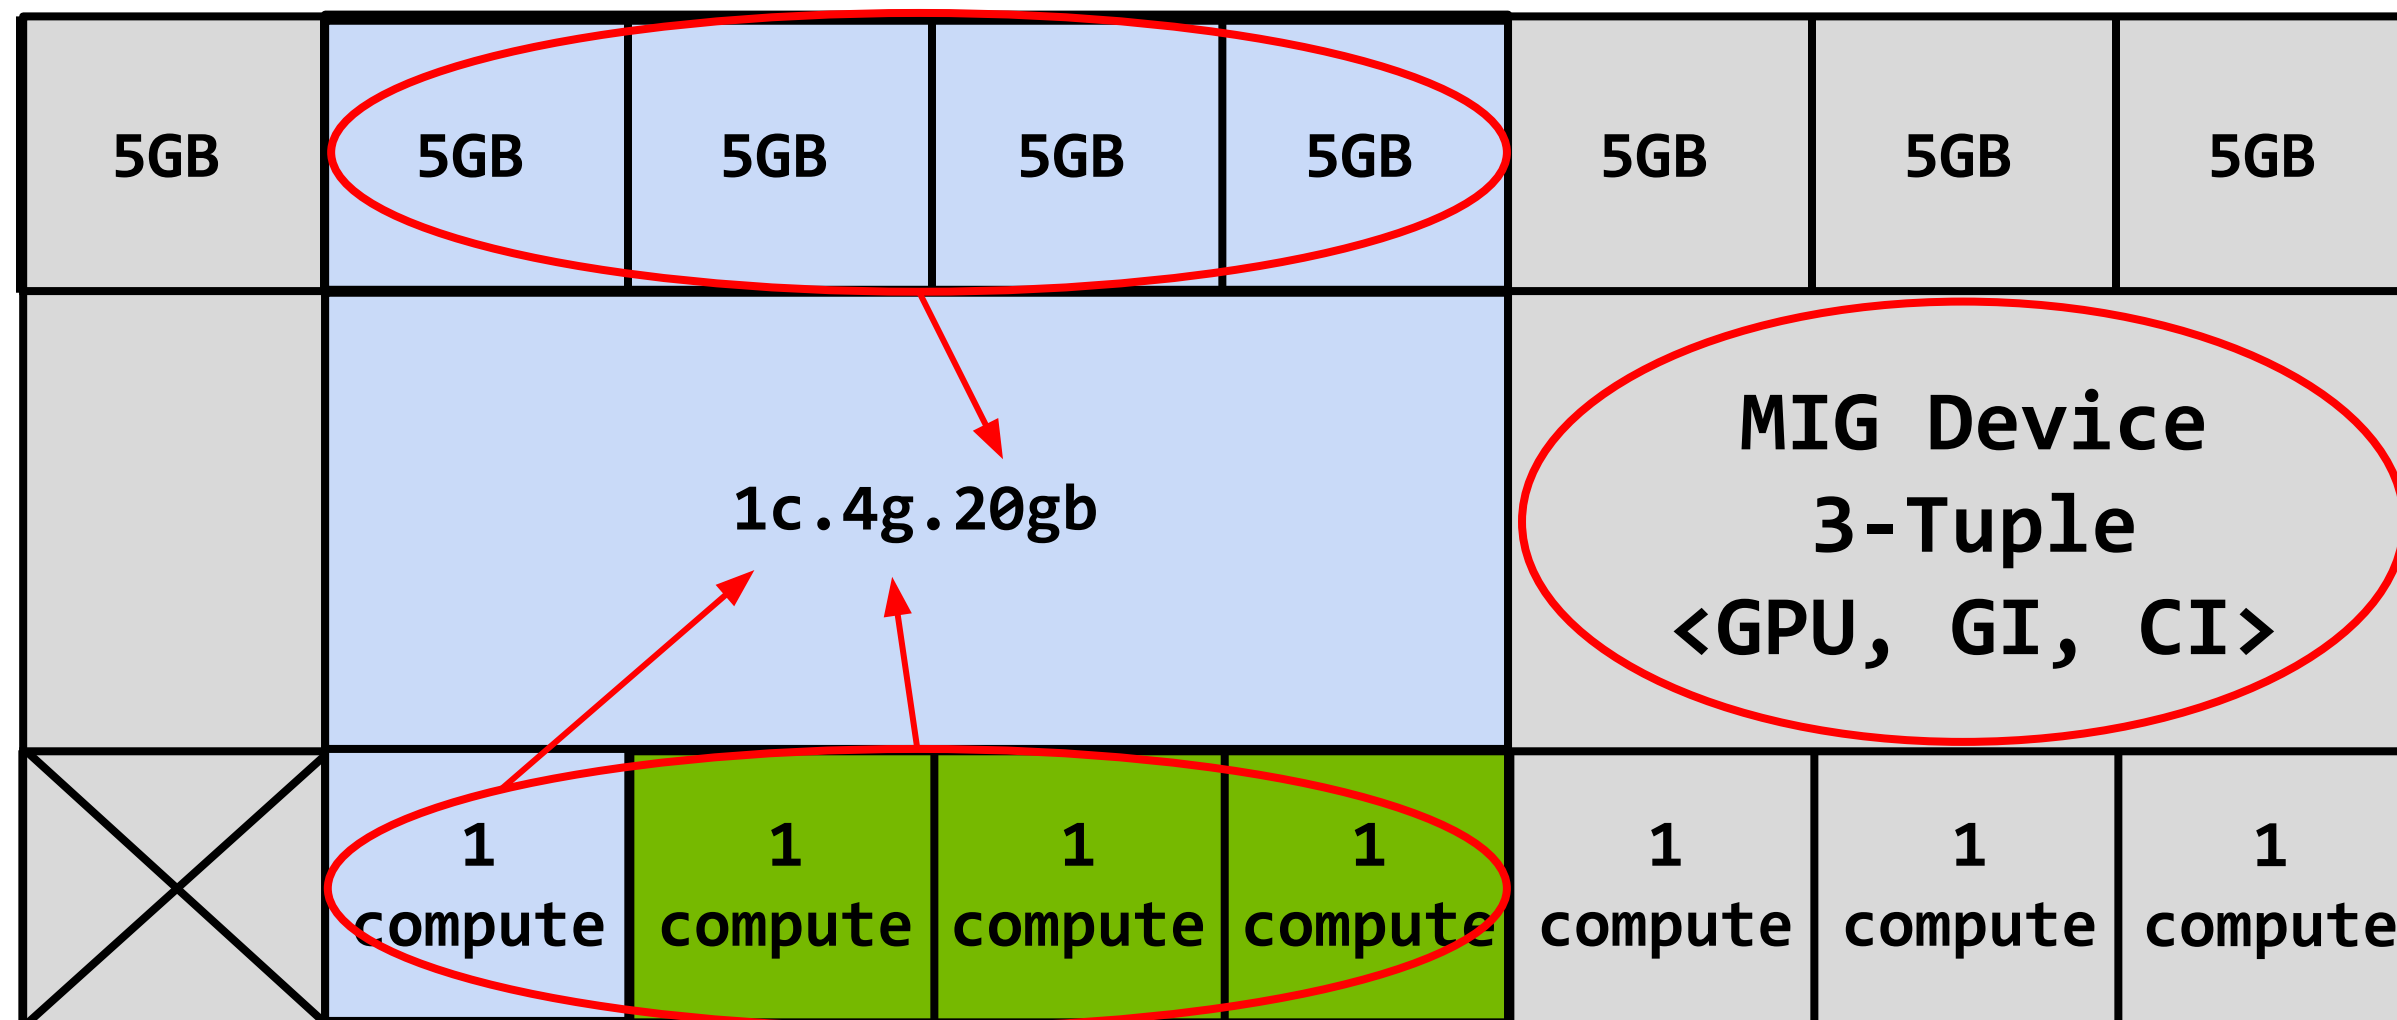


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

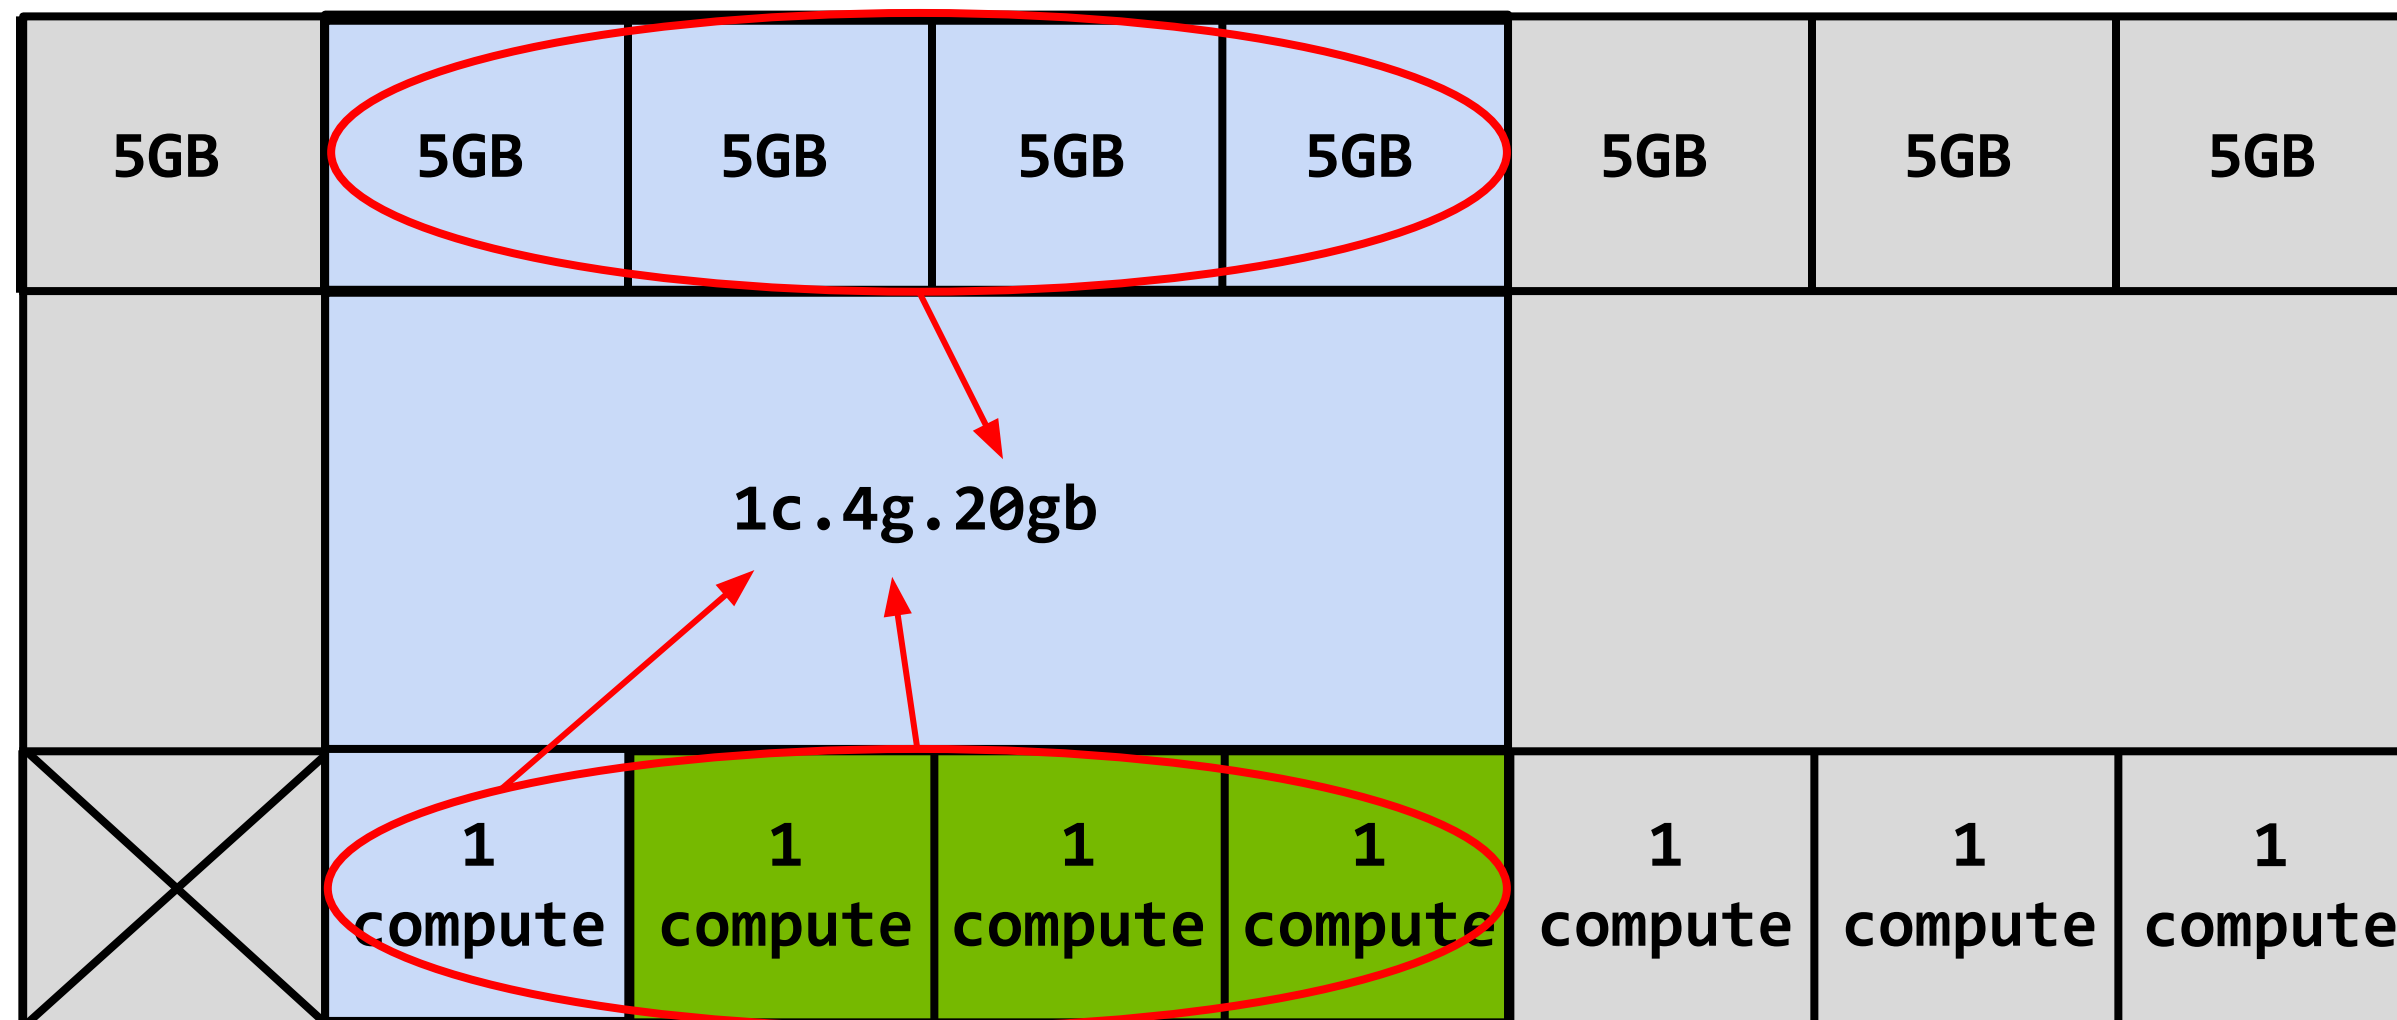


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

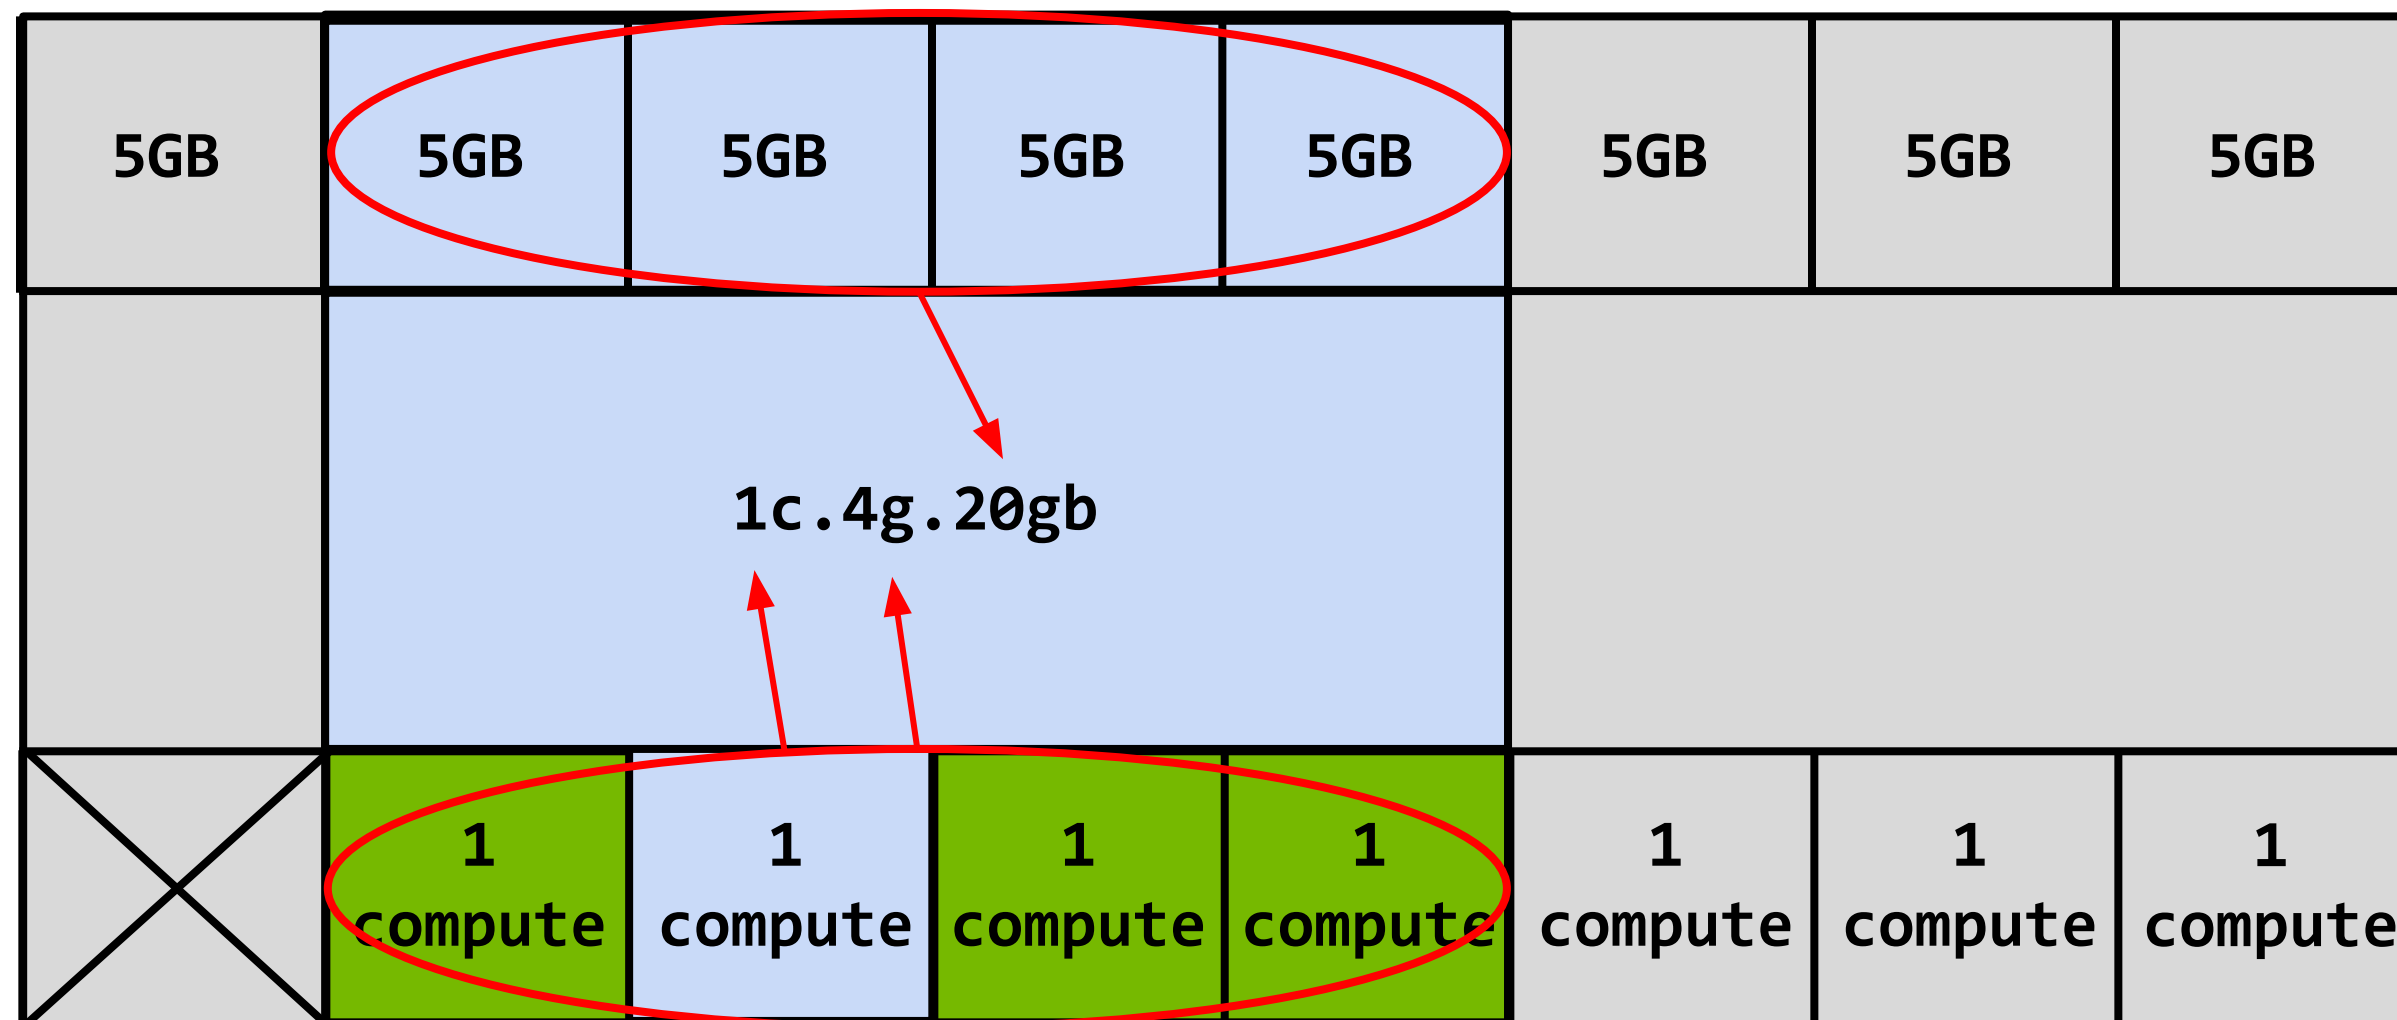


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

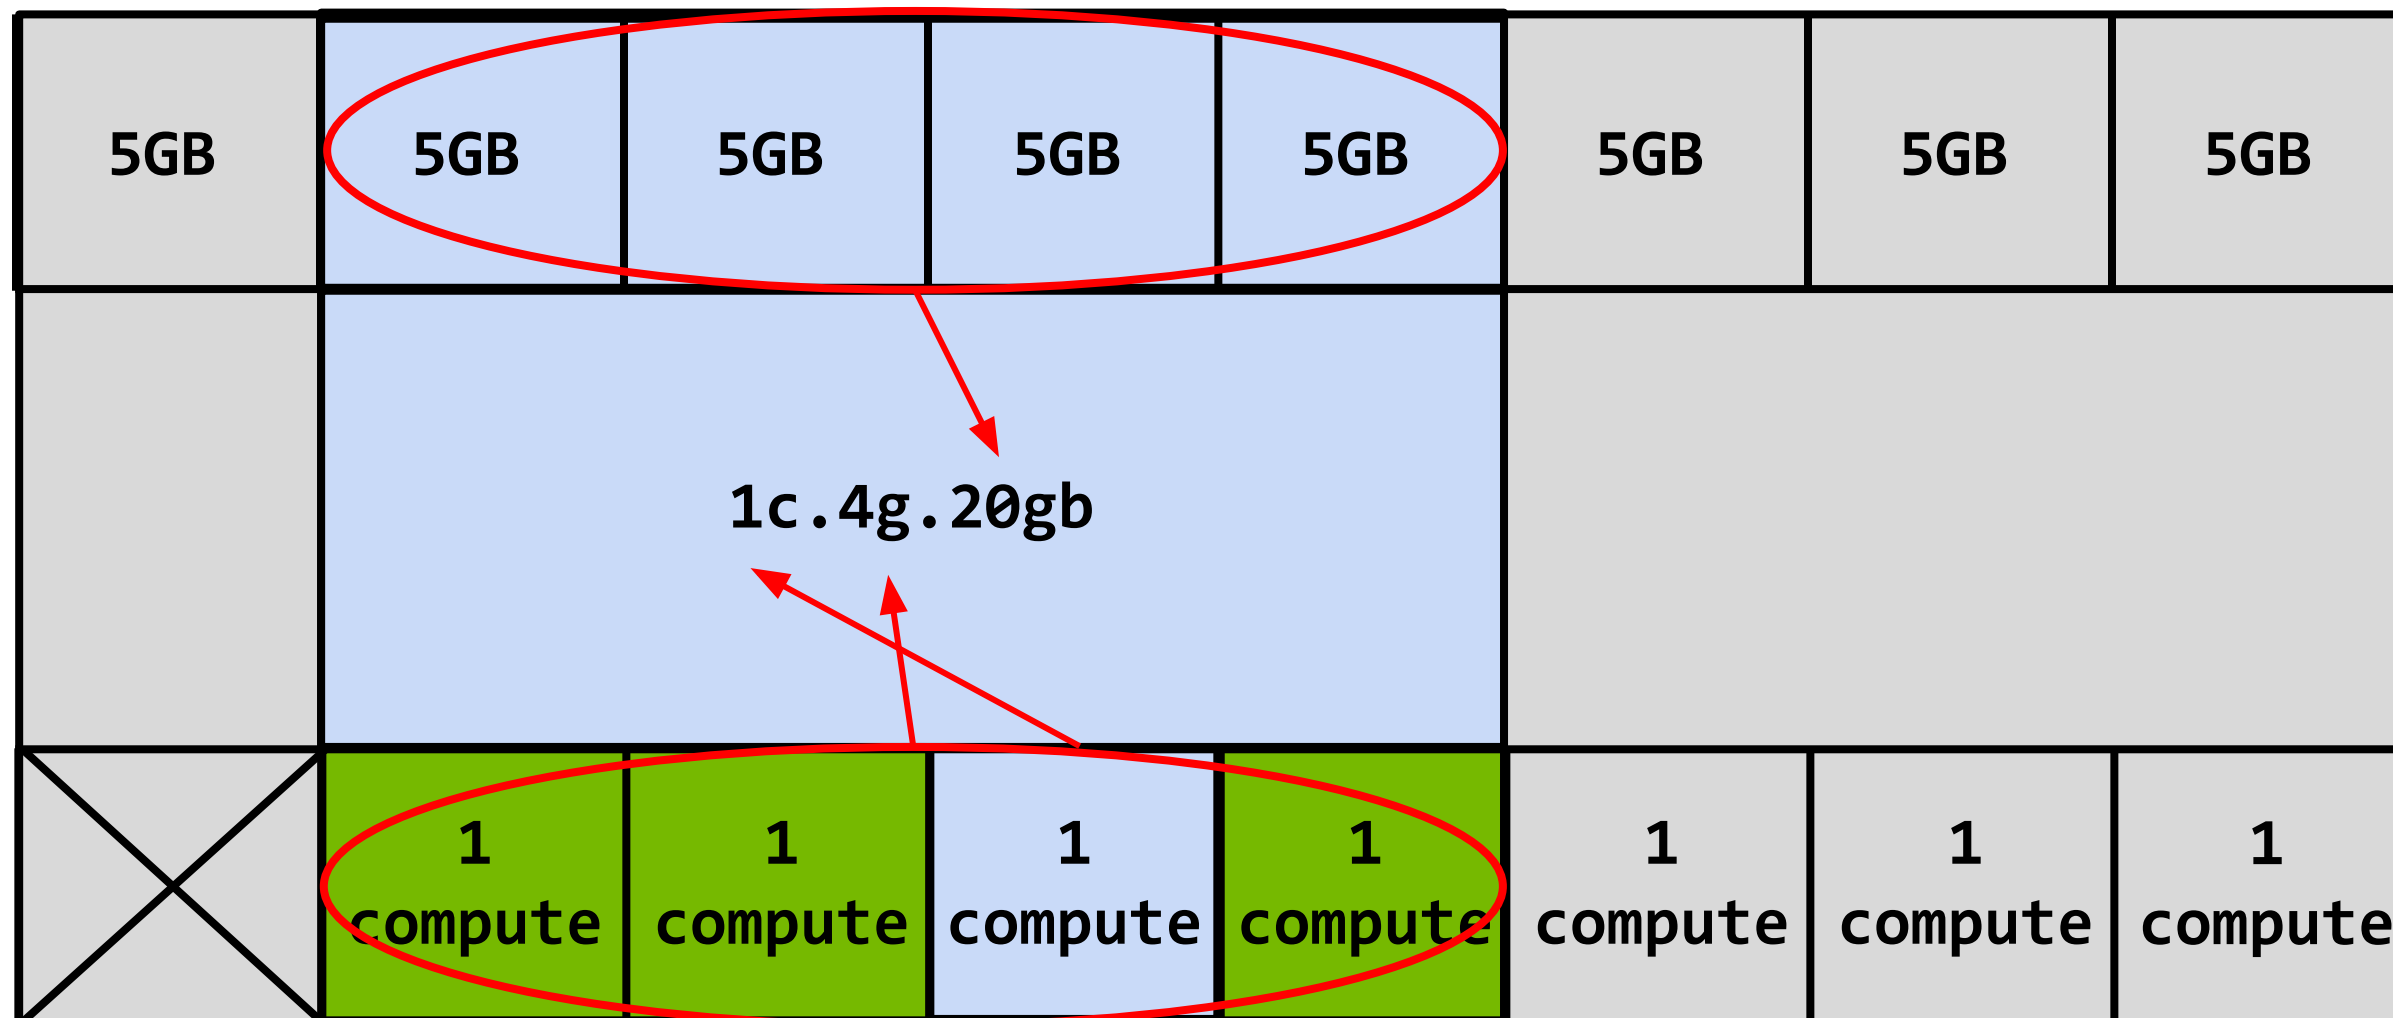


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

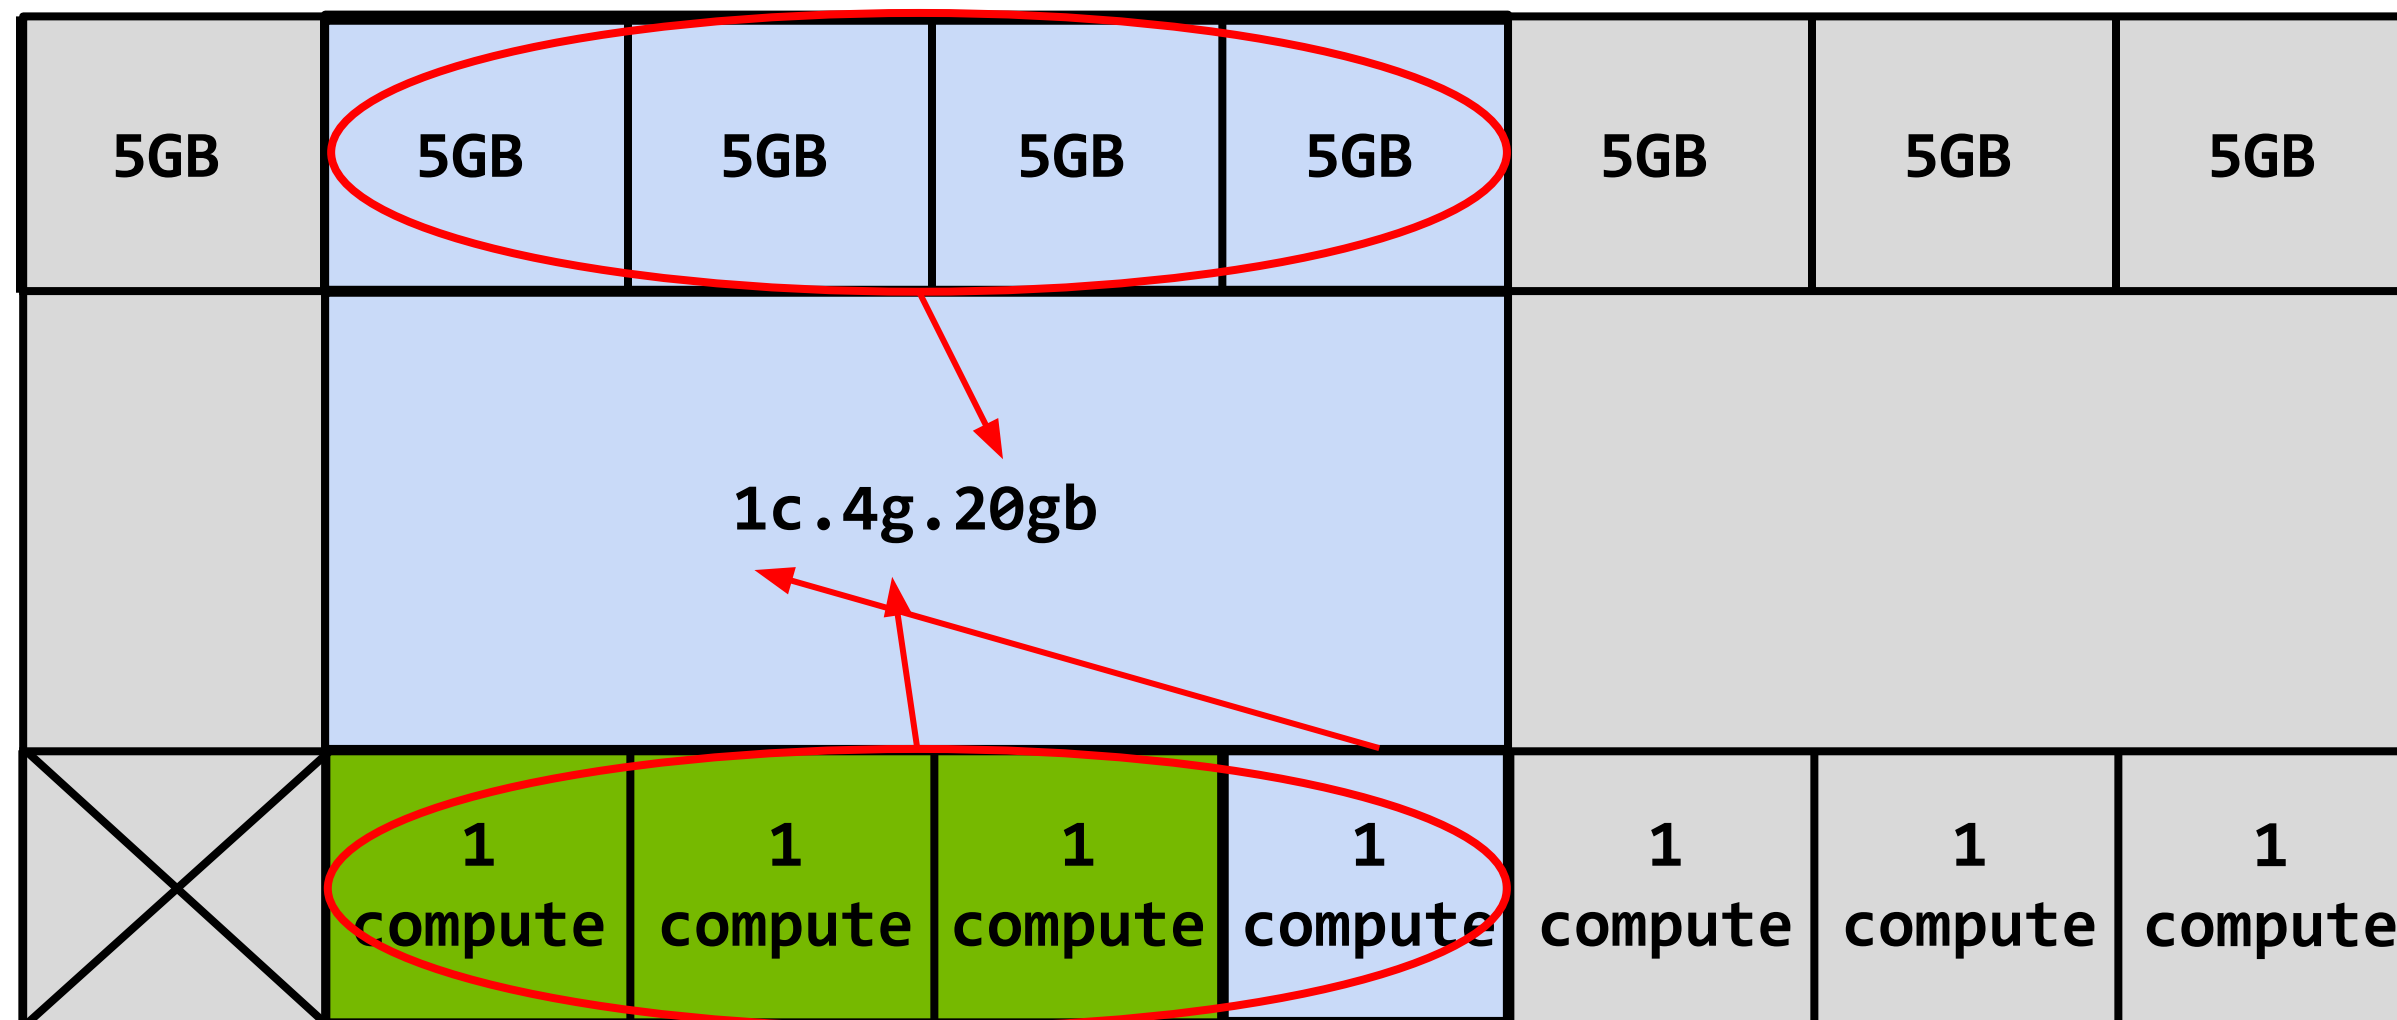


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

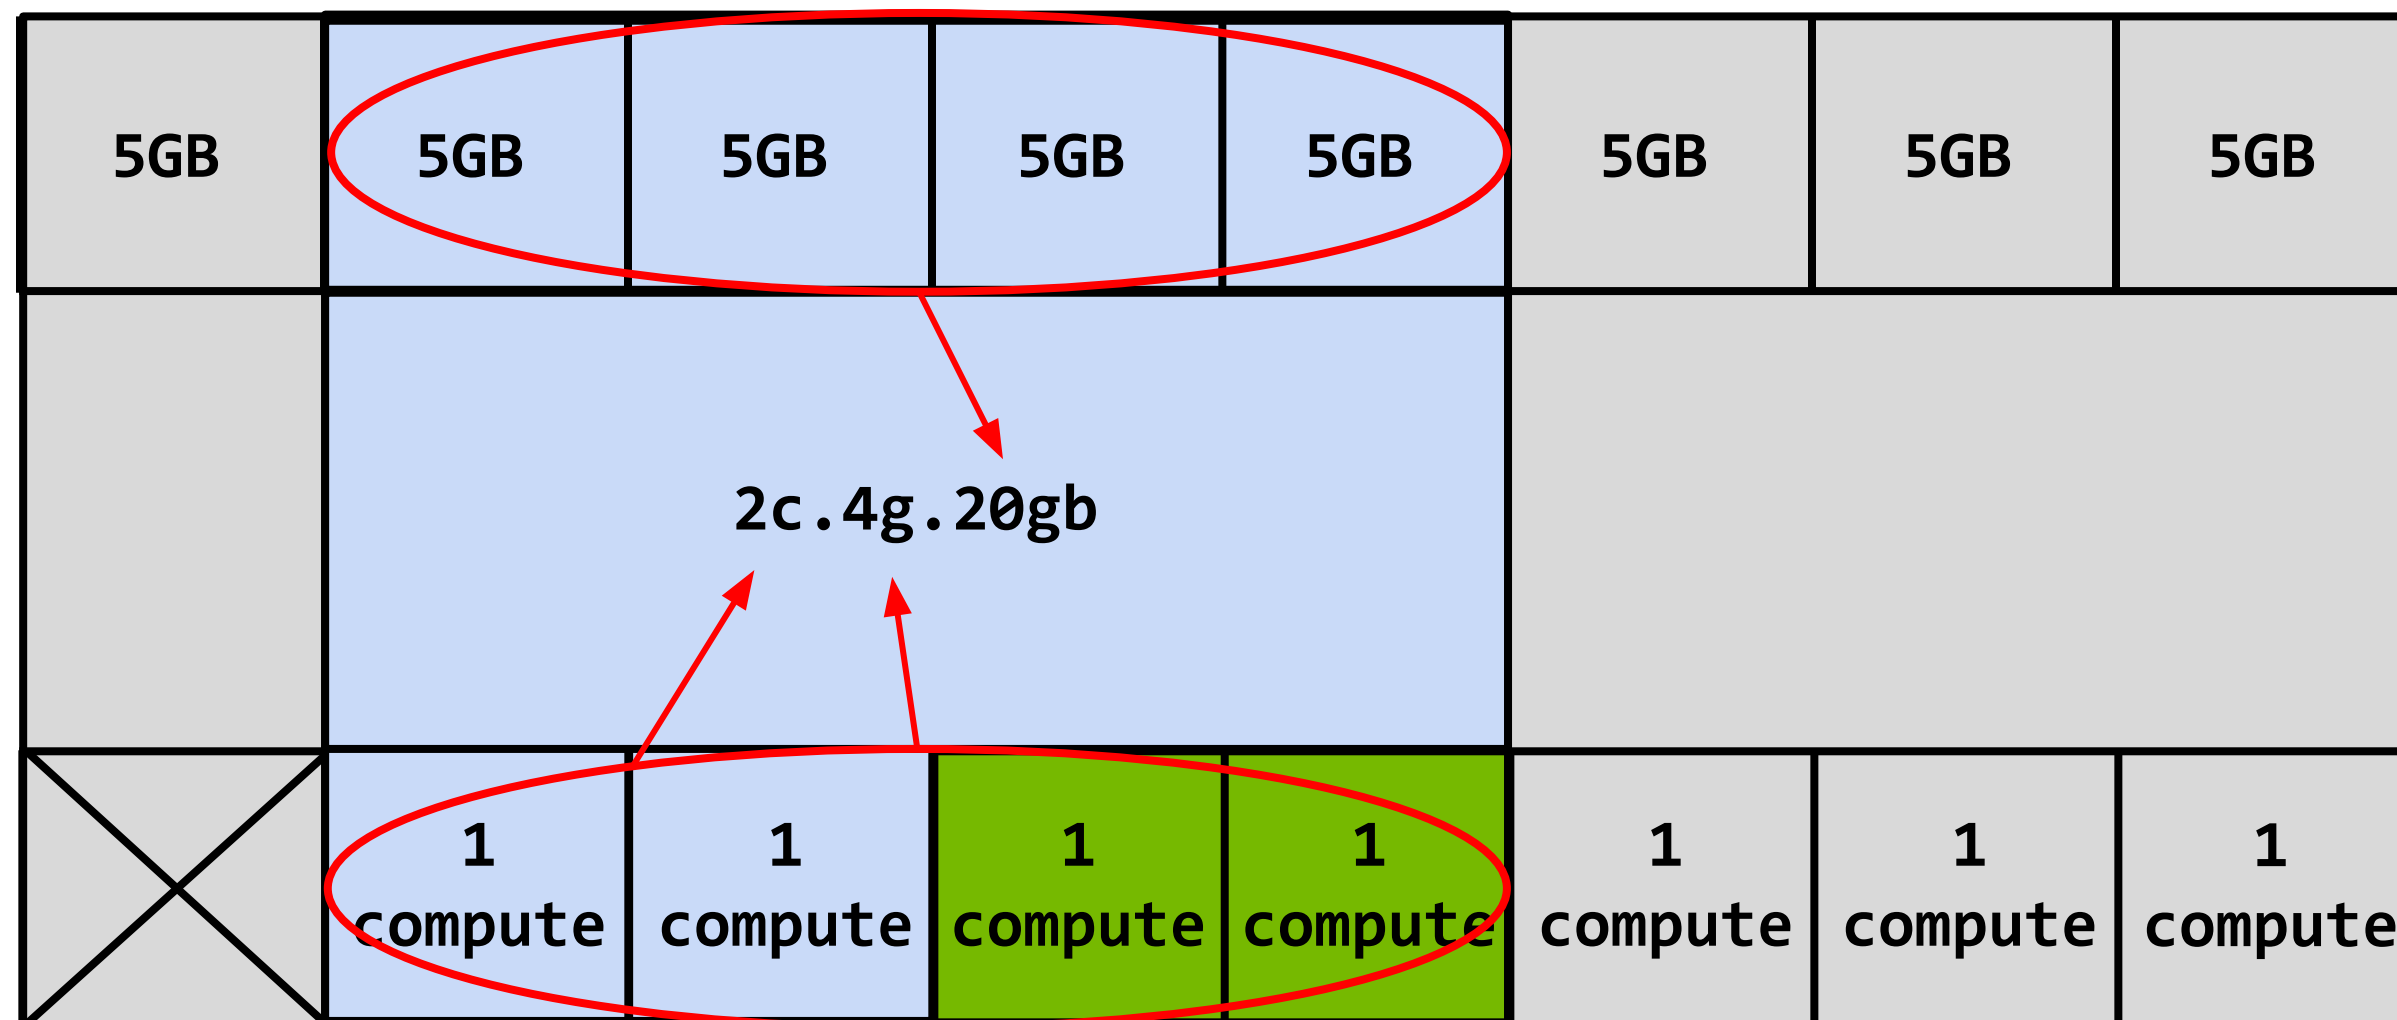


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

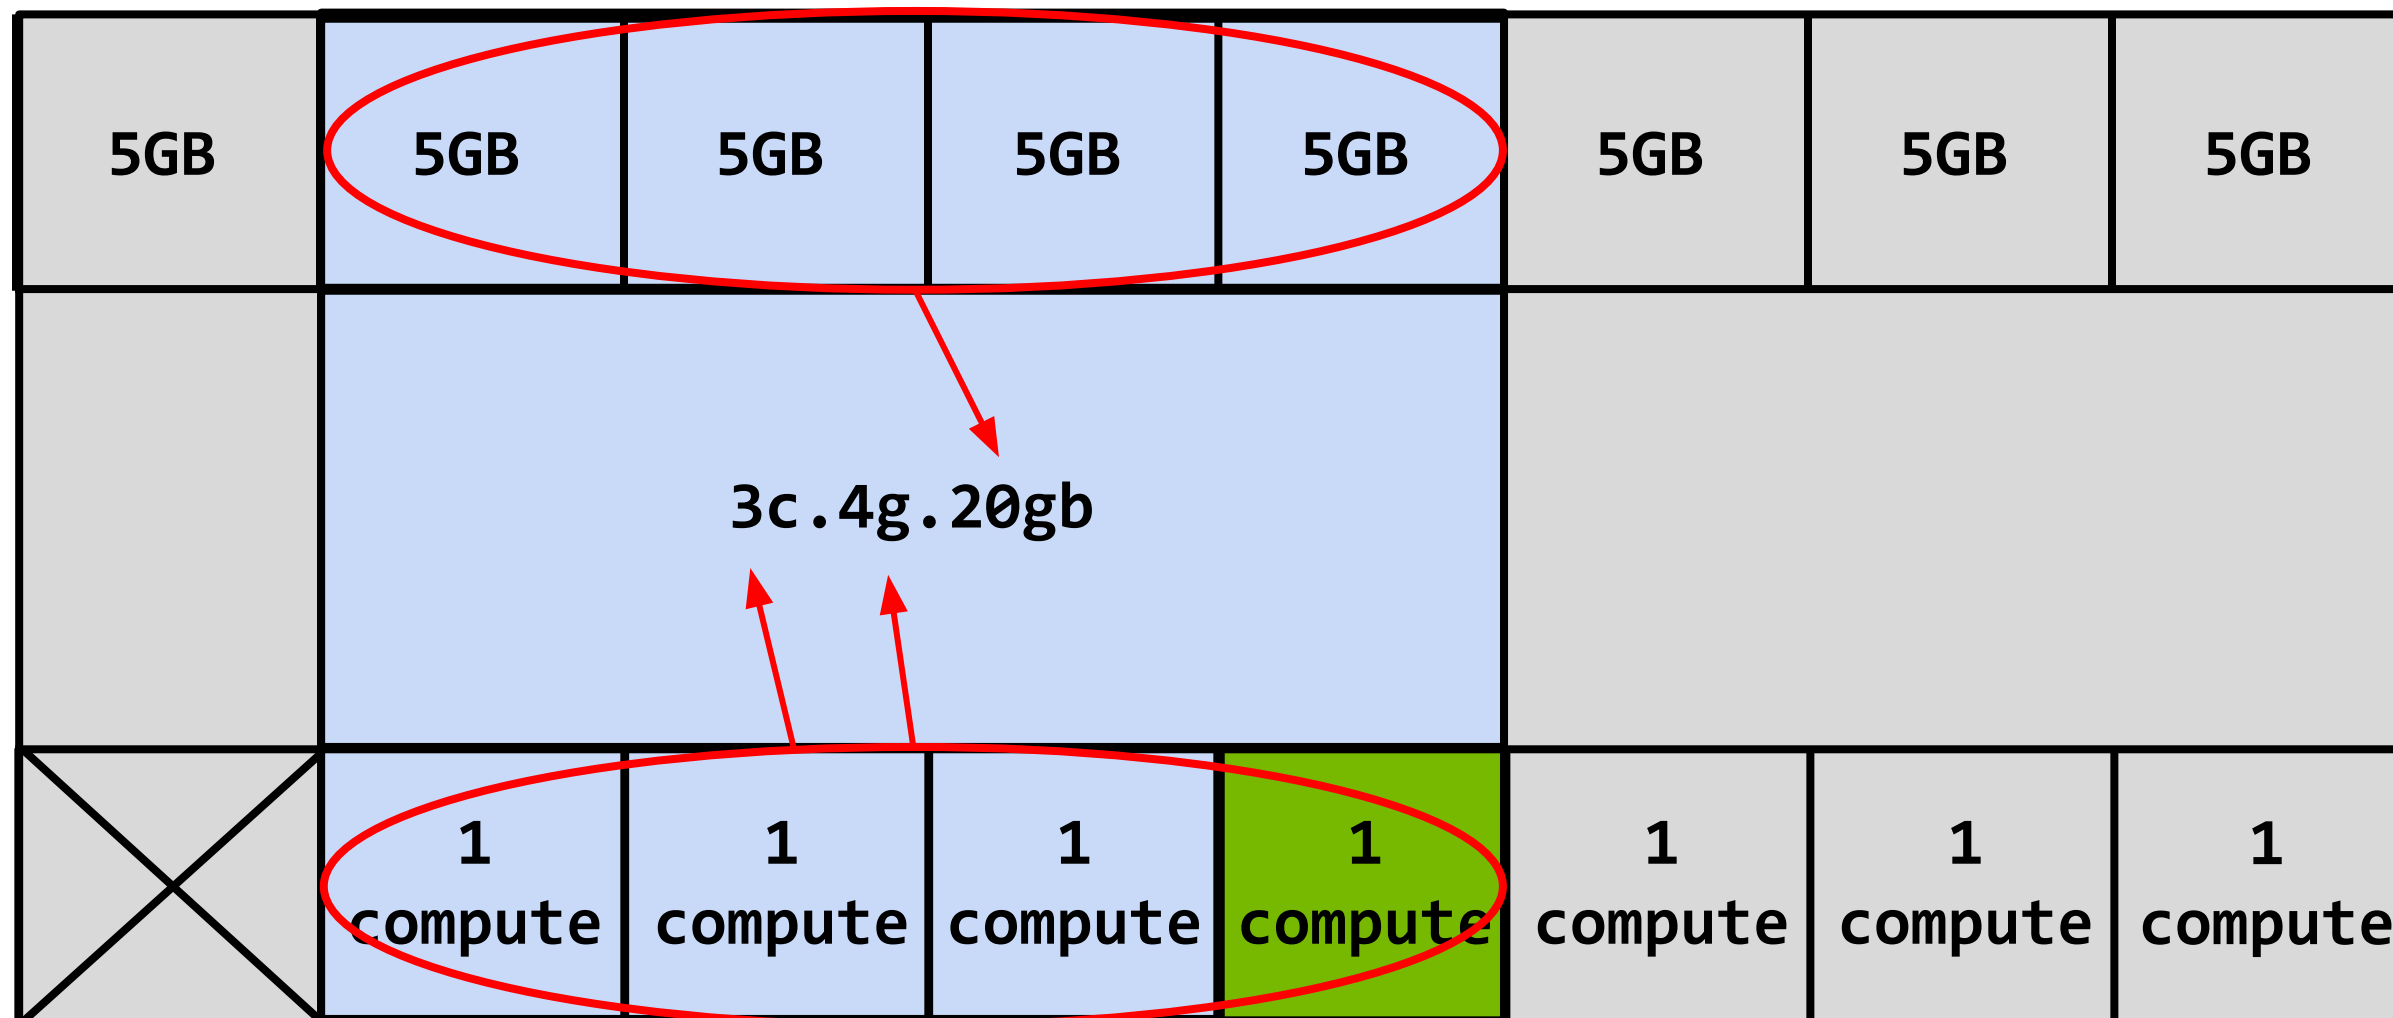


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

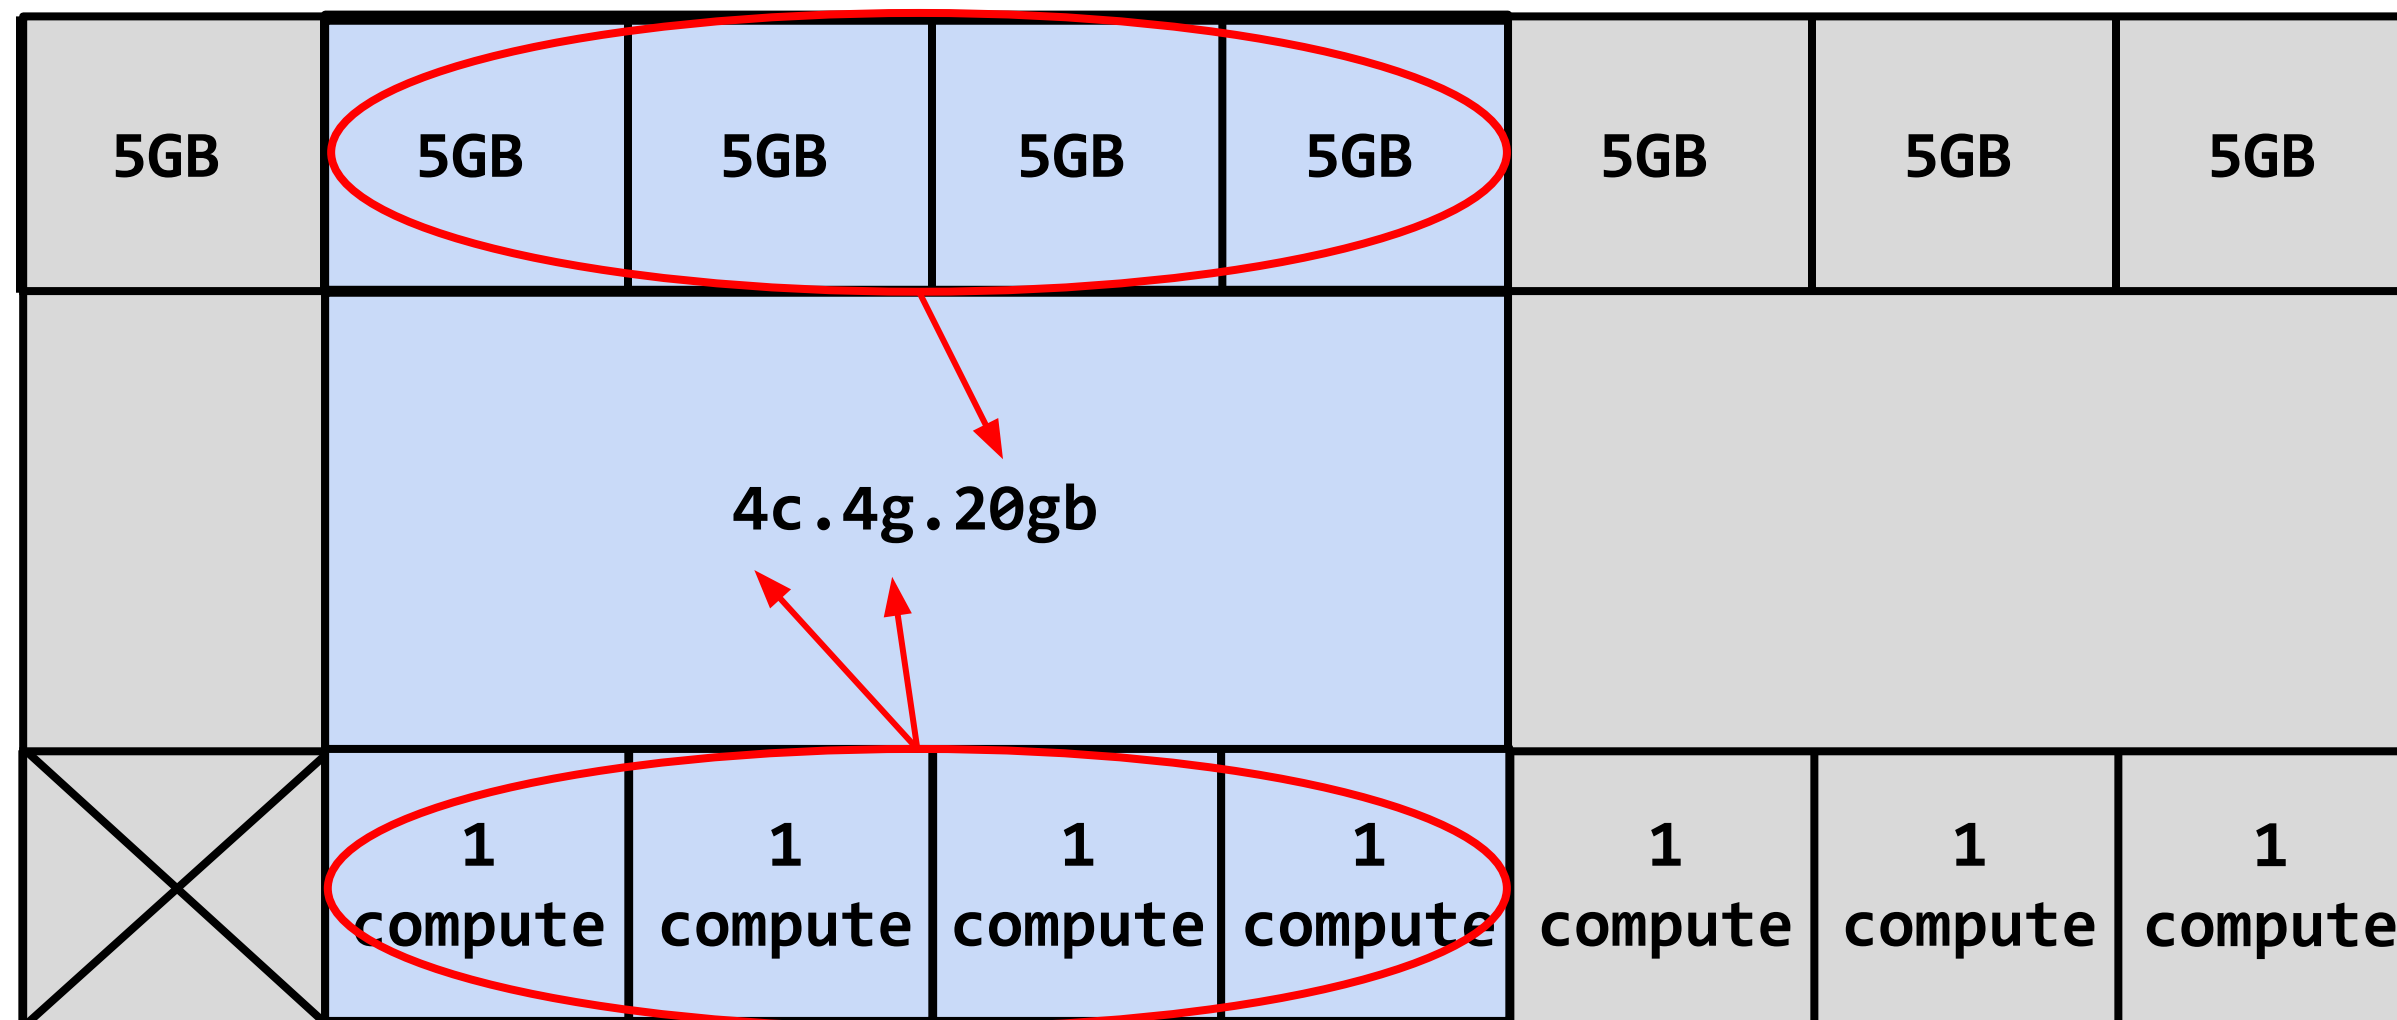


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

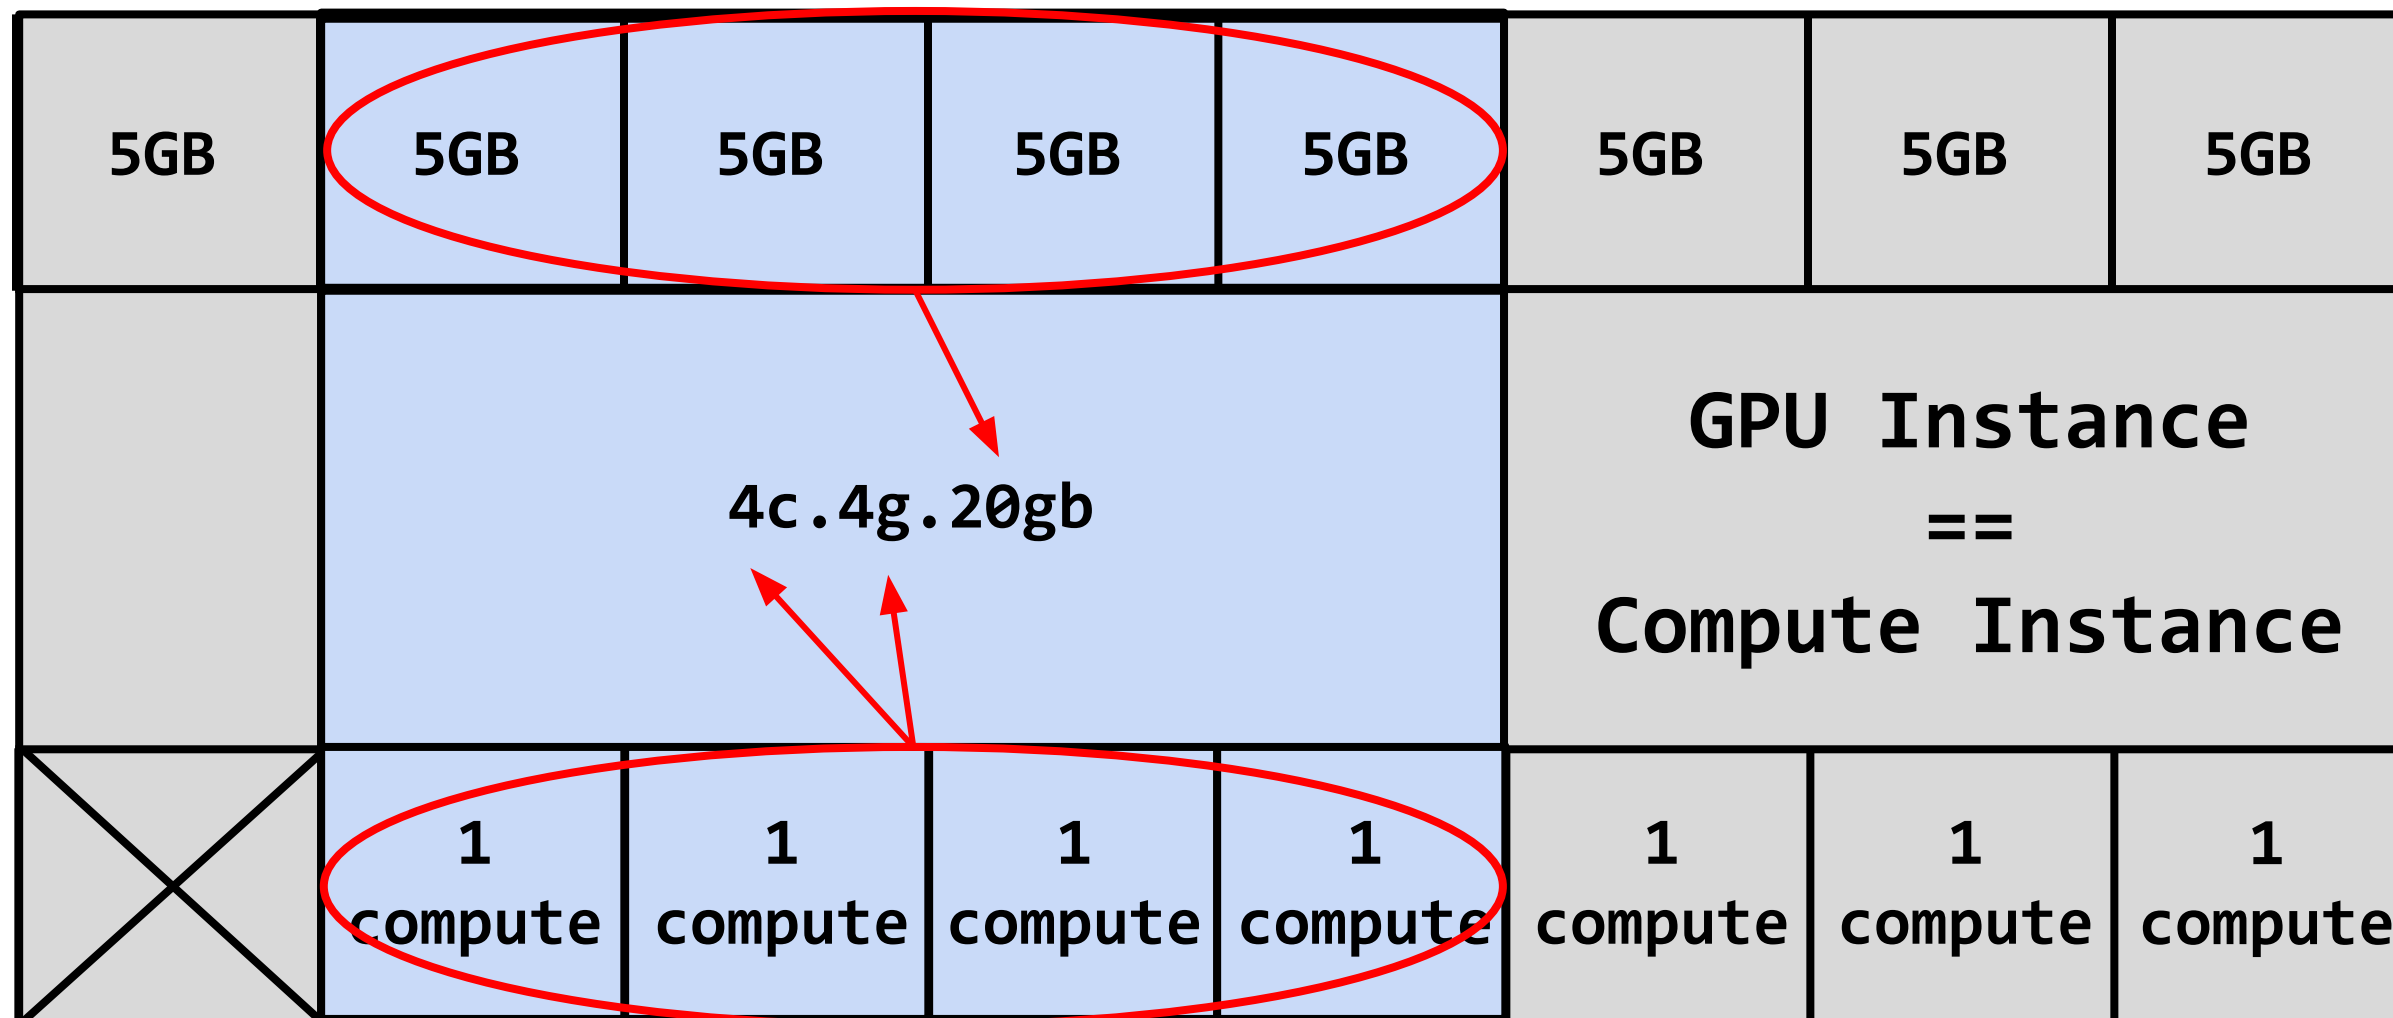


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

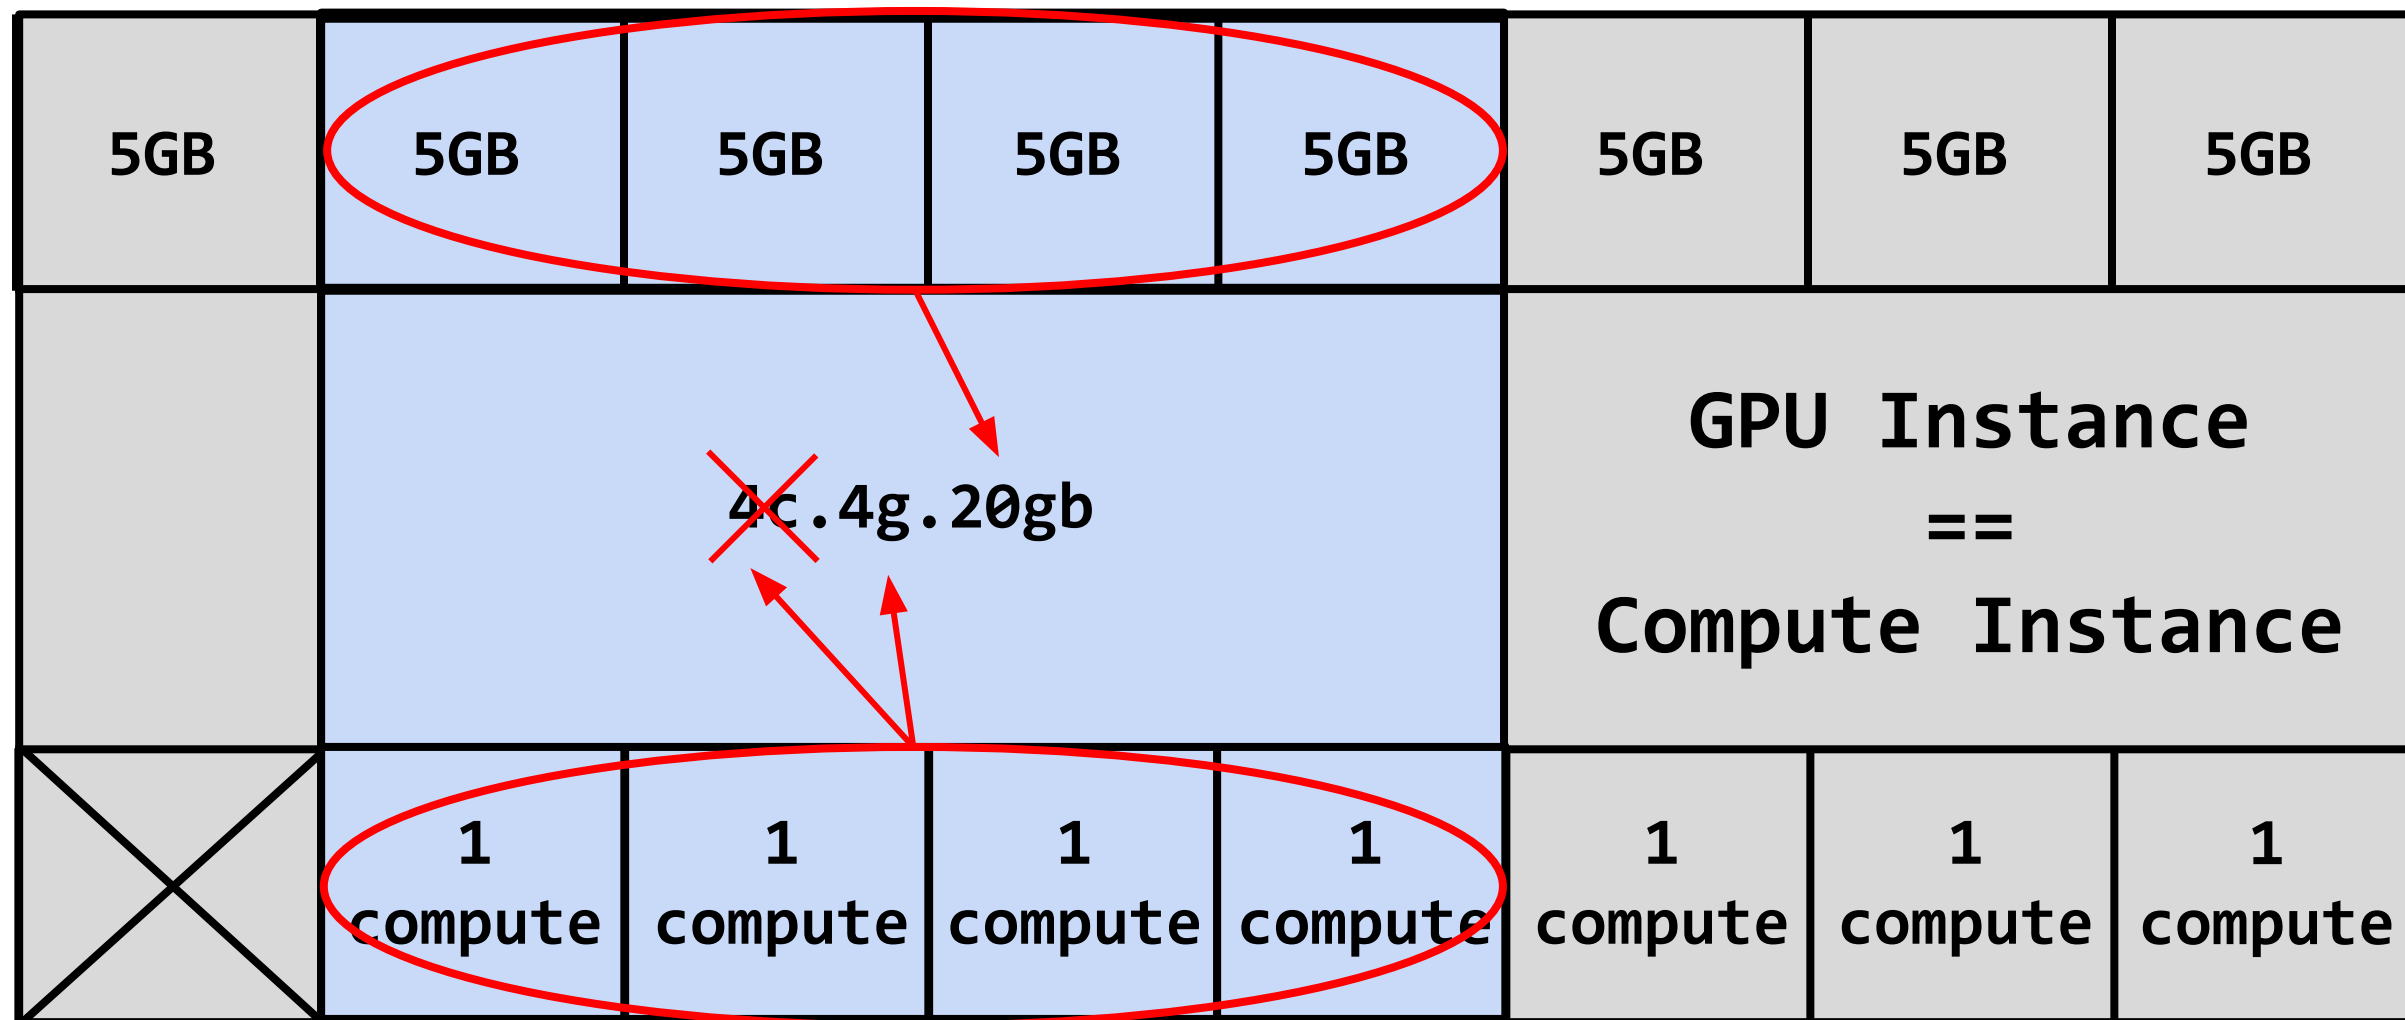


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

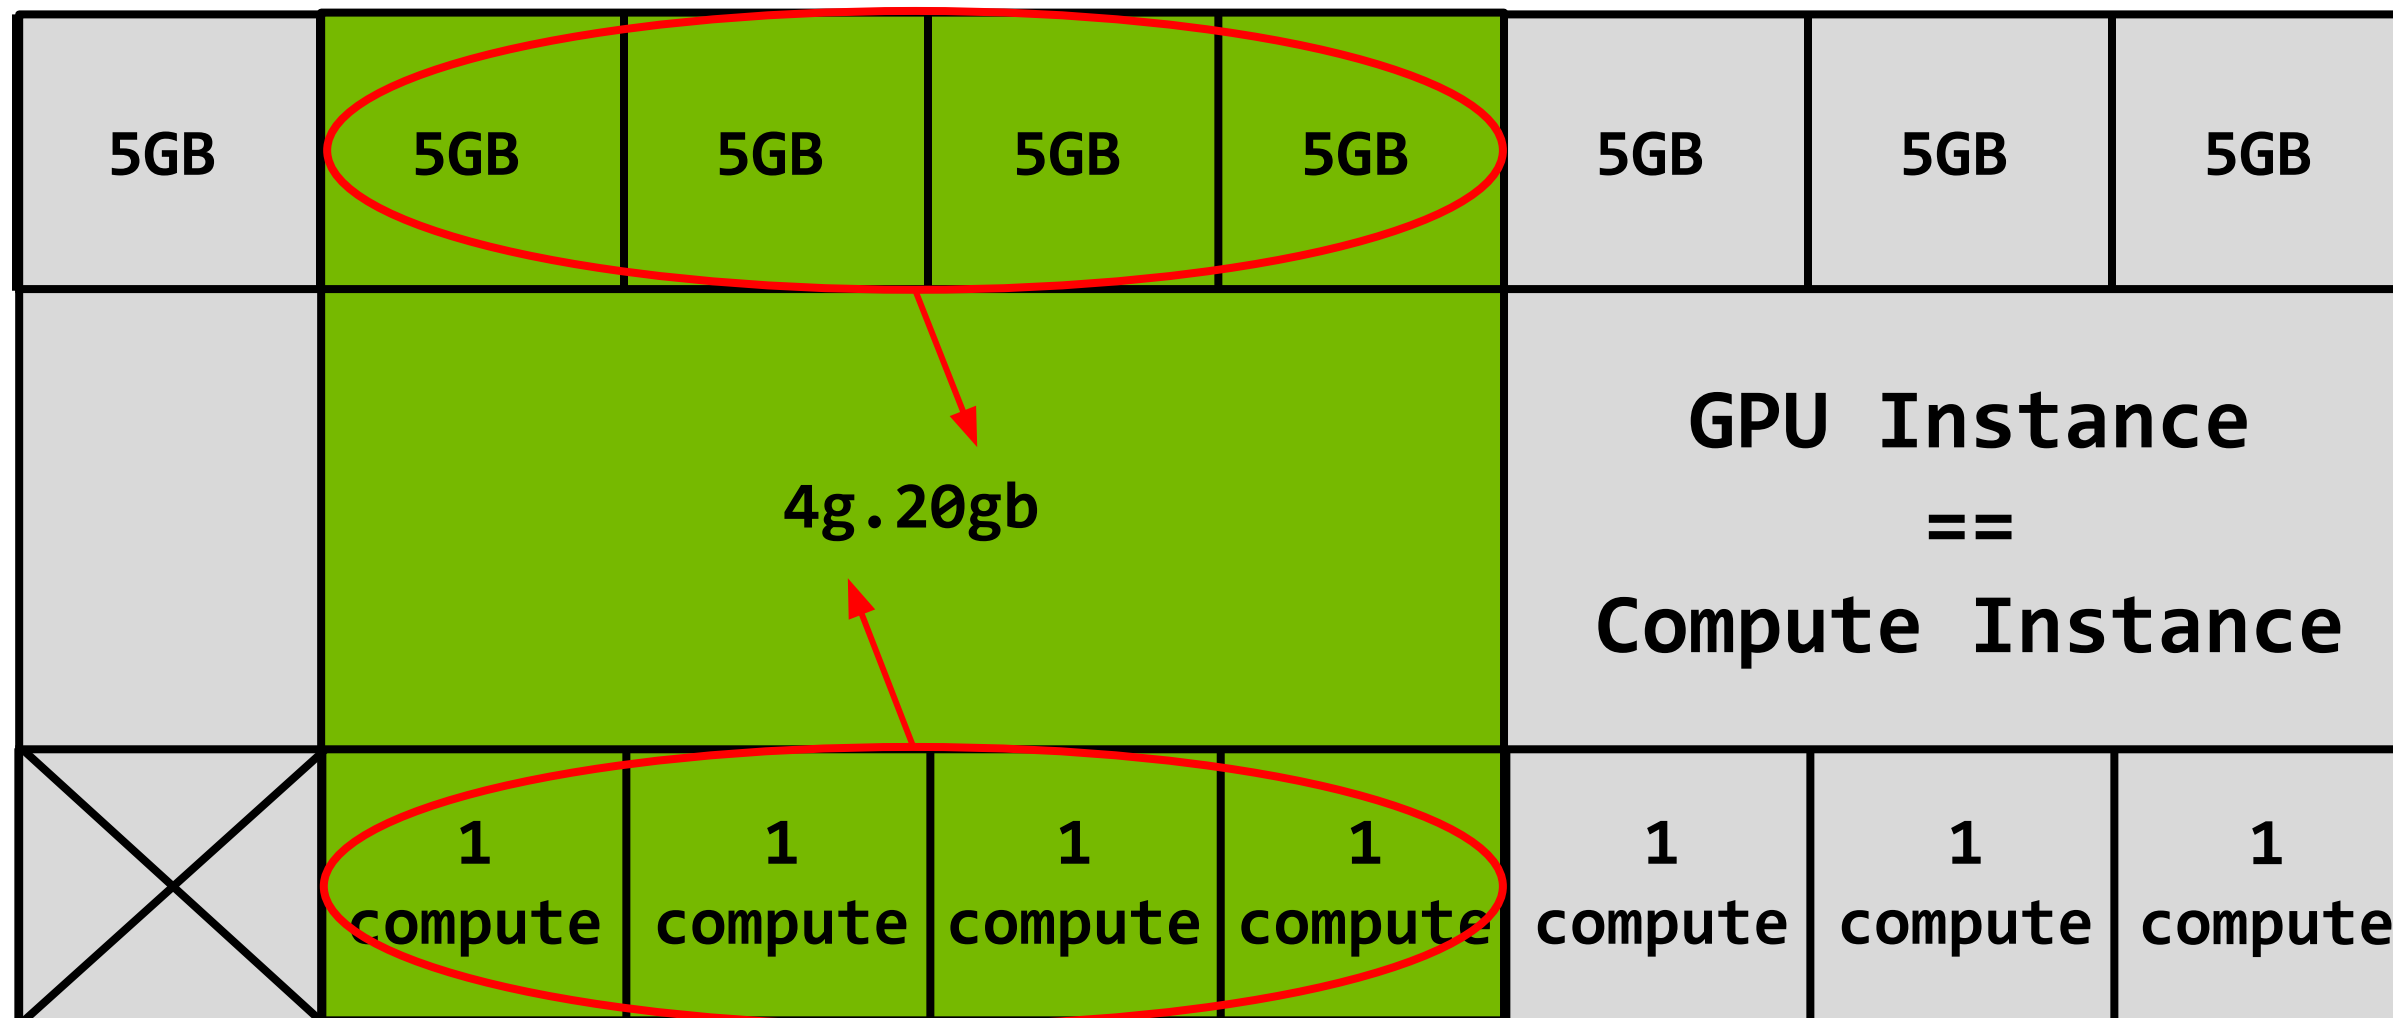


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

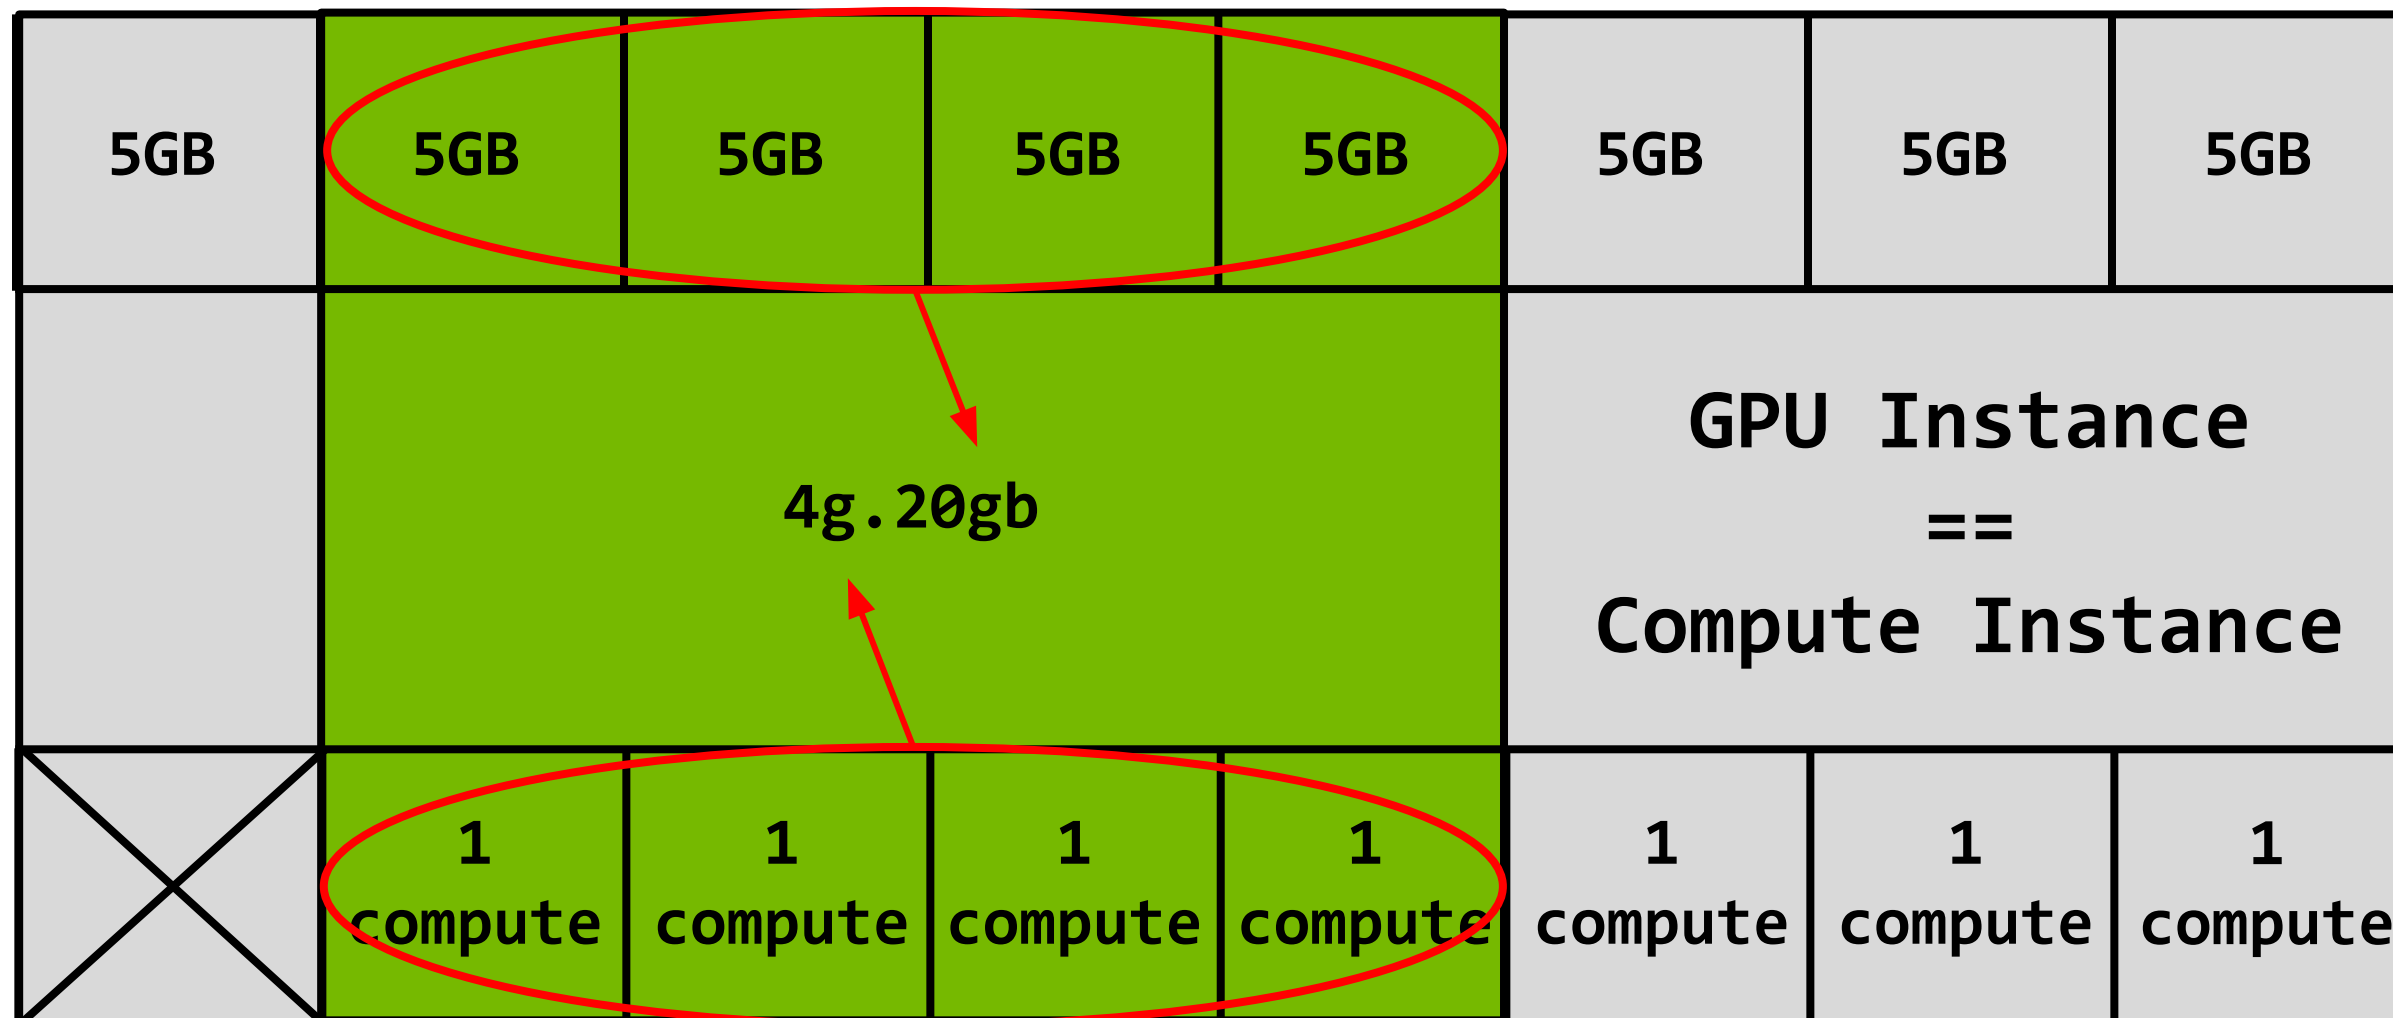


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices



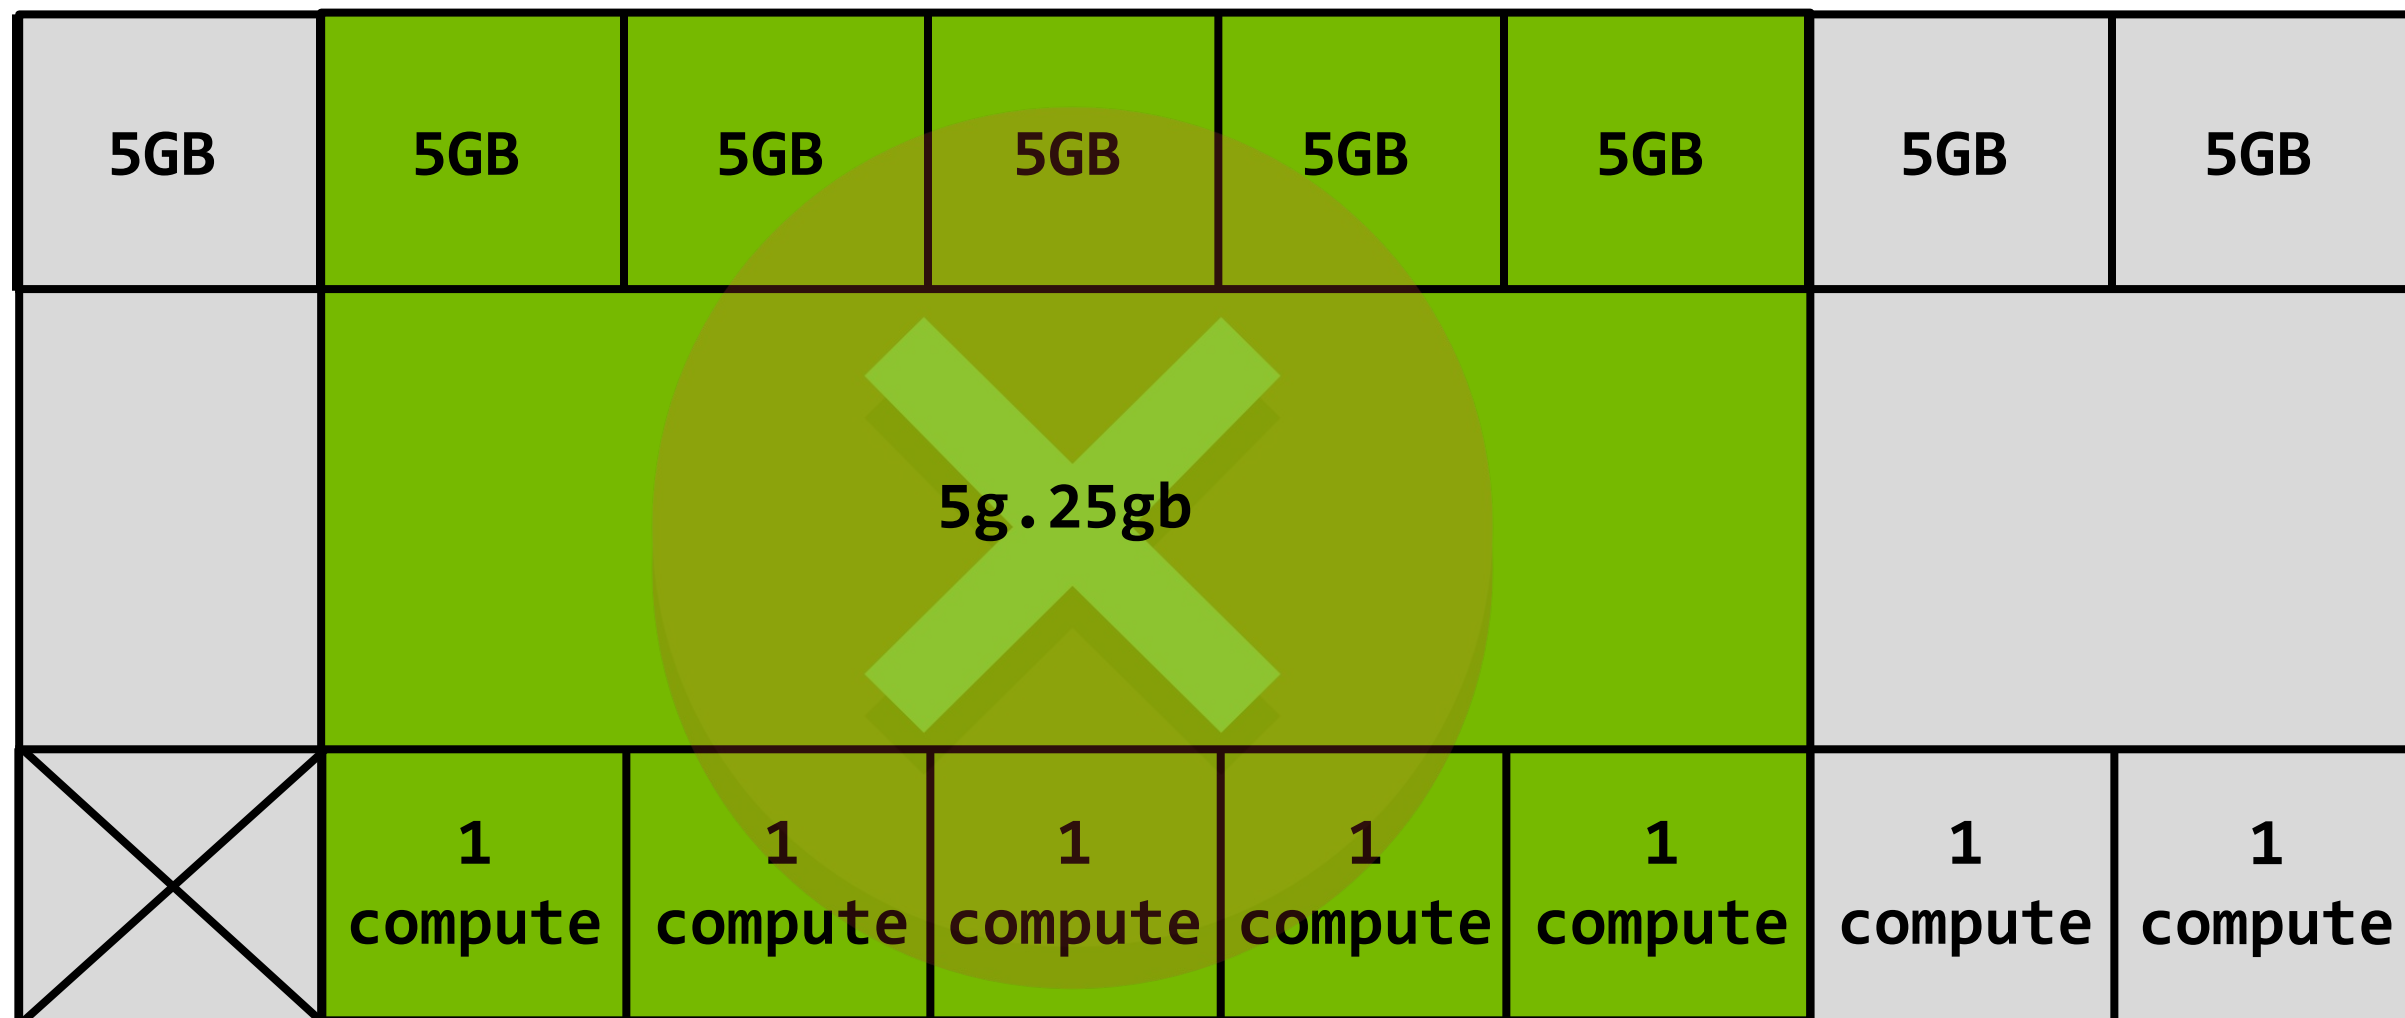
NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices



MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

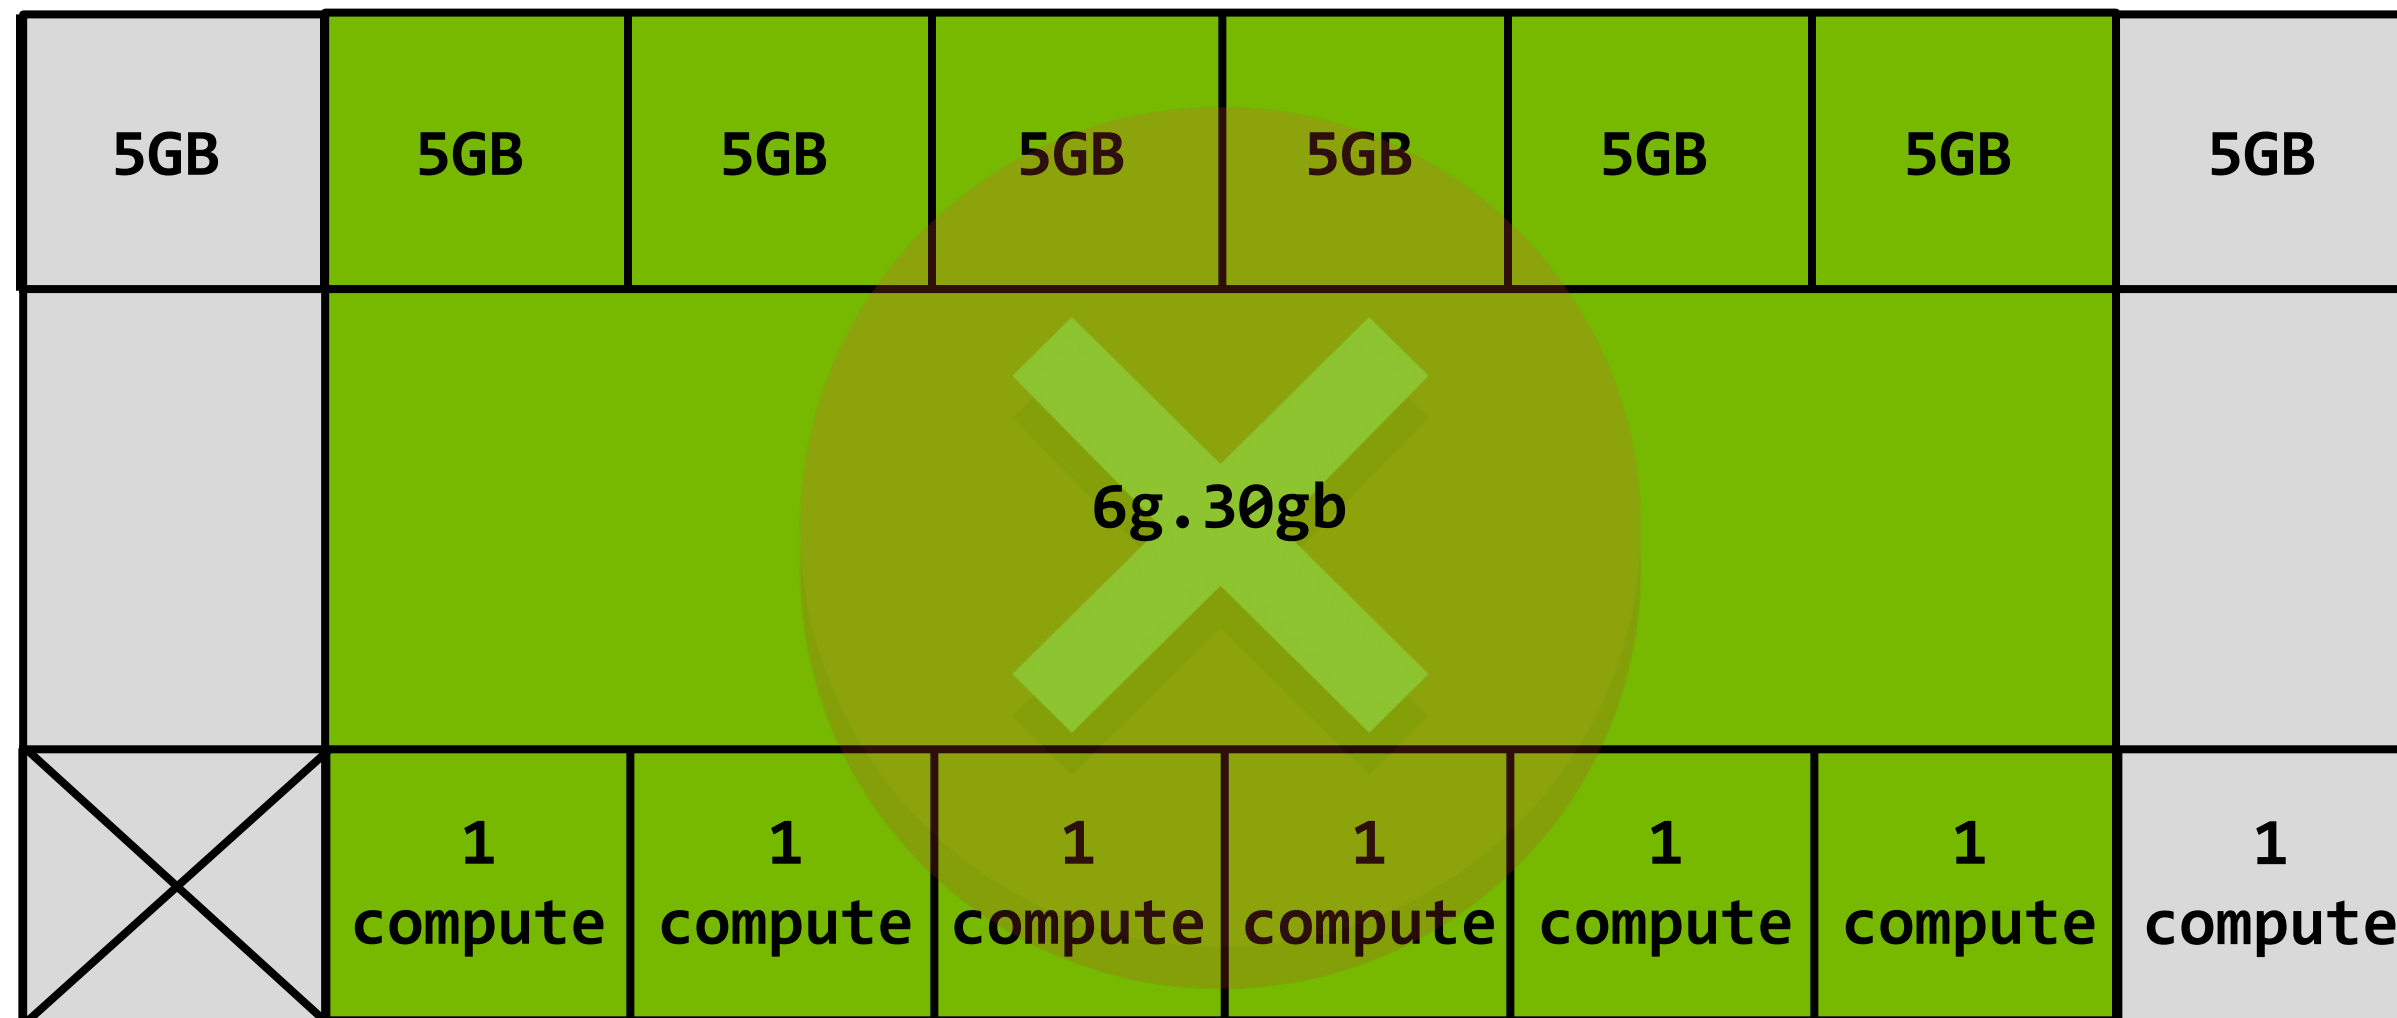


NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices



NVIDIA A100 (40GB)

- 8 x 5GB Memory Slices
- 7 Compute Slices

MULTI-INSTANCE GPUs (MIG)

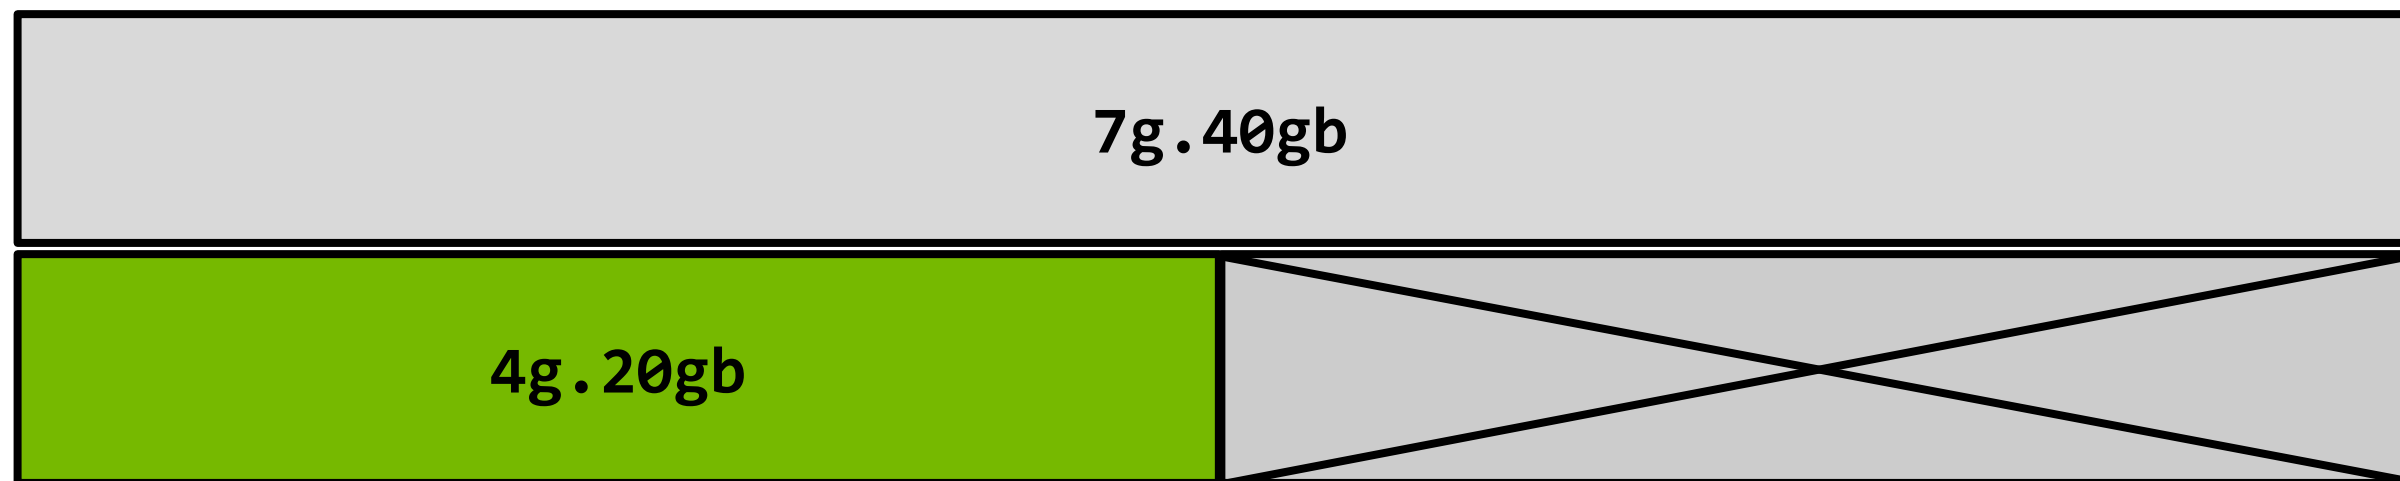
GPU Instances, Compute Instances, and MIG Devices

7g.40gb

- 1 x 7g.40gb

MULTI-INSTANCE GPUs (MIG)

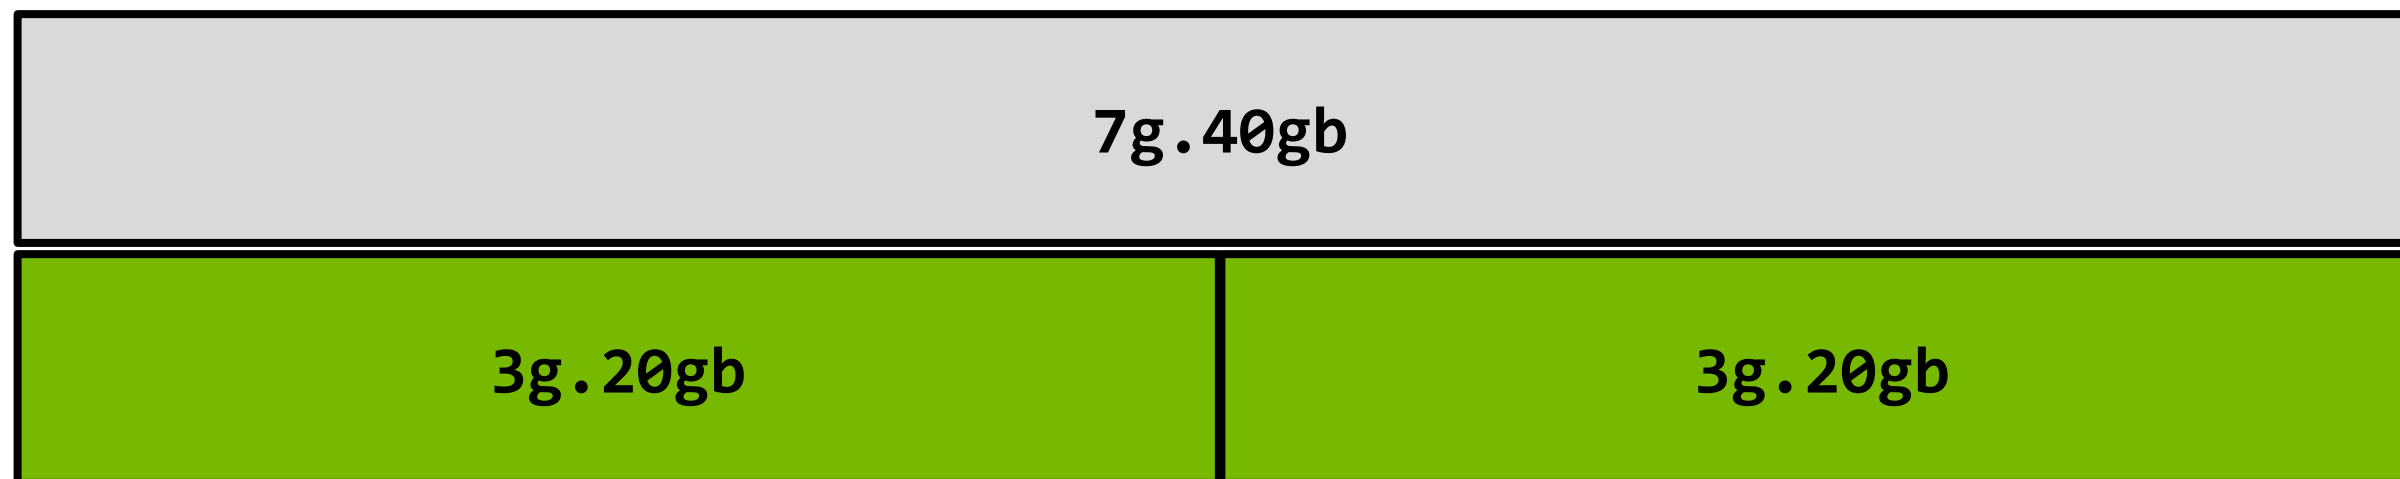
GPU Instances, Compute Instances, and MIG Devices



- 1 x 7g.40gb
or
- 1 x 4g.20gb

MULTI-INSTANCE GPUs (MIG)

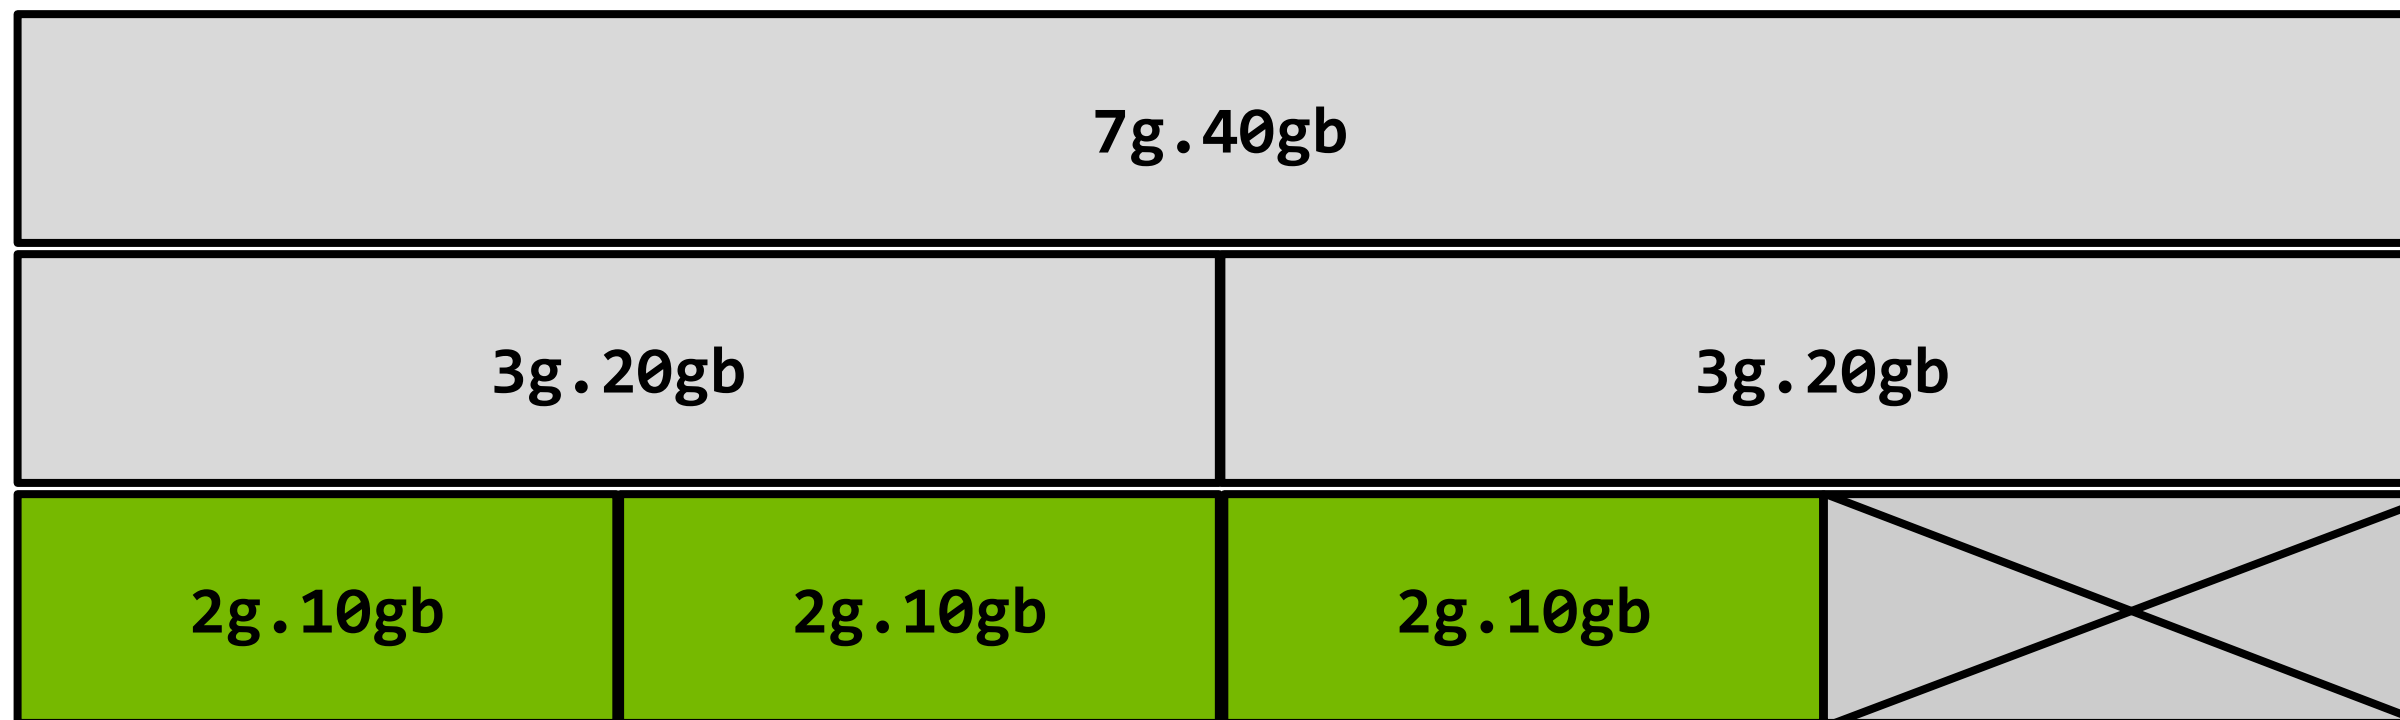
GPU Instances, Compute Instances, and MIG Devices



- 1 x 7g.40gb
or
- 2 x 3g.20gb

MULTI-INSTANCE GPUs (MIG)

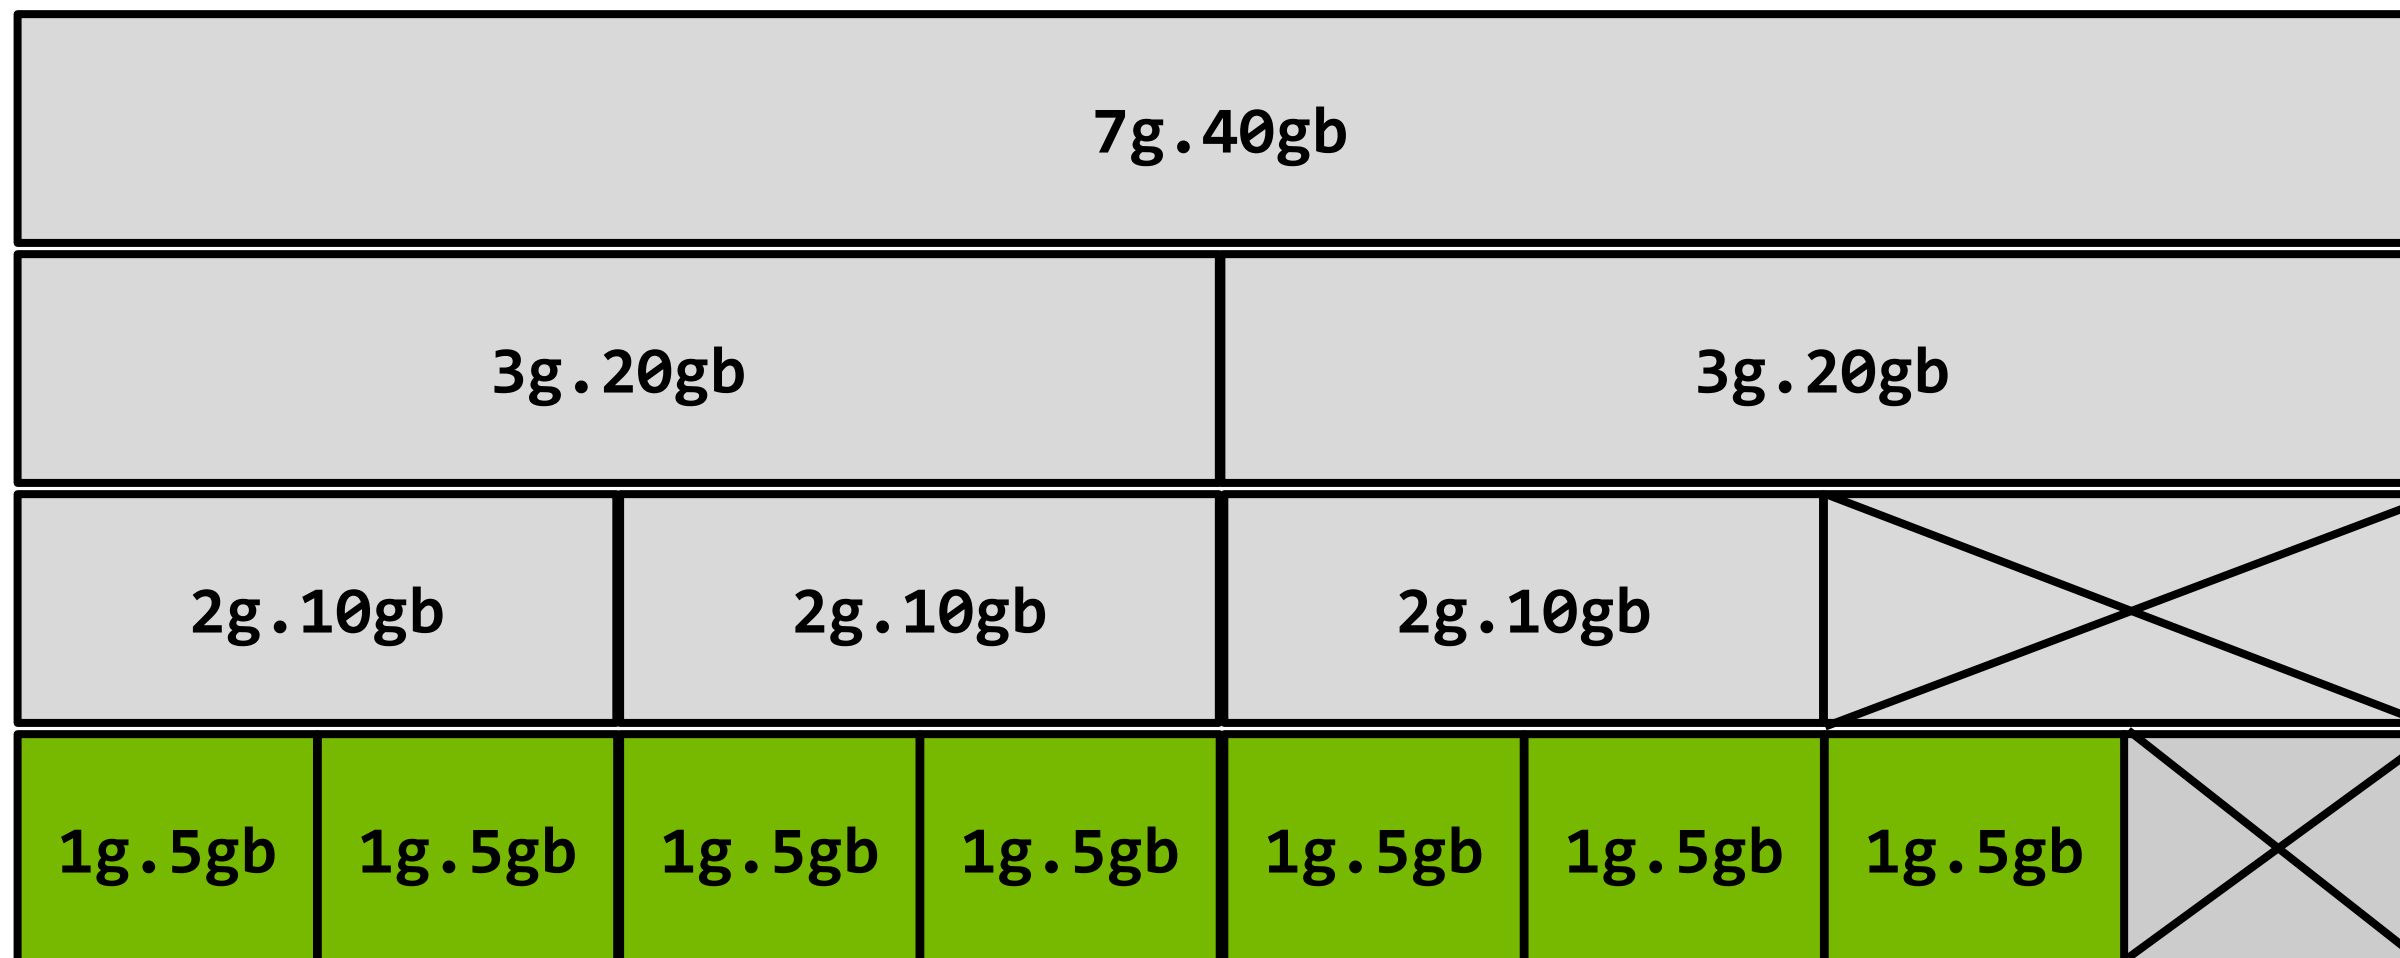
GPU Instances, Compute Instances, and MIG Devices



- **1 x 7g.40gb**
or
- **2 x 3g.20gb**
or
- **3 x 2g.10gb**

MULTI-INSTANCE GPUs (MIG)

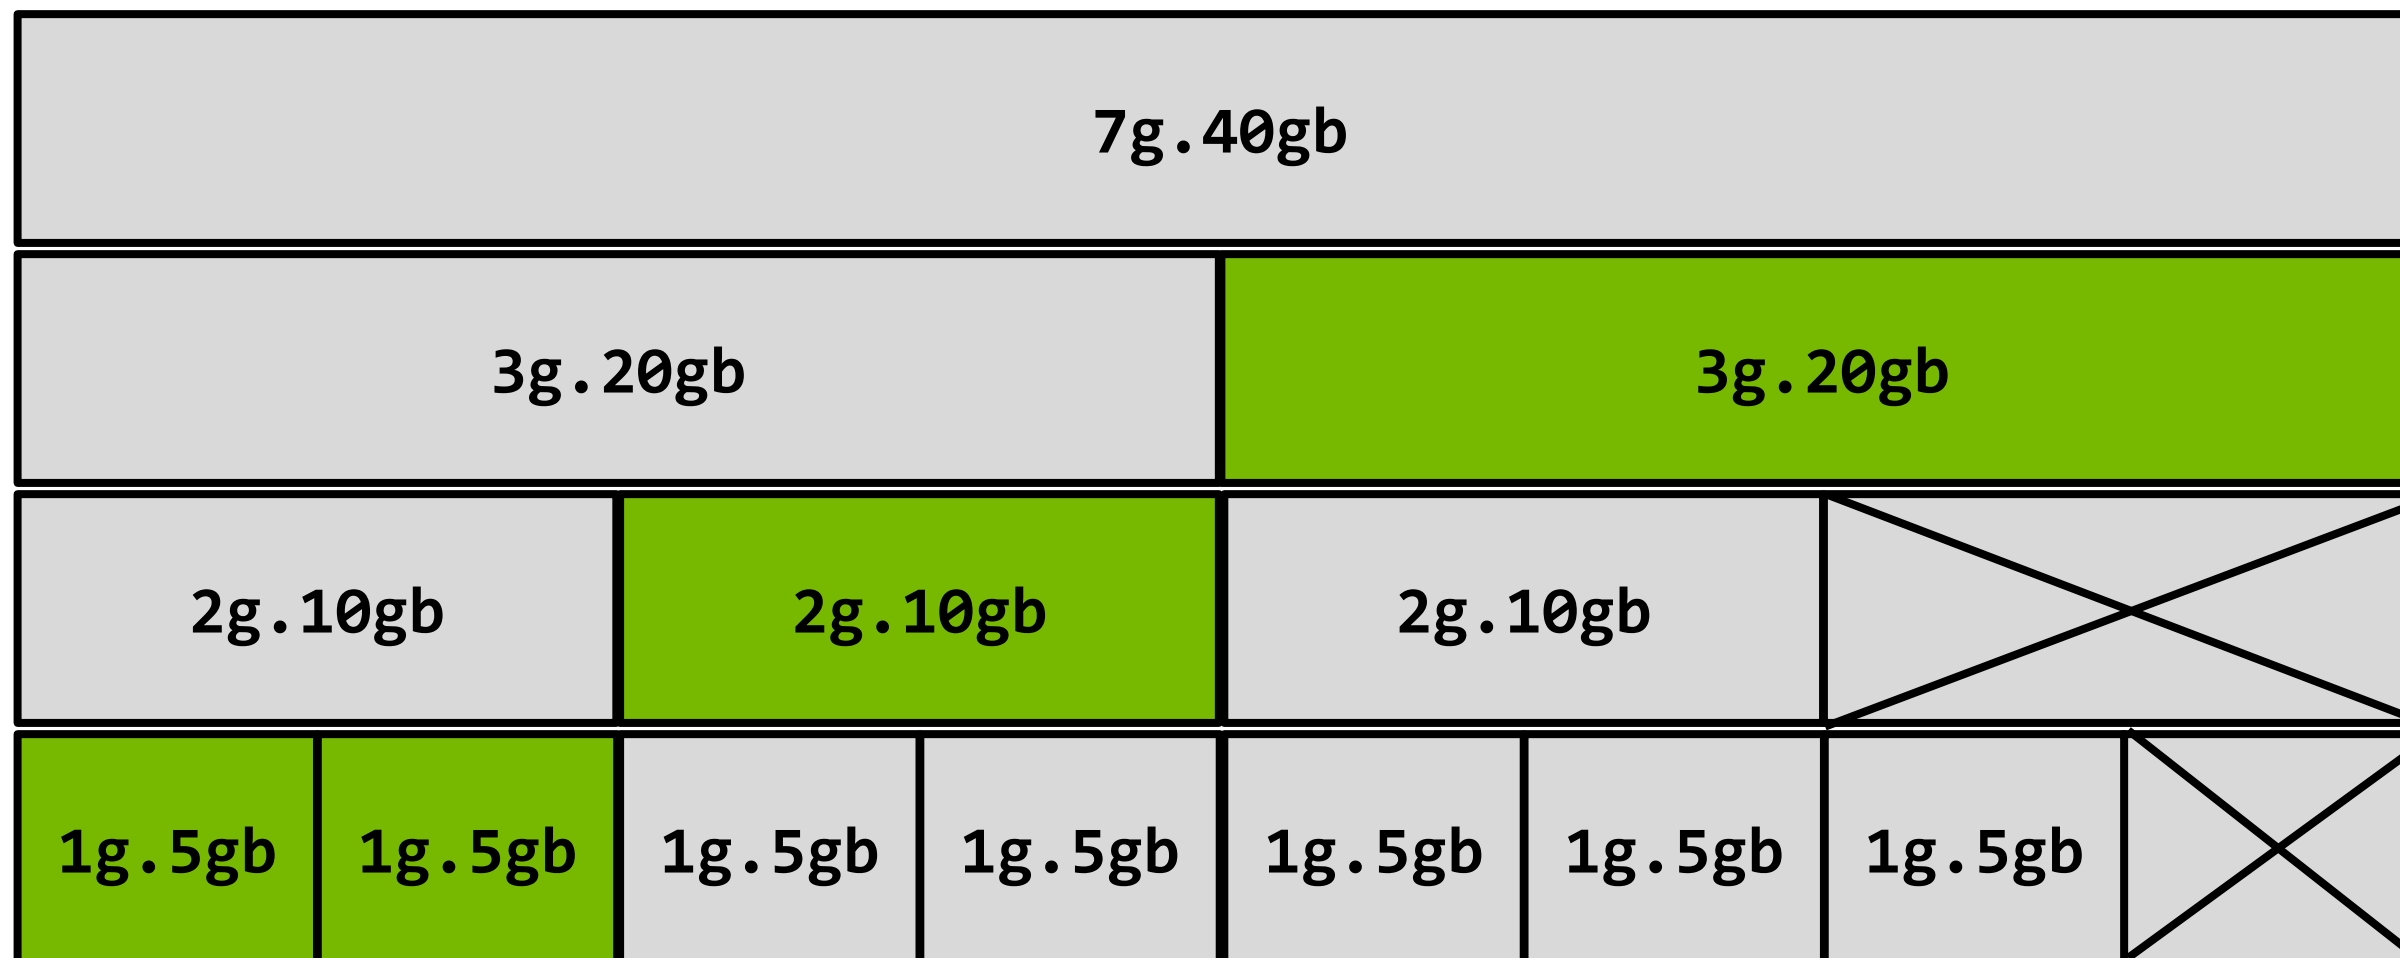
GPU Instances, Compute Instances, and MIG Devices



- 1 x 7g.40gb
or
- 2 x 3g.20gb
or
- 3 x 2g.10gb
or
- 7 x 1g.5gb

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices



- 1 x 3g.20gb
and
- 1 x 2g.10gb
and
- 2 x 1g.5gb

MULTI-INSTANCE GPUs (MIG)

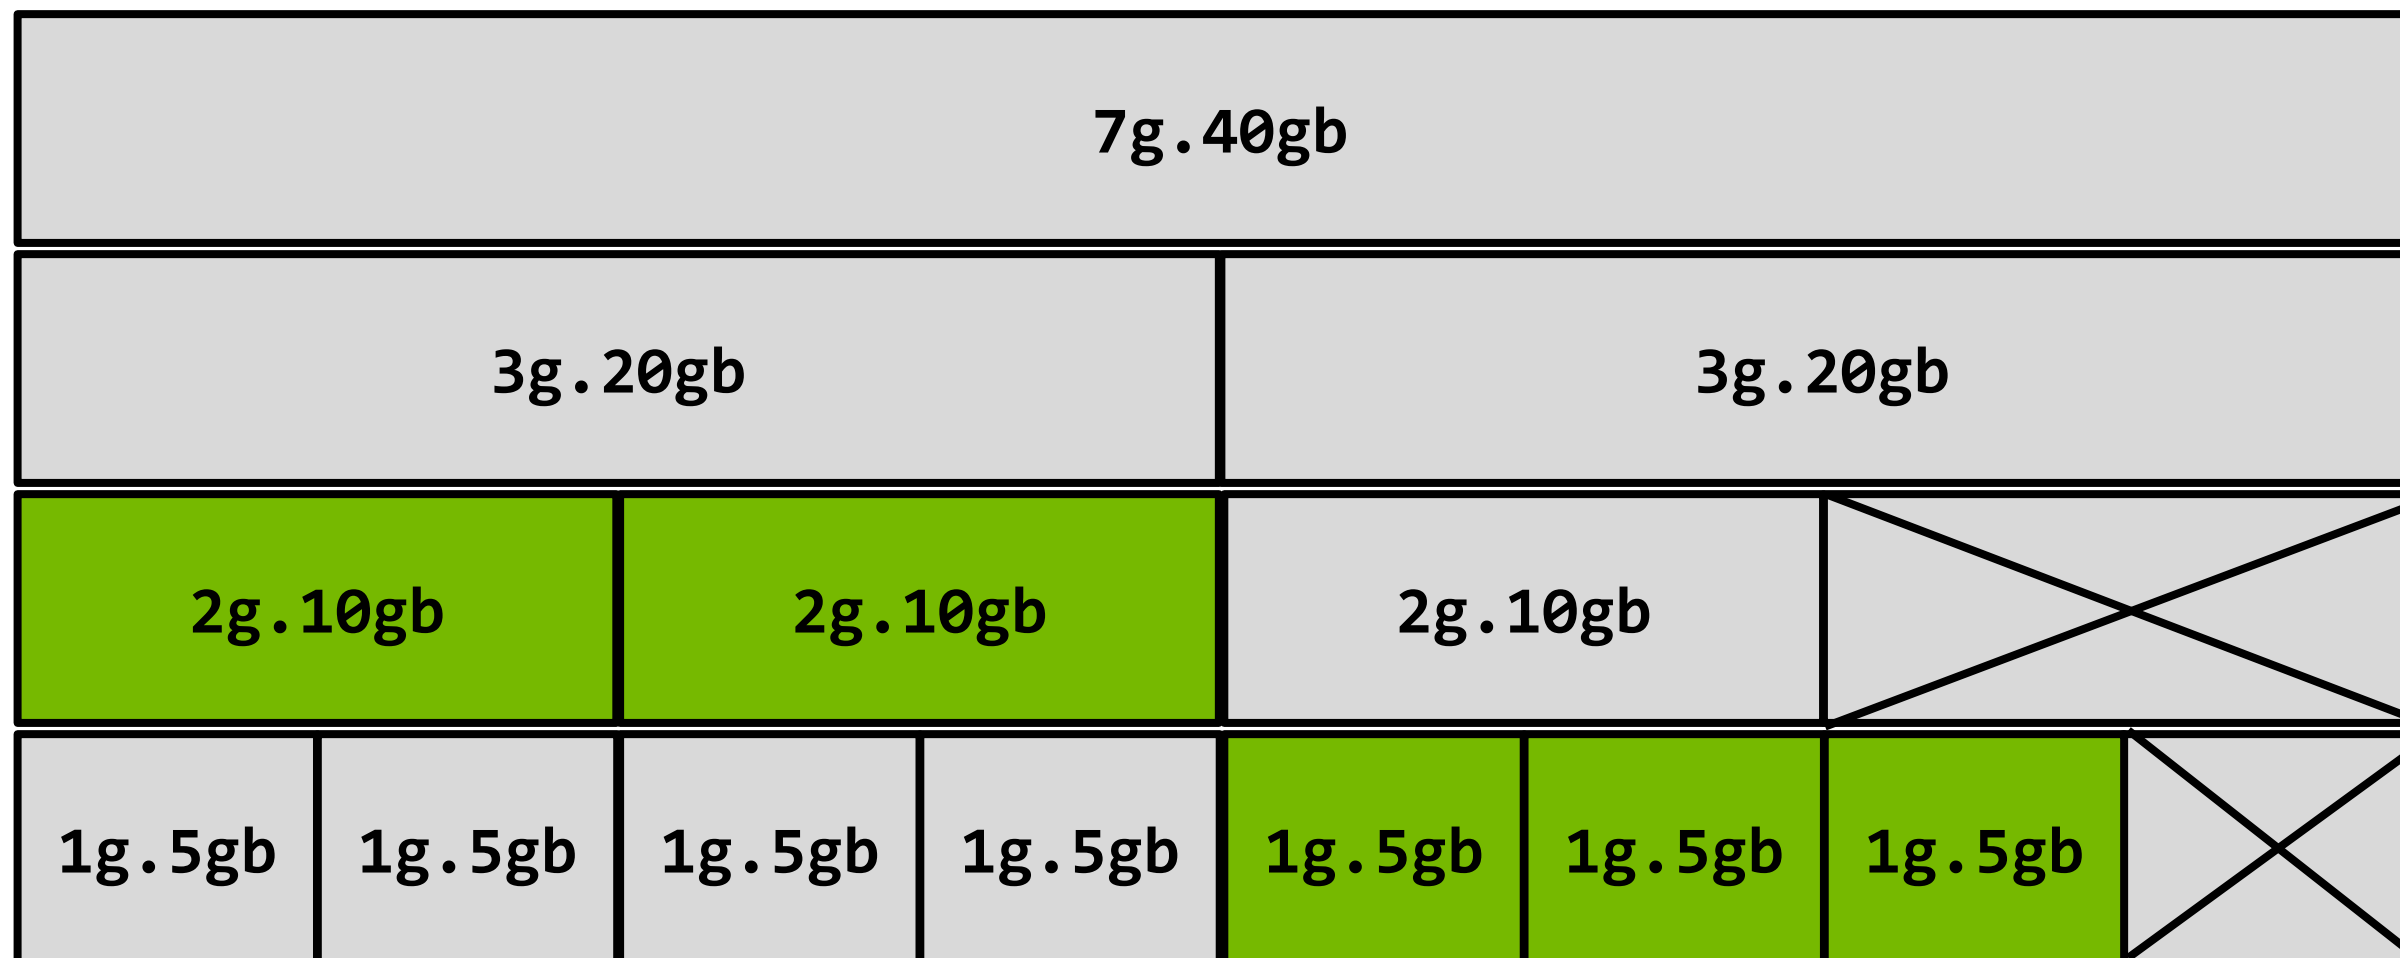
GPU Instances, Compute Instances, and MIG Devices



- 1 x 3g.20gb
and
- 2 x 2g.10gb

MULTI-INSTANCE GPUs (MIG)

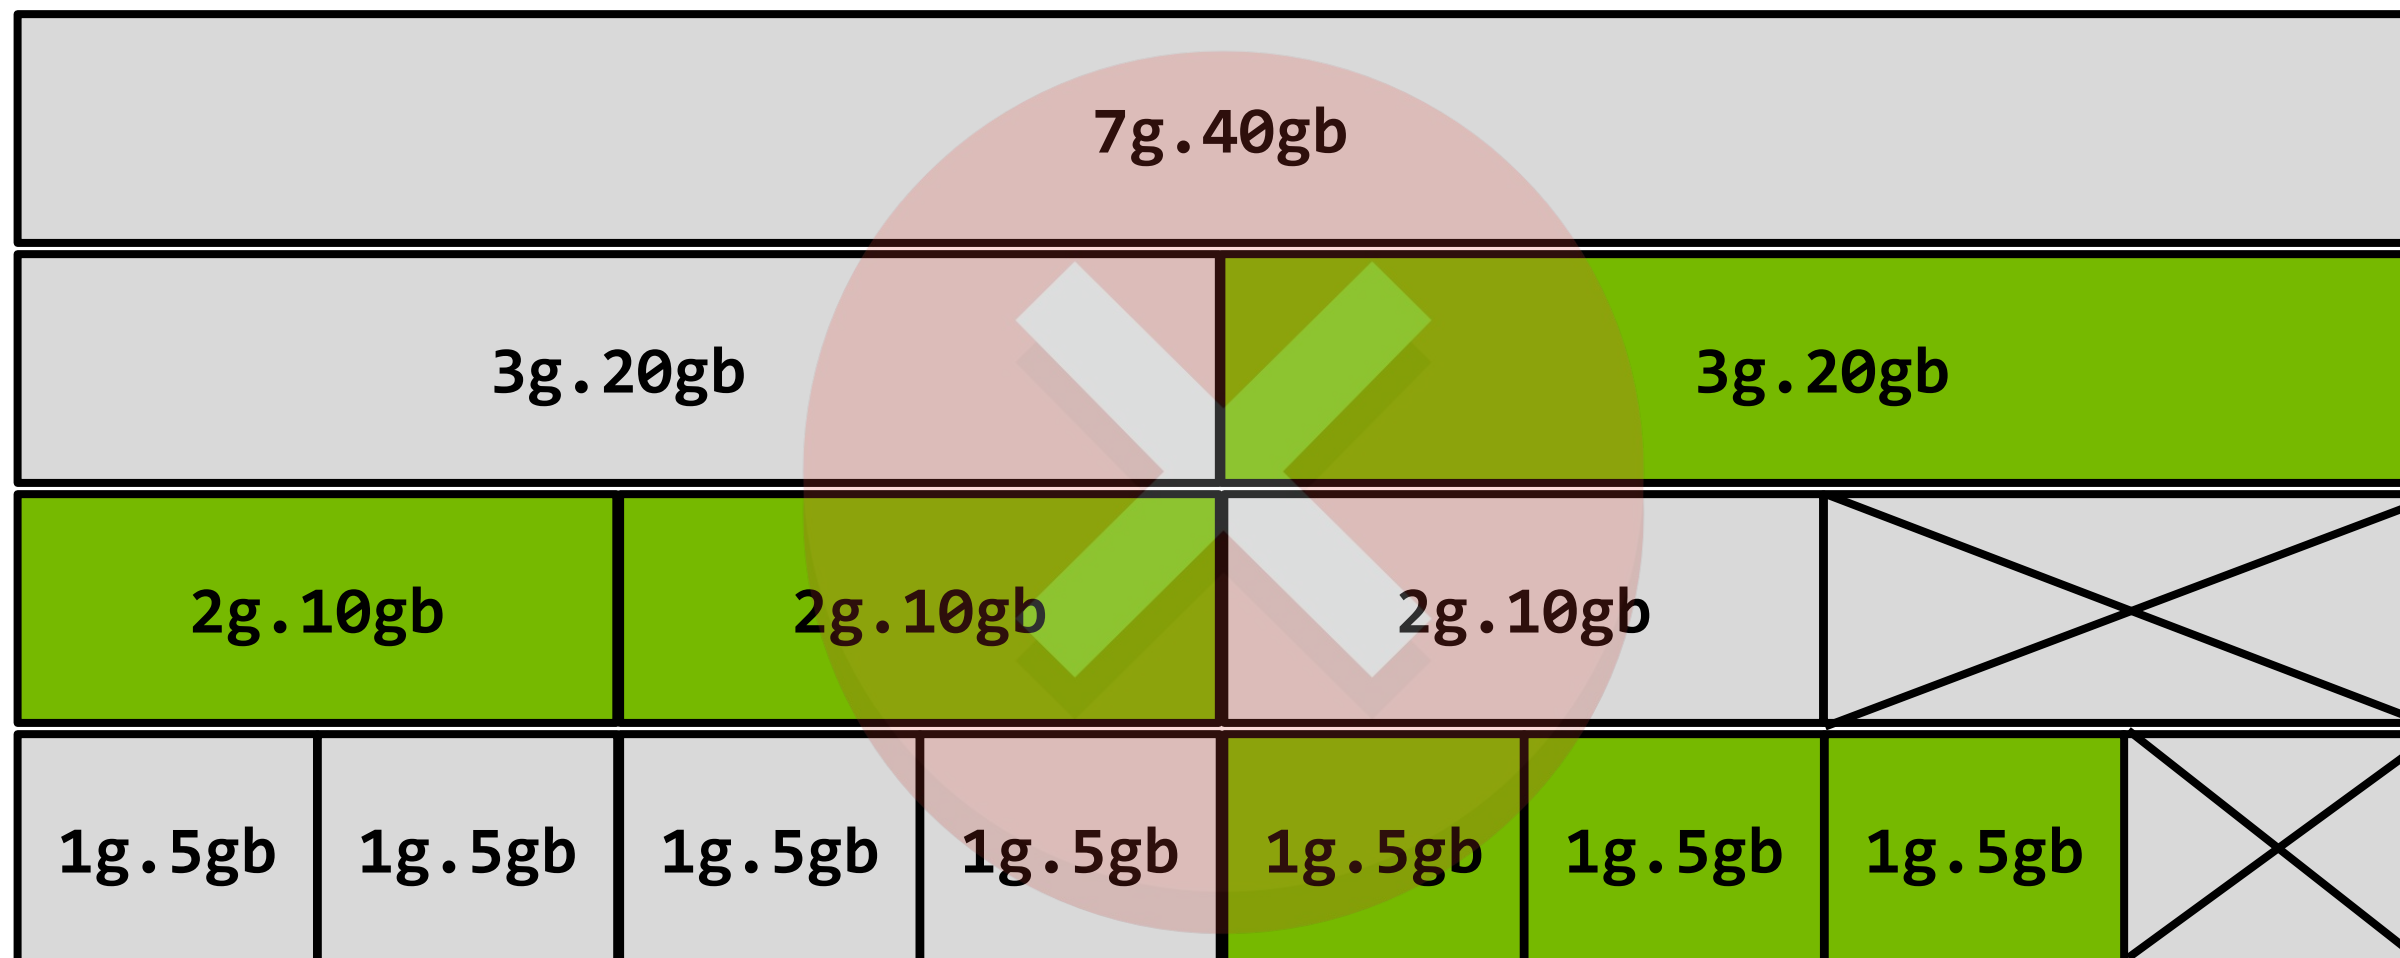
GPU Instances, Compute Instances, and MIG Devices



- 2 x 2g.10gb
and
- 3 x 1g.5gb

MULTI-INSTANCE GPUs (MIG)

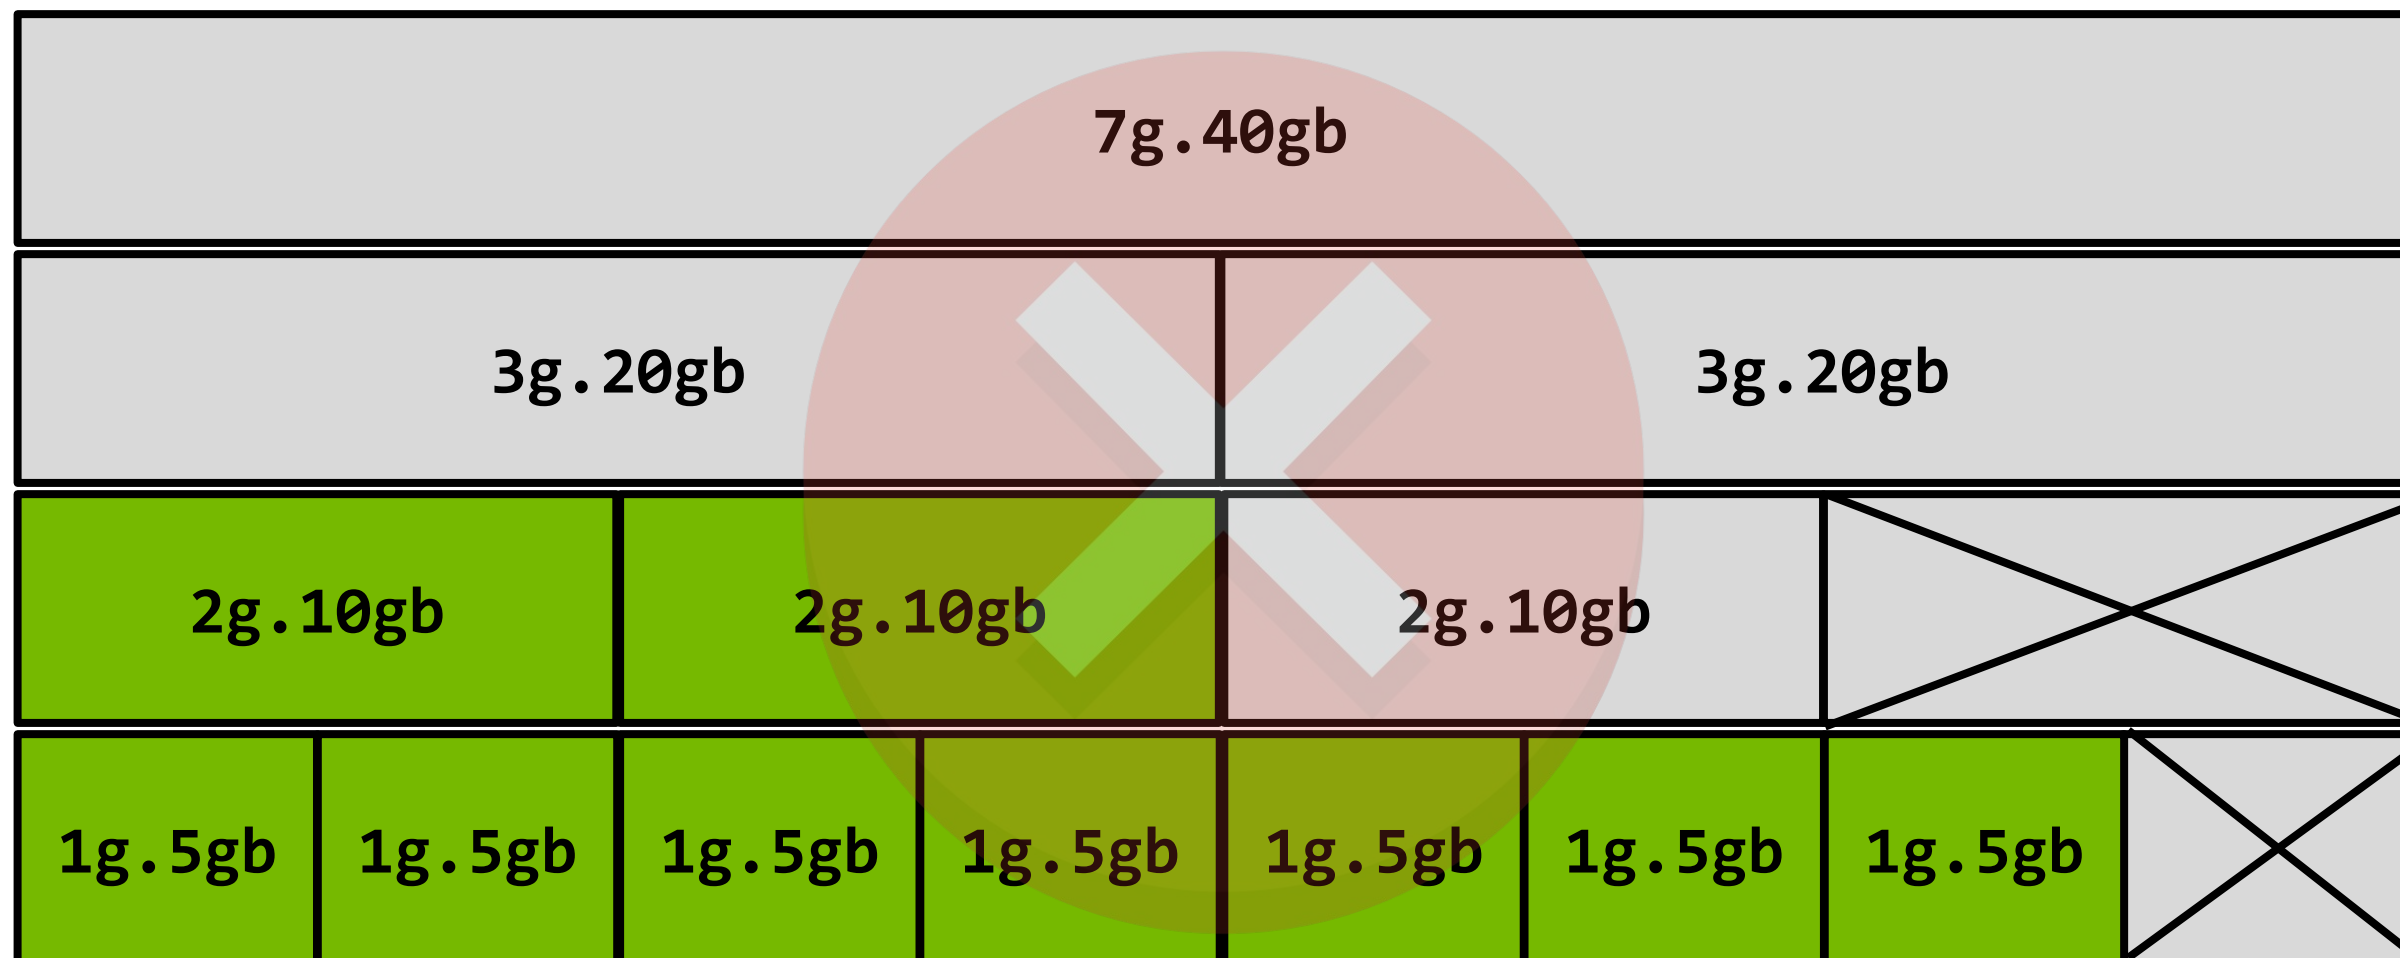
GPU Instances, Compute Instances, and MIG Devices



- 1 x 3g.30gb
and
- 2 x 2g.10gb
and
- 3 x 1g.5gb

MULTI-INSTANCE GPUs (MIG)

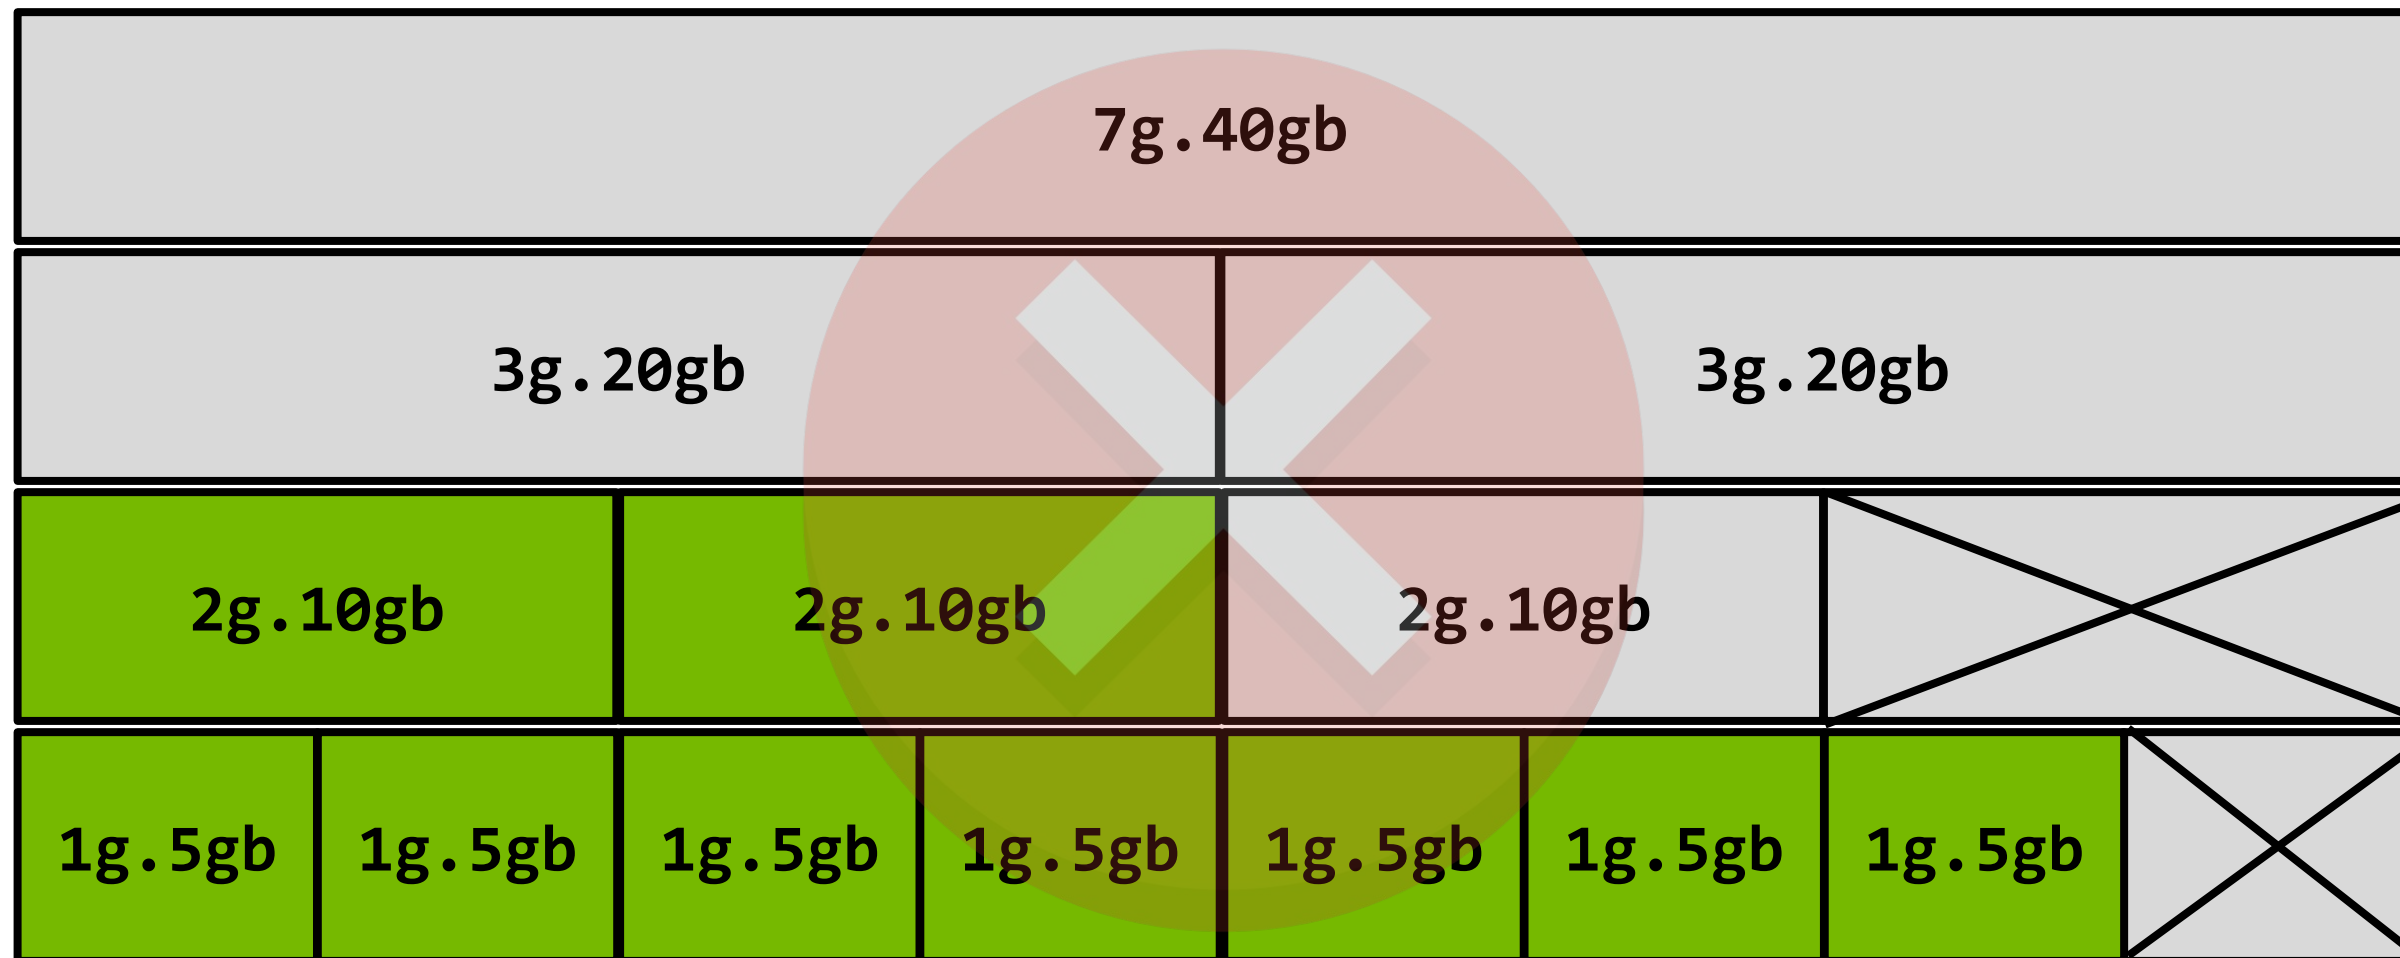
GPU Instances, Compute Instances, and MIG Devices



- 2 x 2g.10gb
and
- 7 x 1g.5gb

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices



- **2 x 2g.10gb**
and
- **7 x 1g.5gb**

→
No Overlapping Verticals

MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

DGX-A100



8 x A100



56 x 1g.5gb
or
24 x 2g.10gb
or
16 x 3g.20gb

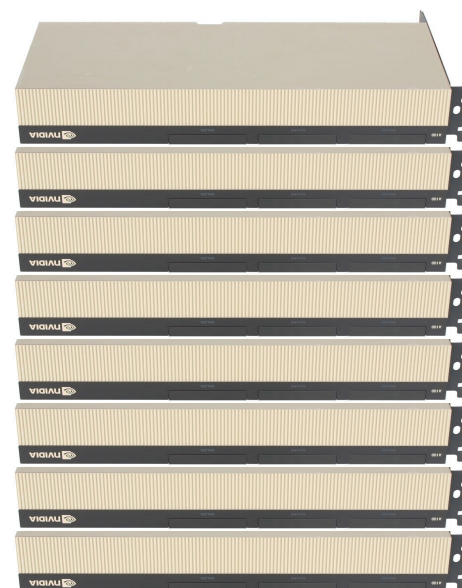
MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

DGX-A100



8 x A100



56 x 1g.5gb

or

24 x 2g.10gb

or

16 x 3g.20gb

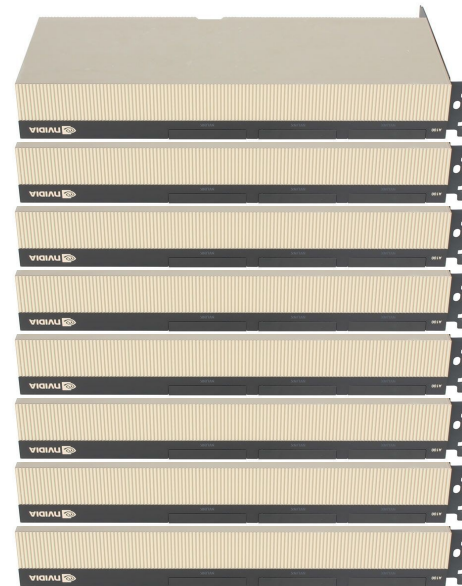
MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

DGX-A100



8 x A100



56 x 1g.5gb

or

24 x 2g.10gb

or

16 x 3g.20gb

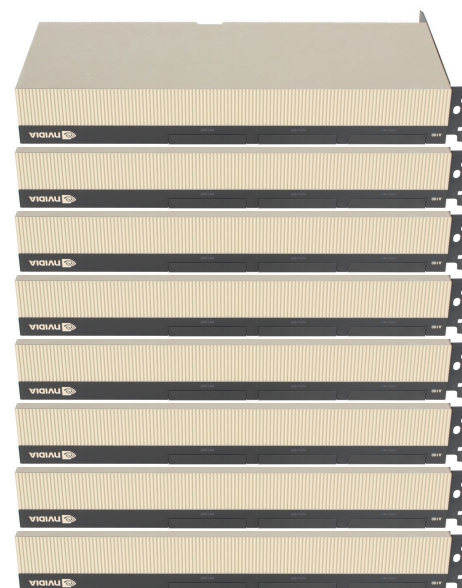
MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

DGX-A100



8 x A100



56 x 1g.5gb

or

24 x 2g.10gb

or

16 x 3g.20gb

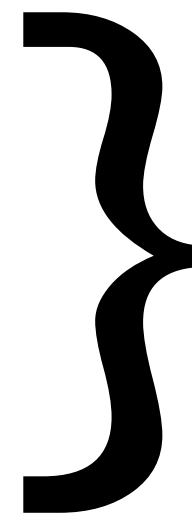
MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

DGX-A100



8 x A100



**16 x 1g.5gb
and
8 x 2g.10gb
and
8 x 3g.20gb**

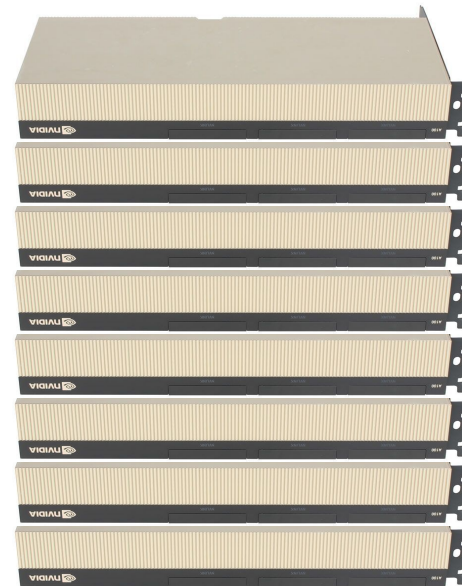
MULTI-INSTANCE GPUs (MIG)

GPU Instances, Compute Instances, and MIG Devices

DGX-A100



8 x A100



4 x Full A100s

and

8 x 1g.5gb

and

4 x 2g.10gb

and

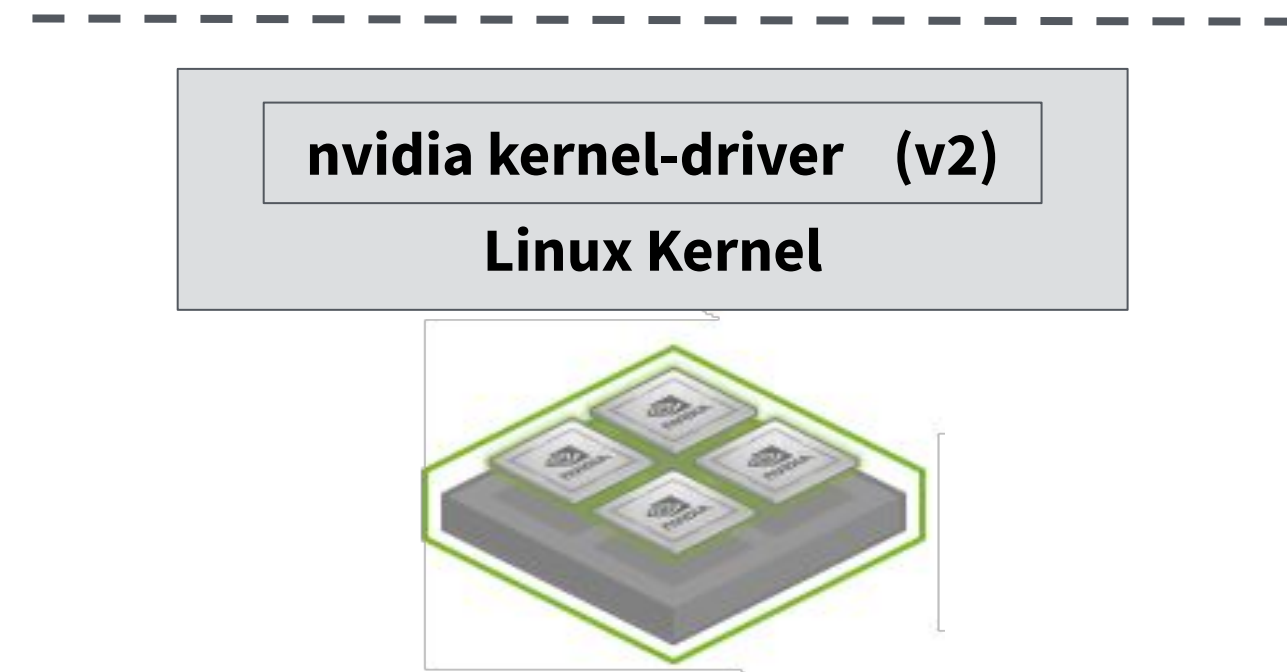
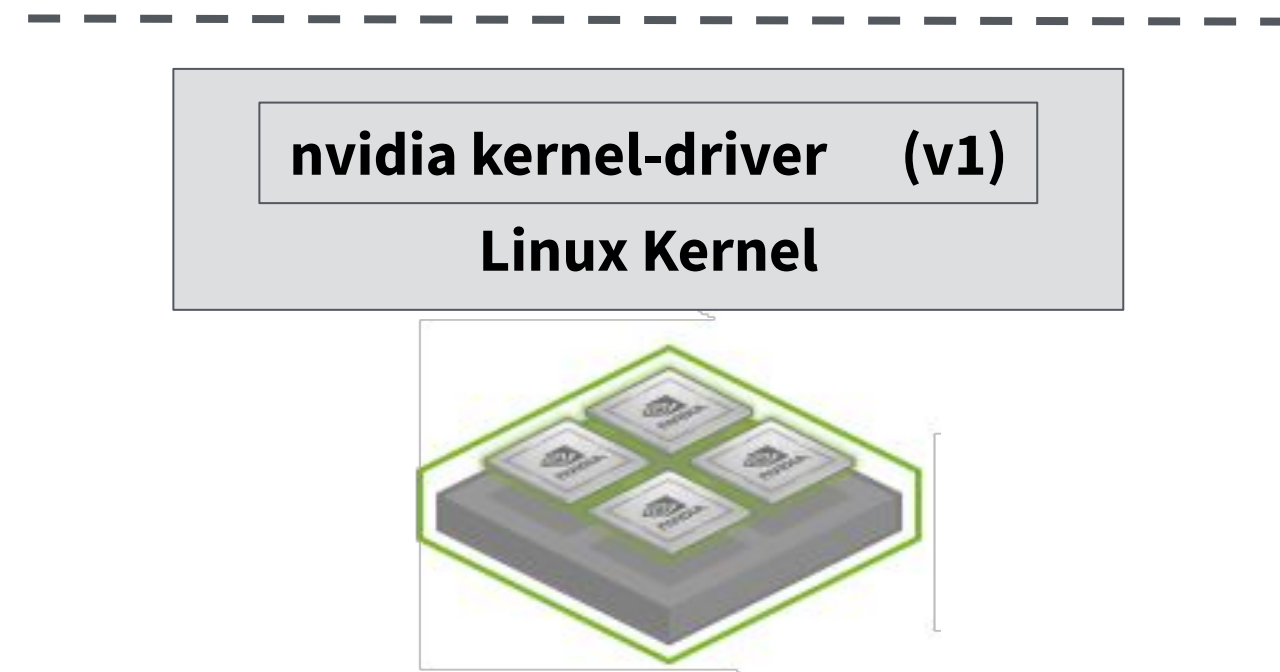
4 x 3g.20gb

GPUs AND CONTAINERS

The NVIDIA Container Toolkit

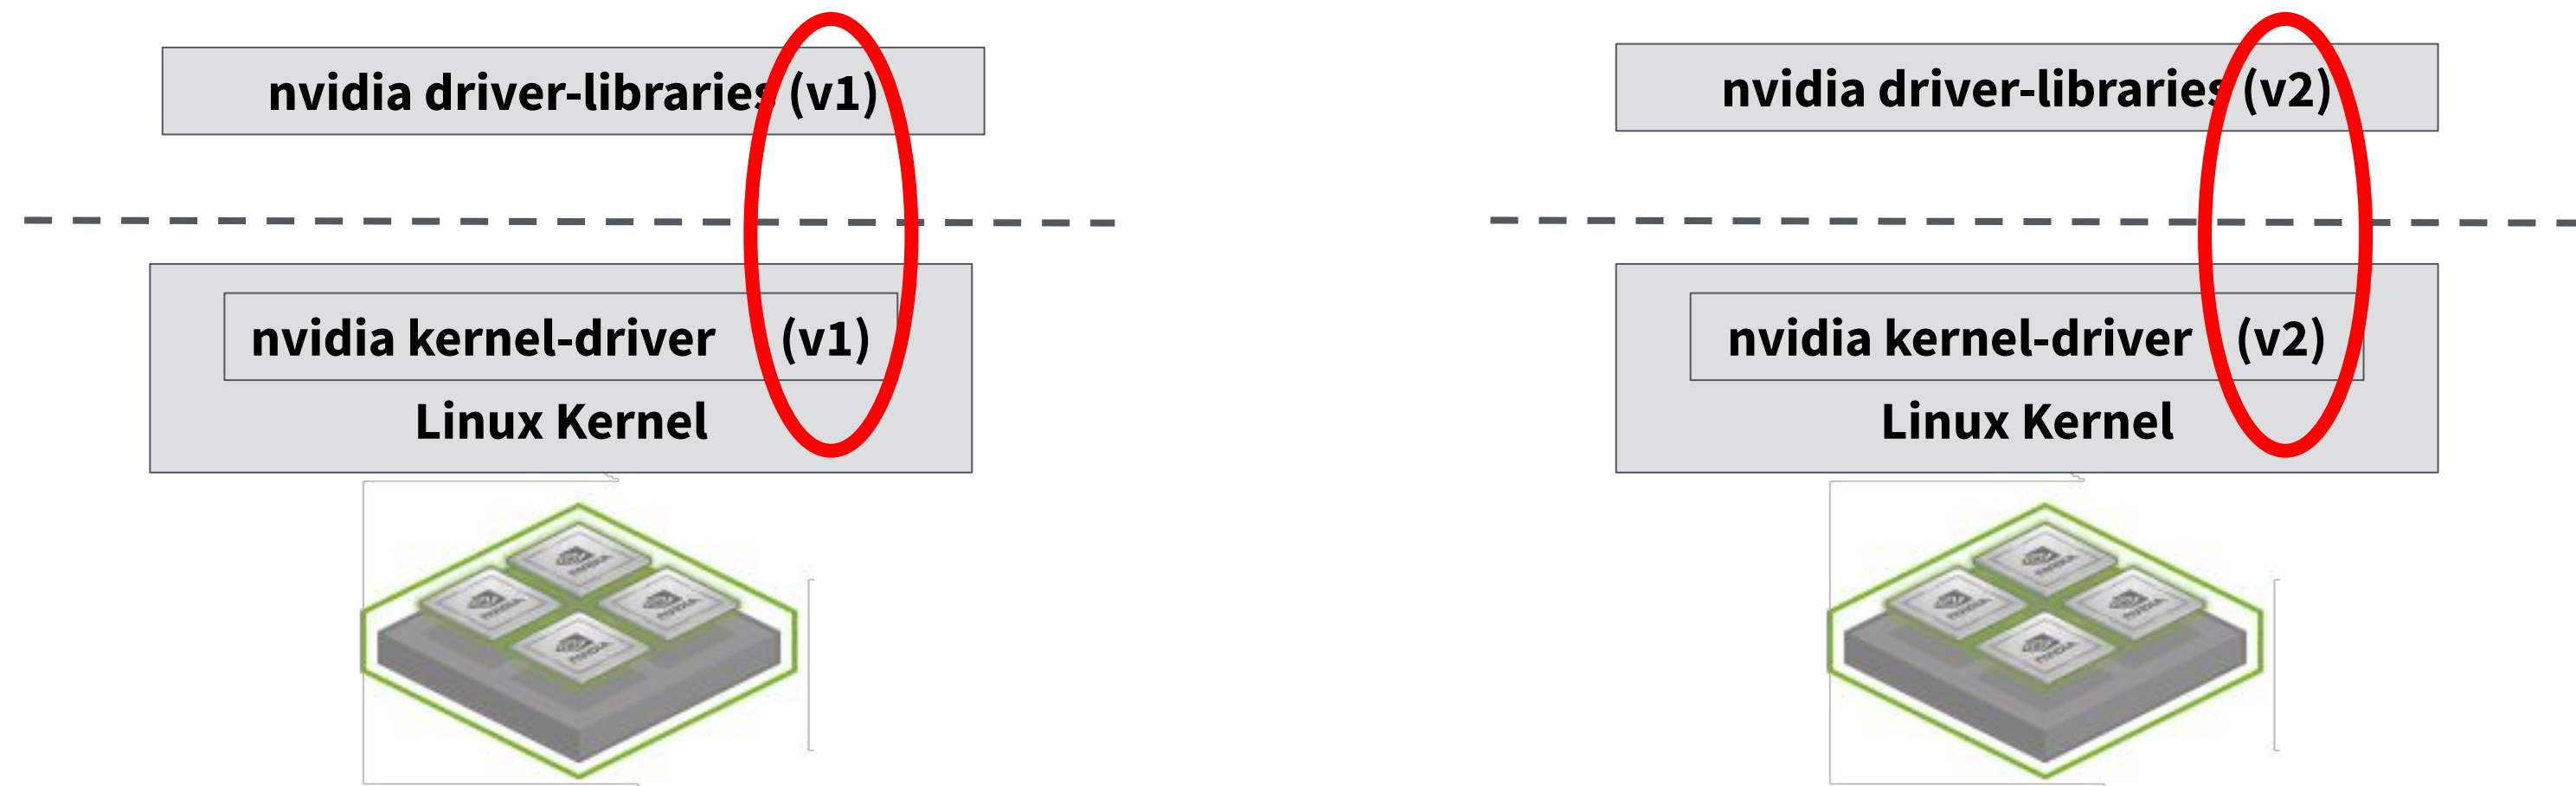
GPUs AND CONTAINERS

The NVIDIA Container Toolkit



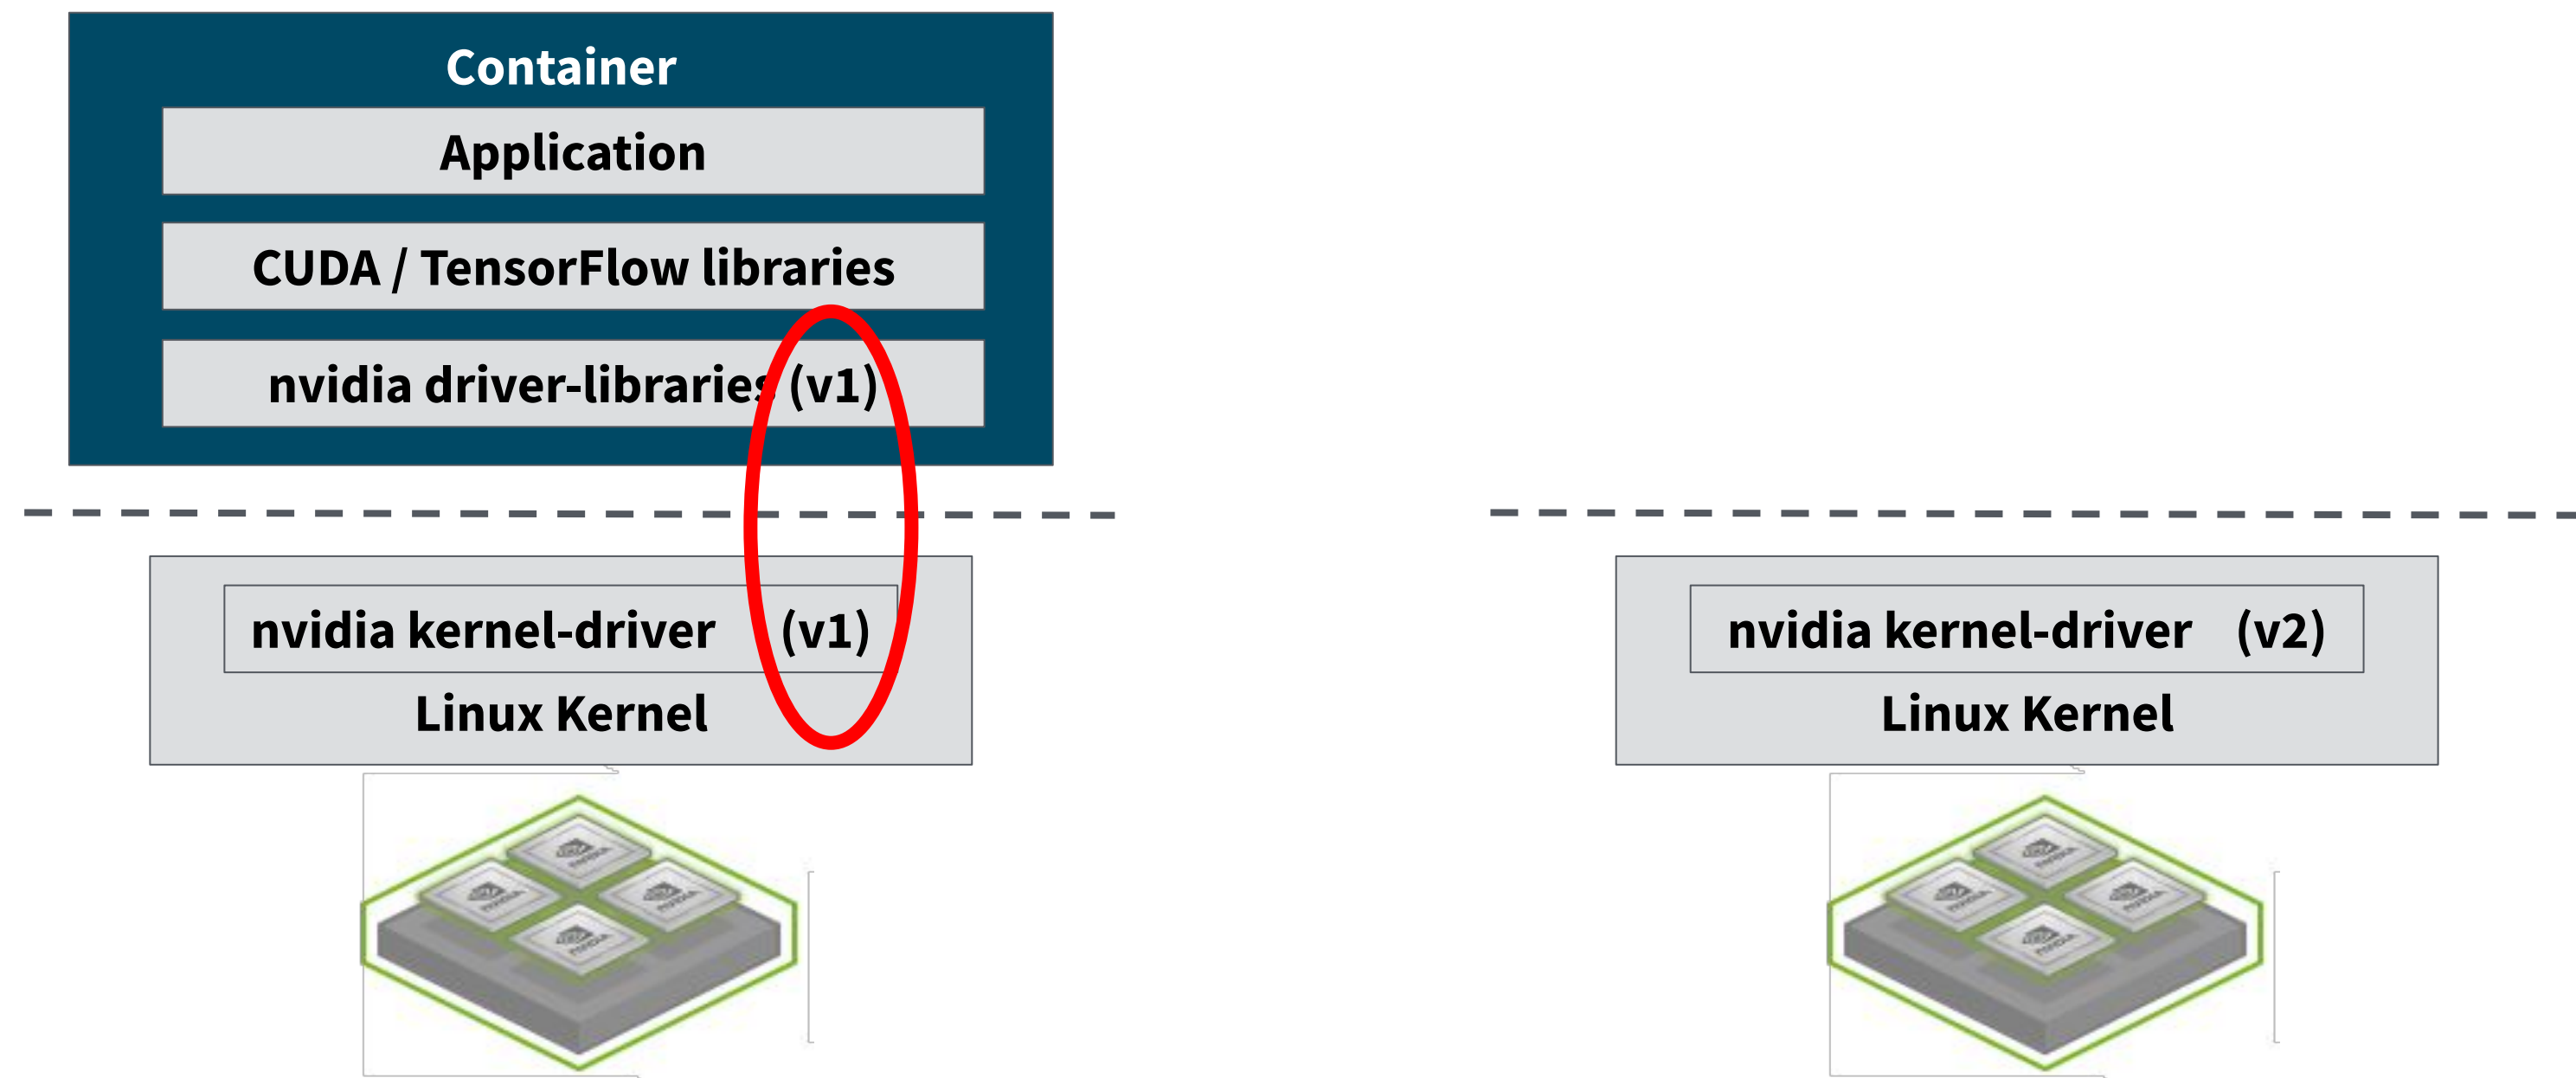
GPUs AND CONTAINERS

The NVIDIA Container Toolkit



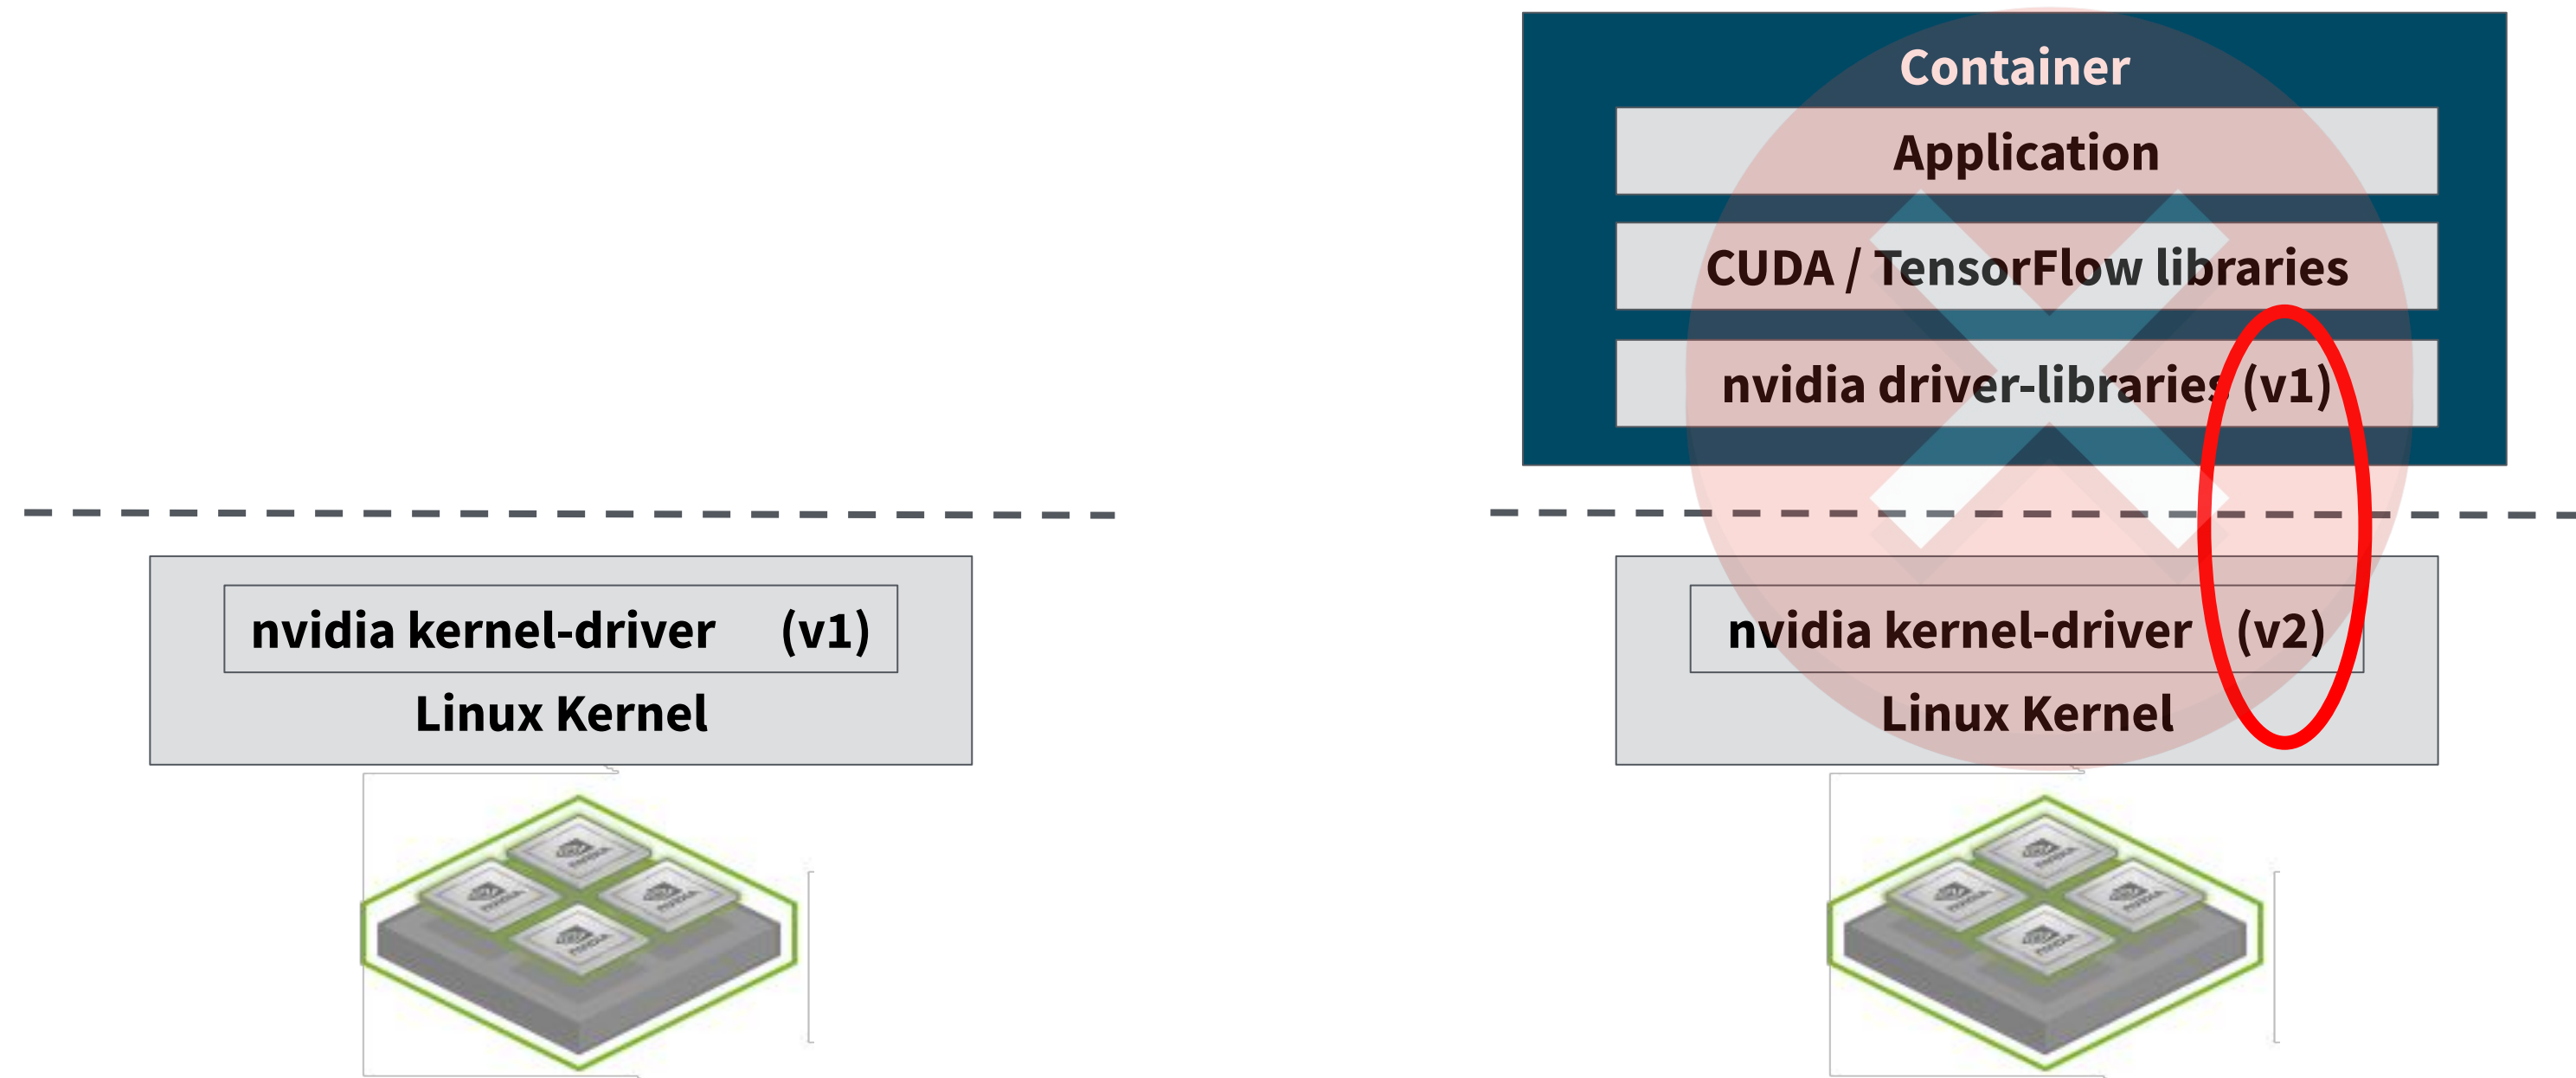
GPUs AND CONTAINERS

The NVIDIA Container Toolkit



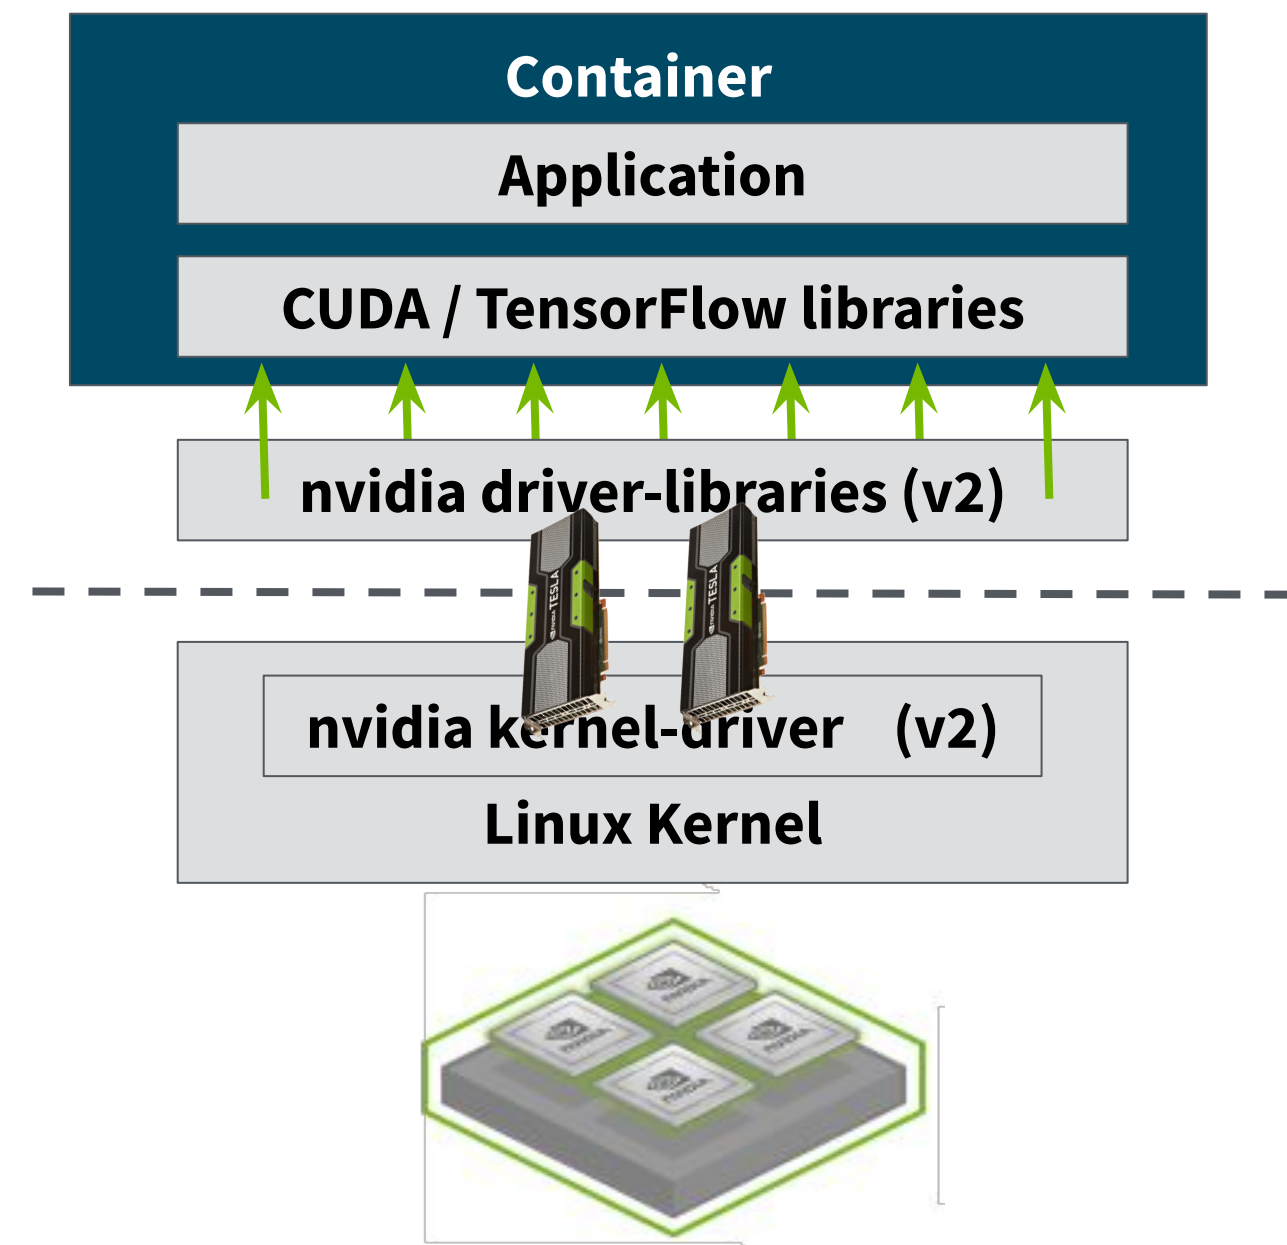
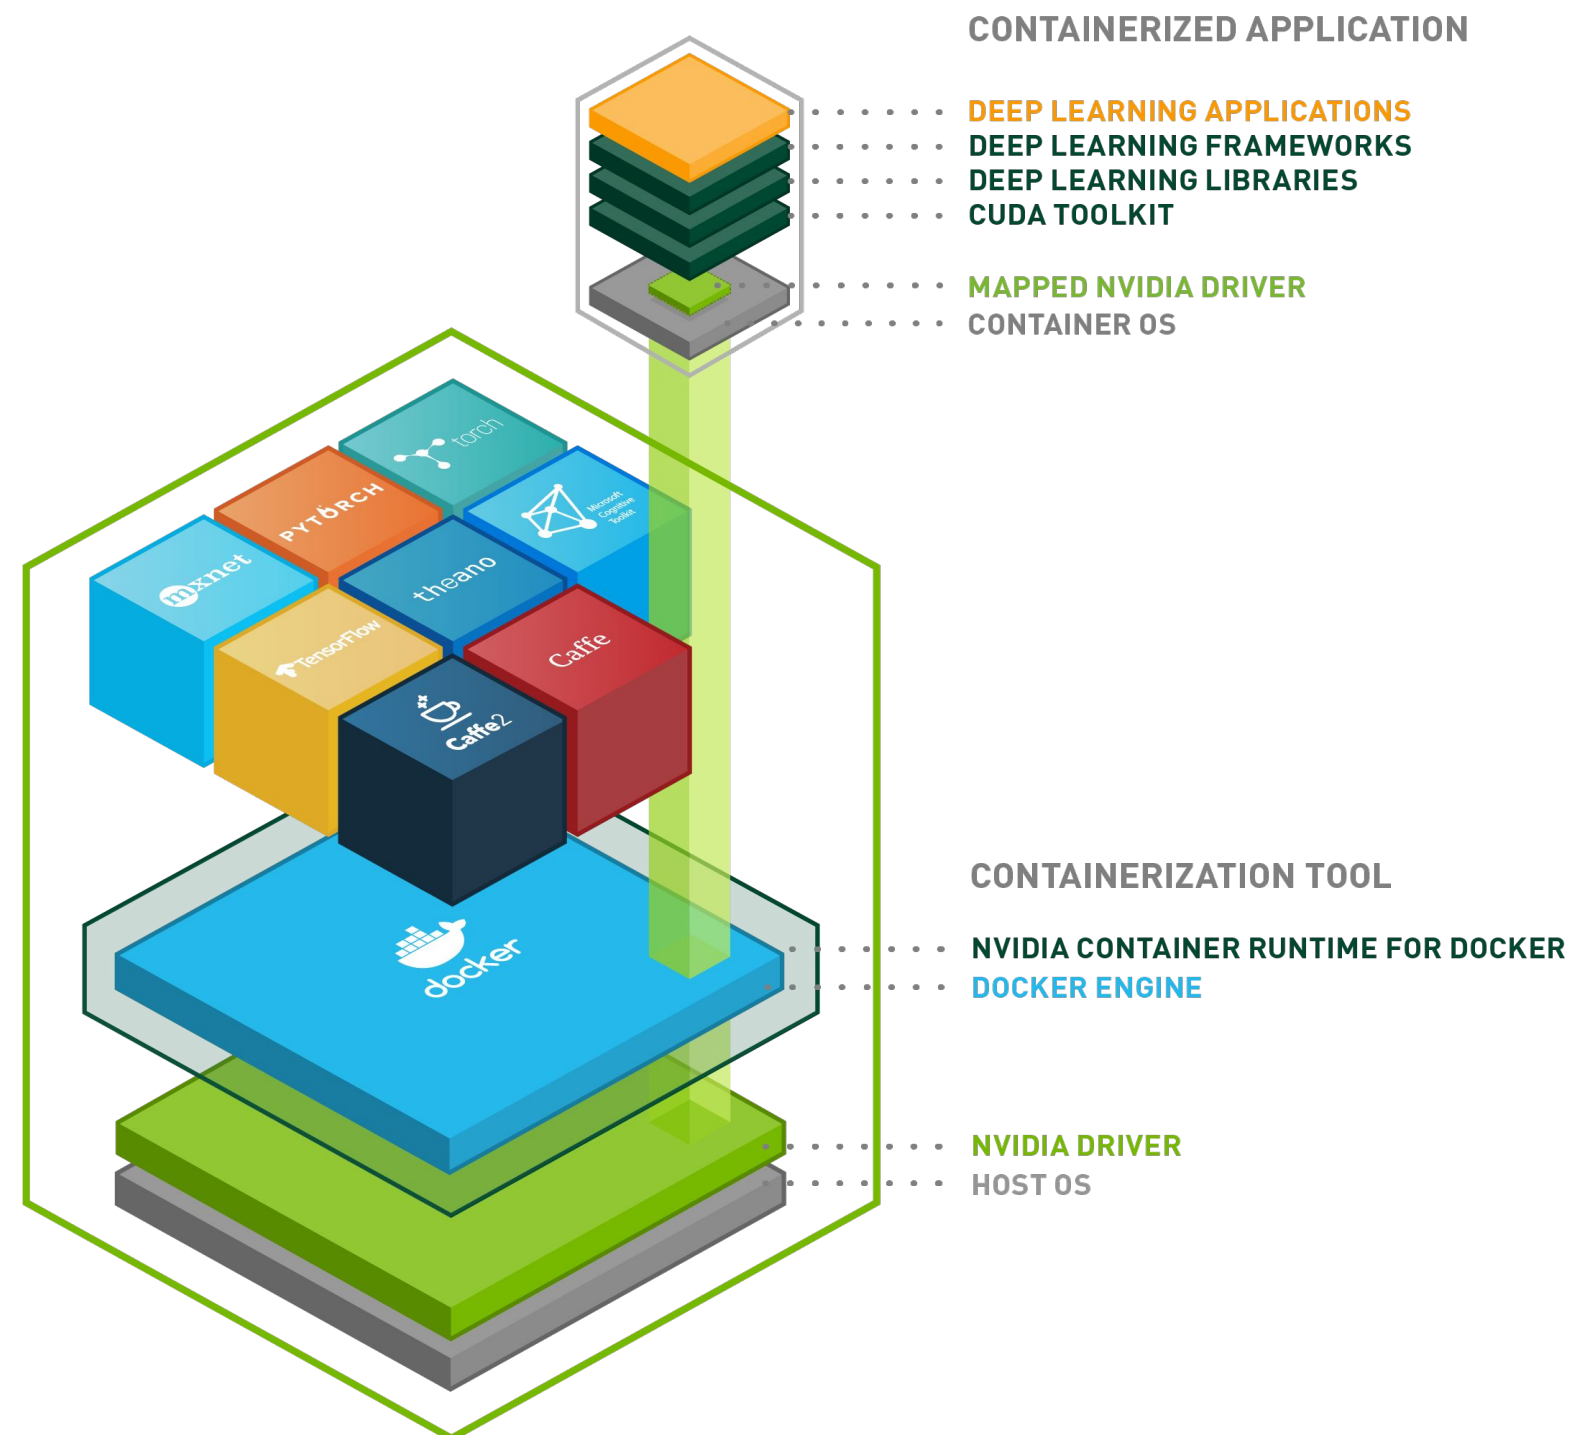
GPUs AND CONTAINERS

The NVIDIA Container Toolkit



GPUs AND CONTAINERS

The NVIDIA Container Toolkit



GPUs AND CONTAINERS

The NVIDIA Container Toolkit

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

```
GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)
```

```
GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)
```

GPUs AND CONTAINERS

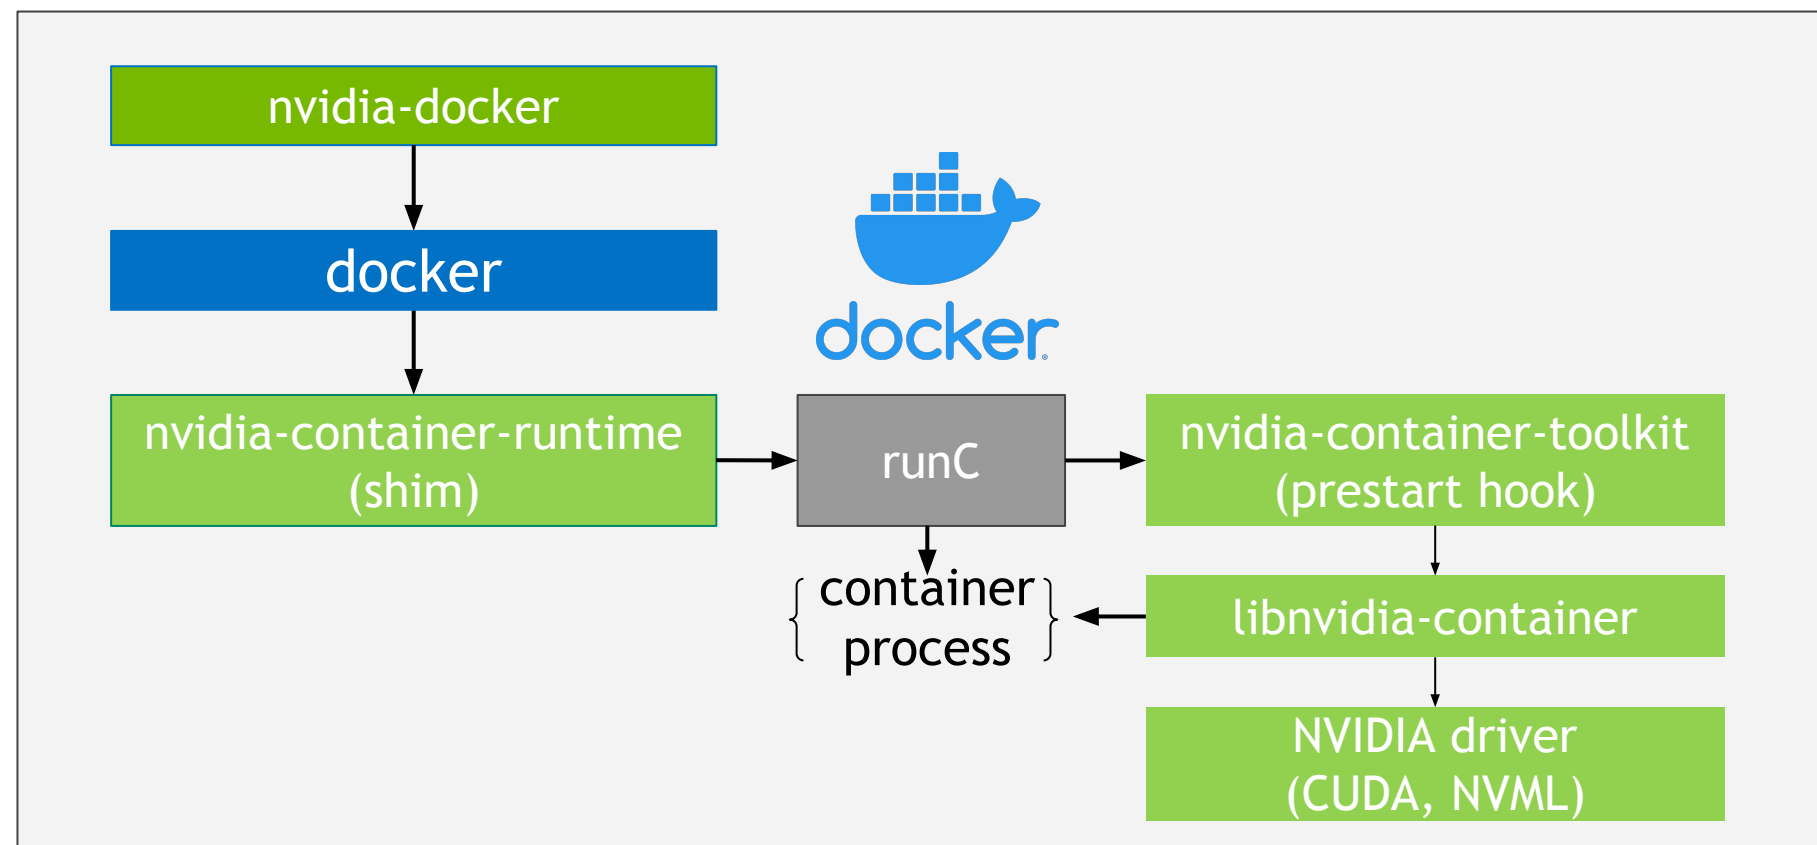
The NVIDIA Container Toolkit

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

- **nvidia-docker**
- **nvidia-container-runtime**
- **nvidia-container-toolkit**
- **libnvidia-container**



GPUs AND CONTAINERS

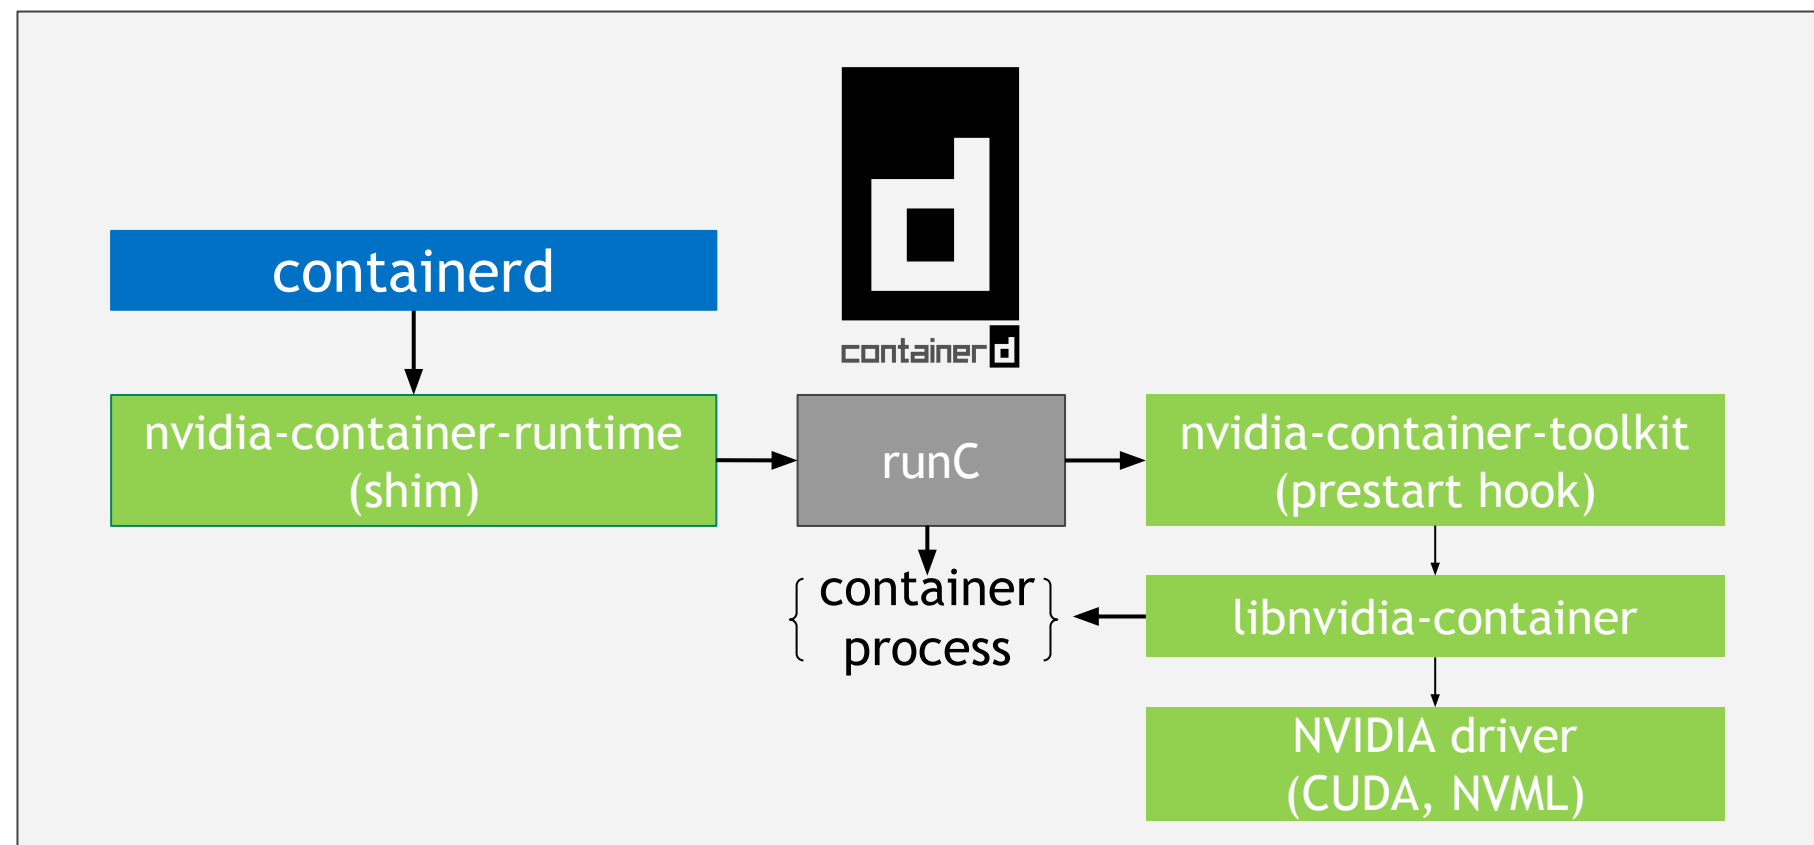
The NVIDIA Container Toolkit

```
$ docker run \
  --gpus '"device=0,1"' \
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

- **nvidia-docker**
- **nvidia-container-runtime**
- **nvidia-container-toolkit**
- **libnvidia-container**



GPUs AND CONTAINERS

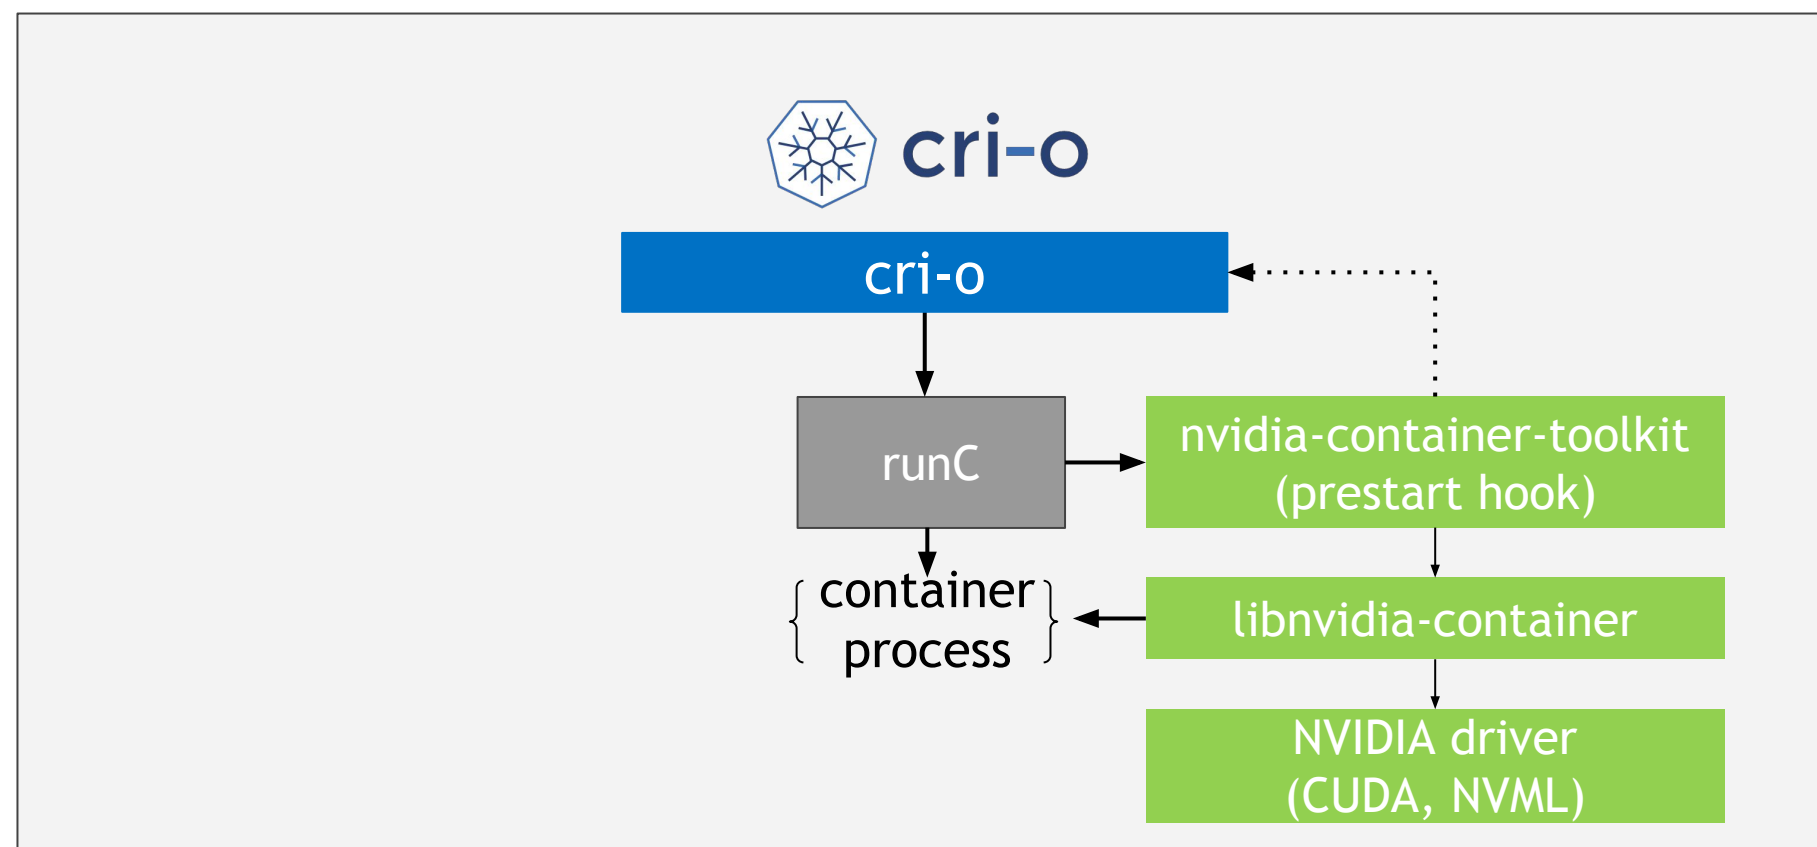
The NVIDIA Container Toolkit

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

- **nvidia-docker**
- **nvidia-container-runtime**
- **nvidia-container-toolkit**
- **libnvidia-container**



GPUs AND CONTAINERS

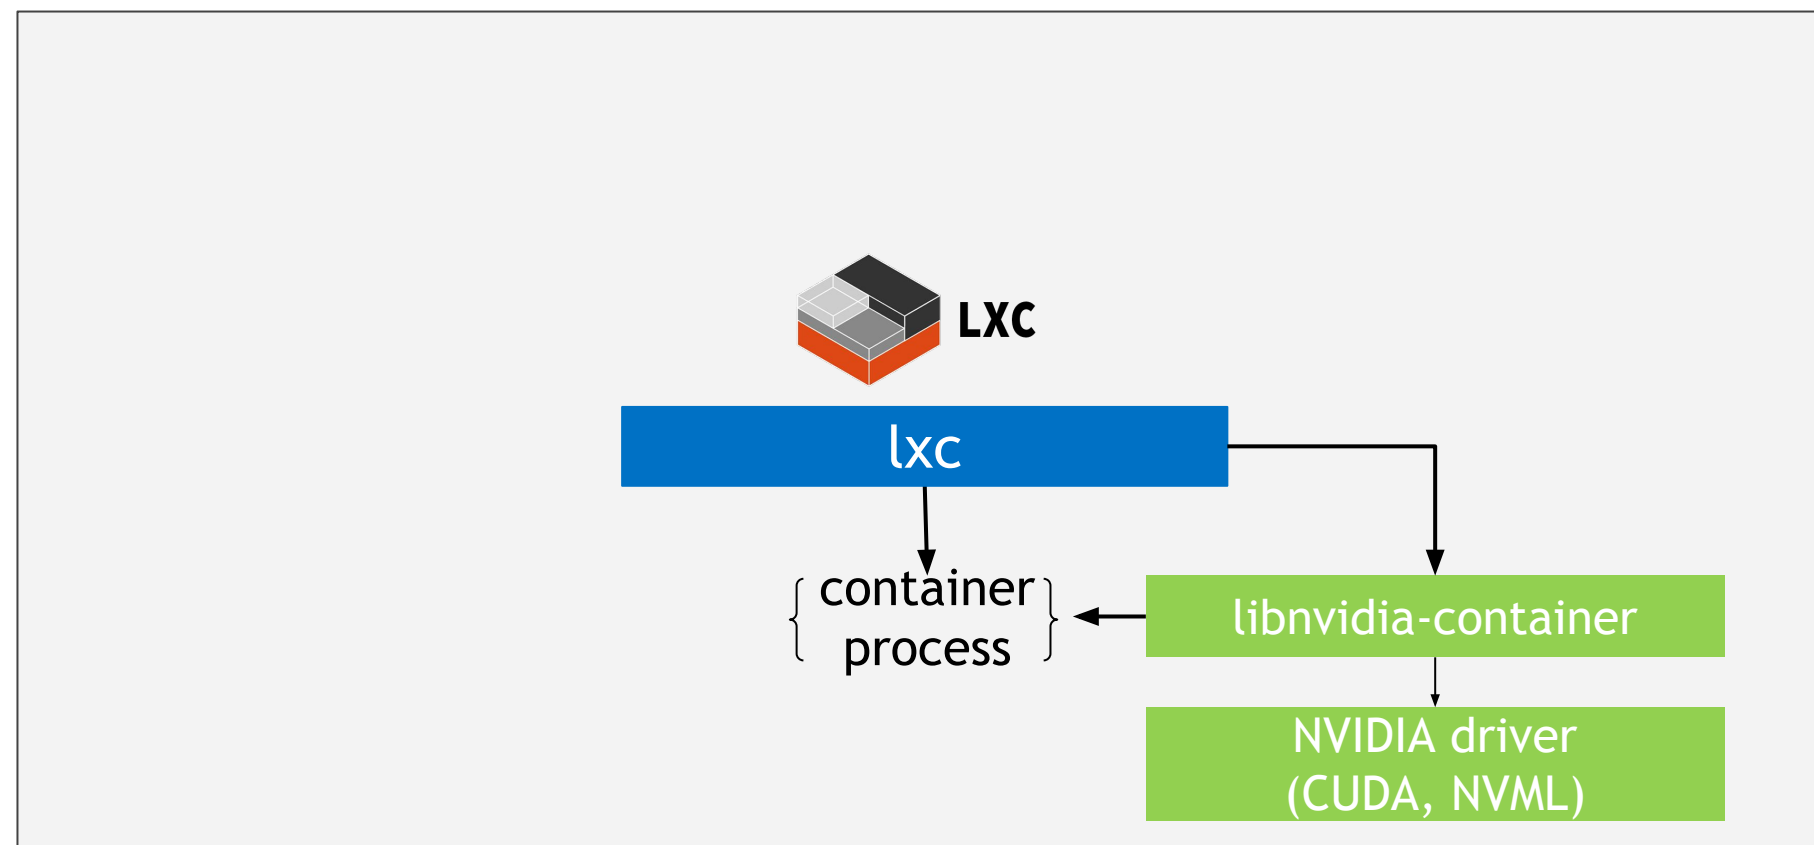
The NVIDIA Container Toolkit

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

- **nvidia-docker**
- **nvidia-container-runtime**
- **nvidia-container-toolkit**
- **libnvidia-container**



GPUs AND CONTAINERS

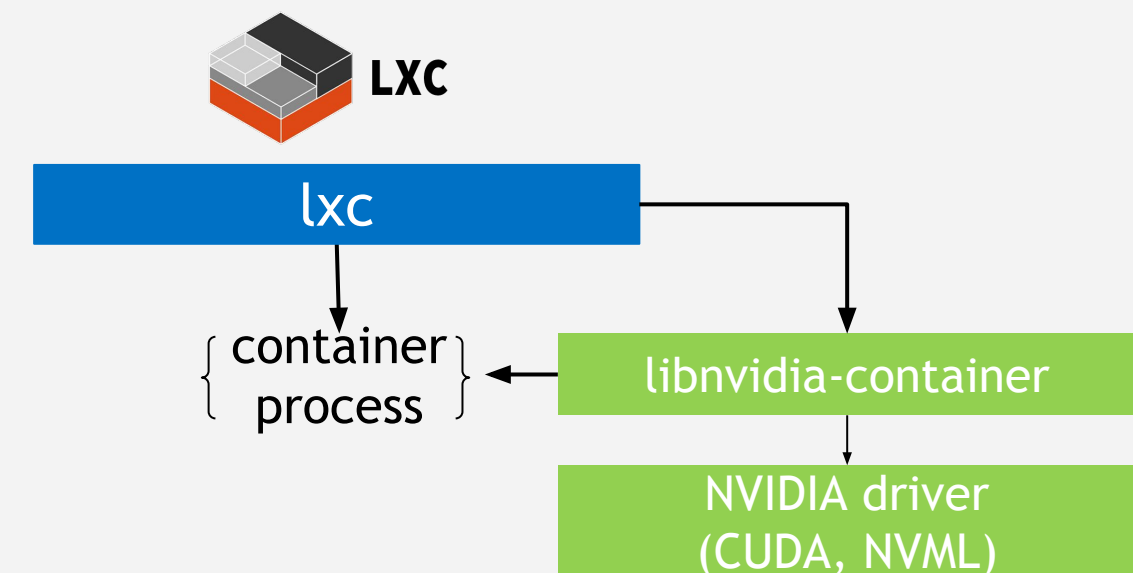
The NVIDIA Container Toolkit

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

- **nvidia-docker**
- **nvidia-container-runtime**
- **nvidia-container-toolkit**
- **libnvidia-container**



GPUs AND CONTAINERS

The NVIDIA Container Toolkit

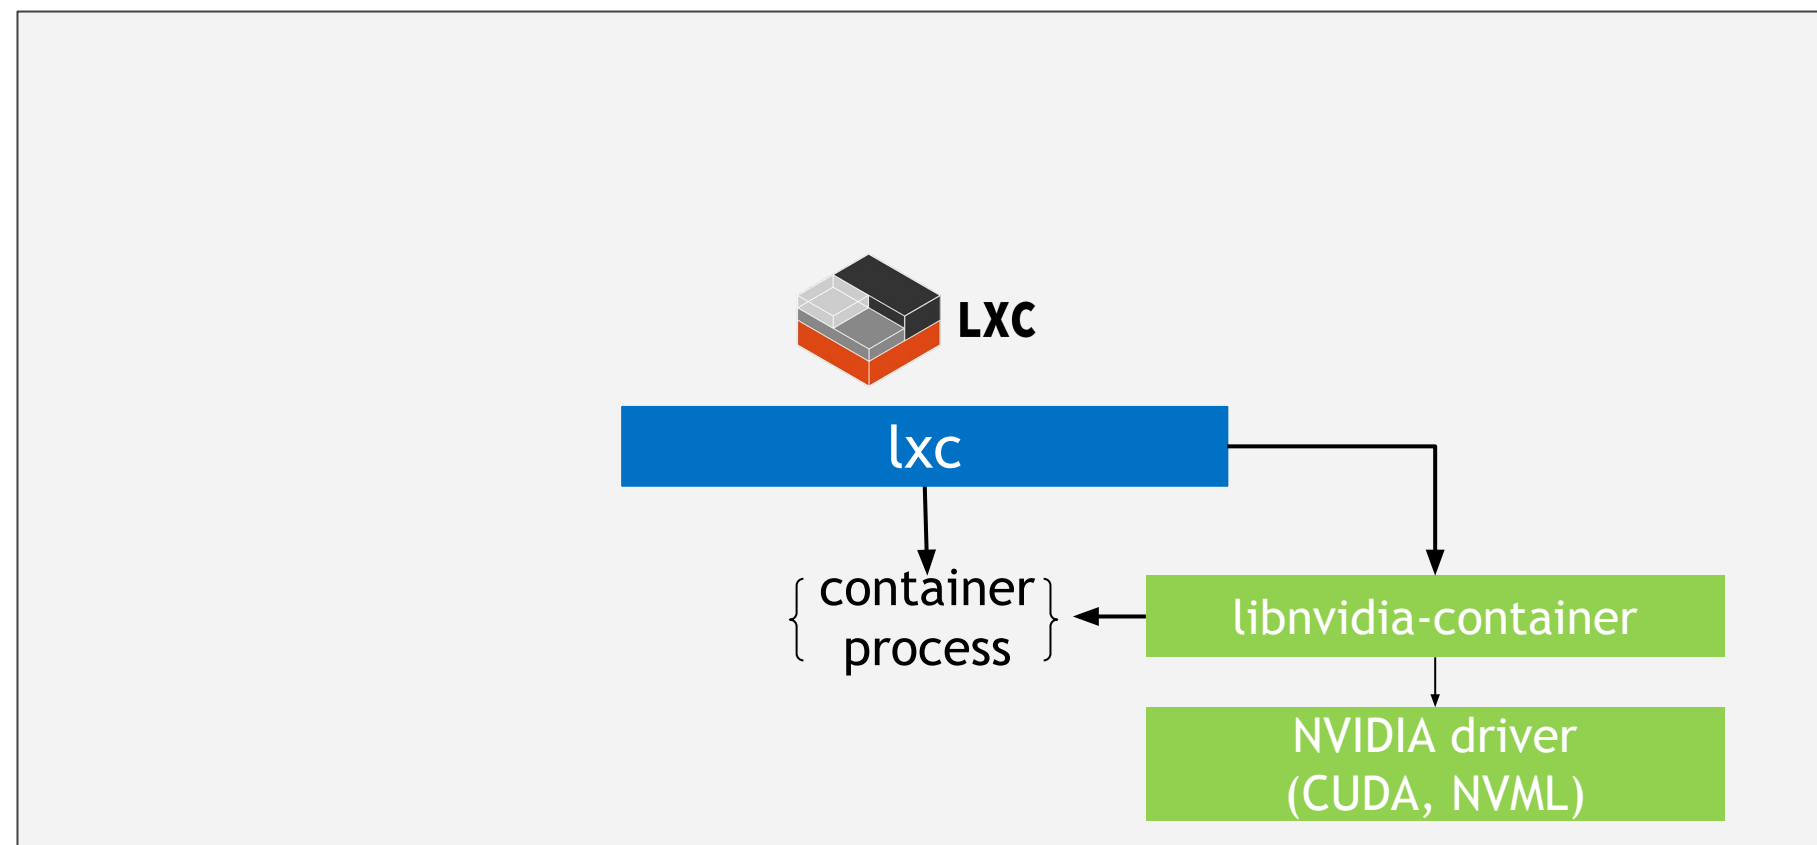
```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

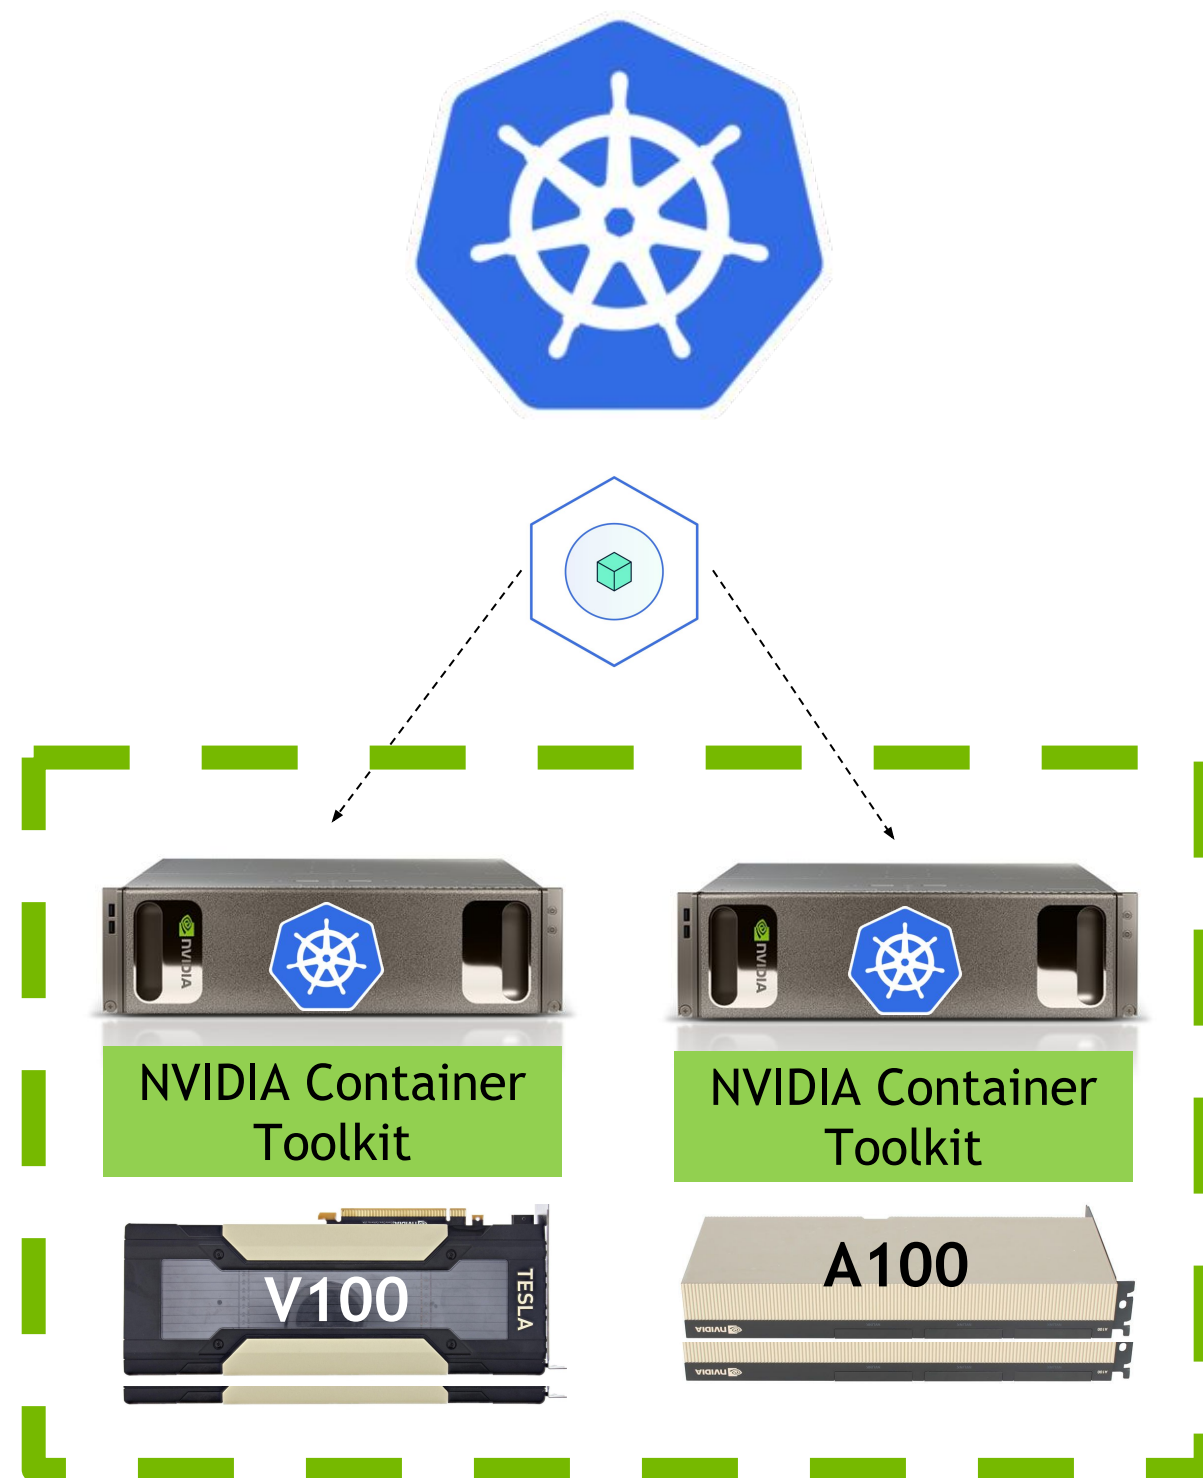
- **nvidia-docker**
- **nvidia-container-runtime**
- **nvidia-container-toolkit**
- **libnvidia-container**

Majority of code for MIG support
in containers added here
v1.3.0+



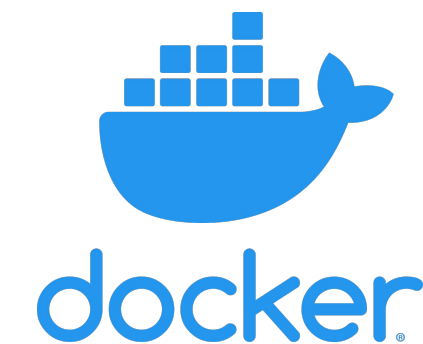
GPUs AND KUBERNETES

Allocate GPUs to pods in a Kubernetes Cluster



GPUs AND KUBERNETES

Allocate GPUs to pods in a Kubernetes Cluster



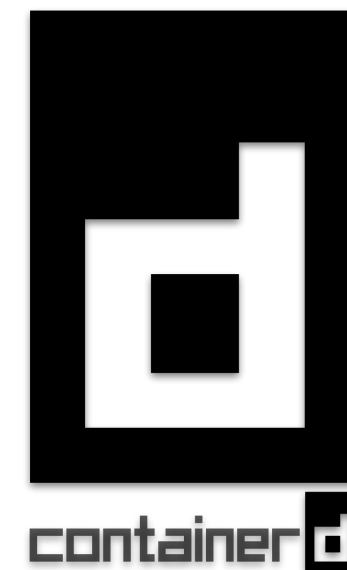
```
$ cat /etc/docker/daemon.json
{
  "default-runtime": "nvidia",
  "runtimes": {
    "nvidia": {
      "path": "/usr/bin/nvidia-container-runtime",
      "runtimeArgs": []
    }
  }
}
```

GPUs AND KUBERNETES

Allocate GPUs to pods in a Kubernetes Cluster

```
$ cat /etc/containerd/config.toml
[plugins]
[plugins."io.containerd.grpc.v1.cri"]
[plugins."io.containerd.grpc.v1.cri".containerd]
    default_runtime_name = "nvidia"

[plugins."io.containerd.grpc.v1.cri".containerd.runtimes]
    [plugins."io.containerd.grpc.v1.cri".containerd.runtimes.nvidia]
        privileged_without_host_devices = false
        runtime_engine = ""
        runtime_root = ""
        runtime_type = "io.containerd.runc.v2"
    [plugins."io.containerd.grpc.v1.cri".containerd.runtimes.nvidia.options]
        BinaryName = "/usr/bin/nvidia-container-runtime"
```



GPUs AND KUBERNETES

Allocate GPUs to pods in a Kubernetes Cluster



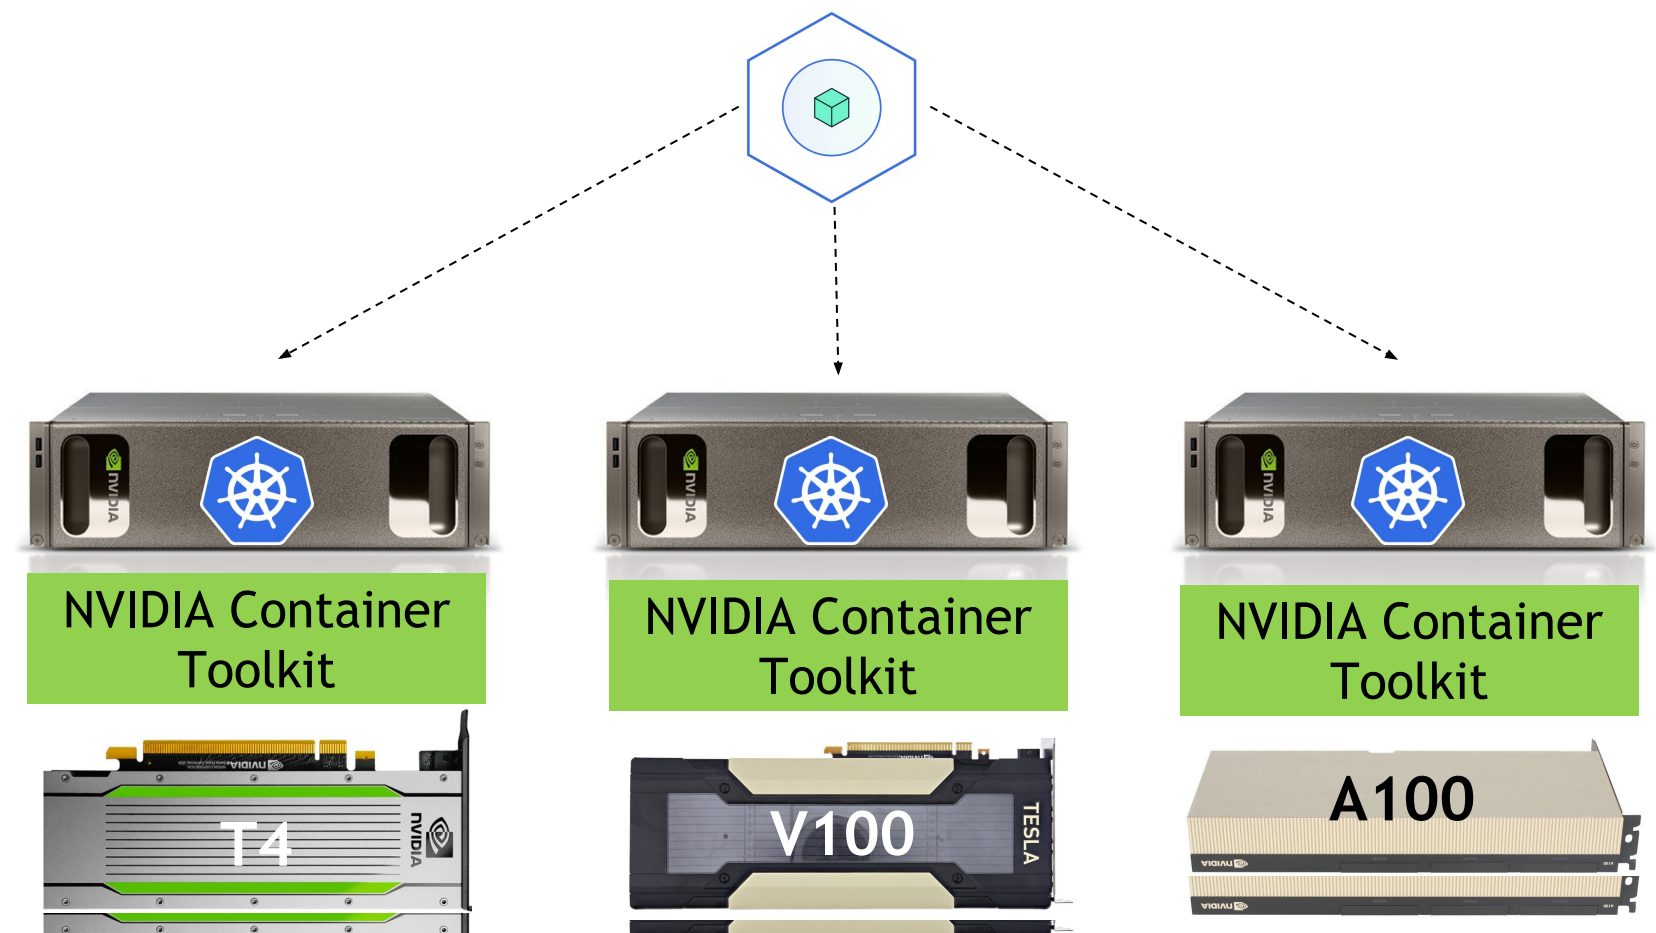
```
$ cat /usr/share/containers/oci/hooks.d/oci-nvidia-hook.json
{
  "version": "1.0.0",
  "hook": {
    "path": "/usr/bin/nvidia-container-toolkit",
    "args": ["nvidia-container-toolkit", "prestart"]
  },
  "when": {
    "always": true,
    "commands": [".*"]
  },
  "stages": ["prestart"]
}
```

GPUs AND KUBERNETES

Allocate GPUs to pods in a Kubernetes Cluster

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/gpu: 4
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
    nvidia.com/cuda.runtime: 11.0
    nvidia.com/cuda.driver: 450.51.06
```

- **k8s-device-plugin**
- **gpu-feature-discovery**

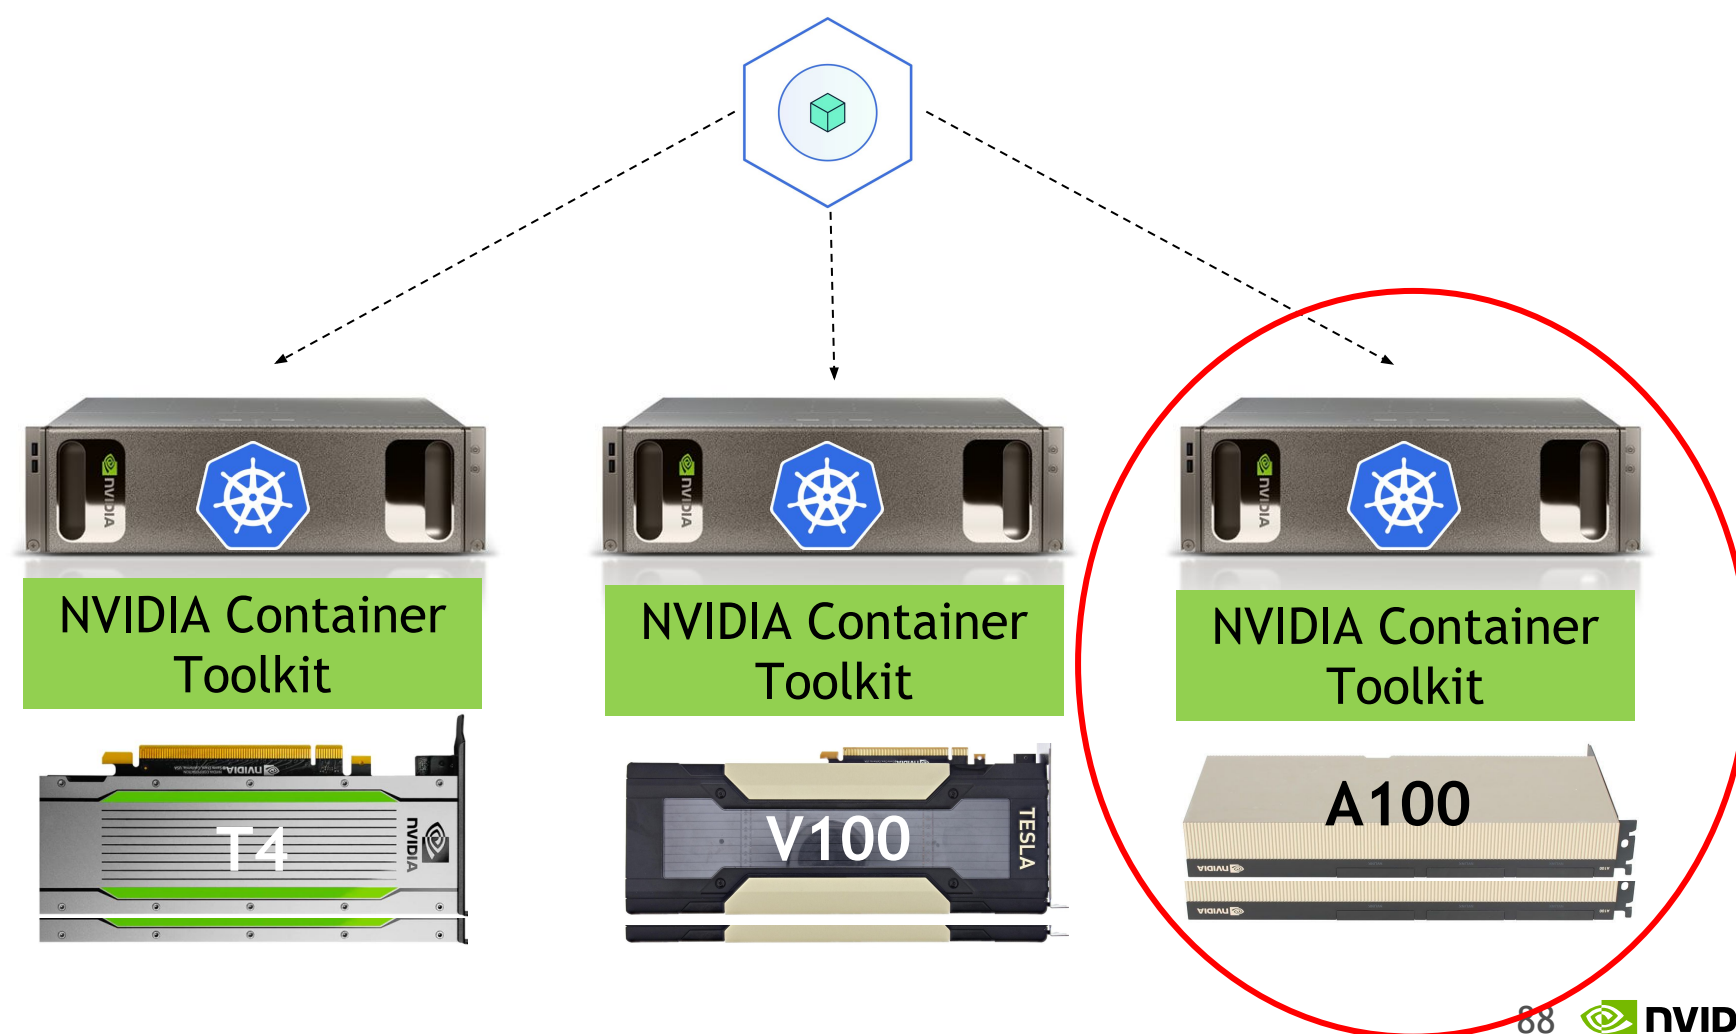


GPUs AND KUBERNETES

Allocate GPUs to pods in a Kubernetes Cluster

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/gpu: 4
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
    nvidia.com/cuda.runtime: 11.0
    nvidia.com/cuda.driver: 450.51.06
```

- **k8s-device-plugin**
- **gpu-feature-discovery**



MIG IN CONTAINERS AND KUBERNETES

Injecting a MIG Device Into a Container

MIG IN CONTAINERS AND KUBERNETES

Injecting a MIG Device Into a Container

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

```
GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)
```

```
GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)
```

MIG IN CONTAINERS AND KUBERNETES

Injecting a MIG Device Into a Container

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

```
GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)
```

```
GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)
```

```
$ docker run \  
  --gpus '"device=0:0,0:1,1:0,1:1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

```
GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)
```

```
  MIG 1g.5gb Device 0: (UUID: MIG-GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa/3/0)
```

```
  MIG 1g.5gb Device 1: (UUID: MIG-GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa/4/0)
```

```
GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)
```

```
  MIG 1g.5gb Device 0: (UUID: MIG-GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125/3/0)
```

```
  MIG 1g.5gb Device 1: (UUID: MIG-GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125/4/0)
```

MIG IN CONTAINERS AND KUBERNETES

Injecting a MIG Device Into a Container

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

```
$ docker run \  
  --gpus '"device=0:0,0:1,1:0,1:1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

MIG 1g.5gb Device 0: (UUID: MIG-GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa/3/0)

MIG 1g.5gb Device 1: (UUID: MIG-GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa/4/0)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

MIG 1g.5gb Device 0: (UUID: MIG-GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125/3/0)

MIG 1g.5gb Device 1: (UUID: MIG-GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125/4/0)

MIG IN CONTAINERS AND KUBERNETES

Injecting a MIG Device Into a Container

```
$ docker run \  
  --gpus '"device=0,1"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

```
$ docker run \  
  --gpus '"device=0:0,0:1,1:0,MIG-GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125/4/0"' \  
  nvidia/cuda:11.1-base nvidia-smi -L
```

GPU 0: A100-SXM4-40GB (UUID: GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa)

MIG 1g.5gb Device 0: (UUID: MIG-GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa/3/0)

MIG 1g.5gb Device 1: (UUID: MIG-GPU-238f350c-0ed9-09cf-9945-fc0649ef02aa/4/0)

GPU 1: A100-SXM4-40GB (UUID: GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125)

MIG 1g.5gb Device 0: (UUID: MIG-GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125/3/0)

MIG 1g.5gb Device 1: (UUID: MIG-GPU-dc36c15c-b7f1-cf33-f79a-bd01c0de7125/4/0)

MIG IN CONTAINERS AND KUBERNETES

Allocate MIG Devices to pods in a Kubernetes Cluster

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/gpu: 1
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
```

MIG IN CONTAINERS AND KUBERNETES

Allocate MIG Devices to pods in a Kubernetes Cluster

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/gpu: 1
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
```

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/mig-1g.5gb: 1
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
```

MIG IN CONTAINERS AND KUBERNETES

Allocate MIG Devices to pods in a Kubernetes Cluster

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/gpu: 1
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
```

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/mig-1g.5gb: 1
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
```

Mixed Strategy

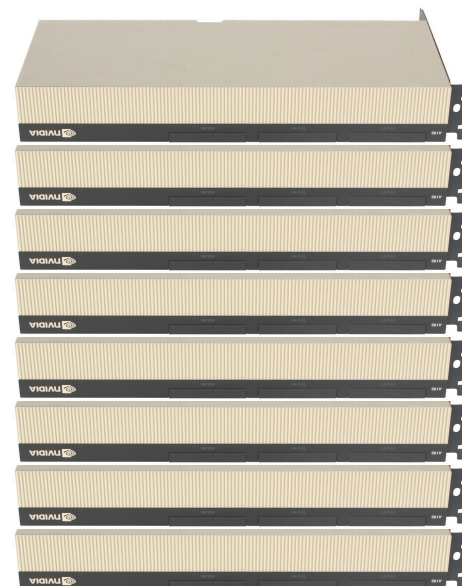
MIG IN CONTAINERS AND KUBERNETES

Allocate MIG Devices to pods in a Kubernetes Cluster

DGX-A100



8 x A100



4 x Full A100s

and

8 x 1g.5gb

and

4 x 2g.10gb

and

4 x 3g.20gb

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
  containers:
    - name: gpu-example
      image: nvidia/cuda
      resources:
        limits:
          nvidia.com/mig-1g.5gb: 1
          nvidia.com/mig-2g.10gb: 1
          nvidia.com/mig-3g.20gb: 1
          nvidia.com/gpu: 1
  nodeSelector:
    nvidia.com/gpu.product: A100-PCIE-40GB
```

Mixed Strategy

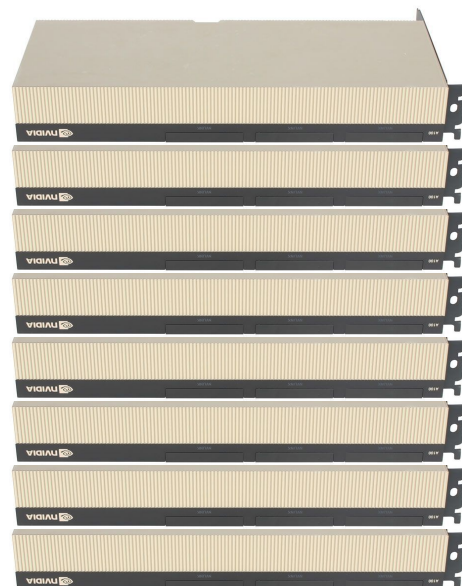
MIG IN CONTAINERS AND KUBERNETES

Allocate MIG Devices to pods in a Kubernetes Cluster

DGX-A100



8 x A100



56 x 1g.5gb
or
24 x 2g.10gb
or
16 x 3g.20gb

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-example
spec:
```

```
  containers:
```

```
    - name: gpu-example
      image: nvidia/cuda
      resources:
```

```
        limits:
```

```
          nvidia.com/gpu: 1
```

```
nodeSelector:
```

```
  nvidia.com/gpu.product: A100-PCIE-40GB MIG 1g.5gb
```

Single Strategy

MIG IN CONTAINERS AND KUBERNETES

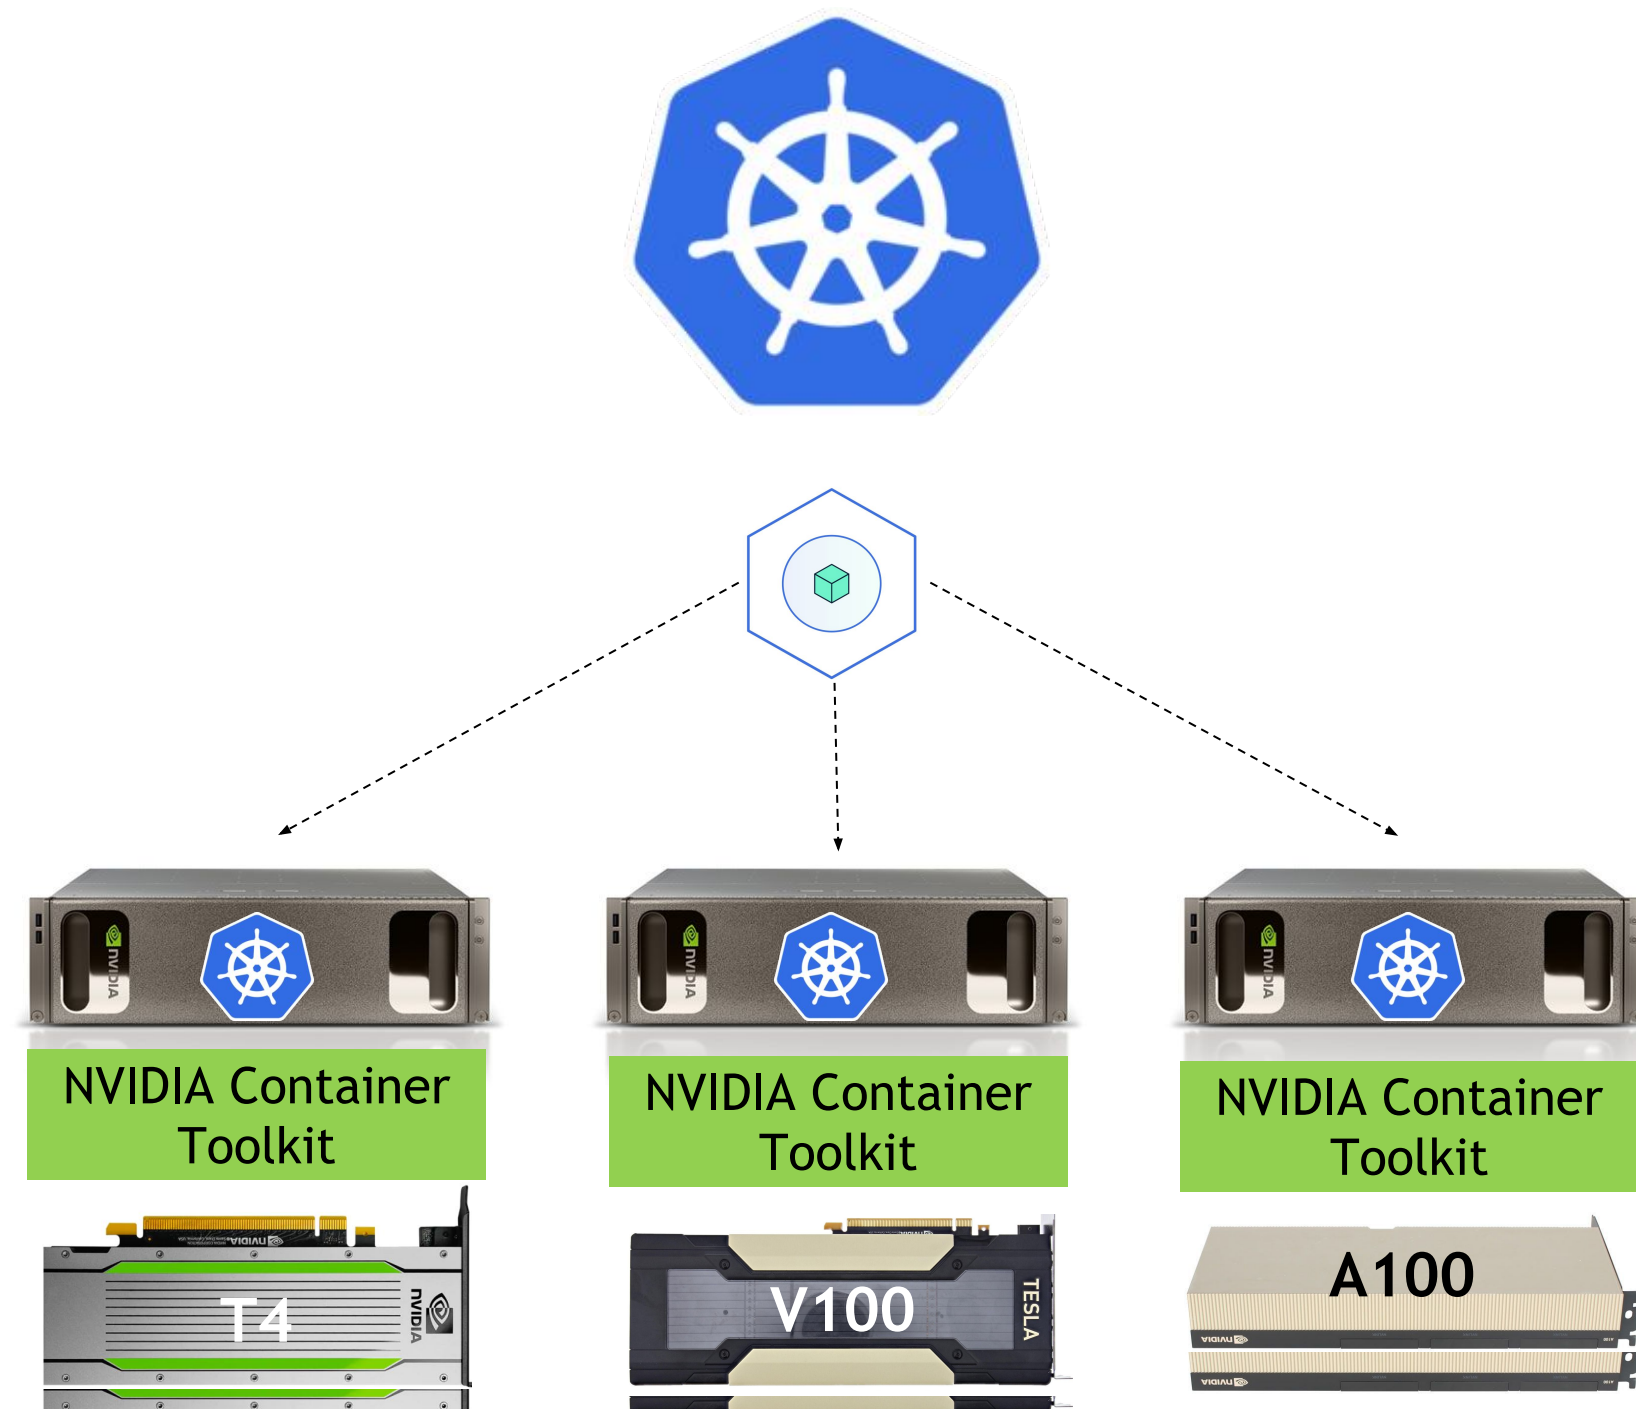
Allocate MIG Devices to pods in a Kubernetes Cluster

NVIDIA Container Toolkit

nvidia-docker	v2.5.0+
nvidia-container-runtime	v3.4.2+
nvidia-container-toolkit	v1.4.2+
libnvidia-container	v1.3.3+

Kubernetes

k8s-device-plugin	v0.8.2+
gpu-feature-discovery	v0.4.1+



SYSTEM LEVEL INTERFACE FOR MIG

Injecting GPUs and MIG devices into a container

SYSTEM LEVEL INTERFACE FOR MIG

Injecting GPUs and MIG devices into a container

For full GPUs:

Inject all of:

/dev

- |— nvidiactl
- |— nvidia-modeset
- |— nvidia-uvm
- |— nvidia-uvm-tools

Selectively inject for isolation:

/dev

- |— **nvidia0**
- |— nvidia1
- |— nvidia2

SYSTEM LEVEL INTERFACE FOR MIG

Injecting GPUs and MIG devices into a container

For full GPUs:

Inject all of:

```
/dev
├── nvidiactl
├── nvidia-modeset
├── nvidia-uvm
└── nvidia-uvm-tools
```

Selectively inject for isolation:

```
/dev
└── nvidia0
    ├── nvidia1
    └── nvidia2
```

Limit /proc view of:

```
/proc/driver/nvidia/gpus
├── 0000:07:00.0
├── 0000:0f:00.0
├── 0000:47:00.0
└── ...
```

Folders represent PCIeBusID of each full GPU whose device node is injected

SYSTEM LEVEL INTERFACE FOR MIG

Injecting GPUs and MIG devices into a container

For MIG devices:

Inject all of:

/dev

- |— nvidiactl
- |— nvidia-modeset
- |— nvidia-uvm
- |— nvidia-uvm-tools

Selectively inject for isolation:

/dev

- |— **nvidia0**
- |— nvidia-caps
- |— **nvidia-cap15**
- |— **nvidia-cap16**

} MIG Device
3-Tuple
<GPU, GI, CI>

SYSTEM LEVEL INTERFACE FOR MIG

Injecting GPUs and MIG devices into a container

For MIG devices:

Inject all of:

```
/dev
├── nvidiactl
├── nvidia-modeset
├── nvidia-uvm
└── nvidia-uvm-tools
```

Selectively inject for isolation:

```
/dev
├── nvidia0
└── nvidia-caps
    ├── nvidia-cap15
    └── nvidia-cap16
```

MIG Device
3-Tuple
<GPU, GI, CI>

SYSTEM LEVEL INTERFACE FOR MIG

nvidia-capabilities

Set of `nvidia-capabilities` in `/proc`


Point to `/dev/nvidia-caps` for access control

SYSTEM LEVEL INTERFACE FOR MIG

nvidia-capabilities

Set of **nvidia-capabilities** in **/proc**

```
/proc/driver
├── nvidia
│   └── capabilities
│       └── gpu0
│           └── mig
│               ├── gi0
│               │   ├── access
│               │   └── ci0
│               │       └── access
│               ├── gi1
│               │   ├── access
│               │   └── ci0
│               │       └── access
```




Point to **/dev/nvidia-caps** for access control

SYSTEM LEVEL INTERFACE FOR MIG

nvidia-capabilities

Set of **nvidia-capabilities** in **/proc**

```
/proc/driver
├── nvidia
│   └── capabilities
│       ├── gpu0
│       │   ├── mig
│       │   │   ├── gi0
│       │   │   │   ├── access
│       │   │   │   ├── ci0
│       │   │   │   └── access
│       │   │   └── gi1
│       │   │       ├── access
│       │   │       ├── ci0
│       │   │       └── access
```



Point to **/dev/nvidia-caps** for access control

```
$ nvidia-smi -L
GPU 0: A100-SXM4-40GB (UUID: GPU-...)
MIG 1g.5gb Device 0: (UUID: MIG-GPU-...)
```

SYSTEM LEVEL INTERFACE FOR MIG

nvidia-capabilities

Set of **nvidia-capabilities** in **/proc**

```
/proc/driver
├── nvidia
│   └── capabilities
│       ├── gpu0
│       │   ├── mig
│       │   │   ├── gi0
│       │   │   │   ├── access
│       │   │   │   └── ci0
│       │   │   │       └── access
│       │   │   └── gi1
│       │   │       ├── access
│       │   │       ├── ci0
│       │   │       └── access
```

Point to **/dev/nvidia-caps** for access control

```
$ nvidia-smi -L
GPU 0: A100-SXM4-40GB (UUID: GPU-...)
MIG 1g.5gb Device 0: (UUID: MIG-GPU-...)

$ cat /proc/.../gpu0/mig/gi0/access
DeviceFileMinor: 15
```

SYSTEM LEVEL INTERFACE FOR MIG

nvidia-capabilities

Set of nvidia-capabilities in /proc

```
/proc/driver
├── nvidia
│   └── capabilities
│       └── gpu0
│           └── mig
│               ├── gi0
│               │   ├── access
│               │   ├── ci0
│               │   └── access
│               └── gi1
│                   ├── access
│                   ├── ci0
│                   └── access
```

Point to /dev/nvidia-caps for access control

```
$ nvidia-smi -L
GPU 0: A100-SXM4-40GB (UUID: GPU-...)
MIG 1g.5gb Device 0: (UUID: MIG-GPU-...)

$ cat /proc/.../gpu0/mig/gi0/access
DeviceFileMinor: 15

$ cat /proc/.../gpu0/mig/gi0/ci0/access
DeviceFileMinor: 16
```


SYSTEM LEVEL INTERFACE FOR MIG

nvidia-capabilities

Set of nvidia-capabilities in /proc

```
/proc/driver
├── nvidia
│   └── capabilities
│       └── gpu0
│           └── mig
│               ├── gi0
│               │   ├── access
│               │   └── ci0
│               │       └── access
│               └── gi1
│                   ├── access
│                   └── ci0
│                       └── access
```

Point to /dev/nvidia-caps for access control

```
$ nvidia-smi -L
GPU 0: A100-SXM4-40GB (UUID: GPU-...)
MIG 1g.5gb Device 0: (UUID: MIG-GPU-...)

$ cat /proc/.../gpu0/mig/gi0/access
DeviceFileMinor: 15

$ cat /proc/.../gpu0/mig/gi0/ci0/access
DeviceFileMinor: 16

$ ls /dev/nvidia-caps
cr----- 1 root root 238, 15 nvidia-cap15
cr--r--r-- 1 root root 238, 16 nvidia-cap16
```

SYSTEM LEVEL INTERFACE FOR MIG

Limiting the view of `/proc/driver/nvidia/capabilities`

For MIG devices:

Additionally limit `/proc` view of:

`/proc/driver/nvidia/capabilities`

└─ `gpu0`

└─ `mig`

└─ `gi0`

└─ `access`

└─ `ci0`

└─ `access`

Access files point to `nvidia-caps` whose devices nodes are injected

CHALLENGES WITH MIG PARTITIONING

How do I create a MIG Device in the first place?

CHALLENGES WITH MIG PARTITIONING

How do I create a MIG Device in the first place?

Two distinct workflows when configuring a GPU for use with MIG

- Enabling MIG mode on a GPU in the first place
- Configuring MIG devices on a GPU that is already in MIG mode

```
# Enable MIG mode
nvidia-smi -mig 1

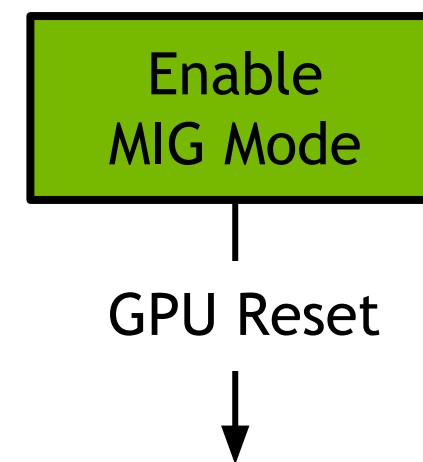
# Create 7 x 1g.5gb MIG devices
nvidia-smi mig -cgi 19,19,19,19,19,19,19 -C
```

CHALLENGES WITH MIG PARTITIONING

How do I create a MIG Device in the first place?

Enabling MIG mode on a GPU

- Complete all running GPU workloads
- Disconnect all driver clients from the GPU
- Enable MIG mode on the GPU
- Perform a GPU reset



Challenges:

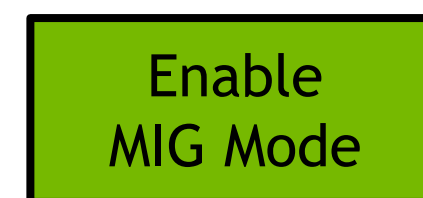
- Hard to enumerate all driver clients and reconnect them after the reset
- Requires a ***full node reboot*** under GPU pass-through virtualization

CHALLENGES WITH MIG PARTITIONING

How do I create a MIG Device in the first place?

Enabling MIG mode on a GPU

- Complete all running GPU workloads
- Disconnect all driver clients from the GPU
- Enable MIG mode on the GPU
- Perform a GPU reset



Very Heavy-weight operation

Should be done very infrequently

Challenges:

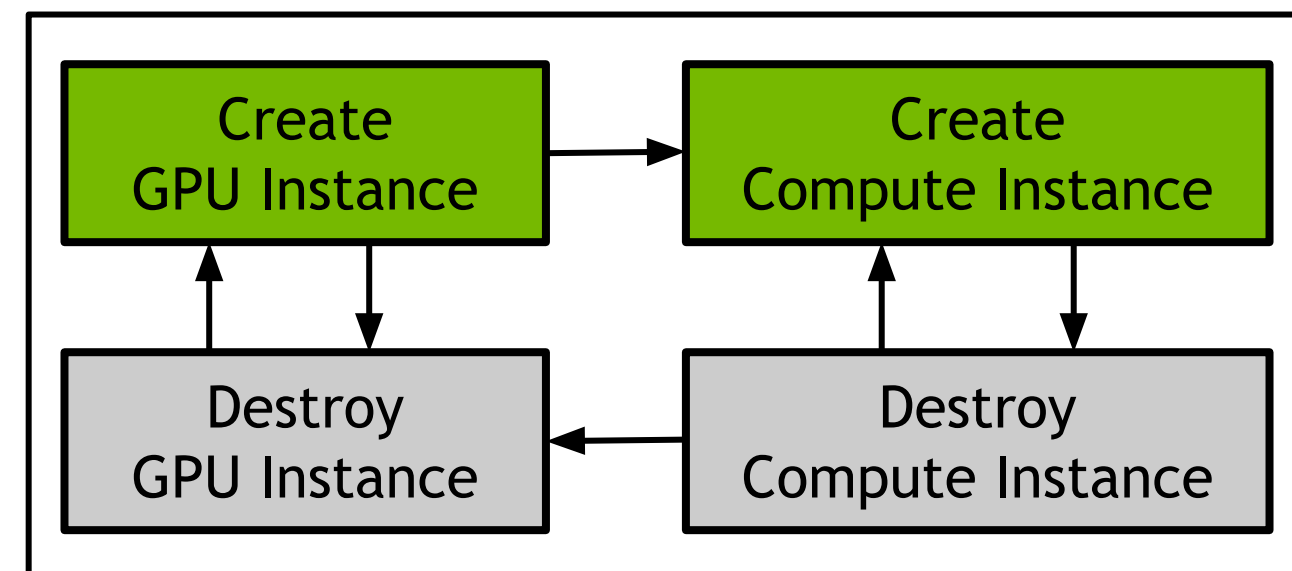
- Hard to enumerate all driver clients and reconnect them after the reset
- Requires a ***full node reboot*** under GPU pass-through virtualization

CHALLENGES WITH MIG PARTITIONING

How do I create a MIG Device in the first place?

Configuring MIG devices on a GPU

- Complete all running GPU workloads
- Perform device reconfiguration



Challenges:

- The order in which MIG devices are created matters
- Need to restart any components caching previous MIG device state
- Does not persist configuration across a GPU reset (or a node reboot)
- Hard to manage MIG device state across a cluster of machines

MIG PARTITION EDITOR

<https://github.com/NVIDIA/mig-parted>

MIG PARTITION EDITOR

<https://github.com/NVIDIA/mig-parted>

```
version: v1
mig-configs:
  all-1g.5gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "1g.5gb": 7

  all-2g.10gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "2g.10gb": 3

  all-3g.20gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "3g.20gb": 2
```

- Declaratively define the set of *possible* MIG partitions you want to create on GPUs throughout your cluster

MIG PARTITION EDITOR

<https://github.com/NVIDIA/mig-parted>

```
version: v1
mig-configs:
  all-1g.5gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "1g.5gb": 7

  all-2g.10gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "2g.10gb": 3

  all-3g.20gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "3g.20gb": 2
```

- Declaratively define the set of *possible* MIG partitions you want to create on GPUs throughout your cluster
- Make this file available to all of their nodes (as an actual file or a **configMap**)

MIG PARTITION EDITOR

<https://github.com/NVIDIA/mig-parted>

```
version: v1
mig-configs:
  all-1g.5gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "1g.5gb": 7
  all-2g.10gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "2g.10gb": 3
  all-3g.20gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "3g.20gb": 2
```

- Declaratively define the set of *possible* MIG partitions you want to create on GPUs throughout your cluster
- Make this file available to all of their nodes (as an actual file or a `configMap`)
- Use `nvidia-mig-parted` to apply one of these configurations on a given node

```
$ nvidia-mig-parted apply -f config.yaml -c all-2g.10gb
```

MIG PARTITION EDITOR

<https://github.com/NVIDIA/mig-parted>

```
version: v1
mig-configs:
  all-1g.5gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "1g.5gb": 7
  all-2g.10gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "2g.10gb": 3
  all-3g.20gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "3g.20gb": 2
```

systemd service wrapper

- Persists MIG configurations across node reboots
- Applies MIG mode changes *without* the NVIDIA driver loaded
- Automatically handles start/stop of GPU clients across configuration

<https://github.com/NVIDIA/mig-parted/tree/master/deployments/systemd>

MIG PARTITION EDITOR

<https://github.com/NVIDIA/mig-parted>

```
version: v1
mig-configs:
  all-1g.5gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "1g.5gb": 7
  all-2g.10gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "2g.10gb": 3
  all-3g.20gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "3g.20gb": 2
```

systemd service wrapper

- Persists MIG configurations across node reboots
- Applies MIG mode changes *without* the NVIDIA driver loaded
- Automatically handles start/stop of GPU clients across configuration

<https://github.com/NVIDIA/mig-parted/tree/master/deployments/systemd>

kubernetes service wrapper (coming soon)

- Provides similar functionality as the **systemd** service wrapper
- But integrated as a Kubernetes service instead
- Will become part of the [GPU Operator](#)

PUTTING IT ALL TOGETHER

Demo

1. Show full GPUs
2. Run **mig-parted**
3. Show mixed strategy
4. Show single strategy

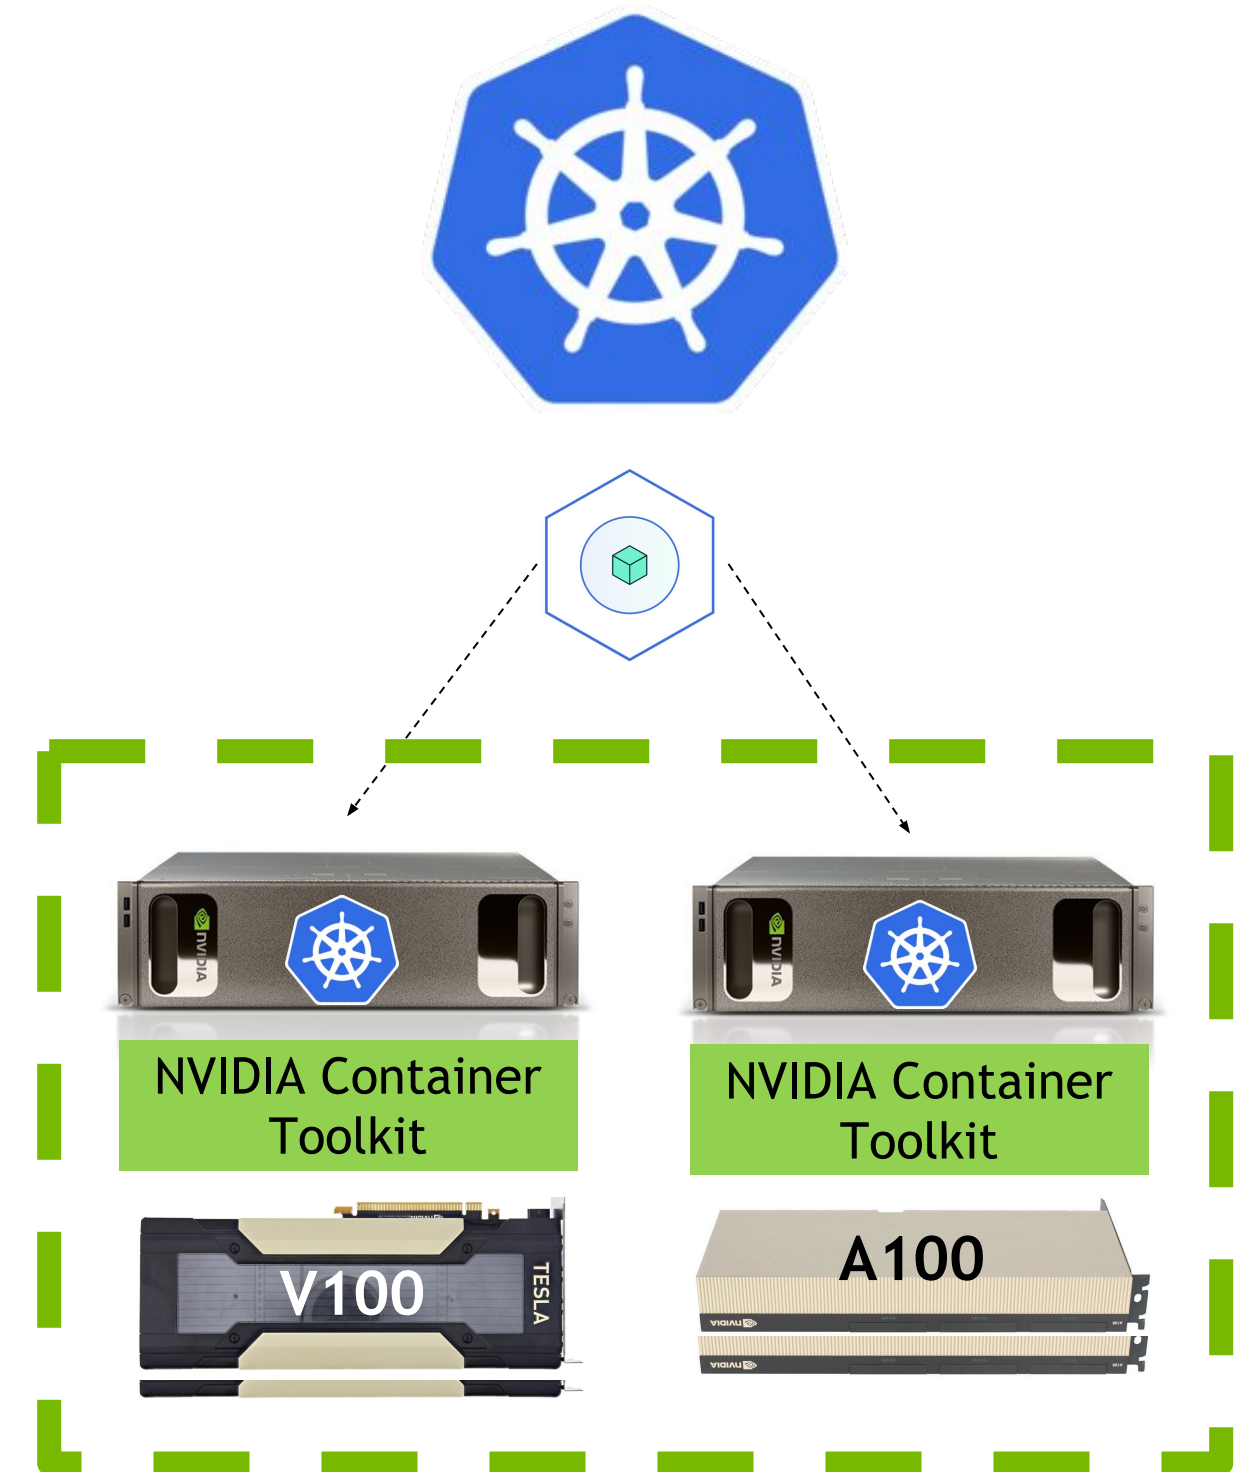
```
kubecon-demo:~# cat /etc/nvidia-mig-manager/config.yaml
version: v1
mig-configs:
  all-disabled:
    - devices: all
      mig-enabled: false
  all-1g.5gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "1g.5gb": 7
  all-2g.10gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "2g.10gb": 3
  all-3g.20gb:
    - devices: all
      mig-enabled: true
      mig-devices:
        "3g.20gb": 2
  all-balanced:
    - devices: all
      mig-enabled: true
      mig-devices:
        "1g.5gb": 2
        "2g.10gb": 1
        "3g.20gb": 1
kubecon-demo:~#
```

SUMMARY AND CONCLUSION

<https://docs.nvidia.com/datacenter/cloud-native/kubernetes/mig-k8s.html>

- MIG provides hardware support for sharing GPUs
- Support exists in both standalone containers and Kubernetes
- The `nvidia-mig-parted` tool simplifies partitioning MIG on a node

Integrated support on
EKS, GKE, and AKS coming very soon





KubeCon

CloudNativeCon

Europe 2021

Virtual



nvidia.