

ISOVALENT

# Better Bandwidth Management with eBPF



**Daniel Borkmann, Cilium Team & eBPF co-maintainer**

**Christopher M. Luciano, Cilium Team**

**Nikolay Aleksandrov, Cilium Team**

# Problem Statement


- Increasing Pod density per node
- Competition for node resources e.g. CPU and memory
- Optimization problem for operators: Resource allocation and efficient use, achieving SLOs, etc

```
apiVersion: v1
kind: Pod
metadata:
  name: frontend
spec:
  containers:
  - name: app
    image: images.my-company.example/app:v4
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
```


What the container is guaranteed to get, e.g. kubelet will only schedule the Pod on a node which can provide this resource.

Hard upper limit, ensures that container never goes above this threshold.

Pods per Node




Median Containers per Host



Source: Sysdig 2022 Cloud Native Security and Usage Report

# Problem Statement

- But what about networking?
- TCP sends AFAP (as fast as possible)
- AFAP output contract, shaping typically implemented by device output queues
- Queue length limit & receive window determines in-flight rate
- “How fast” implicit in queue drain rate
- Who limits a Pod’s network usage in Kubernetes?



Source: <https://hpbn.co/building-blocks-of-tcp/>

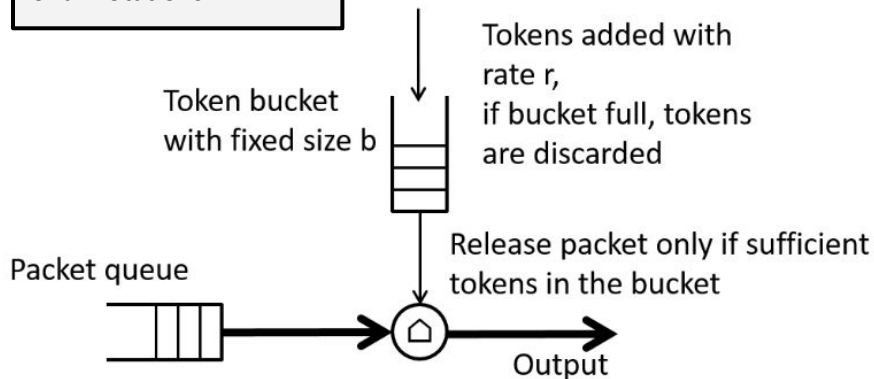


# Problem Statement

- Kubernetes bandwidth enforcement has only been experimental so far :- (
- Support for Pod annotations with 'outsourced' bandwidth meta plugin

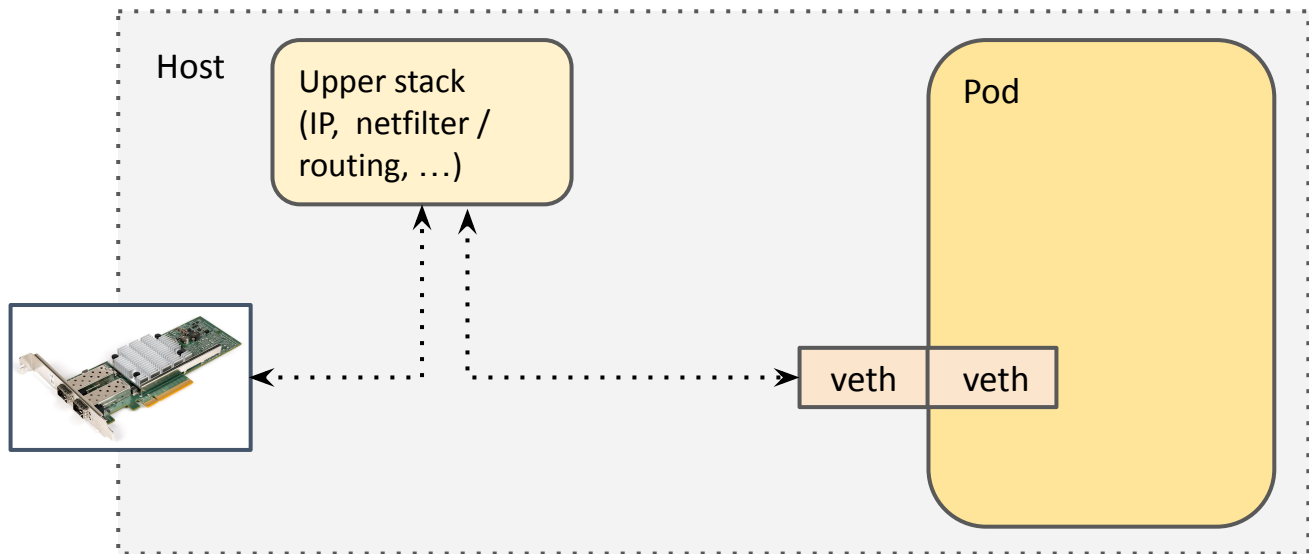
```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubernetes.io/ingress-bandwidth: 1M
    kubernetes.io/egress-bandwidth: 1M
...
```

Adds rudimentary **token bucket filters (TBF)** to implement enforcement of annotations.



# Problem Statement

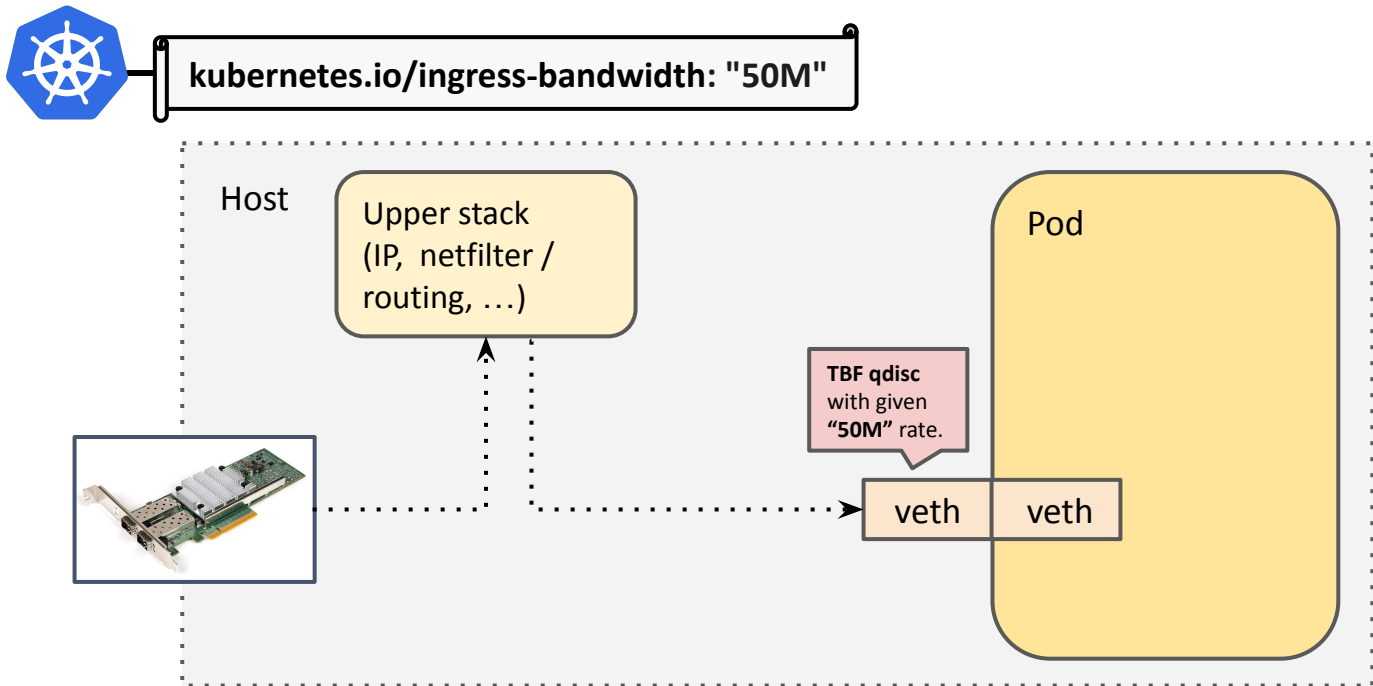
- bandwidth meta plugin not scalable for production use
- TBFs are attached to the Pod's veth devices



# Problem Statement



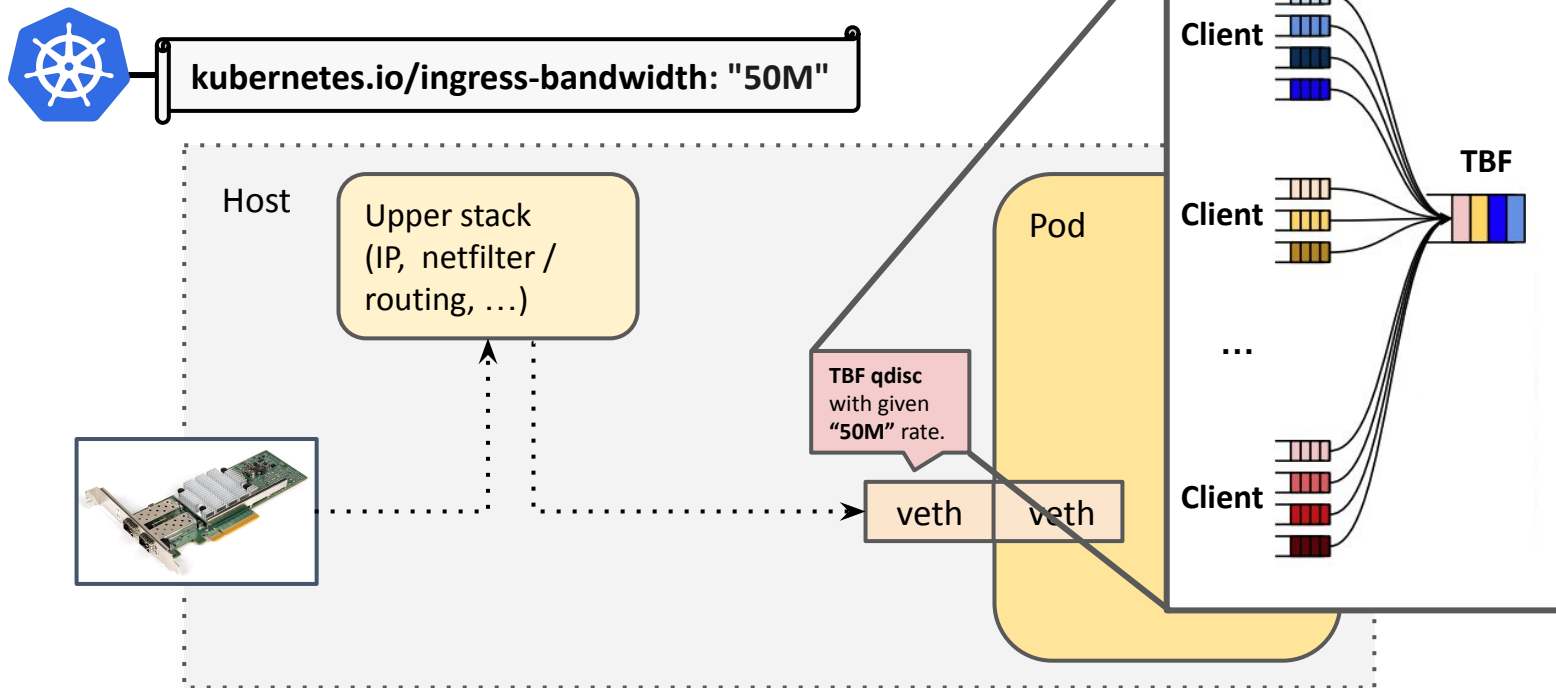
→ bandwidth meta plugin ingress example:



# Problem Statement



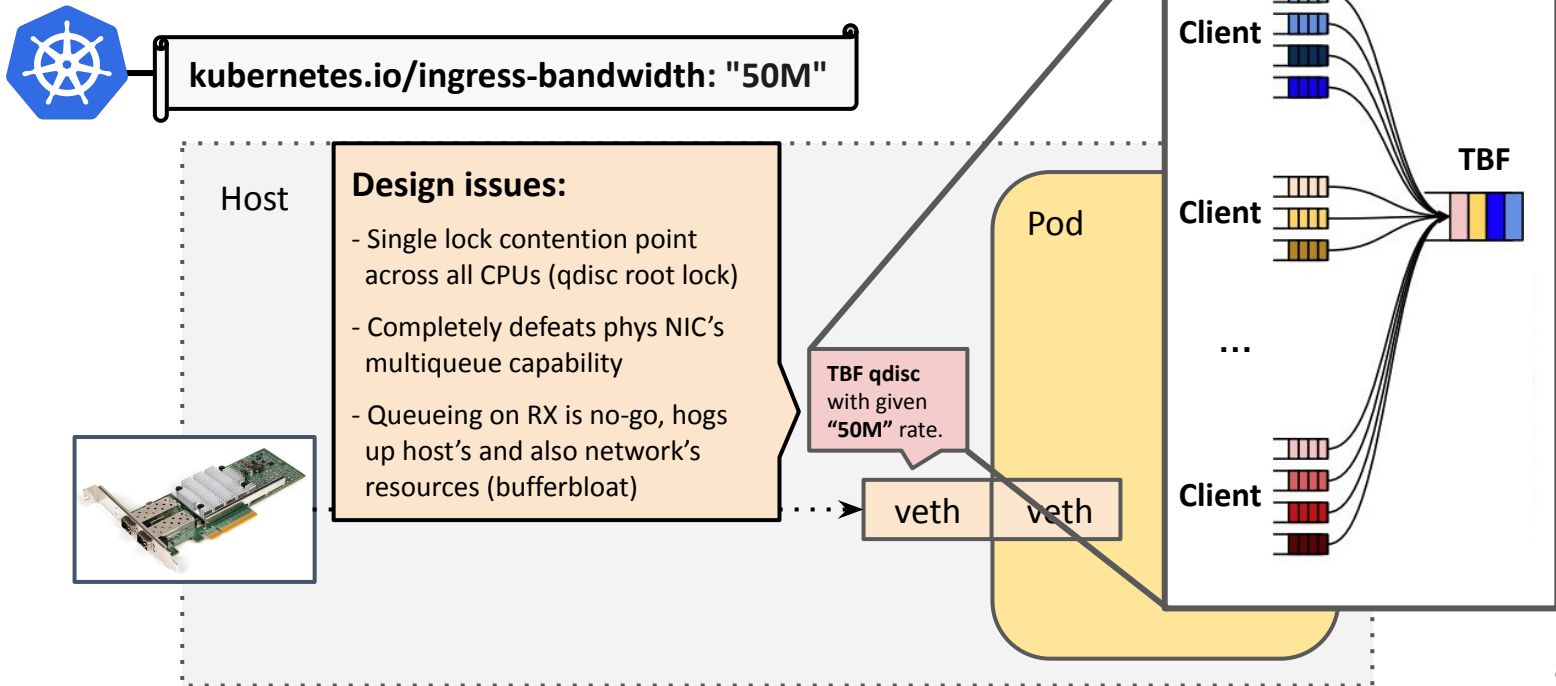
→ bandwidth meta plugin ingress example:



# Problem Statement



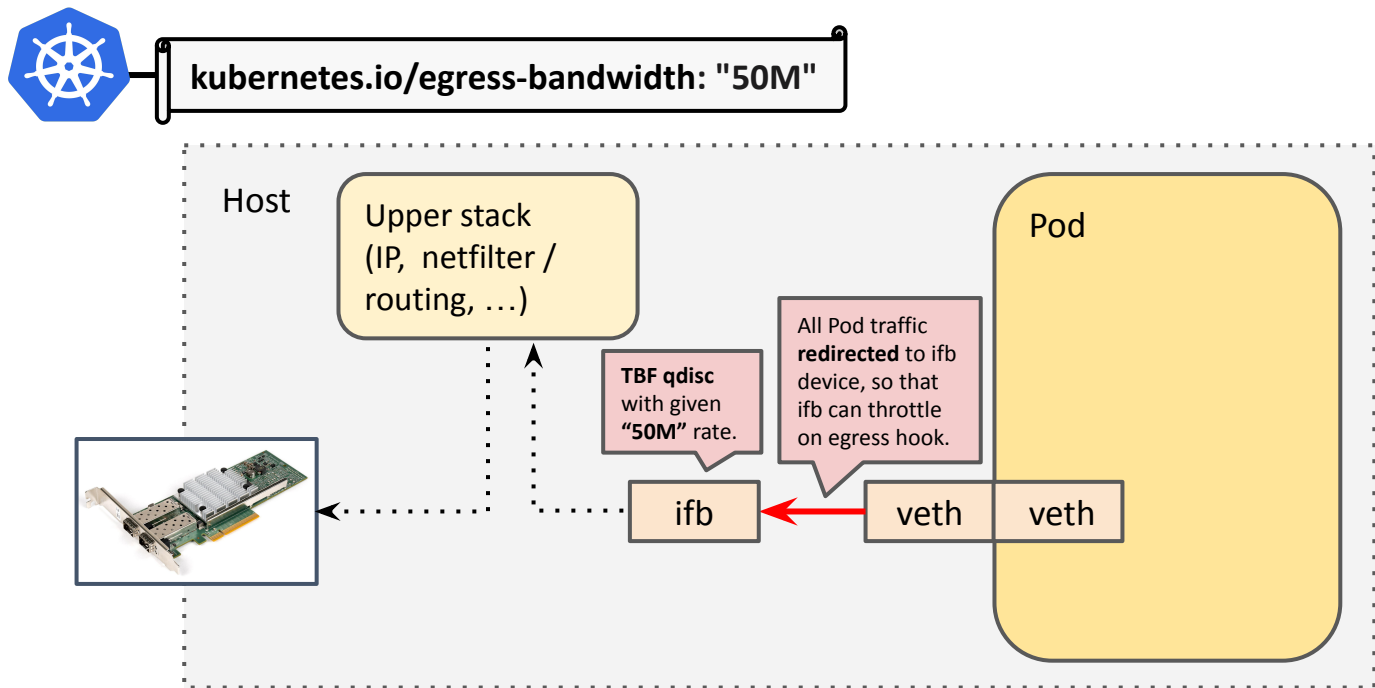
→ bandwidth meta plugin ingress example:





# Problem Statement

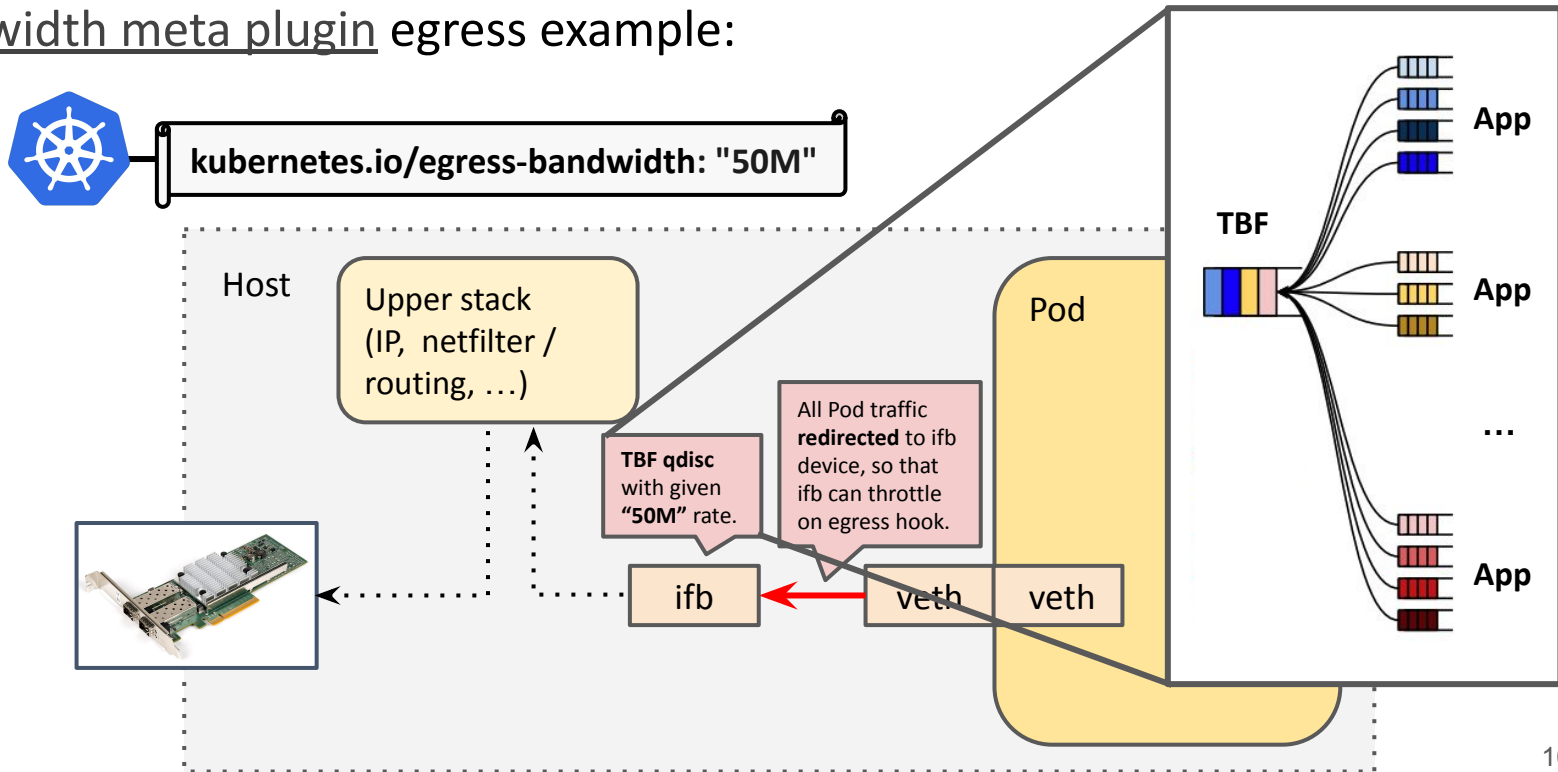
→ bandwidth meta plugin egress example:



# Problem Statement



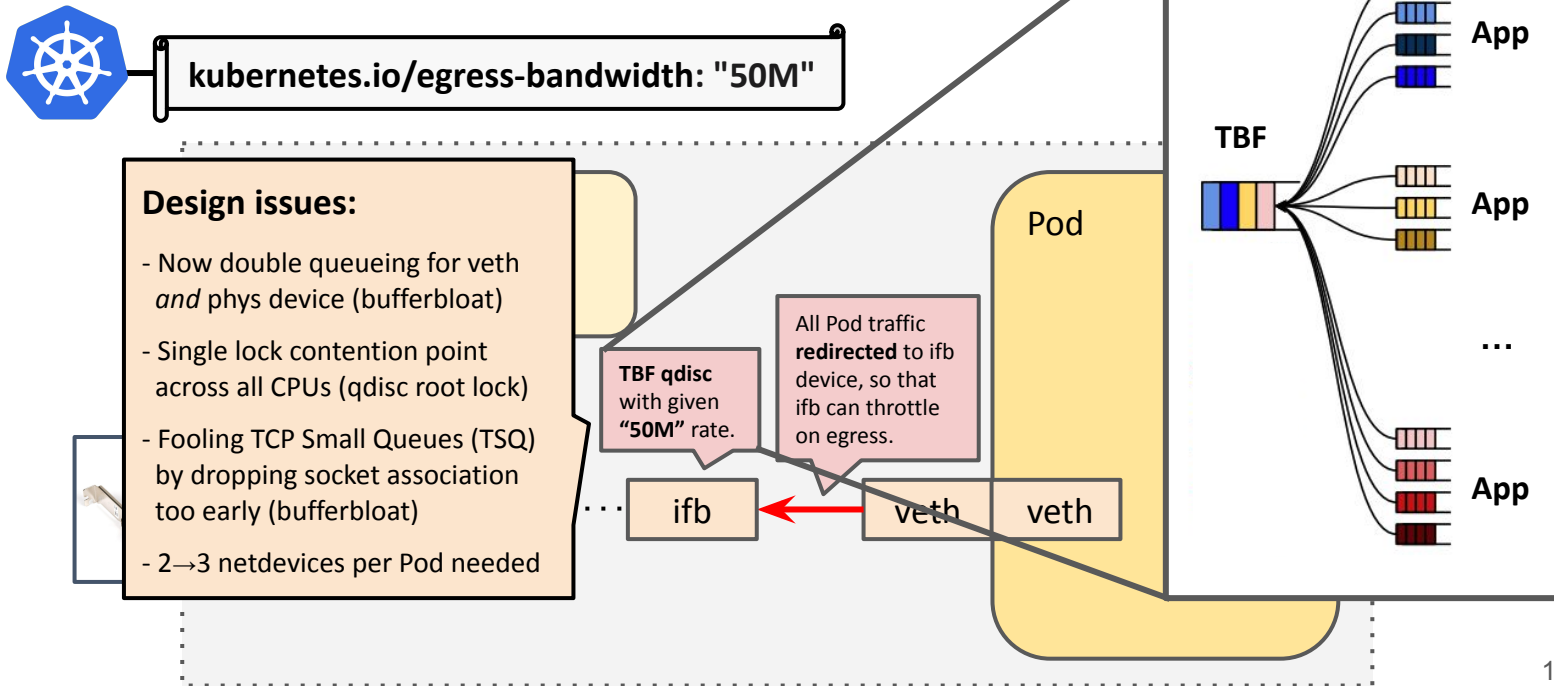
→ bandwidth meta plugin egress example:



# Problem Statement



→ bandwidth meta plugin egress example:



# Problem Statement



- bandwidth meta plugin aka “latency killer”
- tl;dr summary:



# Problem Statement

**... but AFAP makes bottleneck run at 100%**

Queuing theory says this is fragile. E.g., for M/D/1:


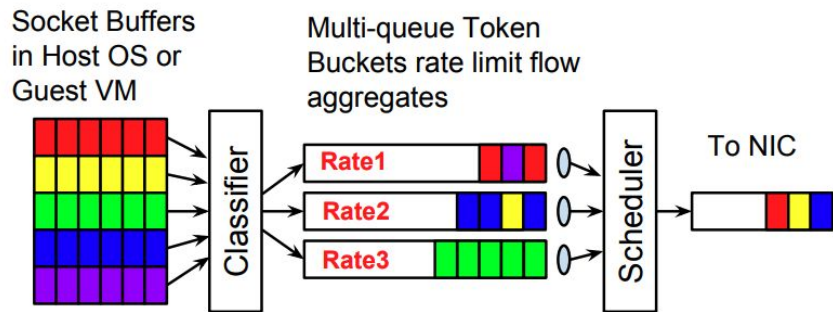


Packet wait-time in queue skyrockets when bottleneck link gets utilized close to 100%.

# From Queues to EDT model

Core Idea to replace queues with two simple pieces:

- Earliest Departure Time (EDT) time stamp in every packet
- Timing-wheel scheduler which replaces the queue



# How can the EDT model be applied to Kubernetes?



**Programmable** and **performant** in-kernel “virtual machine” that **safely** executes native code on certain events/hooks (aka “JavaScript for the kernel”).

# How can the EDT model be applied to Kubernetes?



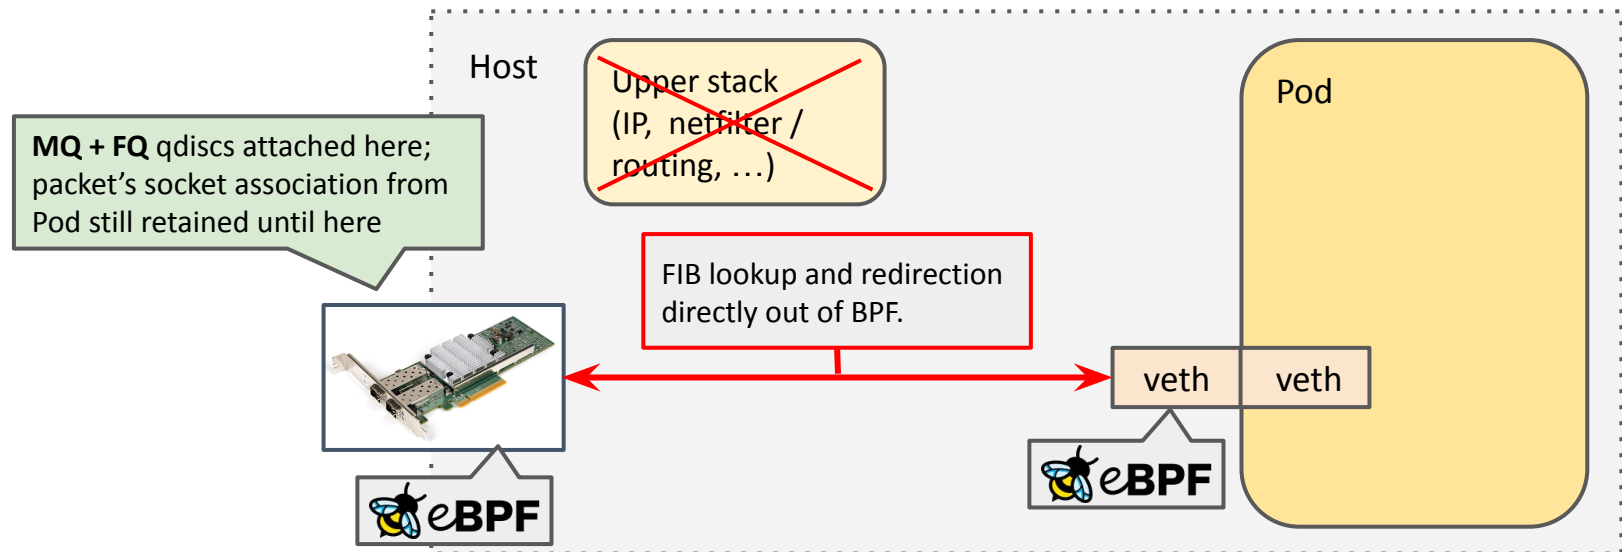
**eBPF-based** CNI / platform which provides Pod connectivity, service load-balancing, network policies, **bandwidth management**, transparent encryption and more.

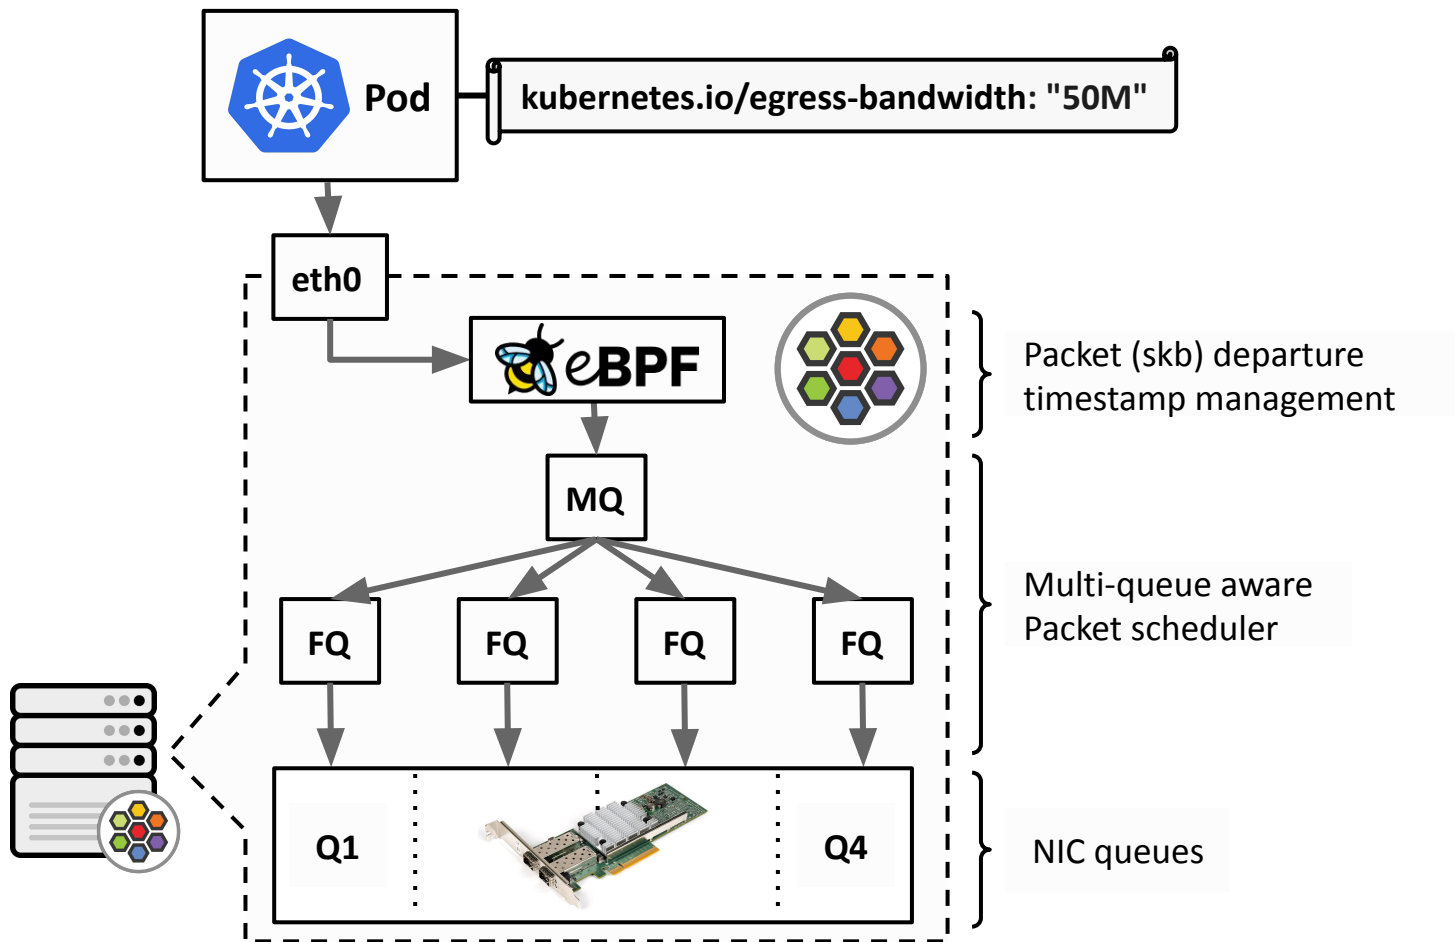


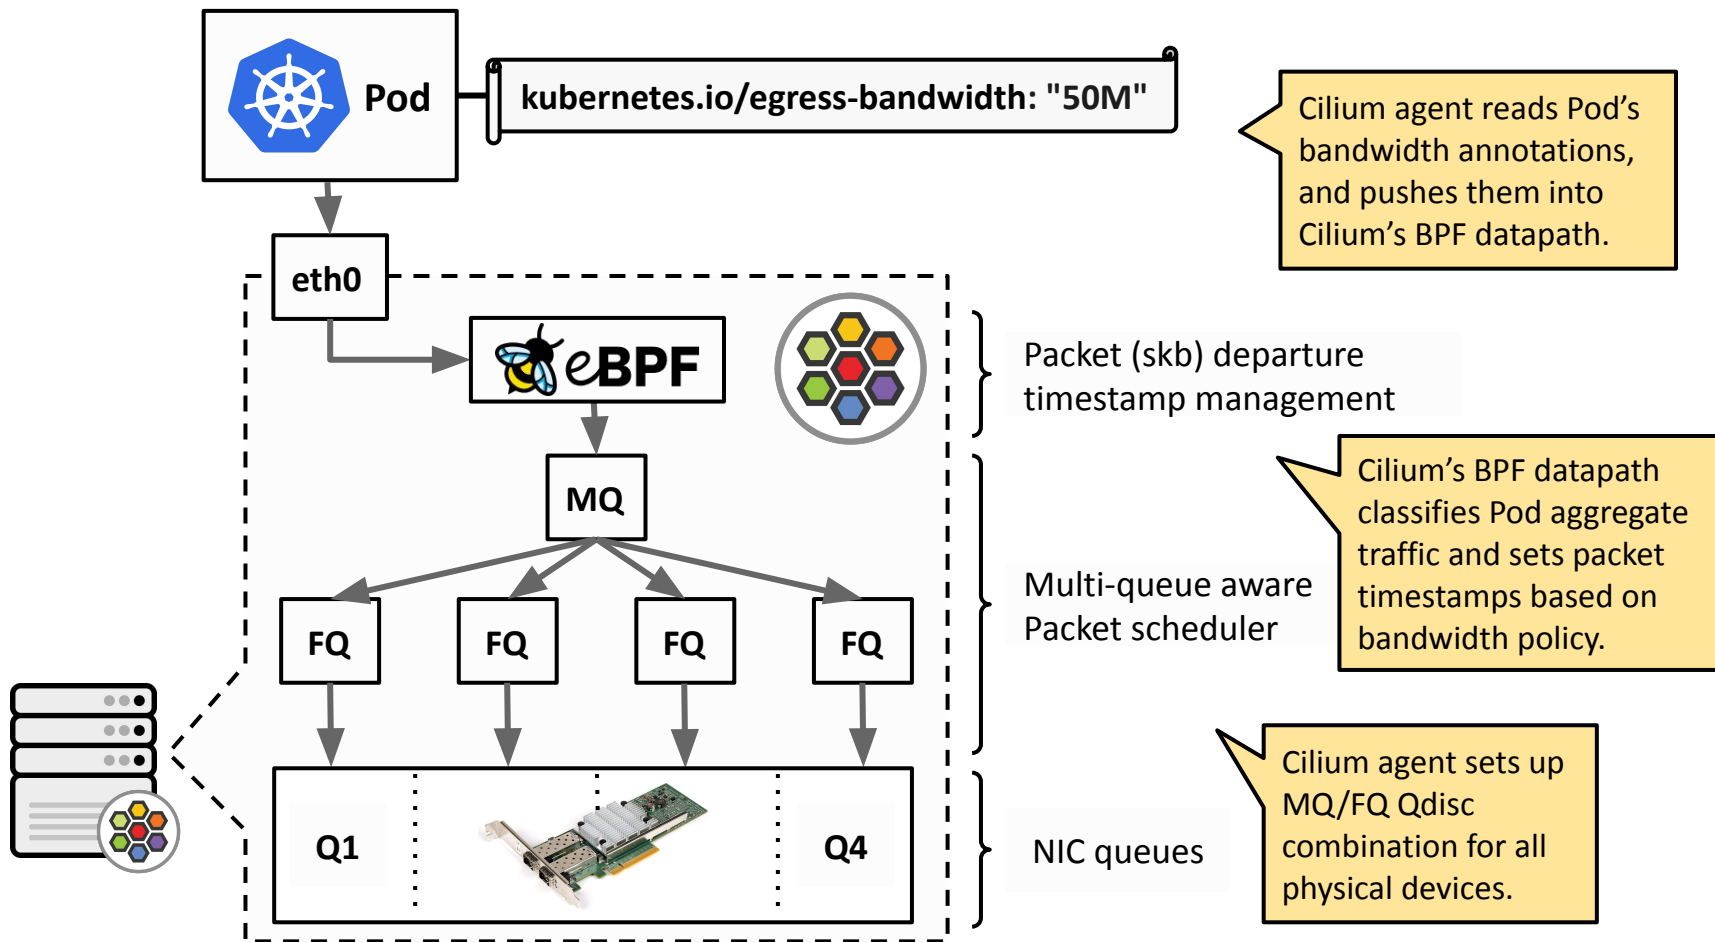


# Cilium's Bandwidth Manager

- Implements lock-less, EDT-based Pod rate-limiting with eBPF
- Enforcement points on phys devices instead of veths to avoid bufferbloat and improve TCP TSQ feedback








# Comparison of Cilium's EDT implementation vs TBF

Single flow latency for EDT and HTB/TBF model (lower is better)



Env: 256 concurrent request/response type flows (TCP\_RR), 100M rate per flow




# Comparison of Cilium's EDT implementation vs TBF

Total transaction rate between EDT and HTB/TBF model (higher is better)

■ EDT ■ TBF/HTB

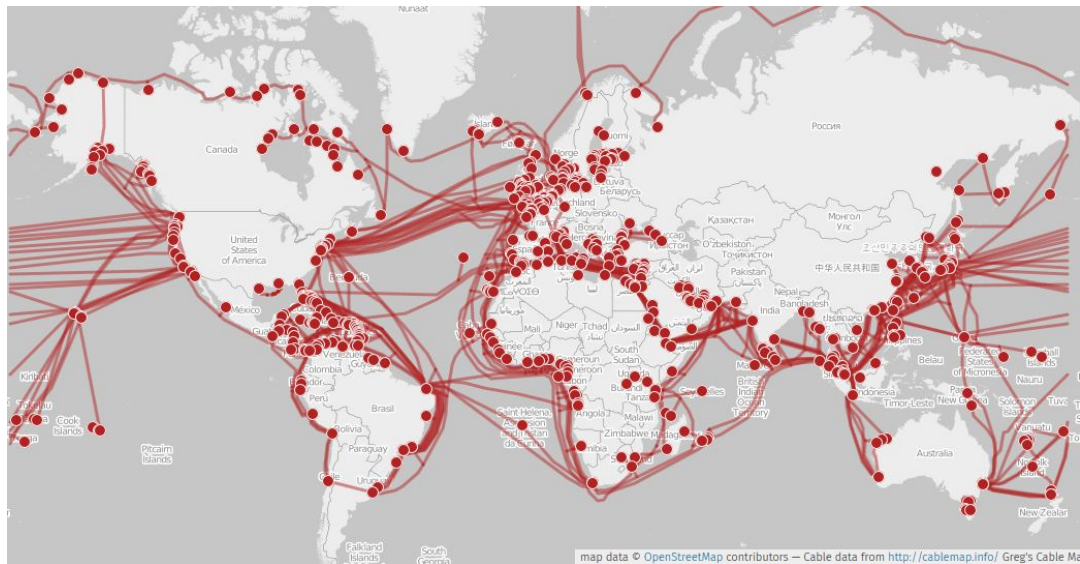
Env: 256 concurrent  
request/response  
type flows (TCP\_RR),  
100M rate per flow





# Cilium's Bandwidth Manager: recap for now


- So far: Cilium's EDT approach allows for scalable bandwidth enforcement
- What about more broadly Internet-level bandwidth management?



# What else does EDT model enable? Enter: BBR


## TCP before BBR

Today's Internet is not moving data as well as it should. TCP sends data at lower bandwidth because the 1980s-era algorithm assumes that packet loss means network congestion.



## TCP BBR

BBR models the network to send as fast as the available bandwidth and is 2700x faster than previous TCPs on a 10Gb, 100ms link with 1% loss. BBR powers google.com, youtube.com, and apps using Google Cloud Platform services.




Source: <https://cloud.google.com/blog/products/networking/tcp-bbr-congestion-control-comes-to-gcp-your-internet-just-got-faster>

# What else does EDT model enable? Enter: BBR

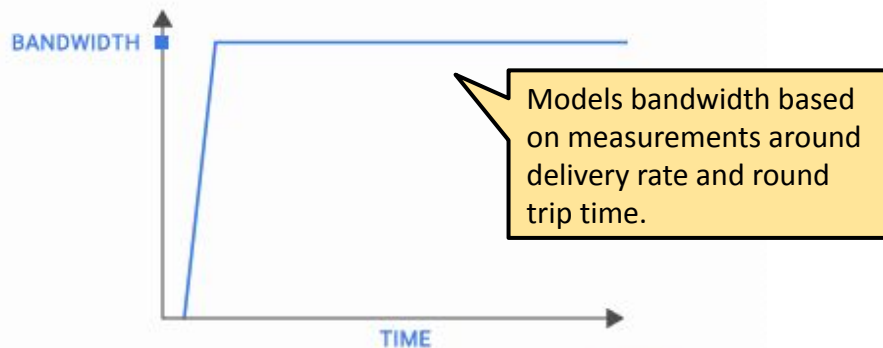
## TCP before BBR

Today's Internet is not moving data as well as it should. TCP sends data at lower bandwidth because the 1980s-era algorithm assumes that packet loss means network congestion.



## TCP BBR

BBR models the network to send as fast as the available bandwidth and is 2700x faster than previous TCPs on a 10Gb, 100ms link with 1% loss. BBR powers google.com, youtube.com, and apps using Google Cloud Platform services.







# When is it useful to consider BBR?

- Kubernetes cluster exposing services to clients over the Internet
- ◆ Significant latency improvements for low-end last-mile networks
  - ◆ Significant throughput improvements for high-speed long-haul links

	CUBIC (default)	BBR (v1)
Model parameters to the state machine	N/A	Throughput, RTT
Loss	Reduce cwnd by 30% on window with any loss	N/A
ECN	RFC3168 (classic ECN)	N/A
Startup	Slow-start until RTT rises (Hystart) or any loss	Slow-start until throughput plateaus

# Example: New York (packet.net) -> Zurich



```
darkstar@linux:~/trees/bpf-... x root@zh-lab-node-1: ~ x darkstar@linux:~/trees/bpf x root@ny-c3-small-x86-01: ~ x root@
root@zh-lab-node-1:~# while [ 1 ]; do iperf3 -c 147.75.66.15 -t 55 -R -i 5 -O 5 ; done
Connecting to host 147.75.66.15, port 5201
Reverse mode, remote host 147.75.66.15 is sending
[ 5] local 192.168.178.91 port 52148 connected to 147.75.66.15 port 5201
[ ID] Interval            Transfer           Bitrate
[ 5]  0.00-5.00      sec  78.4 MBytes      132 Mbits/sec
[ 5]  0.00-5.00      sec  98.2 MBytes      165 Mbits/sec
[ 5]  5.00-10.00     sec  99.1 MBytes      166 Mbits/sec
[ 5] 10.00-15.00     sec  113 MBytes      189 Mbits/sec
[ 5] 15.00-20.00     sec  159 MBytes      267 Mbits/sec
[ 5] 20.00-25.00     sec  257 MBytes      431 Mbits/sec
[ 5] 25.00-30.00     sec  153 MBytes      256 Mbits/sec
[ 5] 30.00-35.00     sec  146 MBytes      245 Mbits/sec
[ 5] 35.00-40.00     sec  148 MBytes      248 Mbits/sec
[ 5] 40.00-45.00     sec  157 MBytes      264 Mbits/sec
[ 5] 45.00-50.00     sec  193 MBytes      323 Mbits/sec
[ 5] 50.00-55.00     sec  272 MBytes      457 Mbits/sec
- - - - -
[ ID] Interval            Transfer           Bitrate      Retr
[ 5]  0.00-55.10     sec  1.76 GBytes      274 Mbits/sec  1501
[ 5]  0.00-55.00     sec  1.75 GBytes      274 Mbits/sec
```

(omitted)

Bandwidth probing, overreaction to loss!  
(Sawtooth pattern nicely visible)

**Default, server runs:**

- TCP CUBIC
- fq\_codel Qdisc

sender  
receiver

# Example: New York (packet.net) -> Zurich



```
darkstar@linux:~/trees/bpf-... x root@zh-lab-node-1: ~ x darkstar@linux:~/trees/bpf x root@ny-c3-small-x86-01: ~ x root@
root@zh-lab-node-1:~# while [ 1 ]; do iperf3 -c 147.75.66.15 -t 55 -R -i 5 -O 5 ; done
Connecting to host 147.75.66.15, port 5201
Reverse mode, remote host 147.75.66.15 is sending
[ 5] local 192.168.178.91 port 52254 connected to 147.75.66.15 port 5201
[ ID] Interval          Transfer    Bitrate
[ 5]  0.00-5.00      sec    152 MBytes  254 Mbits/sec
[ 5]  0.00-5.00      sec    258 MBytes  433 Mbits/sec
[ 5]  5.00-10.00     sec    240 MBytes  403 Mbits/sec
[ 5] 10.00-15.00     sec    255 MBytes  427 Mbits/sec
[ 5] 15.00-20.00     sec    247 MBytes  414 Mbits/sec
[ 5] 20.00-25.00     sec    255 MBytes  428 Mbits/sec
[ 5] 25.00-30.00     sec    255 MBytes  428 Mbits/sec
[ 5] 30.00-35.00     sec    238 MBytes  400 Mbits/sec
[ 5] 35.00-40.00     sec    255 MBytes  428 Mbits/sec
[ 5] 40.00-45.00     sec    239 MBytes  401 Mbits/sec
[ 5] 45.00-50.00     sec    253 MBytes  425 Mbits/sec
[ 5] 50.00-55.00     sec    242 MBytes  407 Mbits/sec
-- -- -- -- --
[ ID] Interval          Transfer    Bitrate    Retr
[ 5]  0.00-55.10     sec    2.68 GBytes  418 Mbits/sec  58812
[ 5]  0.00-55.00     sec    2.67 GBytes  418 Mbits/sec
```

No overreaction, stable.

**Updated, server runs:**

- TCP BBR
- FQ Qdisc (for EDT)

(results also reproduce with netperf)

sender  
receiver

# Example: New York (packet.net) -> Zurich



```
darkstar@linux:~/trees/bpf-... x root@zh-lab-node-1: ~ x darkstar@linux:~/trees/bpf x root@ny-c3-small-x86-01: ~ x root@
```

```
root@zh-lab-node-1:~# while [ 1 ]; do iperf3 -c 147.75.66.15 -t 55 -R -i 5 -O 5 ; done
```

```
Connecting to host 147.75.66.15, port 5201
```

```
Reverse mode, remote host 147.75.66.15 is sending
```

```
[ 5] local 192.168.178.91 port 52254 connected to 147.75.66.15 port 5201
```

[ ID]	Interval		Transfer	Bitrate	
[ 5]	0.00-5.00	sec	152 MBytes	254 Mbits/sec	(omitted)
[ 5]	0.00-5.00	sec	258 MBytes	433 Mbits/sec	
[ 5]	5.00-10.00	sec	240 MBytes	403 Mbits/sec	
[ 5]	10.00-15.00	sec	255 MBytes	427 Mbits/sec	
[ 5]	15.00-20.00	sec	247 MBytes	414 Mbits/sec	
[ 5]	20.00-25.00	sec	255 MBytes	428 Mbits/sec	
[ 5]	25.00-30.00	sec	255 MBytes	428 Mbits/sec	
[ 5]	30.00-35.00	sec	238 MBytes	400 Mbits/sec	
[ 5]	35.00-40.00	sec	255 MBytes	428 Mbits/sec	
[ 5]	40.00-45.00	sec	239 MBytes	401 Mbits/sec	
[ 5]	45.00-50.00	sec	253 MBytes	425 Mbits/sec	
[ 5]	50.00-55.00	sec	242 MBytes	407 Mbits/sec	

[ ID]	Interval		Transfer	Bitrate	Retr
[ 5]	0.00-55.10	sec	2.68 GBytes	418 Mbits/sec	58812
[ 5]	0.00-55.00	sec	2.67 GBytes	418 Mbits/sec	

No overreaction, stable.

**Updated**, server runs:

- TCP BBR
- FQ Qdisc (for EDT)

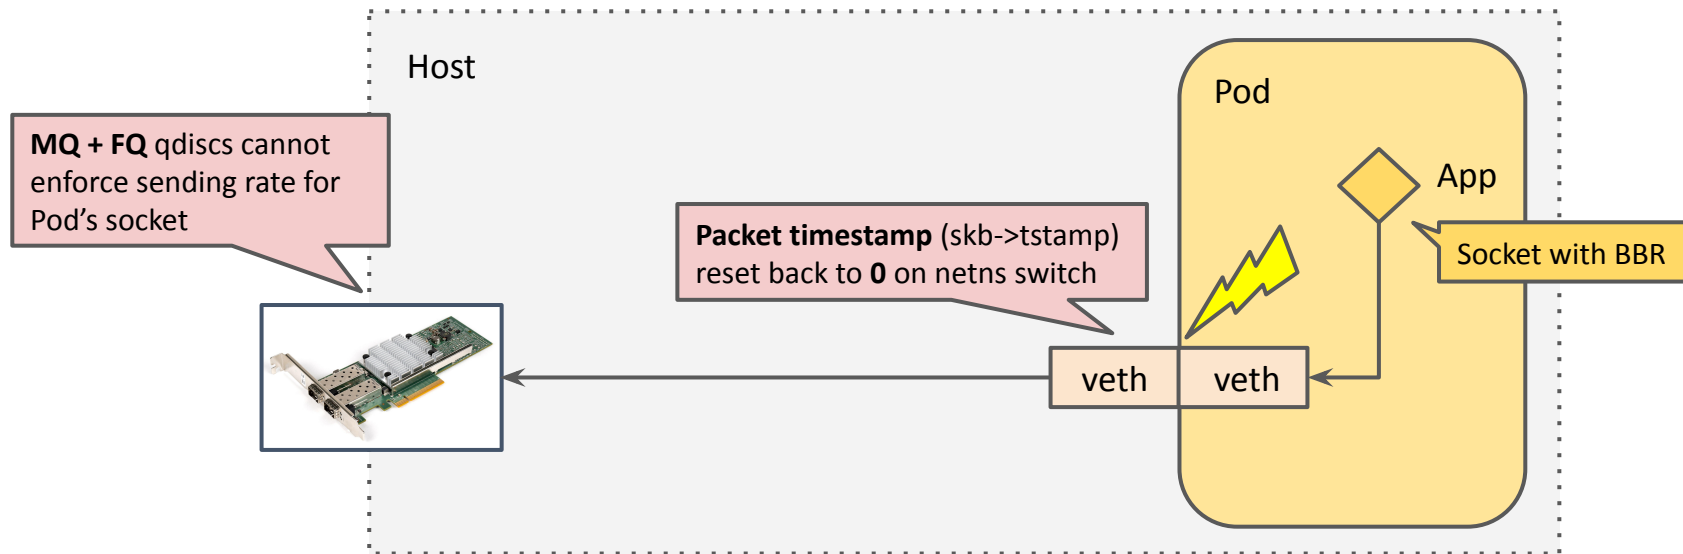
**CUBIC** → **BBR** bumps:

**274** → **418 Mbit/s**




# Can BBR be used for Kubernetes Pods?

- BBR works in conjunction with FQ and sets packet delivery timestamps
- Kernel clears timestamp for packets leaving Pods (== netns)
  - ◆ Usage of BBR for Pods not possible/broken in general today



## Rationale on today's timestamp reset ([see our LPC talk](#))



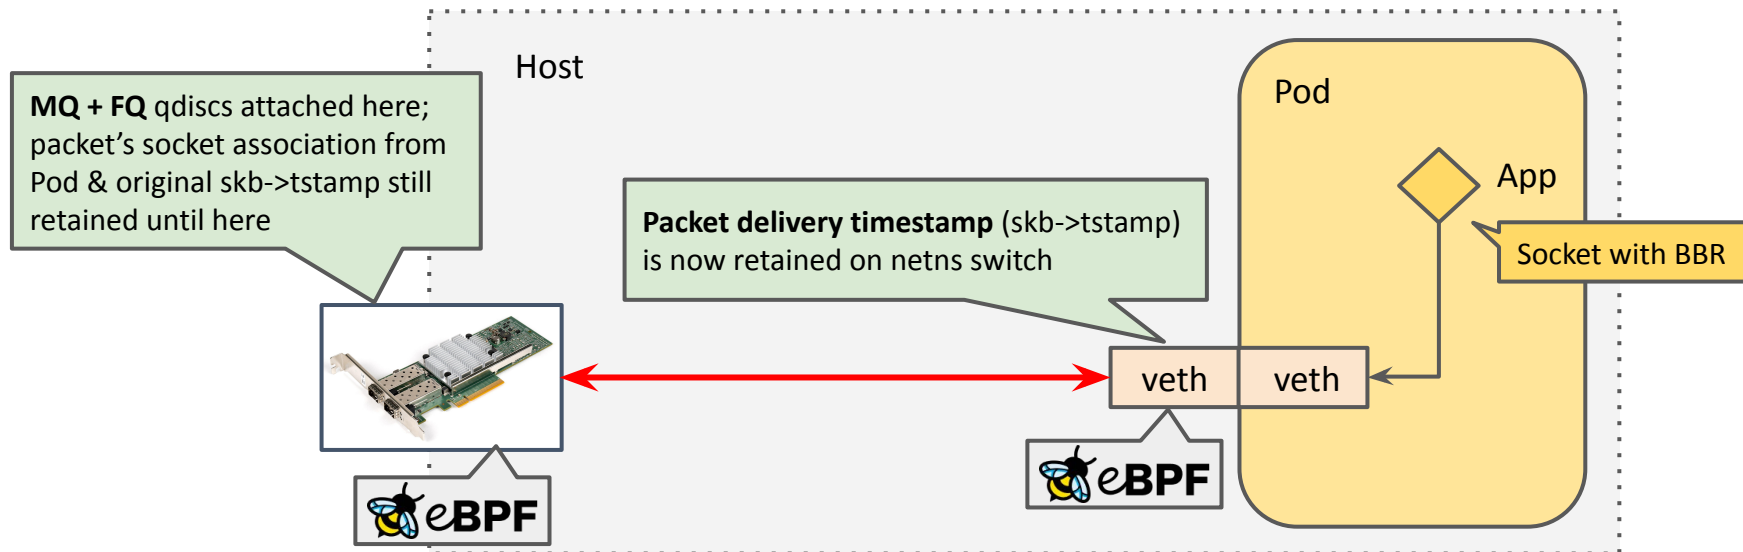
Kernel uses different clock bases for `skb->timestamp`:

- Ingress is `CLOCK_TAI`, egress is `CLOCK_MONOTONIC` (as is FQ)
- Forwarding from RX to TX would cause drop in FQ due to overreaching FQ's drop horizon given different clock's offsets
- No means to figure out clock base from `skb->timestamp`, hence reset



# Cilium's Bandwidth Manager: BBR

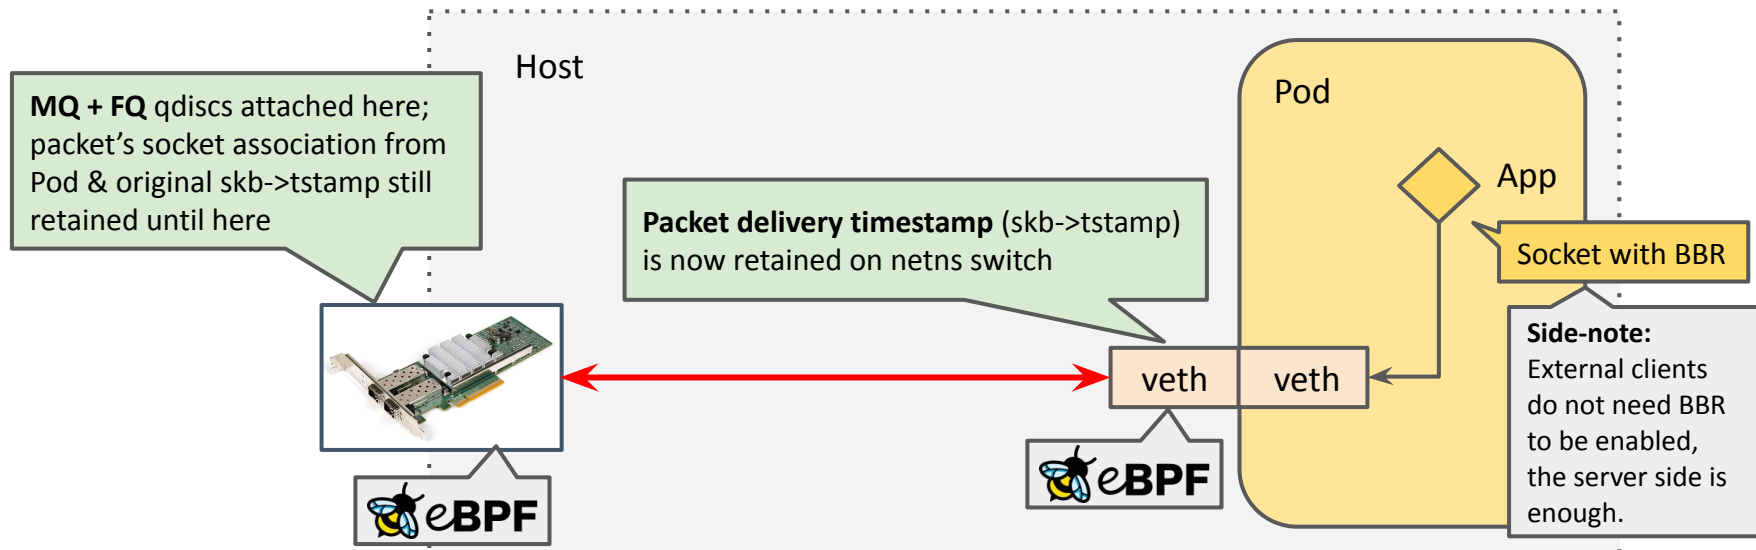
- We helped fixing networking stack in Linux v5.18+ to retain timestamps
- Bandwidth Manager plumbs the appropriate underlying infrastructure
  - ◆ Receives new knob for switching whole cluster over to BBR by default





# Cilium's Bandwidth Manager: BBR

- We helped fixing networking stack in Linux v5.18+ to retain timestamps
- Bandwidth Manager plumbs the appropriate underlying infrastructure
  - ◆ Receives new knob for switching whole cluster over to BBR by default






ISOVALENT

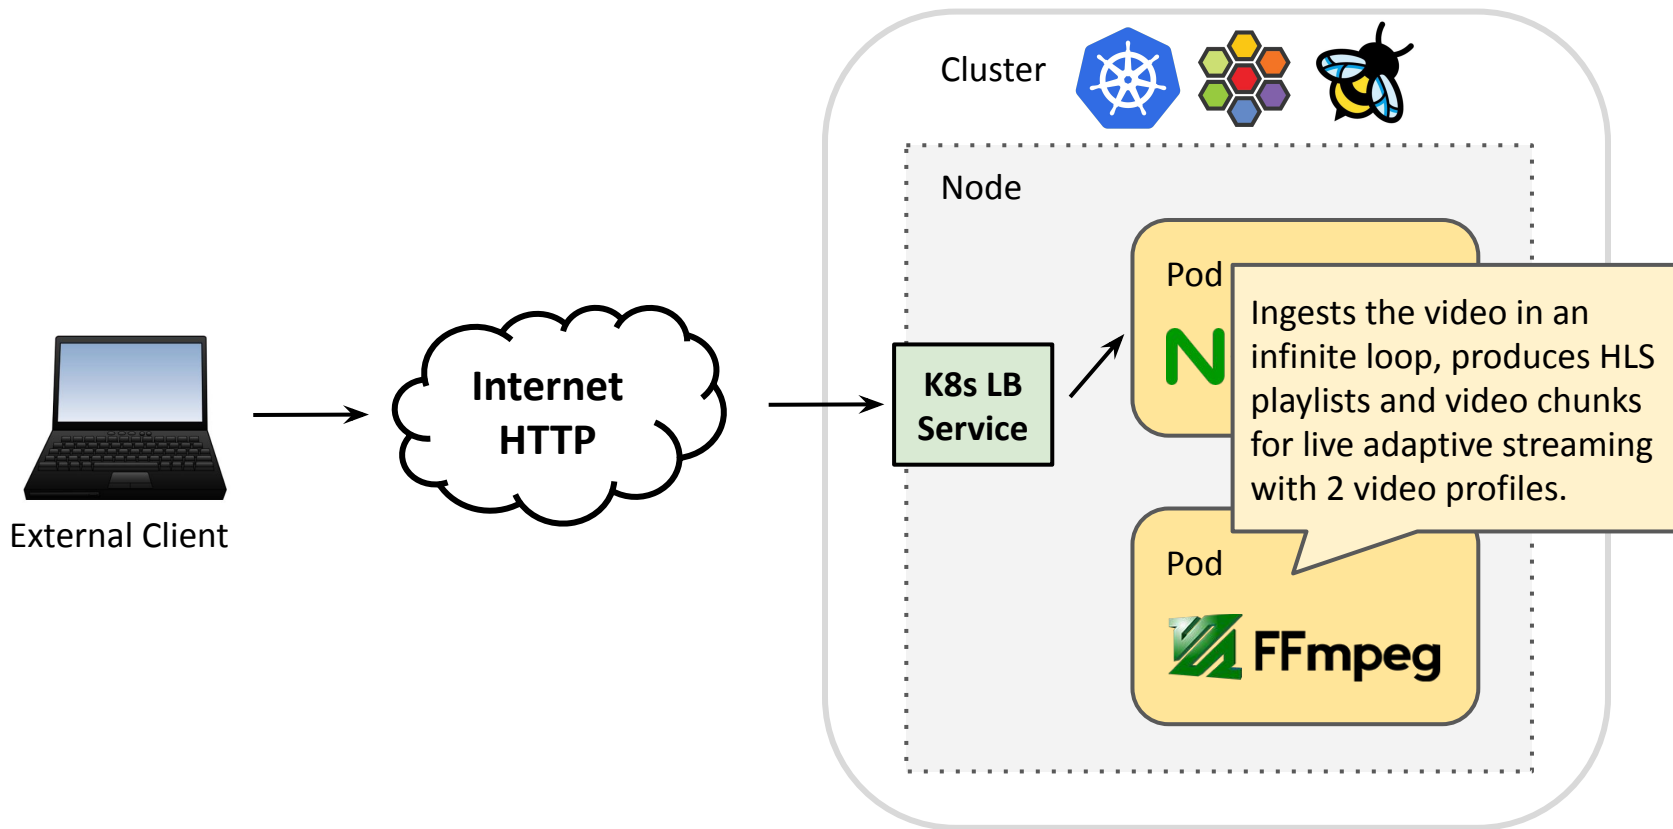
# Demo: BBR for Pods

(K8s/Cilium-backed video streaming service: CUBIC versus BBR)

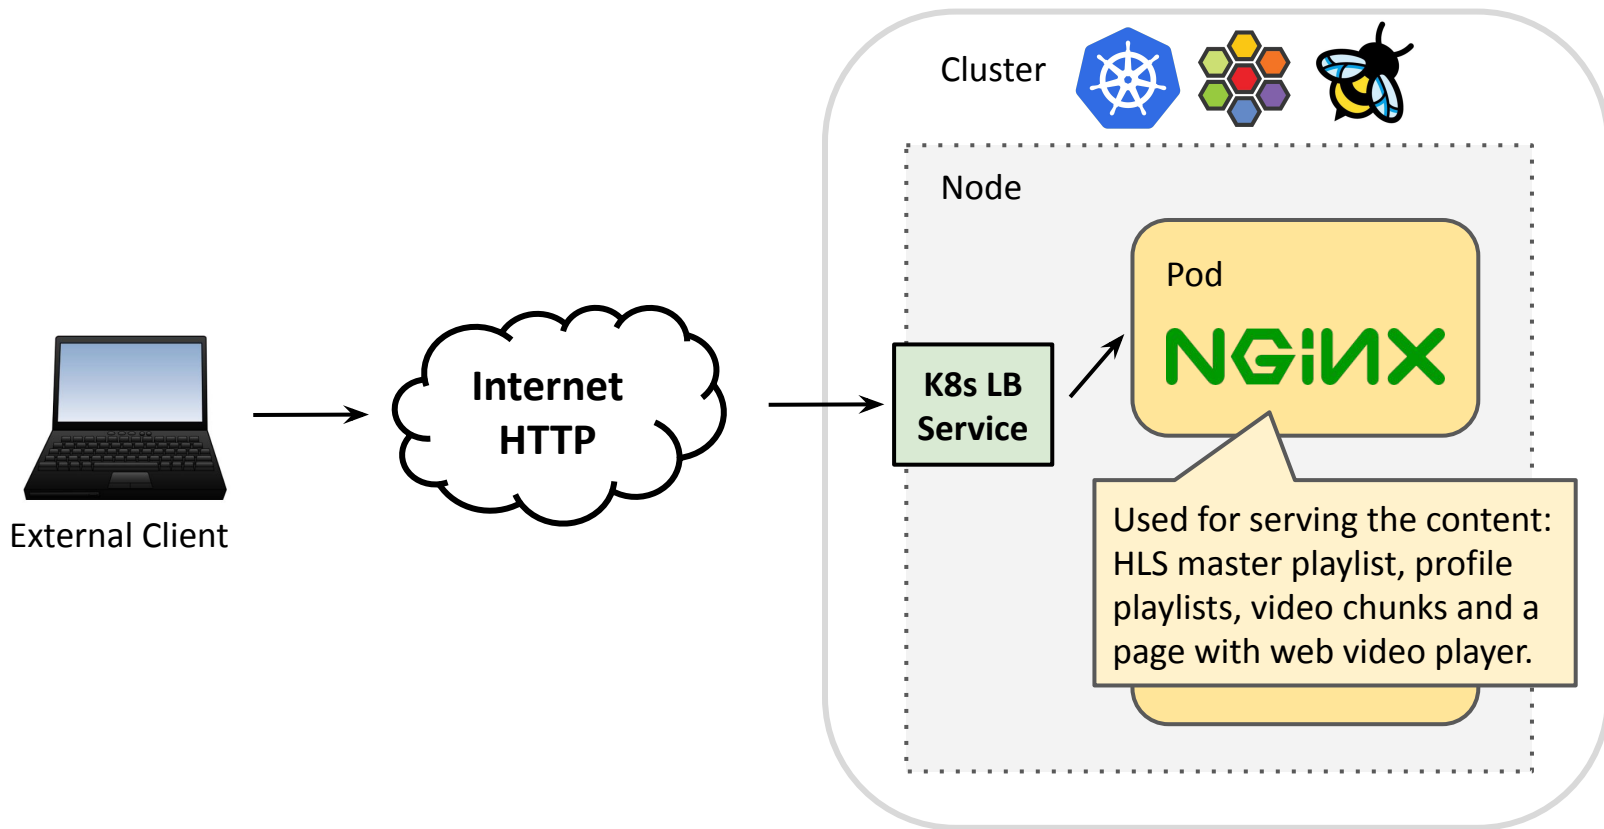
# Video Streaming Service Demo Setup: CUBIC vs BBR



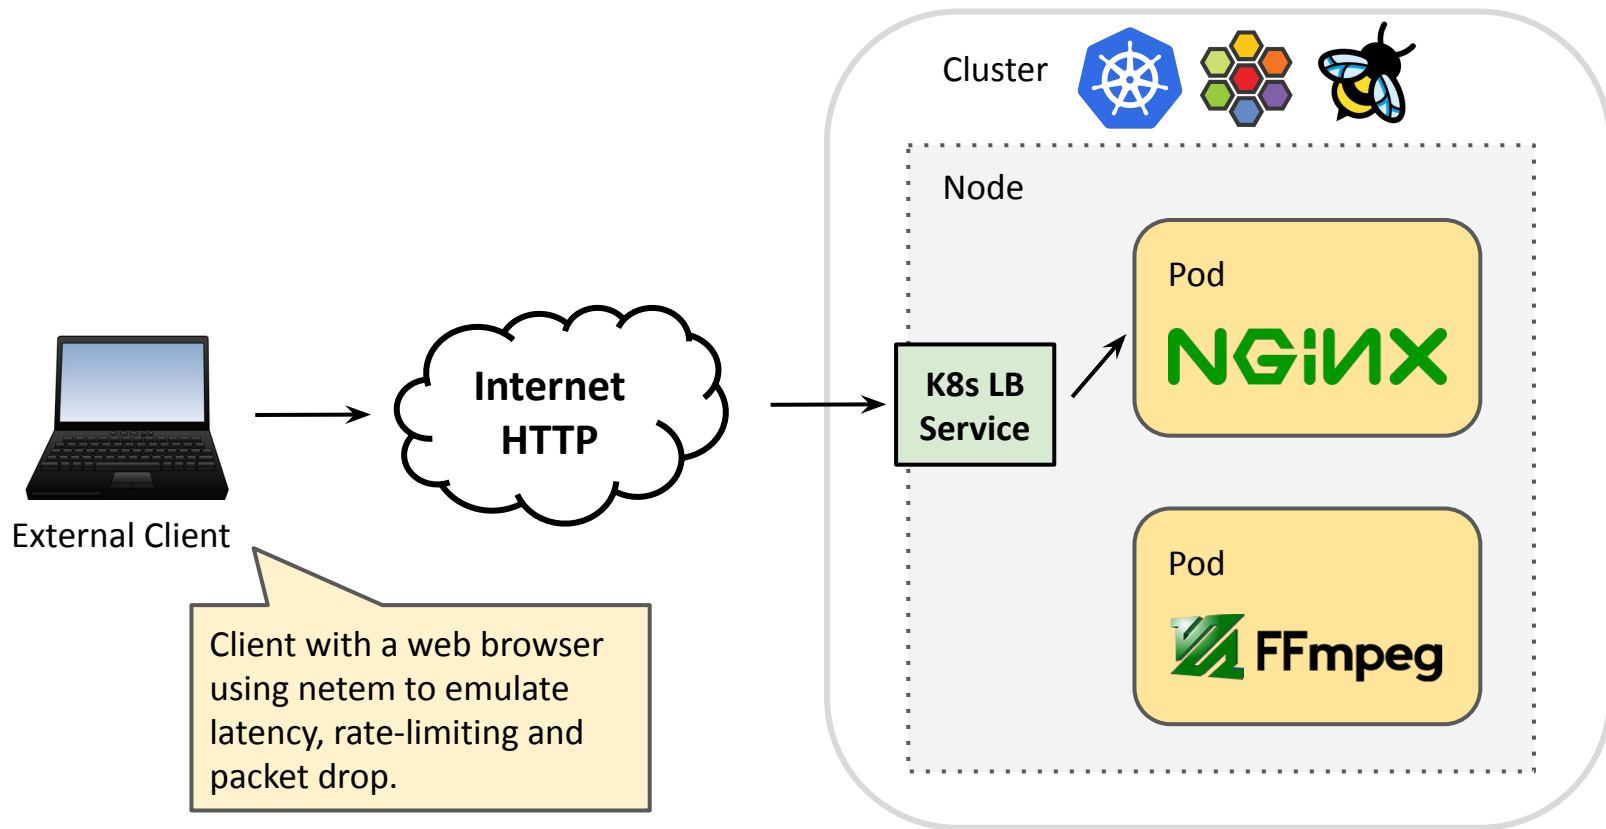
# Video Streaming Service Demo Setup: CUBIC vs BBR



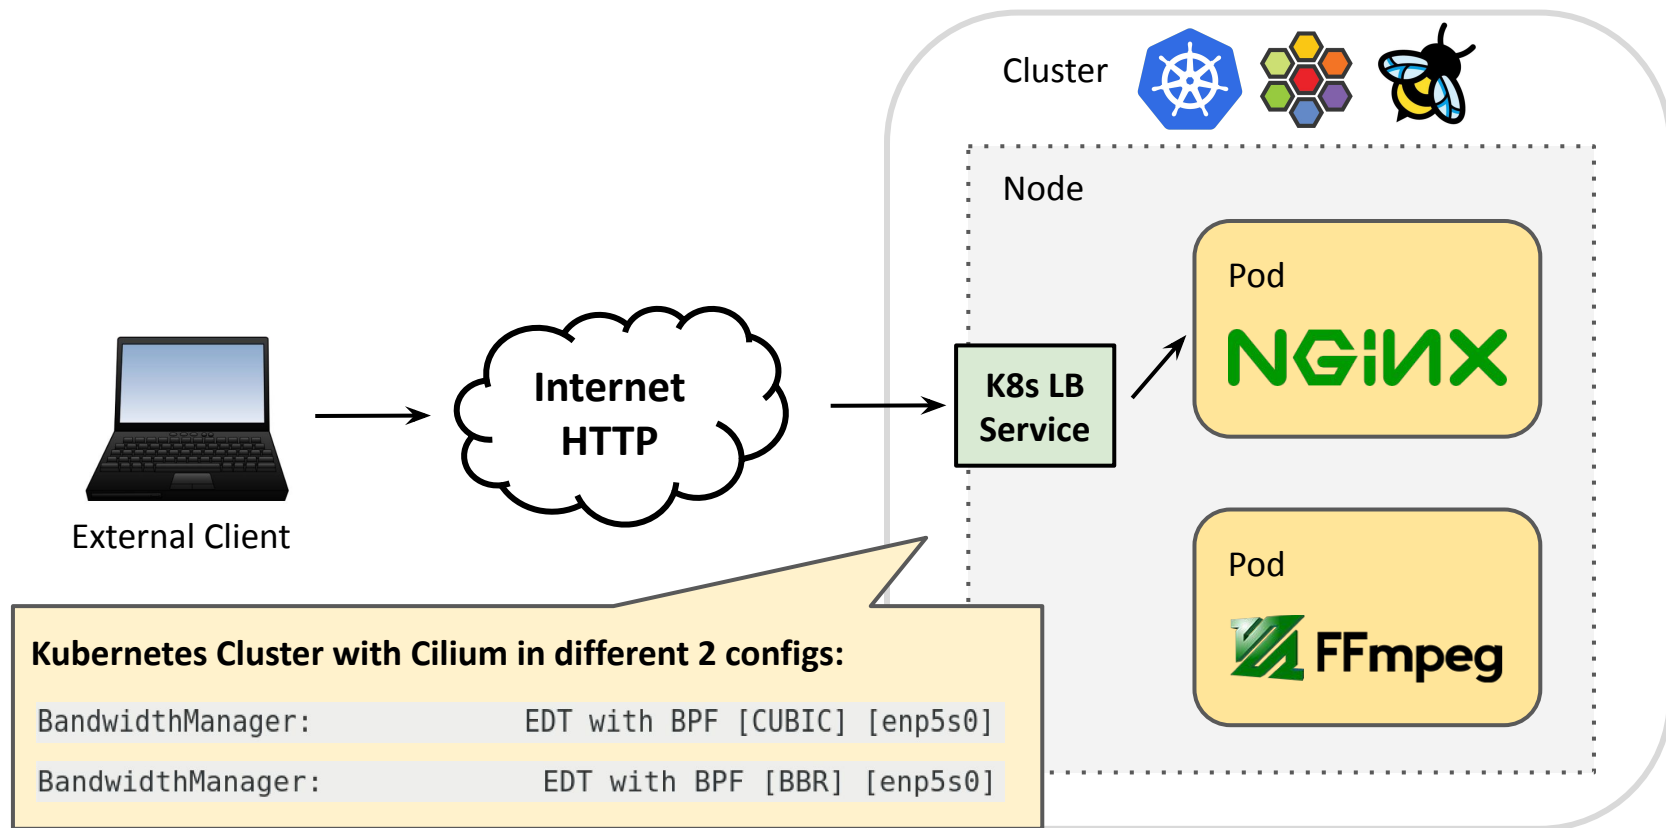
# Video Streaming Service Demo Setup: CUBIC vs BBR



# Video Streaming Service Demo Setup: CUBIC vs BBR



# Video Streaming Service Demo Setup: CUBIC vs BBR



ISOVALENT

A photograph of a theater interior. In the foreground, rows of plush red seats are visible, receding into the distance. The stage is at the far end, covered by a large, dark, vertically-pleated curtain. The lighting is dim, creating a sense of anticipation. The text "Now to the stream ..." is overlaid in the center of the image.

**Now to the stream ...**



# What needs to be considered with use of BBR?

- BBR has potential unfairness issues towards CUBIC when env uses both
- BBR will trigger a higher TCP retransmission rate (more aggressive probing)
- BBRv2 in the works to overcome them







## Revisiting earlier Problem Statement

- Kubernetes bandwidth enforcement does not need to be in a poor state
- Native implementation via Cilium's Bandwidth Manager (GA since v1.12)
  - ◆ Efficient, eBPF-based bandwidth enforcement via EDT model
  - ◆ First CNI to support BBR (& socket pacing) for Pods
  - ◆ Side-note: Realizing such architecture only possible with eBPF



### Getting Started Guide for Bandwidth Manager:

```
helm upgrade cilium ./cilium \  
  --namespace kube-system \  
  --reuse-values \  
  --set bandwidthManager.enabled=true \  
  --set bandwidthManager.bbr=true  
kubectl -n kube-system rollout restart ds/cilium
```

(needs Linux kernel v5.1+)

(needs Linux kernel v5.18+)

# Acknowledgements

- Van Jacobson
- Eric Dumazet
- Vytautas Valancius
- Stanislav Fomichev
- Martin Lau
- John Fastabend
- Cilium, BPF & netdev kernel community

ISOVALENT

# Thank you! Questions?

[github.com/cilium/cilium](https://github.com/cilium/cilium)

[cilium.io](https://cilium.io)

[ebpf.io](https://ebpf.io)

Isovalent: booth S21

Cilium: kiosk 12 (@ project pavilion)




# Appendix: Latency & TPS for 64/128/256 flows

# Comparison of Cilium's EDT implementation vs TBF

Single flow latency for EDT and HTB/TBF model (lower is better)

■ EDT ■ TBF/HTB




Env: 64 concurrent request/response type flows (TCP\_RR), 100M rate per flow



# Comparison of Cilium's EDT implementation vs TBF

Single flow latency for EDT and HTB/TBF model (lower is better)




Env: 128 concurrent request/response type flows (TCP\_RR), 100M rate per flow



# Comparison of Cilium's EDT implementation vs TBF

Single flow latency for EDT and HTB/TBF model (lower is better)

EDT TBF/HTB



Env: 256 concurrent request/response type flows (TCP\_RR), 100M rate per flow




# Comparison of Cilium's EDT implementation vs TBF

Total transaction rate between EDT and HTB/TBF model (higher is better)

■ EDT ■ TBF/HTB

Env: 64 concurrent  
request/response  
type flows (TCP\_RR),  
100M rate per flow






# Comparison of Cilium's EDT implementation vs TBF

Total transaction rate between EDT and HTB/TBF model (higher is better)

■ EDT ■ TBF/HTB



Env: 128 concurrent request/response type flows (TCP\_RR), 100M rate per flow



# Comparison of Cilium's EDT implementation vs TBF

Total transaction rate between EDT and HTB/TBF model (higher is better)

■ EDT ■ TBF/HTB

