



KubeCon



CloudNativeCon


Europe 2022

Distributing PromQL for Fast and Efficient Kubernetes Fleet Monitoring





Moad Zardab
Senior Software Engineer
Red Hat



Filip Petkovski
Senior Production Engineer
Shopify

Agenda

- **Why is PromQL difficult to scale?**
- **Query pushdown**
- **Query sharding**
- **Sharding in practice**
- **Outlook**



KubeCon



CloudNativeCon

Europe 2022

Prometheus Overview

- Ubiquitous for Cloud Native monitoring
- Effective for real time monitoring
- Powerful query language

Ideal for **single cluster monitoring**

- Scraping across cluster boundaries is unreliable
- Relies on disk and memory for retention
- Lacking good scalability mechanisms



Larger Scale, Different Problems


- Projects like Thanos and Cortex gain mass adoption, enabling:
 - **Cheap, long term 'infinite retention'**
 - **Write replication** for redundancy
 - **Multi-tenancy** support
 - Prometheus-compliant **query API**



But...

- **PromQL engine** is still **single-threaded**
 - Needs all of your time-series data present before query evaluation
 - Exasperated by longer-term retention

Thanos Query Path Bottleneck




Bottlenecks on the with the Thanos query path?

- *Thanos-query* has to **stream all series into memory** in a single process
 - Query components need to **scale resource utilization with metric retention**
 - **GiBs worth of data over the network and into memory** for a **single query** that might return a single dimension (e.g. *topk*, *min/max*)
 - Makes component susceptible to **OOM**
- *Thanos-query* has to fanout queries to all targets
 - Time series can be **duplicated**
 - Wasted resources on **deduplicating blocks**

The community was calling out for..

- More efficient load-balancing
- Better query component availability
- Reduce overall bandwidth required for query evaluation



PromCon
North America 2021

Query pushdown; avoid streaming series all-together

Query Pushdown

`max (container_memory_working_set_bytes)`



`max (container_memory_working_set_bytes)`



`max (container_memory_working_set_bytes)`




Query Pushdown

Implemented in Thanos here <https://github.com/thanos-io/thanos/pull/4917>

- Lower latency by
 - Processing data at-rest
 - Significantly reduced bandwidth
- Can be enabled with `--enable-feature=query-pushdown`

Limitations

- Currently applicable to: **max, min, topk, bottomk** and **group**
- Common operations like *sum/rate, avg, histogram_quantile* cannot be pushed down
- Not all store components can execute queries
- Query execution is not free



PromCon
North America 2021

Query sharding; split a *single* query into *multiple distinct* and *disassociated* queries

Query Sharding

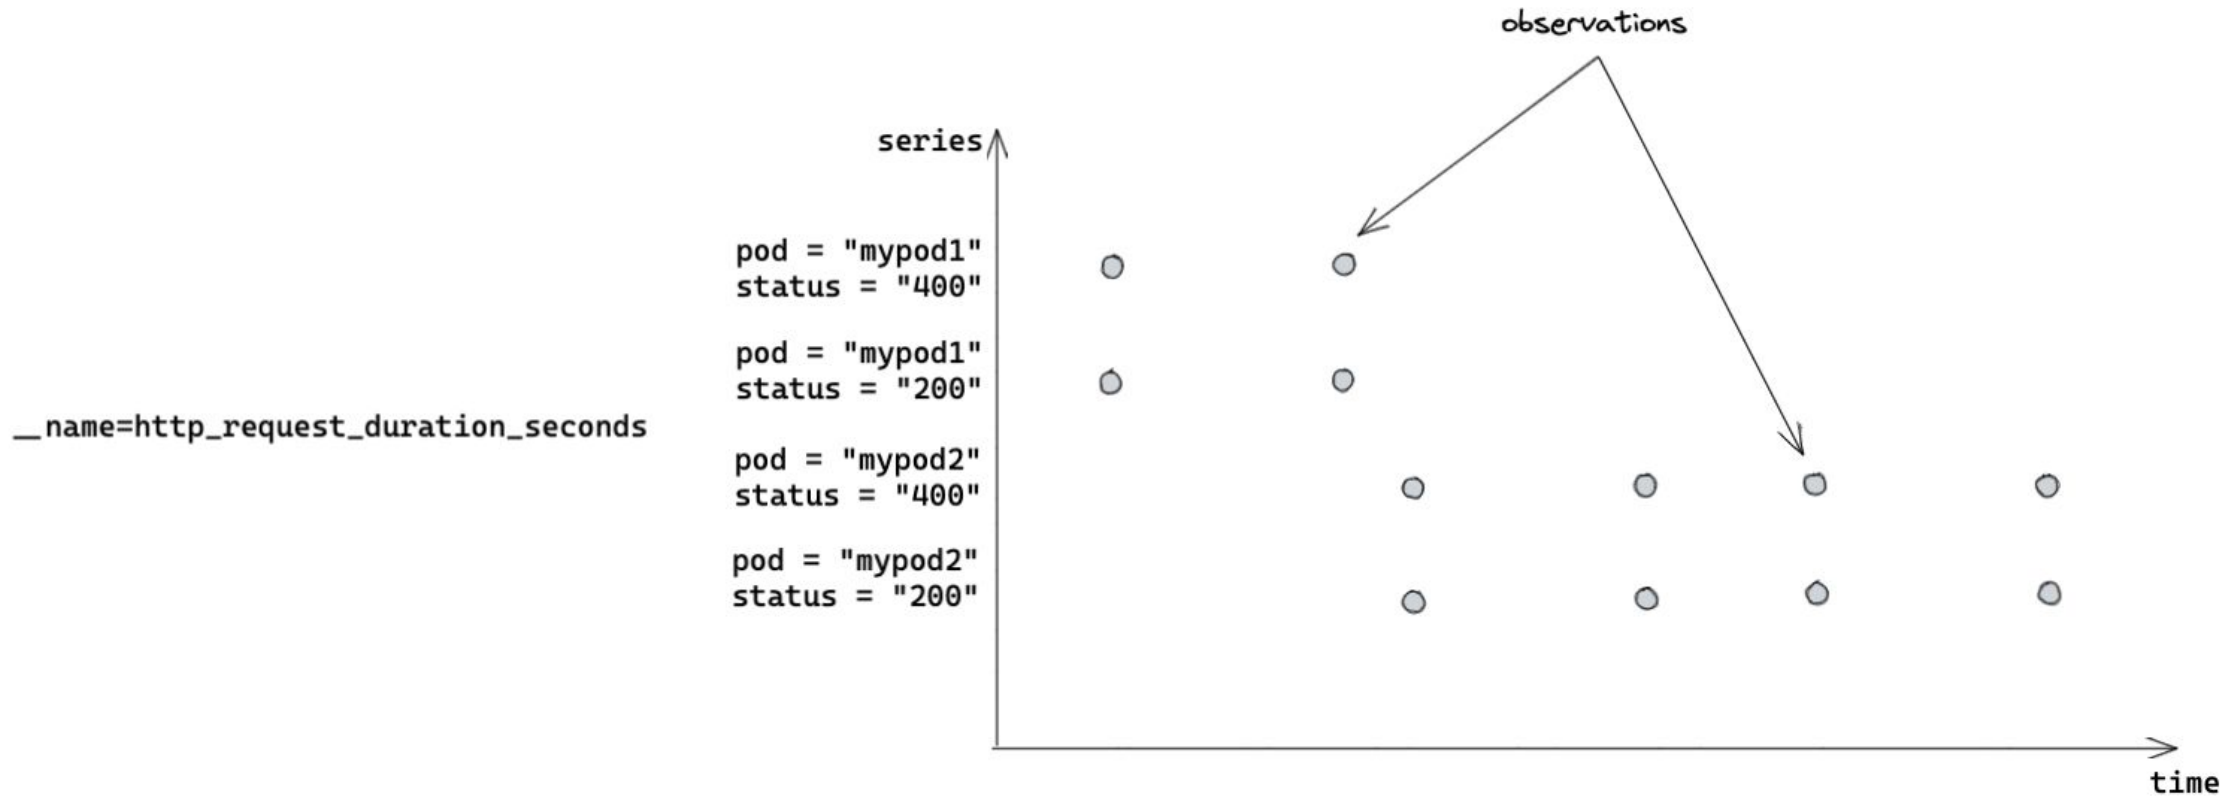


KubeCon



CloudNativeCon

Europe 2022



Query Sharding

Thanos already partially addresses this via **horizontal sharding**

- **Time-range** query splitting
 - e.g. query spanning **four days** becomes **four, one day** queries
- **Only works** for range queries
- Even for short time intervals, metrics can still **have high cardinality**

Query Horizontal Sharding




KubeCon



CloudNativeCon

Europe 2022

```
sum by (pod) (rate(http_request_duration_seconds[1m]))
```



Query Horizontal Sharding



KubeCon



CloudNativeCon

Europe 2022

```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t3"
```




```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t1"
```

```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t1"  
end="t2"
```

```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t2"  
end="t3"
```

- **One query** has been split into **N equivalent range queries** (*where N is the total query time range / query range interval*)
- Each query represents a **disassociated time range**
- *But...* new queries can still have **unbounded cardinality** for the given time range



PromCon
North America 2021

How do we manage metric cardinality?

Query Sharding

We worked on extending this with **vertical sharding**:

- Sharding based on **time series**
- Each shard is **dissociated**
- Can be implemented **without** query planning ahead-of-time
 - Leaves are responsible for **deterministically** returning sharded subset
- No changes to the Thanos write path

We wanted to test this first by sharding **grouped sum/rate expressions**

Query Vertical Sharding



KubeCon

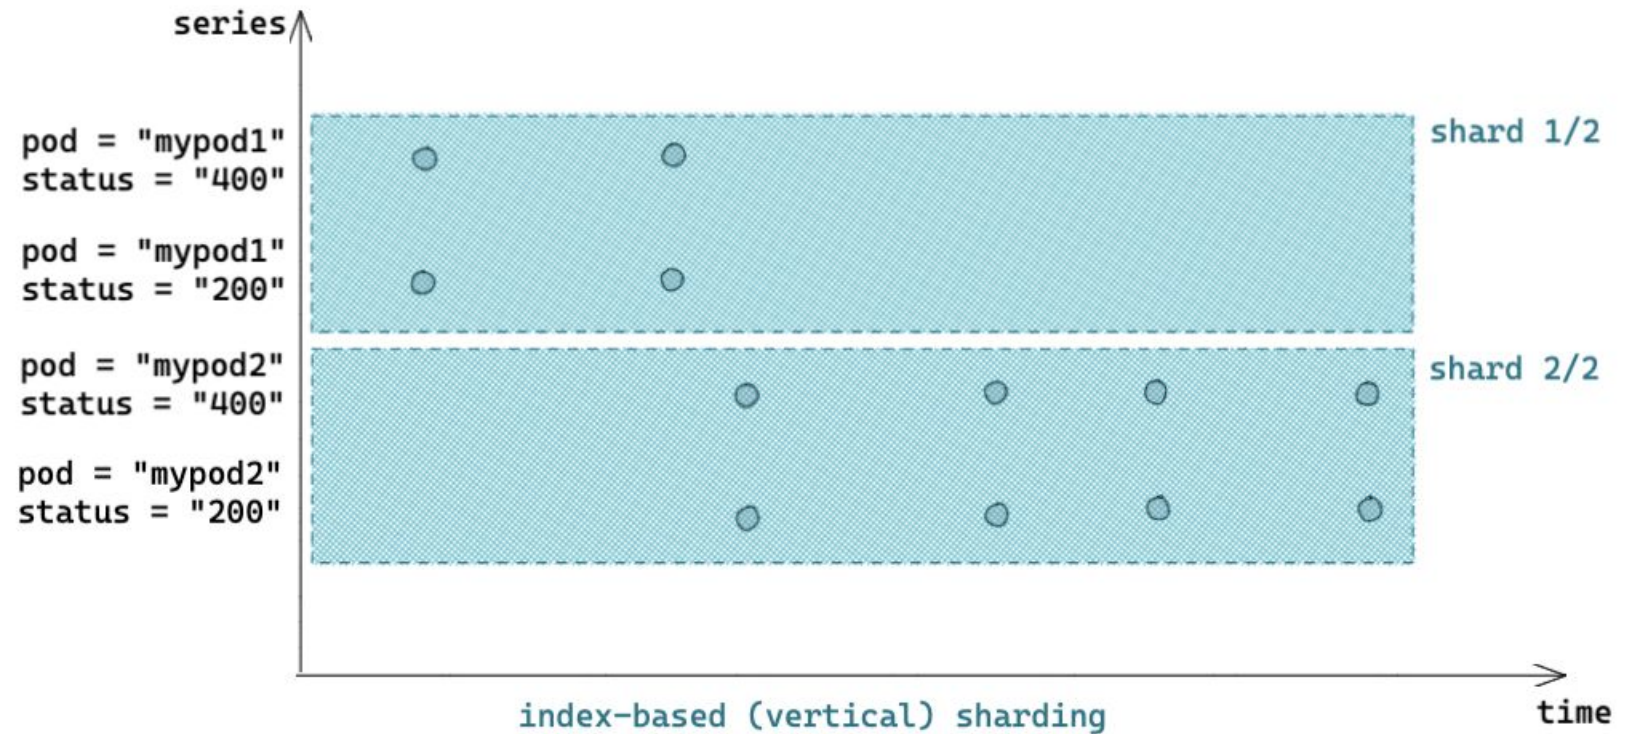


CloudNativeCon

Europe 2022

```
sum by (pod) (rate(http_request_duration_seconds[1m]))
```

__name=http_request_duration_seconds



Query Vertical Sharding



KubeCon



CloudNativeCon

Europe 2022

```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t3"
```



```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t3"  
shardIndex=1  
totalShards=2  
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t3"  
shardIndex=2  
totalShards=2
```

•

Query Vertical & Horizontal Sharding




KubeCon



CloudNativeCon

Europe 2022



Query Vertical & Horizontal Sharding



KubeCon



CloudNativeCon

Europe 2022

```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t3"
```

```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t1"  
  
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t1"  
end="t2"  
  
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t2"  
end="t3"
```

```
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t1"  
shardIndex=1  
totalShards=2  
  
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t0"  
end="t1"  
shardIndex=2  
totalShards=2  
  
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t1"  
end="t2"  
shardIndex=1  
totalShards=2  
  
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t1"  
end="t2"  
shardIndex=2  
totalShards=2  
  
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t2"  
end="t3"  
shardIndex=1  
totalShards=2  
  
expr="sum by (pod) (rate(http_request_duration_seconds[1m]))"  
start="t2"  
end="t3"  
shardIndex=2  
totalShards=2
```

- **One query** has been split into **N x M equivalent vertically** and **horizontally sharded queries** (where N is the total time range/ range interval, and M is the total number of vertical shards)
- Each subquery can be **pre-evaluated in parallel**
- Result is simply concatenation or re-evaluation of intermediate queries in the root querier

Vertical sharding performance

- Benchmark setup:
 - **100K** series, simulating **100 clusters** with **1000 pods** each
 - Executed one query consecutively for a fixed interval against 5 Queriers
 - Used a **sharding factor of 3**
 - Single node setup
- What we aimed to measure:
 - **User experience impact** (query latency)
 - Impact on **overall resiliency** (peak and average memory usage)

Query Latency (*without* sharding)




KubeCon



CloudNativeCon

Europe 2022

P90 latency



Start latency ~9s

End latency ~19s

Query Latency (*with* sharding)




KubeCon



CloudNativeCon

Europe 2022

P90 latency



Start latency ~5.2s (-43%)

End latency ~5.7s (-70%)

Memory Usage (*without* sharding)




KubeCon



CloudNativeCon

Europe 2022

Memory usage



Peak ~1.5GB
Average ~1.2GB

Memory Usage (*with sharding*)



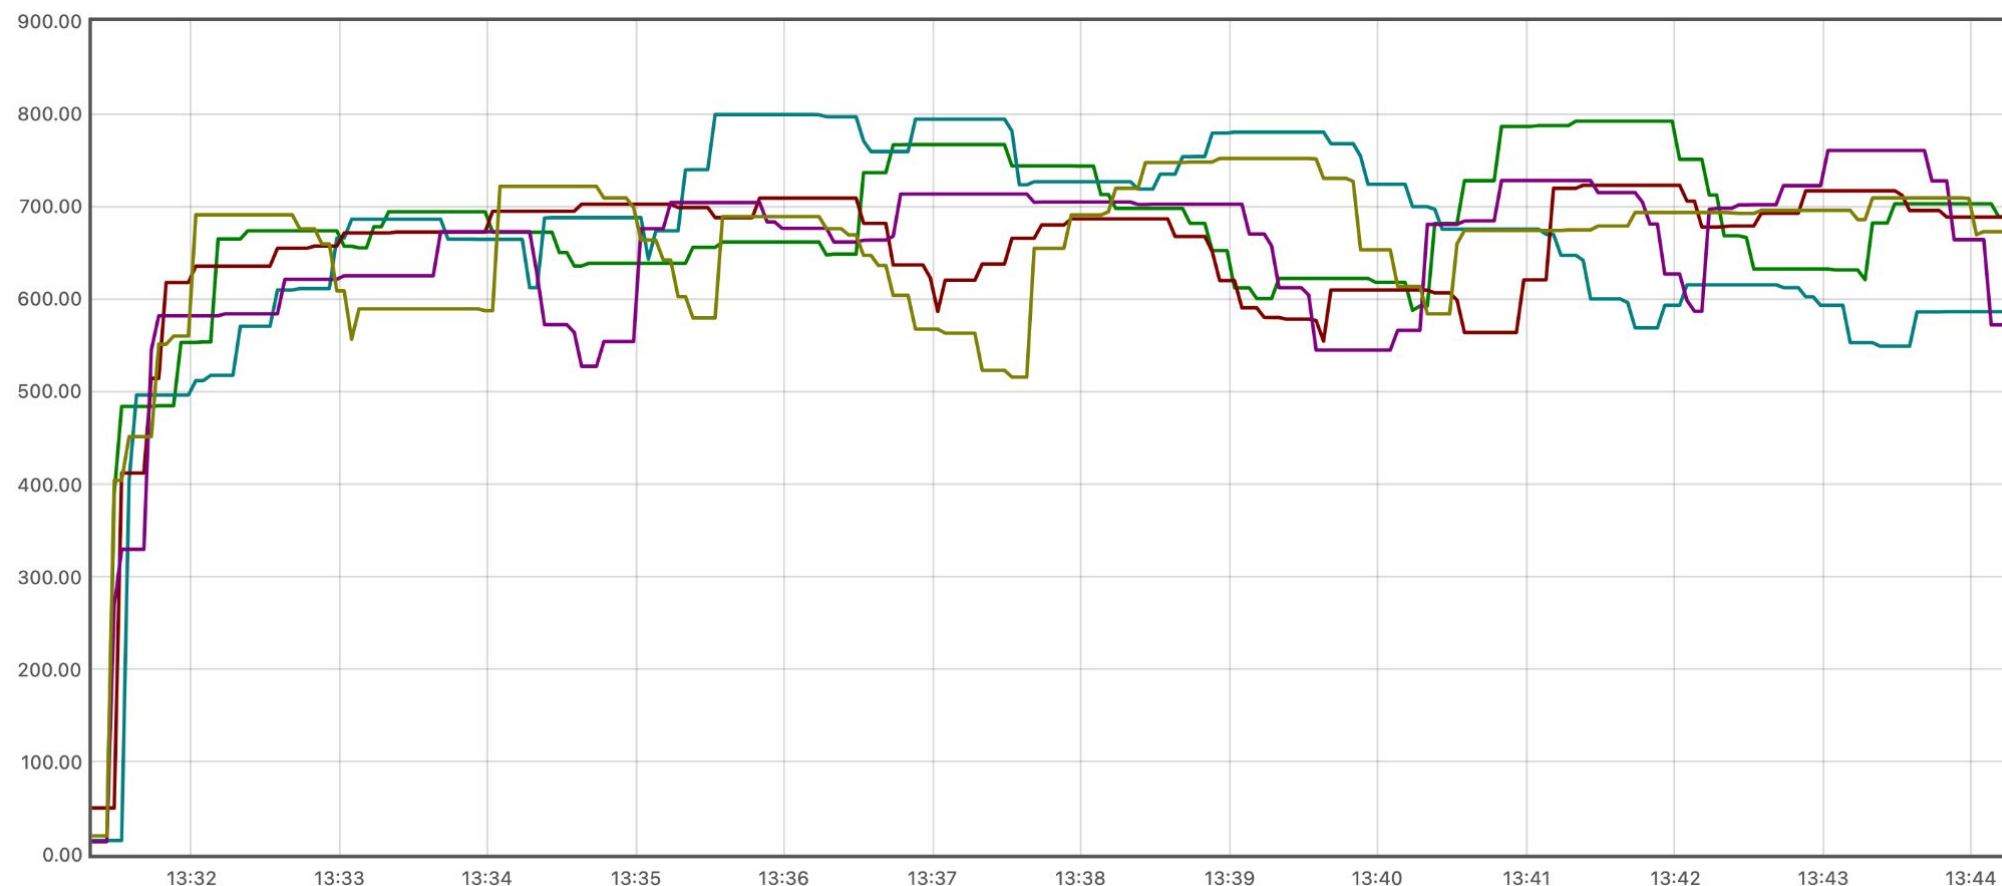
KubeCon



CloudNativeCon

Europe 2022

Memory usage



Peak ~800MB (-47%)
Average ~700MB (-42%)

Vertical Sharding in Summary

Benefits

- Greatly reduced peak memory utilization
- Query latency reduced by double digit percentage
- Applicable to instant queries (including alerting and recording rules)

Caveats

- Currently only implemented for PromQL aggregations
- Increased request volume throughout the system

Demonstration: Fleet Monitoring

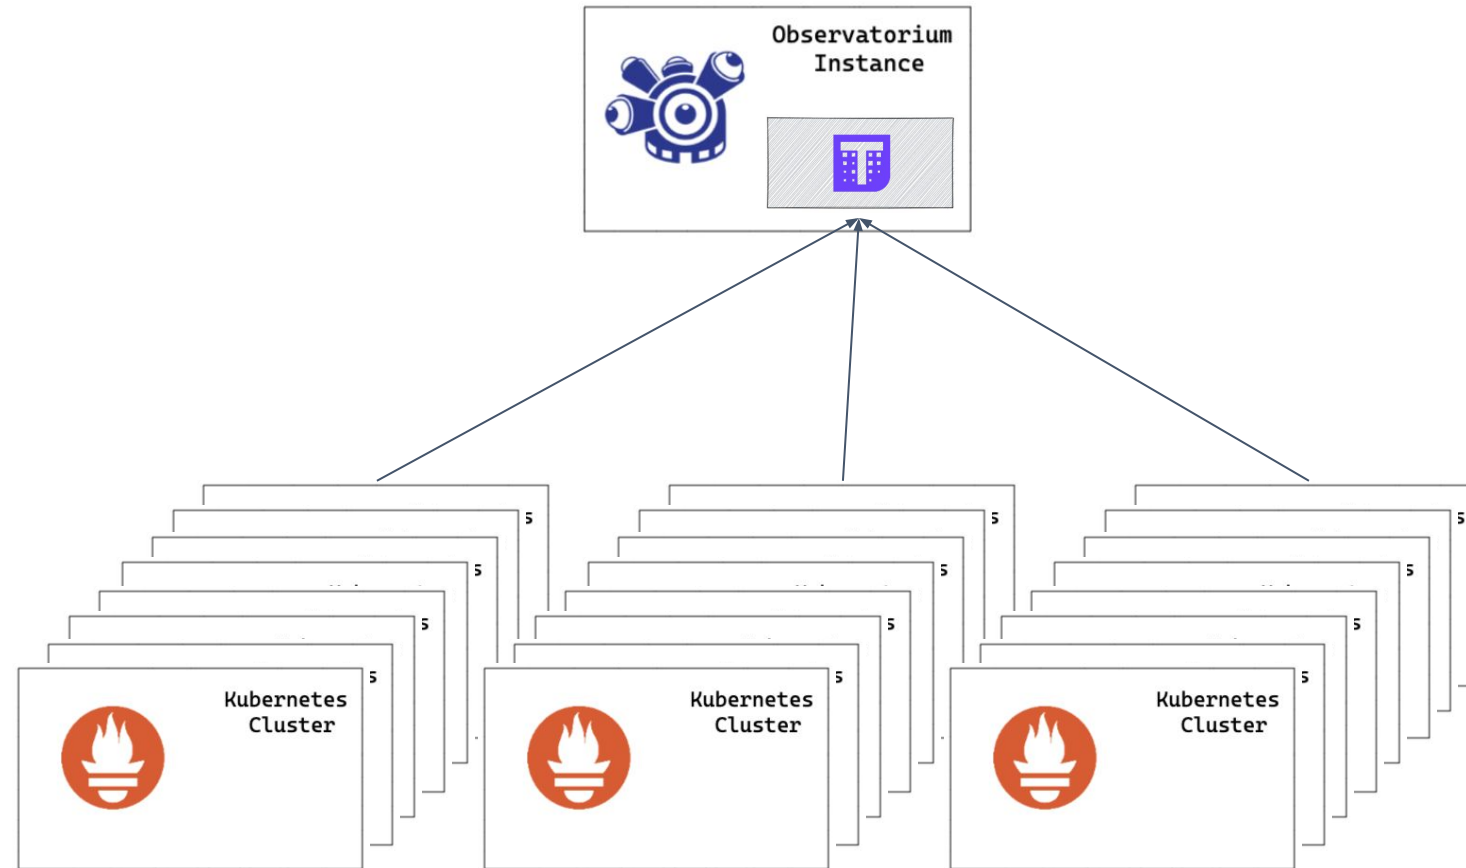


KubeCon



CloudNativeCon

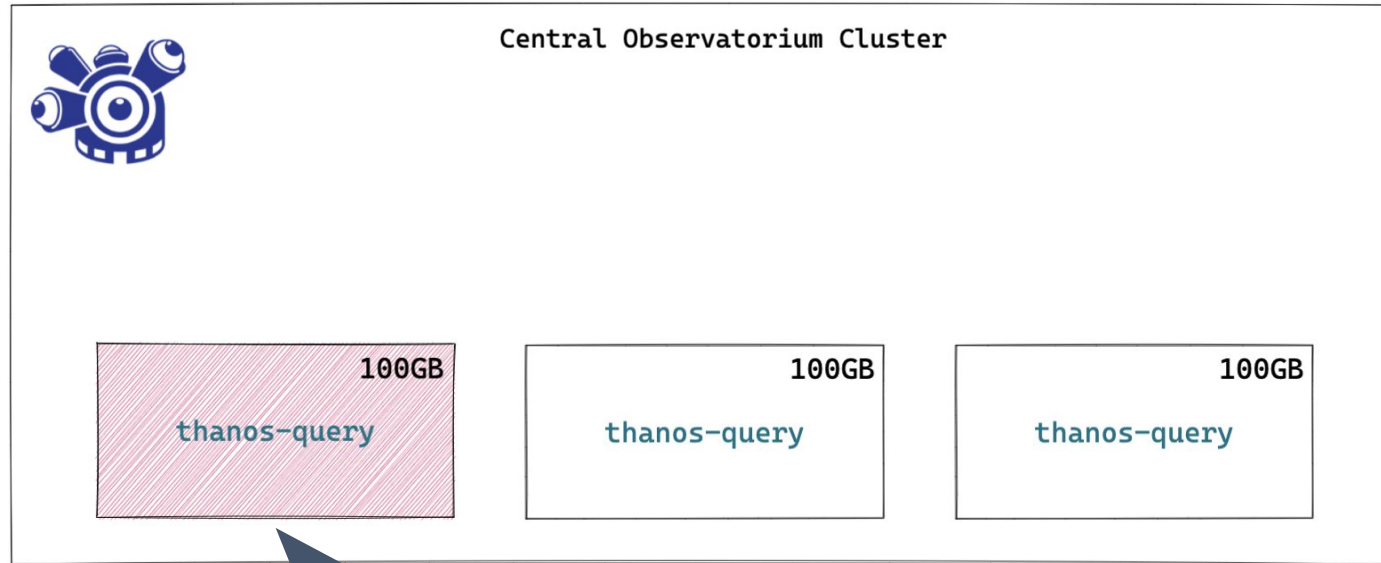
Europe 2022



20,000 clusters, 1000 series/cluster, 30 day retention

Case Study: Observability-as-a-Service Query

Load-balancing



```
sum by (cluster_id) (container_memory_working_set_bytes)
```

Case Study: Observability-as-a-Service Query Load-balancing

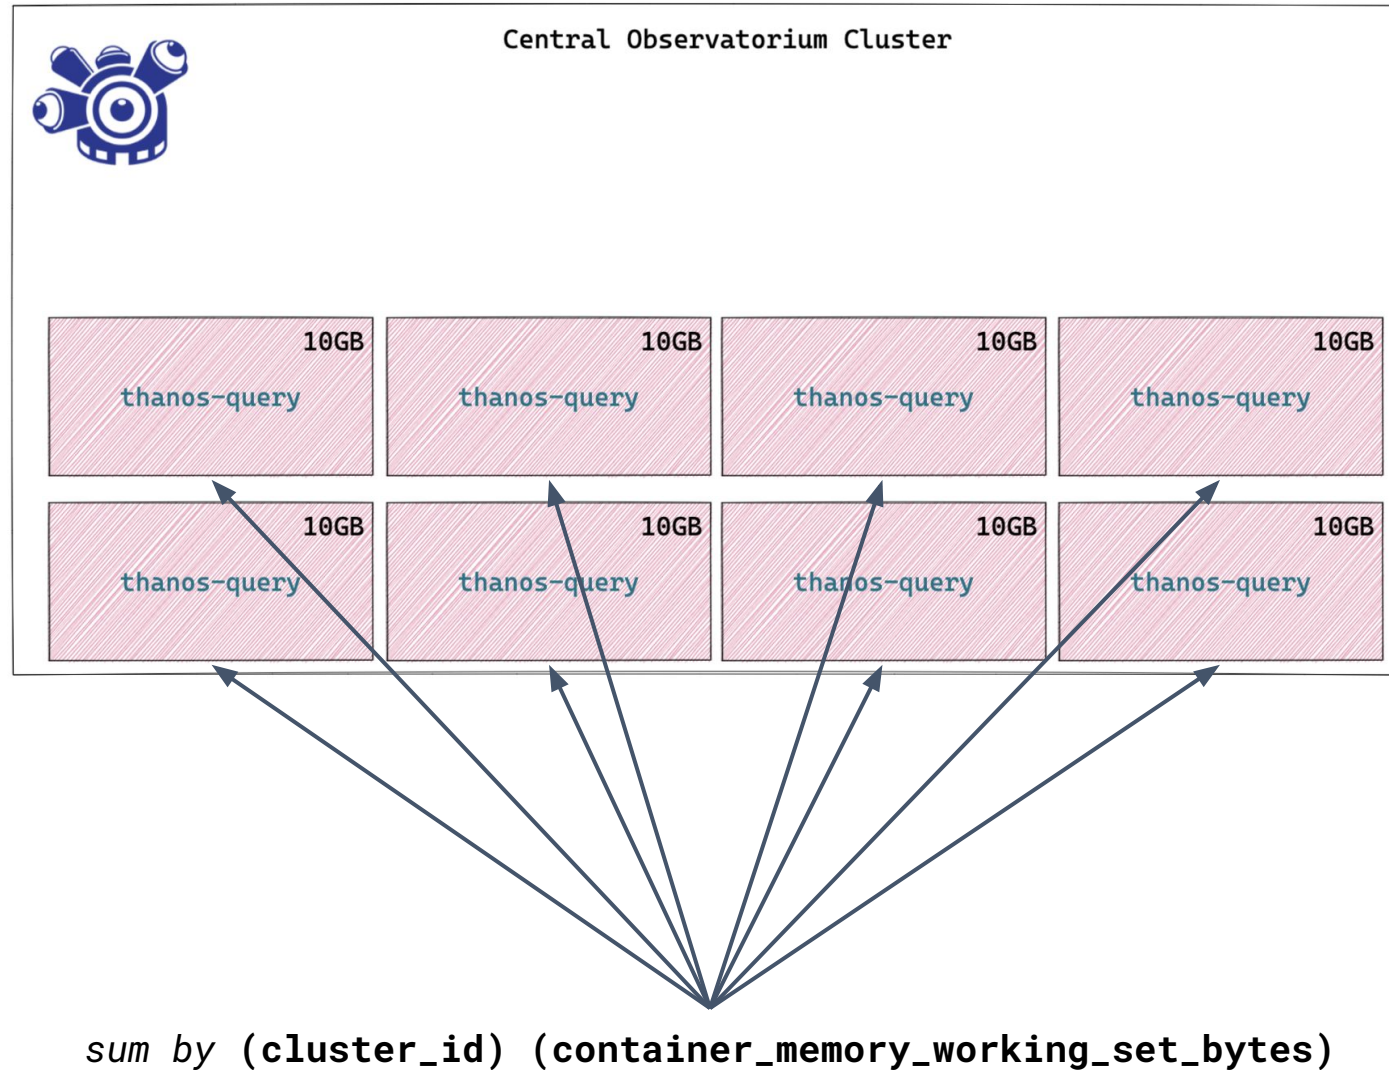


KubeCon



CloudNativeCon

Europe 2022



Query Sharding Future

- Proposal in upstream Thanos <https://github.com/thanos-io/thanos/pull/5350>
- Reference implementation in upstream: Thanos <https://github.com/thanos-io/thanos/pull/5342>
- Expanding PromQL support to cover more expressions

More context:

- [The Prometheus TSDB, Björn Rabenstein](#)
- [Intro to Thanos: Scale Your Prometheus Monitoring With Ease, Lucas Servén & Dominic Green](#)
- [Using Thanos to gain a unified way to query over multiple clusters, Wiard van Rijn](#)
- [Thanos documentation](#)

Thanks!

- Thanos community + maintainers



KubeCon



CloudNativeCon

Europe 2022

fin

