



**KubeCon**



**CloudNativeCon**

**Europe 2022**

**WELCOME TO VALENCIA**





KubeCon



CloudNativeCon

Europe 2022

# Kubectl Said What?!?

Christopher Hanson, RX-M, llc.



# whoami

**Christopher Hanson** (Chris)

Senior Cloud Native Engineer at RX-M, llc.

**Instructor** – Kubernetes, Helm, Prometheus, Argo, Spinnaker, Docker, OpenShift, Concourse, Ansible, OpenStack (among others)

**Consultant / Advisor** – monolith to containers/K8s app migrations, K8s Cassandra DBaaS, monitoring/logging pipeline using the EFK stack

## Speaker

KubeCon / CloudNativeCon Europe Virtual 2021:

*Choose Wisely: Understanding Kubernetes Selectors*

(CNCF YouTube: <https://youtu.be/dLe0TZEGhxo>)

KubeCon / CloudNativeCon Europe Virtual 2020:

*Zero Downtime Deployments: Controlling Application Rollouts and Rollbacks*

(CNCF YouTube: <https://youtu.be/rh6EtXiNOj4>)

San Francisco Meetup: Managing Resources in K8s apps

(RX-M YouTube channel: <https://youtu.be/ijWZinxdHEA>)

Exam Developer: Certified Kubernetes Security Specialist (CKS)

Certified Kubernetes Administrator (CKA-1700-0131-0100)

Certified Kubernetes Application Developer (CKAD-1900-0994-0100)



KubeCon



CloudNativeCon

Europe 2022



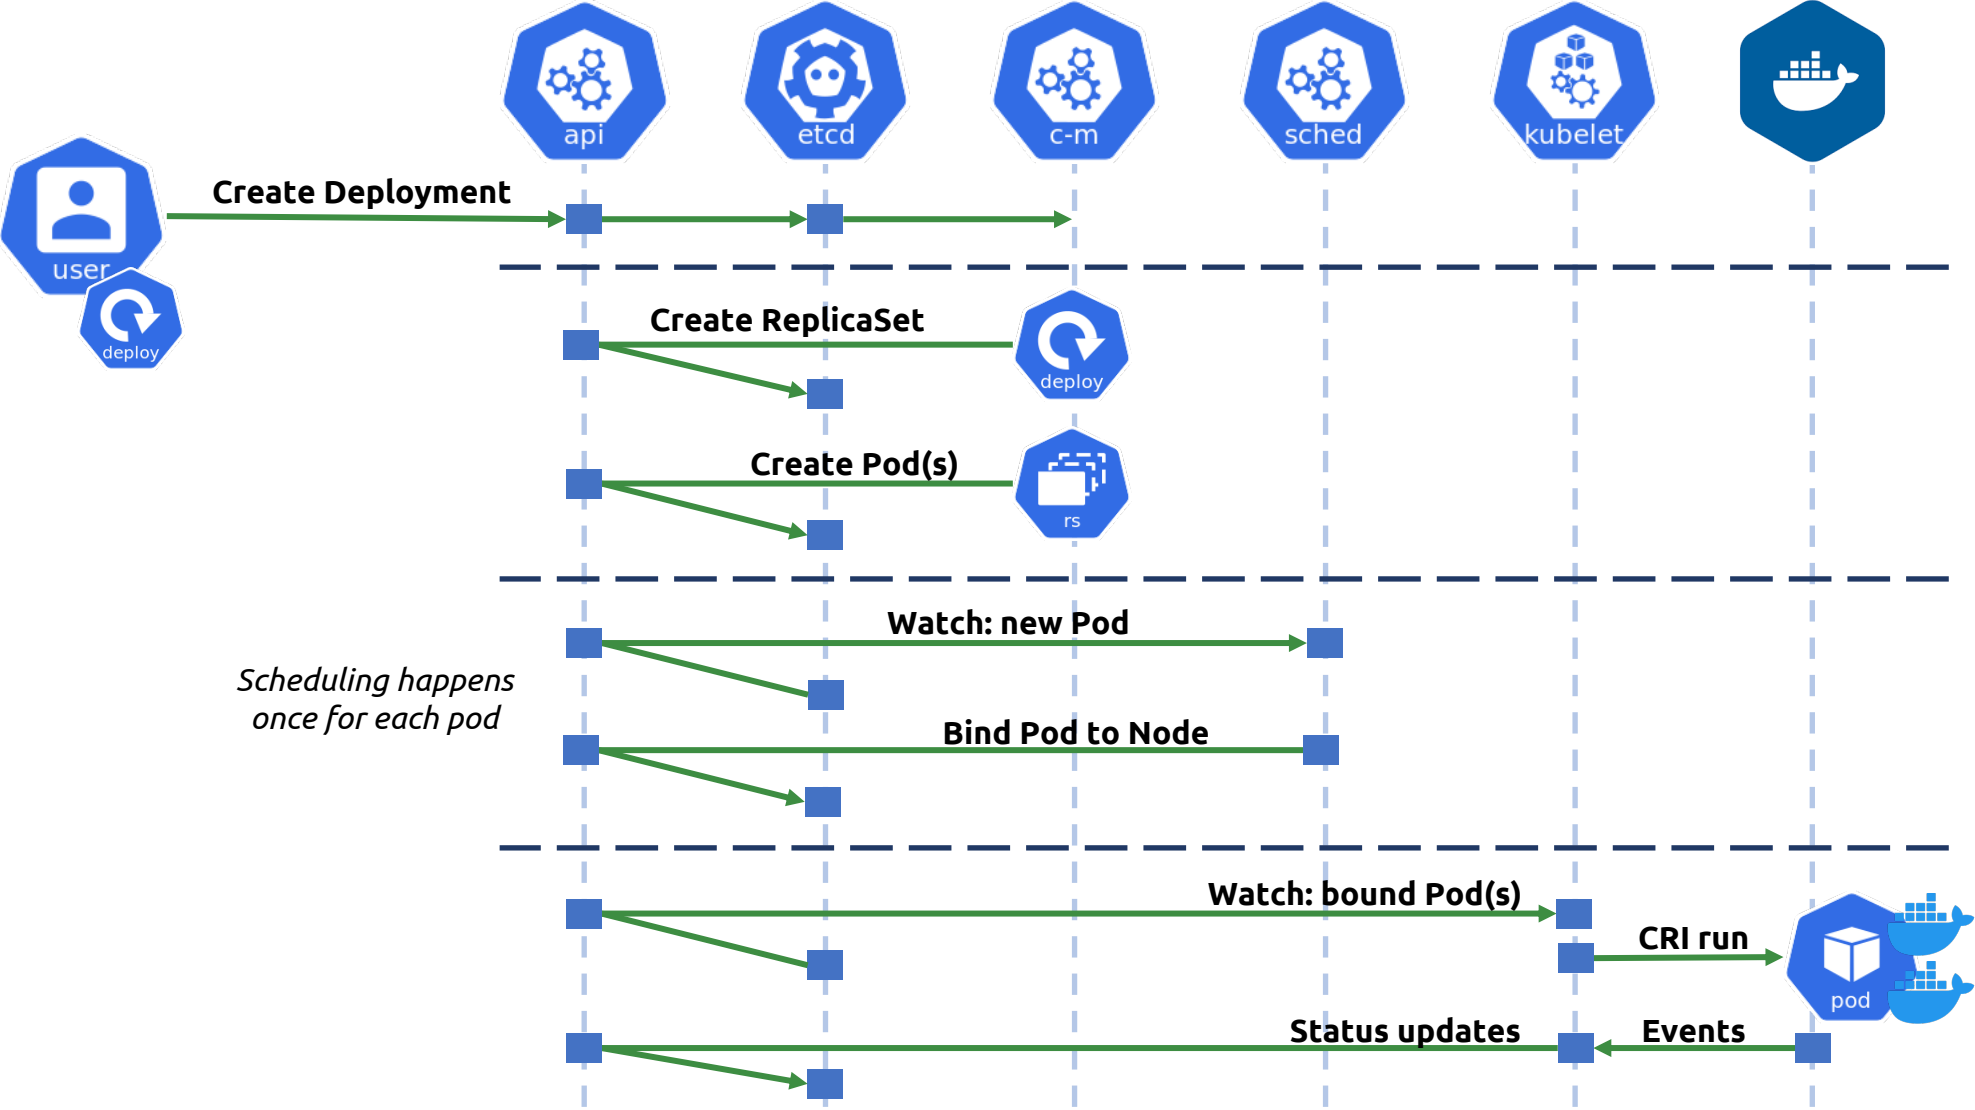
 @CloudNativChris

chris@rx-m.com

**rx-m** cloud native  
training &  
consulting

[rx-m.com](https://rx-m.com) / [info@rx-m.com](mailto:info@rx-m.com)

# Resource Creation Flow



# Deciphering kubectl's Messages

Kubernetes resources have "conditions", "phases", "states", and "status"

- Some are summaries, others are detailed
- All can provide clues to the state of your resources

When first learning K8s, it can be difficult (even frustrating) to find the information you need, when you need it

- Where do I find it?
- How is it presented?
- Some output can be noisy—how do I avoid red herrings?

*"I just dropped in to see what condition my condition was in"*

Information about resource states can be viewed using commands like:

- `kubectl get events | <object>`
- `kubectl get` with modifiers like `--output (-o)`:
  - `kubectl get -o yaml | json`
  - `kubectl get -o jsonpath under status.conditions[*].message`
- `kubectl describe`

Let's examine what these commands present with some practical examples...



# Controllers: Conditions

Controllers track conditions related to pods under their control

In *some cases*, can be seen with `kubectl describe`

- Conditions can be transient

## Deployment

- Available – minimum replicas are up and running for at least `minReadySeconds`
- Progressing – new RS(es) created, scale up/down
- ReplicaFailure – added when one of its pods fails to be created/or deleted

## ReplicaSet


- ReplicaFailure – added when one of its pods fails to be created/or deleted

## StatefulSet

- `statefulset.status.conditions` – is part of the object but is not presented by `kubectl`?

## Job

- `conditions.type` – Complete or Failed based on the pod(s) exit code(s)
  - If still running, "conditions" will not be present



# Controllers: Example

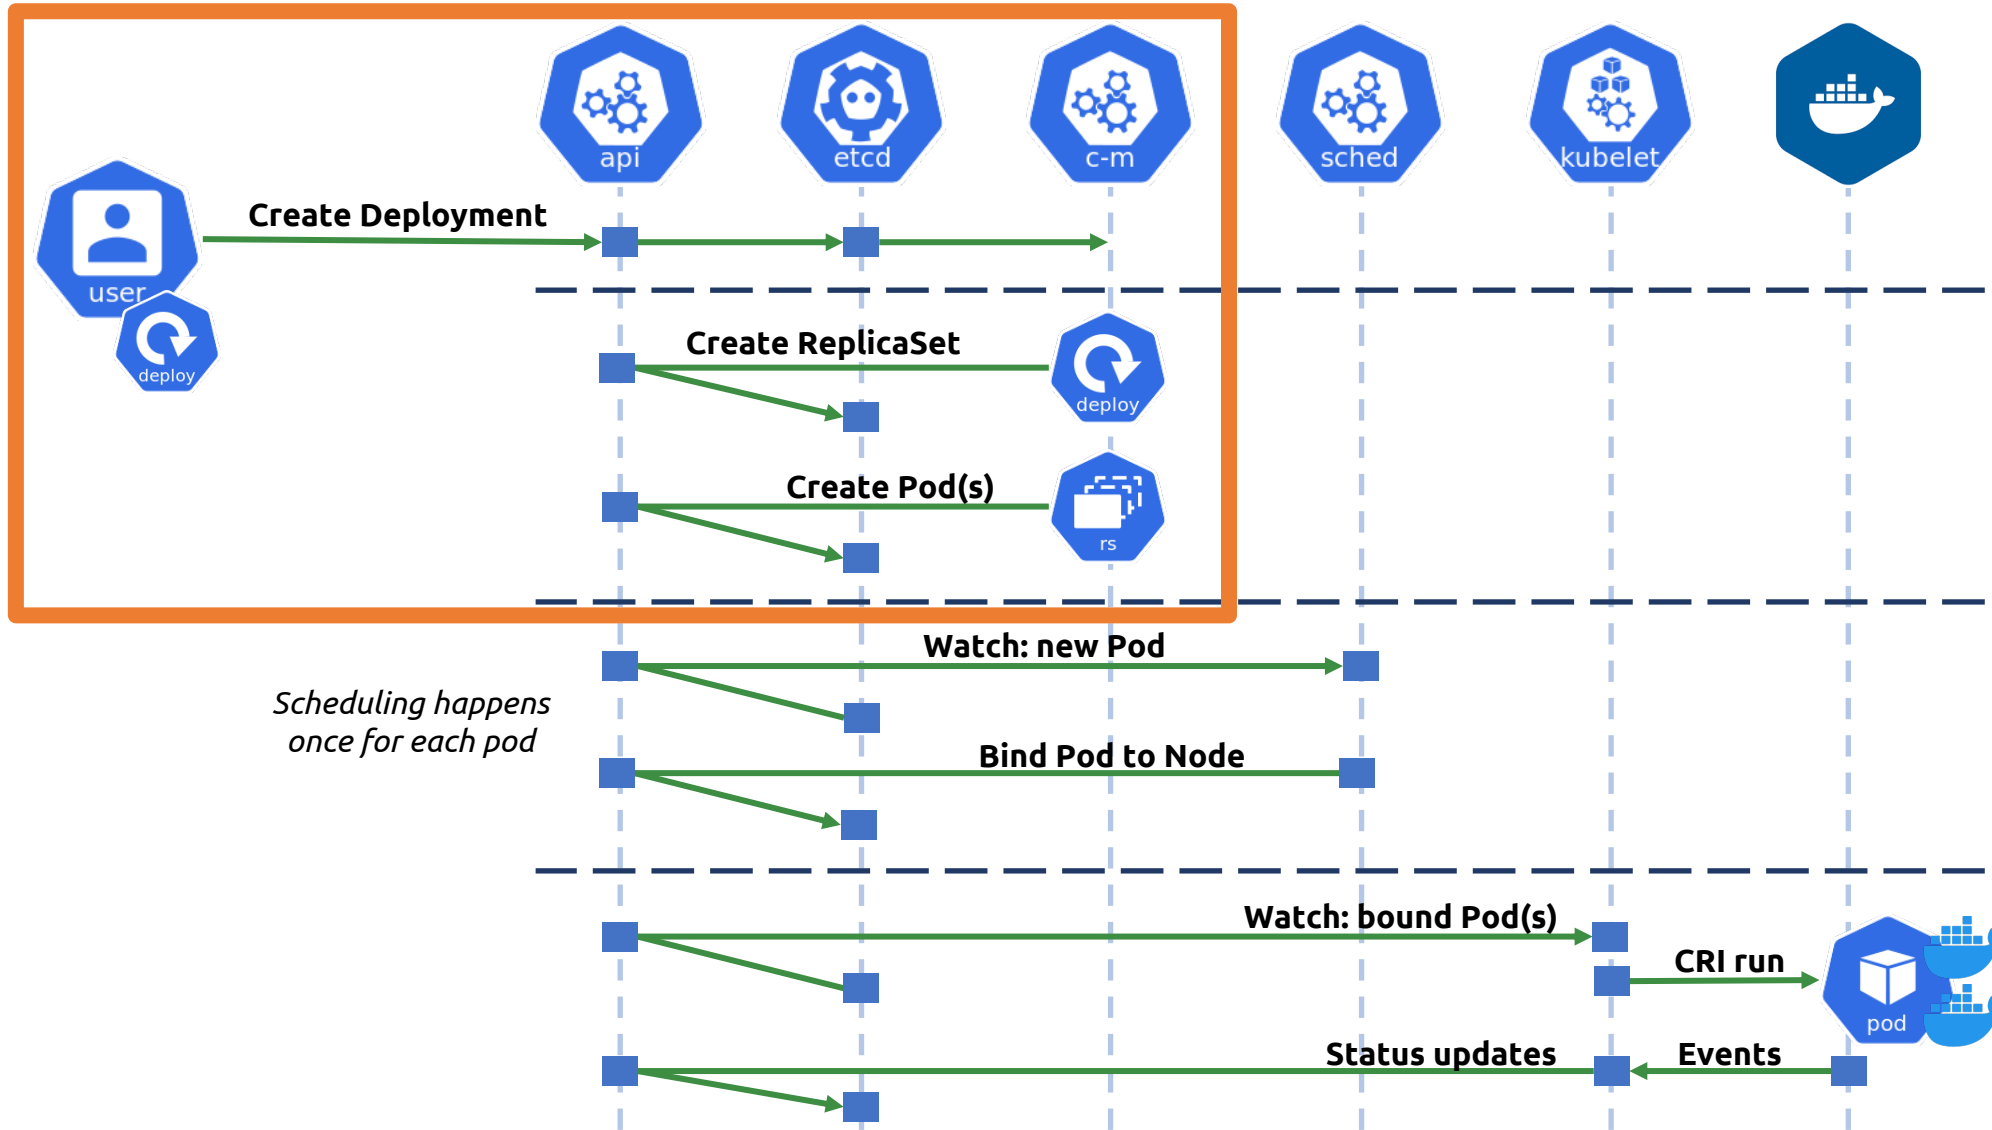


KubeCon



CloudNativeCon

Europe 2022





# Failing to Create Pods



KubeCon



CloudNativeCon

Europe 2022

## Symptoms:

Deployments / StatefulSets will have mismatches between available/ready replicas

- Will look identical to other issues – ***may be misleading!***

ReplicaSets will reflect accurate desired/current/ready replicas

- If no other problems exist, ready = current  $\neq$  desired

**Pods are missing** – pods that can't be created will not have any pod object to query for more info

Errors will populate a related condition "message" field (depending on controller)

## Common Cause:

**Quota** – namespace where a controller is creating pods has a resource quota in effect, blocking pod(s) from being created





# Pods & Containers

When the controller *can* create pod(s), the pod object exists and can be queried for more info

# Pods: Phases

Pod Phases are **macro states** in the lifecycle of a Kubernetes pod

Seen when using `kubectl get pod` but also available from a pod's `status.phase`

**Pending** – pod accepted by the system

- $\geq 1$  of the containers has *not* been created
- Includes time before scheduling and time downloading images

**Running** – pod has been bound to a node (scheduled)

- All of a pod's containers have been created
- $\geq 1$  container is running, starting, or restarting

**Succeeded** – all container(s) in a pod have terminated in success (zero exit) and will not be restarted

**Failed** –  $\geq 1$  container(s) terminated in failure (non-zero exit or terminated by the cluster)

**Unknown** – state of the pod cannot be obtained

- Most likely a communication error with the kubelet on the host where the pod is scheduled/running



KubeCon



CloudNativeCon

Europe 2022

# Pods: Conditions

Array of conditions the pod has / has not passed in its lifecycle

Can be easily found/viewed with `kubectl describe`

**Ready** – able / not able to serve client/peer requests, should be added to service/routing mesh

**Initialized** – all init containers completed successfully (zero exit) or are still running/restarting

**ContainersReady** – all containers in the Pod are / are not ready

**PodScheduled** – the pod has / has not been scheduled to a node

# Pending Pods: Examples

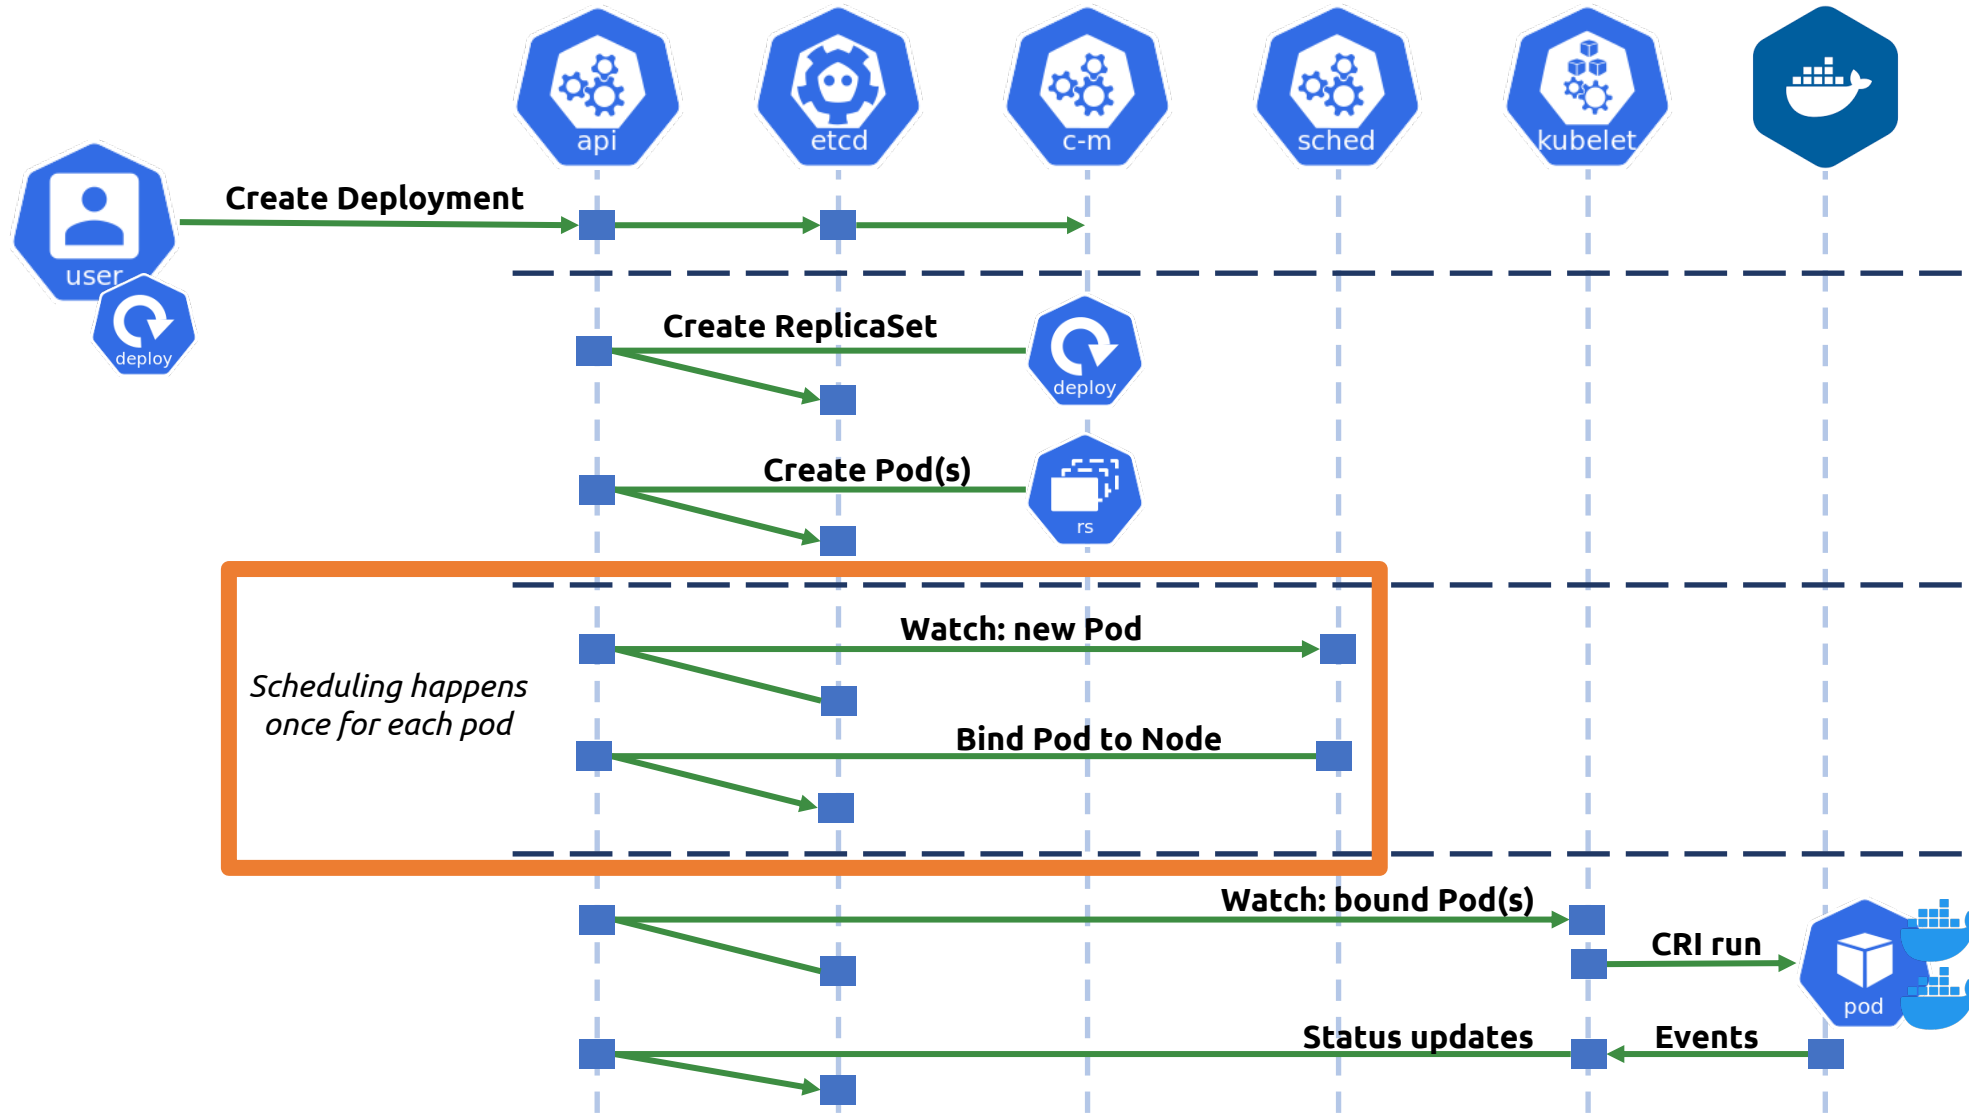


KubeCon



CloudNativeCon

Europe 2022



# Pending: Symptoms

Controllers will have mismatches between desired/current/ready replicas

- Will look identical to other issues – ***may be misleading!***

Pods will exist and `kubectl get pod` will return **Pending** (duh)

K8s will report the pod condition **PodScheduled** as "false"

- No other pod conditions will be present

Pod condition "message" field and related events will provide more details about the cause

Container states will be null – no containers are created unless a pod is scheduled to a node



KubeCon



CloudNativeCon

Europe 2022

# Pending: Common Causes

**Insufficient resources** – container request(s) cannot be met (CPU, memory, disk)

- Ignore "limits" as these *are not* used by the scheduler

**Selectors** – pod specifies nodeName, node label selector (nodeSelector), or node/pod affinity/anti-affinity that doesn't match any nodes or pods

**Volumes** – a pod's PVC cannot be fulfilled by any PVs, missing ConfigMap(s) and/or Secret(s)

- ConfigMaps & Secrets have additional messaging (covered later)

**Priority Class Preemption** – a running pod may be evicted in favor of another pod of a higher priority class

- Controller will attempt to replace the evicted pod with a replacement which will sit in pending state because resources are scarce
- **Let's examine this special case before moving on...**

# Running Pods, Containers: States

Each container in a pod has its state tracked

Includes containers, initContainers, and ephemeralContainers (if enabled)

Can be viewed—under each container's name—with `kubectl describe`

- The `containerStatuses` list of `kubectl get -o yaml` is a better presentation/grouping of this info

**Waiting** – container is running operations to complete start up

- Includes image pulls and ConfigMap or Secret mounting
- Back-off time during CrashLoopBackOff

**Running** – container processes executing

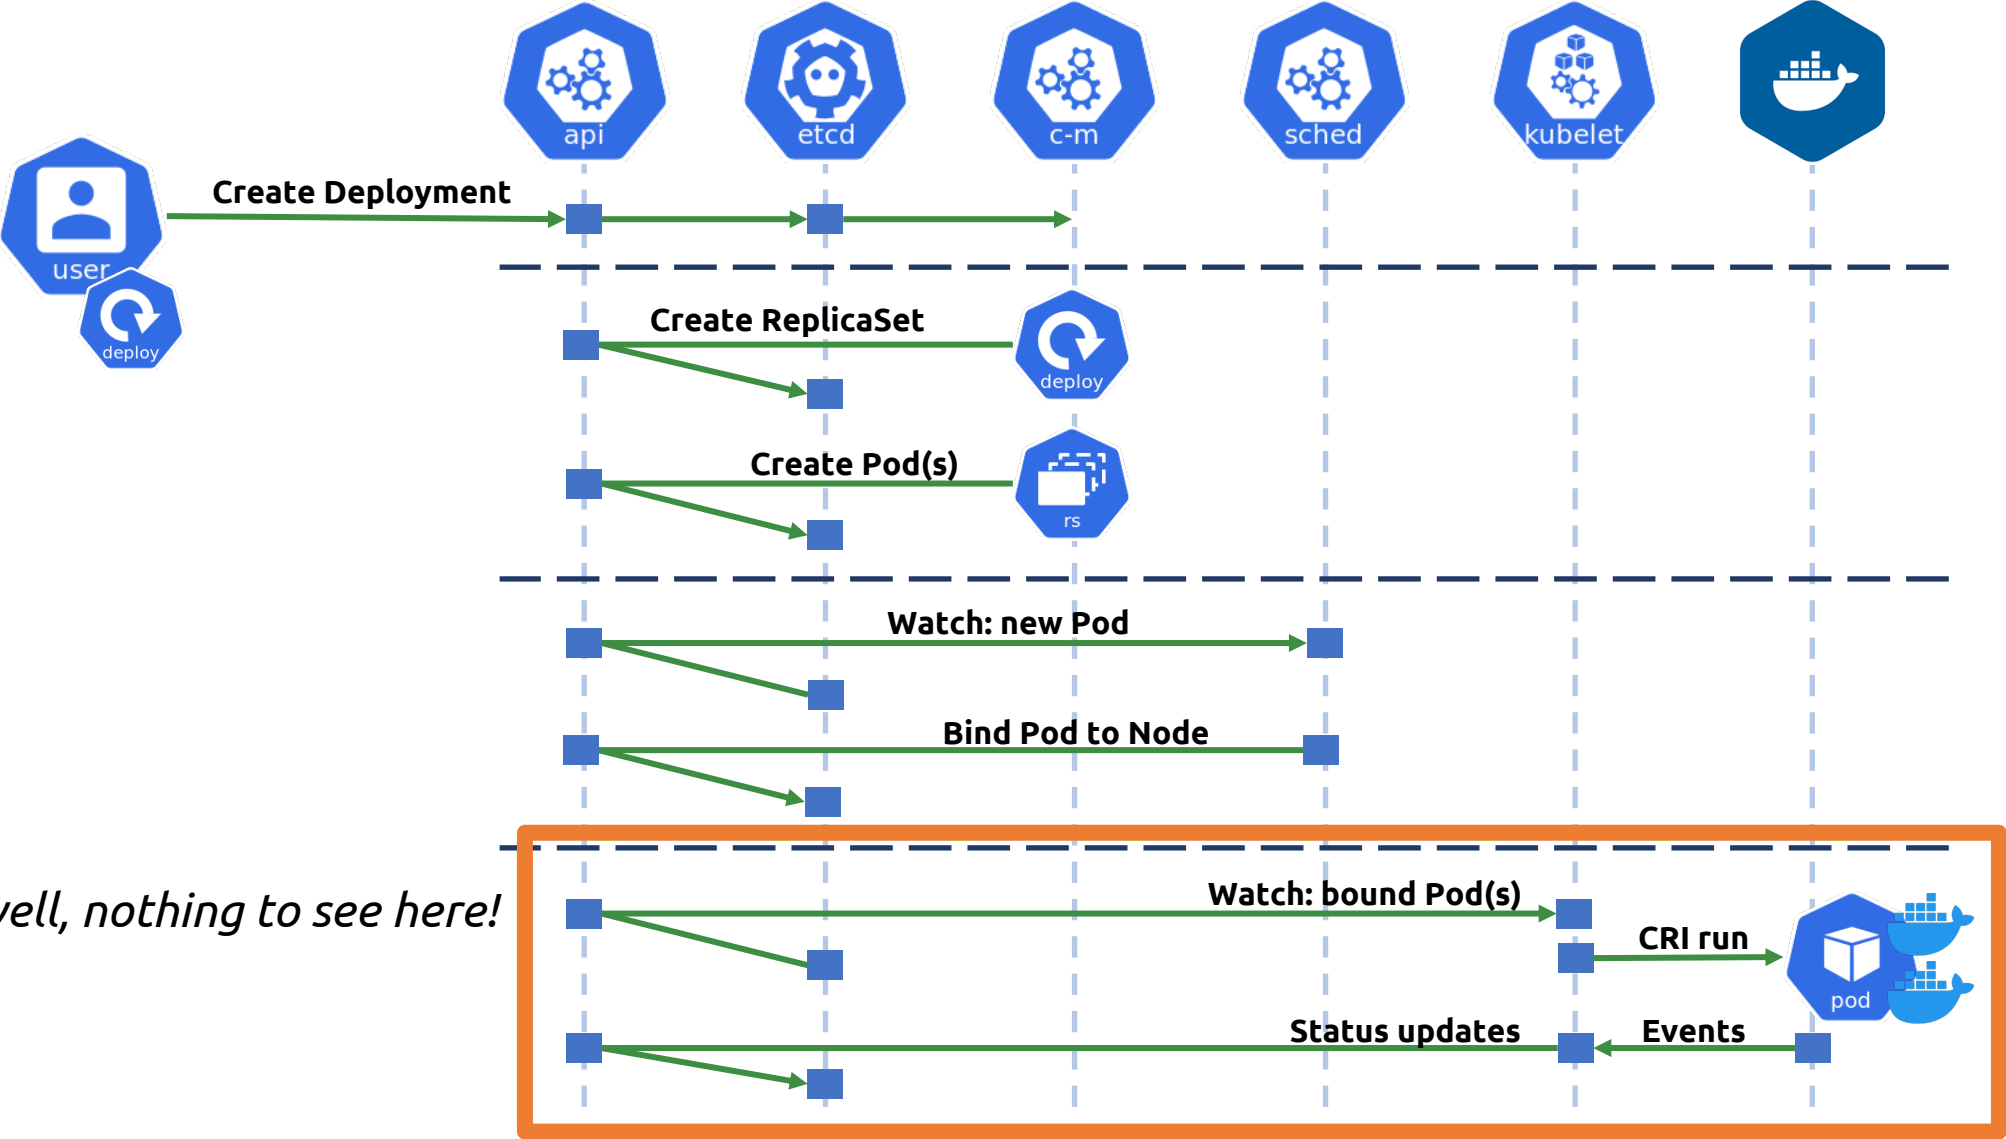
- Any postStart hooks completed successfully

**Terminated** – container ran and failed or finished successfully

- Any preStop hooks executed before a container enters this state



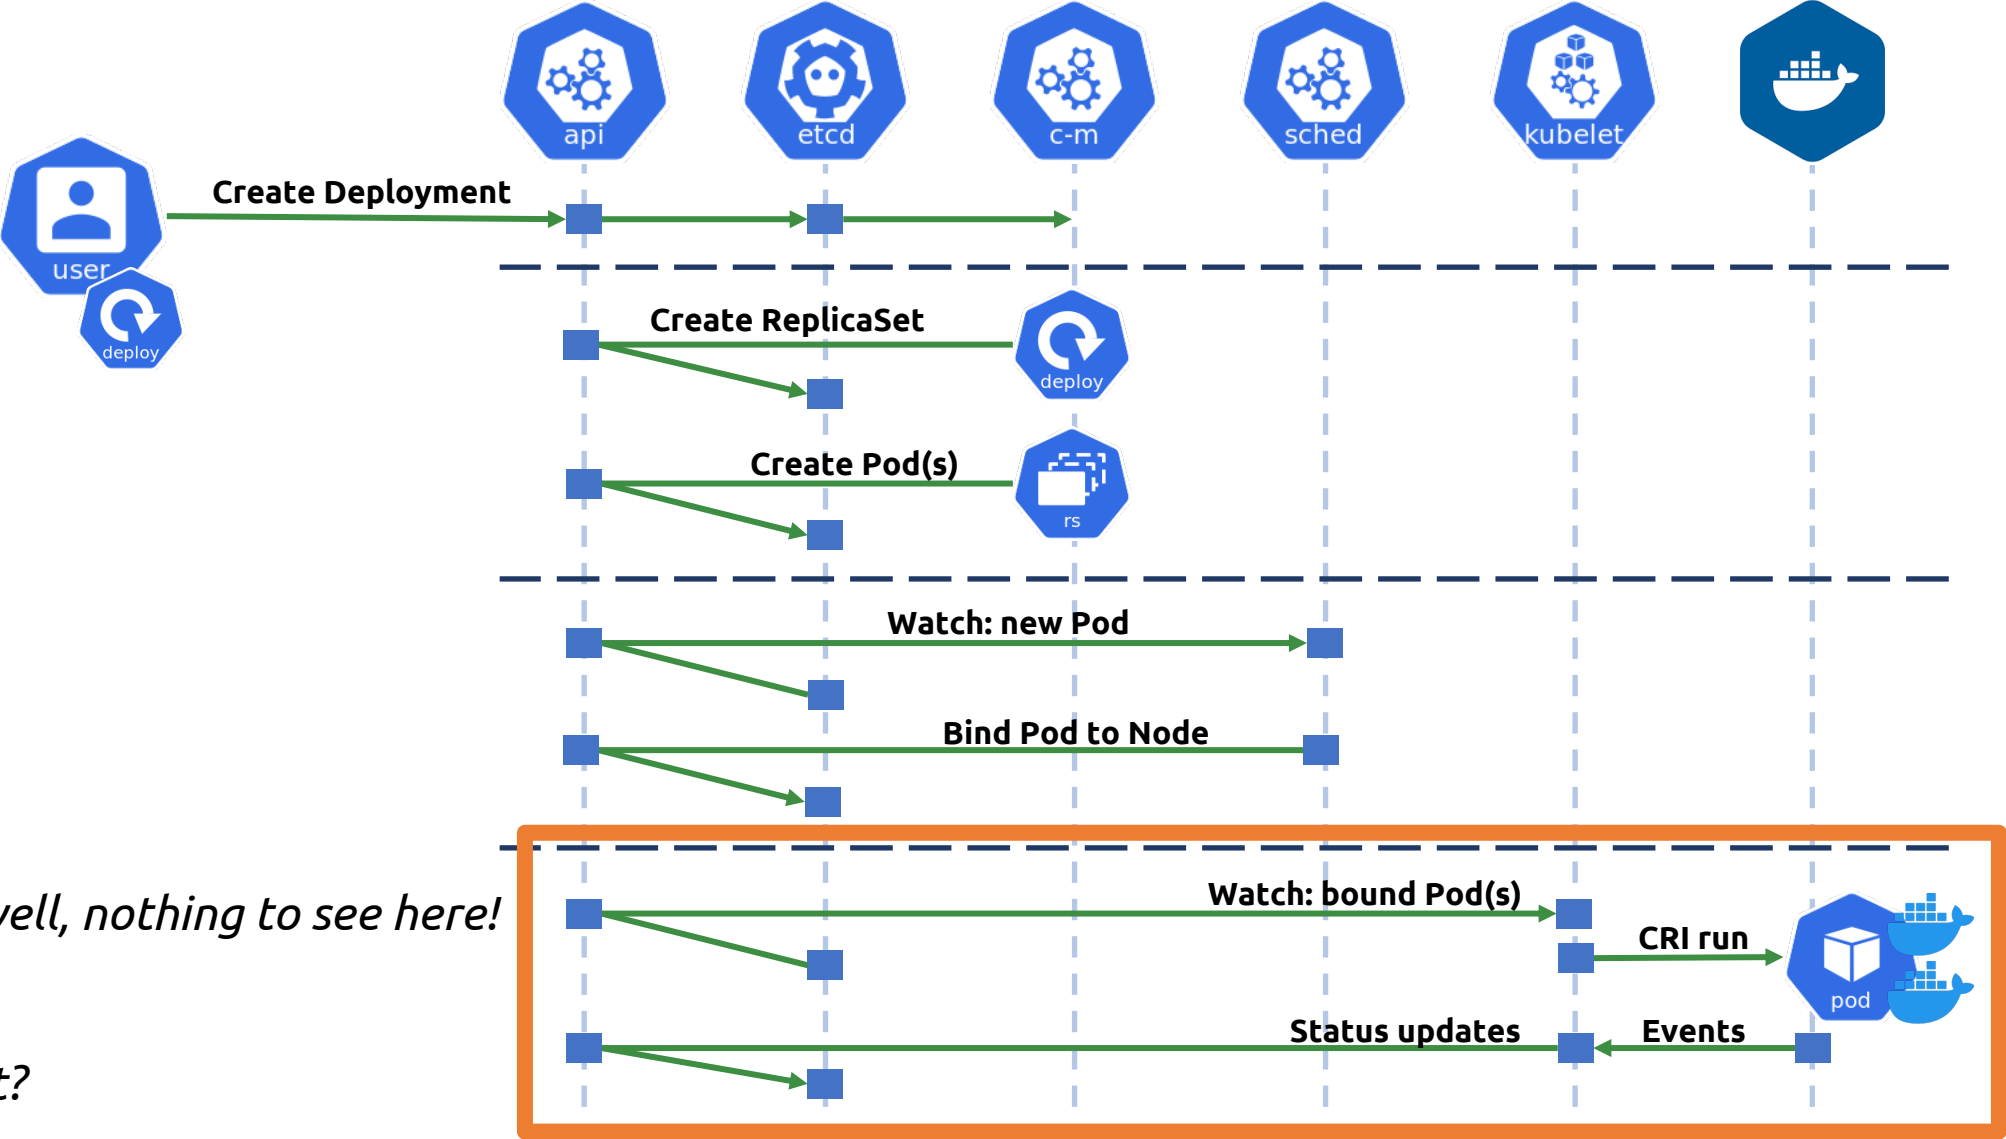
# Running: Examples



*All is well, nothing to see here!*



# Running: Examples



All is well, nothing to see here!

...right?



# Running (but not Ready): Symptoms

Controllers will have mismatches between desired/current/ready/available replicas

- Will look identical to other issues – ***may be misleading!***

`kubectl get pod` will return **Running** but **Ready** will not have equal values for current / desired

K8s will report the pod condition **PodScheduled** as "true"

K8s will report the pod conditions **Ready** and **ContainersReady** as "false"

- Condition messages will indicate which container(s) are problematic
  - Ex: 'containers with unready status: [container-names]'

Container states will vary based on cause

Related: Services will have fewer endpoints than pods that are running

# Running (but not Ready): Common Causes

**Probes** – one or more of the probes has failed

- Container state will be "**running**" but "**ready**" will be false

**Errors / Crashing / Restarting** – because  $\geq 1$  container is running, starting, or restarting the pod remains in the "Running" phase

- When happening frequently, Running is replaced by CrashLoopBackoff
  - Container state will be "**waiting**" and "**ready**" will be false

# Running and Ready (but not available)

## Symptoms:

Controllers will look normal—at first!

- After the kubelet heartbeat timestamp decays (~40 seconds), Ready and Available will mismatch with Desired
  - Will look identical to other issues – *may be misleading*

kubectl get pod returns **Running** and **Ready** until the pod is terminated—*will be misleading*

K8s will report the pod conditions **PodScheduled** and **ContainersReady** as "true"

K8s will report the pod condition **Ready** as "false" only *after* the node is considered not ready

## Common Cause:

**Node failure** – the Running phase will not change but the node controller will indicate that the pod is not ready

- Status updates no longer being delivered by the kubelet (which is down)



KubeCon



CloudNativeCon

Europe 2022

# Other Examples

When a pod's phase does not reflect:

- Pending
- Running
- Succeeded
- Failed
- Unknown



# Pods with other Statuses

Even though a pod's `status.phase` will always be one of:

- Pending
- Running
- Succeeded
- Failed
- Unknown

`kubectl get pod` will show the value of `.status.containerStatuses[*].state.*.reason`

Pending:

- ContainerCreating
- CreateContainerConfigError
- ErrImagePull / ImagePullBackOff

Running:

- Error / CrashLoopBackOff

To help with debugging commonly occurring issues



Now you know what kubectl said!



# Thank You!

**Cloud Native Short Takes:** ~5 minute tips on a myriad of subjects: K8s, containers, Golang, etc. on the RX-M YouTube channel

**Certification Training:** KCNA / CKAD / CKA / CKS – self-study & virtual/in-person instructor-led training by a certified instructor

**RX-M Cloud Native Expertise:**  
**Training – Consulting / Advisory – Staffing**

- Large breadth of cloud native / open source / digital transformation training curriculum
- Custom consultative training programs aligned to job roles and business goals

Contact us at: [rx-m.com](https://rx-m.com) / [info@rx-m.com](mailto:info@rx-m.com)

