Grafana Labs

# Prometheus Sparse High-Resolution Histograms in Action

Ganesh Vernekar

KubeCon EU, May 19 2022

# About Me


## Ganesh Vernekar

Senior Software Engineer at Grafana Labs

Prometheus Team Member

Prometheus TSDB Maintainer

@_codesome

# What is a histogram?

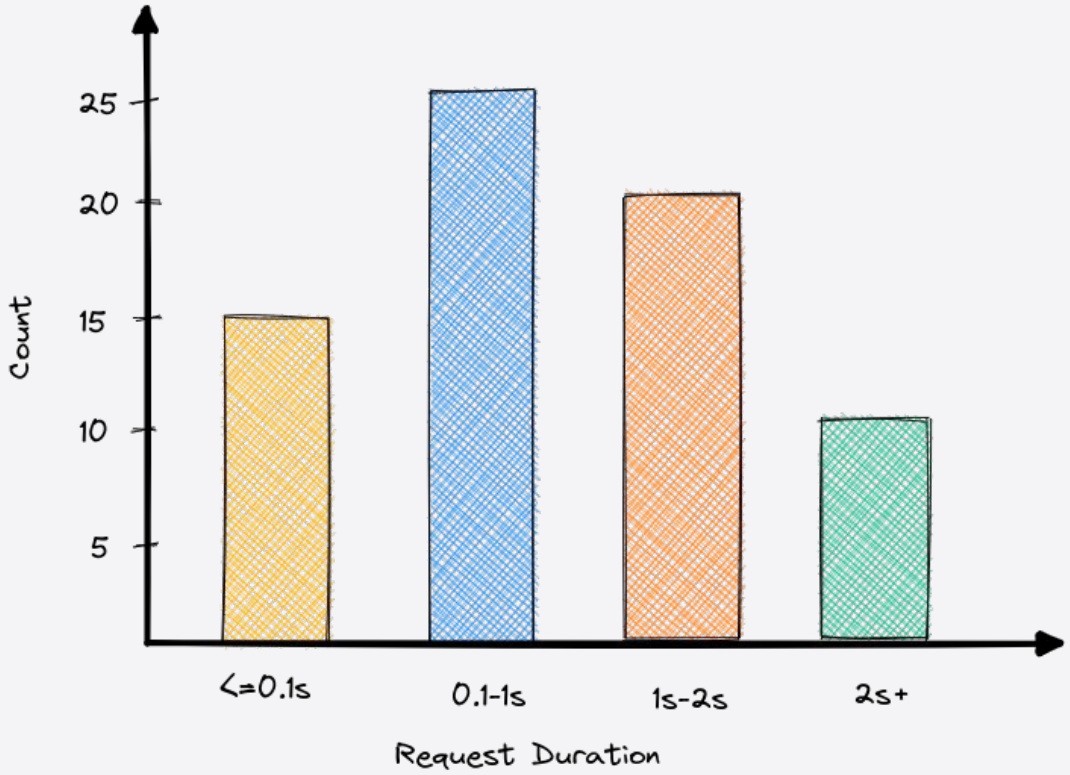

# This is how it is stored in Prometheus

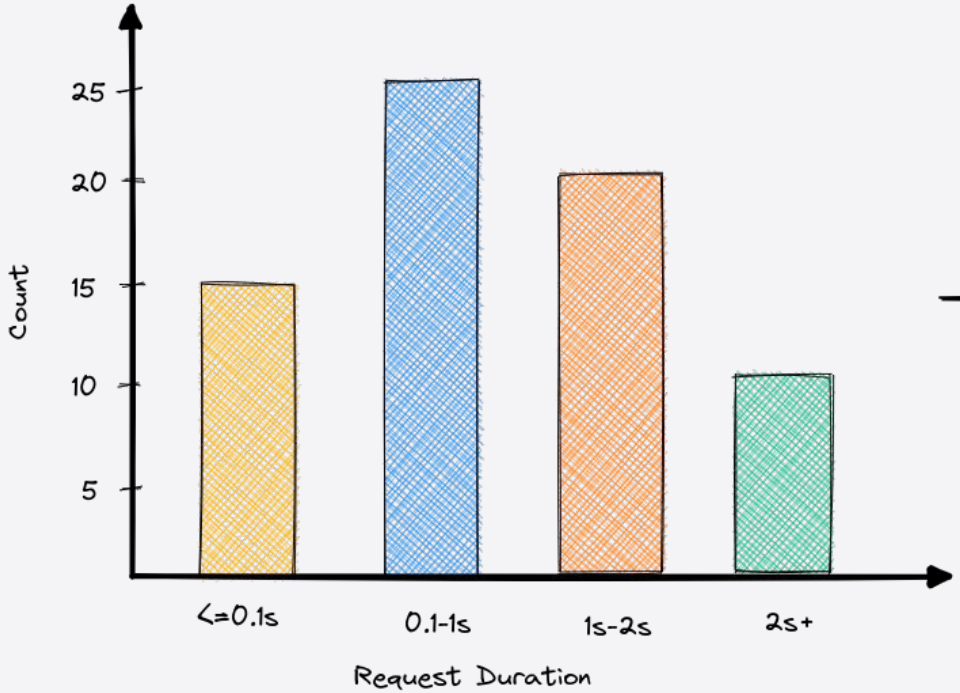


```
series_bucket{le="0.1"}  15
series_bucket{le="1.0"}  40
series_bucket{le="2.0"}  60
series_bucket{le="+Inf"} 70
series_count 70
series_sum 60.75
```

# Problem with this design

Pre-defined, cumulative, non-sparse, buckets.

Got the buckets wrong? Re-instrument and re-deploy it everywhere

```
series_bucket{le="0.1"}  15
series_bucket{le="0.25"} 15
series_bucket{le="0.5"}  15
series_bucket{le="0.75"} 15
series_bucket{le="1.0"}  40
series_bucket{le="1.25"} 40
series_bucket{le="1.50"} 40
series_bucket{le="1.75"} 40
series_bucket{le="2.0"}  60
series_bucket{le="+Inf"} 70
series_count 70
series_sum 60.75
```

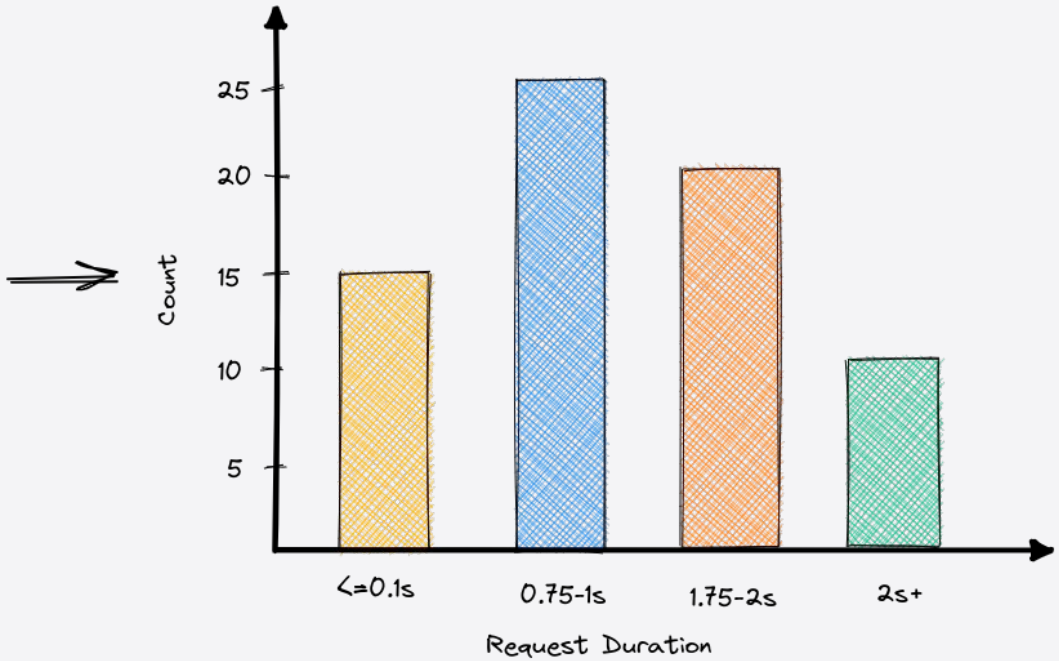

# Problem with this design

Changing bucket layout is painful, can’t correlate anymore in many cases

```
series_bucket{le="0.1"}  15
series_bucket{le="0.25"} 15
series_bucket{le="0.5"}  15
series_bucket{le="0.75"} 15
series_bucket{le="1.0"}  40
series_bucket{le="1.25"} 40
series_bucket{le="1.50"} 40
series_bucket{le="1.75"} 40
series_bucket{le="2.0"}  60
series_bucket{le="+Inf"} 70
series_count 70
series_sum 60.75
```

```
series_bucket{le="0.33"}  15
series_bucket{le="0.66"} 25
series_bucket{le="1.0"}  40
series_bucket{le="1.33"} 45
series_bucket{le="1.66"} 52
series_bucket{le="2.0"}  60
series_bucket{le="+Inf"} 70
series_count 70
series_sum 60.75
```

# Problem with this design

`total_series_per_histogram = num_buckets + 3`

```
series_bucket{le="0.33"}  15
series_bucket{le="0.66"} 25
series_bucket{le="1.0"}  40
series_bucket{le="1.33"} 45
series_bucket{le="1.66"} 52
series_bucket{le="2.0"}  60
series_bucket{le="+Inf"} 70
series_count 70
series_sum 60.75
```

# Here comes the sparse high resolution histograms

Let’s build it step by step

@_codesome





Result of
multi year study and research

# First property: Fixed bucket boundaries

Factor of $2^1$

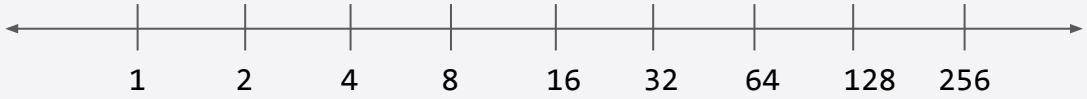


$2^2$, same as merging consecutive buckets from above

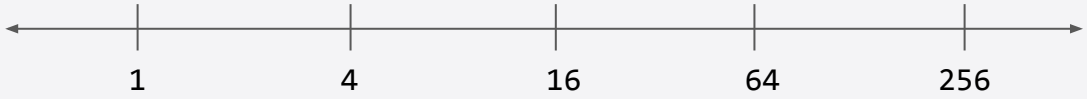


$2^4$

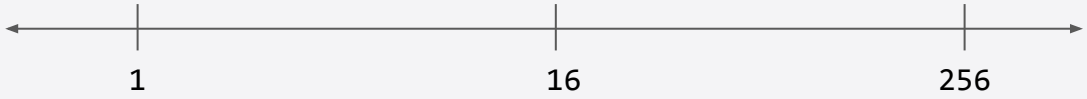


$2^8$

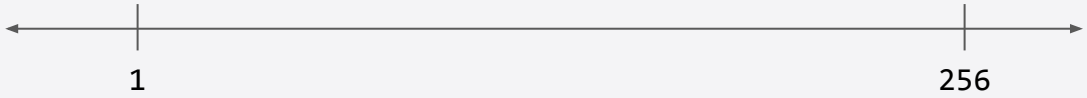


# Fixed bucket boundaries, higher resolution

Factor of $2^1$

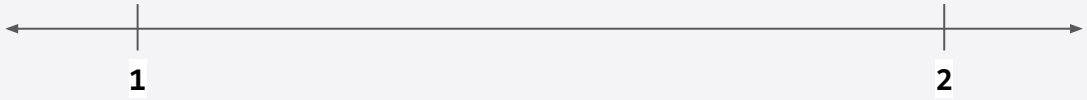


$2^{1/2}$, gives 2 buckets between 1 and 2.

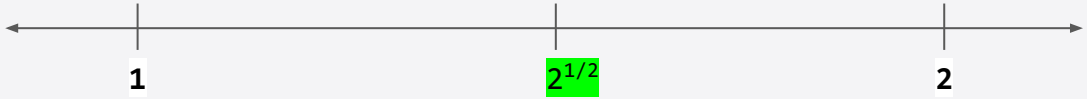


$2^{1/4}$, gives 2 buckets between $1 \leftrightarrow 2^{1/2}$, and $2^{1/2} \leftrightarrow 2$

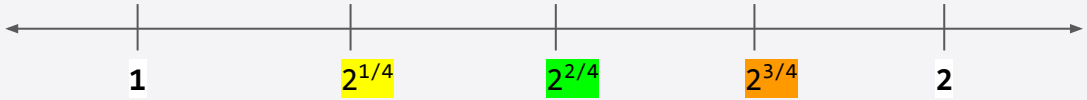


$2^{1/8}$, same as above, 2 buckets between above boundaries

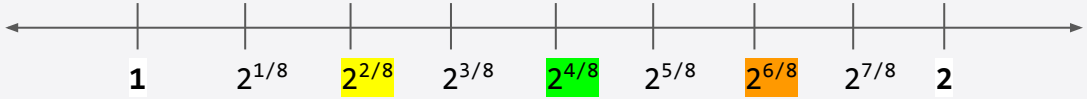


# Fixed bucket boundaries

**$2^1$:**

…| 0.0625 | 0.125 | 0.25 | 0.5 | **1** | 2 | 4 | 8 | 16 |…

**$2^2$:**

…| 0.0625 | 0.25 | **1** | 4 | 16 | 64 |…

**$2^{1/2}$:**

…| **0.25** | $0.5/2^{1/2}$ | **0.5** | $1/2^{1/2}$ | **1** | $2^{1/2}$ | **2** | $2*2^{1/2}$ | **4** | $4*2^{1/2}$ | **8** |…

**$2^{1/4}$:**

…| **0.5** | $1/2^{3/4}$ | $1/2^{2/4}$ | $1/2^{1/4}$ | **1** | $2^{1/4}$ | $2^{2/4}$ | $2^{3/4}$ | **2** | $2*2^{1/4}$ | $2*2^{2/4}$ | $2*2^{3/4}$ | **4** |…

@_codesome

# Why is it like this?

# Convert to a lower resolution

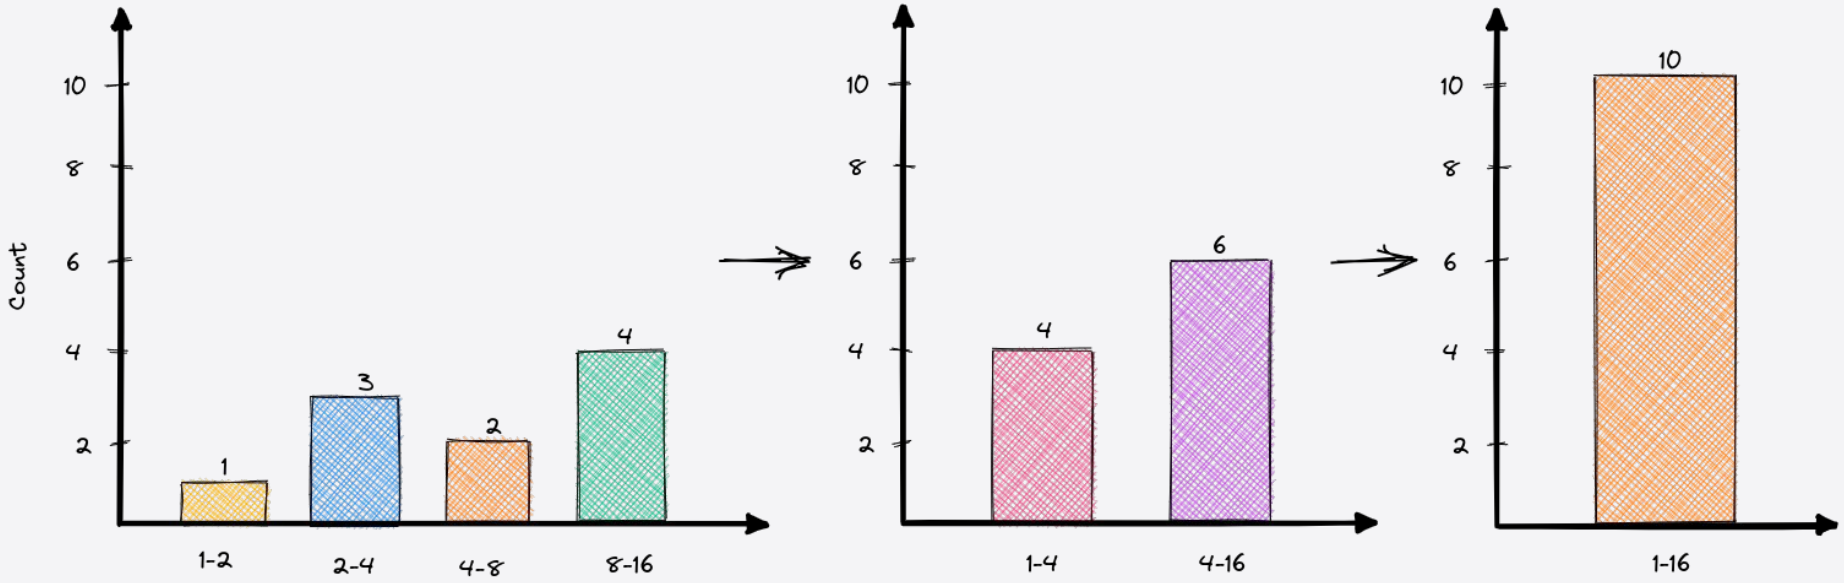

# Combine different resolutions

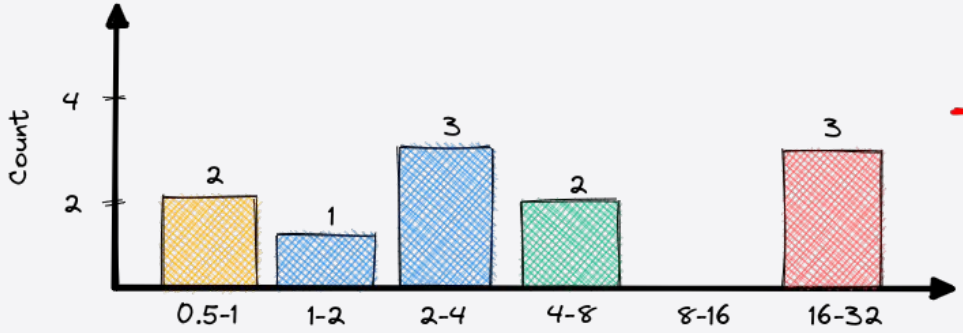


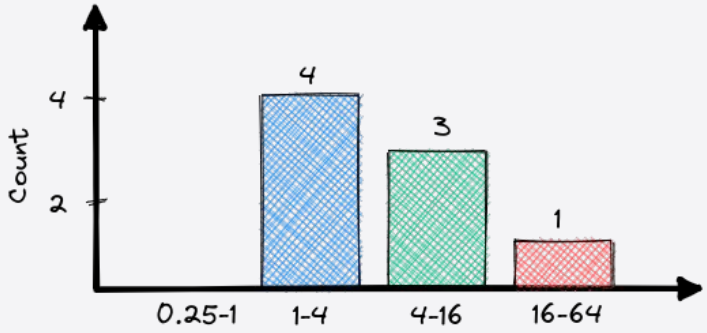


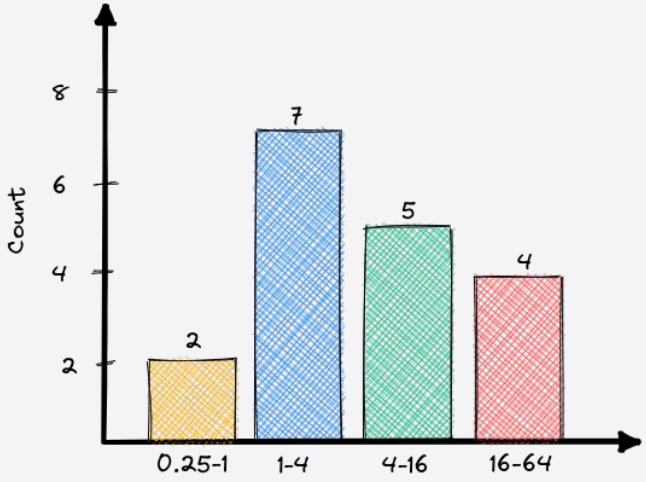

# Unique ID per bucket

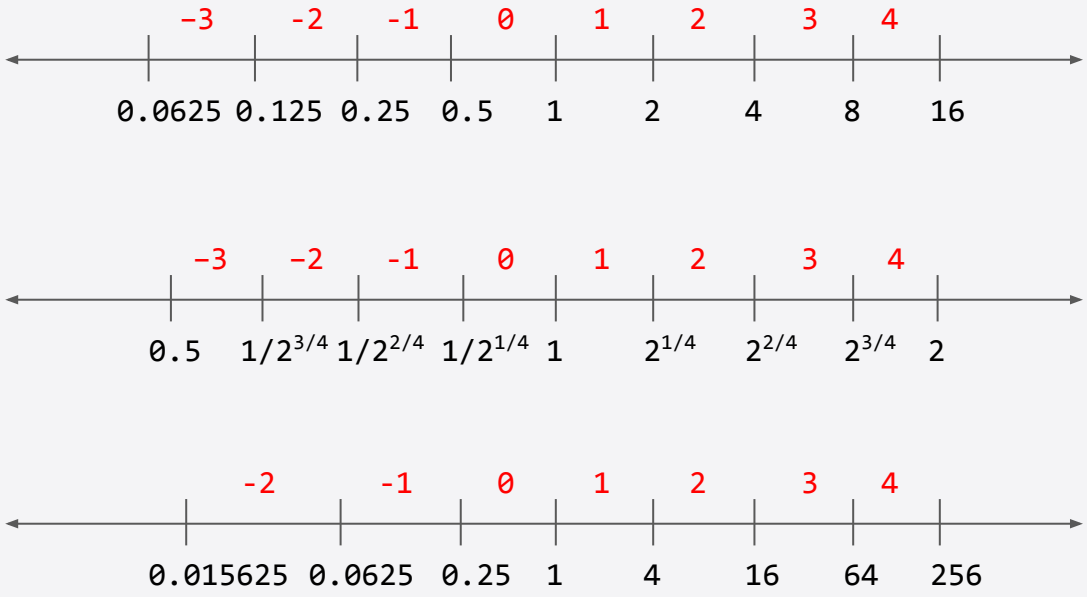


# Finally...encoding these histograms

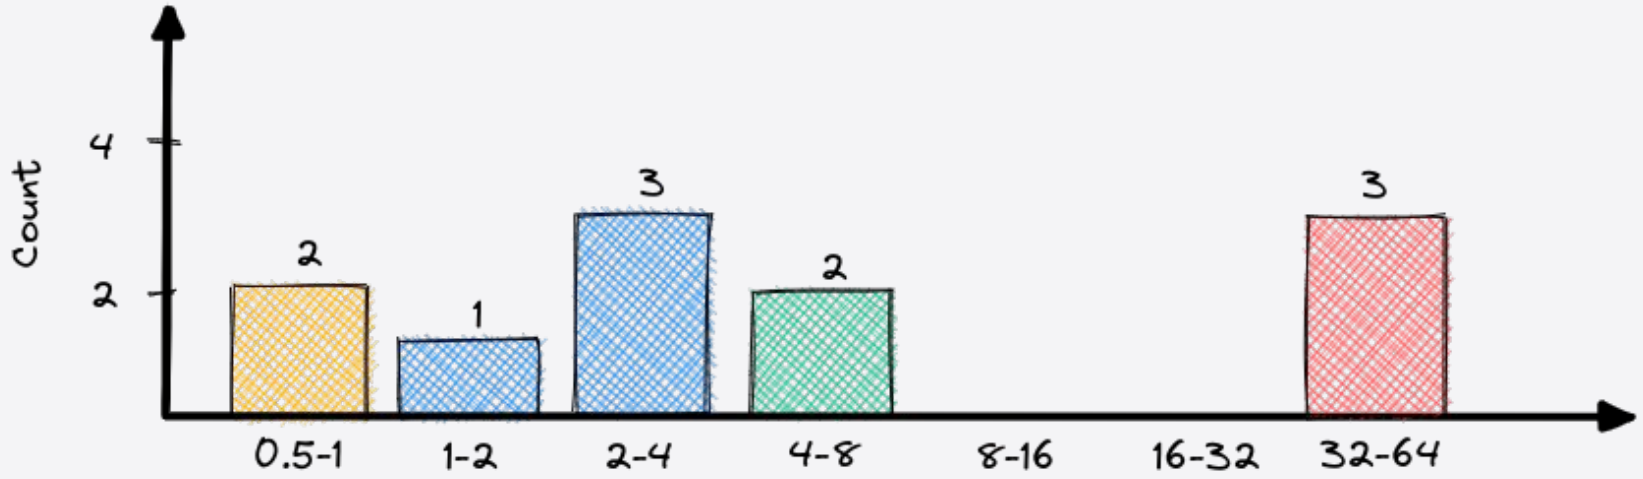


```
<metadata> // Resolution, sum, count, etc

[ (0, 4), (2, 1) ] // “Spans”, bucket layout

[ 2, 1, 3, 2, 3 ] // The bucket values
```

# One time series per histogram

```
my_histogram{foo=”bar”} =>
```

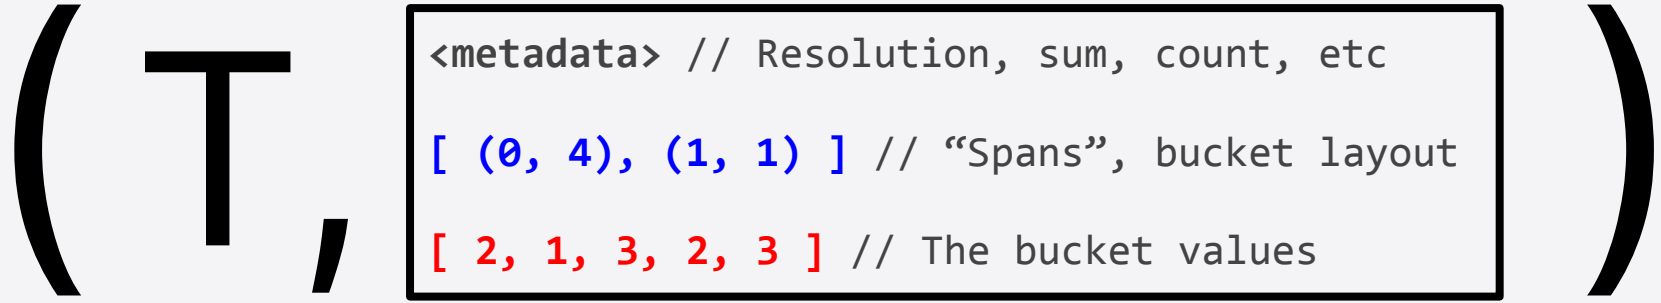


# How do I instrument this?

# Instrumentation

```
mediumRes = promauto.With(reg).NewHistogram(prometheus.HistogramOpts{
        Name:                "kubeconeu2022_demo",
        Help:                "Values obs      during the demo.",
        SparseBucketsFactor: 2,
        ConstLabels:         map[st                 e": "med_res"},
})
lowRes = promauto.With(reg).NewHist                 s.HistogramOpts{
        Name:                "kubecone              ",
        Help:                "Values observed during the demo.",
        SparseBucketsFactor: 4,
        ConstLabels:         map[string]string{"type": "low_res"},
})
```

# Scraping in this PoC

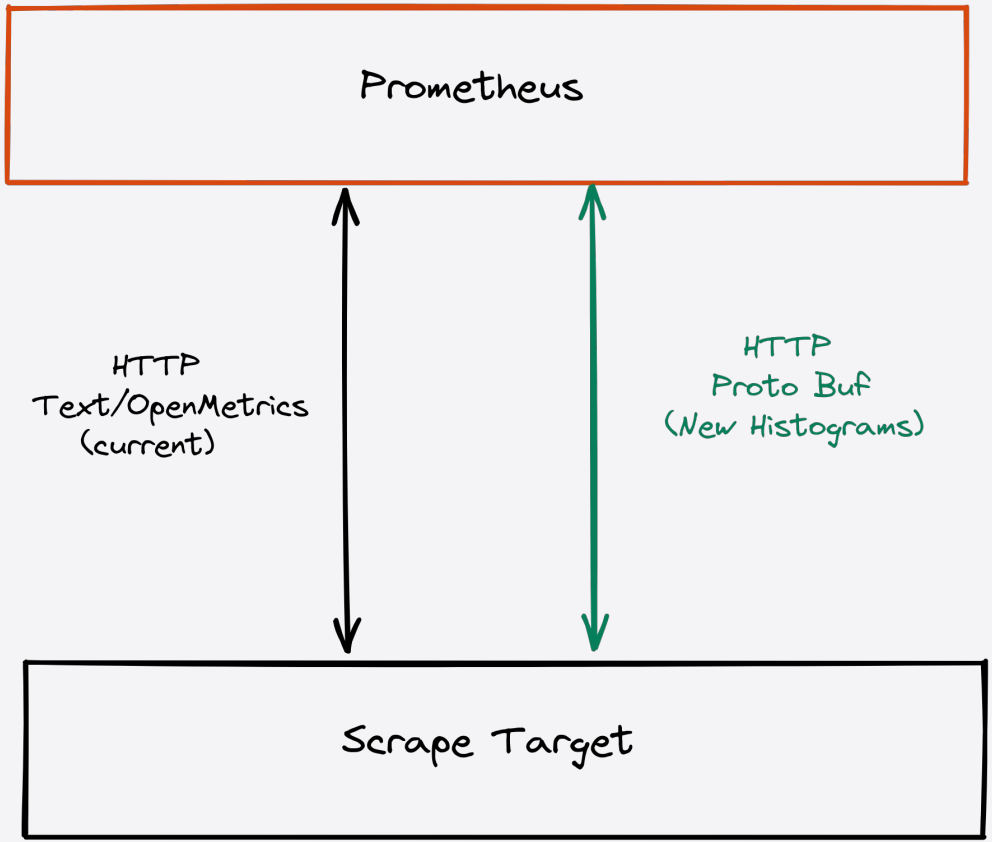

# DEMO TIME!

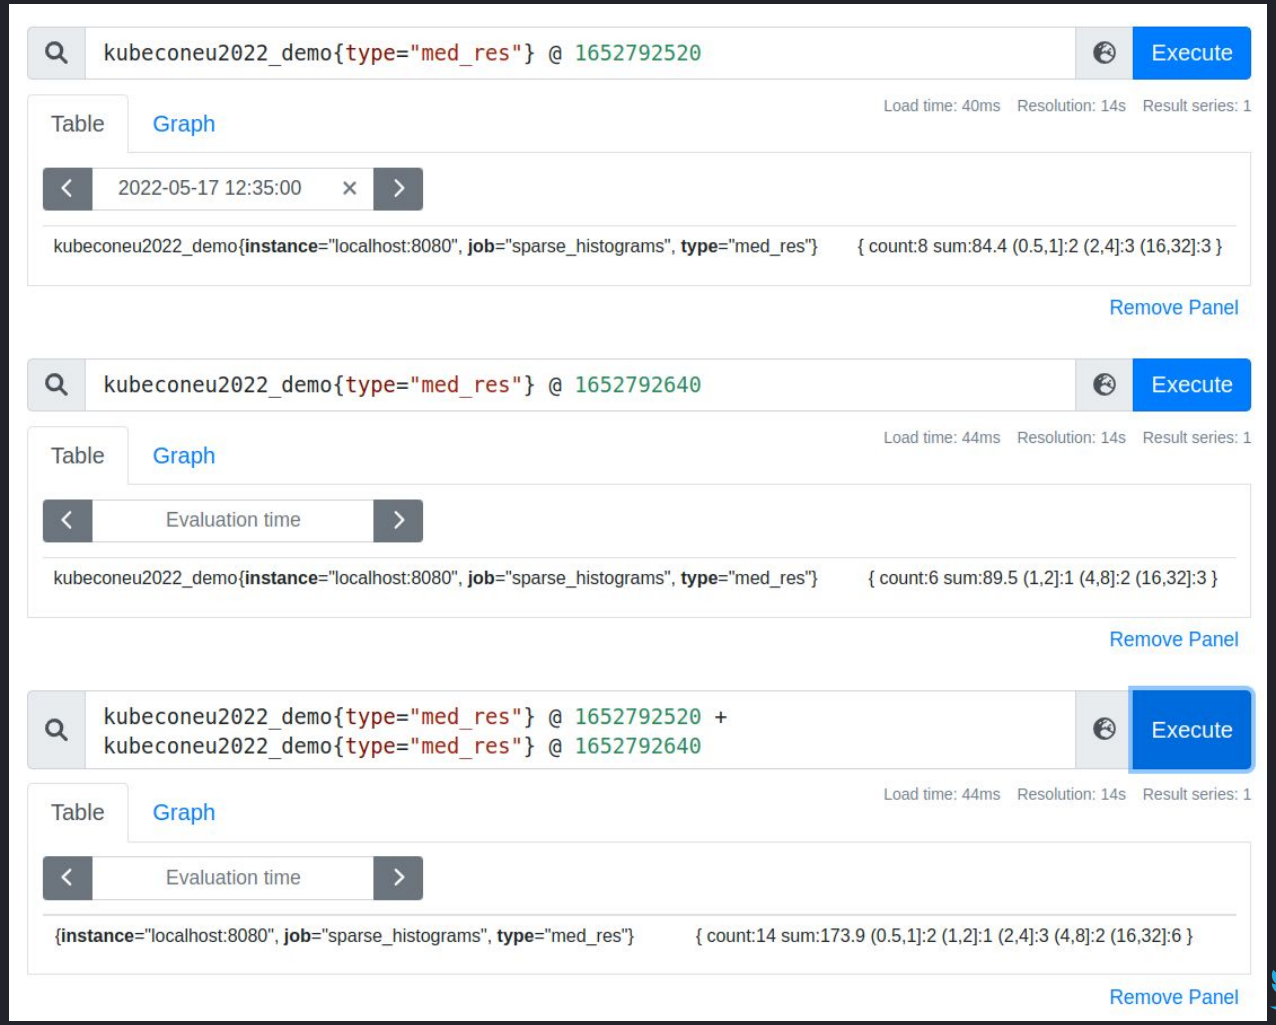
kubeconeu2022_demo{type="med_res"} @ 1652792520
Execute
Load time: 40ms Resolution: 14s Result series: 1
Table
Graph
2022-05-17 12:35:00
kubeconeu2022_demo{instance="localhost:8080", job="sparse_histograms", type="med_res"} { count:8 sum:84.4 (0.5,1]:2 (2,4]:3 (16,32]:3 }
Remove Panel
kubeconeu2022_demo{type="med_res"} @ 1652792640
Execute
Load time: 44ms Resolution: 14s Result series: 1
Table
Graph
Evaluation time
kubeconeu2022_demo{instance="localhost:8080", job="sparse_histograms", type="med_res"} { count:6 sum:89.5 (1,2]:1 (4,8]:2 (16,32]:3 }
Remove Panel
kubeconeu2022_demo{type="med_res"} @ 1652792520 +
kubeconeu2022_demo{type="med_res"} @ 1652792640
Execute
Load time: 44ms Resolution: 14s Result series: 1
Table
Graph
Evaluation time
{instance="localhost:8080", job="sparse_histograms", type="med_res"} { count:14 sum:173.9 (0.5,1]:2 (1,2]:1 (2,4]:3 (4,8]:2 (16,32]:6 }
Remove Panel

kubeconeu2022_demo @ 1652792760

Execute

Load time: 58ms Resolution: 14s Result series: 2

Table Graph

2022-05-17 12:35:00

| | |
|---|---|
| kubeconeu2022_demo{**instance**="localhost:8080", **job**="sparse_histograms", **type**="low_res"} | { count:8 sum:61.33 (1,4]:4 (4,16]:3 (16,64]:1 } |
| kubeconeu2022_demo{**instance**="localhost:8080", **job**="sparse_histograms", **type**="med_res"} | { count:11 sum:98.9 (0.5,1]:2 (1,2]:1 (2,4]:3 (4,8]:2 (16,32]:3 } |

Remove Panel

sum(kubeconeu2022_demo @ 1652792760)

Execute

Load time: 63ms Resolution: 14s Result series: 1

Table Graph

Evaluation time

| | |
|---|---|
| {} | { count:19 sum:160.23000000000002 (0.25,1]:2 (1,4]:8 (4,16]:5 (16,64]:4 } |

Remove Panel

kubeconeu2022_demo

Execute

Table | Graph

Load time: 12ms Resolution: 14s Result series: 2

Evaluation time

| | |
|---|---|
| kubeconeu2022_demo{**instance**="localhost:8080", **job**="sparse_histograms", **type**="low_res"} | { count:517224 sum:3965168.4900032305 (1,4]:258612 (4,16]:193959 (16,64]:64653 } |
| kubeconeu2022_demo{**instance**="localhost:8080", **job**="sparse_histograms", **type**="med_res"} | { count:711183 sum:6394181.6999964 (0.5,1]:129306 (1,2]:64653 (2,4]:193959 (4,8]:129306 (16,32]:193959 } |

Remove Panel

rate(kubeconeu2022_demo[1m])

Execute

Table | Graph

Load time: 33ms Resolution: 14s Result series: 2

Evaluation time

| | |
|---|---|
| {**instance**="localhost:8080", **job**="sparse_histograms", **type**="low_res"} | { count:390.10169491525426 sum:2990.617118647693 (1,4]:195.05084745762713 (4,16]:146.28813559322035 (16,64]:48.76271186440678 } |
| {**instance**="localhost:8080", **job**="sparse_histograms", **type**="med_res"} | { count:536.3898305084746 sum:4822.632203362598 (0.5,1]:97.52542372881356 (1,2]:48.76271186440678 (2,4]:146.28813559322035 (4,8]:97.52542372881356 (16,32]:146.28813559322035 } |

Remove Panel


histogram_quantile(0.99, rate(kubeconeu2022_demo[1m]))
Execute
Table
Graph
Load time: 45ms
Resolution: 14s
Result series: 2
Evaluation time
{instance="localhost:8080", job="sparse_histograms", type="low_res"}
60.16000000000001
{instance="localhost:8080", job="sparse_histograms", type="med_res"}
31.413333333333334
Remove Panel

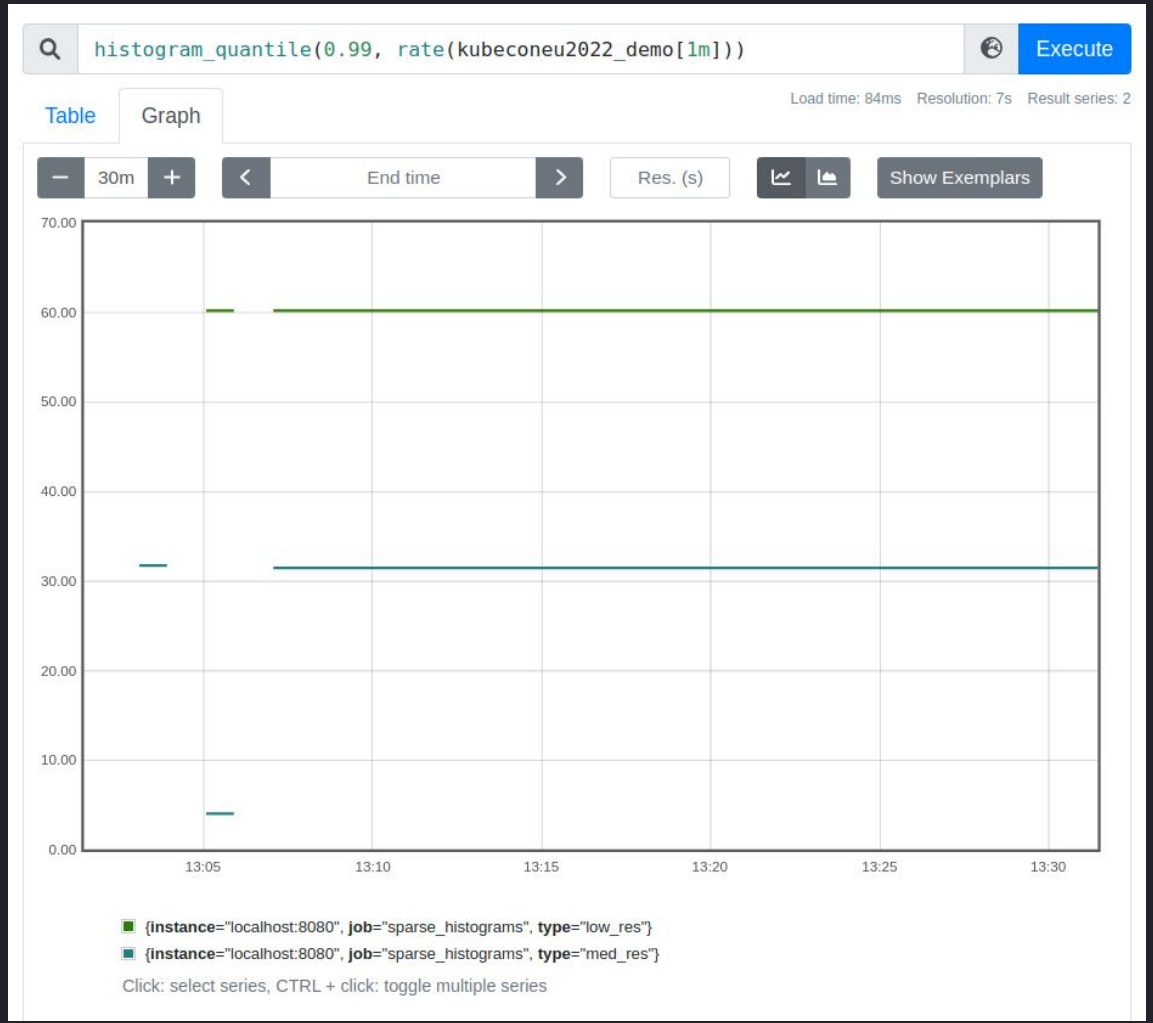
histogram_quantile(0.99, rate(kubeconeu2022_demo[1m]))
Execute
Table
Graph
Load time: 84ms
Resolution: 7s
Result series: 2
30m
End time
Res. (s)
Show Exemplars
70.00
60.00
50.00
40.00
30.00
20.00
10.00
0.00
13:05
13:10
13:15
13:20
13:25
13:30
{instance="localhost:8080", job="sparse_histograms", type="low_res"}
{instance="localhost:8080", job="sparse_histograms", type="med_res"}
Click: select series, CTRL + click: toggle multiple series


@_codesome

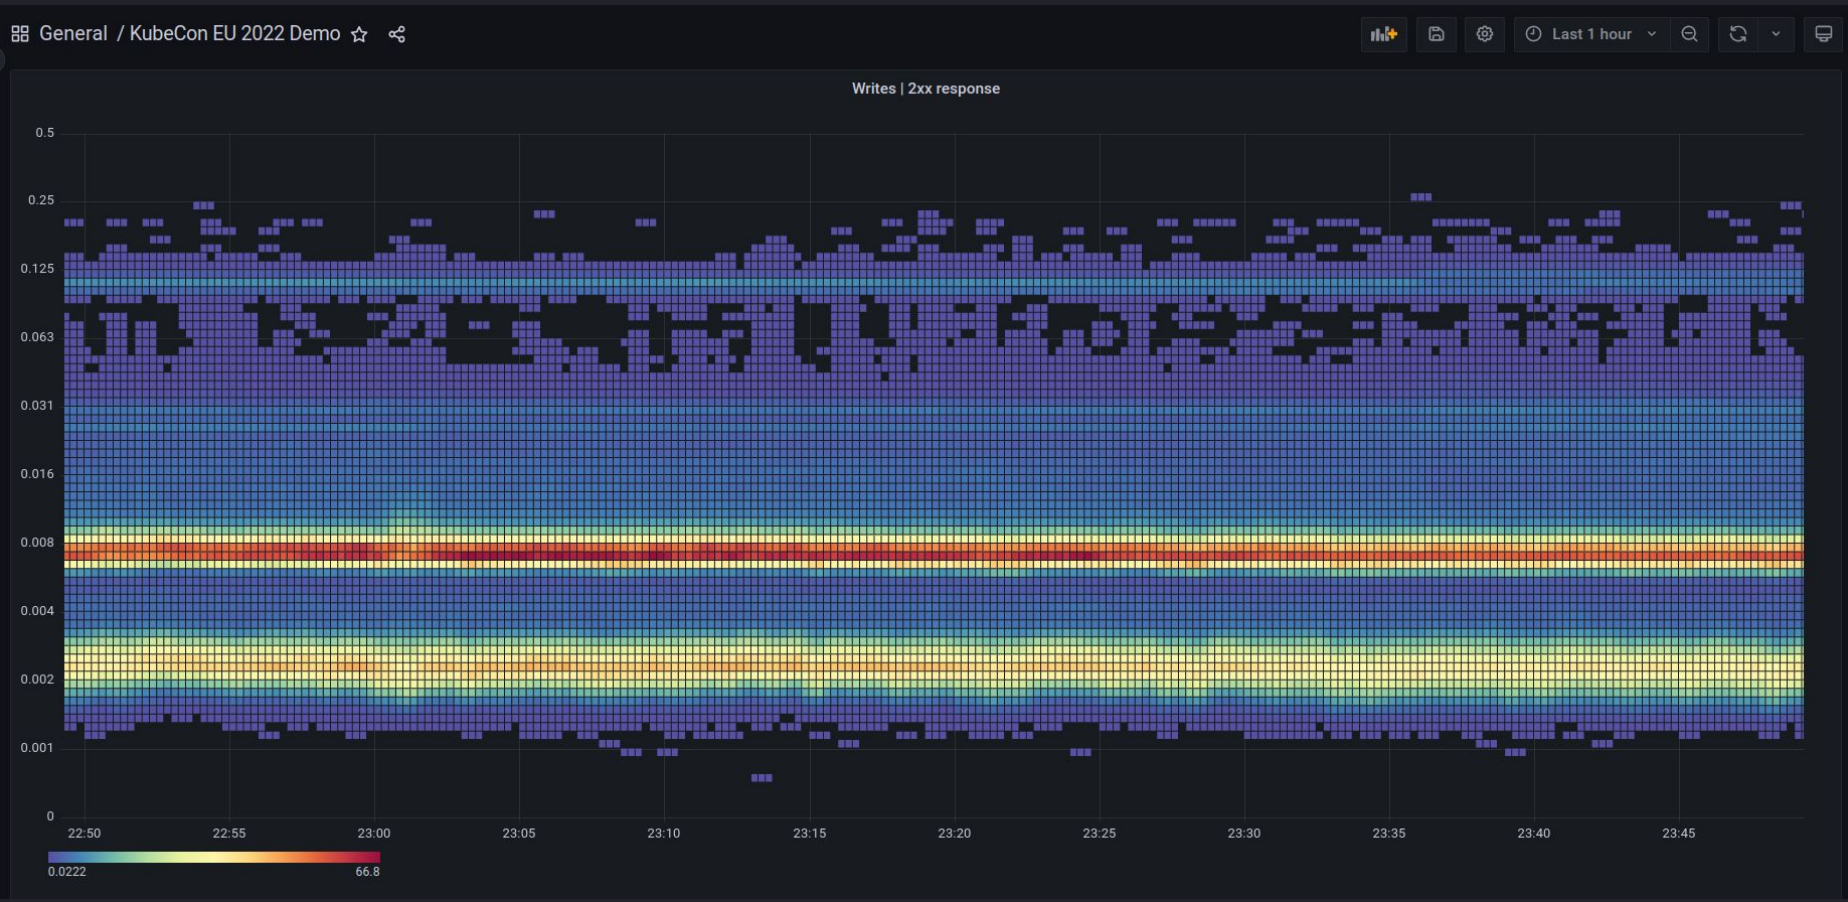
General / KubeCon EU 2022 Demo
Last 1 hour
Writes | 2xx response
0.5
0.25
0.125
0.063
0.031
0.016
0.008
0.004
0.002
0.001
0
22:50
22:55
23:00
23:05
23:10
23:15
23:20
23:25
23:30
23:35
23:40
23:45
0.0222
66.8

@_codesome


General / KubeCon EU 2022 Demo
Last 6 hours
Writes | 2xx response
8
4
2
1
0.5
0.25
0.125
0.063
0.031
0.016
0.008
0.004
0.002
0.001
0
18:00
18:15
18:30
18:45
19:00
19:15
19:30
19:45
20:00
20:15
20:30
20:45
21:00
21:15
21:30
21:45
22:00
22:15
22:30
22:45
23:00
23:15
23:30
23:45
0.0222
66.8

@_codesome


General / KubeCon EU 2022 Demo - Storyteller
2022-05-17 22:57:30 to 2022-05-17 23:07:20
Story of latencies
1
0.5
0.25
0.125
0.063
0.031
0.016
22:58:00
22:58:30
22:59:00
22:59:30
23:00:00
23:00:30
23:01:00
23:01:30
23:02:00
23:02:30
23:03:00
23:03:30
23:04:00
23:04:30
23:05:00
23:05:30
23:06:00
23:06:30
23:07:00
0.0345
8.62

# How can I use it?

It is still in a Proof of Concept stage (but likely gonna stick).

Available here:

Instrumentation: https://github.com/prometheus/client_golang/tree/sparsehistogram

Prometheus server: https://github.com/prometheus/prometheus/tree/sparsehistogram

New heatmaps: main branch of grafana/grafana

Grafana Labs

# THANK YOU! Questions?