



KubeCon



CloudNativeCon

Europe 2022

WELCOME TO VALENCIA





KubeCon



CloudNativeCon

Europe 2022

Rook Intro & Ceph Deep Dive

Rook maintainers

Blaine Gardner, Red Hat

Travis Nielsen, Red Hat

Satoru Takeuchi, Cybozu, Inc.

Sébastien Han, Red Hat



Agenda

- Kubernetes Storage Challenges
- What is Rook?
- What is Ceph?
- Rook key features
- Rook v1.9 new features
- Demo
- Q&A

Kubernetes Storage Challenges



- Kubernetes is a platform to manage distributed apps
 - Ideally stateless
- Reliance on external storage
 - Not portable
 - Deployment burden
 - Day 2 operations - who is managing the storage?
- Reliance on cloud provider managed services
 - Vendor lock-in



What is Rook?

What is Rook?



- Makes storage available inside your Kubernetes cluster
- Consume like any other K8s storage
 - Storage Classes, Persistent Volume Claims
- Kubernetes Operators and Custom Resource Definitions
- Automated management of Ceph
 - Deployment, configuration, upgrades
- Open Source (Apache 2.0)

Rook Resources



Website	https://rook.io/
Documentation	https://rook.io/docs/rook/v1.9/
Slack	https://rook-io.slack.com/
Contributions	https://github.com/rook/rook
Twitter	@rook_io
Community Meeting	https://github.com/rook/rook#community-meeting
Training Videos (new!)	https://kubebuyexample.com/ → Learning Paths → Storage for Kubernetes with Rook

What is Ceph?



- Open Source
- Scalable, fault-tolerant storage service
 - Block
 - Shared File System
 - Object (S3 compliant)
- Favors consistency
- First release in July 2012
- <https://ceph.io/>


“ceph” alopod →



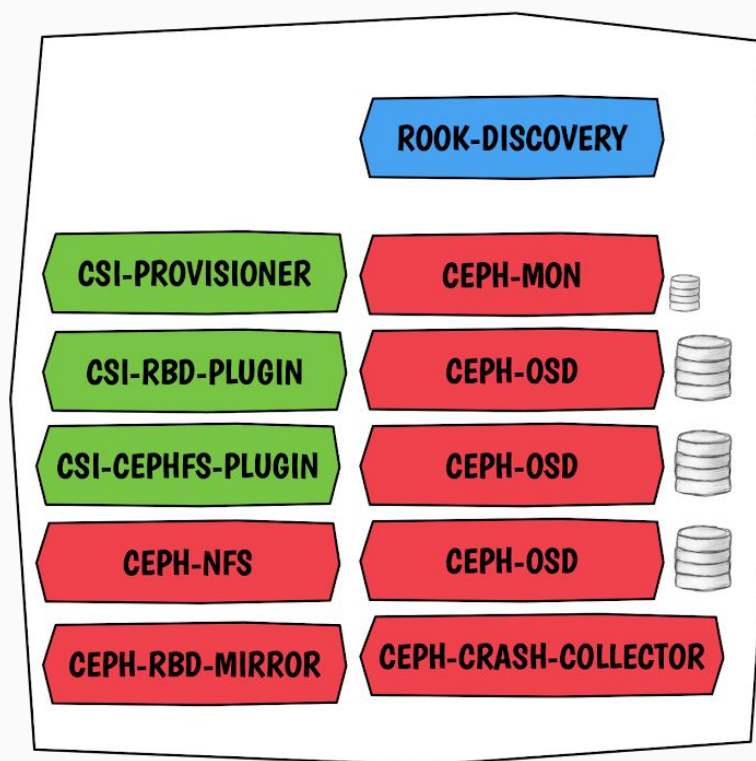
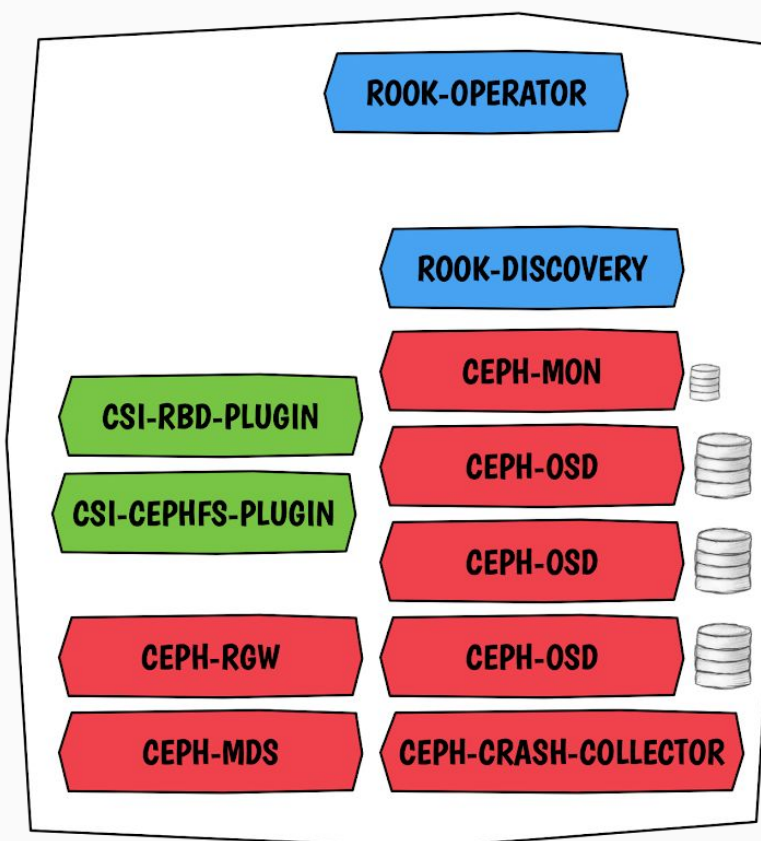
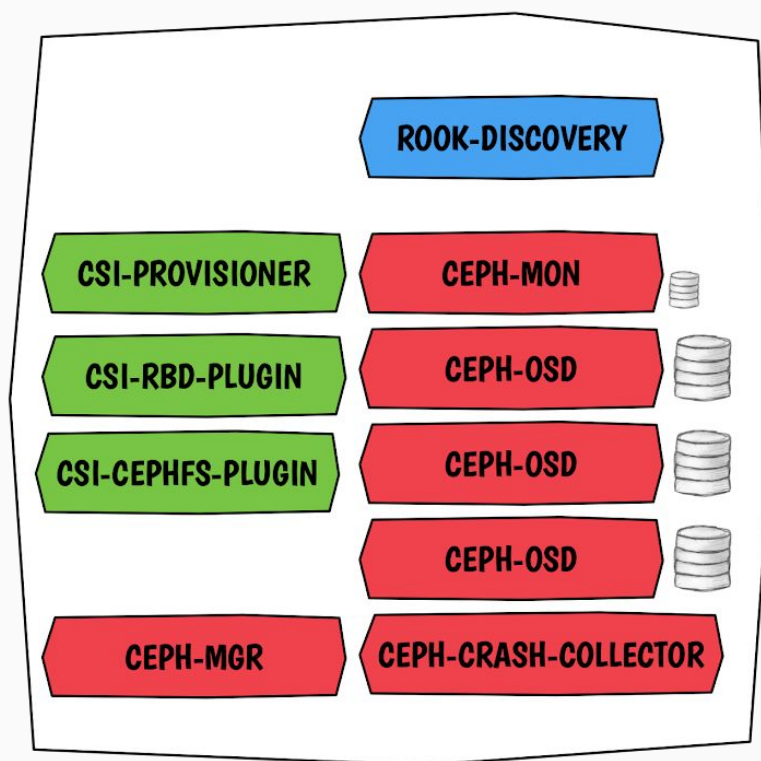
Architectural Layers



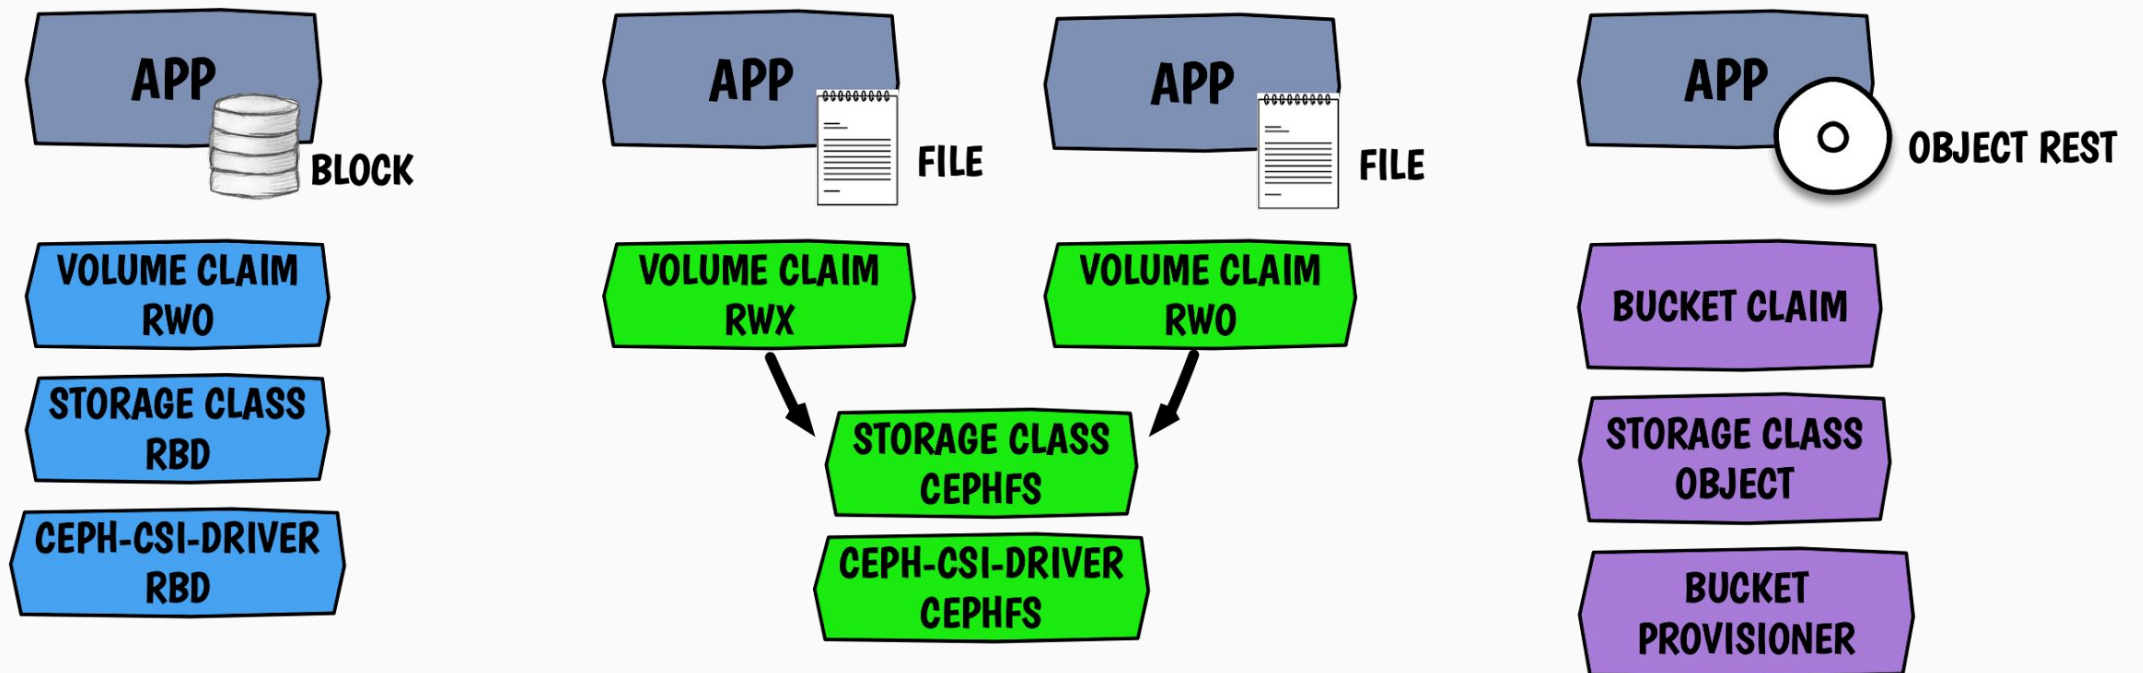
- Rook
 - Operator owns the **deployment** and **management** of Ceph and Ceph CSI (Container Storage Interface) driver
- Ceph-CSI
 - CSI driver dynamically **provisions** and **mounts** Ceph storage to user application Pods
- Ceph
 - **Data** layer



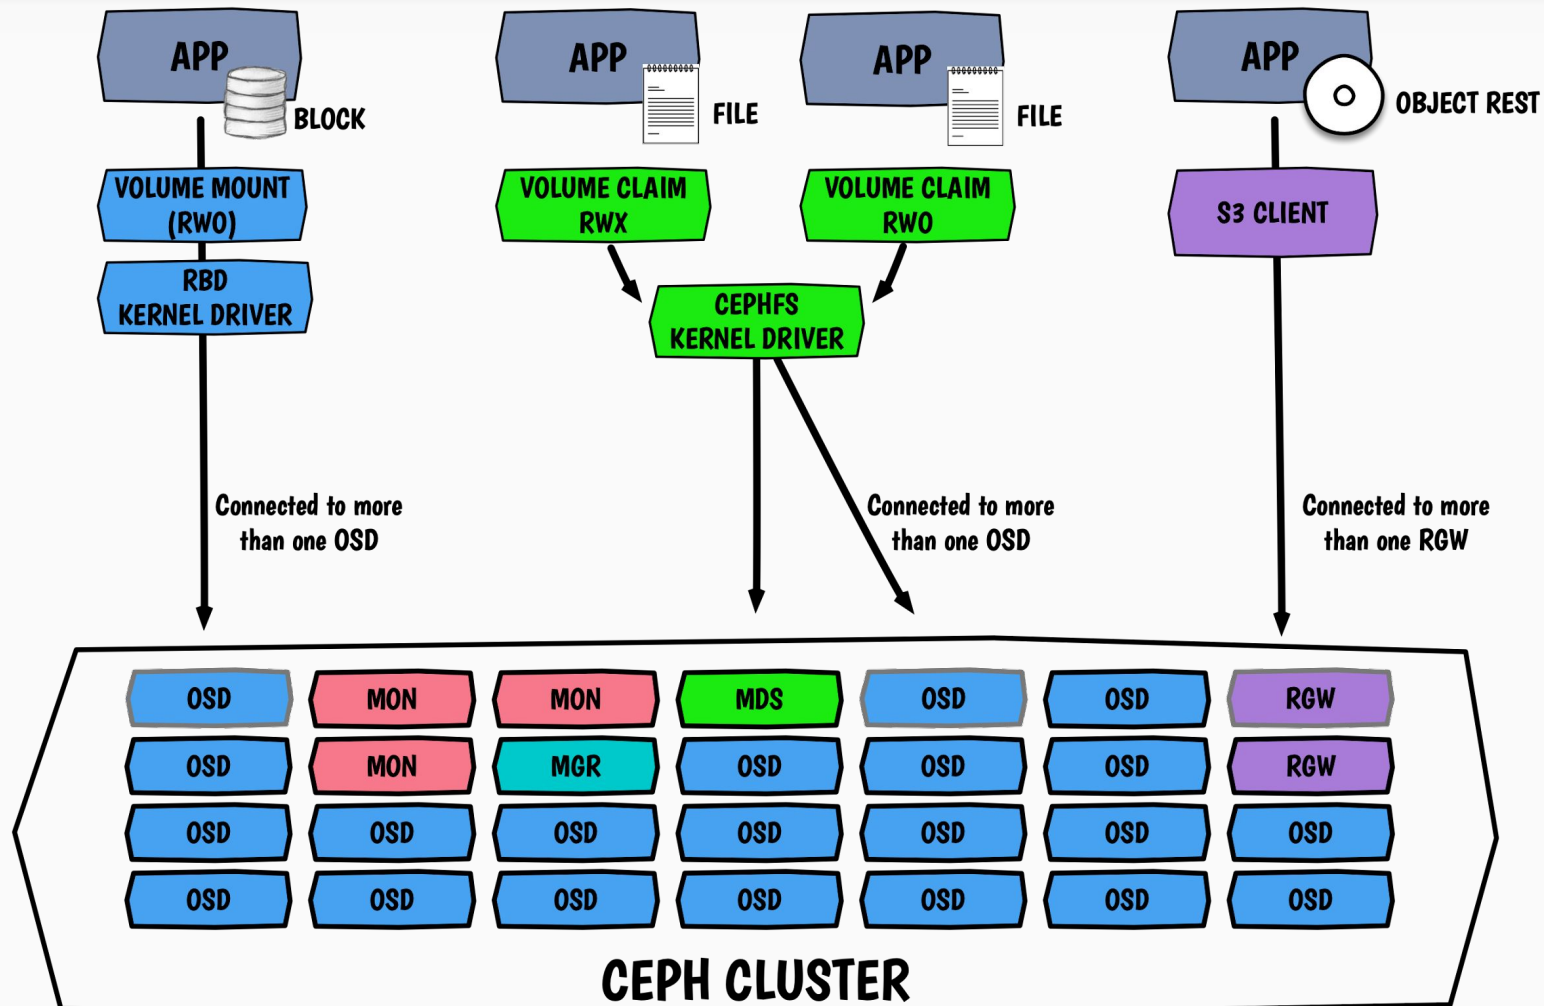
Rook: View of Pod Management



Provisioning storage with CSI



Ceph: Data path





Key Features

Installing Ceph is simple!




- Create Custom Resource Definitions
 - `kubectl create -f crds.yaml`
- Create authorization (RBAC)
 - `kubectl create -f common.yaml`
- Create the Rook-Ceph Operator
 - `kubectl create -f operator.yaml`
- Create the Ceph cluster resource
 - `kubectl create -f cluster.yaml`

```
apiVersion: ceph.rook.io/v1
kind: CephCluster
metadata:
  name: rook-ceph
  namespace: rook-ceph
spec:
  cephVersion:
    image: quay.io/ceph/ceph:v16.2.5
  mon:
    count: 3
  storage:
    useAllNodes: true
    useAllDevices: true
```


Ceph CSI driver features



- Dynamic provisioning for Block and File storage
- Volume expansion
- Snapshots and Clones



Environments



Bare metal

- Bring your own hardware

Cloud providers

- Overcome cloud provider storage limitations

Rook in a Cloud Environment



- Overcome shortcomings of the cloud provider's storage
 - Storage across availability zones (AZs)
 - Faster failover times (seconds instead of minutes)
 - Greater number of PVs per node (many more than ~30)
 - Use storage with better performance:cost ratio
- Consistent storage platform wherever K8s is deployed
- Ceph uses PVCs as underlying storage
 - No need for direct access to local devices

Configure for any cluster topology




- Customizable across/within cluster topologies
- High availability and durability
 - Spread Ceph daemons and data across failure domains
- Deployable on specific nodes if desired
 - Node affinity, taints/tolerations, etc.



Updates are automated



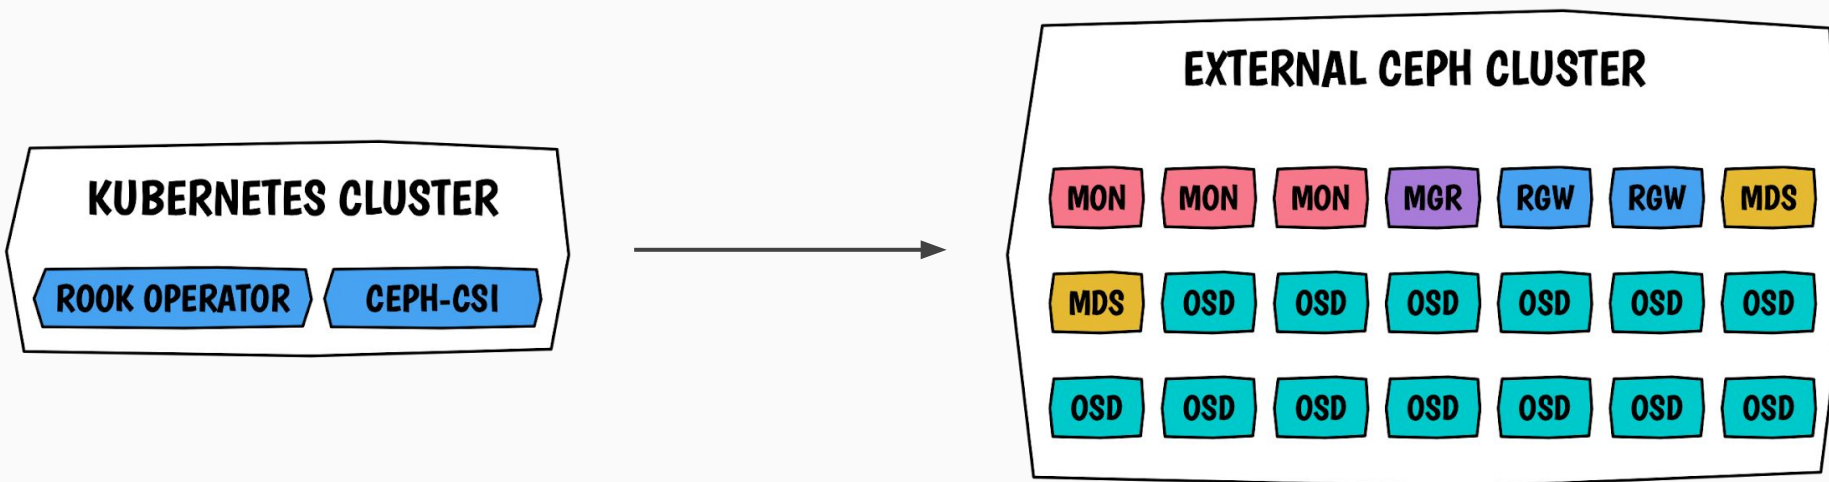
- Ceph updates and even major upgrades are fully automated!
 - Rook handles *everything*
- Rook patch *updates* are fully automated
- Rook minor *upgrades*
 - Take advantage of latest features
 - Occasional K8s/Ceph/CSI/Rook feature deprecations
 - <https://rook.io/docs/rook/latest/ceph-upgrade.html>



Connect to an external Ceph cluster



- Connect to a Ceph cluster outside of the current K8s cluster
- Dynamically create Block/File/Object storage consumable by K8s applications



Provision object storage buckets



- Define a Storage Class for Ceph object storage
- Create an Object Bucket Claim (OBC)
 - Similar pattern to a Persistent Volume Claim (PVC)
 - Rook operator creates a bucket when requested
 - Give access via K8s Secret
- Container Object Storage Interface (COSI)
 - Kubernetes Enhancement Proposal
 - CSI but for object storage







Rook v1.9 features

April 2022

Ceph Quincy Support




- Ceph Quincy (v17) is now supported
- Also released April 2022



CSI Driver Updates




- Ceph-CSI 3.6 release
- Fuse mount recovery: Detection of corrupt Ceph fuse mounts will be detected and remounted automatically
- AWS KMS encryption: CSI can be configured to use Amazon STS



NFS Provisioning




- Create NFS exports via PVCs
- Ceph-CSI driver provisioning
- Mount the volumes with the K8s community NFS driver



RADOS Namespace CRD




- Create RADOS namespaces in a pool
- Isolation/multi-tenancy without creating separate pools



Network Features




- Network features exposed with a simple CephCluster CR setting:
 - Encryption on the wire
 - Compression on the wire
- Recent kernel (5.11) is required



And Much More...



- Admission controller enabled by default if cert manager is available
- Multus networking support
- Updated Prometheus alerts
- ...





Demo

Environment

- OpenShift v4.9.15 (Kubernetes v1.22.3)
- 3 control nodes, 3 worker nodes
- Amazon Web Services m5.8xlarge nodes
 - Run storage with about ~50% room left over for user applications
- Using gp2 for backing volumes
- Rook v1.9.0
- Ceph v17.1.0 (pre-release)

Two types of Rook/Ceph clusters



- Host-based cluster
 - Use disks attached to a node for backing storage
- PVC-based cluster
 - Use Persistent Volume Claims to get backing storage
 - Can be dynamic or local volumes

Host-based cluster

- Suitable for simple cluster
- Storage configuration gets complicated when...
 - Not all nodes/devices are used
 - Using heterogeneous nodes
 - Customizing device layout per-node

```
# ...  
storage:  
  useAllNodes: true  
  useAllDevices: true
```

```
# ...  
storage:  
  nodes:  
    - name: "foo"  
    - devices:  
      - name: "sdb"  
  # ...
```

PVC-based cluster

- No need to describe hardware configuration
- Easy to expand
 - Increase the **count**
 - Increase the **resources.storage** size

```
# ...
storage:
  storageClassDeviceSets:
    - name: set1
      count: 1
      volumeClaimTemplates:
        - spec:
            resources:
              requests:
                storage: 10Gi
            storageClassName: gp2
# ...
```

Create a PVC-based cluster

- Steps
 1. Create the Rook operator
 2. Create a Rook-Ceph cluster
 3. Use rook-ceph Krew plugin to see cluster details
 4. Expand the Ceph cluster's OSD size
 5. Expand the Ceph cluster's OSD count
- Using some recommended configs for production

Create the Rook operator

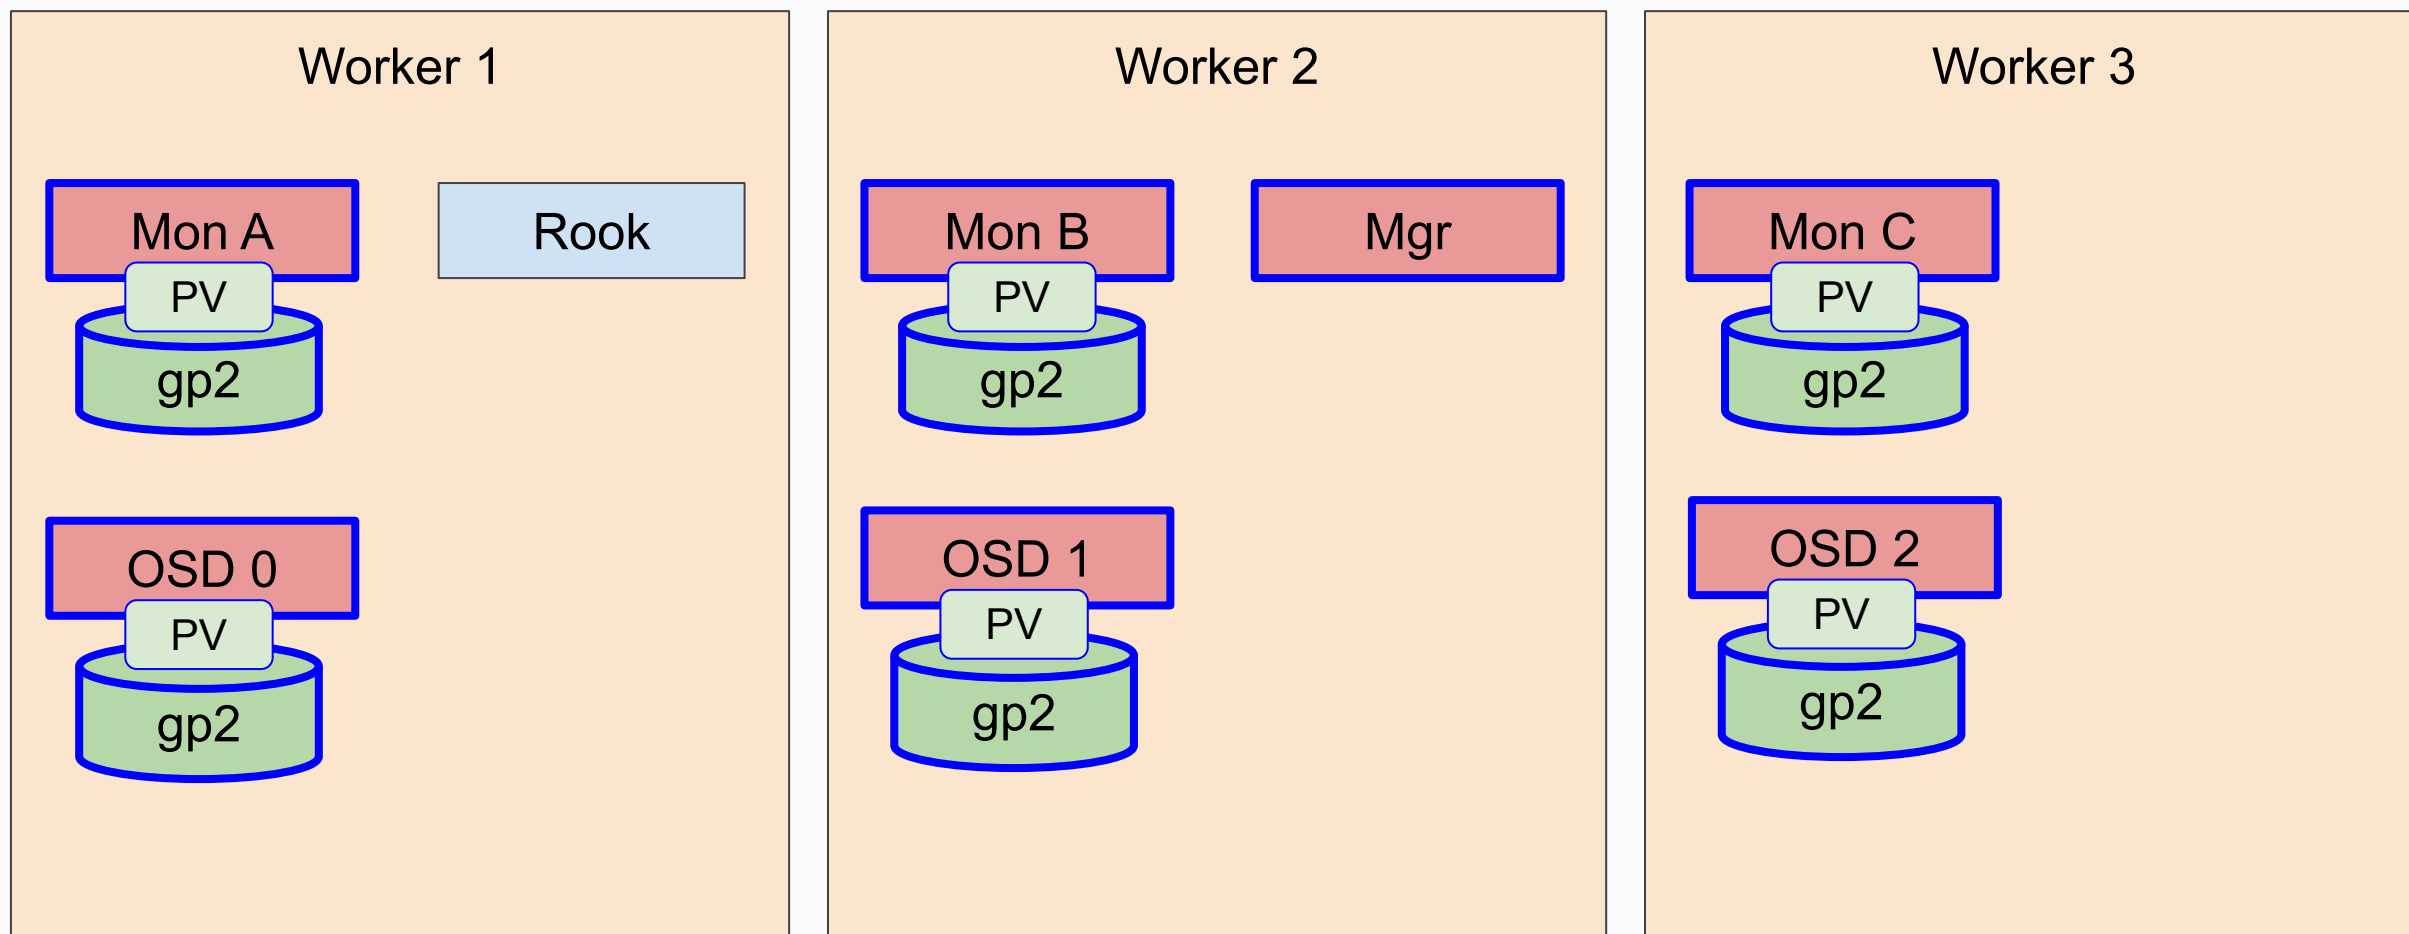
Worker 1

Rook

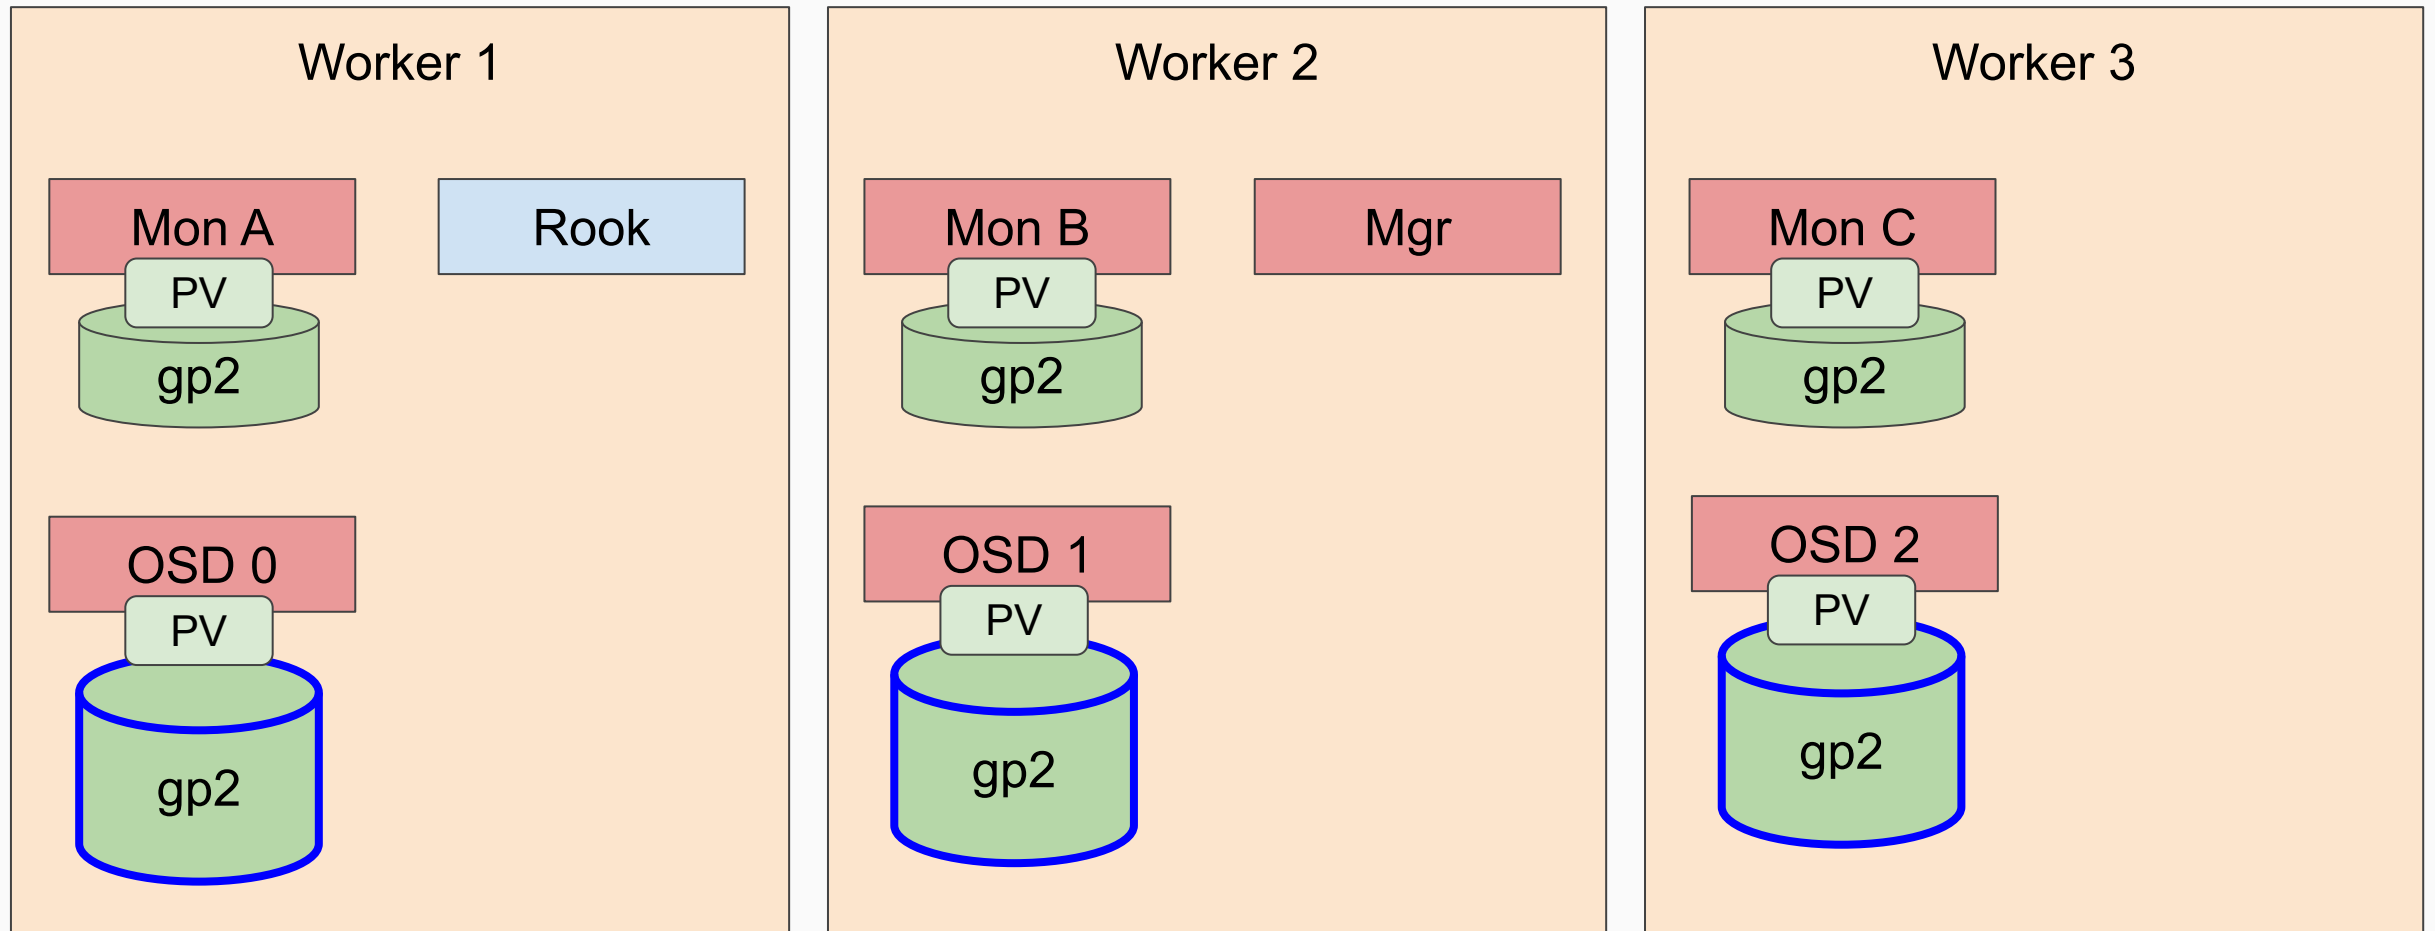
Worker 2

Worker 3

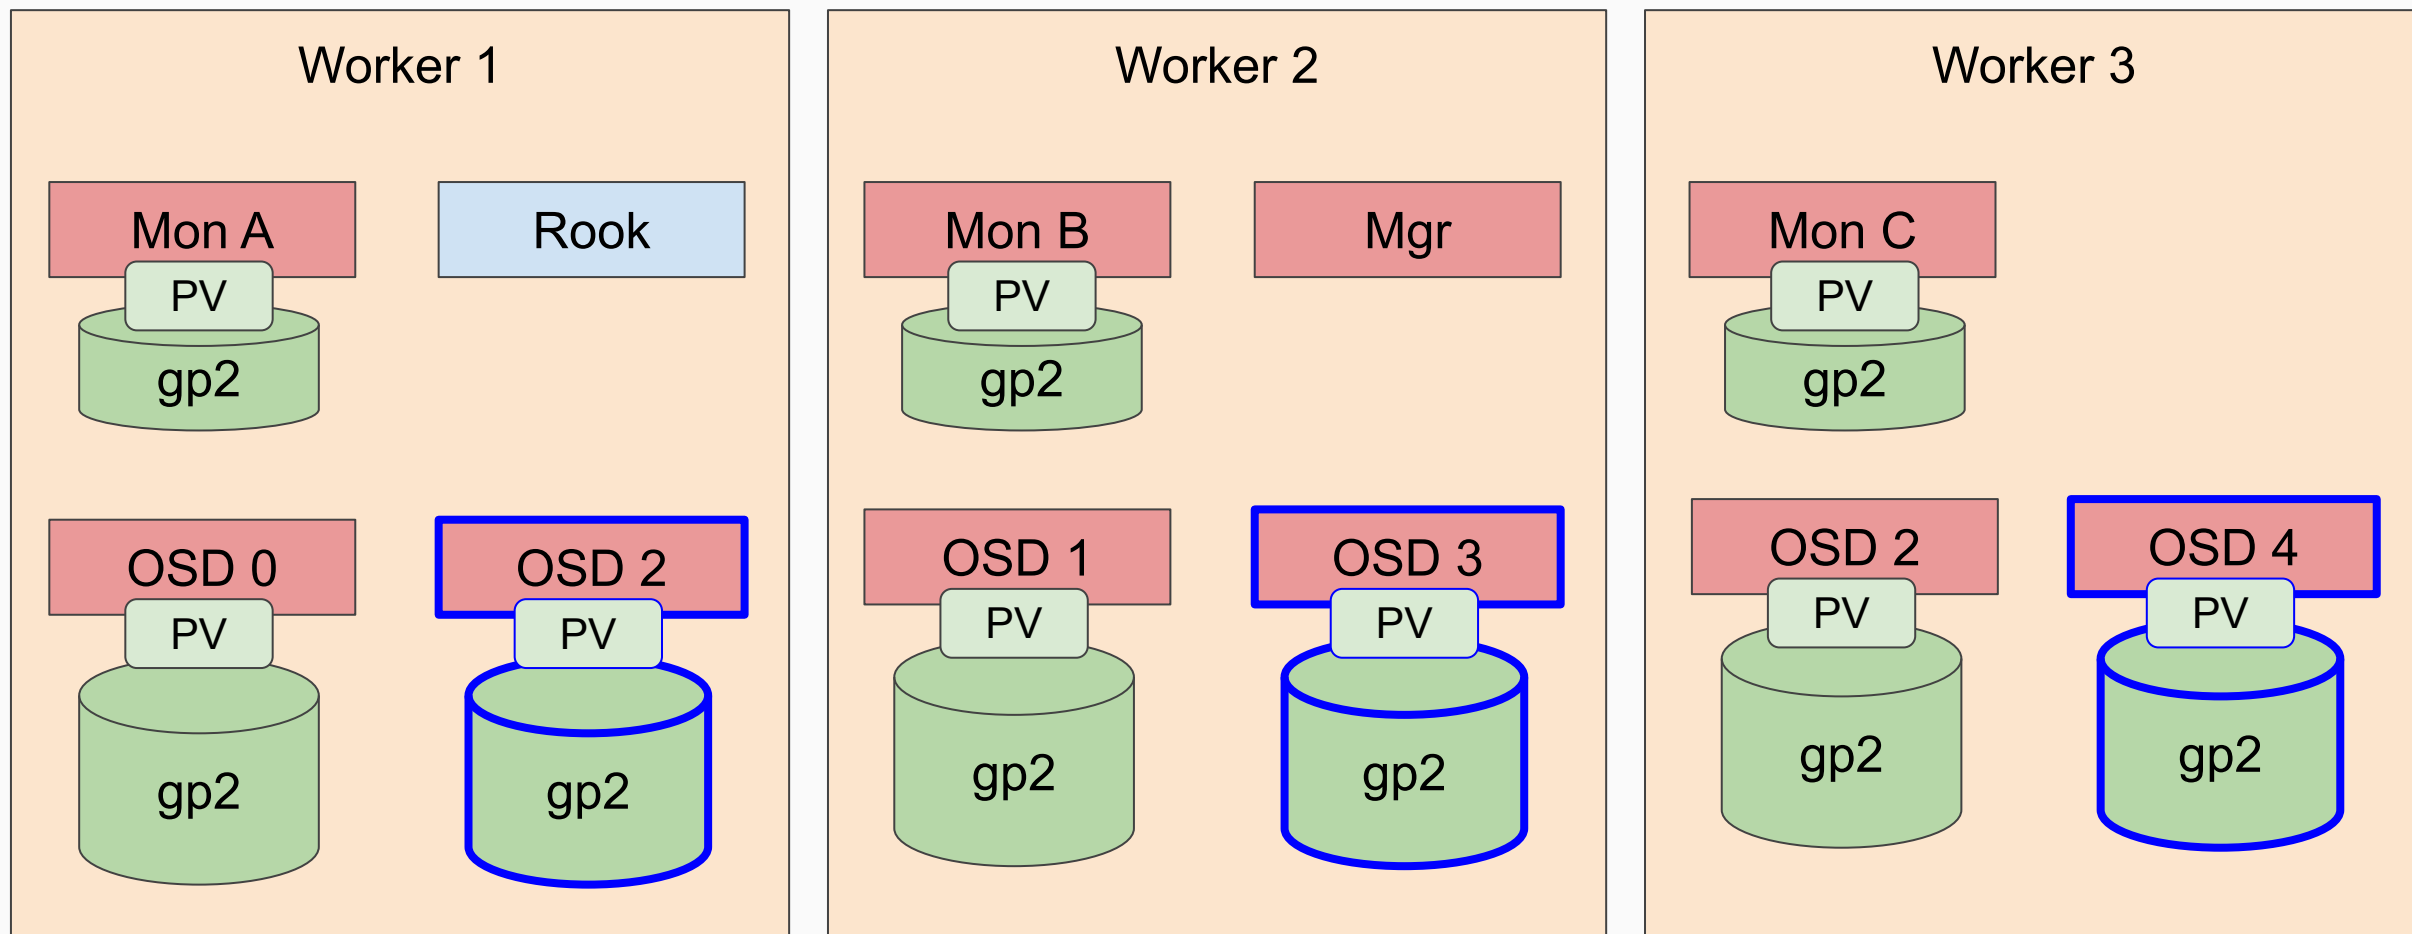
Create a Rook-Ceph cluster



Expand the Ceph cluster's OSD size



Expand the Ceph cluster's OSD count





KubeCon



CloudNativeCon

Europe 2022

Questions?

Website	<u>https://rook.io/</u>
Documentation	<u>https://rook.io/docs/rook/v1.9/</u>
Slack	<u>https://rook-io.slack.com/</u>
Contributions	<u>https://github.com/rook/rook</u>
Twitter	@rook_io
Community Meeting	<u>https://github.com/rook/rook#community-meeting</u>
Training Videos (new!)	<u>https://kubeyexample.com/</u> → Learning Paths → Storage for Kubernetes with Rook

