



Spark on Kubernetes - The Elastic Story

Bowen Li, Huichao Zhao | KubeCon + CloudNativeCon Europe 2022 | 16 May - 20 May 2022

Agenda

Benefits of Cloud

Design Principles & Architecture of Cloud-Native Spark Service

Autoscaling Spark – Cost Saving & Elasticity Needs

Design of Reactive AutoScaling

Productization of Reactive Autoscaling

Learnings & Future Work

Benefits of Cloud

Agile

Resources are on-demand, pay as you go

Elastic & Scalable

Almost infinite scale of compute and storage

Strong Resource Isolation

Container-native on K8S

Privacy-First

Leverage cloud security techniques to enforce security

Operation Friendly

Our developers can focus on building and improving services to achieve higher ROI

Design Principles

Leverage Cloud Infra

Full Containerizing - Elastic, Agile, Lightweight

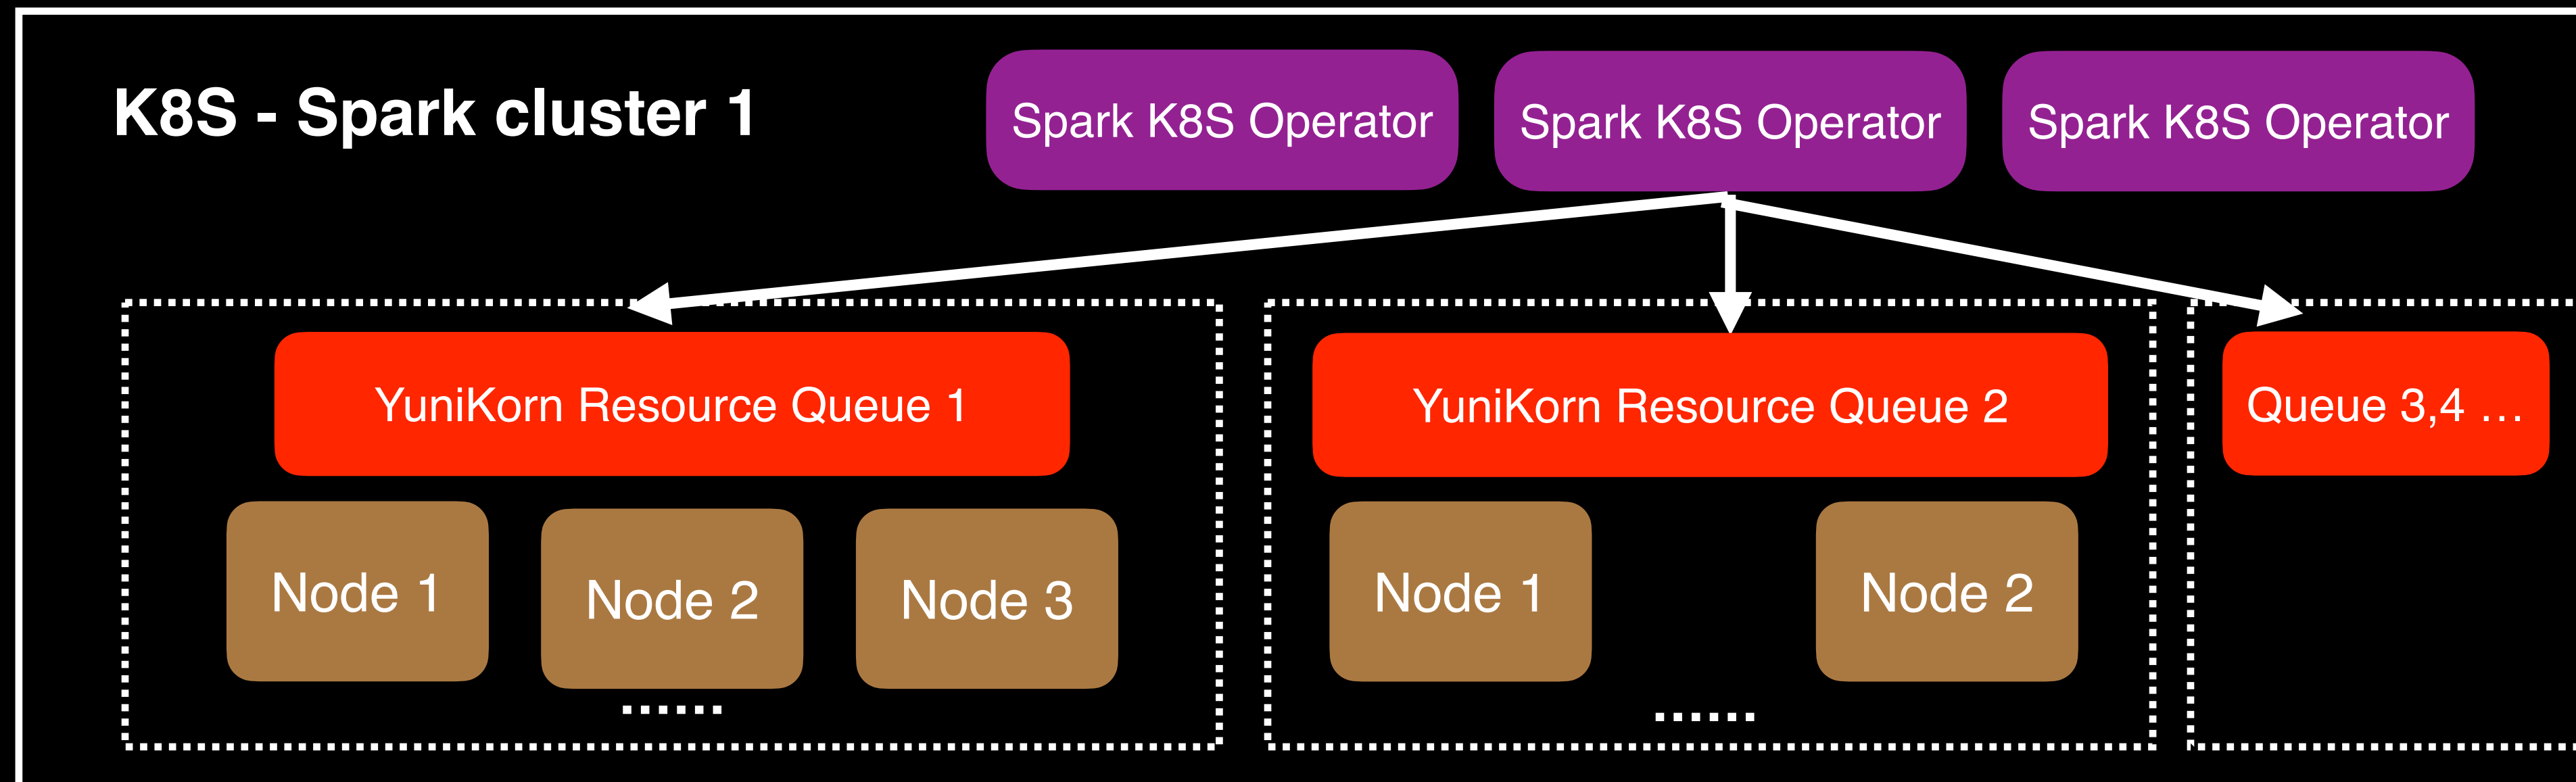
Decouple Compute/Storage, Scale Independently

Developer-Friendly, API Centric

Security & Privacy as First Class Citizen

Use Apple Internal Spark distribution

Architecture of Cloud-Native Spark Service



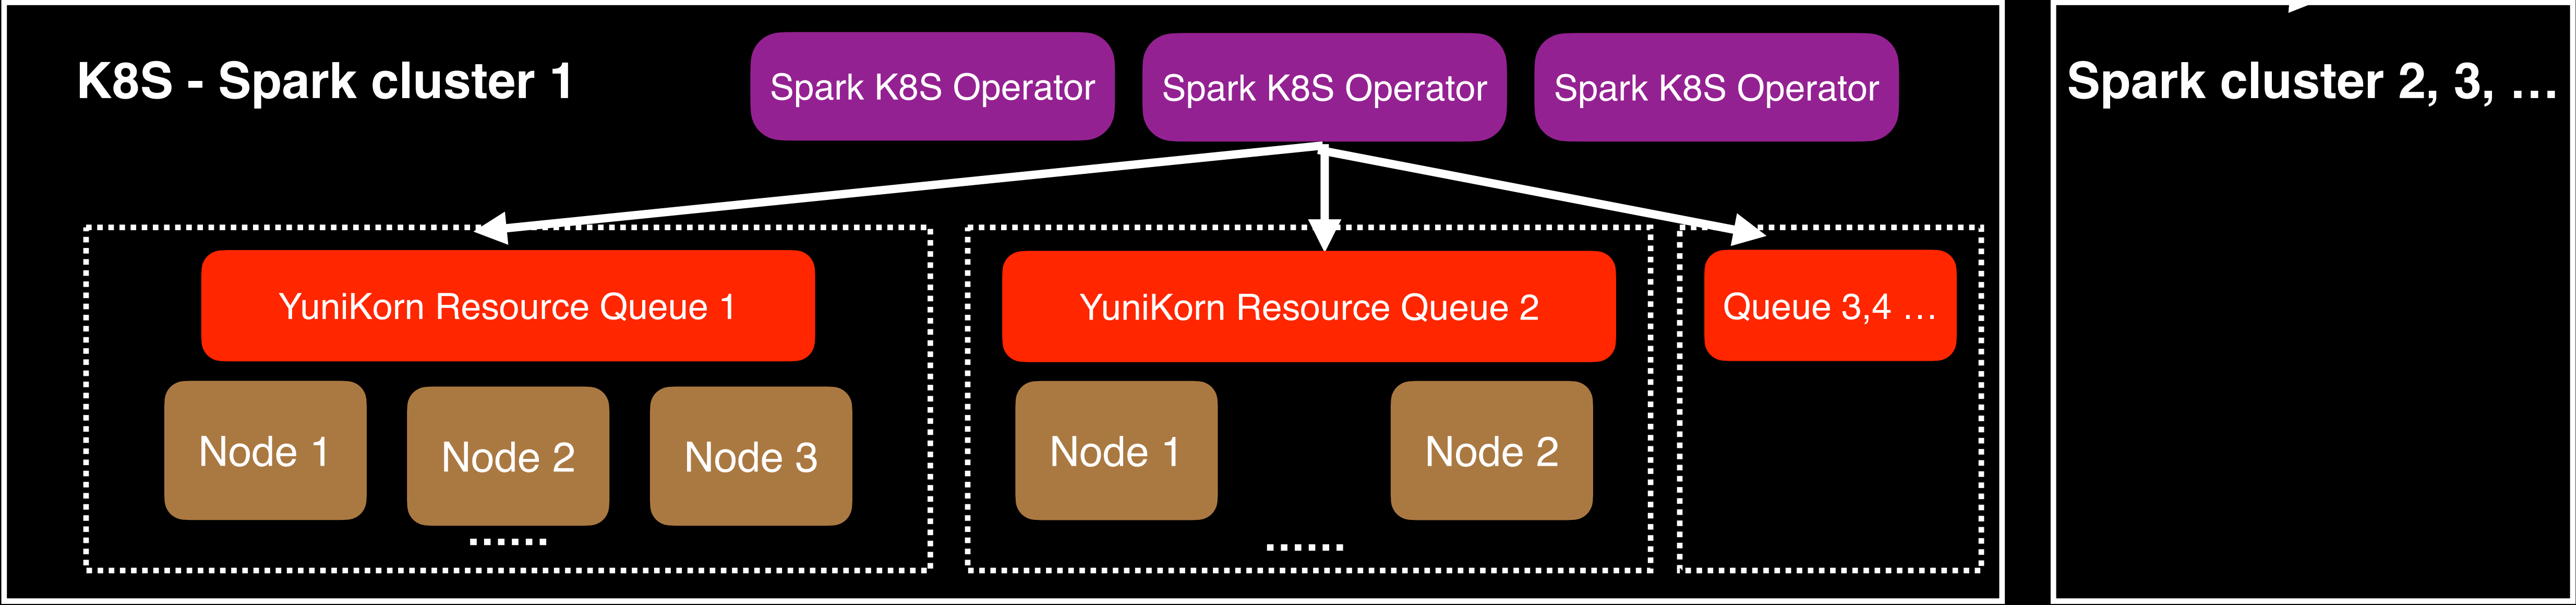
Architecture of Cloud-Native Spark Service

Spark Batch Processing

API/CLI/Airflow Operator

HTTPS request, to submit and manage jobs

Skate (Spark service gateway)



Architecture of Cloud-Native Spark Service

Spark Batch Processing

API/CLI/Airflow Operator

HTTPS request, to submit and manage jobs

Skate (Spark service gateway)

Interactive Spark

Jupyter Notebook

Interactive Spark Gateway

K8S - Spark cluster 1

Spark K8S Operator

Spark K8S Operator

Spark K8S Operator

YuniKorn Resource Queue 1

Node 1

Node 2

Node 3

YuniKorn Resource Queue 2

Node 1

Node 2

Queue 3,4 ...

Spark cluster 2, 3, ...

Architecture of Cloud-Native Spark Service

Spark Batch Processing

API/CLI/Airflow Operator

HTTPS request, to submit and manage jobs

Skate (Spark service gateway)

Interactive Spark

Jupyter Notebook

Interactive Spark Gateway

K8S - Spark cluster 1

Spark K8S Operator

Spark K8S Operator

Spark K8S Operator

YuniKorn Resource Queue 1

Node 1

Node 2

Node 3

YuniKorn Resource Queue 2

Node 1

Node 2

Queue 3,4 ...

Spark cluster 2, 3, ...

Observability Infra

Security & Privacy Infra

Cost saving and Elasticity Needs

Varying workload pattern: fluctuating within a day and/or a week

Different use cases: daily/weekly scheduled jobs, ad hoc jobs, scheduled + adhoc, backfill

Fixed amount of resources must account for max usage, which causes resource waste

Agenda

Benefits of Cloud

Design Principles & Architecture of Cloud-Native Spark Service

Autoscaling Spark – Cost Saving & Elasticity Needs

Design of Reactive AutoScaling


Productization of Reactive Autoscaling

Learnings & Future Work

Reactive AutoScaling Cluster NodeGroups Layout

- Physical isolation: Minimize potential impact


Multi-tenant Autoscaling K8S Cluster (Admin View)



Reactive AutoScaling Cluster NodeGroups Layout

- Physical isolation: Minimize potential impact


Multi-tenant Autoscaling K8S Cluster (Admin View)



Reactive AutoScaling Cluster NodeGroups Layout

- Physical isolation: Minimize potential impact

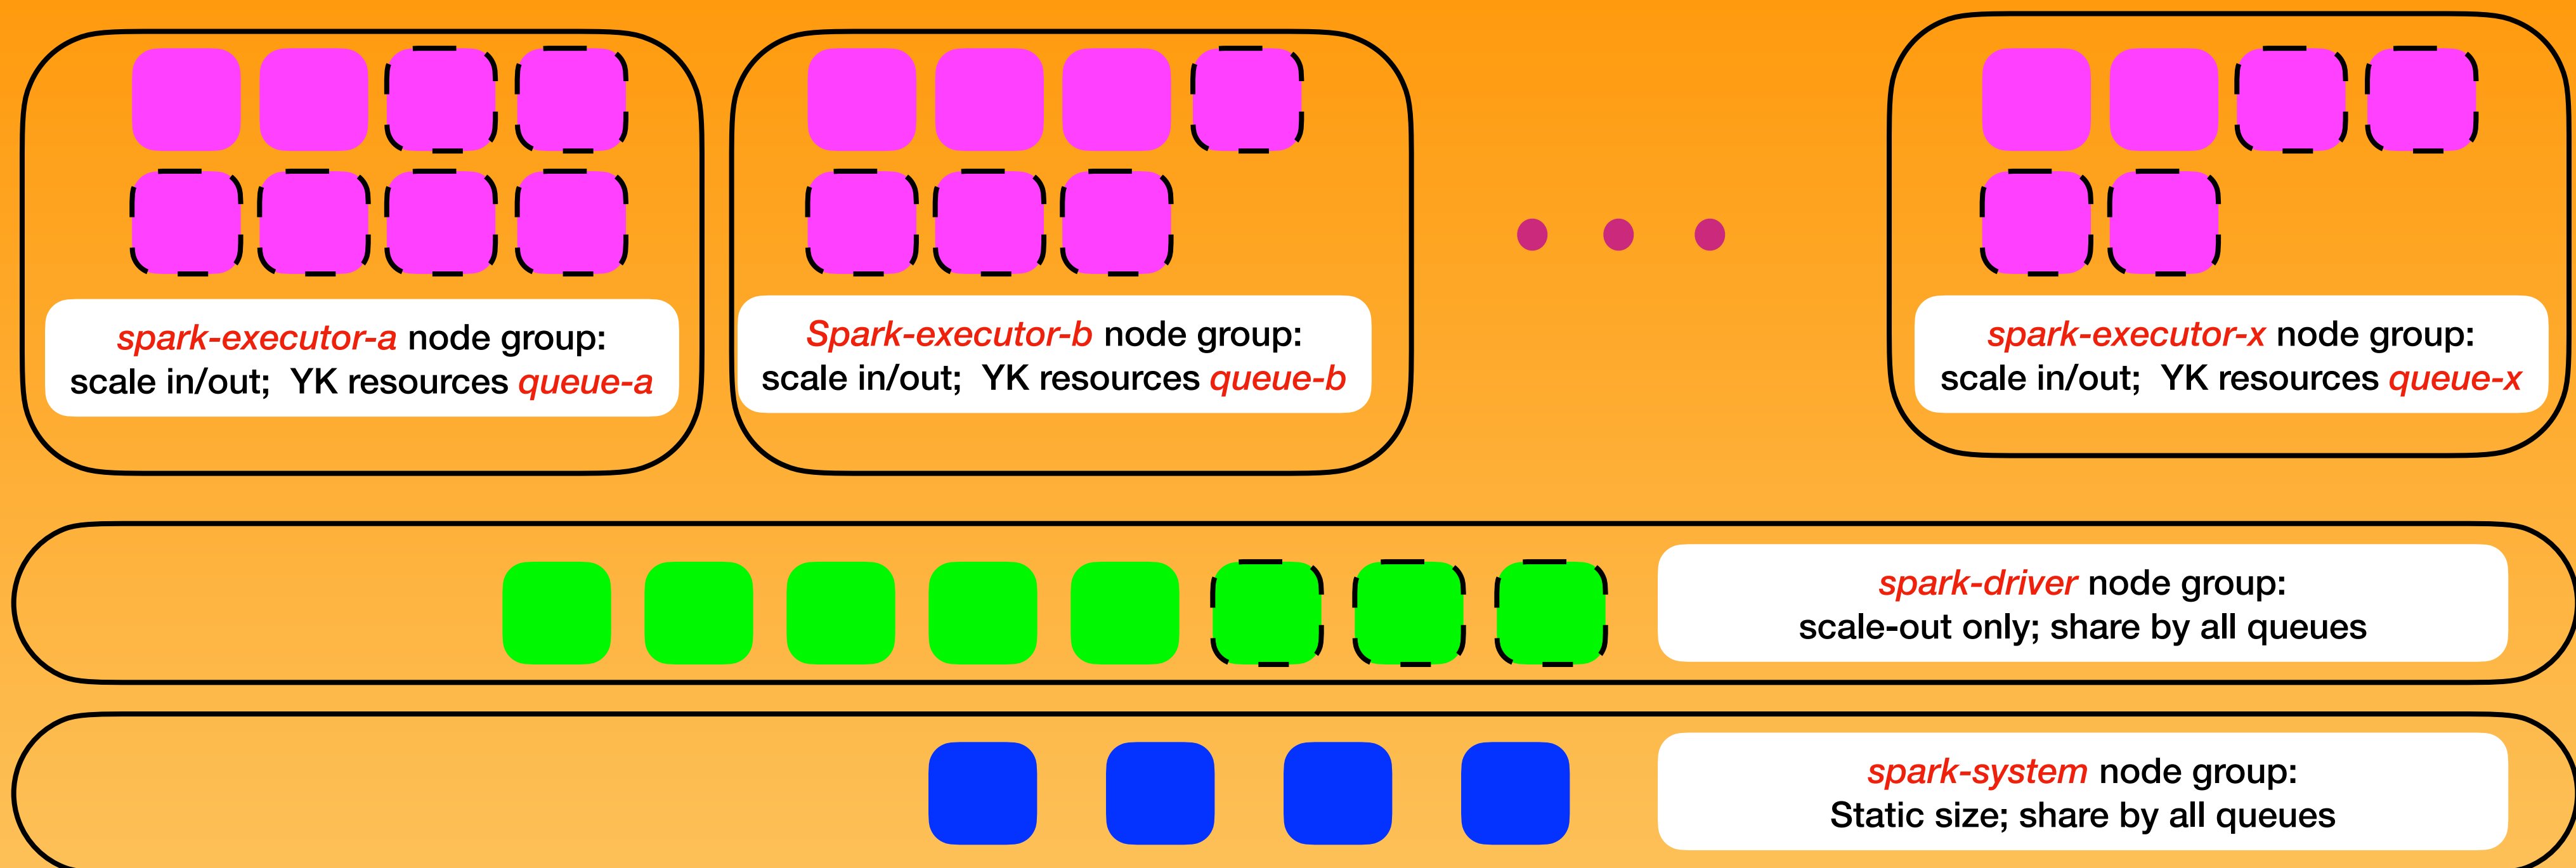
Multi-tenant Autoscaling K8S Cluster (Admin View)



Reactive AutoScaling Cluster NodeGroups Layout

- Physical isolation: Minimize potential impact

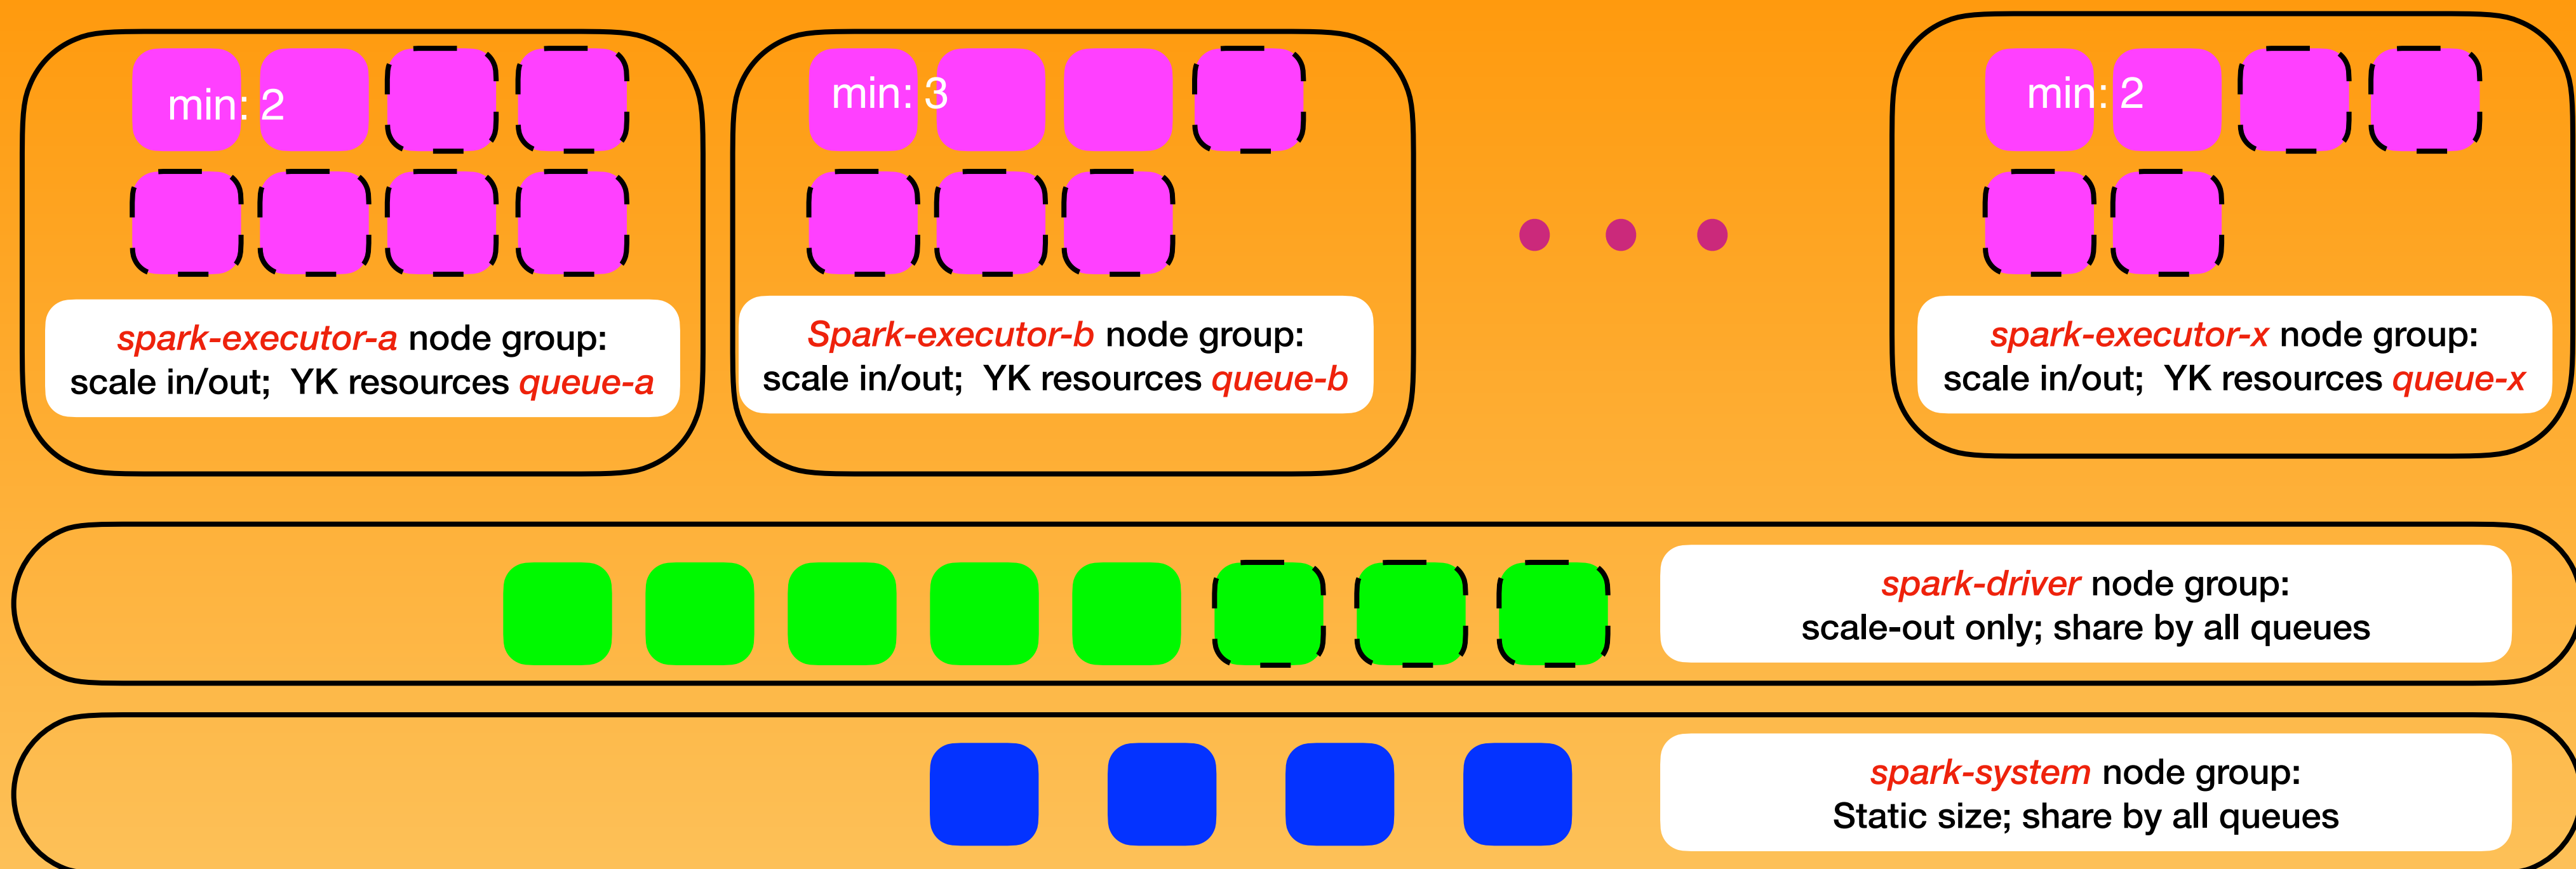
Multi-tenant Autoscaling K8S Cluster (Admin View)



Reactive AutoScaling Cluster NodeGroups Layout

- Physical isolation: Minimize potential impact
- Minimum capacity: Guaranteed **at anytime**

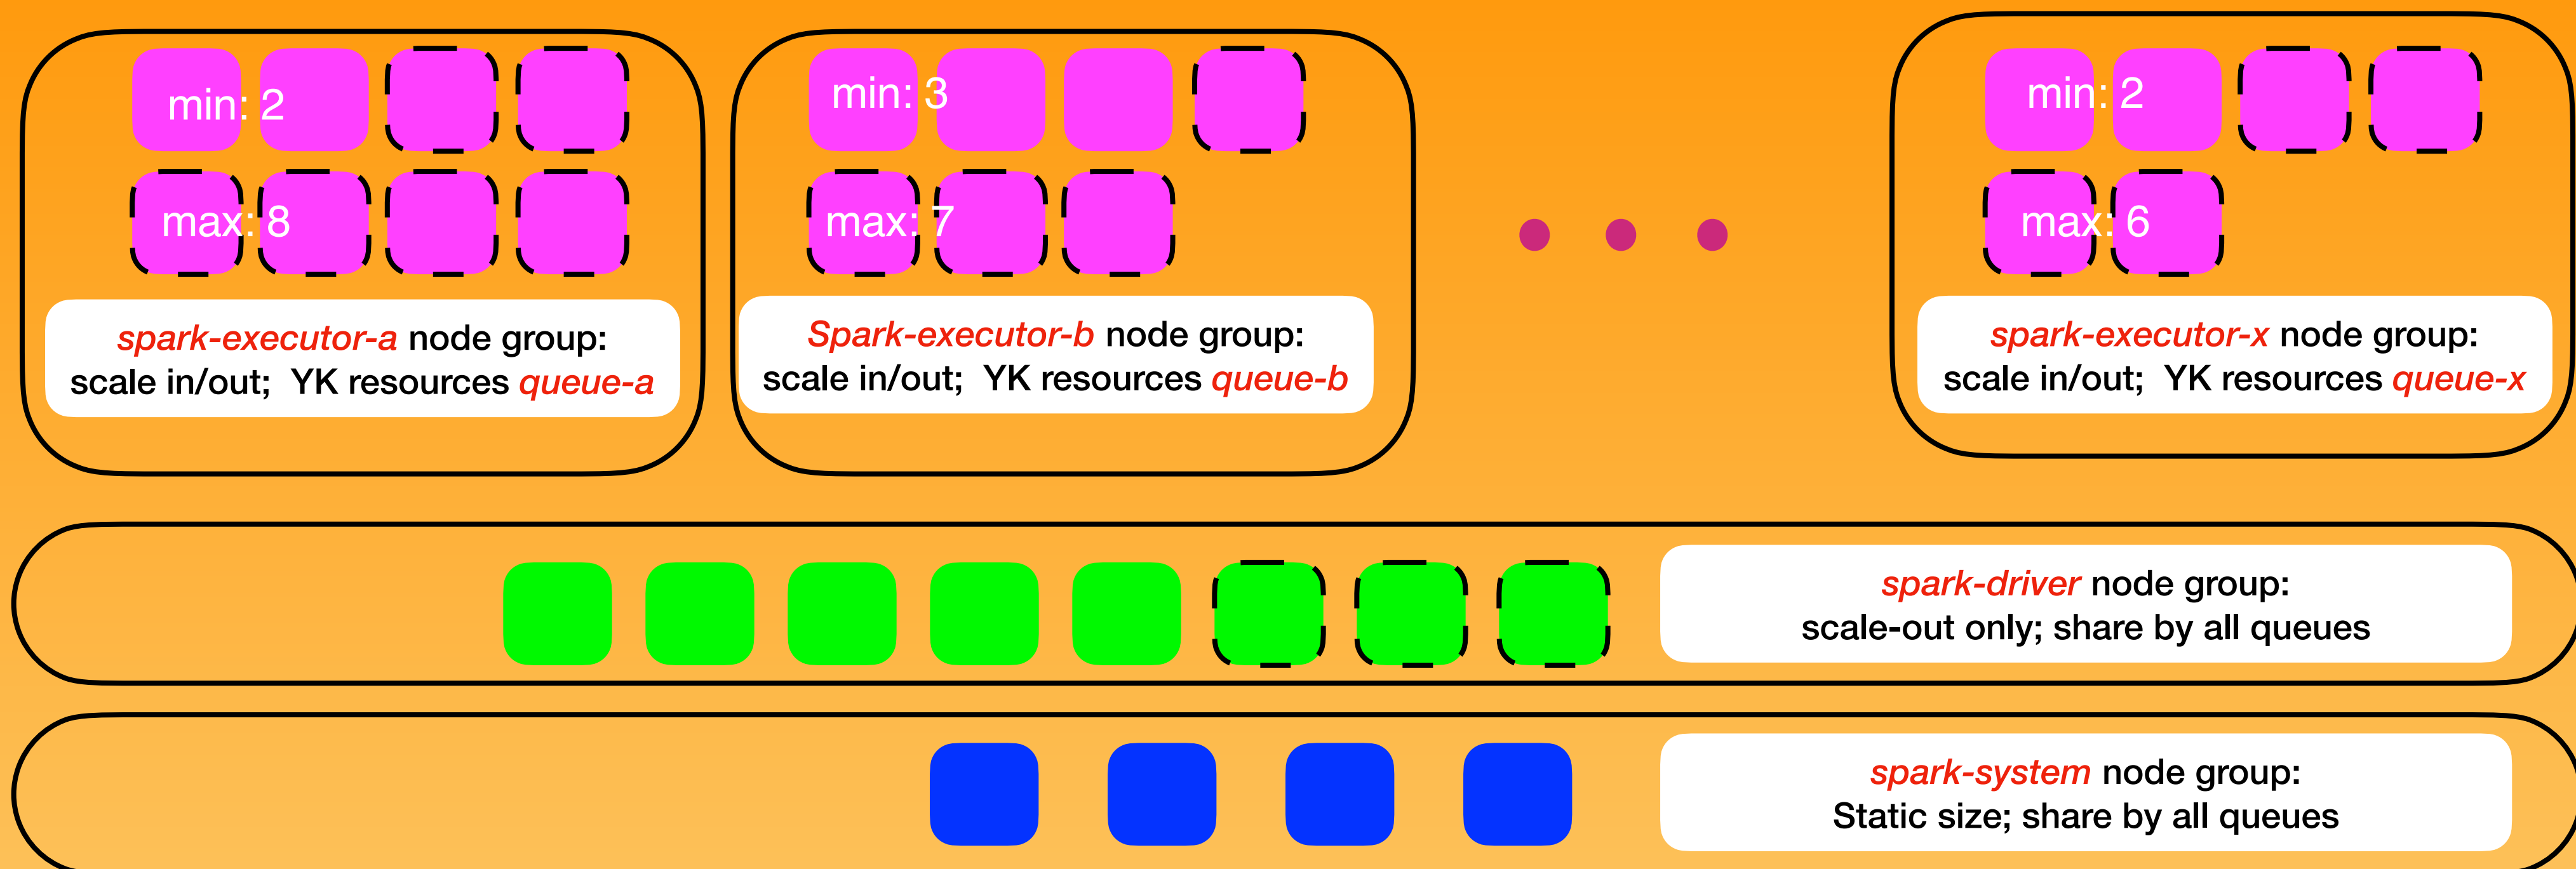
Multi-tenant Autoscaling K8S Cluster (Admin View)



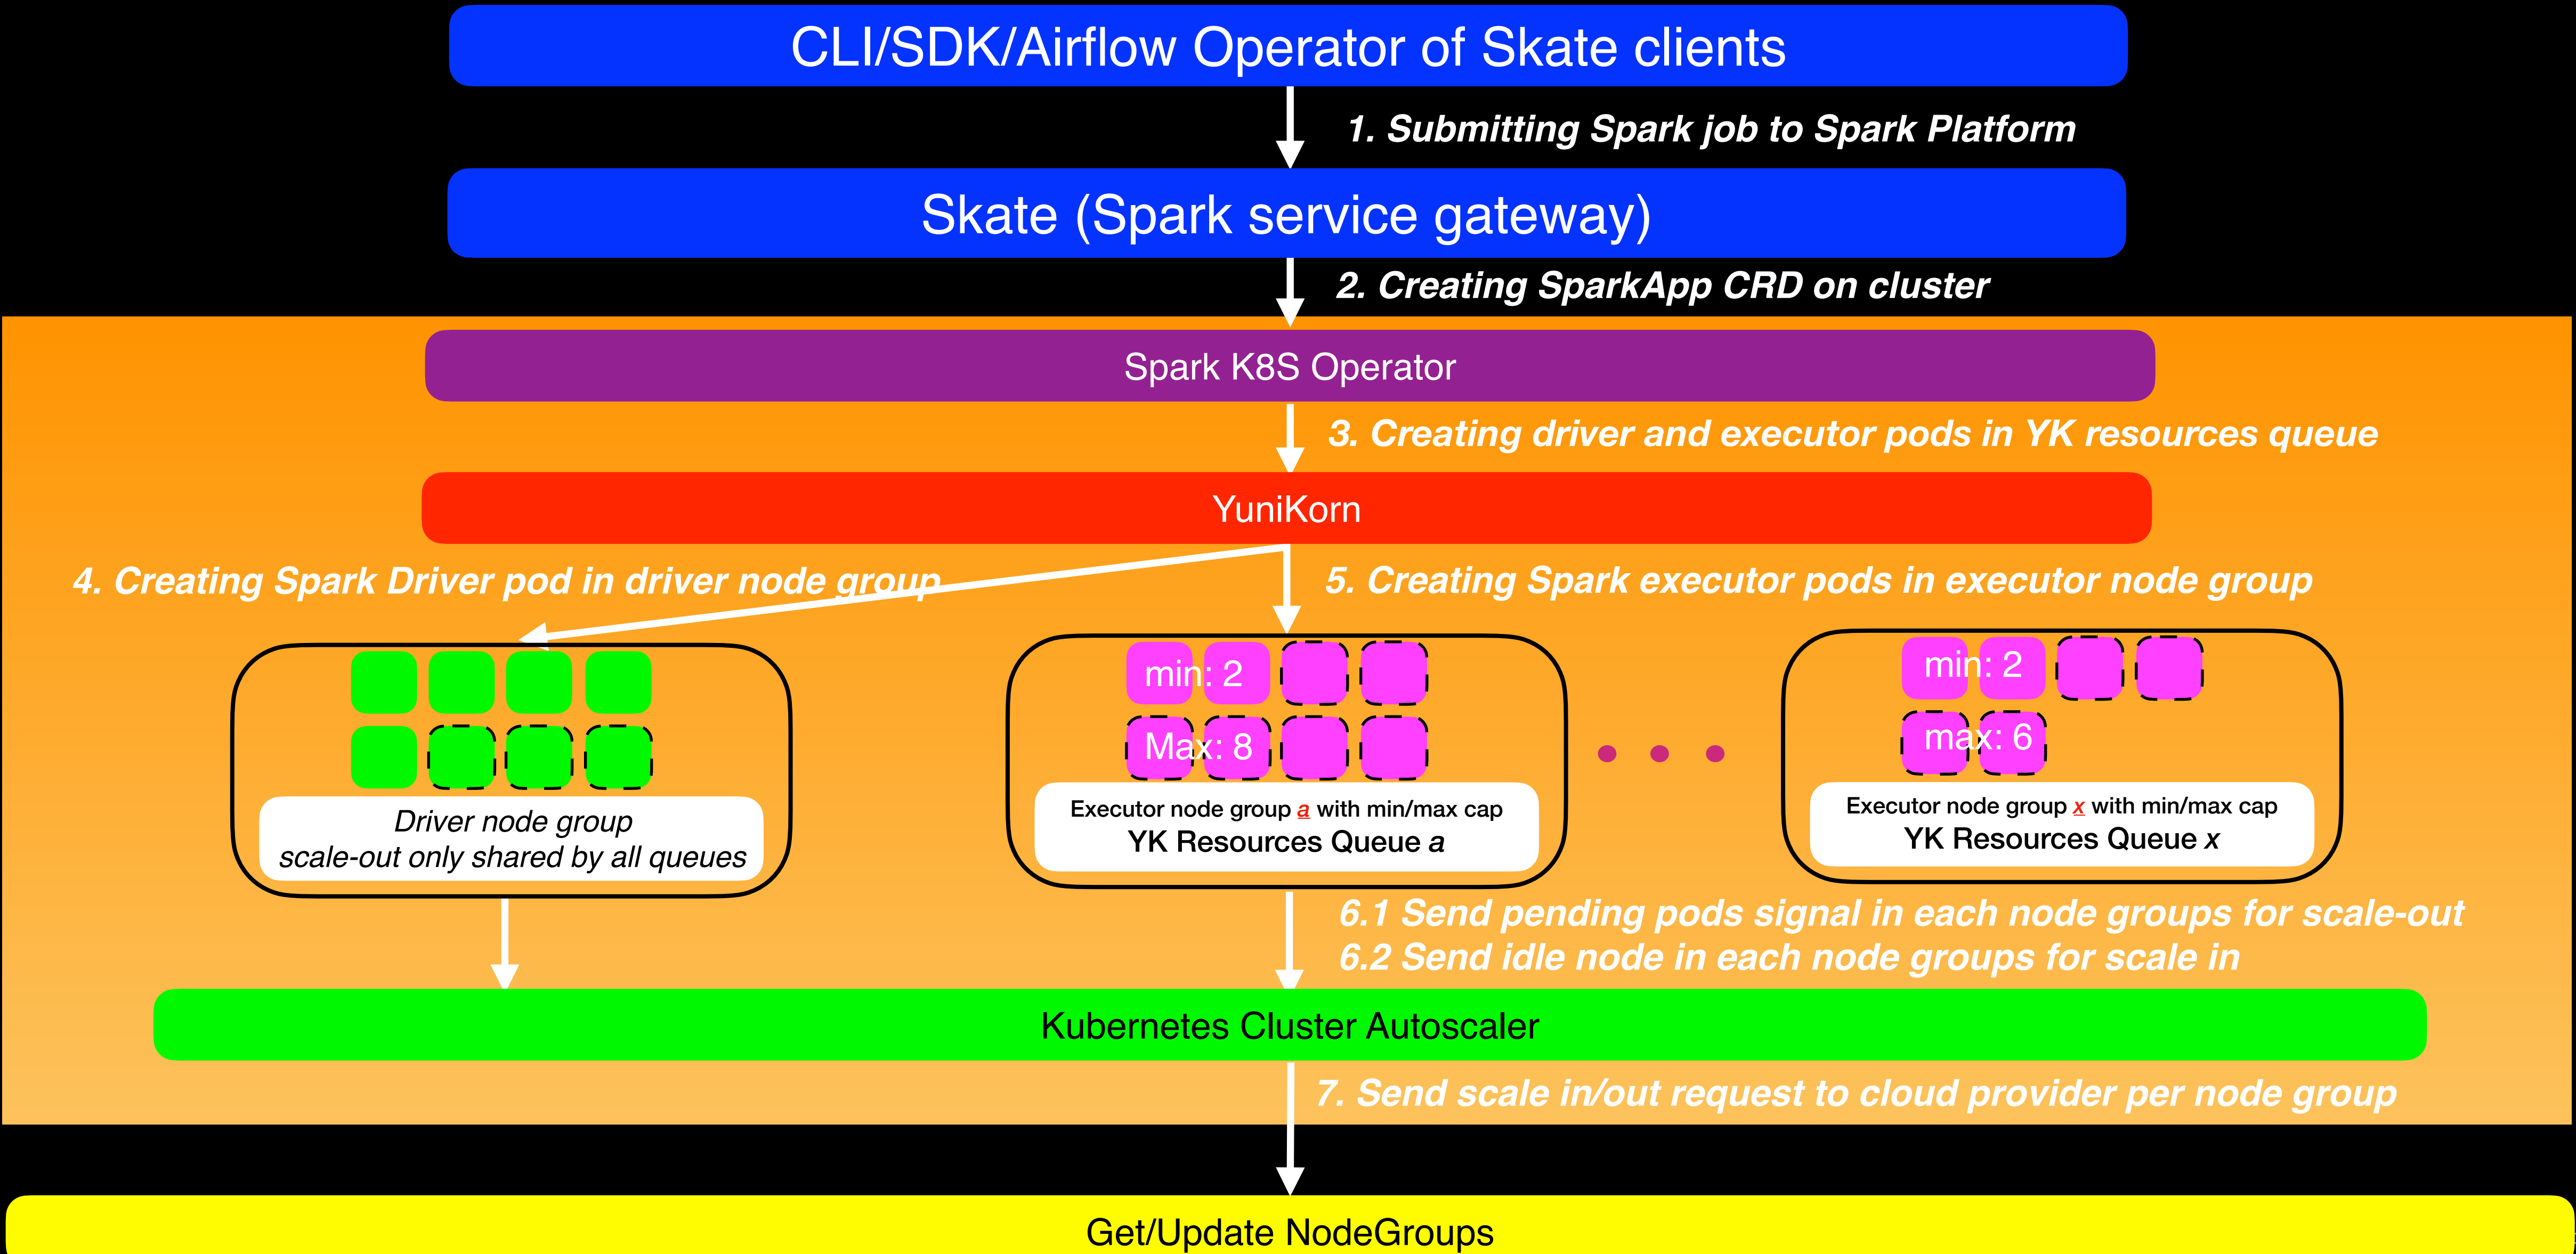
Reactive AutoScaling Cluster NodeGroups Layout

- Physical isolation: Minimize potential impact
- Minimum capacity: Guaranteed **at anytime**
- Maximum capacity: Jobs will be queued if exceed

Multi-tenant Autoscaling K8S Cluster (Admin View)



Reactive AutoScaling Workflow



Reactive AutoScaling Scale in/out Controls

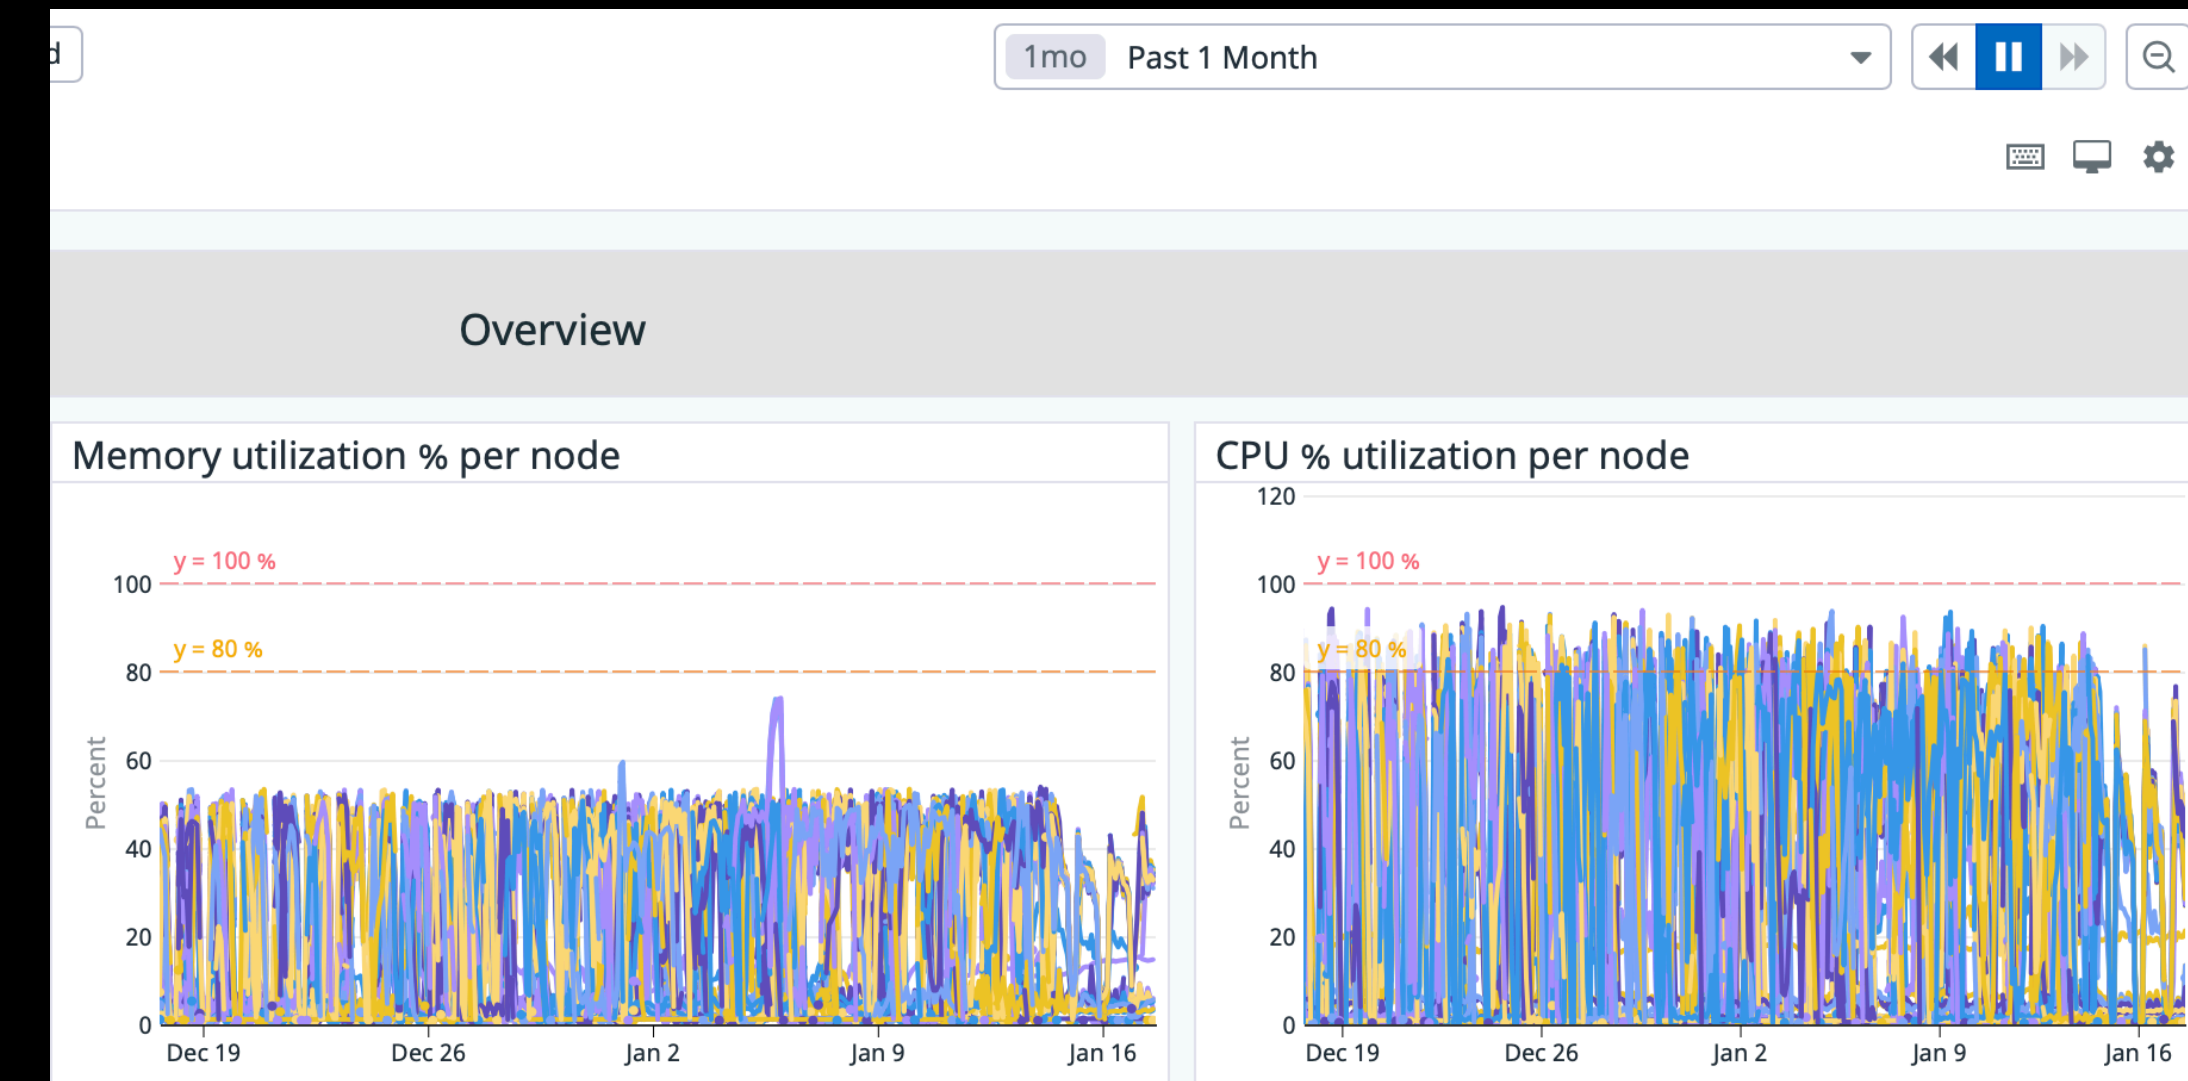


Scale in Controls

- Only when no running executor pods on the node

Reactive AutoScaling Scale in/out Controls

Scale in Controls

- Only when no running executor pods on the node
- Enabled Apache YuniKorn bin-packing in resource scheduling



Reactive AutoScaling Scale in/out Controls

Scale out Controls

- Spark-driver node group scale out only
- Speed up executor pods allocation size config of Spark

Production Status

In production for 3+ months

Cost saving report

- Cost saving percentage per queue is located in 20% - 70%

Migration findings 1

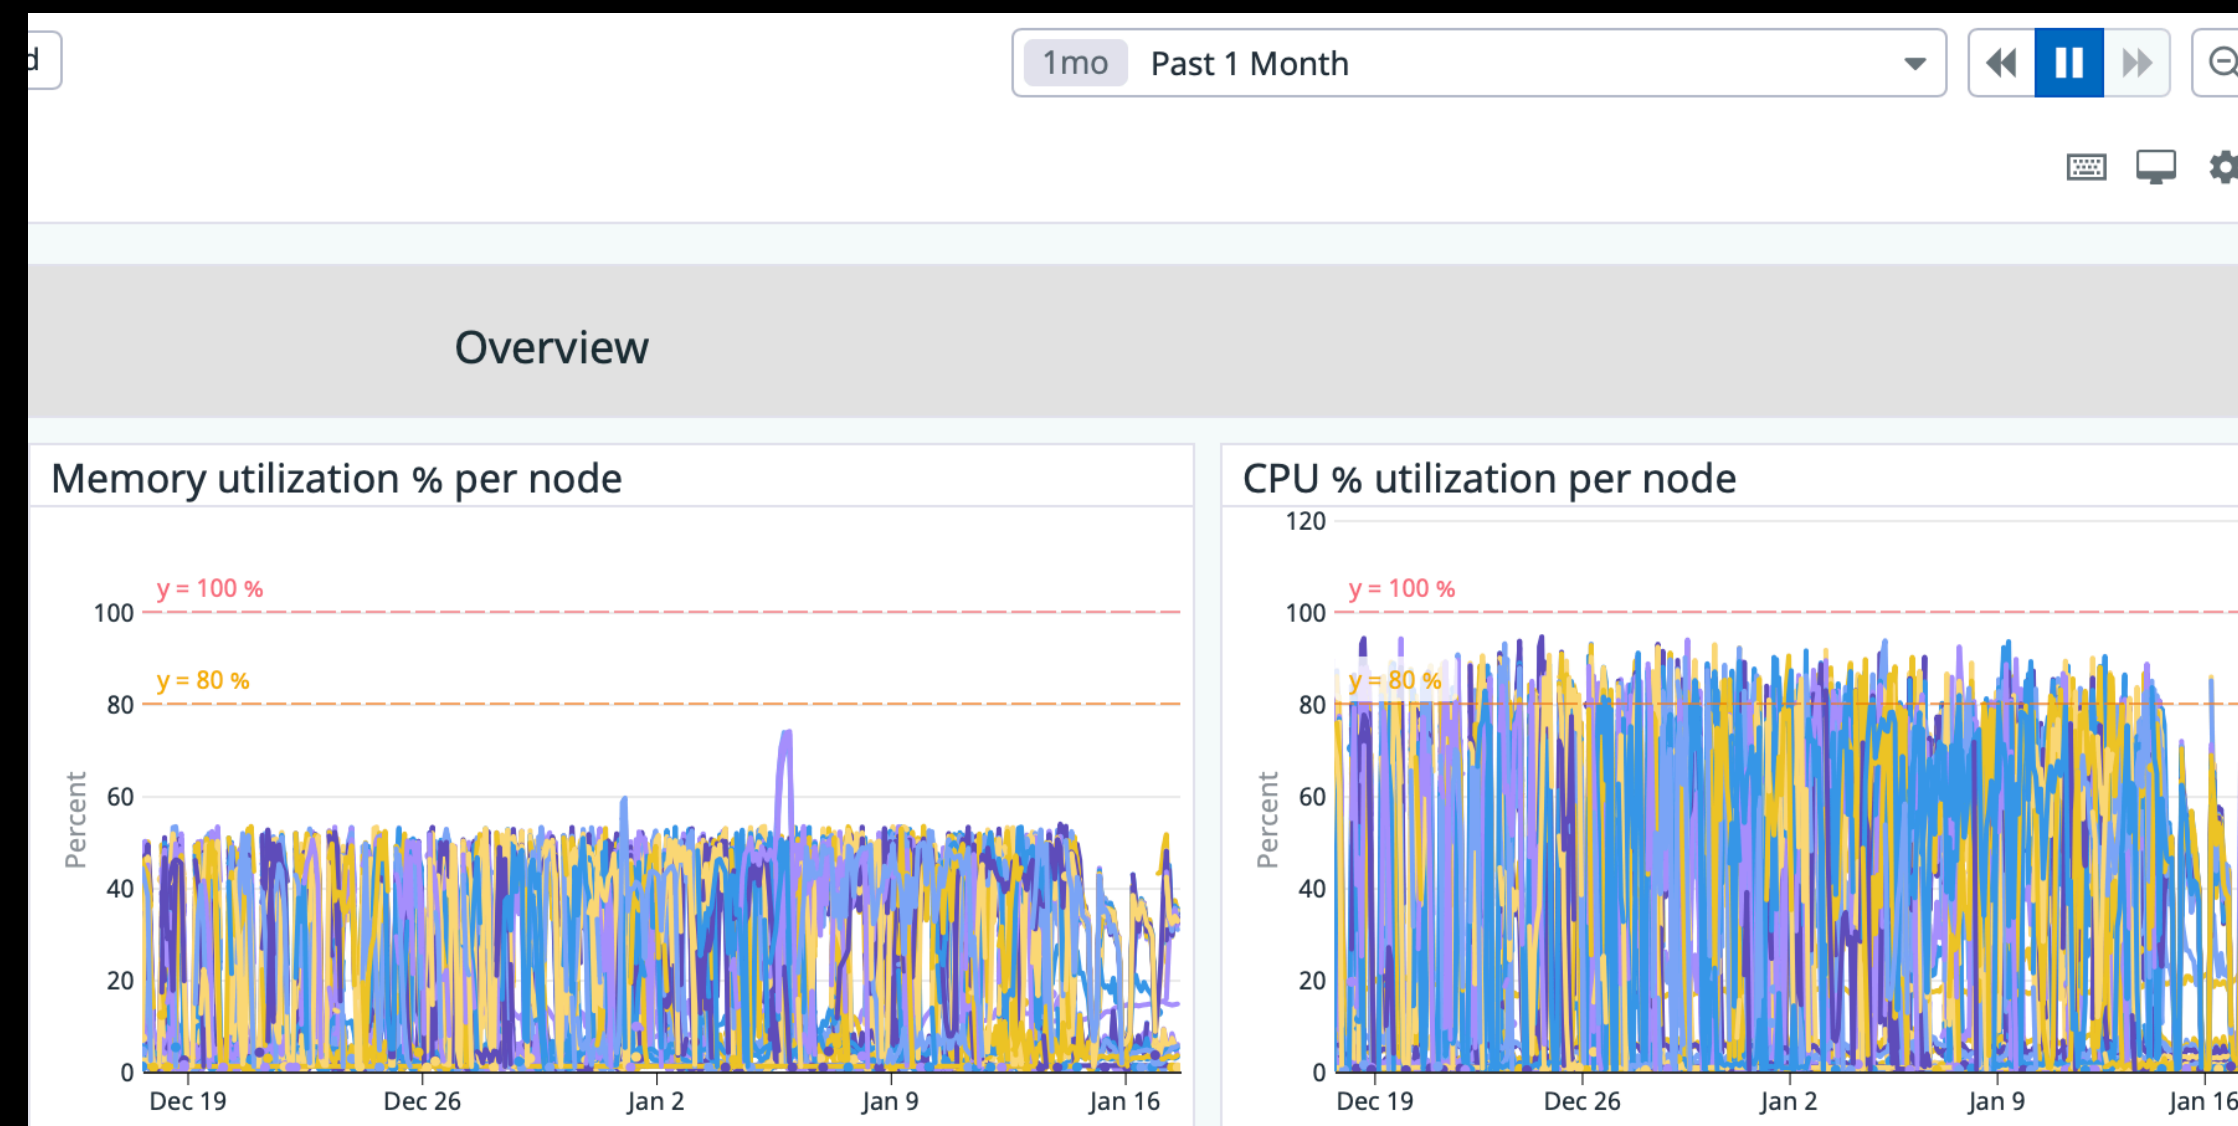
Scale in/out events are stable

- Tens of thousands of job per week running successfully
- All scale-in events works as expected
- Scale out latency is consistent (≤ 5 mins from 2 to 200)

Migration findings 2

Compared to massive over-provisioning approach before, runtime of workloads with autoscaling enabled may increase

- Mostly negligible, a couple jobs increased ~20%
- Users need to take this into consideration and optimize jobs if there's strict data delivery time SLO



Challenges, Solutions, Learnings

- Physical Isolation and min/max capacity setting
- How to guarantee no impact to existing Spark jobs when scale-in
- How to speed up scale-out latency and always allow Spark driver getting start
- Monitoring autoscaling performance

Top Community Feature Requests

- Mixed instance type support
- Dynamic Allocation support
- Spot instance support with Remote Shuffle Service
- Predictive Autoscaling leveraging the platform

