



# The Hitchhiker's Guide to Pod Security



@LachlanEvenson

# My goal for this session?



...everyone here  
enables Pod Security  
on the clusters they manage!



@LachlanEvenson

# What you'll leave with today

- Pod Security Concepts 
- Pod Security in Action 
- Next Steps 



@LachlanEvenson

# About me



Hi! I'm Lachlan (Lachie) Evenson

- Product Manager @ Azure Upstream
- CNCF Ambassador
- CNCF Governing Board member
- Emeritus Kubernetes Steering Committee
- Emeritus Kubernetes 1.16 release lead

I love

- Travel
- Language
- Hiking
- Experiencing different cultures



@LachlanEvenson

# Pod Security Concepts



@LachlanEvenson

# What is Pod Security?



- Pod Security is a **built-in** admission controller
- Evaluates **pod specifications** against a predefined set of **Pod Security Standards** (we will cover this later)
- Applied at the **namespace** level when pods are **created**
- **Beta** as of Kubernetes v1.23
- **Planned** for stable in v1.25

<https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/>



@LachlanEvenson



# Why Pod Security?

- Provides policy standards to **restrict** Pod privileges
- **Simple** and **easy-to-use**
- **Predefined** Pod Security levels that a cluster administrator can configure to meet the desired security posture
- **Supports** and encourages Kubernetes best practices
- Performs **validation** only (no mutation)



@LachlanEvenson

# What about Pod Security Policy?



- PodSecurityPolicy API is deprecated and will be removed from Kubernetes in v1.25
- Pod Security Policy had a few shortcomings
- Pod Security **does not** have feature parity with Pod Security Policy
- Pod Security does not support mutation



@LachlanEvenson



# Migration from Pod Security Policy to Pod Security

- There is a **migration path** from Pod Security Policy to Pod Security
- If you're running Pod Security Policy today and would like to know how to migrate to Pod Security
- <https://kubernetes.io/docs/tasks/configure-pod-container/migrate-from-psp/>
- To summarize the process:
  - Update all existing PSPs to be non-mutating
  - Eliminate options not covered by the Pod Security Standards
  - Apply Pod Security policies in warn or audit mode
  - Upgrade Pod Security policies to enforce mode
  - Remove PodSecurityPolicy from --enable-admission-plugins



@LachlanEvenson



# Other similar ecosystem tooling

- The goal of Pod Security is to provide a built-in set of capabilities to improve the security of your workloads
- These capabilities all fall under Kubernetes policy and governance
- What if you need mutation or other more complex functionality
  - <https://github.com/open-policy-agent/gatekeeper>
  - <https://github.com/kyverno/kyverno>
- Pod Security was designed to be **composable** and work with other solutions

***NB: Other options exist - these are examples to show you where to start investigating***



@LachlanEvenson

# Elements of Pod Security



- Built-in **admission controller** (may be run as a standalone admission webhook)
- There are three different policy levels known as **Pod Security Standards** that range from permissive to restrictive
- Policies are applied in a specific **mode** per namespace
- Multiple modes (with different policy levels) can be set on the same namespace.



@LachlanEvenson



# Elements of Pod Security

- **Pod Security Standards** levels are as follows
  - **privileged** — open and unrestricted
  - **baseline** — Covers known privilege escalations while minimizing restrictions
  - **restricted** — Highly restricted, hardening against known and unknown privilege escalations. May cause compatibility issues



# Elements of Pod Security continued

- Each of these policy levels define which fields are restricted within a pod specification and the allowed values.
  - Some of the fields restricted by these policies include:
    - `spec.securityContext.sysctls`
    - `spec.hostNetwork`
    - `spec.volumes[*].hostPath`
    - `spec.containers[*].securityContext.privileged`
- Applied via labels on Namespace resources, which allows for granular per-namespace policy selection



# Elements of Pod Security

- Policies are applied in a specific mode
- Multiple modes (with different policy levels) can be set on the same namespace.
- List of modes:
  - **enforce** — Any Pods that violate the policy will be rejected
  - **audit** — Violations will be recorded as an annotation in the audit logs, but don't affect whether the pod is allowed.
  - **warn** — Violations will send a warning message back to the user, but don't affect whether the pod is allowed.



@LachlanEvenson



# Elements of Pod Security continued

- In addition to modes you can also pin the policy to a specific Kubernetes version (for example v1.22).
- Pinning to a specific version allows the behavior to remain consistent if the policy definition changes in future Kubernetes releases.



@LachlanEvenson

# Enabling Pod Security



- For Kubernetes >= v1.23.0 it will be enabled by default
- For Kubernetes == v1.22.0 you will need to set the PodSecurity=true feature gate on kube-apiserver
- For Kubernetes version > v1.22.0 that support admission webhooks there is separate documentation on the install <https://github.com/kubernetes/pod-security-admission/tree/master/webhook>

***NB. Please don't test this in Production! Get familiar with Pod Security in a development or lab cluster***



@LachlanEvenson



# Configuring Pod Security

- Policies are applied to a namespace via labels
- These labels are as follows:
  - `pod-security.kubernetes.io/<MODE>: <LEVEL>` (required to enable pod security)
  - `pod-security.kubernetes.io/<MODE>-version: <VERSION>` (optional, defaults to latest)
- <OPTIONAL> cluster-wide policies and exemptions using the AdmissionConfiguration\*.
- The possible <MODE(S)> are **enforce**, **audit** and **warn**.

\* AdmissionConfiguration must be configured via apiserver flags



@LachlanEvenson



# Configuring Pod Security continued

- A specific version can be supplied for each enforcement mode.
- The typical uses for **warn** are to get ready for a future change where you want to enforce a different policy.
- **enforce** is only applied to pod resource. **warn** and **audit** are also applied to workload resources



@LachlanEvenson



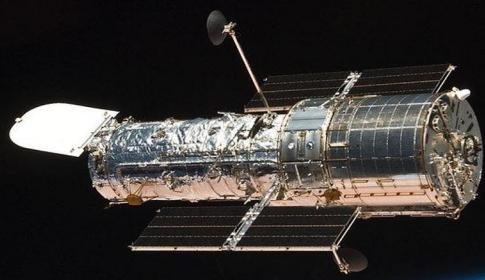
# Configuring Pod Security continued

- The audit mode logs to the cluster audit log (example audit log below)

```
{"authorization.k8s.io/decision": "allow", "authorization.k8s.io/reason": "", "pod-security.kubernetes.io/audit": "allowPrivilegeEscalation != false  
(container \"busybox\" must set securityContext.allowPrivilegeEscalation=false), unrestricted capabilities (container \"busybox\" must set securityContext.capabilities.drop=[\"ALL\"]), runAsNonRoot != true  
(pod or container \"busybox\" must set securityContext.runAsNonRoot=true), seccompProfile (pod or container \"busybox\" must set securityContext.seccompProfile.type to \"RuntimeDefault\" or \"Localhost\")"}}
```



# Pod Security in Action



@LachlanEvenson

•••

demo1 - hitchhikers

□ □ | 08



! demo1 ×

```
! demo1
1 // Demo 1 - Confirm Pod Security is enabled
2
3 // Confirm Kubernetes version and access to the API server
4 kubectl version
5 kubectl get nodes
6
7 // Check the API server flags (if you have access)
8 kubectl -n kube-system exec kube-apiserver-kind-control-plane -it -- kube-apiserver -h | grep 'PodSecurityPolicy'
9
10 // OPTIONAL Run a quick test to confirm Pod Security is enabled
11
```

PROBLEMS TERMINAL OUTPUT DEBUG CONSOLE

bash + □ □ ^ x

```
laevenso@feu [* |kind-kind:default] in ~/sandbox/hitchhikers
$ 
```



# Confirm Pod Security is enabled

&lt;

X 0 △ 0

Live Share kind-kind

Ln 24, Col 12

Spaces: 2

UTF-8

LF

YAML

✓ Spell

No JSON Schema



@LachlanEvenson

demo2.txt -- hitchhikers

demo2.txt

```
1 // Demo 2 - Privileged level and workload
2
3 // Create a namespace
4 kubectl create namespace verify-pod-security
5
6 // Label namespace to enforce a "restricted" security policy and audit on restricted
7 kubectl label --overwrite ns verify-pod-security \
8   pod-security.kubernetes.io/enforce=restricted \
9   pod-security.kubernetes.io/audit=restricted
10
11 // Deploy privileged workload in the namespace
```

PROBLEMS TERMINAL OUTPUT DEBUG CONSOLE

laevenso@feu [(\*) | kind-kind:default] in ~/sandbox/hitchhikers

\$

Ln 5, Col 1 Spaces: 2 UTF-8 LF Plain Text ✓ Spell 🔍 🔔



# Privileged level and workload



@LachlanEvenson

•••

demo3.txt — hitchhikers



...



≡ demo3.txt ×



≡ demo3.txt

```
1 // Demo 3 - Restricted level and workload
2
3 // Create a namespace
4 kubectl create namespace verify-pod-security
5
6 // Label namespace to enforce a "restricted" security policy and audits on restricted
7 kubectl label --overwrite ns verify-pod-security \
8   pod-security.kubernetes.io/enforce=restricted \
9   pod-security.kubernetes.io/audit=restricted
10
11 // Deploy restricted workload in the namespace
```



PROBLEMS 2

TERMINAL

OUTPUT

DEBUG CONSOLE



laevenso@feu [(\*) |kind-kind:default)] in ~/sandbox/hitchhikers  
\$

## Restricted level and workload

X 0 △ 0 ⌂ 2

Live Share kind-kind

Ln 5, Col 1

Spaces: 2

UTF-8

LF

Plain Text ⌂ 2 Spell



@LachlanEvenson

•••

demo4.txt — hitchhikers

□ □ □ 08



demo3.txt

demo4.txt ×

□ ...

```
demo4.txt
1 // Demo 4 – Review Pod Security evaluation metrics
2
3 // Query the metrics endpoint
4
5 kubectl get --raw /metrics | grep pod_security_evaluations_total
6
```

PROBLEMS 2

TERMINAL

OUTPUT

DEBUG CONSOLE

bash + ↻ □ ⌂ ^ ×

laevenso@feu [(\*) | kind-kind:default] in ~/sandbox/hitchhikers  
\$

# Review Pod Security evaluation metrics



&lt;

⊗ 0 △ 0 ① 2

Live Share kind-kind

Ln 3, Col 30

Spaces: 2

UTF-8 LF

Plain Text ✓ Spell



@LachlanEvenson



# Next Steps



@LachlanEvenson



# Next Steps

- Go experiment to start learning!
- Use warn and audit on existing namespaces
  - Use dry-run to evaluate all resources in a namespace when applying an enforce label to a namespace
- Set warn to the same level as enforce so that users get early feedback
- Make a goal to get all workload namespaces to baseline standard
  - Start with audit and move towards enforce



@LachlanEvenson



Thanks for hitchhiking Pod Security

with me! Now it's your turn...

So long, and thanks for all the fish! 

Special thanks to Jim Angel, Tim Allclair, and Bridget Kromhout for their review!