



KubeCon



CloudNativeCon

Europe 2023





KubeCon



CloudNativeCon

Europe 2023

Archetypes for Reliable Systems

Steve McGhee

Reliability Advocate

Google

Ameer Abbas

Product Manager

Google

Who are we?



Steve McGhee

Reliability Advocate
Google



Ameer Abbas

Product Manager
Google

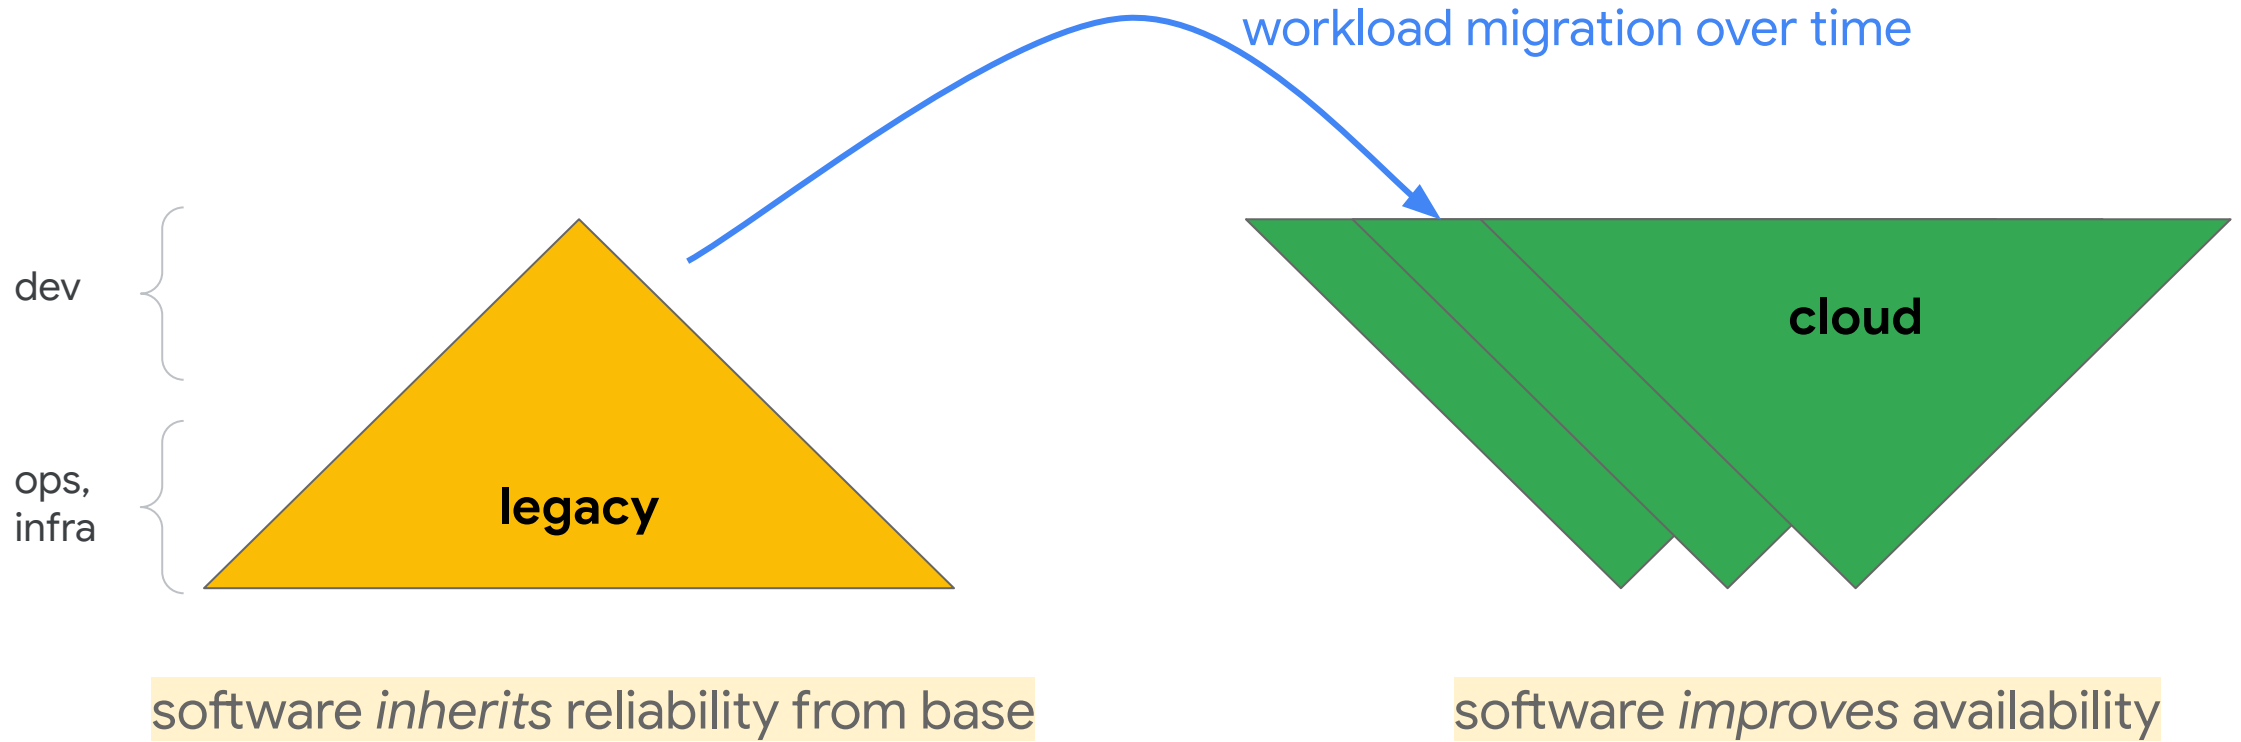
Agenda

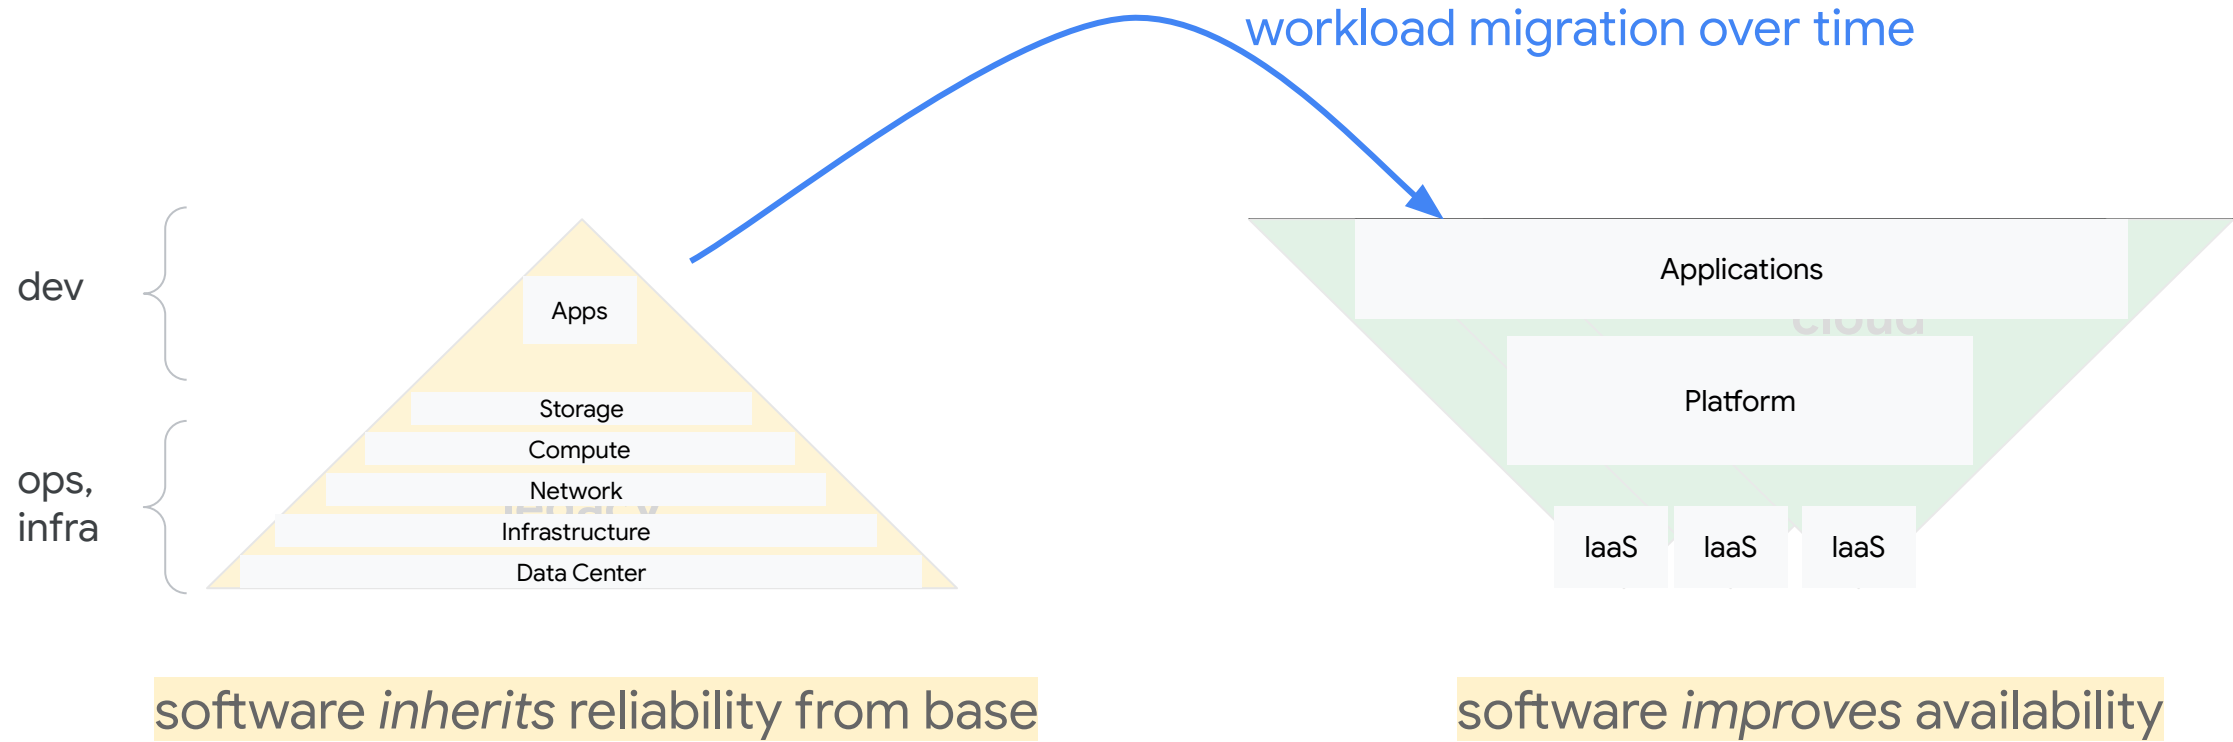
- 01** Reliability Terms
- 02** Application Archetypes
- 03** SLO Math
- 04** Reliable Architectures

Can you build
99.99 services
on
99.9 infrastructure?

Pyramids of Reliability







Application Archetypes



Archetypes to Reliability

- **Archetype** - Abstract model
Replication, redundancy, RTO/RPO, DR, cost
- **Architecture** - Products and Service design
K8s, Mesh, CI/CD, DBs, Storage and backup
- **App/Service** & Footprint - *always changing*
- **SLOs** - expectations, guardrails

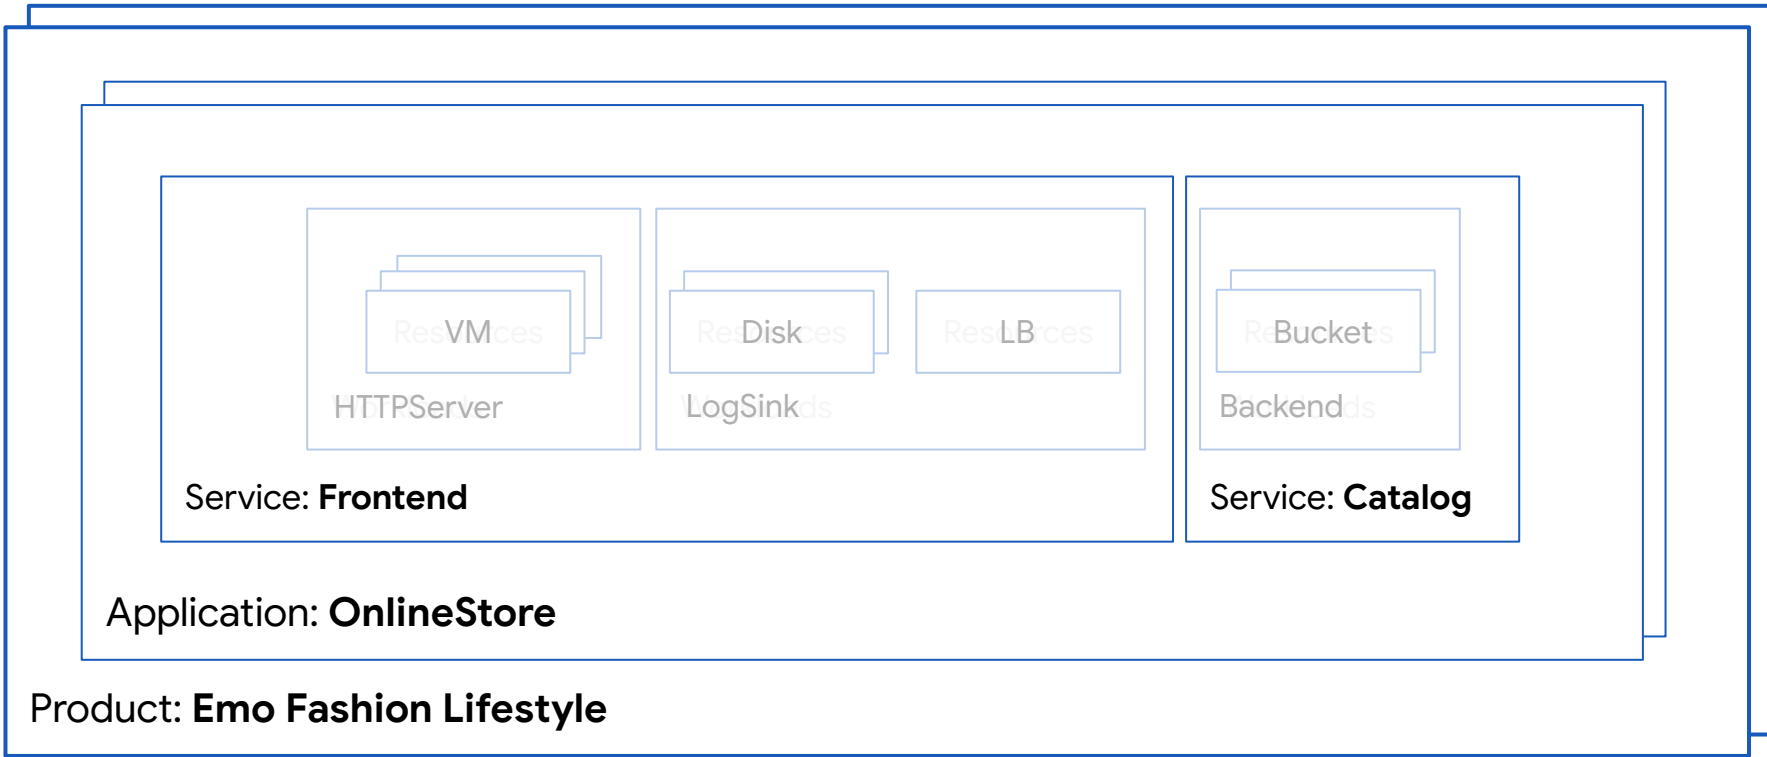
Platforms and Applications





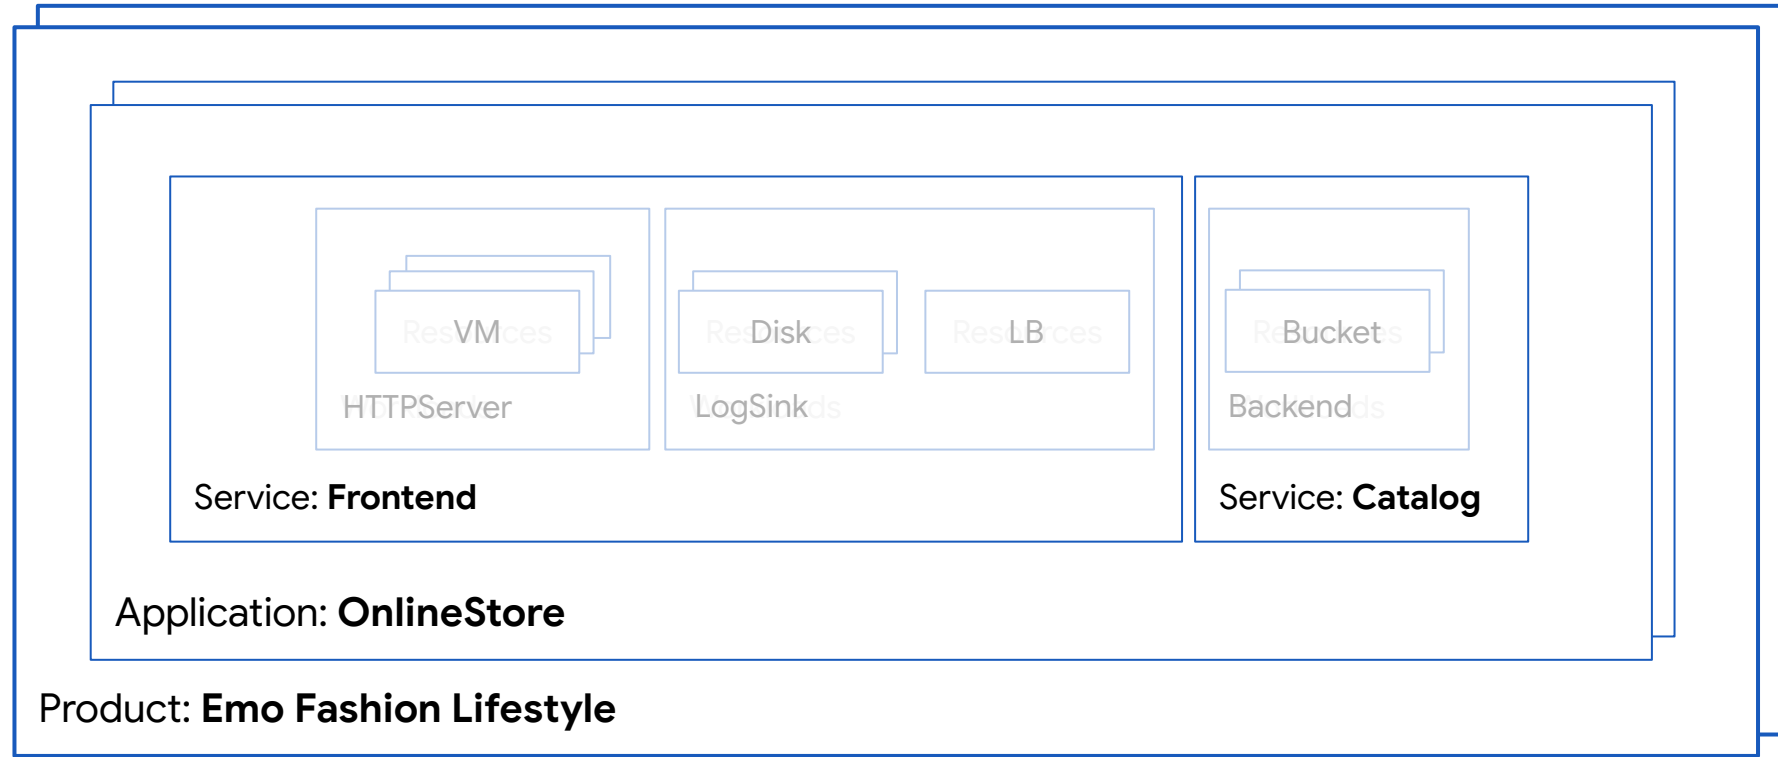
Supports N
Application Archetypes






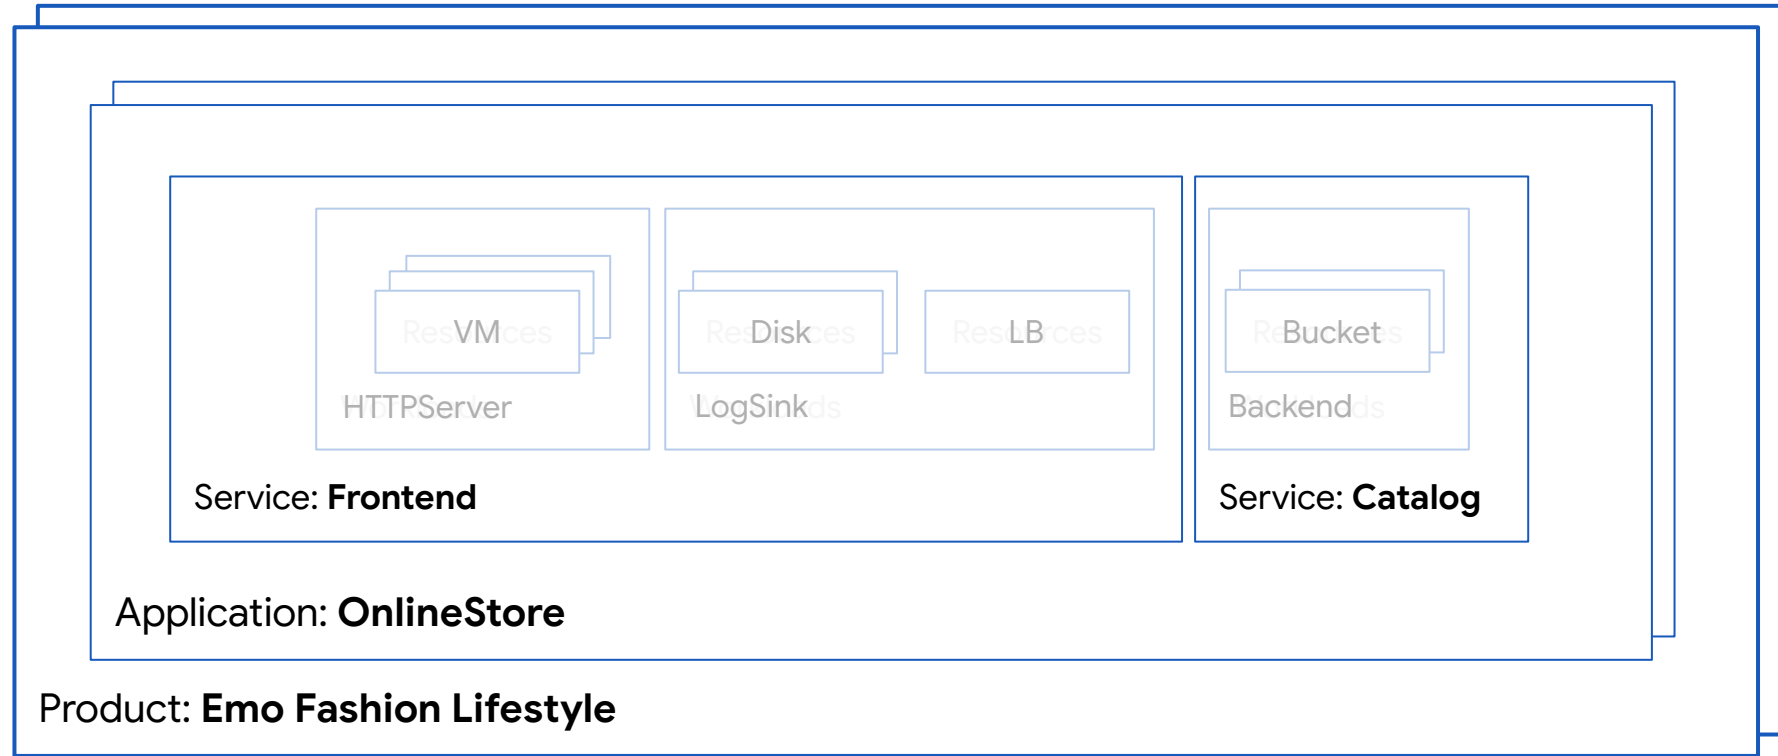
Supports N
Application Archetypes





Supports N
Application Archetypes





Supports 3
Application Archetypes

Global

Active-Passive Region

Active-Passive Zone

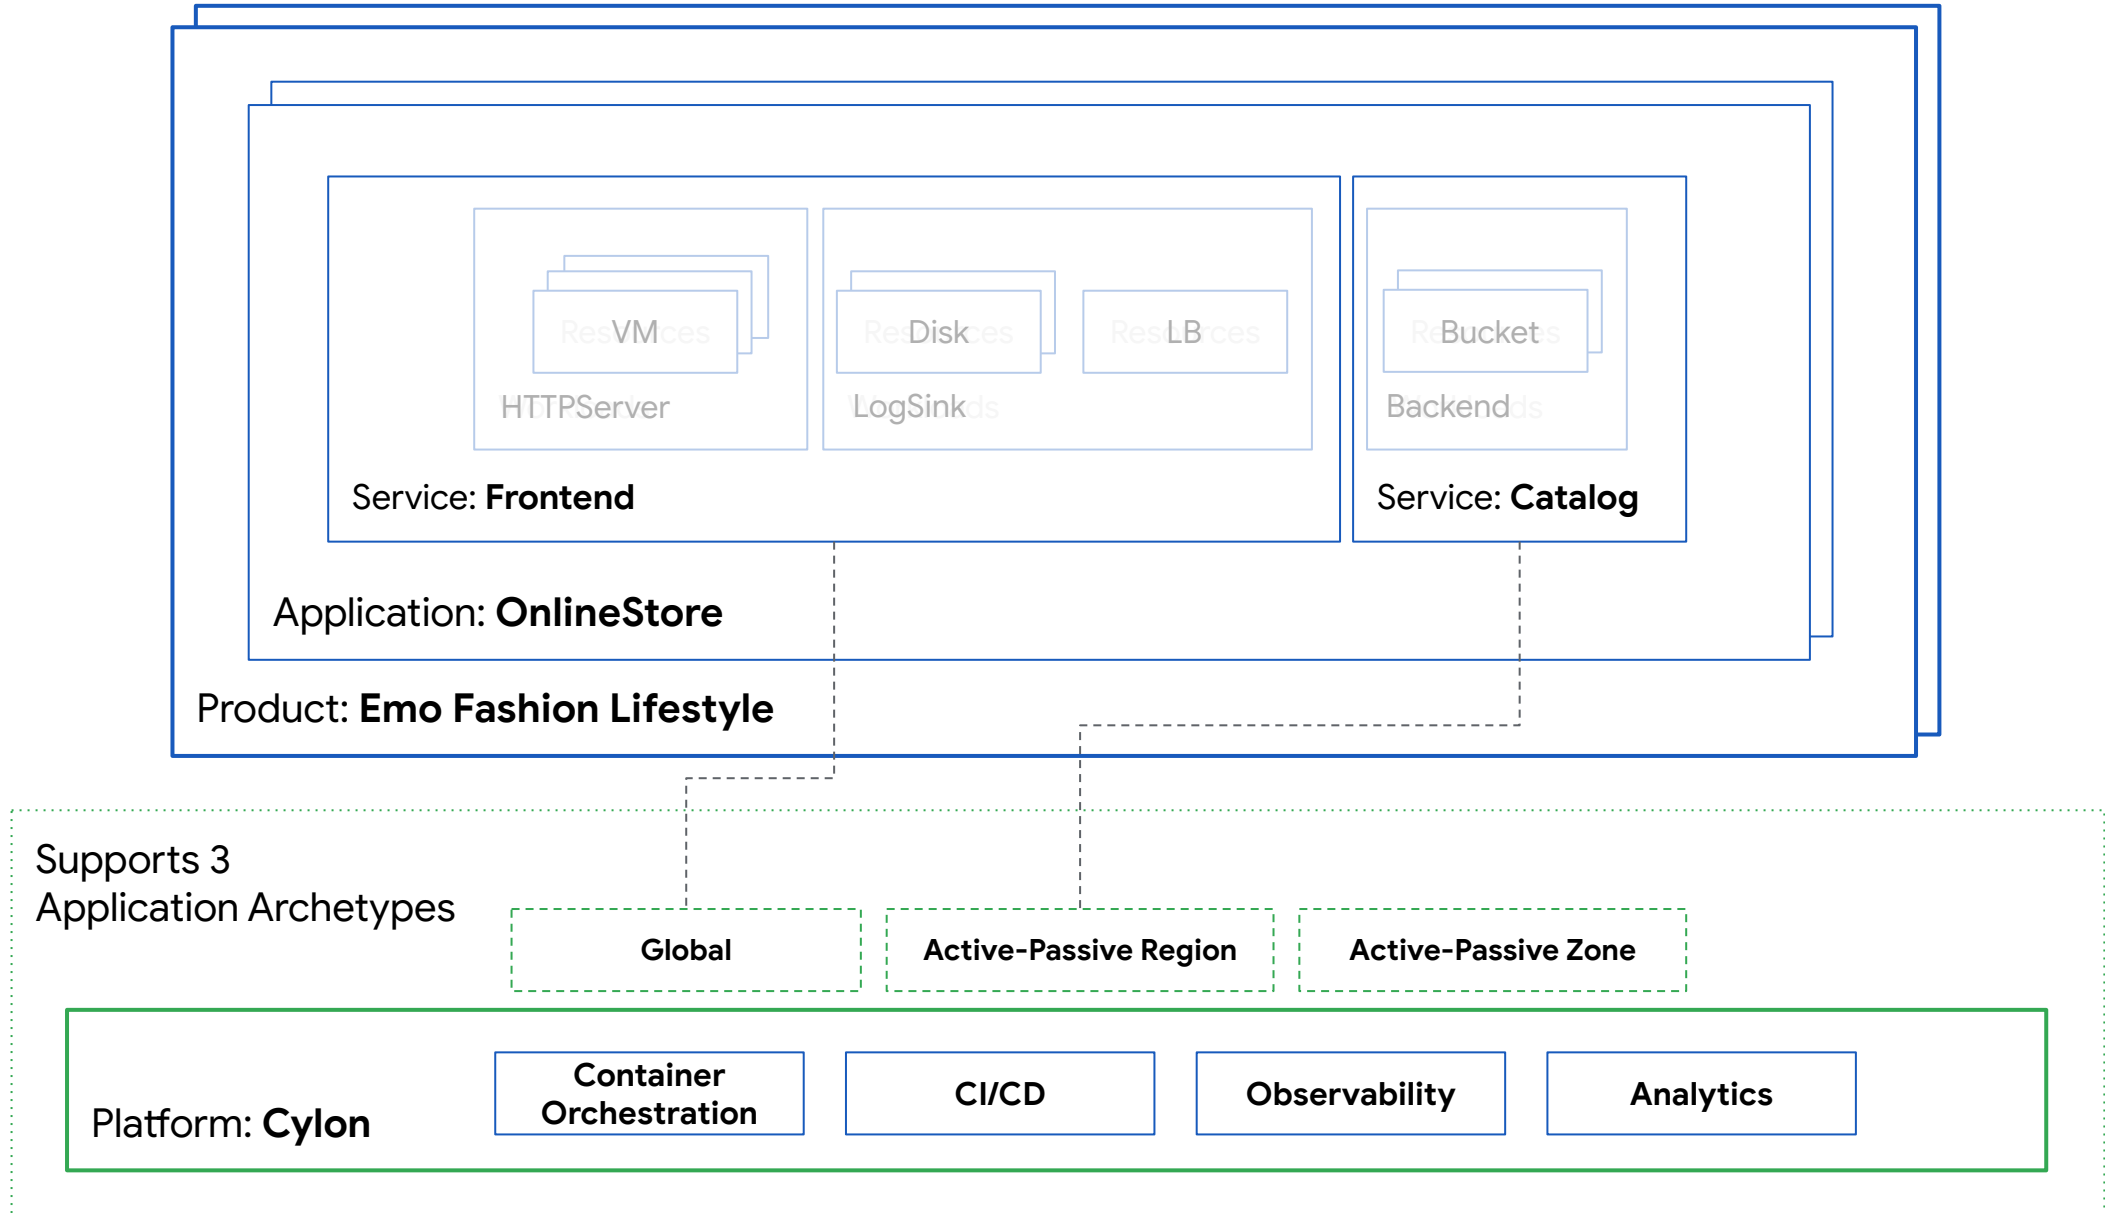
Platform: **Cylon**

Container
Orchestration

CI/CD

Observability

Analytics

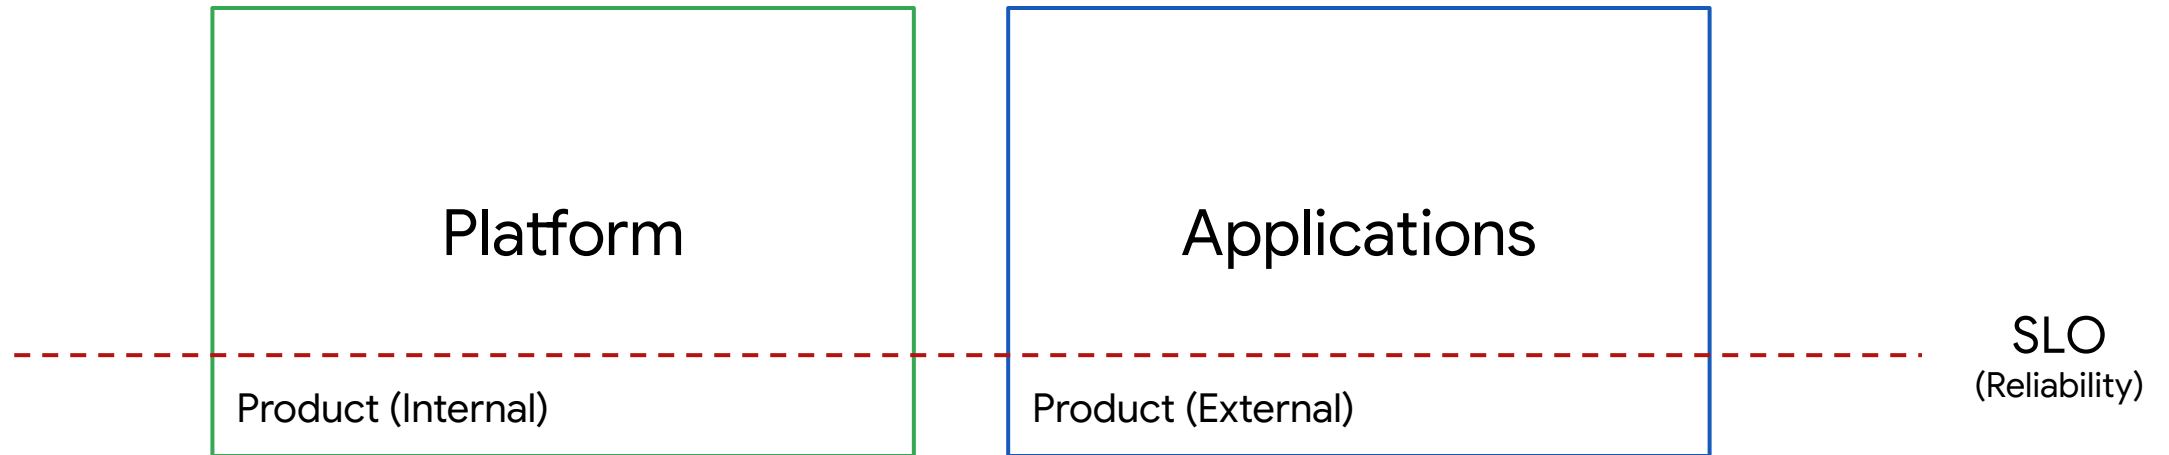


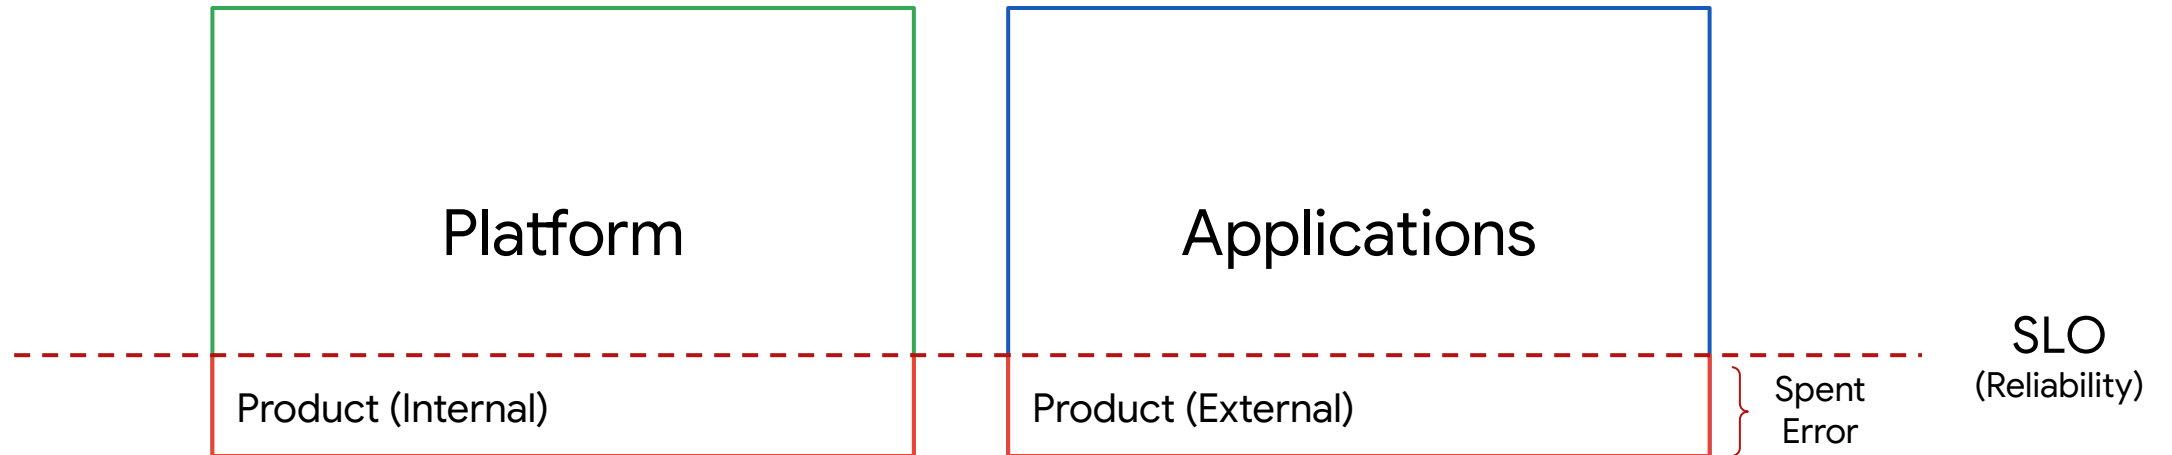
Platform

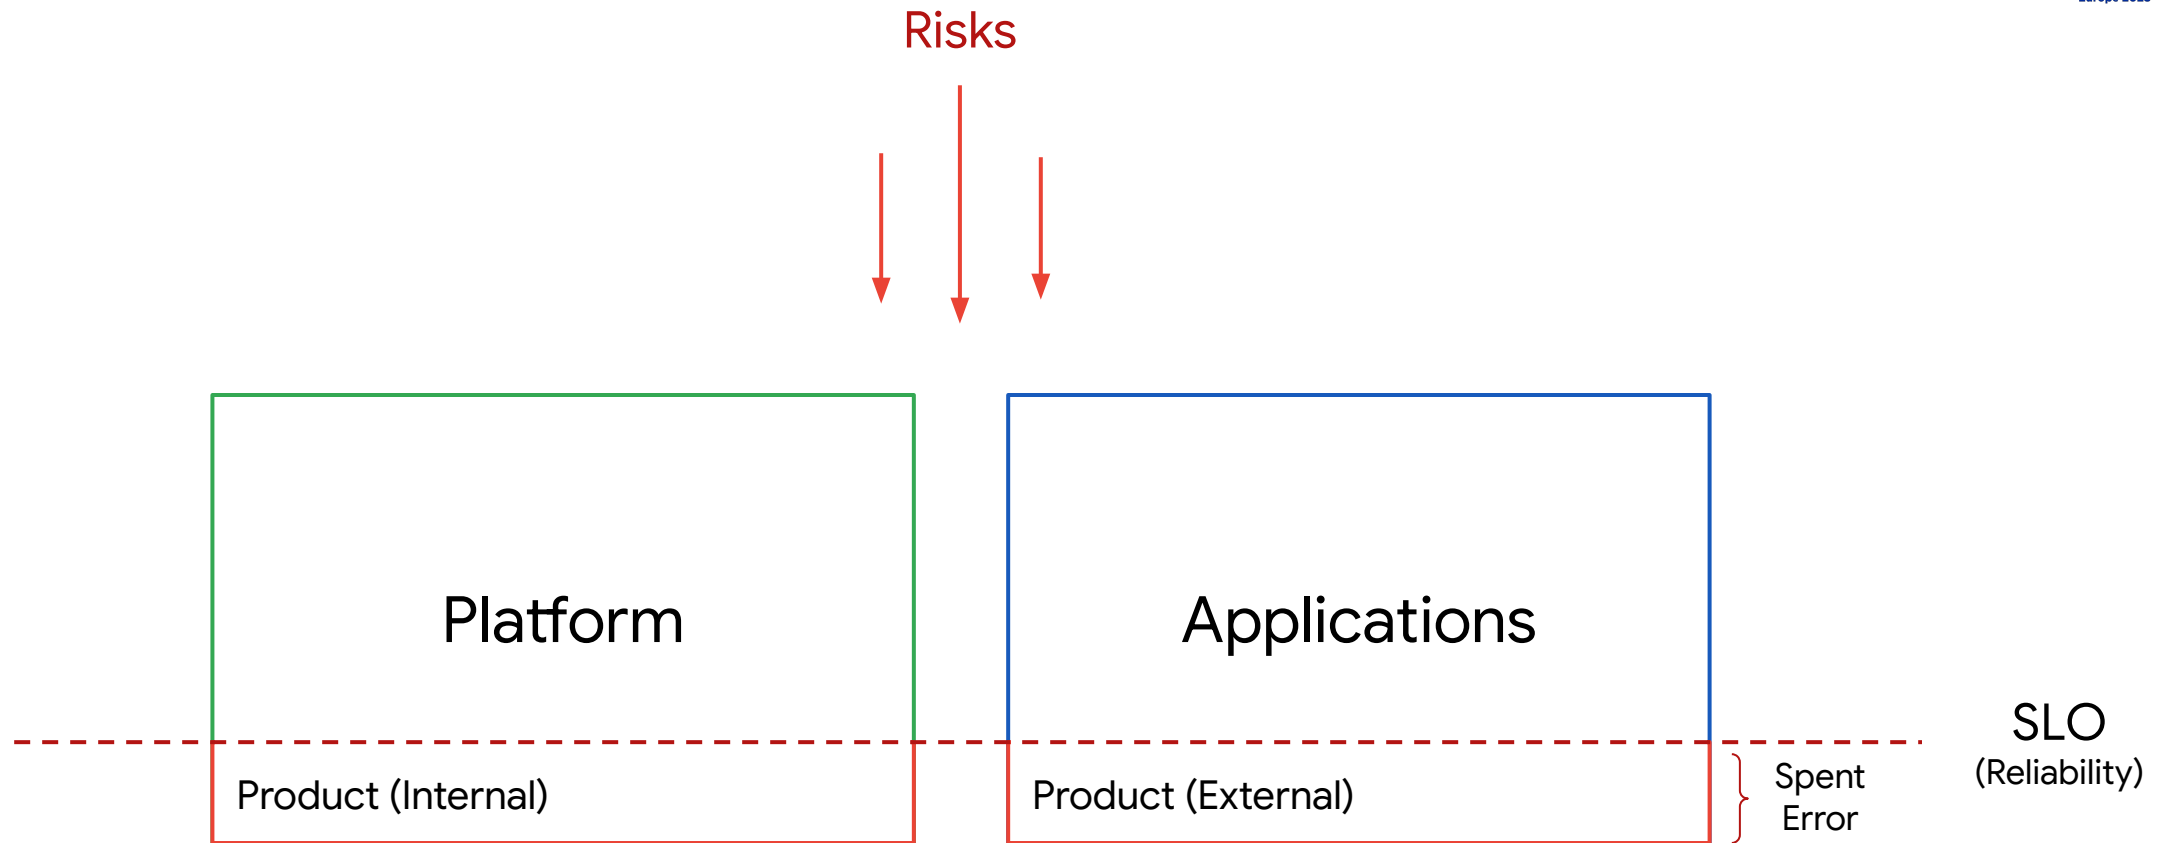
Product (Internal)

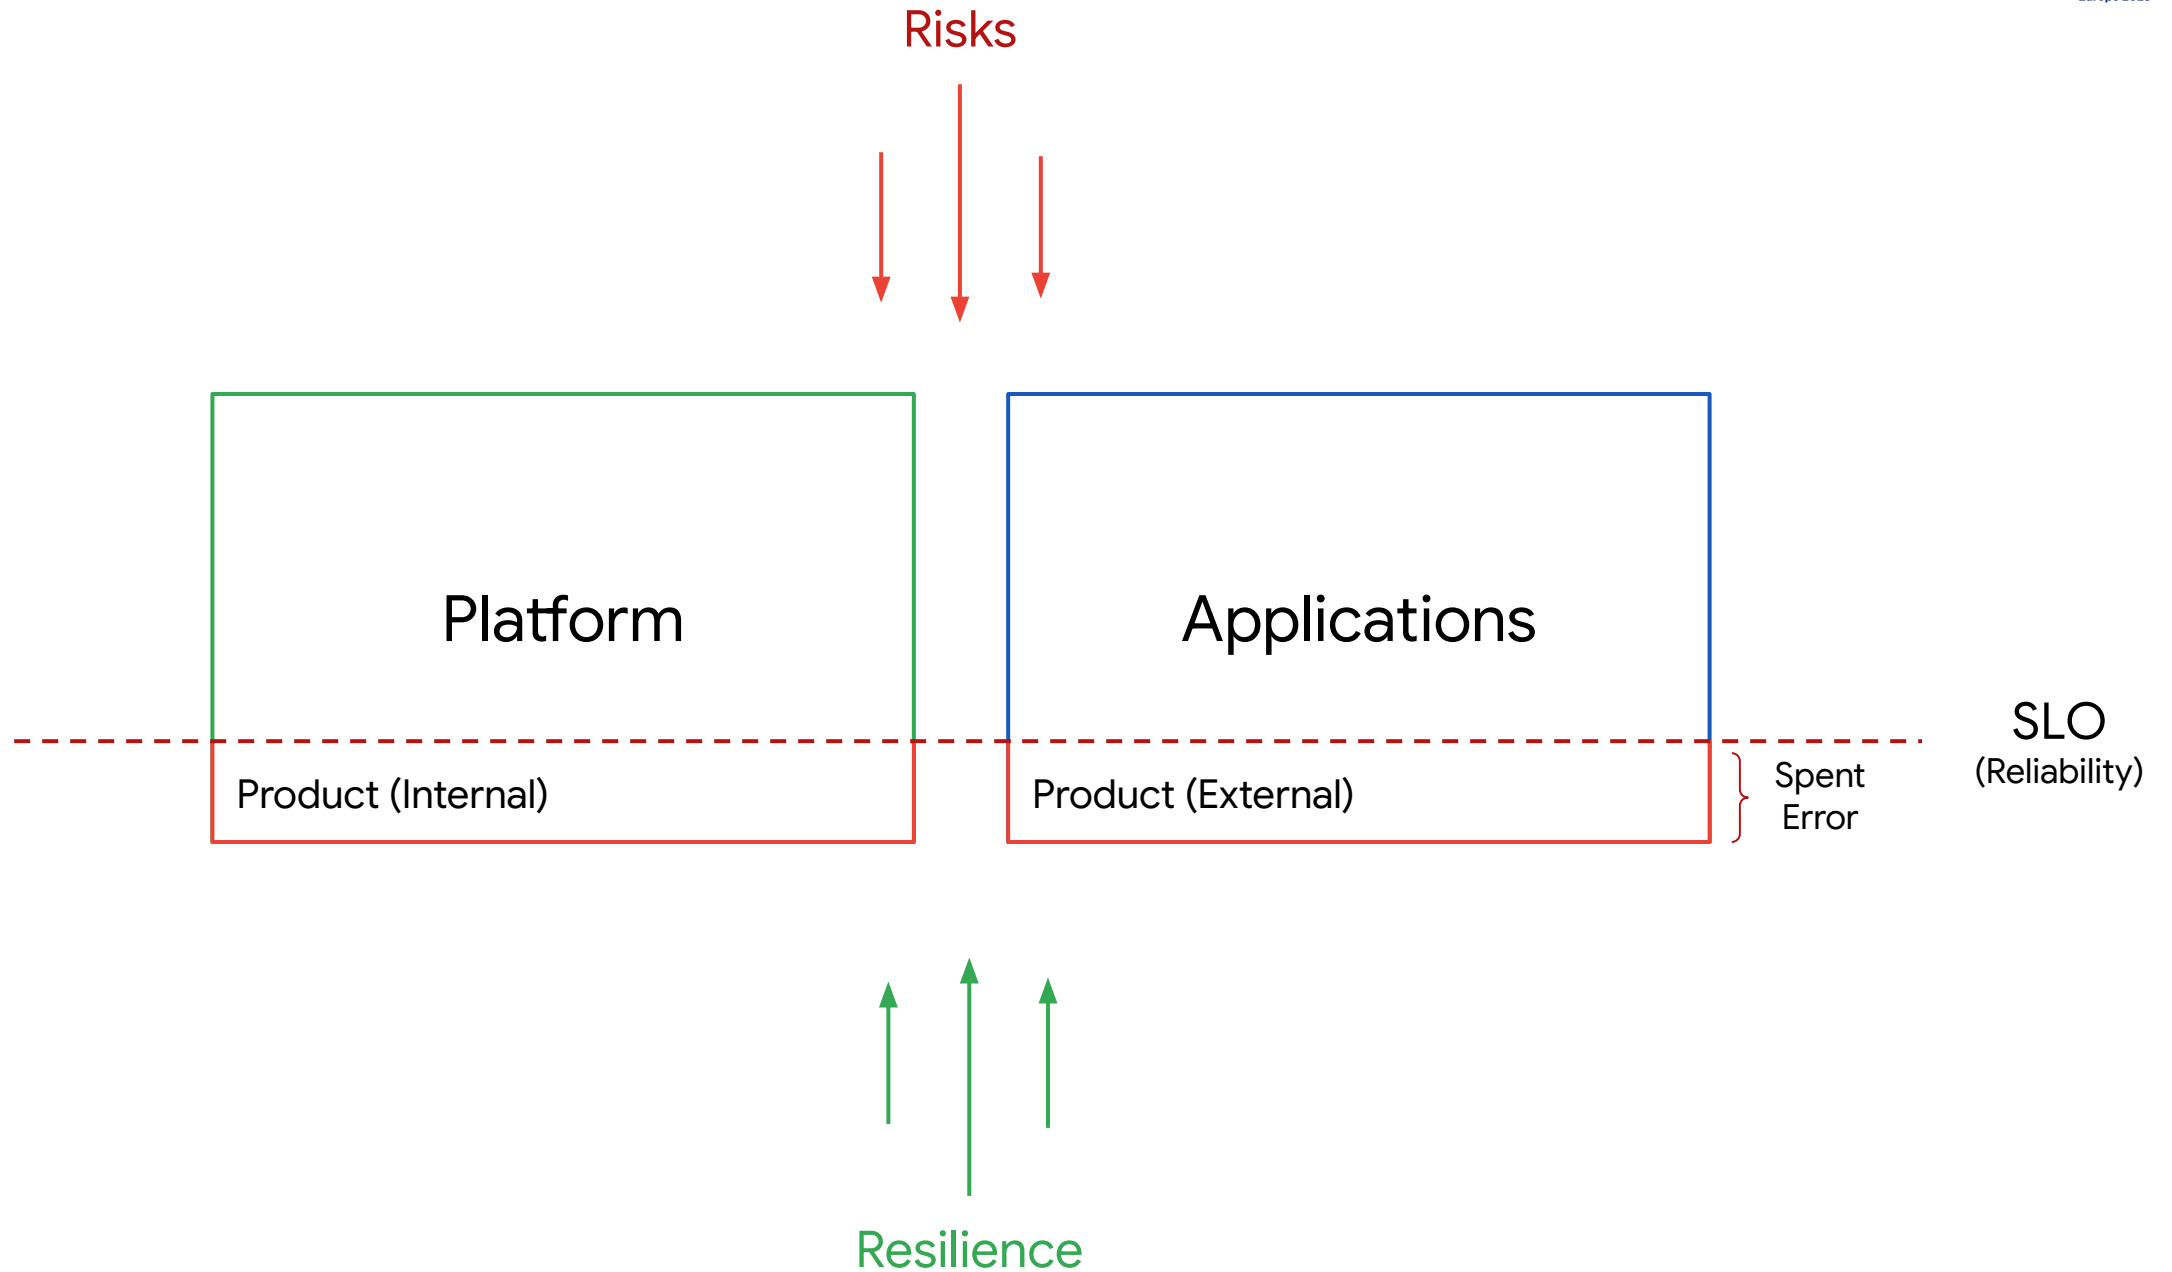
Applications

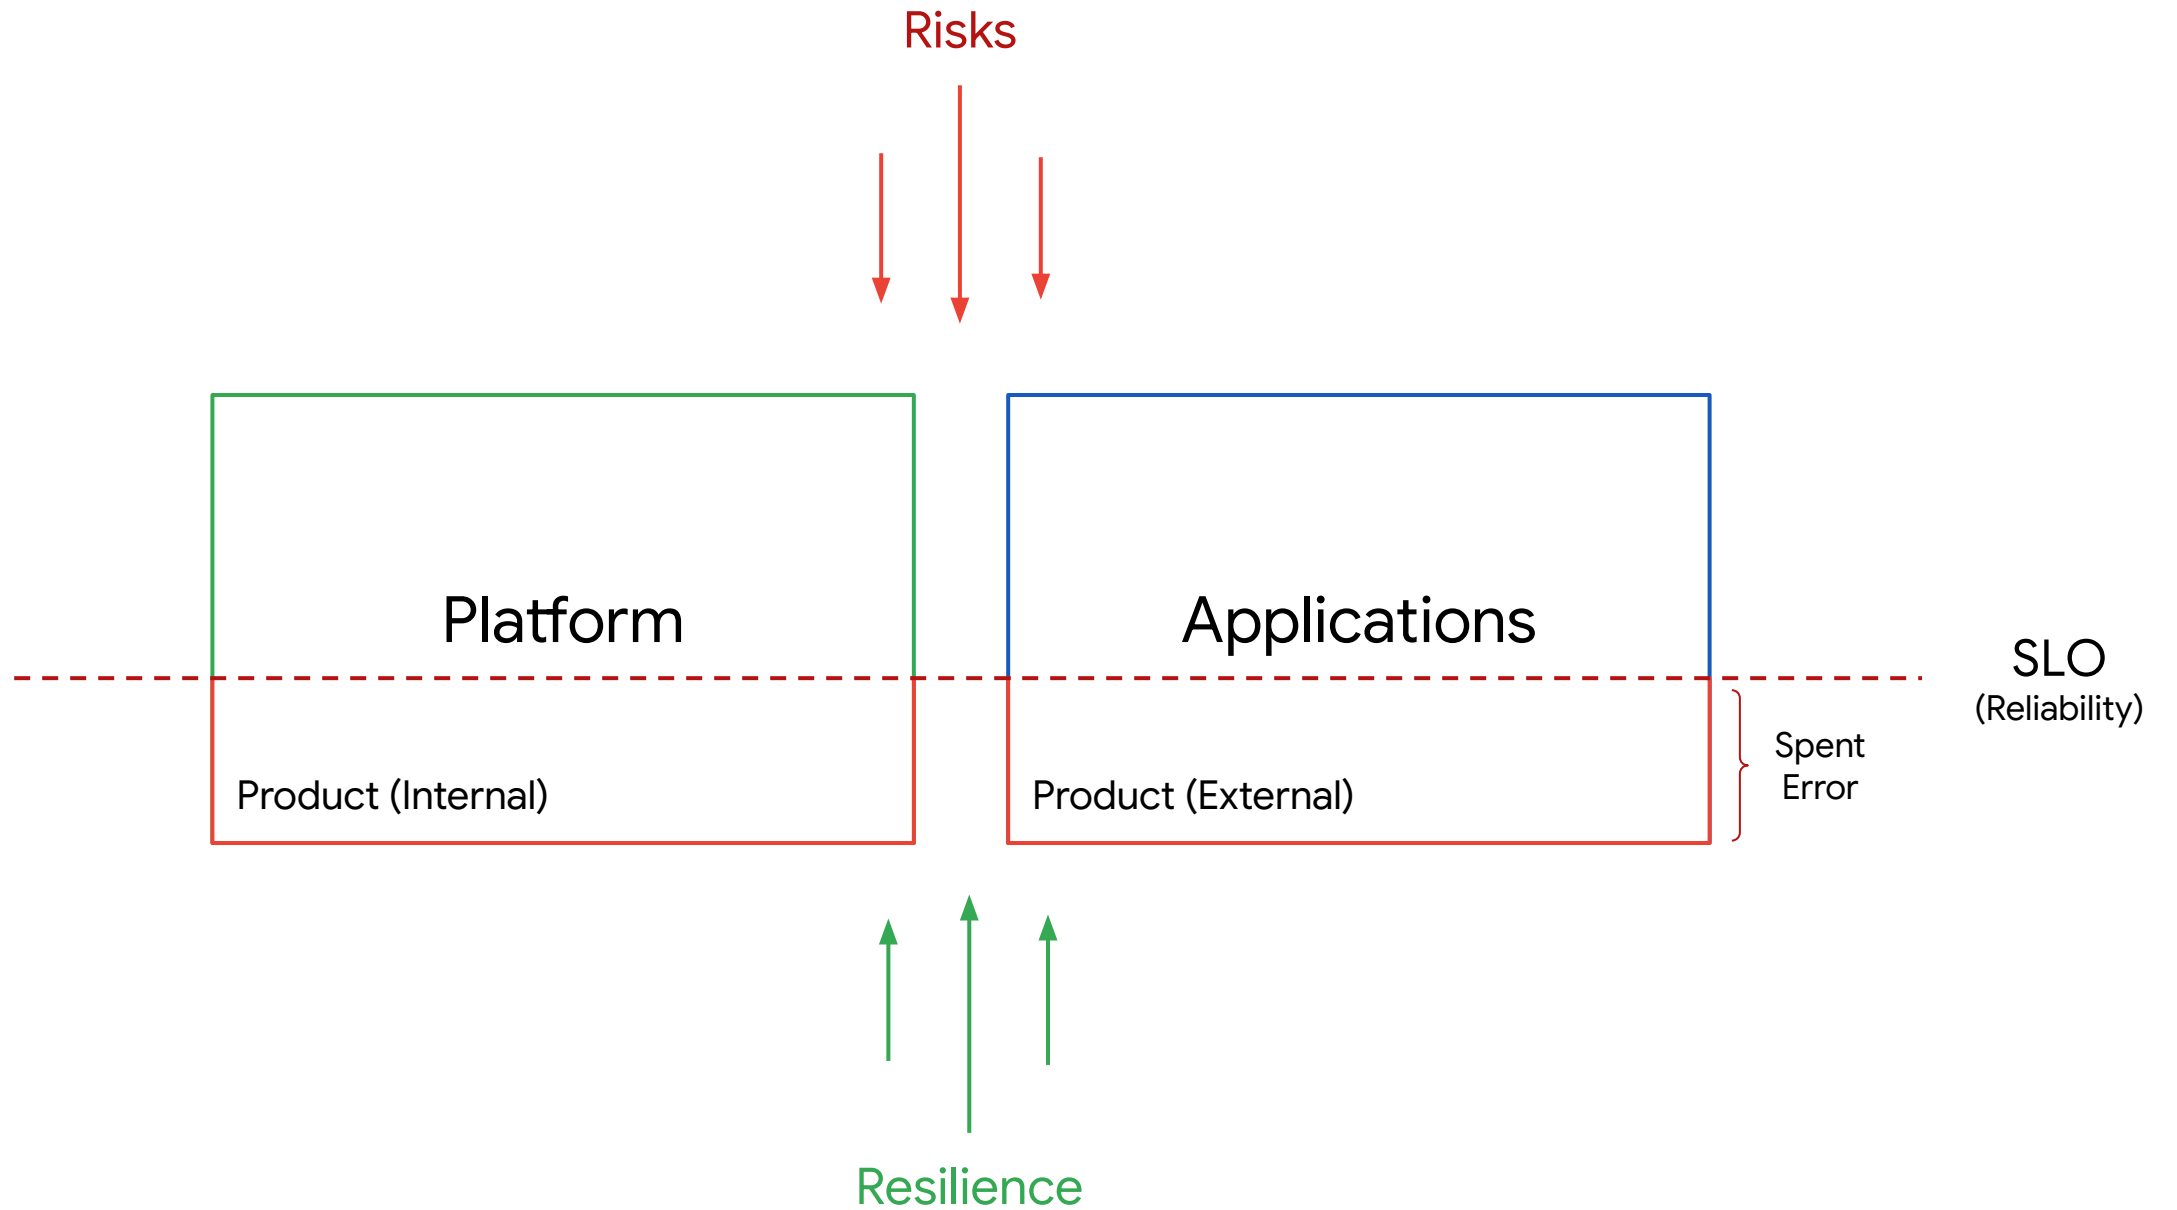
Product (External)

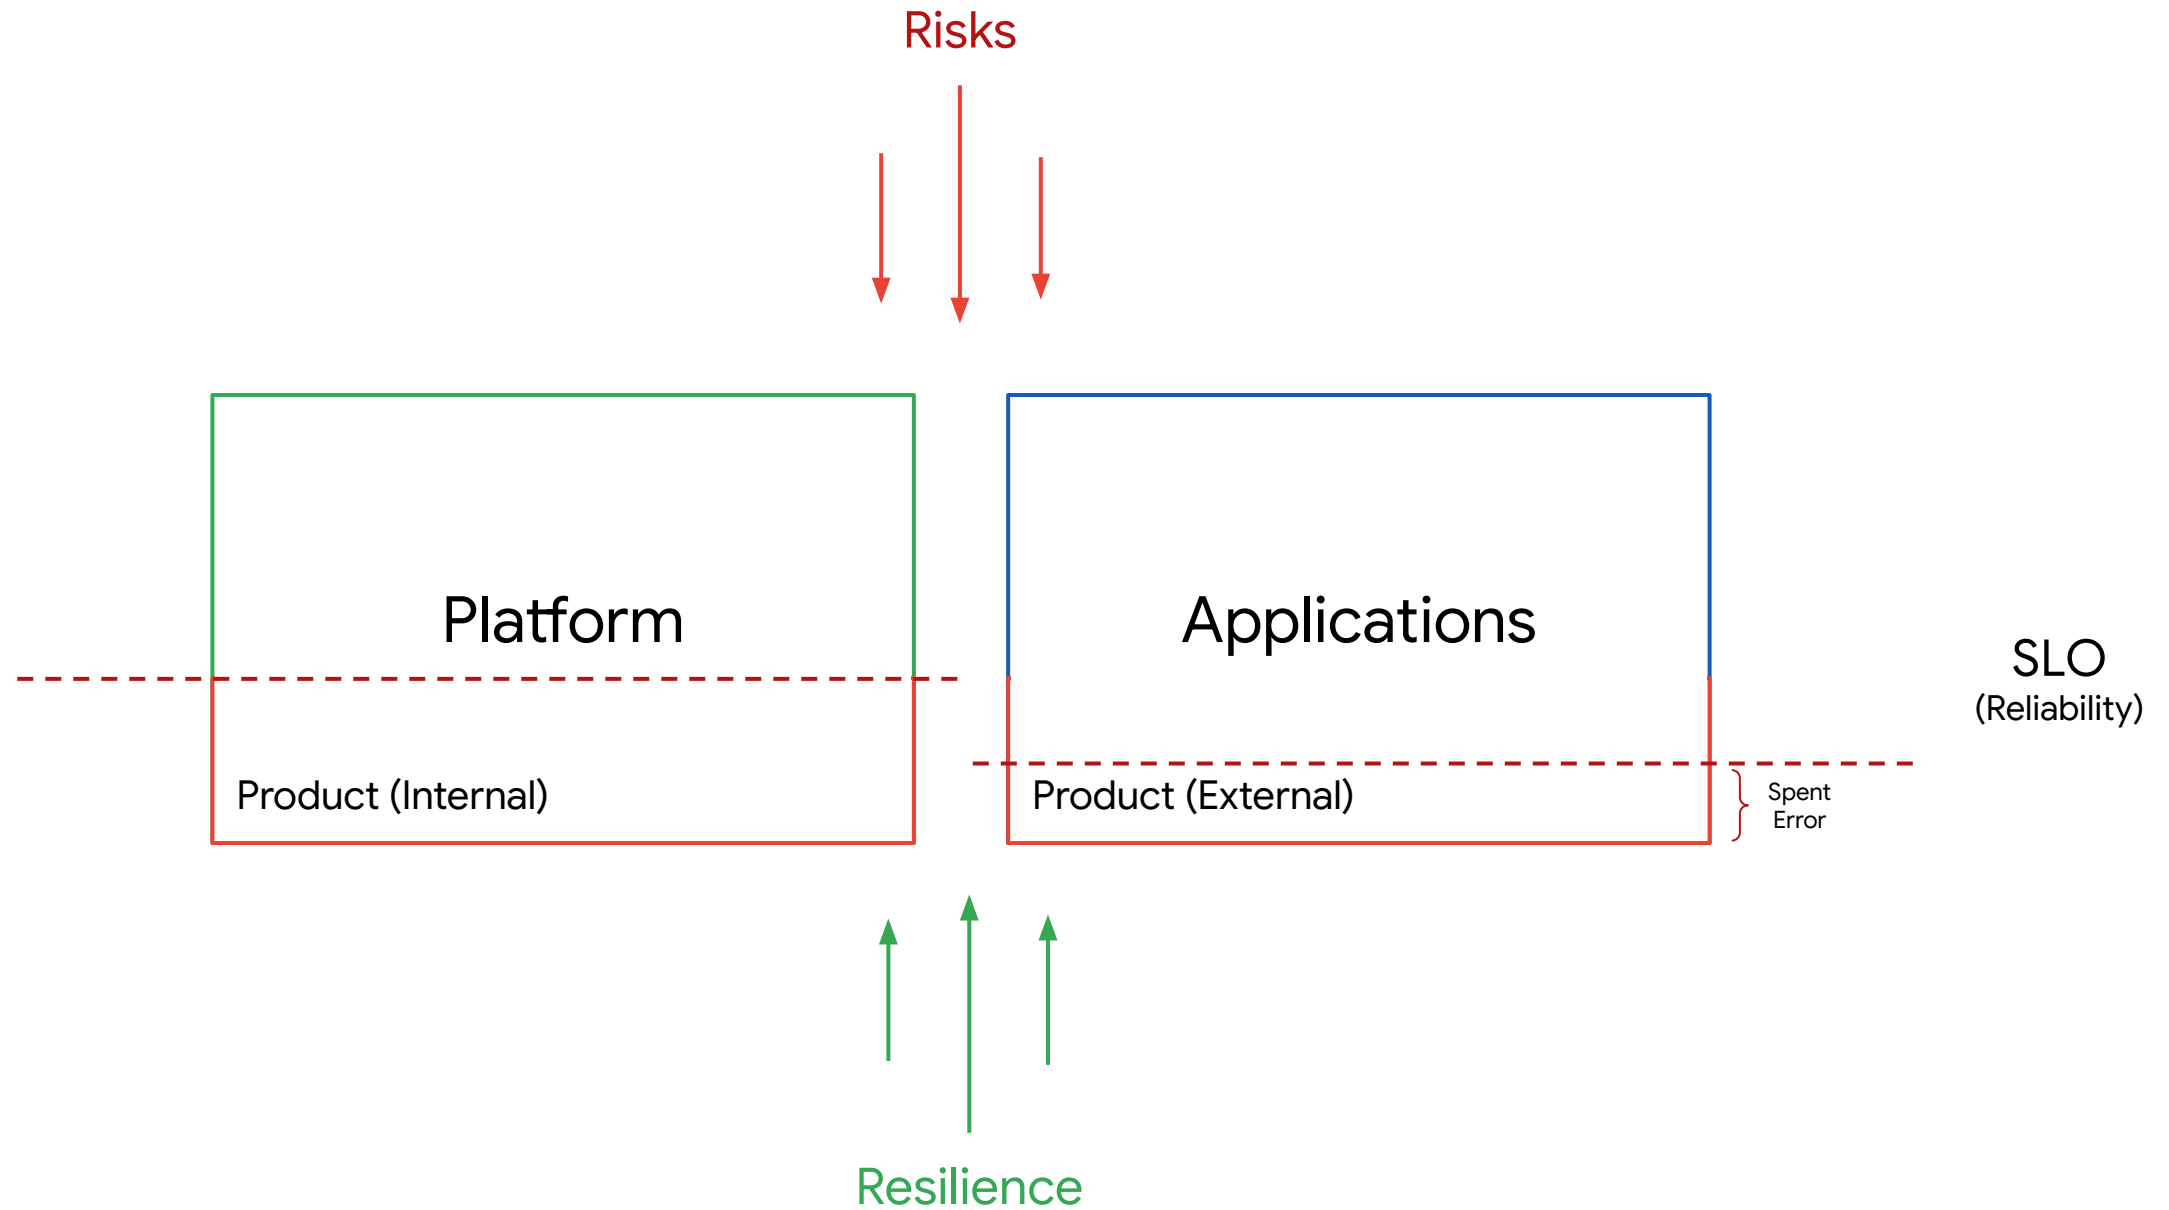












5 Application Archetypes

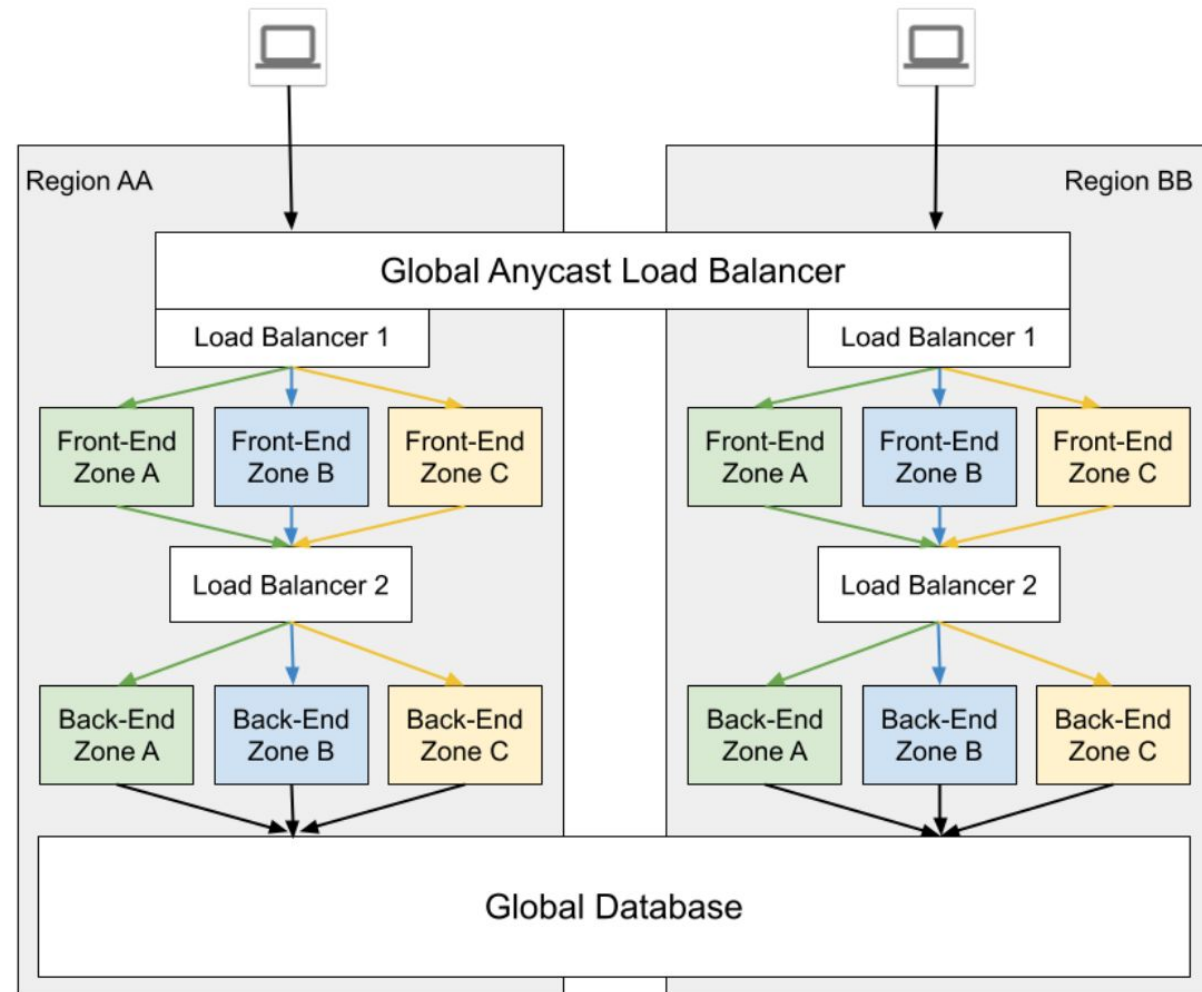


5 Application Archetypes

bit.ly/cloudarchetypes

Anna Berenberg, Brad Calder

ACM Computing Surveys, vol. 55 (2022), pp. 1-48

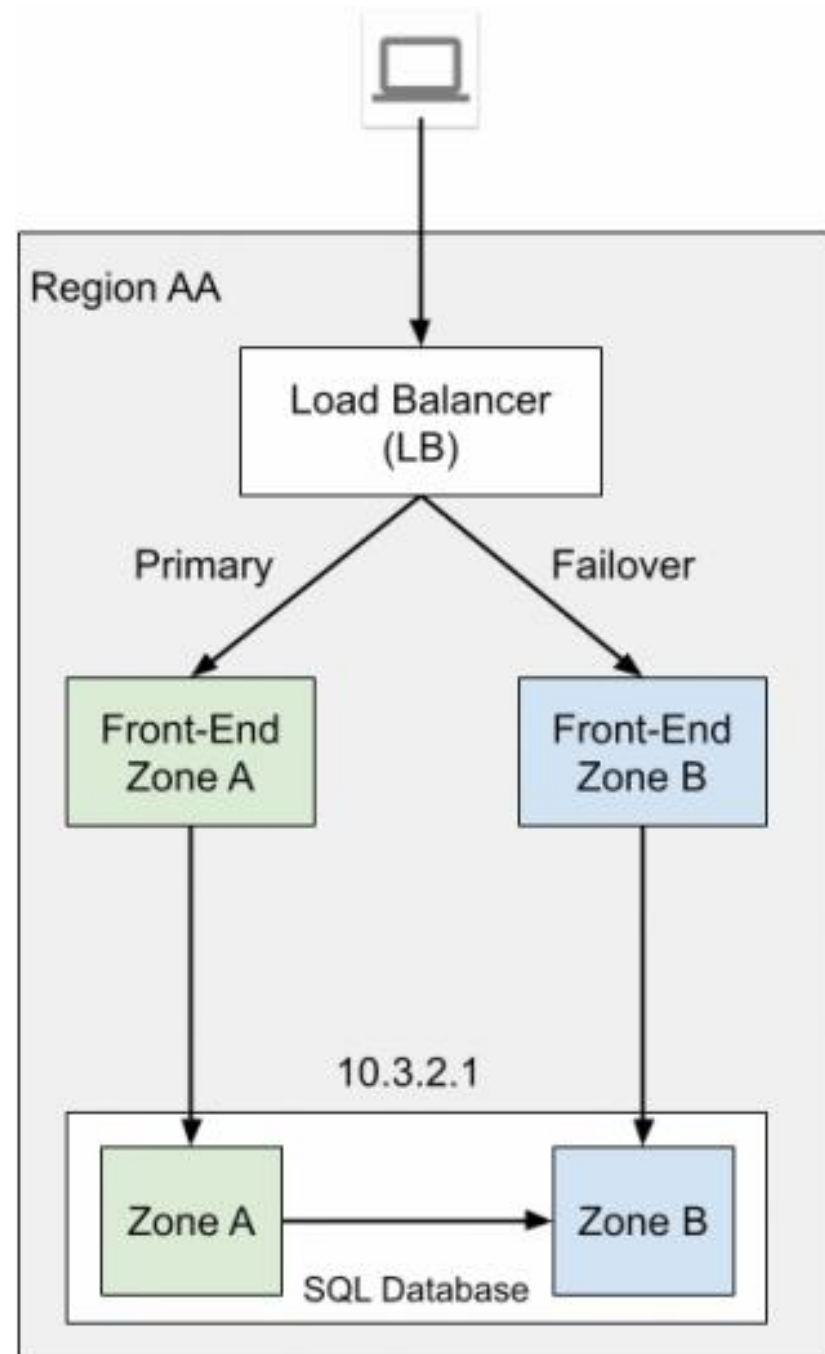


Archetype 5.2: Global Anycast with regional isolated stacks and global database deployment model

Archetype 2.2

Active Passive Zones

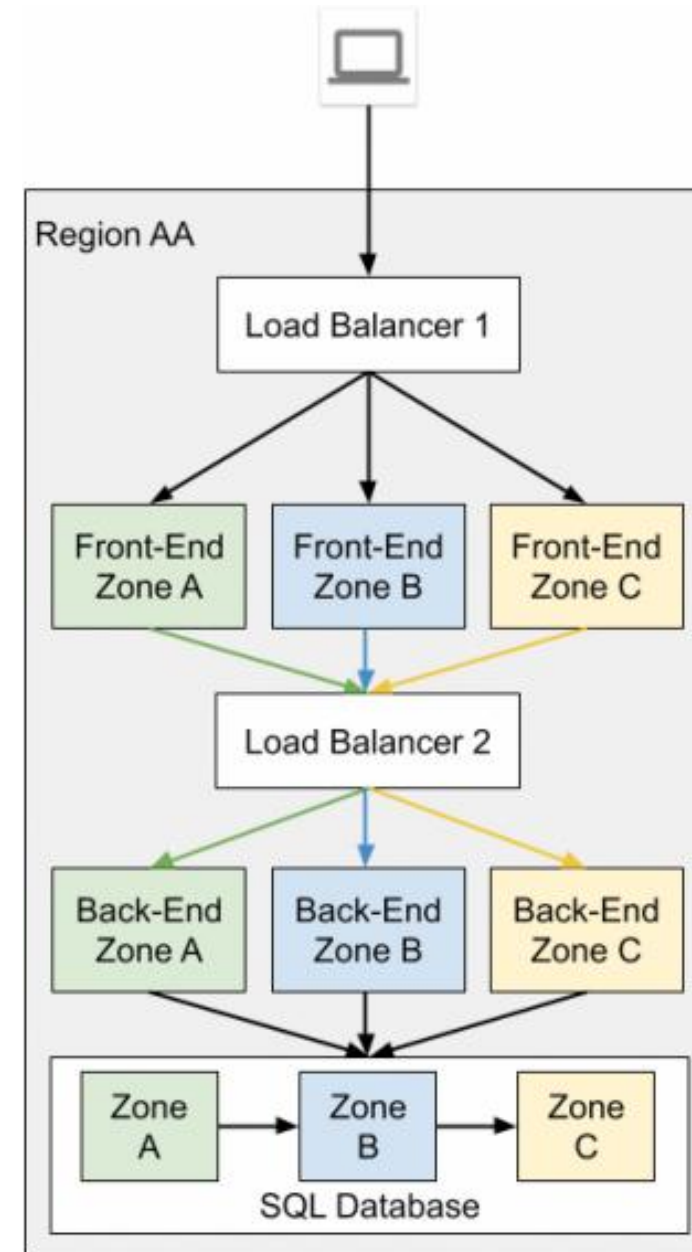
- **Deploy** all services of app to two zones in one region
- **Data** in SQL with a [read replica](#)
- **L4 LB** with one backend group
- **Survives zone** failure. Does not survive region failure.
- **Fail-Ops:** Change LB backend, [promote](#) read replica
- **Cost:** 2x serving + 2x data (1 replica)
- **Complexity:** Low
- **App Refactoring:** None (lift and shift)
- **Type:** COTS, licensing



Archetype 3.1

Multi Zonal


- **Deploy** all services of app to all three zones in one region
- **Data:** Use HA SQL
- Use **Global LB or Reg LB** with 3 backend groups
- **Survives zone** failure. Does not survive region failure.
- **Fail-Ops:** Initiate DB failover
- **Cost:** 1.5x serving + 2x data (HA SQL)
- **Complexity:** Medium
- **App Refactoring:** Low (multi instance)
- **Type:** Web services



Archetype 3.2

Active Passive Region

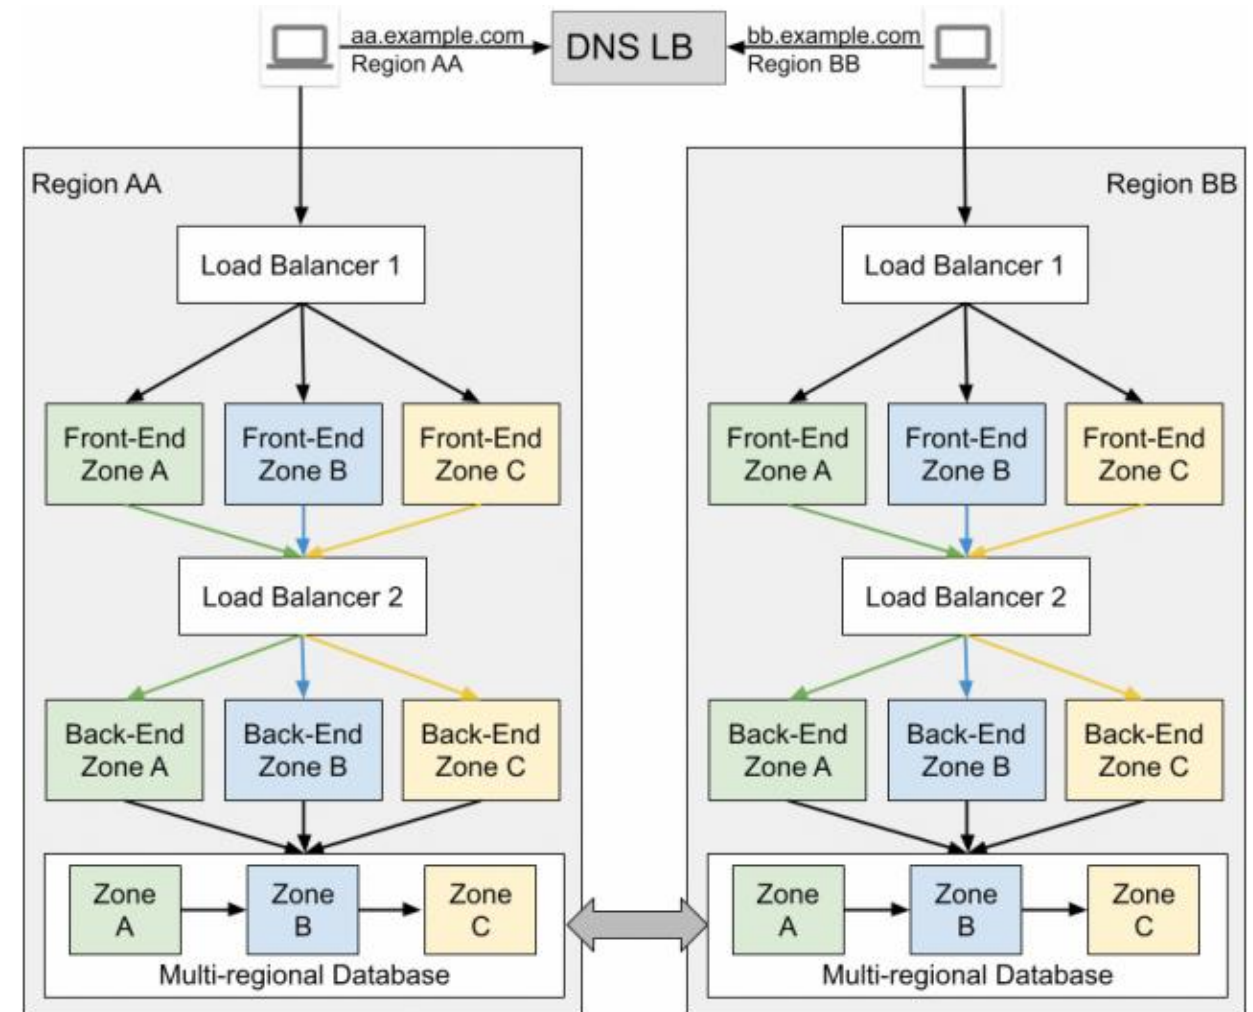
- **Deploy** all services of app to all three zones in each of two regions
- **Data:** SQL with cross-region replication
- **DNS** points at one LB (until disaster)
- **Survives zone and region** failures
- **Fail-Ops:** No action for zone failure.
 - Update DNS to point at standby LB
 - Cross region DR failover process for DB
- **Cost:** 3x serving + 2x data (HA SQL)
- **Complexity:** Medium
- **App Refactoring:** Medium (multi instance, multi regional data)
- **Type:** HA web services



Archetype 4.3

Isolated Regions

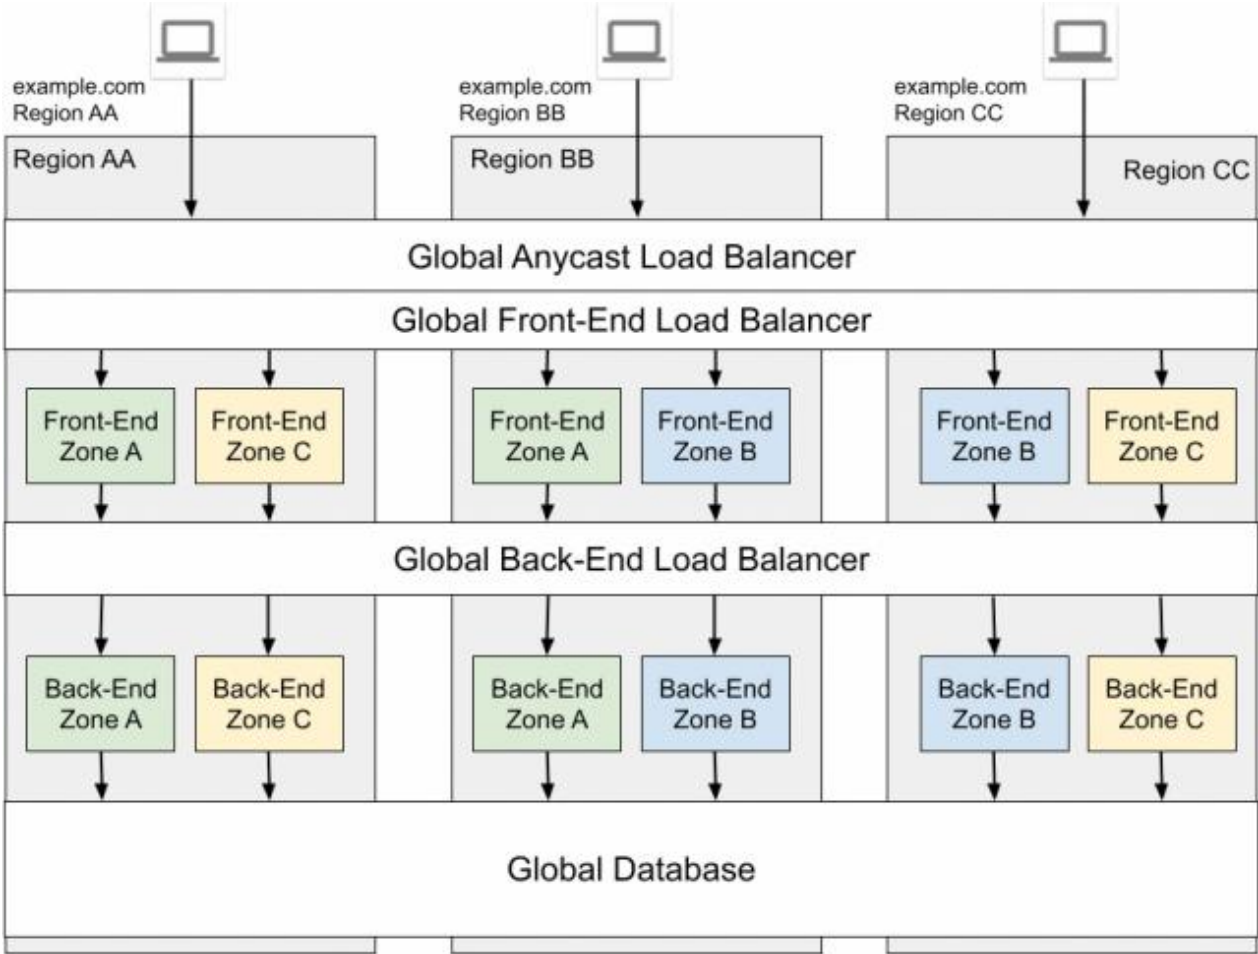
- **Deploy** all services of app to all three zones in each of two regions
- **Data:** Spanner or CockroachDB
- **DNS** points at two Regional LBs
- **Survives zone and region** failures. No impact for ½ consumers. Possible manual failover
- **Fail-Ops:** No action for zone failure. Optional regional failover like Arch 3.2
- **Cost:** 1.5 cost per region for zone failure
- **Complexity:** Medium/High
- **App Refactoring:** Medium (multi instance, multi regional data)
- **Type:** Regulated HA services



Archetype 5.2

Global


- **Deploy** all services of app to all three zones in each of two or more regions
- **Data:** Spanner or CockroachDB
- **Global LB** points at regional backend groups
- **Survives zone and region** failures
- **Fail-Ops:** None
- **Cost:** N+m cost modelling. Global DBs are more expensive
- **Complexity:** High
- **App Refactoring:** High (multi instance, global DBs)
- **Type:** Global consumer services



How to use Archetypes?

01


Services can be
deployed to a **single**
archetype



How to use Archetypes?

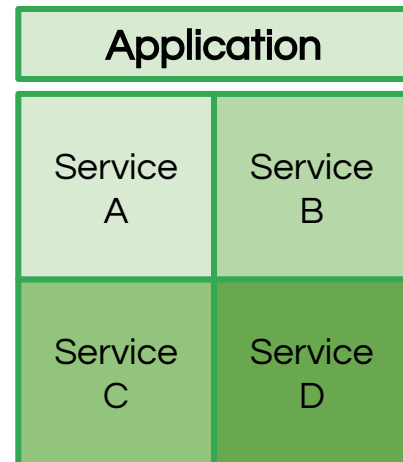
01

Services can be
deployed to a **single**
archetype



02

Application can use
services across
multiple archetypes



How to use Archetypes?

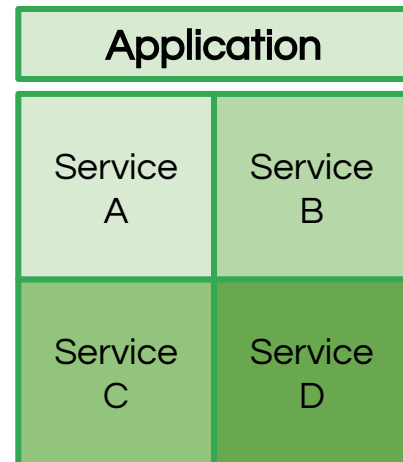
01

Services can be deployed to a **single archetype**




02

Application can use services across **multiple archetypes**



03

Applications should be designed for **graceful degradation**



SLOs and SLIs



SLI

Quantitative measure of some
aspect of the level of service

aka

latency, throughput, availability

SLI

Quantitative measure of some aspect of the level of service

aka

latency, throughput, availability

SLO

a target value or range of values for a service level that is measured by an SLI

aka

99% of **Get RPC calls will complete in less than 100 ms**

SLOs in one slide


A **ratio-rate** of **good/total**, measured over a time duration.


If too much non-good, for too long, tell a human.

SLI is the squiggly line

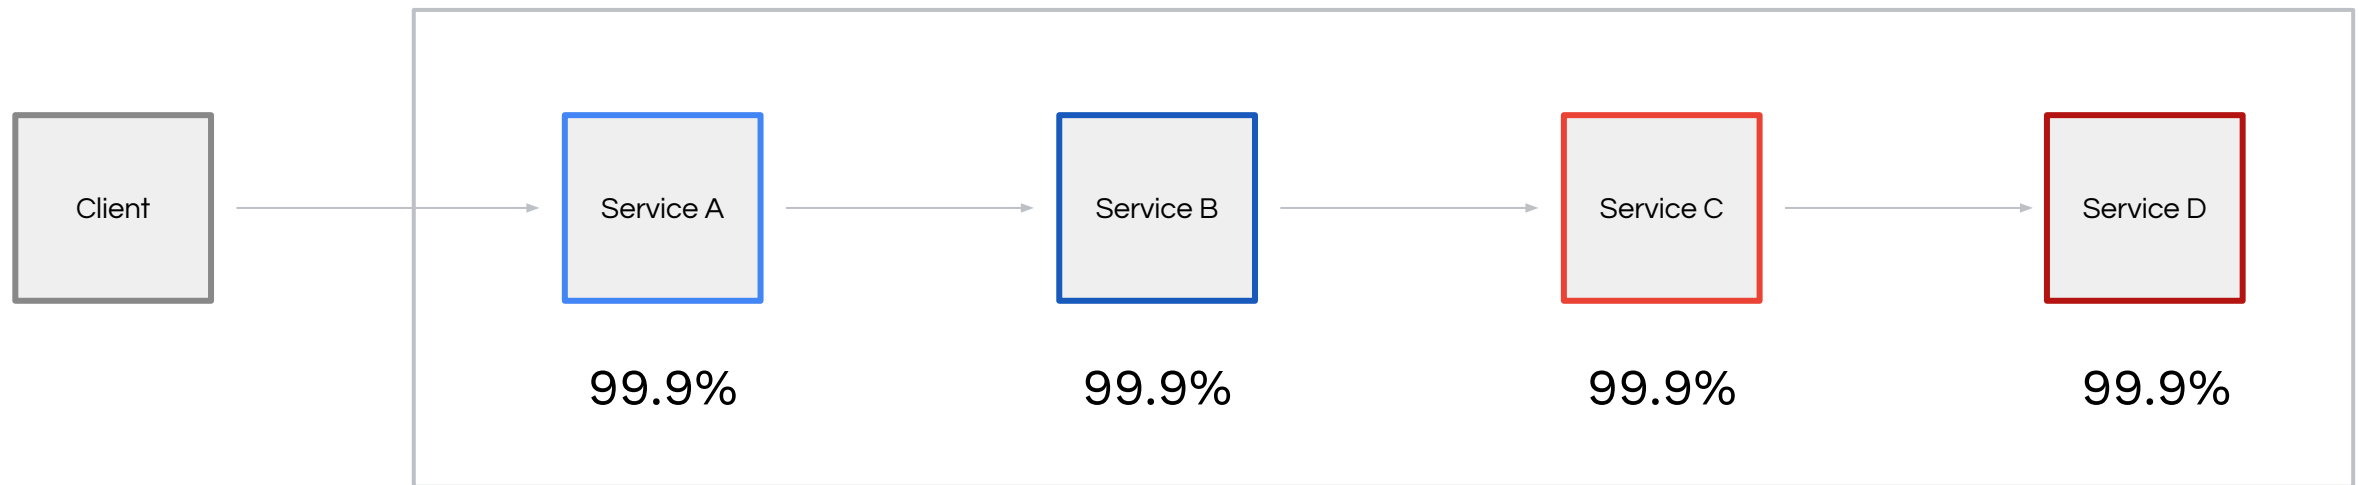
SLO is the straight one

Area is time **exceeding SLO**



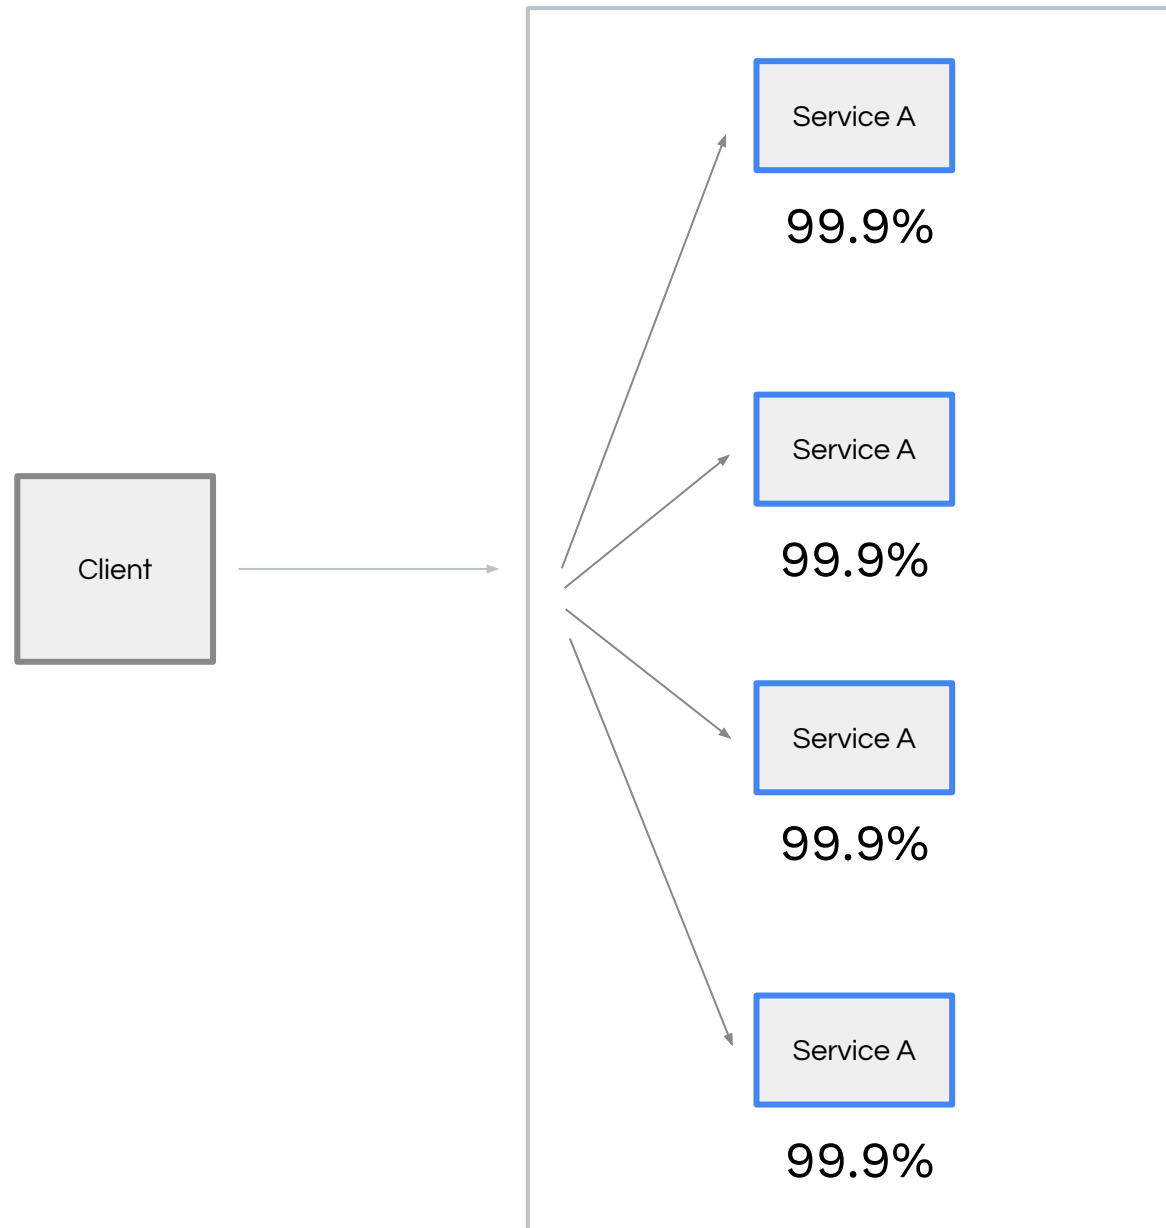


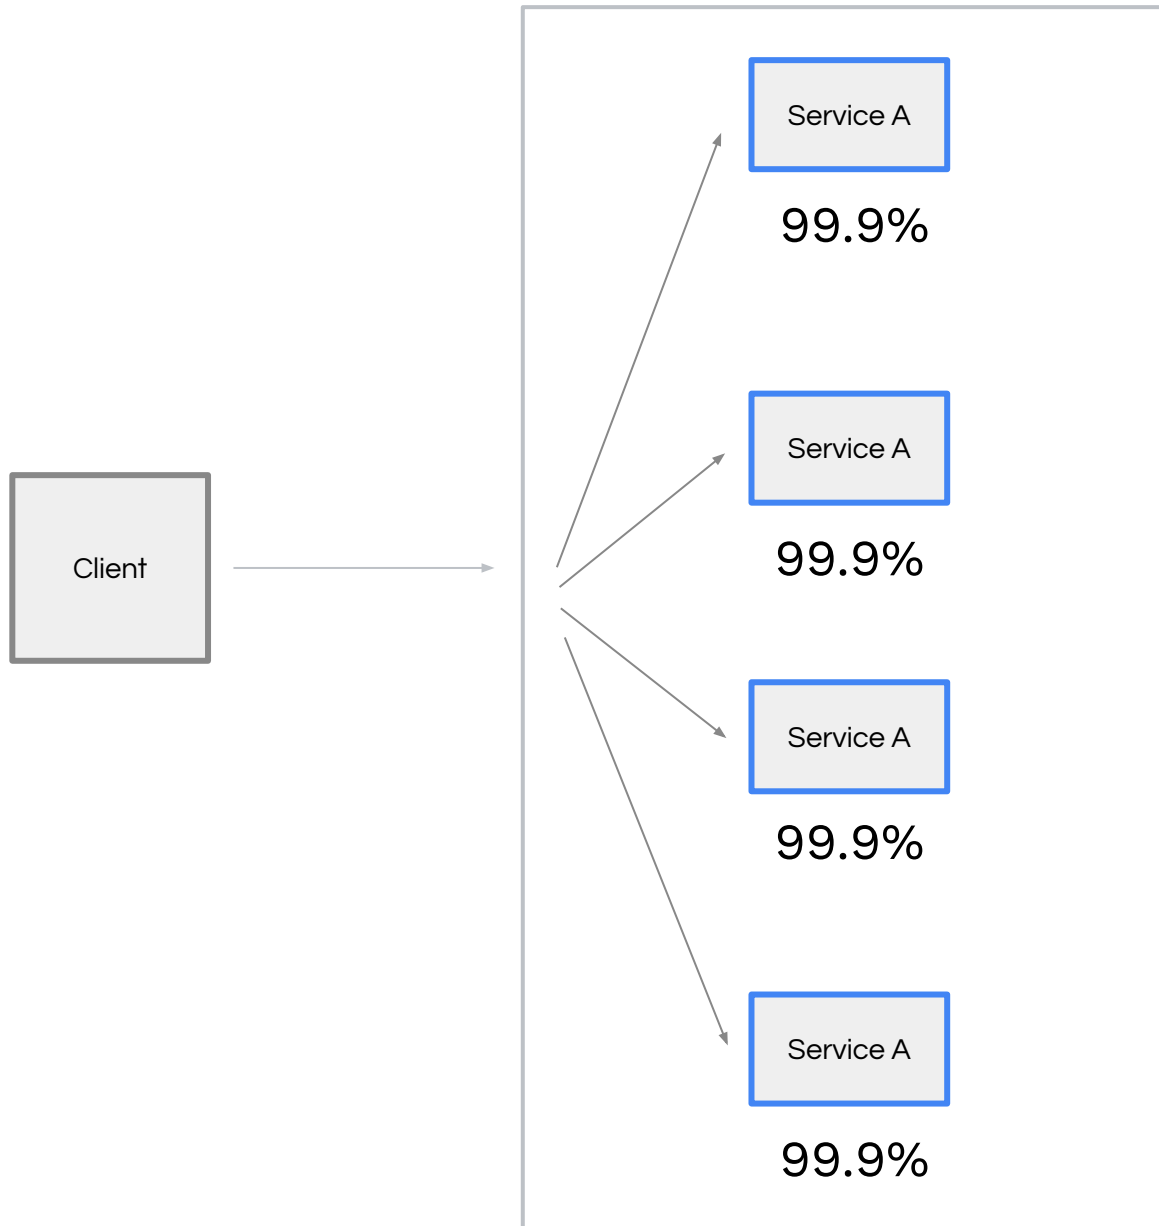
Intersection (or serial)



$$0.999 \times 0.999 \times 0.999 \times 0.999$$

99.6% SLO





Union (aka parallel)

$$1 - (0.001)^4$$

99.9999999999% SLO

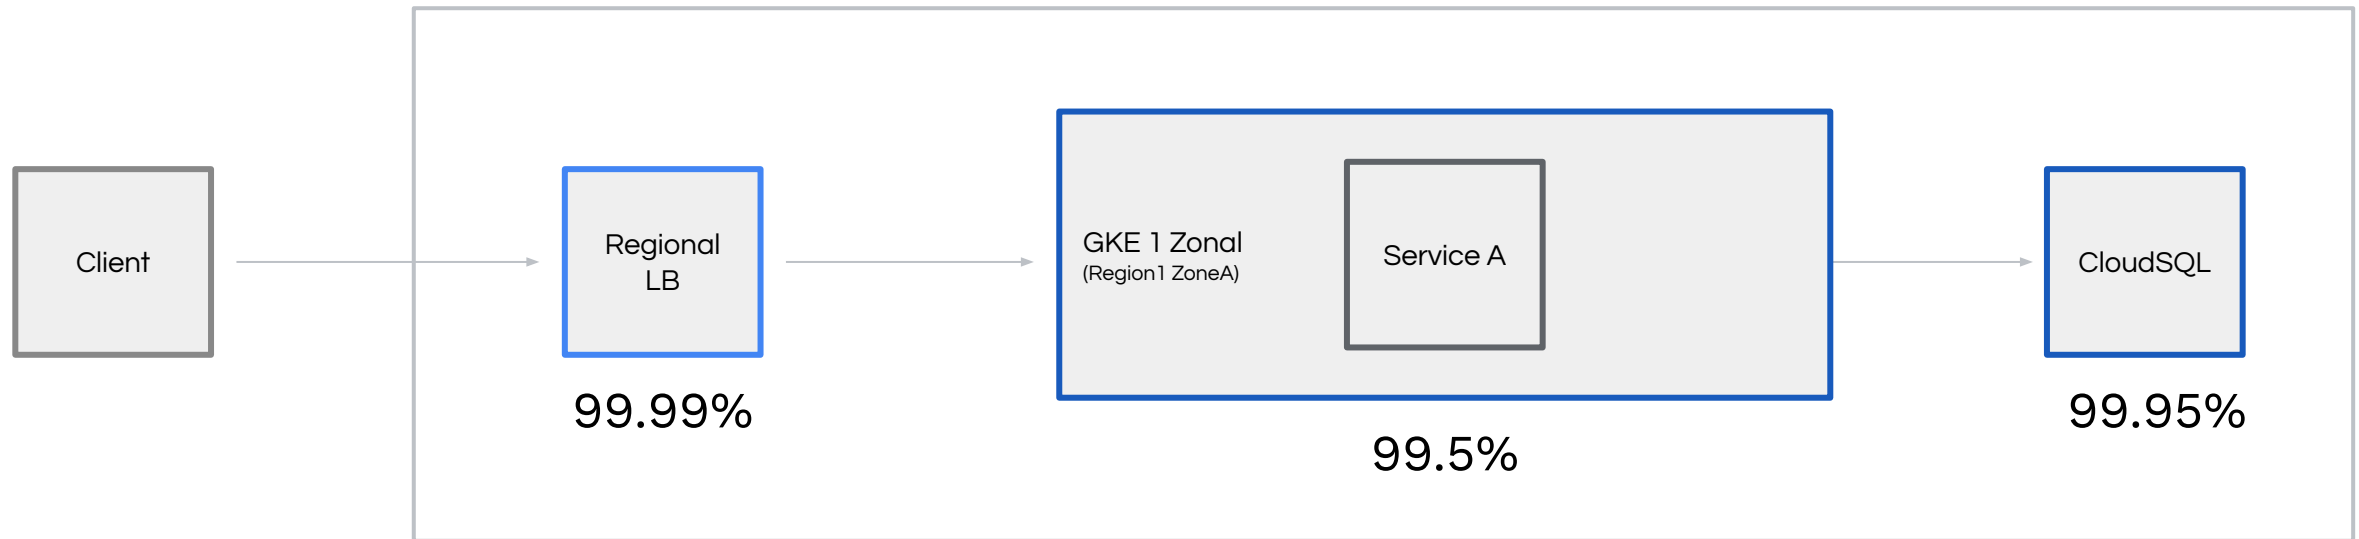
or 11 nines

This is strictly mathematical and does not include any dependent variables like network, LBs, capacity planning, connectivity, and other dependent services

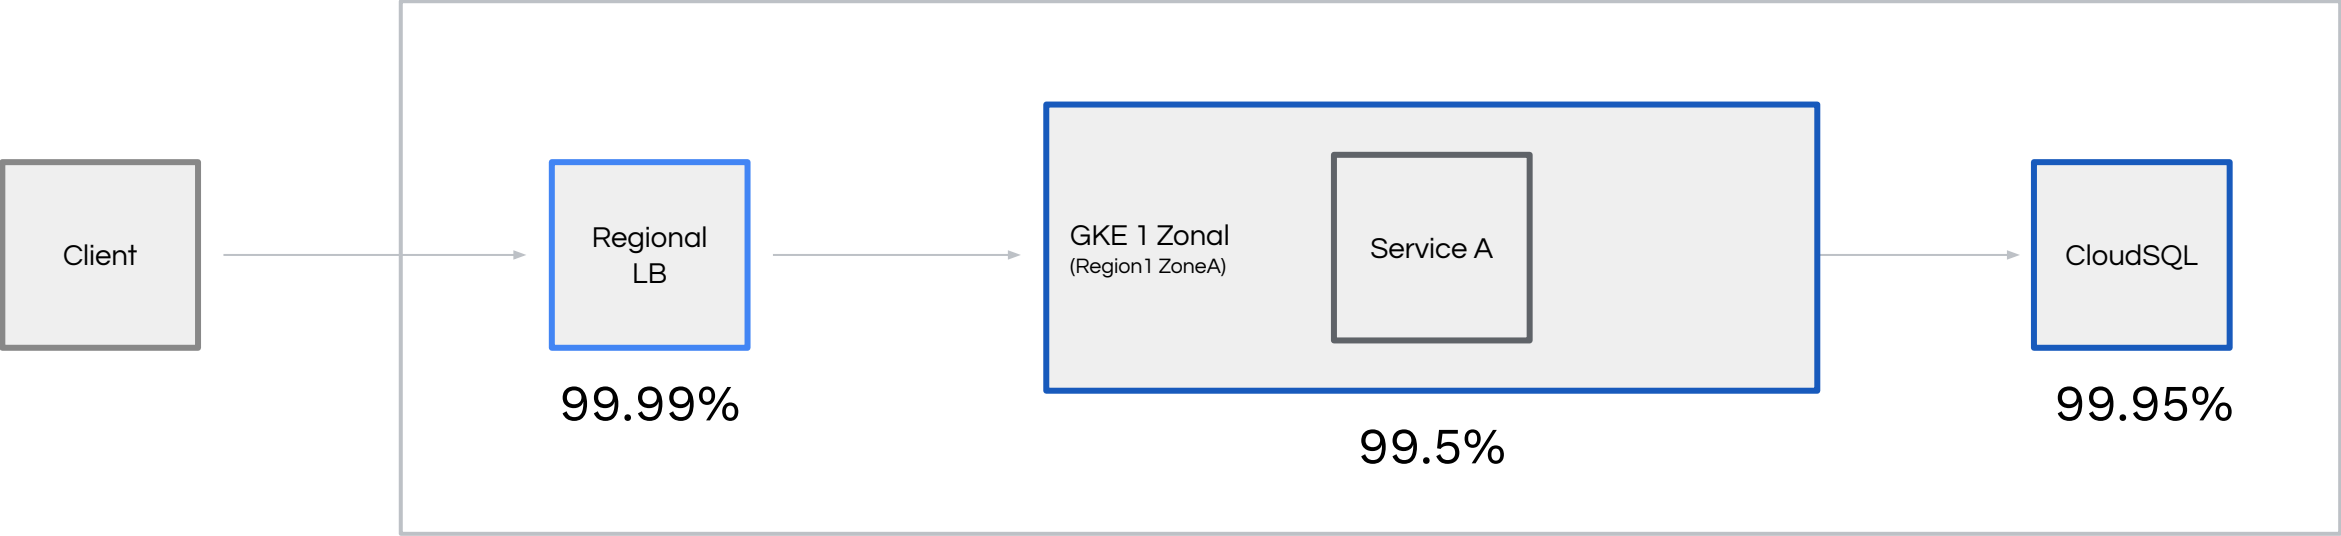
Building Reliable Platforms on Kubernetes



Archetype 2.1 Single zonal GKE cluster with Cloud SQL



Archetype 2.1 Single zonal GKE cluster with Cloud SQL



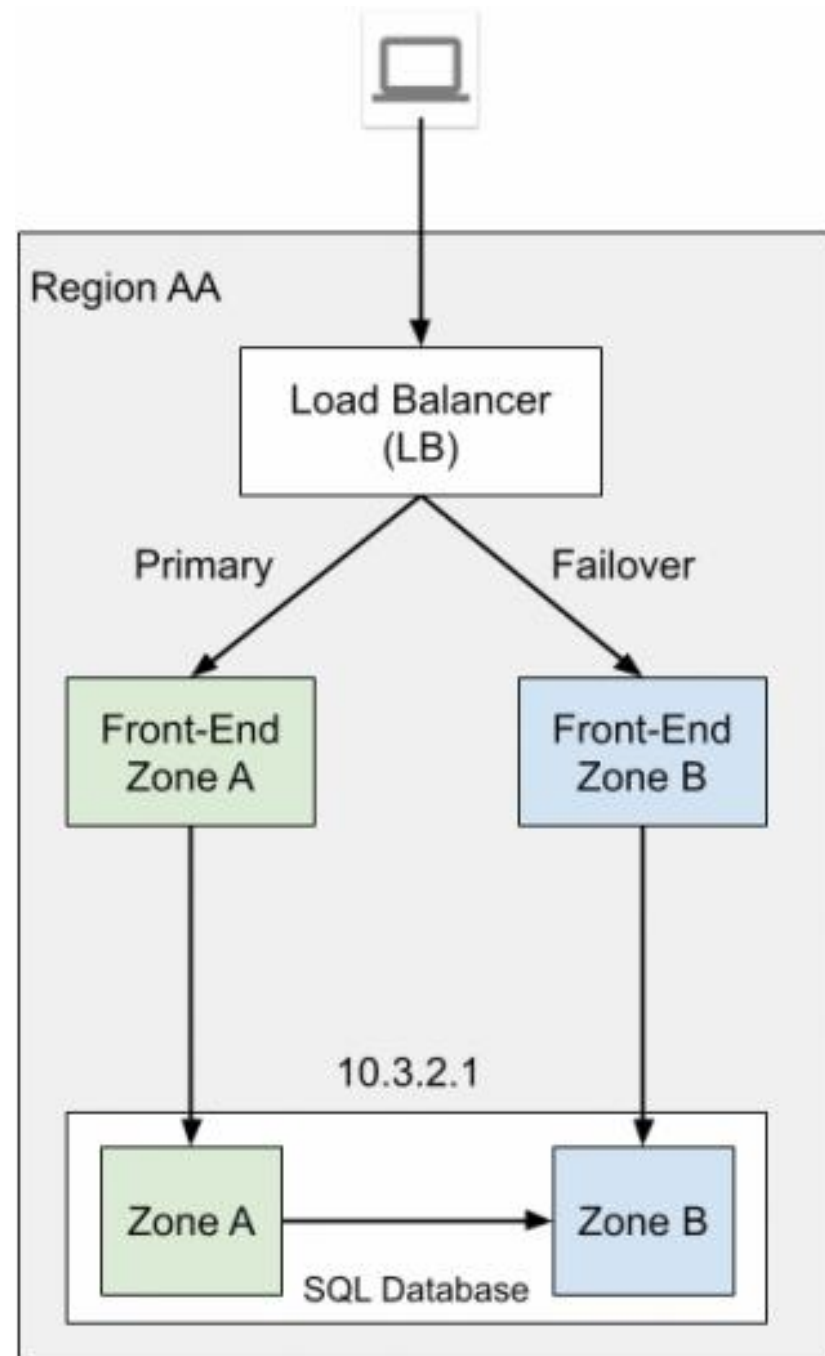
$0.9999 \times 0.995 \times 0.9995$

99.44% SLO

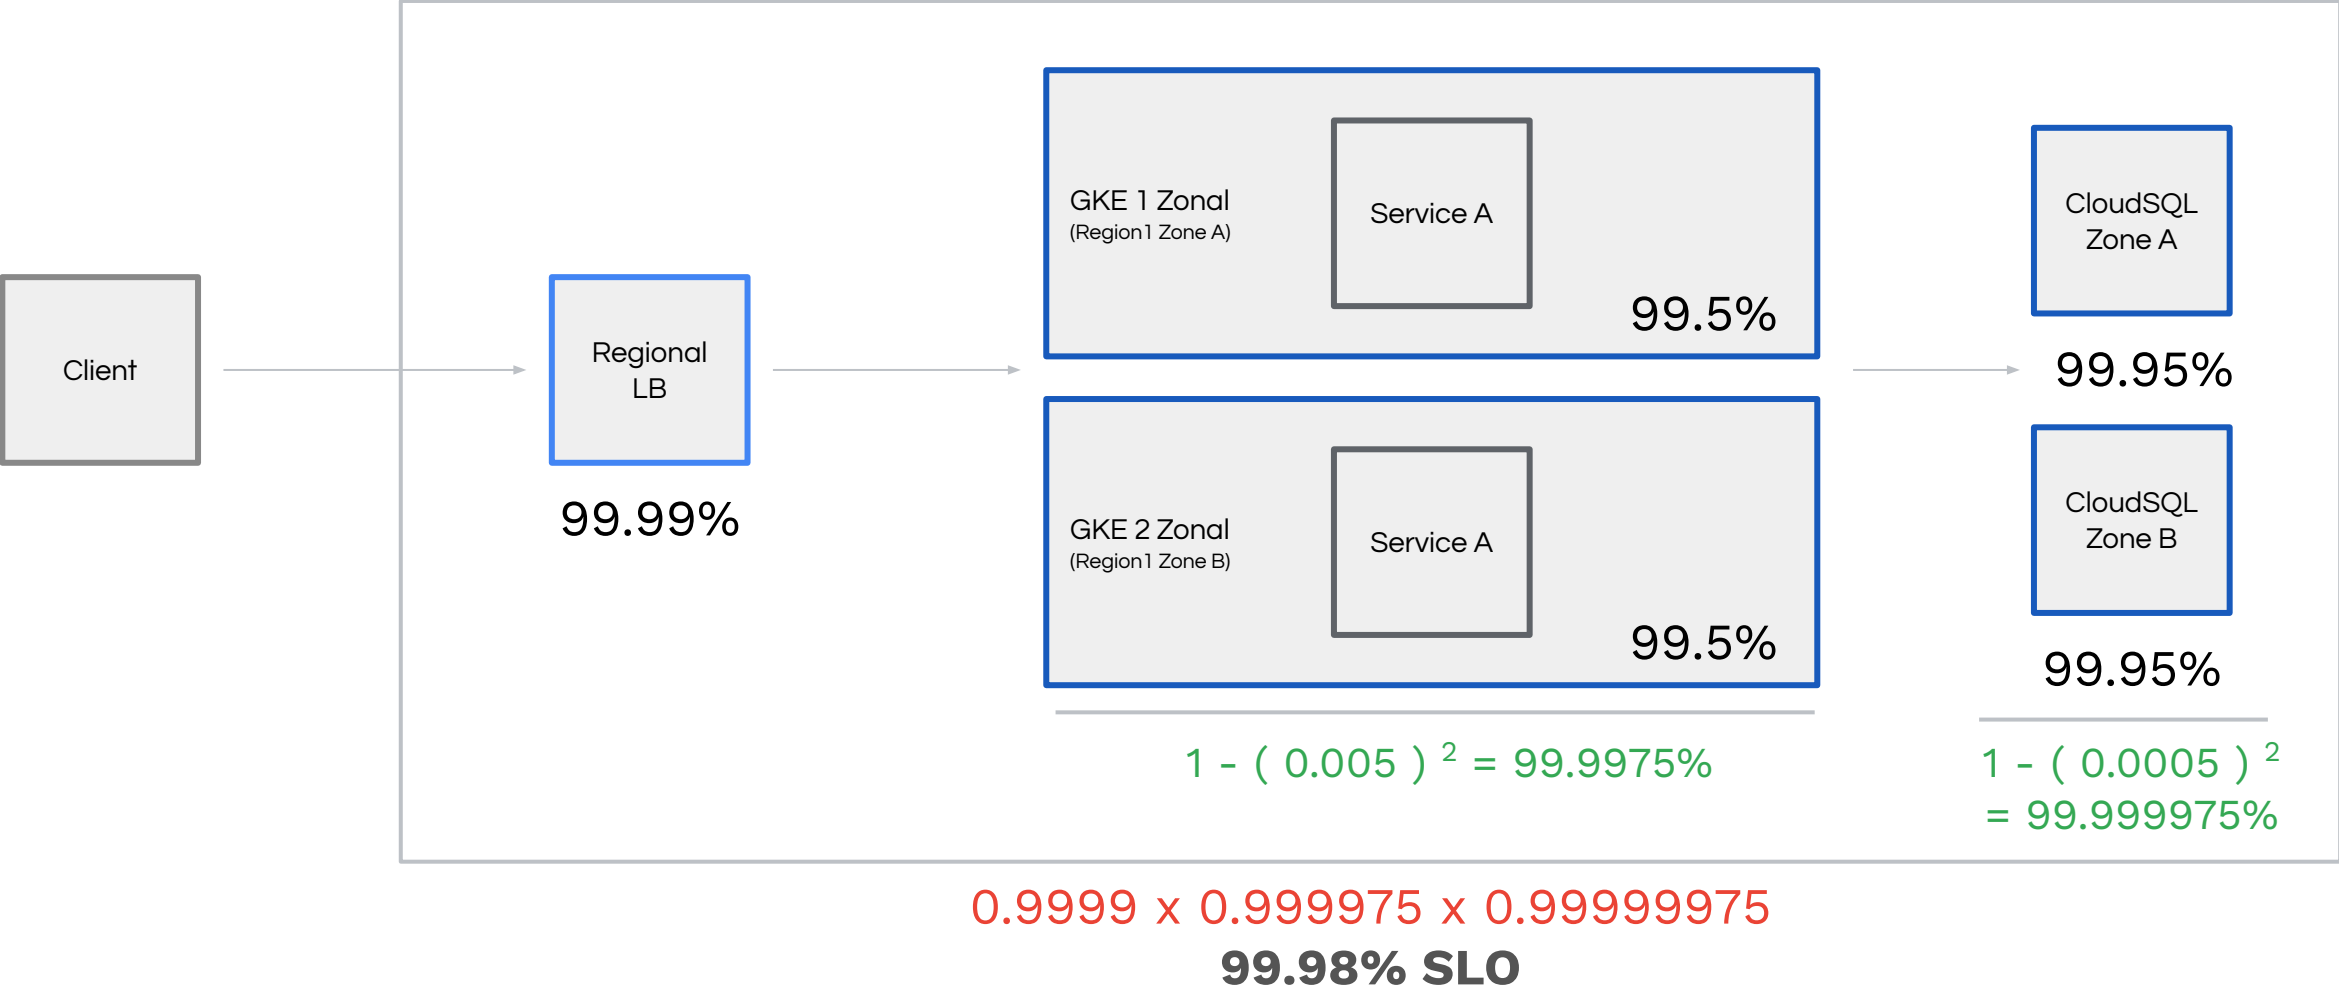
Archetype 2.2

Active Passive Zones

- Deploy all services of app to two zones in one region
- Data in Cloud SQL with a [read replica](#)
- L4 LB with one backend group
- Survives zone failure. Does not survive region failure.
- Fail-Ops: Change LB backend, [promote](#) read replica
- Cost: 2x serving + 2x data (1 replica)
- Complexity: Low
- App Refactoring: None (lift and shift)

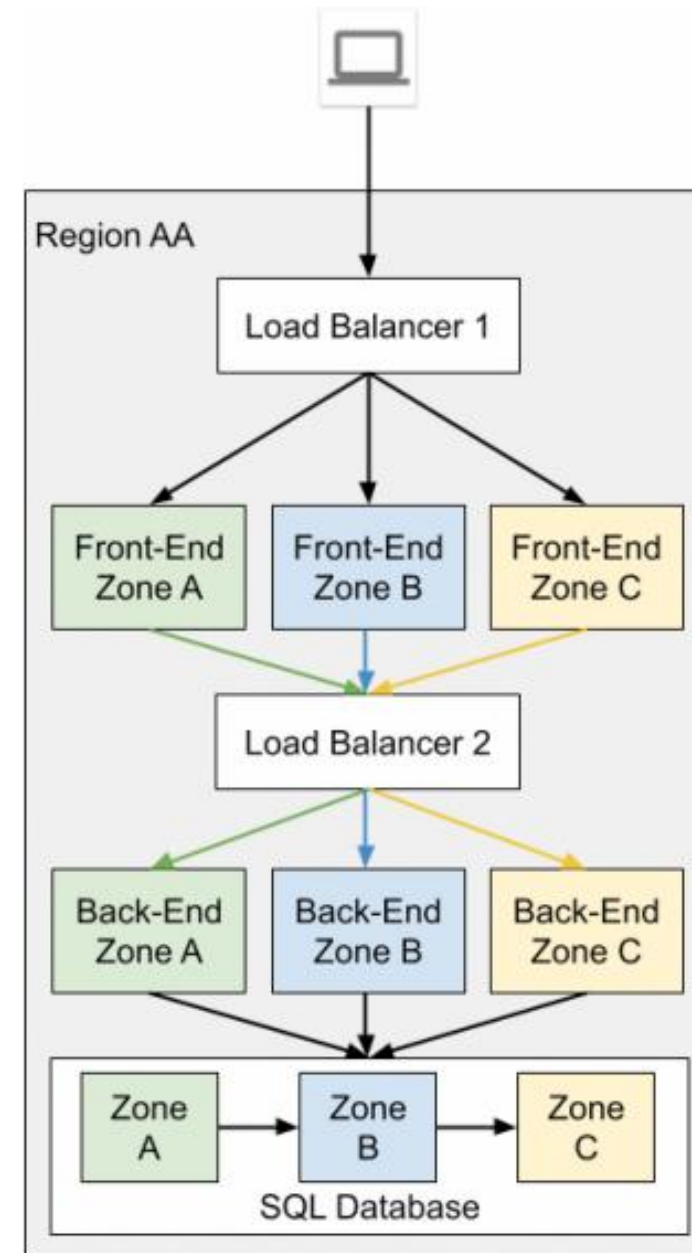


Archetype 2.2 Active Passive Zone with Cloud SQL HA

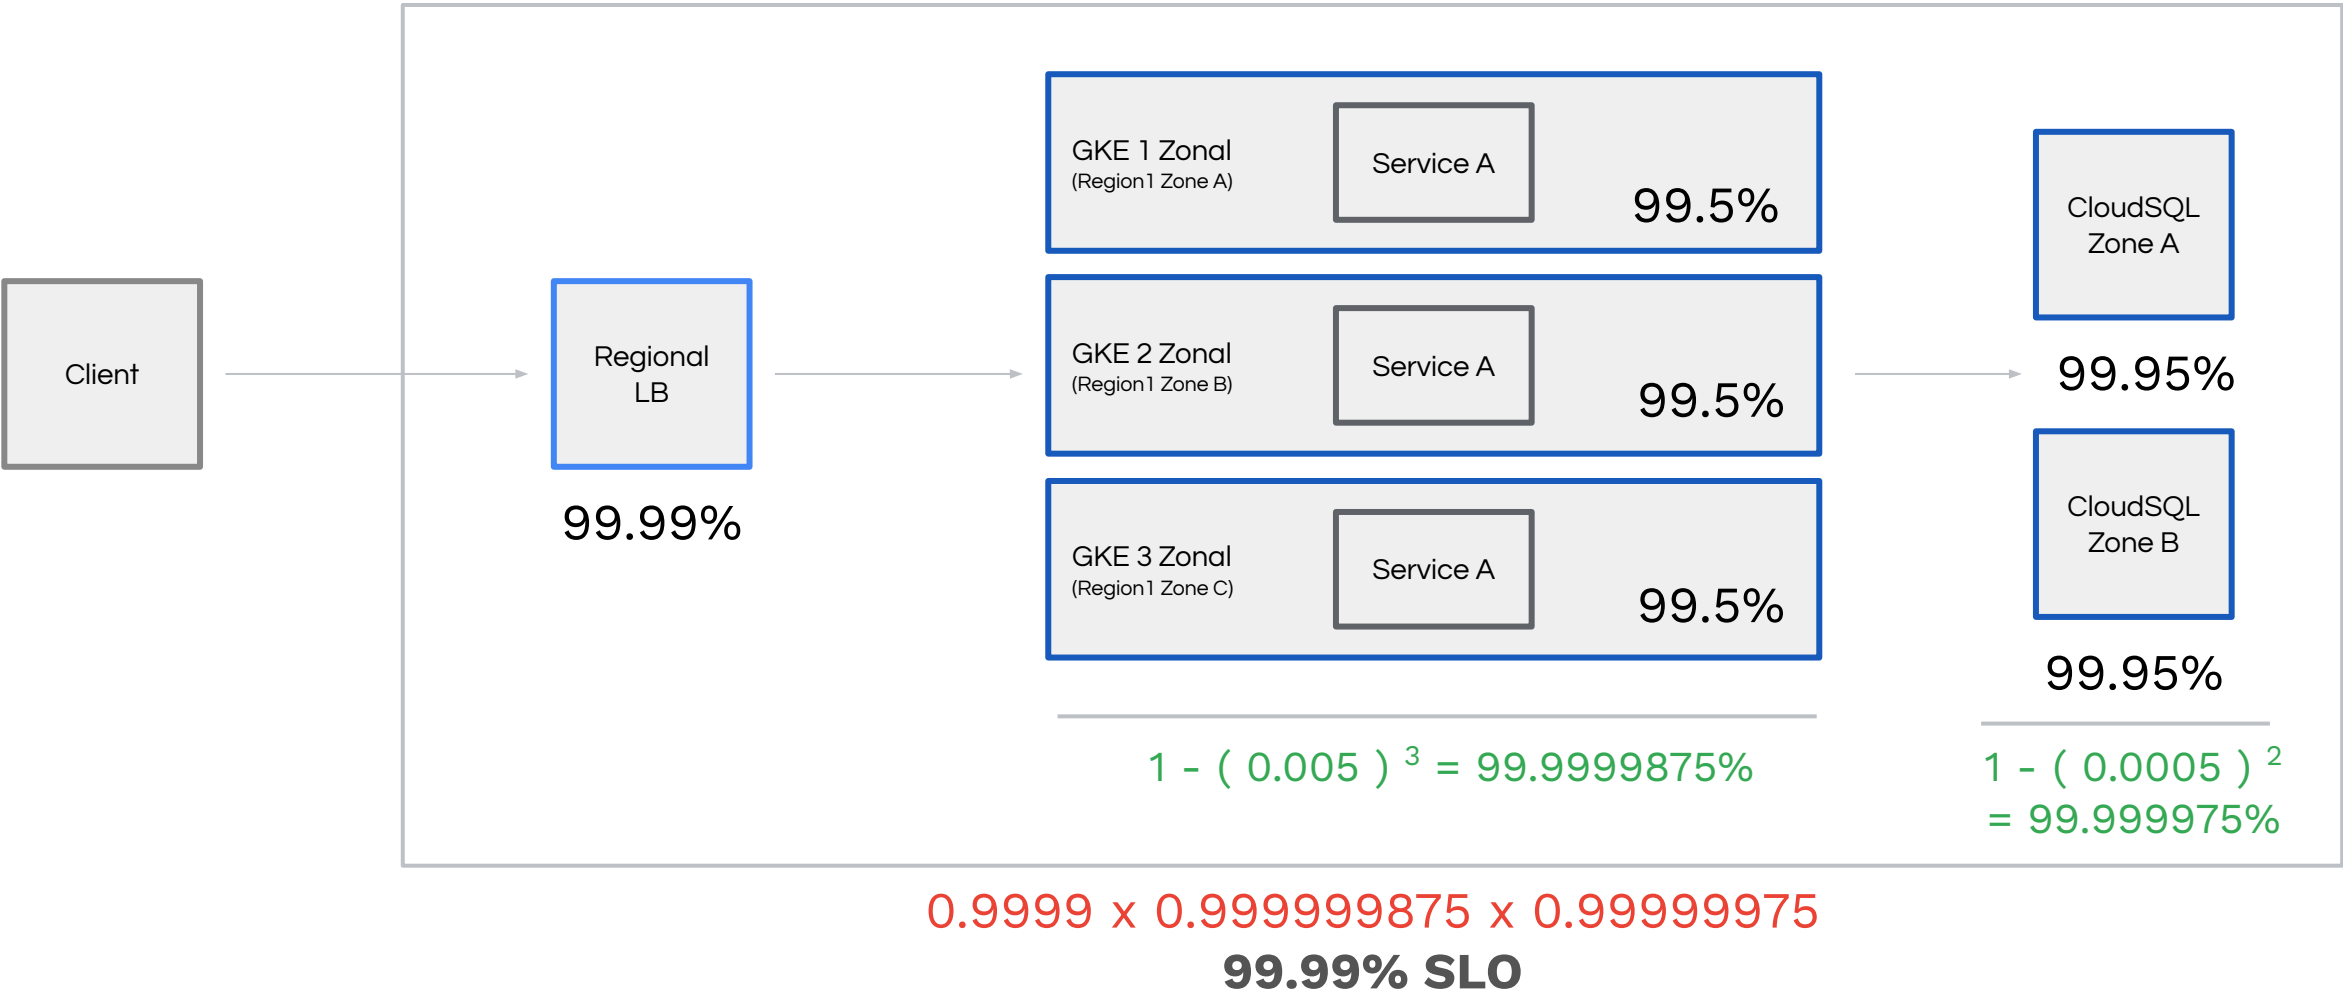


Archetype 3.1 Multi Zonal

- Deploy all services of app to all three zones in one region
- Data: Use [HA Cloud SQL](#)
- Use GLB or RLB with 3 backend groups
- Survives zone failure. Does not survive region failure.
- Fail-Ops: [Initiate DB failover](#) (testable)
- Cost: 1.5x serving + 2x data (SQL HA)
- Complexity: Medium
- App Refactoring: Low (multi instance)



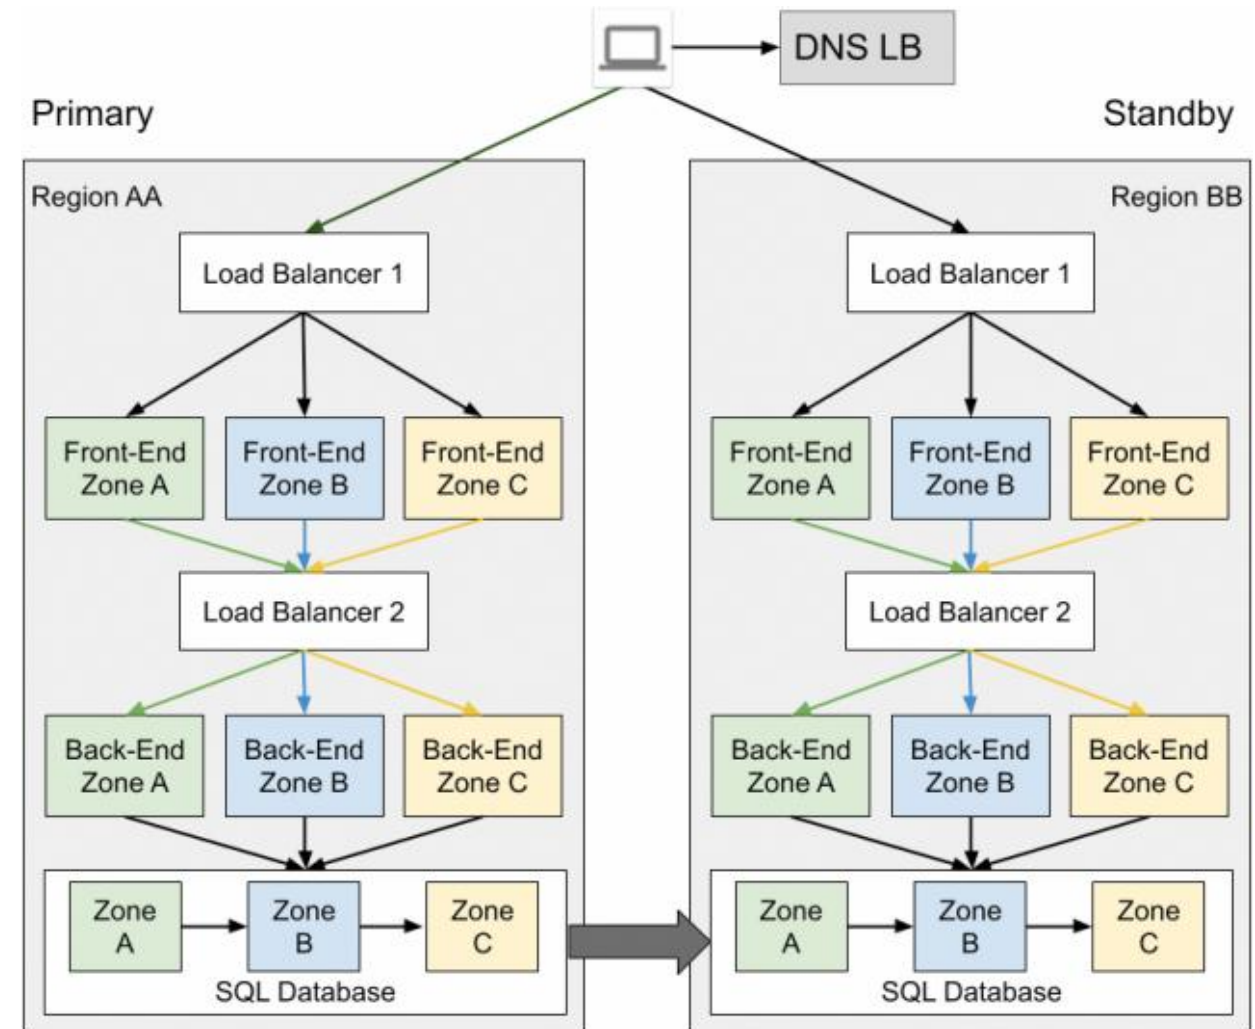
Archetype 3.1 Multi Zonal with Cloud SQL HA



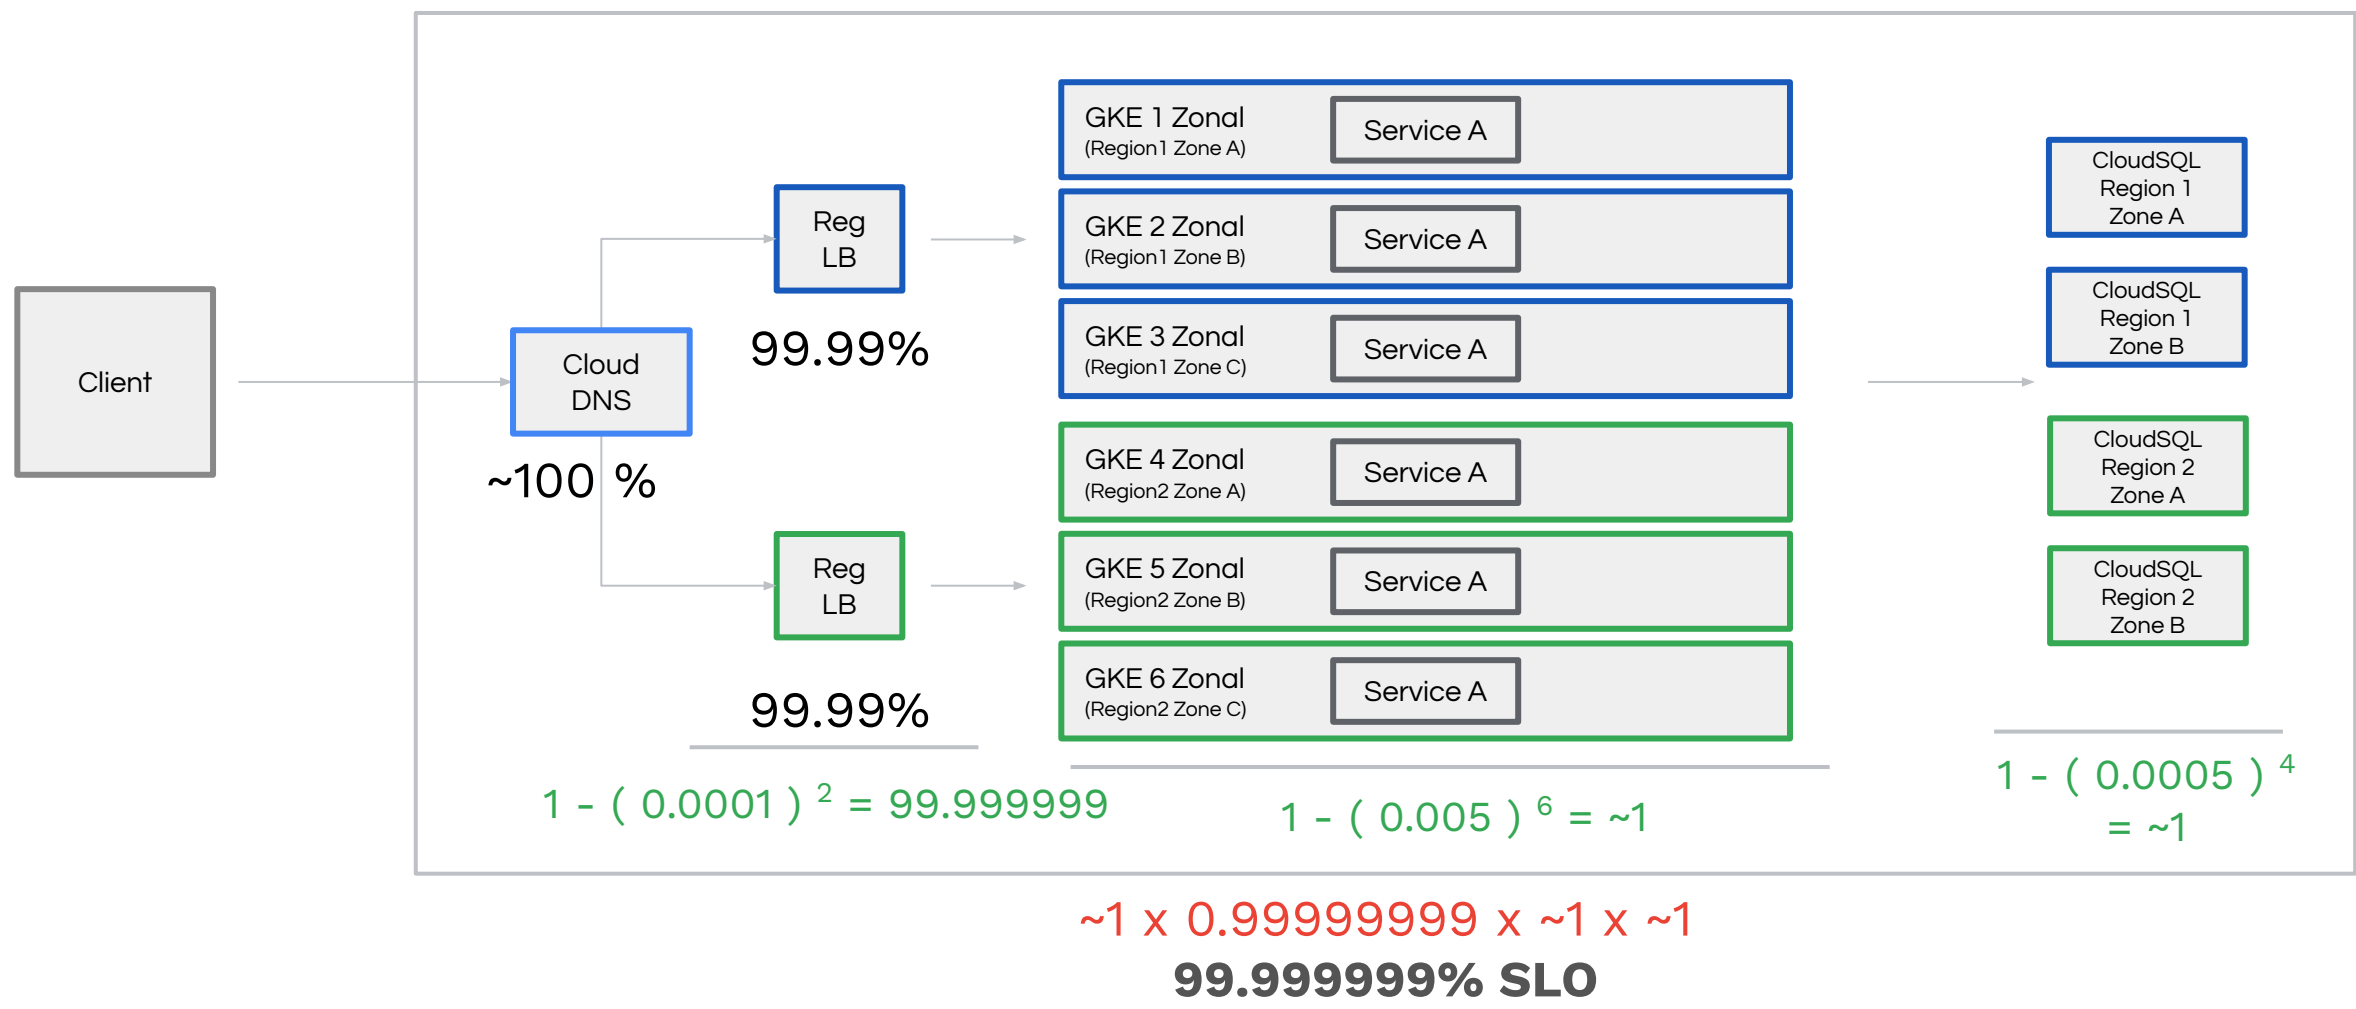
Archetype 3.2

Active Passive Region

- Deploy all services of app to all three zones in each of two regions
- Data: Cloud SQL with [cross-region replication](#)
- Cloud DNS points at one LB (until disaster)
- Survives zone and region failures
- Fail-Ops: No action for zone failure.
 - Update DNS to point at standby LB
 - [Cross region DR failover](#) process for DB
- Cost: 2x serving + 2x data (SQL HA)
- Complexity: Medium
- App Refactoring: Medium (multi instance, multi regional data)



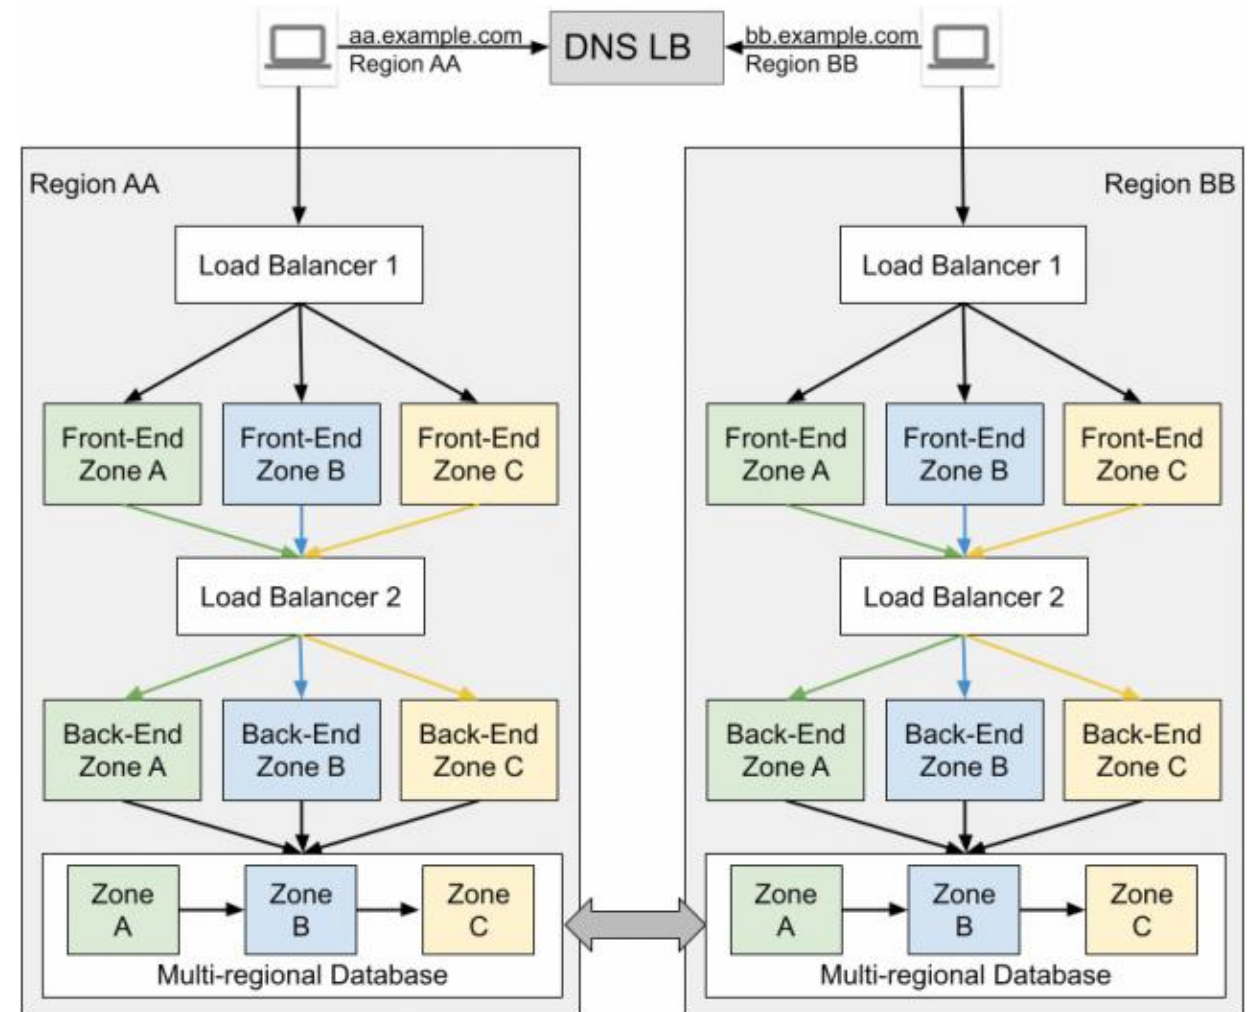
Archetype 3.2 Active Passive Regions with Cloud SQL HA



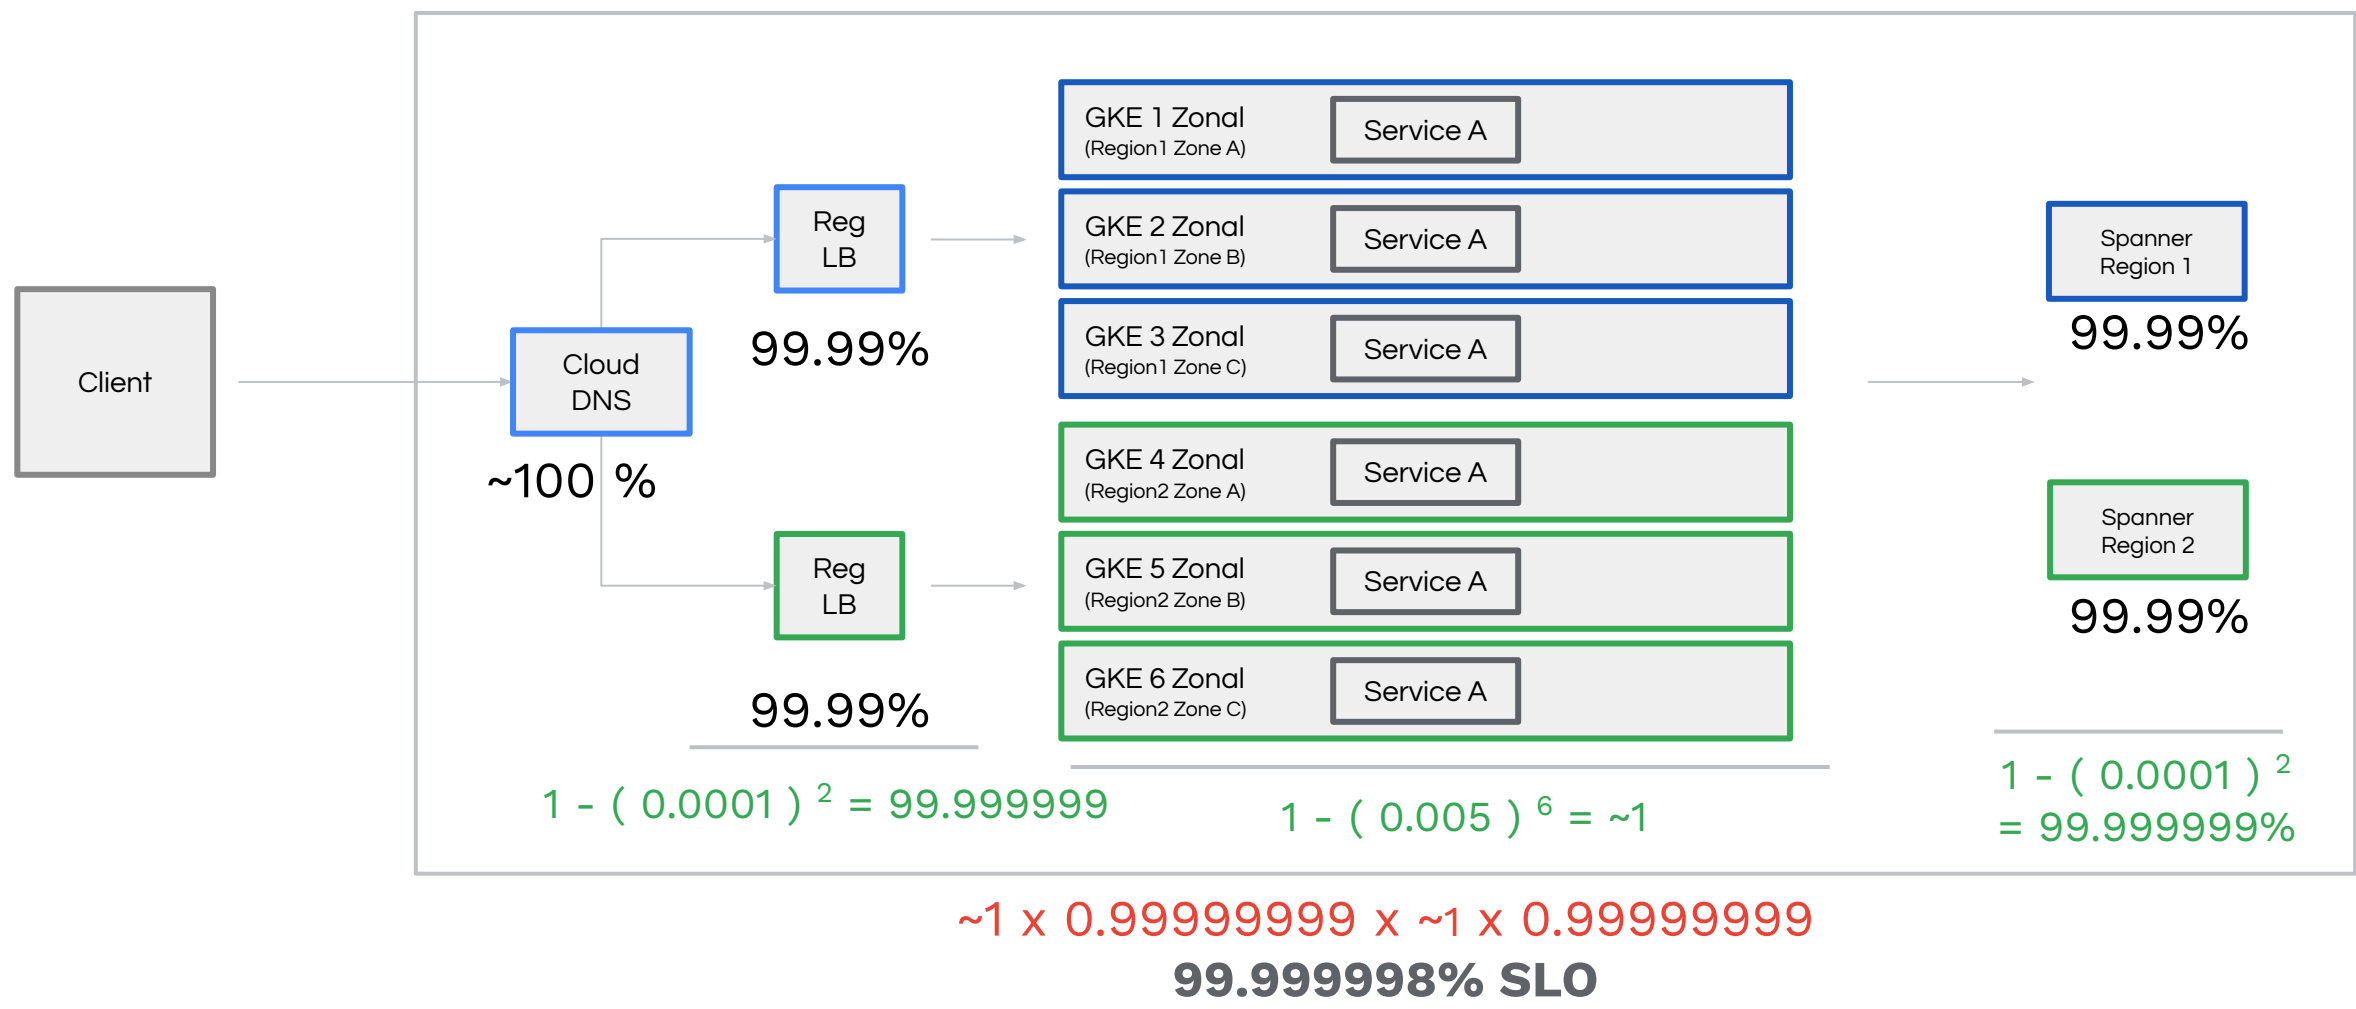
Archetype 4.3

Isolated Regions

- Deploy all services of app to all three zones in each of two regions
- Data: Spanner or CockroachDB
- Cloud DNS points at two Regional LBs
- Survives zone and region failures. No impact for ½ consumers. Possible manual failover
- Fail-Ops: No action for zone failure. Optional regional failover like Arch 3.2
- Cost: 1+½ cost for zone failure. Scale-up required if optional regional failover expected
- Complexity: Medium/High
- App Refactoring: Medium (multi instance, multi regional data)



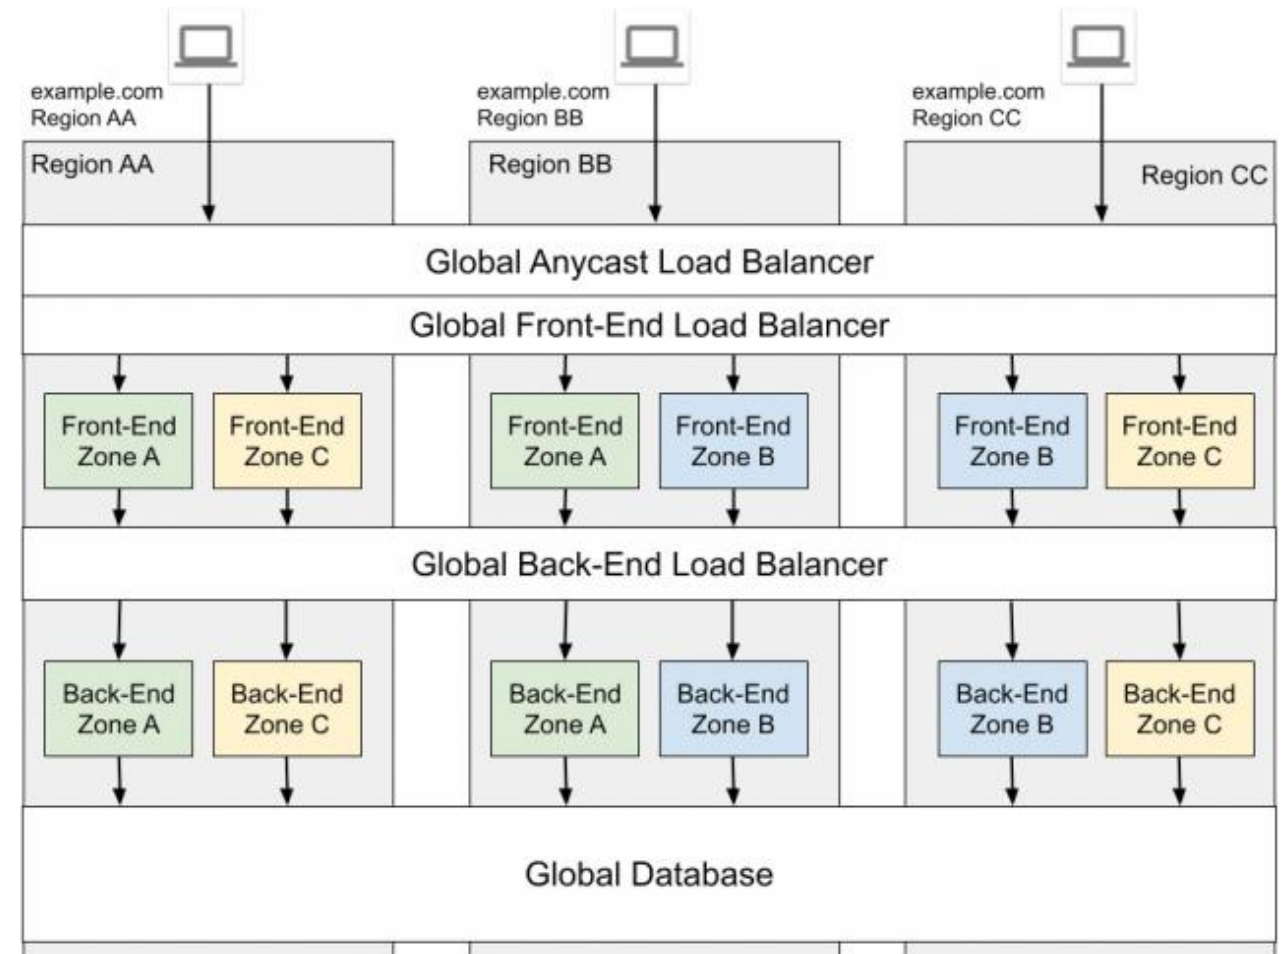
Archetype 4.3 Isolated Regions with Cloud Spanner



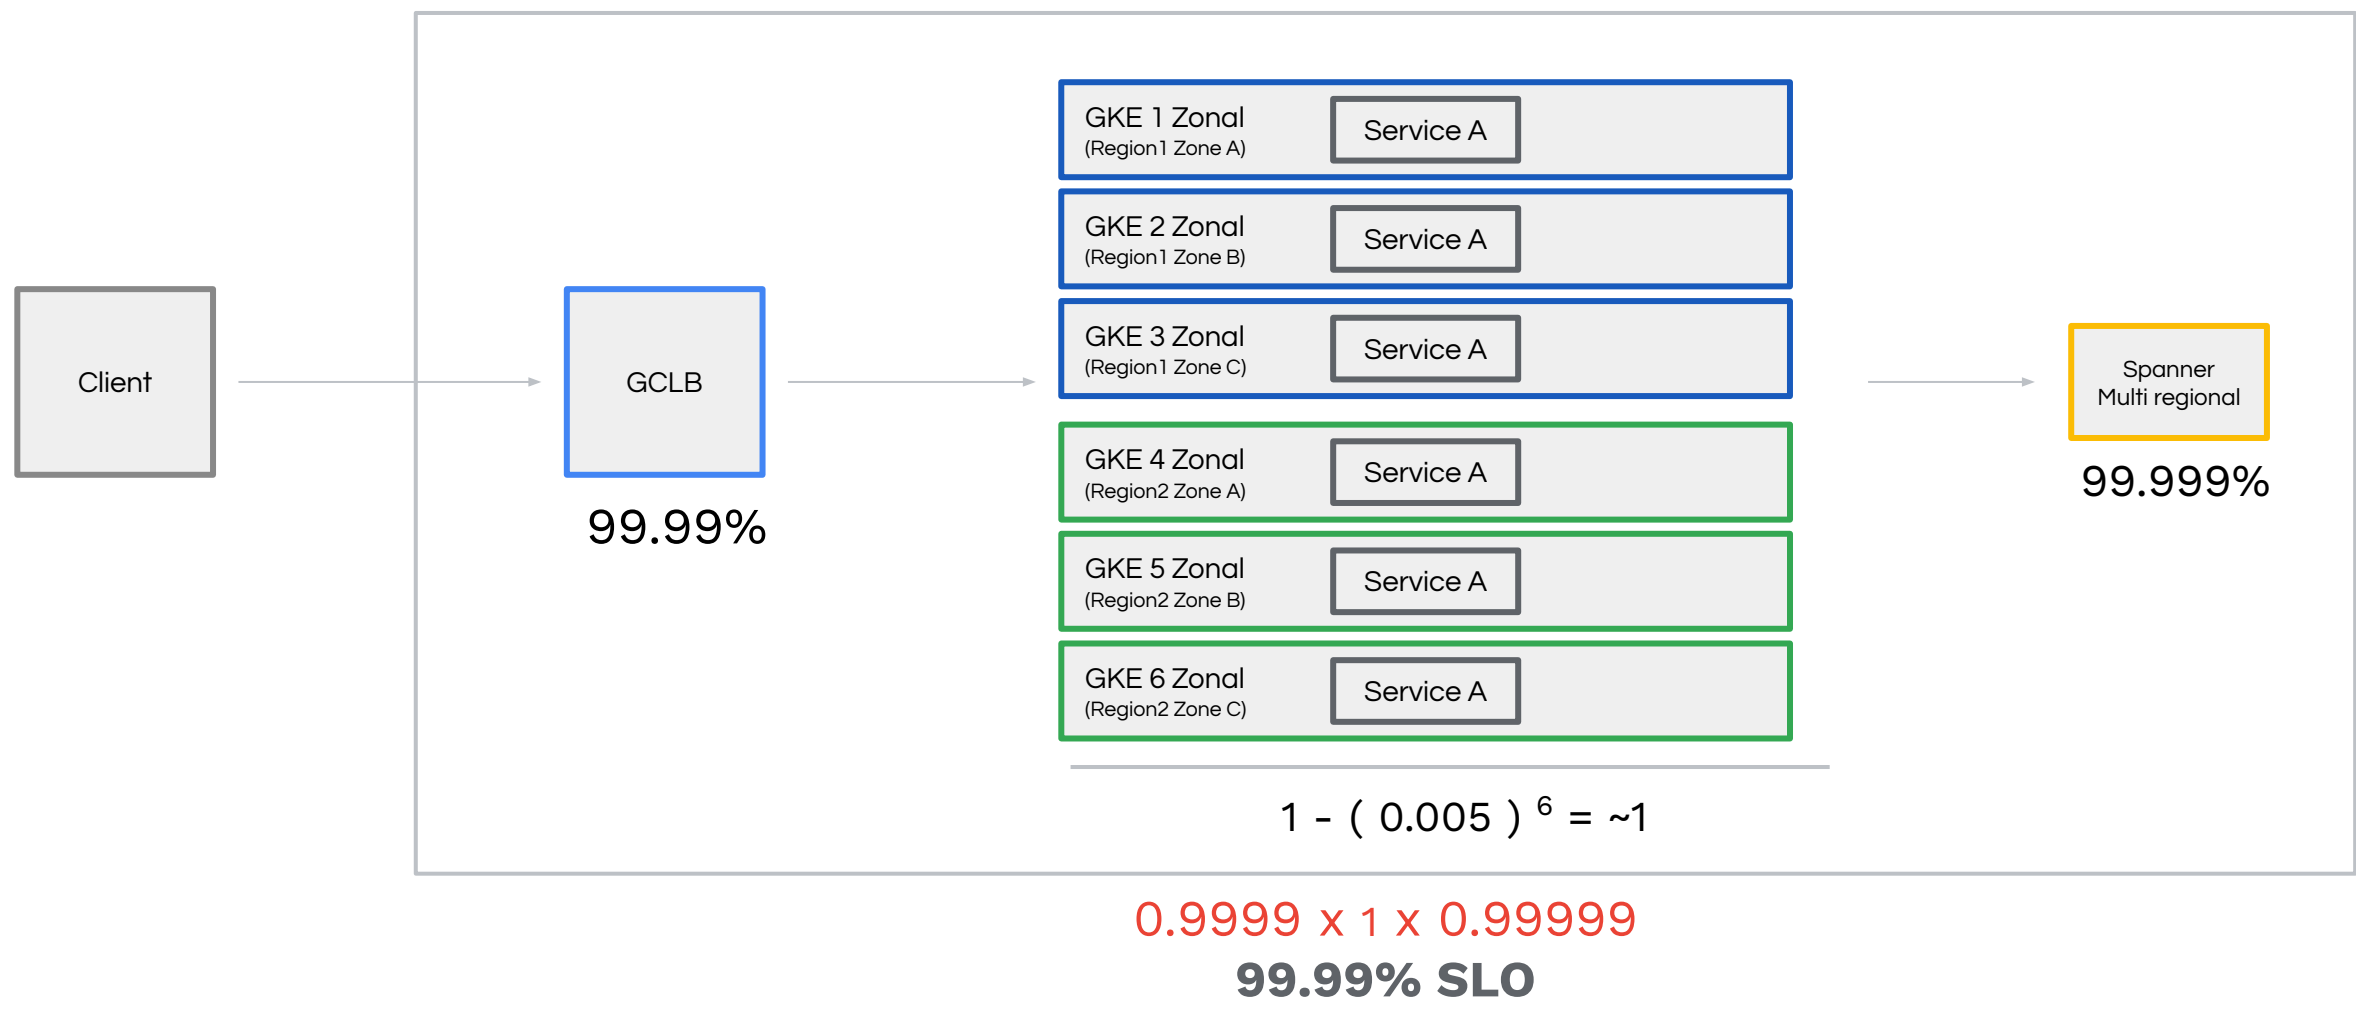
Archetype 5.2

Global

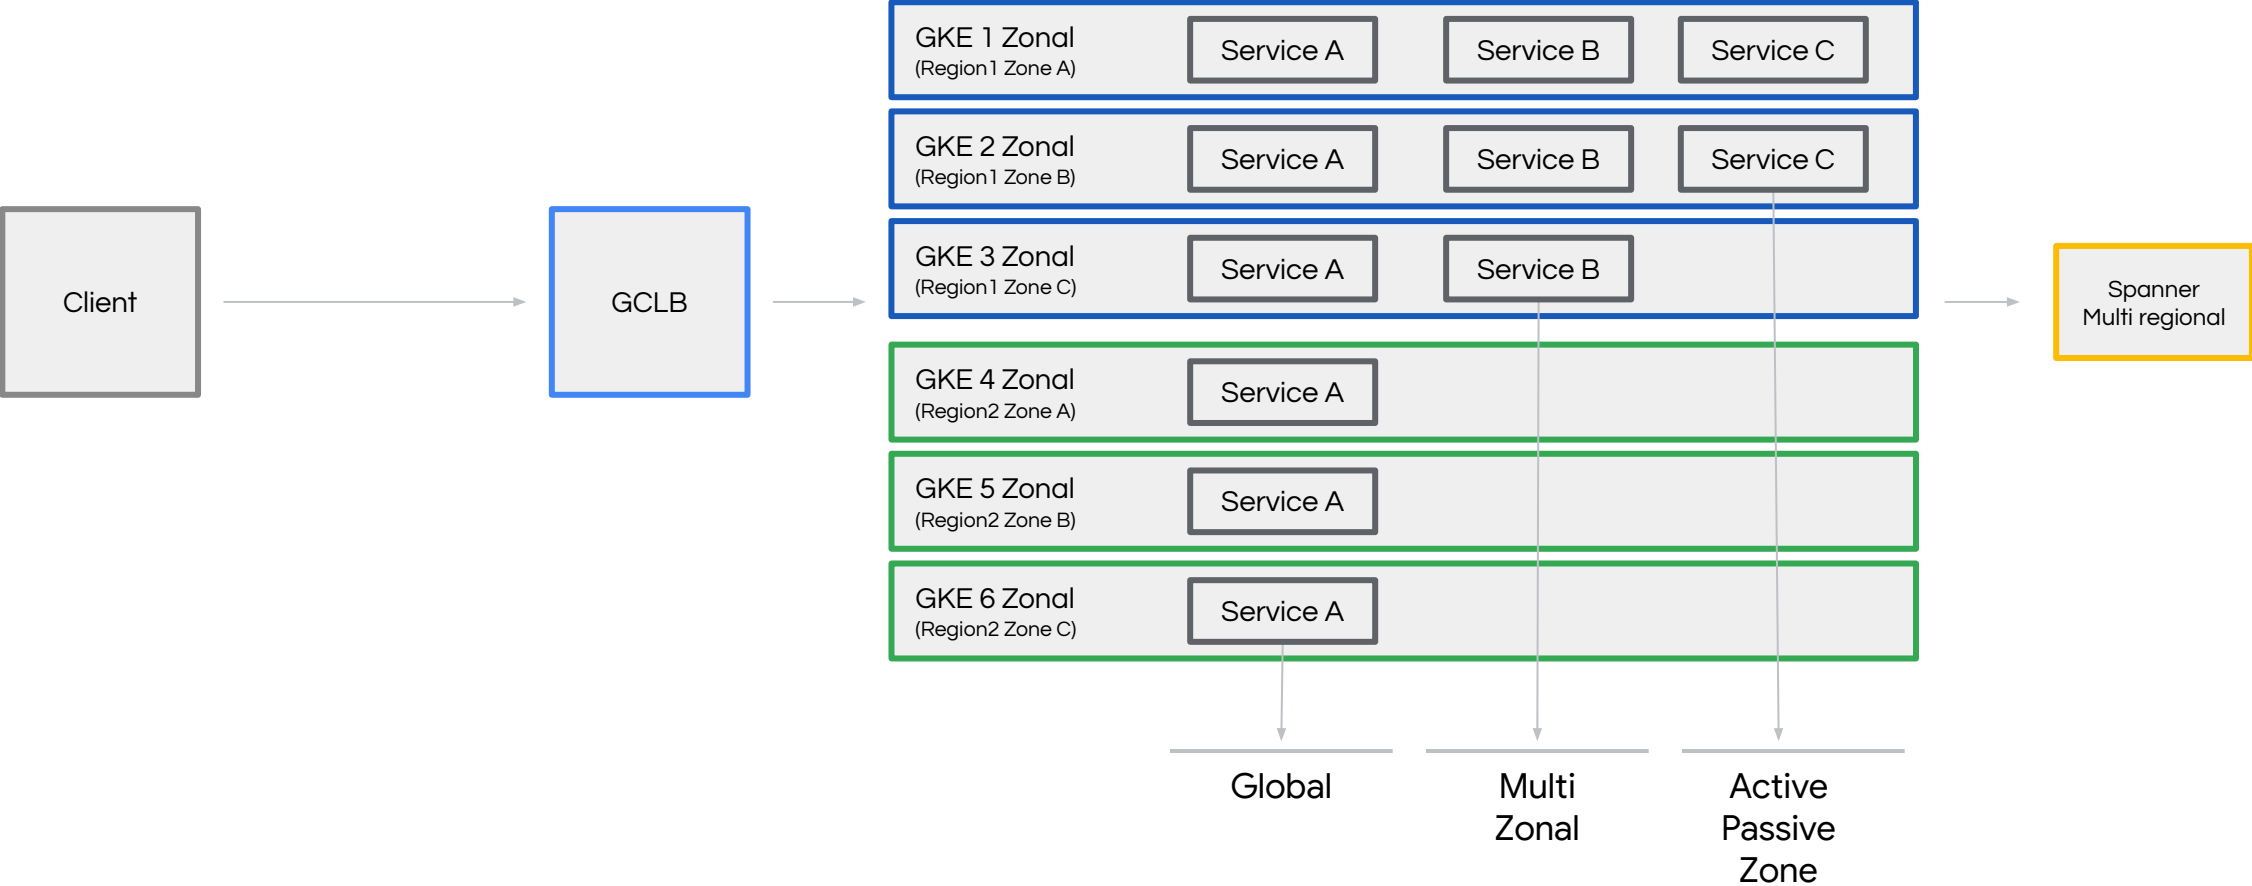
- **Deploy** all services of app to all three zones in each of two or more regions
- **Data:** Spanner or CockroachDB
- **Global LB** points at regional backend groups
- **Survives zone and region failures**
- **Fail-Ops:** None
- **Cost:** N+m cost modelling. Global DBs are more expensive
- **Complexity:** High
- **App Refactoring:** High (multi instance, global DBs)





Archetype 5.2 Global with Cloud Spanner (Multi regional)





Application with Services using multiple Archetypes











Conclusions

- Start with **Archetypes**
- Compose **Services into Applications**
that can **degrade gracefully**
- Develop **resilient teams**
robust platforms
reliable products



Please scan the QR Code above
to leave feedback on this session

Thank you

