



# Continuous Profiling Your Production Kubernetes Deployments At Scale

Nick Pordash – Lead Engineer, Observability

Vijay Samuel – Architect, Observability

# Agenda

Pillars of Observability

What is Continuous Profiling

Architecture

Examples

Potential

Questions

# Pillars of Observability

Logs

Unstructured/ structured strings (logs). Ex: sl4j, zap logger generated logs

Metrics

Continuous time series data. Ex: cpu usage, request count, latency etc.

Traces

Record of request execution through highly distributed microservices

Profiles

Code-level consumption numbers (e.g., memory used, CPU time spent) for various resources across different program functions over time.

# Profiles

Performance metrics at the most granular level possible.  
Historically viewed as non-production friendly (expensive).

## Types

- CPU
- Heap
- Mutex
- IO
- GPU
- Language specific (Goroutines/JVM)

## Information Available

- Why part of server resources are being consumed
- Statistics about execution down to the line number

# Continuous Profiling

Advent of sampling profilers made profiling a lot cheaper.  
A time axis to information collected as part of profiling can be viewed on a chart.



[Image Credit](#)

## Available Open Source Products

- [Pyroscope](#)
- [CNCF Pixie](#)
- [Parca](#)
- [OTEL Profiles](#) - In progress

## Methodology

- Scrape - Collect profiles at a periodic interval by scraping something like a PPROF endpoint
- Push - Instrument with a client that can collect profiles and periodically push them
- eBPF - Use features on the Linux Kernel to extract profiling information.

# Our Setup

## **Pyroscope as a candidate**

Of all the choices, Pyroscope had a nice UX for us to be able to visualize profiling data.

Easy to set up

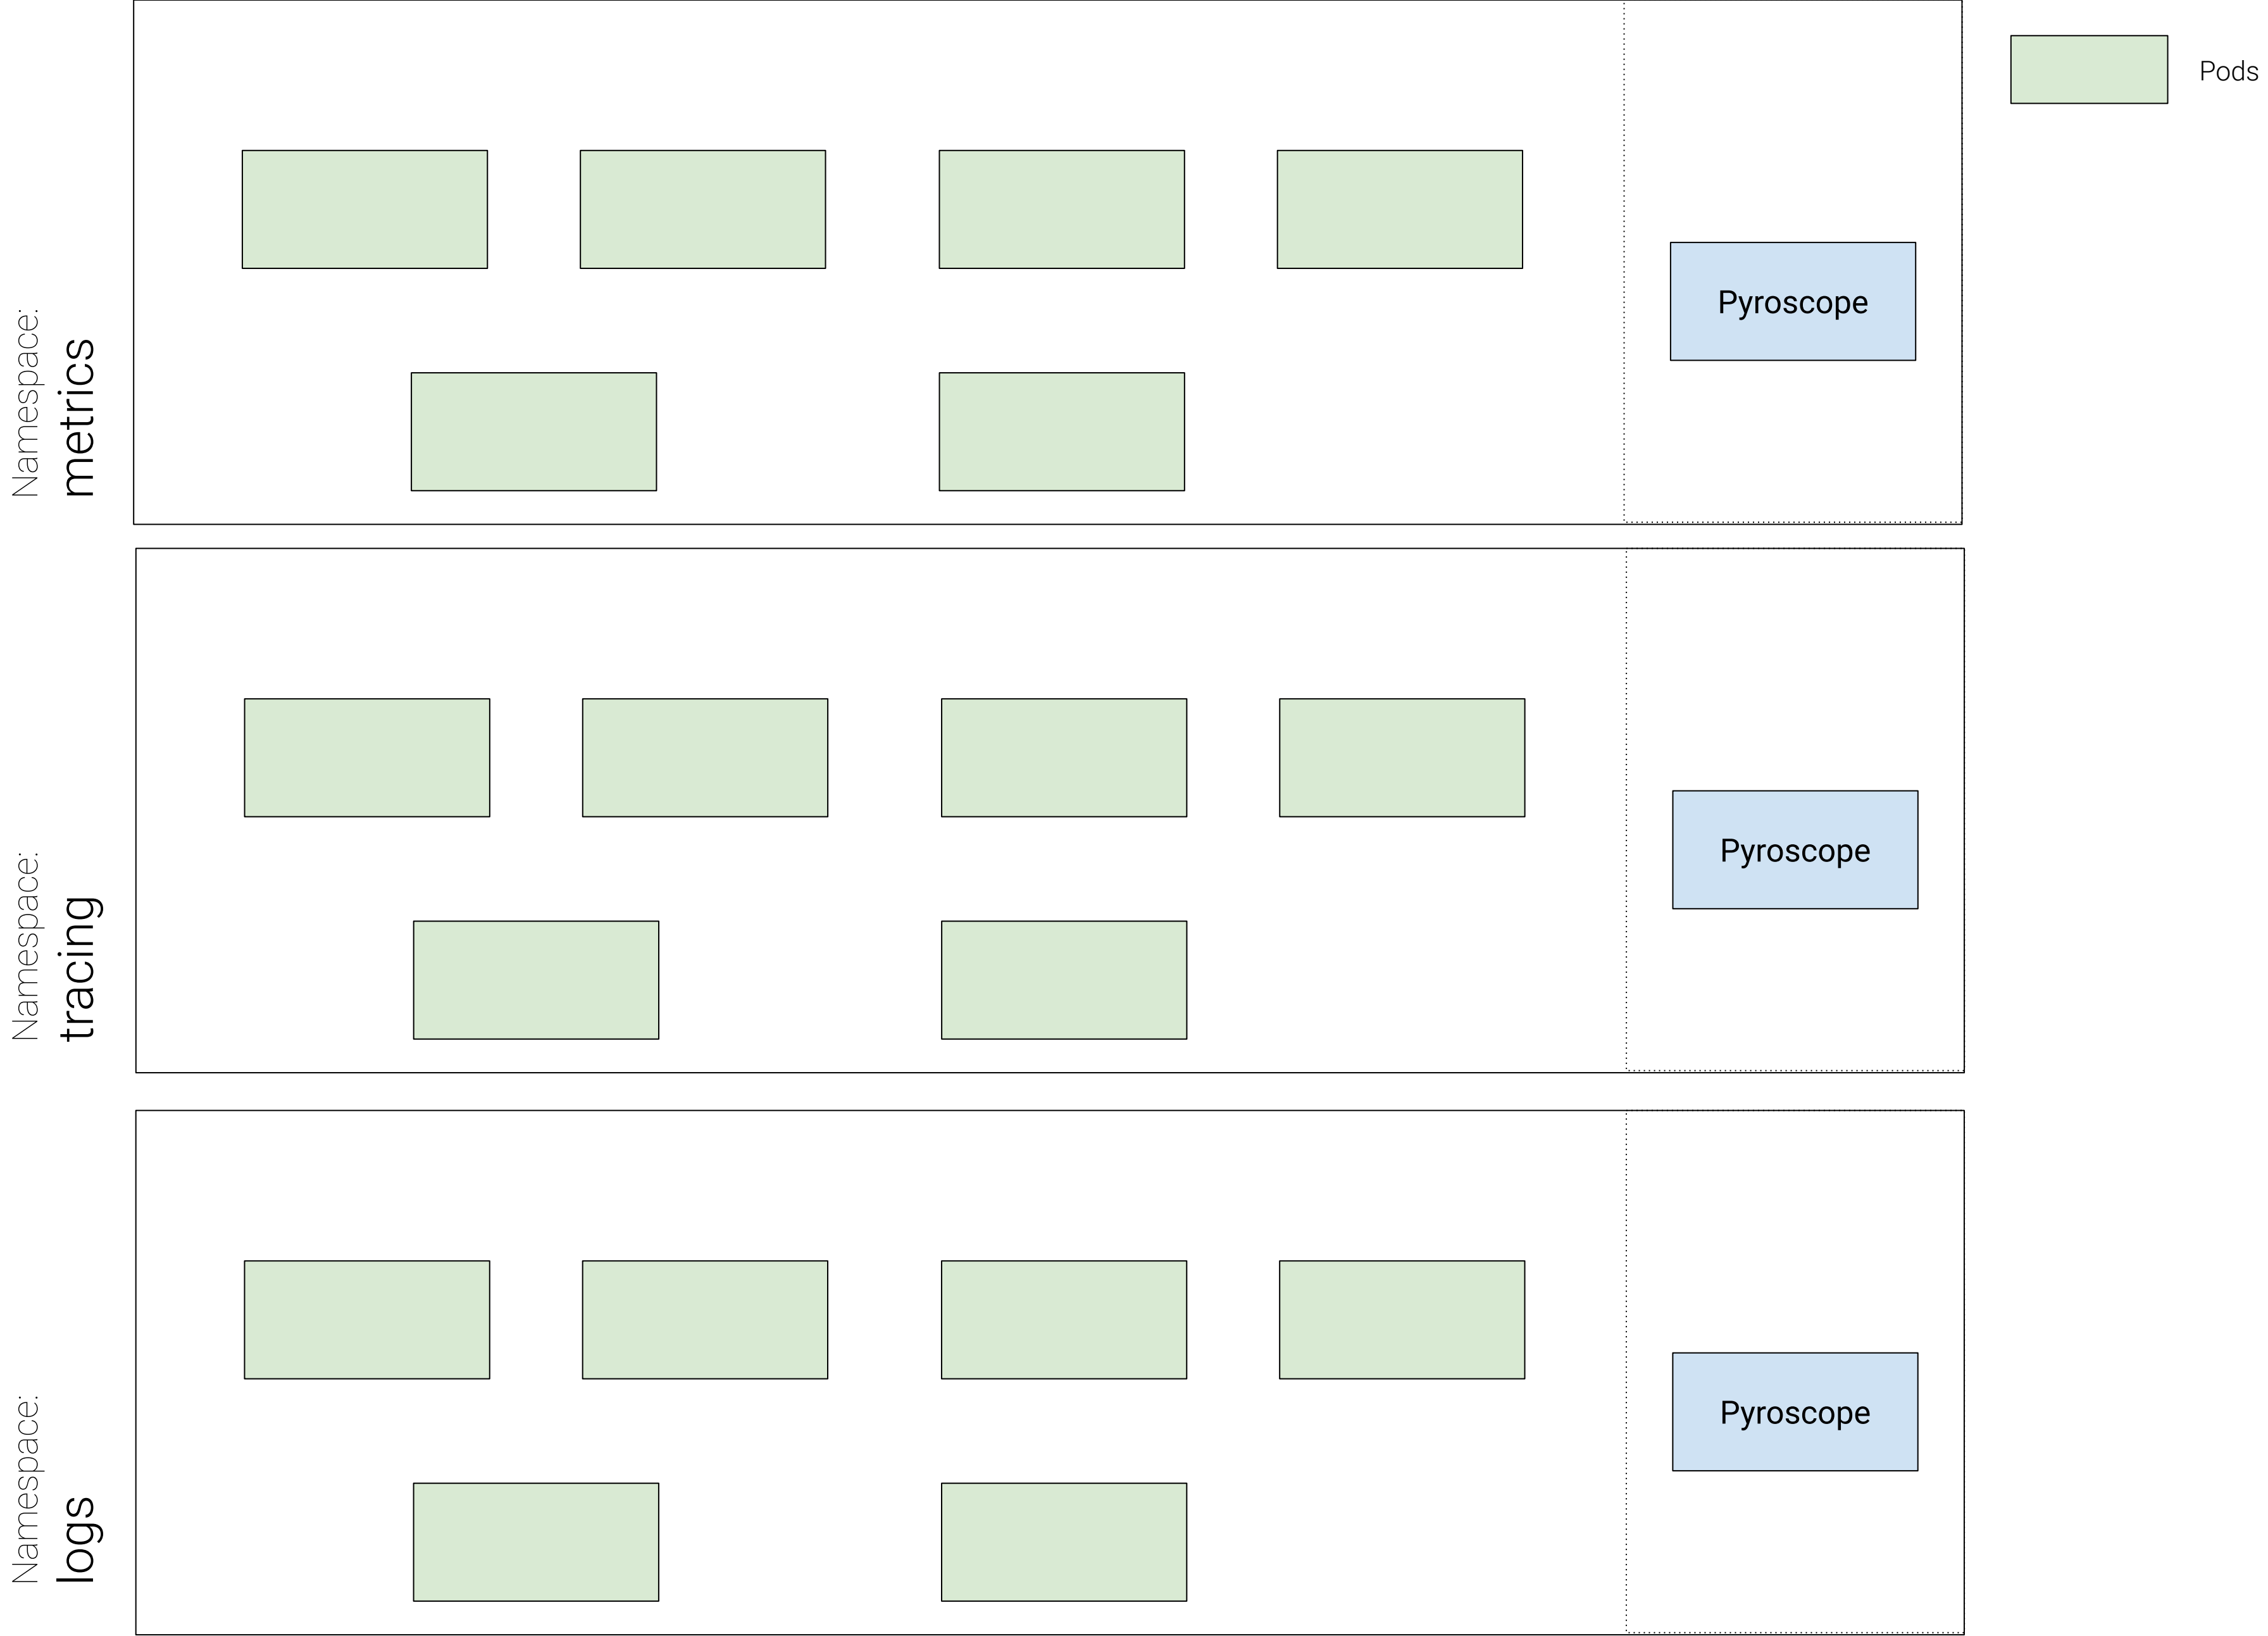
## **Deployed against most of our platform**

Environments that we actively use for L&P analysis were configured with Pyroscope installations.

Subsequently production namespaces were configured too.

Pyroscope kubernetes discovery uses Pod annotations to scrape profiles every 10 seconds.

# Deployment Architecture



# Obvious Questions

## **Are there performance degradations?**

No! Golang pprof is very cheap. 10s polling cause less than 2% increase in CPU usage.

## **Are Profiles expensive to store?**

Sort of. Raw samples do indeed take up a lot of space.

Pyroscope recommends storing few hours of raw samples. Aggregated data retained longer.

## **Did we find good use?**

YES! Few examples next.

A solid red triangle pointing towards the bottom right corner of the slide.



# Example - Missed Compiler Optimization

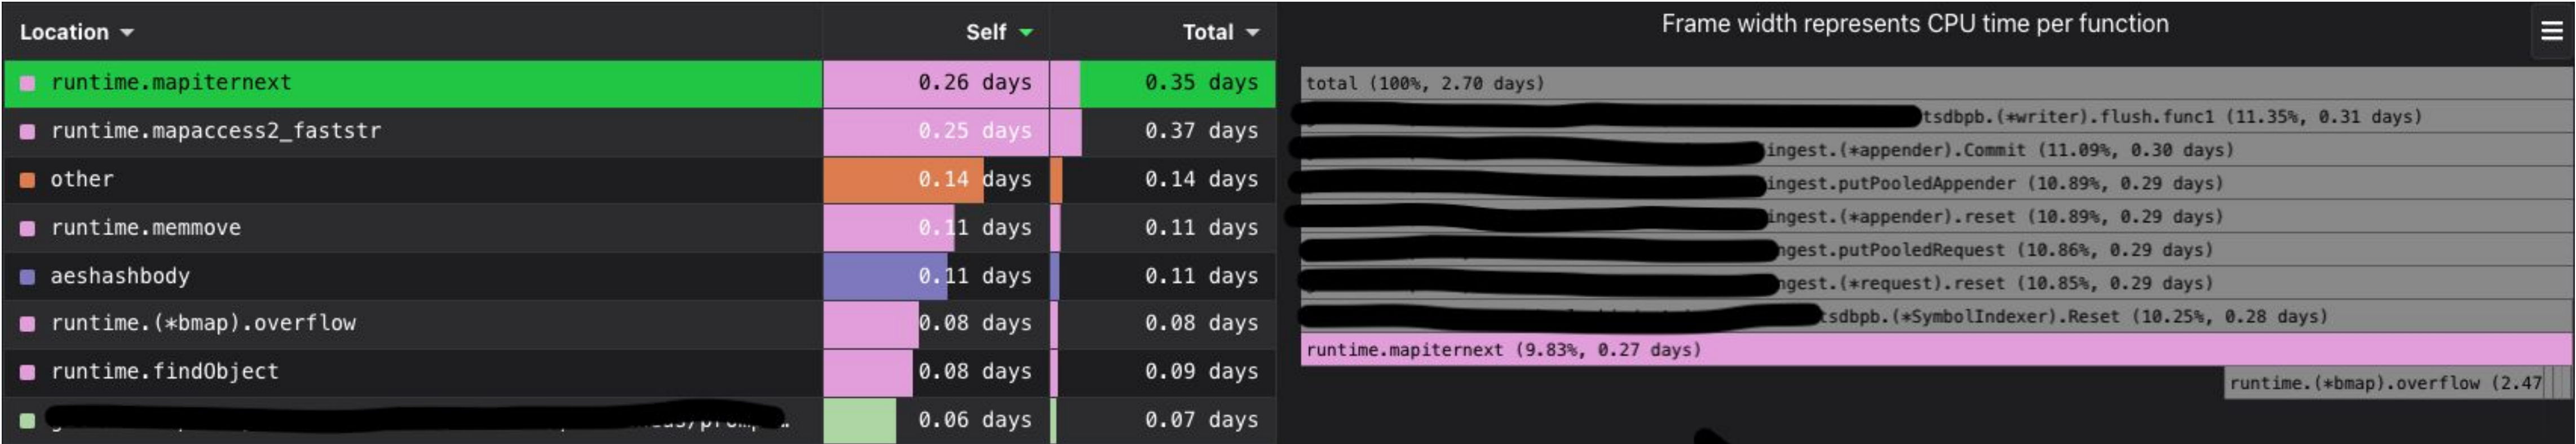
## Context

Golang compiler optimizes the following idiom to perform a map clear instead of iteratively deleting each key:

```
for k := range m {
    delete(m, k)
}
```

## Problem

Highest cpu consumer is caused by an unoptimized map clear. Expected to see this code path call `runtime.mapclear` instead of `runtime.mapiternext`.



# Example - Missed Compiler Optimization

## Let's check the assembly instructions

```
$ go build -gcflags='-S' tsdb/tsdbpb/symbols.go 2>&1 | grep runtime.mapclear
0x0023 00035 (/tsdbpb/symbols.go:38)      CALL    runtime.mapclear(SB)
```

Huh? It says that `runtime.mapclear` is used!

## Let's check the assembly instructions from a caller

```
$ go build -gcflags='-S' ingest/*.go 2>&1 | grep symbols.go | grep runtime.mapiternext
0x00cb 00203 (/tsdb/tsdbpb/symbols.go:38)      CALL    runtime.mapiternext(SB)
```

**Theory: *is function inlining preventing the optimization?***



```
tsdb/tsdbpb/symbols.go
@@ -35,6 +35,11 @@ func (idx *SymbolIndexer) Strings() []string {
35
36 // Reset resets the index to be empty, but it retains the underlying storage for future use.
37 func (idx *SymbolIndexer) Reset() {
38     +     idx.reset()
39     + }
40     +
41     + //go:noinline
42     + func (idx *SymbolIndexer) reset() {
43         for s := range idx.sm {
44             delete(idx.sm, s)
45         }
46     }
47 }
```

# Example - Missed Compiler Optimization

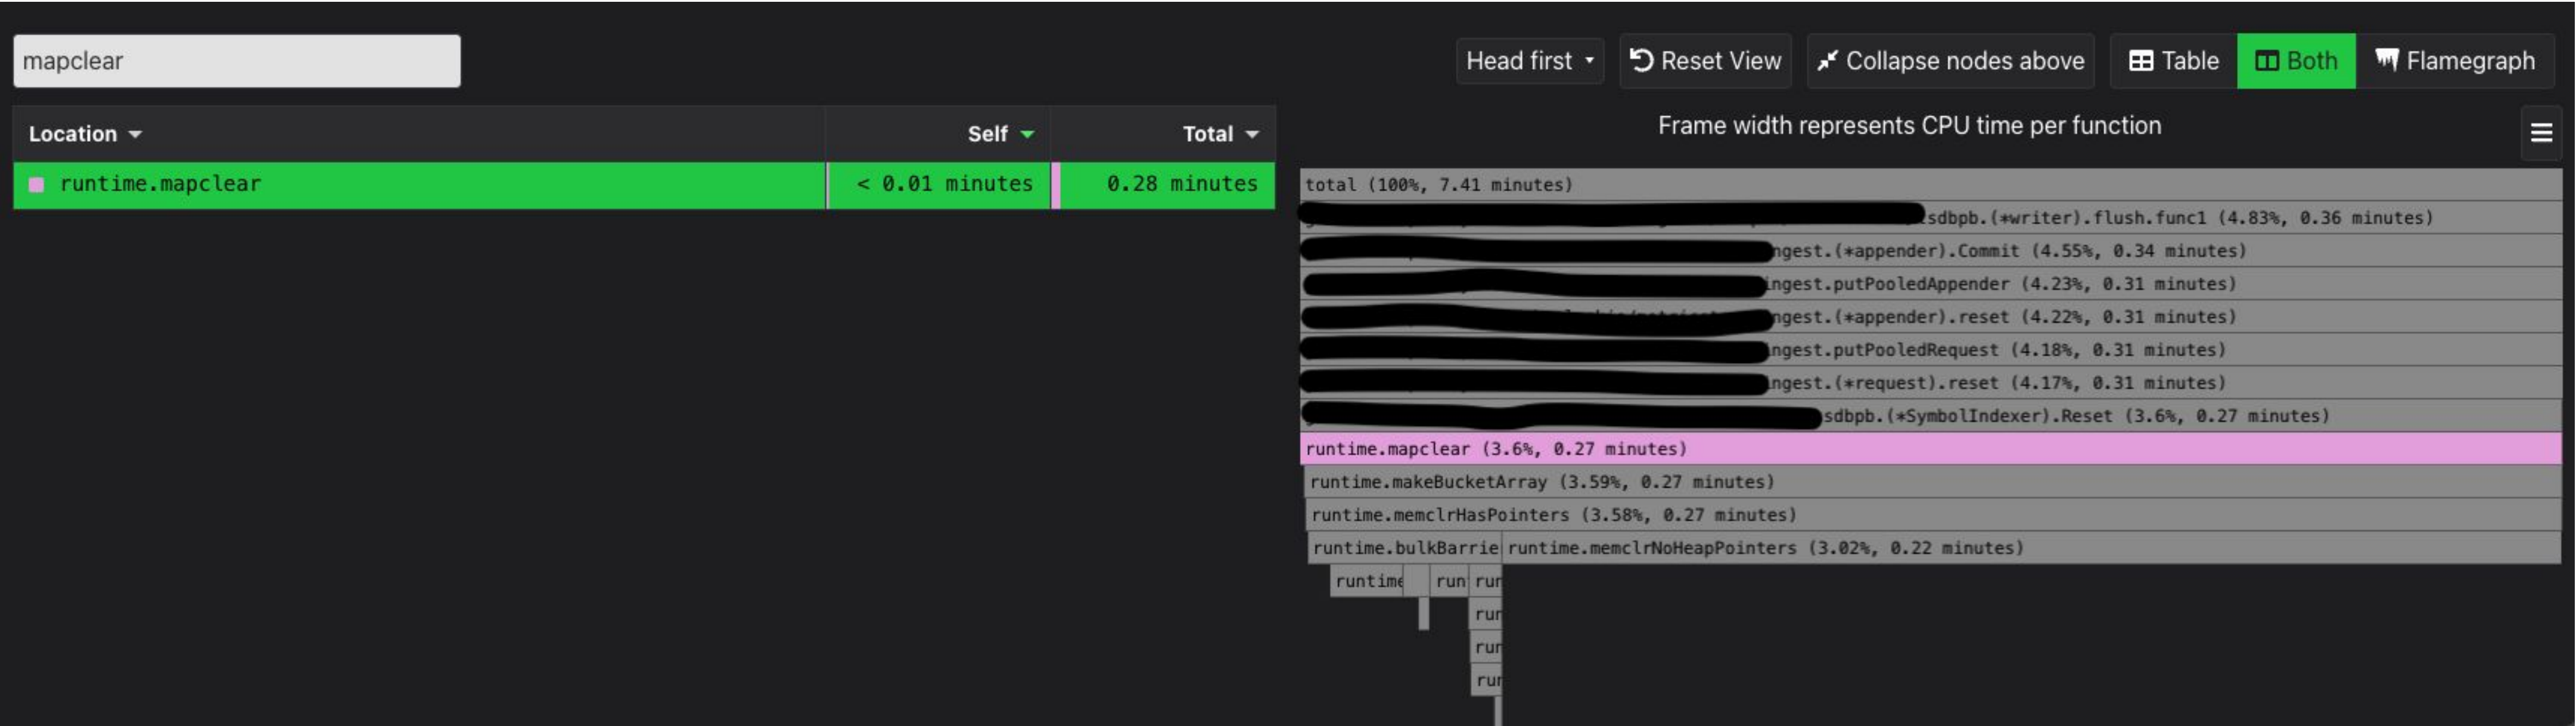
Verify runtime.mapclear is still called for the updated method

```
$ go build -gcflags='-S' tsdb/tsdbpb/symbols.go 2>&1 | grep runtime.mapclear
0x0023 00035 (/tsdb/tsdbpb/symbols.go:43)      CALL      runtime.mapclear(SB)
```

Verify callers do not inline the reset method

```
$ go build -gcflags='-S' ingest/*.go 2>&1 | grep symbols.go | grep reset
0x0056 00086 (/tsdbpb/symbols.go:38)          CALL      /tsdb/tsdbpb.(*SymbolIndexer).reset(SB)
```

Success!










# Example - Removing a Harmful Buffer

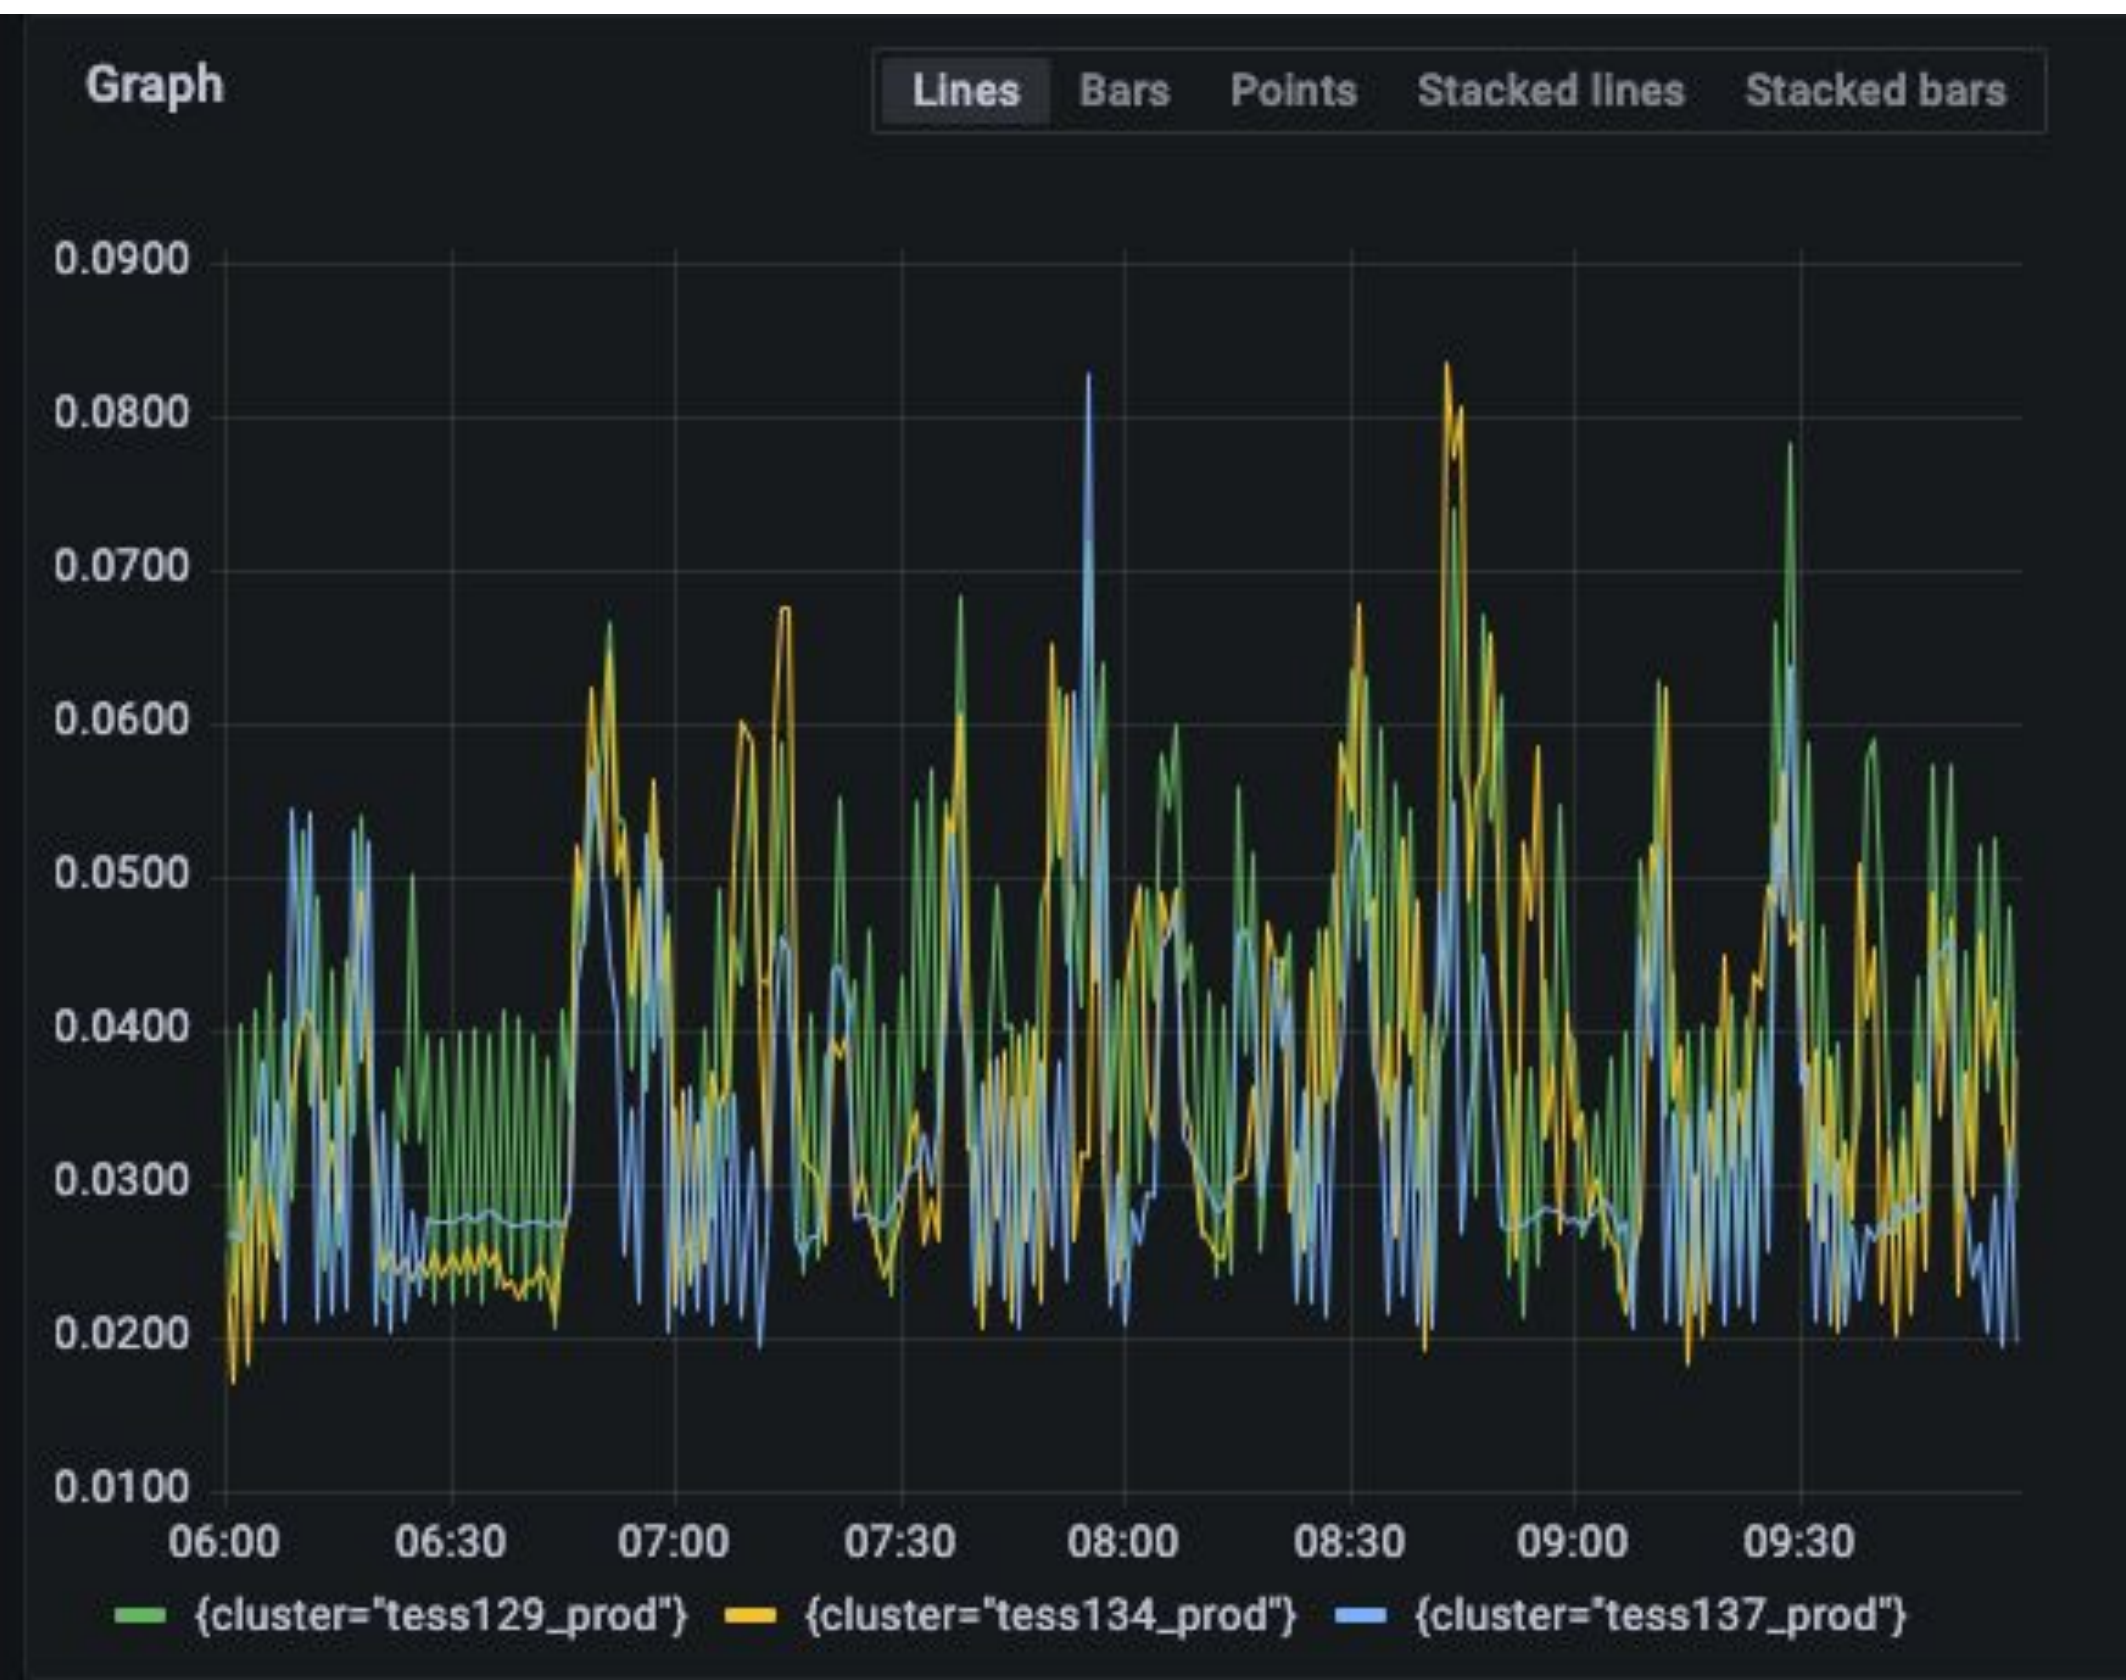
Let's compare cpu and heap usage metrics for additional validation

Cluster 129 is running without the buffer

Memory



CPU



# Demo

# Potential

## Memory Leak/CPU Spike Detection

Allow automation to look at profiles and detect leaks/spikes.

Compare previous and new build to spot anomalous resource usage.

Never have to run a bad build to analyze leaks. Roll it back immediately!!

## 3-Click RCA

Starting from an alert, view metrics, traces and finally profiles to identify method/line level root cause analysis.

Helps drive reduction in time-to-triage.

## Continuous Cost Savings

Analyse performance measures and proactively optimize CPU and memory utilization.

Never waste resources ever again!

Easy to have information readily available than manually profile!

# Questions?



# Thank you