



KubeCon



CloudNativeCon

Europe 2023





KubeCon



CloudNativeCon

Europe 2023

Cortex: How to Run a Rock Solid Multi-Tenant Prometheus

Friedrich Gonzalez, Adobe & Alan Protasio, AWS

Agenda

- Cortex Introduction
- What is the latest on cortex?
- Reliability with Users in mind
- Secret Sauce
- Coming up next!
- Questions






Horizontally Scalable, highly available, multi-tenant, long term storage for Prometheus

- Created in 2016
- CNCF Incubating project
- Apache 2 license

Jun 19, 2016 – Mar 27, 2023

Contributions: Commits ▾

Contributions to master, excluding merge commits and bot accounts



Speakers & Maintainers



Friedrich Gonzalez

Software Engineer
Adobe

@friedrichg



Alan Protasio

Senior Software
Development Engineer
AWS

@alanprot



Other Maintainers



Alvin Lin
Software Dev Manager
AWS

@alvinlin123



Ben Ye
Software Engineer
AWS

@yeya24



Helm Chart Maintainers



Tom Hayward
Infoblox

@kd7lxl

Niclas Schad
STACKIT
@nschad



Other Key Contributors

Harry John
@harry671003

Daniel Blando
@danielblando

Alex Le
@alexqyle

Xiaochao Dong
@damnever

Leon Wang
@wgliang

songjiang
@songjiang

Roy C
@roystchiang

Kama Huang
@kama910


Mengmeng Yang
@mengmengy

Matthew Ames
@SuperMatt



Cortex: Why use Multitenancy?

A. Prometheus scraping a single application (job)





Scenarios:

- High cardinality
- High churn
- Increasing Data Retention



Cortex: Why use Multitenancy?

B. Prometheus scraping multiple applications (multiple jobs) for single owner




Scenarios:

- High cardinality
- High churn
- Increasing Data Retention




Cortex: Why use Multitenancy?

C. Prometheus scraping multiple applications from different teams (multiple jobs with different owners)




Scenarios:

- High cardinality
- High churn
- Increasing Data Retention



Cortex: Why use Multitenancy?

D. Cortex provides long term retention while providing tenant isolation


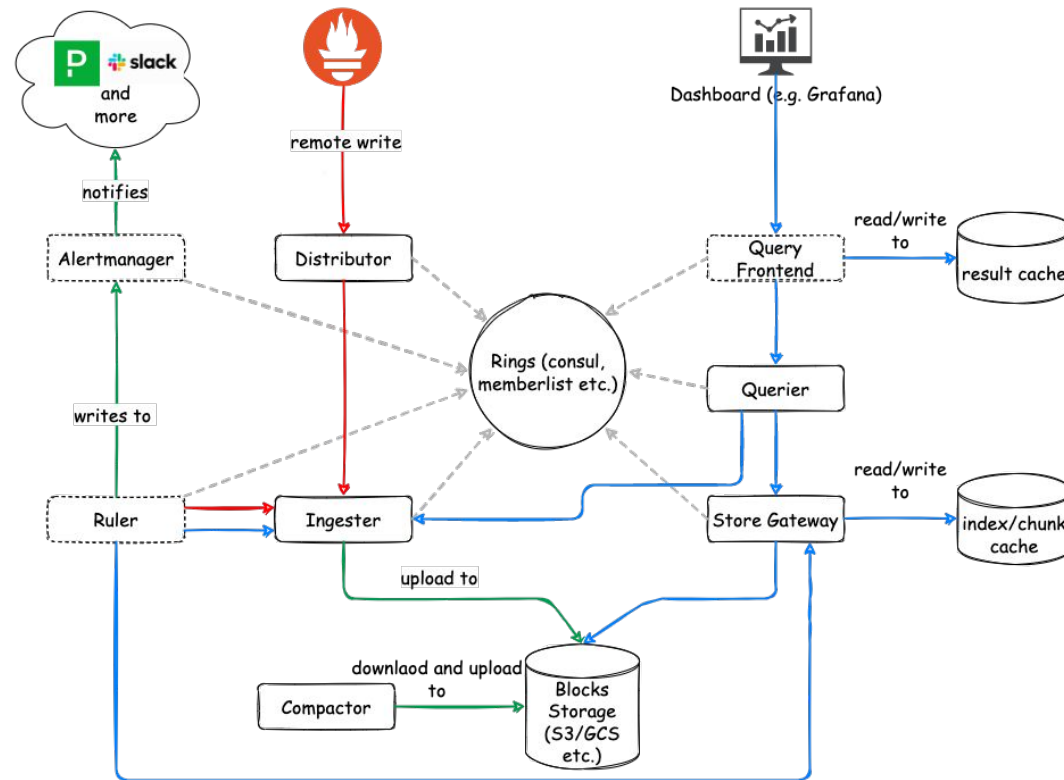


Scenarios:

- High cardinality
- High churn
- Increasing Data Retention



Cortex's Architecture



What is the latest on cortex?

Dec 3, 2022

 alanprot

 v1.14.0

 06d7313 

Compare ▼

Cortex 1.14.0

This release contains 115 contributions from 28 contributors. Thank you!

Some notable changes release are:

- Remove support for chunks storage
- Experimental support for vertical query sharding
- Enable PromQL `@` modifier with negative offset always
- Added configurations for Azure MSI in blocks-storage
- New limits (Querier/QueryFrontend)
- OpenTelemetry Bridge for Tracing
- Multiples performance improvements and bug fixes



What is the latest on cortex?

13 hours ago

 yeya24

 v1.15.0

 92fcee2 

Compare ▼

Cortex 1.15.0

Latest



This release contains 177 contributions from 24 contributors. We also have 13 new contributors. Thank you all for the contribution!

Some notable changes release are:

- Out of order samples ingestion
- MultiKey KV ring for DynamoDB
- Snappy-block gRPC compression
- Redis as index cache and caching bucket backend
- Arm images support
- Thanos PromQL engine support
- Multiples performance improvements and bug fixes



New Cortex Maintainer



Ben Ye
Software Engineer
AWS
Thanos Maintainer

@yeya24



What is the latest on cortex?

Add support for new thanos promql engine #5093

 Merged yeya24 merged 4 commits into `cortexproject:master` from `yeya24:thanos-engine`  on Jan 12

 Conversation 1

 Commits 4

 Checks 13

 Files changed 99



yeya24 commented on Jan 11 • edited by alanprot ▾

Member ...

Signed-off-by: Ben Ye benye@amazon.com

What this PR does:

This pr allows to user to use the experimental [Thanos promql engine](#).

PromQL Query Engine

A multi-threaded implementation of a PromQL Query Engine based on the [Volcano/Iterator model](#).

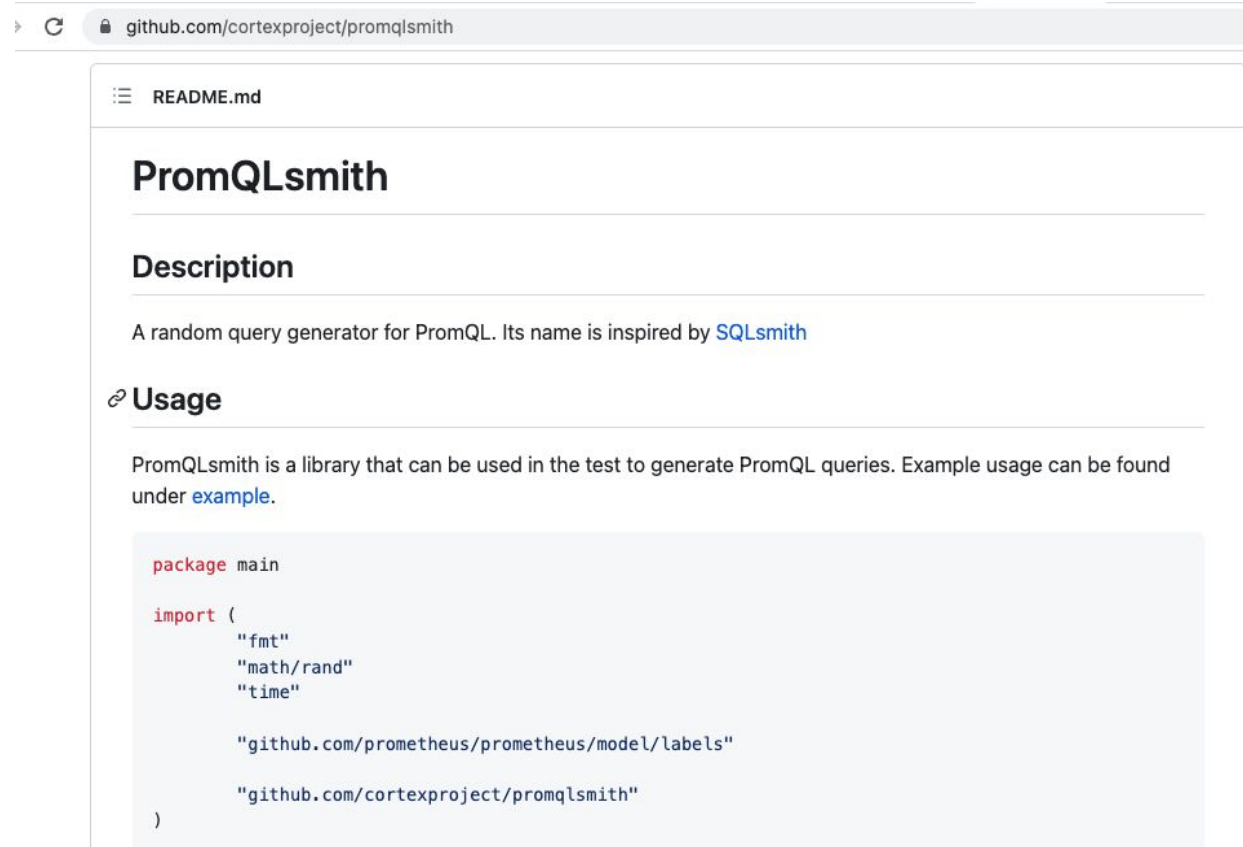
Checklist

- ☒ Tests updated
- ☒ Documentation added
- ☒ `CHANGLOG.md` updated - the order of entries should be `[CHANGE]` , `[FEATURE]` , `[ENHANCEMENT]` , `[BUGFIX]`



What is latest on cortex?

<https://github.com/cortexproject/promqlsmith>



The screenshot shows the GitHub repository page for `cortexproject/promqlsmith`. The page displays the README content, which includes the repository name, a description, and usage instructions. The description states that it is a random query generator for PromQL, inspired by `SQLsmith`. The usage section provides an example of how to use the library in a Go program, showing the necessary imports and the package name.

```
package main

import (
    "fmt"
    "math/rand"
    "time"

    "github.com/prometheus/prometheus/model/labels"
    "github.com/cortexproject/promqlsmith"
)
```



Reliability with Users in mind

What is the single thing users want from monitoring?



Adobe Firefly (Beta)
Image Not for Commercial Use



Reliability with Users in mind

Measure it

- Service Level Indicators with errors and latency

```
sum by (job) (  
  rate(  
    cortex_request_duration_seconds_bucket{le="1.0",route="/cortex.Ingestor/Push",status_code!~"5.."}[5m]  
  )  
)  
/  
sum by (job) (  
  rate(  
    cortex_request_duration_seconds_count{route="/cortex.Ingestor/Push"}[5m]  
  )  
)
```



Reliability with Users in mind

Service Level Objectives

- Be realistic.
- Refine your promise over time
- Page on SLO breaches

2019-11-20 *How to Include Latency in SLO-based Alerting* KubeCon + CloudNativeCon, San Diego, CA, USA




Reliability with Users in mind

Empathize with users and help them achieve their objectives



Secret Sauce



- Kubernetes Ready
- Prometheus based
- Thanos based



Tenant overrides & limits

```
medium_user:: {  
  
    max_series_per_query: 100000,  
  
    max_global_series_per_user: 3000000, // 3M  
    max_global_series_per_metric: 300000, // 300K  
  
    ingestion_rate: 350000, // 350K  
    ingestion_burst_size: 3500000, // 3.5M  
  
    // 1800 rules  
    ruler_max_rules_per_rule_group: 20,  
    ruler_max_rule_groups_per_tenant: 90,  
},
```

```
big_user:: {  
  
    max_series_per_query: 100000,  
  
    max_global_series_per_user: 6000000, // 6M  
    max_global_series_per_metric: 600000, // 600K  
  
    ingestion_rate: 700000, // 700K  
    ingestion_burst_size: 7000000, // 7M  
  
    // 2200 rules  
    ruler_max_rules_per_rule_group: 20,  
    ruler_max_rule_groups_per_tenant: 110,  
},
```



Tenant overrides & limits

```
ingestion_rate: 700000
ingestion_rate_strategy: local
ingestion_burst_size: 7000000
accept_ha_samples: false
ha_cluster_label: cluster
ha_replica_label: __replica__
ha_max_clusters: 0
drop_labels: []
max_label_name_length: 1024
max_label_value_length: 2048
max_label_names_per_series: 30
max_labels_size_bytes: 0
max_metadata_length: 1024
reject_old_samples: false
reject_old_samples_max_age: 2w
creation_grace_period: 10m
enforce_metadata_metric_name: true
enforce_metric_name: true
ingestion_tenant_shard_size: 60
max_exemplars: 0
max_series_per_query: 100000
max_series_per_user: 0
max_series_per_metric: 0
max_global_series_per_user: 6000000
max_global_series_per_metric: 600000
max_metadata_per_user: 8000
max_metadata_per_metric: 10
max_global_metadata_per_user: 0
max_global_metadata_per_metric: 0
out_of_order_time_window: 0s
```

```
max_fetched_chunks_per_query: 0
max_fetched_series_per_query: 0
max_fetched_chunk_bytes_per_query: 0
max_fetched_data_bytes_per_query: 0
max_query_lookback: 0s
max_query_length: 0s
max_query_parallelism: 14
max_cache_freshness: 1m
max_queriers_per_tenant: 0
query_vertical_shard_size: 0
max_outstanding_requests_per_tenant: 100
ruler_evaluation_delay_duration: 0s
ruler_tenant_shard_size: 0
ruler_max_rules_per_rule_group: 20
ruler_max_rule_groups_per_tenant: 110
store_gateway_tenant_shard_size: 0
compactor_blocks_retention_period: 0s
compactor_tenant_shard_size: 0
s3_sse_type: ""
s3_sse_kms_key_id: ""
s3_sse_kms_encryption_context: ""
alertmanager_receivers_firewall_block_cidr_networks: ""
alertmanager_receivers_firewall_block_private_addresses: false
alertmanager_notification_rate_limit: 0
alertmanager_notification_rate_limit_per_integration: {}
alertmanager_max_config_size_bytes: 0
alertmanager_max_templates_count: 0
alertmanager_max_template_size_bytes: 0
alertmanager_max_dispatcher_aggregation_groups: 0
alertmanager_max_alerts_count: 0
alertmanager_max_alerts_size_bytes: 0
```



Instance limits - Ingesters

instance_limits:

- # Max ingestion rate (samples/sec) that ingester will accept. This limit is
- # per-ingester, not per-tenant. Additional push requests will be rejected.
- # Current ingestion rate is computed as exponentially weighted moving average,
- # updated every second. This limit only works when using blocks engine. 0 =
- # unlimited.

- # CLI flag: -ingester.instance-limits.max-ingestion-rate
- [max_ingestion_rate: <float> | default = 0]

- # Max users that this ingester can hold. Requests from additional users will
- # be rejected. This limit only works when using blocks engine. 0 = unlimited.
- # CLI flag: -ingester.instance-limits.max-tenants
- [max_tenants: <int> | default = 0]

- # Max series that this ingester can hold (across all tenants). Requests to
- # create additional series will be rejected. This limit only works when using
- # blocks engine. 0 = unlimited.
- # CLI flag: -ingester.instance-limits.max-series
- [max_series: <int> | default = 0]

- # Max inflight push requests that this ingester can handle (across all
- # tenants). Additional requests will be rejected. 0 = unlimited.
- # CLI flag: -ingester.instance-limits.max-inflight-push-requests
- [max_inflight_push_requests: <int> | default = 0]



Instance limits - Distributors

instance_limits:

Max ingestion rate (samples/sec) that this distributor will accept. This
limit is per-distributor, not per-tenant. Additional push requests will be
rejected. Current ingestion rate is computed as exponentially weighted
moving average, updated every second. 0 = unlimited.

CLI flag: -distributor.instance-limits.max-ingestion-rate
[max_ingestion_rate: <float> | default = 0]


Max inflight push requests that this distributor can handle. This limit is
per-distributor, not per-tenant. Additional requests will be rejected. 0 =
unlimited.

CLI flag: -distributor.instance-limits.max-inflight-push-requests
[max_inflight_push_requests: <int> | default = 0]



Replication Factor and Quorum

- Succeed write request when receiving at least 2 out of 3 positive responses.
- Succeed read when receiving 2 out of 3 positive responses.

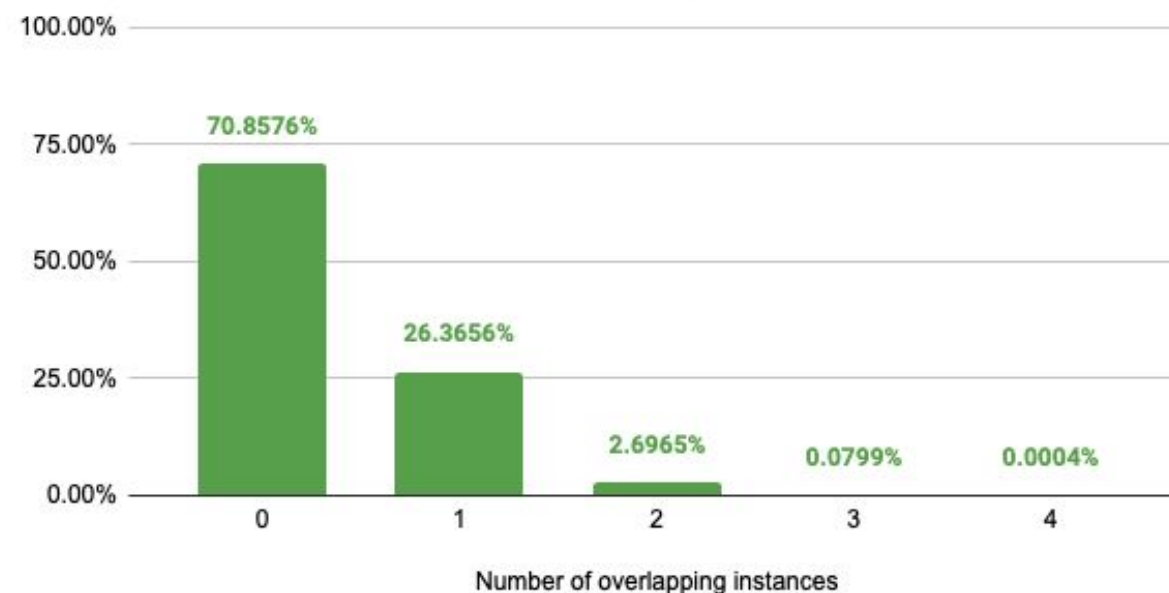


Shuffle Sharding

- Known best practice for workload isolation
- Reduce the blast radius of an outage
- Reduce impact of a misbehaving tenant affecting other tenants.

Probability of overlapping instances between two tenants

50 instances, each tenant with a shard composed by 4 instances



Consistent Hashing: Why?

- Ease scaling: Add new nodes without much rebalancing
- Requires a key/value database.

Rings:

- Consul (2016)
- Etcd (2019)
- Memberlist (2019)
- Dynamodb (2023)

Now also with multikey the hot key problem is gone



Zone aware replication

One replication zone can fail and ingestion and reading stays unaffected



Coming Up Soon


- Auth-Gateway (LFX mentoring project)
- Importing Prometheus TSDB (LFX mentoring project)
- Downsampling
- Federated rules
- Native Histograms



Community

- Join Cortex CNCF channel and let's talk
- Vote on issues, tell us what is important to you
- Help is welcomed!





Please scan the QR Code above
to leave feedback on this session

