



**KubeCon**



**CloudNativeCon**

**Europe 2023**



# Demystifying IPv6 Kubernetes



KubeCon



CloudNativeCon

Europe 2023






**kubernetes**



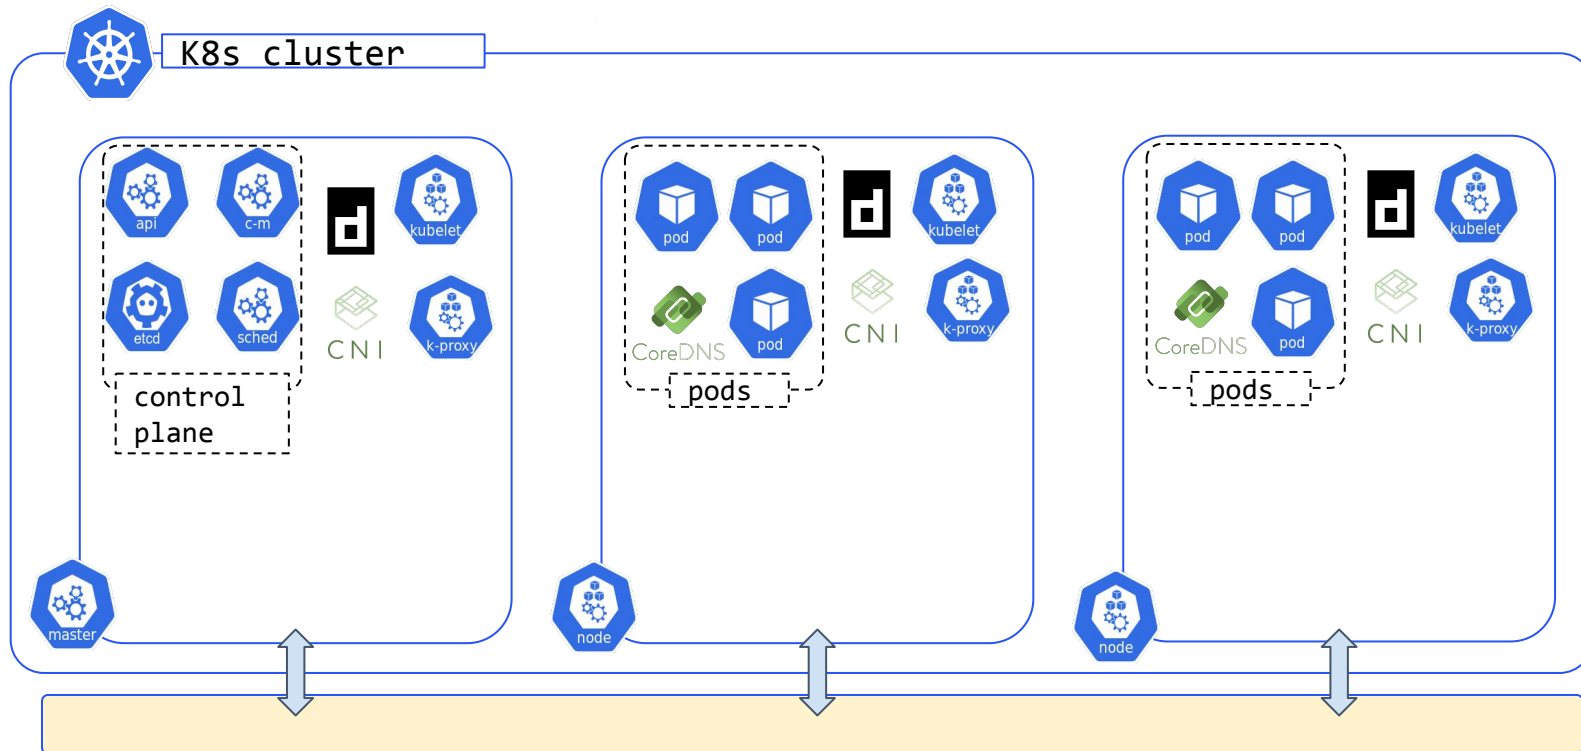
# Kubernetes Networking Architecture



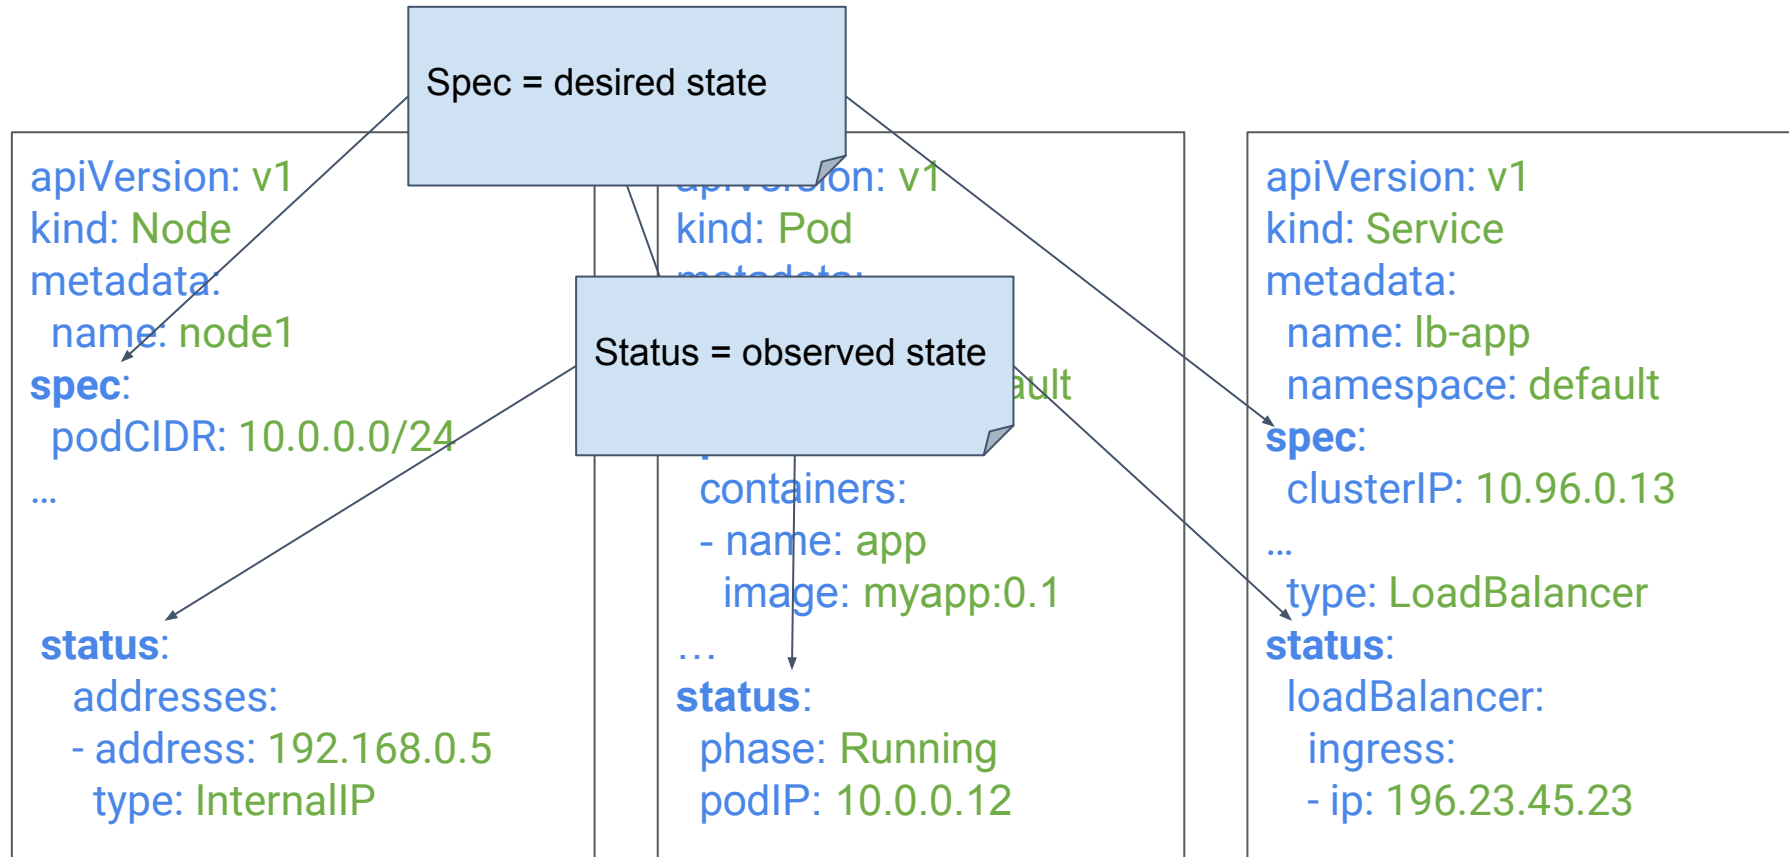
# Kubernetes reference architecture



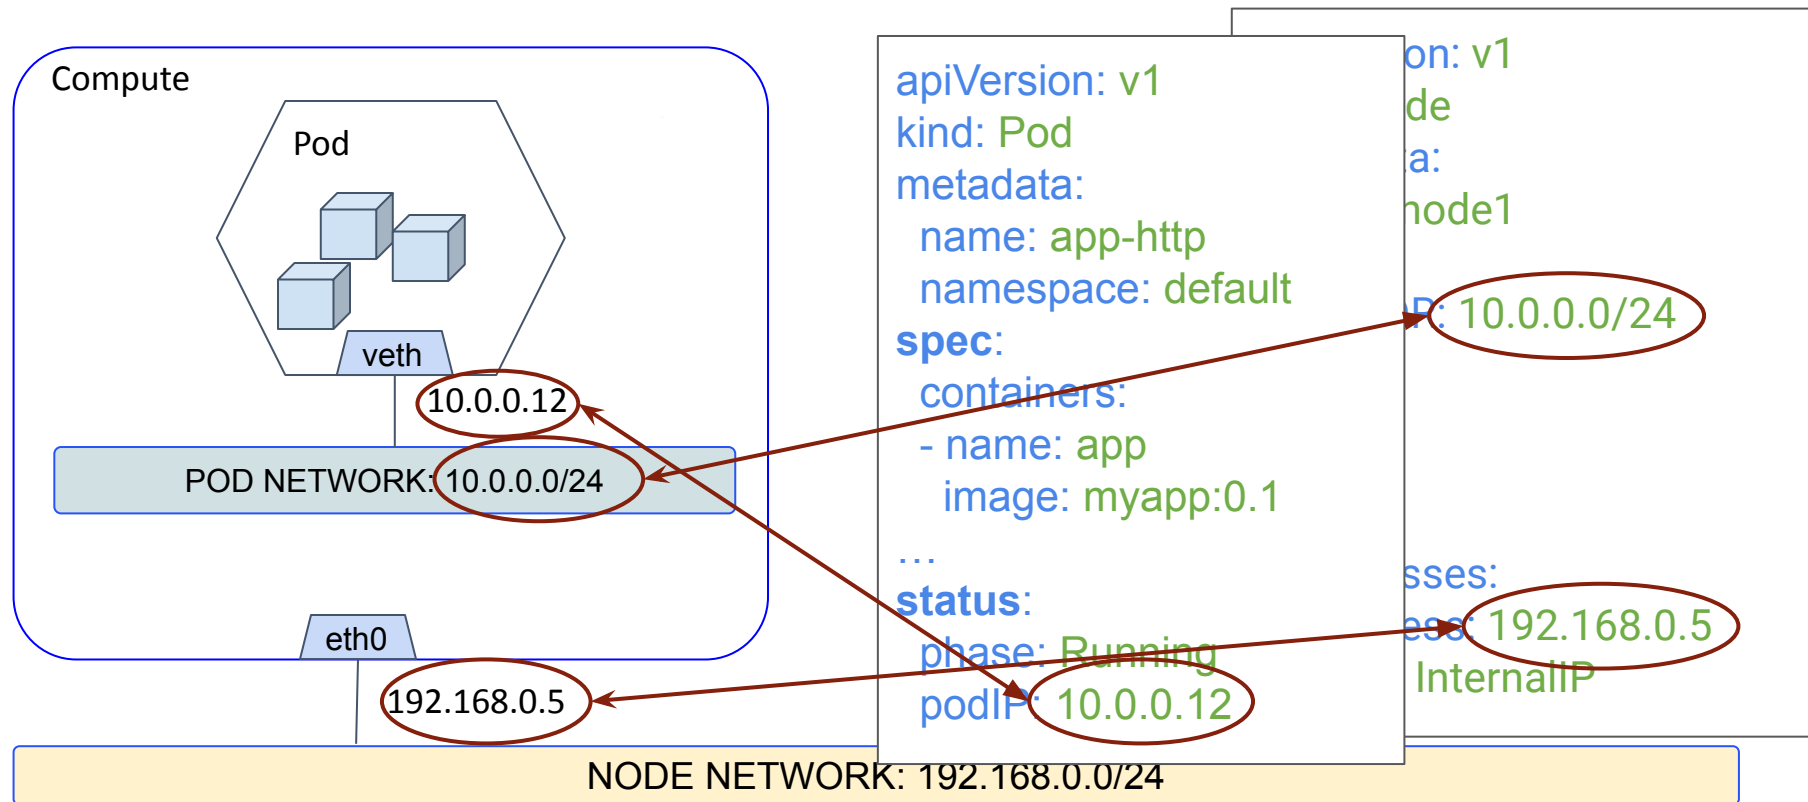
# Kubernetes reference implementation



# Kubernetes API

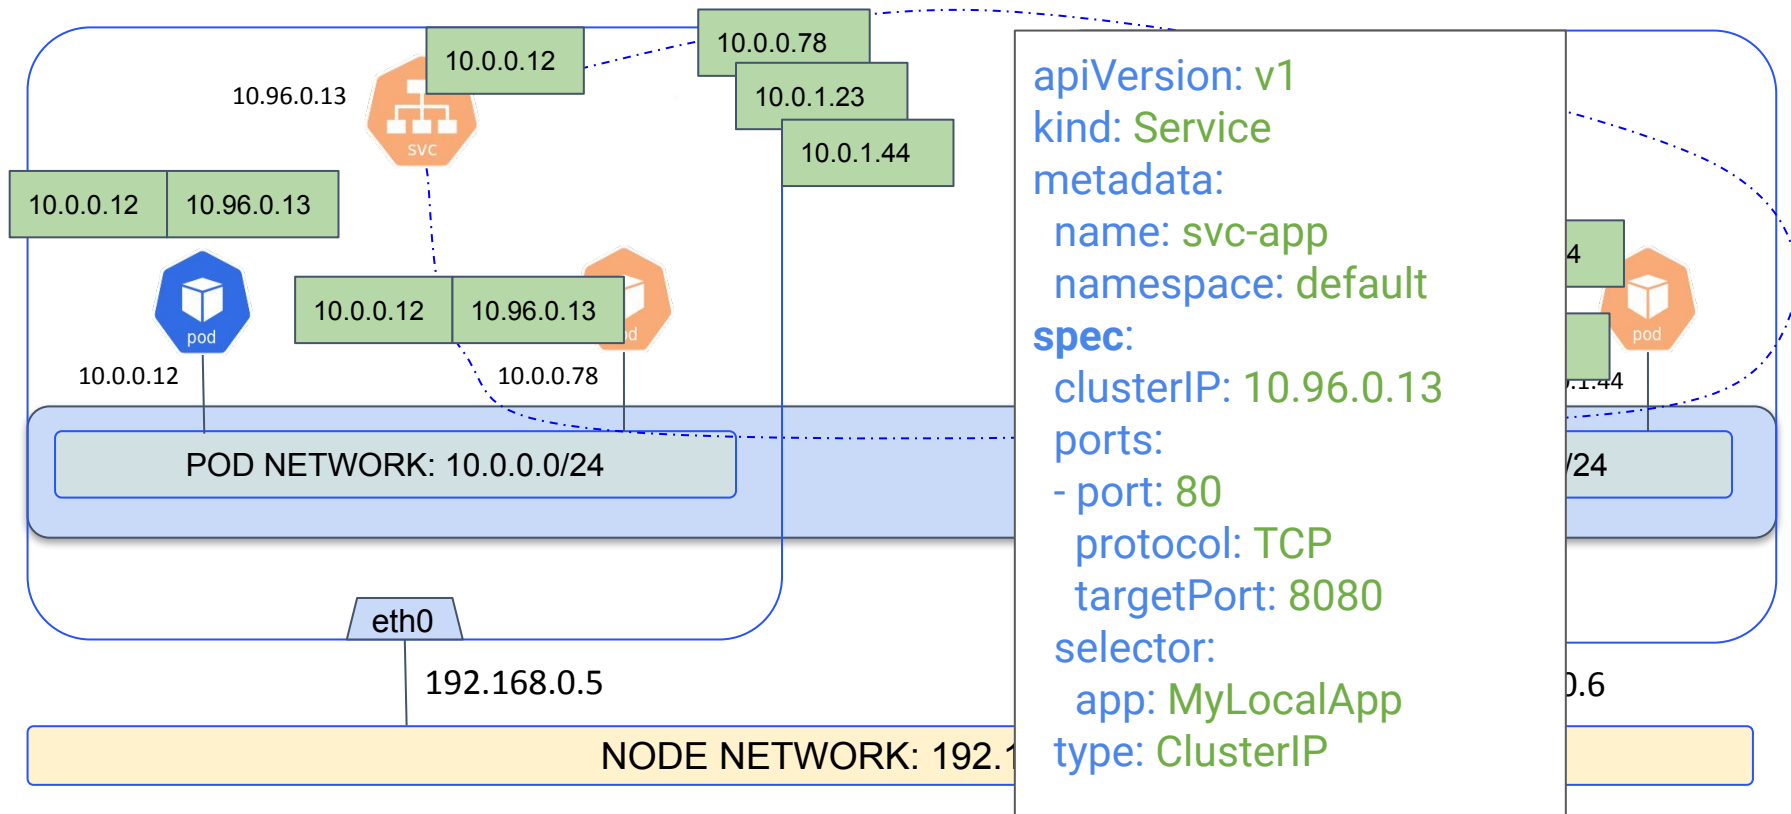


# Kubernetes Networking: Pods





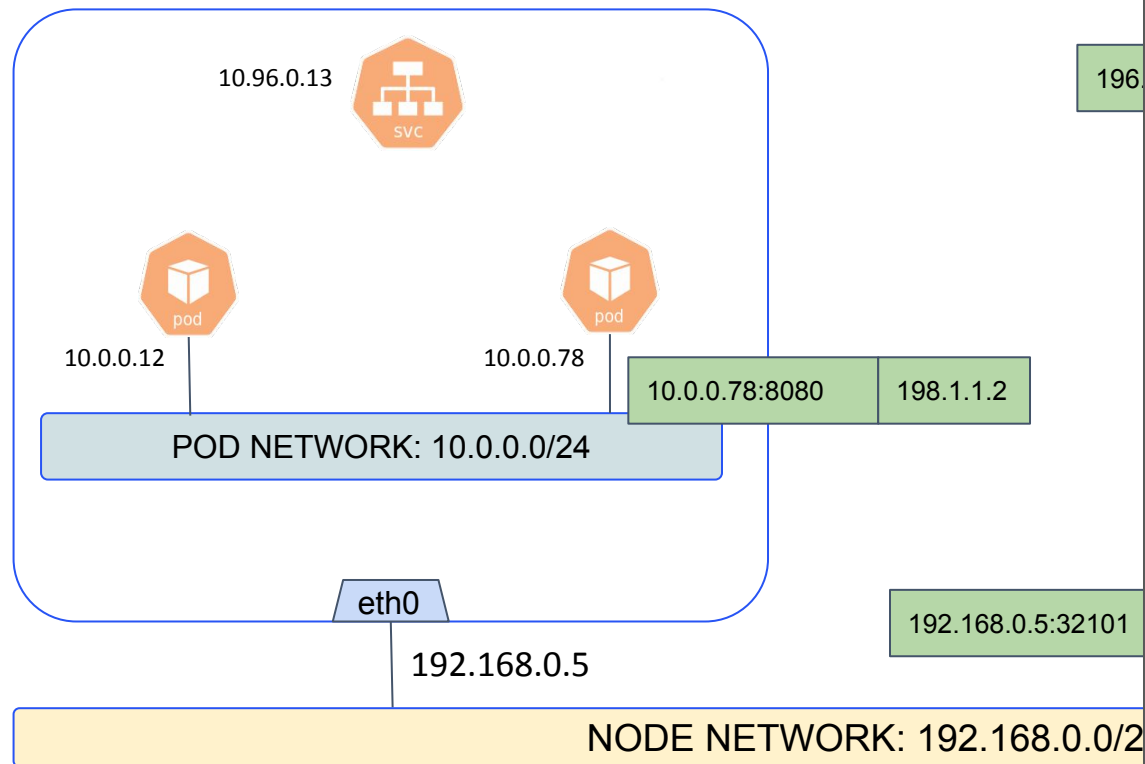
# Kubernetes Networking: Services



# Kubernetes Networking: Services



NativeCon



```
apiVersion: v1
kind: Service
metadata:
  name: lb-app
  namespace: default
spec:
  clusterIP: 10.96.0.13
  ports:
    - nodePort: 32101
      port: 80
      protocol: TCP
      targetPort: 8080
  type: LoadBalancer
status:
  loadBalancer:
    ingress:
      - ip: 196.23.45.23
```

# **Internet Protocol version 6 (IPv6)**




# Brief history of IPv6

- The Internet relies on unique IP addresses for host communication
- But 20+ years ago it was evident IPv4 addresses would be exhausted
- Short-term solutions (stop-gaps):
  - Classless Internet Domain Routing (CIDR)
  - Variable Length Subnet Masking (VLSM)
  - NAT (industry-driven)
- Long term solution:
  - IPv6

# IPv6 in a Nutshell

- Designed to tackle the problem of IPv4 address exhaustion
- Backwards **incompatible** with IPv4
- Original transition/deployment model: dual-stack



# IPv6 Addressing Overview


- IPv6 addresses are 128-bit long
- Similarly to IPv4:
  - Addresses can be aggregated into prefixes (e.g. for routing)
  - Multiple address types (unicast, multicast...)
  - Multiple address scopes (link-local, unique-local, global)
- But IPv6 hosts typically employ **multiple addresses** of different properties
- IPv6 subnets are typically a /64

# IPv6 Unique Local Addresses (ULAs)


- Provide private address space for IPv6
  - IPv6-version of RFC1918 addresses
  - ULA prefix is randomized to avoid address space overlaps
  - Controversial feature in the IPv6 addressing architecture
- Expected use case:
  - Internal communications within a network
  - **Not** meant to be used in combination with IPv6 NAT

# Deployment models: IPv4 vs. IPv6

IPv4



IPv6





**Kubernetes & IPv6**



+



# Kubernetes API: IPv4

apiVersion: v1

kind: Node

metadata:

name: node1

spec:

podCIDR: 10.0.0.0/24

...

status:

addresses:

- address: 192.168.0.5

type: InternalIP

apiVersion: v1

kind: Pod

metadata:

name: app-http

namespace: default

spec:

containers:

- name: app

image: myapp:0.1

...

status:

phase: Running

podIP: 10.0.0.5

apiVersion: v1

kind: Service

metadata:

name: lb-app

namespace: default

spec:

clusterIP: 10.96.0.13

...

type: LoadBalancer

status:

loadBalancer:

ingress:

- ip: 196.23.45.23

# Kubernetes API: IPv6

```
apiVersion: v1
kind: Node
metadata:
  name: node1
spec:
  podCIDR: fd00:1234:1::/64
...

status:
  addresses:
    - address: fd01:789a:1::1
      type: InternalIP
```

```
apiVersion: v1
kind: Pod
metadata:
  name: app-http
  namespace: default
spec:
  containers:
    - name: app
      image: myapp:0.1
...

status:
  phase: Running
  podIP: fd00:1234:1::1
```

```
apiVersion: v1
kind: Service
metadata:
  name: lb-app
  namespace: default
spec:
  clusterIP: fd12:789a:1::1
...
  type: LoadBalancer
status:
  loadBalancer:
    ingress:
      - ip: 2001:db8:7:8::4
```

# Kubernetes API: Dual Stack


```
apiVersion: v1
kind: Node
metadata:
  name: node1
spec:
  podCIDR: 10.0.0.0/24
  podCIDRs:
    - 10.0.0.0/24
    - fd00:1234:1::/64
...
status:
  addresses:
    - address: fd01:789a:1::1
      type: InternalIP
    - address: 192.168.0.1
      type: InternalIP
```

```
apiVersion: v1
kind: Pod
metadata:
  name: app-http
  namespace: default
spec:
  containers:
    - name: app
      image: myapp:0.1
...
status:
  phase: Running
  podIP: 10.0.0.5
  podIPs:
    - 10.0.0.5
    - fd00:1234:1::1
```

```
apiVersion: v1
kind: Service
metadata:
  name: lb-app
  namespace: default
spec:
  clusterIP: 10.96.0.13
  clusterIPs:
    - 10.96.0.13
    - fd12:789a:1::1
...
  type: LoadBalancer
status:
  loadBalancer:
    ingress:
      - ip: 192.168.21.8
      - ip: 2001:db8:7:8::4
```

# IPv6 support in Kubernetes

- One common implementation strategy is replicate the IPv4 architecture with IPv6:
  - Use Unique Local Addresses (private address space) virtually everywhere
  - Expose services with a global (public) address via a load-balancer
- Well-understood model: IPv6 architecture matches the IPv4 counterpart



# Kubernetes & IPv6: The Future Ahead



# IPv6 support in Kubernetes (revisited)


- Kubernetes IPv6 support mimics its IPv4 counterpart
- Are we missing a chance to leverage IPv6 capabilities?
- How could we possibly evolve IPv6 support?
- What if we were to leverage Global Unicast Addresses?

# Leveraging global IPv6 addresses

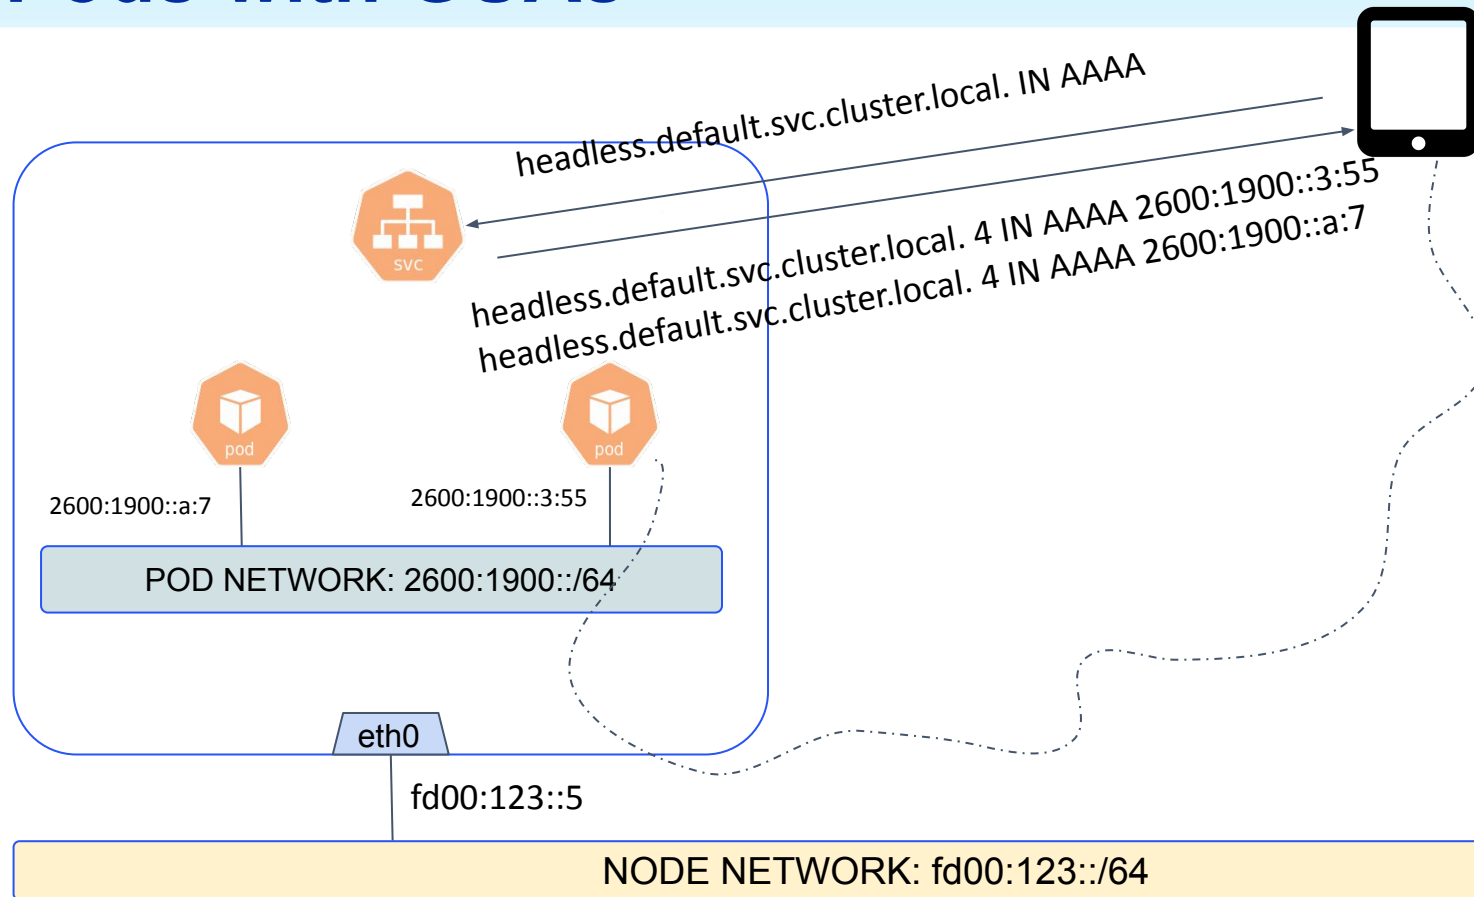
- Services with Global Unicast Addresses (GUAs)
  - Remove one layer of translation!
- Pods with GUAs, and leverage the DNS
  - Remove two layers of translation!



# Services with GUAs



# Pods with GUAs



**TO BE  
CONTINUED...** ➡

**TO BE  
CONTINUED...** 



*To be continued...*



**KubeCon**



**CloudNativeCon**

Europe 2023

# Questions?





Please scan the QR Code above  
to leave feedback on this session