



KubeCon



CloudNativeCon

Europe 2023





KubeCon



CloudNativeCon

Europe 2023

Distributing and Decentralizing Pod Disruption Budget (PDB)

Illya Chekrygin



Disruption* Sucks

* disturbance or problems which interrupt an event, activity, or process

Types of Disruption

Involuntary

Hardware, VM, Network Failures

Machine failure

VM failure

- accidental deletion
- Kernel panic

Network partition

Kubelet Eviction
(out of resources)

Taint based Eviction

Voluntary

Caused or consented by the workload owner

Deletion
Update
Automation

Caused by Infrastructure or Admin

Scheduler Eviction

Disruption Protection - PDB

```
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: test
  namespace: default
spec:
  minAvailable: 3
  selector:
    matchLabels:
      app: test
status:
  currentHealthy: 6
  desiredHealthy: 3
  disruptionsAllowed: 3
  expectedPods: 6
```

Per Namespace

“Healthy” pod requirements (absolute or %-age)

Pod Selector

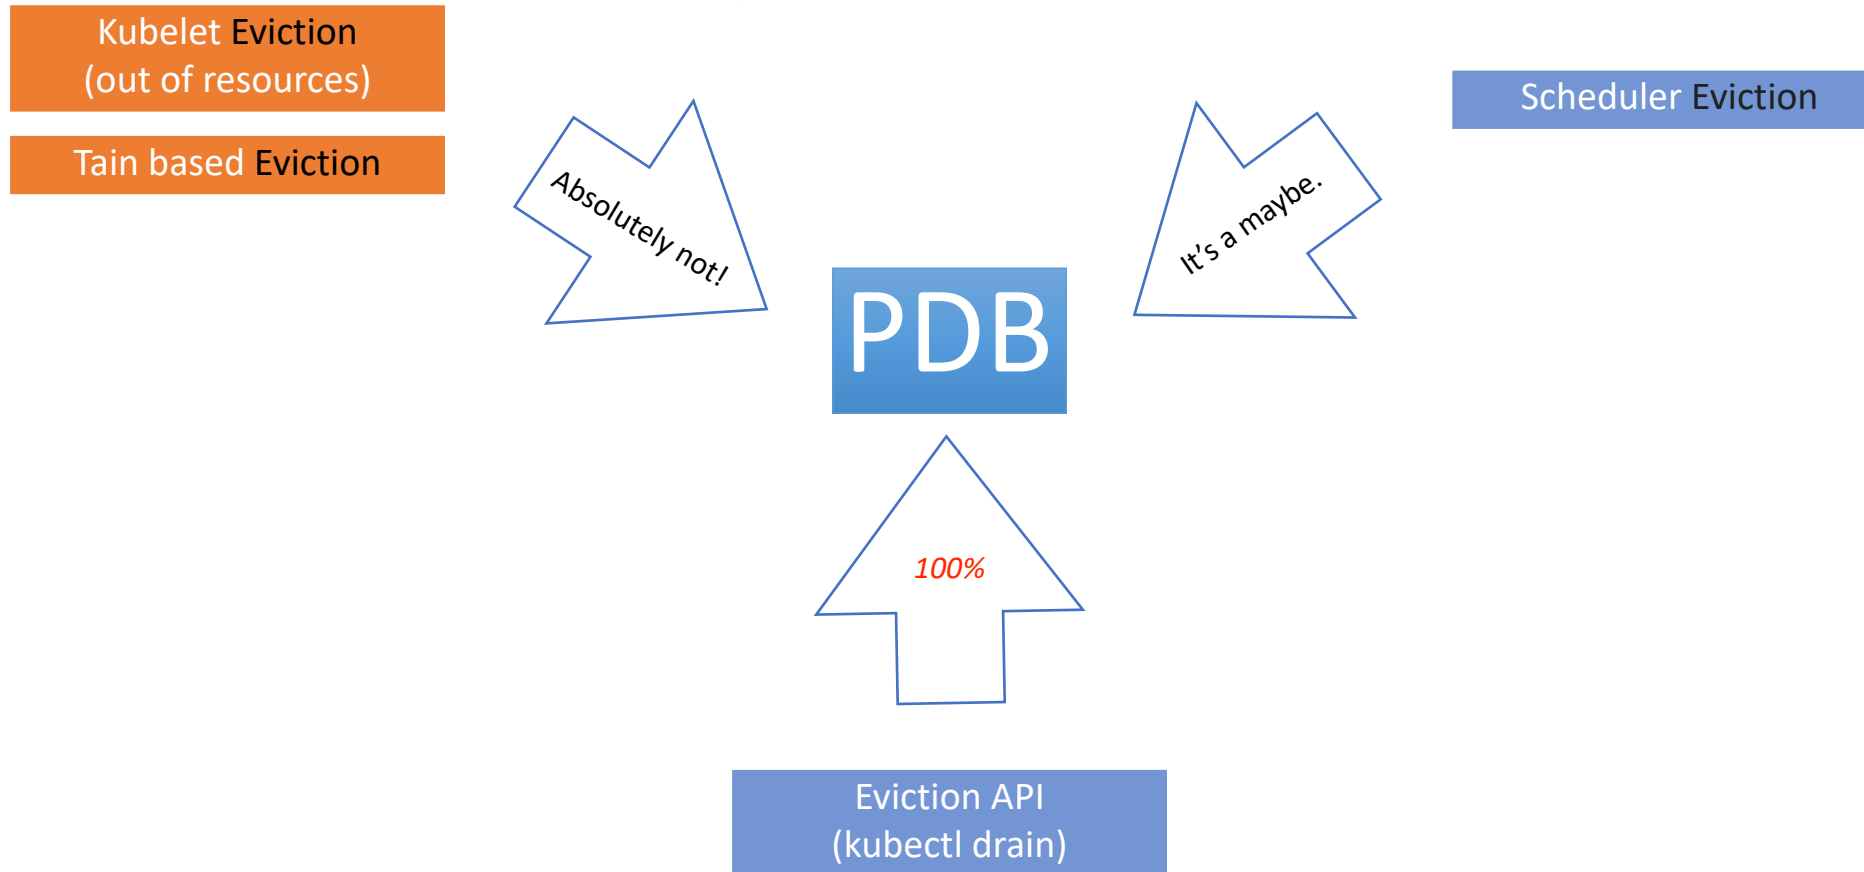
How many “healthy” pods are there at this moment?

How many “healthy” pods are required?

How many can be disrupted at this moment?

How many “healthy” pods are expected?

PDB - Support in K8s



The Good

Simple

Eviction Protection

The Not So Good

Rudimentary (too simple)

Namespace Scoped

Not Universal

**Does not support
complex identity pods**


Not extensible

The Very Not So Good


Err500 on multiple PDBs per Pod

Not extensible

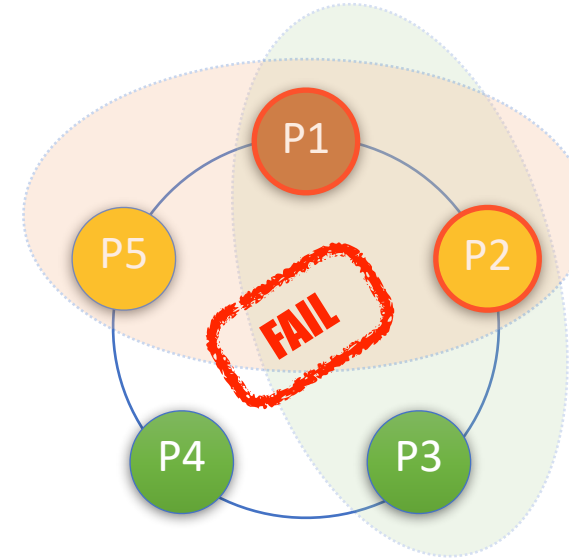
Distributed NoSQL DB Use case




P1: P5,P2
P2: P1,P3
P3: P2,P4
P4: P3,P5
P5: P4,P1



P1: P5,P2
P2: P1,P3
P3: P2,P4
P4: P3,P5
P5: P4,P1




PDB1[P1,P5,P2]
PDB2[P2,P1,P3]
...




[P1,P5,P2] x N
...

Federated PDB

Unidirectional




Bidirectional

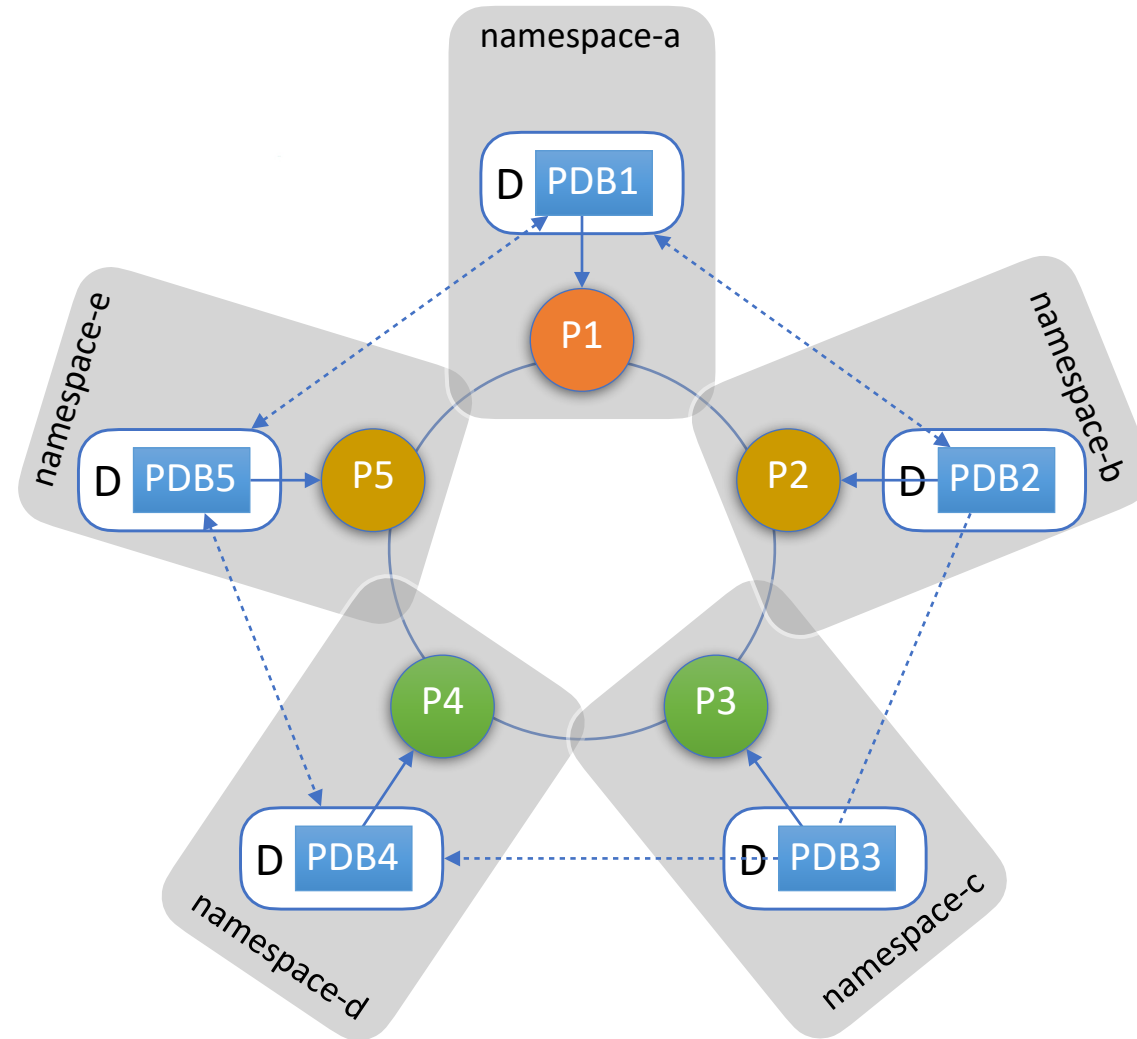


Distributed DB Federated PDB

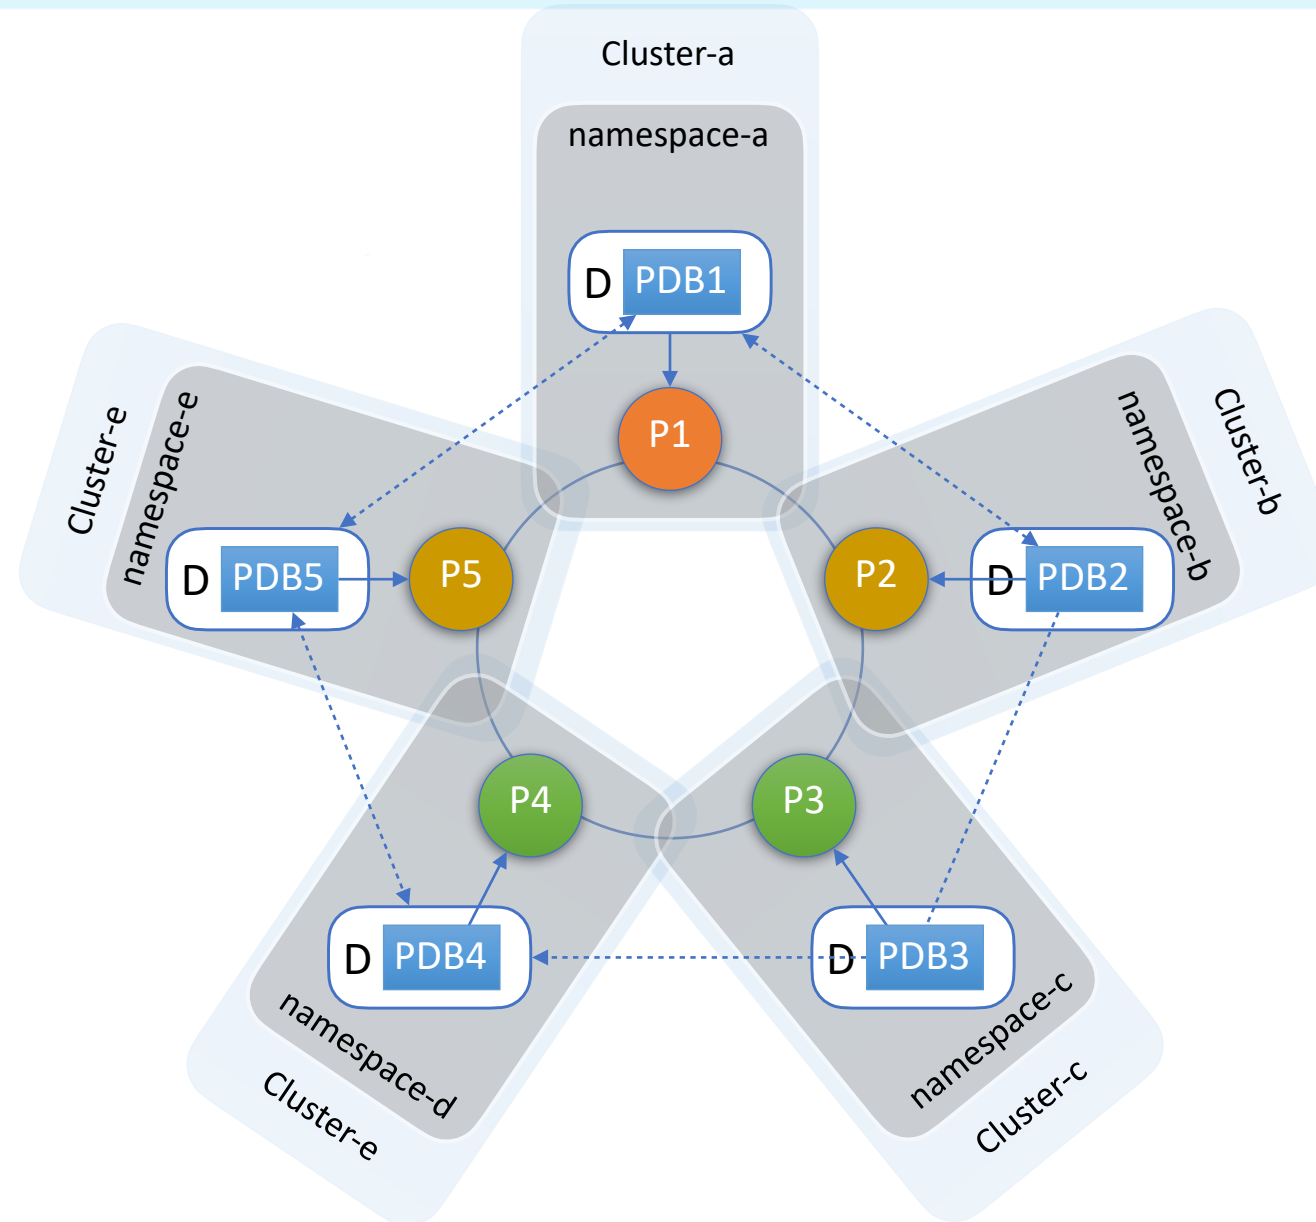
```
apiVersion: demo.illya.com/v1alpha1
kind: DistributedPodDisruptionBudget
metadata:
  name: instance-1
spec:
  minAvailable: 2
  selector:
    matchLabels:
      app: instance-1
  federation:
    - name: instance-5
    - name: instance-2
```



Multi-Namespace PDB



Multi-Cluster PDB





KubeCon



CloudNativeCon

Europe 2023

DEMO TIME



Extensibility

Or the lack thereof for built-in types.

Your Story

Thank you!

<https://github.com/apple/distributed-disruption>

<https://github.com/ichekrygin>

Slack:

Kubernetes - ichekrygin

CNCF: Illya Chekrygin

<https://github.com/apple/distributed-disruption>



<https://github.com/ichekrygin>

Slack:

Kubernetes - ichekrygin

CNCF: Illya Chekrygin

[to leave feedback on this session](#)