



KubeCon




CloudNativeCon

Europe 2023

Image Signing and Runtime Verification at Scale: Datadog's Journey

Ethan Lowman, Datadog



Observability and security platform

- Tens of trillions of events per day
- Millions of hosts
- Over 600 integrations
- Almost 5,000 employees

Runs on self-hosted Kubernetes

- Dozens of clusters
- Tens of thousands of nodes
- Hundreds of thousands of pods



About Me

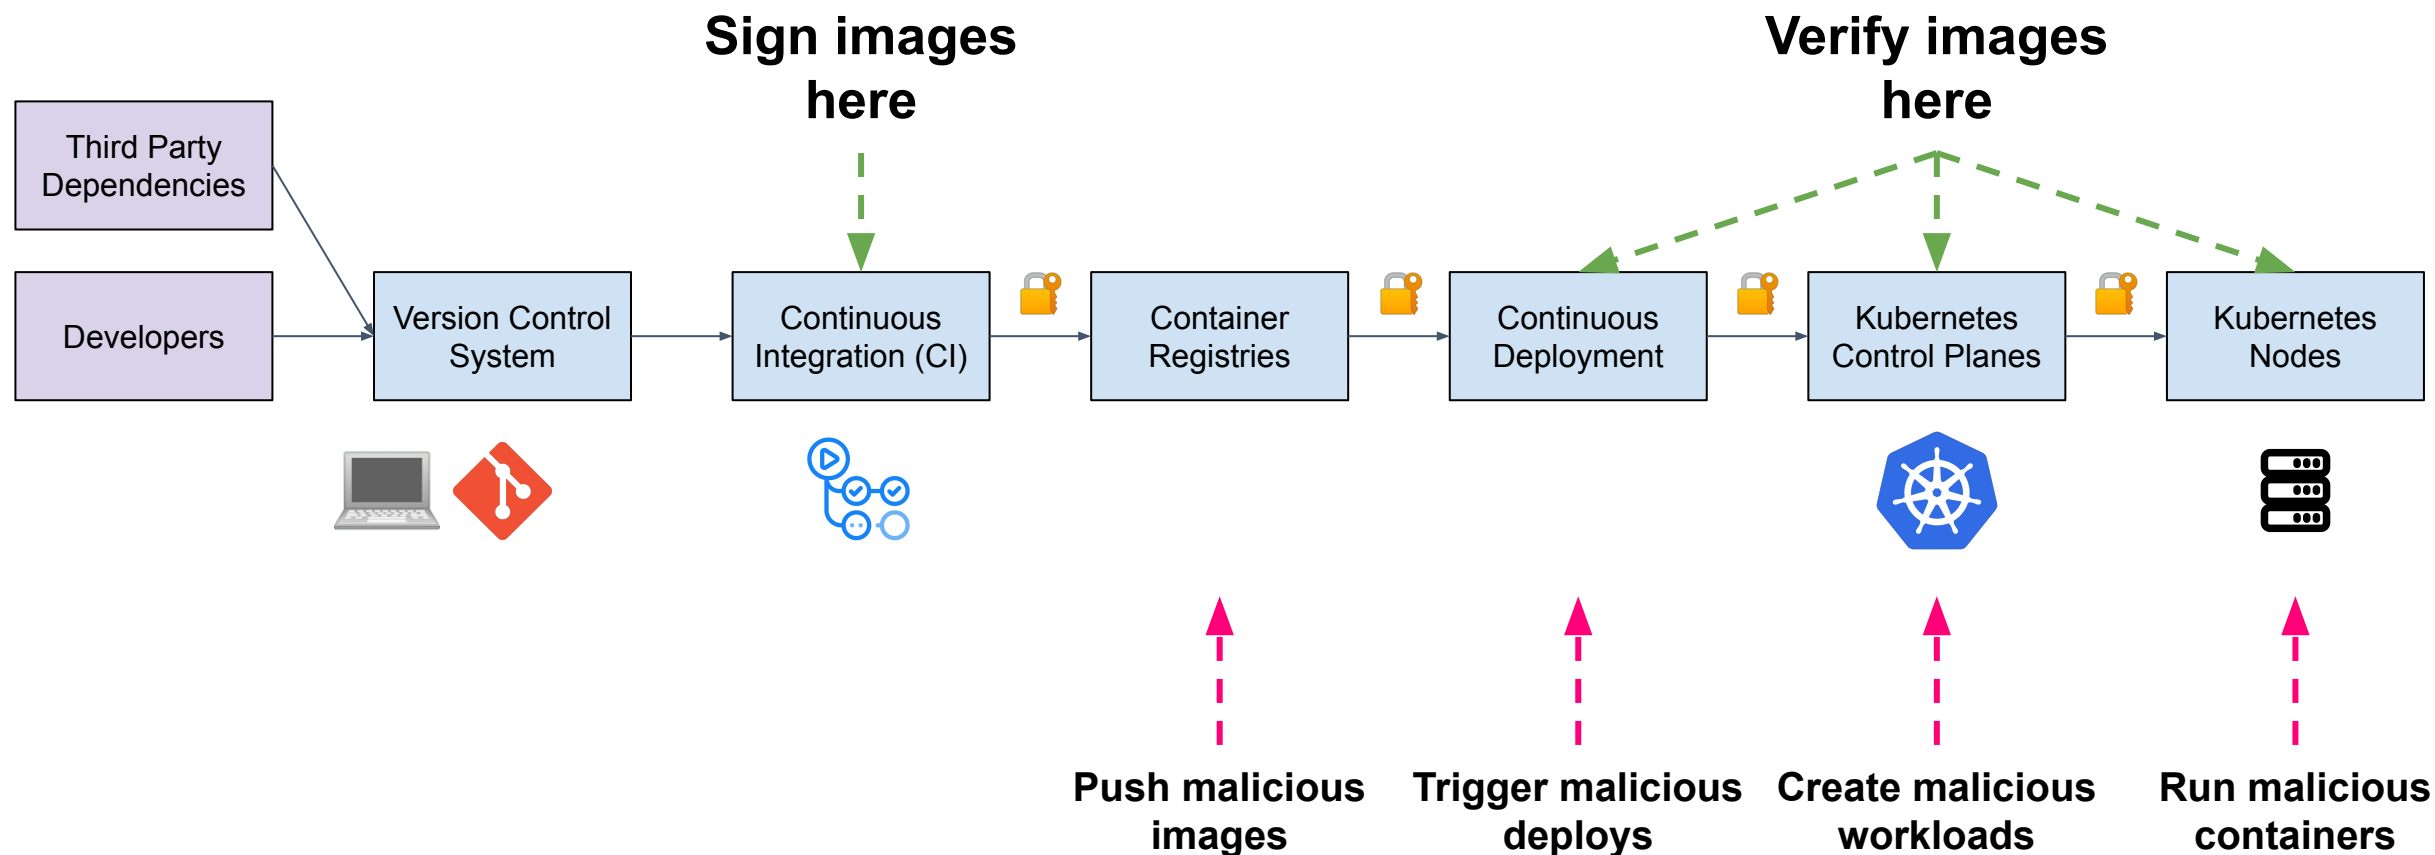


Ethan Lowman

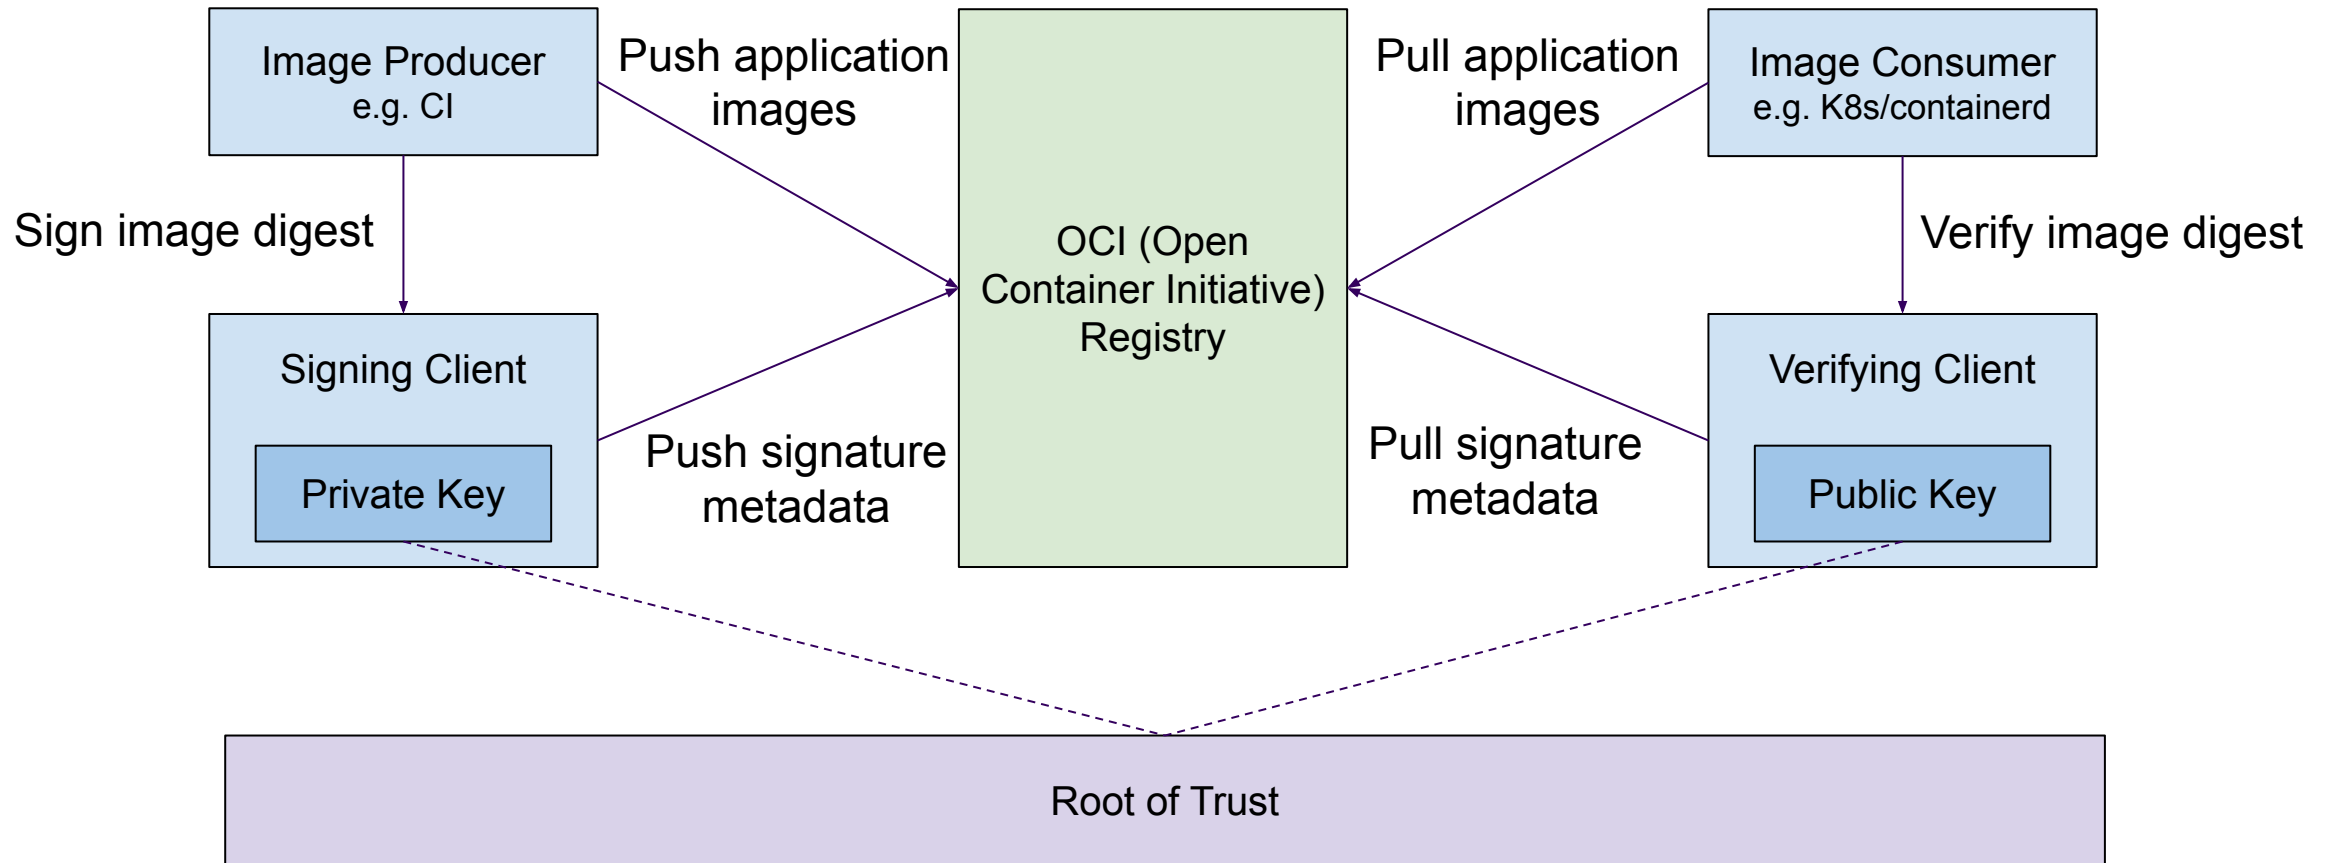
Senior Software Engineer
Datadog

ethan.lowman@datadoghq.com

Why sign & verify images?



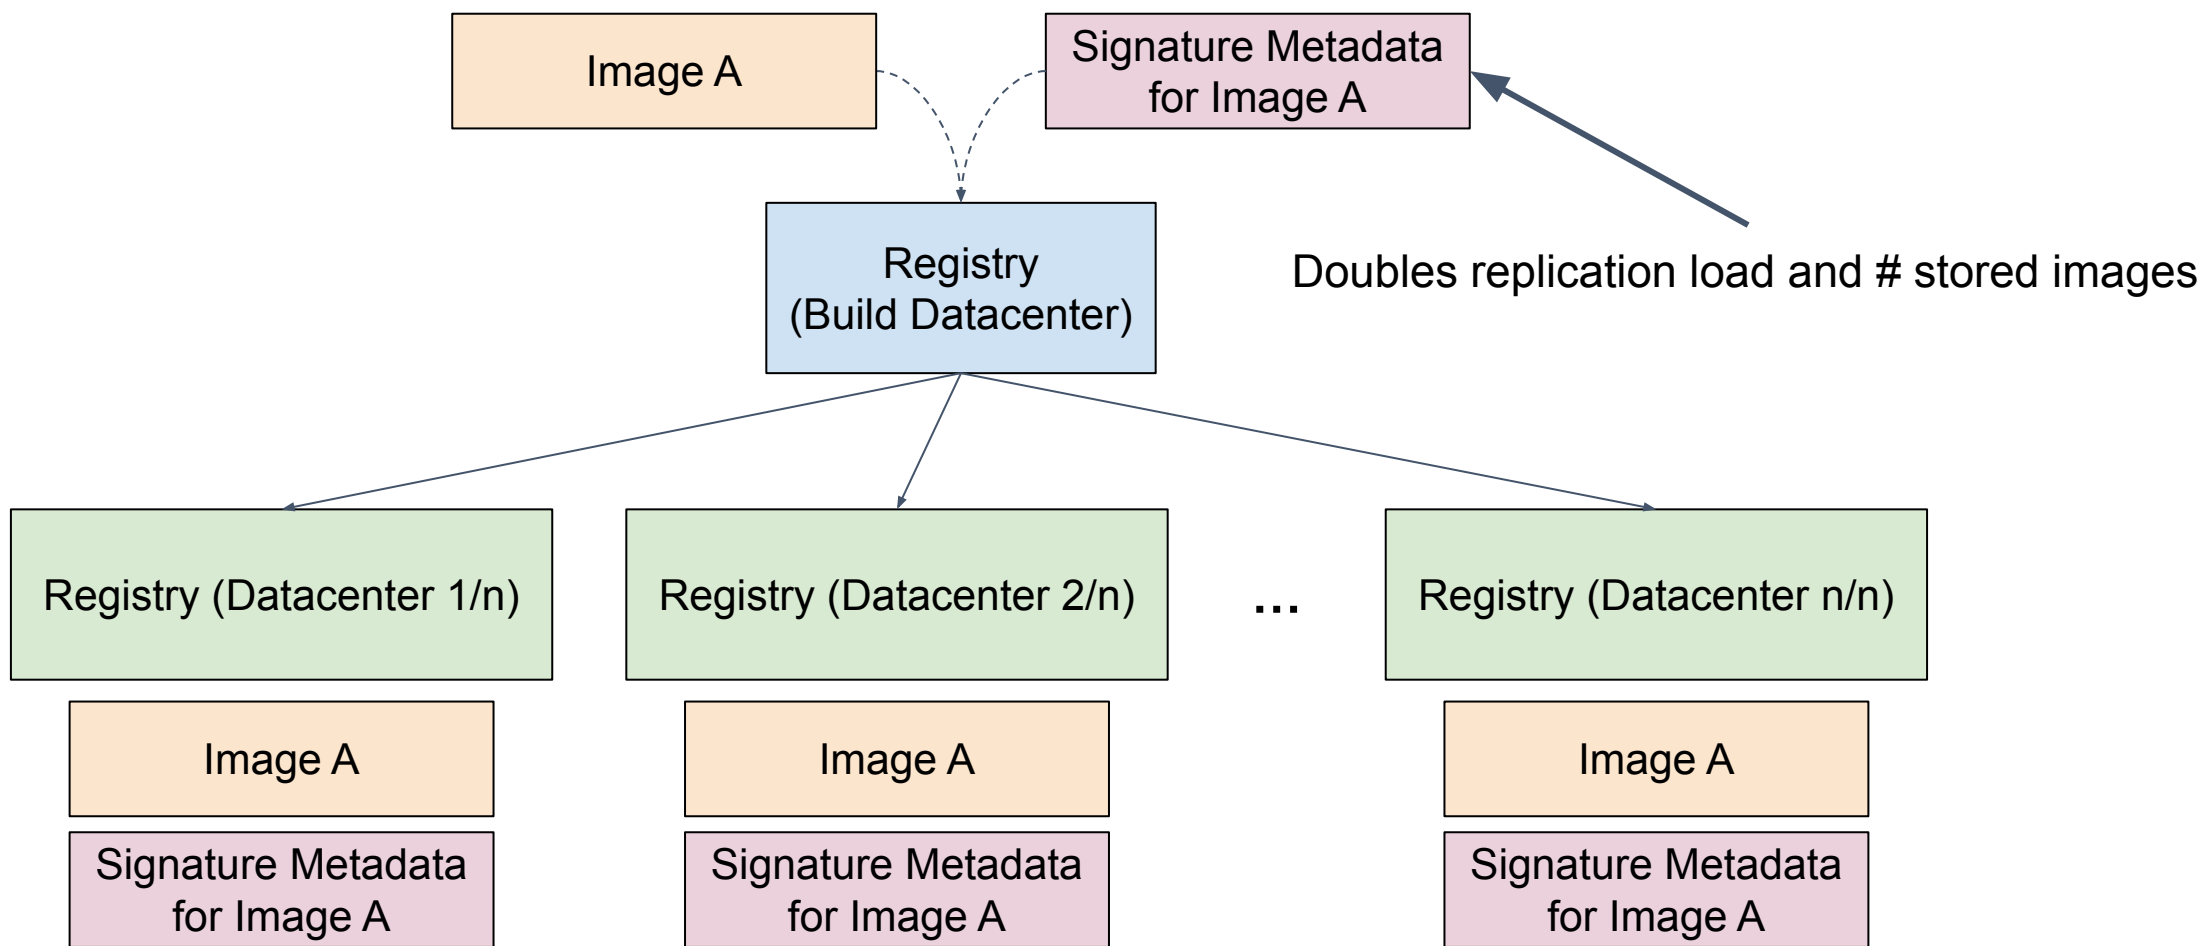
Modern Consensus on Image Signing



Examples: Sigstore/cosign, Notary v2

Signature Metadata in a Registry

- The registry is not a new runtime dependency
- Signature replication is a solved problem if you're already running images in multiple datacenters



Our design is loosely based on Sigstore's cosign signatures.

Caveat

Strongly consider using an open standard if one is available, especially if you need interoperability. Cosign might meet our requirements if we started over today.

Signature Format: Payload

OCI Descriptor with custom annotations for signature metadata

```
{
  "annotations": {
    "com.datadoghq.image-integrity.signer.claims.timestamp": "2023-02-03T21:48:33-04:00",
    "com.datadoghq.image-integrity.signer.claims.client_subject": "0638c9b7-43b3-a2fe-...",
    "com.datadoghq.image-integrity.signer.claims.client_email": "example@internal.service.identity",
    // etc.
  },
  "digest": "sha256:adab3844f497ab9171f070d4cae4114b5aec565ac772e2f2579405b78be67c96",
  "mediaType": "application/vnd.docker.distribution.manifest.list.v2+json",
  "size": 1638
}
```

Digest, media type, and size of
the protected artifact

Custom signature metadata fields
(e.g. signer identity & environment)

Signature Format: Envelope

Signature Algorithm: Ed25519

Envelope Format: Dead Simple Signing Envelope (DSSE)

```
{
  "payload": "<base64-encoded OCI Descriptor>",
  "payloadType": "application/vnd.oci.descriptor.v1+json",
  "signatures": [
    {
      "keyid": "SHA256:TfImnpH/n1t3tBnj0Nv8XGcejhvxFye4vdrk+TwyKnE",
      "sig": "<base64-encoded Ed25519 signature>"
    }
  ]
}
```

Signature Format: OCI Layers

Each OCI layer stores one DSSE blob, annotated by key ID for direct lookup

```
{
  "schemaVersion": 2,
  "mediaType": "application/vnd.oci.image.manifest.v1+json",
  "config": {
    "mediaType": "application/vnd.oci.image.config.v1+json",
    "size": 233,
    "digest": "sha256:c4a08471155229261aa4944de79b3b7c3206105f64299338cd055f9f4c59a2bf"
  },
  "layers": [
    {
      "mediaType": "application/vnd.datadog.image-integrity.dsse.v1+json",
      "size": 1360,
      "digest": "sha256:94047bef27d03f89ff17184a9d00a5e38463a0d869dcf4b676ba8c8b8352c083",
      "annotations": {
        "com.datadoghq.image-integrity.v1/key-id": "SHA256:TfImnpH/n1t3tBnj0Nv8XGcejhvxFye4v..."
      }
    }
  ]
}
```

One layer = one signing envelope for one key

Signature Format: Registry Layout

Resolve signature reference from signed digest

Signed Artifact Location:

registry.example.com/**my/image**@**sha256:abc123**...

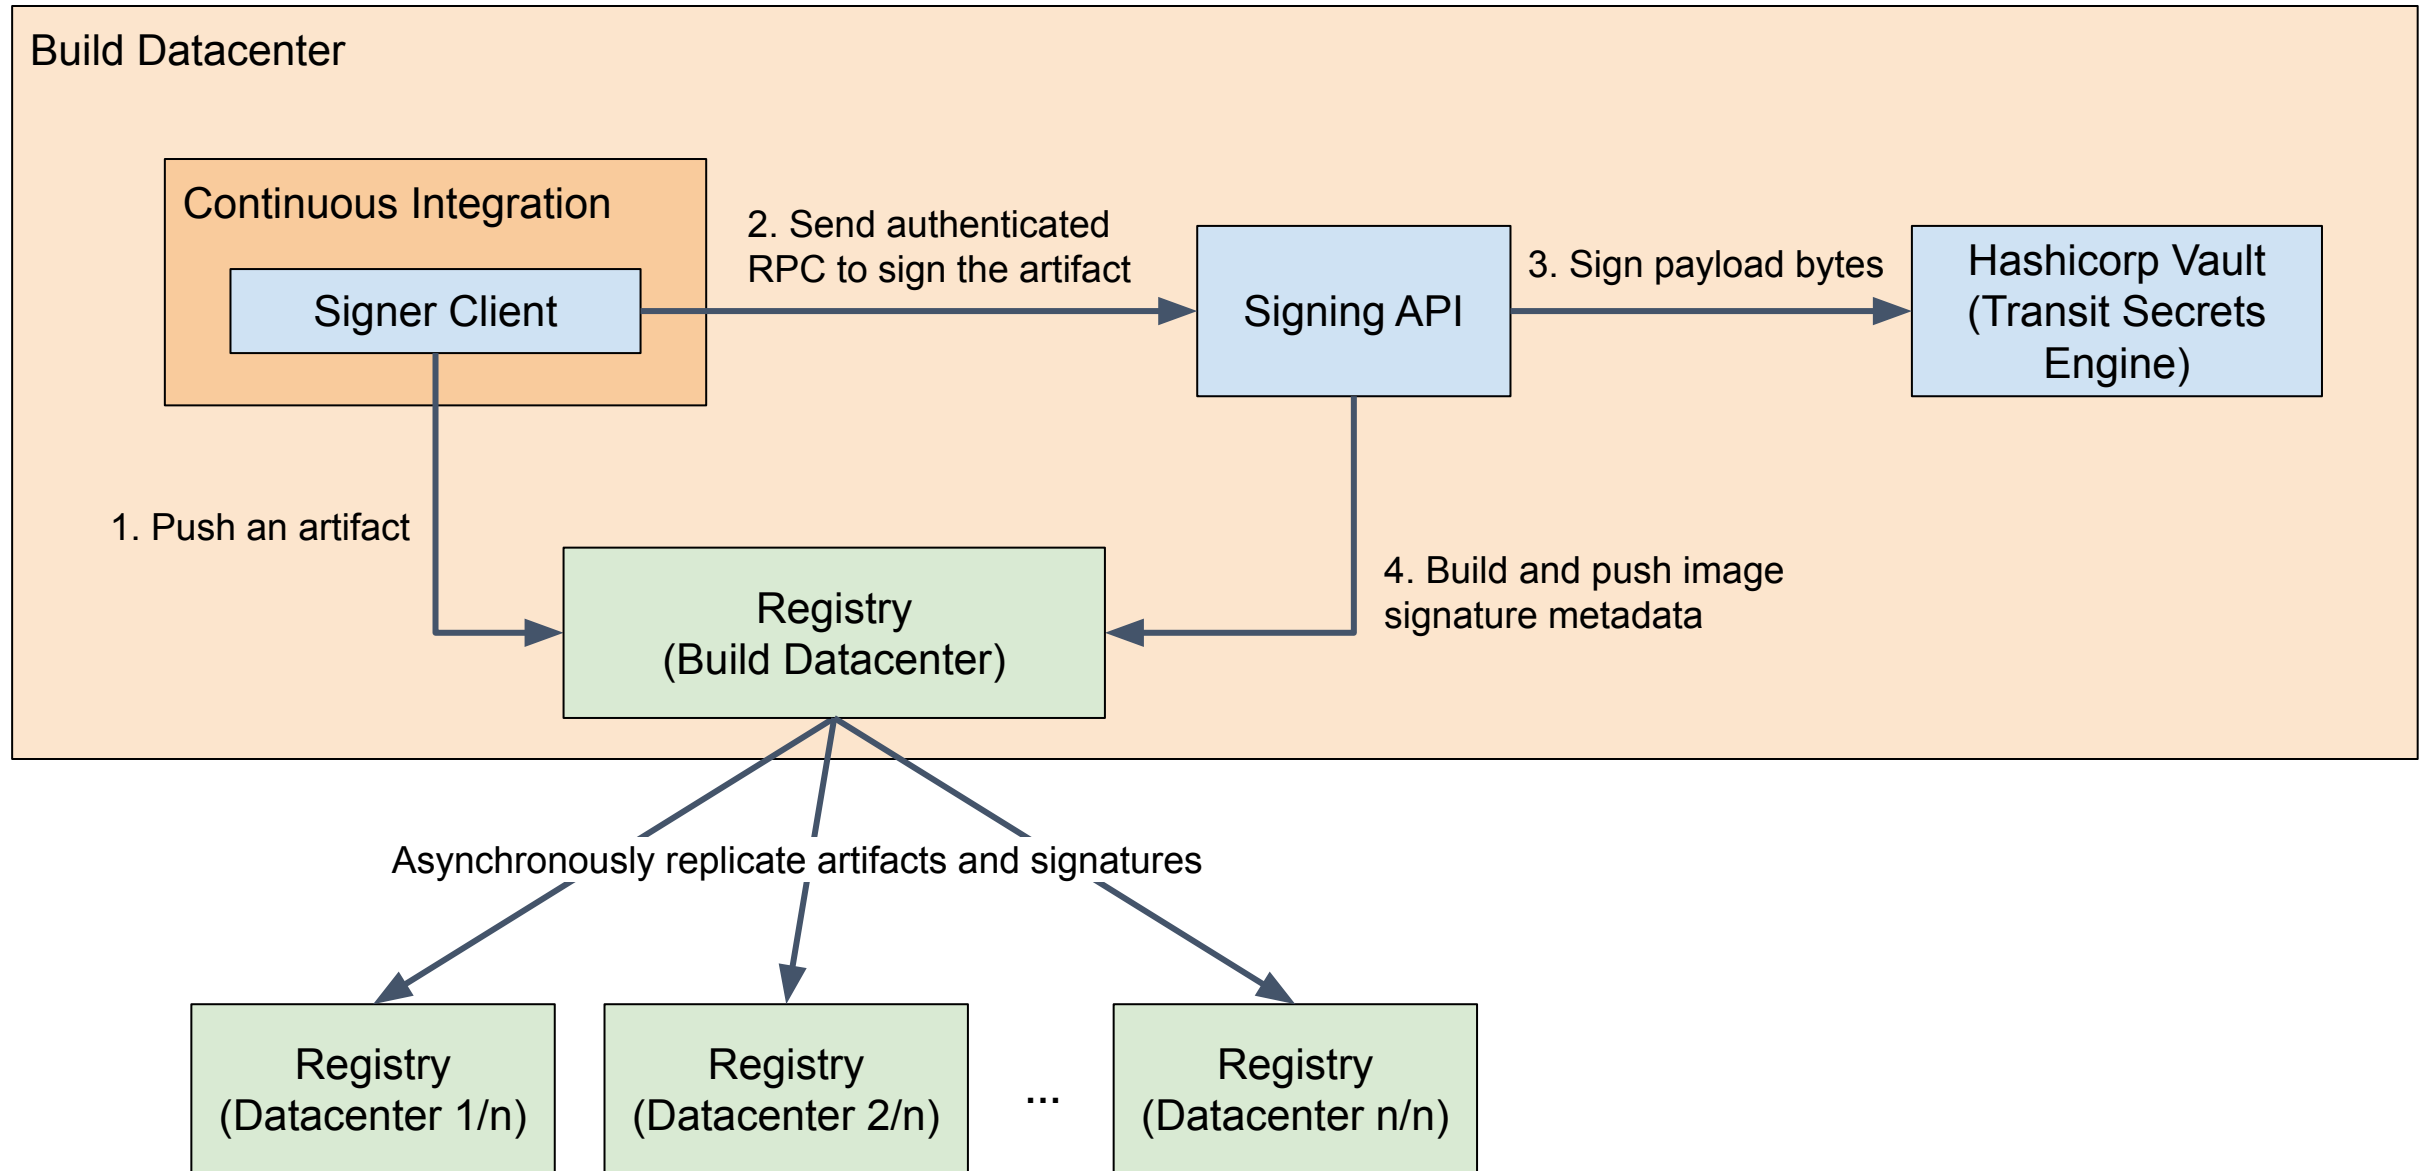
Signature Location:

registry.example.com/**signatures/my/image**:**sha256-abc123**...

Property:

Signatures have dedicated cloud registry storage quotas, independent from the signed images.

Signing as a Service



Signing Thin Client

General Purpose (Shell)

```
ddsign sign registry.internal/my-image:v3@sha256:abc123...
```

Docker-Built Images (Shell)

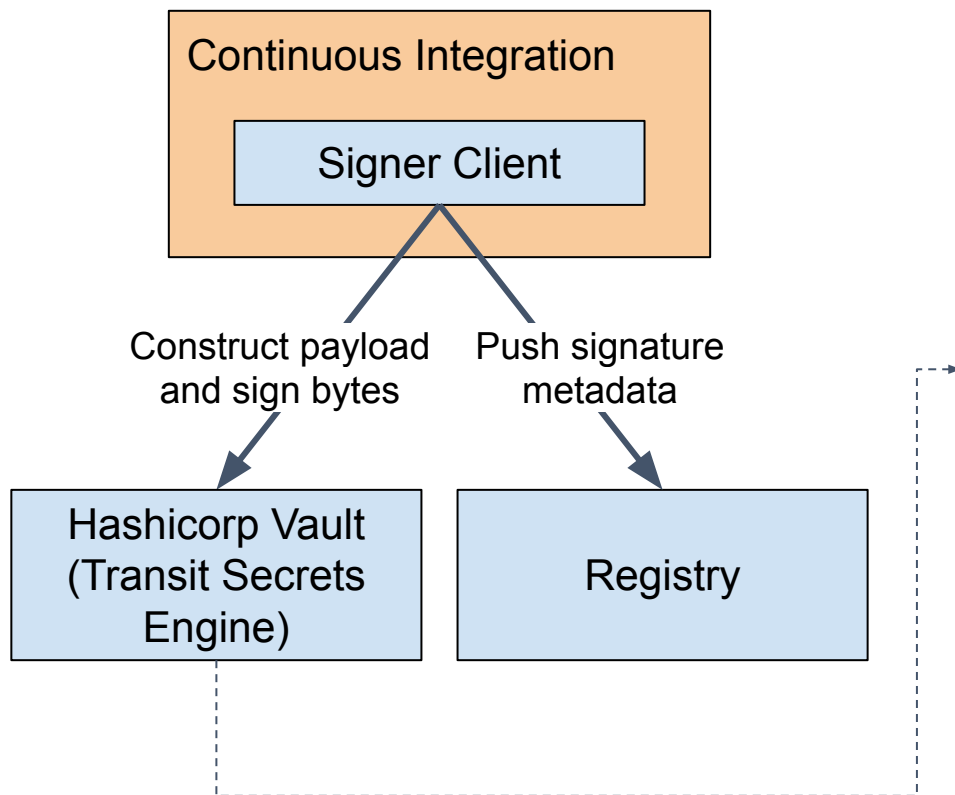
```
ddsign sign registry.internal/my-image:v3 --digest-from-docker
```

Bazel

```
dd_container_sign(  
    name = "sign-my-image",  
    pushed_image = ":pushed-my-image",  
)
```

Signing Service: Least Privilege & Auditability

Suppose CI signed with Hashicorp Vault directly:



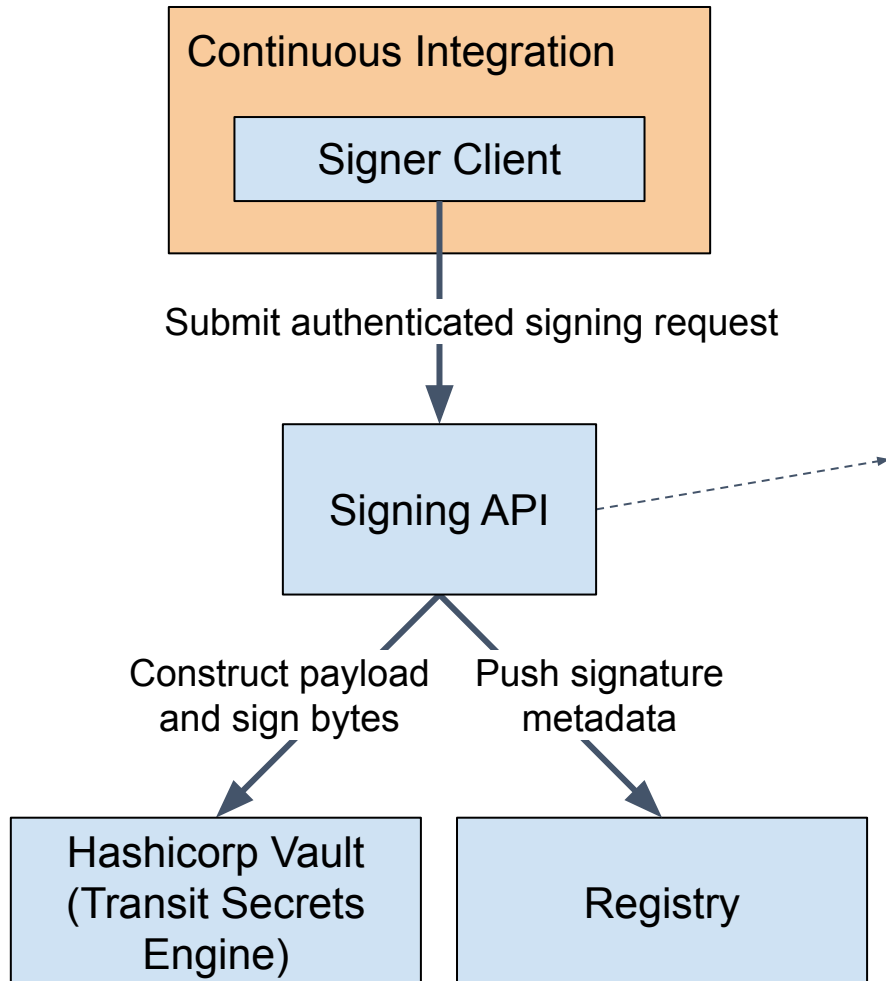
Hashicorp Vault audit logs

```
data {
  input      hmac-sha256:18c77d3538c17171ff34a4d6e5f63e93ed5874307fd302922fc3dd3d5d8dbe2b
  key_version hmac-sha256:5b3469b7db63fb80eff7656a94609f5af41fe5a24290081a1326e0e30e4078ac
}
id          bbe9faa7-0196-e989-940b-8636d9755b0e
mount_accessor transit_b71e90b6
mount_type   transit
namespace {
  id root
}
operation   update
path        crypto/k8s/sign/k8s_image-integrity_image-integrity-signer_artifact_signing_1
```

Audit logs are opaque and hard to interpret.

Signing Service: Least Privilege & Auditability

Mediating Vault & Registry access with the API:



Signing API audit logs

```
ci_commit_branch  main
ci_commit_sha     f108225984de153086c77b2f95ff7cd5e71bc3a6
ci_job_id         246064299
ci_pipeline_id    14429574
ci_project_path   DataDog/image-integrity-testing


■ parsed_artifact {
  digest          sha256:01c1e3796cc29194f84f0ebf3b54a0e0d36f1178ec20a537b6031c46aa9b9bdf
  is_digested     true
  is_tagged       false
  registry        registry.ddbuild.io
  repository      image-integrity-test/hello-world
  valid           true
}
```

Audit logs have full request context.
Minimal trusted compute base for direct Vault access.

Signing Service: Encapsulation


- Rarely need to update signing clients in all CI jobs
- Central control of load against registry & key management provider (Vault)
- Central key management transparent to clients

Artifacts Signed: New vs. Duplicate



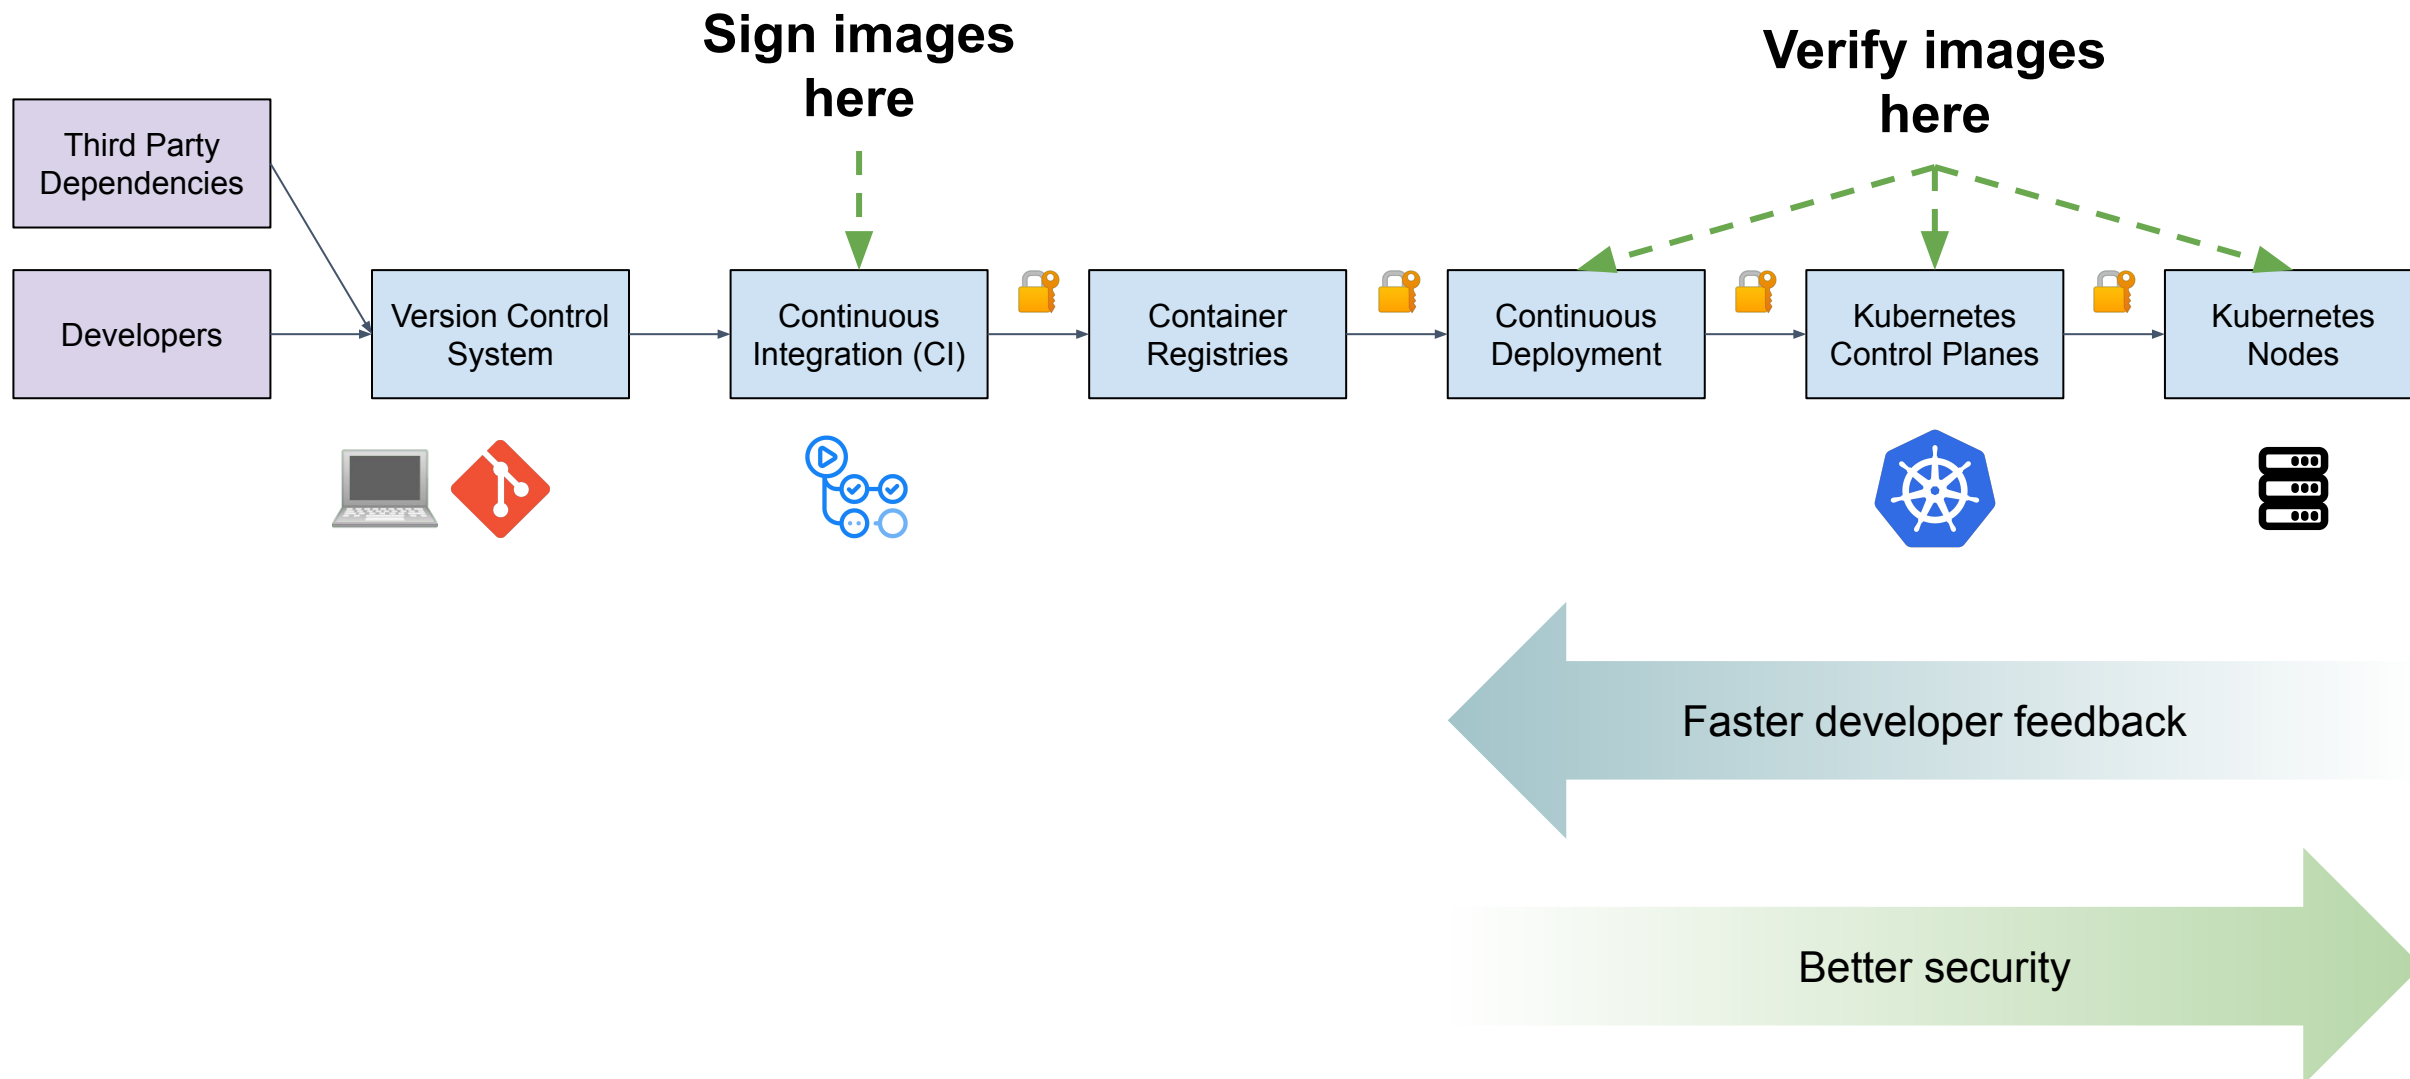
Reproducible builds → only ~3% of image signing requests are for new images.

SignArtifact Response Latency

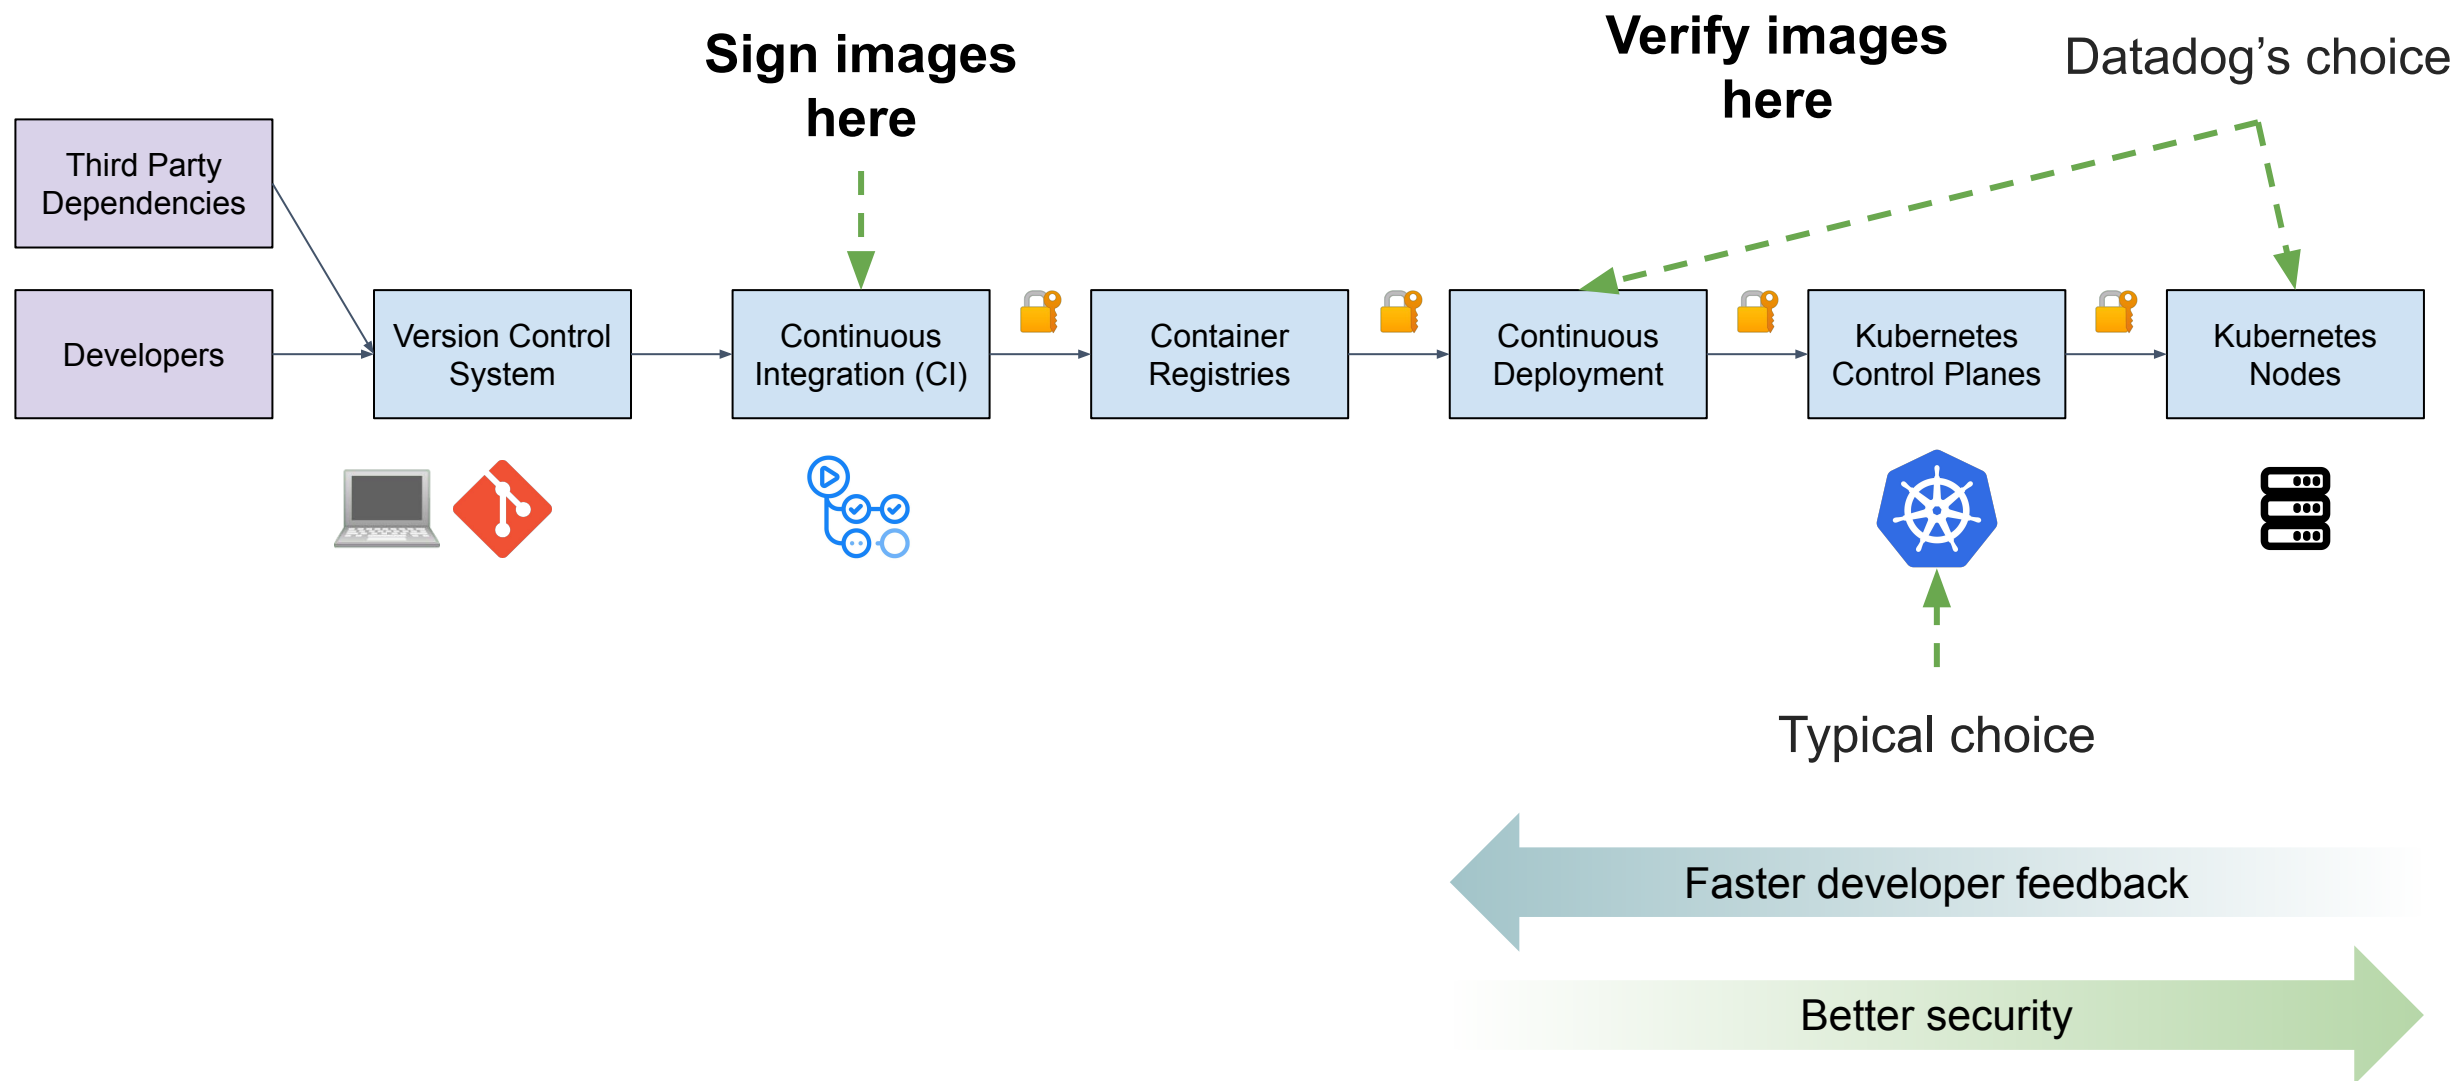


Deduplication deployed

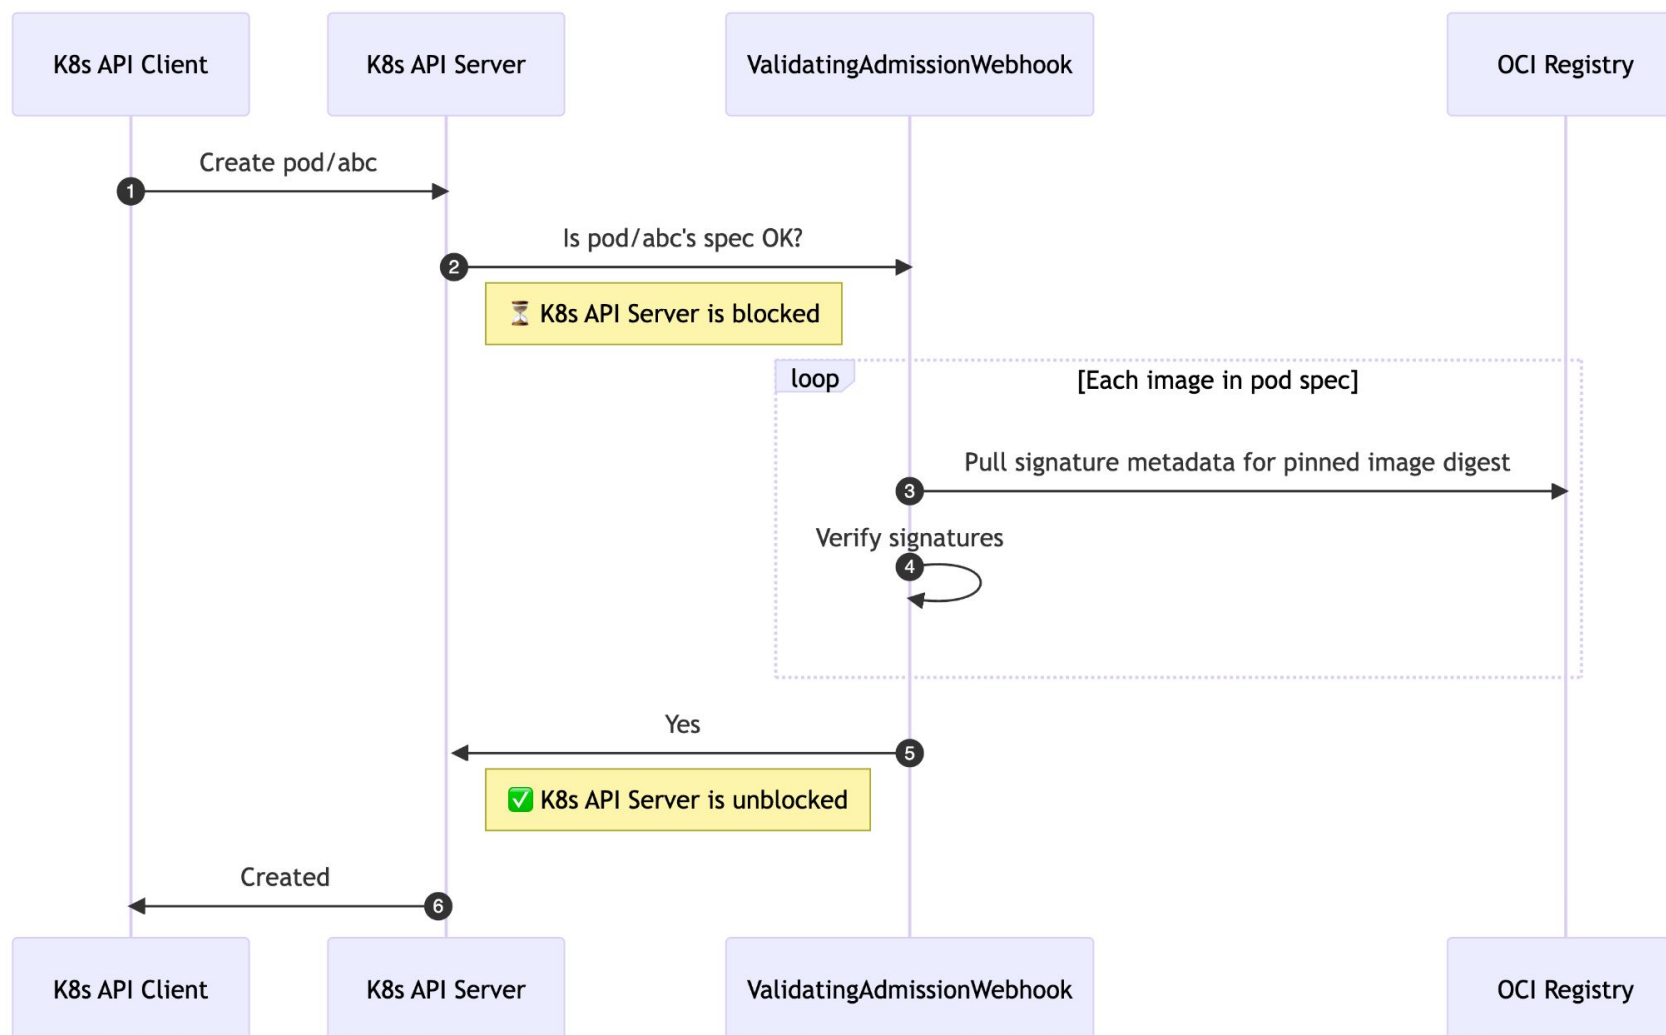
Where to Verify Signatures?



Where to Verify Signatures?



Validating Admission Webhooks



Disadvantages

Fetching registry metadata doesn't fit in webhook latency budget:

Goal: **p99 < 10ms**

Actual: **p50 > 200ms**

Registry would be in the control plane hot path (new dependency).

Workaround

Could use an ImagePolicyWebhook instead (for caching and retries), but relying too heavily on a central cache can lead to metastable failures^[1].

[1] <https://sigops.org/s/conferences/hotos/2021/papers/hotos21-s11-bronson.pdf>

Image Verification in containerd

Why verify in containerd?

- Ultimately responsible for resolving an image digest
- Closest to runtime we can verify image signatures
- Minimal trusted compute base after signature verification

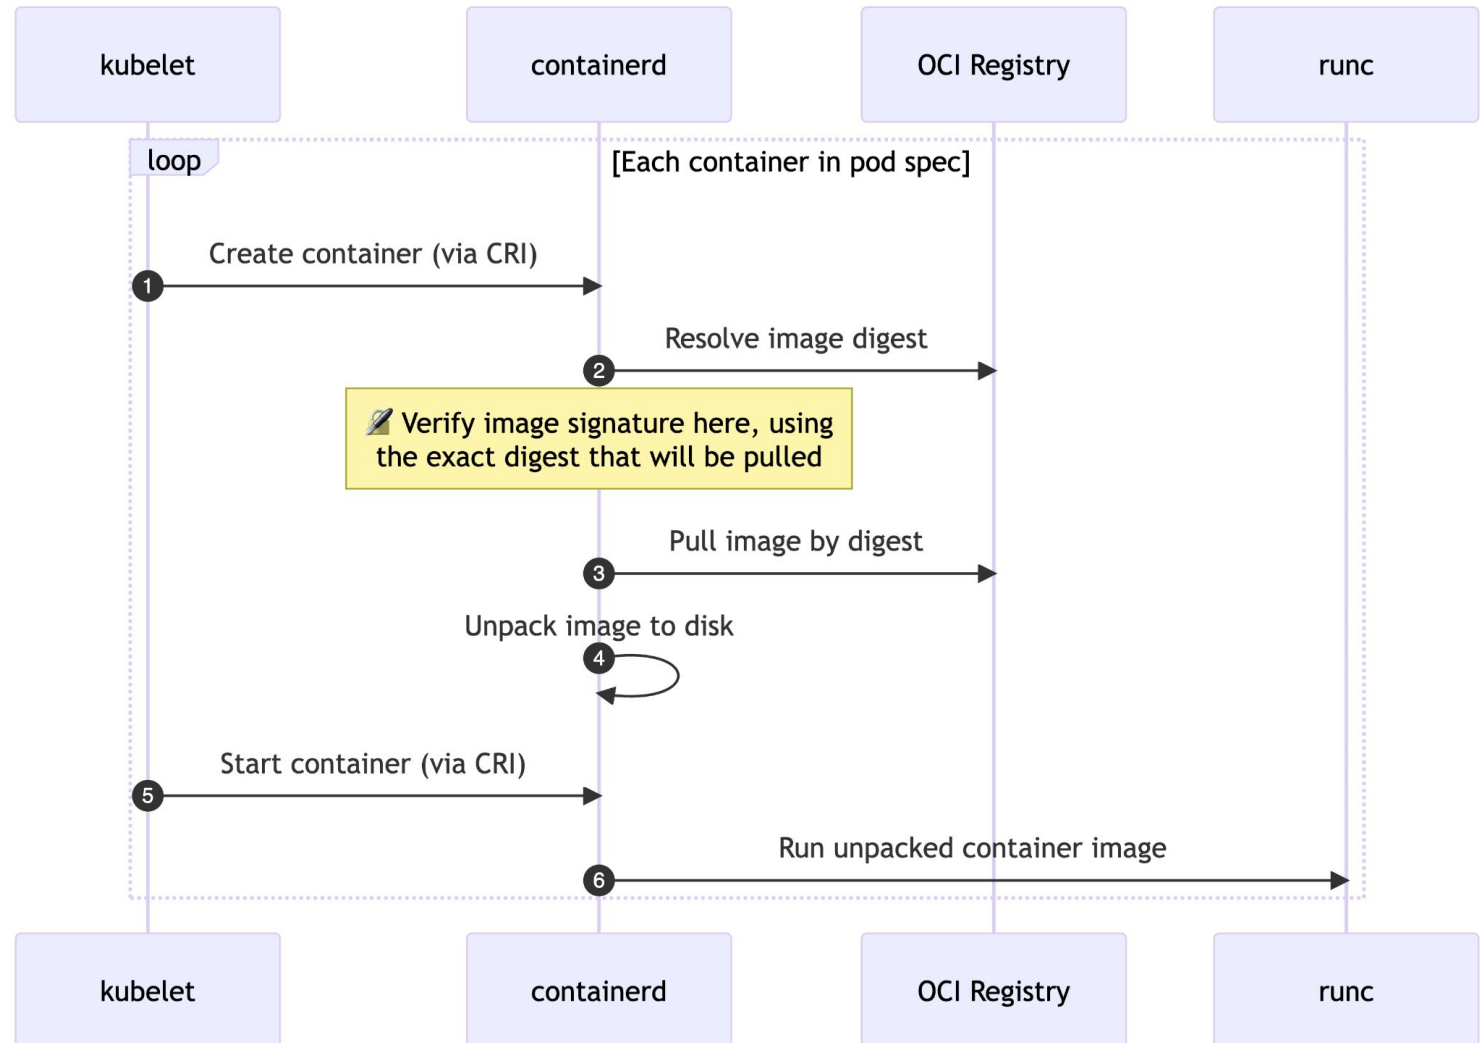
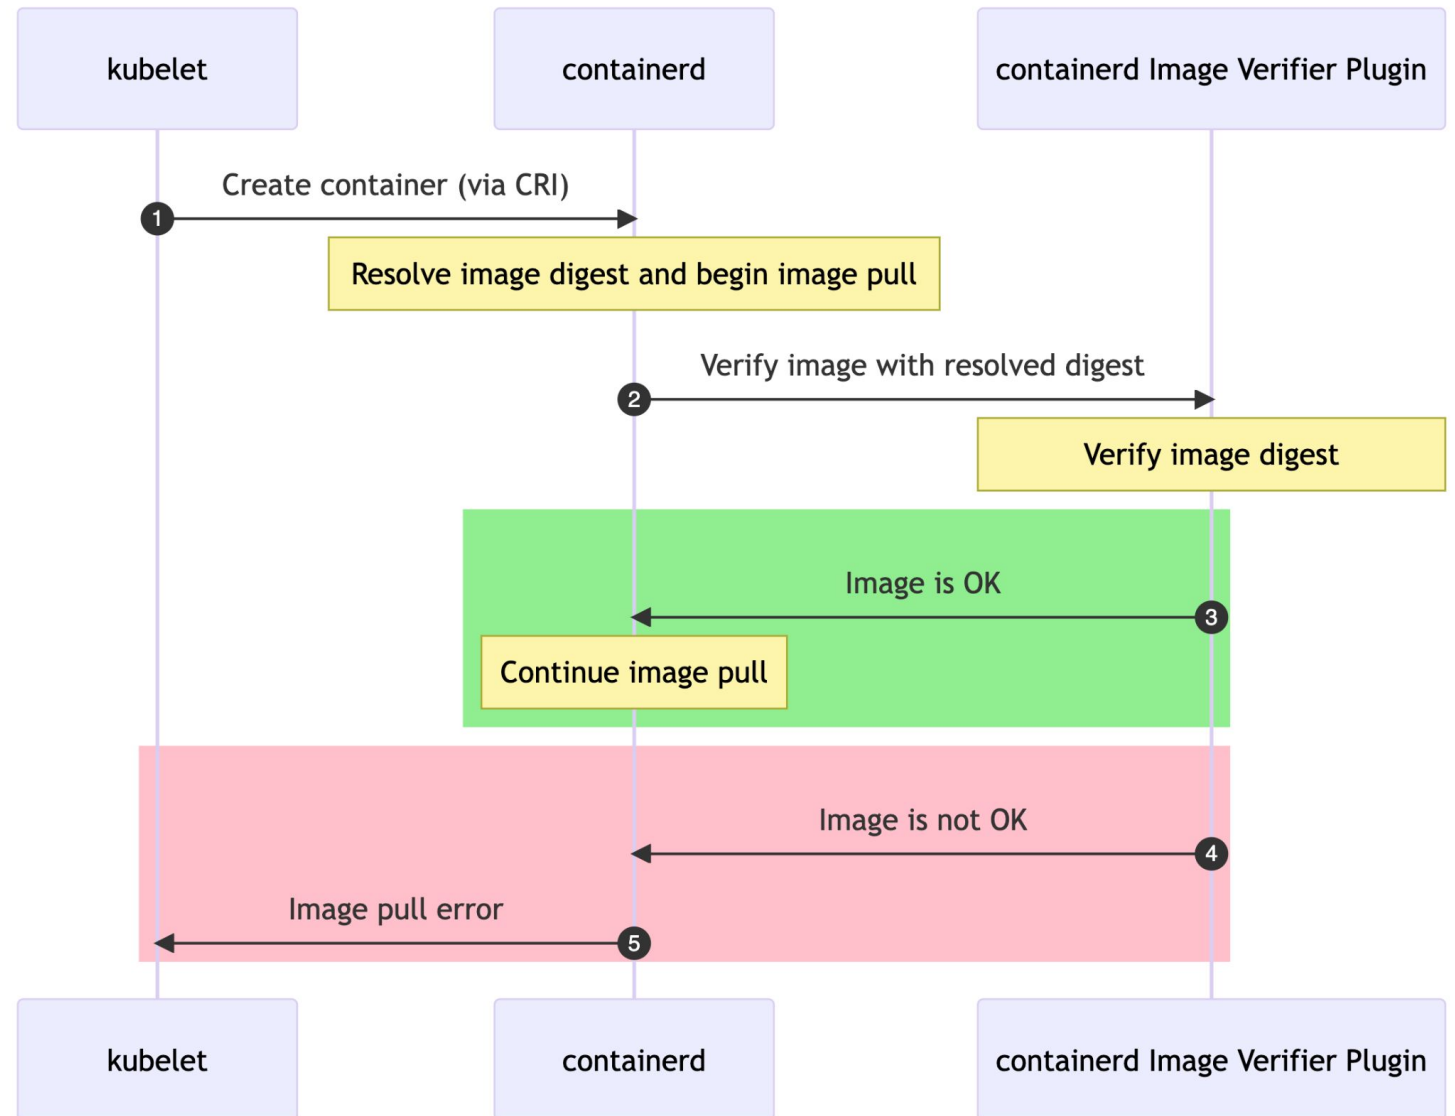


Image Verification in containerd

- Adding a plugin system for image verification
- Benefit from kubelet's reliability features for pulling images
- Running a temporary fork & working with maintainers to get the image verification plugin system into containerd 2.0. See [containerd#6691](https://github.com/containerd/containerd/pull/6691).



Developer Perspective

```
$ kubectl get pods
```


NAME	READY	STATUS	RESTARTS	AGE
example-pod-images-signed	1/1	Running	0	50s
example-pod-images-unsigned	0/1	ImagePullBackOff	0	48s

```
$ kubectl get events
```

```
...
2m Warning Failed pod/example-pod-images-unsigned Failed to pull image ... image verifier blocked
pull of registry.example.com/example-signed with digest
sha256:43292eabb01bf6c9dacf41b249248bc445e16698a91358f4dda6ef67629818f2 for reason: ...
(Need help? https://datadoghq.atlassian.net/l/cp/Fn2mj3vS)
...
```

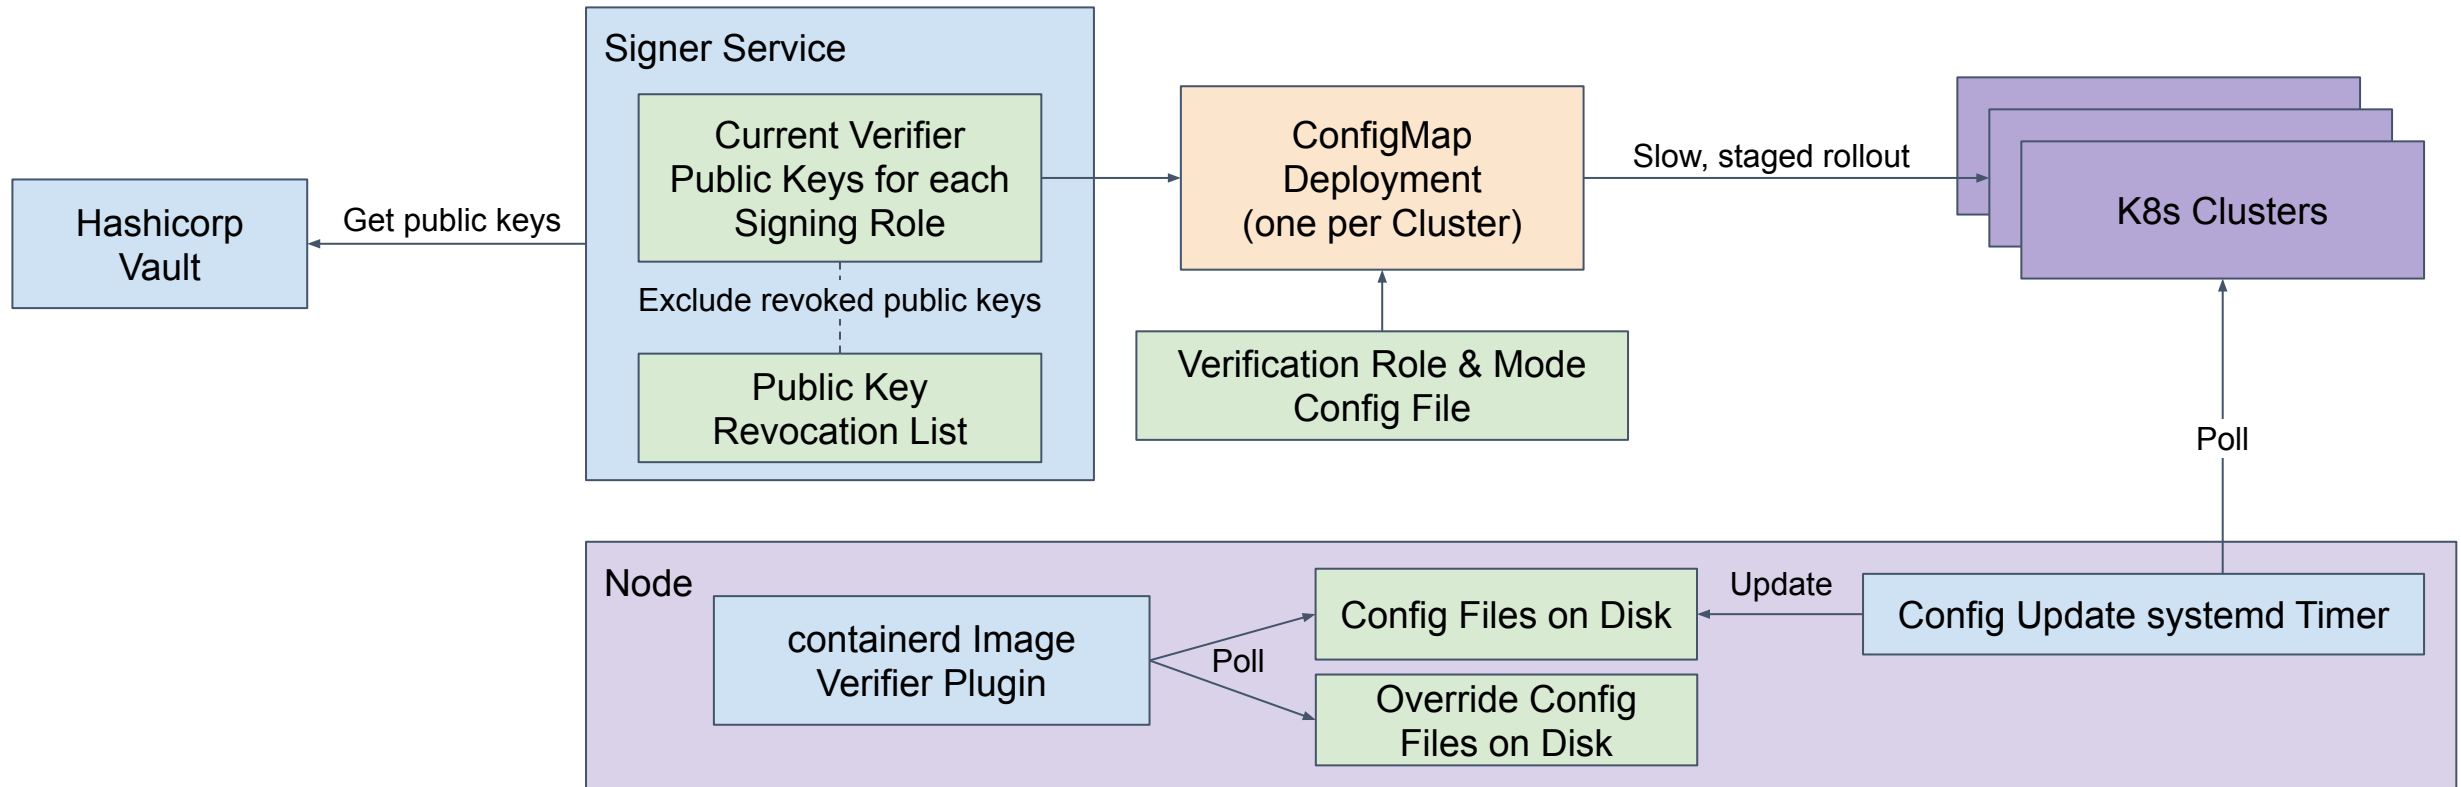
Distributing Verifier Config

- The verifier plugin on each node needs:
 - Public keys
 - Verification mode (audit, block, or disabled)
 - Image digest revocation list
- Requirements:
 - No new node dependencies
 - Slow, staged rollout of config
 - Multiple fallback mechanisms



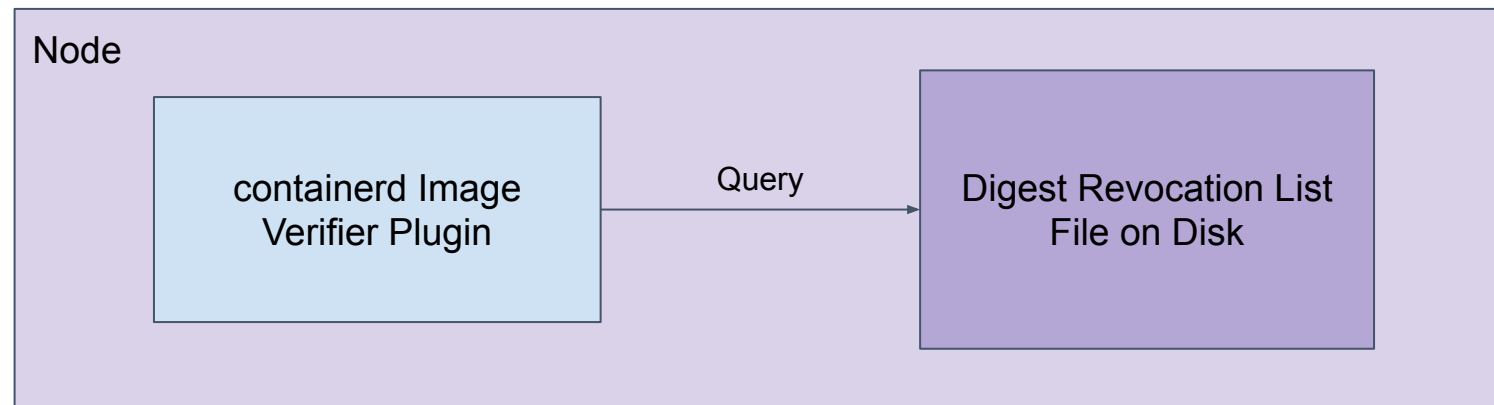
Distributing Public Keys & Mode

- Layered approach:
 - Defaults baked into node image
 - Periodic dynamic updates pulled from a ConfigMap in each cluster
 - Override config layer on disk



Distributing Image Revocation List

- Baked into the node image
- Updated via node image rollout automation
- Dynamic updates wouldn't be useful since containerd caches images
 - Prioritize draining nodes that have run a revoked image



Challenges & Recommendations

- Can't easily configure verification mode by Kubernetes namespaces
 - Harder to turn on blocking mode in large multi-tenant clusters
 - Dedicated clusters for sensitive workloads → easier to sign all images and enable blocking

Challenges & Recommendations

- Integrating image signing into many CI configurations is a lot of work, even with a thin client
 - Monorepos or consistent build tooling (e.g. Bazel) make it easier
 - Roll out audit mode verification before you're done rolling out signing

Challenges & Recommendations

- Node-level image verification is uncharted territory
 - We believe the reliability properties are worth the extra effort

Takeaways

1. Evaluate image signing as a service for security and scalability.
2. Think about the reliability risks of Kubernetes admission webhooks and the advantages of node-level image verification.
3. Join the conversation and give feedback on the design at <https://github.com/containerd/containerd/issues/6691>



KubeCon



CloudNativeCon

Europe 2023

ethan.lowman@datadoghq.com

