



SIG-Auth Deep Dive

David Eads (Red Hat), Jordan Liggitt (Google), Rita Zhang (Microsoft)


Code of Conduct

Remember the **Golden Rule**: Treat others as you would want to be treated - with kindness and respect

Scan the QR code to access and review the **CNCF Code of Conduct**:



Virtual Audience Closed Captioning

Closed captioning for the virtual audience is available during each session through [Wordly](#).  The [Wordly](#) functionality can be found under the “Translation” tab on the session page.

[Wordly](#) will default to English. If another language is needed, simply click the dropdown at the bottom of the “Translation” tab and choose from one of 26+ languages available so you don’t miss a beat from our presenters.

*Note: Closed captioning is ONLY available during the scheduled live sessions and will not be available for the recordings on-demand within the virtual conference platform.

- Virtual attendees may submit questions to speakers through the CNCF Slack channel: **#2-Kubecon-sessions**
- Please create a thread and tag the speaker(s) with questions about their talk.
- Questions will be answered by the speaker and/or other community members after the session concludes.

**Thank you to our Session Recording
Sponsor:**





KubeCon



CloudNativeCon

Europe 2023

Enjoy the Session!





KubeCon



CloudNativeCon

Europe 2023

SIG-Auth Deep Dive

David Eads, Red Hat
Jordan Liggitt, Google
Rita Zhang, Microsoft



What do we do?

SIG-Auth is responsible for features in Kubernetes that control and **protect access** to the API and other core components. This includes **authentication** and **authorization**, but also encompasses features like **auditing** and some **security policy**.

<https://github.com/kubernetes/community/blob/master/sig-auth/charter.md>

Sub-Projects

- Audit Logging
- Authenticators
- Authorizers
- Certificates
- Encryption at rest
- Multi Tenancy
- Node Identity and Isolation
- Policy Management
- Service Accounts
- Secrets Store CSI Driver

Stable

- [2799](#) Reduce legacy service account token attack surface area
 - Stop auto-generating legacy tokens (beta v1.24, stable v1.26)

Implementable

- [2799](#) Reduce legacy service account token attack surface area
 - Track use of legacy tokens (beta v1.27, targeting stable v1.28)
 - Clean up unused legacy tokens (targeting alpha v1.28)
- [3325](#) API to get current user attributes, `kubectl whoami` (beta v1.27, targeting stable v1.28)
- [3299](#) KMS v2 encryption at rest (beta v1.27, targeting stable v1.29)
- [3257](#) Cluster Trust Bundles (API alpha in v1.27, targeting volume mount alpha in v1.28)

Provisional

- [3221](#) Structured Authorization Configuration (in design, targeting alpha v1.28)
- [3331](#) Structured OIDC Configuration (in design, targeting alpha v1.28)
- [3766](#) ReferenceGrant (in design)
- [2718](#) Client Exec Proxy (in design)




KubeCon



CloudNativeCon

Europe 2023

Graduating Enhancements



Legacy Token Reduction

KEP-2799

Move away from long lived secret based SA tokens

Use ephemeral tokens via the token request API instead

- In v1.24+, secret based tokens default to not being auto-generated
(**LegacyServiceAccountTokenNoAutoGeneration**)
- In v1.26, feature is stable




KubeCon



CloudNativeCon

Europe 2023

Implementable Enhancements



Legacy Token Reduction

KEP-2799

Move away from long lived secret based SA tokens

Use ephemeral tokens via the token request API instead

- v1.26: **LegacyServiceAccountTokenTracking** alpha feature gate
 - Use of auto-generated secret based tokens emits warnings
 - Use of secret based tokens labels the secret with the last used date
- v1.27: **LegacyServiceAccountTokenTracking** feature is beta
- v1.28: Targeting alpha of cleanup of unused auto-generated token secrets

SelfSubjectReview

[KEP-3325](#) - Beta in v1.27, targeting stable in v1.28

```
kubectl auth whoami -o json
```

```
{
  "kind": "SelfSubjectReview",
  "apiVersion": "authentication.k8s.io/v1beta1",
  "metadata": {
    "creationTimestamp": "2023-04-21T17:05:55Z"
  },
  "status": {
    "userInfo": {
      "username": "alice",
      "groups": [
        "system:authenticated"
      ]
    }
  }
}
```

```
--runtime-config=authentication.k8s.io/v1beta1=true
--feature-gates=APISelfSubjectReview=true
```

Encryption at Rest Improvements - v1 & v2

- Change algorithm for KMS data encryption to AES-GCM. No user action required.
 - v1.25: KMS v1 write using AES-GCM, read AES-GCM with fallback to AES-CBC
 - KMS v2 only allows AES-GCM

- Dynamic reload of *EncryptionConfiguration* file - v1.26
 - Does not require a kube-apiserver restart

- Custom resource encryption - v1.26
 - Add custom resources to *EncryptionConfiguration*

- Encrypt all resources - v1.27
 - `*.*` to encrypt all resources in all groups
 - `*.<group>` to encrypt all resources in one group

```
kind: EncryptionConfiguration
apiVersion: apiserver.config.k8s.io/v1
resources:
```

```
- resources:
  - secrets
```

```
providers: ...
```

```
- resources:
  - pandas.awesome.bears.com
```



```
providers: ...
```

```
- resources:
  - "*.*"
```

```
providers: ...
```


KMS v1 Gaps

- **Performance**

- a new Data Encryption Key (DEK) is generated for each encryption
- slow cluster startup times due to the large number of requests made to the remote vault
- rate limits on external KMS
- 160ms per request

- **Key Rotation**

- manual and error-prone
- hard to determine what keys are in-use

- **Health check & status**

- kube-apiserver has to make encrypt and decrypt calls to determine KMS plugin health

- **Observability**

- unable to correlate events across kube-apiserver, KMS plugin, and external KMS

What's new in KMS v2?

[KEP-3299](#) - Alpha v1.25, v1.26, Beta v1.27, Stable targeted v1.29

- Benefits
 - Performance - [DEK Re-use in API Server](#) (80μs per request)
 - Health check & status - new [status API](#) health of the KMS plugin
 - Observability - [new UID](#) generated for each envelope operation.
 - New proto format for stored data in etcd
 - Key Rotation - [key_id from status API is used to track current KMS key and KEK rotation](#) allowing rotation without API server restart
- kubernetes.io/docs/tasks/administer-cluster/kms-provider
- Get involved [#sig-auth-kms-dev](#) on slack

New proto format for stored data in etcd

For storage, a new structured protobuf format.

The prefix for the new format is ``k8s:enc:kms:v2:<config name>:``.

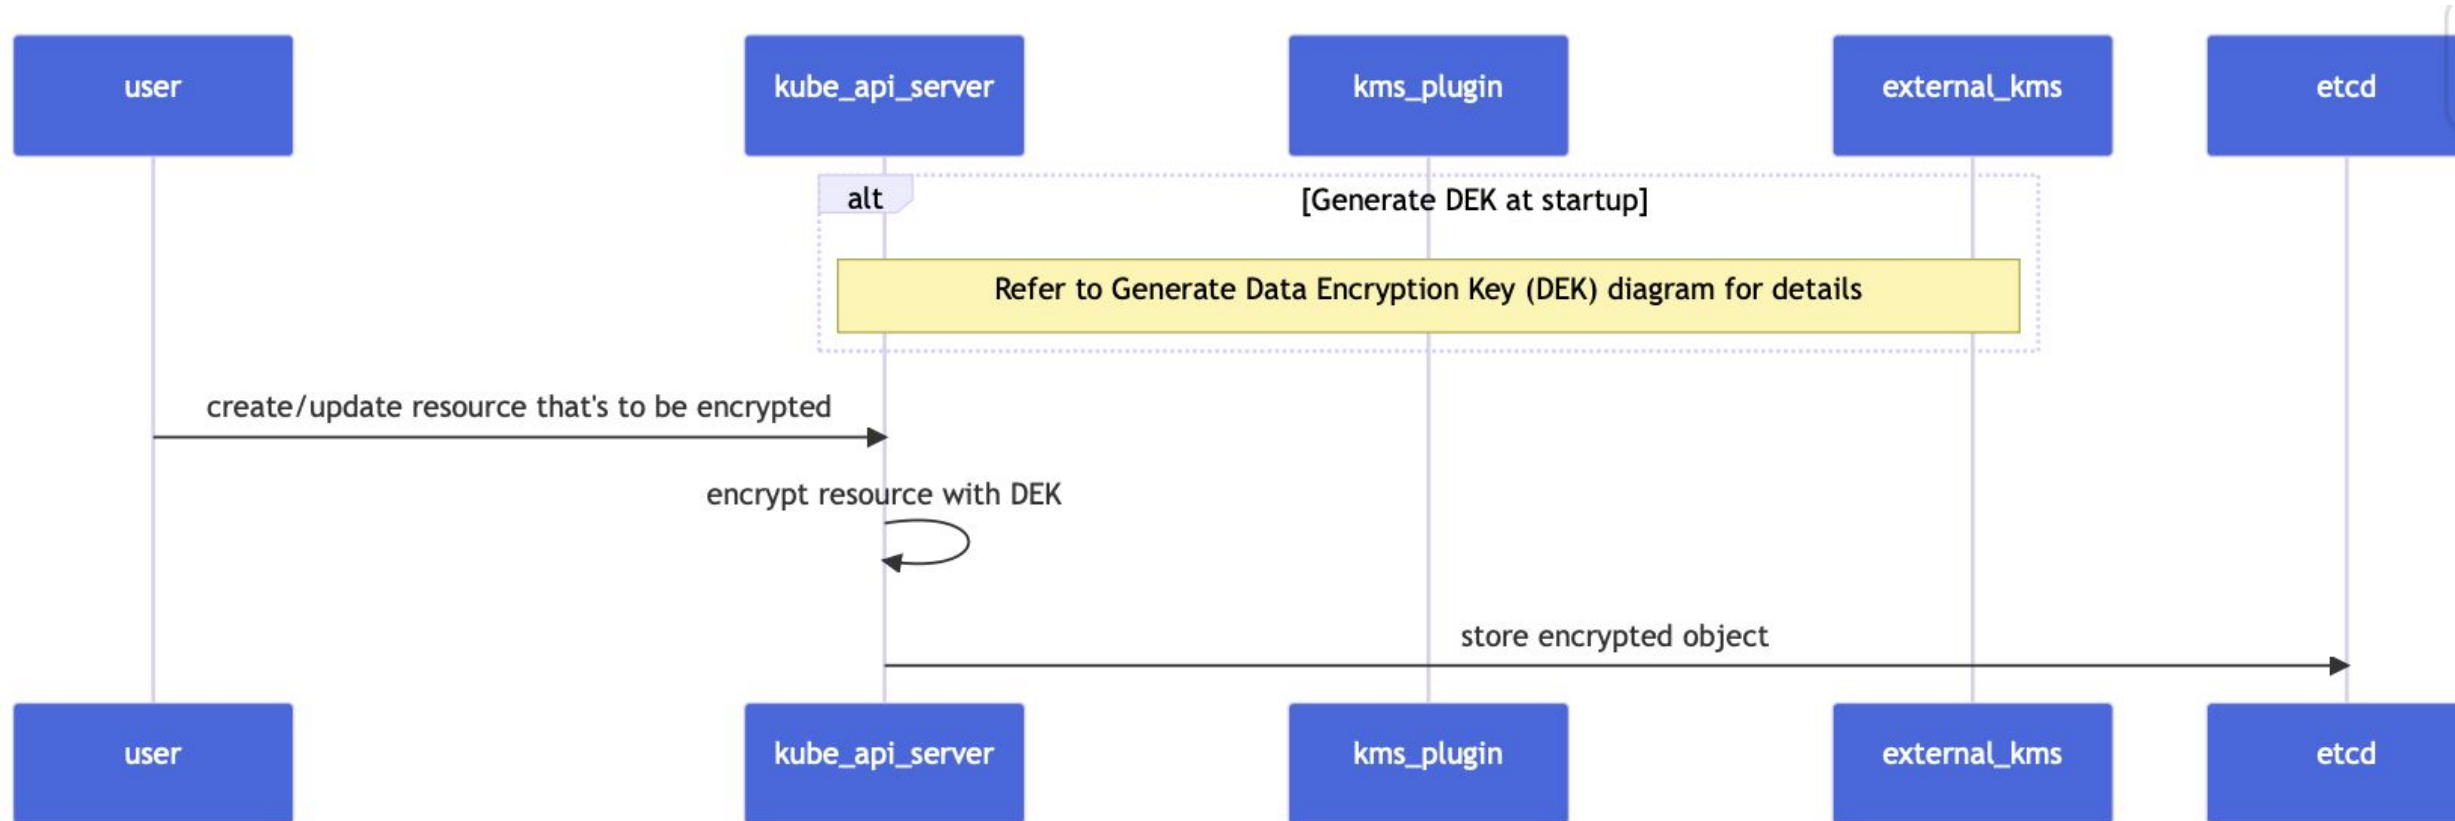
```
// EncryptedObject is the representation of data stored in etcd after envelope encryption.
type EncryptedObject struct {
    // EncryptedData is the encrypted data.
    EncryptedData []byte `protobuf:"bytes,1,opt,name=encryptedData,proto3" `

    // KeyID is the KMS key ID used for encryption operations.
    KeyID string `protobuf:"bytes,2,opt,name=keyID,proto3" `

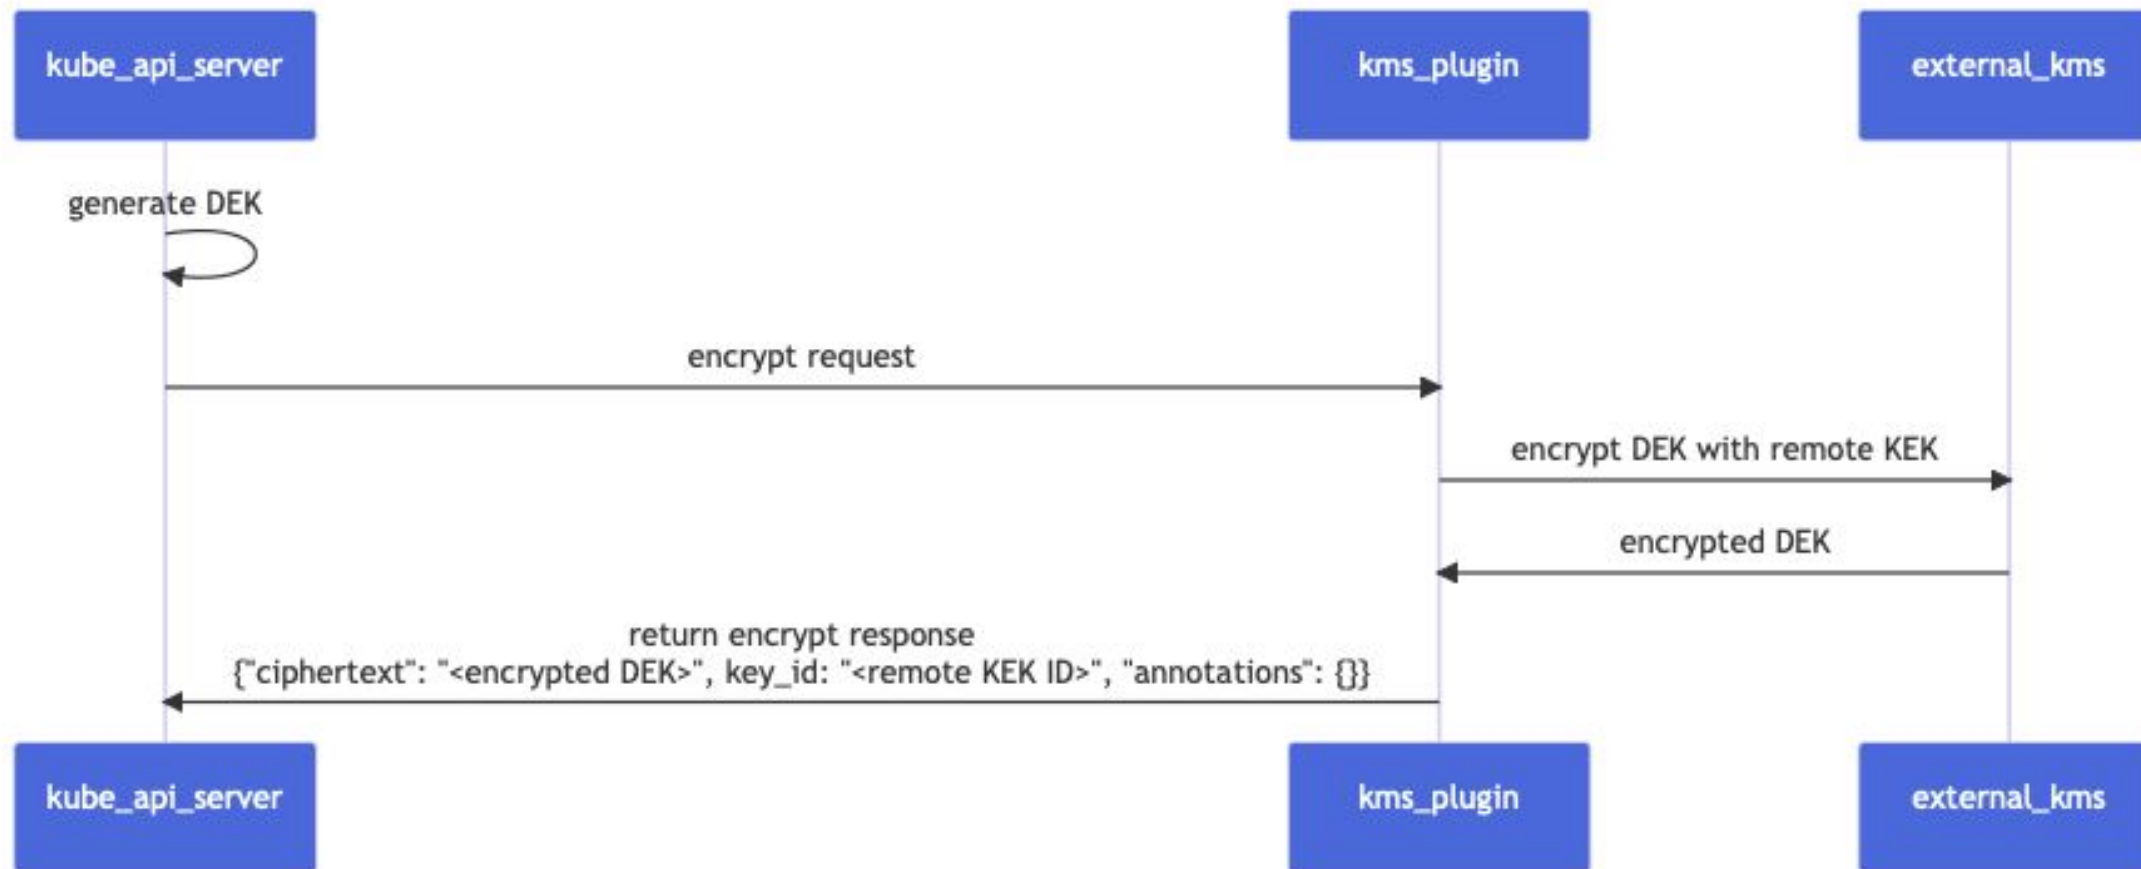
    // EncryptedDEK is the encrypted DEK.
    EncryptedDEK []byte `protobuf:"bytes,3,opt,name=encryptedDEK,proto3" `

    // Annotations is additional metadata that was provided by the KMS plugin.
    Annotations map[string][]byte `protobuf:"name=annotations" `
}
```

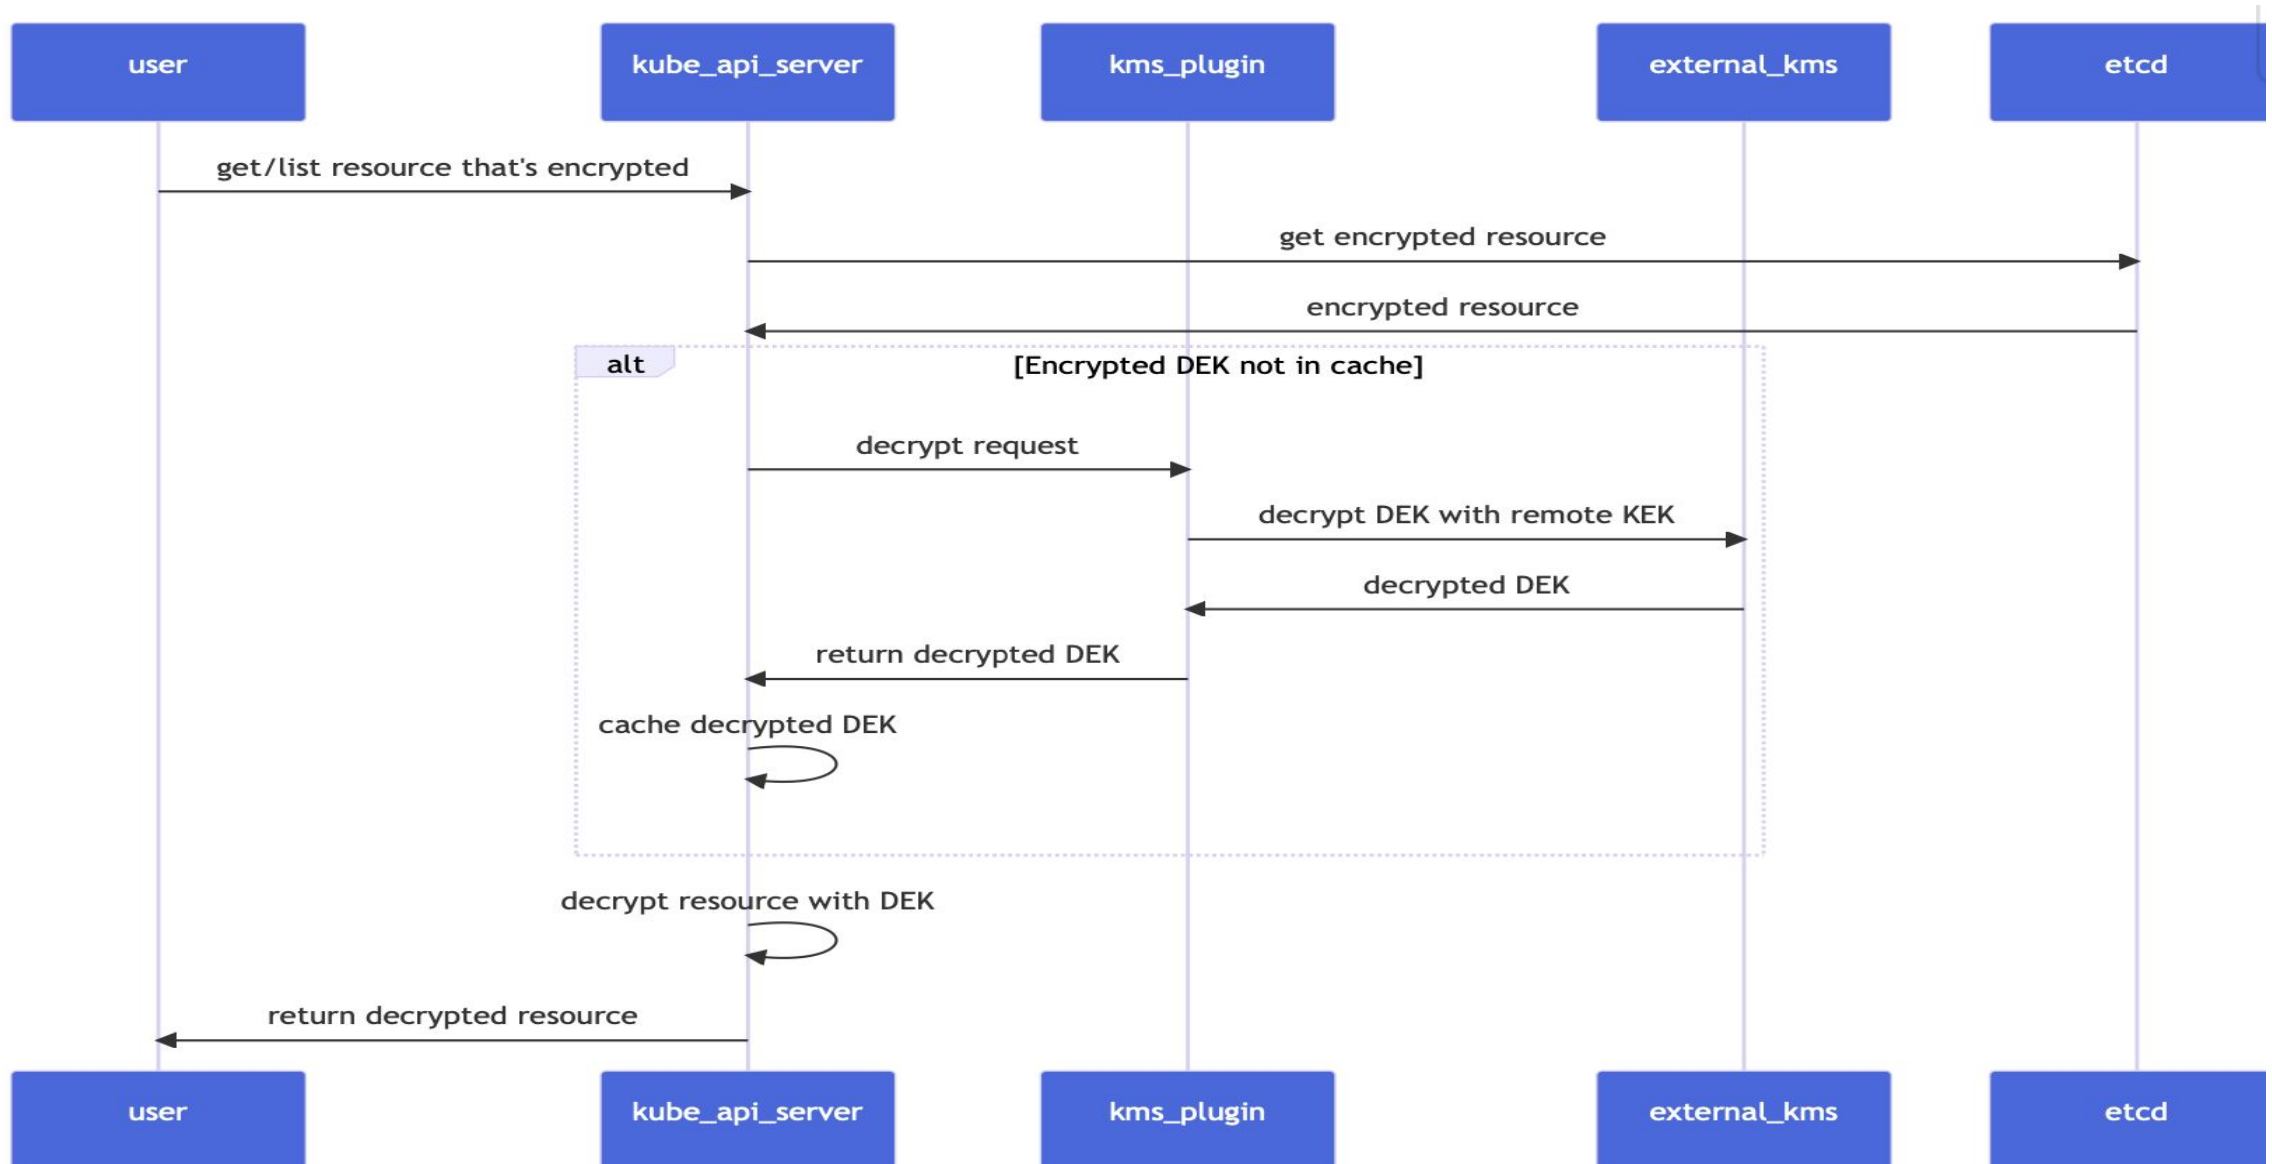
KMS v2 - Encrypt Request



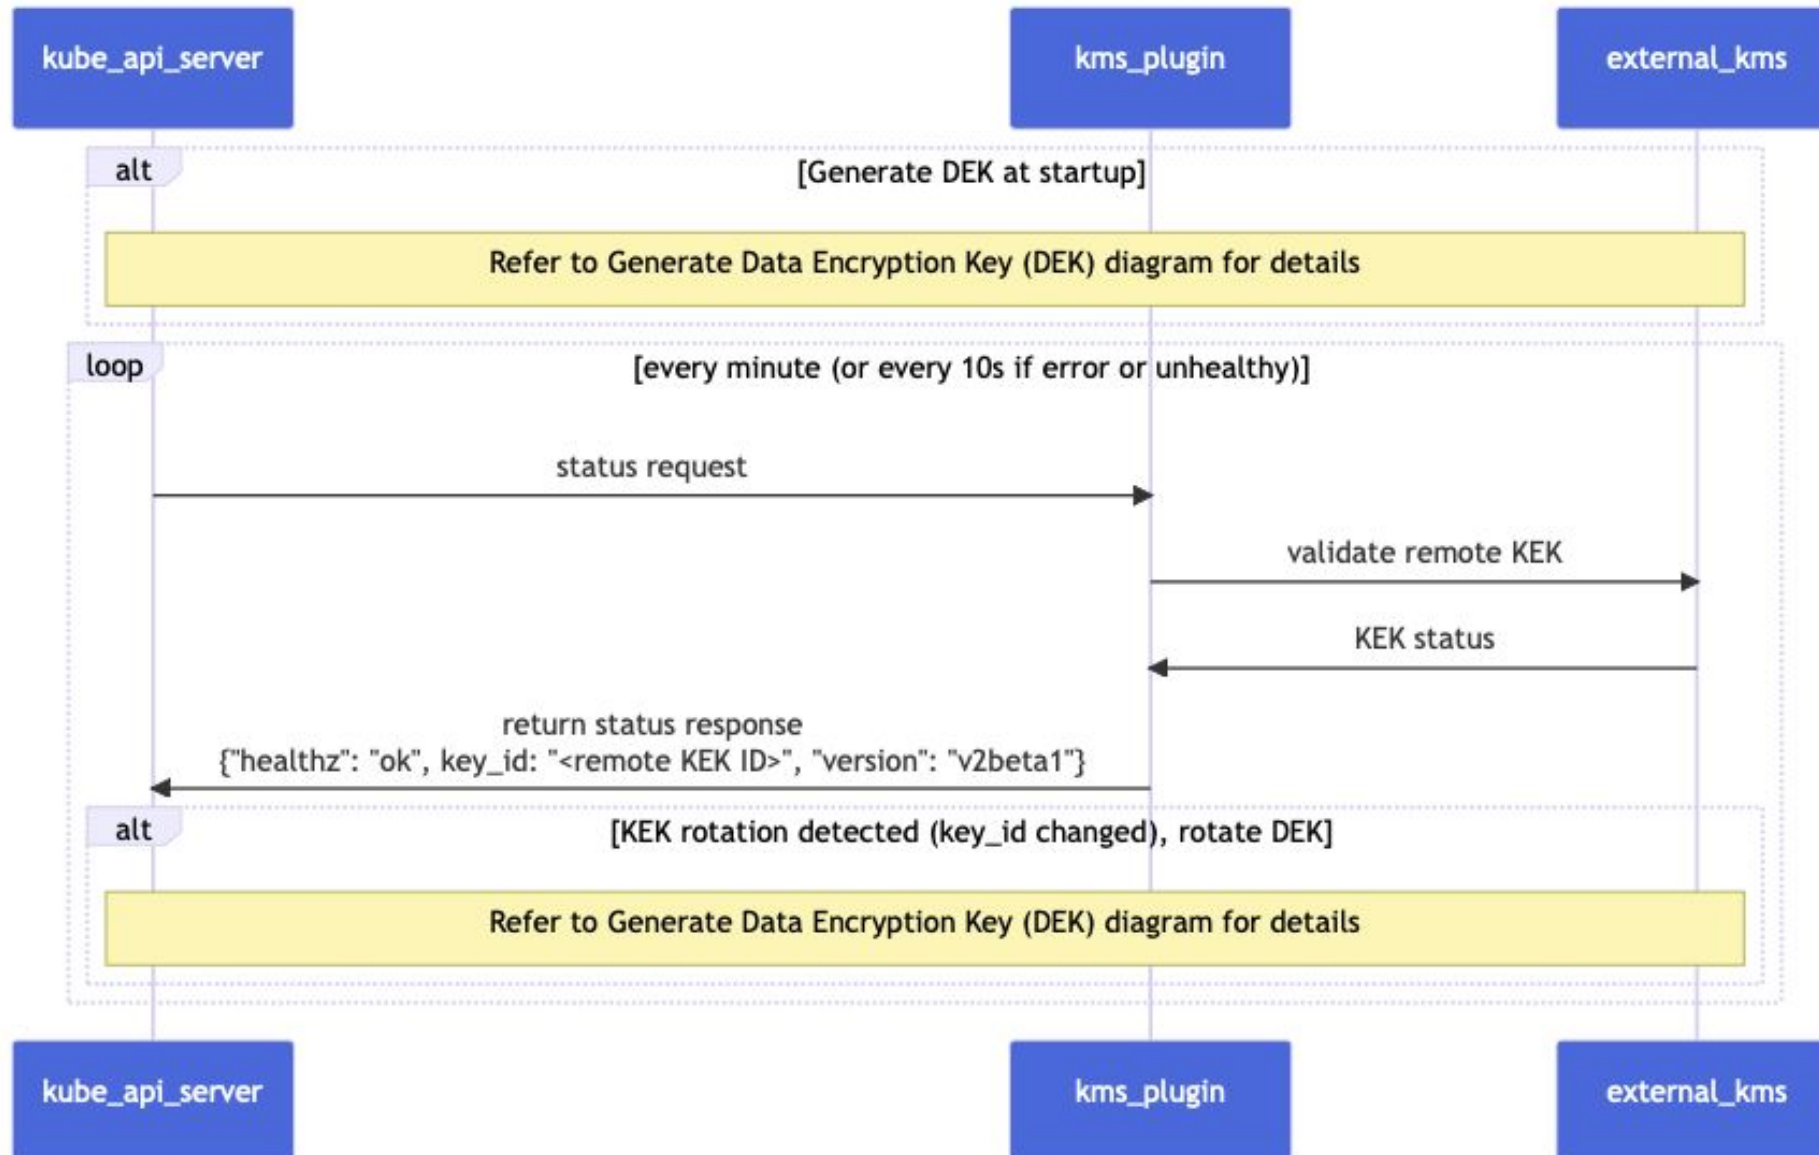
KMS v2 - Generate DEK



KMS v2 - Decrypt Request



KMS v2 - Status Request



Cluster Trust Bundles

KEP-3257

```
type ClusterTrustBundleSpec struct {  
    // The name of the associated signer.  
    // +optional  
    SignerName string `json:"signerName,omitEmpty"`  
  
    // The individual trust anchors for this bundle.  
    // A PEM bundle of PEM-wrapped,  
    // DER-formatted X.509 certificate.  
    // The order of certificates has no meaning.  
    PEMTrustAnchors string `json:"pemTrustAnchors"`  
}
```

DRAFT

```
apiVersion: v1  
kind: Pod  
metadata:  
  namespace: client  
  name: client  
spec:  
  containers:  
    name: main  
    image: my-image  
    volumeMounts:  
      - mountPath: /var/run/example-com-server-tls-trust-anchors  
        name: example-com-server-tls-trust-anchors  
        readOnly: true  
  volumes:  
    - name: example-com-server-tls-trust-anchors  
      projected:  
        sources:  
          + - pemTrustAnchors:  
            +   signerName: example.com/server  
            +   path: ca_certificates.pem
```


Cluster Trust Bundles

Trust the kube-apiserver, find your CA bundle

```
apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
spec:
  signerName: example.com/foo
```

```
apiVersion: certificates.k8s.io/v1alpha1
kind: ClusterTrustBundle
metadata:
  name: example.com:foo:v12
spec:
  signerName: example.com/foo
  trustBundle: [content here]
```

Current

```
kubectl get clustertrustbundle
--field-selector=spec.signerName=example.com/foo
```

Coming Soon

```
volumes:
+ - name: example-com-server-tls-trust-anchors
+   projected:
+     sources:
+       - pemTrustAnchors:
+         signerName: example.com/foo
+         path: ca_certificates.pem
```

DRAFT

Future

- Projected volume for workload client certificates? [Pre-KEP proposal](#)




KubeCon



CloudNativeCon

Europe 2023

In-Design Enhancements



Structured Authorization Configuration

KEP-3221

This proposal would allow:

- Dynamic config reload
- Multiple webhooks
- Configurable failure policy
- CEL-based filtering

DRAFT

```
apiVersion: apiserver.config.k8s.io/v1alpha1
kind: AuthorizationConfiguration
authorizers:
- type: Webhook
  webhook:
    unauthorizedTTL: 30s
    timeout: 3s
    subjectAccessReviewVersion: v1
    onError: Deny
    kubeConfigFile: /kube-system-authz-webhook.yaml
    matchConditions:
      - expression: |
          request.resourceAttributes.namespace == 'kube-system'
      - expression: |
          !('system:serviceaccounts:kube-system' in request.user.groups)
- type: Node
- type: RBAC
- type: Webhook
  webhook:
    authorizedTTL: 5m
    unauthorizedTTL: 30s
    timeout: 3s
    subjectAccessReviewVersion: v1
    onError: NoOpinion
    kubeConfigFile: /authz-webhook.yaml
```

Structured OIDC Configuration

KEP-3331

This proposal would allow:

- Dynamic config reload
- Multiple OIDC providers
- CEL-based authentication validation rules
- CEL-based claim extraction for user/group attributes
- Support for non-OIDC ID JWT credentials

claimValidationRules:

- rule: `'claims.aud == "charmander" || claims.aud == "bulbasaur"'`
message: clients other than charmander or bulbasaur are not allowed
- rule: `'claims.roles.split(",").exists(r, r == "kubernetes-user")'`
message: only kubernetes-user group members can access the cluster
- rule: `'!claims.username.startsWith("system:")'`
message: system identities are not allowed

userValidationRules:

- rule: `'!user.username.startsWith("myidp.")'`
message: system identities are not allowed

claimMappings:

- username: `'claims.username + ":external-user"'`
- groups: `'claims.roles.split(",")'`
- uid: `'claims.sub'`
- extra:
 - key: `'"client_name"'`
expression: `'claims.aud'`

DRAFT

KEP-3766

- Expansion of narrow API used by sig-network, sig-storage
- Goals
 - Users express intent to allow use of a resource
 - Controllers verify that intent was granted
 - Admins avoid broad permission grants to controllers
- Example uses
 - Gateway cross-namespace TLS secrets
 - PersistentVolumeClaim cross-namespace data sources



KubeCon



CloudNativeCon

Europe 2023

Cross-SIG Work



Cross SIG Work

- CEL for Admission Control (api-machinery)
 - [kubernetes/enhancements#3488](#)
- Storage Version API (api-machinery)
 - [kubernetes/enhancements#2339](#)
- Kube API Server Identity (api-machinery)
 - [kubernetes/enhancements#1965](#)

Shout outs!

 THANK YOU! 

<p>Anish @aramase KMSv2</p>	<p>Krzysztof @ibihim KMSv2 / kube-rbac-proxy</p>	<p>Nilekh @nilekhc KMSv2</p>	<p>Maksim @nabokihms whoami</p>
<p>Taahir @ahmedtd Cluster Trust Bundle</p>	<p>Standa @stlaz kube-rbac-proxy</p>	<p>Roma @r-erema OIDC tests</p>	

Where can you find us?

Slack channel: [#sig-auth](#)

Home page: <https://github.com/kubernetes/community/tree/master/sig-auth>

Mailing list: <https://groups.google.com/forum/#!forum/kubernetes-sig-auth>

Bi-weekly meetings Wednesday at 11PT (agenda/recordings links on home page)