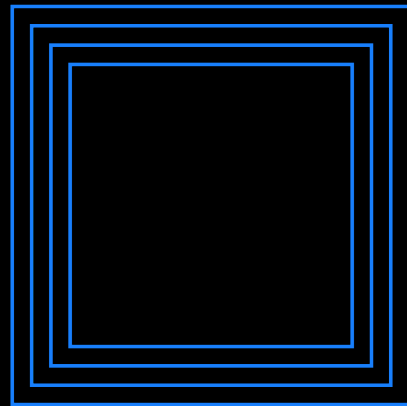
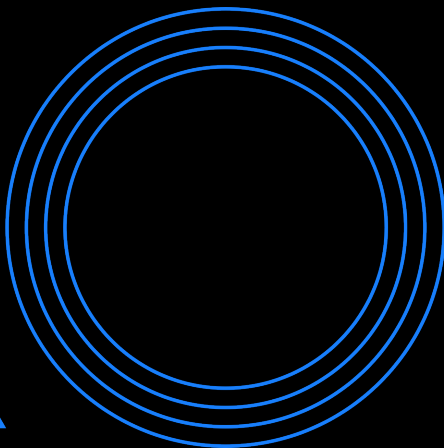
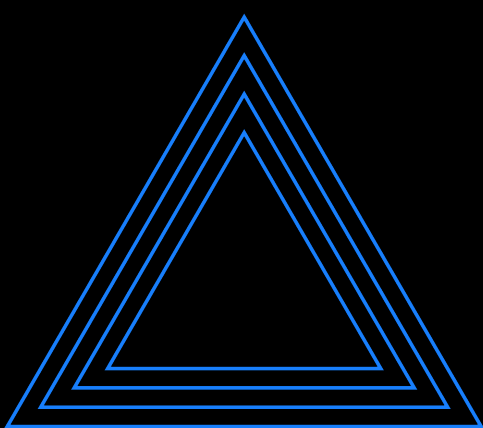




# PLAYSTATION AND KUBERNETES: HOW TO SOLVE A PROBLEM LIKE REAL-TIME



Joseph Irving



# LEVEL SELECT:

**TUTORIAL: REAL-TIME GAMSERVERS**

**LEVEL 1: RUNNING IN KUBERNETES**

**LEVEL 2: AUTOSCALING**

**LEVEL 3: MULTI-REGION**

# LEVEL SELECT:

**TUTORIAL: REAL-TIME GAMSERVERS**

**LEVEL 1: RUNNING IN KUBERNETES**

**LEVEL 2: AUTOSCALING**

**LEVEL 3: MULTI-REGION**

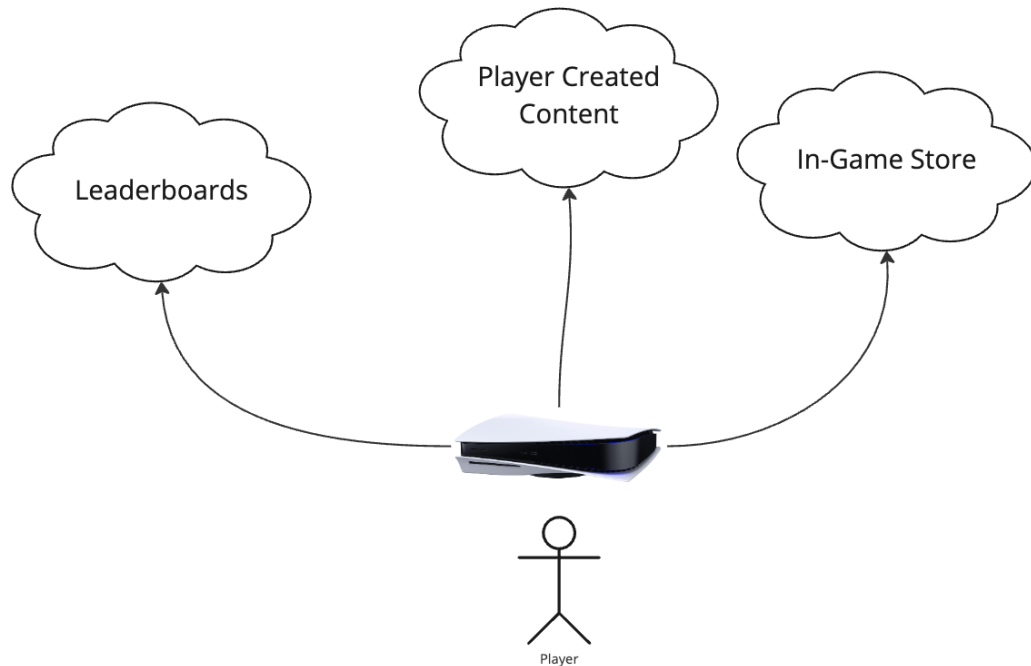
# What isn't Real-time?

Async

Request Based

HTTPS/GRPC

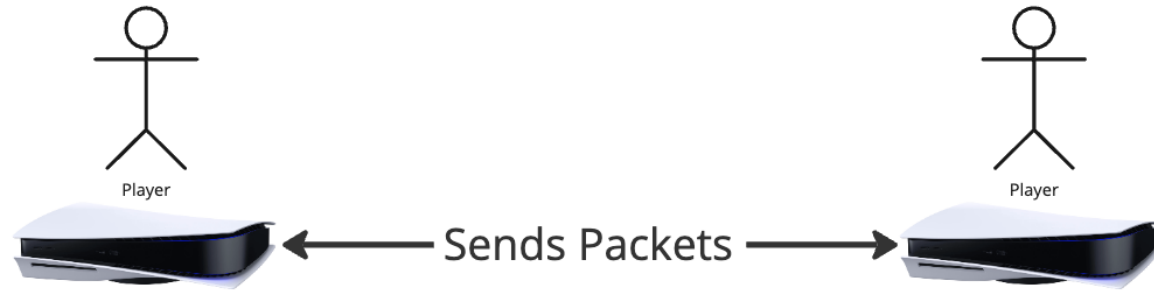
Not Latency Sensitive



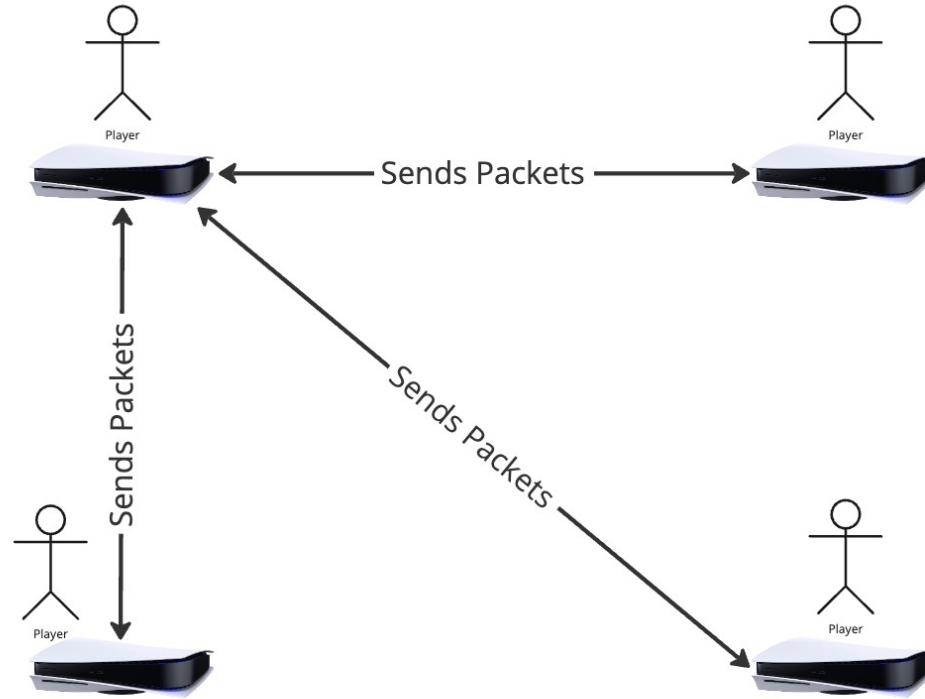
# Real-time



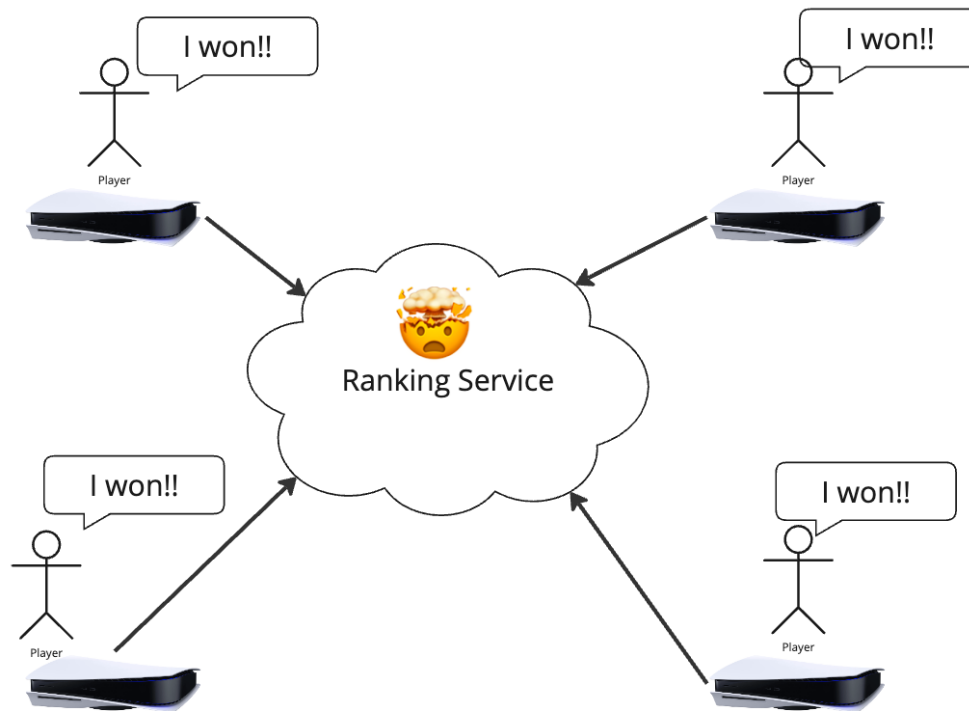
# Peer to Peer



# Peer to Peer

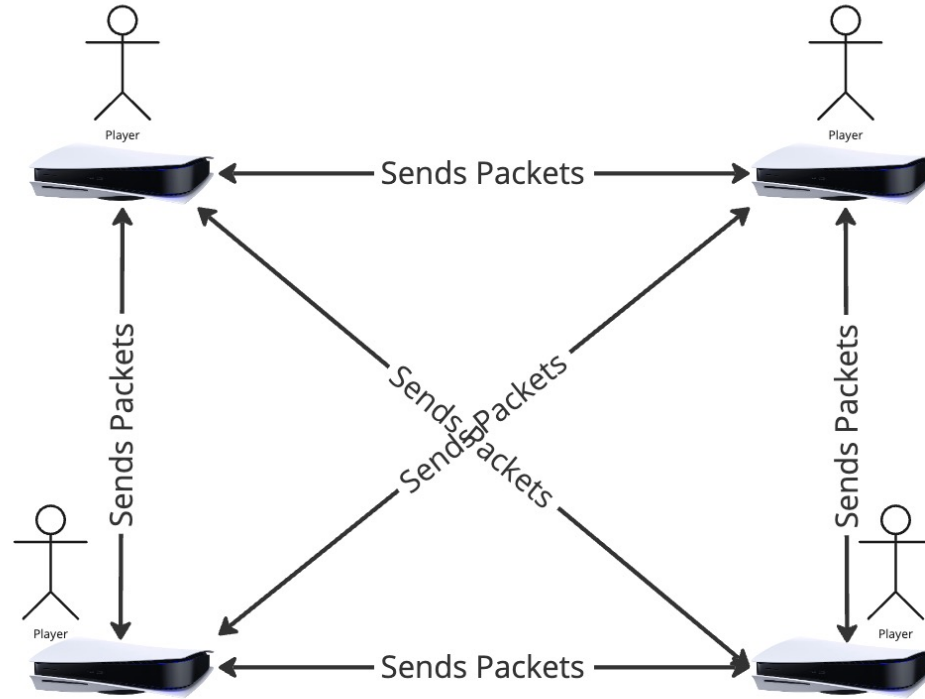


# Peer to Peer

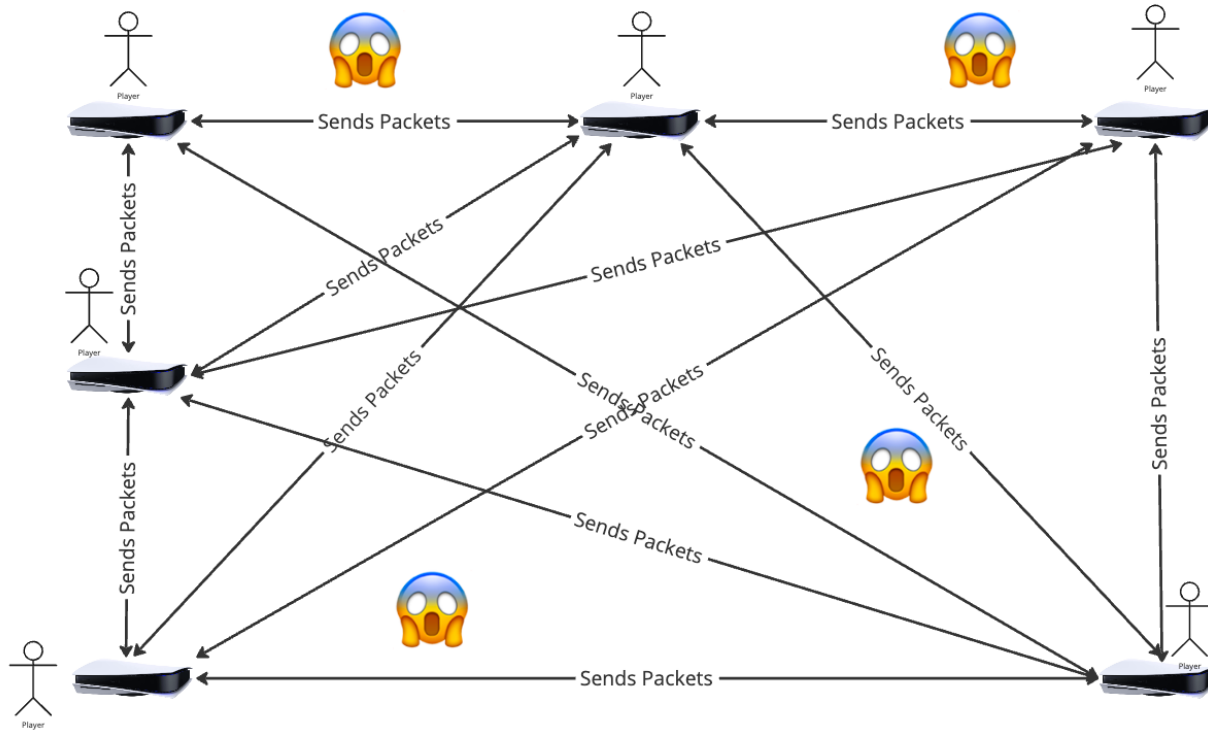




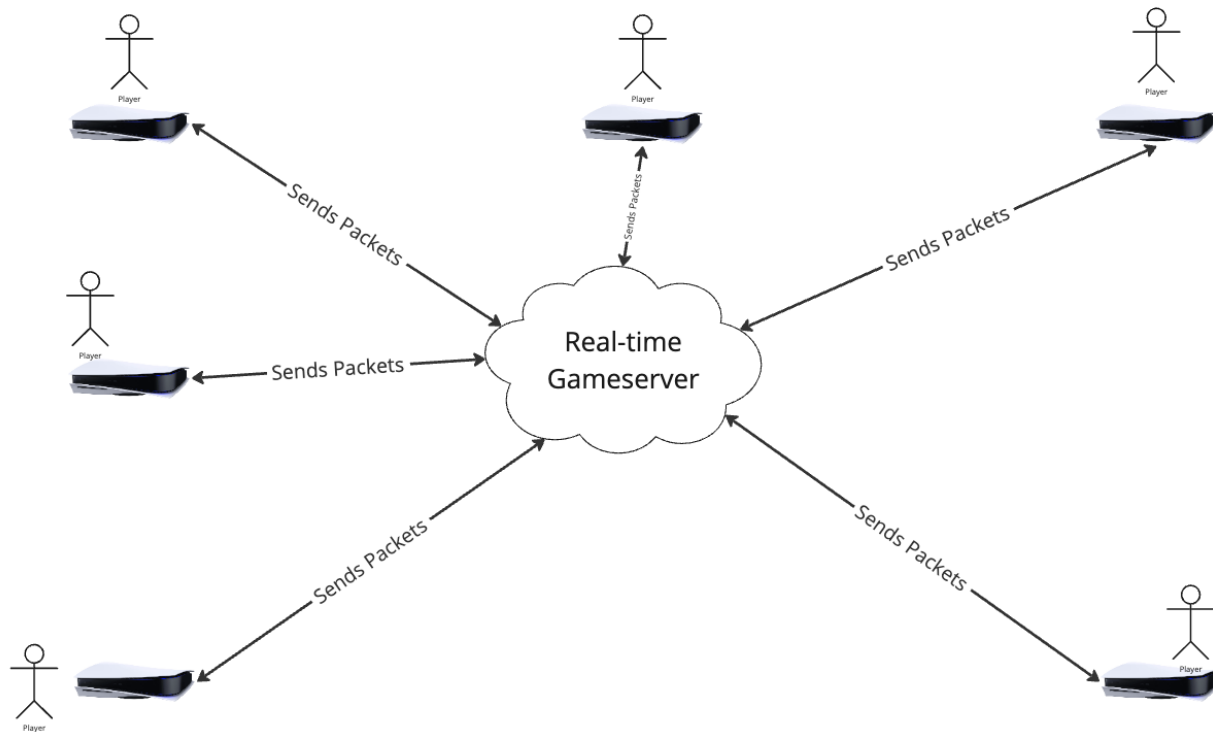
# Peer to Peer



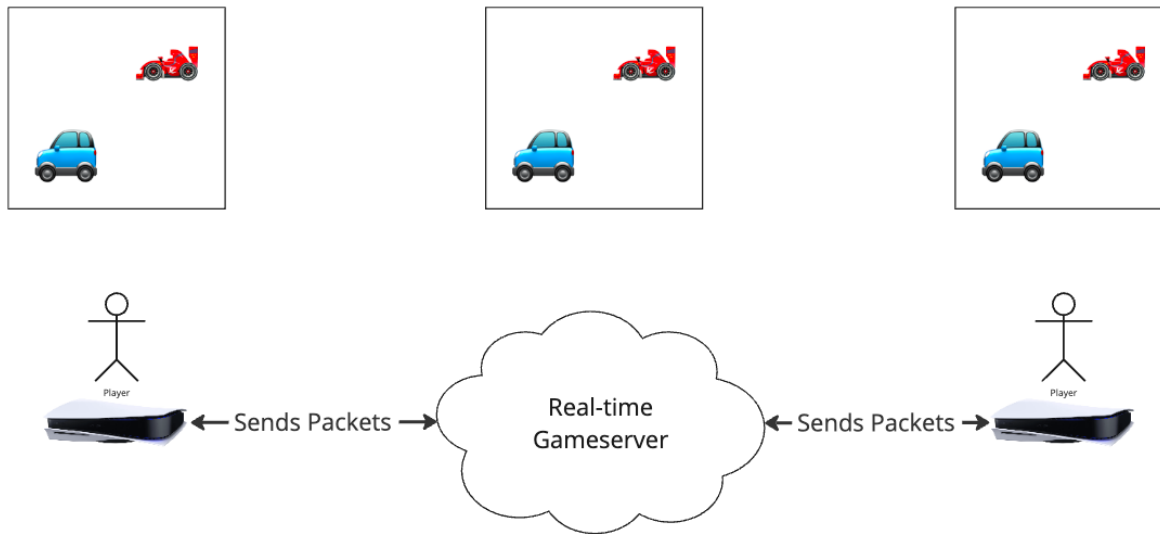
# Peer to Peer



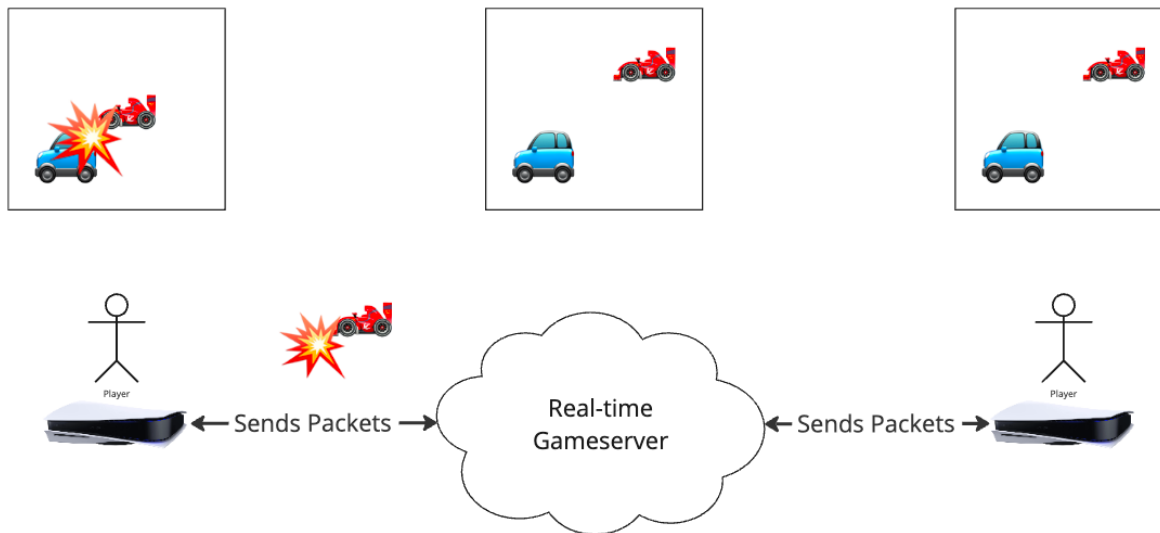
# Dedicated Game Server



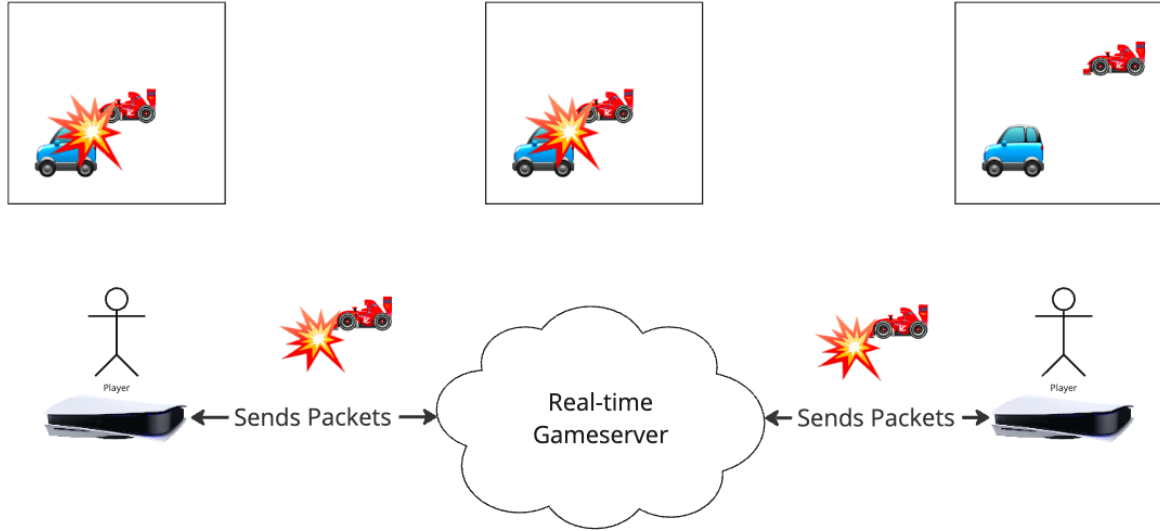
# Dedicated Game Server



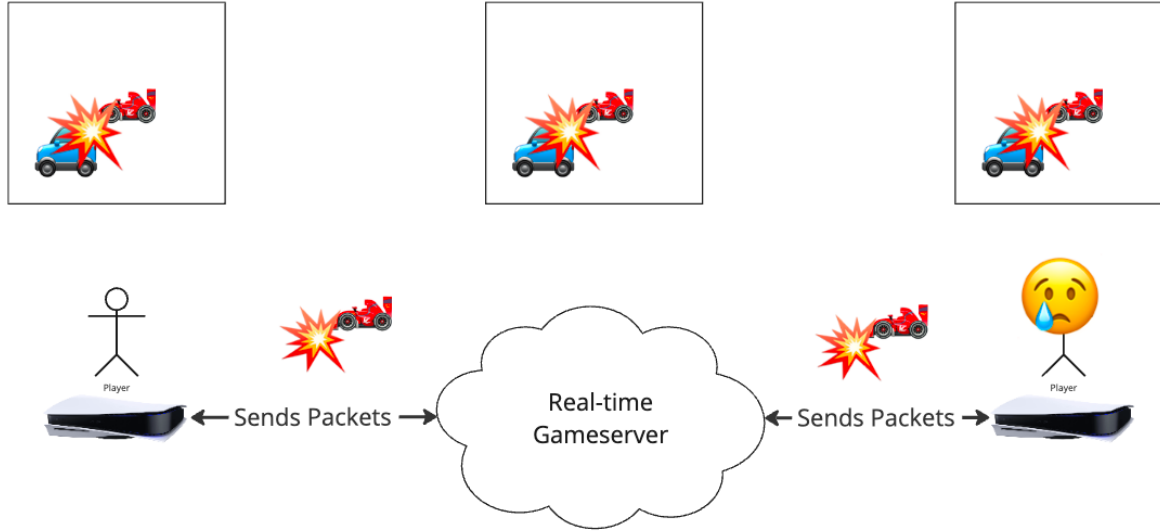
# Dedicated Game Server



# Dedicated Game Server



# Dedicated Game Server



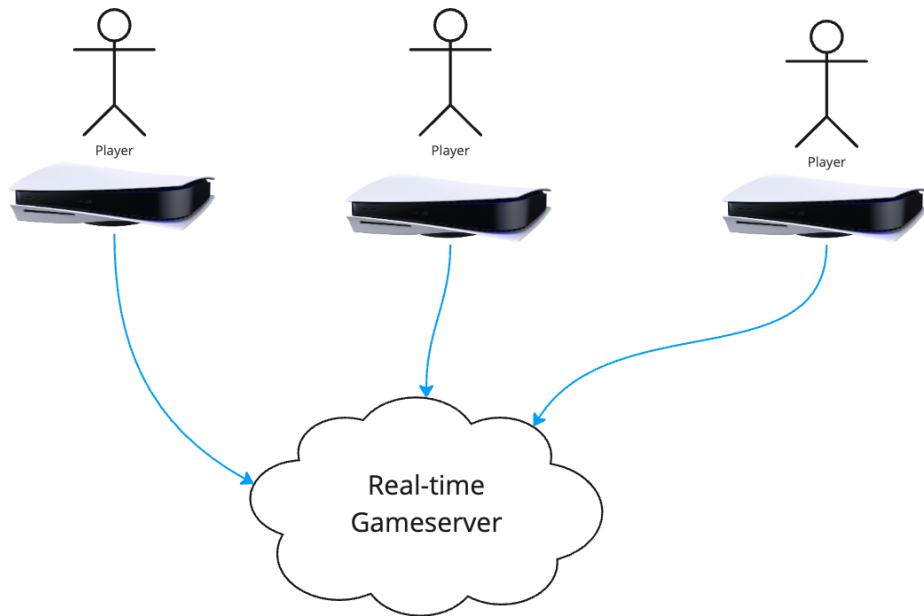
# Real-time

Moment to moment

UDP

Persistent Connection

Latency Sensitive





# LEVEL SELECT:

TUTORIAL: REAL-TIME GAMSERVERS

LEVEL 1: RUNNING IN KUBERNETES

LEVEL 2: AUTOSCALING

LEVEL 3: MULTI-REGION

# Agones



Open Source Project for running  
game servers in Kubernetes

Made by Google

`googleforgames/agones`



# Pod Termination



Cluster Autoscaler

Deployment Rollouts

# GameServers

```
apiVersion: agones.dev/v1
kind: GameServer
metadata:
  name: us-west-2-game-servers-p2qvj-6tj9k
  namespace: dev-game-server
spec:
  container: simple-game-server
  ports:
  - container: simple-game-server
    containerPort: 7654
    hostPort: 7025
    portPolicy: Dynamic
    protocol: UDP
  scheduling: Packed
  template:
    spec:
      containers:
      - image: gcr.io/agones-images/simple-game-server:0.13
        name: simple-game-server
```

```
kubectl get gameserver
```

| NAME                               | STATE | ADDRESS   | PORT | NODE                                    | AGE   |
|------------------------------------|-------|---|------|---|-------|
| us-west-2-game-servers-p2qvj-669m8 | Ready | ec2-35-89-190-116.us-west-2.compute.amazonaws.com | 7140 | ip-10-0-6-88.us-west-2.compute.internal | 3d18h |
| us-west-2-game-servers-p2qvj-6tj9k | Ready | ec2-35-89-190-116.us-west-2.compute.amazonaws.com | 7025 | ip-10-0-6-88.us-west-2.compute.internal | 3d18h |

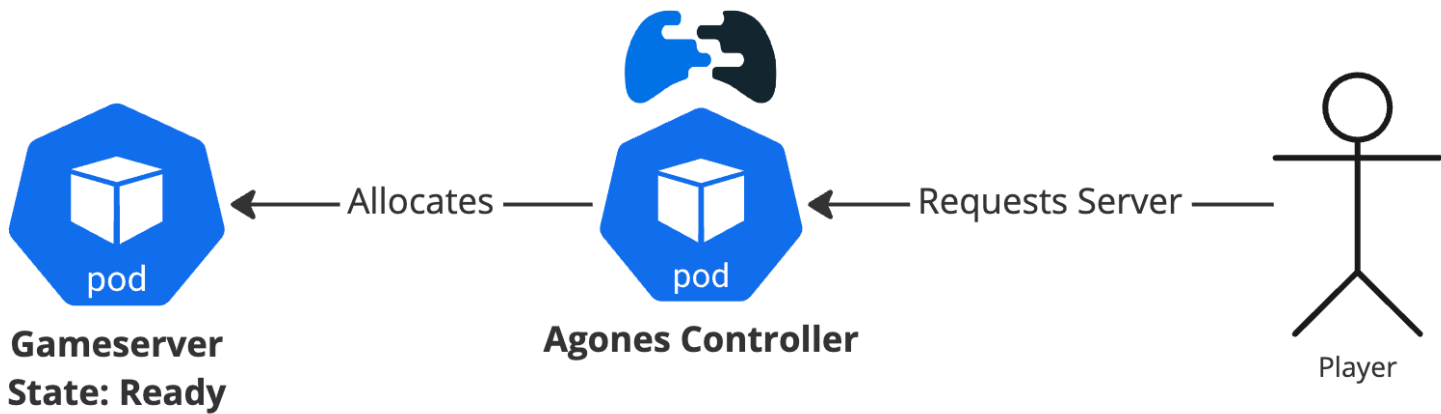
# GameServers

```
apiVersion: agones.dev/v1
kind: GameServer
metadata:
  name: us-west-2-game-servers-p2qvj-6tj9k
  namespace: dev-game-server
spec:
  container: simple-game-server
  ports:
  - container: simple-game-server
    containerPort: 7654
    hostPort: 7025
    portPolicy: Dynamic
    protocol: UDP
  scheduling: Packed
  template:
    spec:
      containers:
      - image: gcr.io/agones-images/simple-game-server:0.13
        name: simple-game-server
```

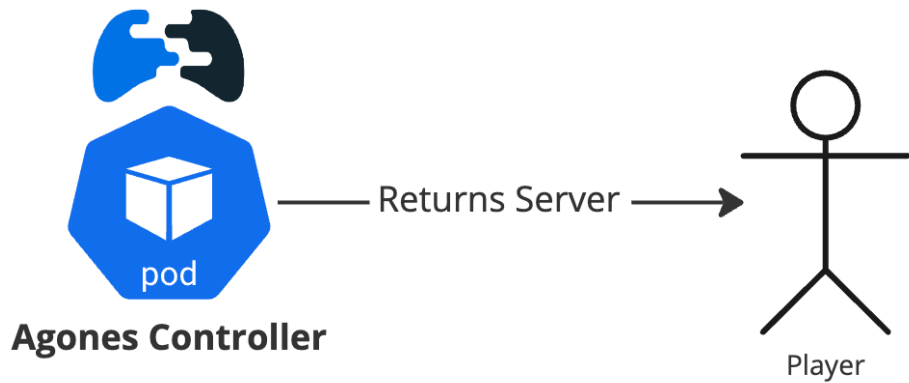
kubectl get gameserver

| NAME                               | STATE | ADDRESS   | PORT | NODE                                    | AGE   |
|------------------------------------|-------|---|------|---|-------|
| us-west-2-game-servers-p2qvj-669m8 | Ready | ec2-35-89-190-116.us-west-2.compute.amazonaws.com | 7140 | ip-10-0-6-88.us-west-2.compute.internal | 3d18h |
| us-west-2-game-servers-p2qvj-6tj9k | Ready | ec2-35-89-190-116.us-west-2.compute.amazonaws.com | 7025 | ip-10-0-6-88.us-west-2.compute.internal | 3d18h |

# GameServers



# GameServers



# GameServer Allocation



```
apiVersion: "allocation.agones.dev/v1"
kind: GameServerAllocation
spec:
  selectors:
    - matchLabels:
        agones.dev/fleet: my-fleet
  players:
    minAvailable: 10
    maxAvailable: 20
```



# GameServer Allocation

```
apiVersion: "allocation.agones.dev/v1"
kind: GameServerAllocation
spec:
  selectors:
    - matchLabels:
        agones.dev/fleet: my-fleet
      players:
        minAvailable: 10
        maxAvailable: 20
```

# GameServer Allocation

```
apiVersion: "allocation.agones.dev/v1"
kind: GameServerAllocation
metadata:
  name: my-allocation
spec:
  selectors:
    - matchLabels:
        agones.dev/fleet: my-fleet
  players:
    minAvailable: 10
    maxAvailable: 20
status:
  address: ec2-54-214-163-180.us-west-2.compute.amazonaws.com
  gameServerName: simple-game-server-ngtq6-ctgjc
  nodeName: ip-10-0-1-113.us-west-2.compute.internal
  ports:
    - name: default
      port: 7044
  state: Allocated
```

# GameServer Allocation

```
apiVersion: "allocation.agones.dev/v1"
kind: GameServerAllocation
metadata:
  name: my-allocation
spec:
  selectors:
    - matchLabels:
        agones.dev/fleet: my-fleet
  players:
    minAvailable: 10
    maxAvailable: 20
status:
  address: ec2-54-214-163-180.us-west-2.compute.amazonaws.com
  gameServerName: simple-game-server-ngtq6-ctgjc
  nodeName: ip-10-0-1-113.us-west-2.compute.internal
  ports:
    - name: default
      port: 7044
  state: Allocated
```

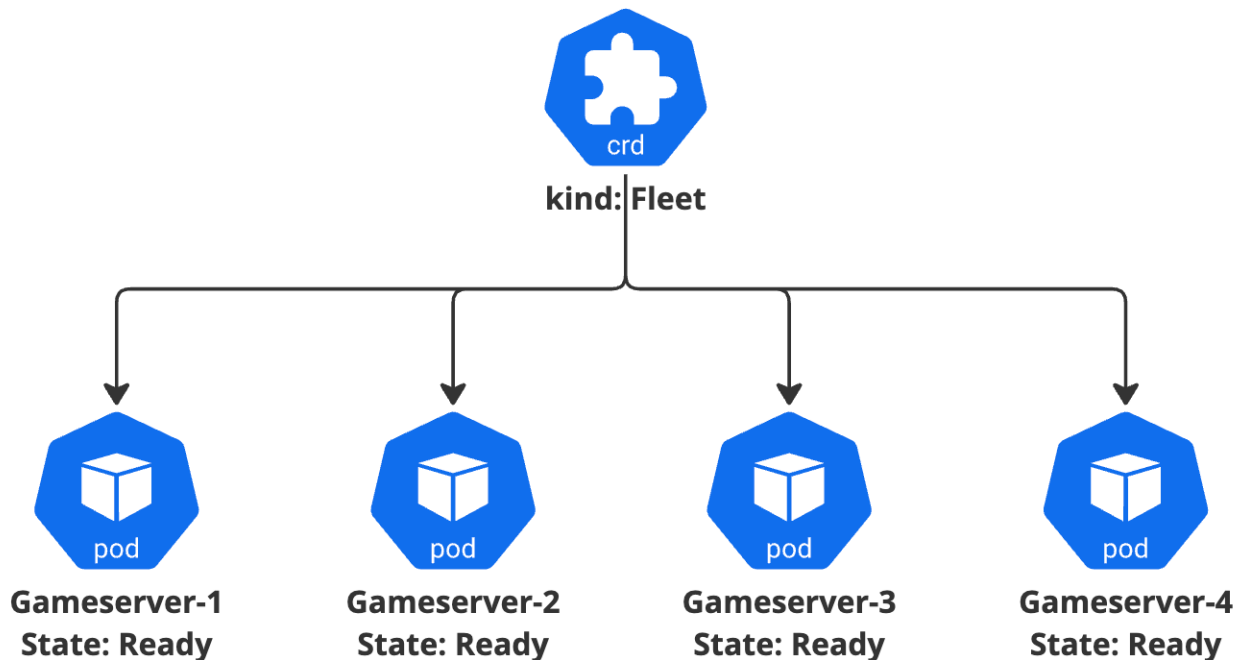
# Fleets

```
apiVersion: agones.dev/v1
kind: Fleet
  name: us-west-2-game-servers
  namespace: dev-game-server
spec:
  replicas: 2
  scheduling: Packed
  template:
    spec:
      ports:
      - containerPort: 7654
        name: default
      template:
        spec:
          containers:
          - image: gcr.io/agones-images/simple-game-server:0.13
            name: simple-game-server
```

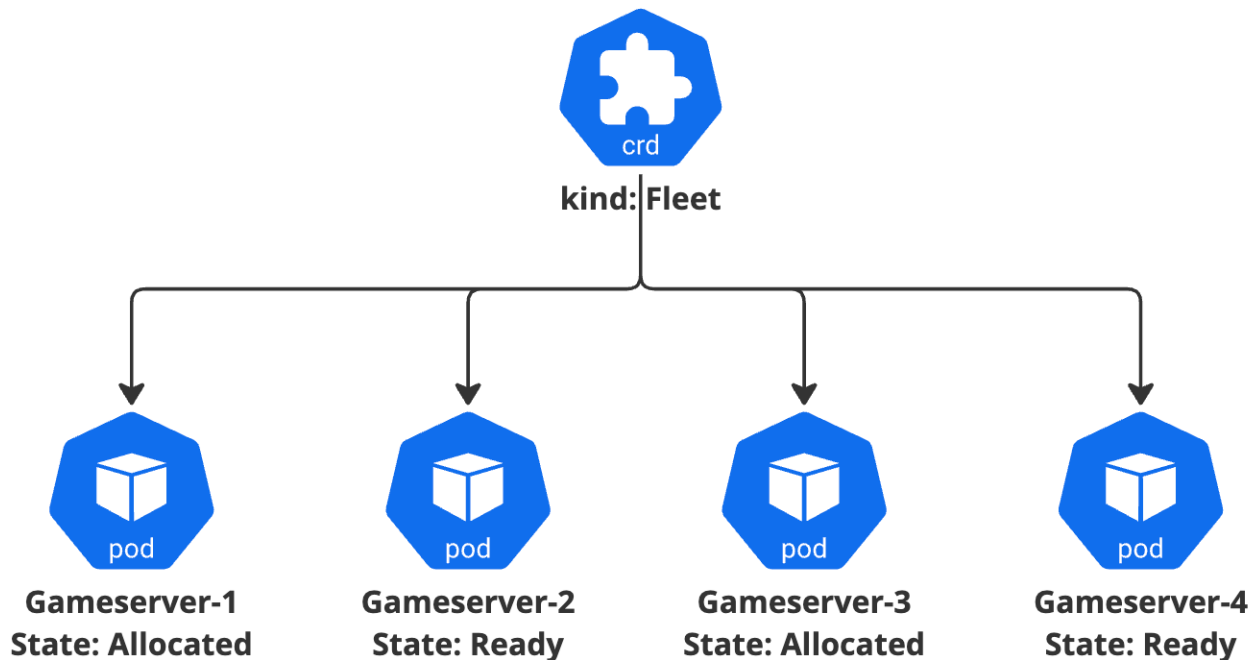
```
kubectl get fleets
```

| NAME                   | SCHEDULING | DESIRED | CURRENT | ALLOCATED | READY | AGE   |
|------------------------|------------|---------|---------|-----------|-------|-------|
| us-west-2-game-servers | Packed     | 2       | 2       | 0         | 2     | 3d22h |

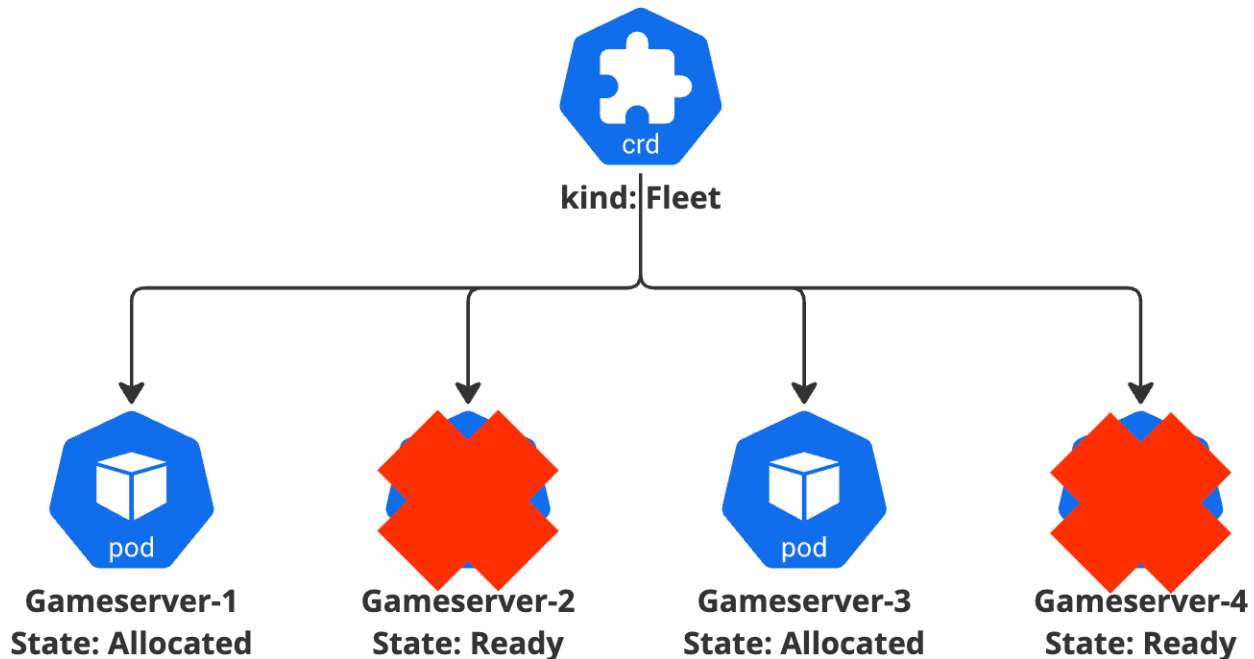
# Fleets



# Kubectl scale fleet --replicas=1



# Kubectl scale fleet --replicas=1



# LEVEL SELECT:

TUTORIAL: REAL-TIME GAMSERVERS

LEVEL 1: RUNNING IN KUBERNETES

LEVEL 2: AUTOSCALING

LEVEL 3: MULTI-REGION



# Autoscaling



HPA based on CPU/Memory doesn't make sense with gameservers

Autoscaling needs to consider the state of the gameservers

# Fleet Autoscaler



```
apiVersion: "autoscaling.agones.dev/v1"
kind: FleetAutoscaler
metadata:
  name: fleet-autoscaler-example
spec:
  fleetName: fleet-example
  policy:
    type: Buffer
    buffer:
      bufferSize: 5
      minReplicas: 10
      maxReplicas: 20
```

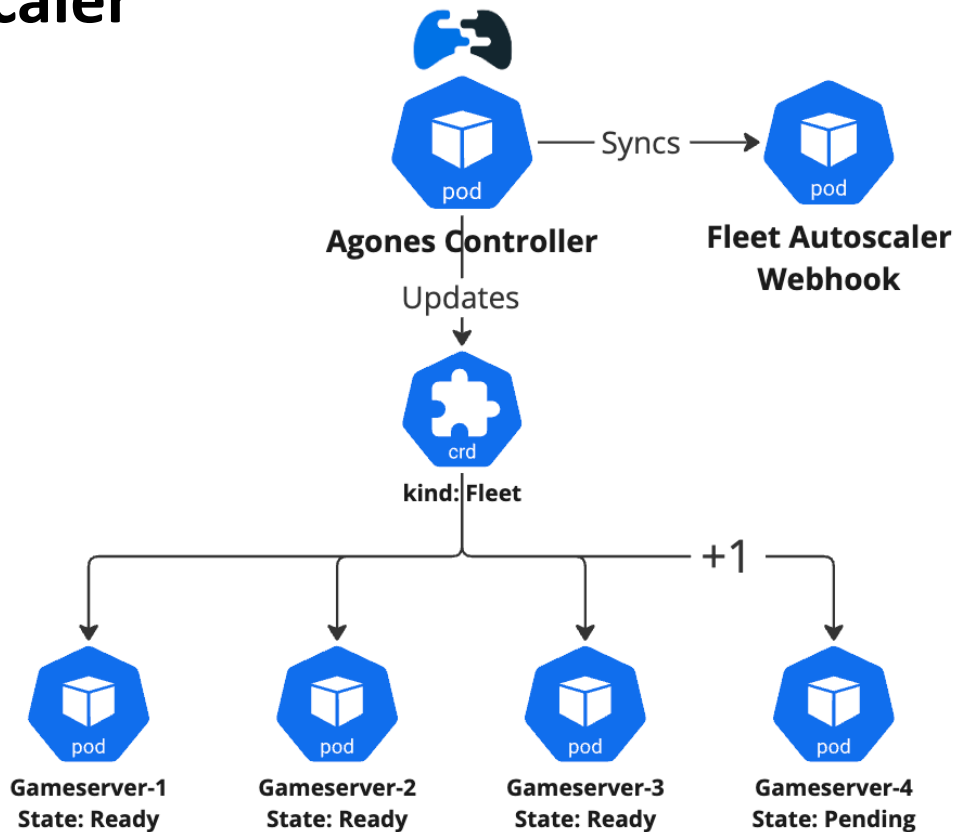
# Fleet Autoscaler

```
apiVersion: "autoscaling.agones.dev/v1"
kind: FleetAutoscaler
metadata:
  name: fleet-autoscaler-example
spec:
  fleetName: fleet-example
  policy:
    type: Buffer
    buffer:
      bufferSize: 5
      minReplicas: 10
      maxReplicas: 20
```

# Fleet Autoscaler

```
apiVersion: "autoscaling.agones.dev/v1"
kind: FleetAutoscaler
metadata:
  name: fleet-autoscaler-example
spec:
  fleetName: fleet-example
  policy:
    type: Webhook
    webhook:
      service:
        name: autoscaler-webhook-service
        namespace: default
        path: scale
```

# Fleet Autoscaler



# LEVEL SELECT:

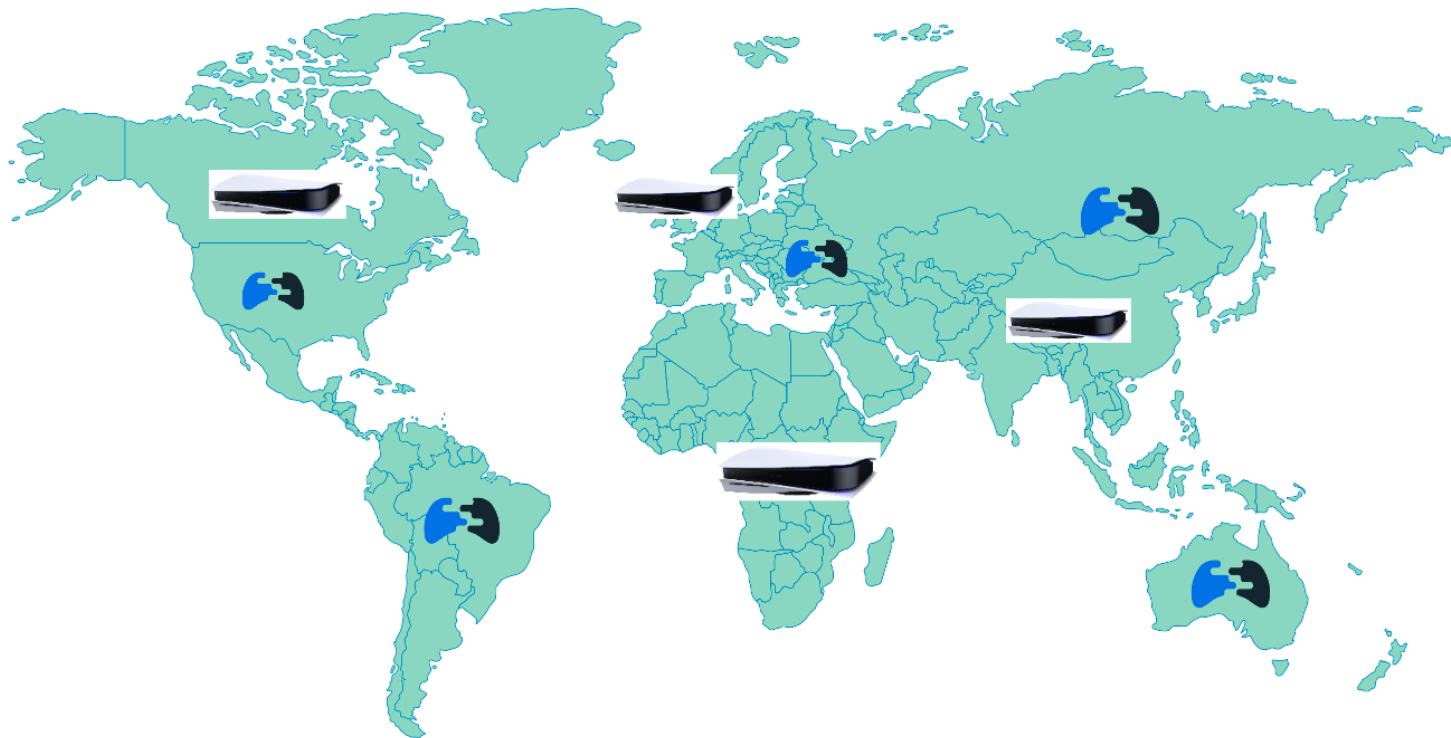
TUTORIAL: REAL-TIME GAMSERVERS

LEVEL 1: RUNNING IN KUBERNETES

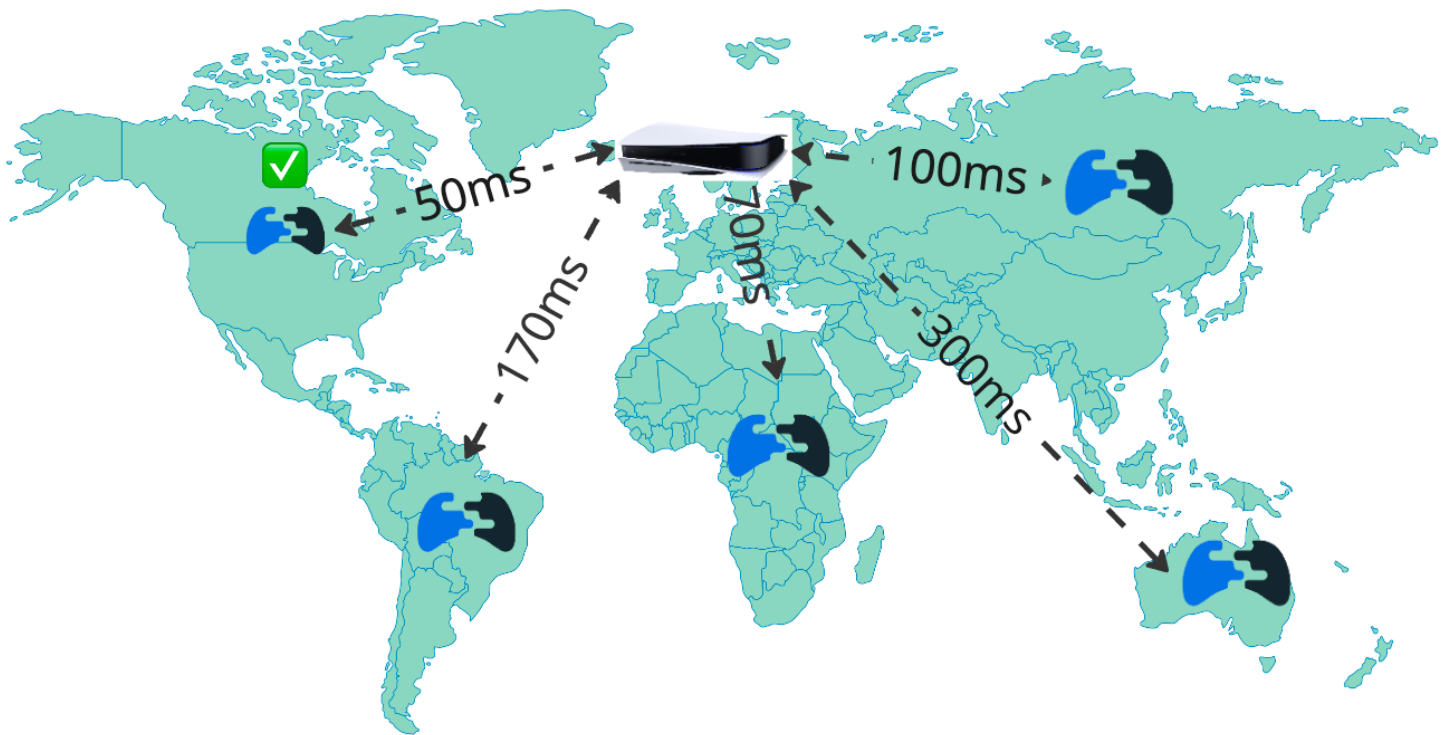
LEVEL 2: AUTOSCALING

LEVEL 3: MULTI-REGION

# Multi-Region Agones



# Multi-Region Agones

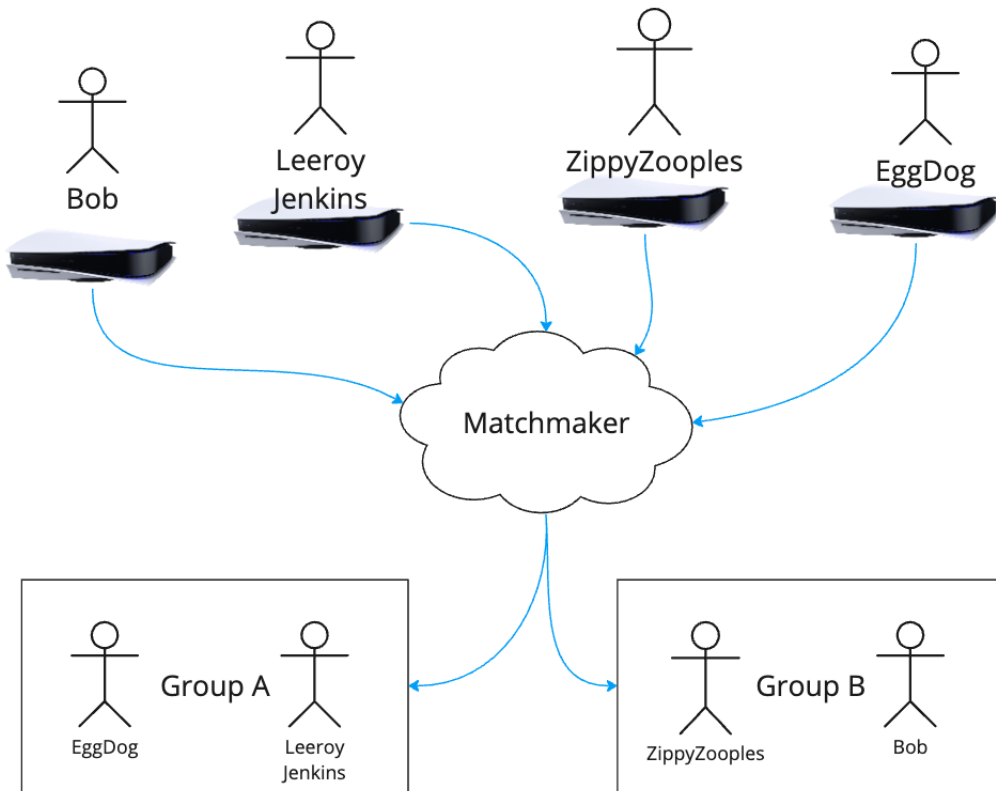




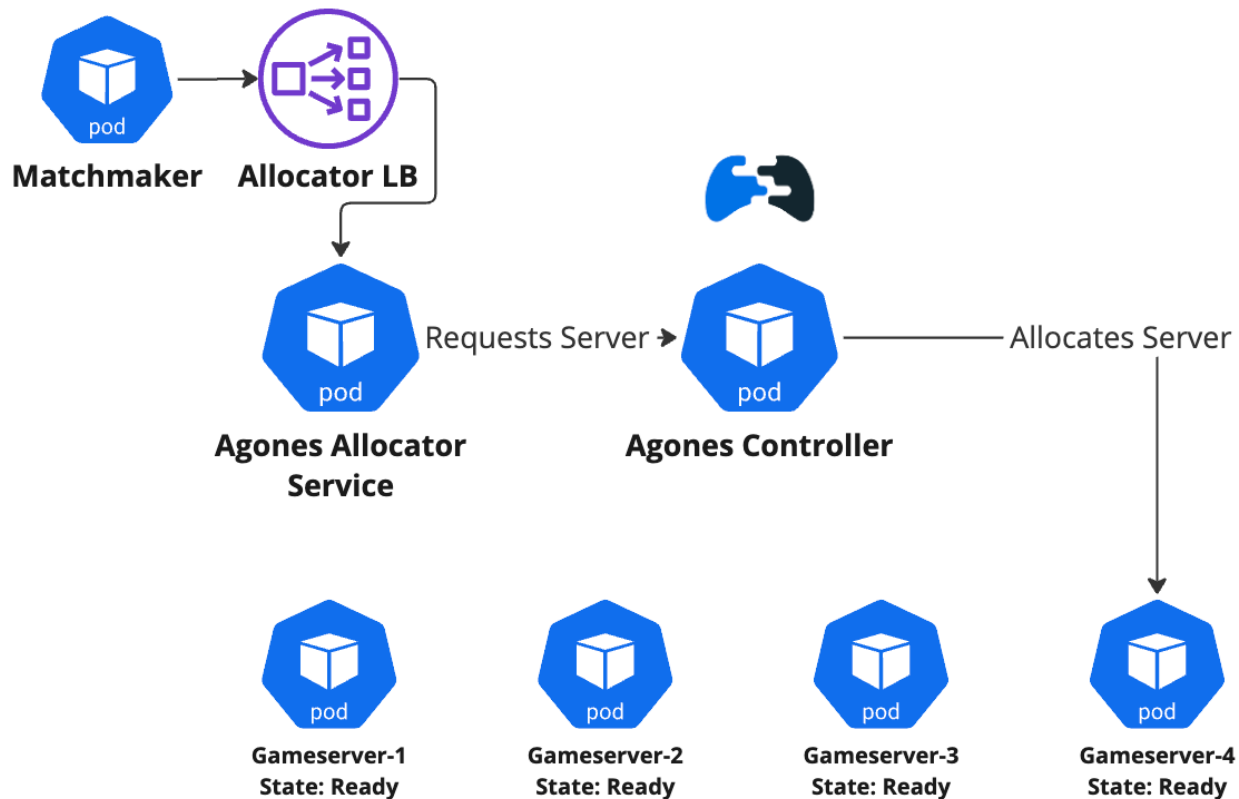
# Matchmaker



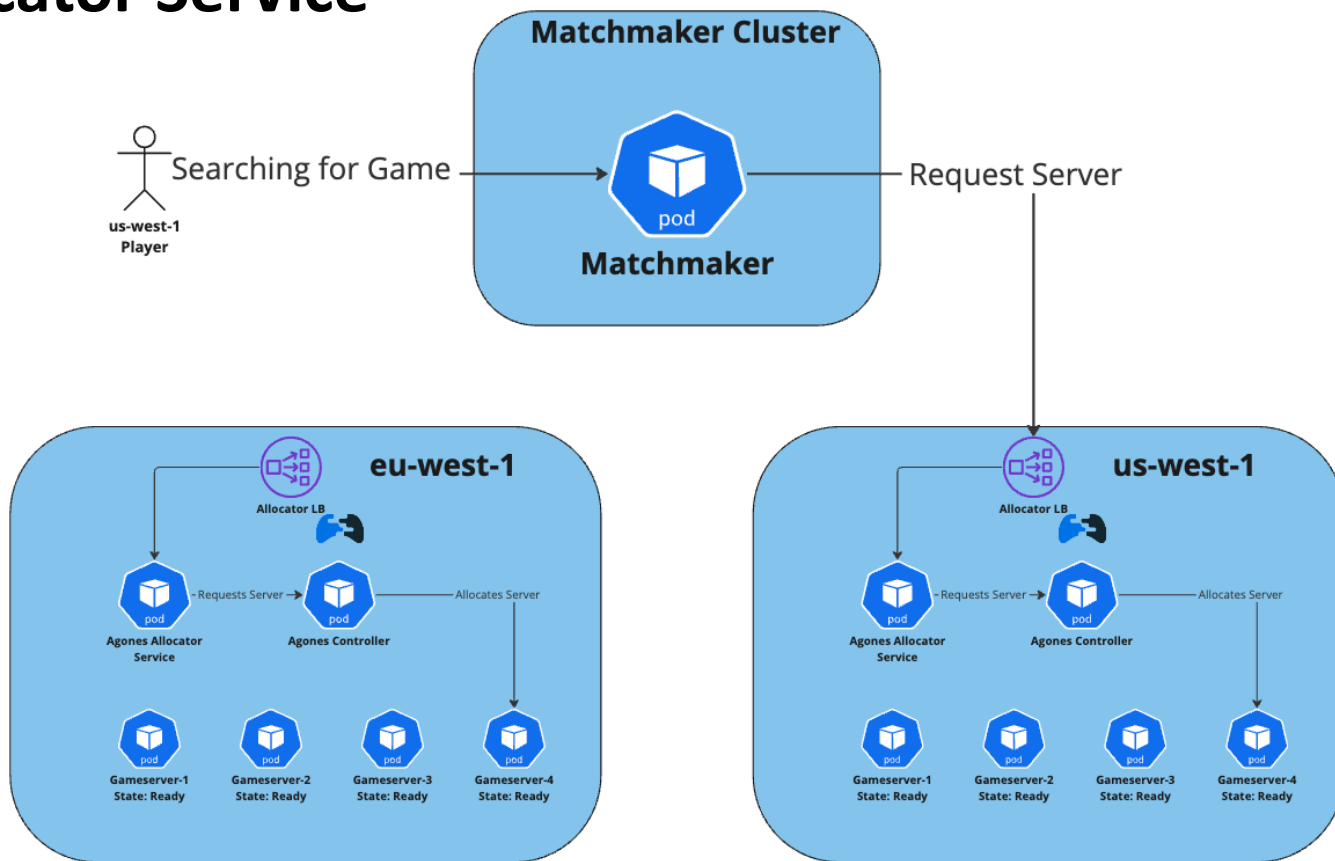
# Matchmaker



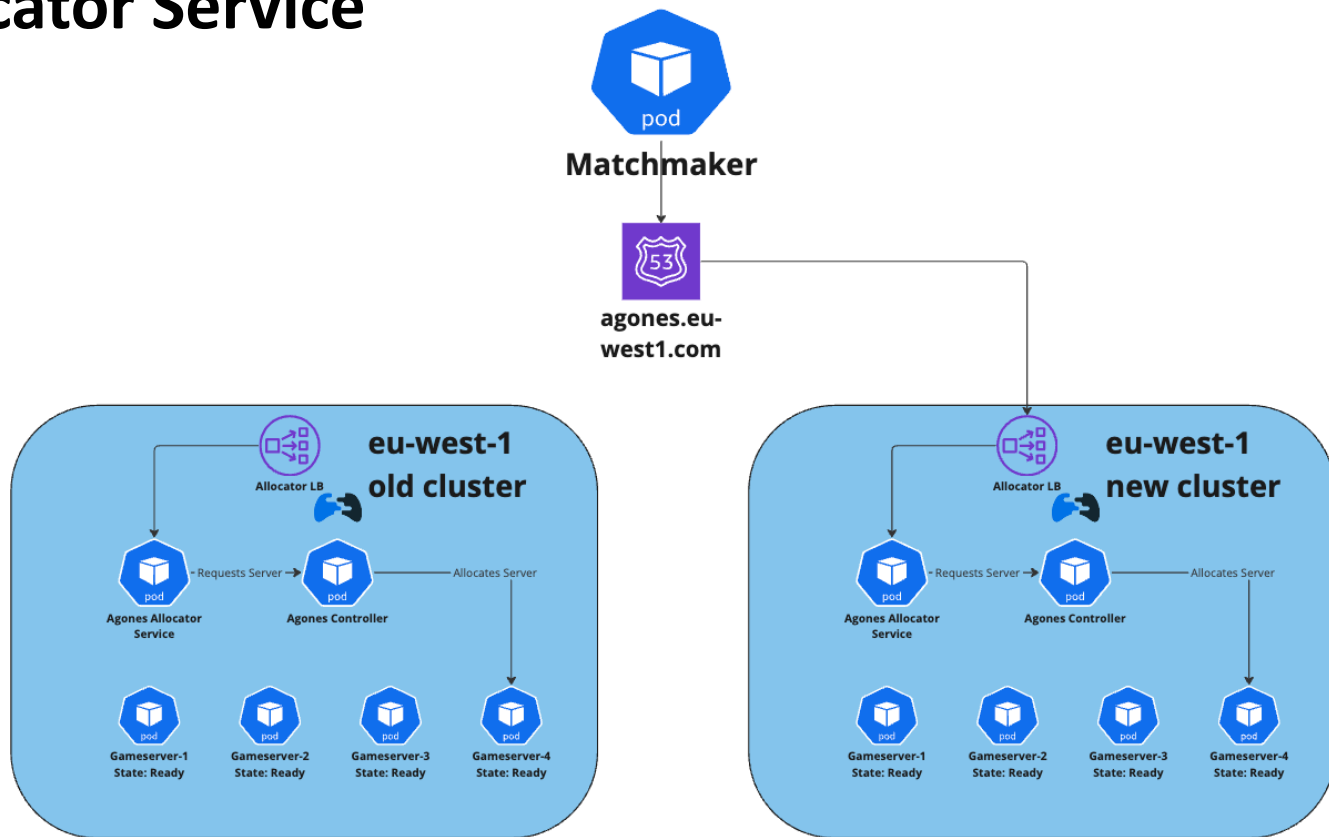
# Allocator Service



# Allocator Service



# Allocator Service



# To Conclude...

Fleets and Gameservers

Fleet Autoscalers

Allocator Service

Multi-Region





Scan the QR  
Code to leave  
feedback and  
see slides!