


# RBAC to the Future: Untangling Authorization in Kubernetes


Jimmy Mesta  
CTO / Co-Founder, KSOC



[tinyurl.com/rbac-to-the-future](https://tinyurl.com/rbac-to-the-future)

# Contents

- › What is Kubernetes again?
- › Identity is the New Perimeter
- › RBAC Terminology
- › AuthN/AuthZ Flow
- › Built-ins
- › Gotchas
- › Monitoring Techniques



**Kubernetes is  
an open-source  
platform built to  
automate  
deployment,  
scaling and  
orchestration of  
containers**

**Identity is the  
new perimeter.**

**Access control in  
Kubernetes can  
be complicated.**



K00 | Introduction

K01 | Insecure Workload Configurations

K02 | Supply Chain Vulnerabilities

K03 | Overly Permissive RBAC Configurations

K04 | Lack of Centralized Policy Enforcement

K05 | Inadequate Logging and Monitoring

K06 | Broken Authentication Mechanisms

K07 | Missing Network Segmentation Controls

K08 | Secrets Management Failures

K09 | Misconfigured Cluster Components

K10 | Outdated and Vulnerable Kubernetes Components

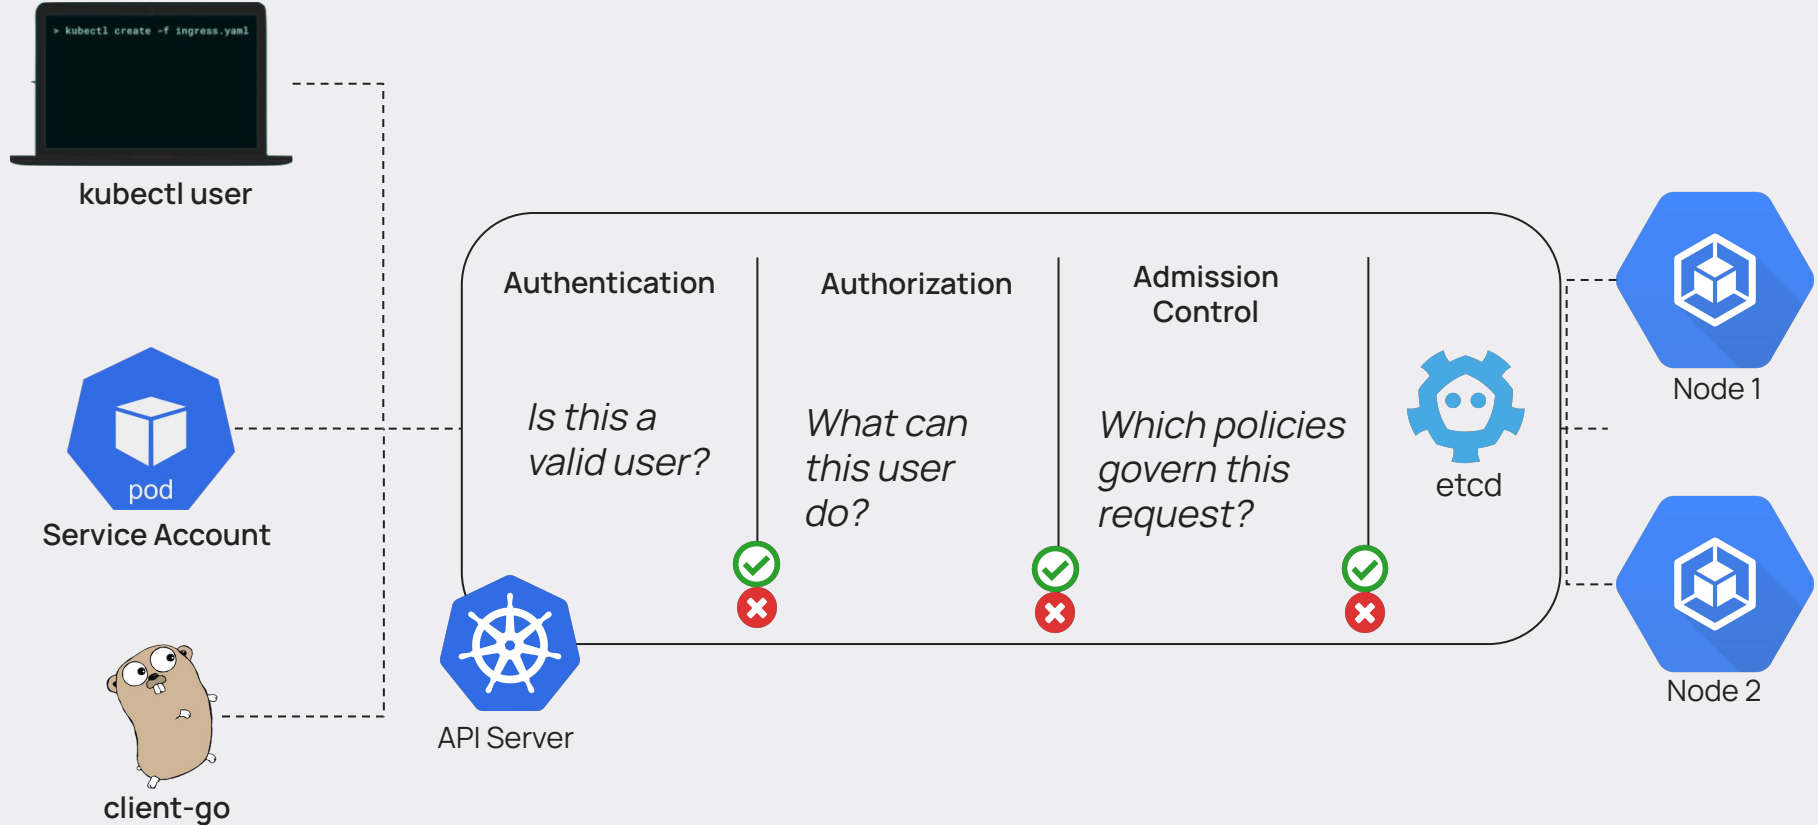
# OWASP Kubernetes Top Ten

[owasp.org/www-project-kubernetes-top-ten](https://owasp.org/www-project-kubernetes-top-ten)

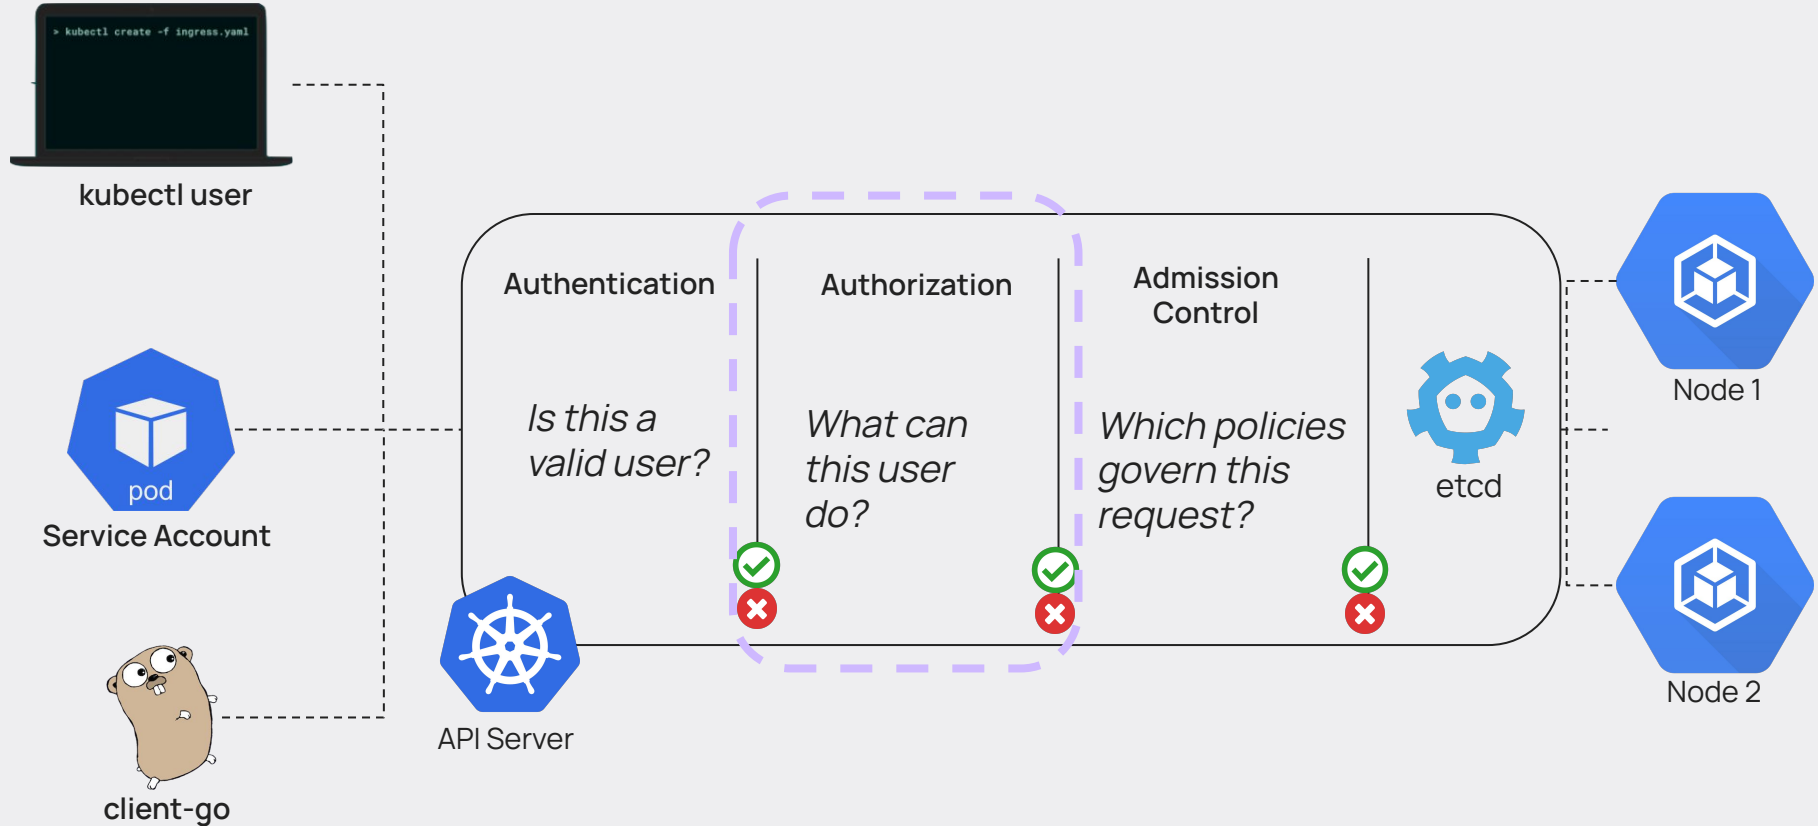


@jimmesta

# API Request Flow



# API Request Flow



---

**Role-Based Access Control**  
enables fine grained access for  
users, groups, and service  
accounts within Kubernetes.  
RBAC can be extremely difficult  
to scope appropriately which  
opens up additional privileges.





---

## Users

you@email.com

Service Account

## API Resources

Namespaces

Pod

Service

Secrets

## Operations

Get

List

Delete

Patch



# RBAC: Role

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  namespace: dev-group-1
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["*"]
  verbs: ["get", "watch", "list"]
```



# RBAC:ClusterRole

```
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  namespace: dev-group-1
rules:
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["get", "watch", "list"]
```



# RBAC: RoleBinding

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: read-pods
  namespace: default
subjects:
  - kind: User
    name: jane
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```



# RBAC: ClusterRoleBinding

```
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: read-pods
subjects:
  -kind: User
    name: jane
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```



# User-facing ClusterRoles

<i>cluster-admin</i>	Allows super-user access to perform any action on any resource.
<i>admin</i>	Allows admin access, intended to be granted within a namespace using a RoleBinding.
<i>edit</i>	Allows read/write access to most objects in a namespace. This role does not allow viewing or modifying roles or role bindings
<i>view</i>	Allows read-only access to see most objects in a namespace. It does not allow viewing roles or role bindings.

<https://kubernetes.io/docs/concepts/security/rbac-good-practices/>



# What's wrong with this picture?

```
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRoleBinding
metadata:
  name: redacted-rbac
subjects:
- kind: ServiceAccount
  name: default
  namespace: default
roleRef:
  kind: ClusterRole
  name: cluster-admin
  apiGroup: rbac.authorization.k8s.io
```



# RBAC Aggregate Permissions

```
# rakless --namespace default
NAME                               LIST  CREATE  UPDATE  DELETE
bindings                           ✓      ✗
configmaps                         ✓      ✓
controllerrevisions.apps           ✓      ✗
cronjobs.batch                     ✓      ✓
daemonsets.apps                    ✓      ✓
daemonsets.extensions              ✓      ✓
deployments.apps                   ✓      ✓
deployments.extensions             ✓      ✓
endpoints                         ✓      ✓
events                             ✗      ✗
events.events.k8s.io               ✗      ✗
horizontalpodautoscalers.autoscaling ✓      ✓
ingresses.extensions               ✓      ✓
jobs.batch                         ✓      ✓
leases.coordination.k8s.io         ✗      ✗
limitranges                       ✓      ✗
localsubjectaccessreviews.authorization.k8s.io ✗      ✗
networkpolicies.extensions         ✓      ✓
```

## api-policy service account

### Assigned Roles and Cluster Roles

Name	Kind	Namespace
api-policy-ksoc-pod-readonly	cluster role	*
api-policy	role	ksoc

### Aggregate Permissions

Namespace	Cluster-wide
Object	Verbs
pods	get, list, watch



[github.com/corneliusweig/rakless](https://github.com/corneliusweig/rakless)

[github.com/alcideio/rbac-tool](https://github.com/alcideio/rbac-tool)

[github.com/appvia/krane](https://github.com/appvia/krane)



# Overly Permissive RBAC Configurations

```
curl http://127.0.0.1:8001/api/v1/namespaces/default/secrets/abc -H "Authorization: Bearer $(kubectl -n default get secrets -ojson | jq '.items[] | select(.metadata.anno
{
  "kind": "Status",
  "apiVersion": "v1",
  "metadata": {
  },
  "status": "Failure",
  "message": "secrets \"abc\" is forbidden: User \"system:serviceaccount:default:only-list-secrets-sa\" cannot get resource \"secrets\" in API group \"\" in the namespa
  "reason": "Forbidden",
  "details": {
    "name": "abc",
    "kind": "secrets"
  },
  "code": 403
}
# Now to get all secrets in the default namespace, despite not having "get" permission
curl http://127.0.0.1:8001/api/v1/namespaces/default/secrets?limit=500 -H "Authorization: Bearer $(kubectl -n default get secrets -ojson | jq '.items[] | select(.metadat
{
  "kind": "SecretList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/namespaces/default/secrets",
    "resourceVersion": "17718246"
  },
  "items": [
    REDACTED : REDACTED
  ]
}
```



# Overly Permissive RBAC Configurations

```
kind: Role
apiVersion:
rbac.authorization.k8s.io/v1
metadata:
  namespace: default
  name: secret-list
rules:
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["list"]
```

```
kubectl get secrets -A -o yaml
```

## GET Request

```
curl http://127.0.0.1:8001/api/v1/namespaces/default/secrets/abc
{
  "kind": "Status",=
  "metadata": {
  },
  "status": "Failure",
  (...)
},
  "code": 403
}
```

## LIST Request

```
curl http://127.0.0.1:8001/api/v1/namespaces/default/secrets?limit=500
{
  "kind": "SecretList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/namespaces/default/secrets",
  },
  "items": [
    foo : bar
  ]
}
```

@jimmesta

# CVE-2023-30512 (CubeFS)

```
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: cfs-csi-cluster-role
rules:
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["get", "list"]
```

CubeFS through 3.2.1 allows Kubernetes cluster-level **privilege escalation**. This occurs because DaemonSet has **cfs-csi-cluster-role** and can thus list all secrets, including the admin secret.



### Minimize wildcard use in Roles and ClusterRoles:Resources

Violating Resource	flux
Violation Message	ClusterRole flux set with resources:* for apiGroups:["*"] and verbs:["*"]. Roles and cluster roles should not use wildcards for resource entries.

#### Violation Details

Description	<p>Kubernetes Roles and ClusterRoles provide access to resources based on sets of objects and actions that can be taken on those objects. It is possible to set either of these to be the wildcard "*" which matches all items. Use of wildcards is not optimal from a security perspective as it may allow for inadvertent access to be granted when new resources are added to the Kubernetes API either as CRDs or in later versions of the product.</p> <p>Remediation: Remove wildcards in Role and ClusterRole resources.</p>
Policy ID	KSOC-K8S-WILDCARD-RESOURCES
Policy Name	policy-0025-role-wildcard-resources
Time First Detected	2022-11-28T12:52:55Z

#### Violated Manifest

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  annotations:
    helm.fluxcd.io/antecedent: flux:helmrelease/flux
    meta.helm.sh/release-name: flux
    meta.helm.sh/release-namespace: flux
  creationTimestamp: "2022-11-28T11:09:59Z"
  labels:
    app: flux
    app.kubernetes.io/managed-by: Helm
    chart: flux-1.11.2
    heritage: Helm
    release: flux
  name: flux
  resourceVersion: "2899"
  uid: 84ac948f-299e-4413-8593-f85a97d58e2a
rules:
- apiGroups:
  - '*'
  resources:
  - '*'
  verbs:
  - '*'
- nonResourceURLs:
  - '*'
  verbs:
  - '*'
```

# Detecting and stopping misconfigured RBAC



## Better Logs == Better RBAC

A Kubernetes Environment has the ability to generate logs at a variety of levels from many different components. When logs are not captured, stored, or actively monitored attackers have the ability to exploit vulnerabilities while going largely undetected.



# Better Logs == Better RBAC

```
{
  "kind": "Event",
  "verb": "get",
  "user": {
    "username": "serviceaccount:inventory:inventory",
    "uid": "3e94bfaf-8edc-4562-b2ed-44e9a9e565fb",
    "groups": [
      "serviceaccounts",
      "serviceaccounts:inventory"
    ]
  },
  "sourceIPs": [
    "10.16.4.17"
  ],
  "userAgent": "server/v0",
  "responseStatus": {
    "metadata": {},
    "code": 200
  },
  "requestReceivedTimestamp": "2021-07-16T14:33:24.429585Z",
  "stageTimestamp": "2021-07-16T14:33:24.434300Z",
  "annotations": {
    "authorization.k8s.io/decision": "allow",
    "authorization.k8s.io/reason": "RBAC: allowed by RoleBinding \\"inventory\\" of Role \\"inventory\\" to ServiceAccount \\"inventory/inventory\\"
  }
}
```

---

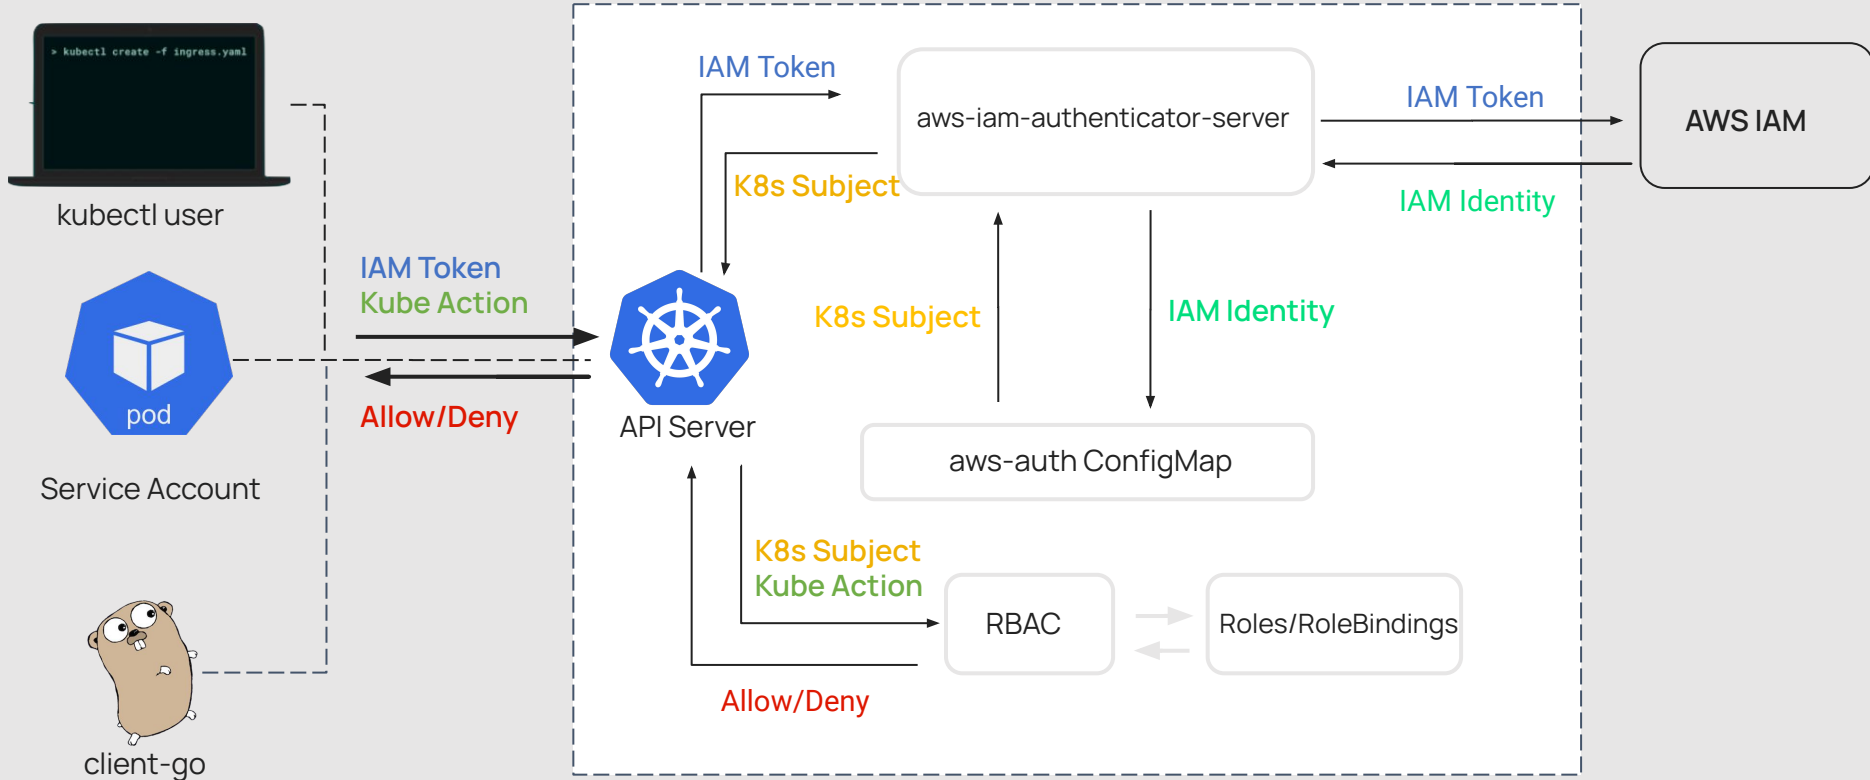
## RBAC Considerations

Watch for the use of the **Escalate**, **Impersonate**, and **Bind** verb as they can lead to RBAC bypass and privilege escalation.


Only grant RBAC permission to create **PersistentVolumes** when absolutely necessary as those workloads can then access the underlying host filesystem.



## AWS EKS AuthN/AuthZ








**Thanks,  
& let's hang out!**



`jimmy@ksoc.com`

<https://github.com/cncf/tag-security/issues/1051>

ksoc.com



KSOC

**Secure your  
clusters in  
real-time**