



KubeCon



CloudNativeCon

Europe 2023





KubeCon



CloudNativeCon

Europe 2023

Revamping Kubernetes with Contextual and Structured Logging, a Deep Dive

Maintainer Track, Kubernetes Structured Logging WG



#wg-structured-logging

Contributor

Kubernetes, Istio

Tetrade

Github: shivanshu1333

Twitter: shivanshu1333




KubeCon



CloudNativeCon

Europe 2023

Shout out to 🎉



Marek Siarkowicz

#wg-structured-logging
lead

Github: serathius
Twitter: serathius



Patrick Ohly

#wg-structured-logging
lead

Github: pohly

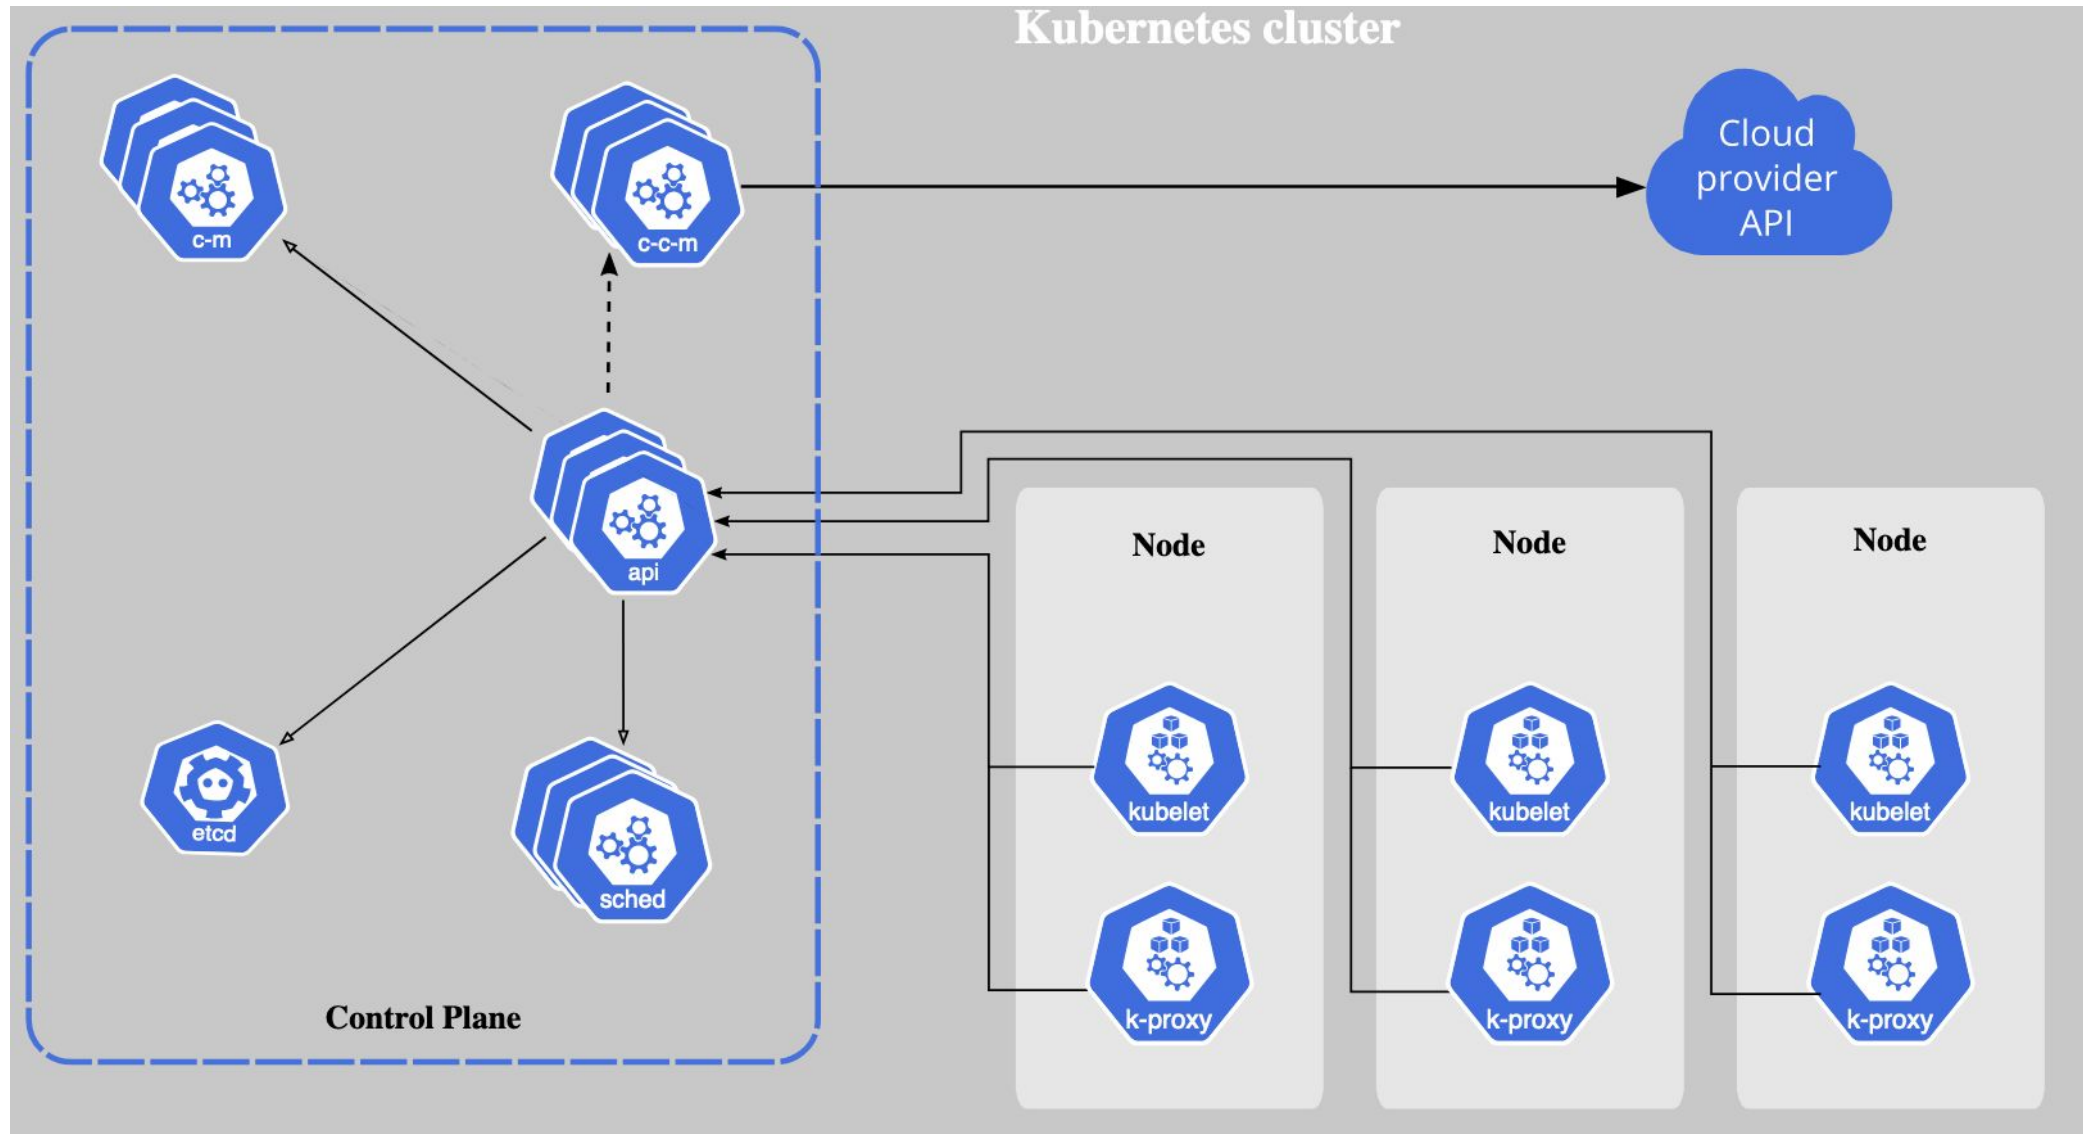
Agenda

- Introduction (~10 mins)
- Deep Dive (~15 mins)
- Migration instructions (~5 mins)
- Questions (~5 mins)

- Kubernetes core contributors
- Contributors to logging agents (fluentd, opentelemetry etc)
- End users of Kubernetes
- New contributors :)

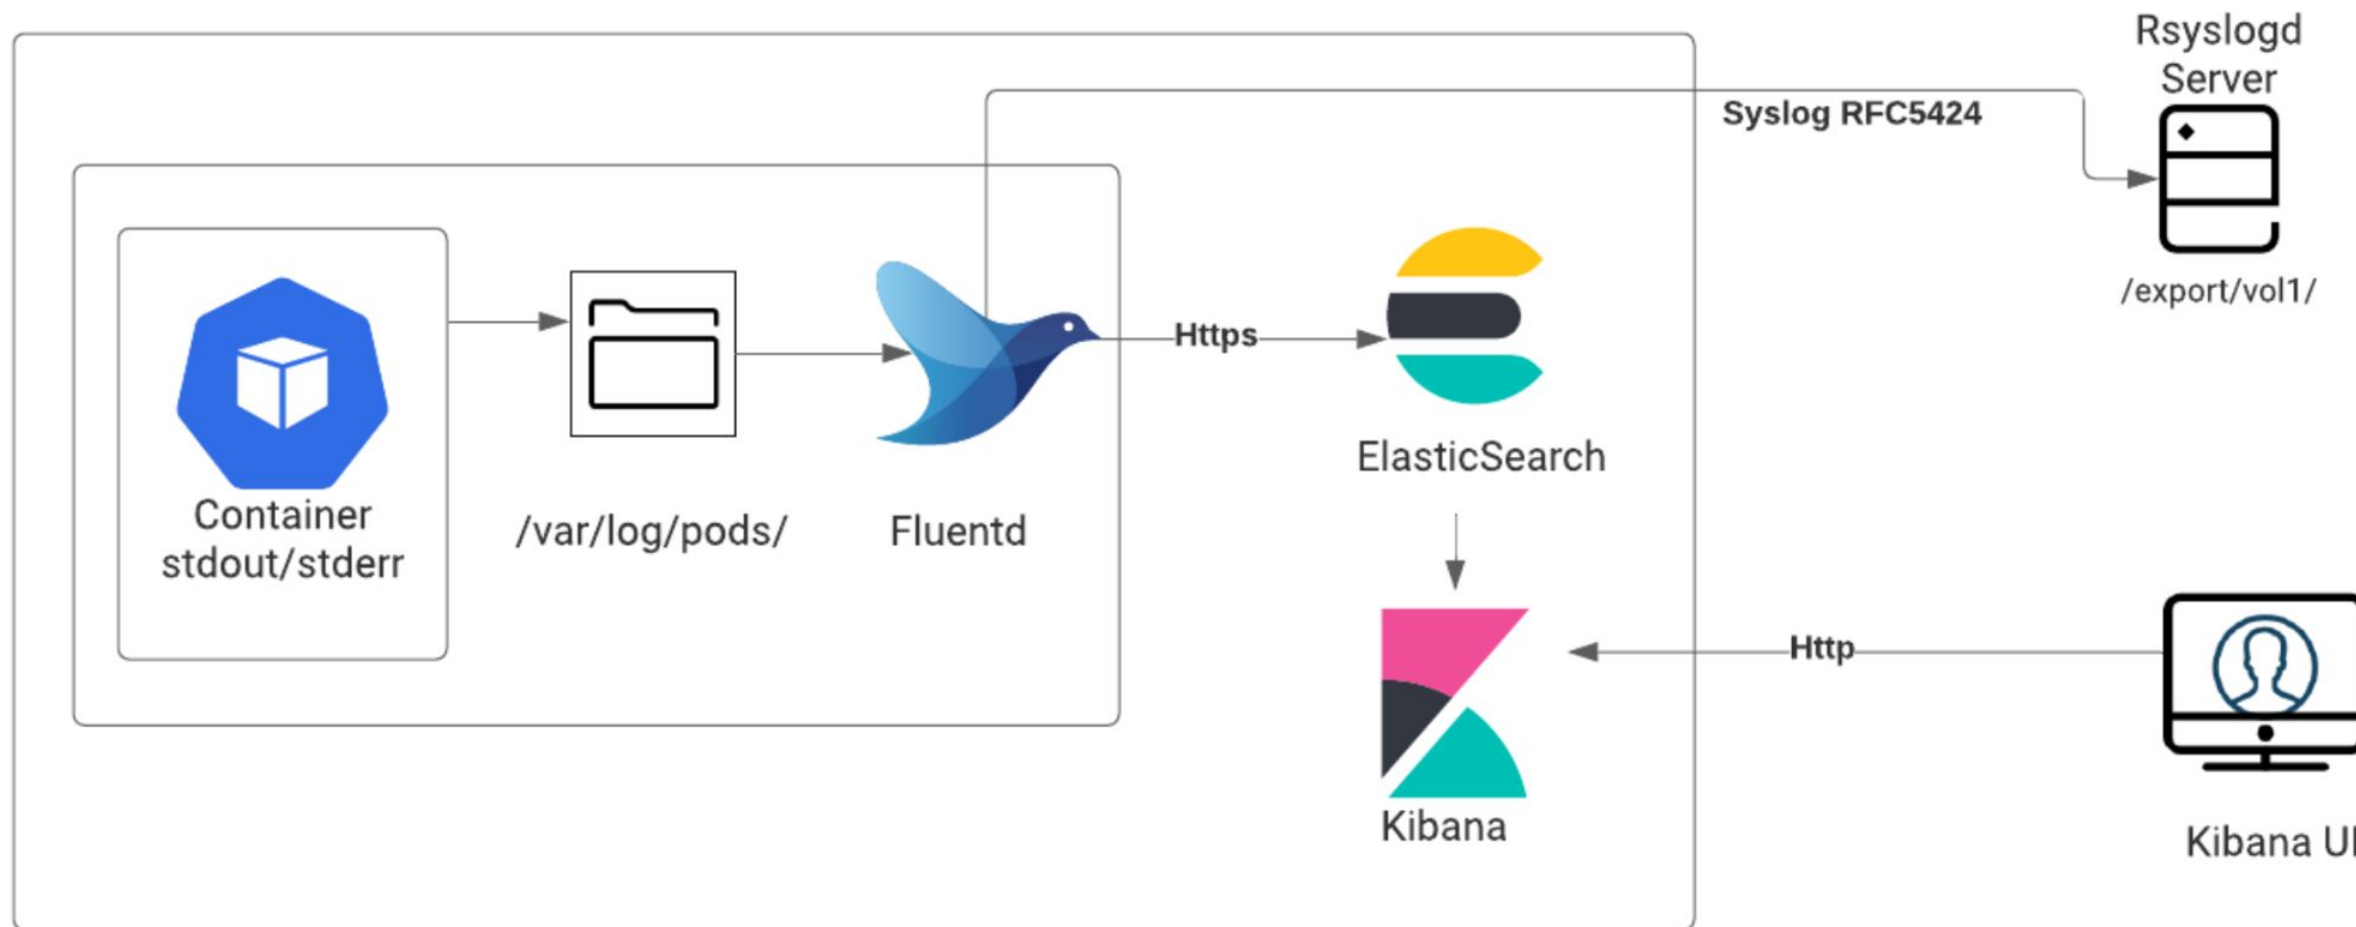
- Motivation
 - Kubernetes logs are messy :|
 - Kubernetes uses klog as its default logger (fork of glog)
 - Easy and standardized log collection
 - Fine grained information of Kubernetes components
 - Easy and automated monitoring

Introduction



Introduction

Standard K8 Approach



- Structured Logging
 - Introduction
 - Proposal
 - Goals/Non Goals
 - Deep dive
 - Implementation details
 - Migration details

- Proposal
 - Define standard structure for Kubernetes log messages
 - Add methods to klog to enforce this structure
 - Add ability to configure Kubernetes components to produce logs in JSON format
 - Initiate migration to structured logging

eg:

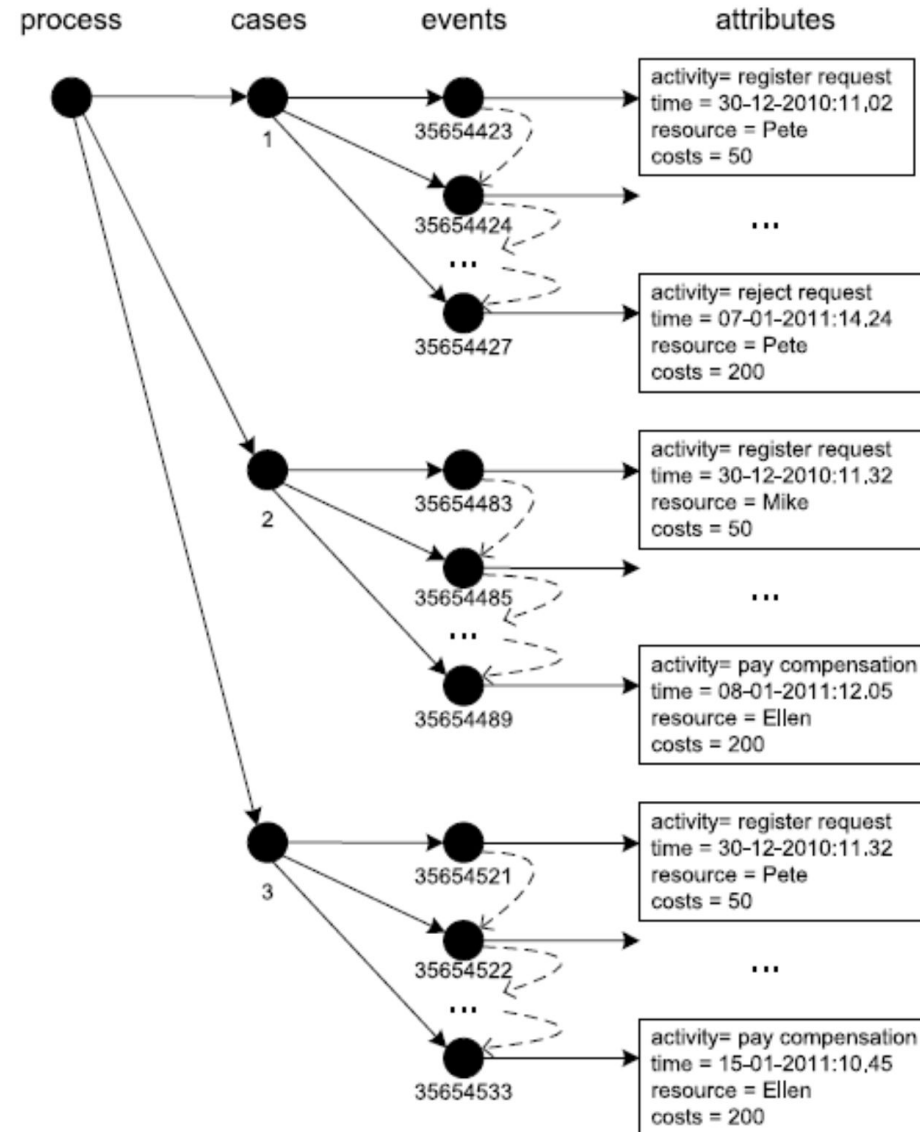
```
I1025 00:15:15.525108    1 controller_utils.go:116] "Pod status updated" pod="kube-system/kubedns" status="ready"
```

- Goals/Non Goals
 - Goals
 - Make most common logs more queryable by standardizing log message and references to Kubernetes objects (Pods, Nodes etc.)
 - Enforce log structure by introduction of new klog methods that could be used to generate structured logs.
 - Simplify ingestion of logs into third party logging solutions by adding an option to output logs in the JSON format

- Goals/Non Goals
 - Non-Goals
 - We are not replacing currently used logging library (klog) or the way in which it is used

- Contextual Logging
 - Introduction
 - Proposal
 - Goals/Non Goals
 - Deep dive
 - Implementation details
 - Migration details

Introduction



- Proposal
 - Replaces the global logger by passing a logr.Logger instance into functions via a context.
 - Adding extended support for klog to support contextual logging (i.e. adding support to use go-logr/logr apis in klog)

eg:

```
I0404 18:00:02.916429 451895 logger.go:94] "example/myname: runtime" foo="bar" duration="1m0s"
```


- Goals/Non Goals
 - Goals
 - Grant the caller of a function control over logging inside that function, either by passing a logger into the function or by configuring the object that a method belongs to.
 - Provide documentation and helper code for setting up logging in unit tests.
 - Change as few exported APIs as possible

- Goals/Non Goals
 - Non-Goals
 - Remove the klog text output format
 - Deprecate klog

- Structured Logging
 - Deep dive
 - Implementation details
 - Log message structure
 - References to Kubernetes objects
 - Introduce JSON output format in klog
 - Logging configuration
 - Performance
 - Migration details

- Log message structure
 - There could be multiple ways of standardising the logging structure in Kubernetes, the one we agreed to implement in the KEP is to have following logging structure

`<message> <key1>=<value1> <key2>=<value2> ...`

e.g.

```
pod := corev1.Pod{Name: "kubedns", Namespace: "kube-system", ...}  
klog.InfoS("Pod status updated", "pod", klog.KObj(pod), "status", "ready")
```

- References to Kubernetes objects:
 - The idea is to use k8s api first approach to get k8s objects and embed the object related information into the logs
 - Correlate between different kubernetes objects

```
func KObj(obj ObjectMeta) ObjectRef  
func KRef(namespace, name string) ObjectRef
```

```
type ObjectRef struct {  
    Name    string `json:"name"`  
    Namespace string `json:"namespace,omitempty"`  
}
```

- References to Kubernetes objects:

Namespaced objects: `<namespace>/<name>`
e.g. kube-system/kubedns

Non-namespaced objects: `<name>`
e.g. node cluster1-vm-72x33b8p-34jz

e.g.

```
klog.InfoS("Pod status updated", "pod", klog.KObj(pod), "status", "ready")
```

```
·  
·
```

```
klog.ErrorS(err, "Failed to update pod status", "pod", klog.KObj(pod))
```

- Introduce JSON output format in klog:
 - Introduction of new methods to klog library to support JSON.
 - With klog v2 we can take further advantage of this fact and add an option to produce structured logs in JSON format.
- Some pros of using JSON:
 - Broadly adopted by logging libraries with very efficient implementations (zap, zerolog).
 - Out of the box support by many logging backends (Elasticsearch, Stackdriver, BigQuery, Splunk, Open Telemetry)
 - Easily parsable and transformable
 - Existing tools for ad-hoc analysis (jq)

- Introduce JSON output format in klog:

```
klog.InfoS("Pod status updated", "pod", klog.KObj(pod), "status", "ready")

{
  "ts": 1580306777.04728,
  "v": 4,
  "msg": "Pod status updated",
  "pod": {
    "name": "nginx-1",
    "namespace": "default"
  },
  "status": "ready"
}
```

Deep Dive (Implementation details)

- Introduce JSON output format in klog:

```
type Request struct {
```

```
    Method string
```

```
    Timeout int
```

```
    secret string
```

```
    Con *Connection
```

```
}
```

```
req := Request{Method: "GET", Timeout: 30, secret: "pony"}
```

```
klog.InfoS("Request finished", "request", Request)
```

Deep Dive (Implementation details)

- Introduce JSON output format in klog:

```
{  
  "ts": 1580306777.04728,  
  "v": 4,  
  "msg": "Request finished",  
  "request": {  
    "Method": "GET",  
    "Timeout": 30  
  }  
}
```

- Logging configuration
 - Implementation of LoggingConfig structure as part of k8s.io/component-base

introduced flag **--logging-format** values:

- a) text: for text-based logging format (default)
- b) json: for new JSON format

- Performance

- Logging performance with the new implementation. Performance is wrt to log volume and the performance impact.

| logger | time [ns/op] | bytes[B/op] | allocations[alloc/op] |
|------------|--------------|-------------|-----------------------|
| Text Infof | 2252 | 248 | 3 |
| Text InfoS | 2455 | 280 | 3 |
| JSON Infof | 1406 | 19 | 1 |
| JSON InfoS | 319 | 67 | 1 |

- InfoS implementation for text is 9% slower than Infof.
- Kubernetes performance as logging takes less than 2% of overall CPU usage.

- Contextual Logging
 - Deep dive
 - User stories:
 - *"kube-scheduler developer Joan wants to know which pod and which operation and scheduler plugin log messages are associated with"*
 - *"wants to increase the verbosity of the scheduler while it processes a certain pod ("per-flow additional log")"*
 - *...many more*
 - Implementation details
 - Removing the dependency on the global klog logger
 - Extend klog for contextual logging
 - Use migrated structured logs and attach context with it
 - Migration details

Deep Dive (Implementation details)

- Removing the dependency on the global klog logger
 - `klog.ErrorS` -> `logger.Error`
 - `logger` is a `logr.Logger` instance. `klog.Logger` is an alias for that type

Deep Dive (Implementation details)

- Extend klog to be for contextual logging
 - Several new klog functions help with that:
 - klog.FromContext
 - klog.Background
 - klog.TODO

```
// FromContext retrieves a logger set by the caller or, if not set,  
// falls back to the program's global logger (a Logger instance or klog  
// itself).  
func FromContext(ctx context.Context) Logger {  
    if logging.contextualLoggingEnabled {  
        if logger, err := logr.FromContext(ctx); err == nil {  
            return logger  
        }  
    }  
  
    return Background()  
}
```

Deep Dive (Implementation details)

- Extend klog to be for contextual logging
 - Several new klog functions help with that:
 - klog.FromContext
 - klog.Background
 - klog.TODO

```
// Background retrieves the fallback logger. It should not be called before
// that logger was initialized by the program and not by code that should
// better receive a logger via its parameters. TODO can be used as a temporary
// solution for such code.
func Background() Logger {
    if logging.loggerOptions.contextualLogger {
        // Is non-nil because logging.loggerOptions.contextualLogger is
        // only true if a logger was set.
        return logging.logger.Logger
    }

    return klogLogger
}
```

- Extend klog for contextual logging
 - Use migrated structured logs and attach context with it
 - With the helper functions that we added in klog, structured logs can also be transformed to more fine grained contextual logs

- Structured and Contextual Logging migration instructions

We Need Your Help!!

- From SIGs and Working Groups
 - Consider a big volume of PRs, we request you to PTAL :)
 - Adopt the new practice for latest PRs
 - LGTM on us, Approval is on you (mostly)
- From New contributors
 - Join [#wg-structured-logging](#) channel on slack, Join the [mailing list](#)
 - [Regular Meeting](#): Thursdays at 15:30 London-UK (biweekly)
 - [Meeting notes and Agenda](#), [Meeting recordings](#).
 - Pick up migration issues (label wg-structured-logging)
 - Please ask questions!!



KubeCon



CloudNativeCon

Europe 2023

That's it from my side please ask your

Questions

I'll try to answer them :)

Thank you!!