



**KubeCon**



**CloudNativeCon**

**Europe 2023**





KubeCon



CloudNativeCon

Europe 2023

# SIG Autoscaling Updates and Feature Highlights

*Chen Wang, IBM*

*Michele Orlandi, IBM*

*Piotr Betkier, Google*

*Jayant Jain, Google*

*Guy Templeton, Skyscanner*



- Sub-Projects:
  - Balancer (New!)
  - Cluster Autoscaler
  - Horizontal Pod Autoscaling
  - Vertical Pod Autoscaling
  - Addon Resizer

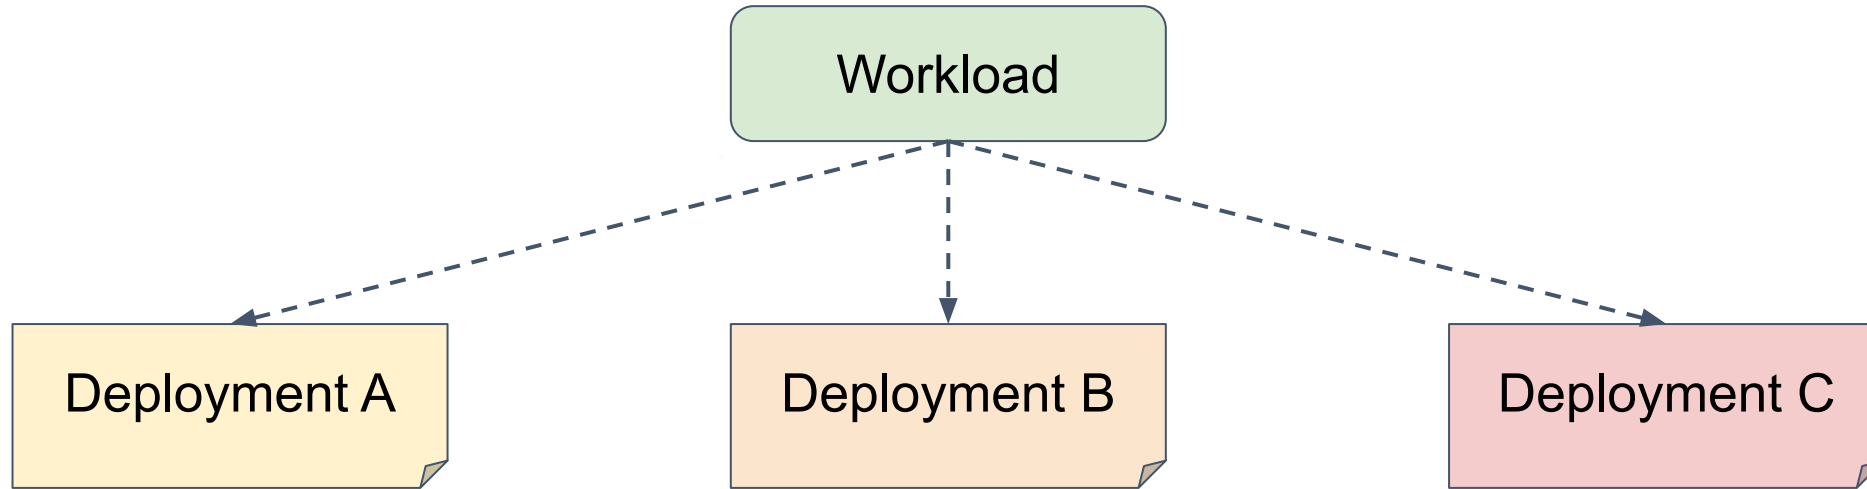
- Other Updates:
  - HPA - Container Resource Scaling - being promoted to beta for 1.27
  - Cluster Autoscaler - Improved release process going forward
  - Vertical Pod Autoscaling - Gathering the work to enable taking advantage of dynamic pod resizing

- Where we need help:
  - Feedback
  - Issue Triage
  - PR Reviewers
  - Contributors

- How to get in touch:
  - [github.com/kubernetes/autoscaler](https://github.com/kubernetes/autoscaler)
  - SIG Meetings - Weekly at 16:00 CET
  - #sig-autoscaling

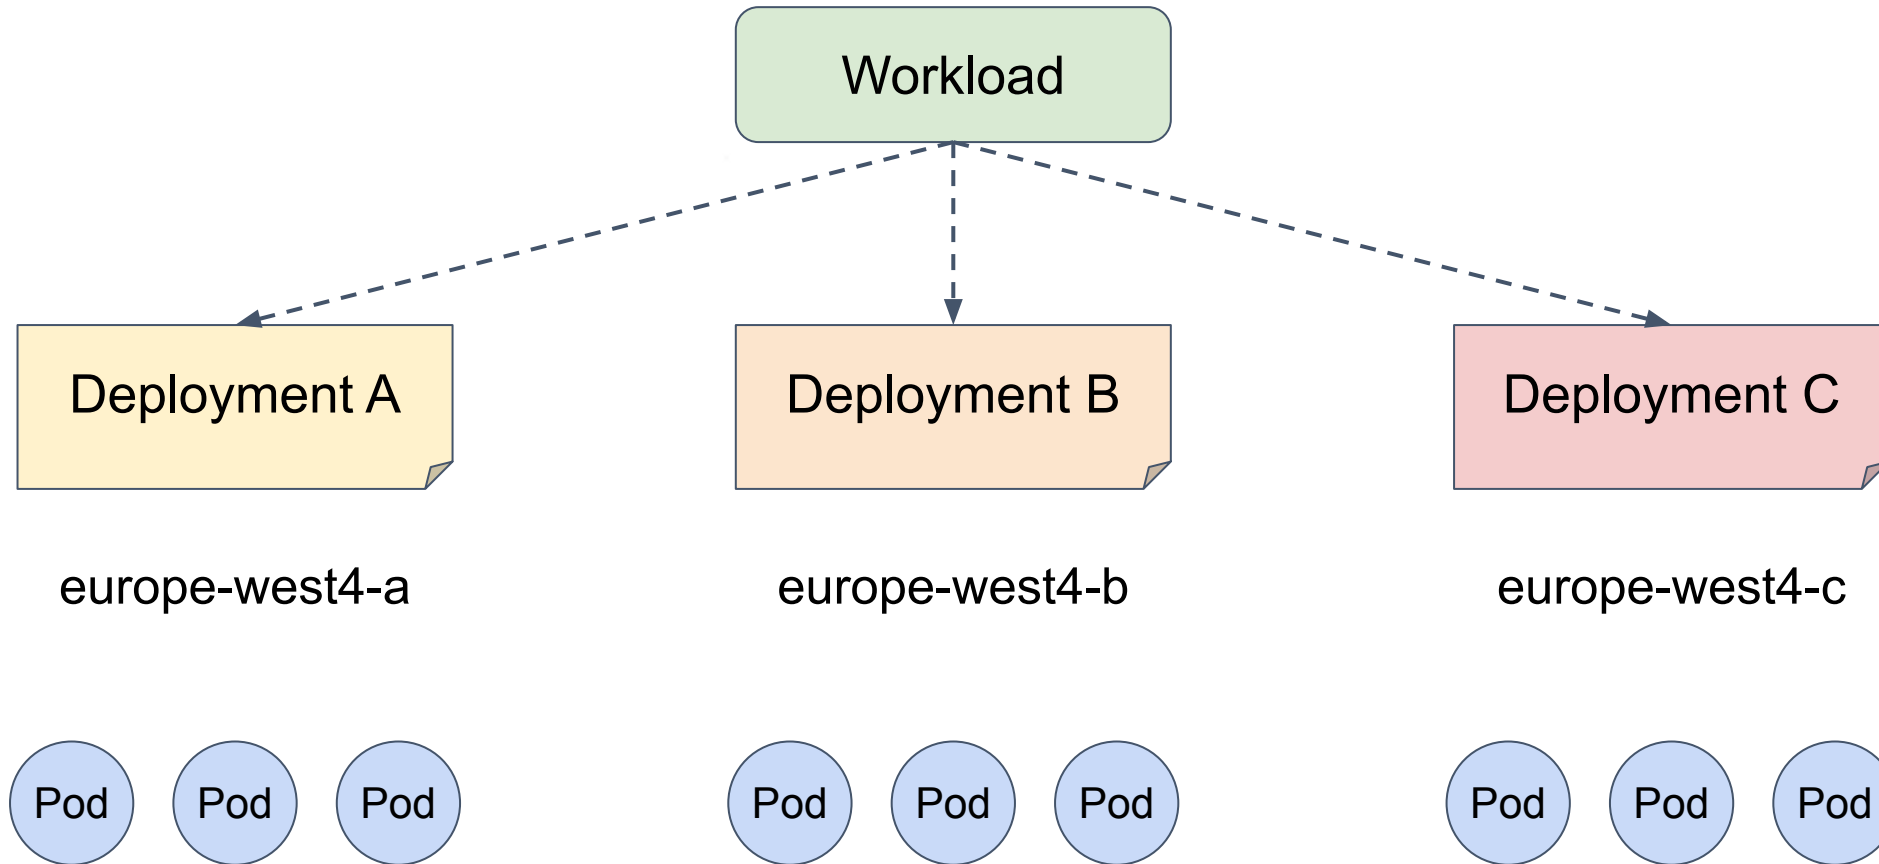
- Balancer allows to
  - Balance Pod distribution across similar deployments
  - Autoscale Pods from similar deployments together
- New resource and its controller
  - Requires installing an optional component Balancer
  - Stage: alpha
  - See [github.com/kubernetes/autoscaler](https://github.com/kubernetes/autoscaler)

# Balancer workload in multiple deployments

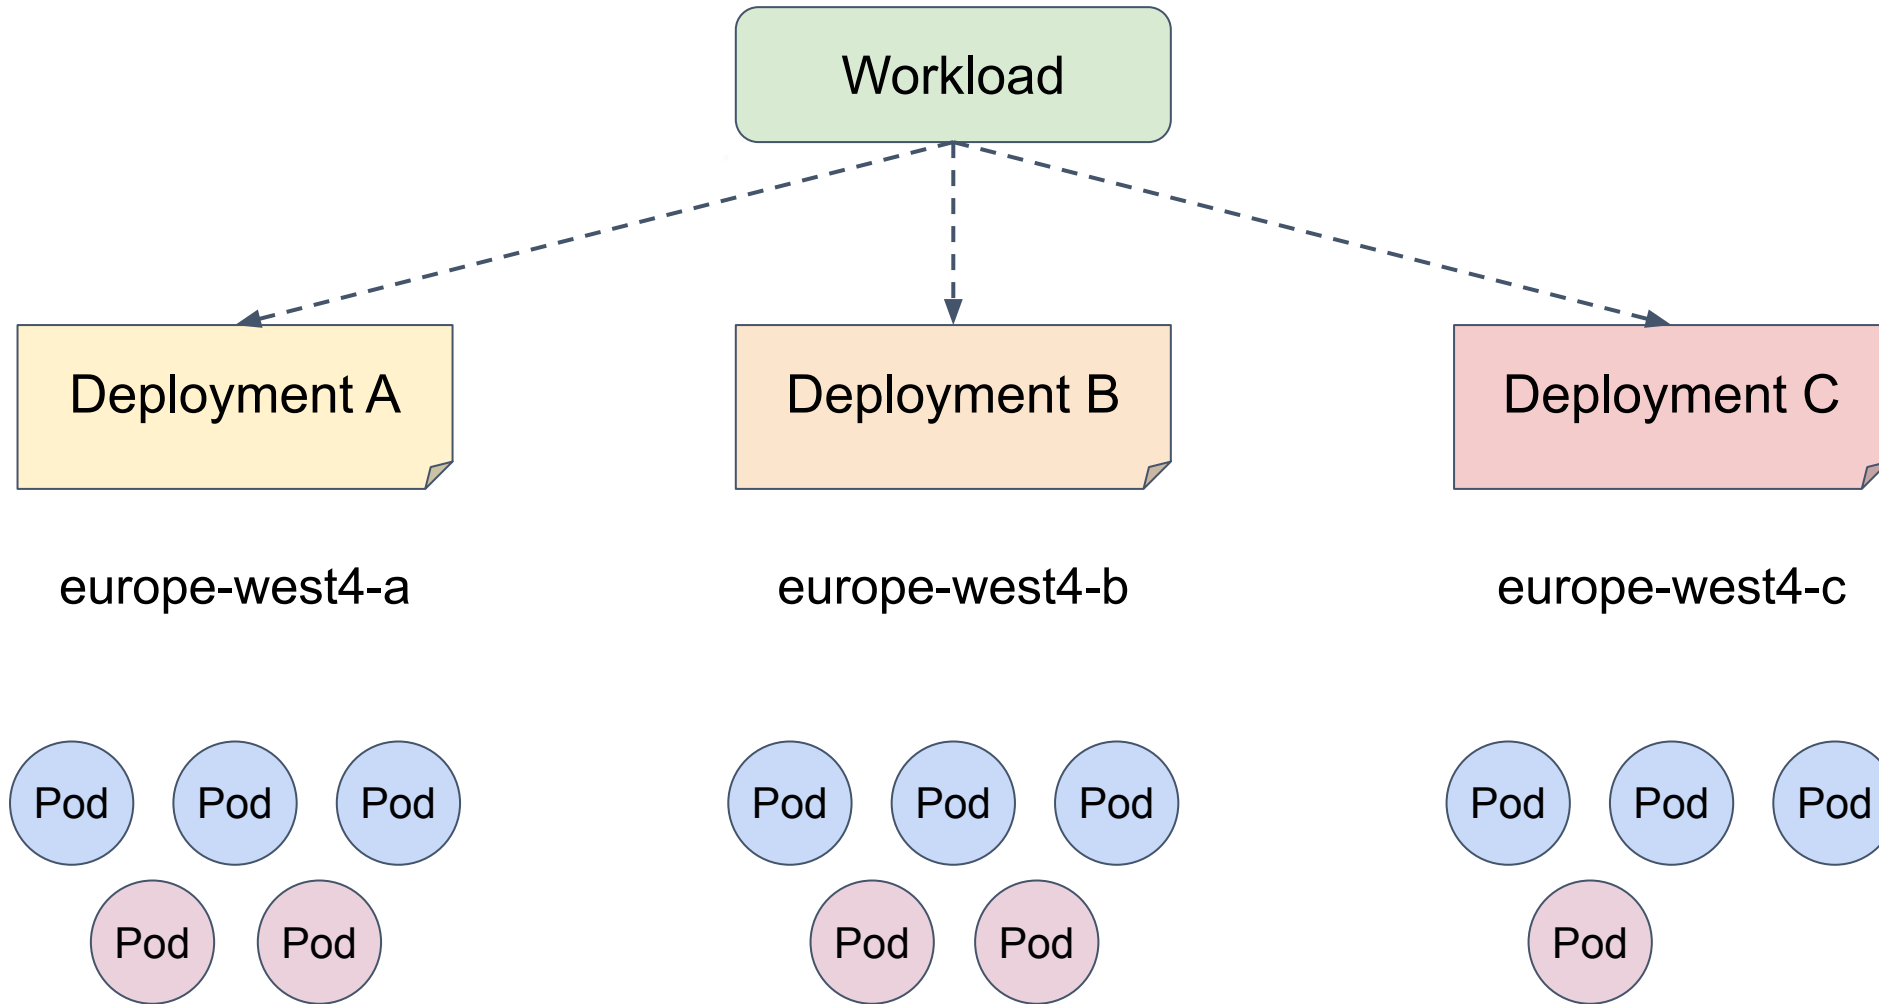




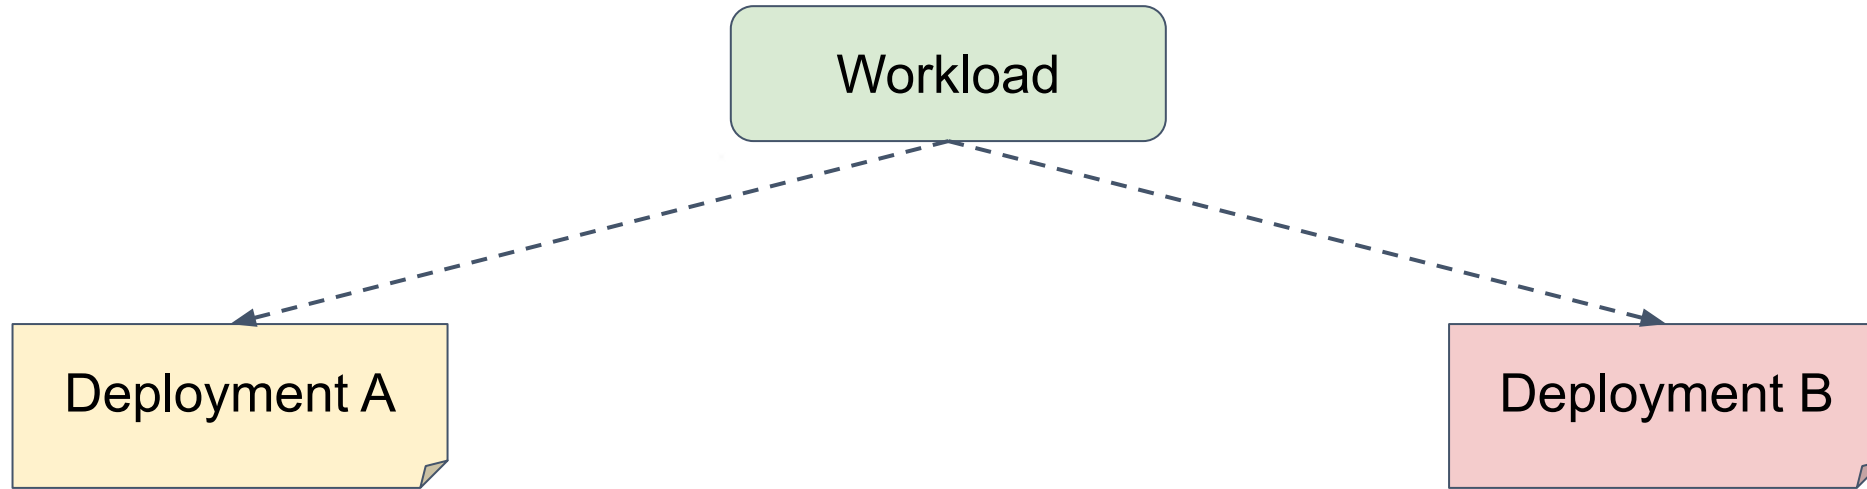
# Balancer evenly distribute pods across zones



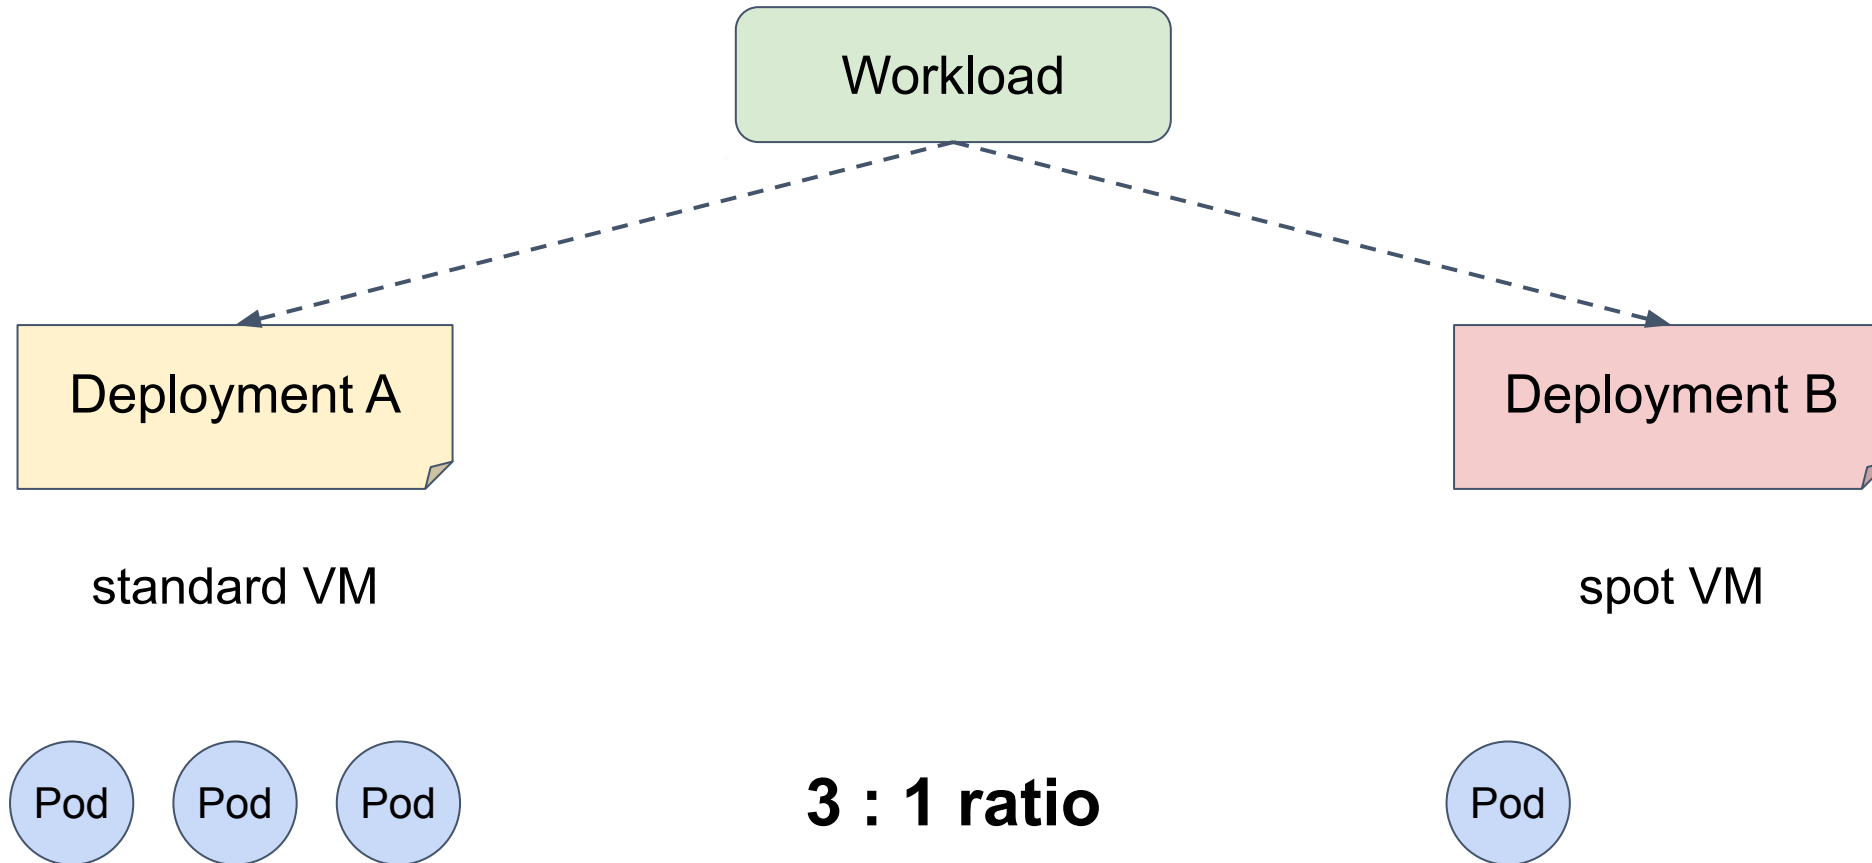
# Balancer evenly distribute pods across zones



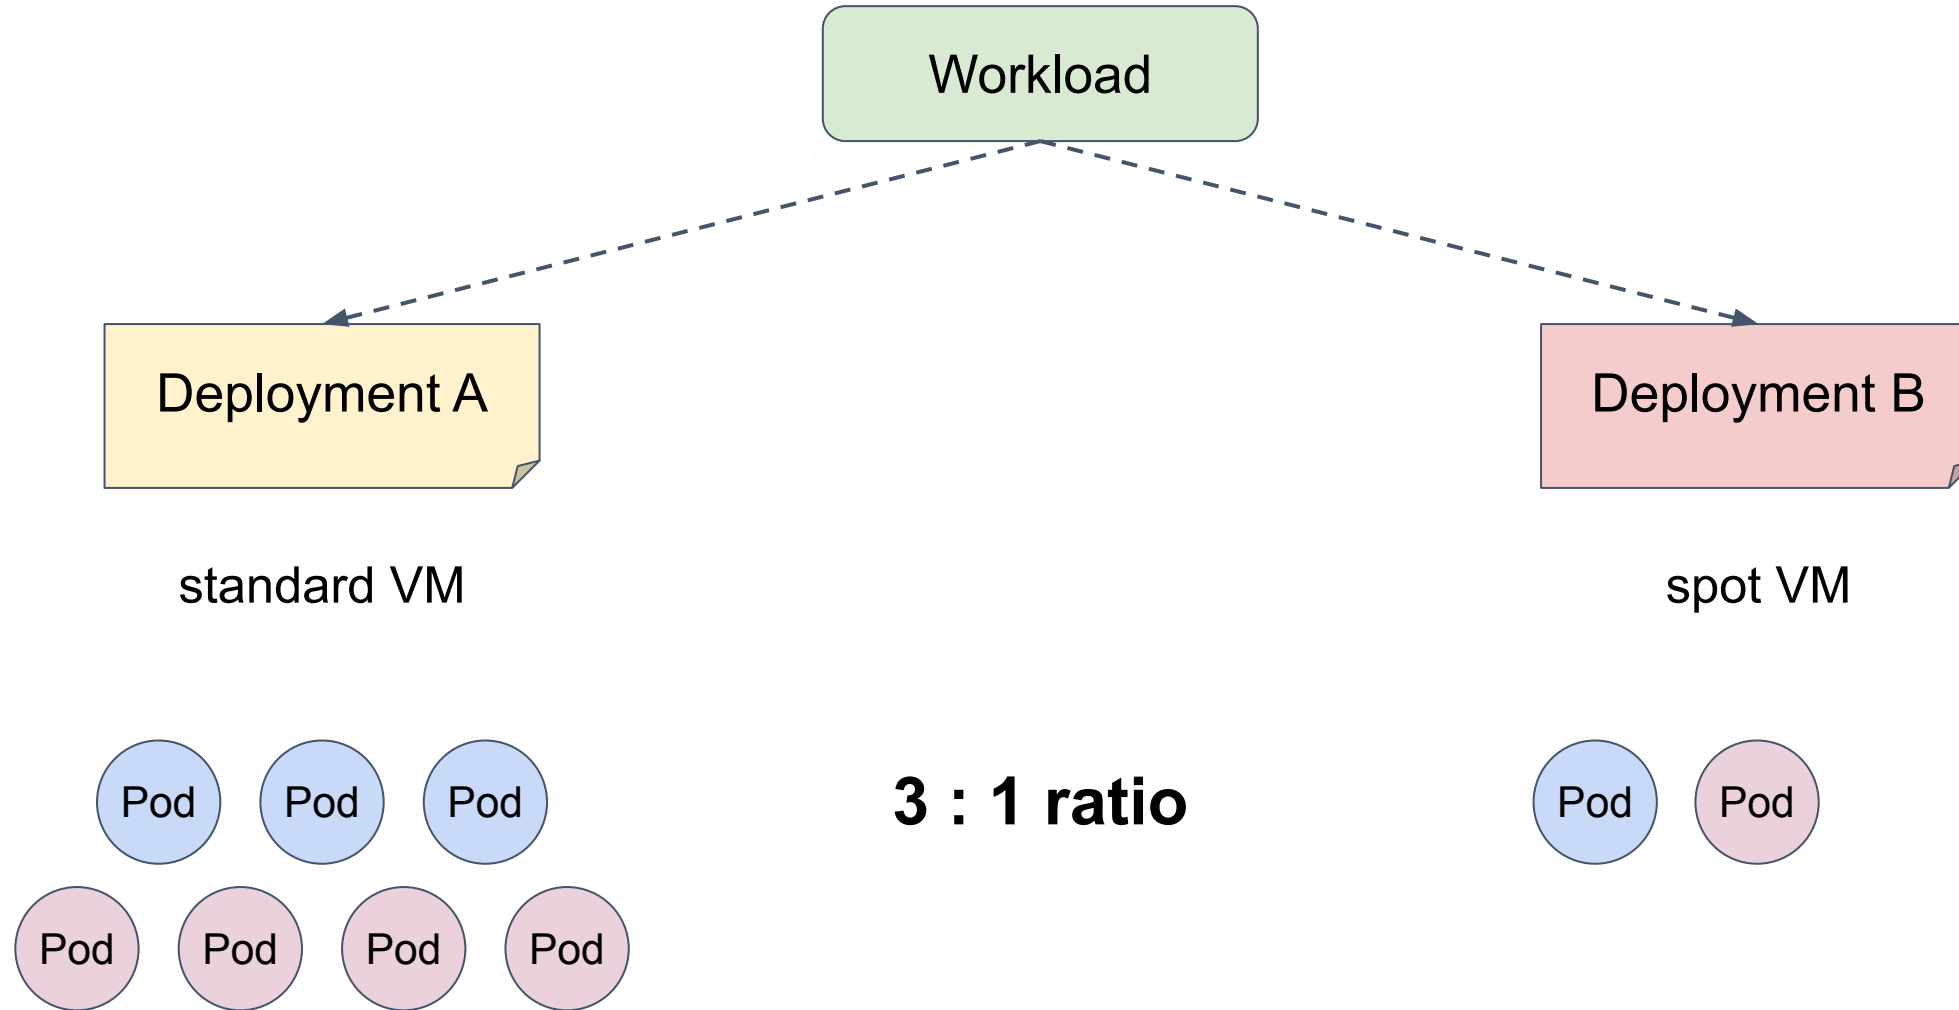
# Balancer maintain pods ratio on spot VMs



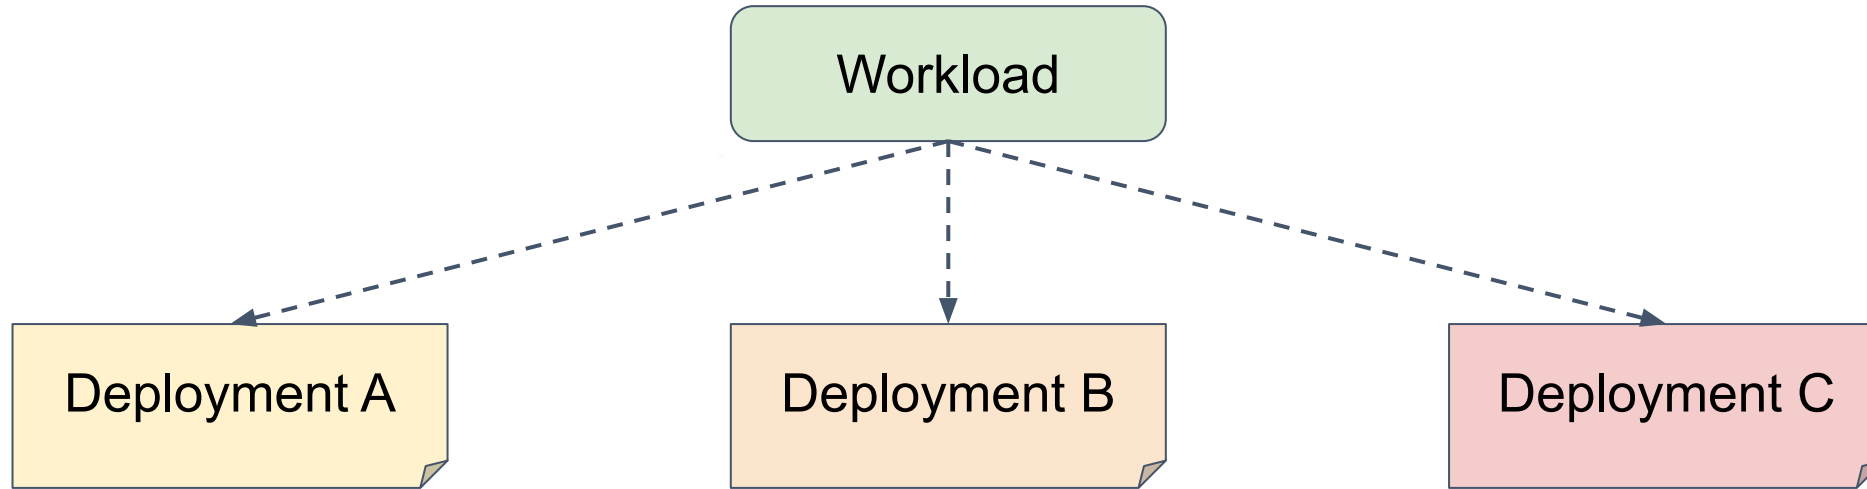
# Balancer maintain pods ratio on spot VMs



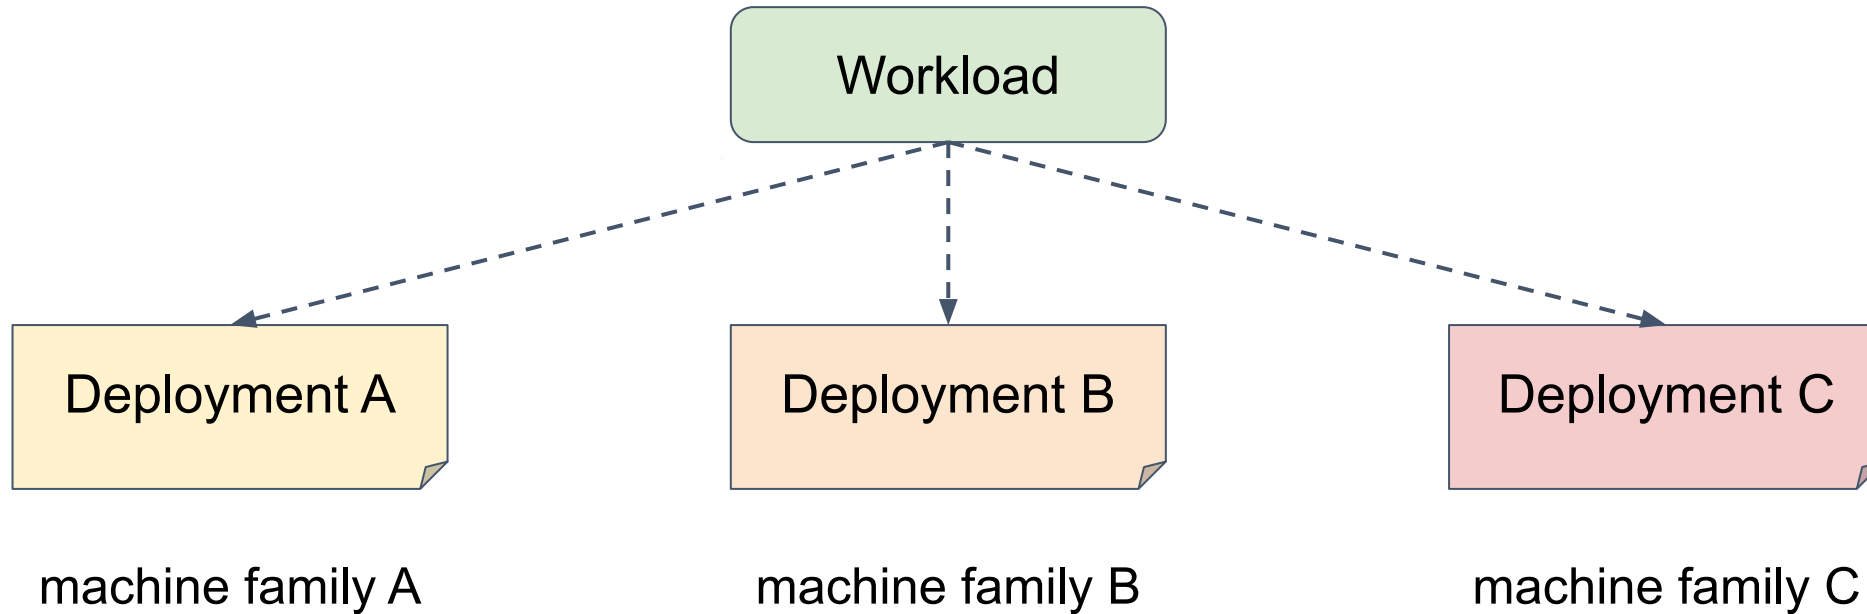
# Balancer maintain pods ratio on spot VMs



# Balancer workload on different machines



# Balancer workload on different machines



# Balancer how to use it



# Balancer how to use it

I want to balance **myapp** between  
deployments **myapp-A** and **myapp-B** in  
3:1 ratio.

# Balancer how to use it

I want to balance **myapp** between  
deployments **myapp-A** and **myapp-B** in  
3:1 ratio.

```
apiVersion: balancer.x-k8s.io/v1alpha1
kind: Balancer
metadata:
  name: myapp-balancer
spec:
  selector:
    app: myapp
  targets:
    - name: myapp-A
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-A
      minReplicas: 1
      maxReplicas: 10
    - name: myapp-B
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-B
      minReplicas: 0
      maxReplicas: 10
  policy:
    type: proportional
    proportions:
      targetRatios: [myapp-A: 3, myapp-B: 1]
```

# Balancer how to use it

I want to balance **myapp** between  
deployments **myapp-A** and **myapp-B** in  
3:1 ratio.

```
apiVersion: balancer.x-k8s.io/v1alpha1
kind: Balancer
metadata:
  name: myapp-balancer
spec:
  selector:
    app: myapp
  targets:
    - name: myapp-A
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-A
      minReplicas: 1
      maxReplicas: 10
    - name: myapp-B
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-B
      minReplicas: 0
      maxReplicas: 10
  policy:
    type: proportional
    proportions:
      targetRatios: [myapp-A: 3, myapp-B: 1]
```

# Balancer how to use it

I want to balance **myapp** between  
deployments **myapp-A** and **myapp-B** in  
3:1 ratio.

```
apiVersion: balancer.x-k8s.io/v1alpha1
kind: Balancer
metadata:
  name: myapp-balancer
spec:
  selector:
    app: myapp
  targets:
    - name: myapp-A
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-A
      minReplicas: 1
      maxReplicas: 10
    - name: myapp-B
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-B
      minReplicas: 0
      maxReplicas: 10
  policy:
    type: proportional
    proportions:
      targetRatios: [myapp-A: 3, myapp-B: 1]
```

# Balancer how to use it

I want to balance **myapp** between  
deployments **myapp-A** and **myapp-B** in  
3:1 ratio.

With horizontal pod autoscaling

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: myapp
spec:
  scaleTargetRef:
    apiVersion: balancer.x-k8s.io/v1beta1
    kind: Balancer
    name: myapp
  minReplicas: 1
  maxReplicas: 20
  metrics:
    - (...)
```

```
apiVersion: balancer.x-k8s.io/v1alpha1
kind: Balancer
metadata:
  name: myapp-balancer
spec:
  selector:
    app: myapp
  targets:
    - name: myapp-A
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-A
      minReplicas: 1
      maxReplicas: 10
    - name: myapp-B
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-B
      minReplicas: 0
      maxReplicas: 10
  policy:
    type: proportional
    proportions:
      targetRatios: [myapp-A: 3, myapp-B: 1]
```

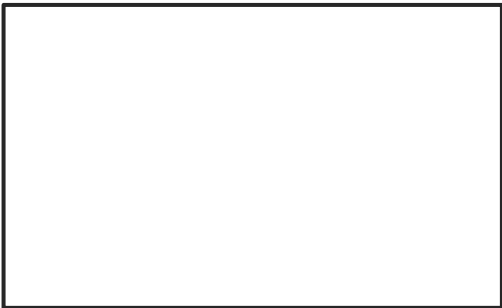
# Balancer how to use it

# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

Deployment A

Deployment B



# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.



Deployment A

Deployment B

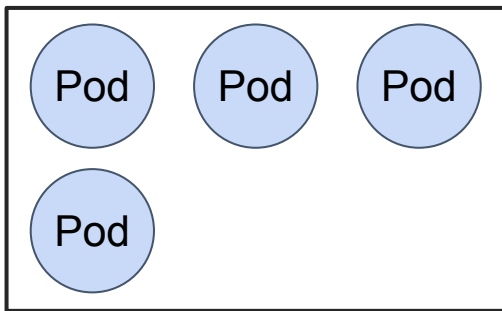


# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

Deployment A

Deployment B

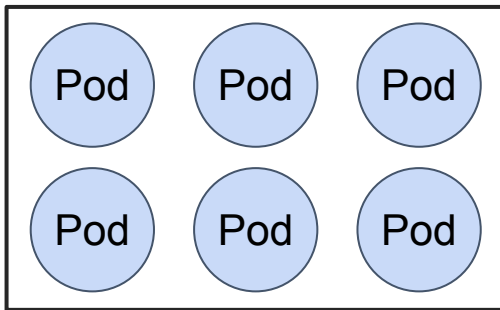


# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

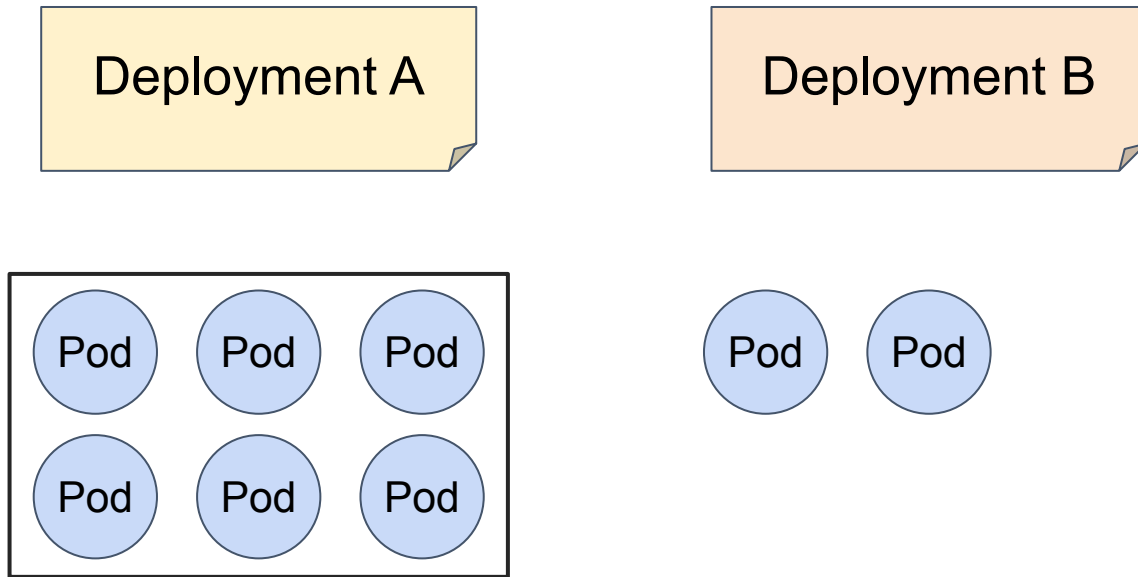
Deployment A

Deployment B



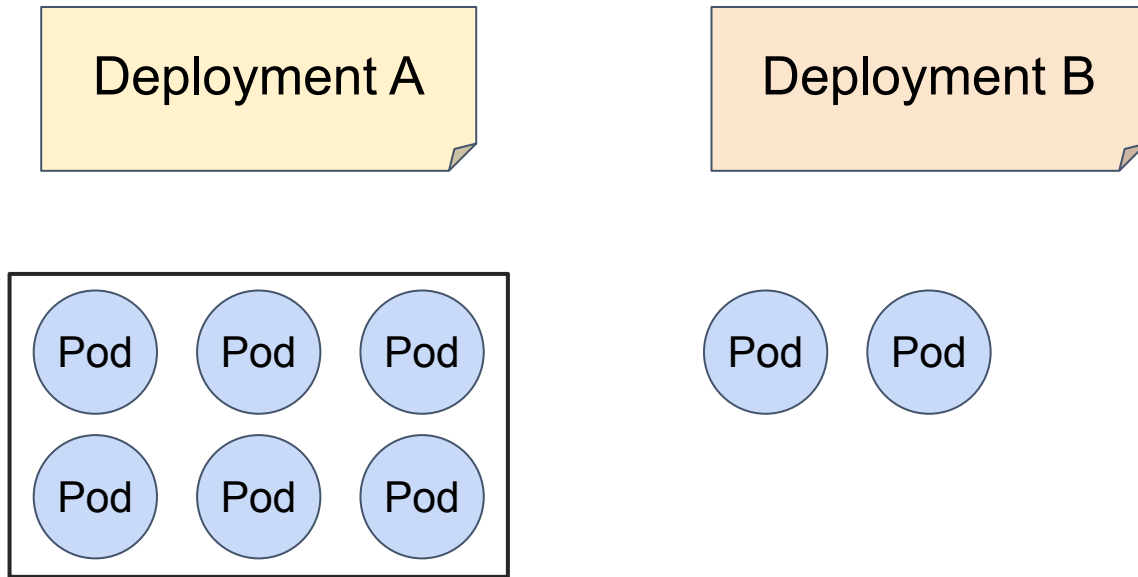
# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.



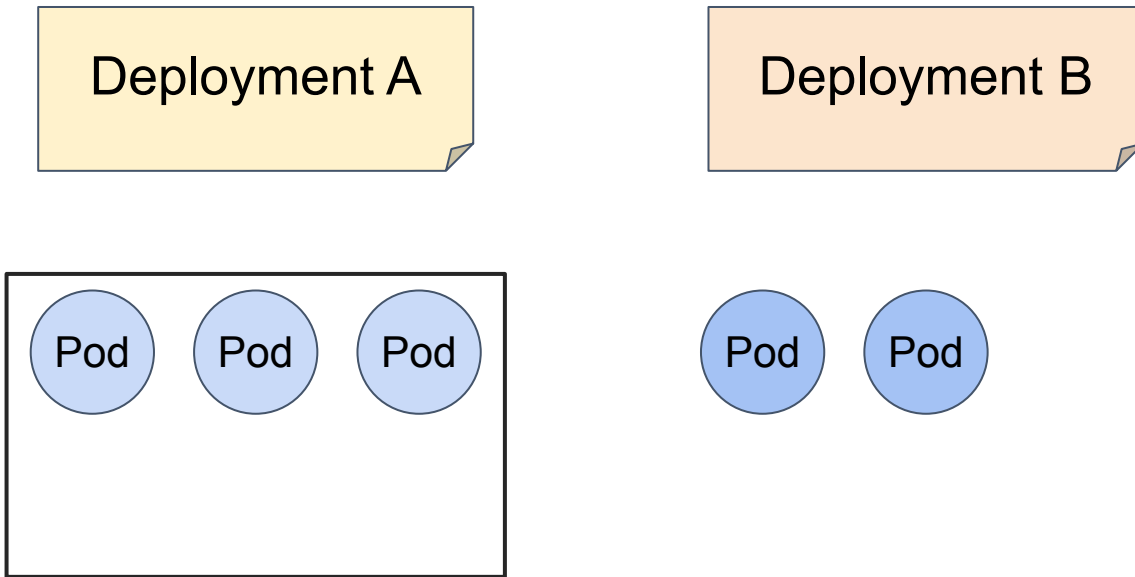
# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.



# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

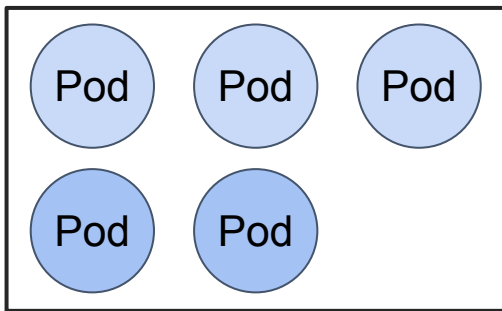


# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

Deployment A

Deployment B

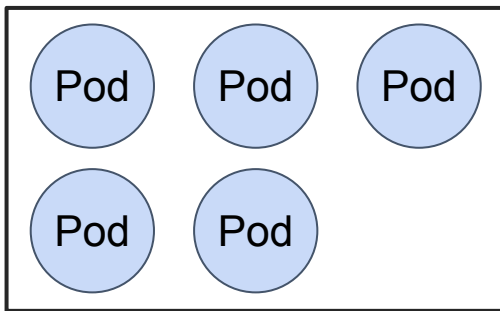


# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

Deployment A

Deployment B

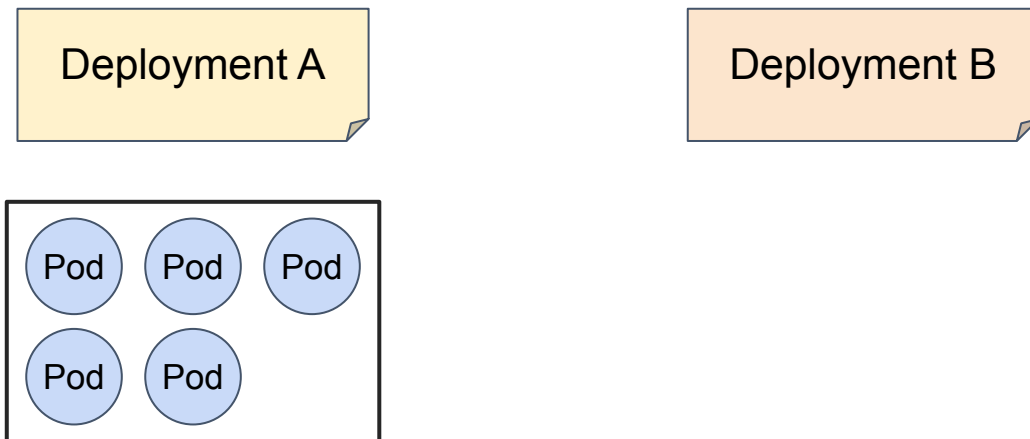


```
apiVersion: balancer.x-k8s.io/v1alpha1
kind: Balancer
metadata:
  name: myapp-balancer
spec:
  selector:
    app: myapp
  targets:
    - name: myapp-A
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-A
      minReplicas: 1
      maxReplicas: 6
    - name: myapp-B
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-B
  policy:
    type: priority
    priority:
      targetOrder: [myapp-A, myapp-B]
```

# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

Also fallback to **myapp-B** if cannot start Pods in **myapp-A**.

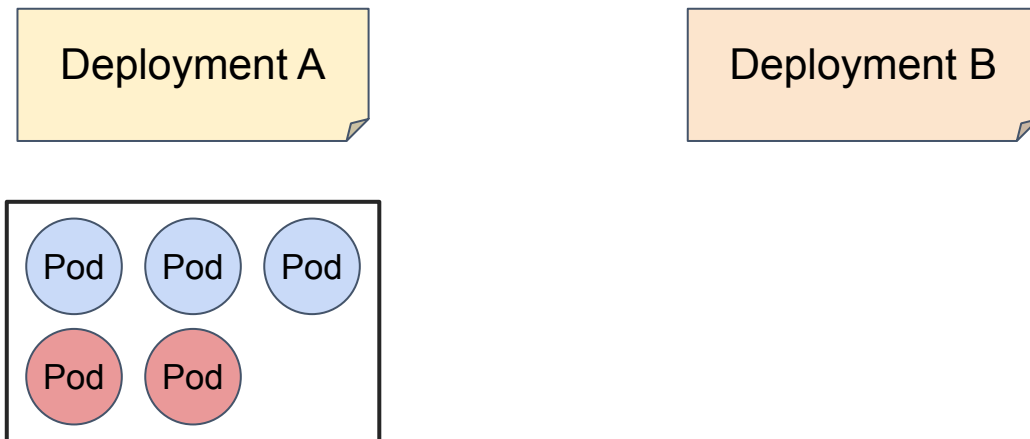




# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

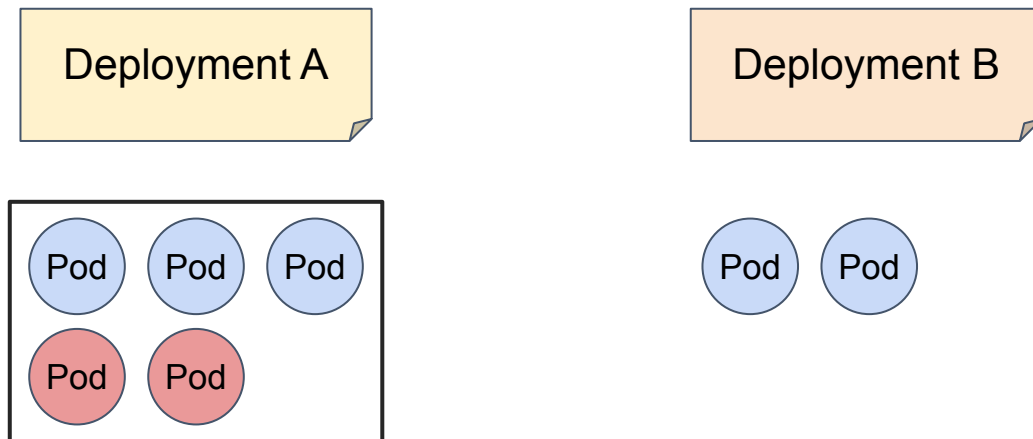
Also fallback to **myapp-B** if cannot start Pods in **myapp-A**.



# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

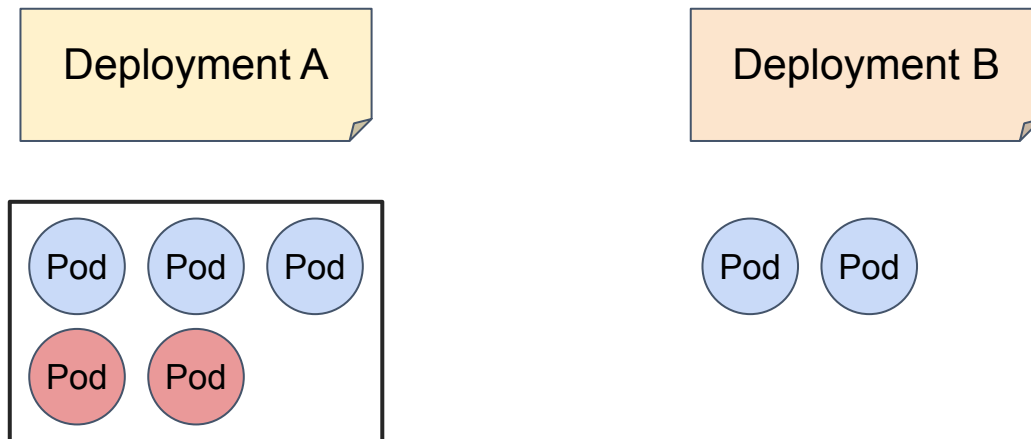
Also fallback to **myapp-B** if cannot start Pods in **myapp-A**.



# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

Also fallback to **myapp-B** if cannot start Pods in **myapp-A**.

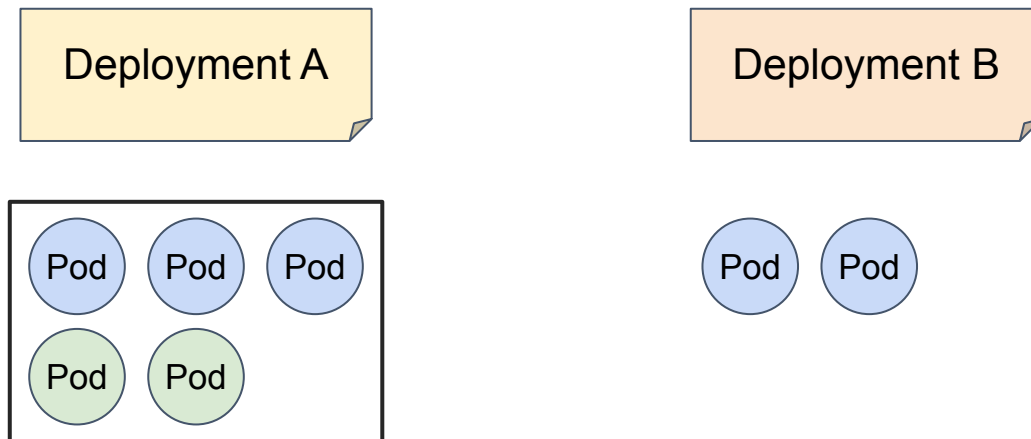


```
apiVersion: balancer.x-k8s.io/v1alpha1
kind: Balancer
metadata:
  name: myapp-balancer
spec:
  selector:
    app: myapp
  targets:
    - name: myapp-A
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-A
      minReplicas: 1
      maxReplicas: 10
    - name: myapp-B
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-B
  policy:
    type: priority
    priority:
      targetOrder: [myapp-A, myapp-B]
    fallback:
      enabled: true
      startupTimeout: 5min
```

# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

Also fallback to **myapp-B** if cannot start Pods in **myapp-A**.

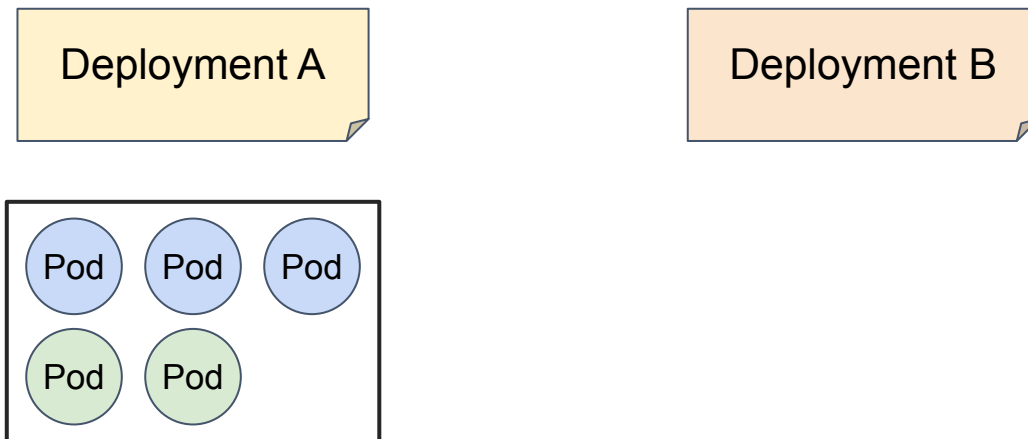


```
apiVersion: balancer.x-k8s.io/v1alpha1
kind: Balancer
metadata:
  name: myapp-balancer
spec:
  selector:
    app: myapp
  targets:
    - name: myapp-A
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-A
      minReplicas: 1
      maxReplicas: 10
    - name: myapp-B
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-B
  policy:
    type: priority
    priority:
      targetOrder: [myapp-A, myapp-B]
    fallback:
      enabled: true
      startupTimeout: 5min
```

# Balancer how to use it

I want to run **myapp** in deployment **myapp-A**, but overflow to **myapp-B** if max replicas in **myapp-A** reached.

Also fallback to **myapp-B** if cannot start Pods in **myapp-A**.



```
apiVersion: balancer.x-k8s.io/v1alpha1
kind: Balancer
metadata:
  name: myapp-balancer
spec:
  selector:
    app: myapp
  targets:
    - name: myapp-A
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-A
      minReplicas: 1
      maxReplicas: 10
    - name: myapp-B
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: myapp-B
  policy:
    type: priority
    priority:
      targetOrder: [myapp-A, myapp-B]
    fallback:
      enabled: true
      startupTimeout: 5min
```

# Balancer summary

- Use Balancer to
  - Balance Pod distribution across similar deployments
  - Autoscale Pods from similar deployments together
- Learn more at [github.com/kubernetes/autoscaler](https://github.com/kubernetes/autoscaler)
- Give feedback on GitHub issues!

## *What are we trying to solve?*

- CA logs information about the decisions taken, but we don't log what data these decisions are based on as CA internally simulates the behaviour of the entire cluster when making decisions. When something goes wrong, we usually need to understand how these decisions were taken.
- RCA of the issue usually takes time, and the internal state of CA changes when we mitigate the issue, thereby making it more difficult to debug the issue.

### Debugging Snapshotter:

It's a tool to visualize the internal state of cluster-autoscaler at a point in time to help debug autoscaling issues.

## *Common use-cases*

- Scale up and scale down not working, with special focus on scale from zero nodes.
- Mismatch between scheduler and cluster-autoscaler decisions.
- Mismatch in resource availability on the node.



## *What data is captured?*

```
// DebuggingSnapshotImpl is the struct used to collect all the data to be output.
// Please add all new output fields in this struct. This is to make the data
// encoding/decoding easier as the single object going into the decoder
type DebuggingSnapshotImpl struct {
    NodeList                []*ClusterNode    `json:"NodeList"`
    UnscheduledPodsCanBeScheduled []*v1.Pod          `json:"UnscheduledPodsCanBeScheduled"`
    Error                    string             `json:"Error,omitempty"`
    StartTimestamp           time.Time          `json:"StartTimestamp"`
    EndTimestamp             time.Time          `json:"EndTimestamp"`
    TemplateNodes            map[string]*ClusterNode `json:"TemplateNodes"`
}
```

[https://github.com/kubernetes/autoscaler/blob/master/cluster-autoscaler/debuggingsnapshot/debugging\\_snapshot.go#L63](https://github.com/kubernetes/autoscaler/blob/master/cluster-autoscaler/debuggingsnapshot/debugging_snapshot.go#L63)

## *How does it work?*

- Captures a set of internal fields (when it receives a request) and returns a well-formed JSON.
- After receiving the request, in the following cluster-autoscaler loop, snapshotter will snapshot the current state, limited to a single loop (no cross-loop data).
- An HTTP request blocks for the duration of the snapshot generation, and returns the JSON as a HTTP response to the request.
- SSH to server running the leader cluster-autoscaler, and curl `http://127.0.0.1:8085/snapshotz`

# Debugging Snapshotter

Demo

# Debugging Snapshotter quick-start

- Enable the snapshotter by adding the following flag to cluster-autoscaler manifest

```
--debugging-snapshot-enabled=true
```

- Available in cluster-autoscaler versions 1.24+
- Detailed instructions in cluster-autoscaler FAQs.  
<https://github.com/kubernetes/autoscaler/blob/master/cluster-autoscaler/FAQ.md#how-to>
- Looking forward to hearing feedback! (and possible extensions to make life easier)

Session QR Codes will be  
sent via email before the event

Please scan the QR Code above  
to leave feedback on this session