



**KubeCon**



**CloudNativeCon**

**Europe 2023**





KubeCon



CloudNativeCon

Europe 2023

# The path to self contained CRDs

*Cici Huang, Google*



# Things to be covered?

- The Journey of CRD in Kubernetes
- The **Common Expression Language** (CEL)
- Make CRD More self contained
- The Power to extend - Policy Enforcement in Admission Chain
- Q & A



KubeCon



CloudNativeCon

Europe 2023

# The Journey of CRD



# The Journey of CRD

- Start as ThirdPartyResource
- Beta in 1.7; Stable in 1.16
- Extension of the Kubernetes API



# The Journey of CRD

- Trying to make CRD as good as built-in types
  - Versioning
  - Subresources
  - OpenAPI schema
  - Structural schema
  - Defaulting
  - Pruning
  - .....



# Validation is critical

If you don't validate the data comes in, **things will break in a way hard to reason!**

- Immutability
- Cross field checks
- Mutually exclusive
- Specific format of a field
- ...



## Before CEL

- Build-in Validation
  - CRD structural schemas
  - OpenAPIV3 validation
    - Format, Regex, Range/Size limit
- Validating Admission Webhook





# The Price of Webhooks

- Development complexity
  - Core logic + monitoring + alerting
  - Packaging/Releasing
- Operational complexity
  - Mis-configure
  - Latency added
  - Upgrade/rollback



Things people wanna do with CRD validation are **simple!**

- 1 Ensure a CRD field is immutable.
- 2 Require that field x has a specific format.
- 3 Cross field validations.

**Can we use something simpler than webhooks?**



KubeCon



CloudNativeCon

Europe 2023

# The Common Expression Language (CEL)



CEL Expression	Purpose
<code>self.minReplicas &lt;= self.replicas &amp;&amp; self.replicas &lt;= self.maxReplicas</code>	Validate that the three fields defining replicas are ordered appropriately
<code>self.components['Widget'].parts[2].isEnabled</code>	Access a deeply nested field.
<code>self.created + self.ttl &lt; self.expired</code>	Validate that 'expired' date is after a 'create' date plus a 'ttl' duration

The **Common Expression Language (CEL)** implements common semantics for expression evaluation, enabling different applications to more easily interoperate.

**Keep it small & fast.**

CEL evaluates in linear time, is mutation free, and not Turing-complete. This limitation is a feature of the language design, which allows the implementation to evaluate orders of magnitude faster than equivalently sandboxed JavaScript.

**Make it extensible.**

CEL is designed to be embedded in applications, and allows for extensibility via its context which allows for functions and data to be provided by the software that embeds it.

**Developer-friendly.**

The language is approachable to developers. The initial spec was based on the experience of developing Firebase Rules and usability testing many prior iterations.

The library itself and accompanying toolings should be easy to adopt by teams that seek to integrate CEL into their platforms.

# CEL Limitations

No `for/while/...`

Instead use comprehension "macros":

`all()`

`exists()`

`exists_one()`

`map()`

`filter()`

No `if/else..`

Instead use the ternary operator:

`<condition> ? <ifTrue> : <ifFalse>`

- CEL "stdlib"
- CEL extended string function library(regex, contains, ...)
- Kubernetes CEL library
  - list library
  - more regex
  - URLs support

# Learn more:

<https://kubernetes.io/docs/reference/using-api/cel/>

[github.com/google/cel-spec](https://github.com/google/cel-spec)





KubeCon



CloudNativeCon

Europe 2023

# Make CRDs more self contained

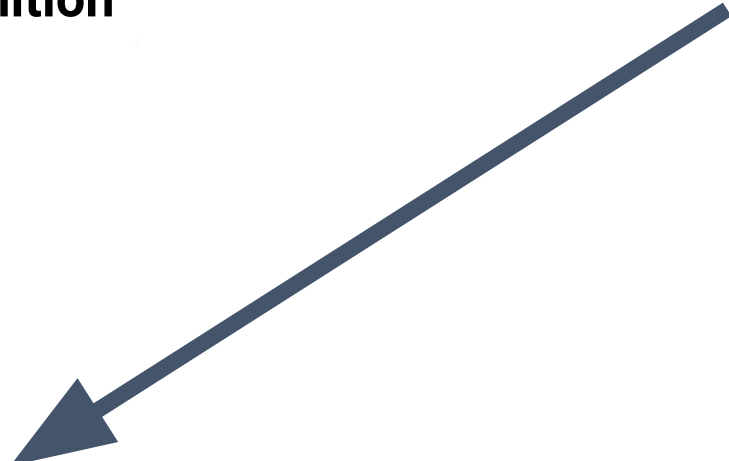


x-kubernetes-validations

# CRD Validation Rules

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
spec:
...
openAPIV3Schema:
  type: object
  properties:
    spec:
      type: object
      x-kubernetes-validations:
        - rule: "self.replicas <=
self.maxReplicas"
      properties:
        ...
        replicas:
          type: integer
        maxReplicas:
          type: integer
```

Can be put anywhere  
under openAPIV3Schema.



# CRD Validation Rules

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
spec:
```

```
...
```

```
  openAPIV3Schema:
```

```
    type: object
```

```
    properties:
```

```
      spec:
```

```
        type: object
```

```
        x-kubernetes-validations:
```

```
          - rule: "self.replicas <= self.maxReplicas"
```

```
        properties:
```

```
          ...
```


```
          replicas:
```

```
            type: integer
```

```
          maxReplicas:
```

```
            type: integer
```

"self" is a CEL variable. It provides access to values, scoped to current schema.



# CRD Validation Rules

apiVersion: apiextensions.k8s.io/v1  
kind: CustomResourceDefinition  
spec:

...

openAPIV3Schema:

type: object

properties:

spec:

type: object

required: [foo]

properties:

foo:

type: integer

x-kubernetes-validations:

- rule: "self == **oldSelf**"

"oldSelf" is another CEL variable.

If you use it, the CEL expression is called a "transition rule".

This transition rule enforces immutability of the foo field.

# CRD Validation Rules

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
spec:
```

```
...
```

```
  openAPIV3Schema:
```

```
    type: object
```

```
    properties:
```

```
      spec:
```

```
        type: object
```

```
        x-kubernetes-validations:
```

```
          - rule: "!has(self.foo) ||
self.foo.startsWith('kube')"
```

```
        properties:
```

```
          ...
```

```
          foo:
```

```
            type: string
```

```
            x-kubernetes-validations:
```

```
              - rule: "self.startsWith('kube')"
```



Same validation but  
scoped differently.

# CRD Validation Rules

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
spec:
```

```
...
```

```
  openAPIV3Schema:
```

```
    type: object
```

```
    properties:
```

```
      spec:
```

```
        type: object
```

```
        x-kubernetes-validations:
```

```
          - rule: "self.replikas <= self.maxReplicas"
```

```
        properties:
```

```
          ...
```

```
          replicas:
```

```
            type: integer
```

```
          maxReplicas:
```

```
            type: integer
```

Type checking while  
creating or updating  
your CRDs:

compilation failed: ERROR:  
<input>:1:5: undefined field 'replikas'

# CRD Validation Rules

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
spec:
```

```
...
```

```
  openAPIV3Schema:
```

```
    type: object
```

```
    properties:
```

```
      spec:
```

```
        type: object
```

```
        x-kubernetes-validations:
```

```
          - rule: "self.replicas <= self.maxReplicas"
```

```
            message: "replicas must be no greater than 3"
```

```
            messageExpression: "replicas must be no greater  
than ' + string(self.maxReplicas)"
```

```
        properties:
```

```
          ...
```

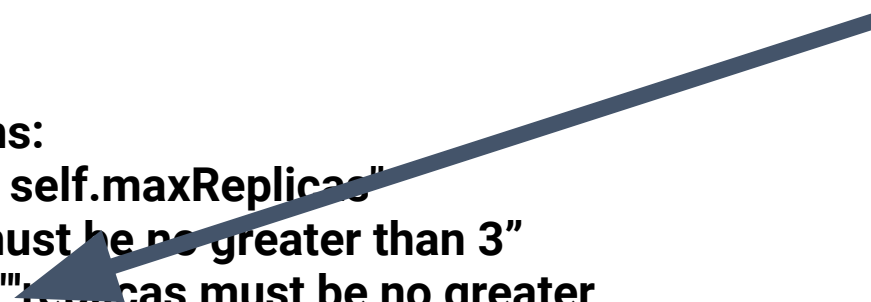
```
          replicas:
```

```
            type: integer
```

```
          maxReplicas:
```

```
            type: integer
```

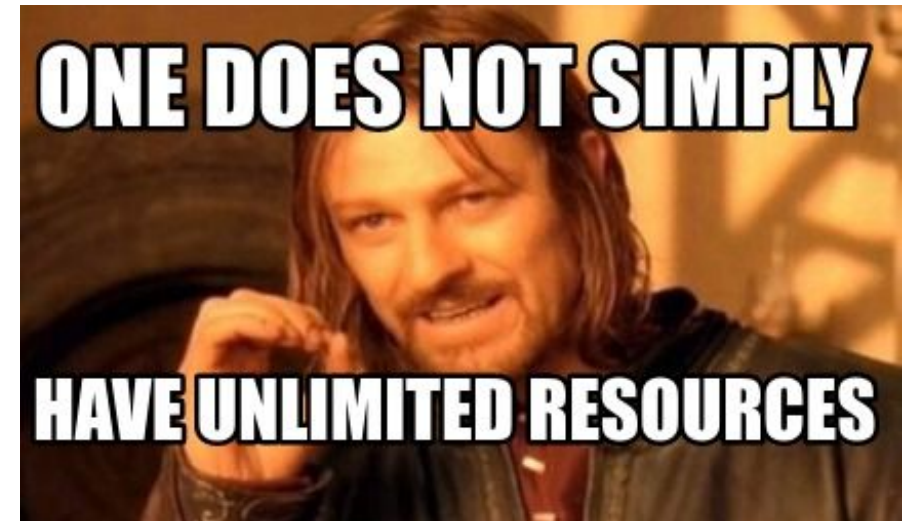
Support CEL expression  
in validation failure  
message.





## Runtime Constraints

- Estimated Cost Limit
- Runtime Cost Limit
- Context Cancelation



Learn More:

<https://kubernetes.io/docs/tasks/extend-kubernetes/custom-resources/custom-resource-definitions/#resource-use-by-validation-functions>

# One more step...



Conversion is required when

- custom resource is requested in a different version than stored version.
- Watch is created in one version but the changed object is stored in another version.
- custom resource PUT request is in a different version than storage version.

**Webhook Conversion** is in Stable since 1.16, but...

# Webhook conversion

apiVersion: apiextensions.k8s.io/v1

kind: CustomResourceDefinition

spec:

...

**conversion:**

# a Webhook strategy instruct API server to call an external  
webhook for any conversion between custom resources.

strategy: **Webhook**

webhook:

# conversionReviewVersions indicates what  
ConversionReview versions are understood/preferred by the  
webhook.

# The first version in the list understood by the API server is  
sent to the webhook.

# The webhook must respond with a ConversionReview  
object in the same version it received.

conversionReviewVersions: ["v1","v1beta1"]

clientConfig:

service:

namespace: default

name: example-conversion-webhook-server

path: /crdconvert

caBundle: "Ci0tLS0tQk...<base64-encoded PEM  
bundle>...tLS0K"

- Hard to configure
- Involve Webhook
- Perform bad

## Next Step: CRD Conversion with CEL

- Self Contained
- Fast
- Sufficient subs conversion we















mentation  
8s.io/v1  
ition

S

Name  
toField: .spec.name.first  
transformation: self

## CRD is **more self contained** than built-in types

	built-in	CRD
API declaration		
OpenAPI schema		
Versioning		
Subresources		
Declarative validation		
Declarative conversion		

## Catch up on Native Types: Declarative Validation







KubeCon



CloudNativeCon

Europe 2023

Extend the power to a larger  
and more impactful area: **policy  
enforcement.**



# Policy Enforcement in Kubernetes

- Internal support such like Pod

Security Admission



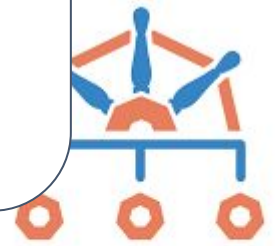
Open Policy Agent

- Policy

Gate

- Build

**Validating Admission Policy**



**Kyverno**

## Roles involved in policy management



Policy Author

Concerns:

- Correctness of policy
- Reuse/Configurability



Usually not the  
same person

Often in  
different  
organizations



Cluster Administrator

Concerns:

- Configuring policy for organization
- Operability, esp. Safe rollouts

# Validating Admission Policy



apiVersion: admissionregistration.k8s.io/v1alpha1  
kind: ValidatingAdmissionPolicy  
metadata:

name: "replicalimit-policy.example.com"

spec:

matchConstraints:

resourceRules:

- apiGroups: ["apps"]
- apiVersions: ["v1"]
- operations: ["CREATE", "UPDATE"]
- resources: ["deployments"]

validations:

- expression: "object.spec.replicas <= int(params.data.maxReplicas)"

paramKind:

apiVersion: v1

kind: ConfigMap

apiVersion: admissionregistration.k8s.io/v1alpha1  
kind: ValidatingAdmissionPolicyBinding  
metadata:

name: "replicalimit-binding-test.example.com"

spec:

policyName: "replicalimit-policy.example.com"

paramRef:

name: "replica-limit-test.example.com"  
namespace: "default"

matchResources:

namespaceSelector:

matchLabels:

environment: test

validationActions: [Deny]

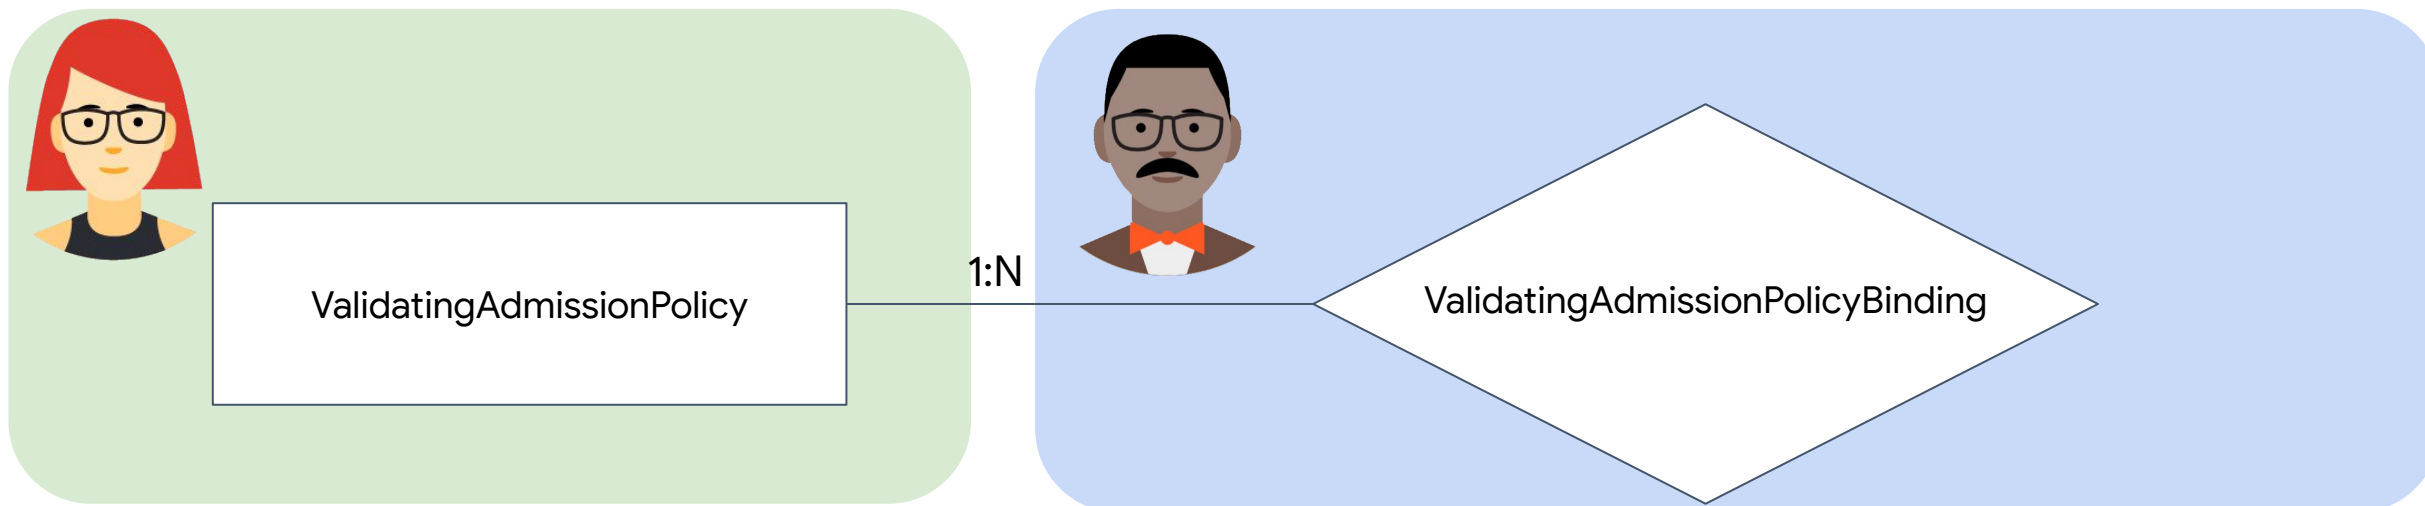
apiVersion: v1  
kind: ConfigMap  
metadata:

name: "replica-limit-test.example.com"

data:

maxReplicas: "3"

## No parameterization needed?



apiVersion: admissionregistration.k8s.io/v1alpha1

kind: ValidatingAdmissionPolicy

metadata:

name: "replicalimit-policy.example.com"

spec:

matchConstraints:

resourceRules:

- apiGroups: ["apps"]

apiVersions: ["v1"]

operations: ["CREATE", "UPDATE"]

resources: ["deployments"]

validations:

- expression: "object.spec.replicas <= 3"

apiVersion: admissionregistration.k8s.io/v1alpha1

kind: ValidatingAdmissionPolicyBinding

metadata:

name: "replicalimit-binding-test.example.com"

spec:

policyName: "replicalimit-policy.example.com"

validationActions: [Deny]

- **Parameterization**
- `policy.matchConstraints` VS `binding.matchResources` VS `policy.matchConditions`
- `failurePolicy` VS `validationActions`
- Leverage `authz` check

- Parameterization
- **policy.matchConstraints VS binding.matchResources**  
**VS policy.matchConditions**
- failurePolicy VS validationActions
- Leverage authz check

# Best Practices - matchConditions

kind: ValidatingAdmissionPolicy

spec:

...

matchConditions:

- name: 'exclude-leases' # Each match condition must have a unique name

expression: '!(request.resource.group == "coordination.k8s.io" && request.resource.resource == "leases")' # Match non-lease resources.

- name: 'exclude-kubelet-requests'  
expression: '!("system:nodes" in request.userInfo.groups)' # Match requests made by non-node users.

- name: 'rbac' # Skip RBAC requests.  
expression: 'request.resource.group != "rbac.authorization.k8s.io"'

...

MatchConditions to fine-grained request filtering:

- Are CEL expressions
- Must evaluate to true
- Has been added into Webhook as alpha feature in 1.27



Learn More:

<https://kubernetes.io/docs/reference/access-authn-authz/validating-admission-policy/#matching-requests-matchconditions>

- Parameterization
- `policy.matchConstraints` VS `binding.matchResources` VS `policy.matchConditions`
- **`failurePolicy` VS `validationActions`**
- Leverage `authz` check



# failurePolicy VS validationActions

## policy.failurePolicy

- Define how to handle failures
  - Compilation errors
  - Runtime errors
  - Mis-configuration
  - ...
- Fail
- Ignore

## binding.validationActions

Define how `validations` are enforced:


- **Deny** - Validation failure results in a denied request.
- **Warn** - Validation failure is reported to the request client as a warning
- **Audit** - Validation failure is included in the audit event for the API request.

- Parameterization
- `policy.matchConstraints` VS `binding.matchResources` VS `policy.matchConditions`
- `failurePolicy` VS `validationActions`
- **Leverage authz check**


# Best Practices - Authz Check

CEL Expression	Purpose
<code>authorizer.path('/healthz').check('get').allowed()</code>	Check if the principal (user or service account) submitting the request is authorized to perform a HTTP GET for the non-resource `/healthz` path.
<code>authorizer.serviceAccount('default', 'myserviceaccount')</code>	Checks if the principal (user or service account) submitting the request is authorized the service account with namespace 'default' and name 'myserviceaccount'.
<code>authorizer.requestResource.check('custom-verb').allowed()</code>	Checks if the principal (user or service account) submitting the request is authorized for `custom-verb` on the resource being handled.

# Adoption

 **Rita Zhang** @ritazzhang · Apr 2  
Replying to @ritazzhang @rawcode and 4 others  
We've recently added multi-engine support to [github.com/open-policy-agent/...](https://github.com/open-policy-agent/frameworks) to enable integration between OPA gatekeeper and CEL ValidatingAdmissionPolicy such that operators can manage policies in the same manner while the rules can be written in either rego or CEL. Coming soon!

open-policy-agent/  
**frameworks**



25 Contributors   4 Used by   107 Stars   43 Forks

github.com  
**GitHub - open-policy-agent/frameworks**  
Contribute to open-policy-agent/frameworks development by creating an account on GitHub.

1   4   21   1,041

kyverno/kyverno Public

<> Code   Issues 270   Pull requests 24   Discussions   Actions   Projects 2   Wiki   Security 2   Insights

## [Feature] K8s ValidatingAdmissionPolicy support in Kyverno #5441

Open   2 of 5 tasks   chipzoller opened this issue on Nov 22, 2022 · 2 comments

## Kubernetes Validating Admission Policies: A Practical Example

Thursday, March 30, 2023

**Authors:** Craig Box (ARMO), Ben Hirschberg (ARMO)

Wanna try the feature early?

**kubernetes/cel-admission-  
webhook**

# Takeaway



# Takeaway

CRD advanced  
validation

Policy Enforcement

Declarative APIs

Deployments

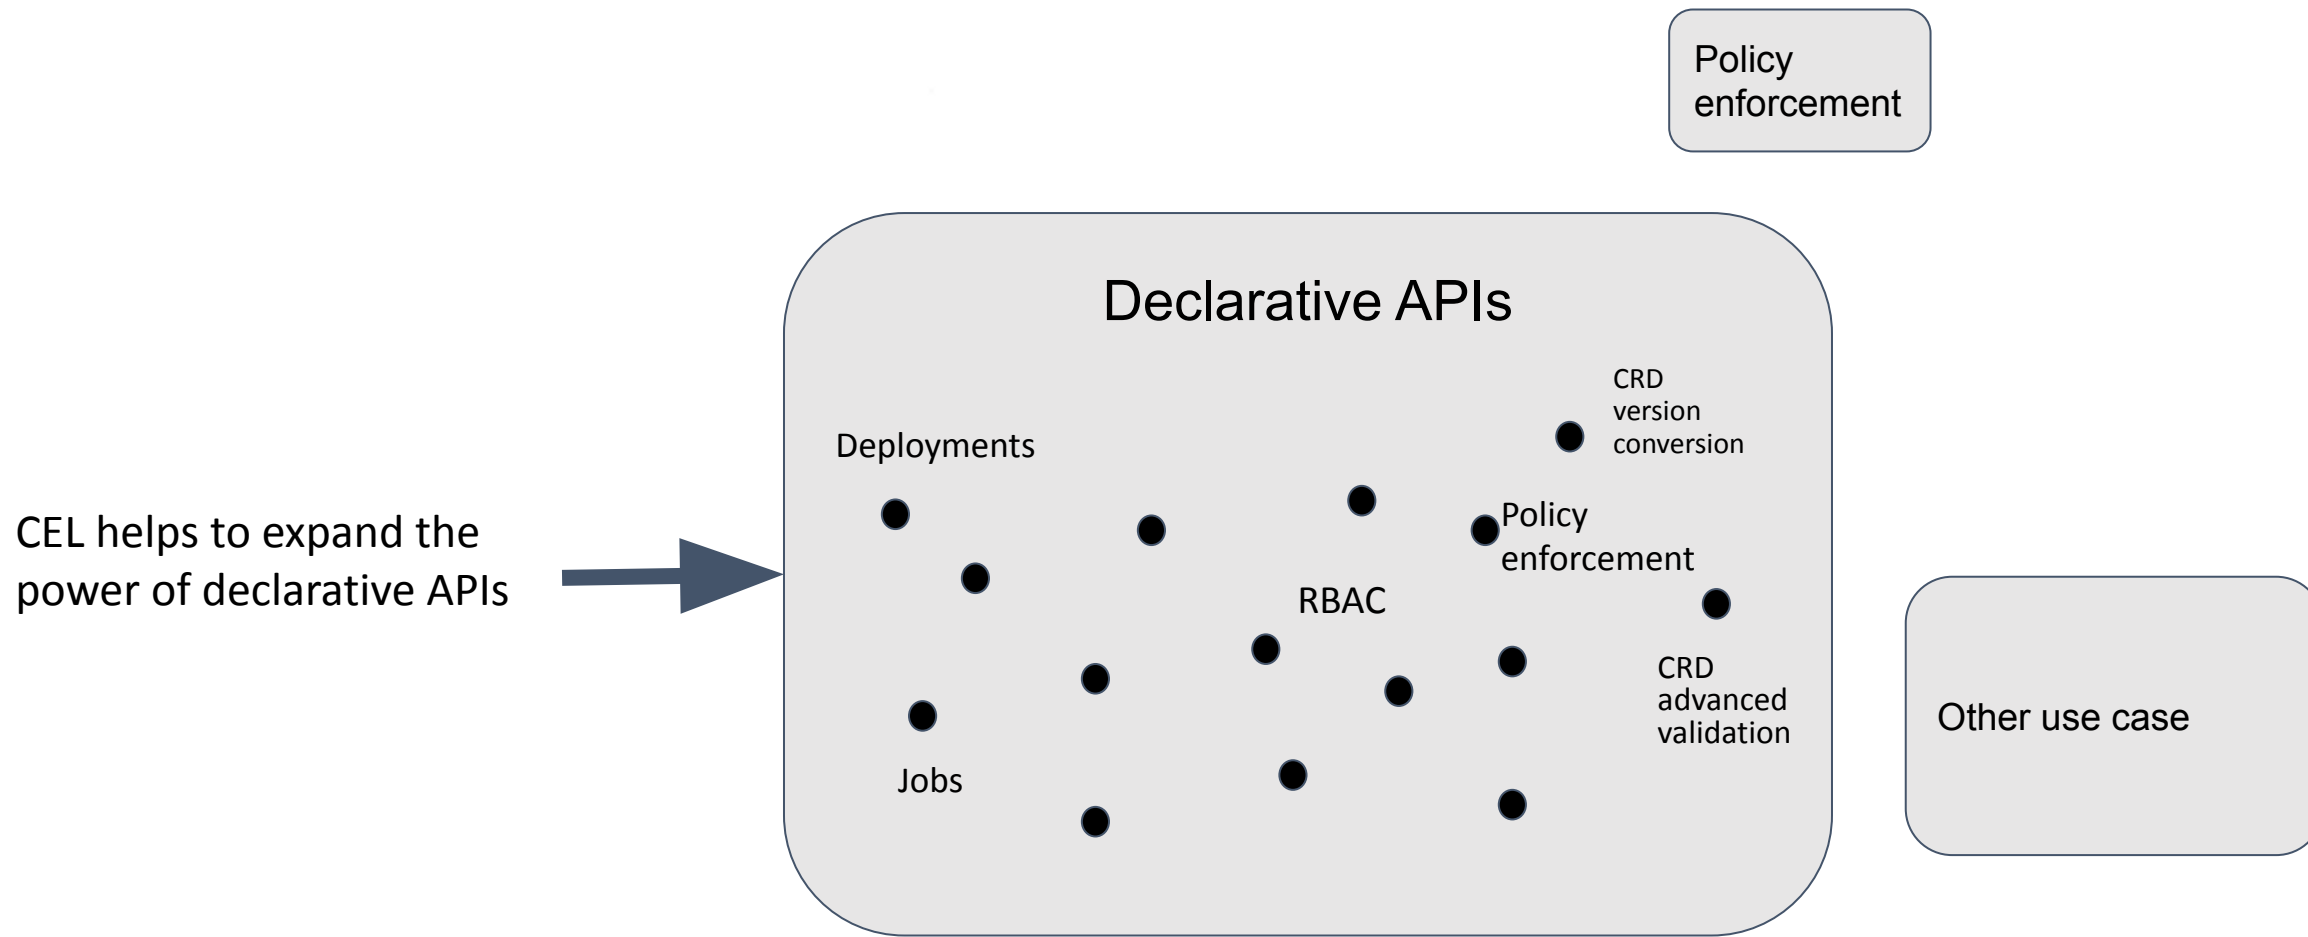
RBAC

Jobs

CRD version  
conversion

Other use cases

# Takeaway





- CRD Conversion with CEL
- CEL in native types: declarative validation
- MutatingAdmissionPolicy
- Client Side Validation Tool

# To Learn More:

## SIG API Machinery CEL group:

- Mailing list:
  - [kubernetes-sig-api-machinery@googlegroups.com](mailto:kubernetes-sig-api-machinery@googlegroups.com)
  - [wg-sig-api-machinery-cel-dev-external](#)
- Slack channel: #sig-api-machinery-cel-dev
- Bi-weekly community meetings



Please scan the QR Code above  
to leave feedback on this session