



KubeCon



CloudNativeCon

Europe 2023





KubeCon



CloudNativeCon

Europe 2023

The State and Future of Cloud Native ML Model Serving

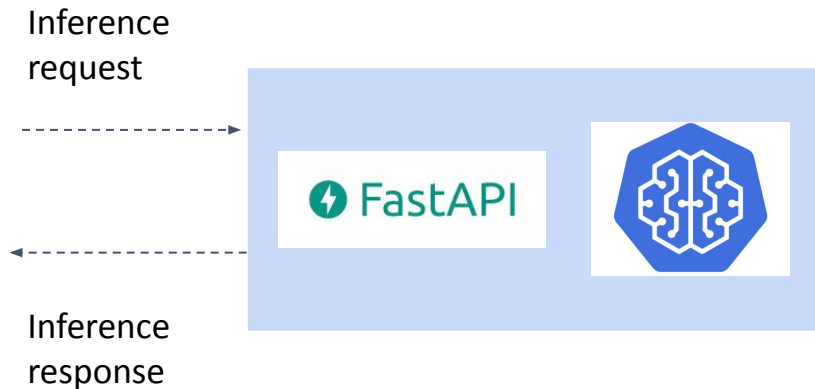
Dan Sun, Bloomberg

Theofilos Papapanagiotou, AWS



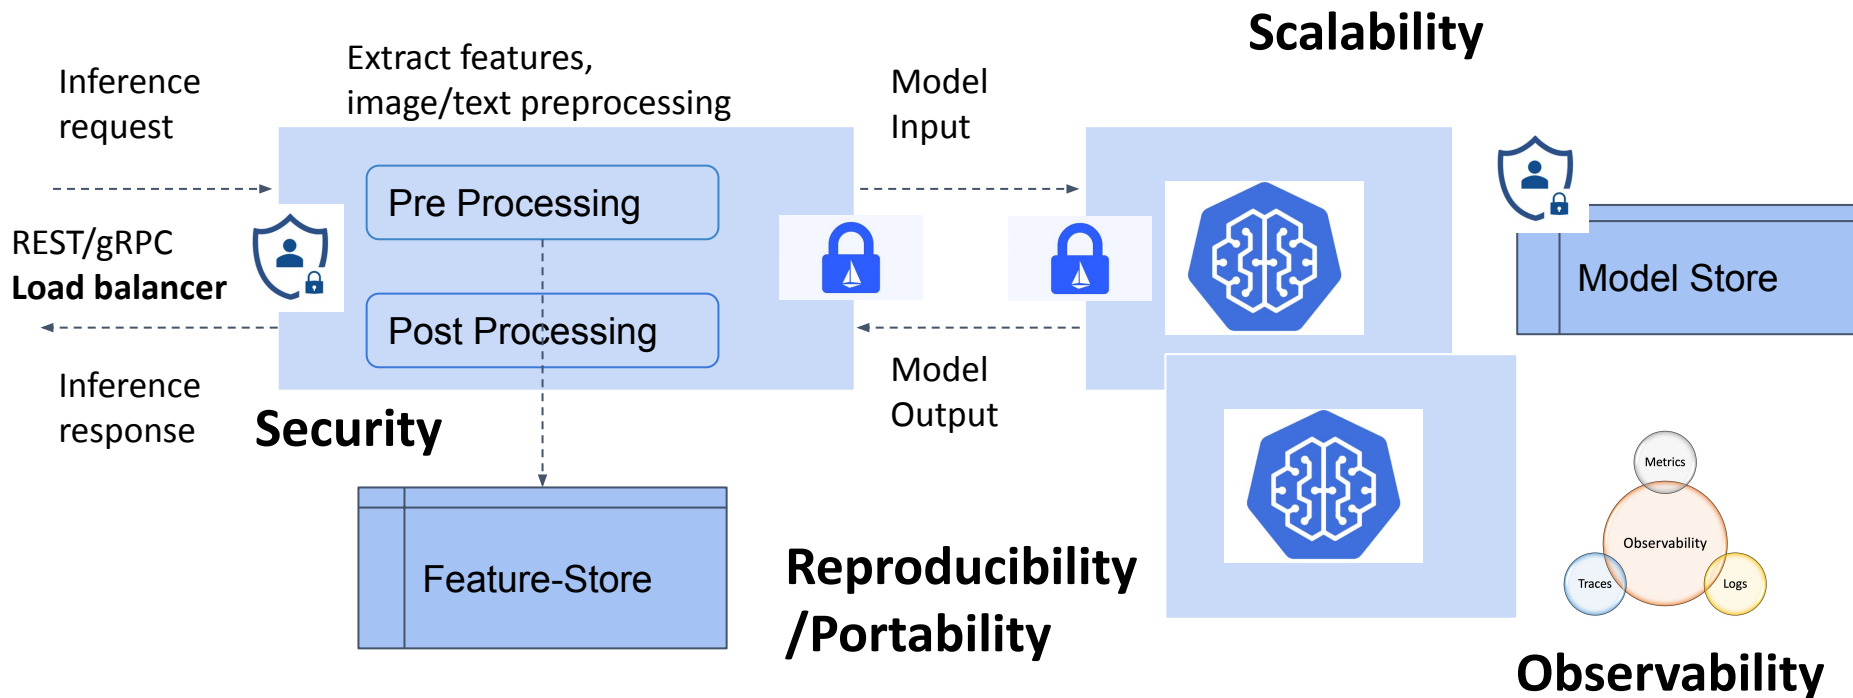
Productionize AI Model is Challenging

“Launching AI application pilots is deceptively easy, but deploying them into production is notoriously challenging.”



Productionize AI Model is Challenging

“Launching AI application pilots is deceptively easy, but deploying them into production is notoriously challenging.”



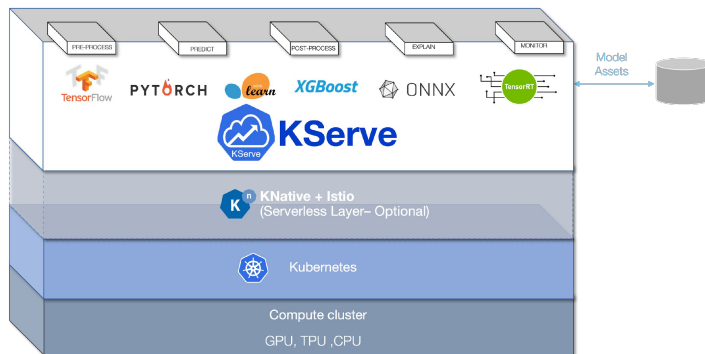
Why Kubernetes is a Great Platform for Serving Models?

- **Microservice:** Kubernetes handles container orchestration, facilitate deployment of microservices with resource management capabilities and load balancing.
- **Reproducibility/Portability:** ML model deployments can be written once, reproduced, and run with declarative yaml in a cloud agnostic way.
- **Scalability:** Kubernetes provides horizontal scaling features for both CPU and GPU workload.
- **Fault-tolerance:** Detect and recover from container failures, more resilient to outages and minimize downtime.



What's KServe?

- **KServe** is a **highly scalable** and **standards-based cloud-native model inference platform** on **Kubernetes** for Trusted AI that encapsulates the complexity of deploying models to production.
- KServe can be deployed **standalone** or as an **add-on** component with **Kubeflow** in the **cloud or on-premises environment**.



KServe History



**IBM contributed
ModelMesh**




2019

2020

2021

2022

**NVIDIA contributed
Inference Protocol**



The Current State

- Features supported as of KServe 0.10 release

Core Inference

- Transformer/Predictor
- Serving Runtimes
- Custom Runtime SDK
- Open Inference Protocol
- Serverless Autoscaling
- Cloud/PVC Storage

Advanced Inference

- ModelMesh for Multi-Model Serving
- Inference Graph
- Payload Logging
- Request Batching
- Canary Rollout

Model Explainability & Monitoring

- Text, Image, Tabular Explainer
- Bias Detector
- Adversarial Detector
- Outlier Detector
- Drift Detector

Serving Runtimes

- **ModelSpec** specifies the model formats or version for trained model
- KServe automatically selects the serving runtime to instantiate the deployment that supports the given model format.















InferenceService (for user)

```
apiVersion: serving.kserve.io/v1beta1
kind: InferenceService
metadata:
  name: example-sklearn-isvc
spec:
  predictor:
    model:
      modelFormat:
        name: sklearn
        version: 1.0
    storageUri: s3://bucket/sklearn/mnist.joblib
    runtime: kserve-sklearnserver # optional
```

Serving Runtime (for KServe admin)

```
apiVersion: serving.kserve.io/v1alpha1
kind: ClusterServingRuntime
metadata:
  name: kserve-sklearnserver
spec:
  supportedModelFormats:
    - name: sklearn
      version: "1"
    autoSelect: true
  containers:
    - name: kserve-container
      image: kserve/sklearnserver:latest
      args:
        --model_name={{.Name}}
        --model_dir=/mnt/models
        --http_port=8080
  resources:
    requests:
      cpu: "1"
      memory: 2Gi
```

Serving Runtime Support Matrix

Serving Runtime/ Model Format	scikit-learn	xgboost	lightgbm	TensorFlow	PyTorch	TorchScript	ONNX	MLFlow	Custom
MLServer (open)									
Triton (open)									
TorchServe (v1, open)									
KServe Runtime (v1, open)									
TFServing (v1)									

Open Inference Protocol

- Open inference protocol enables a standardized high performance inference data plane.
- Allows interoperability between serving runtimes.
- Allows building client or benchmarking tools that can work with a wide range of serving runtimes
- It is implemented by [KServe](#), [MLServer](#), [Triton](#), [TorchServe](#), [OpenVino](#), [AMD Inference Server](#).



<https://github.com/kserve/open-inference-protocol>

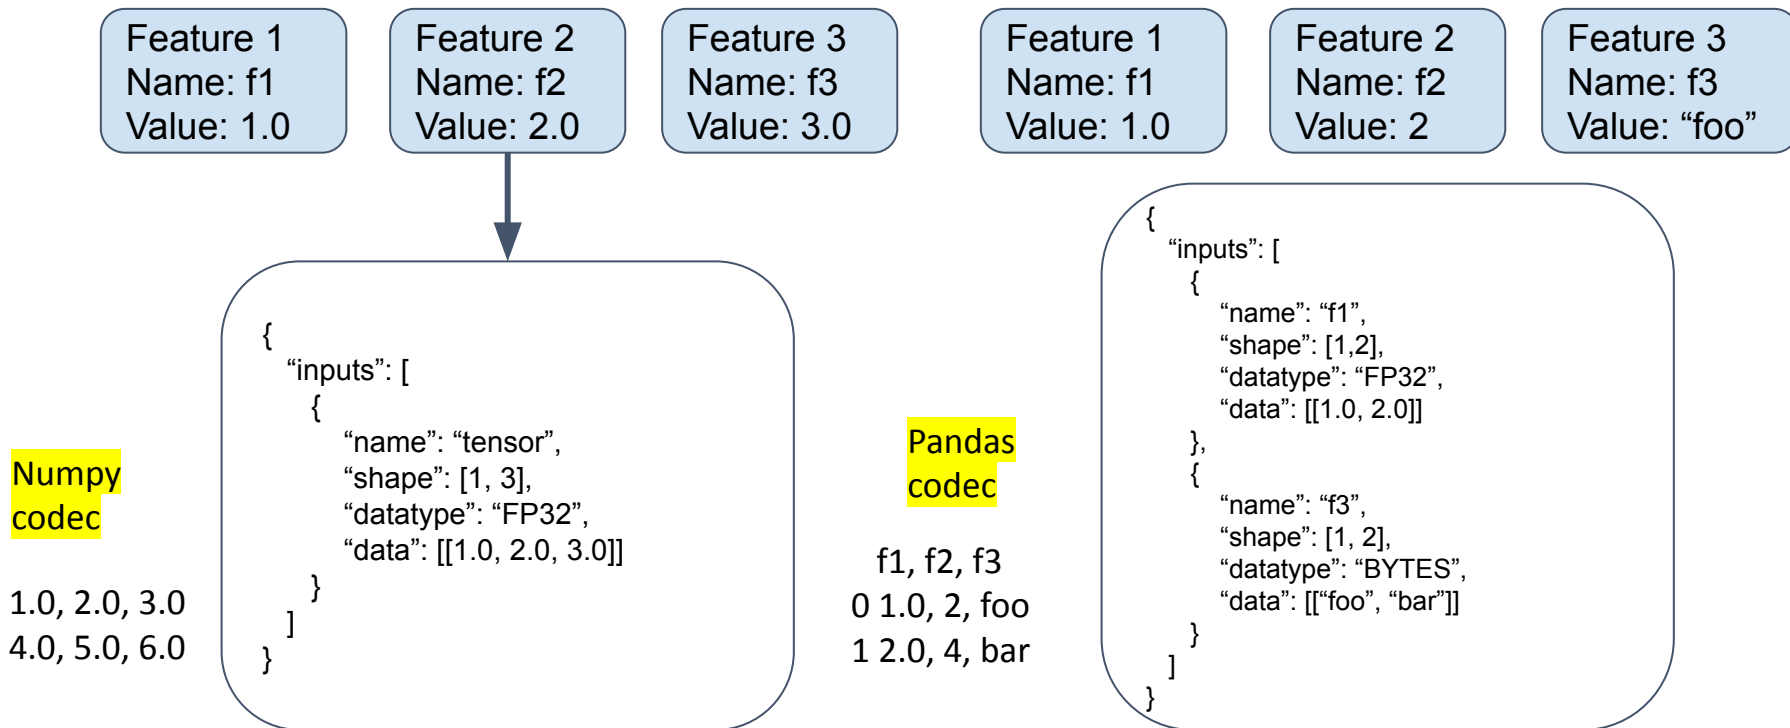
Open Inference Protocol

- REST vs. gRPC: Ease of use vs high performance

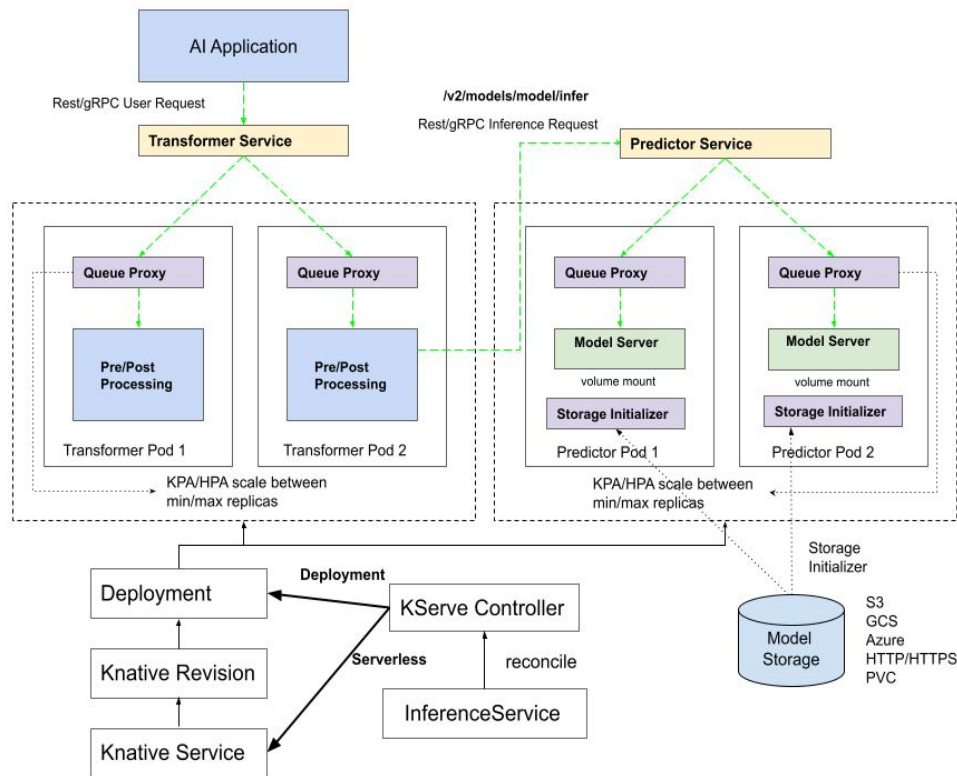
REST	gRPC
GET v2/health/live	rpc ServerLive(ServerLiveRequest) returns (ServerLiveResponse)
GET v2/health/ready	rpc ServerReady(ServerReadyRequest) returns (ServerReadyResponse)
GET v2/models/{model_name}/ready	rpc ModelReady(ModelReadyRequest) returns (ModelReadyResponse)
GET v2/models/{model_name}	rpc ModelMetadata(ModelMetadataRequest) returns (ModelMetadataResponse)
POST v2/models/{model_name}/infer	rpc ModelInfer(ModelInferRequest) returns (ModelInferResponse)

Open Inference Protocol Codec

- Single Input vs. Multiple Inputs with numpy/pandas codec



Inference Service Deployment

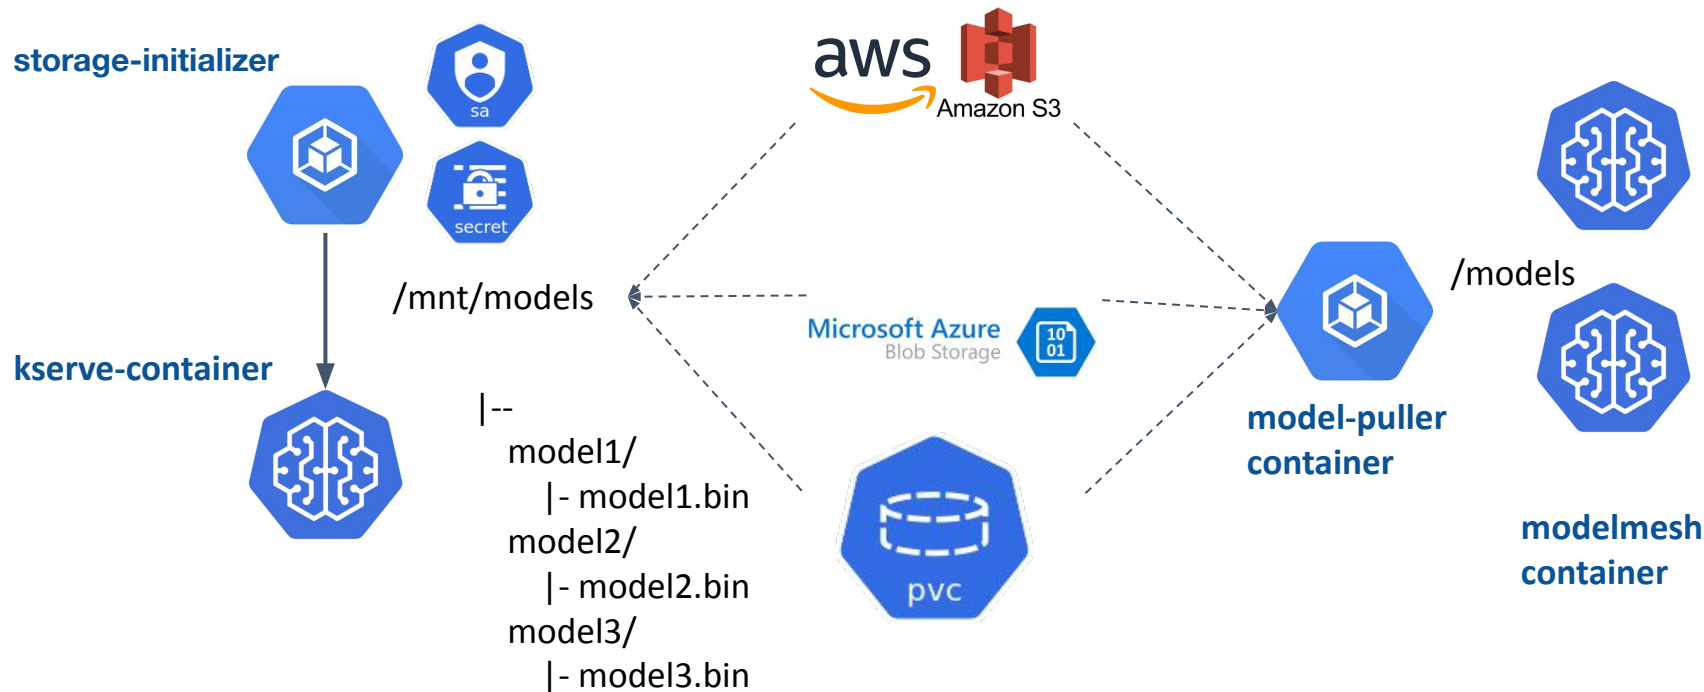


```
apiVersion: "serving.kserve.io/v1beta1"
kind: "InferenceService"
metadata:
  name: "sklearn-feast-transformer"
spec:
  transformer:
    containers:
      - image:
        kserve/driver-transformer:latest
        name: kserve-container
        command:
          - "python -m driver_transformer"
        args:
          - --entity_ids
          - driver_id
          - --feature_refs
          - driver_hourly_stats:acc_rate
          - driver_hourly_stats:avg_daily_trips
          - driver_hourly_stats:conv_rate
  predictor:
    model:
      modelFormat:
        name: sklearn
      storageUri: "gs://pv/driver"
```

<https://kserve.github.io/website/0.10/model-serving/v1beta1/transformer/feast/>


Model Storage Patterns

- Runs as init container to download the models before starting KServe container.
- Runs as sidecar to pull models as needed in ModelMesh.




Multi-Model Serving

- Single model single container usually under-utilize the resources.
- You can share the container resources by deploying multiple models, especially for the GPU sharing cases.
- When scaling up the same set of models need to be scaled out together.




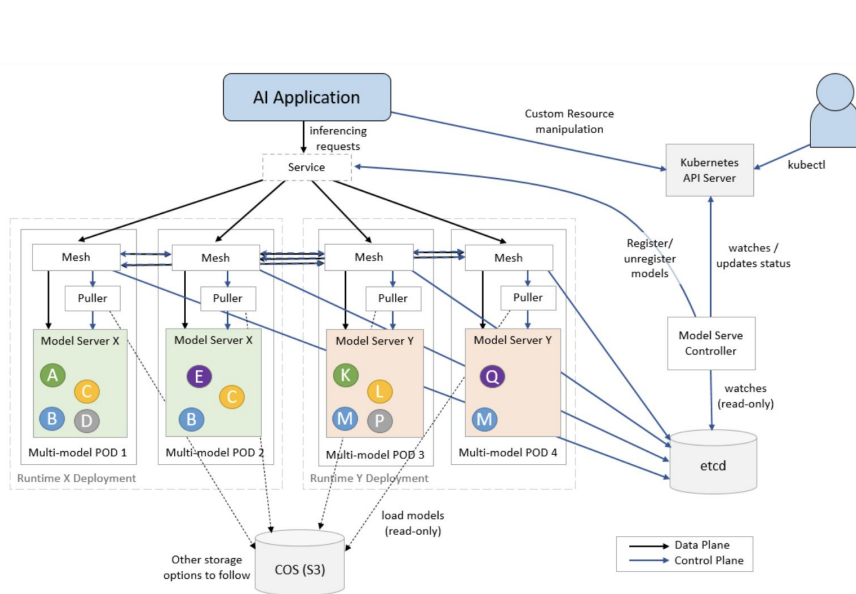
Inference Service Replica 1



Inference Service Replica 2

ModelMesh

- Models can be scaled independently while sharing the container resource
- ModelMesh is a scale-out layer for model servers that can load and serve multiple models concurrently which uses the same KServe InferenceService API.
- Designed for high volume, high density, high churn production serving, it has powered IBM production AI services for a number of years and manages thousands of models.

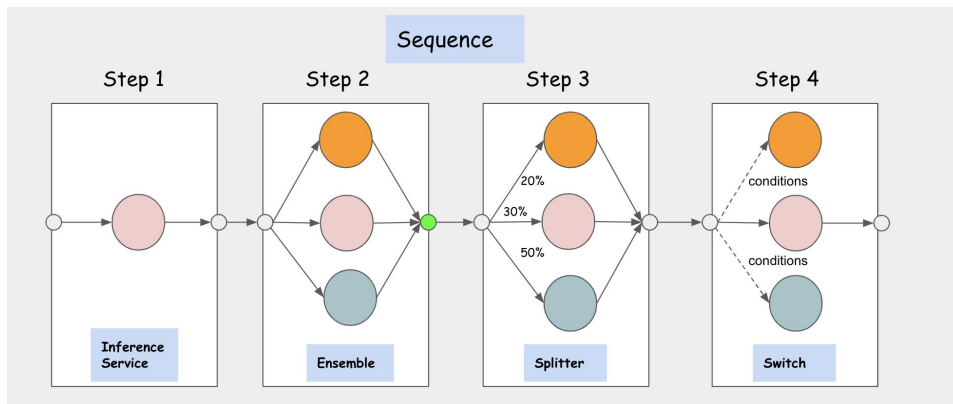


Kubernetes schedules pod onto nodes

ModelMesh schedules models onto container

Inference Graph

- Inference Graph is built for more complex **multi-stage inference pipelines**.
- Inference Graph is deployed in a **declarative way and highly scalable**.
- Inference Graph supports **Sequence, Switch, Ensemble and Splitter** nodes.
- Inference Graph is **highly composable**. It is made up with a list of routing nodes and each node consists of a set of routing steps which can be either route to an InferenceService or another node.



```
apiVersion: "serving.kserve.io/v1alpha1"
kind: "InferenceGraph"
metadata:
  name: "dog-breed-pipeline"
spec:
  nodes:
    root:
      routerType: Sequence
      steps:
        - serviceName: cat-dog-classifier
          name: cat_dog_classifier # step name
        - serviceName: dog-breed-classifier
          name: dog_breed_classifier
          data: $request
          condition: "[@this].#(predictions.0==\"dog\")"
```

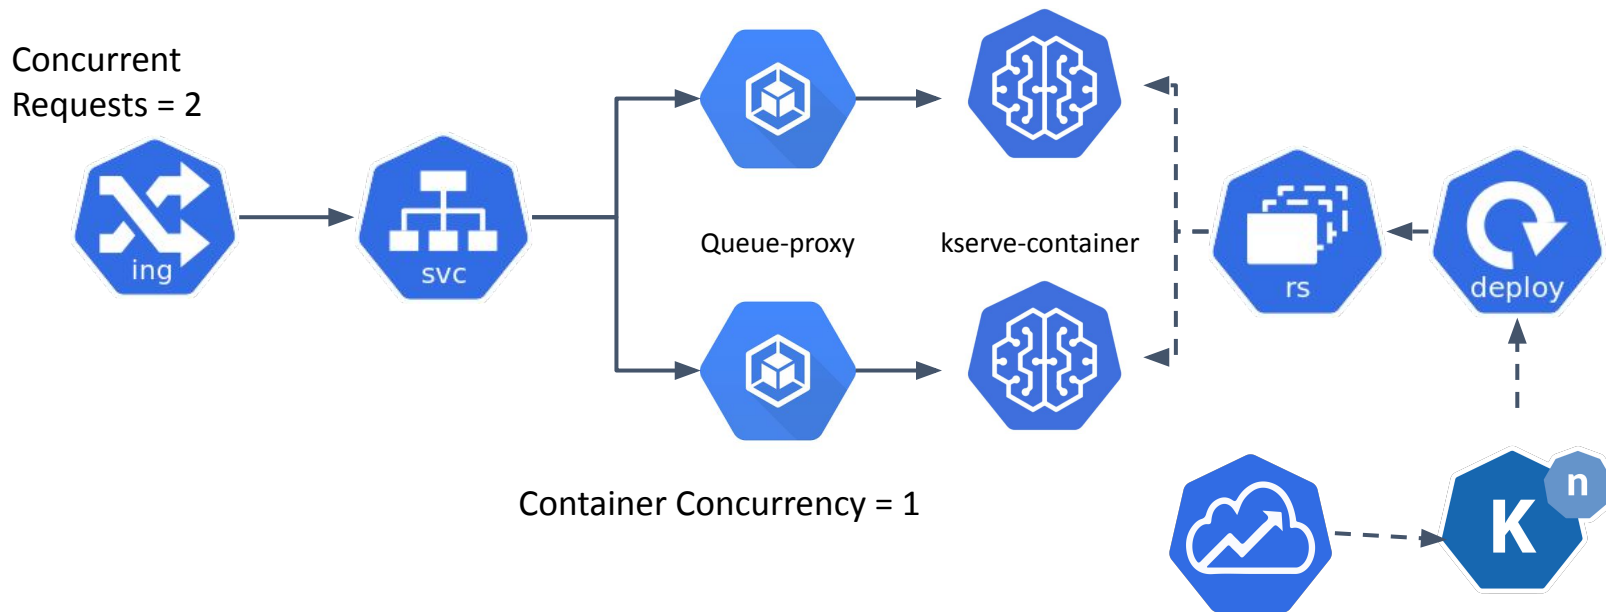
KServe Cloud Native Ecosystem

- **KServe** offers seamless integration with many CNCF projects to empower the production model deployments for **security**, **observability**, and **serverless capabilities**.



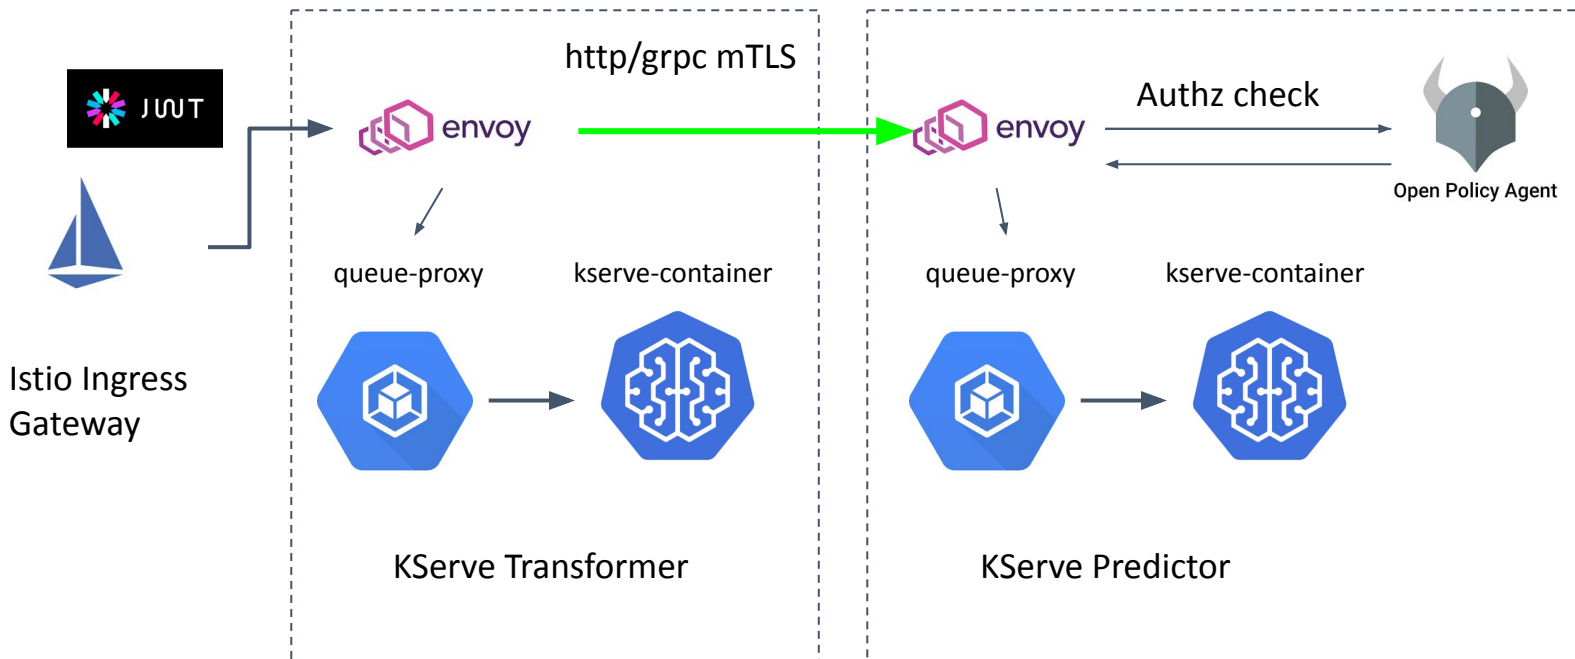
Serverless Inference

- Autoscaling based on incoming request RPS or Concurrency.
- Unified autoscaler for both CPUs and GPUs.



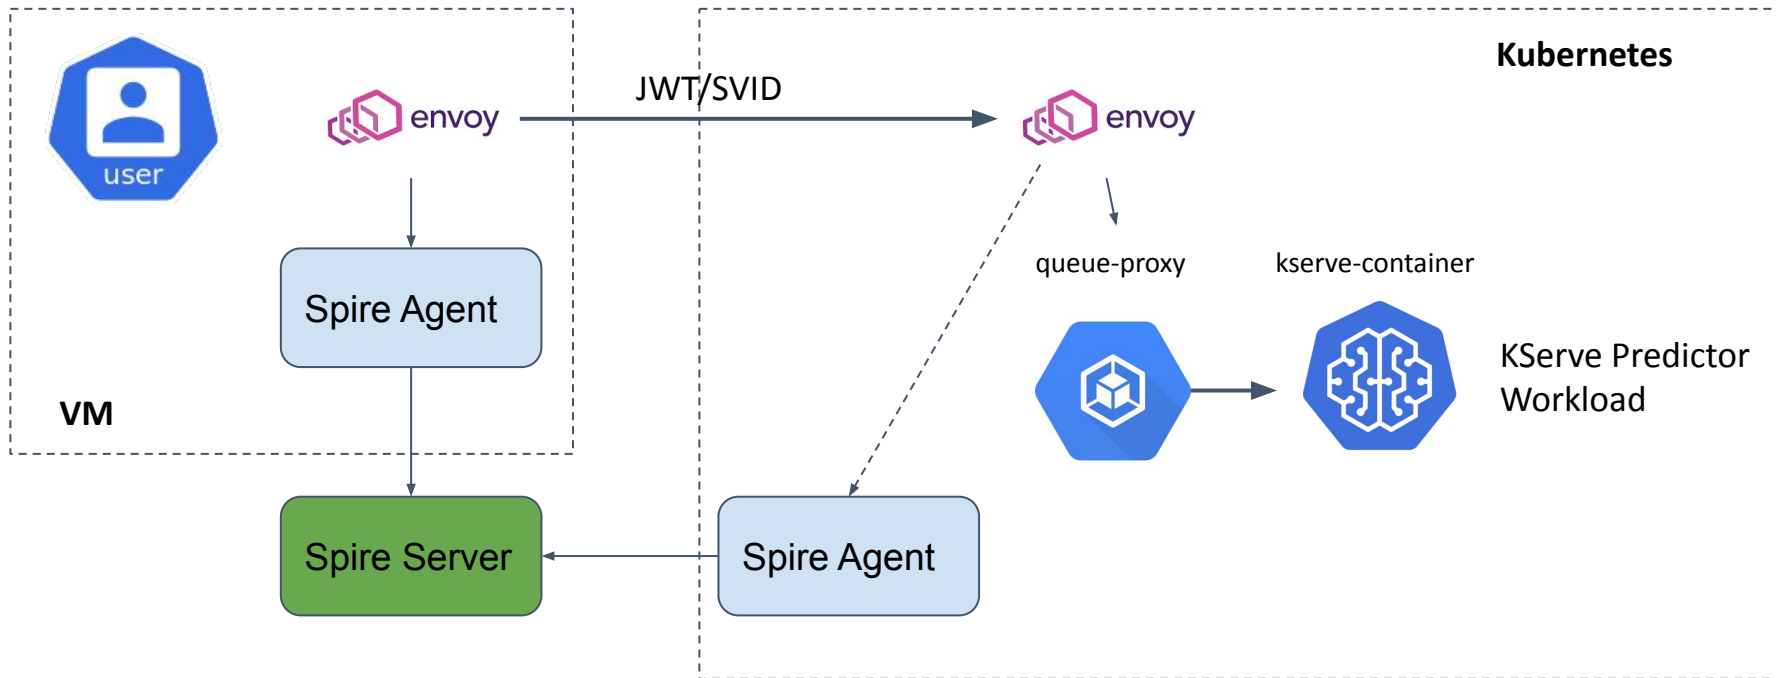
ServiceMesh: Secure Inference Service

- Istio control plane mounts the client certificates to sidecar proxy to allow automatic authentication between transformer and predictor and encrypt service traffic.
- Authorization can be implemented with Istio Authorization Policy or Open Policy Agent.



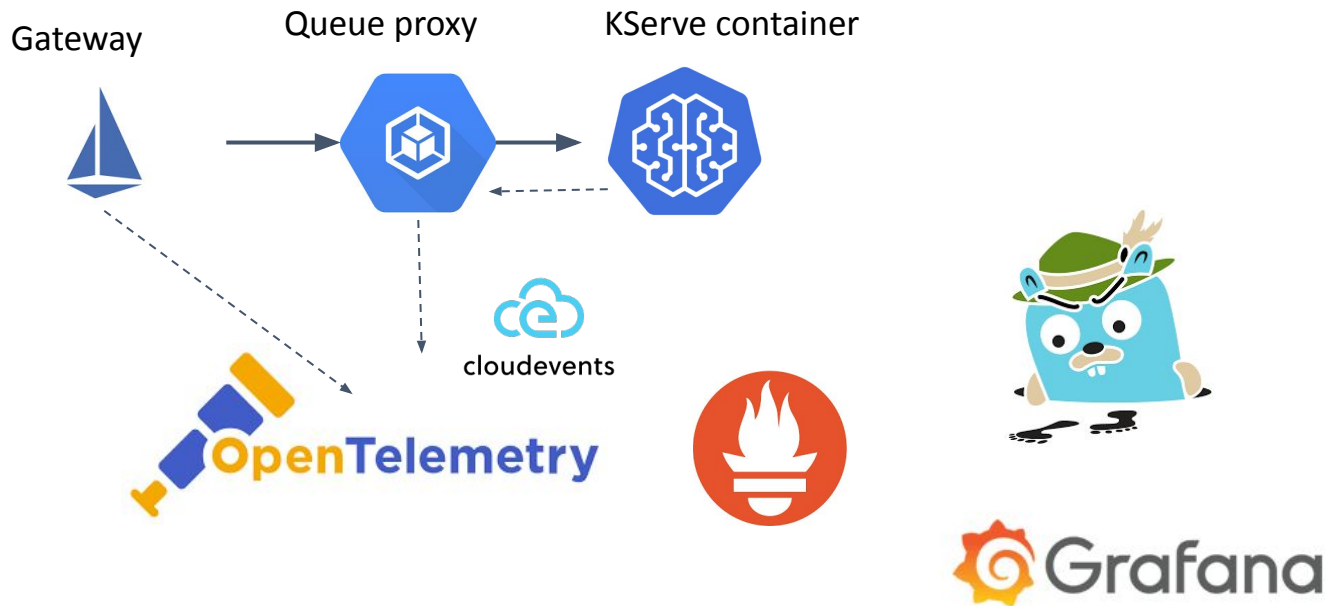
SPIFFE and SPIRE: Strong Identity

- **SPIFFE** is a secure identity framework that can be federated across heterogeneous environments.
- **SPIRE** distinguishes from other identity providers, such as API keys or secrets, with an attestation process before issuing the credential.



KServe Observability


- **OpenTelemetry** can be used to instrument, collect and export metrics, log and tracing data that you can analyze to understand the production Inference Service performance.



Cloud Native Buildpacks

- Cloud Native Buildpacks can turn your custom transformer or predictor into a container image without Dockerfile and run anywhere.
- Ensures meeting the security and compliance requirements.

```
class ImageTransformer(Model):  
    def __init__(self, name: str, predictor_host: str, protocol: str):  
        super().__init__(name)  
        self.predictor_host = predictor_host  
        self.protocol = protocol  
        self.model_name = name  
  
    def preprocess(self, request: InferRequest, headers: Dict[str, str] = None) -> InferRequest:  
        input_tensors = [image_transform(instance) for instance in request.inputs[0].data]  
        input_tensors = np.asarray(input_tensors)  
        infer_inputs = [InferInput(name="INPUT__0", datatype='FP32', shape=list(input_tensors.shape),  
                                data=input_tensors)]  
  
        infer_request = InferRequest(model_name=self.model_name, infer_inputs=infer_inputs)  
        return infer_request
```



Looking Forward




KServe 1.0 Roadmap

- Graduate KServe core inference capabilities to stable/GA
- Open Inference Protocol coverage for all supported serving runtimes
- Graduate KServe SDK to 1.0
- Graduate ModelMesh and Inference Graph
- Large Language Model (LLM) Support
- KServe Observability and Security improvements
- Update KServe 1.0 documentation


<https://github.com/kserve/kserve/blob/master/ROADMAP.md>

Large Language Models



Distributed Inference for LLM


- Large transformer based models with **hundreds of billions of parameters** are **computationally expensive**.
- NVIDIA **FasterTransformer** implements an accelerated engine for inference of transformer based model spanning many GPUs and nodes in distributed manner.
- **Huggingface Accelerate** allows distributed inference with sharded checkpoints using less memory.
- **Tensor parallelism** allows running inference on multiple GPUs.
- **Pipeline parallelism** allows on multi-GPU and multi-node environment.



Large Language Model Challenges

- **Inference hardware requirements**
 - Cost: Requires significant computational resources with high end GPUs and large amount of memory
 - Latency: long response time up to tens of seconds
- **Model blob/file sizes in GBs(BLOOM):**
 - 176bln params = 360GB
 - 72 splits of 5GB which needs to mapped to multiple GPU devices
- **Model loading time**
 - From network (S3, minio) to instance disk
 - From instance disk to CPU RAM
 - From CPU RAM to GPU RAM
- **Model Serving Runtime:** FasterTransformer-Triton (32GB)
- **Model data can be sensitive and private for inference**

LLM(BLOOM 176bIn) Inference Demo



```
curl https://fastertransformer.default.kubecon.theofpa.people.aws.dev/v1/models/fastertransformer:predict -d '{
  "inputs":
  [
    {
      "input": "Kubernetes is the best platform to serve your models because",
      "output_len": "40"
    }
  ]
}'
```

[["Kubernetes is the best platform to serve your models because it is a cloud-based service that is built on top of the Kubernetes API. It is a great platform for developers who want to build their own Kubernetes application"]]

<https://github.com/kserve/kserve/tree/master/docs/samples/v1beta1/triton/fastertransformer>

LLM Download and Loading Time

```
$ git lfs clone https://huggingface.co/bigscience/bloom-560m
```

```
$ ls -lh bloom-560m
```

```
total 3.2G
-rw-r--r-- 1 theofpa domain^users 693 Apr 13 11:16 config.json
-rw-r--r-- 1 theofpa domain^users 1.1G Apr 13 11:22 flax_model.msgpack
-rw-r--r-- 1 theofpa domain^users 16K Apr 13 11:16 LICENSE
-rw-r--r-- 1 theofpa domain^users 1.1G Apr 13 11:22 model.safetensors
-rw-r--r-- 1 theofpa domain^users 1.1G Apr 13 11:22 pytorch_model.bin
-rw-r--r-- 1 theofpa domain^users 21K Apr 13 11:16 README.md
-rw-r--r-- 1 theofpa domain^users 85 Apr 13 11:16 special_tokens_map.json
-rw-r--r-- 1 theofpa domain^users 222 Apr 13 11:16 tokenizer_config.json
-rw-r--r-- 1 theofpa domain^users 14M Apr 13 11:16 tokenizer.json
```

Huggingface transformer to FasterTransformer conversion based on tensor parallelisms and target precision.

```
python3 FasterTransformer/examples/pytorch/gpt/utils/huggingface_bloom_convert.py\  
-o fastertransformer/1 -i ./bloom-560m/
```

```
$ ls -lhS fastertransformer/1 | head
```


```
total 2.1G
-rw-r--r-- 1 theofpa domain^users 980M Apr 13 11:28 model.wte.bin
-rw-r--r-- 1 theofpa domain^users 16M Apr 13 11:28 model.layers.0.mlp.dense_4h_to_h.weight.0.bin
-rw-r--r-- 1 theofpa domain^users 16M Apr 13 11:28 model.layers.0.mlp.dense_h_to_4h.weight.0.bin
-rw-r--r-- 1 theofpa domain^users 16M Apr 13 11:28 model.layers.10.mlp.dense_4h_to_h.weight.0.bin
-rw-r--r-- 1 theofpa domain^users 16M Apr 13 11:28 model.layers.10.mlp.dense_h_to_4h.weight.0.bin
-rw-r--r-- 1 theofpa domain^users 16M Apr 13 11:28 model.layers.11.mlp.dense_4h_to_h.weight.0.bin
-rw-r--r-- 1 theofpa domain^users 16M Apr 13 11:28 model.layers.11.mlp.dense_h_to_4h.weight.0.bin
-rw-r--r-- 1 theofpa domain^users 16M Apr 13 11:28 model.layers.12.mlp.dense_4h_to_h.weight.0.bin
-rw-r--r-- 1 theofpa domain^users 16M Apr 13 11:28 model.layers.12.mlp.dense_h_to_4h.weight.0.bin
```

storage-initializer:

```
2023-04-13 19:28:34.151 1 root INFO [downl  
Copying contents of s3://kubeccon-models/mo  
2023-04-13 19:29:28.922 1 root INFO [downl  
Successfully copied s3://kubeccon-models/mo  
/mnt/models
```

predictor:

```
I0413 19:29:34.205155 1 libfastertransforme  
TRITONBACKEND_ModelInitialize: fastertrans  
1)  
I0413 19:29:45.448307 1 model lifecycle.cc:  
successfully loaded 'fastertransformer' ve
```

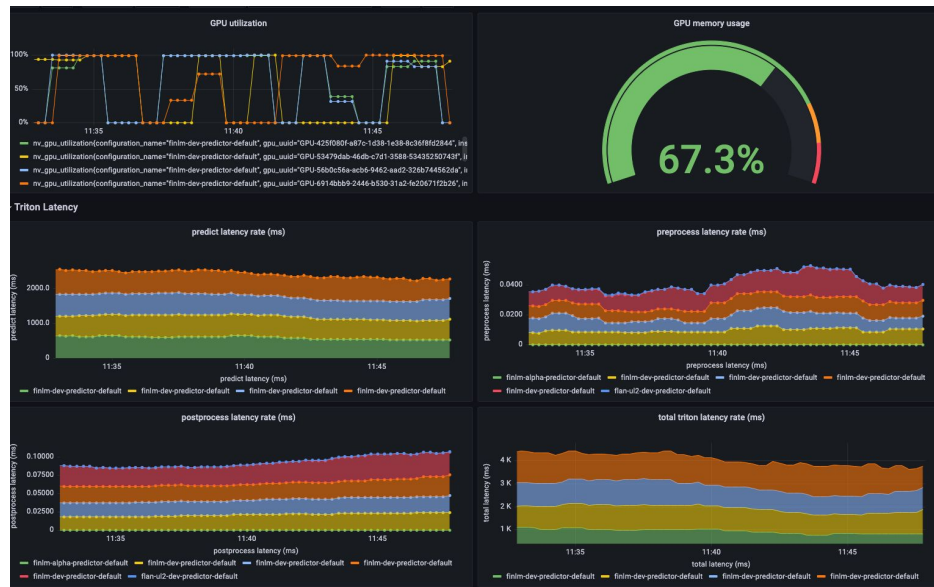
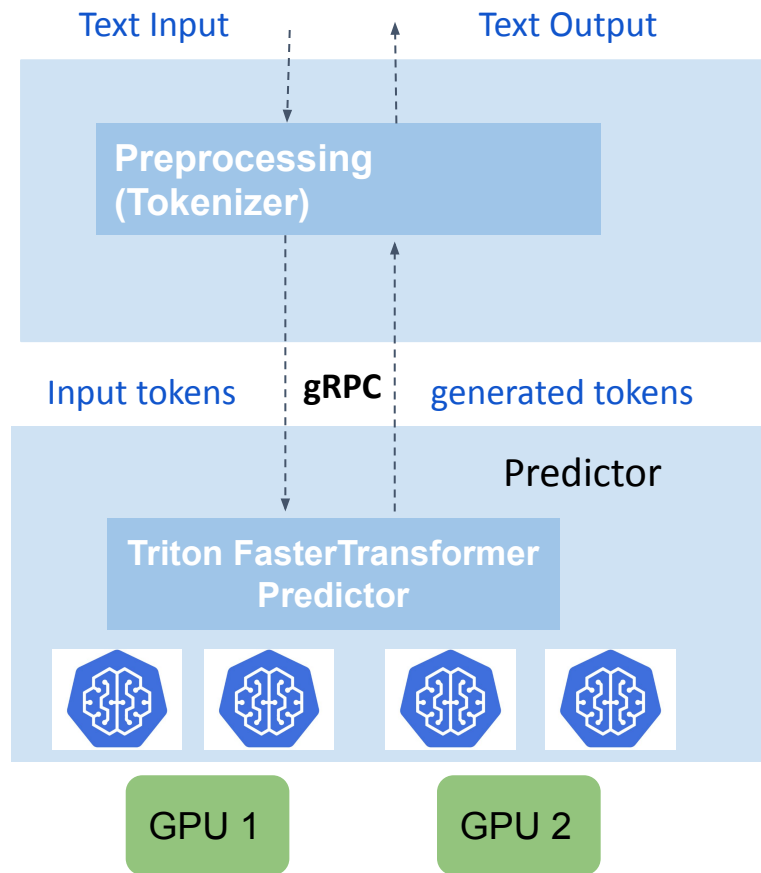


Introducing BloombergGPT, Bloomberg's 50-billion parameter large language model, purpose-built from scratch for finance

- Trained approximately 53 days on 64 servers, each with 8 A100 GPUs on AWS Sagemaker.
- HF checkpoints are converted into FasterTransformer format setting target precision to BF16, tensor parallelism 2 and pipeline parallelism 1.
- Deployed on KServe on-prem for internal product development with NVIDIA Triton FasterTransformer Serving Runtime on 2 A100 GPU (80G memory) for each replica.

<https://www.bloomberg.com/company/press/bloomberggpt-50-billion-parameter-llm-tuned-finance/>
BloombergGPT Paper: <https://arxiv.org/abs/2303.17564>

BloombergGPT - KServe Deployment



KServe Community

- [KServe Website](#)
- [KServe on GitHub](#)
- [KServe Community](#)





Please scan the QR Code above
to leave feedback on this session