



KubeCon



CloudNativeCon

Europe 2023





KubeCon




CloudNativeCon

Europe 2023

Deploying Cloud-Native Applications Using Kubevela and OAM

Daniel Higuero



Deploying Cloud-Native Applications Using Kubevela and OAM




KubeCon



CloudNativeCon

Europe 2023

18-21 April



Daniel Higuero
CTO

About me

- CTO @  **NAPPTIVE**
- OAM and KubeVela maintainer
- Work in the past on
 - Big Data
 - Streaming ML models
 - Edge computing



Questions on OAM/KubeVela?
⇒ **KubeVela CNCF kiosk #27**

- Open Application Model
- KubeVela
- Installing KubeVela
- Basic operations
- VelaUX
- Application workflows
- Multi-cluster deployment
- GitOps
- Extra content




<https://github.com/napptive/kubecon-23-oam-kubevela-tutorial>

Open Application Model

Open Application Model (OAM)


Open Application Model (OAM) is a set of standard yet higher level abstractions for modeling cloud native applications on top of today's hybrid and multi-cloud environments.



<https://oam.dev/>

Open Application Model

- Application as the top level entity
 - Born on 2019
 - Latest revision is v0.3.1
- Infrastructure agnostic
- Provider agnostic
- Focus on
 - Reusability
 - Understandability
- Can easily be extended



<https://oam.dev/>

 Alibaba Cloud

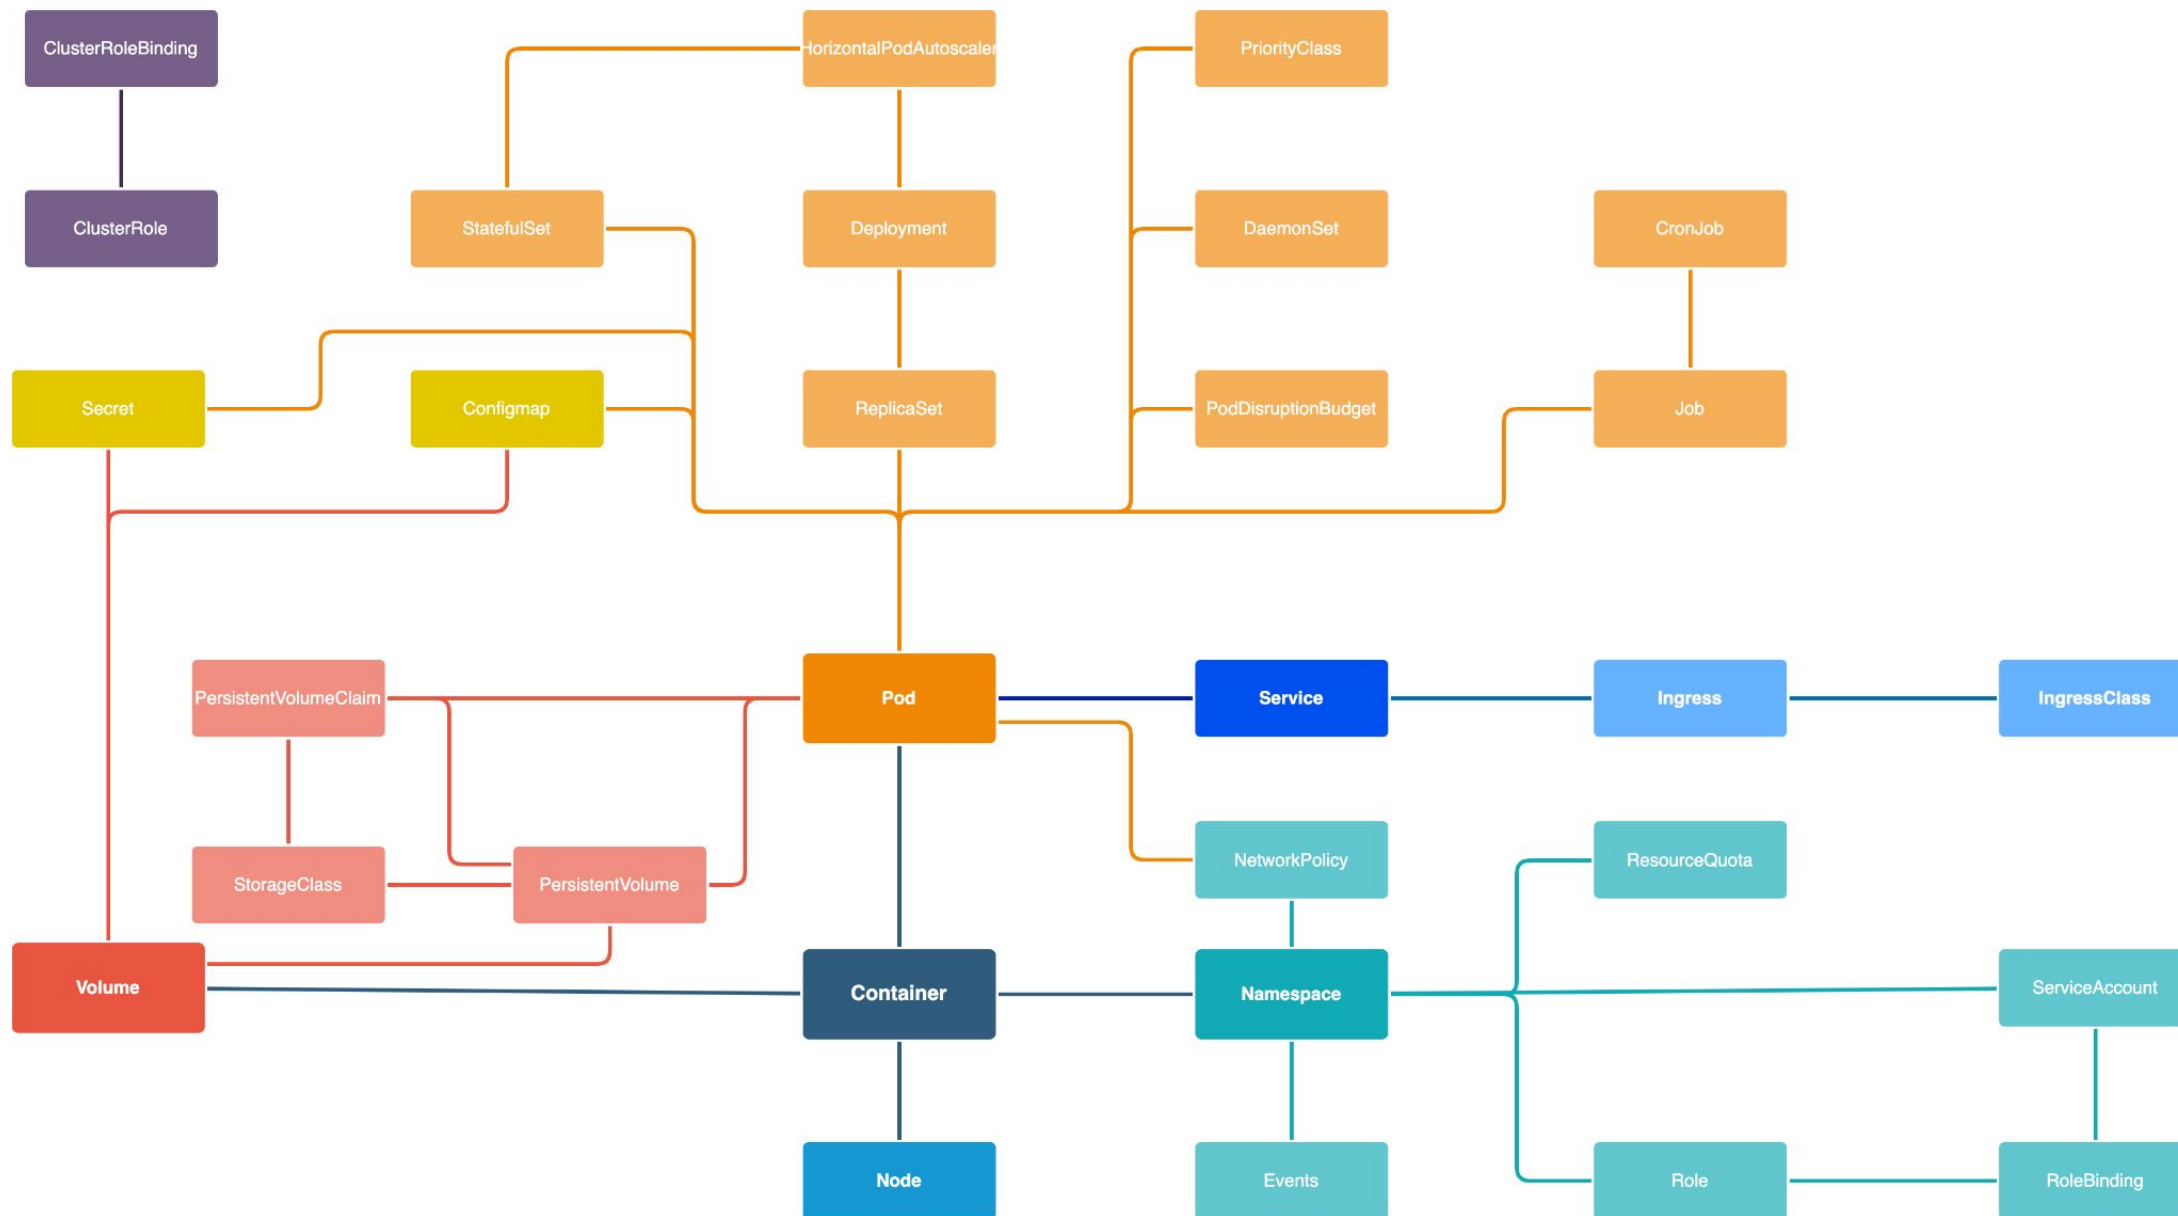
 **NAPPTIVE**

 Microsoft

Which problem are we addressing?

- Kubernetes is a high-performance, battle-tested framework to run our applications
 - Can be used to run any application
 - Can be adapted to any use case
- But
 - It is **difficult to learn**
 - It works with low-level entities
 - It requires making several entities work coherently together

Kubernetes entity map

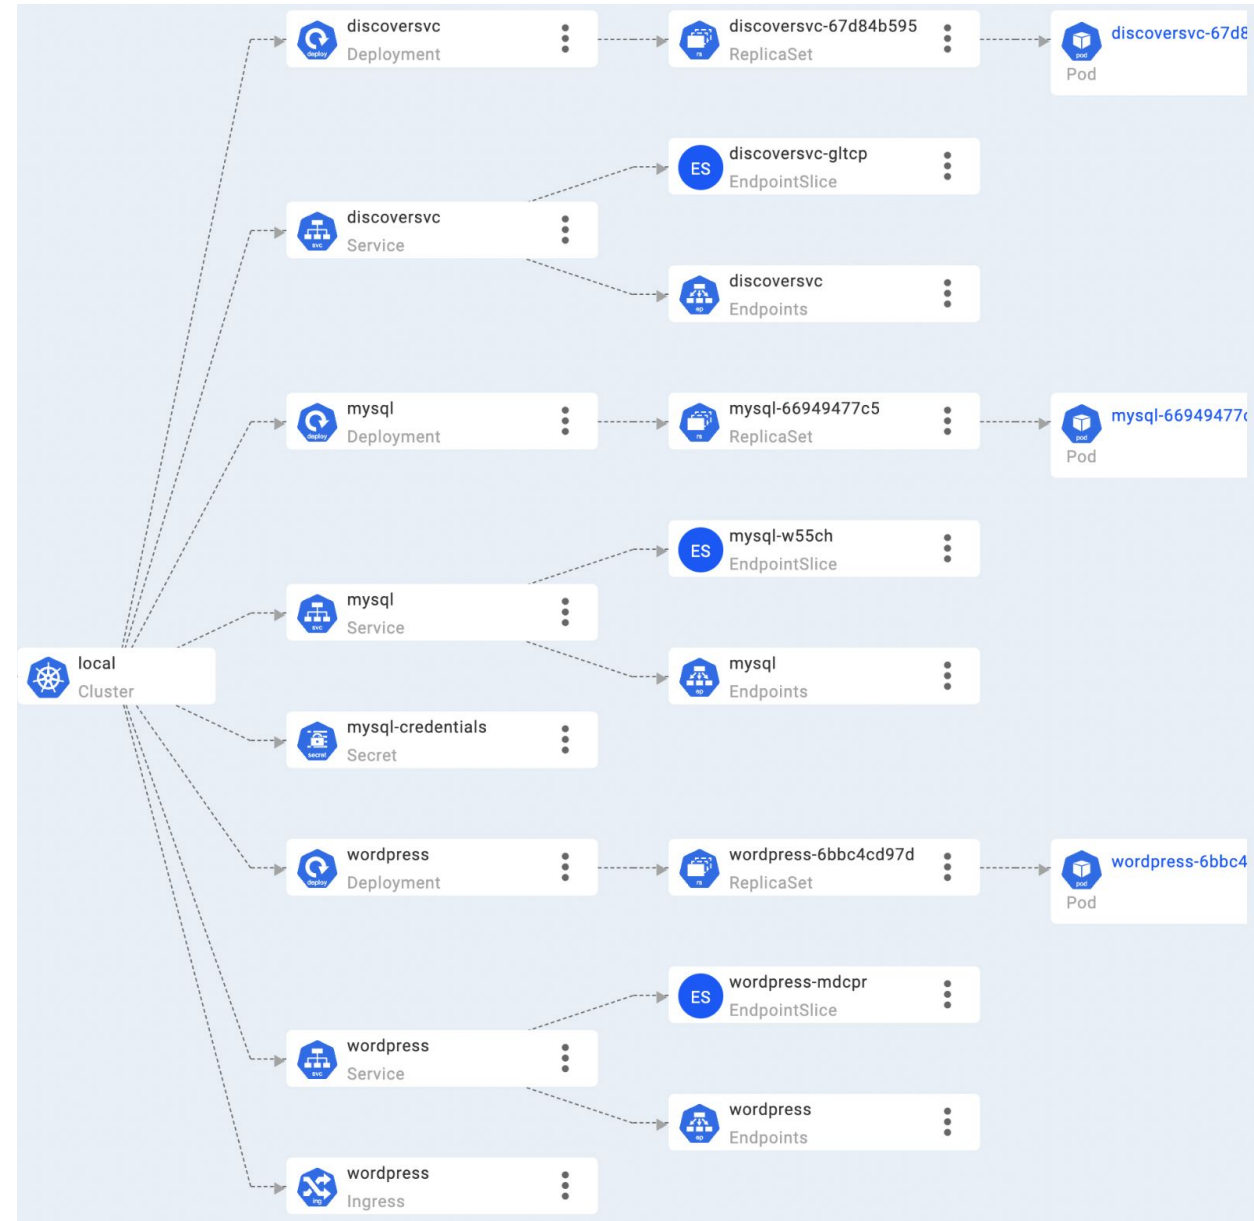


Is my app running?

- The low level entities allow us to
 - Represent the majority of use-cases
 - Deploy any application
 - Extract the required performance from the application
 - Adapt the cluster resources to the deployed workload
- But follows a bottom-up approach
 - Determining if an application is running means identifying and checking all the base K8s entities
 - Some entities are provider dependent

Is my app running?

- Typical mitigation approach
 - Labeling all resources
 - But we may depend on existing labeling approaches for existing applications



Why OAM?

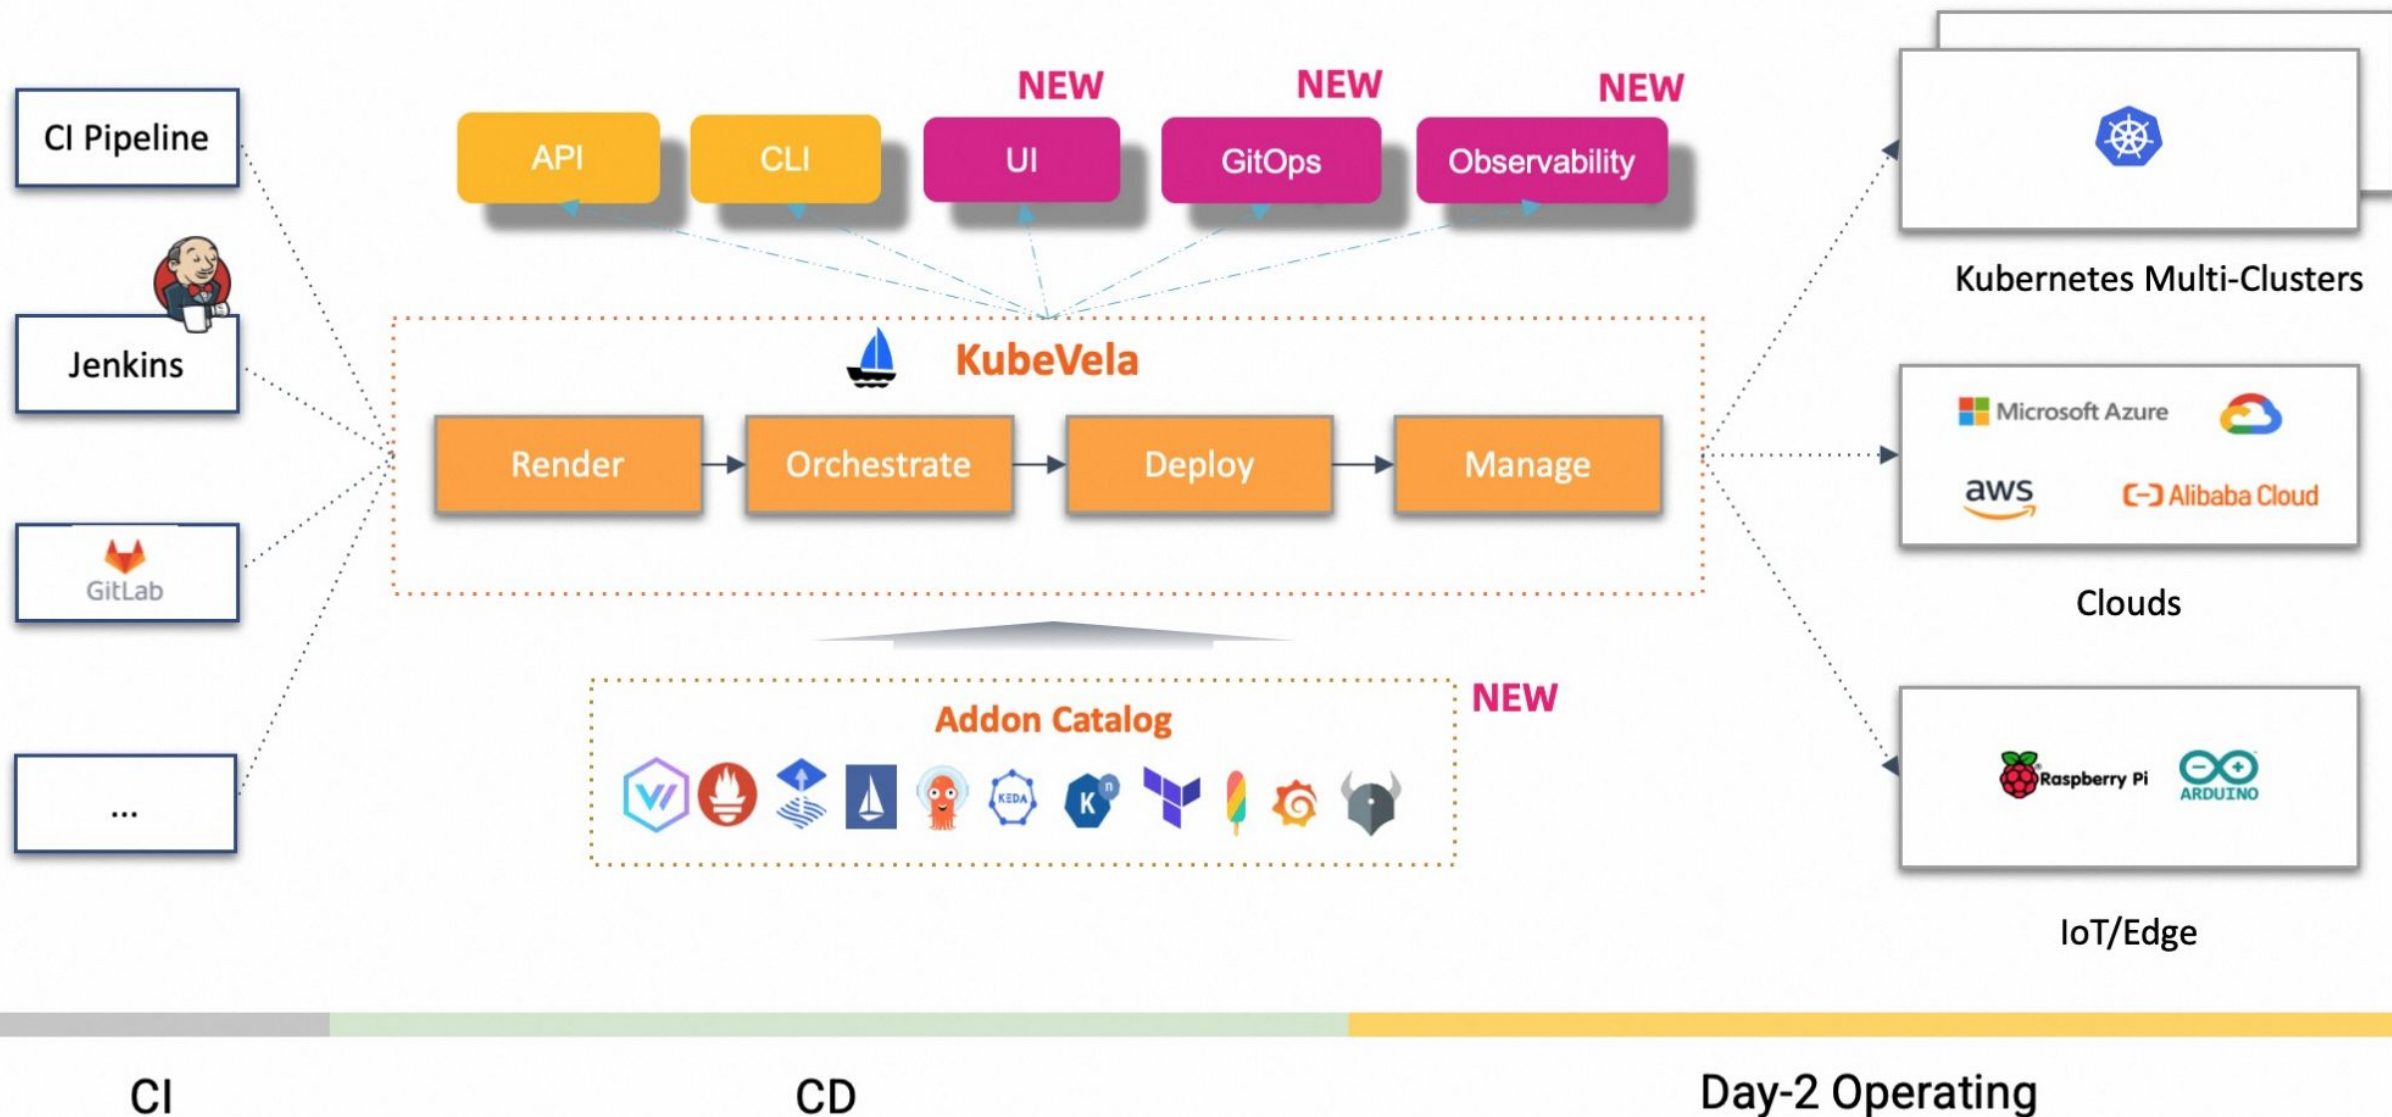
- Offer a top-down approach to reason about applications
- An application becomes the focus of the system
 - All resources related to an application are automatically identified
 - Status information can be easily aggregated

KubeVela


- K8s OAM runtime
- CNCF Incubating project
- Multi-tenant
- Multi-cluster
- Extensible
 - Custom definitions
 - Addons



<https://kubvela.io>



Project timeline



Anatomy of an application

- Application is the top-level entity
- Application components are associated with the different microservices
- Traits modify the behaviour of the components

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: nginx-app
spec:
  components:
    - name: nginx
      type: webservice
      properties:
        image: nginx:1.20.0
        ports:
          - port: 80
            expose: true
      traits:
        - type: gateway
          properties:
            http:
              "/": 80
```

- Typically associated with a microservice
 - But could be: configuration, resources, or even infra...
- Default types
 - webservice
 - worker
 - daemon
 - cron-task
 - task
- Each type defines their own parameters

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: nginx-app
spec:
  components:
    - name: nginx
      type: webservice
      properties:
        image: nginx:1.20.0
      ports:
        - port: 80
          expose: true
```

- Modifies or augment the component functionality
 - Adding extra elements linked with the deployment
 - Ingresses, secrets, etc.
 - Patching the underlying deployment
 - Request quotas, default labels, extra annotations, etc.
 - Add sidecars
 - Export logs

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: nginx-app
spec:
  components:
    - type: webservice
      properties:
        ...
  traits:
    - type: gateway
      properties:
        http:
          "/": 80
```

- Similar in concept to traits, but are applied to the whole application
 - Global configuration
 - Multi-cluster deployment configuration
 - Health checks
 - Integration with third party applications (e.g., monitoring)

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: example-app-policy
spec:
  components:
    - name: hello-world-server
      type: webservice
      properties:
        ...
  policies:
    - name: health-policy-demo
      type: health
      properties:
        probeInterval: 5
        probeTimeout: 10
```

- Define how the application should be deployed
- Composed of WorkflowSteps
- Supports basic if-then-else logic
- Use case
 - Component dependency
 - Preload data before processing
 - Interaction with other subsystems upon deployment
 - And much more

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: app-with-workflow
spec:
  components:
  - name: express-server
    type: webservice
    properties:
      ...
    traits:
      ...
  workflow:
    steps:
    - name: express-server
      type: apply-component
      properties:
        component: express-server
```

- Easy definition of custom definitions of any type
 - Components
 - Traits
 - Policies
 - WorkflowStep
- Defined with CUE
 - Context information
- Can replace simple operators
- Adapt OAM to our use cases

```
apiVersion: core.oam.dev/v1beta1
kind: TraitDefinition
metadata:
  name: myscaler
spec:
  workloadRefPath: spec.workloadRef
  schematic:
    cue:
      template: |
        outputs: scaler: {
          apiVersion: "core.oam.dev/v1alpha2"
          kind: "ManualScalerTrait"
          spec: {
            replicaCount: parameter.replicas
          }
        }
      parameter: {
        //+short=r
        //+usage=Replicas of the workload
        replicas: *1 | int
      }
    }
```


Installing KubeVela

- Create a kind cluster

```
kind create cluster --config=installation/basic_cluster/kind_basic_cluster_config.yaml --name=kubevela
kubectl --context kind-kubevela wait --for=condition=Ready nodes --all --timeout=600s
kubectl --context kind-kubevela apply -f
https://raw.githubusercontent.com/kubernetes/ingress-nginx/main/deploy/static/provider/kind/deploy.yaml
```

- Install KubeVela

```
helm repo add kubevela https://charts.kubevela.net/core
helm repo update
helm install --create-namespace -n vela-system kubevela kubevela/vela-core --wait
```

How to get started with OAM

`doc/01.install_kubevela.md`



<https://github.com/napptive/kubecon-23-oam-kubevela-tutorial>

Vela CLI

- Tool to facilitate the interaction with OAM entities

```
curl -fsSl https://kubeverla.net/script/install.sh | bash  
vela comp
```

- You can still use kubectl to retrieve the information

```
kubectl --context kind-kubeverla -n vela-system get componentdefinitions.core.oam.dev
```

Basic operations

How to get started with OAM

`doc/02.deploy_basic_app.md`



<https://github.com/napptive/kubecon-23-oam-kubevela-tutorial>

Deploying the first application

```
vela env init kubecon --namespace kubecon  
  
vela up -f scenarios/basic_app/nginx-app.yml
```

```
apiVersion: core.oam.dev/v1beta1  
kind: Application  
metadata:  
  name: nginx-app  
spec:  
  components:  
    - name: nginx  
      type: webservice  
      properties:  
        image: nginx:1.20.0  
        ports:  
          - port: 80  
            expose: true  
      traits:  
        - type: gateway  
          properties:  
            http:  
              "/": 80
```


- Check the application status

```
vela status nginx-app -n kubecon
```

```
kubectl --context kind-kubevela -n kubecon get app
```

- Access the logs

```
vela logs nginx-app -n kubecon
```

- Deleting the app

```
vela delete nginx-app
```

```
kubectl --context kind-kubevela -n kubecon delete app nginx-app
```

`doc/03.deploy_complex_app.md`



<https://github.com/napptive/kubecon-23-oam-kubevela-tutorial>

Multi-component applications

```
$ vela up -f scenarios/complex_app/wordpress.yml
```

```
$ kubectl -n kubecon get all -l app.oam.dev/name=my-wordpress
```

NAME	READY	STATUS	RESTARTS	AGE
pod/mysql-77f7db94b7-j9s4n	1/1	Running	0	5m
pod/wordpress-7b4b47ff64-97rw9	1/1	Running	1 (3m39s ago)	5m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/wordpress	ClusterIP	10.43.239.80	<none>	8080/TCP	5m
service/mysql	ClusterIP	10.43.61.129	<none>	3306/TCP	5m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/mysql	1/1	1	1	5m
deployment.apps/wordpress	1/1	1	1	5m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/mysql-77f7db94b7	1	1	1	5m
replicaset.apps/wordpress-7b4b47ff64	1	1	1	5m

Defining component dependencies

- Option 1
 - Use the dependsOn clause on the components
 - But it is too coupled with the component definition

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: my-wordpress-dep
  annotations:
    version: v1.0.0
spec:
  components:
    - name: wordpress
      type: webservice
      dependsOn:
        - mysql
      properties:
        image: bitnami/wordpress:latest
        ...
    - name: mysql
      type: webservice
      properties:
        image: mysql:8.0.32
        ...
```

VelaUX

Continuous Delivery

 Applications

 Environments


 Pipelines

Resources

 Clusters

 Targets

Extension

 Addons

 Definitions

Platform

 Users

 Roles

 Projects

 Configs

 [Projects](#) / [Default](#)

Summary

Applications

Pipelines

Roles

Members

Search by Environment

Search by Target

Search by name and descriptor



Card

Table

New Application



remote-deployment



Automatically converted from KubeVela
Application in Kubernetes.

2023/04/11 16:54:41

doc/04.install_velaux.md



<https://github.com/napptive/kubecon-23-oam-kubevela-tutorial>

- Provided as an addon

```
vela addon enable velaux
```

- Open the dashboard

Initialized admin username and password: `admin / VelaUX12345`

To open the dashboard directly by port-forward:

```
vela port-forward -n vela-system addon-velaux 9082:80
```

Select `"local | velaux | velaux"` from the prompt.

Application workflows

doc/05.app_workflows.md



<https://github.com/napptive/kubecon-23-oam-kubevela-tutorial>

Defining component dependencies

- Option 2
 - Define a deployment workflow
 - Manage not only dependencies but intermediate processes
 - Pre-loading data
 - Send requests to other systems
 - Slack notifications
 - And much more

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: my-wordpress-wf
spec:
  components:
    - name: wordpress ...
    - name: mysql ...
    - name: discoversvc ...
  workflow:
    steps:
      - name: deploy-fake-discover-svc
        type: apply-component
        properties:
          component: discoversvc
      - name: deploy-mysql
        type: apply-component
        properties:
          component: mysql
      - name: deploy-wordpress
        type: apply-component
        properties:
          component: wordpress
    ...
```

Application workflows

my-wordpress-wf-v1(14 hours ago) ▾

Launch Workflow Studio

my-wordpress-wf-v1

Started at: 2023/04/11 21:15:10


Duration: 13s

Mode: StepByStep-DAG

Revision: my-wordpress-wf-v1



✓ SUCCEEDED



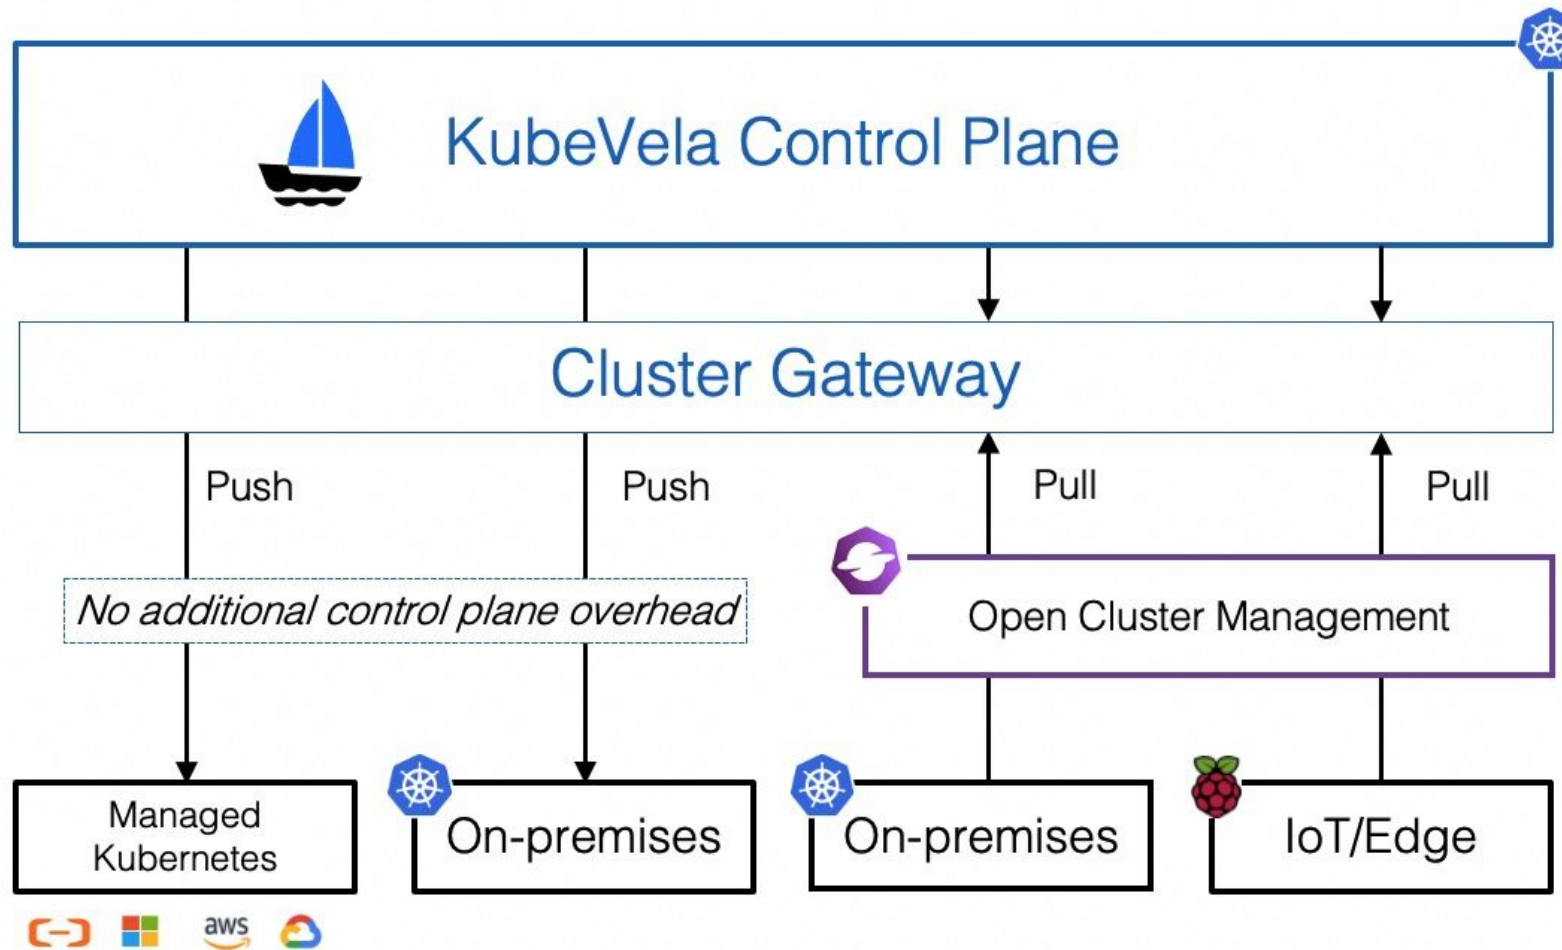
Multi-cluster

Multi-cluster deployments

- Oriented to continuous delivery use cases
- Set the target clusters in the application
- When combined with workflows and addons we can provision clusters and managed services before deployment

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: remote-deployment
spec:
  components:
    # app components
  policies:
    - name: target-clusters
      type: topology
      properties:
        clusters: ["cluster-1", "cluster-2"]
```

Multi-cluster deployments

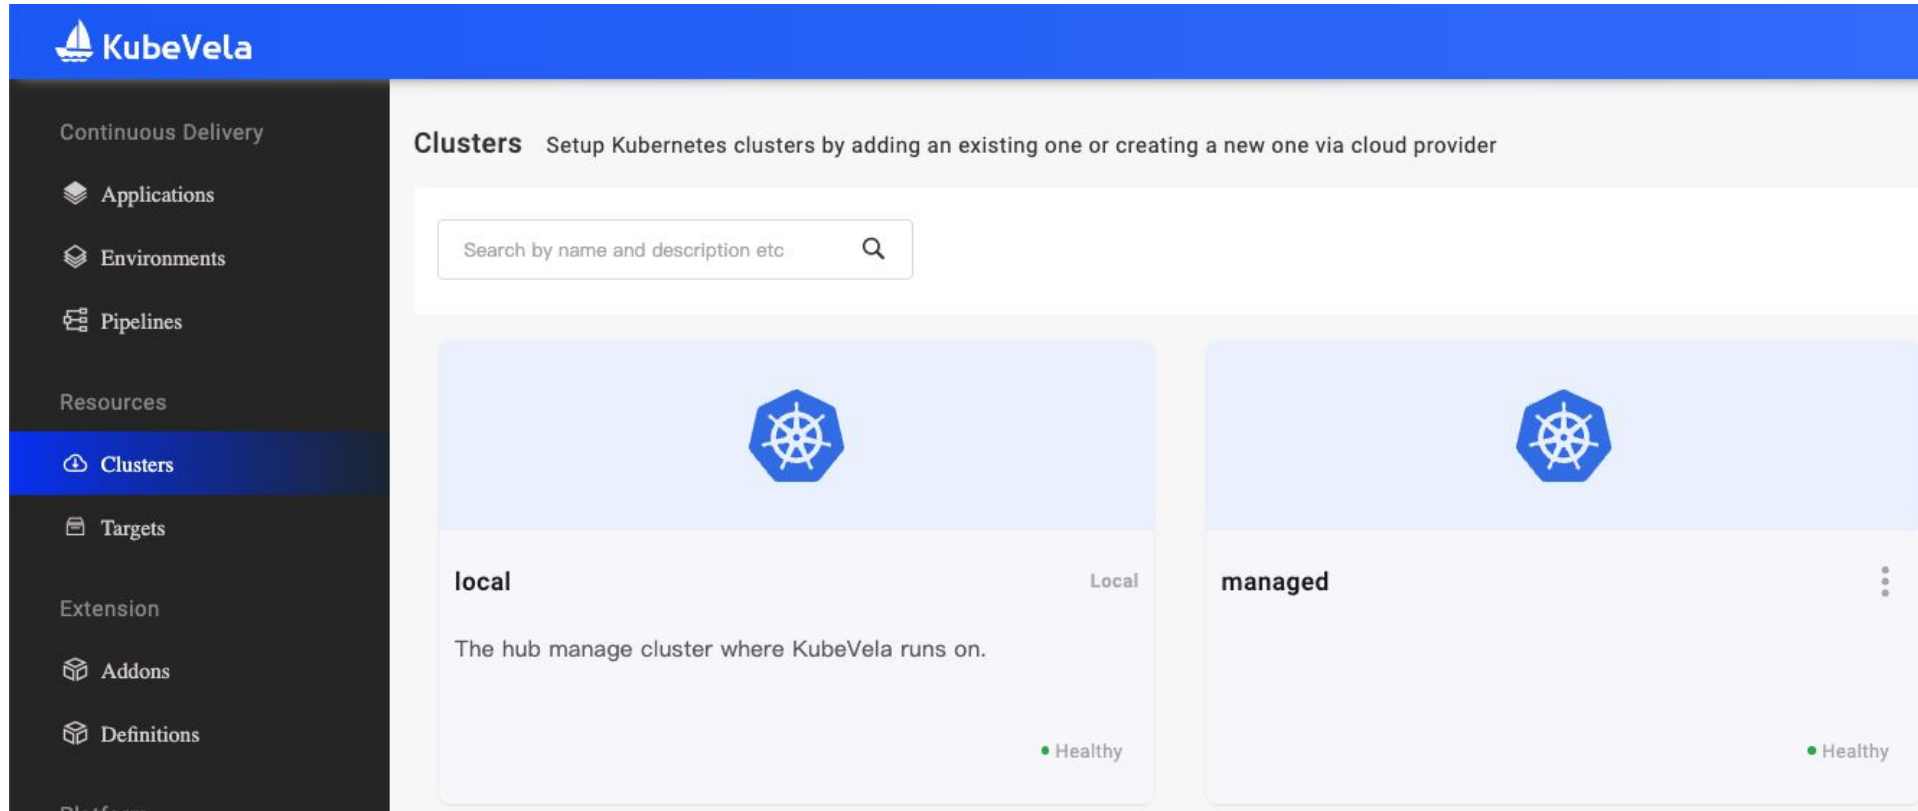


doc/06.multicluster.md



<https://github.com/napptive/kubecon-23-oam-kubevela-tutorial>

Multi-cluster install



The screenshot shows the KubeVela web interface. On the left is a dark sidebar with navigation links: Continuous Delivery, Applications, Environments, Pipelines, Resources, Clusters (highlighted), Targets, Extension, Addons, Definitions, and Platforms. The main content area has a blue header with the KubeVela logo. Below the header, the 'Clusters' section is active, showing a subtitle: 'Setup Kubernetes clusters by adding an existing one or creating a new one via cloud provider'. A search bar is present. Two cluster cards are displayed: 'local' (Local) and 'managed'. Both cards show a Kubernetes logo, a description, and a 'Healthy' status indicator.

KubeVela

Continuous Delivery

Applications

Environments

Pipelines

Resources

Clusters

Targets

Extension

Addons

Definitions

Platforms

Clusters Setup Kubernetes clusters by adding an existing one or creating a new one via cloud provider

Search by name and description etc

local Local

The hub manage cluster where KubeVela runs on.

Healthy

managed

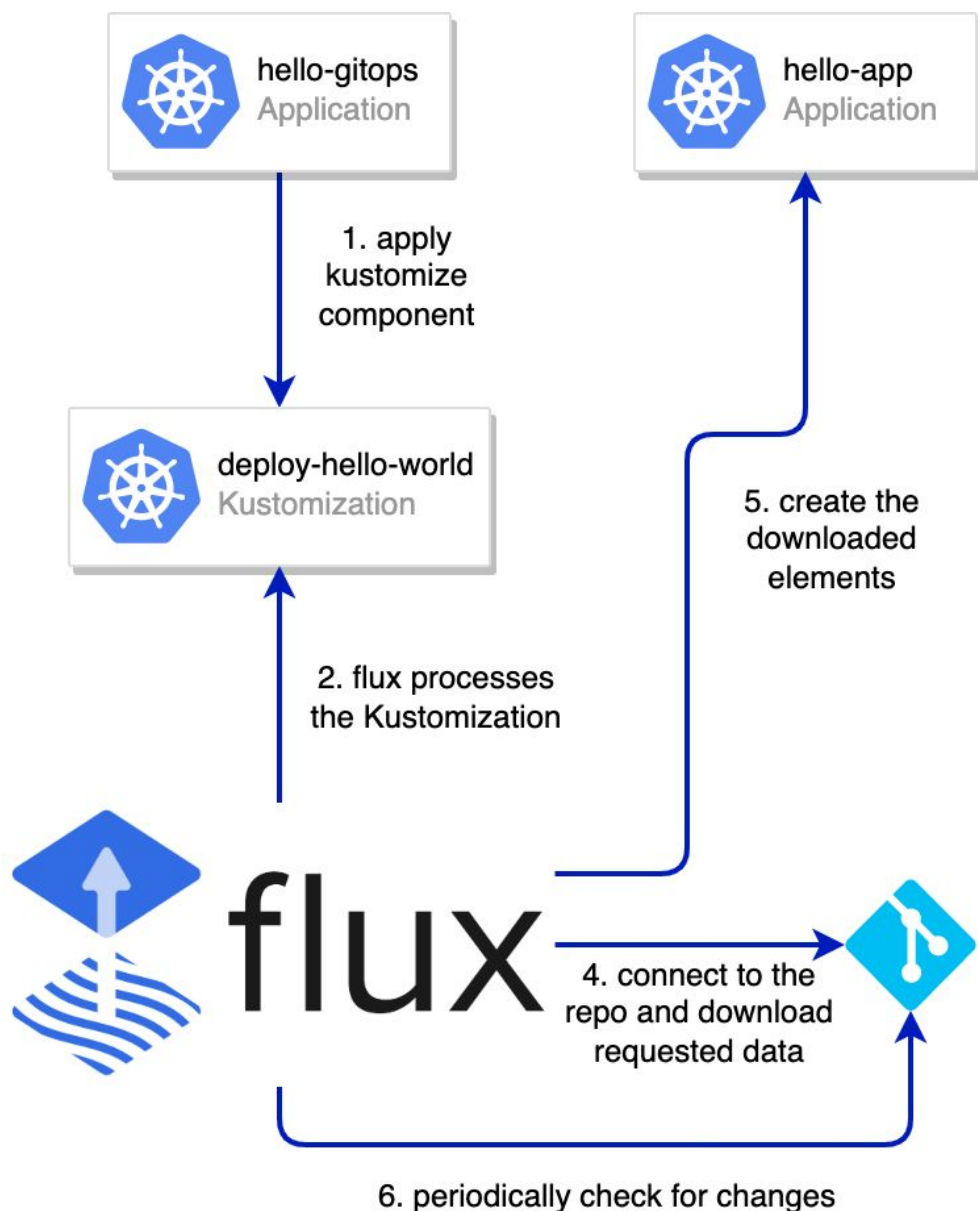
Healthy

GitOps with FluxCD

- The FluxCD addon provides a set of new component definitions
 - Kustomize
 - Helm
- Helm components are a good way to progressively migrate apps to OAM
- Kustomize components enables GitOps deployments

```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: helm-redis
spec:
  components:
    - name: redis
      type: helm
      properties:
        repoType: "helm"
        url: "https://charts.bitnami.com/bitnami"
        chart: "redis"
        version: "16.8.5"
        values:
          master:
            persistence:
              size: 16Gi
          replica:
            persistence:
              size: 16Gi
```

GitOps hello world



```
apiVersion: core.oam.dev/v1beta1
kind: Application
metadata:
  name: hello-gitops
spec:
  components:
  - name: deploy-hello-world
    type: kustomize
    properties:
      targetNamespace: kubecon
      repoType: git
      # replace it with your own repo url to explore further
      url: https://github.com/napptive/kubecon-23-oam-kubevela-tutorial
      # replace it with your git secret if it's a private repo
      # secretRef: git-secret
      # the pull interval time, set to 30s for demo purposes
      pullInterval: 30s
      git:
        # the branch name
        branch: initial_commit
        # the path to sync
        path: ./scenarios/gitops/target
```

Extra content

Addon catalog

Addons Manages and extends platform capabilities

Addon Registries

All

Search by name and de



Tags ☐ Official ☐ Cloud Resource ☐ Terraform Provider ☐ Observability ☐ Kubernetes ☐ Metrics ☐ Prometheus



velaux

KubeVela User Experience (UX). An extensible, application-oriented...

Official

v1.7.6

Enabled



fluxcd

Extended workload to do continuous and progressive delivery

Of...

Extended W...

Gi...

H...

2.3.4

Enabled



cert-manager

Automatically provision and manage TLS certificates in Kubernetes

Let's En...

T...

Certifi...

AC...

2.1.0



chartmuseum

ChartMuseum is an open-source and easy to deploy Helm Chart Repositor...

h...

reg...

helm_...

chartm...

4.1.0



cloudshell

A Friendly Kubernetes CloudShell (Web Terminal)

Official

0.0.8



dex

Enable dex for login

Official

OAuth

0.6.6



flink-kubernetes-operator

A Kubernetes operator for Apache Flink

apache/flink-kubernetes-operator

1.3.1



grafana

Grafana is a multi-platform open source analytics and interactive...

Of...

Observ...

Kuber...

Gr...

v0.4.2



ingress-nginx

gateway for serving ingress traffic

network

gateway

1.1.1



keda

An AutoScaler addon for KubeVela.

Scaler

AutoScaler

2.8.2



kruise-rollout

Rollout workload by kruise controller

Official

Rollout

1.5.0-beta.1



kube-state-metrics

A simple service that listens to the Kubernetes API server and generates...

O...

Obs...

Ku...

M...

Pro...

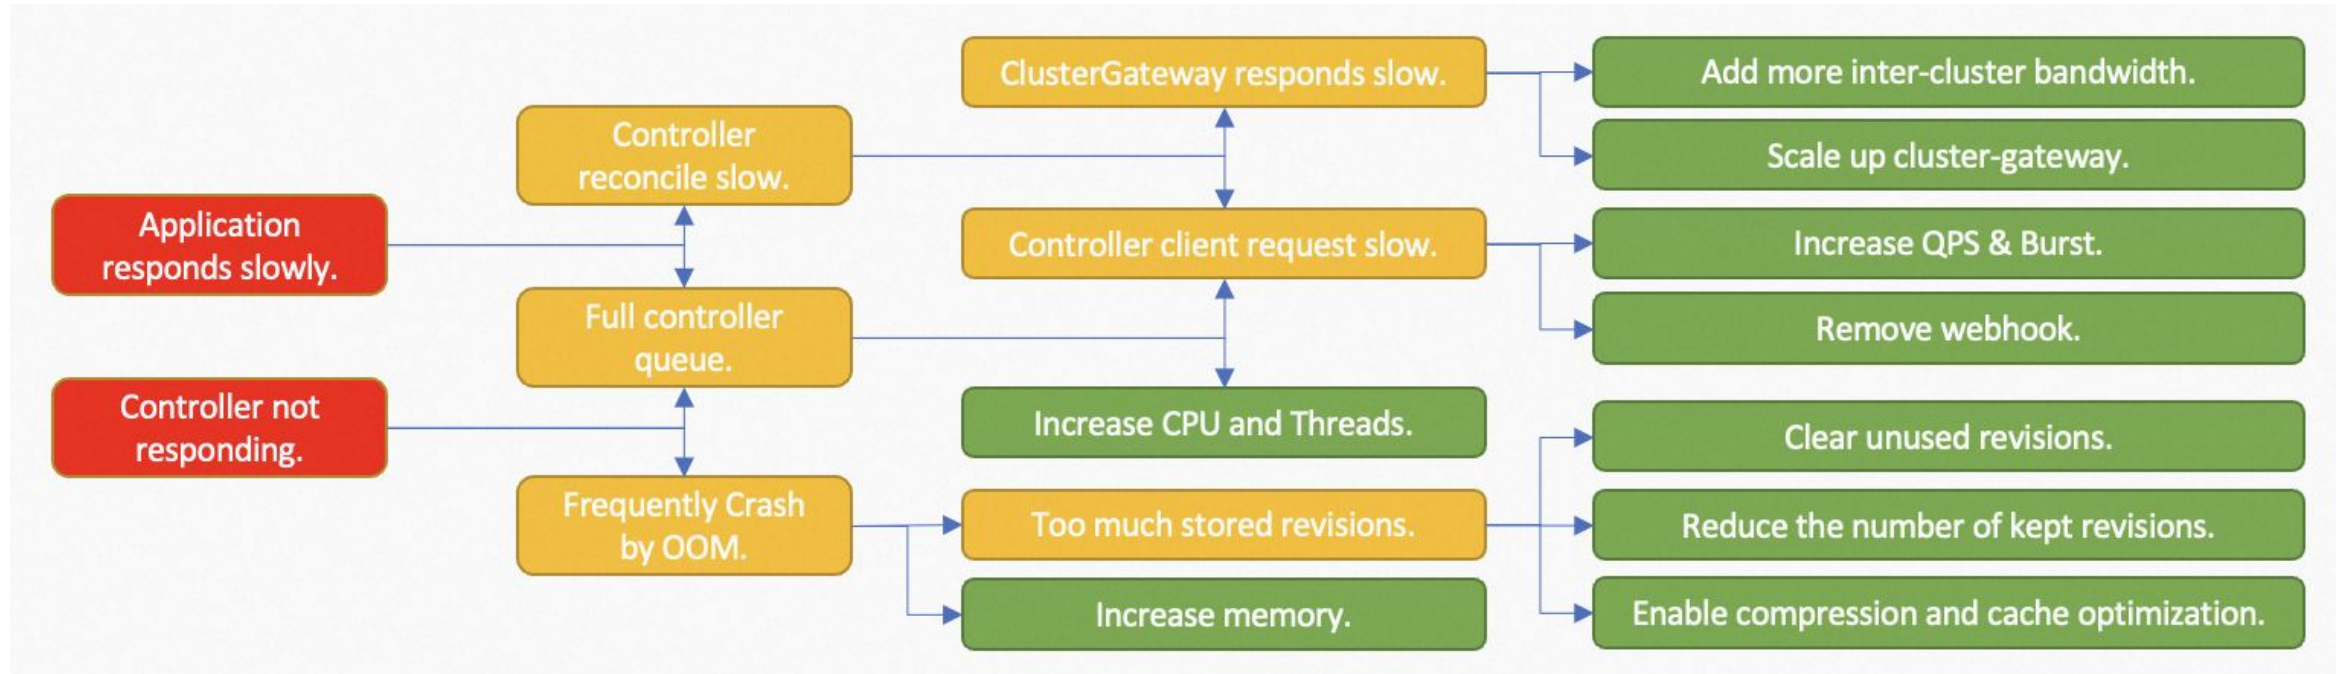
v0.3.0

- KubeVela can be integrated with the underlying K8s RBAC

```
helm upgrade --install kubevela kubevela/vela-core  
--create-namespace -n vela-system  
--set authentication.enabled=true --set authentication.withUser=true  
--wait
```

- Applications will be deployed impersonating the requesting user to guarantee that role limitations are enforced

KubeVela performance troubleshooting



- OAM
 - <https://oam.dev/>
 - <https://github.com/oam-dev/spec>
- KubeVela
 - <https://kubvela.io/>
 - <https://github.com/kubvela/kubvela>
 - #kubvela on the CNCF Slack
 - BiWeekly meetings English & Chinese
- Napptive
 - <https://napptive.com>
 - <https://docs.napptive.com>



KubeCon



CloudNativeCon

Europe 2023

Thanks!





Please scan the QR Code above
to leave feedback on this session



KubeCon



CloudNativeCon

Europe 2023

