



KubeCon



CloudNativeCon

Europe 2023





KubeCon



CloudNativeCon

Europe 2023

Understand Systems with OpenTelemetry: A Hybrid Telemetry Data Backend

Ran Xu, Huawei Cloud & Xiaochun Yang, Northeastern University China



Xu Ran

A database expert in Huawei Cloud Database Innovation Lab and openGemini time-series database architect

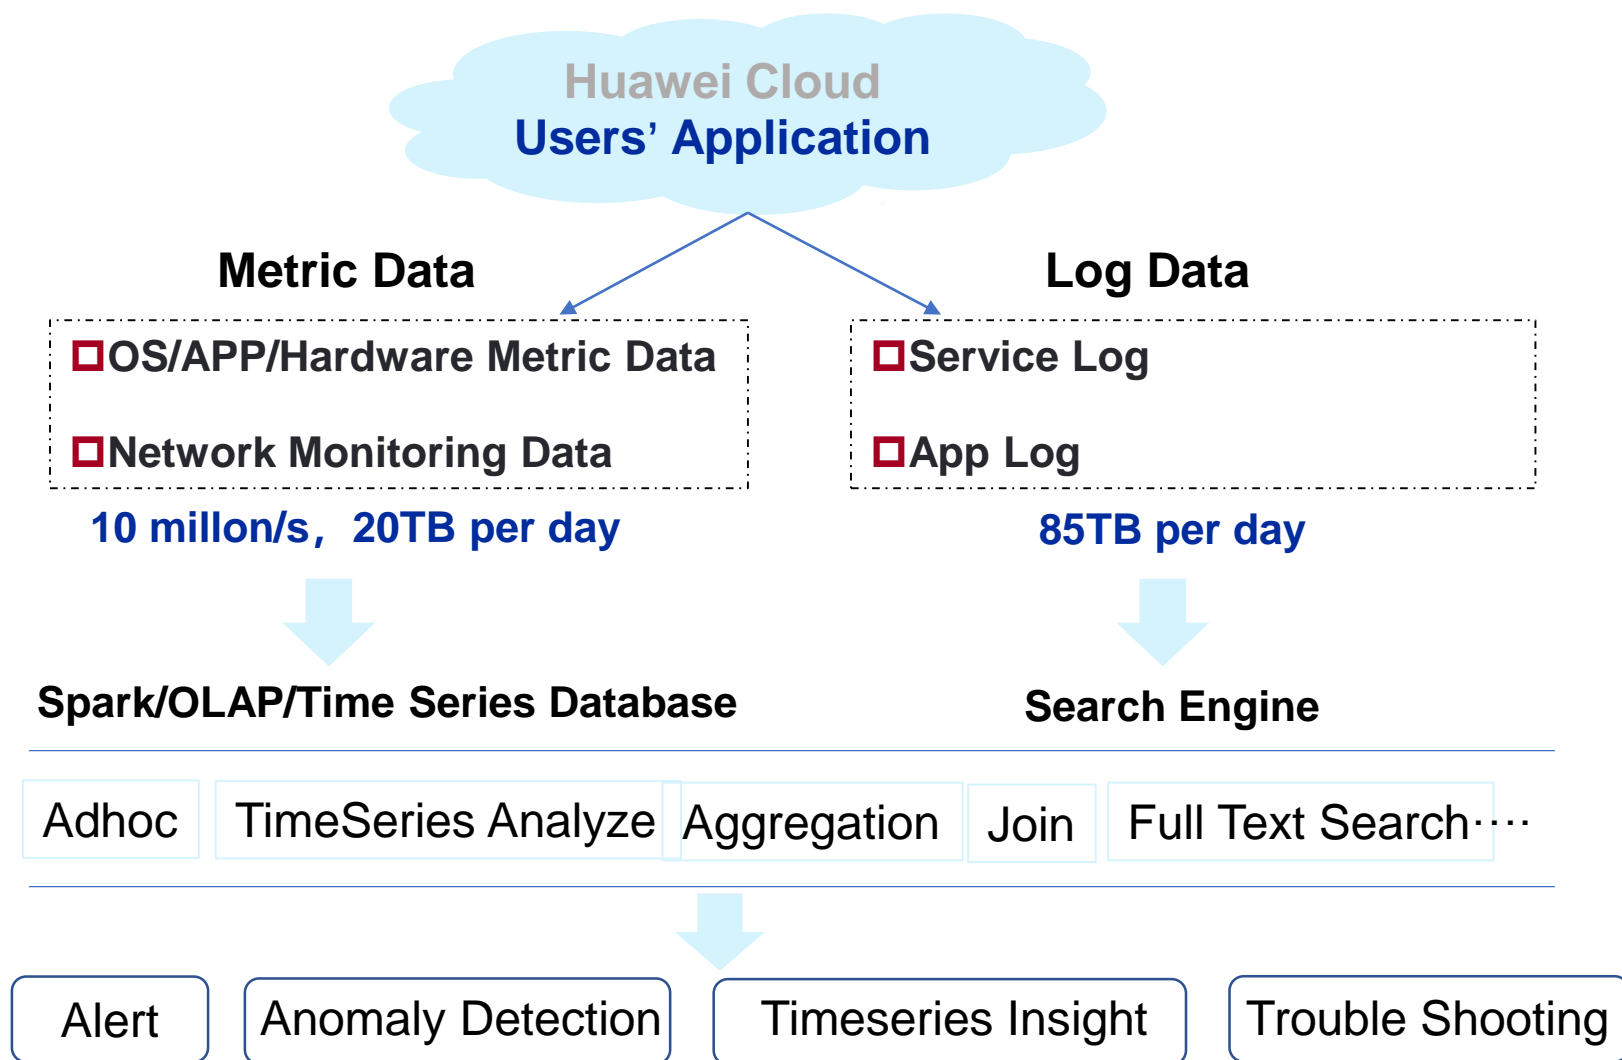


Xiaochun Yang

A professor of Northeastern University in China, has been engaged in teaching and research in the field of data management and analysis for a long time.

- **Challenge of Observability**
- **openGemini as the Backend Storage Solution of OpenTelemetry**
- **How openGemini Works for Storing Logs, Metrics, and Traces**


Challenge: Massive Telemetry Data with Different Technologies



Technical Challenges

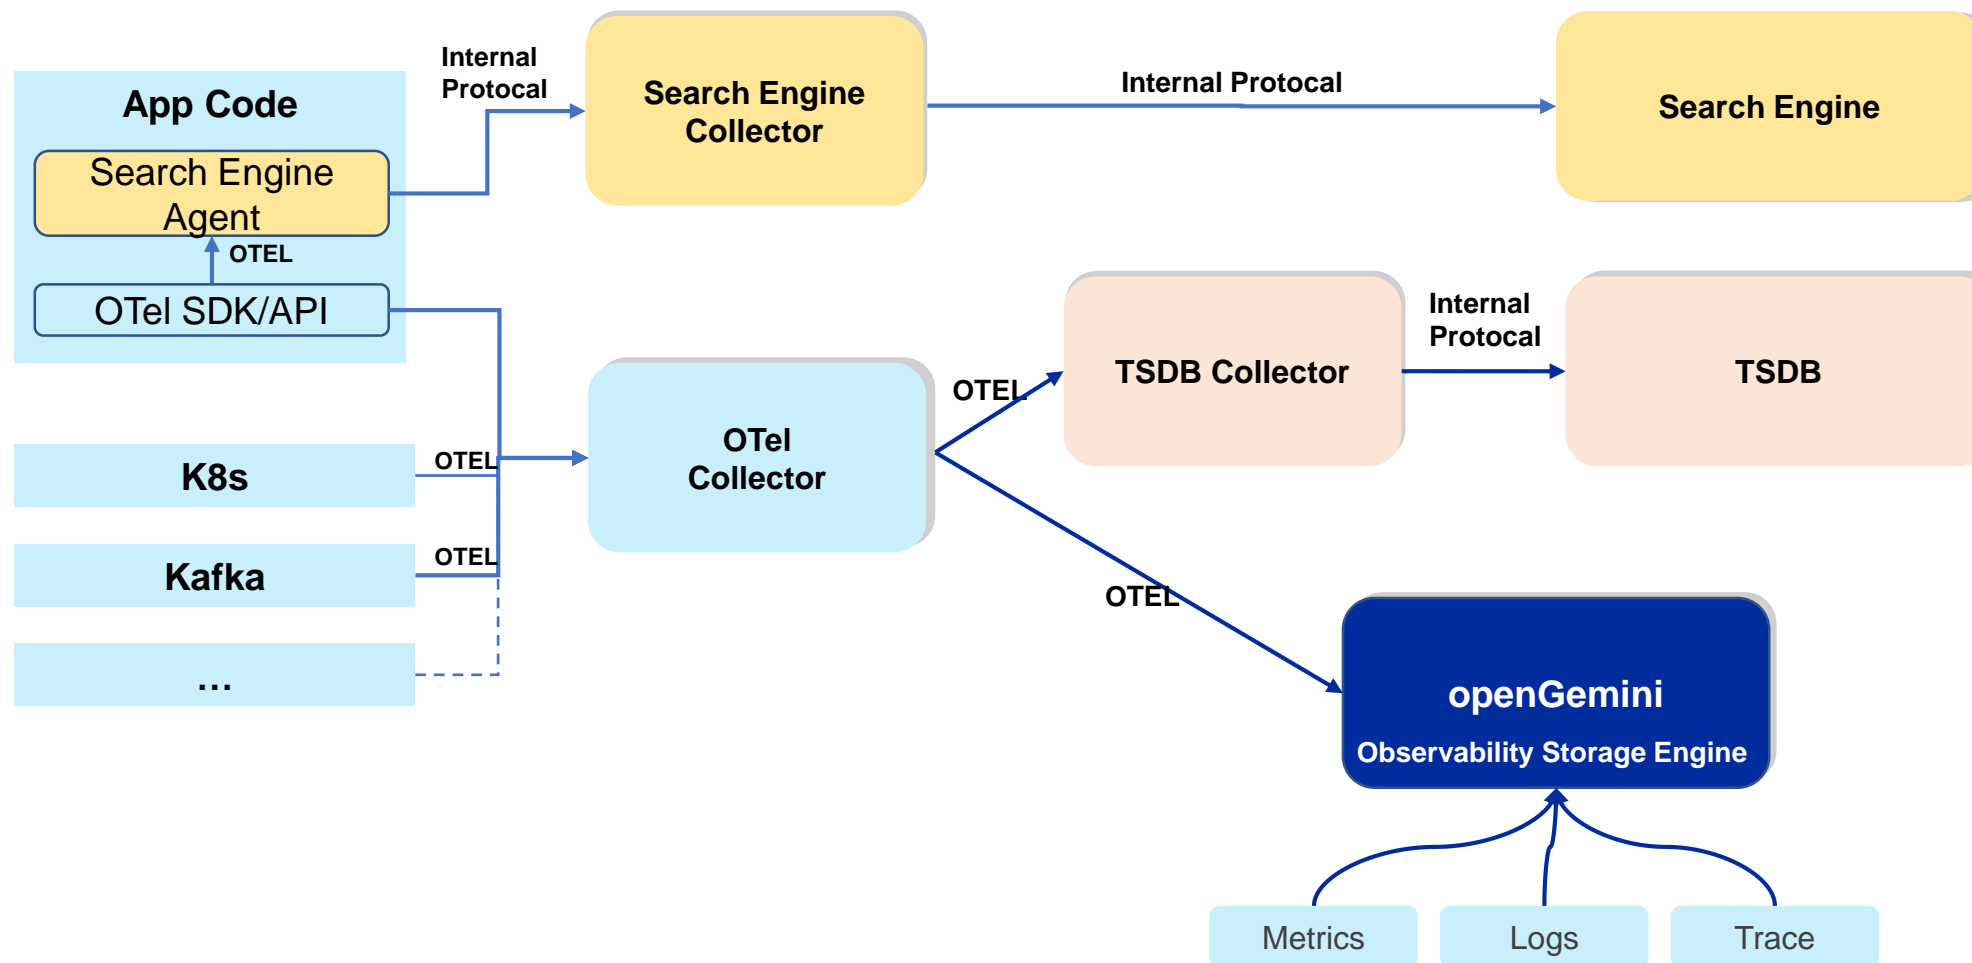
- **Massive Data**
- **Cost Sensitive**
- **Operation**
- **Correlation Analysis**

OTel: Unified Standards, Correlation Between Metrics, Logs, and Traces



- **Unified Access Protocols:** metrics, traces, and logs.
- **Standard Data models**
- **Various SDKs:** 11 common languages, such as C++ and Go, are supported.
- **Open Source**

openGemini: High-performance Backend Storage of openTelemetry

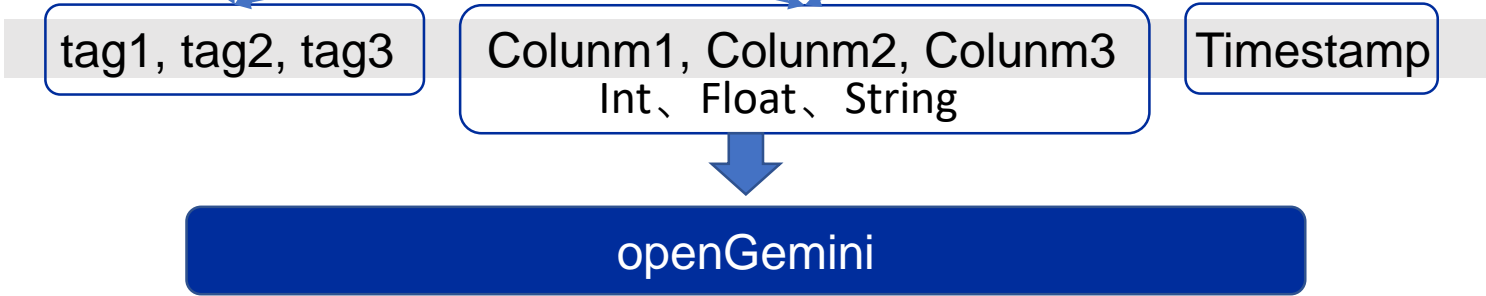


Converts OTEL Format to Storage Schema



OTEL data format sample

Metrics	Traces	Logs
<p>Resource Attributes</p> <p>host.name = "Str(servername)"</p> <p>process.command_args = Slice(["./main"])</p> <p>process.pid: Int(30653)</p> <p>...</p> <p>Data point attributes</p> <p>...</p> <p>http.status_code: Int(200)</p> <p>http.user_agent: Str(Go-http-client/1.1)</p> <p>...</p> <p>Value : 26</p>	<p>Span #0 {</p> <p>trace_id : 467e3765765a170e881f9c765fb94</p> <p>span_id : bf611bd6810518f2</p> <p>parentSpanId: "7ab66566adaca823",</p> <p>name: "HTTP GET",</p> <p>kind: 3,</p> <p>startTimeUnixNano: "1679552089108338000",</p> <p>endTimeUnixNano: "1679552089653340238",</p> <p>Attributes: -> [json]</p> <p>-> http.method: Str(GET)</p> <p>-> http.flavor: Str(1.1)</p> <p>-> http.url: Str(http://0.0.0.0:7080/hello)</p> <p>-> net.peer.name: Str(0.0.0.0)</p> <p>-> ...</p>	<p>Resource attributes</p> <p>-> host.name: Str(xxx-Pro.local)</p> <p>-> process.command_args: Slice(["./main"])</p> <p>-> process.executable.name: Str(main)</p> <p>...</p> <p>-> process.owner: Str(xiangyu)</p> <p>-> service.name: Str(demo-server)</p> <p>-> process.pid: Int(30653)</p> <p>-> telemetry.sdk.language: Str(go)</p> <p>-> telemetry.sdk.name: Str(opentelemetry)</p> <p>-> telemetry.sdk.version: Str(1.13.0)</p> <p>trace_id : 467e3a170e881f9d89639bb5c765fb94</p> <p>span_id : bf611bd6810518f2</p> <p>Body: "Error: connect time out"</p>

unified data schema




Metric Store Evaluation




Performance

Write



Cardinality

Write




Simple Query latency


Complex Query latency

Log Store Evaluation

Index Size




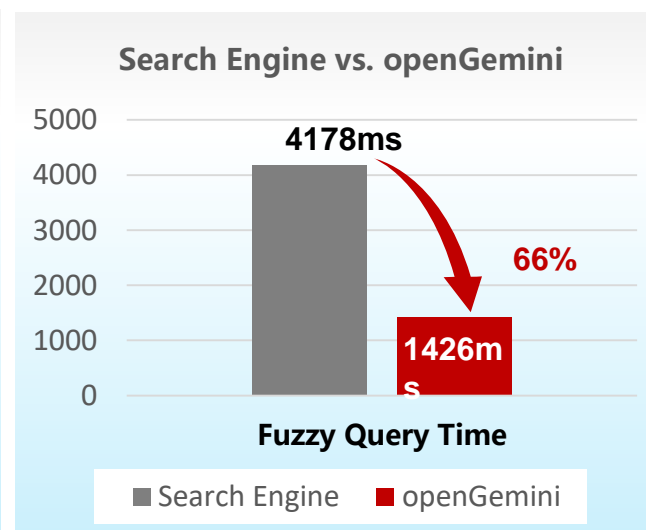
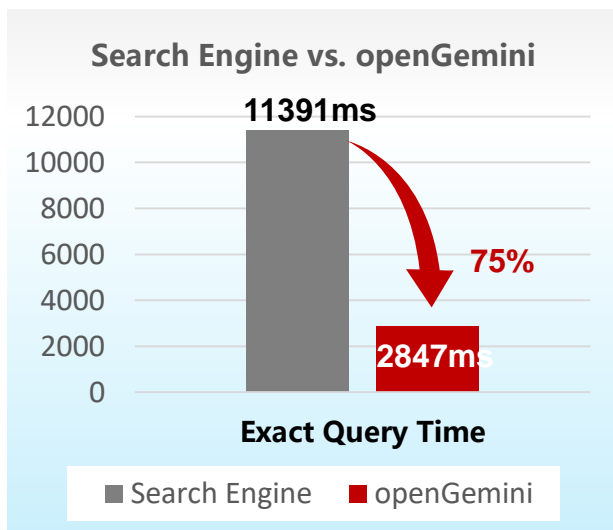
Index Build Time



1 billion data test results

Reduces memory index space by **65%** and index build time by **60%**

Under end-to-end scenarios, for average per query, exact query time reduce by **75%** and fuzzy query time reduce by **66%** and regex query time reduce by **63%**



Query Types For Logs

Exact query

Fuzzy query

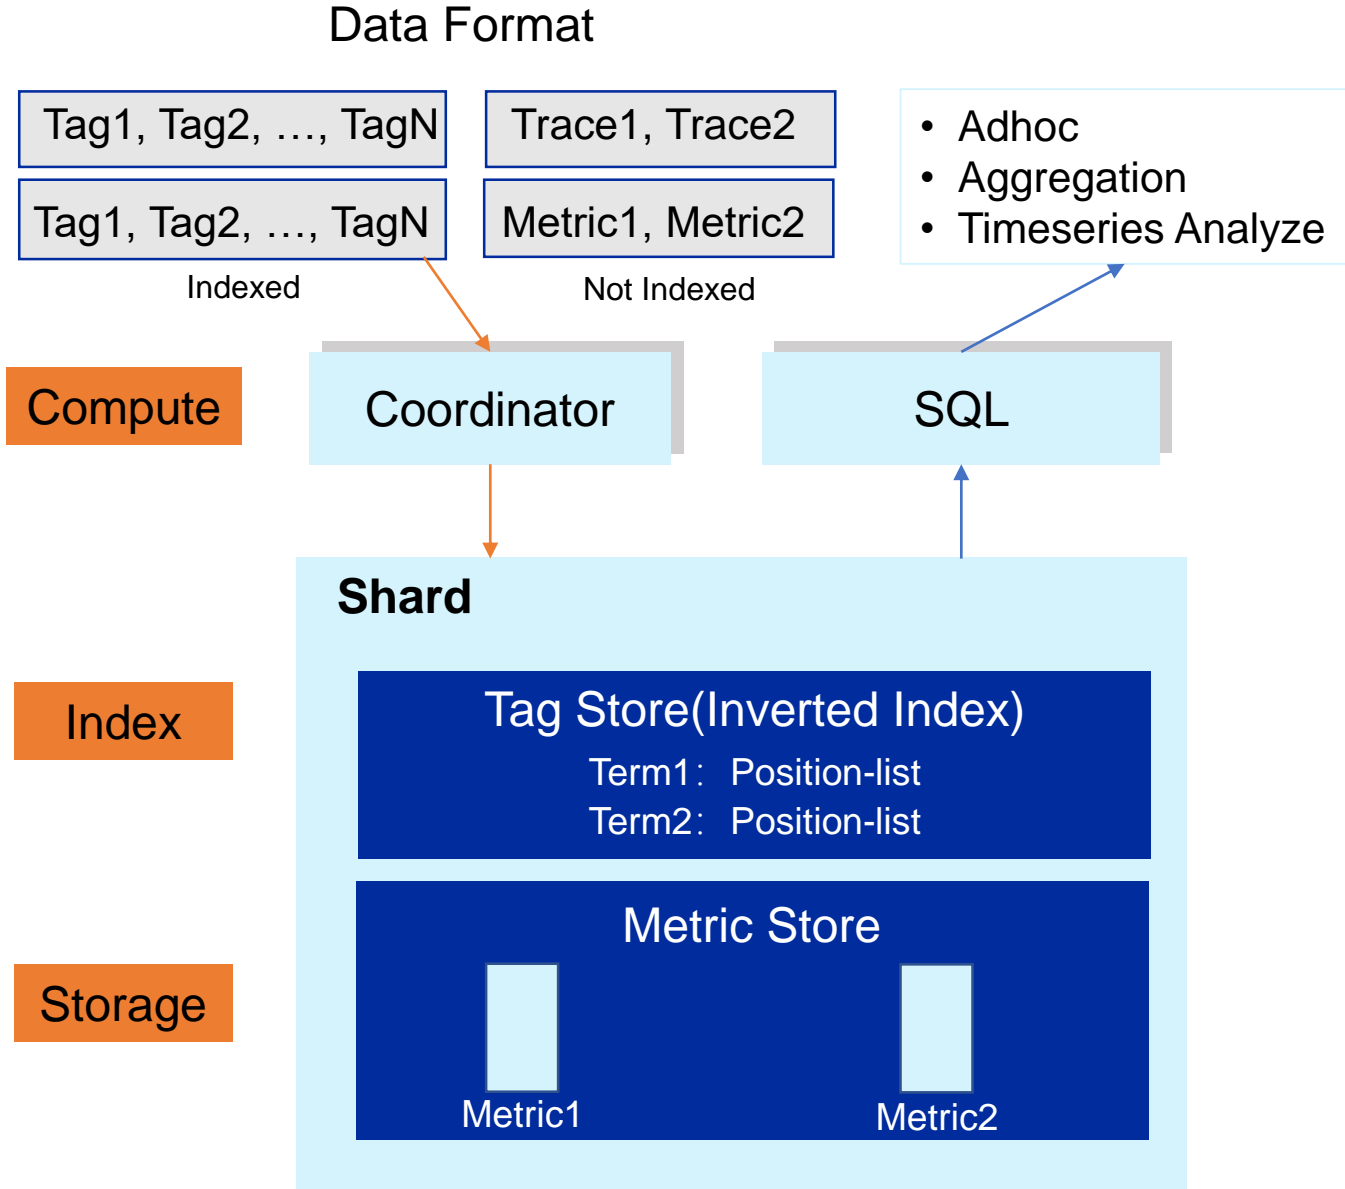
Regex

Prefix/ suffix

score

Average Per Query Time (End-to-end Performance)

1st Metric Data Store



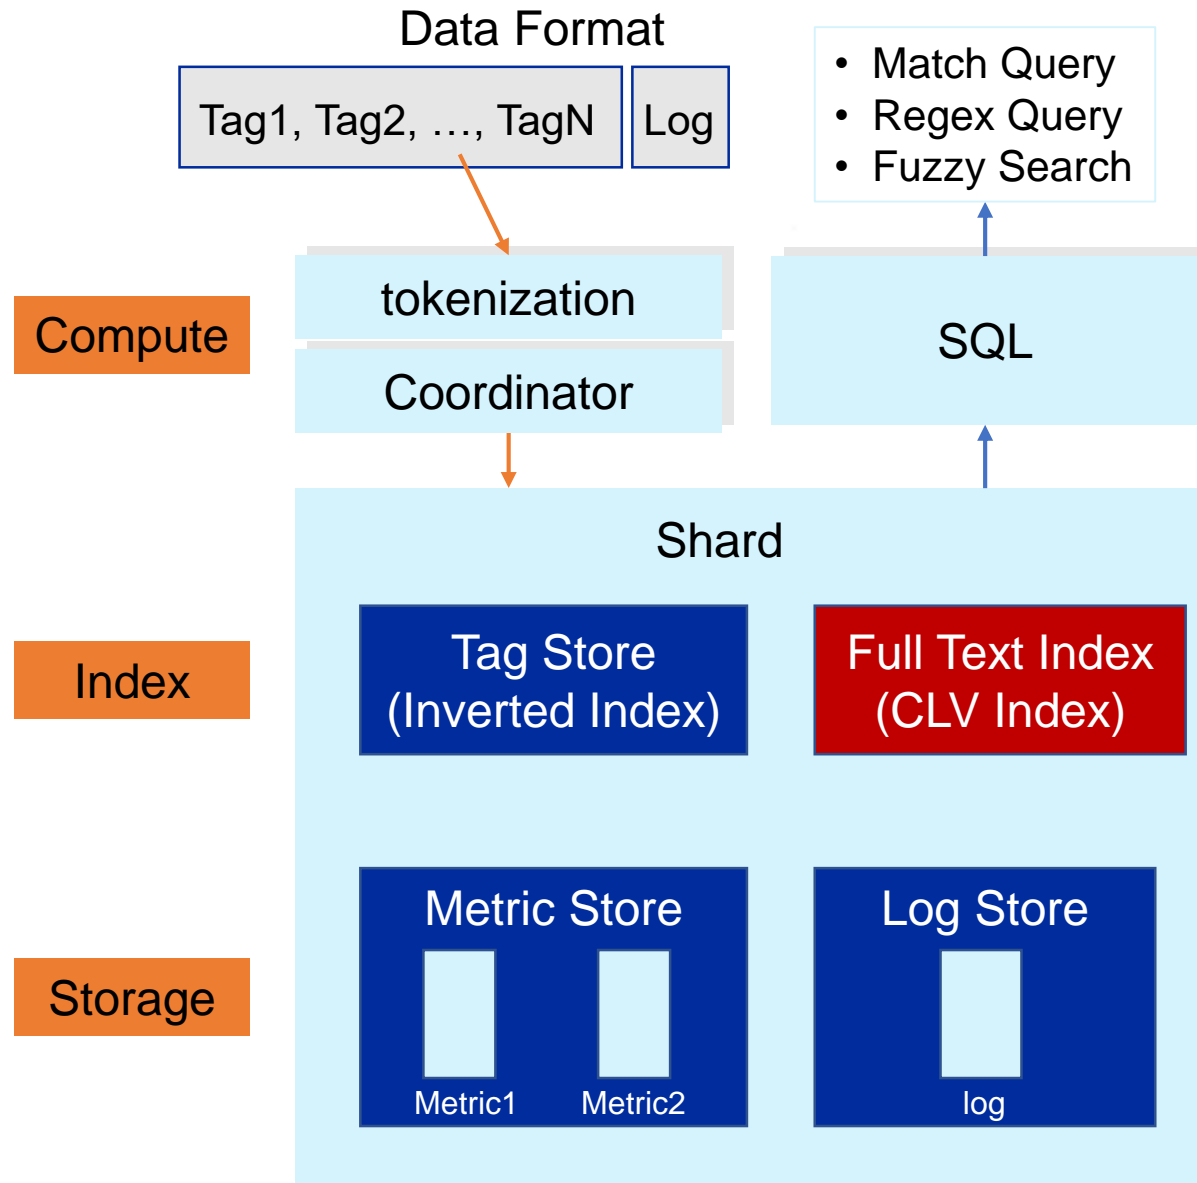
Pros

- Powerful Selectivity Capability
- Easy data compression

Cons

- High Cardinality Tag
- Capability for Logs

2st Log Store

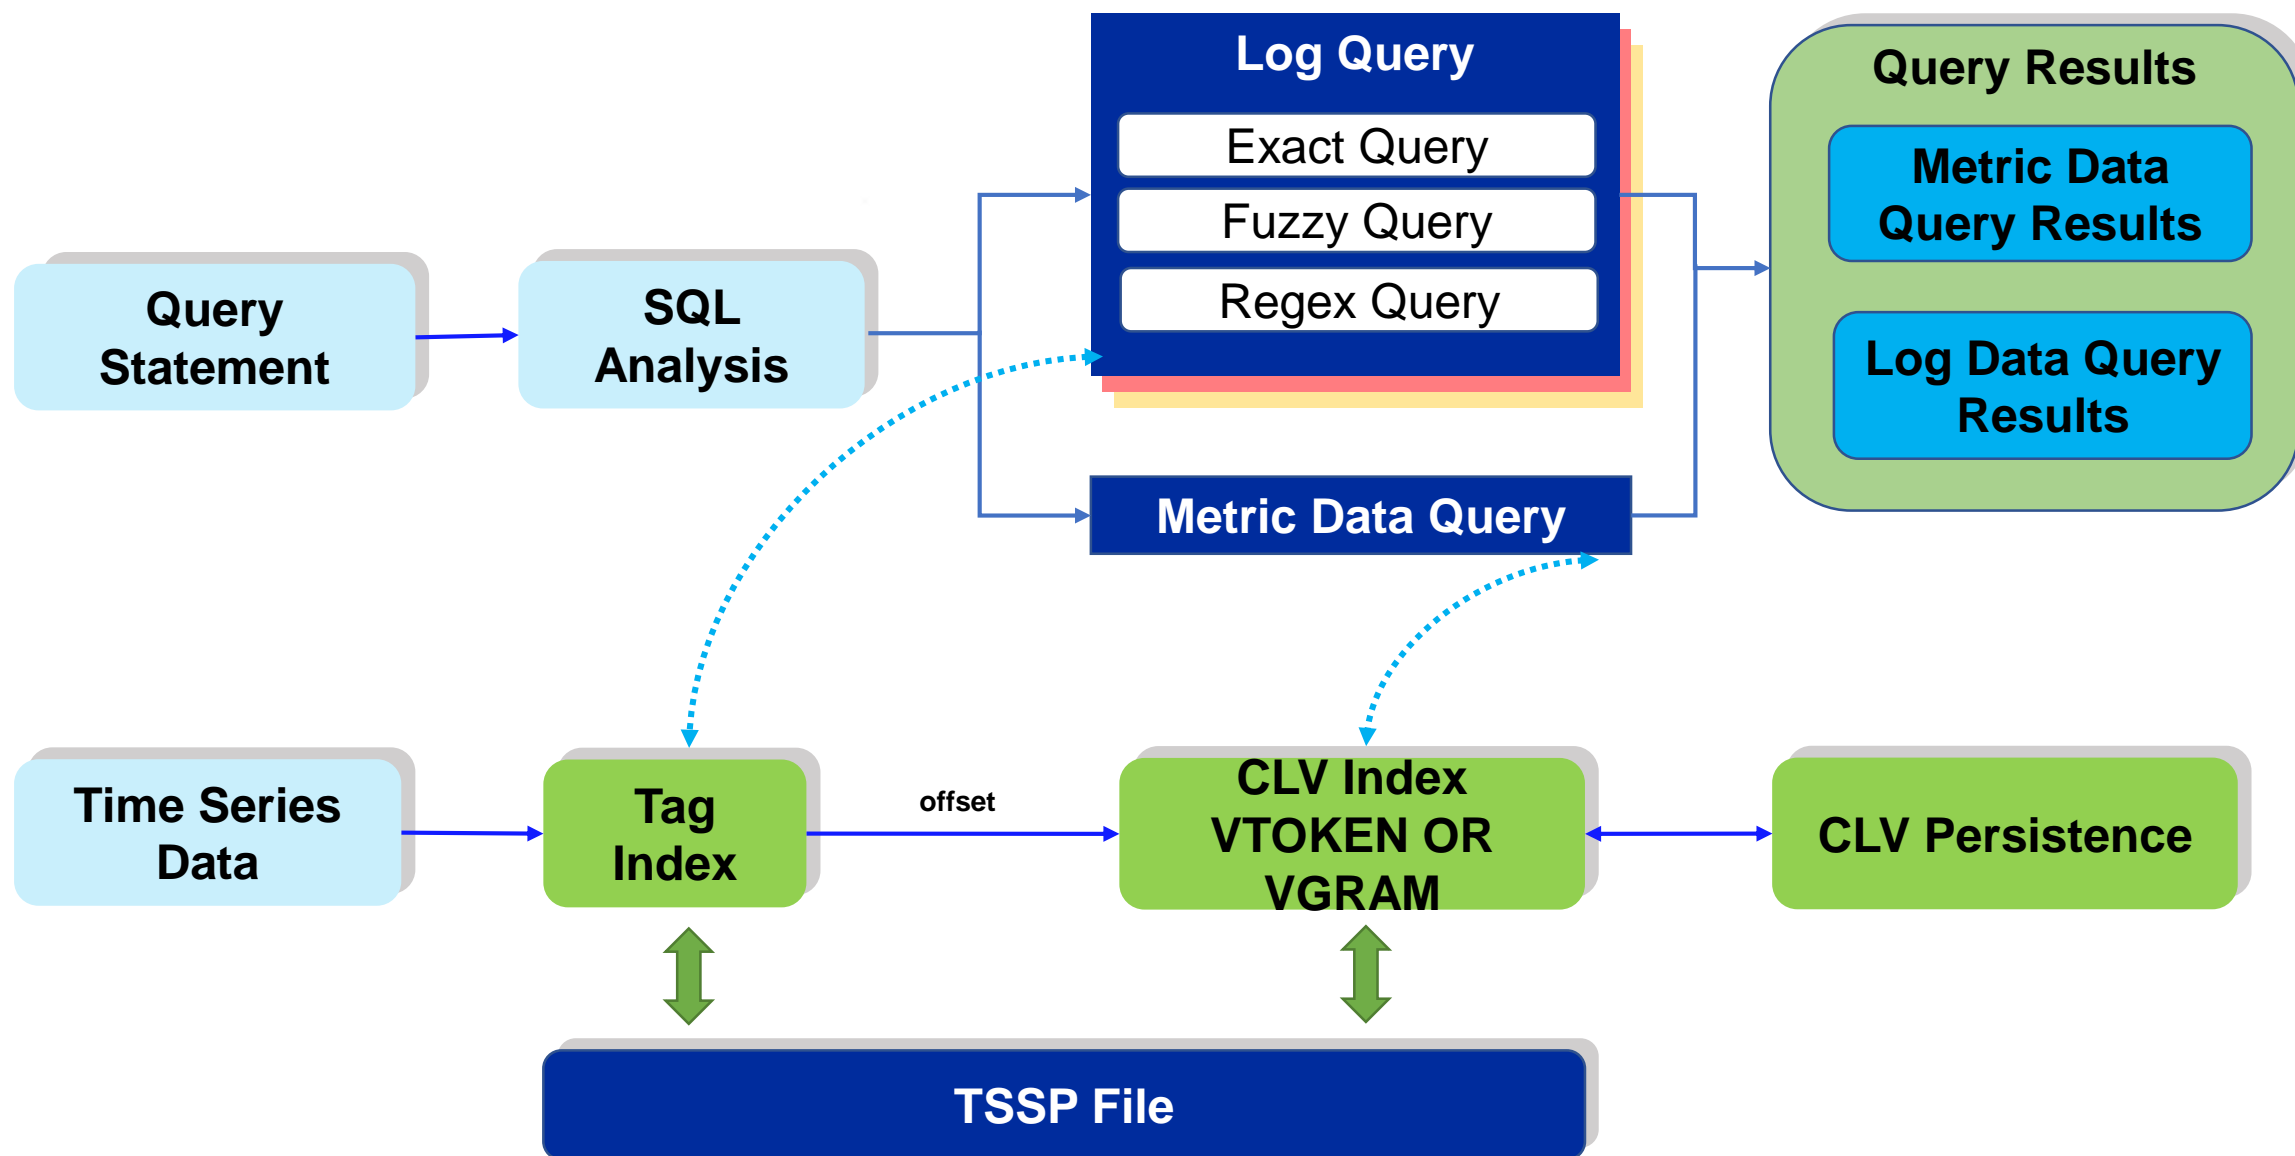


Tag Index + Full Text Index

Column Store

One Engine For Metrics And Logs

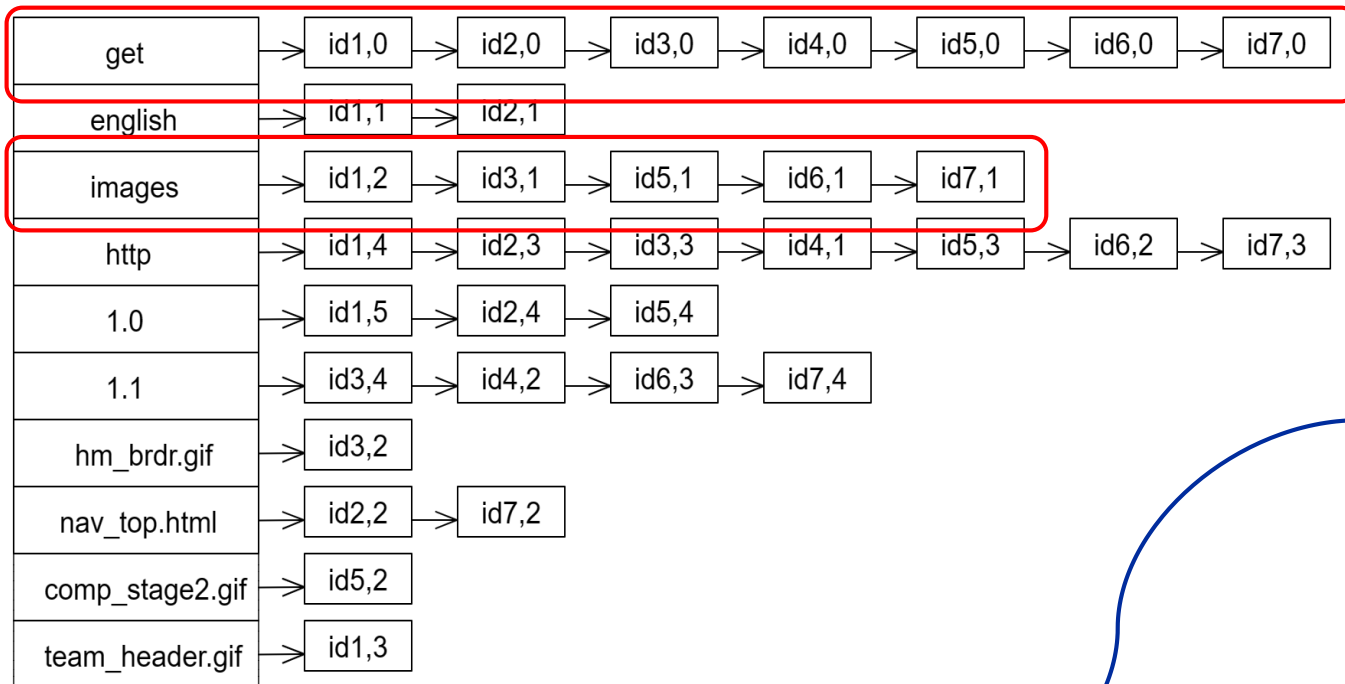
CLV Index: Full Text Search Index For Log



CLV Index: Compressed Indices for Efficient Search

Inverted index

Number of keys: **10** Index list size: **33**

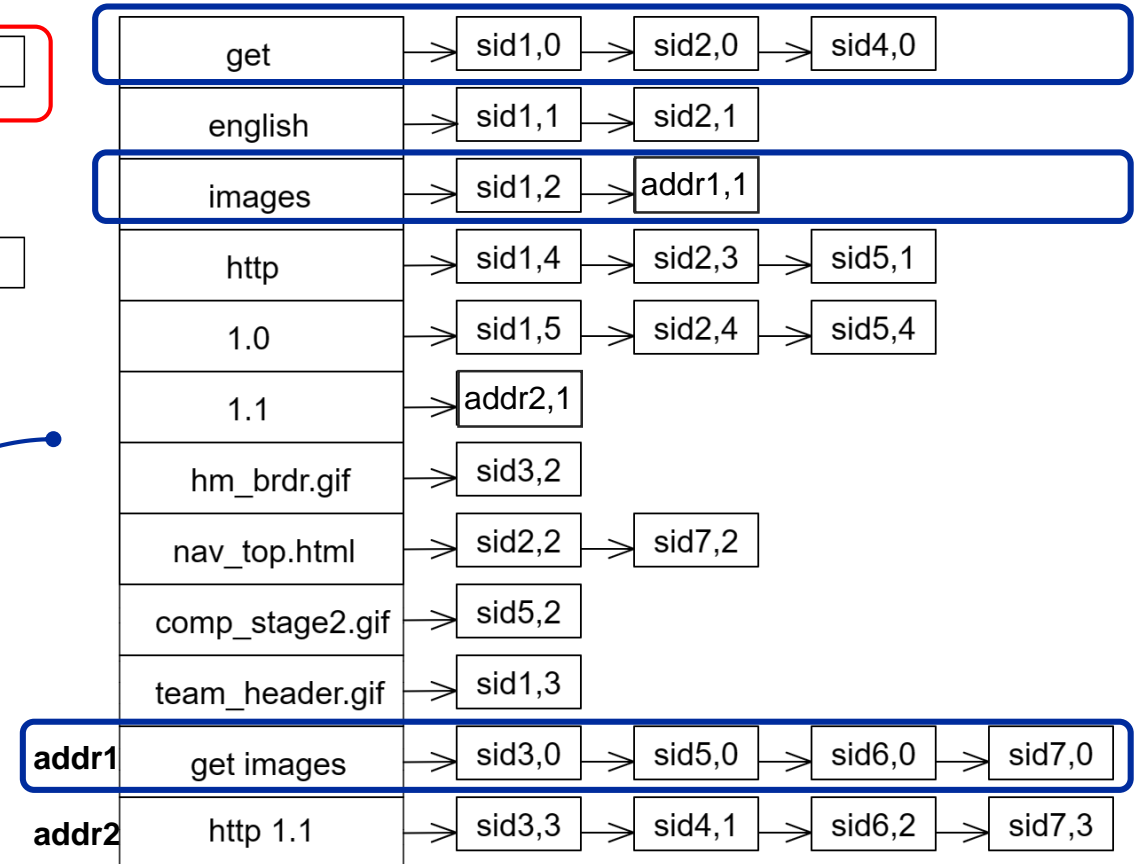


Compressed Inverted Index

- Less index space
- Efficient query time

Compressed inverted index

Number of keys: **12** Index list size: **25 + 2** pointers



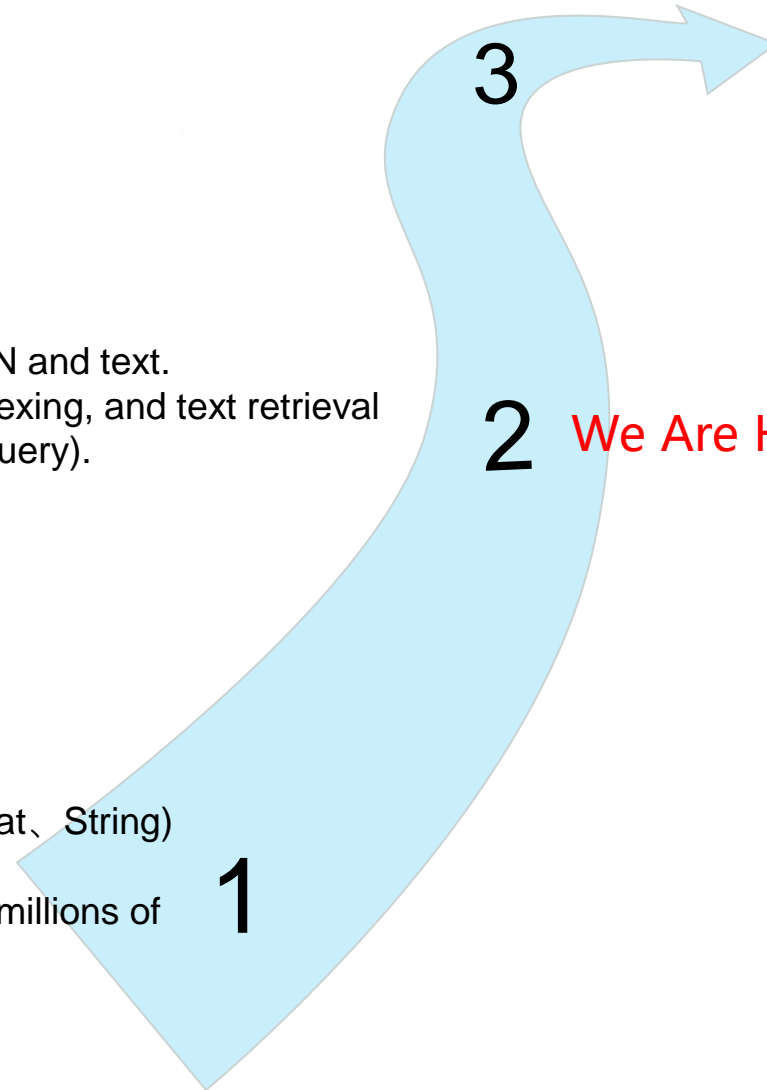
OpenGemini Evolution Roadmap

Log Store

- More Data types such as JSON and text.
- Supports tokenization, text indexing, and text retrieval (exact match 、 fuzzy match query).
- Supports the OTEL protocol.

Metric Store

- Metric data storage and query.(Int、 Float、 String)
- SQL-Like.
- PB-level data storage and hundreds of millions of time series management.



Observability Backend Store

- Tracing data storage
- High cardinality optimization
- Metric、 Log、 Tracing Correlation Analysis

Welcome to try and Give Feedback

```
docker run -d -p 8086:8086 --name openGemini-dev opengemini-db/opengemini-server
```



<https://github.com/openGemini>



https://join.slack.com/t/opengemini/shared_invite/zt-1sncd7nwn-9Mmq2bPVt3b0DP3ng2NOHA



<https://twitter.com/openGemini>



<http://www.openGemini.org>

Learn more about openTelemetry and openGemini

Welcome to Booth : D7



Please scan the QR Code above
to leave feedback on this session