



What Can Go Wrong When You Trust Nobody?

Threat Modelling Zero Trust



James Callaghan

Ric Featherstone

About ControlPlane (Hall 5, Booth SU57)

A Cloud Native security consultancy established in 2017. Our diverse culture empowers and develops individuals with talent and integrity.

We regularly participate in the largest worldwide security conferences.

- Security specialists in cloud, Kubernetes, and containers
- Clients include government, financial services, and regulated industries
- 50 people across the UK, Northern Europe, and Australasia

Overview

How can we:

- derive a **Zero Trust philosophy** from a high-level Threat Model?
- build a **secure architecture** based on these principles?
- carry out a **detailed threat model** of our technical design?
- **understand the risks** posed to our system?
- **define additional controls** to mitigate our risks?
- **demonstrate** our control design using a **prototype**?

This talk will focus on **Zero Trust for workloads**

What is Threat Modelling?


Threat Modelling is:

- Identifying and enumerating **threats and vulnerabilities**
- Devising **mitigations**
- Prioritising **residual risks**
- **Escalating** the most important risks

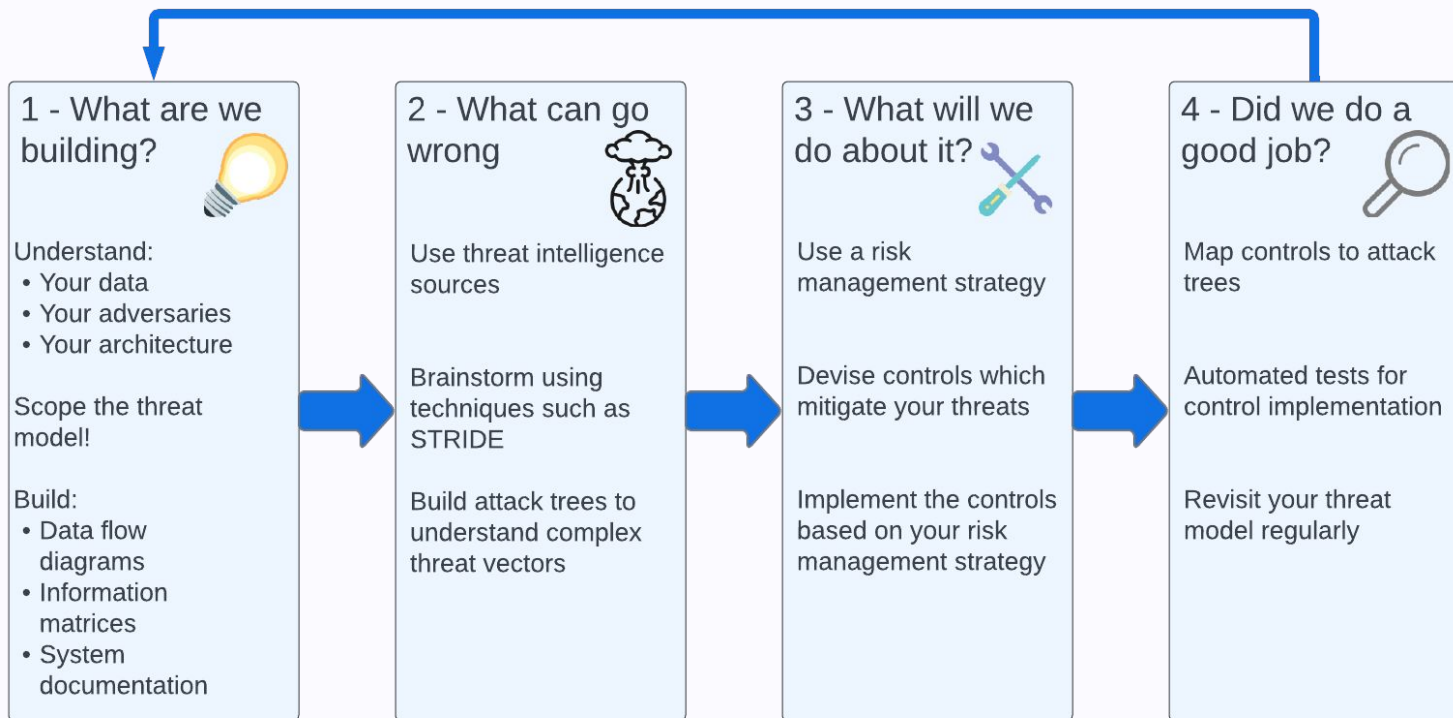
Why Threat Model?

- Identify **security flaws early**
- **Save money and time** consuming redesigns
- Focus your **security requirements**
- Identify **complex risks and data flows** for critical assets

Everyone can (and should!) Threat Model - **not just security teams**




Threat Modelling is Iterative



Making Sense of Zero Trust

Deriving Zero Trust Principles via Threat Modelling

Sky-High-Level Threat Model



Spoofing

Tampering

Repudiation

Information Disclosure

Denial of Service

Escalation of Privilege

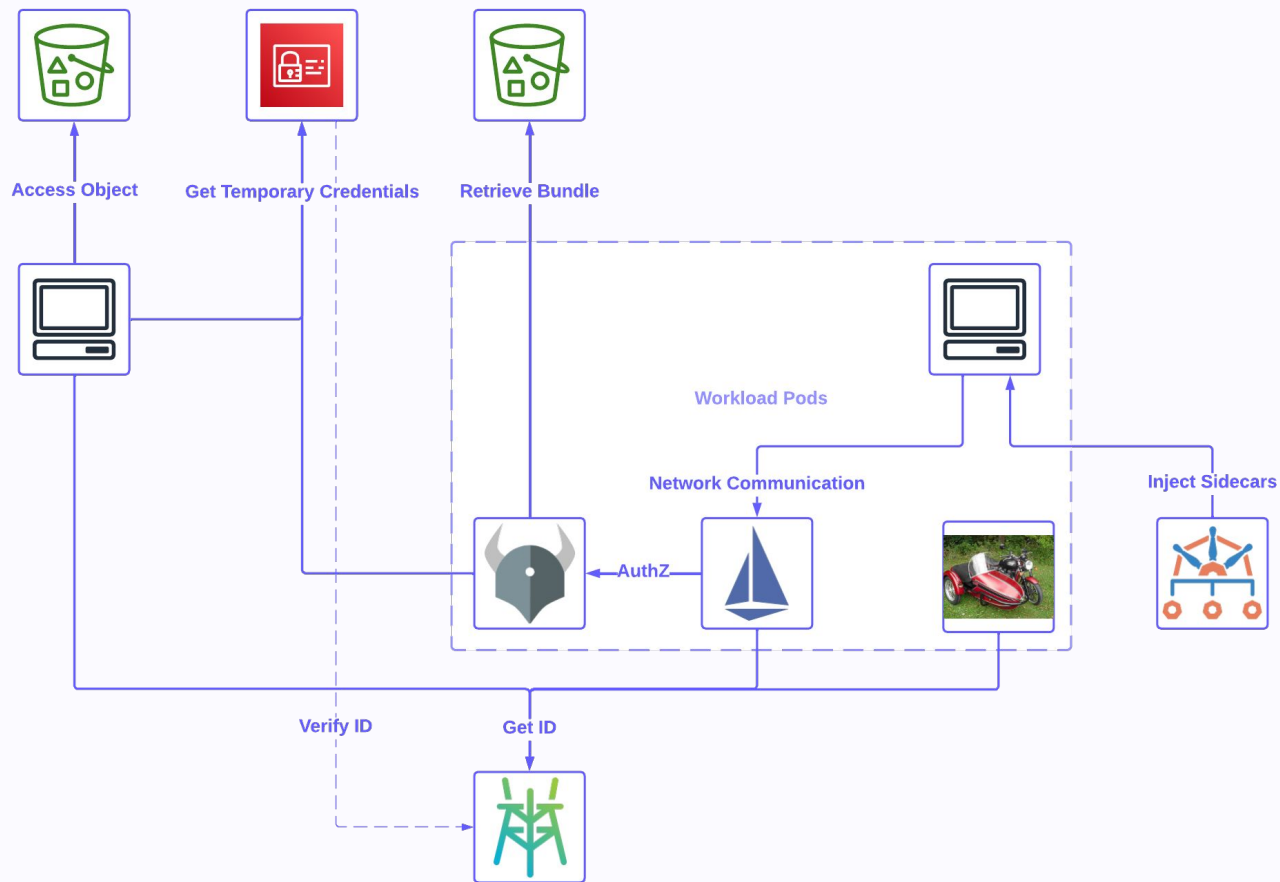
Deriving our architectural principles

STRIDE	Threat	Architectural Control
S	User Impersonation	User AuthN best practices IdP controls, e.g. MFA
S	Workload spoofing (client or server side)	Cryptographically strong workload identity mTLS
T	Alter information in transit	mTLS
T	Tamper with stored data	Strict AuthZ everywhere
R	Malicious action not attributable to an identity	Audit logs tied back to cryptographically strong identity
I	Exfiltrate data	Egress control Policy as versioned code
D	Prevent workloads from communicating	Highly available workload identity mechanism
E	Compromised workload pivots	Least Privilege AuthZ policies

What are we Building?

Step 1 - Draw and annotate architecture diagrams

Architecture Diagram



What are we Building?

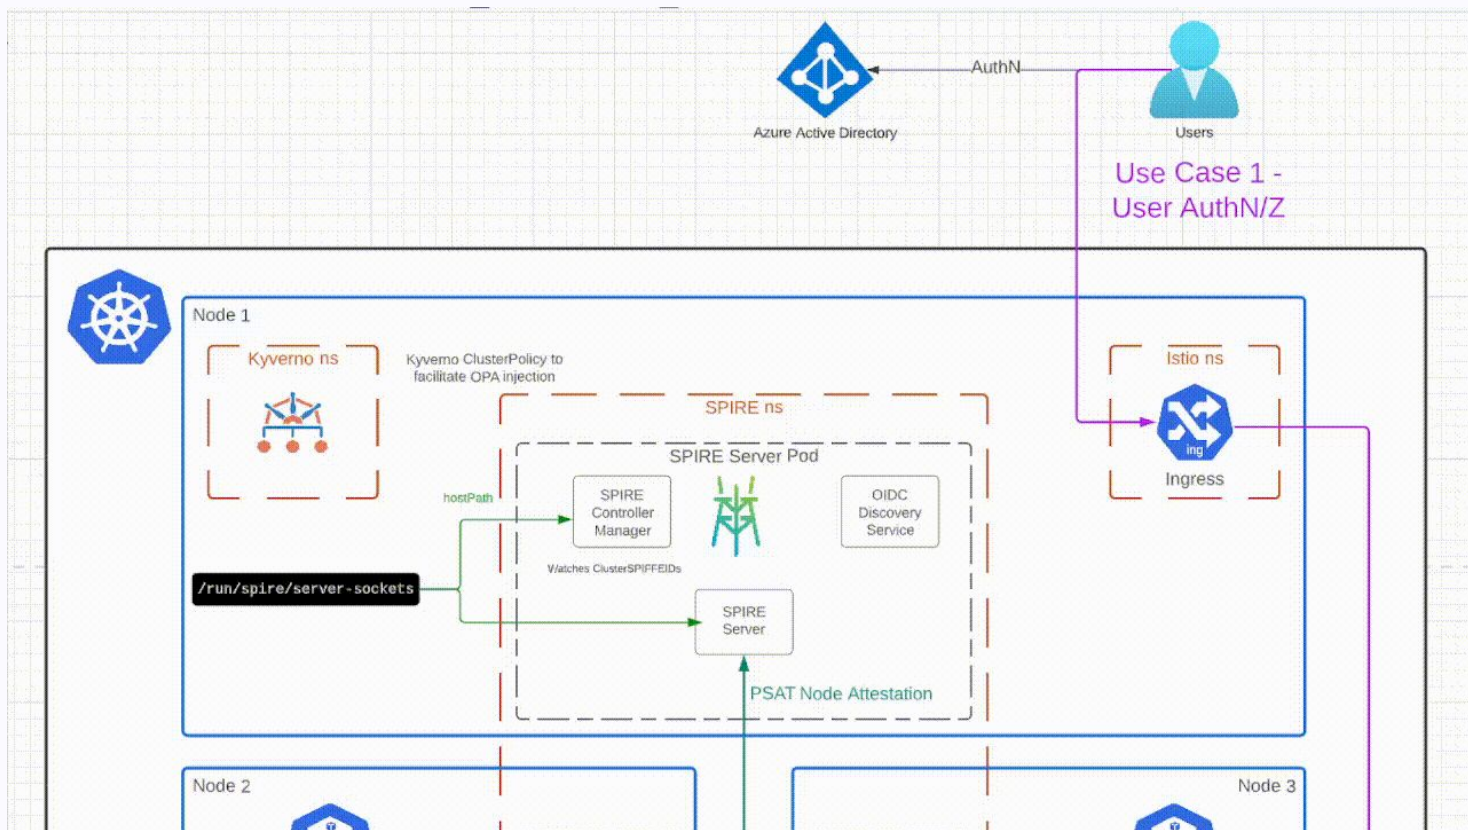
Step 2 - Build a Prototype

Prototypes Help us Threat Model!

- **Prototype goals:**
 - Understand **how technologies integrate**
 - Helps us to understand **what can go wrong**
 - **Quick and easy** to spin up:
 - We use a **local kind cluster** with a **minimal set of AWS resources** to show the OIDC provider setup
- In **Production** we would expose the JWKS of our SPIRE server by making an **OIDC discovery service** available
- To avoid this setup, we ship OIDC discovery info to an **S3 bucket**

Let's have a quick look
...and make sure it works

Draw Data Flow Diagrams!



What can go Wrong?

Threat Modelling the Detail

Draw Attack Trees

- Useful way to brainstorm and document **what can go wrong**
- Kelly Shortridge shows us how to create **attack trees as code**:
<https://www.deciduous.app/>
- From now, we use the following key:

OR Node

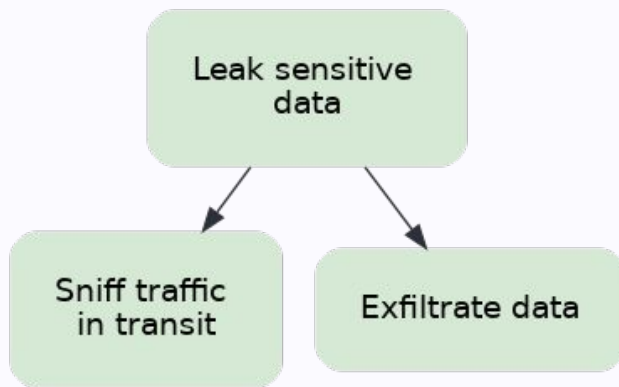
AND Node

Single Node

Out of Scope

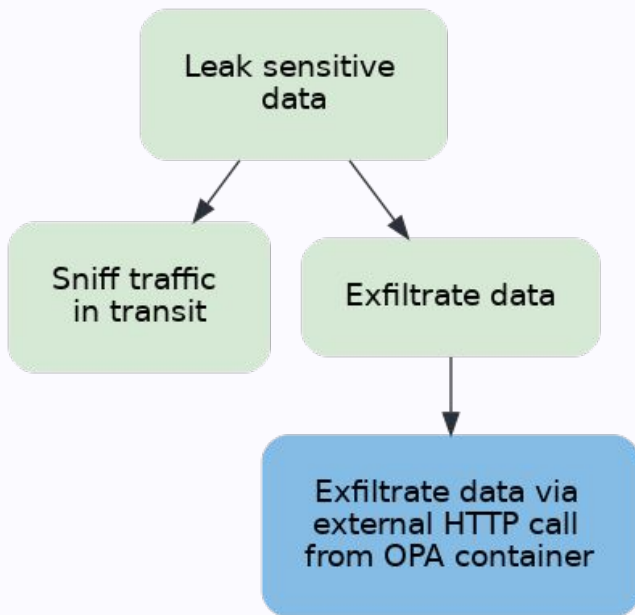
Building an Example Tree

- Start with a **bad outcome**
 - Could be a compromise of **Confidentiality, Integrity or Availability**
- **Confidentiality example:**

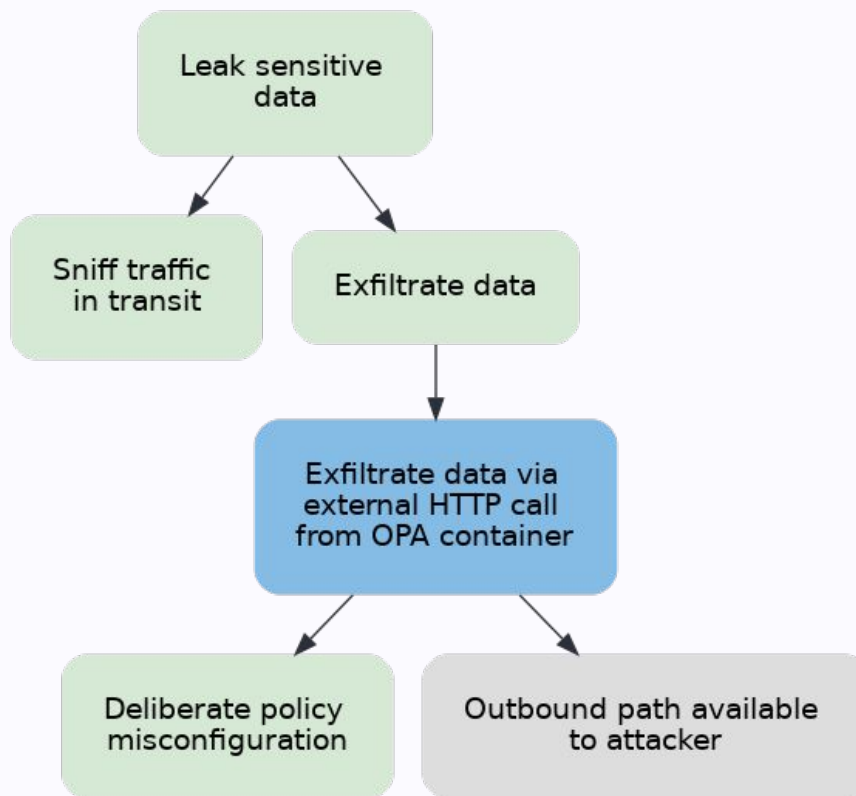


Adding a level

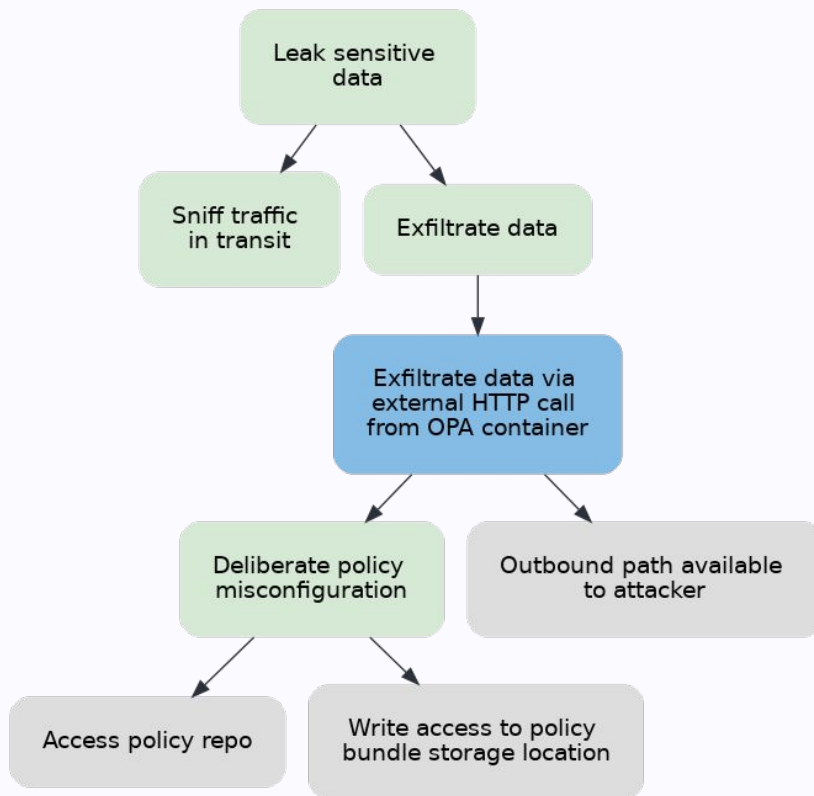
- For each node, brainstorm what an attacker would need to do to achieve that particular goal



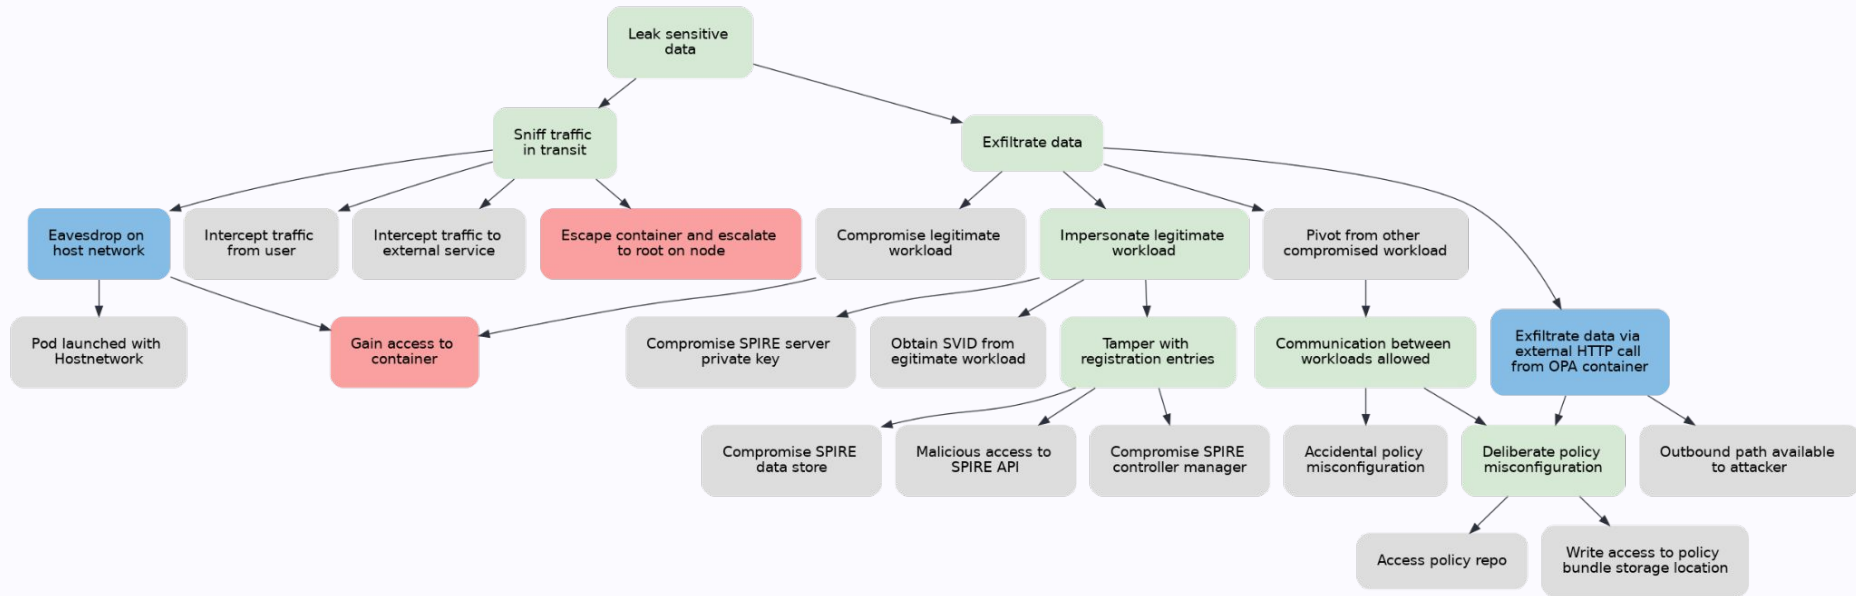
One More Time!



One More Time!



Putting it all Together



What will we do about it?

Designing Controls

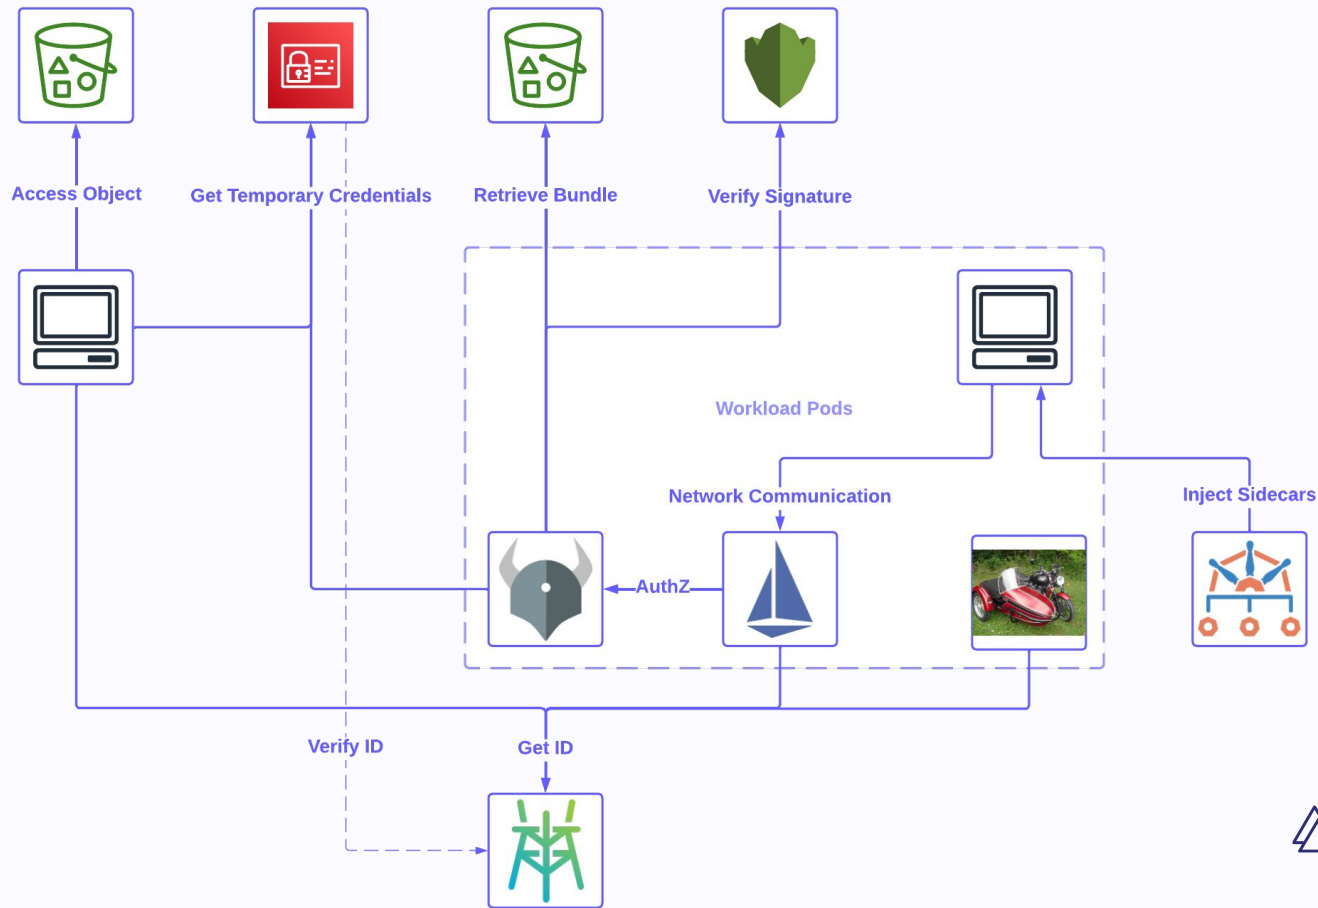
Devise Mitigating Controls

Confidentiality Threat	Control
T01 - Eavesdrop on Host Network	C01 - Validating Admission Control Policy C02 - Seccomp and Apparmor profiles
T02 - Intercept traffic in transit between users/workloads/services	C03 - TLS enforced for ingress and access to cloud services C04 - mTLS between workloads
T03 - Exfiltrate data	C05 - Egress control policies
T04 - Pivot from compromised workload	C06 - Network Policy C07 - Cryptographically strong workload IDs from SPIRE C08 - Istio External AuthZ with OPA Policy Engine
T05 - Accidental policy misconfiguration	C09 - Policy linting C10 - Policy unit tests
T06 - Access policy repo	C11 - Least Privilege repo access control C12 - Protected branches/enforced PRs with mandated review
T07 - Overwrite policy bundle	C13 - Cloud Provider RBAC least privilege & audit C14 - Policy bundle signing and verification
etc.	

Building Custom Controls

- Some controls will need **further architectural work** to be precisely defined
 - e.g. '**compromise SPIRE data store**' threat
 - requires data storage **design decision**
 - **Lower-level attack tree** can then be created
- We have looked at one **example attack path**:
 - **malicious internal actor** exploits **misconfigured IAM** policy to **overwrite policy bundle**
- We currently have a mitigating control for this:
 - C14 - **Policy bundle signing and verification**
- However, we still need to design the control implementation...

Iterating the Architecture



OPA Custom Bundle Signing

- What's in the **signature**?
 - <https://www.openpolicyagent.org/docs/latest/management-bundles/#signing>
- Where do you even **start**?
 - **OPA contrib repo**
 - https://github.com/open-policy-agent/contrib/tree/main/custom_bundle_signing

Kick the tyres!

...but not in production!

[https://github.com/controlplaneio/
threat-modelling-zero-trust-talk](https://github.com/controlplaneio/threat-modelling-zero-trust-talk)

Summary

- Even well designed systems should be threat modelled:
 - Threat landscapes change
 - Technologies evolve
- Adopting a Zero Trust mindset is becoming increasingly crucial
- Cryptographically strong workload identity is key
 - SPIRE
- Existing integrations can be used
 - e.g. Istio External AuthZ
- Custom controls can be designed if risk decisions require them
 - e.g. OPA custom bundle signing

Did we do a good job?
You tell us!
Thank you for listening!!

