

Kubernetes Operators: Safety First Through Model Checkers

Neven Miculinic
grid.ai

Motivation

With kubernetes building operators is easy...right?

Motivation

Our failure story with Redis operator for K8s (+ a brief look at Redis data analysis tools)



FLANT

Flant staff

Follow

Mar 13 · 18 min read



Motivation

[Release Notes for MongoDB Enterprise Kubernetes Operator](#) > Known Issues in the MongoDB Enterprise Kubernetes Operator

Known Issues in the MongoDB Enterprise Kubernetes Operator

Difficulties with Updates

In some cases, the Kubernetes Operator can **stop receiving change events**. As this problem is **hard to reproduce**, the recommended workaround is to delete the operator pod.

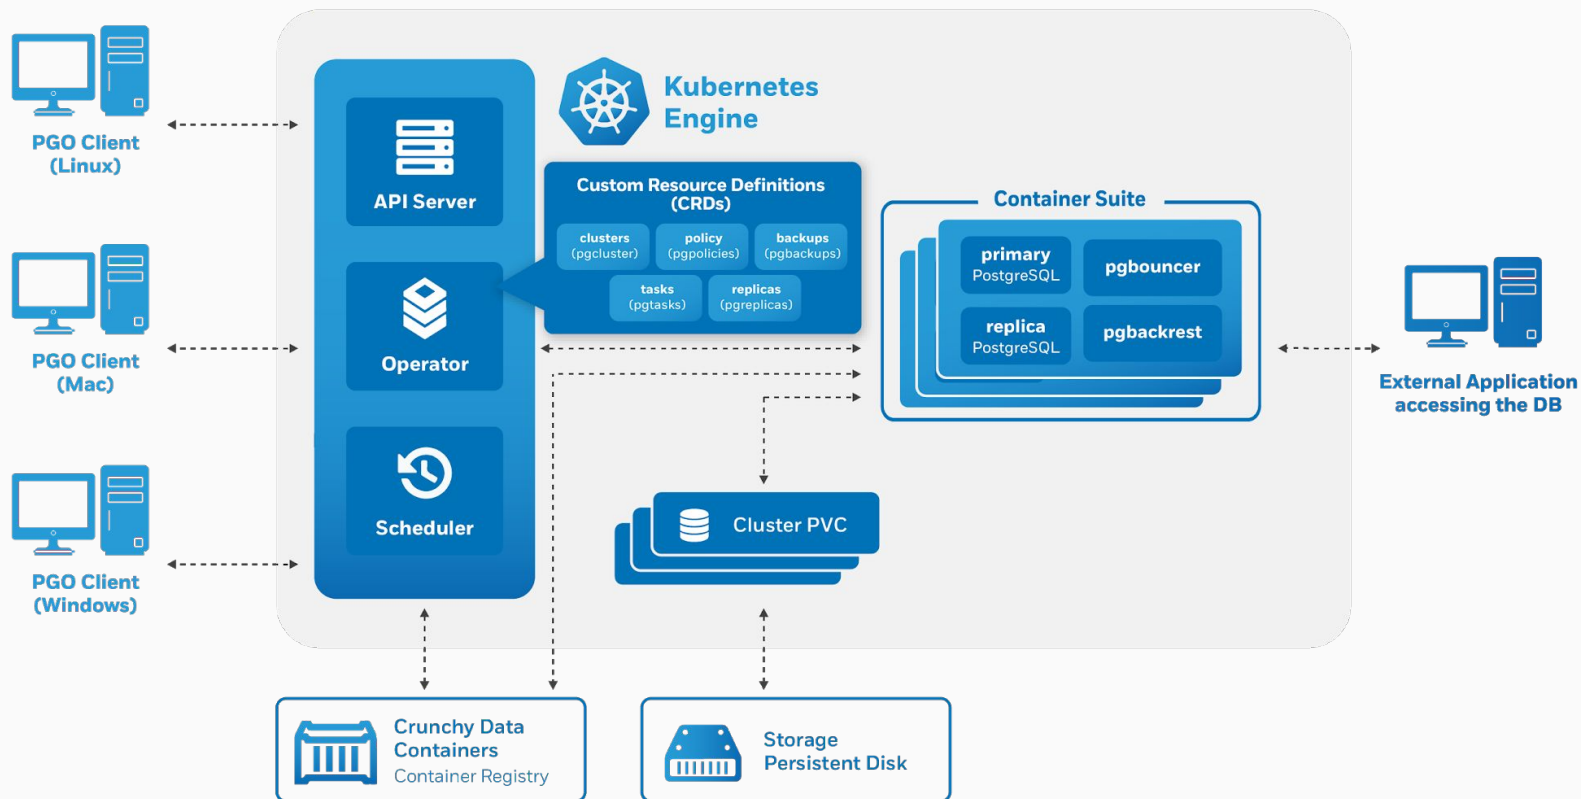
Kubernetes starts the new Kubernetes Operator automatically and starts working correctly:


The Kubernetes Operator will not be able to apply the **following change** on a MongoDB Deployment **simultaneously**:


- The TLS configuration is disabled (security.tls.enabled: false)
- The number of members in a Replica Set is increased

If both operations are applied at the same time, the MongoDB Resource could go into a **unrecoverable state**.

Tools

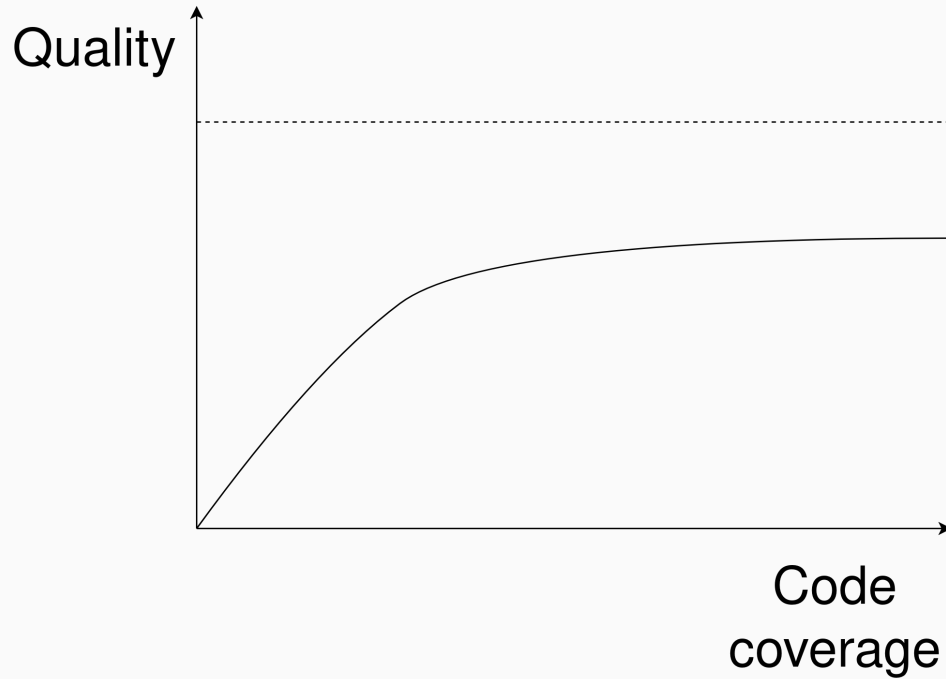






Better:

- Tools
- Process
- Externalities



- Unit tests
- Law of diminishing marginal utility

Conventional approach

- Deep design review
- Static code analysis
- Stress testing
- Fault-injection testing
- Unit tests
- Code reviews
- ...

A “New” tool appears

Model checkers:

- We specify the system model, its transitions functions, invariants, rules and properties we want checking
- The system searches the whole solution space for contradictions

Wrong tool for the job

- Real time constraints
- It's better expressing safety or liveness properties than real time constraints

Use of Formal Methods at Amazon Web Services

2014-09-29

Adoption at AWS

A crime thriller!

Starting
Chris Newcombe as C.N.
Tim Rath as T.R.

“To a first approximation, we can say that accidents are almost always the result of incorrect estimates of the likelihood of one or more things.” -C. Michael Holloway, NASA

- C.N. was dissatisfied with the quality
- C.N. dismissed formal methods due to myths

Misconceptions

- Takes long time to learn
- Only a small fraction of real-life problems fit this paradigm
- A low return on investment
- They are impractical

- A mystic paper appears
- '10-year test of time' award at SIGCOMM 2011

Using Lightweight Modeling To Understand Chord

Pamela Zave
AT&T Laboratories—Research
Florham Park, New Jersey USA
pamela@research.att.com

ABSTRACT

Correctness of the Chord ring-maintenance protocol would mean that the protocol can eventually repair all disruptions in the ring structure, given ample time and no further disruptions while it is working. In other words, it is “eventual

ambitious way. The breezy, informal reasoning seems prone to the all-too-human misconception that complex systems will work exactly the way we expect them to.

Distributed systems frequently do not work the way we expect them to—but the gap between human intuition and reality can now be bridged by powerful and convenient tools.

- C.N. investigates the first tool, a unsung hero called Alloy
- Limited expressivity



- C.N. looked further for the holy grail

Fast Paxos

Leslie Lamport

14 July 2005

Revised 18 January 2006

Minor revision 14 April 2006

MSR-TR-2005-112

- And found the TLA+ spec at the end...

```
----- MODULE FastPaxos -----  
module imports two standard modules. Module Naturals defines the set Naturals and the ordinary arithmetic operators; module FiniteSets defines IsFiniteSet to be true iff S is a finite set and defines Cardinality(S) to be the number of elements of S if S is finite.  
EXTENDS Naturals, FiniteSets
```

Constants

$\max(S)$ is defined to be the maximum of a nonempty finite set *S* of numbers.

$\text{max}(S) \triangleq \text{CHOOSE } i \in S : \forall j \in S : j \leq i$

The next statement declares the specification's constant parameters, which have the following meanings:



Industry usage -- AWS



Industry usage -- AWS





DynamoDB

- T.R. enters the story with DynamoDB
- Performed classical testing approach
- A couple of weeks with TLA+
- Found a bug in his design requiring 35 high level steps + 2 more bugs

Industry usage -- AWS




“Debugging designs”

“Exhaustively testable
pseudo-code”


Applying TLA+ to some of our more complex systems

System	Components	Line count (excl. comments)	Benefit
S3	Fault-tolerant low-level network algorithm	804 PlusCal	Found 2 bugs. Found further bugs in proposed optimizations.
	Background redistribution of data	645 PlusCal	Found 1 bug, and found a bug in the first proposed fix.
DynamoDB	Replication & group-membership system	939 TLA+	Found 3 bugs, some requiring traces of 35 steps
EBS	Volume management	102 PlusCal	Found 3 bugs.
Internal distributed lock manager	Lock-free data structure	223 PlusCal	Improved confidence. Failed to find a liveness bug as we did not check liveness.
	Fault tolerant replication and reconfiguration algorithm	318 TLA+	Found 1 bug. Verified an aggressive optimization.

- Leslie Lamport & Microsoft Research



2002



Engineers use TLA+ to prevent serious but subtle bugs from reaching production.

BY CHRIS NEWCOMBE, TIM RATH, FAN ZHANG, BOGDAN MUNTEANU, MARC BROOKER, AND MICHAEL DEARDEUFF

How Amazon Web Services Uses Formal Methods

- April 2015
- December 26th email

- Service Fabric
 - Azure Batch
 - Azure Storage
 - Azure Networking
 - Azure IoT Hub
- uncovered a safety violation even in our most confident implementation

- Intel
- Microsoft
 - XBOX
 - Azure
- AWS
- OpenComRTOS → 10x code reduction with TLA+ help
- [ElasticSearch 7.0+](#)
- [MongoDB](#)
- ...

Benefits

- Improved design quality
- Less bugs
- Performance optimizations
- improved time-to-market
- Documentation

Benefits

Informal
model

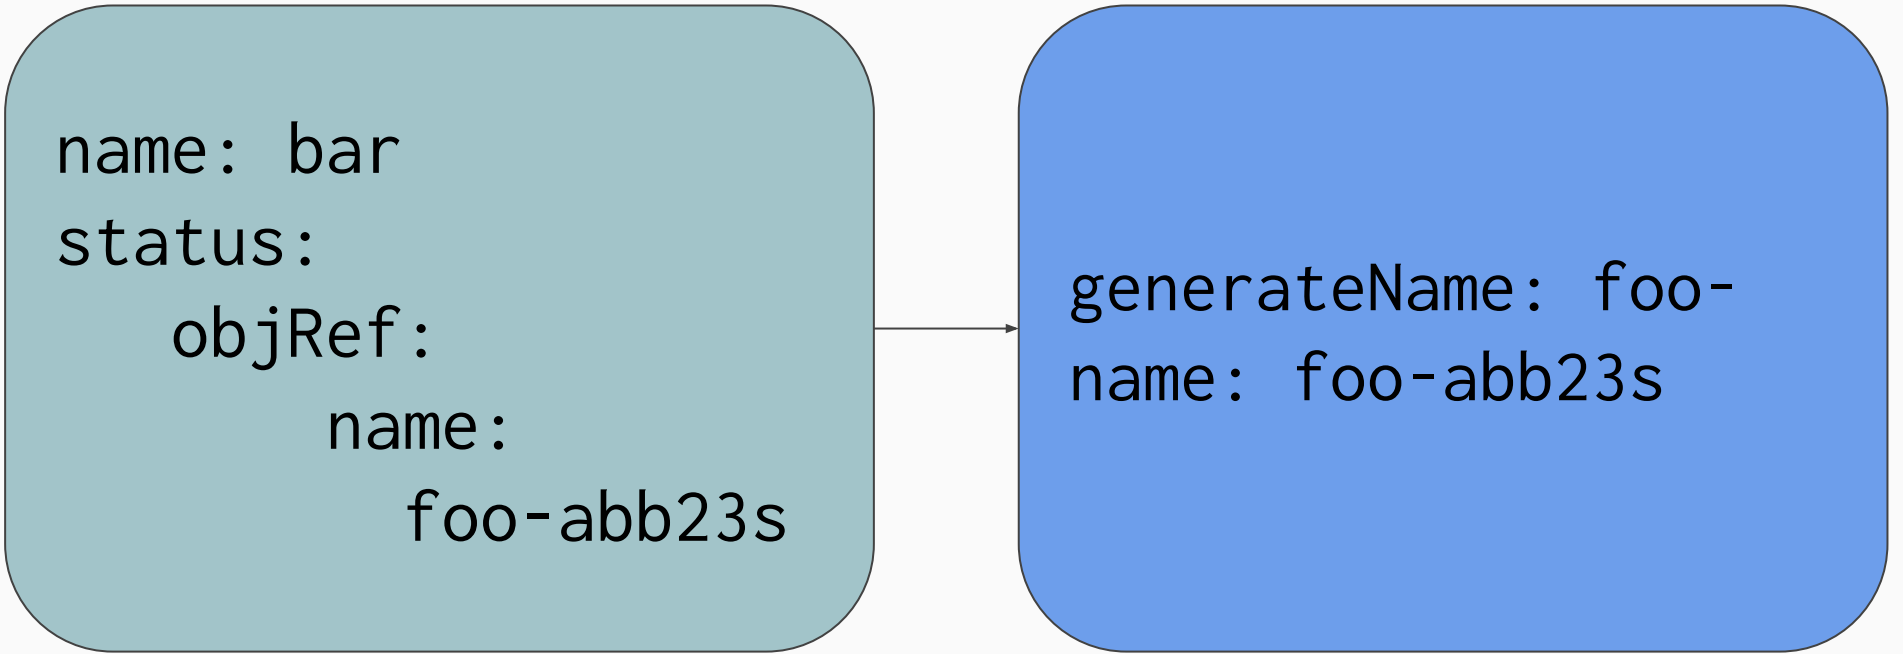
Model
spec

Code

complexity & precision



Example TLA+



```
graph LR; A["name: bar  
status:  
  objRef:  
    name:  
      foo-abb23s"] --> B["generateName: foo-  
name: foo-abb23s"]
```

name: bar
status:
 objRef:
 name:
 foo-abb23s

generateName: foo-
name: foo-abb23s

```
----- MODULE ns -----  
EXTENDS TLC, Integers, Sequences
```

```
(*--algorithm ns
variables
  fooObj = <<>>,
  barObj = [name |-> "foo", created
|-> FALSE],
```


- Safety properties
- Liveness properties

```
define
  LimitedFoo0bjs == Len(foo0bjs) <= 1
end define;
```

```
begin
  Start: while ~barObj.created do
    ...
  end while;
  assert(Len(fooObjs) = 1);
end algorithm;*)
```

Example TLA+ (PlusCAL)

```
either
  CreateObject: fooObjs := Append(fooObjs,
    [generateName |-> barObj.name \o "-"]);
  either
    MarkCreated: barObj.created := TRUE;
  or
    RebootDuringMarkingCreation: skip;
  end either;
or
  RebootDuringCreation: skip;
end either;
```

Example TLA+ (PlusCAL)

```
either
  CreateObject: fooObjs := Append(fooObjs,
    [generateName |-> barObj.name \o "-"]);
  either
    MarkCreated: barObj.created := TRUE;
  or
    RebootDuringMarkingCreation: skip;
  end either;
or
  RebootDuringCreation: skip;
end either;
```

Example TLA+ (PlusCAL)

Error Trace

▼ 1: Initial predicate

▶ barObj (2)	[created -> FALSE, name -> "foo"]
fooObjs (0)	<<>>
pc	"Start"

▼ 2: Start in ns >>

▶ barObj (2)	[created -> FALSE, name -> "foo"]
fooObjs (0)	<<>>
pc M	"CreateObject"

Example TLA+ (PlusCAL)

▼ 3: CreateObject in ns >>

▶ barObj (2)

▶ fooObjs (1) M

pc M

[created |-> FALSE, name |-> "foo"]

<<[generateName |-> "foo-"]>>

"RebootDuringMarkingCreation"

▼ 4: RebootDuringMarkingCreation in ns >>

▶ barObj (2)

▶ fooObjs (1)

pc M

[created |-> FALSE, name |-> "foo"]

<<[generateName |-> "foo-"]>>

"Start"

Example TLA+ (PlusCAL)

▼ 5: Start in ns >>

► barObj (2)

[created |-> FALSE, name |-> "foo"]

► fooObjs (1)

<<[generateName |-> "foo-"]>>

pc M

"CreateObject"

▼ 6: CreateObject in ns >>

► barObj (2)

[created |-> FALSE, name |-> "foo"]

► fooObjs (2) M

<<[generateName |-> "foo-"], [generateName |-> "foo-"]>>

pc M

"MarkCreated"

Thanks!

Q&A

Neven Miculinic

neven.miculinic@gmail.com

References

[Use of Formal Methods at Amazon Web Services](#) Chris Newcombe, Tim Rath, Fan Zhang, Bogdan Munteanu, Marc Brooker, Michael Deardeu

[Why Amazon Chose TLA +](#)

[Industrial use of TLA+ \(Lamport\)](#)

[Our failure story with Redis operator for K8s \(+ a brief look at Redis data analysis tools\)](#)

[Known Issues in the MongoDB Enterprise Kubernetes Operator — MongoDB Kubernetes Operator 1.4](#)

[How Amazon Web Services Uses Formal Methods | April 2015 | Communications of the ACM](#)

[TLA+ Conf: Program](#)

[\(PDF\) A Brief History of Formal Methods](#)