

# Navigating the app delivery landscape while solving everyday problems

*Lei Zhang, Alois Reitbauer*



# Session abstract



KubeCon



CloudNativeCon

North America 2020

Virtual

Recently there was a lot of buzz around the CNCF landscape getting overly complex. There are more and more options to choose from which is good to see a growing ecosystem. However, this can sometimes also become overwhelming. In this session we walk through major challenges that people are facing and show how they can be solved with projects available today. We will look into scenarios from defining your application all the way to running it in production. Whether you are just about to start running applications on Kubernetes or want to improve your skills. This session will help you to identify what to consider when building app platforms, share good practices, compare different approaches and give you an interactive and quick tour through the CNCF landscape with a focus on delivering and running applications. You will also get some demo code samples to try everything hands on and have a reference to get started with your own projects.

---

## Speakers



### Lei Zhang

Staff Engineer - OAM/KubeVela, Alibaba

Lei is a co-maintainer of Kubernetes community, and co-chair of CNCF App Delivery SIG. Lei is co-leading engineering effort in Alibaba including Kubernetes and large-scale cluster management system. Before it, Lei worked for Hyper\_ and Microsoft Research (MSR). Lei is a popular speaker... [Read More →](#)



### Alois Reitbauer

Chief Technical Strategist, Dynatrace

Alois is an executive member of the technical staff at Dynatrace. He has been building monitoring and application management solutions for more than 15 years. Alois has successfully brought multiple products to market that are used by the biggest companies on the planet. His current... [Read More →](#)

# about SIG App-Delivery



KubeCon



CloudNativeCon

North America 2020

*Virtual*

## CNCF App Delivery SIG

The Application Delivery SIG focuses on delivering cloud native applications which involves multiple phases including building, deploying, managing, and operating. Additionally, the SIG produces supporting material and best practices for end-users, and provide guidance and coordination for CNCF projects working within the SIG's scope.

See our full charter here: <https://github.com/cncf/toc/blob/master/sigs/app-delivery.md>

### Chairs

Alois Reitbauer, Bryan Liles and Lei Zhang.

### Community

- Slack channel: #sig-app-delivery in CNCF workspace - <https://cloud-native.slack.com/messages/CL3SL0CP5>
- Mailing list: <https://lists.cncf.io/g/cncf-sig-app-delivery/topics>

### Meetings

We have meetings every 1st and 3rd Wednesday of the month at 8:00 a.m. PST.

- Agenda and Notes:  
<https://docs.google.com/document/d/1OykvqvhSG4AxEdmDMXilrupsX2n1qCSJUWwTc3I7AOs/edit#>
- Zoom Meeting: <https://zoom.us/j/7276783015>
- Recordings of previous meetings: <https://www.youtube.com/playlist?list=PLj6h78yzYM2OHd1Ht3jiZuucWzvouAAci>

### Mission

Consistent with the [CNCF SIG definition](#), the mission of CNCF SIG App Delivery is:

- To collaborate on areas related to developing, distributing, deploying, managing and operating secure cloud-native applications with the target of delivering application in manner of cloud native.
- To develop informational resources including guides, tutorials and white papers to give the community an understanding of best practices, trade-offs, and value-adds regarding to application delivery.
- To identify suitable projects and gaps in the landscape, periodically updating the TOC with suggested actions in a structured manner. This includes helping the TOC with assessments and due diligence of prospective new projects.

### Areas considered in Scope

SIG Application Delivery focuses on the following topics of the lifecycle of cloud-native applications:

- Application definition, including description, parameter and configuration
- Guidance and practice for application design and development
- Application bundling and deployment
- Package management
- Application delivery workflow and strategy
- Configuration source driven workflow
- Release management

The SIG will work on developing best practices, fostering collaboration between related projects, working on improving tool interoperability as well as proposing new initiatives and projects when blank spots in the current landscape are identified.

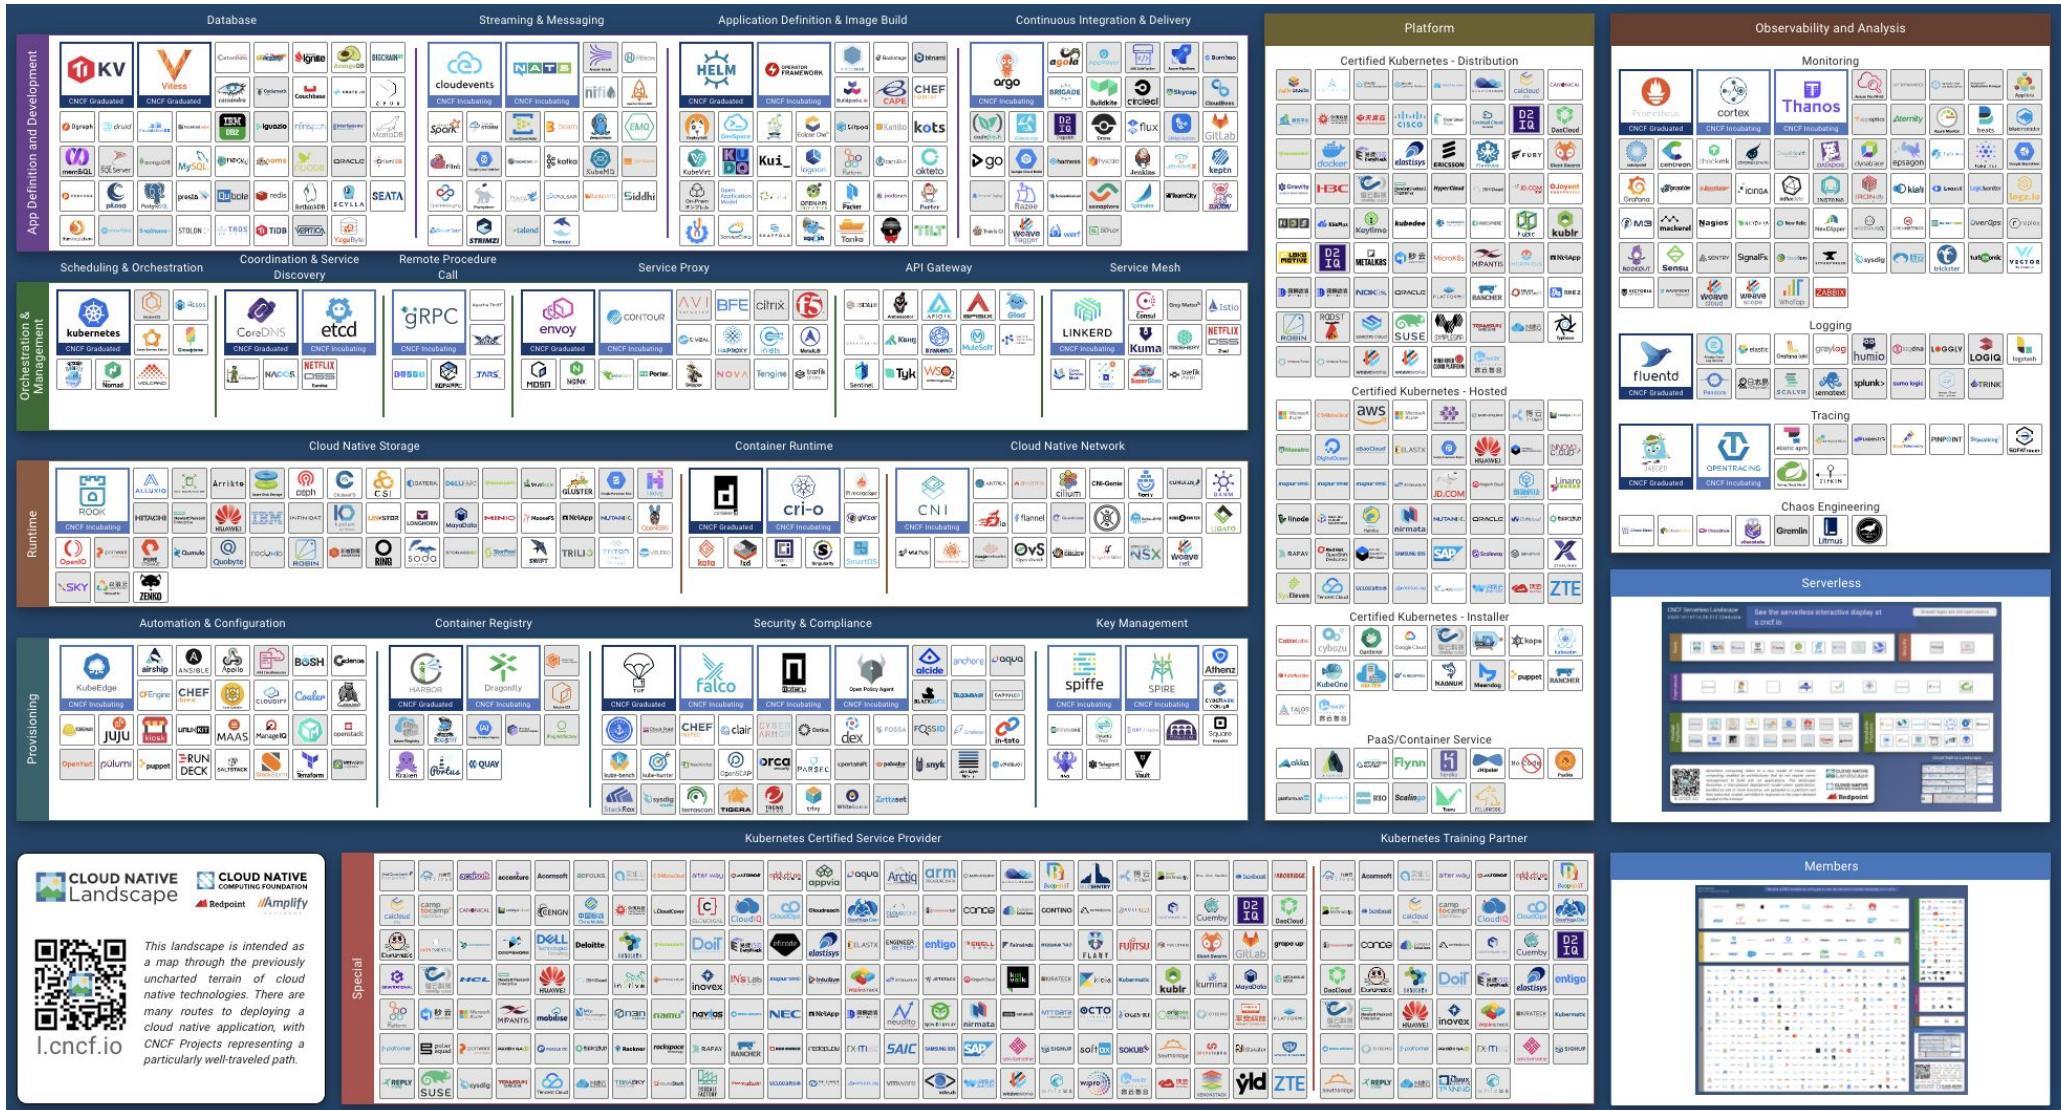
# The CNCF landscape ...



CloudNativeCon

North America 2020

*Virtual*



# There is even a puzzle ...



KubeCon



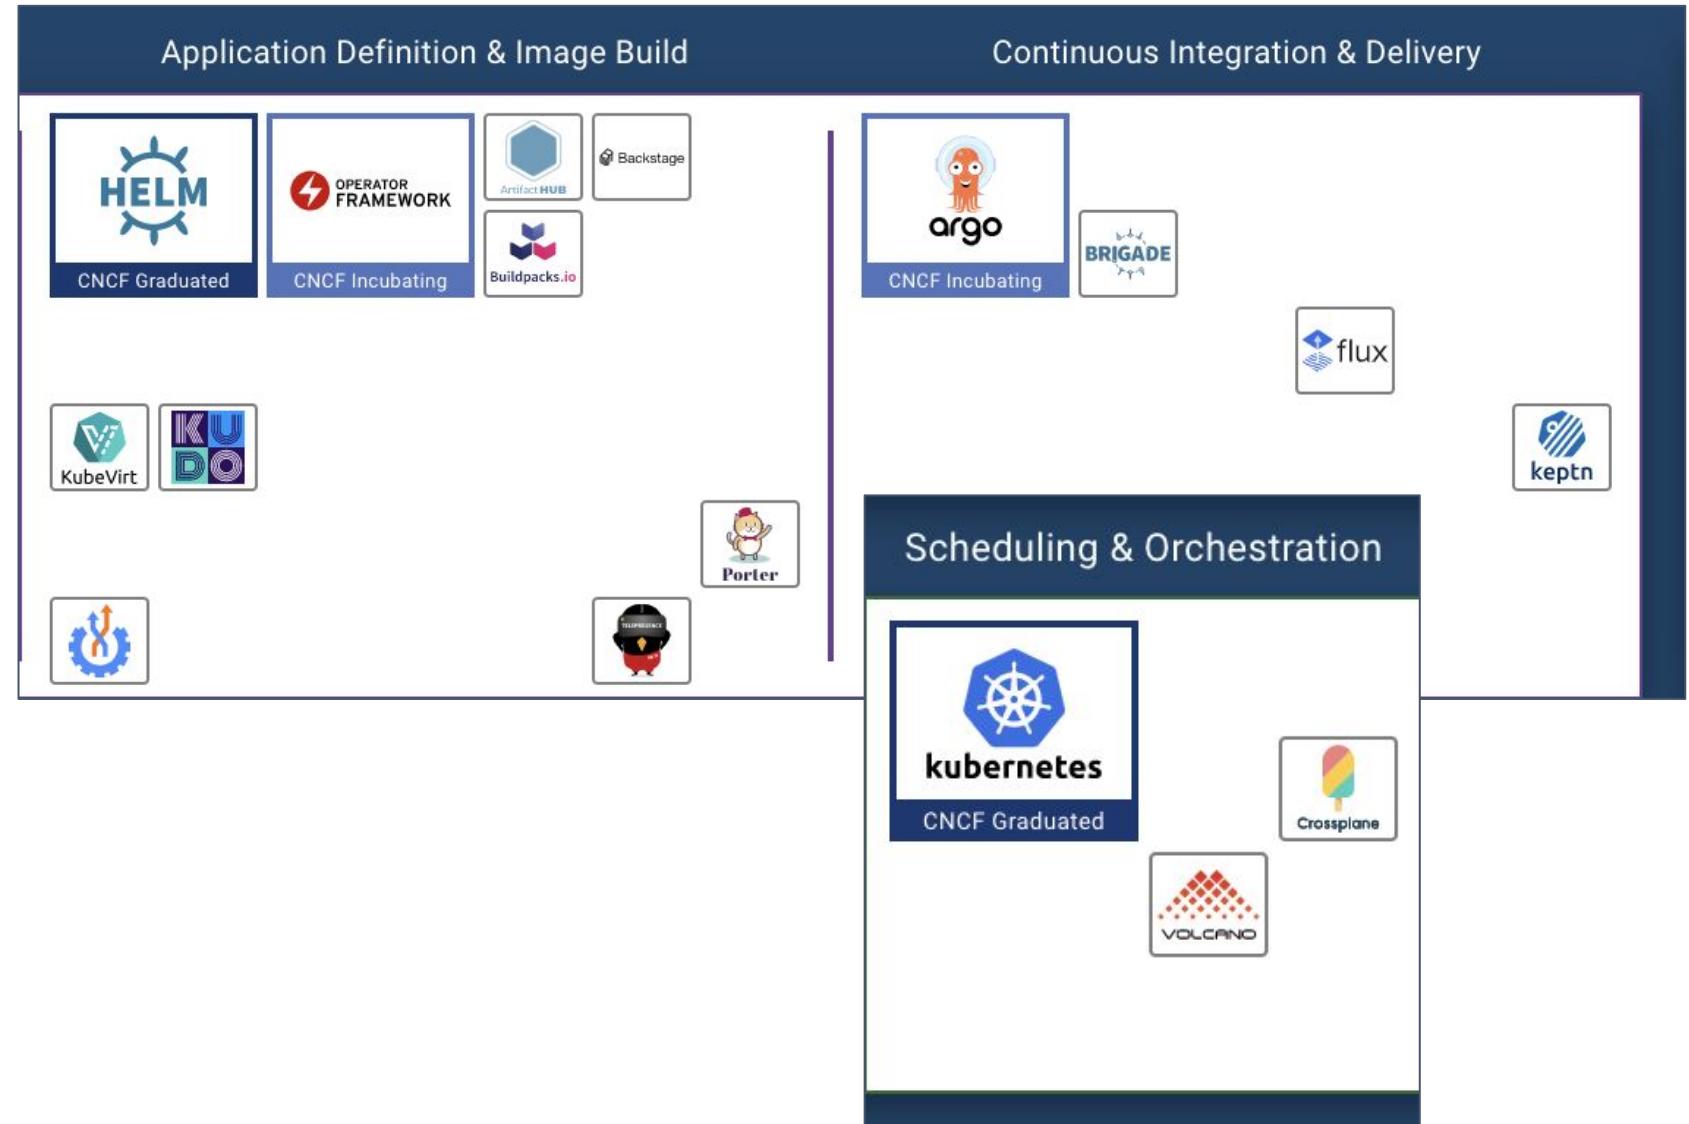
CloudNativeCon

North America 2020

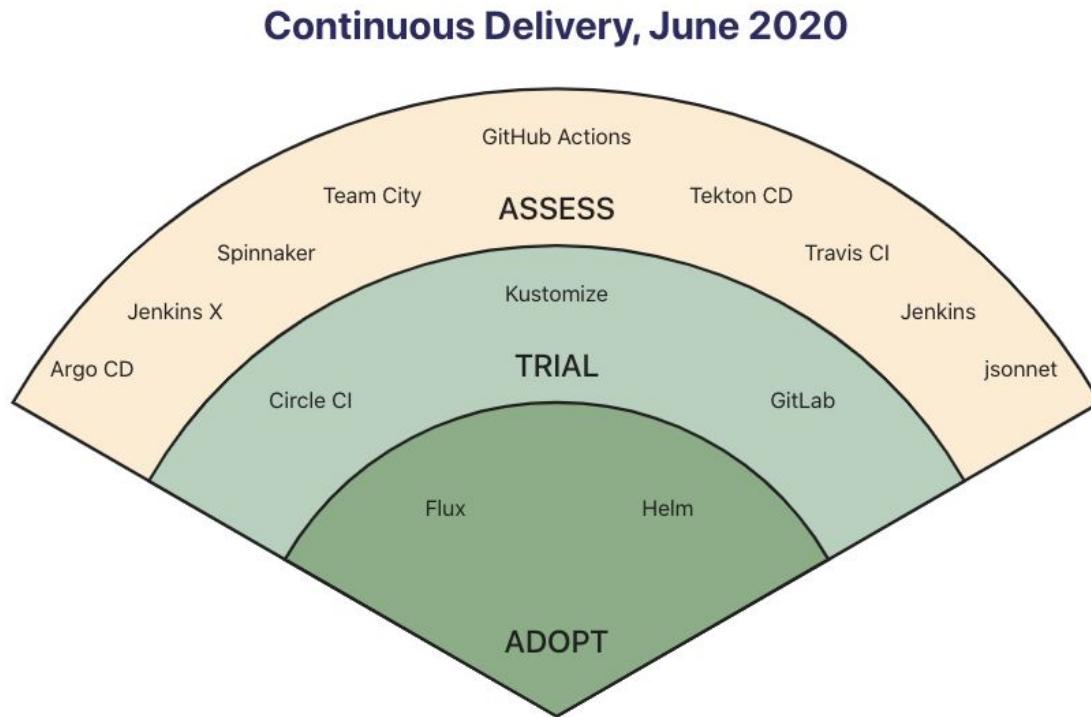
*Virtual*



# ... let us look at CNCF app-delivery only



# The CNCF tech radar



## The Themes

### 1. Publicly available solutions are combined with in-house tools

Many end users had tried up to 10 options and settled on adopting 2-4. Several large enterprise companies have built their own continuous delivery tools and open sourced components, including LunarWay's release-manager, Box's kube-applier, and stackset-controller from Zalando. The public cloud managed solutions on the CNCF landscape were not suggested by any of the end users, which may reflect the options available a few years ago.

### 2. Helm is more than packaging applications

While Helm has not positioned itself as a Continuous Delivery tool (it's the Kubernetes package manager first), it's widely used and adopted as a component in different CD scenarios.

### 3. Jenkins is still broadly deployed, while cloud native-first options emerge.

Jenkins and its ecosystem tools (Jenkins X, Jenkins Blue Ocean) are widely evaluated and used. However, several end users stated Jenkins is primarily used for existing deployments, while new applications have migrated to other solutions. Hence end users who are choosing a new CD solution should assess Jenkins alongside tools that support modern concepts such as GitOps (for example, Flux).

# The App Delivery Reference Model



CloudNativeCon

North America 2020

Virtual

## Topic 1: Application Definition

- app descriptor
- app model

## Topic 1.5: Application Packaging

- app packaging
- app parameter & configuration

## Topic 2: Application Operation & Rollout

- app lifecycle mgmt & config src driven workflow
- app rollout strategies: blue-green, canary etc
- app traffic mgmt

## Topic 3: Workload Instance Automation & Operation

- workload instance healing, scale in/out, sharding
- workload instance lifecycle mgmt

## Topic 4: Platform

- resource mgmt & scheduling
- container lifecycle mgmt, healing and runtime
- networking, logging, monitoring, mesh

CNCF SIG App Delivery

- Kubernetes Operator
- Kubernetes SIG Apps

- Kubernetes
- FaaS
- Cloud Services

# Introducing podTato Head



The screenshot shows the GitHub README.md page for the `appdelivery-demo` repository. The page includes a commit history from AloisReitbauer, a note about 1 contributor, and a file size of 36 lines (26 sloc) / 1.41 KB. A large image of Mr. Potato Head is displayed, and a note at the bottom states: "This project is in its very early stages, so please - be kind." Another note below says: "What you are getting This project consists of the smallest possible application to demo cloud native application delivery. It - for sure - will grow over time. Right."

**podtato Head is:**

- the simplest cloud-native project
- a way to explore different delivery tools
- a funny name for a GitHub project

<https://github.com/cncf/podtato-head>



KubeCon



CloudNativeCon

North America 2020

*Virtual*

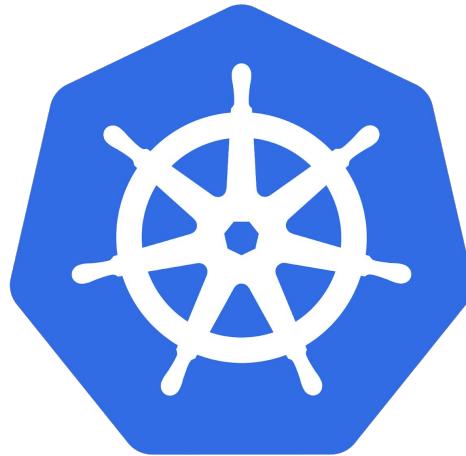
Let us start super, super simple

# Using Manifest yaml files



*Virtual*

The most basic way to deploy an application is by using Kubernetes manifest files



```
apiVersion: v1
kind: Service
metadata:
  name: helloservice
  namespace: demospace
spec:
  selector:
    app: helloservice
  ports:
  - name: http
    port: 9000
    protocol: TCP
    targetPort: 9000
  type: LoadBalancer
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: helloservice
  namespace: demospace
spec:
  selector:
    matchLabels:
      app: helloservice
  template:
    metadata:
      labels:
        app: helloservice
    spec:
      terminationGracePeriodSeconds: 5
      containers:
      - name: server
        image: aloisreitbauer/hello-server:v0.1.1
        imagePullPolicy: Always
        ports:
        - containerPort: 9000
        env:
        - name: PORT
          value: "9000"
```



KubeCon



CloudNativeCon

North America 2020

*Virtual*

# What if we want more flexibility

# Using Helm templates



KubeCon



CloudNativeCon

North America 2020

Virtual



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: helloservice
  namespace: {{ .Release.Namespace }}
spec:
  selector:
    matchLabels:
      app: helloservice
  template:
    metadata:
      labels:
        app: helloservice
    spec:
      terminationGracePeriodSeconds: 5
      containers:
        - name: server
          image: {{ .Values.imageName}}
          ports:
            - containerPort: 9000
          env:
            - name: PORT
              value: "9000"
imageName: aloisreitbauer/hello-server:v0.1.2
```

```
apiVersion: v1
kind: Service
metadata:
  name: helloservice
  namespace: {{ .Release.Namespace }}
spec:
  type: ClusterIP
  selector:
    app: helloservice
  ports:
    - name: http
      port: 9000
      protocol: TCP
      targetPort: 9000
  type: LoadBalancer
```



KubeCon



CloudNativeCon

North America 2020

*Virtual*

We want changes to be applied automatically

# Moving towards GitOps with Flux



*Virtual*





KubeCon



CloudNativeCon

North America 2020

*Virtual*

We want to see what happens

# GitOps with ArgoCD



KubeCon



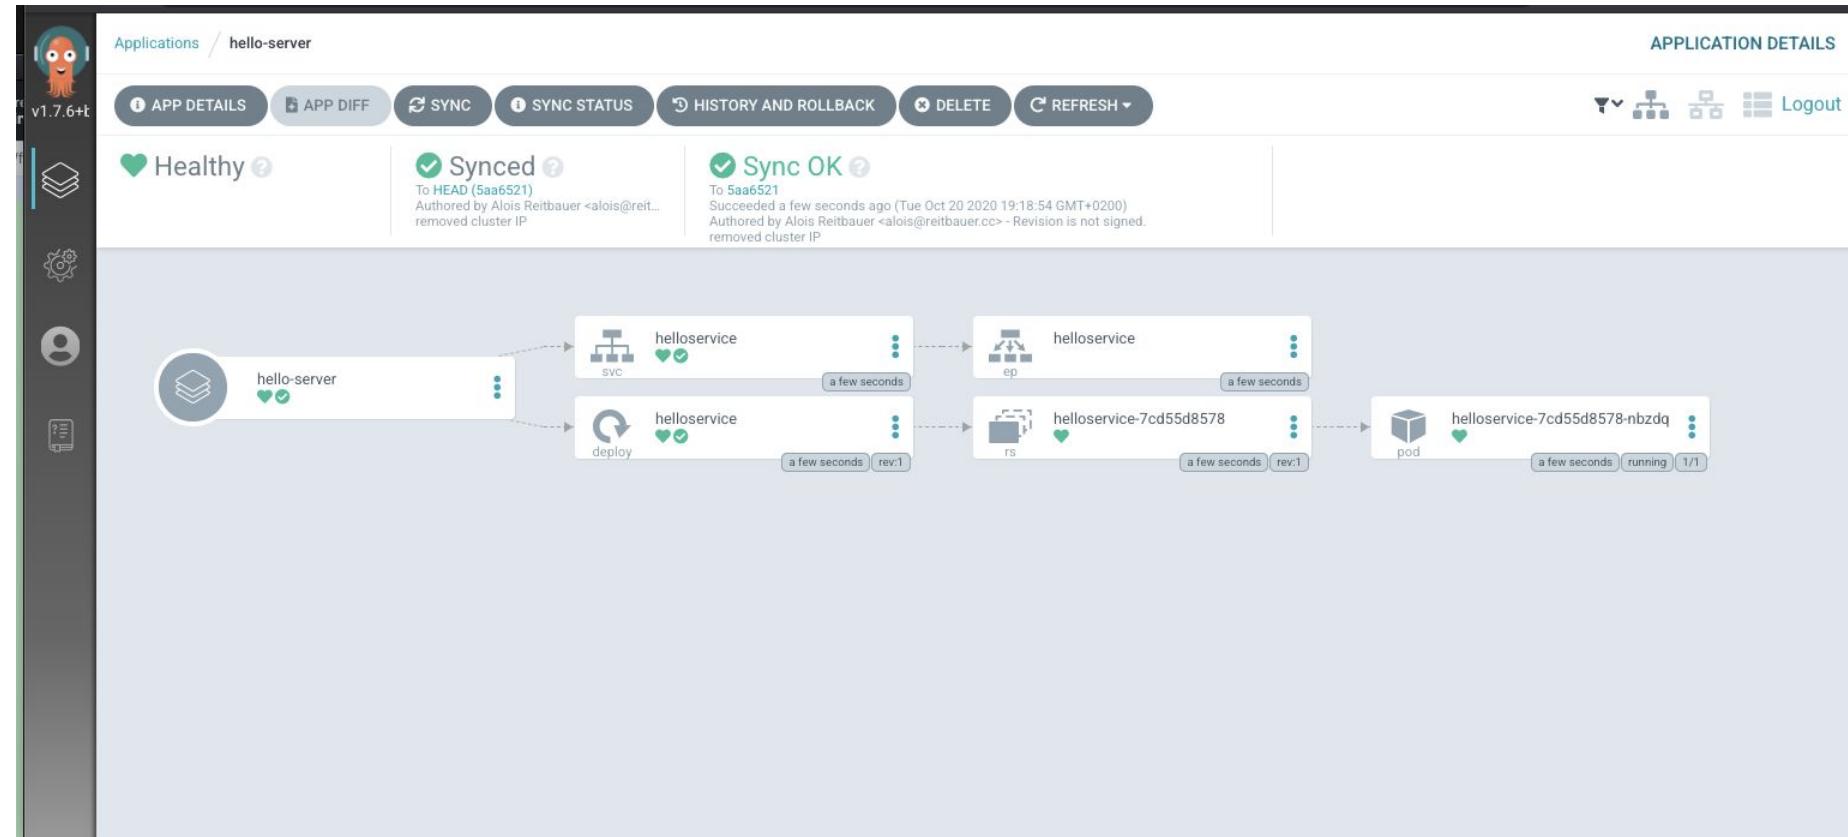
CloudNativeCon

North America 2020

Virtual



# argo





KubeCon



CloudNativeCon

North America 2020

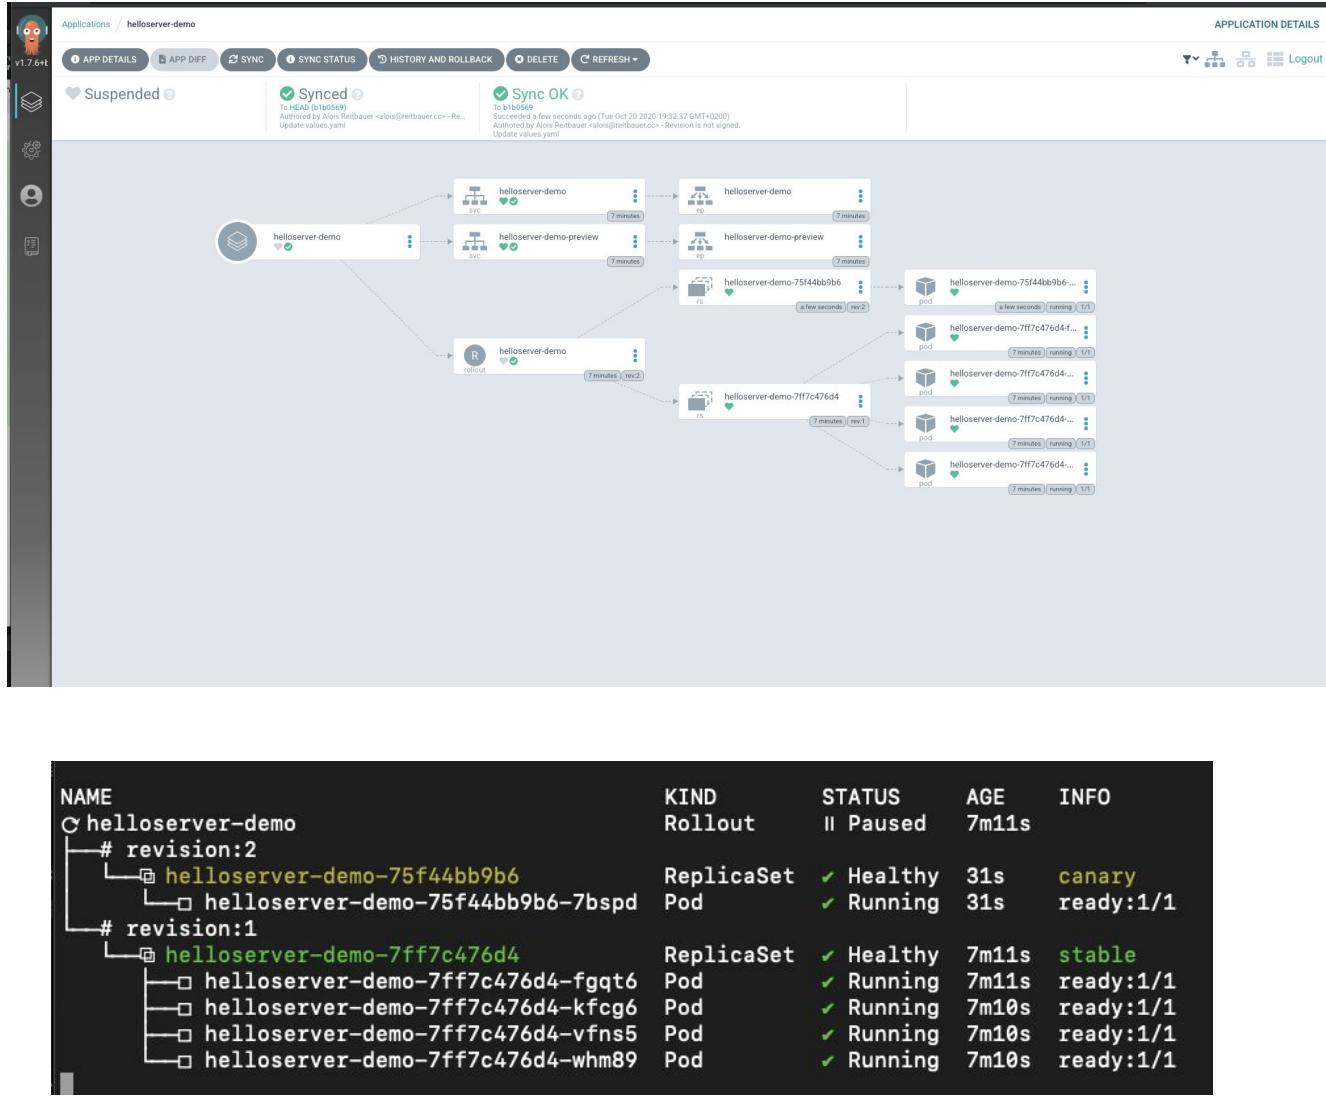
*Virtual*

We want to get more control over the rollout

# Using Argo Rollouts



```
apiVersion: argoproj.io/v1alpha1
kind: Rollout
metadata:
  name: {{ template "helloservice-demo.fullname" . }}
  labels:
    app: {{ template "helloservice-demo.name" . }}
    chart: {{ template "helloservice-demo.chart" . }}ku
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  replicas: {{ .Values.replicaCount }}
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: {{ template "helloservice-demo.name" . }}
      release: {{ .Release.Name }}
  strategy:
    canary:
      #blueGreen:
        # activeService: {{ template "helloservice-demo.fullname" . }}
        # previewService: {{ template "helloservice-demo.fullname" . }}-preview
        steps:
          - setWeight: 20
          - pause: {duration: 10}
          - setWeight: 40
          - pause: {duration: 10}
          - setWeight: 60
          - pause: {duration: 10}
          - setWeight: 80
          - pause: {duration: 10}
  template:
    metadata:
      labels:
        app: {{ template "helloservice-demo.name" . }}
        release: {{ .Release.Name }}
    spec:
      containers:
        - name: {{ .Chart.Name }}
          image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
          imagePullPolicy: {{ .Values.image.pullPolicy }}
      ports:
        - name: http
          containerPort: 9000
```





KubeCon



CloudNativeCon

North America 2020

*Virtual*

# What if I want to orchestrate a multi-stage environment

# Complex releases with Keptn



North America 2020

Virtual



keptn / pod-tato-head ▾

Environment

3 Stages

Services

Integration

dev

staging

production

hello-serverv0.1.0

hello-serverv0.1.0

hello-serverv0.1.0

```
stages:
  - name: "dev"
    approval_strategy:
      pass: "automatic"
      warning: "automatic"
  - name: "staging"
    approval_strategy:
      pass: "automatic"
      warning: "automatic"
    deployment_strategy: "blue_green_service"
  - name: "production"
    approval_strategy:
      pass: "manual"
      warning: "manual"
    deployment_strategy: "blue_green_service"
```

The screenshot displays the Keptn UI across several panels:

- Top Right Panel:** Shows a GitHub repository interface with branches (dev, 4 branches), tags (0 tags), and commit history (5bf3bf9, 35 minutes ago). It includes sections for About, Releases, and Packages.
- Middle Left Panel:** Shows the Keptn environment configuration with three stages: dev, staging, and production, each associated with a hello-server service.
- Middle Right Panel:** Shows the execution of a release. It lists events: Approval triggered (Source: gatekeeper-service), Approval finished (Source: gatekeeper-service), Deployment finished (Source: helm-service), Tests finished (Source: jmeter-service), Evaluation done (Source: lighthouse-service), Evaluation of test on staging (Result: pass, Evaluation timeframe: 2020-10-20 16:00 - 2020-10-20 16:00 (8 seconds)), and Configuration changed (Source: lighthouse-service).
- Bottom Left Panel:** Shows the helloservice deployment details, including the artifact aloisreitbauer/hello-server:v0.1.0 and deployment time Today at 15:59:40.
- Bottom Right Panel:** Shows a heatmap visualization of deployment scores over time, with a legend for pass (green), warning (yellow), fail (red), and info (light gray).



KubeCon



CloudNativeCon

North America 2020

*Virtual*

[Hand over to Harry]



KubeCon



CloudNativeCon

North America 2020

*Virtual*

# What if I want to ship hiding all the details?

# Package as a Helm-Operator



*Virtual*

```
operator-sdk init --plugins=helm --domain=helloservice --group=helloservice-demo --version=v1alpha1 --helm-chart ../charts/|
```



OPERATOR  
FRAMEWORK

```
apiVersion: helloservice-demo.helloservice/v1alpha1
kind: Helloserver
metadata:
  name: helloserver-sample
spec:
  # Default values copied from <project_dir>/helm-charts/helloserver/values.yaml
  imageName: aloisreitbauer/hello-server:v0.1.1
```



KubeCon



CloudNativeCon

North America 2020

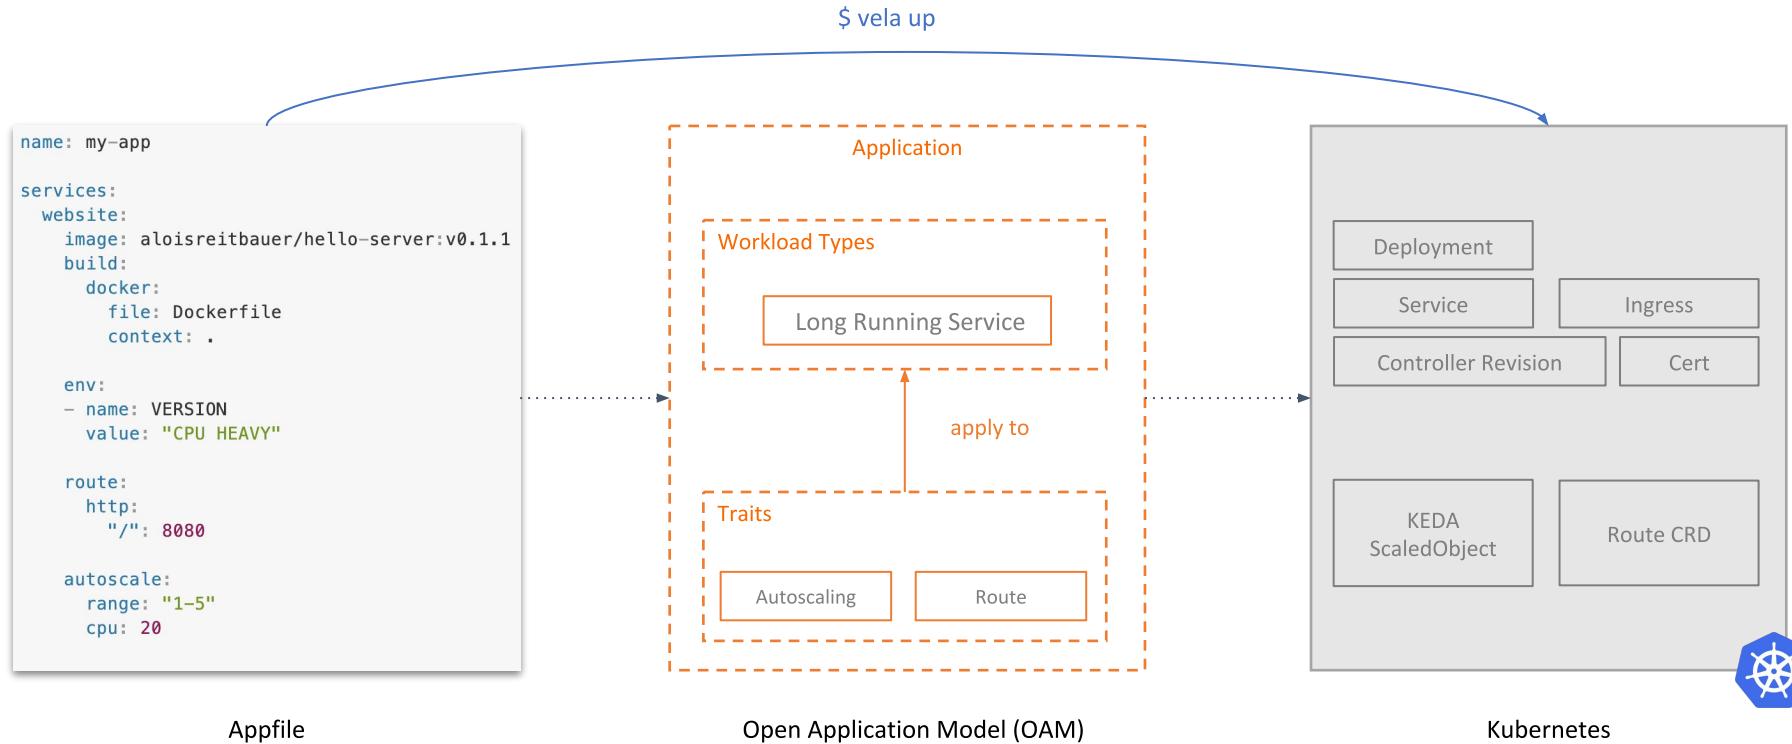
*Virtual*

# What if I want a Heroku but on Kubernetes?

# Ship app using KubeVela



- KubeVela
  - The extensible app platform based on Kubernetes
- Open Application Model (OAM)
  - The model to define applications in developer-centric primitives



# Ship app using KubeVela

- KubeVela is **highly extensible!**

Add a new capability!

```
apiVersion: core.oam.dev/v1alpha2
kind: TraitDefinition
metadata:
  name: canaries.flagger.app
spec:
  appliesToWorkloads:
    - podspecworkload.standard.oam.dev
    - deployments.apps
  definitionRef:
    name: canaries.flagger.app
  workloadRefPath: spec.targetRef
  revisionEnabled: true
```

Done!



```
$ kubectl apply -f trait-def.yaml
```

```
name: my-app

services:
  website:
    image: aloisreitbauer/hello-server:v0.1.1
    build:
      docker:
        file: Dockerfile
        context: .

  env:
    - name: VERSION
      value: "CPU HEAVY"

  route:
    http:
      "/": 8080

  autoscale:
    range: "1-5"
    cpu: 20

  rollout:
    strategy: canary
    analysis:
      interval: 30s
      threshold: 10
      maxWeight: 50
      stepWeight: 10
      maxReplicas: 4
```

```
$ vela up
```

Your newly added capability!

# What is next for podTato head ?



*Virtual*



<https://github.com/cncf/podtato-head>

Extending **podTato head** for more complex use case:

- multiple services
- stateful workloads
- databases
- cross-stage secrets
- external dependencies
- .....

Add additional tools and use cases

- (Chaos) Testing
- App operations automation

# What is next for SIG App Delivery ?



*Virtual*



- The community: <https://github.com/cncf/sig-app-delivery>
- We have meetings every 1st and 3rd Wednesday of the month at 8:00 a.m. PST.

**“Show me the code, not whitepapers!”**

- More real-world demonstrations to answer the question of:
  - *What project X is used for?*
- Series tech blog about:
  - *Deep dive project Y*
  - *How company Z is using project X to solve real-world problems?*
- Other plans:
  - **CNCF Radar for application delivery**

