

**RESILIENCE**  
**REALIZED**



**KubeCon**



**CloudNativeCon**

North America 2021



KubeCon



CloudNativeCon

North America 2021

RESILIENCE

REALIZED

# Cloud Native Storage: Storage TAG Intro, Projects, Landscape & Technology

*Alex Chircop, StorageOS*

*Xing Yang, VMware*

*Raffaele Spazzoli, Red Hat*

# Agenda



KubeCon



CloudNativeCon

North America 2021

- Overview of the TAG, how to join and how to help
- Overview of storage projects in the CNCF
- Projects that are currently being reviewed
- Overview of the CNCF Storage Landscape document
- Overview of the Performance and Benchmarking document
- Overview of the DR document



## TAG Storage

CNCF SIGs were renamed TAGs (Technical Advisory Groups)

Meetings are on the 2nd and 4th Wednesday of every month at 8AM PT (USA Pacific)

- Home: <https://github.com/cncf/tag-storage>
- Conf call: <http://bit.ly/cncf-storage-tag-call>
- Agenda: <http://bit.ly/cncf-storage-tag-minutes>
- Recordings: <http://bit.ly/cncf-storage-tag-recordings>
- Mail list: <https://lists.cncf.io/g/cncf-tag-storage>

*Our calls and membership are open!*

# Who are we

- We are a diverse set of users & developers of Cloud Native technologies with a storage focus
- We are leaders & early adopters
- We are organized as:

## Co-Chairs

- Alex Chircop
- Quinton Hoole
- Xing Yang

## Tech Leads

- Raffaele Spazzoli
- Luis Pabón
- Sheng Yang
- Nick Connolly

## TOC Liaisons

- Saad Ali
- Erin Boyd

**“Scale contributions by the CNCF technical and user community, while retaining integrity and increasing quality in support of the CNCF mission (to make cloud native computing ubiquitous).”**

*...this means we*

- Educate
- Review Storage Project Proposals
- Engage with the user community
- Work directly with the TOC & other TAGs as subject matter experts

## End User Education

- Provide up-to-date, high quality, unbiased and easy-to-consume material to help end users to understand and effectively adopt cloud-native technologies and practises within the TAG's area, for example:
  - White papers, presentations, videos, or other forms of training clarifying terminology, comparisons of different approaches, available projects or products, common or recommended practices, trends, illustrative successes and failures, etc.

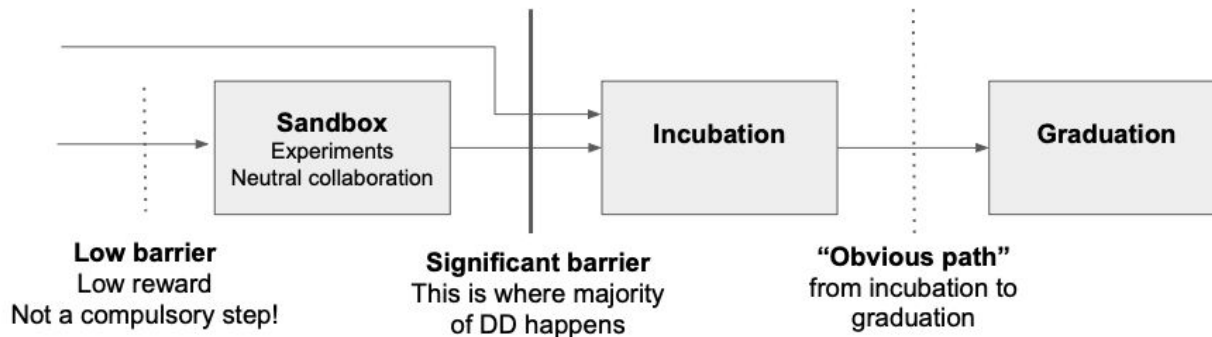
[CNCF Storage White Paper \(https://bit.ly/cncf-storage-whitepaperV2\)](https://bit.ly/cncf-storage-whitepaperV2)

[CNCF TAG Storage Performance & Benchmarking Whitepaper \(https://bit.ly/cncf-tag-storage-performance-benchmarking\)](https://bit.ly/cncf-tag-storage-performance-benchmarking)

- As far as possible, information should be based on research and fact gathering

# Project Stages

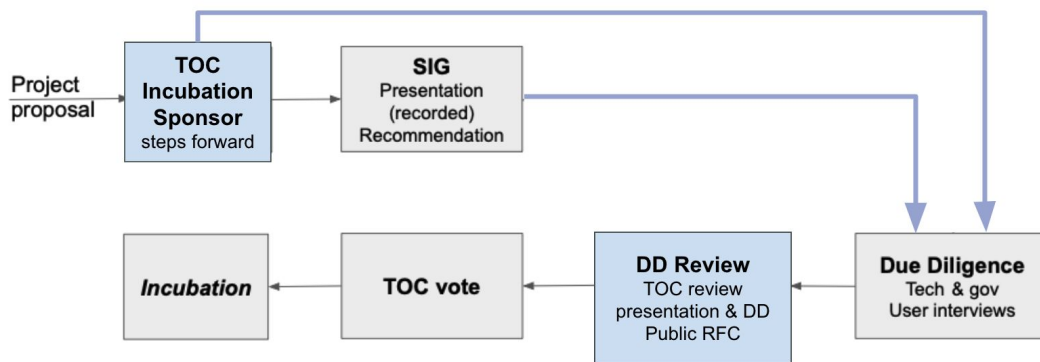
<https://www.cncf.io/projects/>



Sandbox	Incubation	Graduation
<ul style="list-style-type: none"><li>■ Experiments</li><li>■ IP Policy</li><li>■ Build Community</li></ul>	<ul style="list-style-type: none"><li>■ Used successfully in production</li><li>■ Healthy number of committers</li><li>■ Project metrics</li></ul>	<ul style="list-style-type: none"><li>■ Mainstream production use</li><li>■ Security audits</li><li>■ Committers from multiple organisations</li></ul>



# Project Review



- Understand and **document** a high level roadmap of projects within this space, including CNCF and non-CNCF projects.
- For projects that fall within the CNCF, perform **health checks**.
- Perform **discovery** of and **outreach** to candidate projects.
- **Help candidate projects** prepare for presentation to the TOC.
- Every CNCF project will be assigned to one suitable TAG by the TOC.

## End User Input Gathering

- Gather useful end user input and feedback regarding expectations, pain points, primary use cases, etc.
- Compile this into easily consumable reports and/or presentations to assist projects with feature design, prioritization, UX, etc.

## Community Enablement

- TAGs are open organizations with meetings, meeting agendas and notes, mailing lists, and other communications in the open.
- The mailing list, TAG meeting calendar, and other communication documents of the TAG will be openly published and maintained.

## Trusted Expert Advisors to the TOC

- Perform **technical due diligence** on new and graduating projects, and advise TOC on findings.
- Be involved with, or **periodically check** in with projects in their areas, and advise TOC on health, status and proposed actions (if any) as necessary or on request.
- Note: in all cases, the TOC make the final decisions

# CNCF Storage Projects




KubeCon



CloudNativeCon

North America 2021

## Graduated



ROOK



Vitess



HARBOR

## Incubating



Dragonfly



etcd



TiKV

CNCF Projects: <https://www.cncf.io/projects/>

Sandbox Projects: <https://www.cncf.io/sandbox-projects/>

# Current Projects in Review



KubeCon



CloudNativeCon

North America 2021

- **Sandbox → Incubation**

- Longhorn

- <https://github.com/longhorn/longhorn>

- ChubaoFS

- <https://github.com/chubaofs/chubaofs>

- OpenEBS

- <https://github.com/openebs/openebs>

# Storage projects listed in TAG repo



KubeCon



CloudNativeCon

North America 2021

- <https://github.com/cncf/tag-storage>

## Current CNCF Storage Projects

---

### Graduated Projects

- [etcd](#)
- [Rook](#)
- [TiKV](#)
- [Vitess](#)

### Incubating Projects

- [Dragonfly](#)

### Sandbox Projects

- [ChubaoFS](#)
- [Longhorn](#)
- [OpenEBS](#)
- [Pravega](#)
- [Piraeus](#)
- [Vineyard](#)

# CNCF Storage Whitepaper

- **Whitepaper:** <https://bit.ly/cncf-storage-whitepaperV2>
- **Definition of the attributes of a storage system**
- **Definition of the layers in a storage solution with a focus on terminology and how they impact the attributes**
- **Definition of the data access interfaces in terms of volumes and application APIs**
- **Definition of the management interfaces**

# Storage Attributes

## Availability

- Failover
- Moving access between nodes
- Redundancy
- Data Protection

## Scalability

- Clients
- Operations
- Throughput
- Components

## Performance

- Latency
- Operations
- Throughput

## Consistency

- Delay to access correct data after a commit
- Delay between commit and data being committed to non-volatile store

## Durability

- Data protection
- Redundancy
- Bit-Rot



**Orchestrator, Host and Operating System**

**Storage Topology**

(centralized, distributed, sharded, hyperconverged)

**Data Protection**

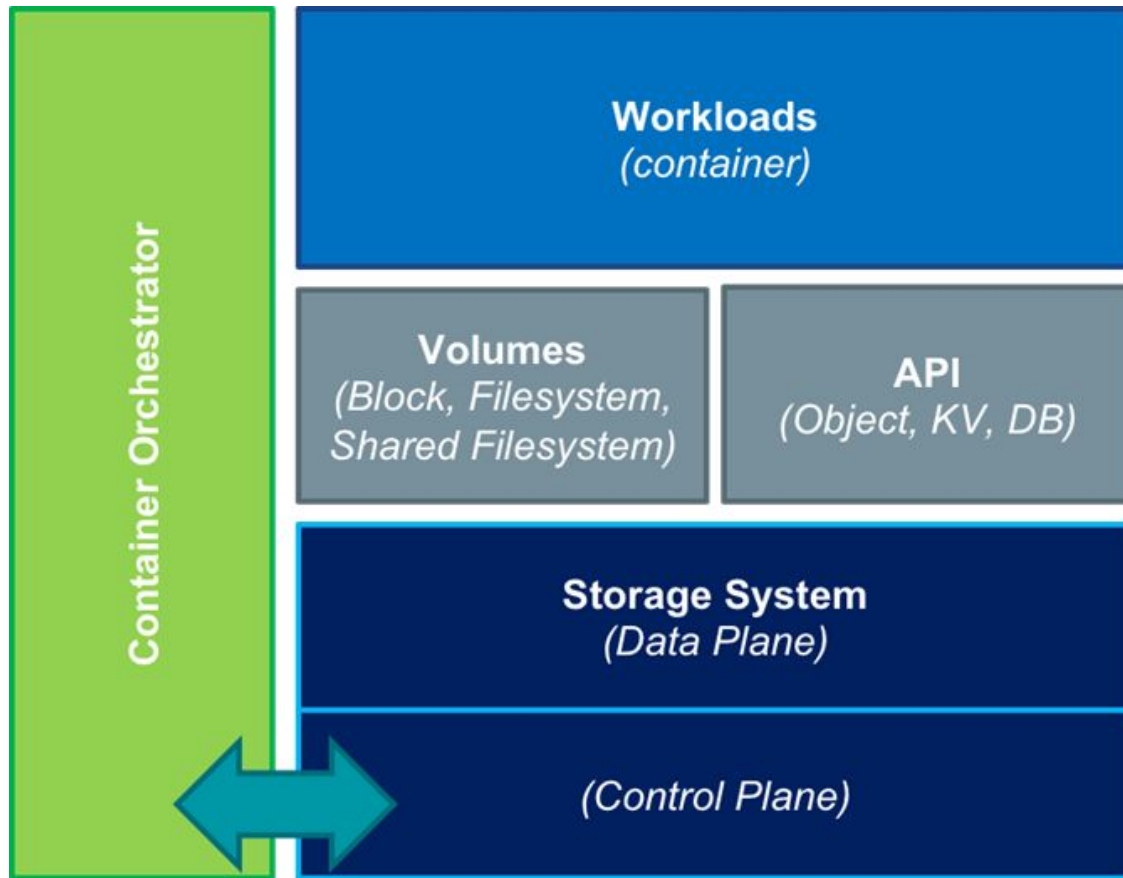
(RAID, Erasure coding, Replicas)

**Data Services**

(Replication, Snapshots, Clones, etc.)

**Physical, Non-Volatile Layer**

# Data Access Interfaces



Storage can be accessed via **Data Access Interfaces**:

- **Volumes** – accessed through a more traditional file interface in a **block** or **filesystem** interface
- **API** – other ways to persist data such as **object stores**, **KV stores** or **databases**

# Management Interfaces

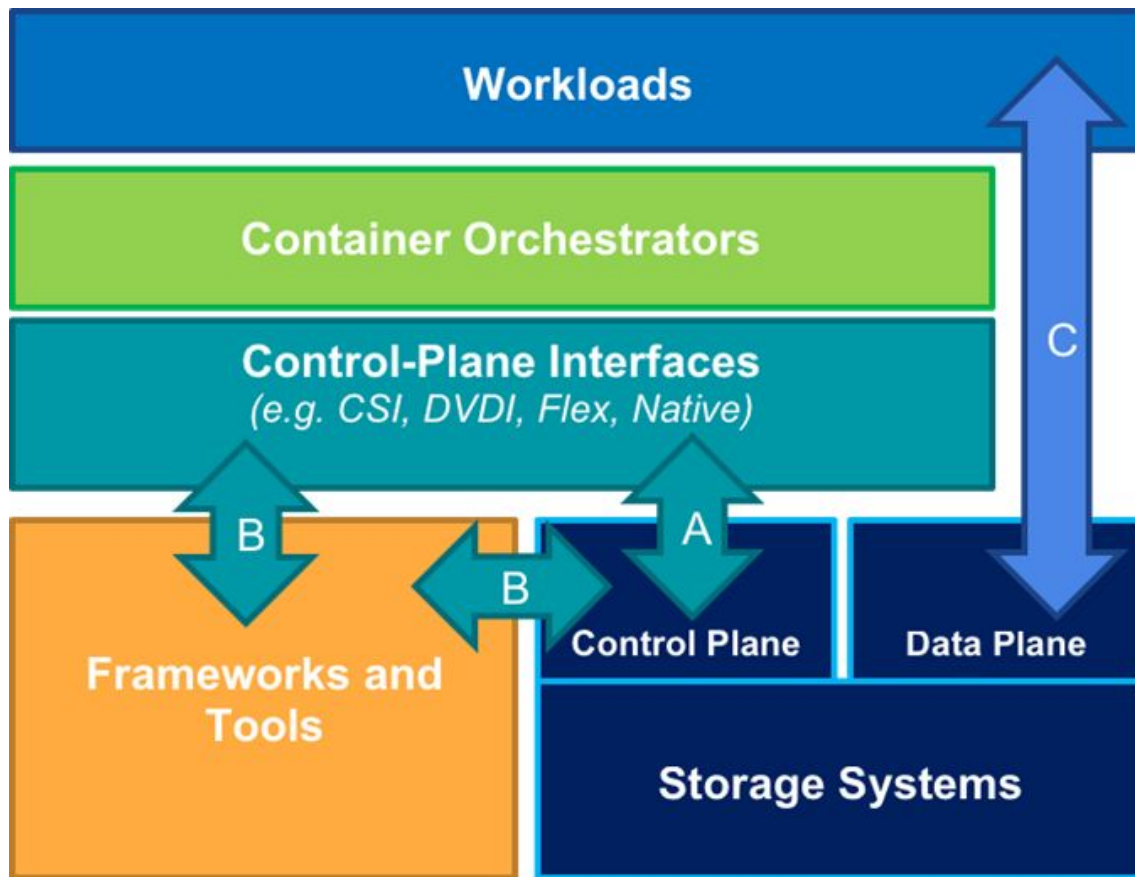


KubeCon



CloudNativeCon

North America 2021



**Container Orchestration system (CO)** uses an interface to interact with a storage system

The storage system can:

**(A)** support control-plane API directly

**(B)** interact via an API Framework layer or other Tools

Workloads consume **(C)** storage via a data access interface

# Performance Whitepaper



KubeCon



CloudNativeCon

North America 2021

- **Whitepaper:**  
<https://bit.ly/cncf-tag-storage-performance-benchmarking>
- **Definition of common concepts for measuring performance and benchmarking for volumes and databases**
- **Definition of common pitfalls and considerations**
  - **Basics: operations vs throughput**
  - **Topology, Data Protection, Data Reduction, Encryption**
  - **Latency**
  - **Concurrency: queue depths, multiple clients and backends**
  - **Caching at multiple layers**
  - **Managing the environment, Cloud, and Client headroom**

# Performance Whitepaper

- **Whitepaper:**  
<https://bit.ly/cncf-tag-storage-performance-benchmarking>
- **Important takeaway:** published results are not useful for making comparisons - it is hard to compare published results without a deep understanding of the test conditions, so it is always important to run **your own test**, on **your own environment** with **your own applications**

# Cloud Native Disaster Recovery

- **Whitepaper:** <http://bit.ly/cncf-cloud-native-DR>

Concern	Traditional DR	Cloud Native DR
Type of deployment	active/passive, rarely active/active	Active / active
Disaster Detection and Recovery Trigger	Human	Autonomous
Disaster Recovery Procedure execution	Mix of manual and automated tasks	Automated
Recovery Time Objective (RTO)	From close to zero to hours	Close to zero
Recovery Point Objective (RPO)	From zero to hours	Exactly zero for strongly consistent deployments. Theoretically unbounded, practically close to zero for eventual consistent deployments.
DR Process Owner	Often the Storage Team	Application Team
Capabilities needed for DR	From storage (backup/restore, volume replication)	From networking (east-west communication, global load balancer)

# CNDR Whitepaper content



KubeCon



CloudNativeCon

North America 2021

- Cloud Native Disaster Recovery (CNDR) definition
- Definitions: Failure Domain, HA, DR
- CAP theorem: Consistency, Availability, Network Partitioning
- Anatomy of Distributed Stateful workloads: Replicas, Shards, API Layer, Storage Layer
- Consensus Protocols: Paxos, Raft, 2PC, 3PC
- CNDR Reference Architectures:
  - Strong Consistency
  - Eventual Consistency

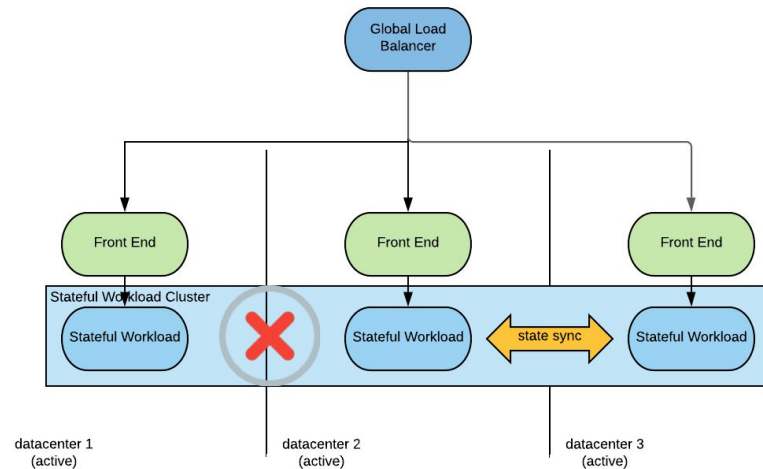
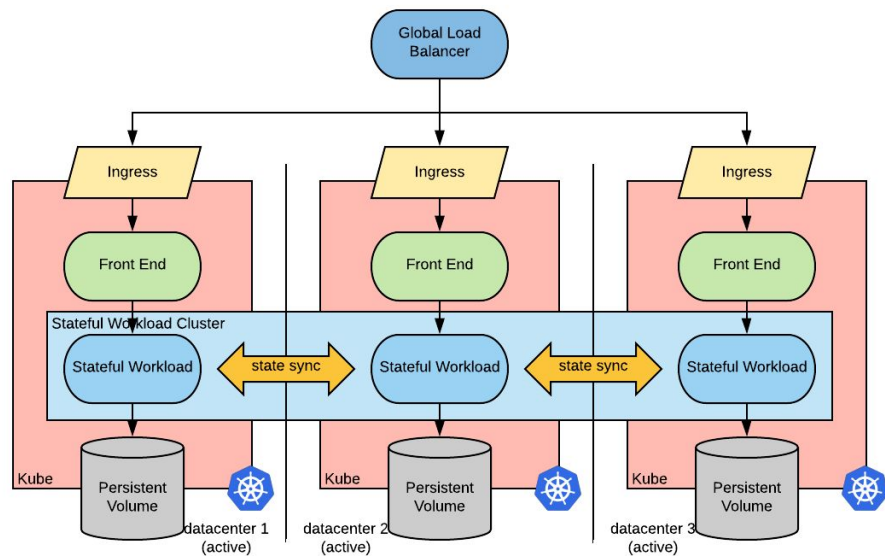
# Examples of Consensus Protocol choices



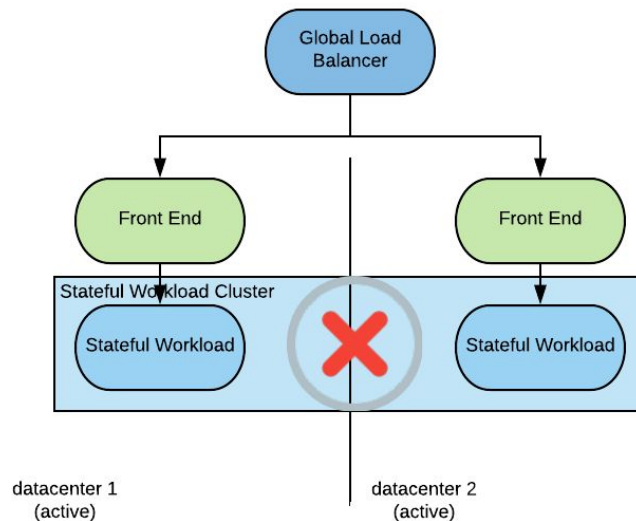
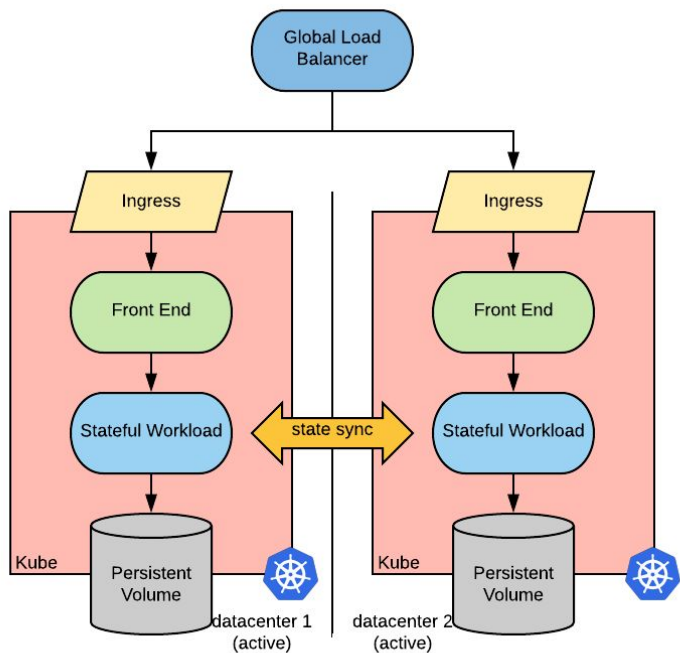
Product	Replica consensus protocol	Shard consensus protocol
Etcd	Raft	N/A (no support for shards)
Consul	Raft	N/A (no support for shards)
Zookeeper	Atomic Broadcast (a derivative of Paxos)	N/A (no support for shards)
ElasticSearch	Paxos	N/A (No support for transactions)
Cassandra	Paxos	Supported, but details are not available.
MongoDB	Paxos	Homegrown protocol.
CockroachDB	Raft	2PC
YugabyteDB	Raft	2PC
TiKV	Raft	Percolator
Spanner	Raft	2PC+high-precision time service
Kafka	A custom derivative of PacificA	Custom Implementation of 2PC



# CNDR -- Strong Consistency - Kubernetes Reference Architecture



# CNDR -- Eventual Consistency - Kubernetes Reference Architecture



- How you can get involved?
  - Join our meeting
    - 2nd & 4th Wednesday each month
  - Submit and help review projects for consideration
    - <https://github.com/cncf/toc/tree/master/process>
- We value community presentations of projects in the cloud native storage space including, but not limited to: *management frameworks, block stores, filesystems, object stores, key-value stores and databases*
- Several projects have presented to the TAG such as: CSI, Rook, REX-Ray, TiKV, Dotmesh, Yugabyte, OpenEBS, Open Services Broker, Vitess, Minio, OpenSDS, Redfish/Swordfish, ChubaoFS, Longhorn, Dragonfly, Harbor, Pravega, Piraeus, Dataset Lifecycle Framework, Linstor, Vineyard



RESILIENCE

REALIZED



KubeCon



CloudNativeCon

North America 2021

*Thank You*