

**RESILIENCE**  
**REALIZED**



**KubeCon**



**CloudNativeCon**

North America 2021



KubeCon



CloudNativeCon

North America 2021

RESILIENCE

REALIZED

# Correlating Signals in OpenTelemetry

*Jaana Dogan, Principal Engineer @ AWS*

*Morgan McLean, Director of PM @ Splunk*



KubeCon



CloudNativeCon

North America 2021



Morgan McLean  
Director of PM @ Splunk



Jaana Dogan  
Principal Engineer @ AWS

# OpenTelemetry: Quick Intro

OpenTelemetry provides:

- A Collector (typically run as an agent / daemonset, can be run as a clustered service)
- SDKs and language auto-instrumentation agents
- A protocol
- A specification and semantic conventions for various use cases
- A community!

OpenTelemetry is opinionated: there is one set of semantic conventions for a given scenario, there is one SDK and / or language agent for a given language, etc.

# OpenTelemetry: Quick Intro



KubeCon



CloudNativeCon

North America 2021

OpenTelemetry captures:

- Traces (GA)
- Metrics (beta, possibly GA by the time that you're watching this!)
- Logs (alpha, **NOT BETA**)
- Resource metadata

- If you don't capture and analyze traces, metrics, logs, and other telemetry, **you are setting yourself up for failure**
- If you capture and analyze traces, metrics, logs, and other telemetry in siloes, **you are setting yourself up for failure**

## WHY ARE YOU SETTING YOURSELF UP FOR FAILURE?

The point of capturing and analyzing telemetry is to give you visibility into the entire stack! Looking at pieces in isolation doesn't achieve this.

Failures 10 levels down in a system will manifest in very strange ways at upper layers!

# Example: Errors Causing Latency

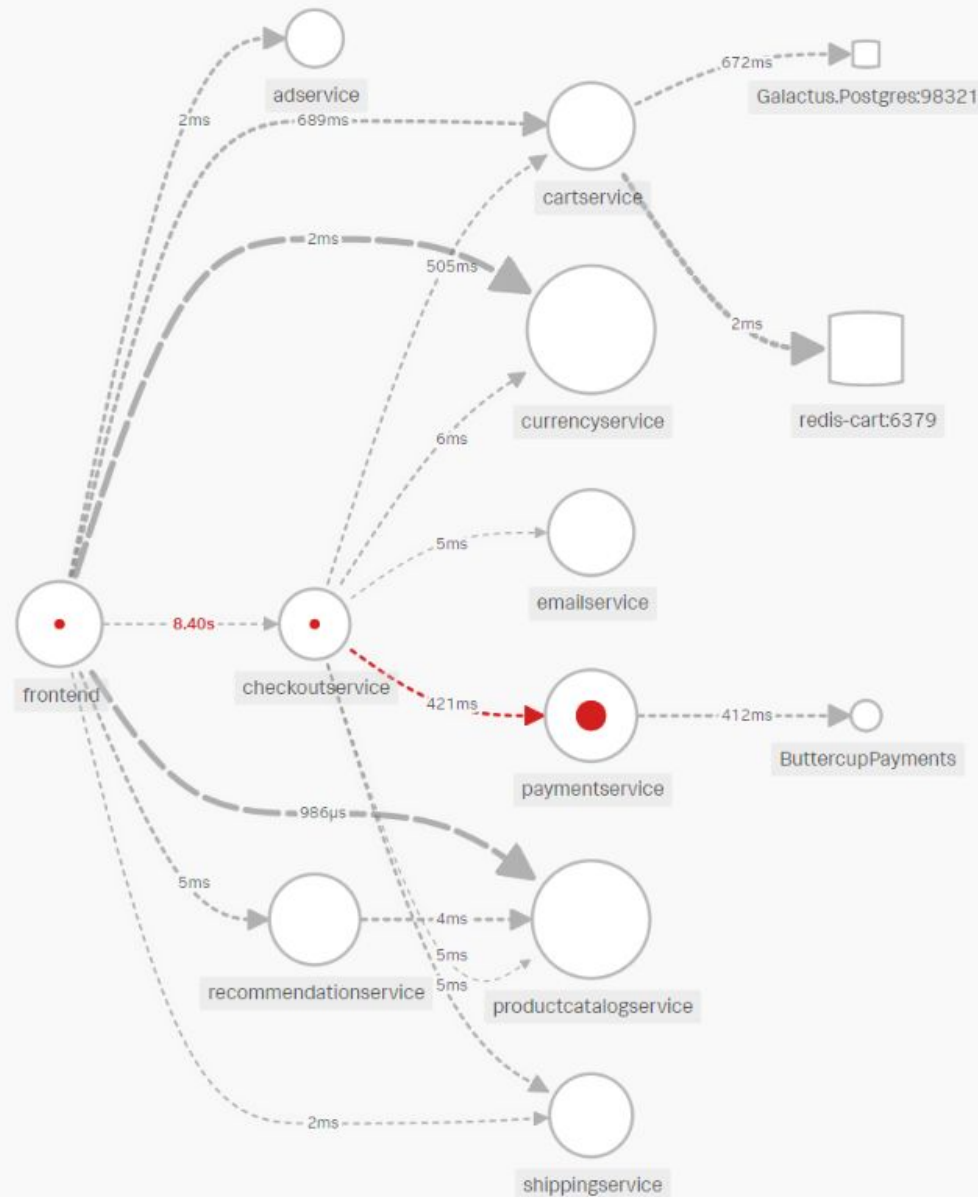


KubeCon



CloudNativeCon

North America 2021



# Example: Errors Causing Latency

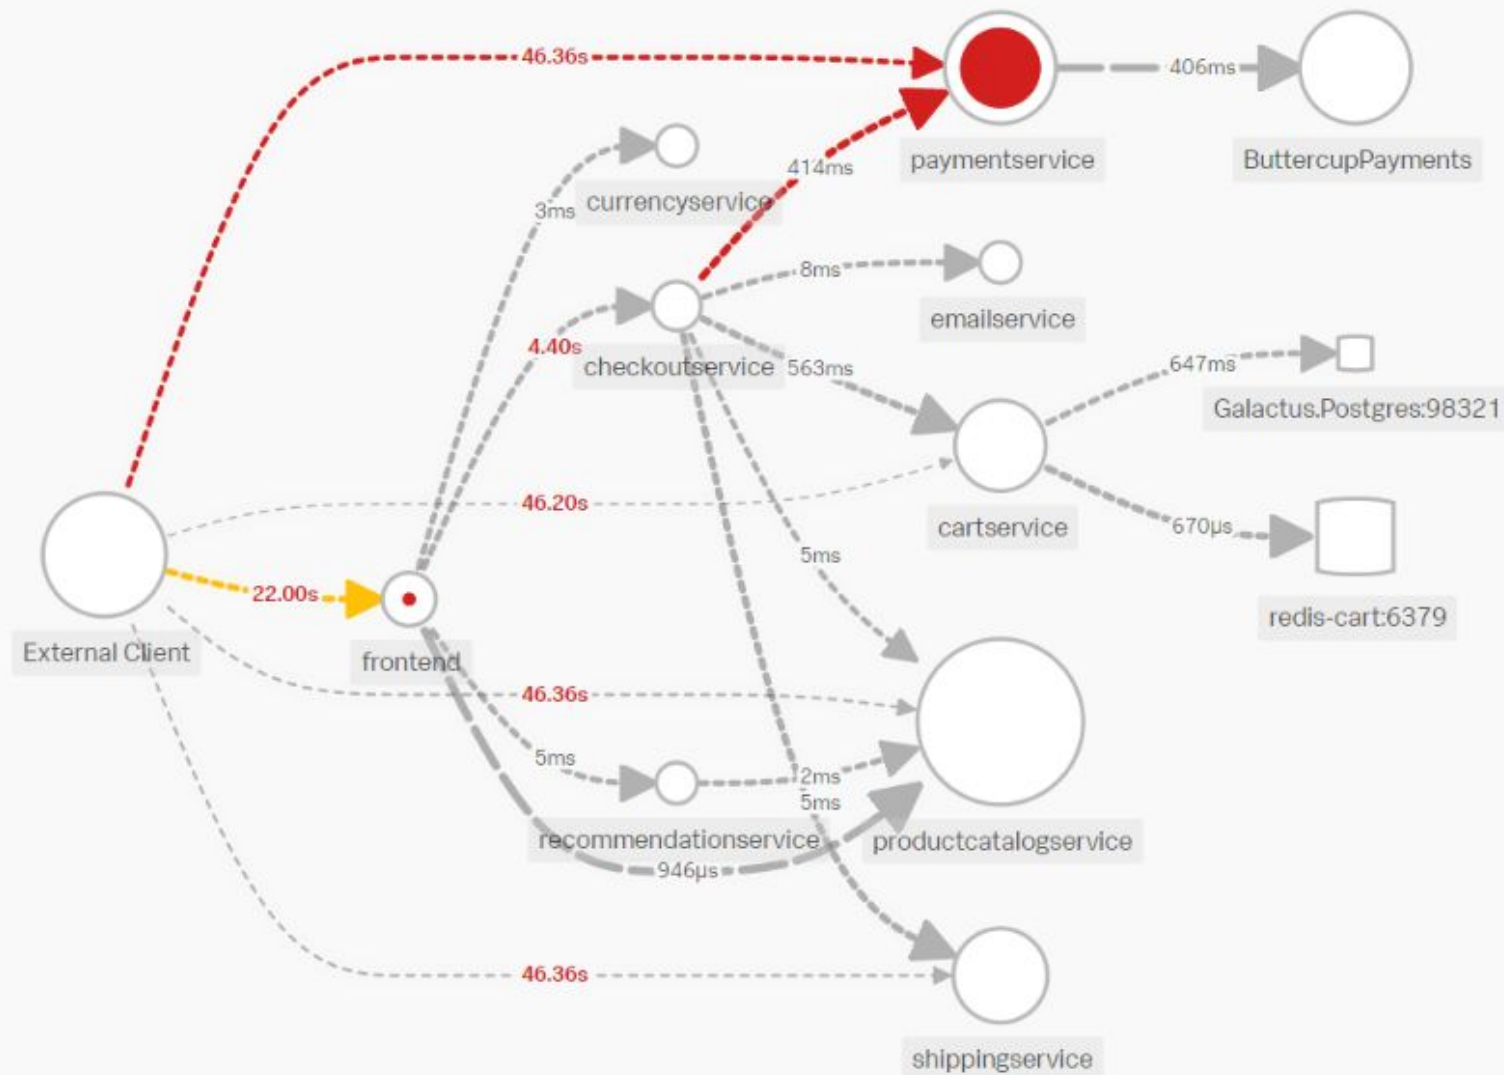


KubeCon



CloudNativeCon

North America 2021





# Example: Errors Causing Latency

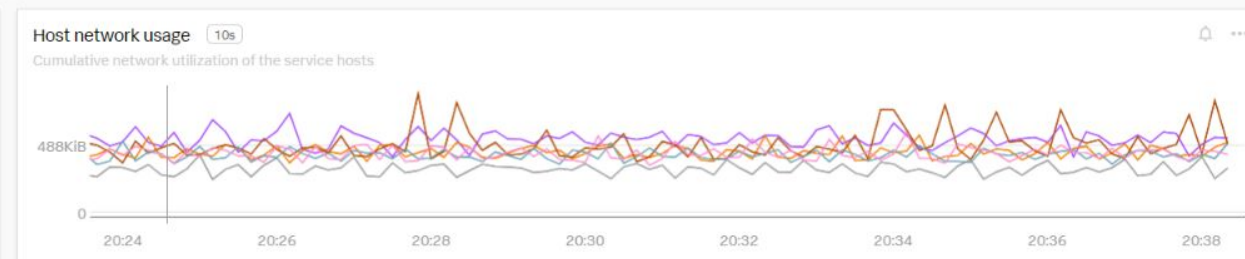
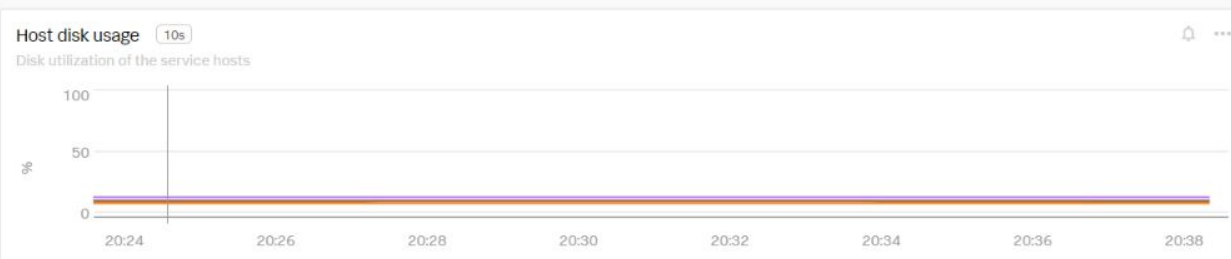
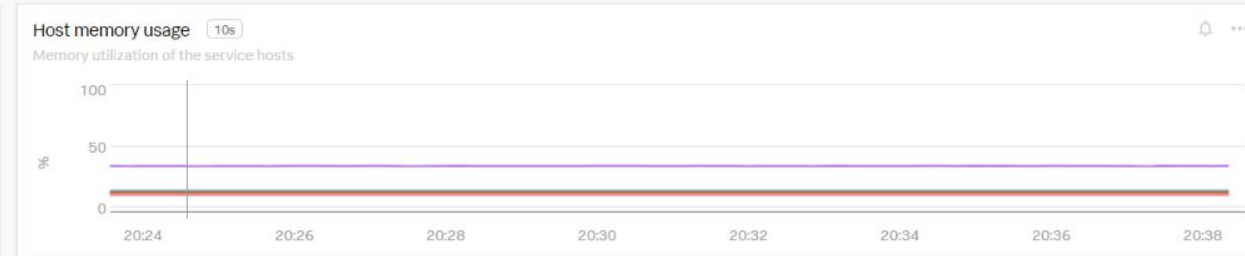
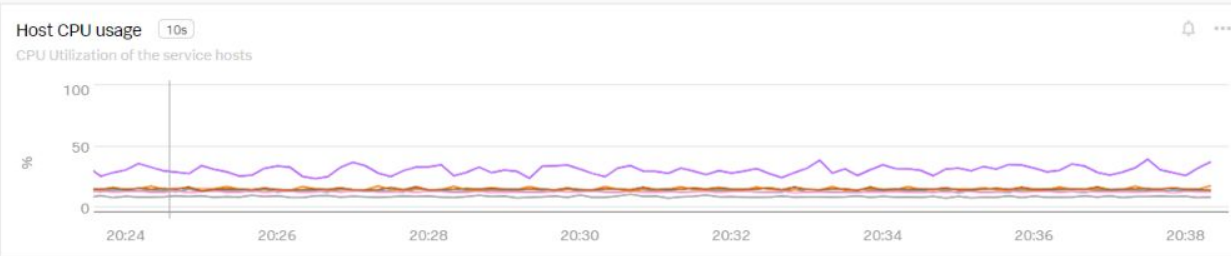


KubeCon



CloudNativeCon

North America 2021



# Example: Errors Causing Latency

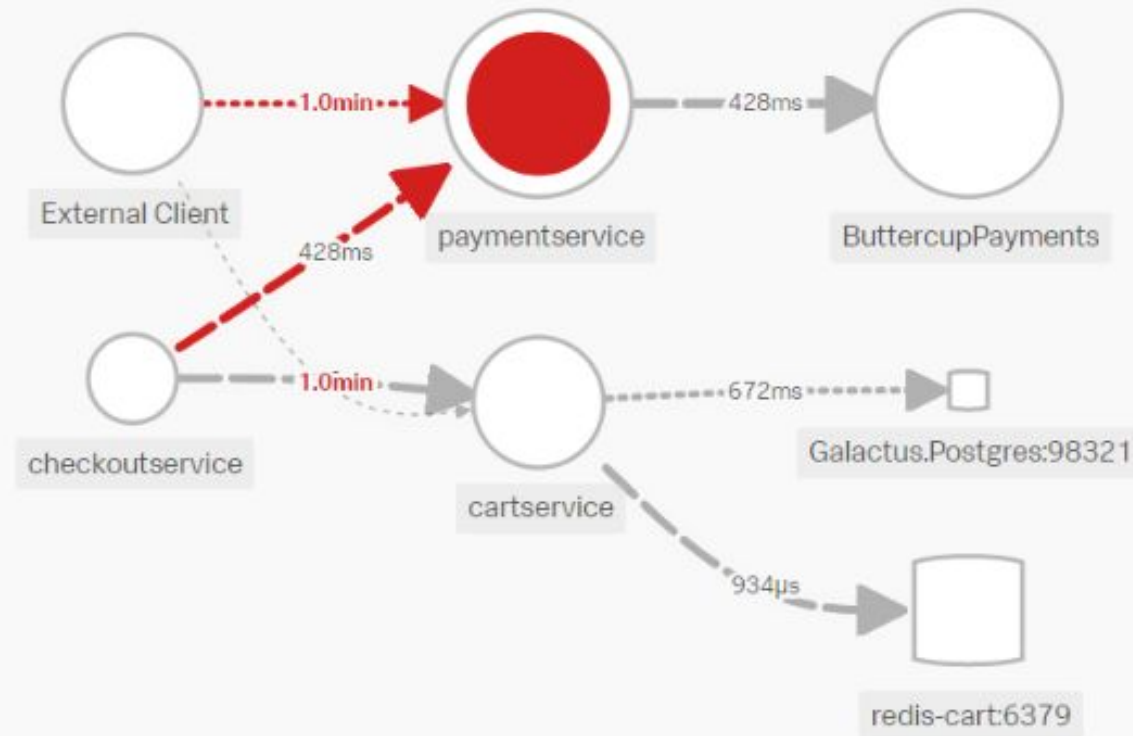


KubeCon



CloudNativeCon

North America 2021



# Example: Errors Causing Latency



KubeCon

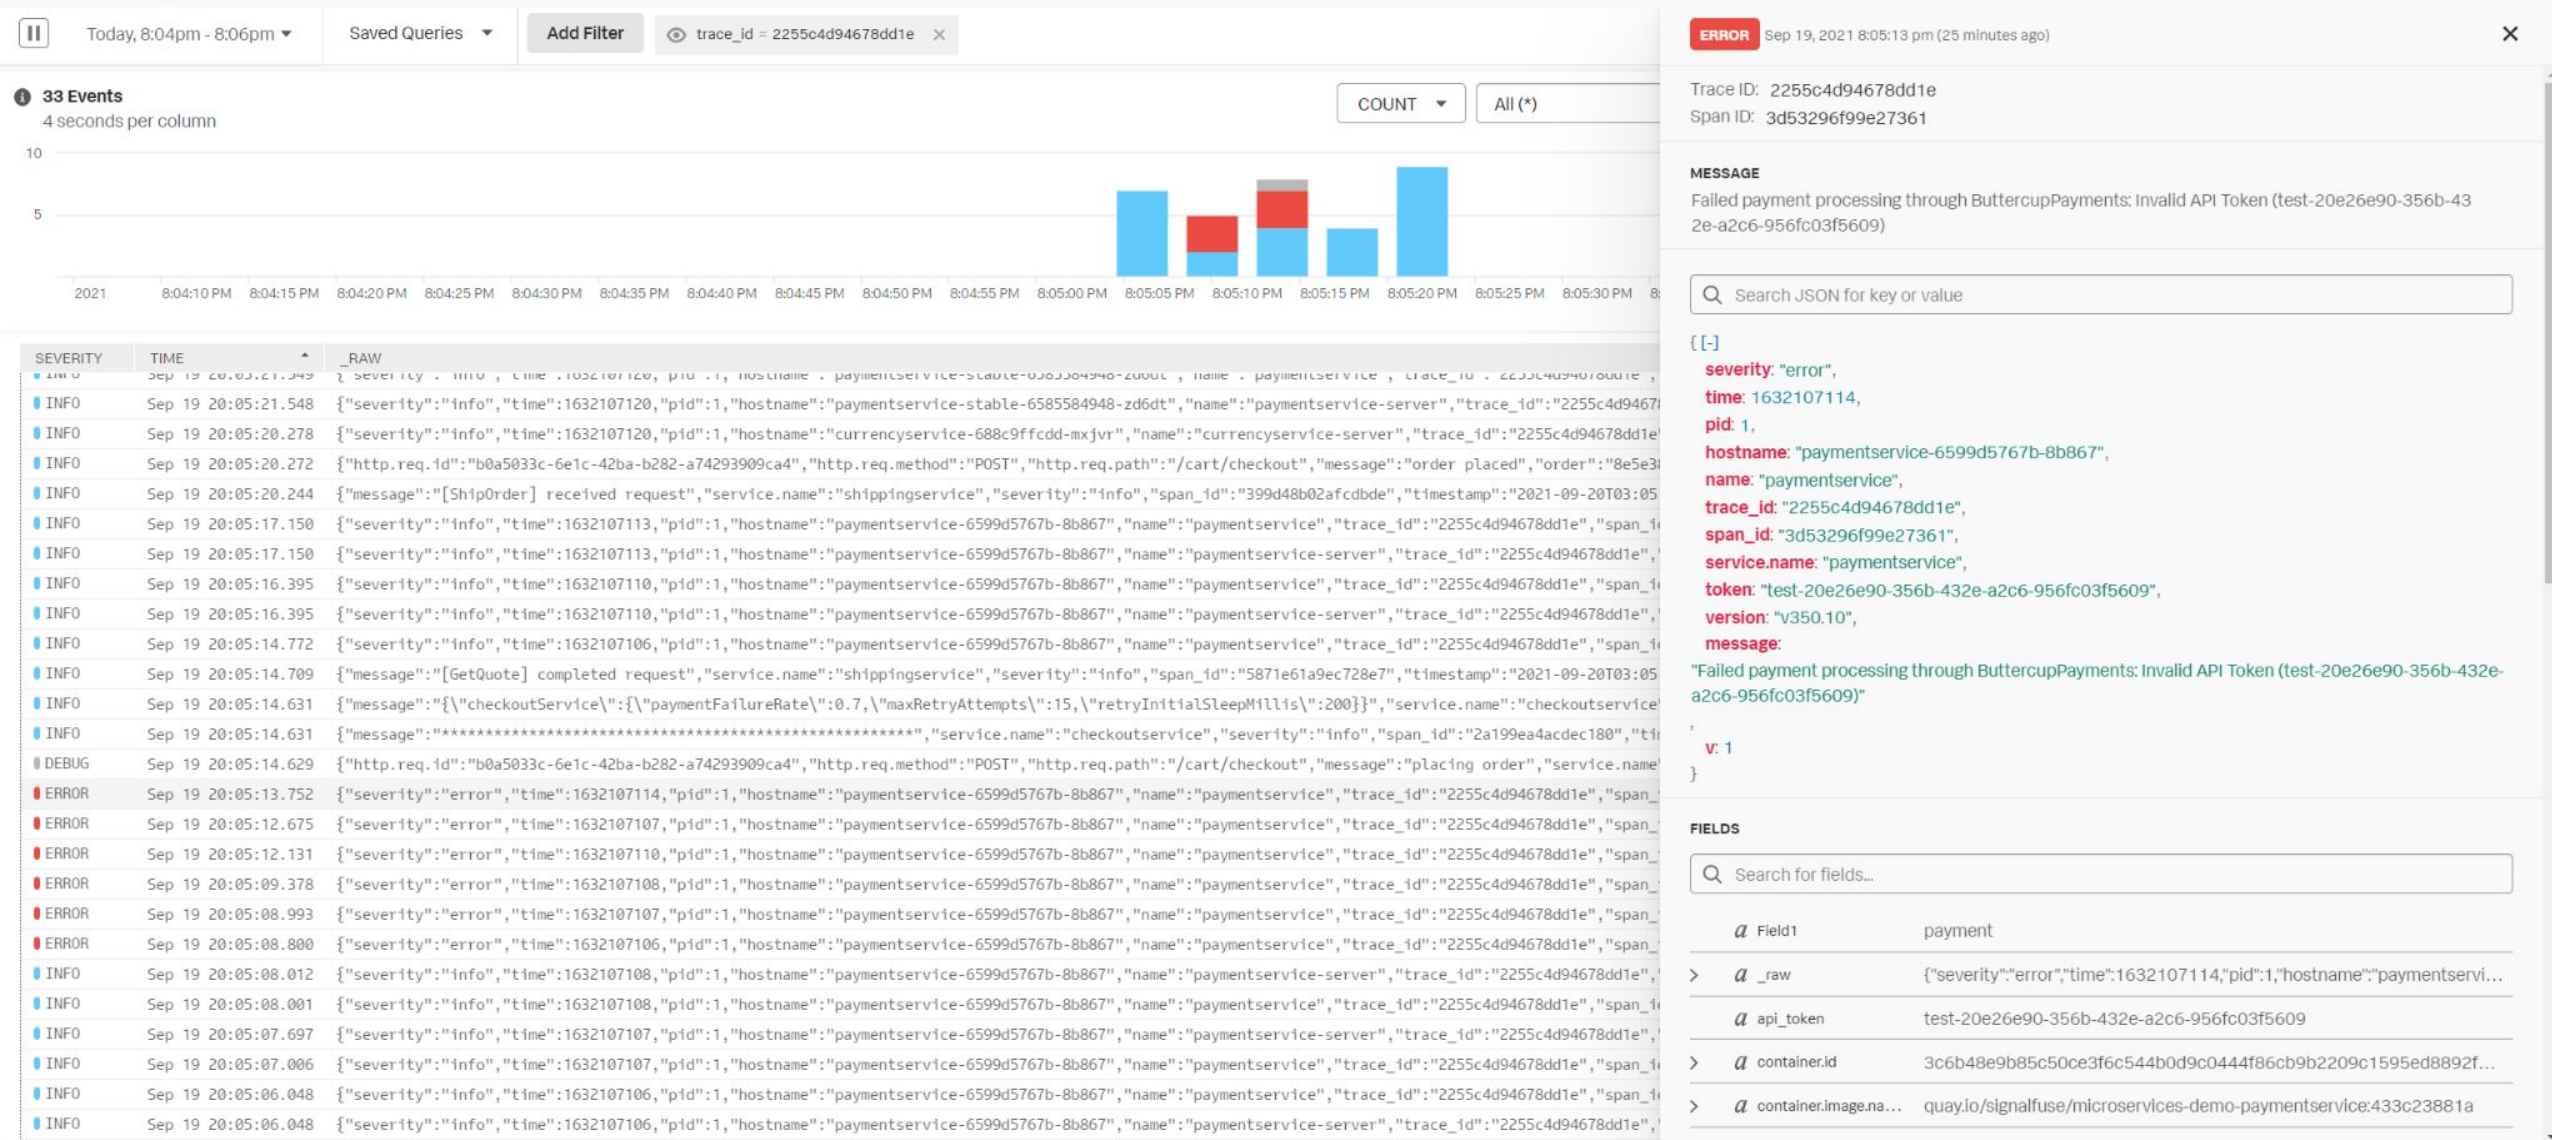


CloudNativeCon

North America 2021



# Example: Errors Causing Latency



# Example: Errors Causing Latency

How did we do this?

- Correlate application metrics with services
- Correlate host metrics with services
- Propagate request data all the way down the stack, correlate it with metrics
- Correlate logs with traces
- Correlate logs with services

# Example: Errors Causing Latency

Ok, but how did we actually do this?

- Use OpenTelemetry SDKs or language instrumentation to capture traces and application metrics
  - Send to the Collector
- Use the OpenTelemetry Collector to capture host metrics, host logs, and application logs

Important configuration

- Set service name in resource metadata
- Stamp trace and span IDs in logs



# Example: Crashes Causing Cost

Service is causing latency but there is nothing new in the traces or logs.

- Elevated CPU usage from nodes
- Breaking down CPU usage reveals simulations are impacted
- Searching for traces and logs reveals nothing new
- Simulation engine RPCs experience slight new latency

# Example: Crashes Causing Cost

Metrics, traces or logs are not useful. What else is available?

- Crash dumps from the nodes reveal consistent crashes
- Crash occurs after serving the RPC
- Each incoming request causes a cold start
- Debugging code locally reveals service panics

Simulation engine was crashing after handling each request and new one was being forked. Although it caused little new latency, it elevated the CPU usage and impacted efficiency.



# Example: Library Causing Latency

Service begins to experience latency and elevated CPU usage.

- A large number of changes have been pushed recently
- Checking the cluster reveals no impact to other services
- Traces show nothing new, no retries or outbound request issues

# Example: Library Causing Latency

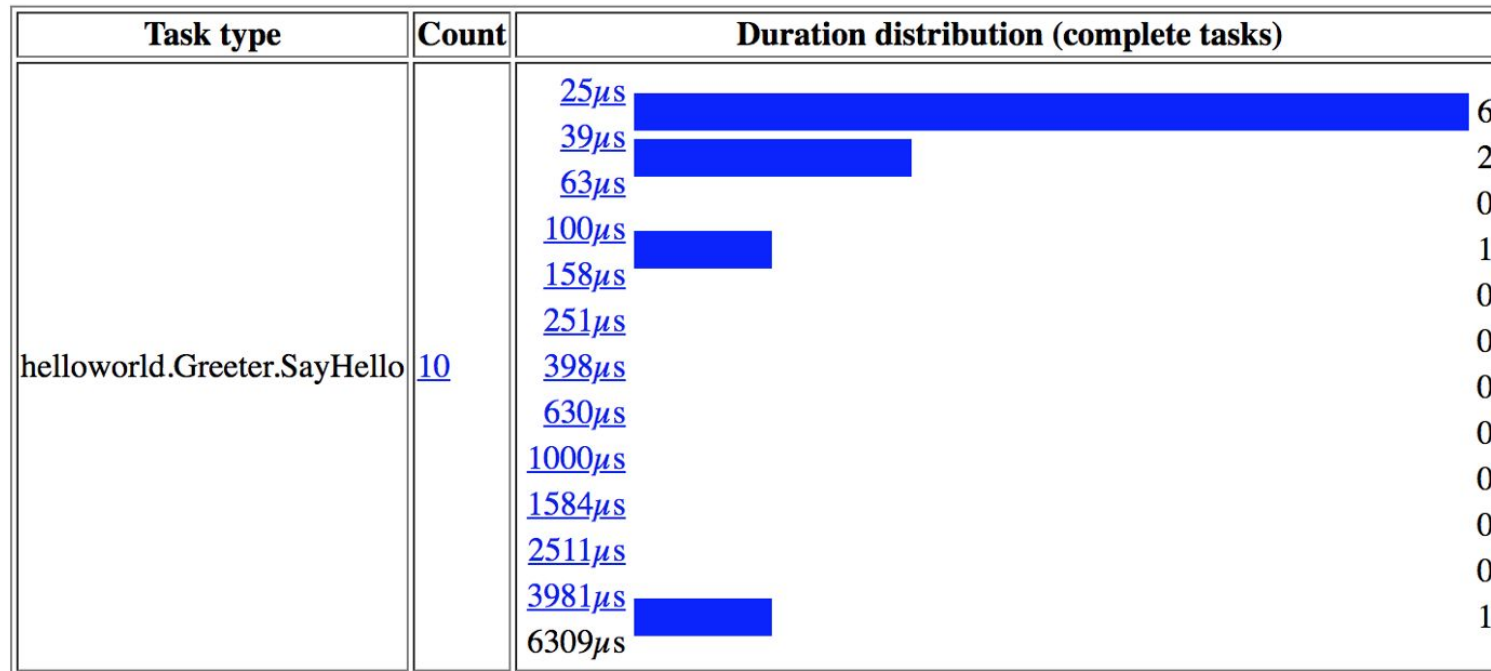
Reviewing diffs reveals new versions of library dependencies.

- Briefly enabling profiling confirms
- Profiles reveal only a certain RPC is impacted
- Reading RPC handler reveals the use a new compression library

Service was 50% more CPU intensive due to the new compression library. Rolling back mitigated the issue rather than running out of CPU capacity.

# Example: Library Causing Latency

Since 1.11, Go runtime tracer can report events in the lifetime of an RPC.



`type=connection.init,pc=13eccb1,latency >= 1ms,latency <= 1.584893ms`

Goroutine	Task	Total	Execution	Network wait	Sync block	Blocking syscall	Scheduler wait	GC sweeping	GC pause
<a href="#">1</a>		1309μs		372μs	0ns	0ns	174μs	204μs	0ns (0.0%) 1015μs (77.6%)

# Correlations

Types of correlations:

- Correlations among metrics, logs, traces
- Resource inferences
- Local and propagated annotations

Correlations don't just make it easier to navigate telemetry but also make it easier to rule out non-issues quickly.

# What's Next?



KubeCon



CloudNativeCon

North America 2021

- Probably profiles
  - Will take a while; need to ship metrics and finish logs first)
  - Can correlate performance data all the way down to code!
- Core dumps, crashes, exception call stacks
- Associations with resource information gained from eBPF
  - Correlate network performance information with traces!
- Correlations with language runtime traces

**RESILIENCE**  
**REALIZED**



**KubeCon**



**CloudNativeCon**

North America 2021