



KubeCon



CloudNativeCon

North America 2021

RESILIENCE

REALIZED

# Faster Container Image Distribution on a Variety of Tools with Lazy Pulling

*Kohei Tokunaga, NTT Corporation*

*Tao Peng, Ant Group*

# Pull is time-consuming

- Large images speeds down cold start of containers
- Build takes long if base images are large
- Not all images are minimizable e.g. language runtimes, frameworks, ..

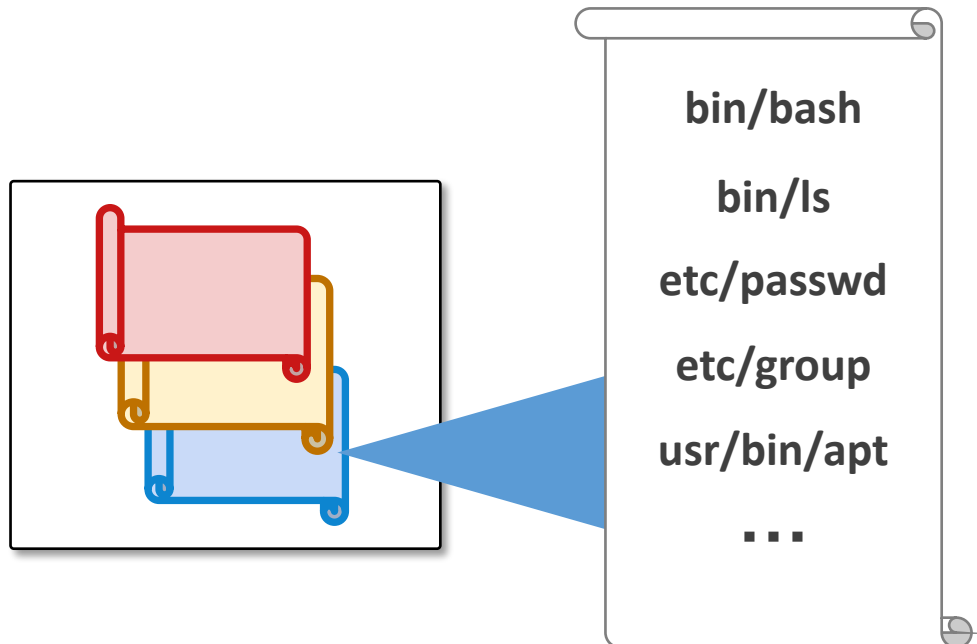
pulling packages accounts for 76% of container start time,  
but only 6.4% of that data is read [Harter et al. 2016]

[Harter et al. 2016] Tyler Harter, Brandon Salmon, Rose Liu, Andrea C. Arpaci-Dusseau, Remzi H. Arpaci-Dusseau. "Slacker: Fast Distribution with Lazy Docker Containers". 14th USENIX Conference on File and Storage Technologies (FAST '16). February 22–25, 2016, Santa Clara, CA, USA

# Problem on the current OCI image

container can't start before the all layers become locally available

Image is a set of **tar (+compression) layers**



- Non seekable
  - Need to scan the entire blob even for extracting single file entry
- No parallel extraction
  - Need to scan sequentially

## Lazy pulling:

Starting up containers without waiting for the pull completion

- **eStargz**

- Lazy pullable format with prefetch optimization + content verification
- Proposed as a backward-compatible extension to OCI Image Spec

- **Nydus**

- Lazy pullable format with prefetch, chunk dedup and e2e data integrity
- Compatible with OCI Distribution Spec and Artifacts Spec
- Proposed as OCI “v2” format (incompatible to current OCI Image Spec)

# eStargz

*Kohei Tokunaga, NTT Corporation*

# eStargz: Standard-compatible lazy pulling



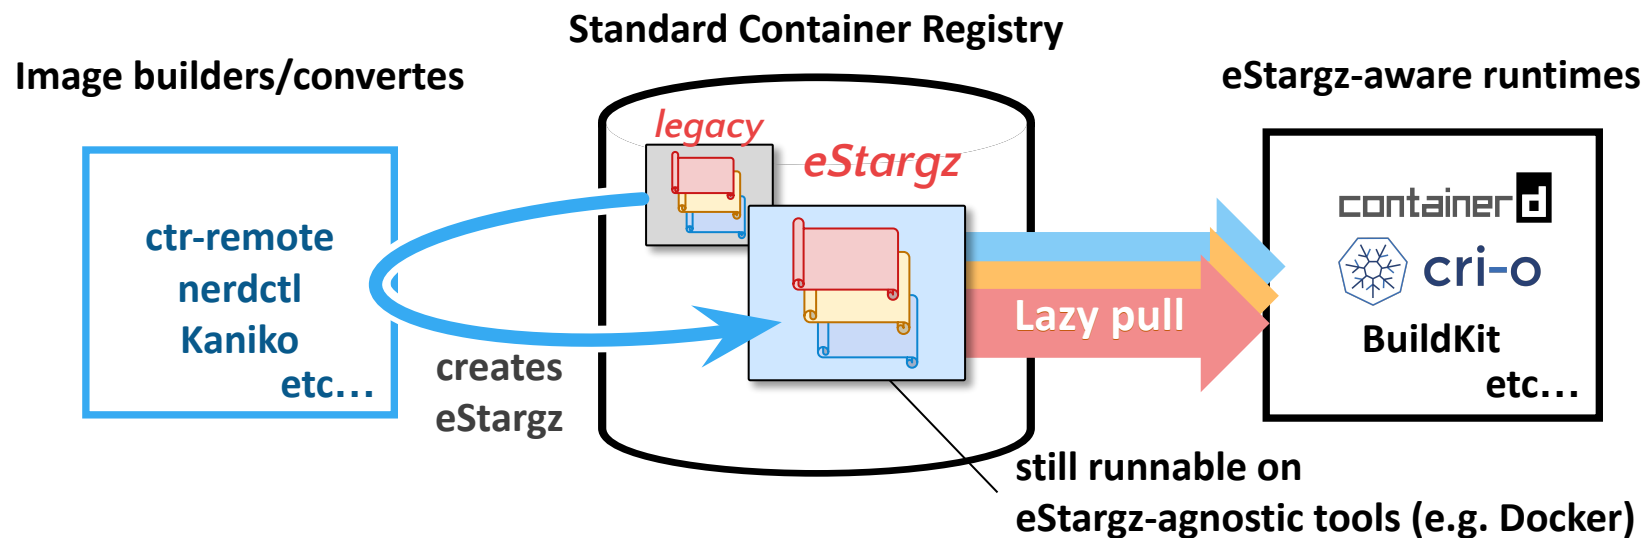
KubeCon



CloudNativeCon

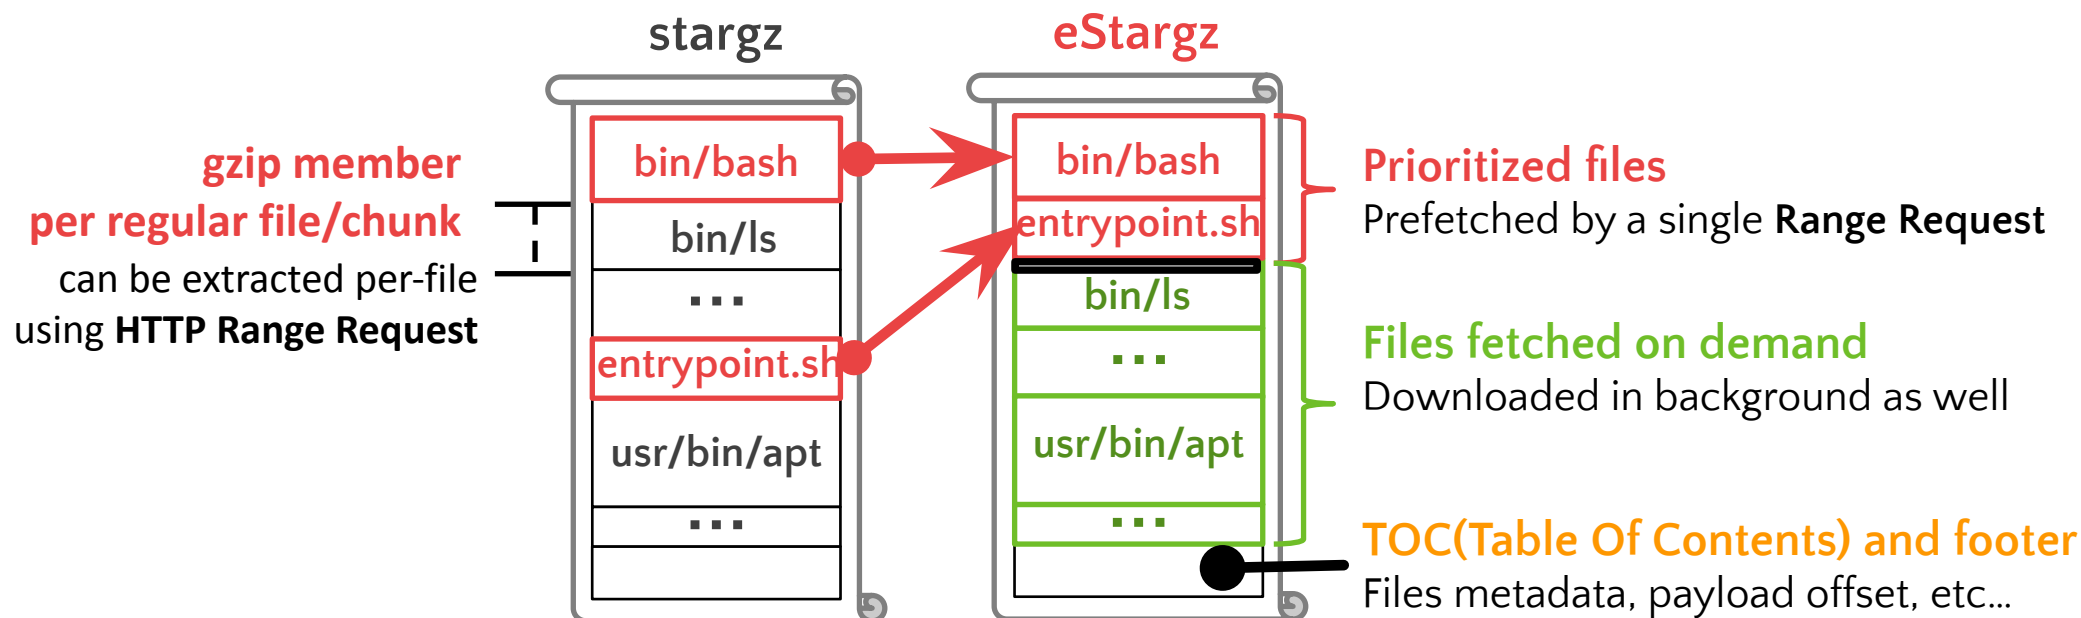
North America 2021

- **100% OCI-compatible**
  - Lazy pullable from standard registries (ghcr.io, docker.io, ...etc)
  - Even legacy (lazy-pulling-agnostic) runtime can run eStargz
- **Usable with a variety of tools** (tracker <https://github.com/containerd/stargz-snapshotter/issues/258>)
  - Kubernetes, k3s, containerd, CRI-O, Podman, BuildKit, Kaniko, ko, buildpacks.io, go-containerregistry...
- **Performance optimization and content verification**
  - Important files can be prefetched to avoid NW overhead
  - Each chunk comes with checksum



# eStargz image layer format

- Discussed in Stargz Snapshotter of containerd: <https://github.com/containerd/stargz-snapshotter>
- Compatible to gzip = usable as a valid OCI/Docker image layer
- Based on stargz by CRFS (<https://github.com/google/crfs>)
  - eStargz comes with performance optimization and content verification



For more details: <https://github.com/containerd/stargz-snapshotter/blob/master/docs/stargz-estargz.md>

# Benchmarking result




KubeCon



CloudNativeCon

North America 2021



Waits for prefetch completion during create

- Method based on Hello Bench [Harter, et al. 2016]
- Takes 95 percentile of 100 operations
- Host: EC2 Oregon (m5.2xlarge, Ubuntu 20.04)
- Registry: GitHub Container Registry (ghcr.io)
- Runtime: containerd
- Stargz Snapshotter commit: 7f45f7438617728dd06bc9853afb5e42c1d3d5a3



# Benchmarking result




KubeCon



CloudNativeCon

North America 2021



- Benchmarking spec is the same as the previous slide

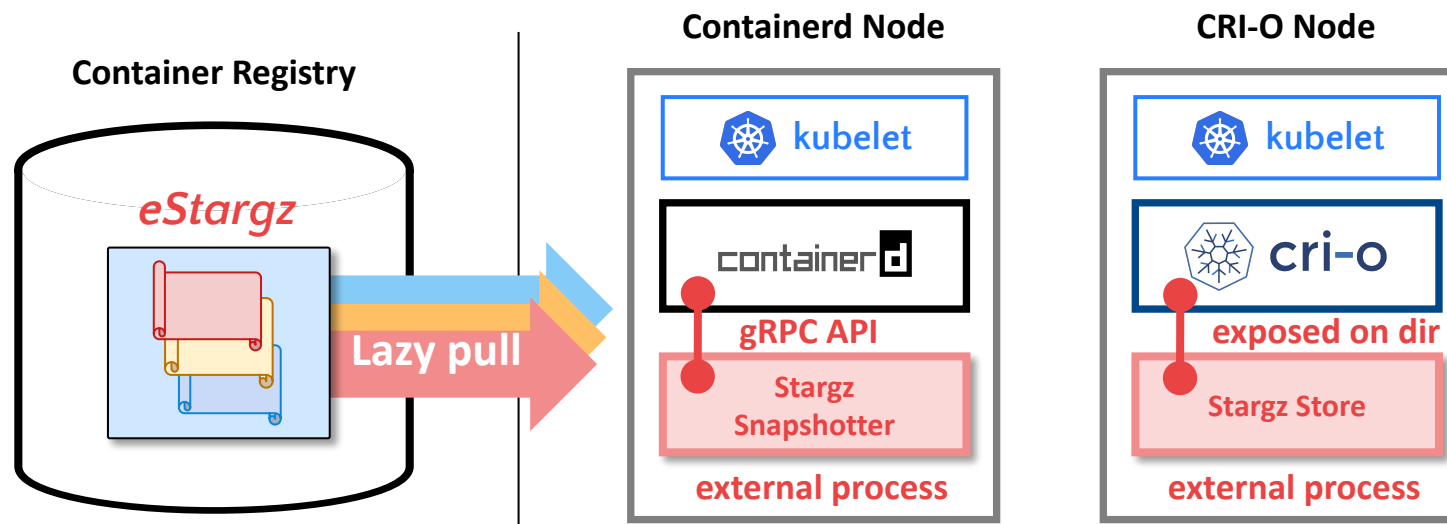
# Lazy pulling on Kubernetes

- CRI runtimes + plugins (discussed later) enable lazy pulling on Kubernetes
  - Containerd + Stargz Snapshotter
  - CRI-O + Stargz Store
- Real-world use-case at CERN for speeding up analysis pipeline [1] (13x faster pull for 5GB image)  
[1] Ricardo Rocha & Spyridon Trigazis, CERN. "Speeding Up Analysis Pipelines with Remote Container Images". KubeCon+CloudNativeCon 2020 NA. <https://sched.co/ekDi>
- k3s supports lazy pulling of eStargz (merged to the main, will be included in k3s v1.22)

```
$ k3s server --snapshotter=stargz
```

- Lazy-pull-enabled KinD image is available on `ghcr.io/stargz-containers` repo

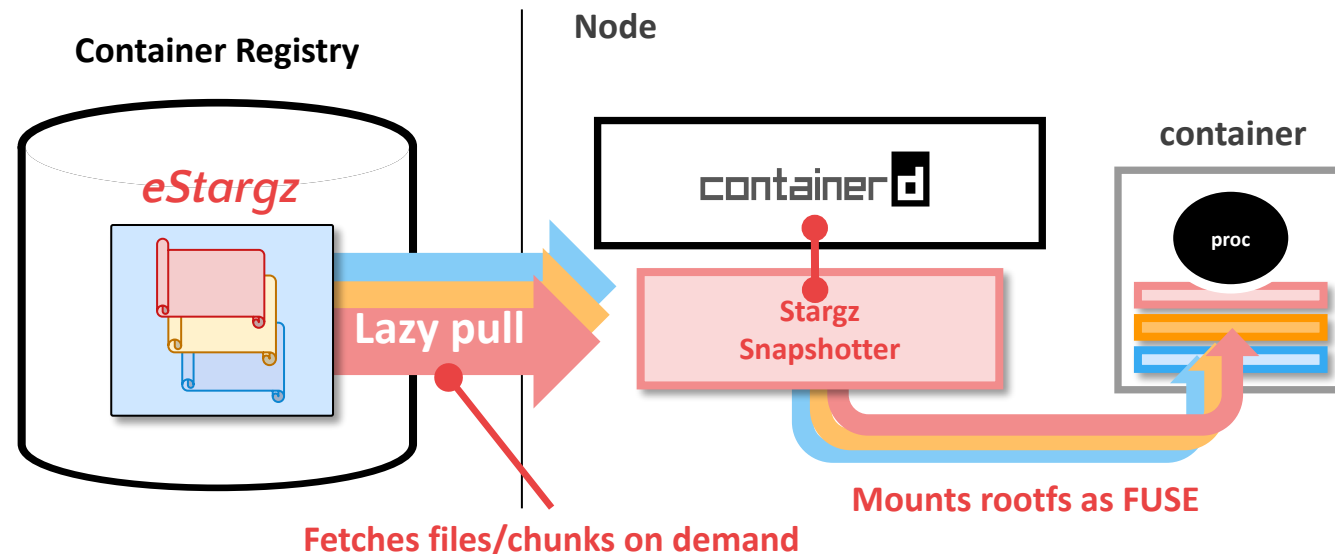
```
$ kind create cluster --name estargz-demo --image ghcr.io/stargz-containers/estargz-kind-node:0.8.0
```



# Lazy pulling on containerd

- Stargz Snapshotter plugin : <https://github.com/containerd/stargz-snapshotter>
  - Implements **remote snapshotter** plugin interface
- **nerdctl** (Docker-compatible CLI for containerd) supports lazy pulling
  - <https://github.com/containerd/nerdctl>

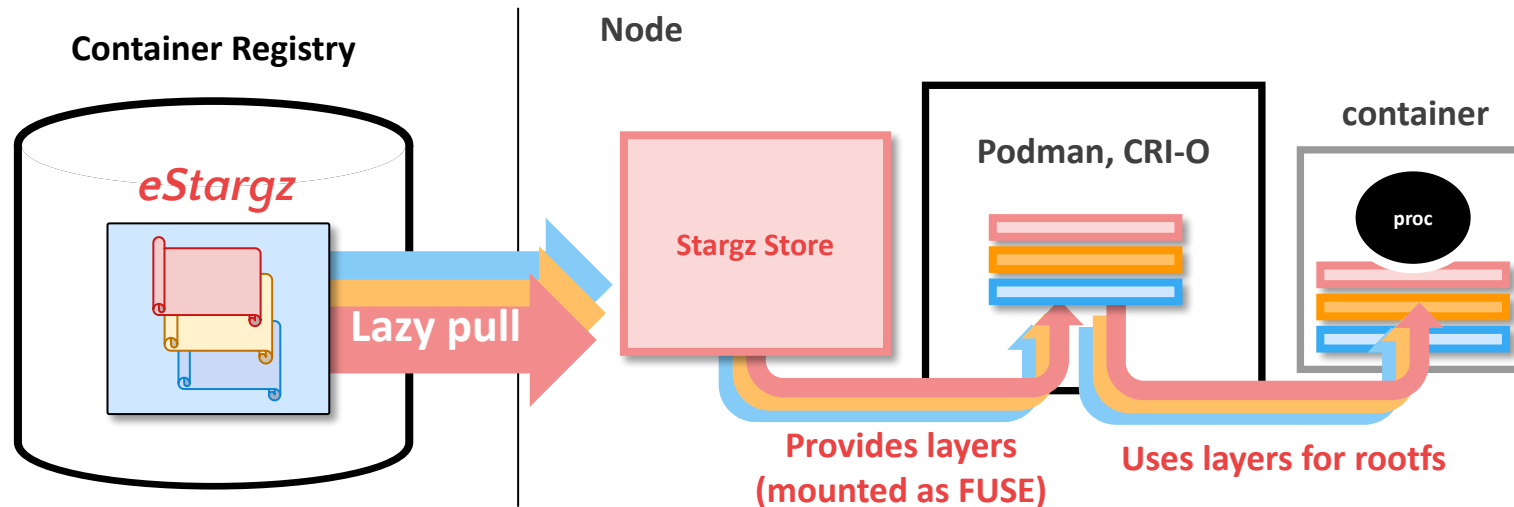
```
$ nerdctl --snapshotter=stargz run ghcr.io/stargz-containers/python:3.9-esgz  
$ nerdctl --snapshotter=stargz compose -f docker-compose.stargz.yaml up
```



# Lazy pulling on Podman/CRI-O

- Stargz Store plugin : <https://github.com/containerd/stargz-snapshotter>
  - Developed in Stargz Snapshotter project
  - Implements **Additional Layer Store** plugin
  - Podman  $\geq$  v3.3.0, CRI-O  $\geq$  v1.22.0

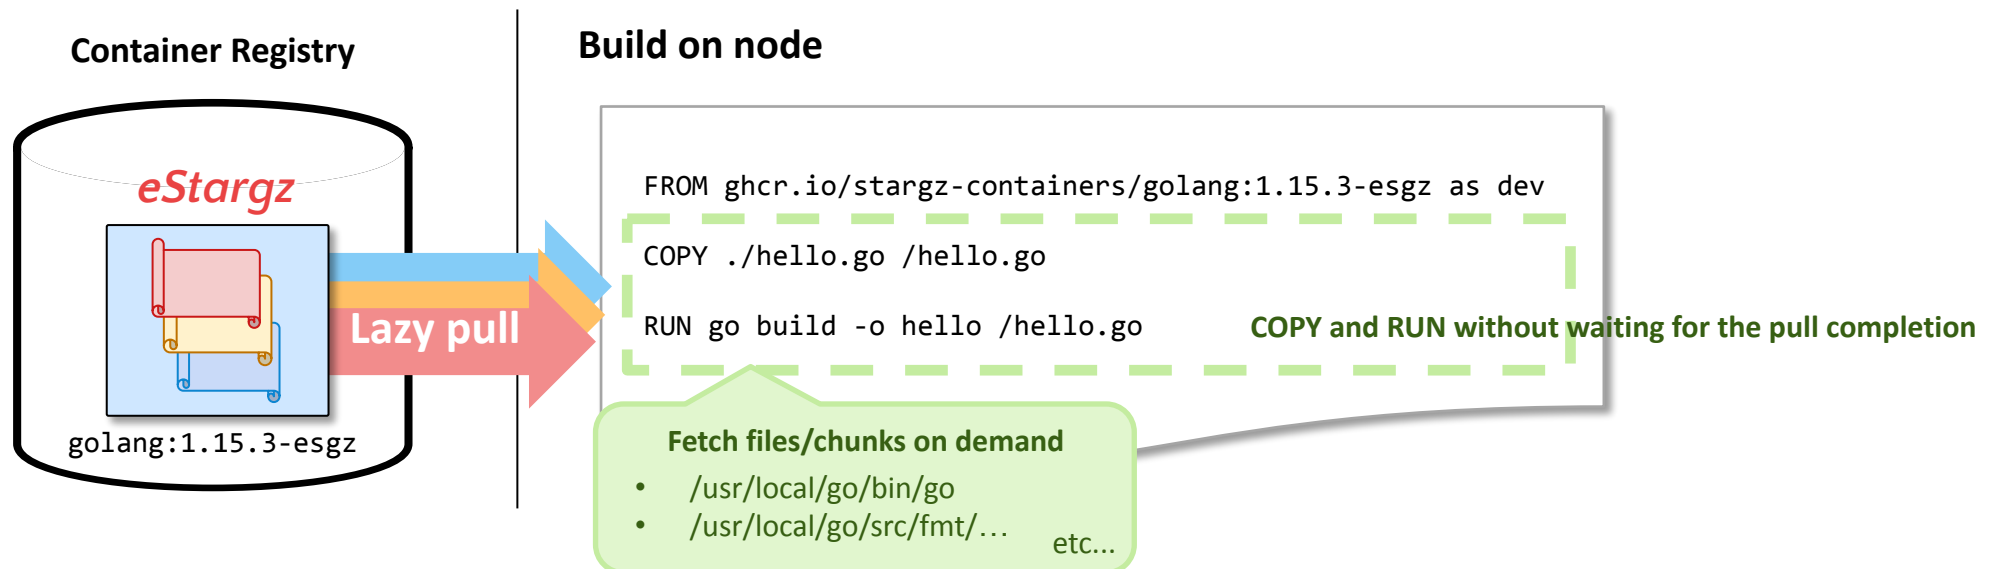
```
$ podman run ghcr.io/stargz-containers/python:3.9-esgz
```



# Lazy pulling on BuildKit

- Experimentally supports lazy pulling of base images since v0.8.0
- Can shorten the time of build that require pull
  - e.g. on temporary (and fresh) CI instances

```
$ docker buildx create --use --name lazy-builder \  
    --driver docker-container --driver-opt image=moby/buildkit:master \  
    --buildkitd-flags '--oci-worker-snapshotter=stargz'  
$ docker buildx inspect --bootstrap lazy-builder  
$ docker buildx build .
```



More details at blog: <https://medium.com/nttlabs/buildkit-lazypull-66c37690963f>

# Building eStargz

**BuildKit:** <https://github.com/moby/buildkit>

- Building eStargz supported on the main branch, hope to come in v0.10.x
- eStargz is supported by Buildx or standalone BuildKit (buildctl + buildkitd)
- Docs: <https://github.com/moby/buildkit/blob/master/docs/stargz-estargz.md#creating-stargzestargz-images>

```
$ docker buildx build \  
  -o type=registry,name=ktokunaga/hello:esgz,oci-mediatypes=true,compression=estargz .
```

**Kaniko:** <https://github.com/GoogleContainerTools/kaniko>

- Image builder runnable in containers and Kubernetes
- Requires GGCR\_EXPERIMENT\_ESTARGZ=1
- Base images need to be eStargz

```
$ docker run --rm -e GGCR_EXPERIMENT_ESTARGZ=1 \  
  -v /tmp/context:/workspace -v ~/.docker/config.json:/kaniko/.docker/config.json:ro \  
  gcr.io/kaniko-project/executor:v1.6.0 --destination "ghcr.io/ktock/sample:esgz"
```

# Converting an image into eStargz

**ctr-remote:** <https://github.com/containerd/stargz-snapshotter/blob/v0.7.0/docs/ctr-remote.md>

- CLI for containerd provided by Stargz Snapshotter project
- Supports prefetch optimization for eStargz

```
$ ctr-remote image optimize --oci ghcr.io/ktock/foo:1 ghcr.io/ktock/foo:1-esgz
```

**nerdctl:** <https://github.com/containerd/nerdctl/blob/v0.11.1/docs/stargz.md>

- Docker-compatible CLI for containerd
- Can be combined with `nerdctl build` command

```
$ nerdctl image convert --estargz --oci ghcr.io/ktock/foo:1 ghcr.io/ktock/foo:1-esgz
```

**go-containerregistry and crane CLI :** <https://github.com/google/go-containerregistry>

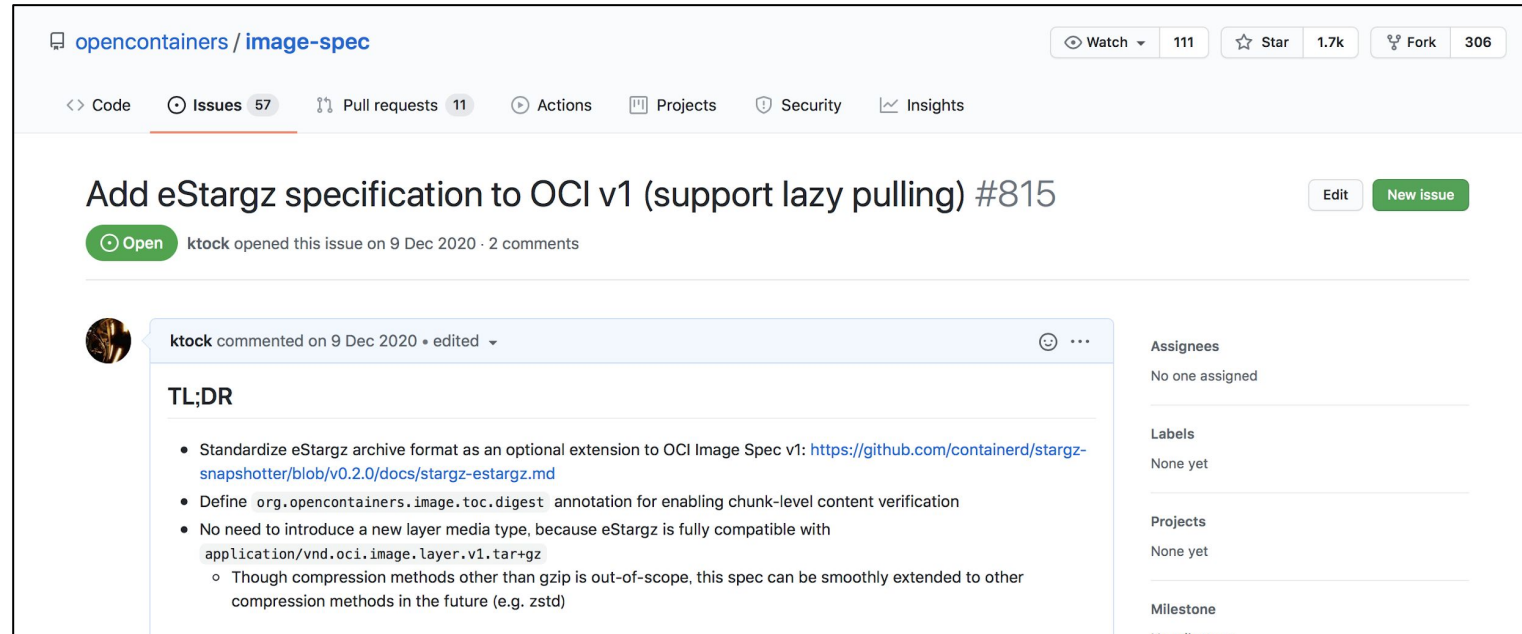
- Library and CLI to interact with registries
- requires `GGCR_EXPERIMENT_ESTARGZ=1`
- downstream tools (e.g. Kaniko, ko and buildpacks.io) supports eStargz creation as well

```
$ go install github.com/google/go-containerregistry/cmd/crane@latest
```

```
$ GGCR_EXPERIMENT_ESTARGZ=1 crane optimize ghcr.io/ktock/foo:1 ghcr.io/ktock/foo:1-esgz
```

# Discussion toward OClv1 extension

<https://github.com/opencontainers/image-spec/issues/815>



- eStargz is now widely usable on tools in the community
- Proposed to OCI Image Spec (discussion is on-going)
- Proposed as backward-compatible extensions
  - Extension for the current gzip layer ( `application/vnd.oci.image.layer.v1.tar+gz` )
  - Additional annotation for content verification: `org.opencontainers.image.toc.digest`





# Nydus Image Service

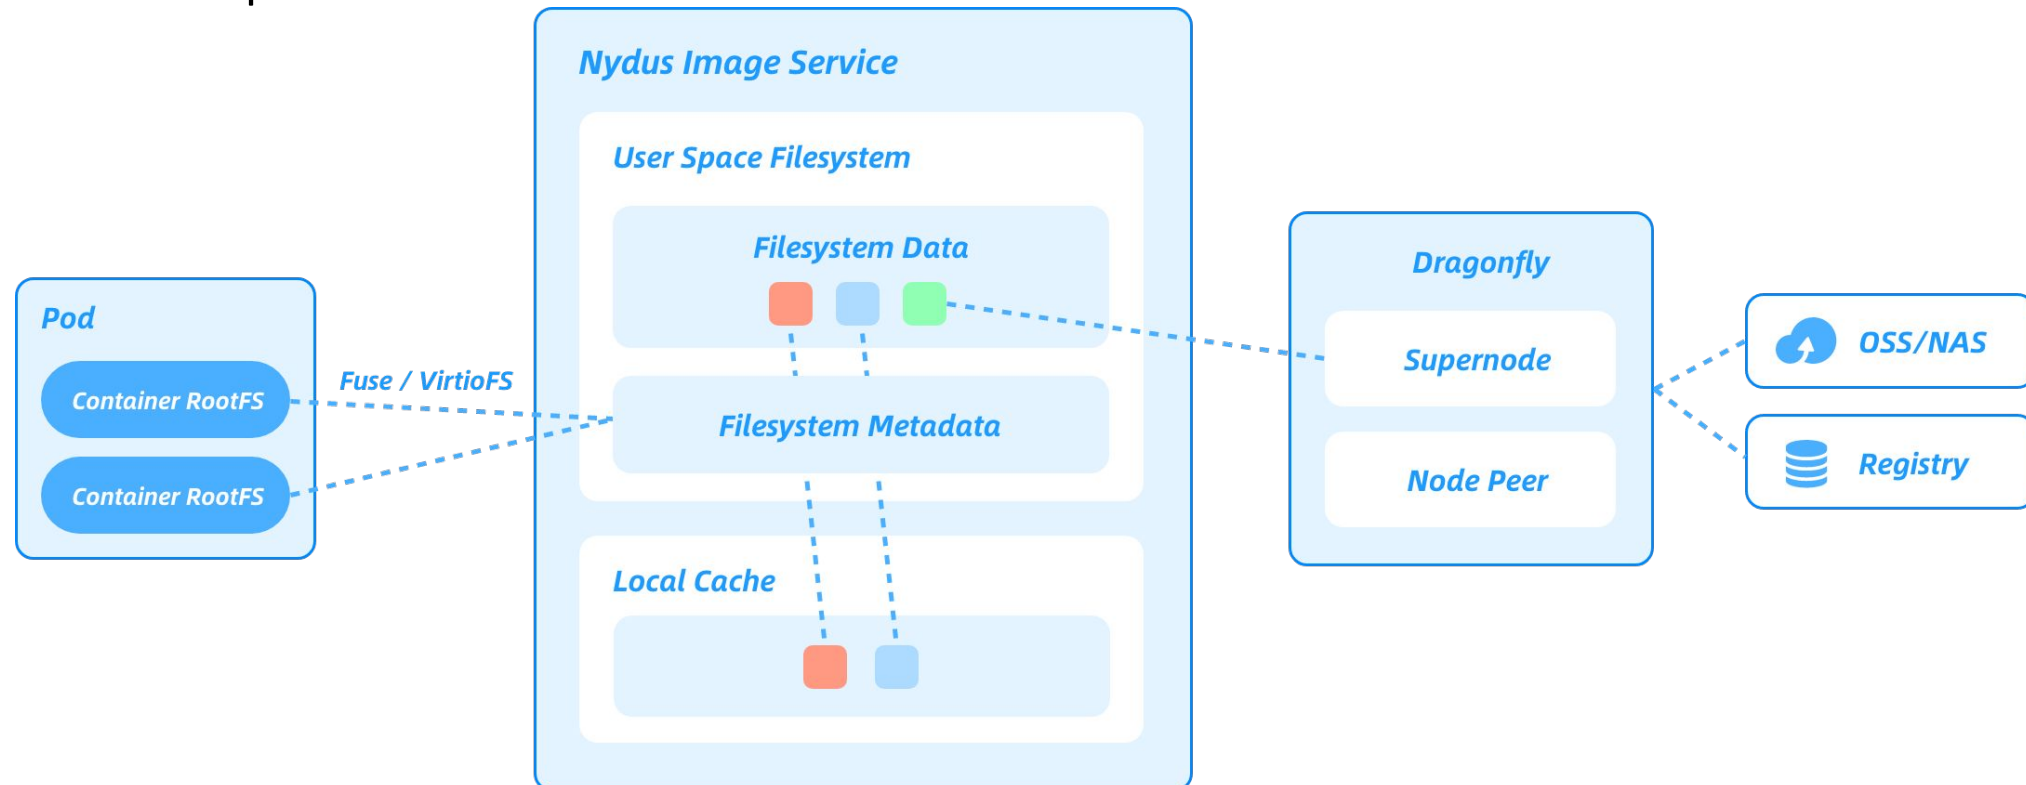
*Tao Peng, Ant Group*

# Nydus: Improved Lazy Pulling and More

- **Standard lazy pulling with prefetch policies**
  - Lazy pullable from standard registries (cloud registries, docker registry, harbor, etc)
  - Semantical prefetch, hinted prefetch, background prefetch all data
- **Chunk level deduplication with layered build cache**
  - Speed up both image conversion and downloading
- **End-to-end data integrity**
  - Runtime data integrity check
- **Reproducible Image Building**
  - Build environment agnostic
- **Compatible with OCI Distribution Spec and Artifacts Spec**
  - Usable with most existing container registry implementations
- **Rich container ecosystem integration**
  - Kubernetes, docker, containerd, buildkit, harbor, dragonfly, runc, kata-containers etc
- **Resource efficient and production ready**
  - Large scale deployment at Ant and available via Alibaba Cloud services

# FUSE/virtiofs Unified Architecture

- **Native container runtime support**
  - One translation layer (FUSE or virtiofs) for both runc and kata containers
- **Shared uncompressed local cache**
  - Download/uncompress once and use all the time
- **Singleton mode with extremely low memory footprint**
  - ~5MB per instance



# Registry Acceleration File System



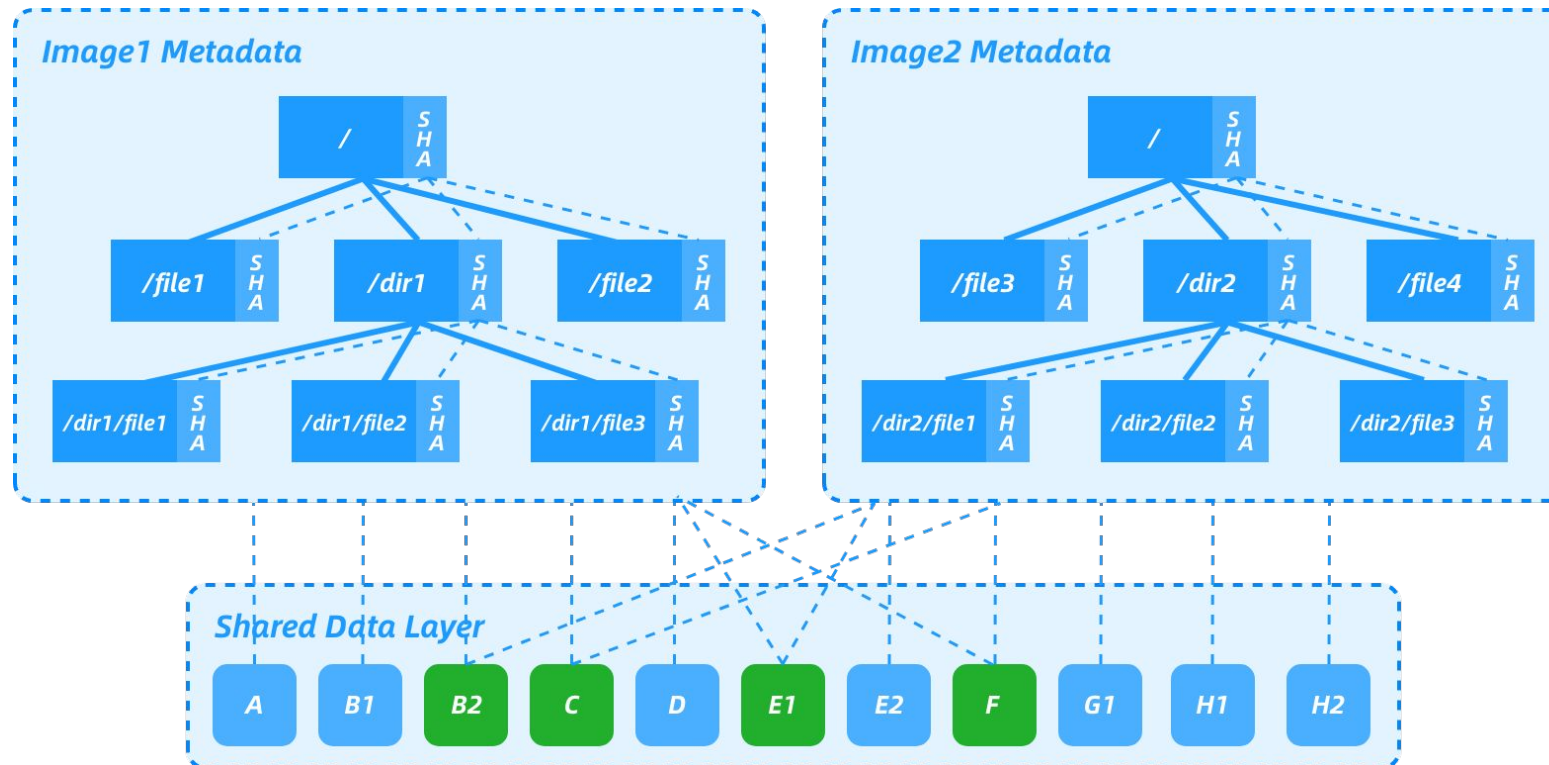
KubeCon



CloudNativeCon

North America 2021

- **Merkle tree metadata layer**
  - self-verifiable
- **Chunk-shared data layer**
  - enabling chunk level deduplication



# OCI Spec Compatibility



KubeCon



CloudNativeCon

North America 2021

- Metadata/data layers are OCI manifest layers
- New media type in the image manifest spec
- Compatible with OCI artifacts/distribution spec
- Compatible registry GC functionality
- Widely working with container registries
  - cloud registries
  - docker registry
  - harbor
  - etc.

```
1 {
2   "schemaVersion": 2,
3   "mediaType": "",
4   "config": {
5     "mediaType": "application/vnd.docker.container.image.v1+json",
6     "size": 981,
7     "digest": "sha256:a27f27be5546ba699ec38344a3fcbeb92ccfe7bdf0ac13d62ce630dea0178bbd"
8   },
9   "layers": [
10    {
11      "mediaType": "application/vnd.oci.image.layer.nydus.blob.v1",
12      "size": 51522,
13      "digest": "sha256:8a44bc8c2e35502f68d1ad692f7bf247eb9e21dca2742b6b0df58ba7b6a96ef3",
14      "annotations": {
15        "containerd.io/snapshot/nydus-blob": "true"
16      }
17    },
18    {
19      "mediaType": "application/vnd.oci.image.layer.nydus.blob.v1",
20      "size": 524,
21      "digest": "sha256:1d51ac9ebde626252c1b02fc2d446a5e328eadcb1ca26942bfbdd482b5e386e49",
22      "annotations": {
23        "containerd.io/snapshot/nydus-blob": "true"
24      }
25    },
26    {
27      "mediaType": "application/vnd.oci.image.layer.v1.tar.gz",
28      "size": 664576,
29      "digest": "sha256:35bdd331b926eccd78440b0060c8484355ad69a8e6f38290fed4d0a3491ba76e",
30      "annotations": {
31        "containerd.io/snapshot/nydus-bootstrap": "true"
32      }
33    }
34  ]
35 }
```

# Lazy Pulling Benchmark




KubeCon



CloudNativeCon

North America 2021



# Lazy Pulling with Variety of Tools




KubeCon



CloudNativeCon

North America 2021

- Nydus Snapshotter plugin
  - <https://github.com/dragonflyoss/image-service/tree/master/contrib/nydus-snapshotter>
- Usable for variety of container management tools
  - kubernetes (kubelet)
  - crictl
  - nerdctl
  - ctr-remote





# Manual Image Conversion -- nydusify

- nydusify
  - <https://github.com/dragonflyoss/image-service/tree/master/contrib/nydusify>
  - download image from remote registry, convert it to nydus image, and push nydus image to remote registry
  - nydusify convert --source <registry>/<repo>:<tag> \ --target <registry>/<repo>:<nydus-tag>

```
[@~]$sudo ./nydusify convert --source docker.io/bergwolf/ubuntu:19.10 --target docker.io/bergwolf/ubuntu:19.10-nydus
INFO[2021-08-25T06:28:59Z] Version: 648f538ab5a32ad501aef6007a34b0ebfc2aec45.20210825.0309
INFO[2021-08-25T06:28:59Z] Parsing image docker.io/bergwolf/ubuntu:19.10
INFO[2021-08-25T06:29:01Z] Converting to docker.io/bergwolf/ubuntu:19.10-nydus
INFO[2021-08-25T06:29:01Z] [SOUR] Mount layer Digest="sha256:3f2411103a12c8e169df7a9ea00ff26ab07501858e3eff315a2e11c219e78ce1" Size="28 MB"
INFO[2021-08-25T06:29:01Z] [SOUR] Mount layer Digest="sha256:354c6da61dcc176c5b363ded571bea1de41b718131658fa0e563f2a749028cd1" Size="164 B"
INFO[2021-08-25T06:29:02Z] [SOUR] Mount layer Digest="sha256:354c6da61dcc176c5b363ded571bea1de41b718131658fa0e563f2a749028cd1" Size="164 B" Time=907.918066ms
<snip...>
INFO[2021-08-25T06:29:09Z] [BLOB] Push blob Digest="sha256:00d151e7d392e68e2c756a6fc42640006ddc0a98d37dba3f90a7b73f63188bbd" Size="7 B" Time=1.227762628s
INFO[2021-08-25T06:29:09Z] [BLOB] Push blob Digest="sha256:471e21f4587c06f589e8ba2aca83806dbc219994ea422048f1e13f8364485fee" Size="40 MB" Time=1.455197822s
INFO[2021-08-25T06:29:09Z] [MANI] Push manifest
INFO[2021-08-25T06:29:13Z] [MANI] Push manifest Time=3.430874134s
INFO[2021-08-25T06:29:13Z] Converted to docker.io/bergwolf/ubuntu:19.10-nydus
```



# Manual Image Conversion -- buildkit



KubeCon



CloudNativeCon

North America 2021

- buildkit
  - <https://github.com/moby/buildkit/pull/2045>
  - build nydus image directly from dockerfile and push it to remote registry
  - buildctl build --frontend=dockerfile.v0 \  
--local context=/build/dir \  
--local dockerfile=/build/DockerfileDir \  
--output type=nydus,name=<registry>/<repo>:<tag>

```
+ buildctl build --frontend=dockerfile.v0 --local context=/tmp/tmp.x0ro44qGoh --local dockerfile=/tmp/tmp.x0ro44qGoh --output type=nydus,name=localhost:5001/ubuntu:nydus,merge-manifest=true,oci-mediatypes=true
[+] Building 1.7s (5/5) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 64B
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for localhost:5001/ubuntu:latest
=> CACHED [1/1] FROM localhost:5001/ubuntu@sha256:1e48201ccc2ab83afc435394b3bf70af0fa0055215c1e26a5da9b50a1ae367c9
=> => resolve localhost:5001/ubuntu@sha256:1e48201ccc2ab83afc435394b3bf70af0fa0055215c1e26a5da9b50a1ae367c9
=> Exporting Nydus image to localhost:5001/ubuntu:nydus
=> => exporting layers
=> => exporting manifest sha256:5f482d4f5541c919f5a48fcb4013df91fa61339a2d713a04766fe3db9ed5569b
=> => exporting config sha256:f5e02ea100de451b267775e5c9517d7ac6e483f511a5562111184bacf75bb450
=> => [SOUR] Mount layer [Digest=sha256:16ec32c2132b43494832a05f2b02f7a822479f8250c173d0ab27b3de78b2f058 Size=104 MB]
=> => [DUMP] Build layer [Digest=sha256:16ec32c2132b43494832a05f2b02f7a822479f8250c173d0ab27b3de78b2f058 Size=104 MB]
=> => [BOOT] Push bootstrap [Digest=sha256:16ec32c2132b43494832a05f2b02f7a822479f8250c173d0ab27b3de78b2f058 Size=688 kB]
=> => [BLOB] Push blob [Digest=sha256:966586db563798466633b14645f6343372b29c76f332f4ab5aa0544e170f49d9 Size=41 MB]
=> => [MANI] Push manifest
```

# Automatic Conversion with Harbor



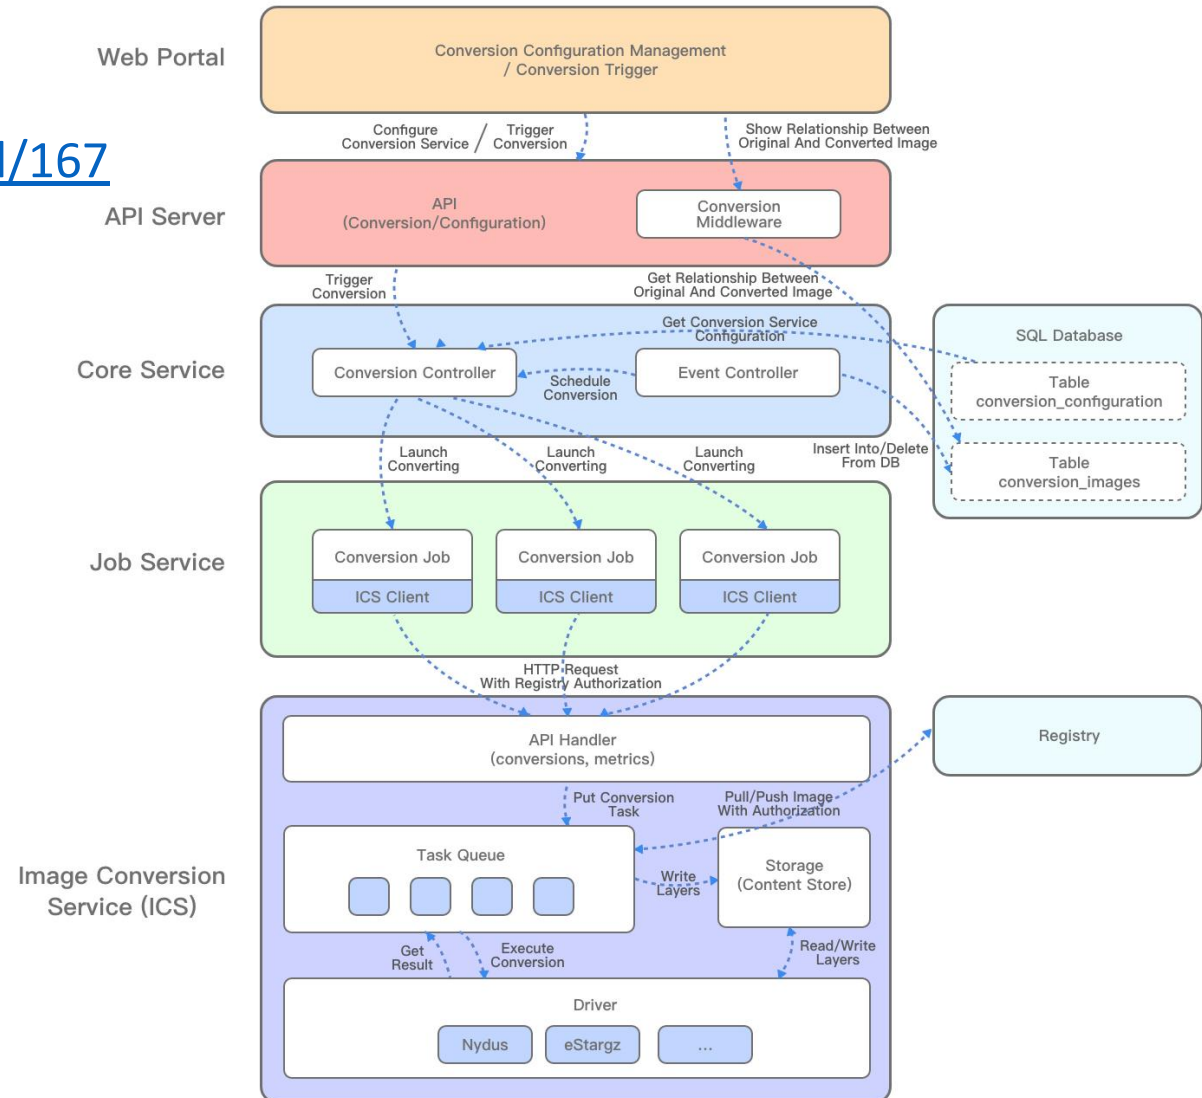
KubeCon



CloudNativeCon

North America 2021

- <https://github.com/goharbor/community/pull/167>
- New image conversion service in harbor
- Started image acceleration working group
- Image format agnostic
  - nydus
  - estargz
  - future more format



# OCI Artifacts Manifest



KubeCon



CloudNativeCon

North America 2021

- <https://github.com/opencontainers/artifacts/pull/29>
- Map OCI image manifest to artifact manifest, to connect artifacts to images
- Beneficial to SBOM, Signatures, Nydus, Scan Results etc.
- Nydus image as an artifact type
  - cncf.nydus.v1-rc1

```
1 {
2   "schemaVersion": 3,
3   "mediaType": "application/vnd.oci.artifact.manifest.v1+json",
4   "artifactType": "cncf.nydus.v1-rc1",
5   "blobs": [
6     {
7       "mediaType": "application/vnd.cncf.nydus.bootstrap.v1.tar+gzip",
8       "digest": "sha256:f6bb0822fe567c98959bb87aa316a565eb1ae059c46fa8bba65b573b4489b44d",
9       "size": 32654
10    },
11    {
12      "mediaType": "application/vnd.cncf.nydus.blob.v1",
13      "digest": "sha256:d35808e58856ef91d07dedf94b93301b6efdccc831477d3b1bb37e6c3e19cce6",
14      "size": 25851449
15    },
16    {
17      "mediaType": "application/vnd.cncf.nydus.blob.v1",
18      "digest": "sha256:dbad66bcfe29ef383157a3e122acbd08cd2ebd40f5658afa2ae62c52ffe26e9f",
19      "size": 226
20    }
21  ],
22  "subjectManifest": {
23    "mediaType": "application/vnd.oci.image.manifest.v1+json",
24    "digest": "sha256:73c803930ea3ba1e54bc25c2bdc53edd0284c62ed651fe7b00369da519a3c333",
25    "size": 16724
26  },
27  "annotations": {
28    "org.cncf.nydus.v1.author": "wabbit-networks.io"
29  }
30 }
```

# Community V2 Image Spec Requirements

- OCI community requirements on V2 Image format
  - <https://hackmd.io/@cyphar/ociv2-brainstorm>

Solution	Reduced Duplication	Reproducible Image Building	Explicit (and Minimal) Filesystem Objects and Metadata	Runtime Data Integrity	Mountable Filesystem Format	Lazy fetch support
OCI Image	Layer	No	No	No	No	No
Dragonfly	Layer	No	No	No	No	No
Kraken	Layer	No	No	No	No	No
CernVM-fs	Chunk	No	No	Partial[1]	Yes	Yes
slacker	Layer	No	No	No	Yes	Yes
eStargz	Layer	Possible	Possible	Possible	Yes	Yes
filegrain	File	Possible	Possible	No	Yes	Yes
umoci	Chunk	Possible	Possible	No	Yes	Yes
nydus	Chunk	Yes	Yes	Yes	Yes	Yes

[1] CernVM-FS incoming data is validated but only on download

# Recap



KubeCon



CloudNativeCon

North America 2021

- OCI images are large and slow
- Image lazy pulling is important for starting containers quickly
- Prefetching is a good friend to lazy pulling
- Ecosystem adoption is crucial for new image formats
- eStargz
  - Backward compatibility
  - Extension to the existing OCI image spec
- Nydus
  - Future looking
  - Proposal to the next generation OCI image spec

# eStargz Project Information



- 100% OCI-compatible image format for lazy pulling
- Subproject of CNCF graduated containerd
- github: <https://github.com/containerd/stargz-snapshotter>
- slack: <https://slack.containerd.io/> (#containerd-dev channel at CNCF slack)
- Pre-converted images are available on [ghcr.io/stargz-containers](https://ghcr.io/stargz-containers)

# Nydus Project Information



- Image acceleration service with various improvements
- CNCF incubator dragonfly sub-project
- github: <https://github.com/dragonflyoss/image-service>
- slack: <https://tinyurl.com/nydus-slack>
- tutorial: <https://tinyurl.com/nydus-tutorial>

**Thank You  
&  
Questions!**