



KubeCon



CloudNativeCon

North America 2021

RESILIENCE

REALIZED

How NOT to start with K8s

Christian Heckelmann

2018 - How everything started...



KubeCon



CloudNativeCon

North America 2021

- We want to use containers!
- PoC with K8s and Docker Swarm
- Needs to run on-prem (CentOS on VMWare and Xen)
- Kubernetes 1.9
- Don't start by your own – You are building a little Datacenter within a Datacenter!



KubeCon



CloudNativeCon

North America 2021

RESILIENCE

REALIZED

Infrastructure and Operations

Do not install K8s from scratch



KubeCon




CloudNativeCon

North America 2021

- Failed hard with my first cluster while upgrading
- Switched to Rancher

Other options (for on-prem)

- VMWare Tanzu
- Kubespray



Automation everywhere



KubeCon



CloudNativeCon

North America 2021

- Try to automate provisioning of clusters
- Provision a cluster with a UI is easy but what if you need to deploy 10 - 20 - 100 Clusters?
- You don't want to run „kubectl apply -f“ X-Times to get your cluster ready



Network CNI and Configuration



KubeCon



CloudNativeCon

North America 2021



- The right CIDR Range for pod and service network
- Network Policies
- Encrypt data in transit
- Know at least the basics (like DNS in K8s)

<https://kubernetes.io/docs/concepts/cluster-administration/networking/>

How we could reach our service?




KubeCon



CloudNativeCon

North America 2021

- Loadbalancer (F5, MetalLB)
- Which protocols (https/tcp)?
- Do you want to split your traffic (internal/external)?
- Security (Web Application Firewall)
- SSL Certificates
 - Certmanager with LetsEncrypt



We need a way to store some files...



KubeCon



CloudNativeCon

North America 2021

- Keep it simple!
- Backup
- What kind of workload do you expect?
- What is your Infrastructure already able to provide? – Ask the experts!

<https://kubernetes.io/docs/concepts/storage/storage-classes/>

<https://medium.com/volterra-io/kubernetes-storage-performance-comparison-v2-2020-updated-1c0b69f0dcf4>



- Use Service Accounts from the beginning
- Use separate accounts for deployment/monitoring and operation tasks
- Don't share credentials across teams
- Don't use "kubeconfig" files in your Pipelines

Logging and Monitoring



Infrastructure

Cluster services like etcd, API-Server, DNS

Kubernetes events like CrashLoopBackOff, ImagePullBackoff, OOMKilled

Requested and used resources

Logs (where should I write my application log?)

All ingresses are down!?



KubeCon



CloudNativeCon

North America 2021

Check for common mistakes and reject changes in K8s before it's getting worse like with Policies:

- Ingress validation
- Pod Security (e.g. Disallow Privileged Containers)
- Required Resource Limits/Requests
- Health checks



kubectl1

Operators (and Developers)
should be familiar with the
common kubectl commands.



Kube Control



Kube Cuddle



KubeCon



CloudNativeCon

North America 2021

RESILIENCE

REALIZED

Development & Deployment

Wrong motivation



KubeCon



CloudNativeCon

North America 2021

“I run it in K8s cause I don’t want to
request a VM”

Local Development



KubeCon



CloudNativeCon

North America 2021

Dev: Kubernetes is down, my deployment is not working!!!!

Ops: Ok your Container is not working at all

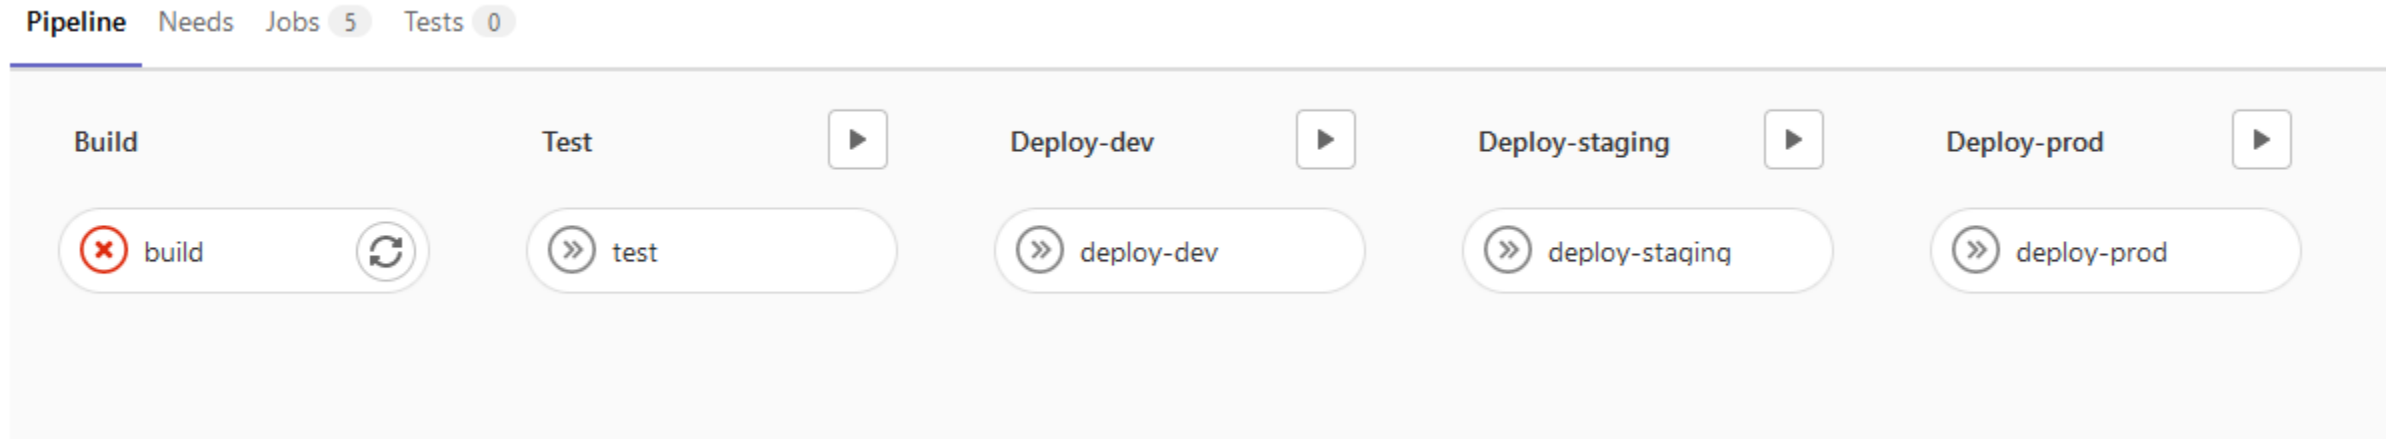
- Install Docker on your local machine and build the image
- Use KinD, MiniKube, K3s to test your application and deployment



:latest - why isn't it working anymore?

Never ever use the latest tag either in your CI/CD Pipeline or in your Kubernetes deployments.

If a new version on an image is getting pulled, it could break the container or workload.



Private Registry & Base images



KubeCon




CloudNativeCon

North America 2021

Create a library of own base images (private registry) for the tools you are using, to have full control

- Check the images for vulnerabilities - Harbor
- Add additional configurations to the image (e.g. Java8 cgroup memory bug)
- No need to pull from docker hub and hit the rate limit (and don't use ImagePullPolicy: Always)
- Block/Allow lists if possible



```
docker pull untrusted/withacryptominer:latest
```

Not using the „power“ of Kubernetes



KubeCon



CloudNativeCon

North America 2021

“Why Autoscaling? I have two Pods!”

Health checks

Make sure your application is configured with a proper health check.

If you get ever asked to “restart” a container within a cluster to mitigate an issue, you must ask them why this is necessary.

Use a policy engine like OPA or Kyverno to block those deployments

[+IT-Ops](#) - Please restart the ECG Algorithm service in Kubernetes.

Thanks,

Resource Limits and Requests



KubeCon



CloudNativeCon

North America 2021

- Kubernetes decides based on your resource requests where to run the pod in your cluster
- If limits are not defined, the pod can “explode” and affect another workload on the worker node
- Add default memory requests and limits on your cluster or namespaces
- Prevent requests like CPU: 12 – memory 128GB
- Dev: But I don’t know how many resources my container is consuming? → Local Development, `kubectl top pod`


utilization

requests


limits

available

requests



memory requests



Env Variables for Config and Secrets?!



KubeCon



CloudNativeCon

North America 2021

- Do not put Secrets in Environment Variables
- Create a standard how your application configuration should be managed
- Use a secret store like vault or sealed secrets
- K8s secrets are not encrypted!
- Use configmaps for configuration



- Over 90% of deployments are the same
- Provide Templates which can be adopted easily
- Preventing common mistakes “forgotten health checks, accidentally exposed services...”

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: api-service
spec:
  selector:
    matchLabels:
      app.kubernetes.io/name: api-service
      app.kubernetes.io/instance: {{ .Release.Name }}
      app.kubernetes.io/managed-by: {{ .Release.Service }}
      app.kubernetes.io/part-of: keptn-{{ .Release.Namespace }}
      app.kubernetes.io/component: {{ include "control-plane.name" . }}
      app.kubernetes.io/version: {{ .Values.apiService.image.tag | default "latest" }}
  replicas: {{ .Values.apiService.replicas }}
  template:
    metadata:
      labels:
        app.kubernetes.io/name: api-service
        app.kubernetes.io/instance: {{ .Release.Name }}
        app.kubernetes.io/managed-by: {{ .Release.Service }}
        app.kubernetes.io/part-of: keptn-{{ .Release.Namespace }}
        app.kubernetes.io/component: {{ include "control-plane.name" . }}
```


Summary




KubeCon



CloudNativeCon

North America 2021

- Invest in Training
- Document how to use the platform
- Provide templates to avoid common mistakes
- Reject everything which does not meet the required standards
- Involve SecOps from the beginning

A person is seen from behind, sitting at a desk and looking at a computer monitor. The monitor displays lines of code in a dark-themed editor. The person is wearing a dark t-shirt. The background is dark, suggesting an office or home workspace at night.

**the “H” in
kubernetes
stands for
"happiness"**

Ping me!



KubeCon



CloudNativeCon

North America 2021



<https://www.linkedin.com/in/christian-heckelmann-82375a25/>



@wurstsalat



heckelmann@gmail.com

