

Testing Kubernetes Clusters

Building Confidence in Your Changes



@gjtempleton @ruio @SkyscannerEng

Agenda

1. **A little about us**
2. **The problem(s) we faced**
3. **Potential solutions**
4. **Our solution**
5. **Gotchas, Tips & Other Options**



Who We Are



Guy Templeton

Principal Software Engineer

- Twitter: **@gjtempleton**
- GitHub: **@gjtempleton**



Matteo Ruina

Senior Software Engineer

- Twitter: [@ruio](#)
- GitHub: [@maruina](#)



We're the travel company who puts you first

- Skyscanner helps millions of people in 52 countries and over 30 languages find the best travel options for flights, hotels and car hire every month.
- Skyscanner is available on desktop, mobile web and its highly rated app has 100 million downloads.
- Working with 1200 travel partners, Skyscanner's mission is to lead the global transformation to modern and sustainable travel.



K8s @ Skyscanner

- Running K8s since 1.6
- 35+ production clusters across 4 AWS regions
- 475+ services
- 40k+ CPU cores
- 150+ TB RAM



The Problem(s) We Faced



The Kubernetes platform is a complex set of distributed components running together

- Kubernetes has a good set of unit, integration and end-to-end (e2e) tests
- Every addons has its own set of unit tests
- The problem is how to test everything in **your** specific environment



Why testing

- Reduce disruption on clusters 😊
- Increase squad velocity
- Build confidence, for the squad and the users
- Give us a baseline where we can compare the cluster



Exploring Potential Solutions



Manual Acceptance Testing

- Doesn't scale as the number of clusters increases
- Easy to make mistakes
- Inconsistent, even if made into runbooks



Custom Service

- Constantly running tests - easier to catch a degradation
- Harder to block update pipelines on
- No (easy) ability to run ad-hoc tests



Building on top of existing frameworks

- Existing infrastructure and Kubernetes testing frameworks
- Would minimise the custom work we had to do
- Number of different options



kubetest

<https://github.com/kubernetes/test-infra/tree/master/kubetest>

- Allows for creation and teardown of test clusters
- Supports the addition of custom ginkgo tests
- Supported by the kubernetes project
- Comes with extra overhead
- Not necessarily representative of your production clusters



The Conformance Suite

"The standard set of conformance tests is currently those defined by the [Conformance] tag in the [kubernetes e2e](#) suite."

- Easy for us to use the community's existing tests
- Tests a huge range of functionality of clusters
- Very time consuming to run the entire suite
- Some disruptive tests if running the full suite



Sonobuoy

<https://sonobuoy.io/>

- Open source tool
- Active community and constantly improving project
- Ability to run custom tests via plug-ins



Our Solution



Requirements

- Easy to run from a local laptop and from a CI/CD pipeline
- Can run tests in parallel
- Should expose and store the test results
- Allow alerting if the tests are failing
- Write as little boilerplate as possible



Solution

- Custom Sonobuoy plugin
- Custom smoke tests leveraging
Kubernetes/Kubernetes test framework
- Not open source yet ☹️



Our setup

- A custom Sonobuoy image with a script to pass a return code and with meaningful output on the failed tests.

```
7  /sonobuoy run "$@"
8
9  RESULTS_FILE="$(/sonobuoy retrieve)"
10 # Output the report in a human readable way for debugging purposes
11 /sonobuoy results "$RESULTS_FILE"
12
13 # Use detailed mode for more detailed analysis
14 FAILED_TESTS=$(/sonobuoy results "$RESULTS_FILE" --mode=detailed | jq 'select (.status=="failed").name')
15
16 if [ -z "$FAILED_TESTS" ]; then
17     echo "All tests were successful!"
18     exit 0
19 else
20     printf "Some tests failed:\n"
21     echo "$FAILED_TESTS"
22     exit 1
23 fi
```

Our setup

- The e2e docker image based off the upstream Kubernetes conformance image.

```
1  # Build stage
2  FROM golang:1.16 AS build-stage
3  WORKDIR /go/src/github.skyscannertools.net/k8s/test-infra
4  COPY . .
5  RUN make build-e2e
6
7  # Create the tests image
8  FROM k8s.gcr.io/conformance:v1.19.8
9  ENV UNREADY_NODES=""
10 COPY --from=build-stage /go/src/github.skyscannertools.net/k8s/test-infra/scripts/run.sh /run.sh
11 RUN chmod +x /run.sh
12 COPY --from=build-stage /go/src/github.skyscannertools.net/k8s/test-infra/bin/e2e.test /usr/local/bin/
13 COPY --from=build-stage /go/src/github.skyscannertools.net/k8s/test-infra/e2e-repo-spec.yaml /usr/local/
```



Our setup

k8s / test-infra Internal

Configure Sourcegraph

Watch ▼ 3

Star 1

Fork 0

Code

Pull requests

Insights

Settings

master ▼

test-infra / e2e /

Go to file

Add file ▼

williamhutcheson ARGO-4361 Add test for custom metric scaling ...

54c2683 2 days ago History

..

common.go

Address PR Feedback

5 months ago

e2e.go

Fix functions

6 months ago

e2e_test.go

Fix parsing flag logic

6 months ago

networking.go

ARGO-4194 - use GET as proxy method

2 months ago

smoke.go

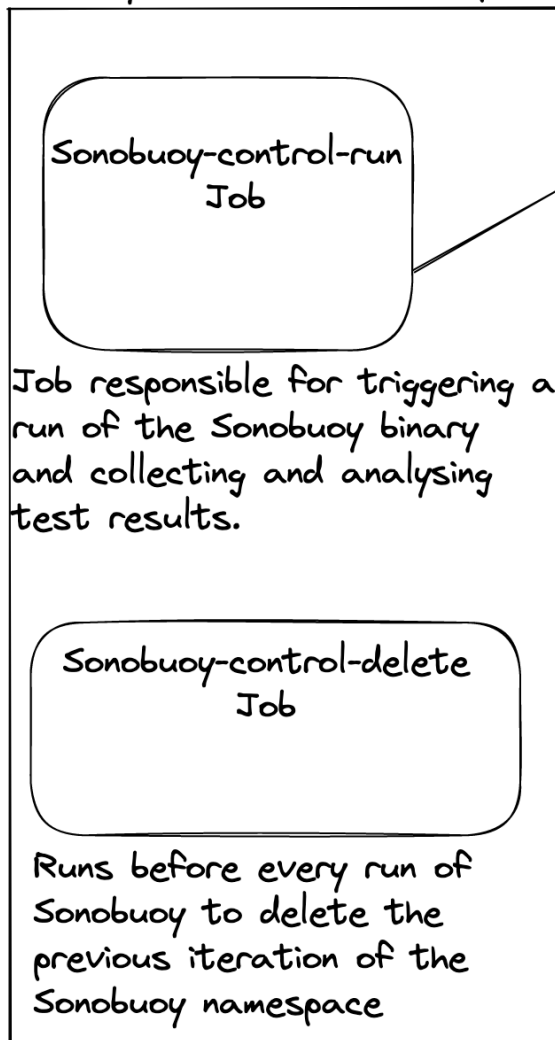
ARGO-4361 Add test for custom metric scaling

yesterday




Our CD setup

Sonobuoy-Control-Job Namespace




Sonobuoy Namespace



Within the Sonobuoy namespace a first pod is created to co-ordinate running the tests and gather results from them.

e2e-test-smoke-<numbers> Namespace



These namespaces are spun up and torn down by our Sonobuoy plugin, each only being used for one test.

Our e2e tests

- Networking: pod-to-pod communication
- Istio: configuration related to our multi-cluster setup
- DNS: internal and external resolution
- HPA and scaling on custom metrics
- Ability of pods to assume their IAM role
- All nodes are in a known state



Example

```
22 var _ = SmokeDescribe("kube-proxy", func() {
23     f := framework.NewDefaultFramework("e2e-test-smoke")
24
25     ginkgo.It("should allow pod to pod traffic", func() {
26         ns := f.Namespace.Name
27         cs := f.ClientSet
28         ctx := context.Background()
29         createOpts := metav1.CreateOptions{}
30
31         ginkgo.By("Creating a source pod")
32         srcName := "src"
33         srcPod := &v1.Pod{
34             ObjectMeta: metav1.ObjectMeta{
35                 Name: srcName,
36                 Labels: map[string]string{
37                     "name": srcName,
38                 },
39             },
40             Spec: v1.PodSpec{
41                 Containers: []v1.Container{
42                     {
43                         Name: srcName,
44                         Image: imageutils.GetE2EImage(imageutils.Agnhost),
45                     },
46                 },
47             },
48         }
49         srcPod, _ = f.ClientSet.CoreV1().Pods(ns).Create(ctx, srcPod, createOpts)
50         framework.ExpectNoError(e2epod.WaitForPodNameRunningInNamespace(cs, srcPod.Name, srcPod.Namespace), "Pod %s failed to run", srcPod.Name)
```

Example

```
52     ginkgo.By("Creating a new destination service and pod")
53     dstName := "dst"
54     jig := service.NewTestJig(cs, ns, dstName)
55     dstPod := &v1.Pod{
56         ObjectMeta: metav1.ObjectMeta{
57             Name:     dstName,
58             Labels: jig.Labels,
59         },
60         Spec: v1.PodSpec{
61             Containers: []v1.Container{
62                 {
63                     Name:     dstName,
64                     Image: imageutils.GetE2EImage(imageutils.NginxNew),
65                     Ports: []v1.ContainerPort{{ContainerPort: 80}},
66                     ReadinessProbe: &v1.Probe{
67                         Handler: v1.Handler{
68                             HTTPGet: &v1.HTTPGetAction{
69                                 Port: intstr.FromInt(80),
70                             },
71                         },
72                     },
73                 },
74             },
75         },
76     }
77     dstPod, _ = f.ClientSet.CoreV1().Pods(ns).Create(ctx, dstPod, createOpts)
78     framework.ExpectNoError(e2epod.WaitForPodNameRunningInNamespace(cs, dstPod.Name, dstPod.Namespace), "Pod %s failed to run", dstPod.Name)
```


Example

```
80     ginkgo.By("Creating a destination service")
81     dstPort := int32(80)
82     _, err := jig.CreateTCPServiceWithPort(nil, dstPort)
83     gomega.Expect(err).NotTo(gomega.HaveOccurred(), "Failed to create service with port %s in namespace %s", dstPort, ns)
84
85     ginkgo.By("Making the request from source to destination via kube-proxy")
86     err = wait.PollImmediate(interval, shortTimeout, func() (bool, error) {
87         // We're adding the || true because we want to retry the curl command if it fails.
88         // Kube-proxy needs some time to propagate the iptables rule across all the nodes.
89         // See https://github.com/skyscannertools/k8s/test-infra/pull/31 for more info.
90         cmd := fmt.Sprintf("curl -o /dev/null -i -q -s -S -w %{http_code} --connect-timeout 10 http://%s.%s || true", dstName, ns)
91         stdout := f.ExecShellInPod(srcPod.Name, cmd)
92         if stdout == "200" {
93             return true, nil
94         }
95         framework.Logf("Expected status code 200, got %s. Retrying...", stdout)
96         return false, nil
97     })
98     framework.ExpectNoError(err, "Failed to get 200 response code within %v seconds.", shortTimeout)
99 })
100 }
```

Gotchas, Tips & Other Options



Gotchas with our approach




liggitt commented on 26 Jun 2019 Member  ...

Copied from the referenced golang issue thread:

k8s.io/kubernetes is not primarily intended to be consumed as a module. Only the published subcomponents are (and go get works properly with those).


If you want to consume k8s.io/kubernetes as a module, you'd probably need to add require directives for matching versions of all of the subcomponents, rather than using go get

/close

  2  51



Gotchas with our approach



abursavich commented on 15 Aug 2019

Contributor

😊 ...

For anyone else who hits this issue, after much weeping and gnashing of teeth, this is the little script I wrote to switch kubernetes versions:

```
#!/bin/sh
set -euo pipefail

VERSION=${1#"v"}
if [ -z "$VERSION" ]; then
    echo "Must specify version!"
    exit 1
fi
MODS=(
    curl -sS https://raw.githubusercontent.com/kubernetes/kubernetes/v${VERSION}/go.mod |
    sed -n 's|.*k8s.io/(.*) => ./staging/src/k8s.io/.*|k8s.io/\1|p'
)
for MOD in "${MODS[@]}; do
    V=$(
        go mod download -json "${MOD}@kubernetes-${VERSION}" |
        sed -n 's|.*"Version": "(.*)".*|\1|p'
    )
    go mod edit "-replace=${MOD}=${MOD}@${V}"
done
go get "k8s.io/kubernetes@v${VERSION}"
```

😊

👍 115

😄 11

🔔 5

❤️ 19

🚀 9

🔗 9

Gotchas with our approach

- [Versioning the e2e code](#)
- Kubernetes version upgrades
- Testing images
- Flaky tests
- May need to modify the [run.sh](#) script



Tips

- Capture the behavior your users care about
- Leverage the community's efforts – not just the code, also the images
- Figure out which tests can safely be run in parallel



Other Projects

- Kubetest2 framework
 - <https://github.com/aws/aws-k8s-tester>
- <https://testinfra.readthedocs.io/en/latest/>
- <https://github.com/vapor-ware/kubetest>
- <https://github.com/kuberhealthy/kuberhealthy>



Thanks



@gjtempleton @ruio @SkyscannerEng