

RESILIENCE  
REALIZED



KubeCon



CloudNativeCon

North America 2021

# Vitess: Introduction and New Features

*Deepthi Sigireddi - PlanetScale @ATechGirl*

*Alkin Tezuysal - PlanetScale @ask\_dba*

*Andrew Mason - Slack Corp. @andrew\_mason1*

*Malcolm Akinje - Slack Corp. @LESkidd113*

# Agenda

- ❖ Introduction - *Deepthi Sigireddi*
- ❖ VTOP (Vitess Operator) + Demo - *Alkin Tezuysal*
- ❖ VTAdmin Overview + Demo - *Andrew Mason*
- ❖ VDiff Coordinator Demo - *Malcolm Akinje*

# Vitess

A database clustering system for horizontal scaling of MySQL / MariaDB

- CNCF graduated project
- Open source, Apache 2.0 licence
- Contributors from around the community
- Written in Golang



# Vitess

Cloud-native distributed database

- Runs in Kubernetes

Scalable

Highly available

Durability guarantees

Illusion of “single database”

- Single dedicated connection
- MySQL 5.7 or 8.0
- Compatible with frameworks / ORMs etc.

# Vitess serves **millions of QPS** in production

 slack

 New Relic®

 Square

*Flipkart*

HubSpot

*peak*

 Pinterest

 YouTube

 nozzle

 weave

GitHub

JD. 京东  
COM

Quiz of Kings

 stitchlabs

 planetscale

# Concepts

## Keyspace

- Logical database

## Shard

- Subset of logical database

## Cell

- Failure domain

# Vitess Architecture basics

A common replicated database cluster with primary and replicas



# Vitess Architecture basics

Each MySQL server is assigned a **vttablet**

- A daemon/sidecar
- Controls the **mysqld** process
- Interacts with the **mysqld** server
- Typically on same host as **mysqld**





# Vitess Architecture basics

In production you have multiple clusters



# Vitess Architecture basics

User and application traffic is routed via **vtgate**

- A smart, stateless proxy
- Speaks the MySQL protocol
- Impersonates a monolithic MySQL server
- Relays queries to **vtablets**



# Vitess Architecture basics

A vitess deployment will run multiple **vtgate** servers for scale out



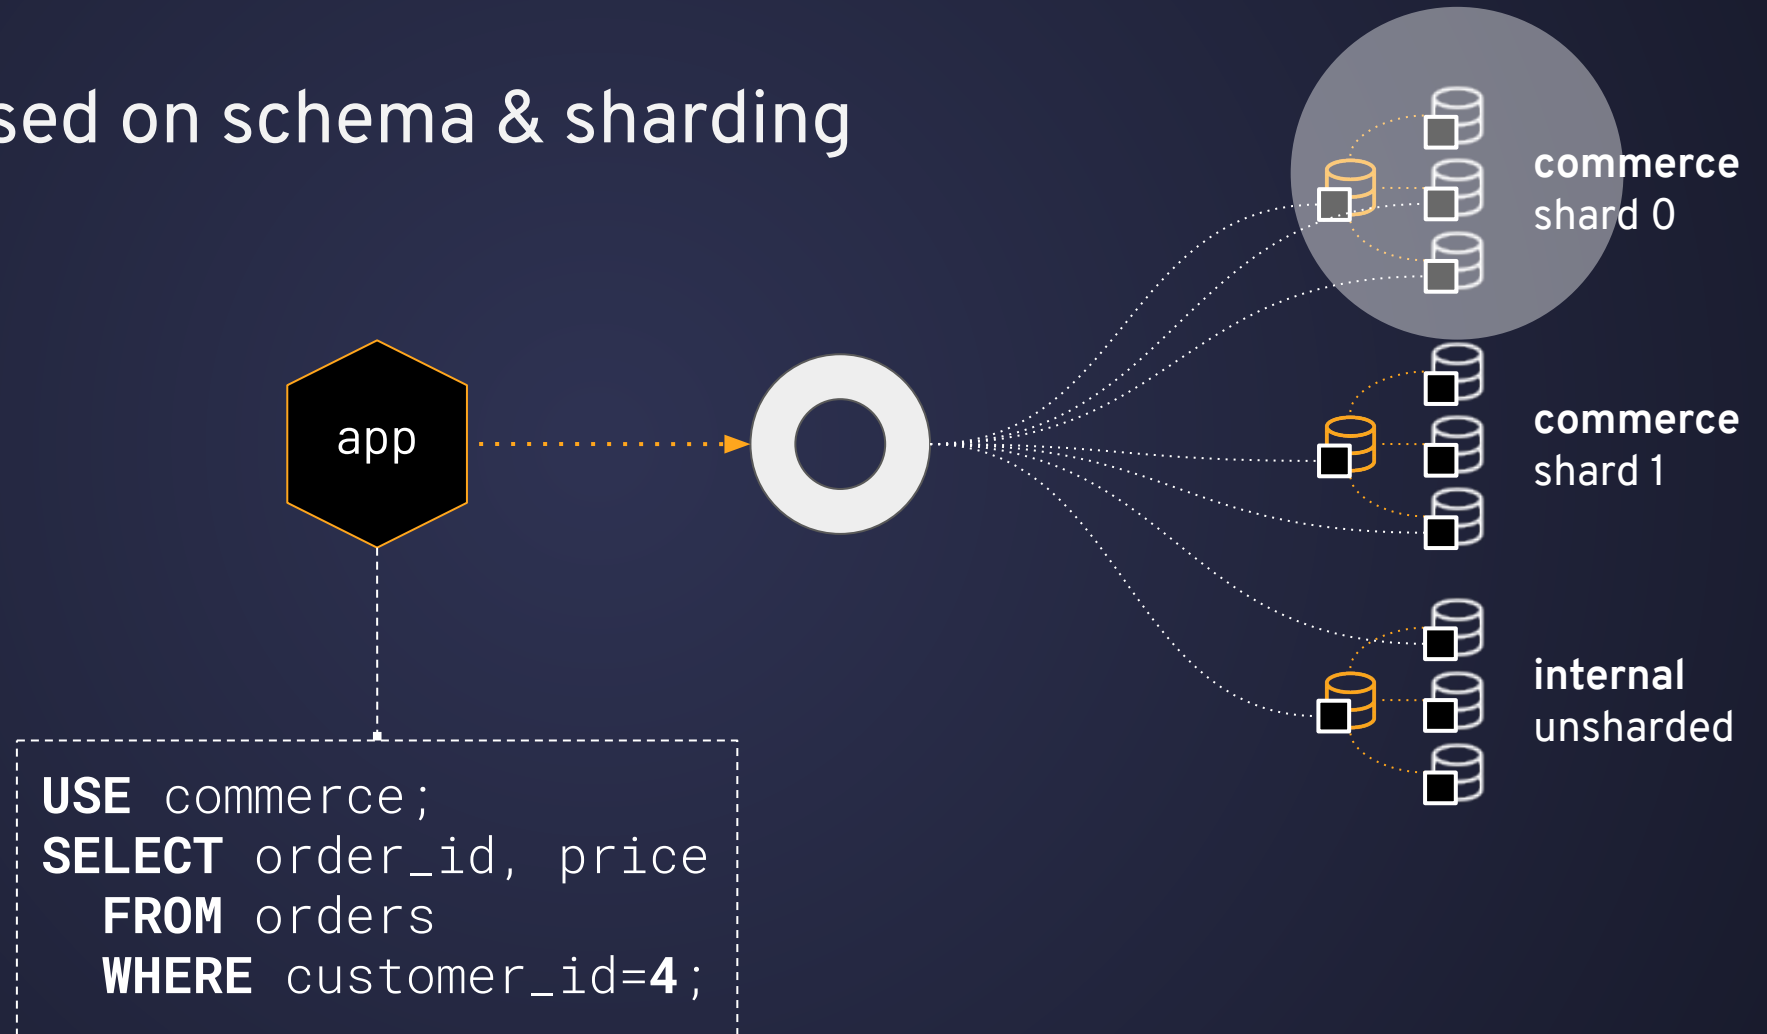
# Vitess Architecture basics

vtgate must transparently route queries to correct clusters, to relevant shards



# Vitess Architecture basics

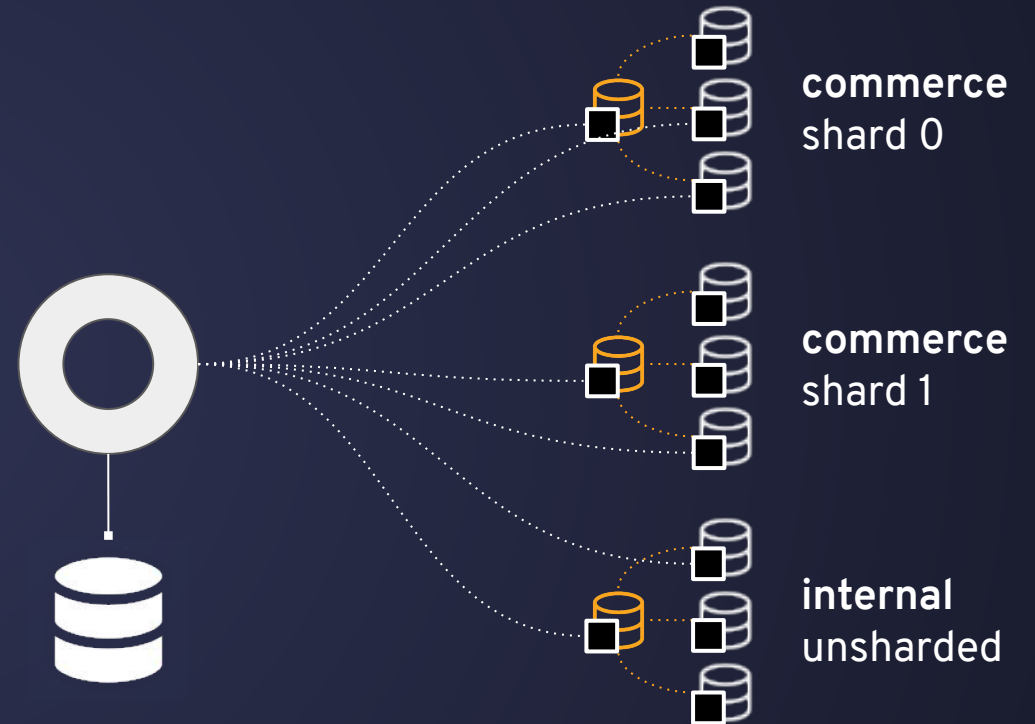
Queries route based on schema & sharding scheme



# Vitess Architecture basics

**topo:** distributed key/value store

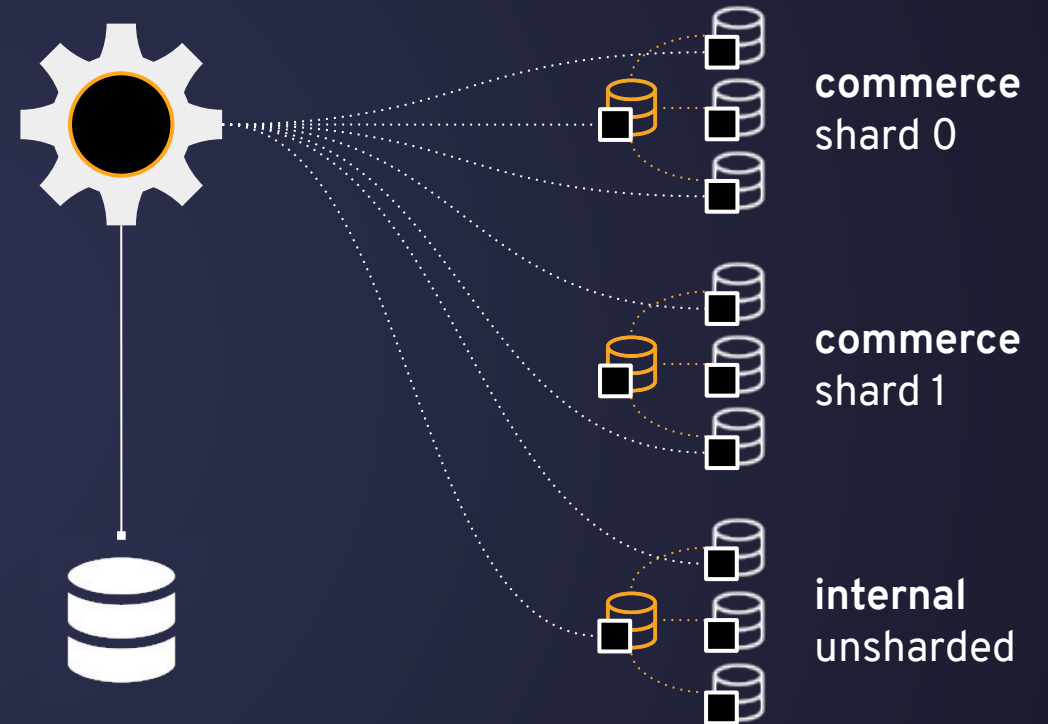
- Stores the state of vitess: schemas, shards, sharding scheme, tablets, roles, etc.
- etcd/zookeeper/kubernetes
- Small dataset, mostly cached by **vtgate**



# Vitess Architecture basics

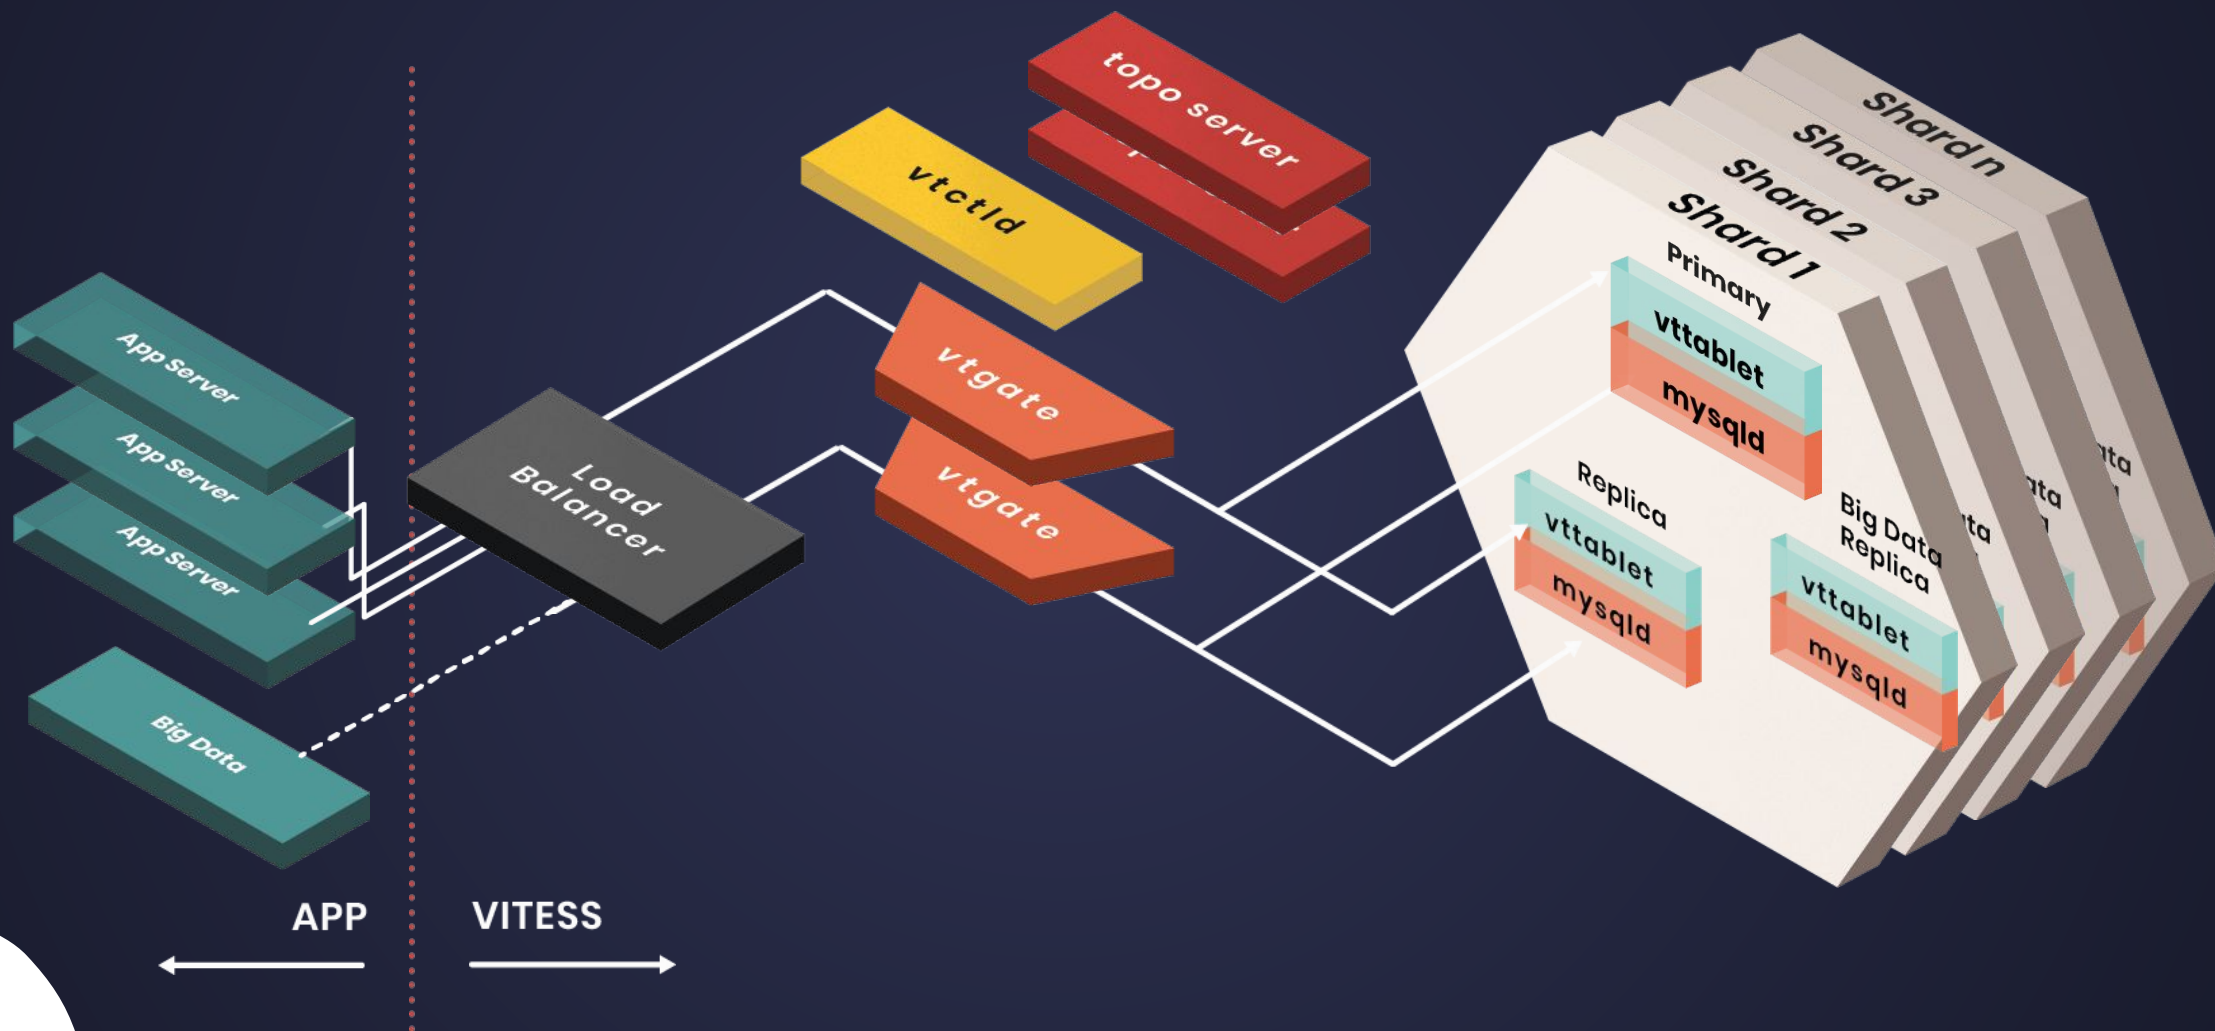
**vtctld**: control daemon

- Runs ad hoc operations
- API server
- Reads/writes **topo**
- Uses locks
- Operates on tablets





# 16 Vitess Architecture Summary





# New and Upcoming Features

- 12.0 release GA Oct 26 2021
- Gen4 planner - experimental in 12.0
- Continuous Benchmarking - since 11.0
- Performance Improvements - ongoing
- More supported query constructs - ongoing
- VTAdmin
- Multi-column vindexes
- Automatic failure detection
- Collations
- Distributed transactions

# VTOP - Vitess Operator for Kubernetes

## Operator Functions and Automation of Vitess functions

- Automation of Vitess functions
- Deploy Clusters
- Deploy Shards
- Failover nodes
- Recover pods
- Replicate data cross AZ



The Vitess Operator is open source on [GitHub](#).

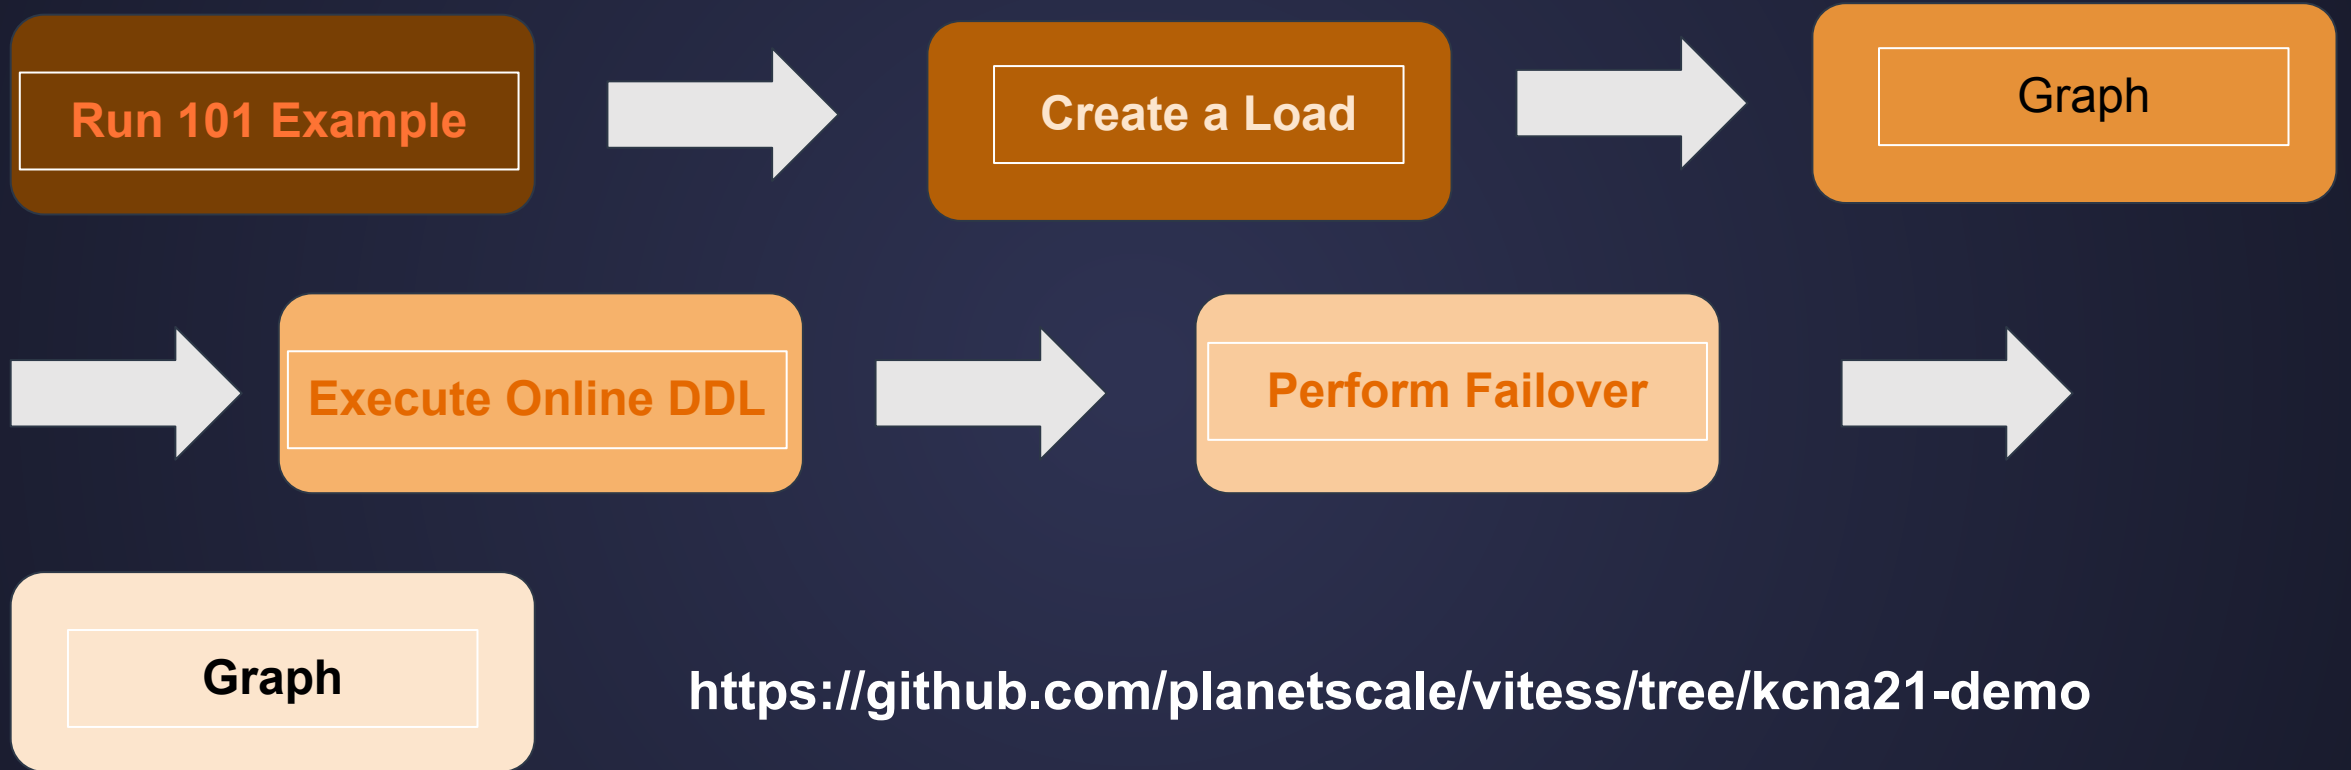
# VTOP - DEMO

## Use of Operator

- Deploying a test cluster with Operator (Already deployed)
- Creating Sharding Workflow (Already created)
- Create a database load (Live)
- Failover during OnlineDDL on Sharded cluster (Live)



# VTOP - DEMO Workflow

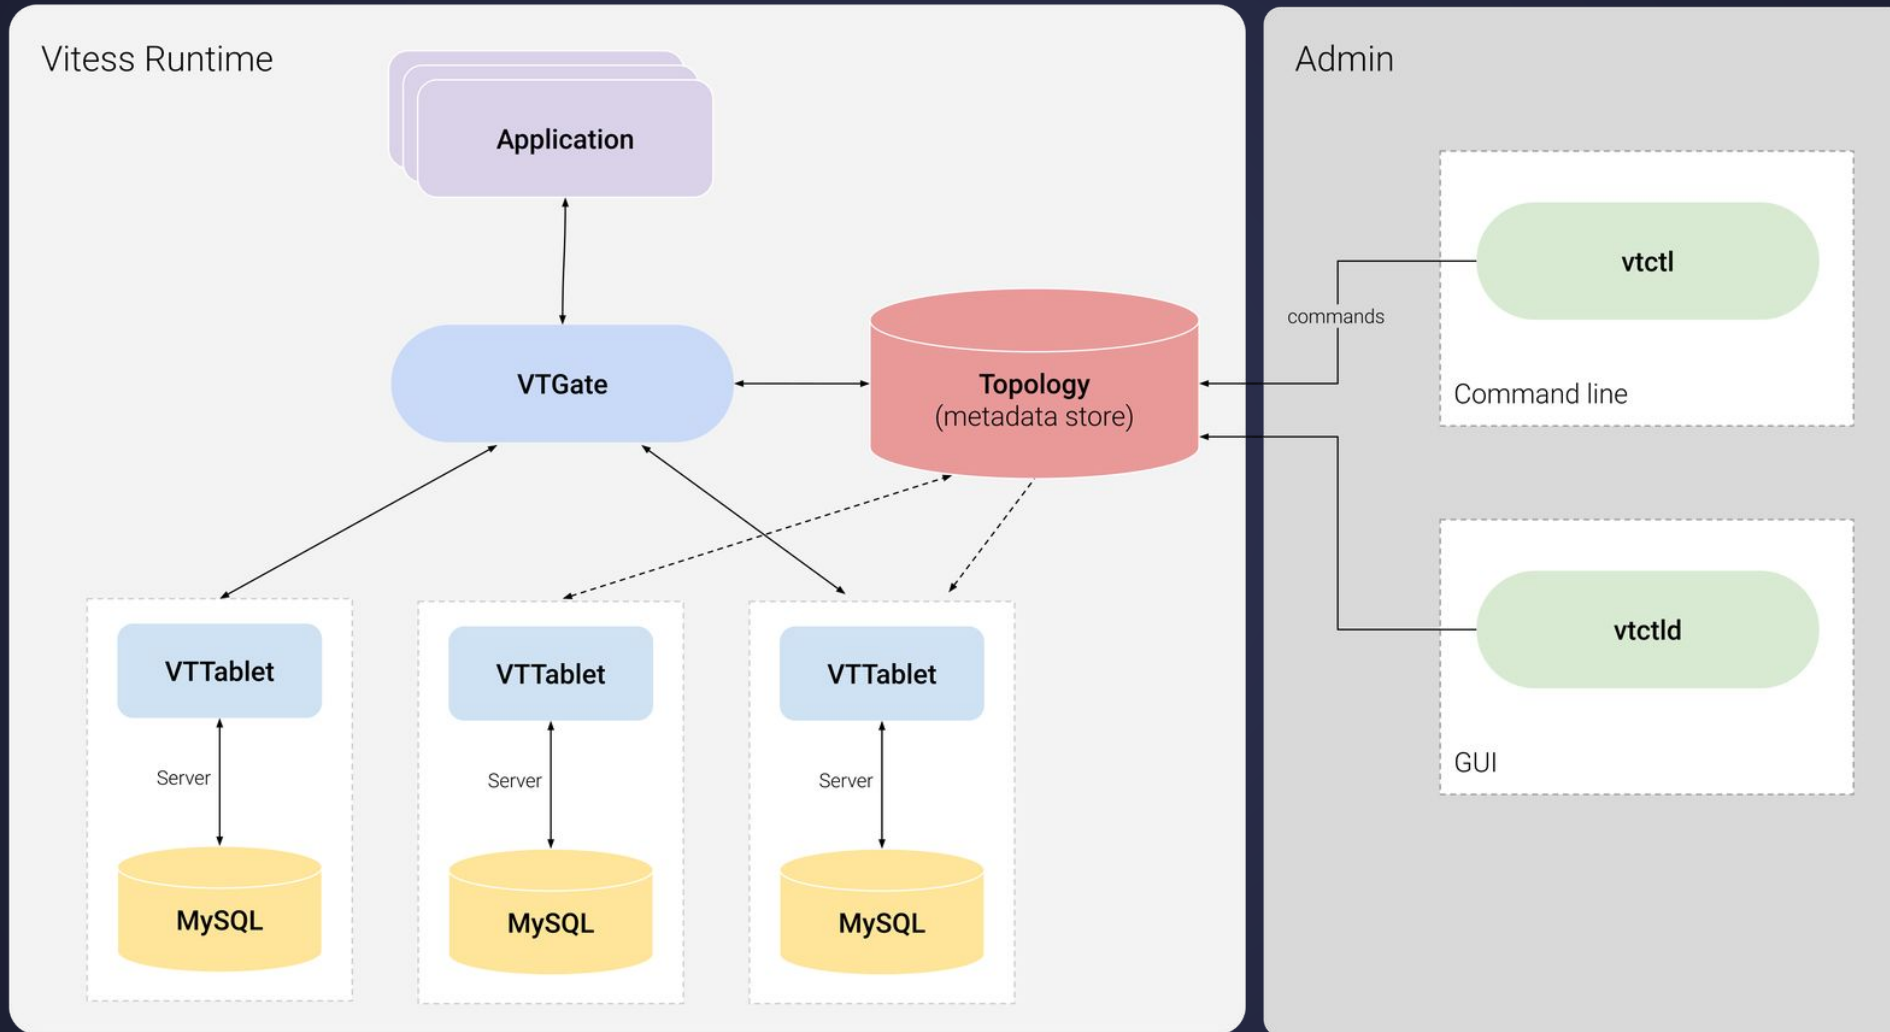


<https://github.com/planetscale/vitess/tree/kcna21-demo>

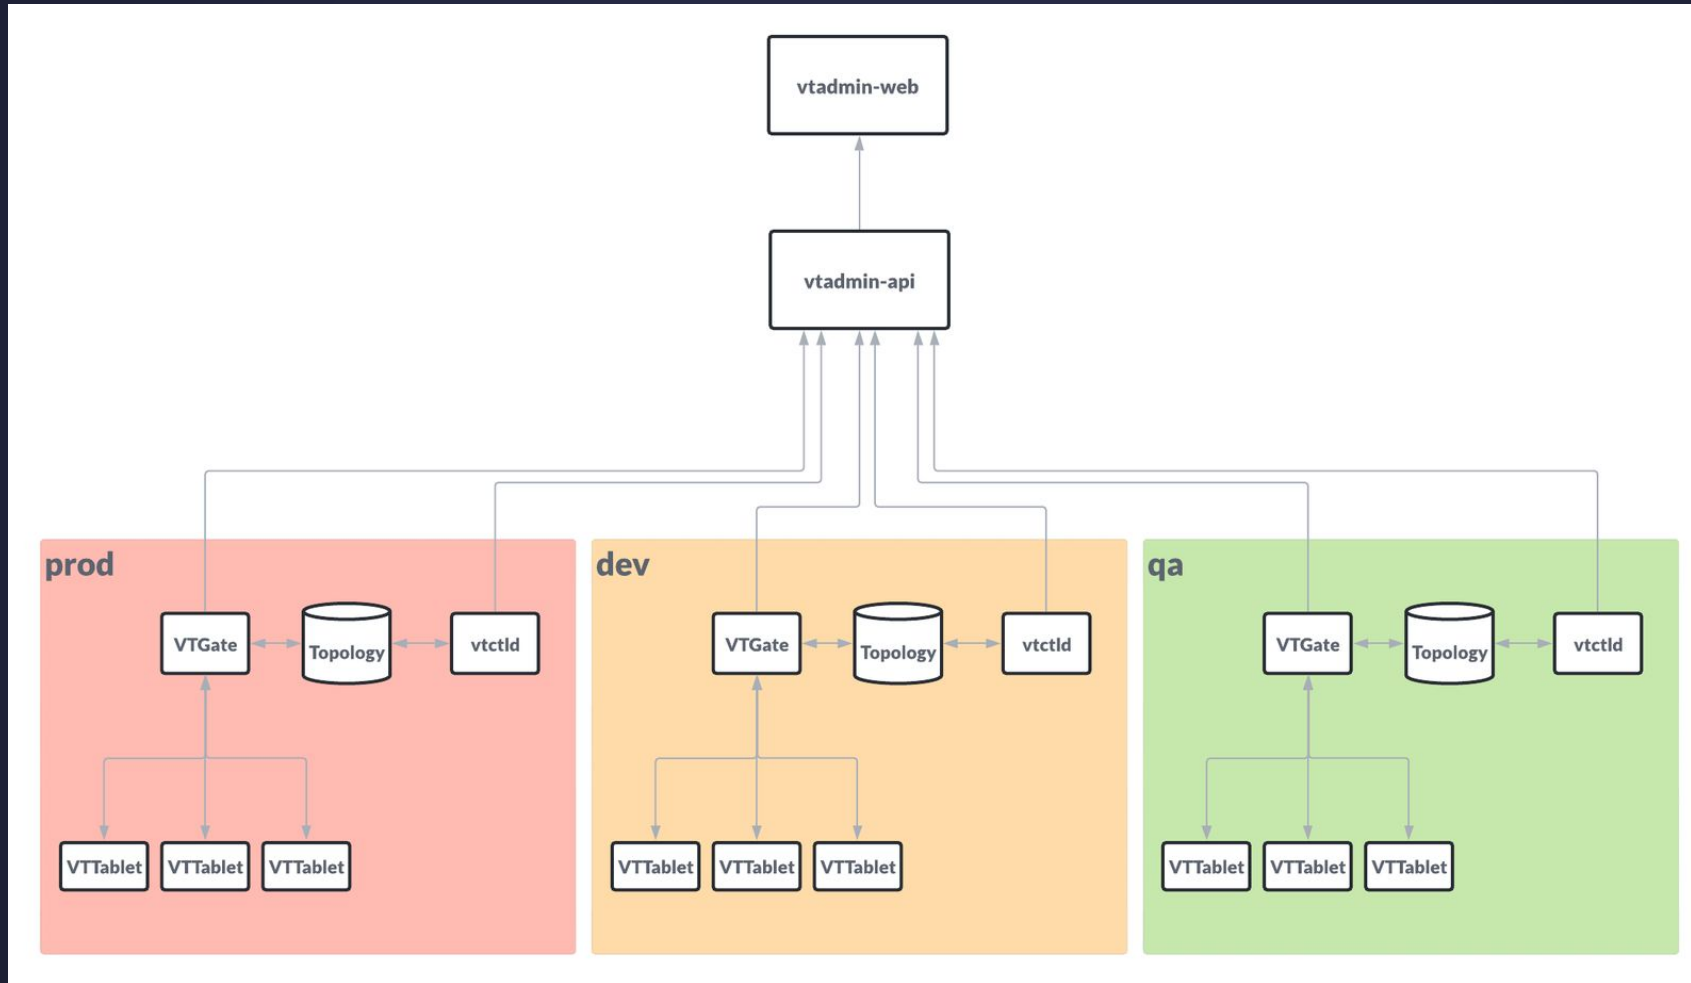
# VTAdmin - Overview

- Single control plane for many Vitess deployments (“clusters”)
- gRPC/http API and UI components
  - Golang backend with react frontend
- (Eventual) replacement for the existing vtctld UI

# VTAdmin - Architecture (single cluster)



# VTAdmin - Architecture (multi cluster)

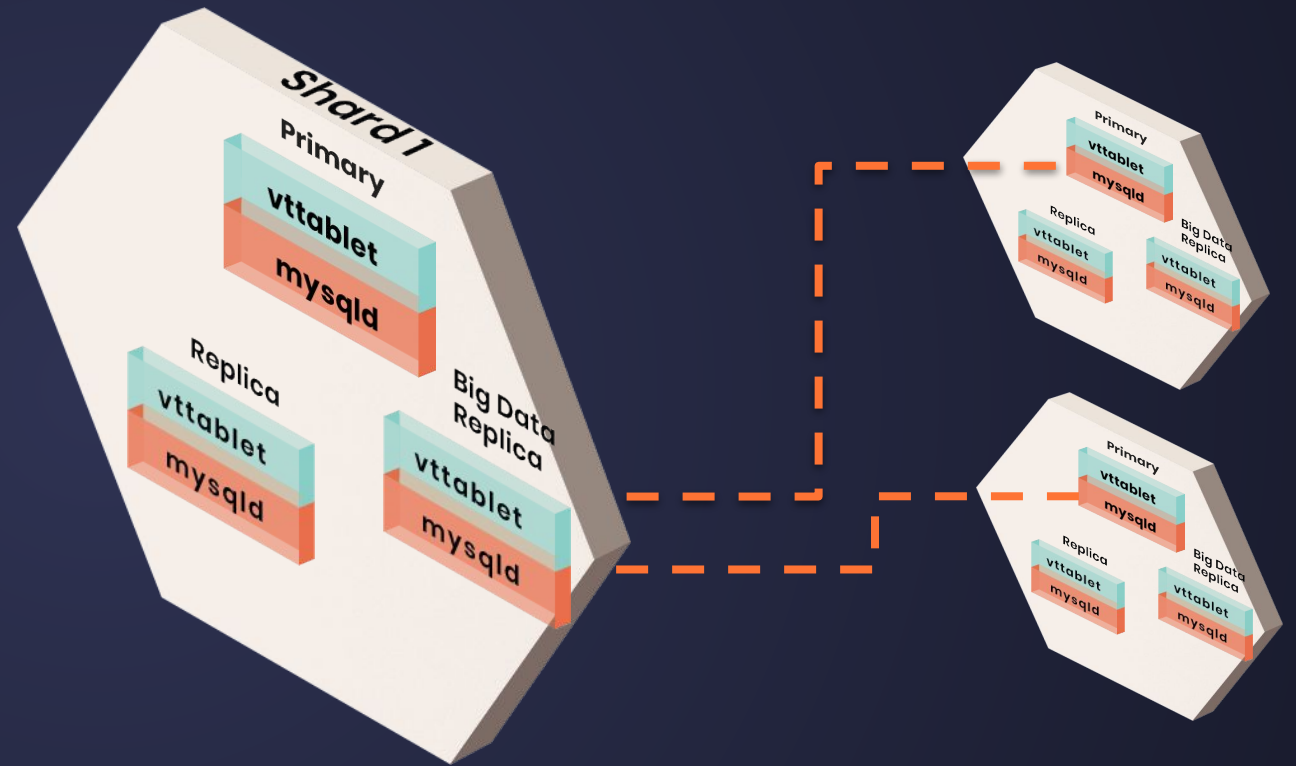


# VTAdmin Demo



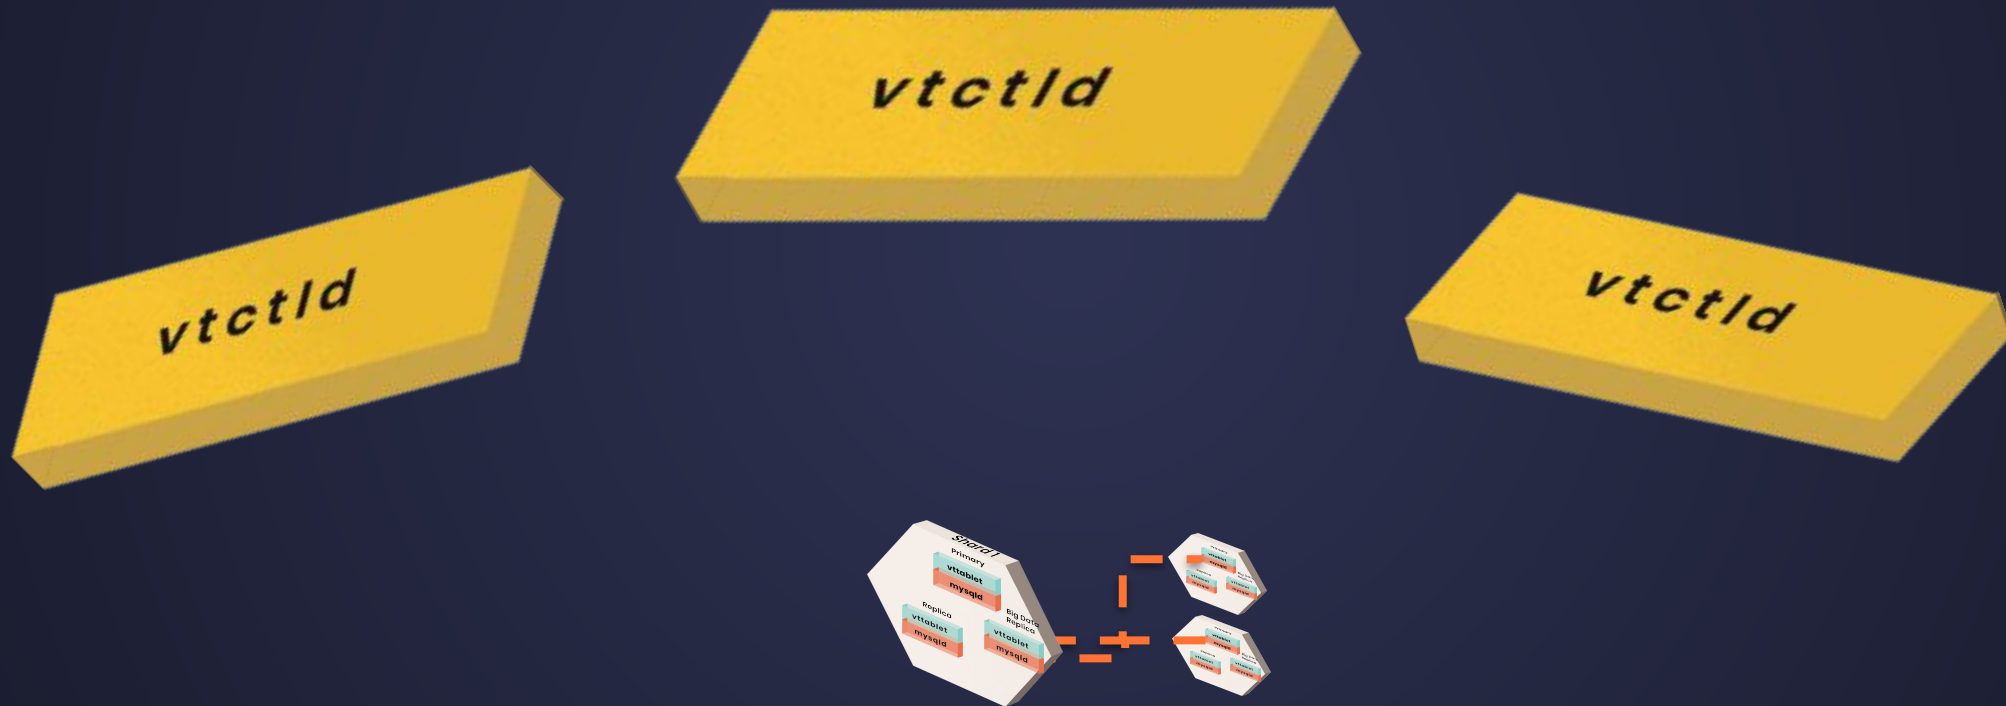
# VDiff

- Runs a diff between the source and target shards
- One VDiff per workflow
- This is a blocking call



# VDiff

(Work mainly happens in a vtctld)

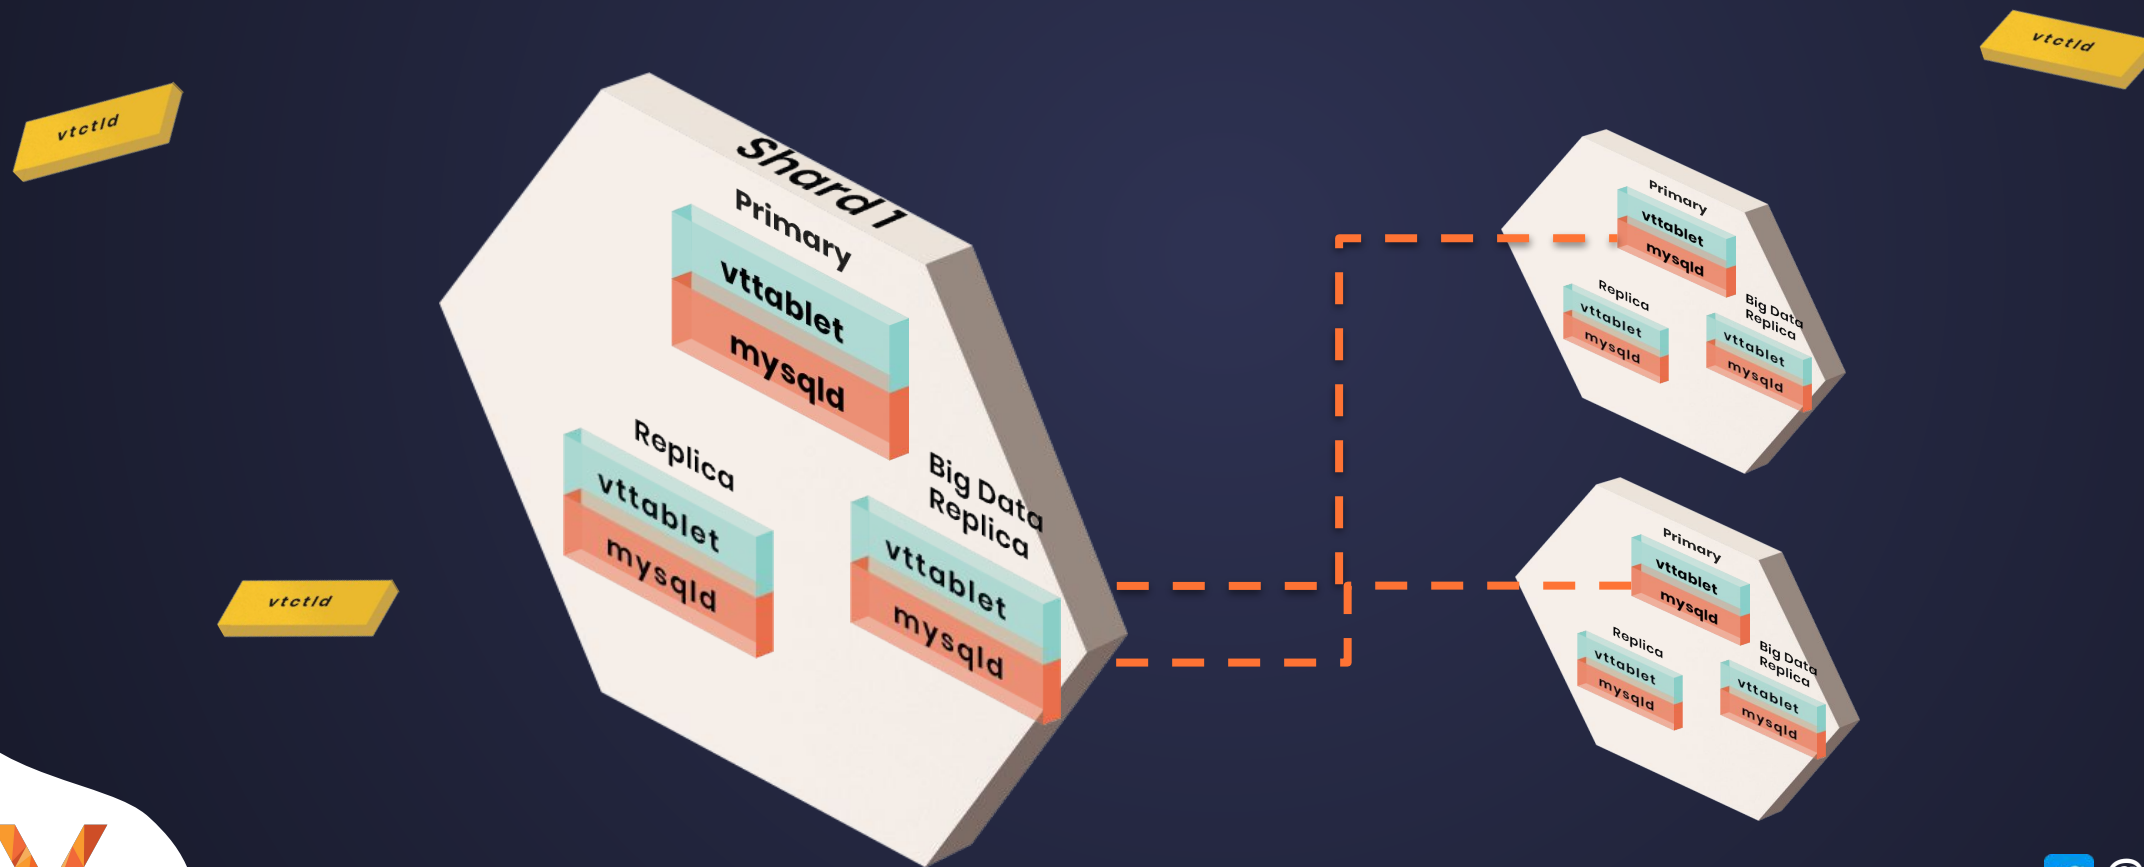


# VDiff

```
makinje@slack-vtctld-1:~$ sudo ./vtops vdiff \  
--scell us_east_1 \  
--tcell us_east_1 \  
--ttype replica \  
--ttype ronly \  
--wf keyspace.reshard_28-30 \  
--wf keyspace.reshard_50-58 \  
--wf keyspace.reshard_58-60  
I 0624 14:02:59.636186 12366 vdiff:106 {}  
I 0624 14:02:59.636221 12366 vdiff:109 {}  
I 0624 14:02:59.636232 12366 vdiff:109 {}  
I 0624 14:02:59.636240 12366 vdiff:109 {}  
I 0624 14:02:59.636250 12366 vdiff:133 {}  
I 0624 14:02:59.636259 12366 vdiff:136 {}  
I 0624 14:02:59.636268 12366 vdiff:142 {}  
I 0624 14:02:59.636284 12366 vdiff:136 {}  
I 0624 14:02:59.636292 12366 vdiff:142 {}  
I 0624 14:02:59.636320 12366 vdiff:136 {}  
I 0624 14:02:59.636329 12366 vdiff:142 {}  
I 0624 14:02:59.636346 12366 vdiff:148 {}  
I 0624 14:02:59.636355 12366 vdiff:151 {}  
I 0624 14:02:59.636368 12366 vdiff:151 {}  
🐞🐞VDiffWorker-0 picked up keyspace.reshard_28-30🐞🐞  
I 0624 14:02:59.636375 12366 vdiff:151 {} keyspace.reshard_58-60  
🐞🐞VDiffWorker-1 picked up keyspace.reshard_50-58🐞🐞  
Retries Left: 3  
Retries Left: 3  
🐞🐞VDiffWorker-2 picked up keyspace.reshard_58-60🐞🐞  
Retries Left: 3  
🟢🟢 Vdiff for keyspace.reshard_50-58 completed 🟢🟢  
🟢🟢 Vdiff for keyspace.reshard_28-30 completed 🟢🟢  
🟢🟢 Vdiff for keyspace.reshard_58-60 completed 🟢🟢  
🔍🔍Found Vtctlds🔍🔍  
slack-vtctld-1  
slack-vtctld-2  
slack-vtctld-3  
👤👤Starting 3 VDiff Workers👤👤  
VDiffWorker-0  
Using: slack-vtctld-1.example.com:15999  
VDiffWorker-1  
Using: slack-vtctld-2.example.com:15999  
VDiffWorker-2  
Using: slack-vtctld-3.example.com:15999  
📁📁Enqueuing 3 workflows📁📁  
keyspace.reshard_28-30  
keyspace.reshard_50-58
```

# VDiff Coordinator Demo

# VDiff Future Goals (diffs done on tablet layer)



# Resources



Docs: [vitepress.dev/](https://vitepress.dev/)

Code: [github.com/vitepress/vitepress](https://github.com/vitepress/vitepress)

Slack: [vitepress.slack.com](https://vitepress.slack.com)

Thank you!