



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Building Container Images in Kubernetes: It's Been a Journey!

Laurent Bernaille & Eric Mountain

 @ericmountain
 @lbernail

About Datadog

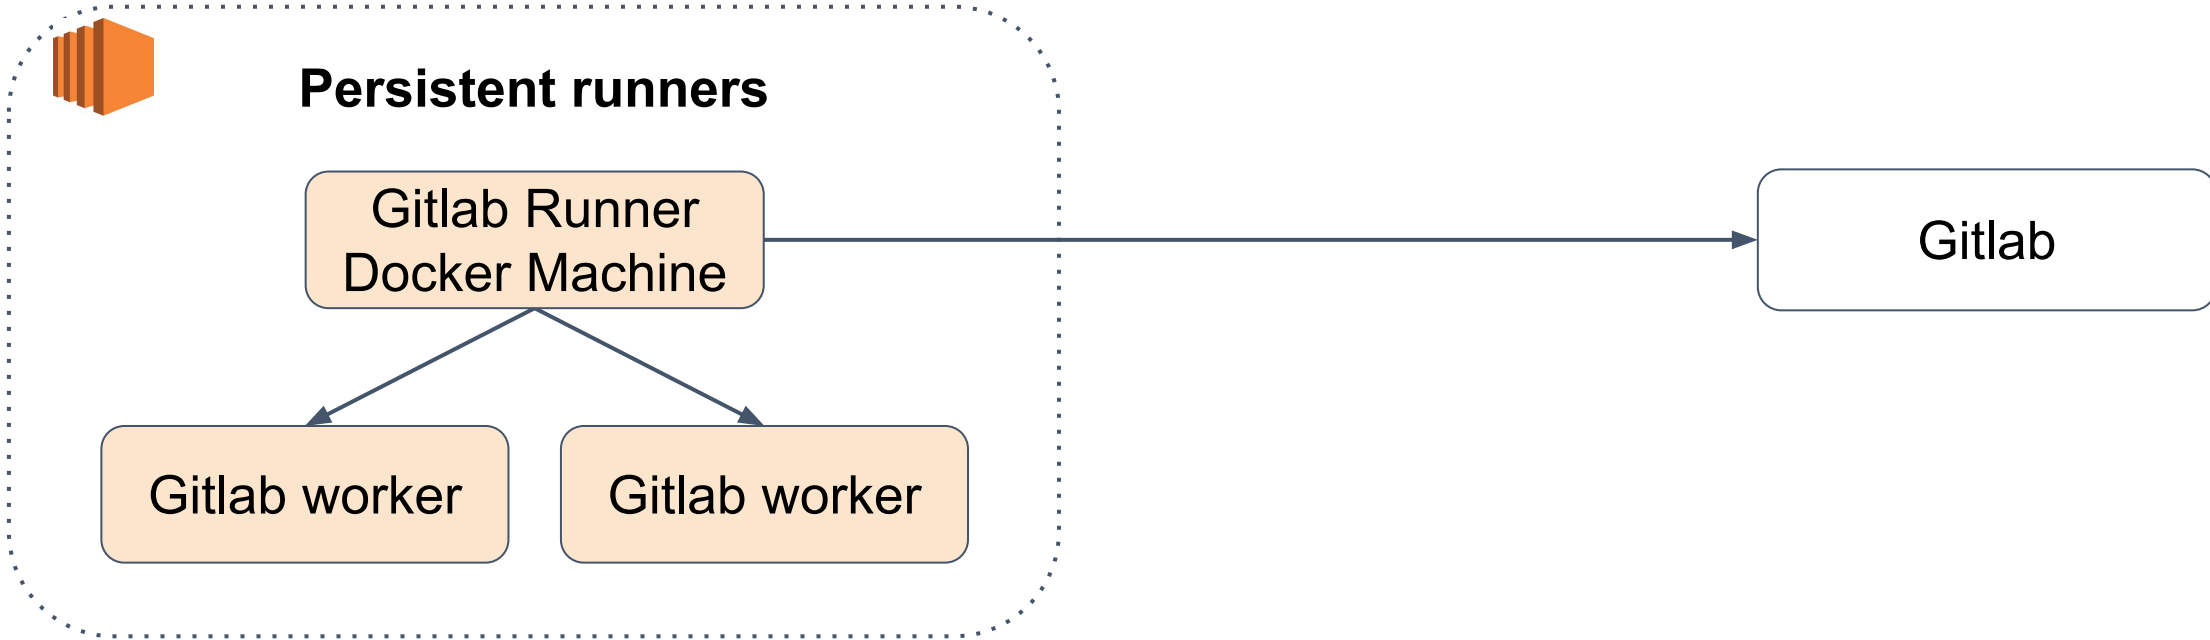
Over 500 integrations
Over 3000 employees
Over 18,500 customers
Millions of hosts reporting
Trillions of data points per day

10000s hosts in our infra
Dozens of k8s clusters
Multi-cloud
Rapid growth

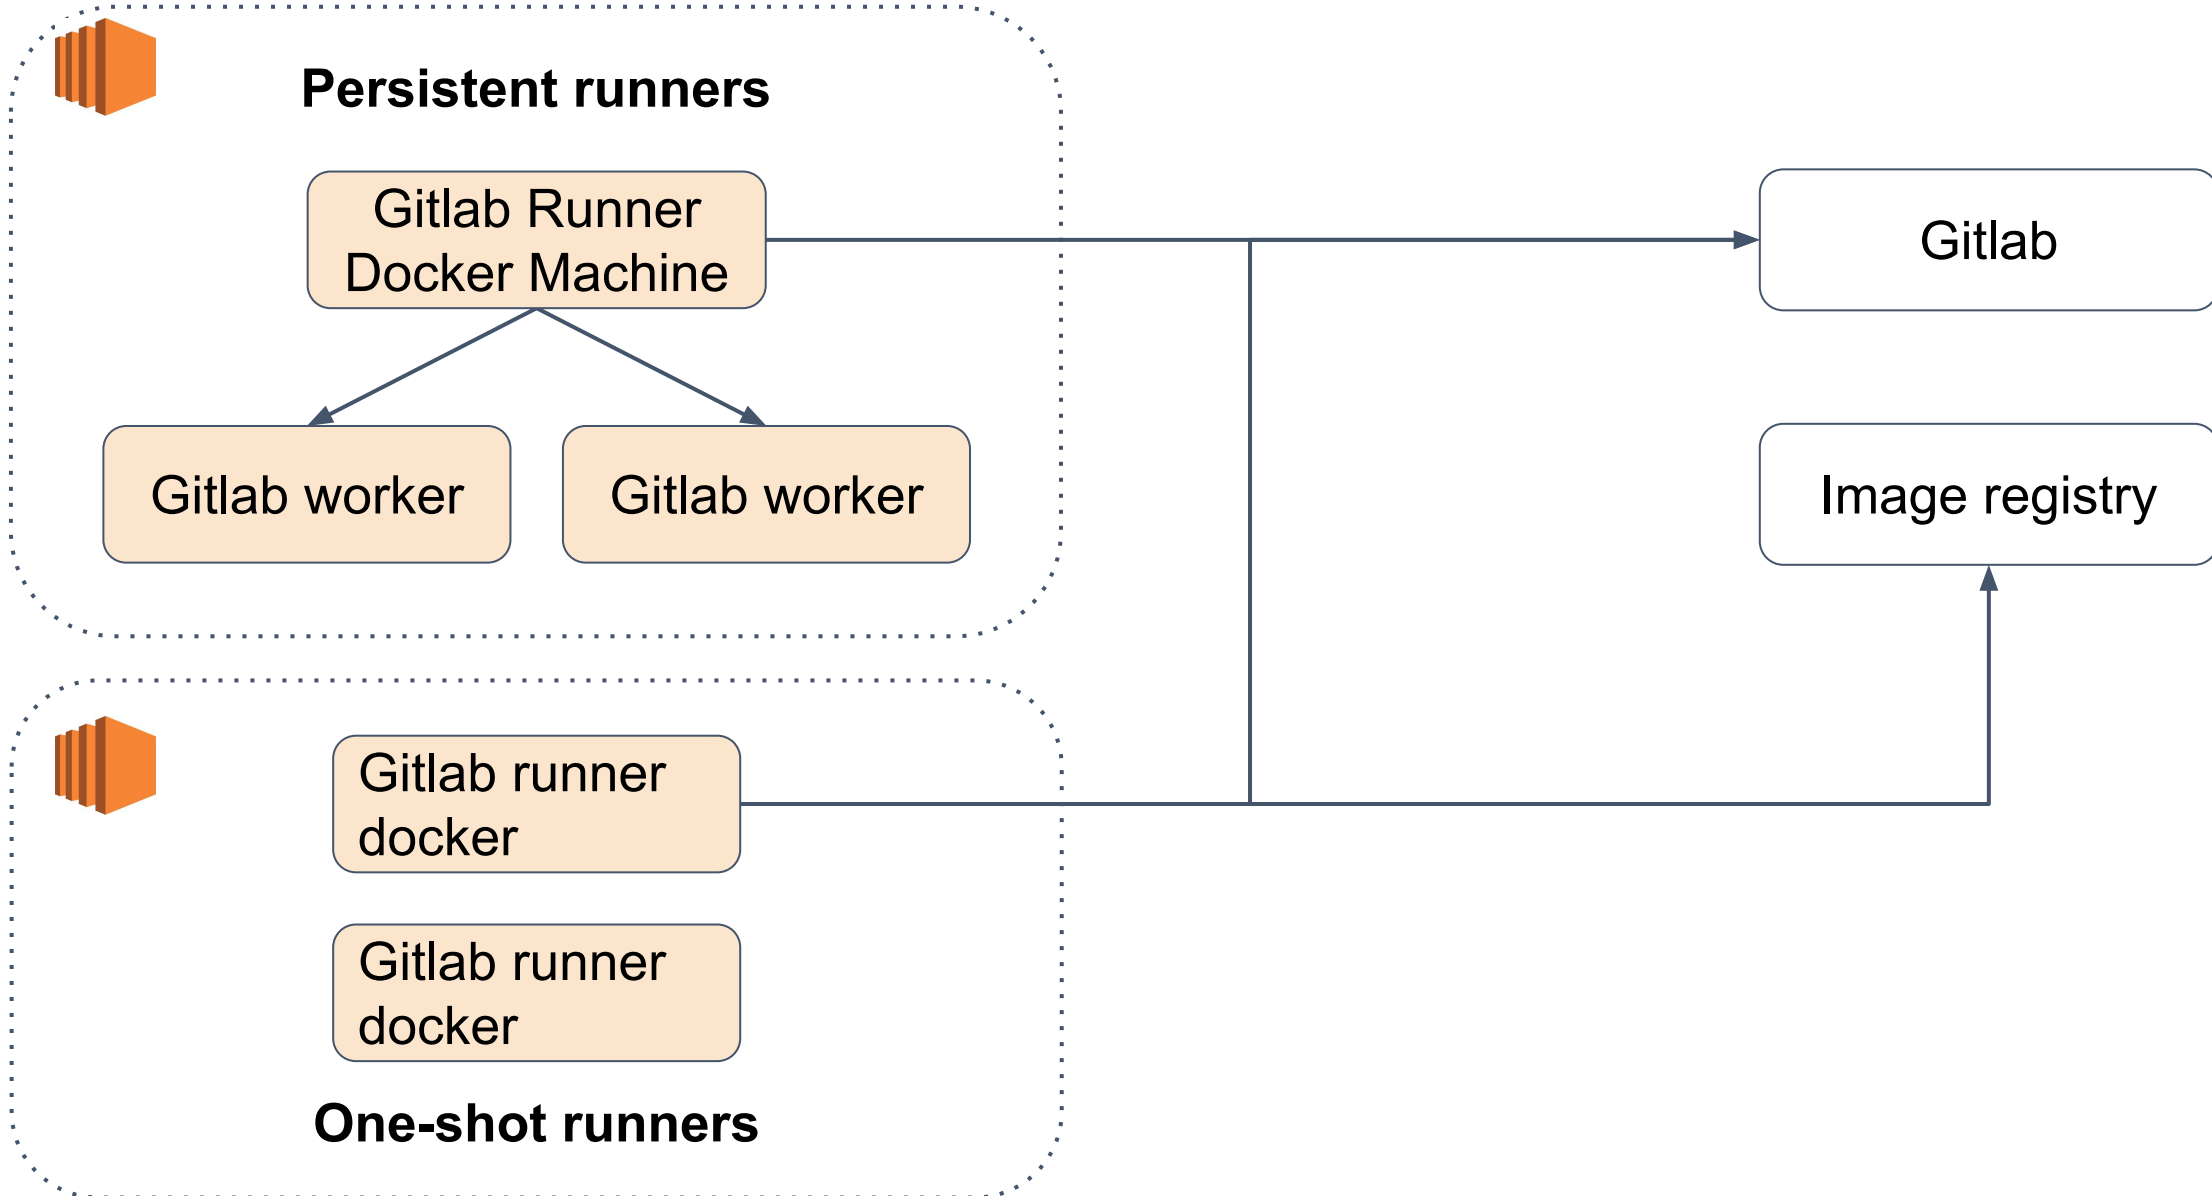


A Quick history of our build system

History: Building our applications




History: Building docker images



History: Using kubernetes builders



Persistent runners in k8s



Gitlab runner
docker

Gitlab runner
docker

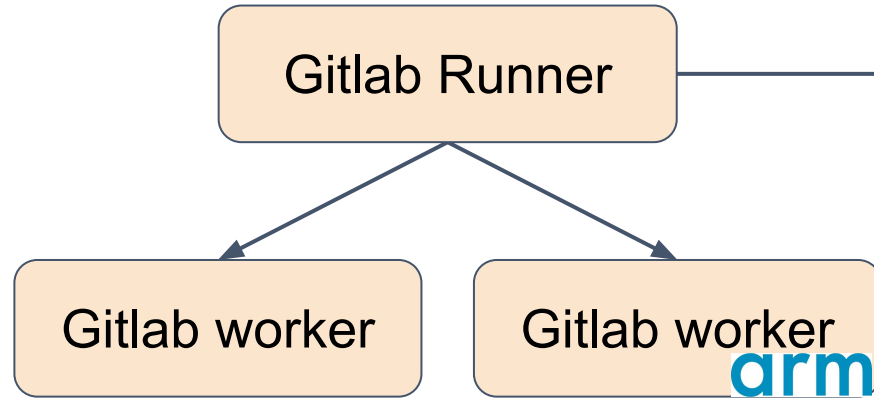
One-shot runners



History: Building ARM binaries



Persistent runners in k8s



Gitlab runner
docker

Gitlab runner
docker

One-shot runners


Gitlab

Image registry

History: Building multiarch images



Persistent runners in k8s



Gitlab runner
buildx  EMU

Gitlab runner
buildx  EMU

One-shot runners

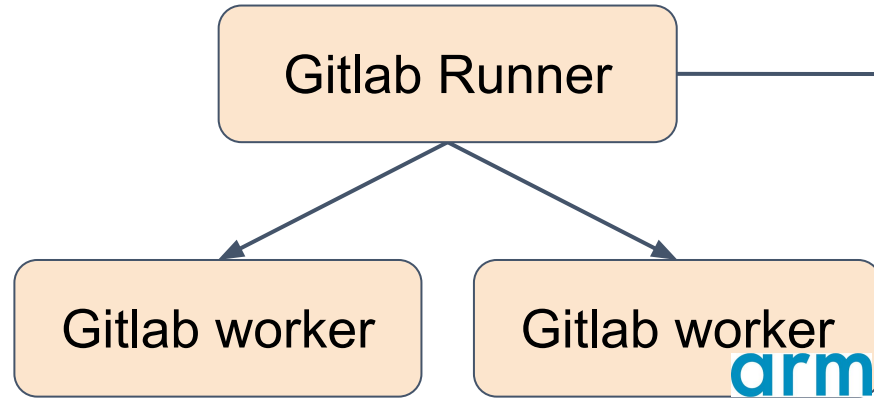
Gitlab

Image registry

History: Building multiarch images



Persistent runners in k8s



Gitlab runner
buildx  QEMU

Gitlab runner
buildx  QEMU

One-shot runners

Gitlab

Image registry

Limits of this system

- Last workload outside of Kubernetes
- Several new needs requiring investment in legacy platform
- Challenging in terms of security

What if we could build images *in Kubernetes?*

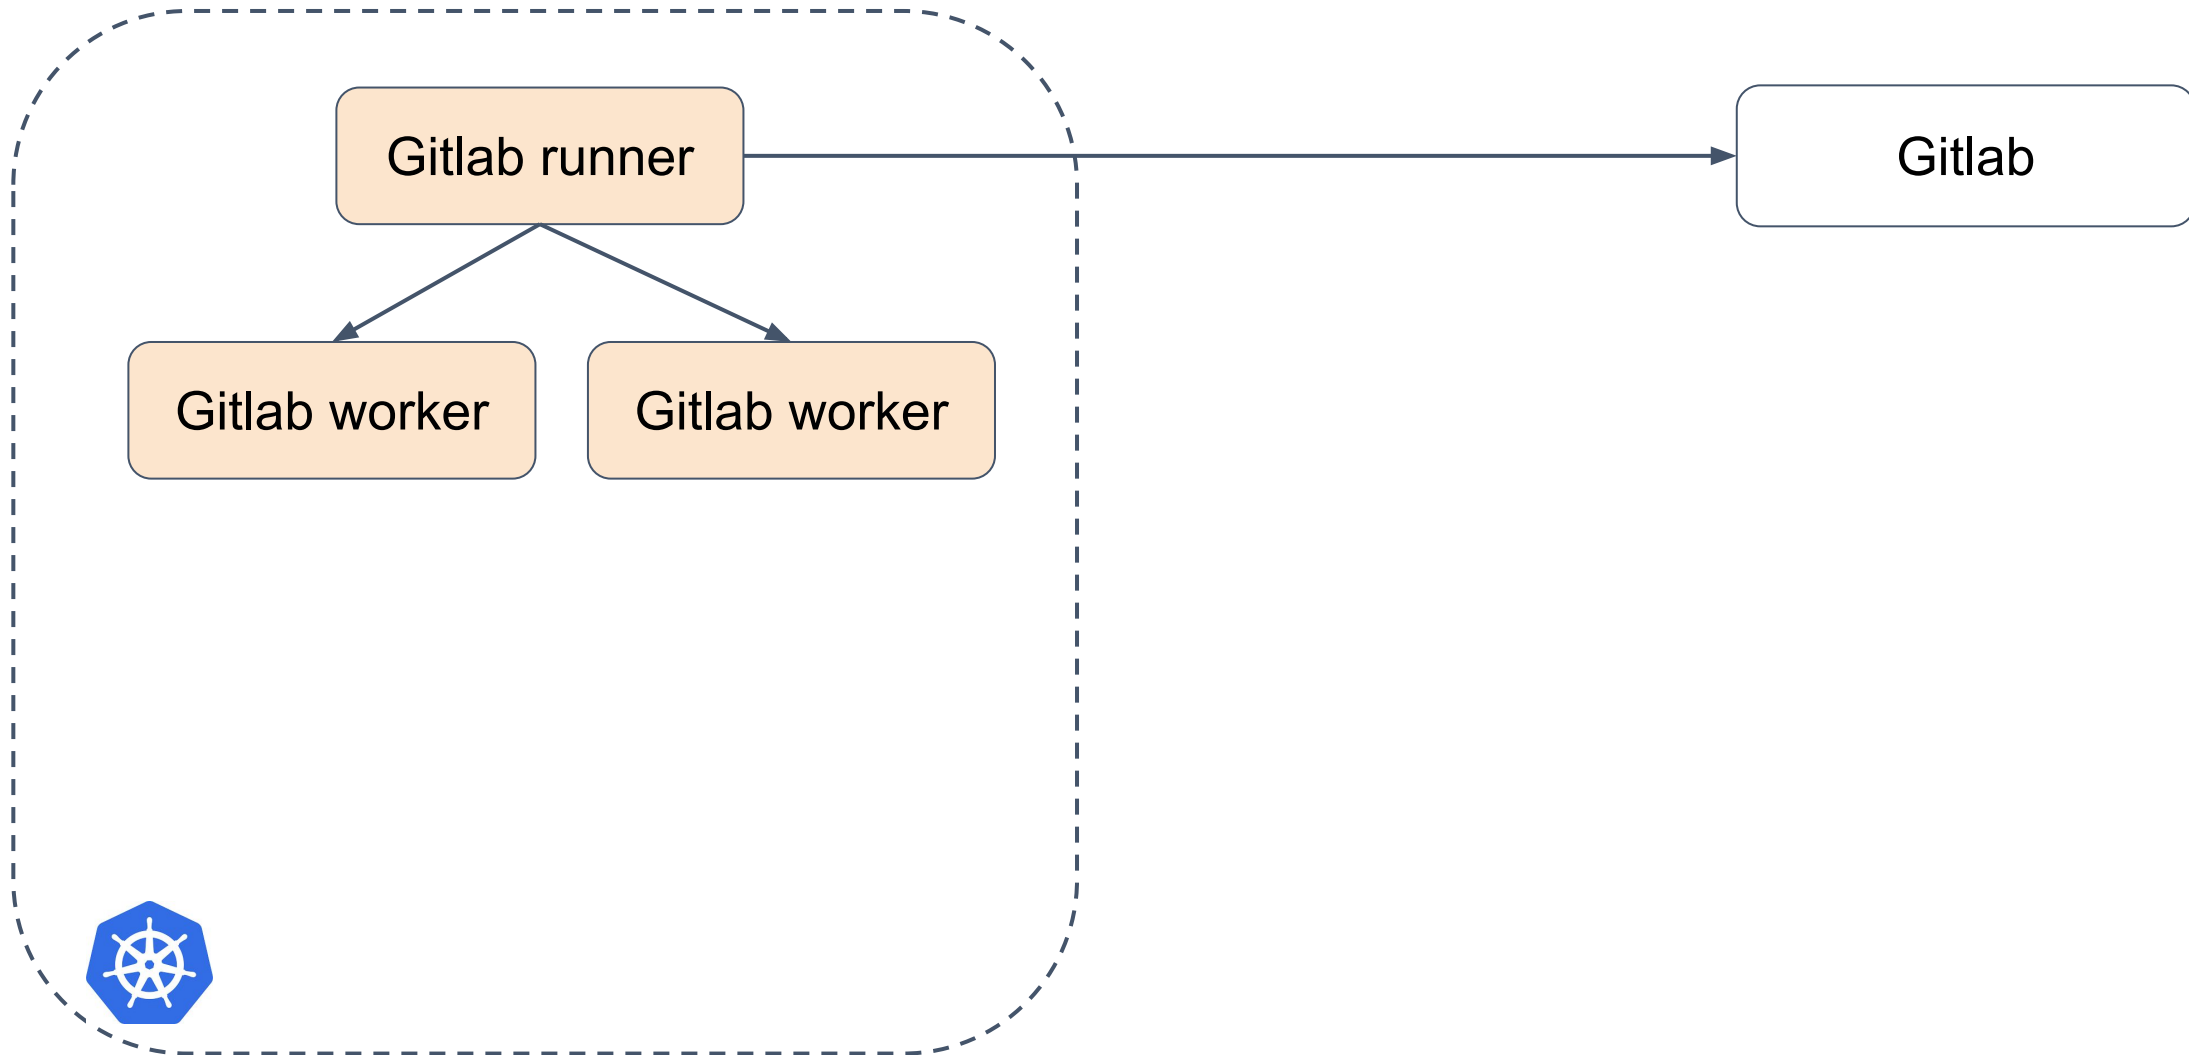
Can we build images in Kubernetes?

- docker-in-docker
- Standalone builders
 - buildah
 - kaniko
 - img
- Dedicated build daemons: buildkitd

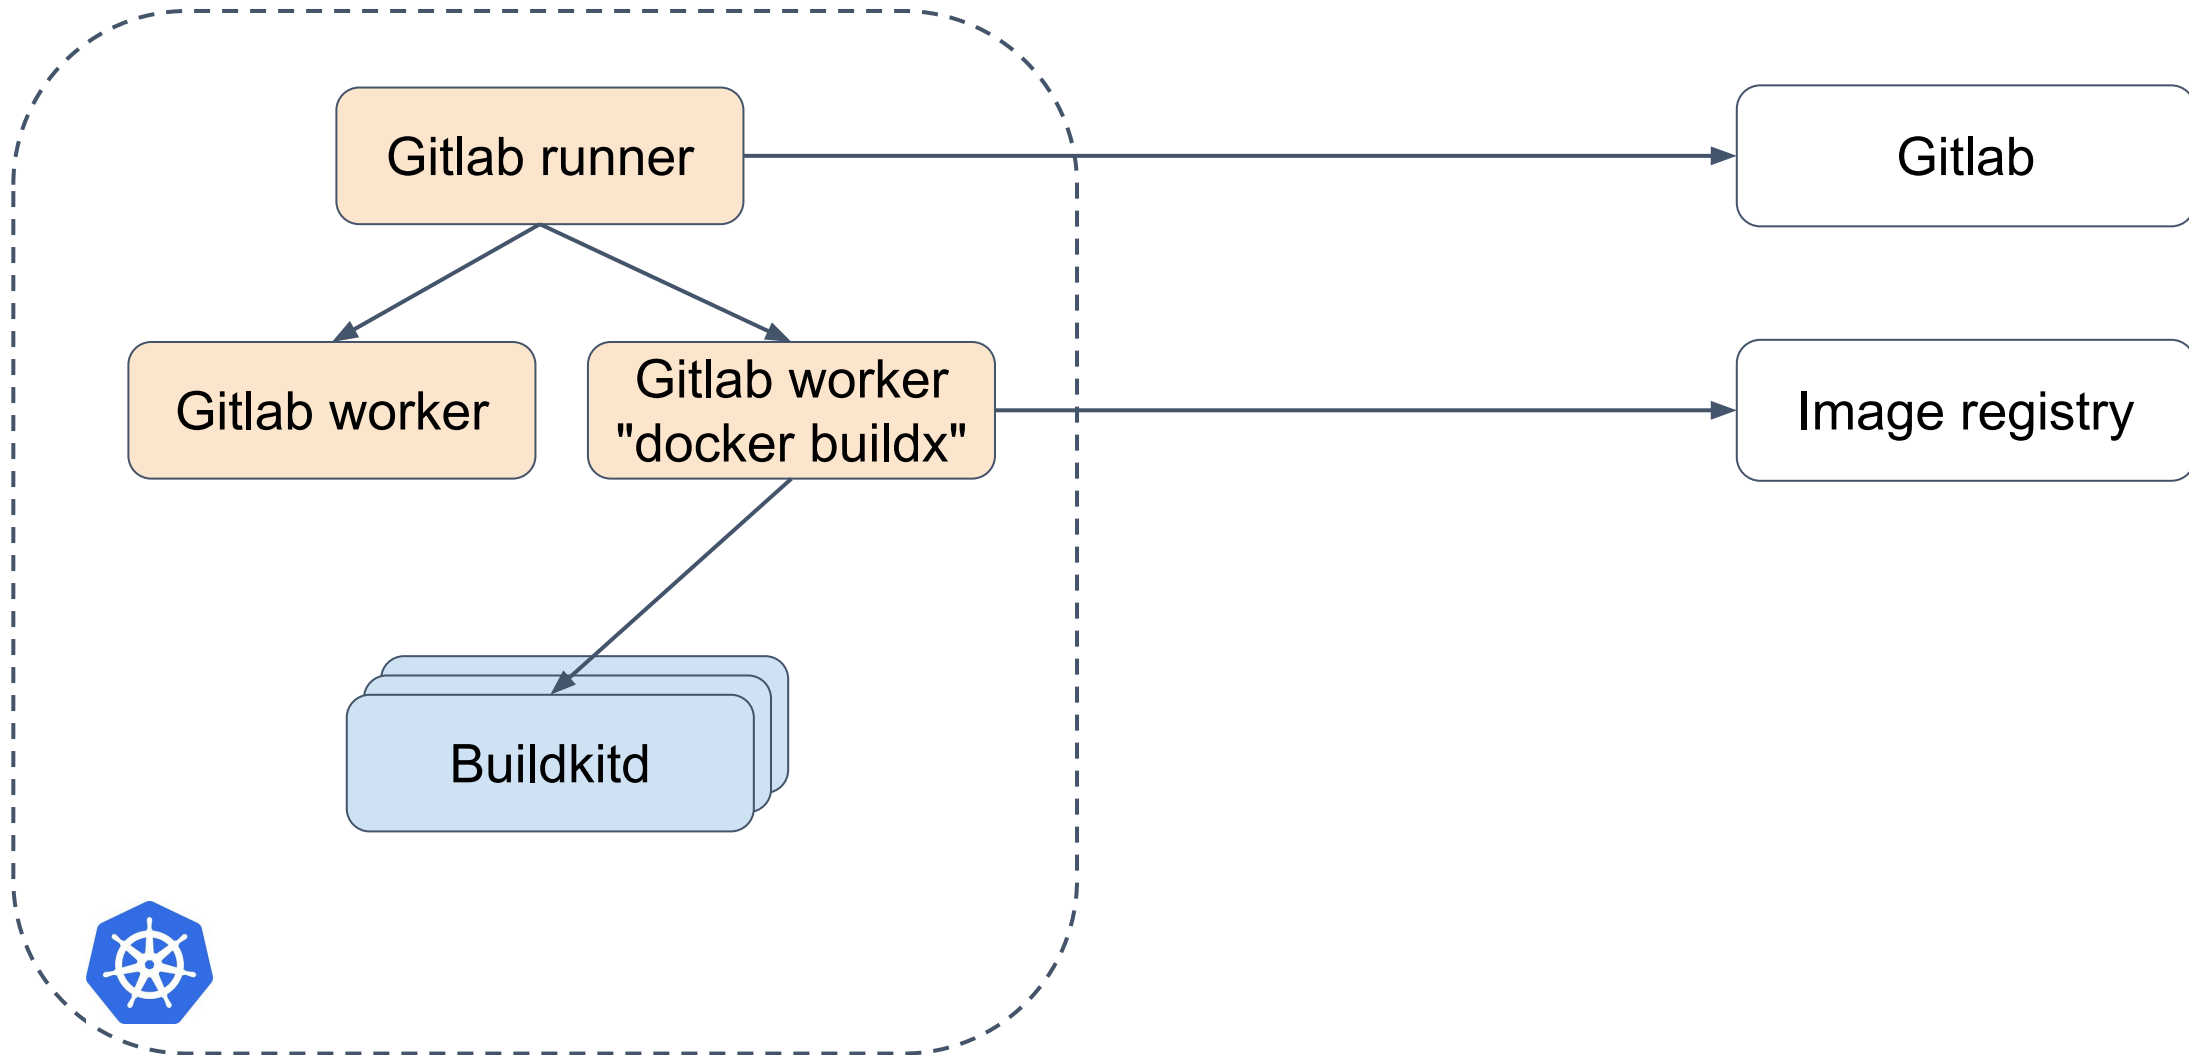
Can we use them?

- **docker-in-docker**
=> **Security implications**
- **Standalone builders**
=> Complex setup: Distribute jobs / Assemble multiarch images
- Dedicated build daemons: **buildkitd**
=> Great UX, promising

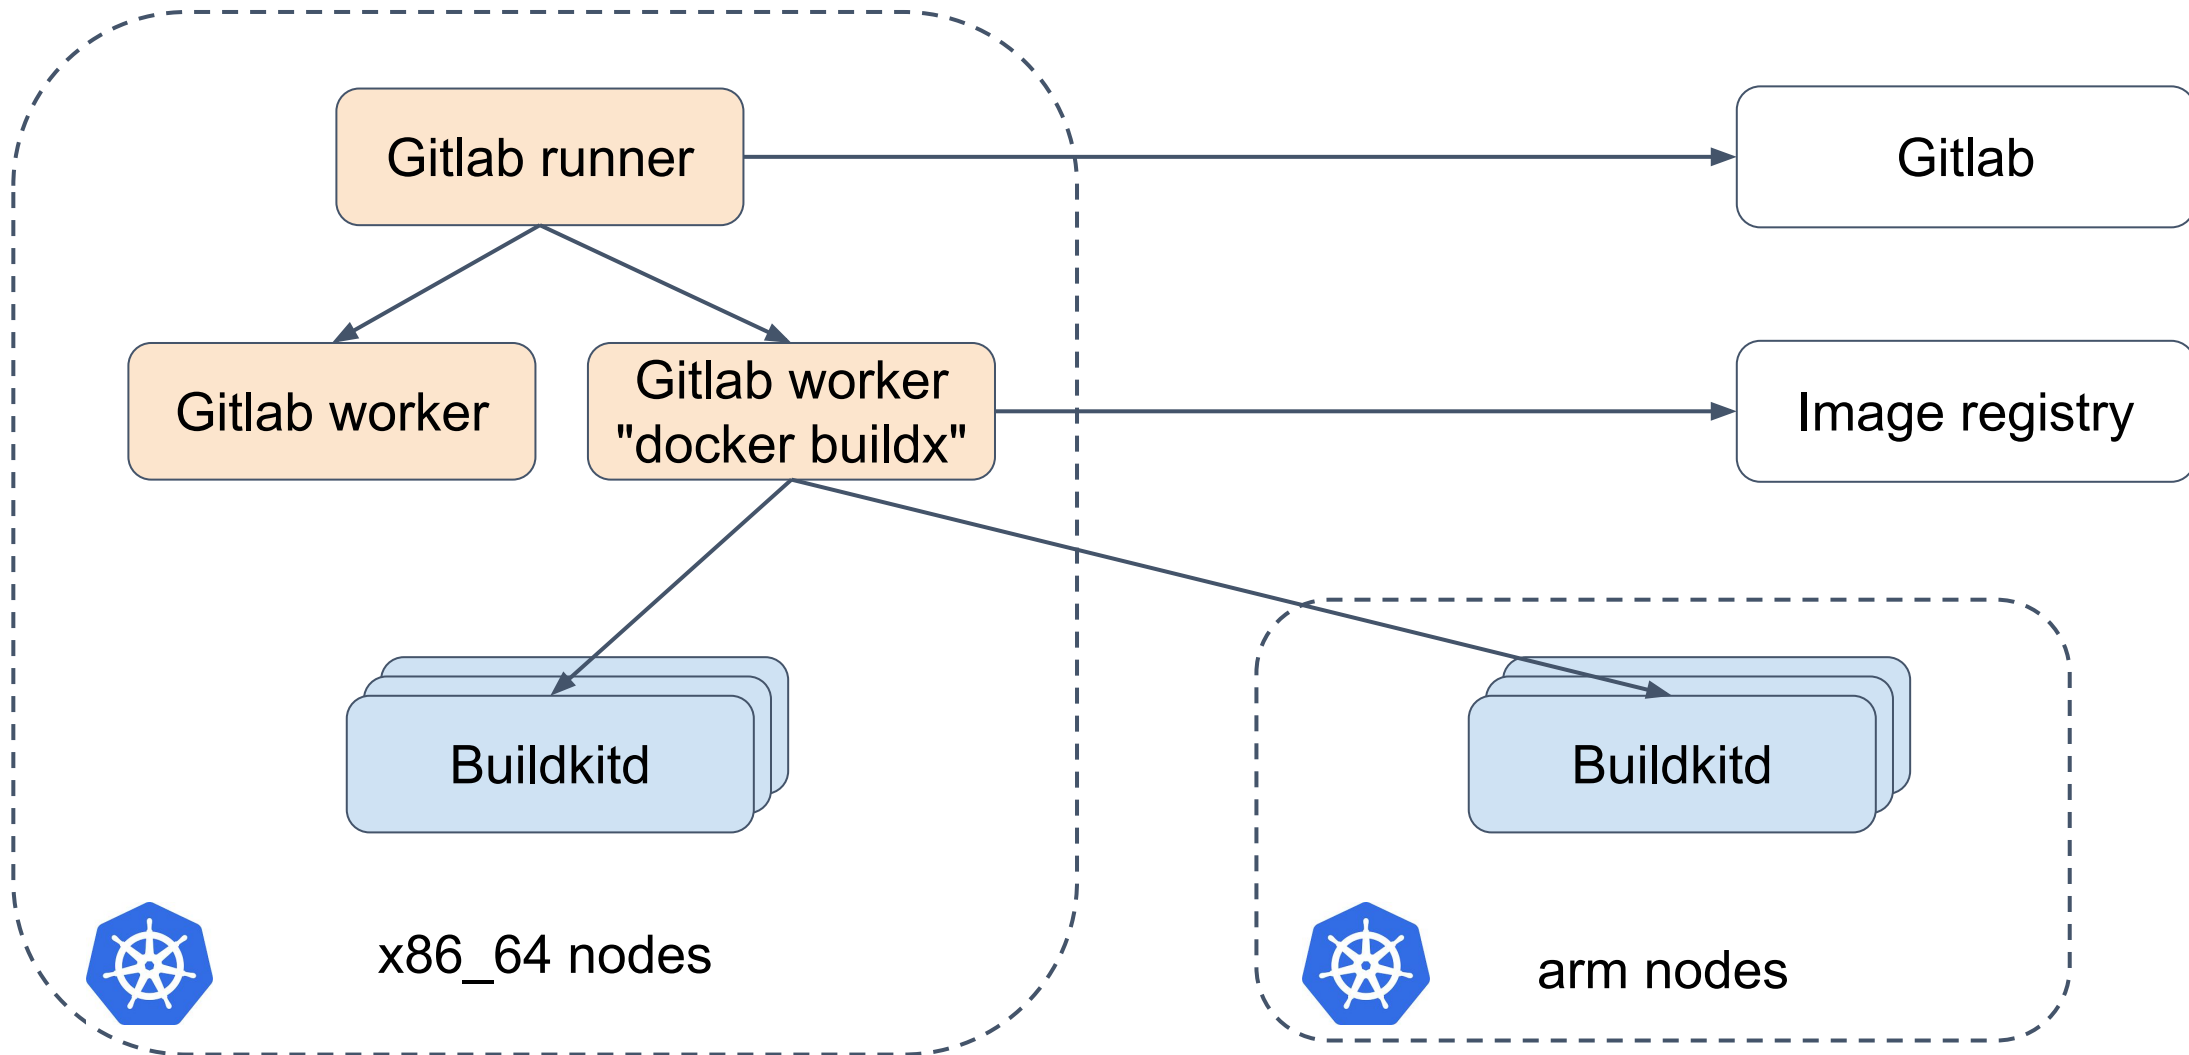
What would it look like?



What would it look like?



What would it look like?



Can we make this safe?

- Image builds require privileges
 - e.g. package installation
- Containers should not run as root
- → “rootless” builds

“rootless”?

- User namespaces
 - root user of host != root user in container
- Affects for instance:
 - Capabilities
 - Mount namespaces, VFS, filesystems
 - Linux Security Modules
- [Building images efficiently and securely on Kubernetes with BuildKit](#)
Akihiro Suda, Kubecon EU, 2019

Overview

BuildKit Worker Container

```
rootlesskit [uid=1000]
  \_buildkitd [uid=0*]
    \_buildkit-runc [uid=0*]
      \_<RUN Steps> [uid=0*]
```

User namespace
uid 0 != uid 0 on the host

Example

```
$ unshare --user --map-root-user bash
```

Example

```
$ unshare --user --map-root-user bash

# id
uid=0(root) gid=0(root) groups=0(root)

# touch /etc/x
touch: cannot touch '/etc/x': Permission denied

# touch /tmp/x

# ls -l /tmp/x
-rw-r--r-- 1 root root 0 Oct 11 08:38 /tmp/x
```

Process user
namespace

Example

```
$ unshare --user --map-root-user bash
```

```
# id
```

```
uid=0(root) gid=0(root) groups=0(root)
```

```
# touch /etc/x
```

```
touch: cannot touch '/etc/x': Permission denied
```

```
# touch /tmp/x
```

```
# ls -l /tmp/x
```

```
-rw-r--r-- 1 root root 0 Oct 11 08:38 /tmp/x
```

```
# exit
```

```
$ ls -l /tmp/x
```

```
-rw-r--r-- 1 1000 1000 0 Oct 11 08:38 /tmp/x
```

Process user
namespace

Before we engage

- Most images built fine on buildkit/kube from the start
- We'll only talk about the tricky bits though...

Before we engage

- Most images built fine on buildkit/kube from the start
- We'll only talk about the tricky bits though...
- **Where's the "fun" otherwise?**

Issue #1: Why is RUN broken?

Time to Fun*: a few hours

*** Time between "Wait what???" and "Oh that's fun"**

A complex Dockerfile

```
FROM gitlab/gitlab-runner-helper:arm64-v14.0.0  
RUN echo "test"
```

A complex Dockerfile

```
FROM gitlab/gitlab-runner-helper:arm64-v14.0.0
RUN echo "test"
```

```
$ docker buildx ...
```

```
#5 [2/2] RUN echo "test"
```

```
#5 ERROR: mount callback failed on  
/run/user/1000/containerd-mount4xx: operation not permitted
```

A complex Dockerfile

```
FROM gitlab/gitlab-runner-helper:arm64-v14.0.0
RUN echo "test"
```

```
$ docker buildx ...

#5 [2/2] RUN echo "test"
#5 ERROR: mount callback failed on
/run/user/1000/containerd-mount4xx: operation not permitted
```

Surprising, right?

Which Operation is Denied?

```
# strace -f -p <buildkitd_pid> -e trace=%file |& grep EPERM
```

```
... lsetxattr(".",  
            "security.selinux",  
            "system_u:..."...,  
            33, 0)  
    = -1 EPERM (Operation not permitted)
```

It's always ~~DNS~~ SELinux!

Delving Into Image Layer

```
00001000: 3637 2053 4348 494c 592e 7861 7474 722e 67 SCHILY.xattr.  
00001010: 7365 6375 7269 7479 2e73 656c 696e 7578 security.selinux  
00001020: 3d73 7973 7465 6d5f 753a 6f62 6a65 6374 =system_u:object  
00001030: 5f72 3a75 6e6c 6162 656c 6564 5f74 3a73 _r:unlabeled_t:s  
00001040: 3000 0a00 0000 0000 0000 0000 0000 0000 0.....
```

⇒ Files in tarball have SELinux extended attributes

SELinux in User Namespaces

```
# cd /tmp  
# touch x  
# chcon system_u:object_r:unlabeled_t:s0 x
```

} Host
User
NS

```
# nsenter -t <buildkitd pid> -U  
# touch y  
# chcon system_u:object_r:unlabeled_t:s0 y  
chcon: failed to change context ...: Operation not permitted
```

} Process
User
NS

Kernel disallows setting SELinux context because
root in user namespace \neq root in initial user namespace

- Upstream issue: <https://github.com/moby/buildkit/issues/2407>
- No real solution but easy enough to mitigate
 - Pull / Push to create an image without SELinux labels
 - Use an image without SELinux labels (gitlab-runner-helper:v14.1.0)

Issue #2: Who stole my port?

Time to Fun: a few days

A more complex Dockerfile

```
RUN curl -OL https://.../some_app.deb \  
    && dpkg -i some_app.deb \  
    && rm some_app.deb
```

- Build times out

A more complex Dockerfile

```
RUN curl -OL https://.../some_app.deb \  
    && dpkg -i some_app.deb \  
    && rm some_app.deb
```

- Build times out
- Let's retry!

A more complex Dockerfile

```
RUN curl -OL https://.../some_app.deb \  
    && dpkg -i some_app.deb \  
    && rm some_app.deb
```

- Build times out
- Let's retry!
- RUN fails but **with an error this time**

```
#2 [2/2] RUN curl...
```

```
#2 ERROR: Bind failed: Address already in use
```

A more complex Dockerfile

```
RUN curl -OL https://.../some_app.deb \  
    && dpkg -i some_app.deb \  
    && rm some_app.deb
```

- Build times out
- Let's retry!
- Let's be scientific and retry **again**

A more complex Dockerfile

```
RUN curl -OL https://.../some_app.deb \  
    && dpkg -i some_app.deb \  
    && rm some_app.deb
```

- Build times out
- Let's retry!
- Let's be scientific and retry **again**

```
#2 [2/2] RUN curl...
```

```
#2 ERROR: Bind failed: Address already in use
```

A more complex Dockerfile

```
RUN curl -OL https://.../some_app.deb \  
    && dpkg -i some_app.deb \  
    && rm some_app.deb
```

- Build times out
- Let's retry!
- Let's be scientific and retry again
- What if we delete the buildkitd pod?

A more complex Dockerfile

```
RUN curl -OL https://.../some_app.deb \  
    && dpkg -i some_app.deb \  
    && rm some_app.deb
```

- Build times out
- Let's retry!
- Let's be scientific and retry again
- What if we delete the buildkitd pod?
- Build times out

A more complex Dockerfile

```
RUN curl -OL https://.../some_app.deb \  
    && dpkg -i some_app.deb \  
    && rm some_app.deb
```

- Build times out
- Let's retry!
- Let's be scientific and retry again
- What if we delete the buildkitd pod?
- Build times out
- Following builds get the error

```
#2 [2/2] RUN curl...
```

```
#2 ERROR: Bind failed: Address already in use
```

Let's debug!

```
RUN netstat -tunlp &&  
    curl -OL https://.../some_app.deb \  
        && dpkg -i some_app.deb \  
        && rm some_app.deb \  
        && netstat -tunlp
```

With a brand new buildkitd pod

- **1st** netstat: no port bound
- **2nd** netstat: **"some_app"** port is bound by **some_appd**
- Build hangs

Let's debug!

```
RUN netstat -tunlp &&  
    curl -OL https://.../some_app.deb \  
        && dpkg -i some_app.deb \  
        && rm some_app.deb \  
        && netstat -tunlp
```

Next build

- **1st netstat: "some_app" port is bound by some_appd**
- **dpkg -i fails with "Address already in use"**

What's happening?

```
RUN netstat -tunlp &&  
    curl -OL https://.../some_app.deb \  
        && dpkg -i some_app.deb \  
        && rm some_app.deb
```

=> some_app.deb starts a daemon on install

=> Looks like the daemon is still running when we do build #2

What's happening?

```
RUN netstat -tunlp &&  
    curl -OL https://.../some_app.deb \  
        && dpkg -i some_app.deb \  
        && rm some_app.deb
```

=> some_app.deb starts a daemon on install

=> Looks like the daemon is still running when we do build #2

Can we reproduce?

Basic Reproducer

```
FROM ubuntu
ADD script.sh /
RUN /script.sh
RUN echo "Dockerfile done"
```

```
=> [3/4] RUN /script.sh
=> => # + echo 'Script done'
=> => # + sleep infinity
=> => # Script done
...hangs...
```

```
#!/bin/bash
set -x
sleep infinity &
echo "Script done"
```

Basic Reproducer

```
FROM ubuntu
ADD script.sh /
RUN /script.sh
RUN echo "Dockerfile done"
```

```
=> [3/4] RUN /script.sh
=> => # + echo 'Script done'
=> => # + sleep infinity
=> => # Script done
...hangs...
```

```
#!/bin/bash
set -x
sleep infinity &
echo "Script done"
```

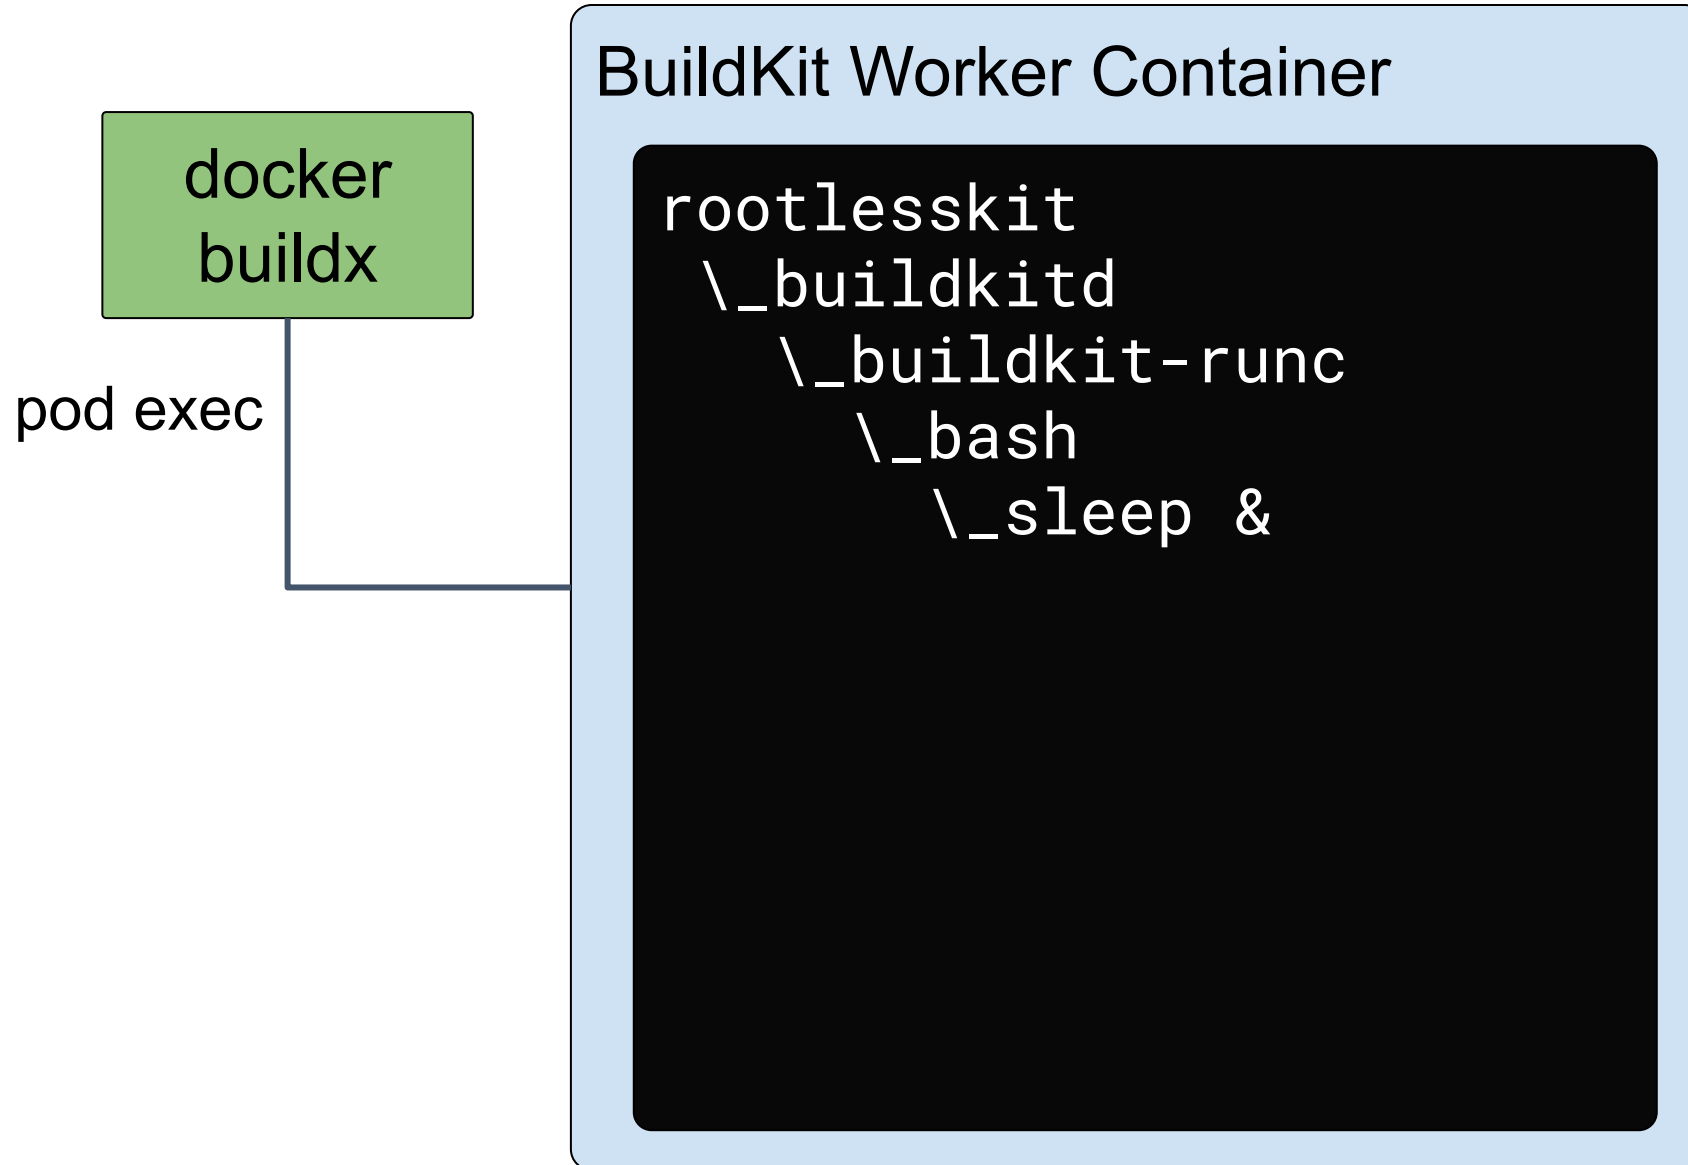
=> Process leaked when build times out or is interrupted

Let's Peek Under the Hood

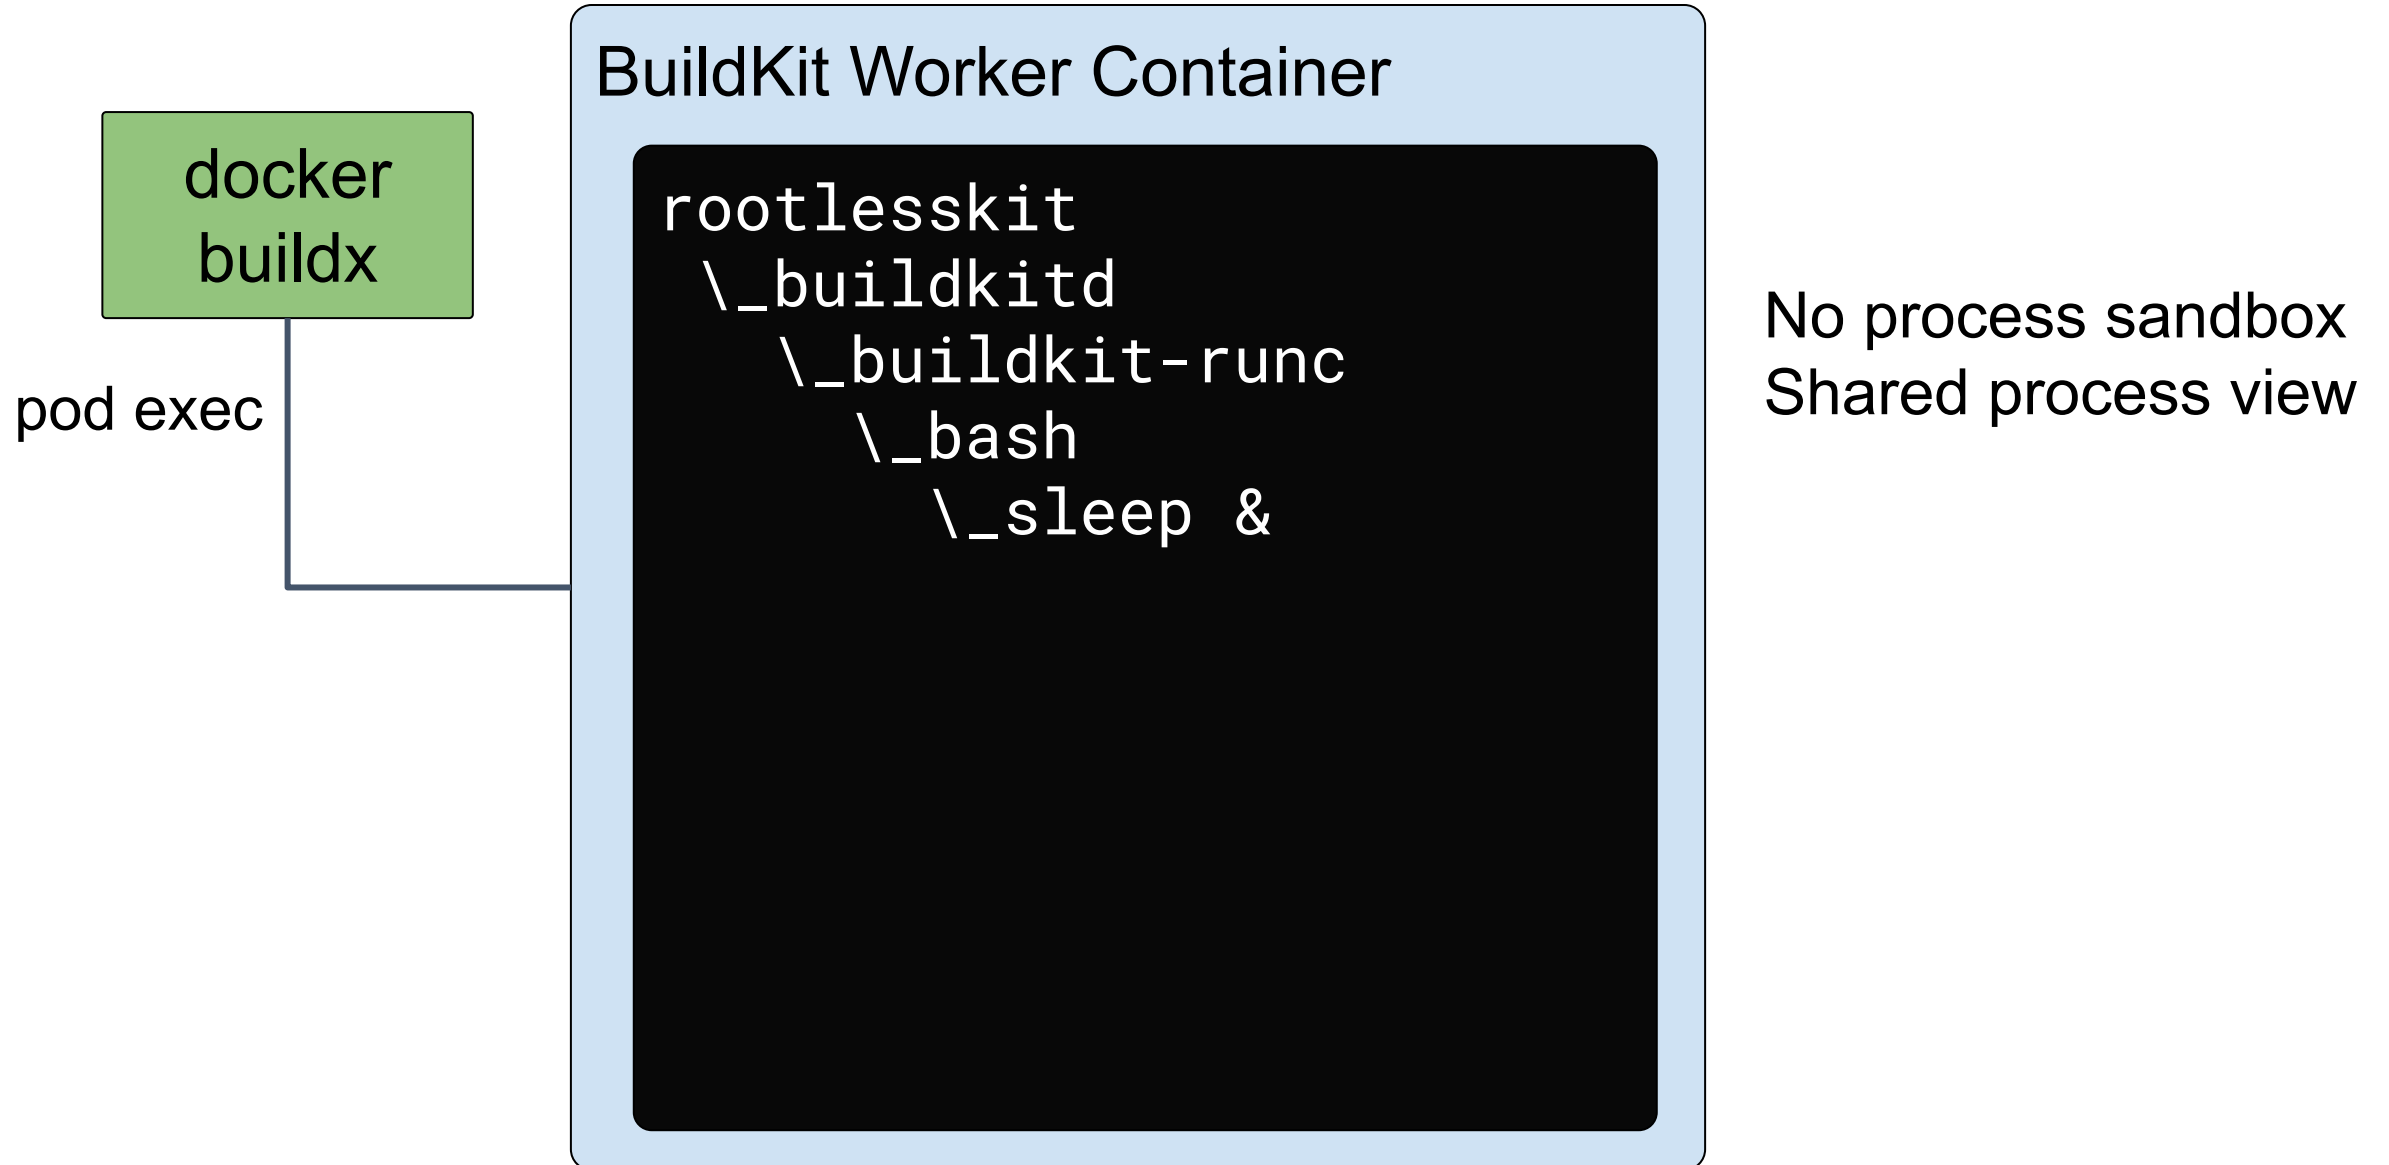
BuildKit Worker Container

```
rootlesskit  
  \_buildkitd
```

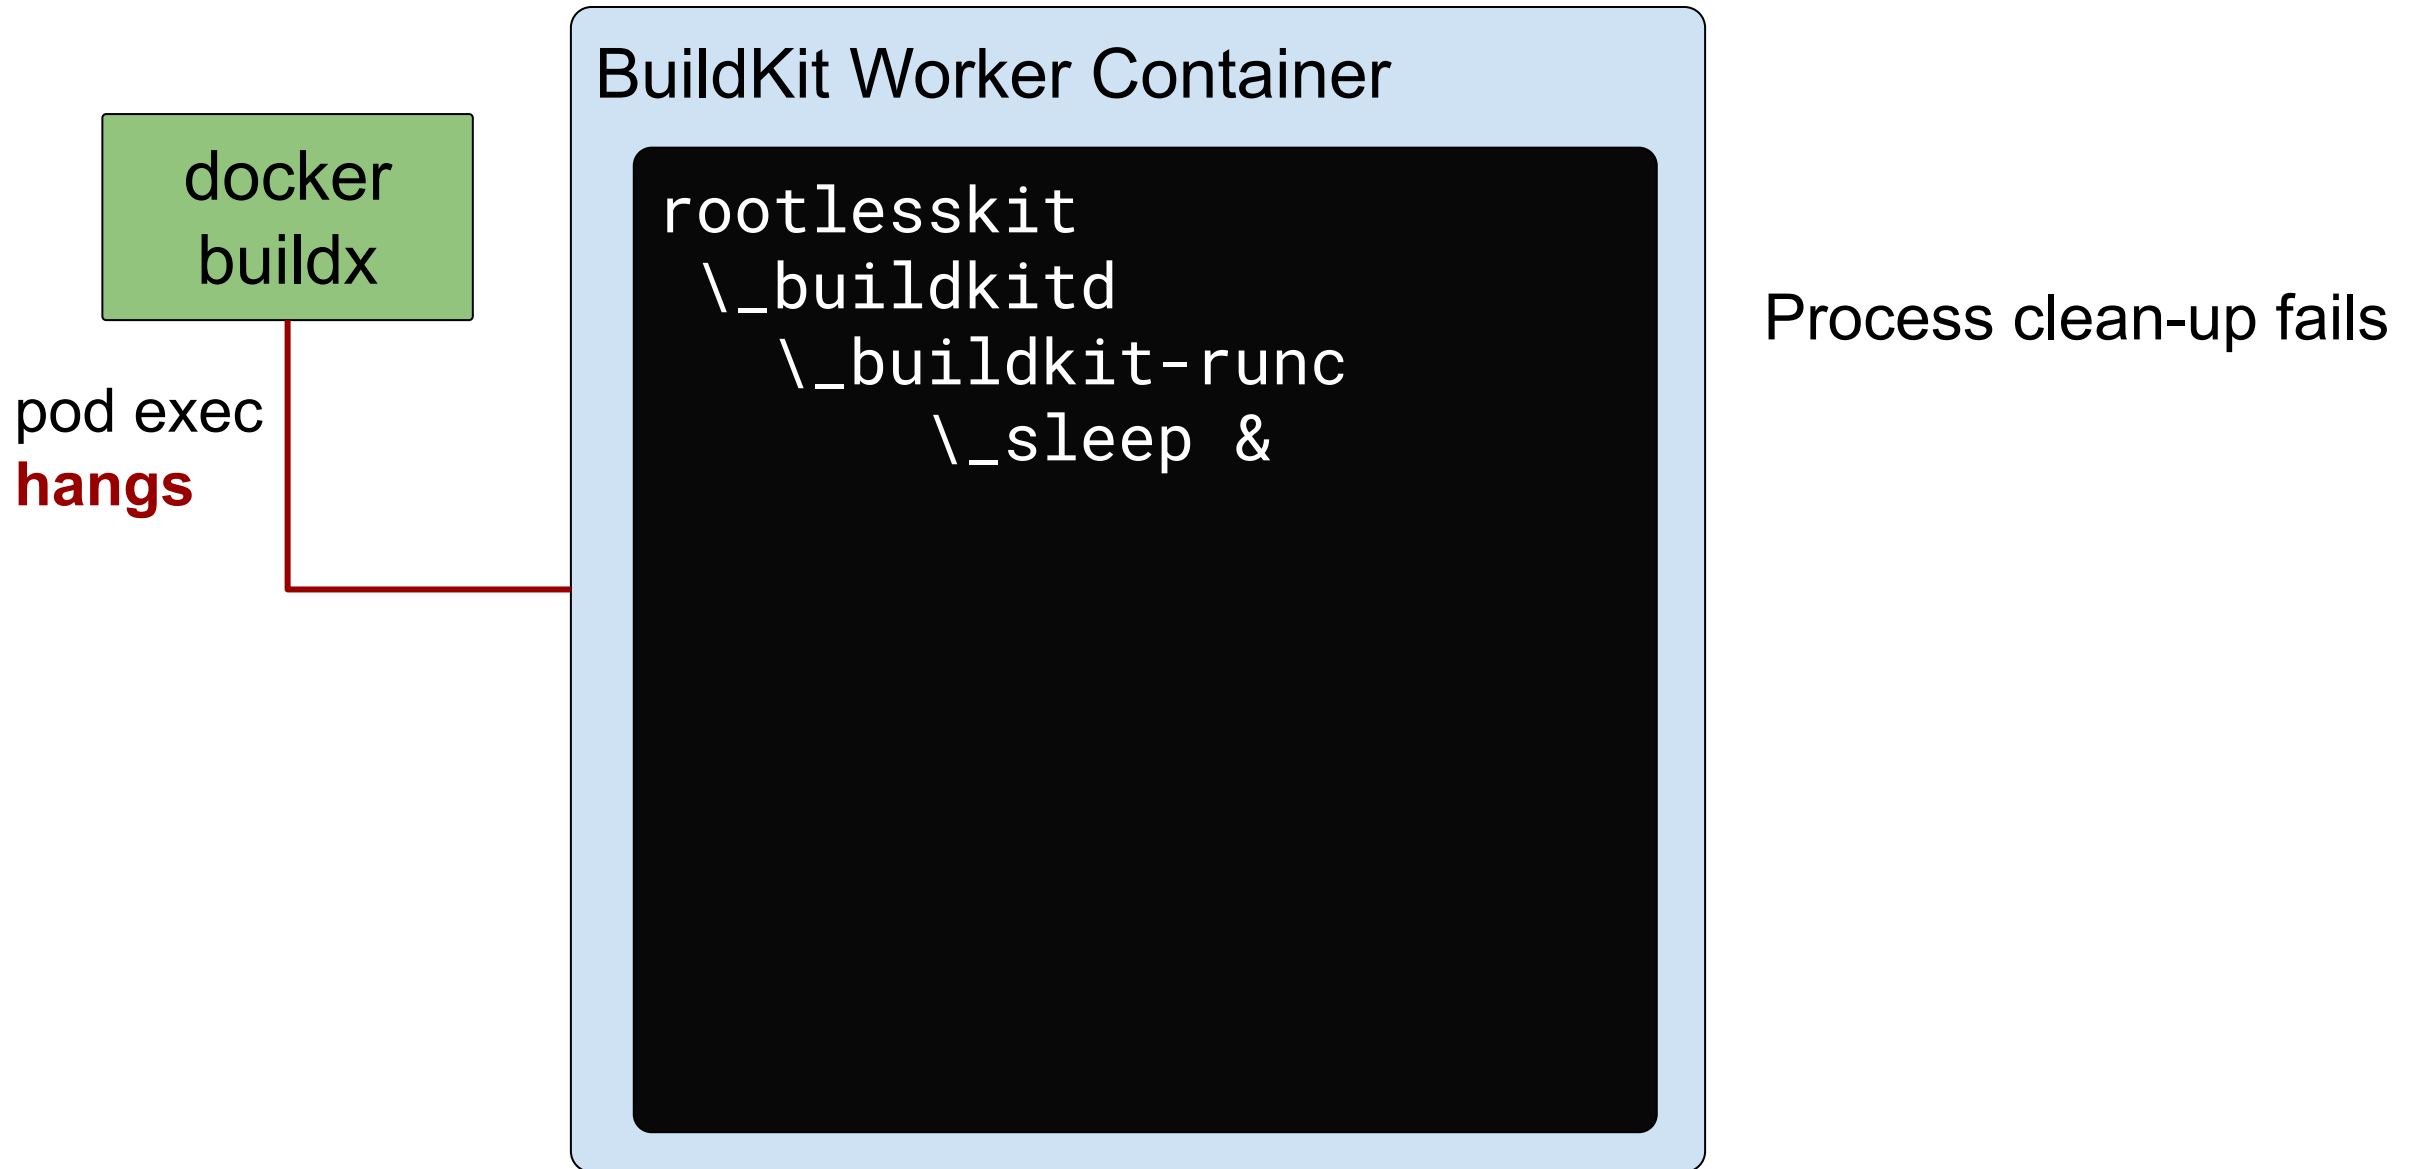
Building our image



Building our image



bash exits



What if we kill sleep?

BuildKit Worker Container

```
rootlesskit  
  \_buildkitd  
  \_sleep <defunct>
```

We get a Zombie

How does it work usually?

- Build steps run in a process sandbox
- When the step finishes, all processes in the sandbox are killed

=> We don't have this in rootless mode (**--oci-worker-no-process-sandbox**)

=> We can't fully keep track of processes started in the build steps

Why do we need this flag?

/proc in containers

buildkit container

```
/proc  
/1  
/2  
/x
```

```
/proc/kcore  
/proc/keys  
...
```

```
/proc/irq  
/proc/sysrq-trigger  
...
```

For security reasons,
we can't fully expose /proc to containers

masked: bind mounted with empty dir
(prevents access to sensitive data)

readonly: bind mounted read-only
(prevents system-wide modifications)

/proc in containers

buildkit container

```
/proc  
/1  
/2  
/x
```

```
/proc/kcore  
/proc/keys  
...
```

```
/proc/irq  
/proc/sysrq-trigger  
...
```

Creating a new procfs in this container fails

=> Kernel check: "mount_too_revealing"
If existing procfs is partially masked, disallow mount

=> Explains need for `--oci-worker-no-process-sandbox`

Details:

<https://github.com/opencontainers/runc/issues/1658>

Conclusion

- No real solution to date: <https://github.com/moby/buildkit/issues/2417>
- Mitigations
 - **Make sure no daemons are started or stop them**
 - Potentially: ProcMountType=Unmasked
 - Use jobs for builds
- Security/stability [implications](#)
 - Leaked processes/zombies
 - Affect subsequent builds (e.g. bound ports)

Issue #3: Ghosts in the filesystem

Time to Fun: more than a week

Let's Build Some Go

```
RUN git clone https://[...]/sig-storage-local-static-provisioner
```

```
RUN git -C ... checkout $TAG
```

```
RUN go build -o ... cmd/local-volume-provisioner/main.go
```

Compiler Error

```
vendor/k8s.io/utils/io/read.go:36:6:  
    ConsistentRead redeclared in this block  
  
previous declaration at  
    vendor/k8s.io/utils/io/consistentread.go:28:61
```

Compiler Error

```
vendor/k8s.io/utils/io/read.go:36:6:  
    ConsistentRead redeclared in this block  
  
previous declaration at  
    vendor/k8s.io/utils/io/consistentread.go:28:61
```

This makes no sense:

Same Dockerfile builds fine with docker

Why would rootless buildkit trigger this???

Let's Check the Directory

```
#3 RUN ls -l vendor/k8s.io/utils/io/  
[...]
```

```
#3 0.121 ... Sep 25 09:58 consistentread.go
```

```
#3 0.121 ... Sep 25 09:58 read.go
```

Let's Check the Directory

```
#3 RUN ls -l vendor/k8s.io/utils/io/  
[...]  
#3 0.121 ... Sep 25 09:58 consistentread.go  
#3 0.121 ... Sep 25 09:58 read.go
```

In master, directory only contains "read.go"

In tag, directory only contains "consistentread.go"

How can we have both files in this directory?

Check Each Layer

```
#1 RUN git clone ... && ls -l .../utils/io/
```

```
#1 ... Sep 25 09:58 read.go
```

Expected

```
#2 RUN git checkout ... && ls -l .../utils/io/
```

```
#2 ... Sep 25 09:58 consistentread.go
```

Expected

```
#3 RUN ls -l vendor/k8s.io/utils/io/
```

```
#3 ... Sep 25 09:58 consistentread.go
```

```
#3 ... Sep 25 09:58 read.go
```

Wait. What?

Check Each Layer

```
#1 RUN git clone ... && ls -l ../utils/io/
```

```
#1 ... Sep 25 09:58 read.go
```

Expected

```
#2 RUN git checkout ... && ls -l ../utils/io/
```

```
#2 ... Sep 25 09:58 consistentread.go
```

Expected

```
#3 RUN ls -l vendor/k8s.io/utils/io/
```


```
#3 ... Sep 25 09:58 consistentread.go
```

```
#3 ... Sep 25 09:58 read.go
```

Wait. What?

Something is very wrong with our filesystem

We use Overlayfs. How does it work?



Reproducing the buildkit steps

```
$ unshare --mount --user --map-root-user bash  
# mkdir layer1 layer2 l2diff work
```

Init build env

```
# git clone <...> layer1
```

Run step #1

```
# mount -t overlay overlay \  
    -olowerdir=layer1,upperdir=l2diff,workdir=work \  
    layer2
```

Prepare step #2

```
# git -C layer2 checkout ...
```

Run step #2

At this point everything looks ok

```
# ls -l layer1/vendor/k8s.io/utils/io  
... Oct 10 15:08 read.go
```

```
# ls -l l2diff/vendor/k8s.io/utils/io  
... Oct 10 15:09 consistentread.go
```

```
# ls -l layer2/vendor/k8s.io/utils/io  
... Oct 10 15:09 consistentread.go
```

What about step #3?

```
# umount layer2
# mkdir l3diff layer3 work3

# mount -t overlay overlay \
    -olowerdir=l2diff,layer1,upperdir=l3diff,workdir=work3 \
    layer3

# ls layer3
... Oct 10 15:08 read.go
... Oct 10 15:09 consistentread.go
```

Prepare step #3

Run step #3

Definitely broken, we have reproduced!

Taking a step back

layer 1

read.go

Step #1: git clone

Taking a step back

mount

layer2

consistentread.go

upper

l2diff

consistentread.go

lower

layer 1

read.go

Step #2: git checkout

Taking a step back



Step #3: ls

Something feels wrong with l2diff

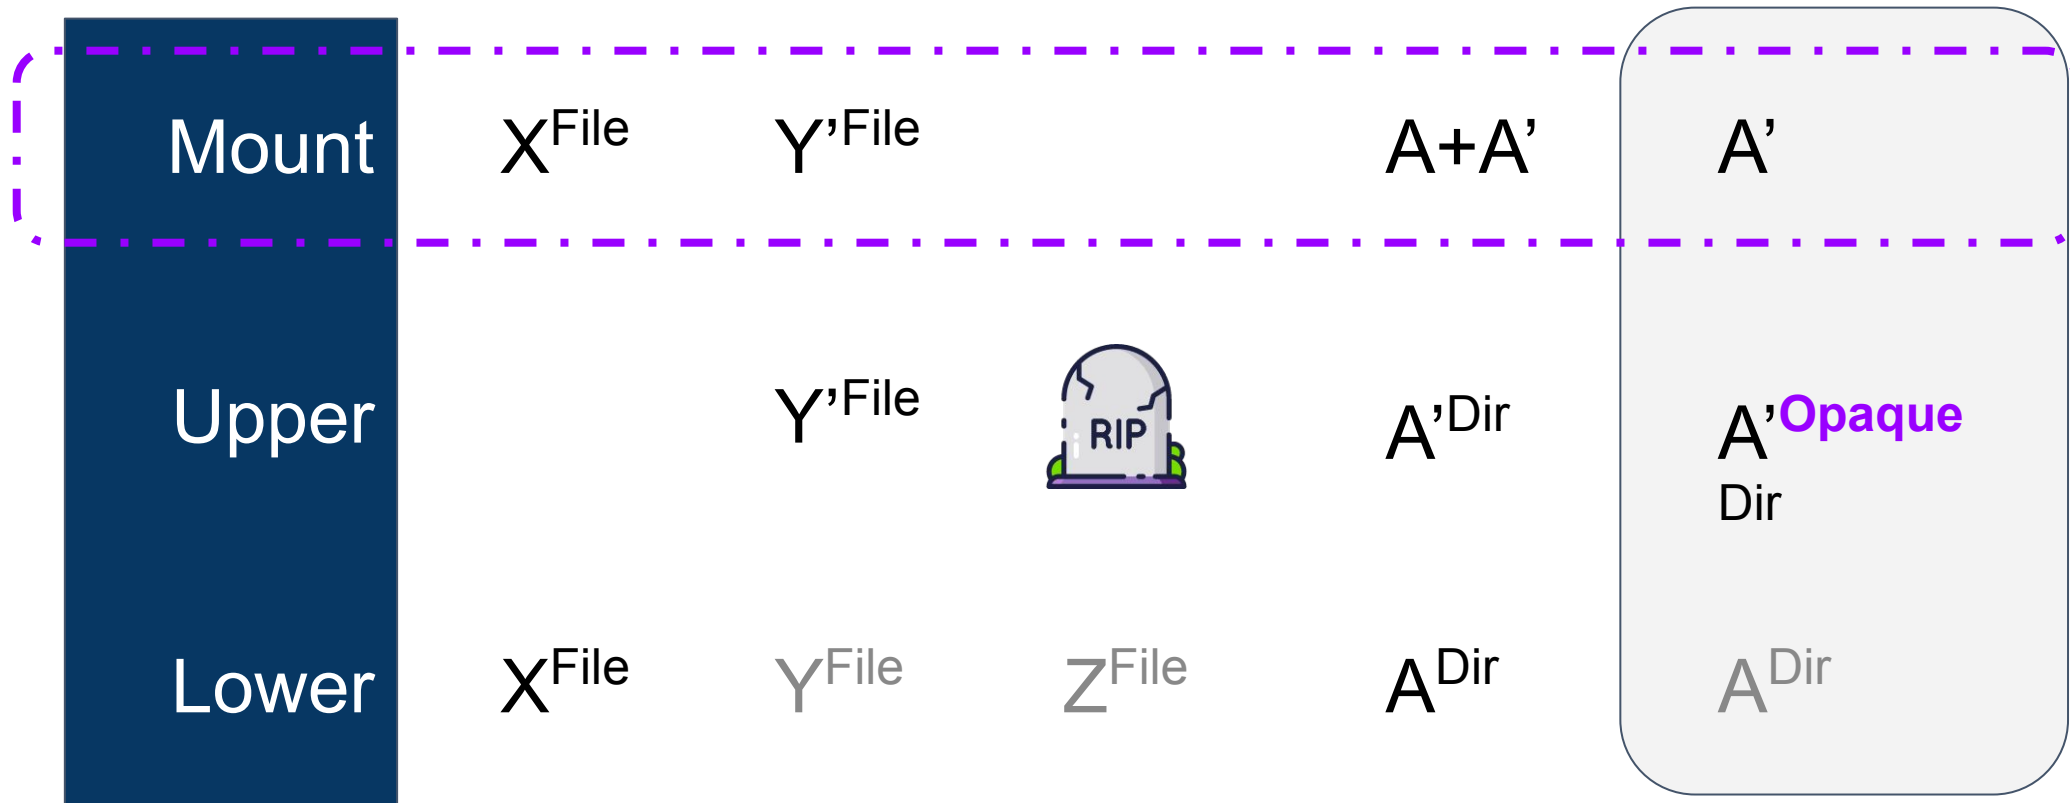


Looking more closely at I2diff



Where is the tombstone for read.go?

Overlayfs: we missed a bit



Opaque directory: mask directory from lower layers

=> Faster and simpler

Opaque directory

mount

layer2

consistentread.go

upper

l2diff

consistentread.go
/utils/io : **Opaque**

lower

layer 1

read.go

Step #2: git checkout

Opaque directory implementation

Uses directory xattr: **trusted.overlay.opaque**

Opaque directory implementation

Uses directory xattr: **trusted.overlay.opaque**

Can we see it?

```
[in user namespace]  
  
# getfattr -R -d -m "" 12diff/.../utils/io  
[nothing]
```

How is this working at step #2?

Opaque directory implementation

Uses directory xattr: **trusted.overlay.opaque**

Can we see it in the initial user namespace?

```
[in initial user namespace]
```

```
$ sudo getfattr -R -d -m "" 12diff/.../utils/io  
trusted.overlay.opaque="y"
```

This makes more sense

Opaque directory implementation

Uses directory xattr: **trusted.overlay.opaque**

Can we see it in the initial user namespace?

```
[in initial namespace]
```

```
$ sudo getfattr -R -d -m "" 12diff/.../utils/io  
trusted.overlay.opaque="y"
```

This makes more sense

But we should not be able to write **trusted.*** xattr from user namespace

Mysteries

- `trusted.overlay.opaque` set despite lacking `CAP_SYS_ADMIN`
- It seems that the issue
 - does **not** reproduce when the opaque directory is “**upper**” (step #2)
 - does reproduce when the opaque directory is a “**lower**” (step #3)

Kernel function tracing

Using ftrace to look at the **git checkout** operation

```
# trace-cmd record -p function -P <PID> -c -e all  
  
    ovl_check_setxattr  
        ...  
        __vfs_setxattr_noperm  
            __vfs_setxattr  
                xattr_resolve_name  
                ext4_xattr_trusted_set  
                ...
```

vfs_setxattr_noperm

Specific Ubuntu patch to allow the use of overlayfs in user namespaces

```
UBUNTU: SAUCE: overlayfs: Skip permission checking for trusted.overlayfs.* xattrs
```

The original mounter had CAP_SYS_ADMIN in the user namespace where the mount happened, and the vfs has validated that the user has permission to do the requested operation. This is sufficient for allowing the kernel to write these specific xattrs, so **we can bypass the permission checks for these xattrs.**

To support this, export `__vfs_setxattr_noperm` and add an similar `__vfs_removexattr_noperm` which is also exported. Use these when **setting or removing trusted.overlayfs.* xattrs.**

vfs_setxattr_noperm

Specific Ubuntu patch to allow the use of overlayfs in user namespaces

```
UBUNTU: SAUCE: overlayfs: Skip permission checking for trusted.overlayfs.* xattrs
```

The original mounter had CAP_SYS_ADMIN in the user namespace where the mount happened, and the vfs has validated that the user has permission to do the requested operation. This is sufficient for allowing the kernel to write these specific xattrs, so **we can bypass the permission checks for these xattrs.**

To support this, export `__vfs_setxattr_noperm` and add an similar `__vfs_removexattr_noperm` which is also exported. Use these when **setting or removing trusted.overlayfs.* xattrs.**

But no equivalent patch for reads!

Back to layers



We mount layer3 in userns
We can't read the Opaque xattr

But wait, it works for step #2

mount

layer2

consistentread.go

upper

l2diff

consistentread.go
/utils/io : **Opaque**

lower

layer 1

read.go

We can write the xattr thanks to the noperm patch
But we should not be able to read it

But wait, it works for step #2



We can write the xattr thanks to the noperm patch
But we should not be able to read it

Filesystems cache heavily, could it be related?

Dropping the Cache

```
$ unshare --mount --user --map-root-user bash
# mkdir layer1 layer2 l2diff work
# git clone <...> layer1
```

Step #1

```
# mount -t overlay overlay \
    -olowerdir=layer1,upperdir=l2diff,workdir=work \
    layer2
# git -C layer2 checkout ...
# ls
consistentread.go
```

Step #2

Dropping the Cache

```
$ unshare --mount --user --map-root-user bash
# mkdir layer1 layer2 l2diff work
# git clone <...> layer1
```

Step #1

```
# mount -t overlay overlay \
    -olowerdir=layer1,upperdir=l2diff,workdir=work \
    layer2
```

Step #2

```
# git -C layer2 checkout ...
```

```
# ls
```

```
consistentread.go
```

```
[In initial user namespace]
```

```
$ sudo /bin/sh -c 'echo 3 > /proc/sys/vm/drop_caches'
```

```
# ls
```

```
consistentread.go
```

```
read.go
```

Summary

- Ubuntu added a patch to make overlayfs work in user namespaces
 - But only patched setxattr and not getxattr
 - Thanks to caching, it works until we have to read from disk
- => Leads to *interesting* behaviors

- New overlayfs mount option “userxattr” in kernel 5.11
 - "user.overlay.*" xattr namespace
- buildkit overlay implementation adds userxattr when available

=> Upgrading Kubernetes nodes to 5.11 just worked!

Details: <https://github.com/moby/buildkit/issues/2381>


Results!

Image Builds on Kube

- Starting with a mono-repo cleared most issues upfront
- Decommissioned dedicated Docker runners
 - Easier node-lifecycle management and tuning
- Native multi-arch builds
 - Emulation is (of course) simply too slow

Image Builds on Kube

- Several hundred distinct images now built on kube
- Handling >1k builds per day reliably



Takeaways

- Buildkit is great
 - Remote builds
 - Multiarch images
- Rootless Buildkit on Kube is bleeding edge
 - Overlayfs work fine in user namespaces with 5.11+
 - No great solution for process sandboxing
 - Maybe with "ProcMount: Unmasked"?
- After the initial hurdles it was really worth the effort for us!

We have many people to thank!

- Mayeul Blanzat for moving image builds to Kubernetes at Datadog
- Tõnis Tiigi and Akihiro Suda for
 - Making rootless buildkit possible
 - Answering our *many* questions
- Jess Frazelle and Alban Crequy for helping us make sense of procs



Please scan the QR Code above to
leave feedback on this session

Thank you

We're hiring!

<https://www.datadoghq.com/careers/>

eric.mountain@datadoghq.com

laurent@datadoghq.com

 **@ericmountain**
 **@lbernail**