

# Building Multi-tenant Routing and Scaling with Envoy



BUILDING FOR THE ROAD AHEAD

**DETROIT 2022**



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

**DETROIT 2022**

**October 24-28, 2022**



**Yiming Peng**

Senior SDE, TLM,  
AWS






KUBECON + CLOUDNATIVECON NORTH AMERICA 2022

# Building multi-tenant routing and scaling with Envoy

**Yiming Peng**

Senior Software Development Engineer  
AWS



# About me

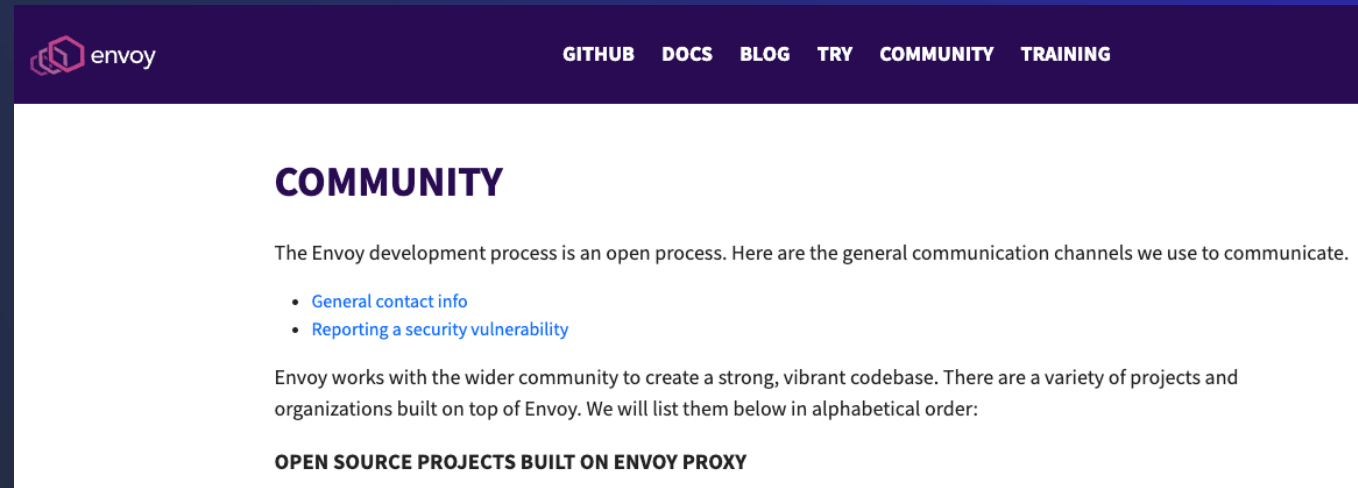
- Senior Software Development Engineer at AWS based in Seattle, WA
- Founding engineer of AWS App Runner – built stealth-mode product from 0 to 1
- Products working on: AWS App Runner, Amazon ECS, AWS Fargate, AWS Elastic Beanstalk
- Expertise areas: Cloud native, containers, serverless, open source
- Open-source enthusiast: Founder and maintainer of **CloudNative & Serverless Meetup**: [github.com/CloudNative-Serverless-Meetup](https://github.com/CloudNative-Serverless-Meetup)
- Contacts: @pymhq (GitHub and Twitter)



**AWS Fargate**

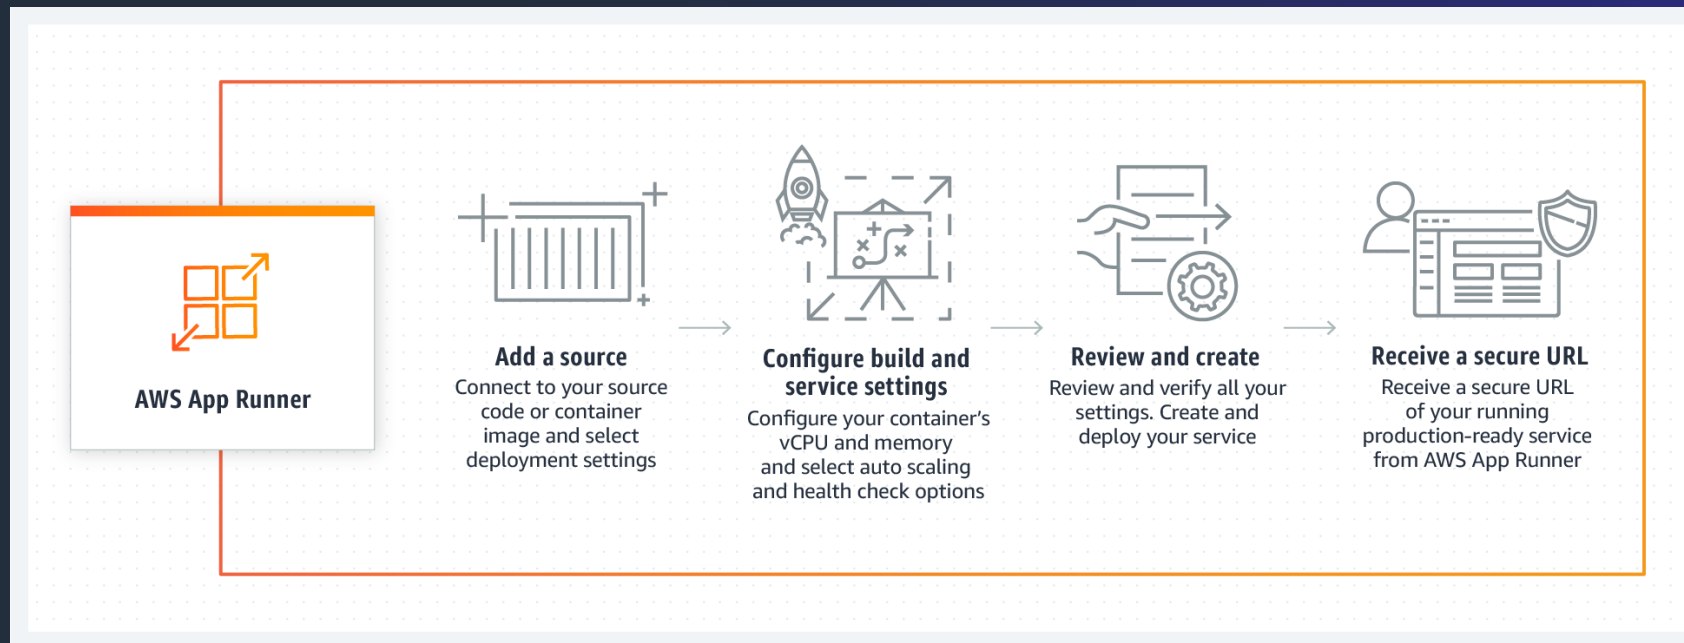
# Motivation

- Share the journey from the end user point of view for Envoy
- Give back to community – share experience and lessons learned; hope brings value and helps community growth
- Appreciate support from Envoy community
- Platform for discussion and communication




# About AWS App Runner

- AWS App Runner (2021 GA) a fully managed service that makes it easy for developers to quickly deploy containerized web applications and APIs – at scale and with no prior infrastructure experience required
- Official website: <https://aws.amazon.com/apprunner/>



# User point of view: App Runner

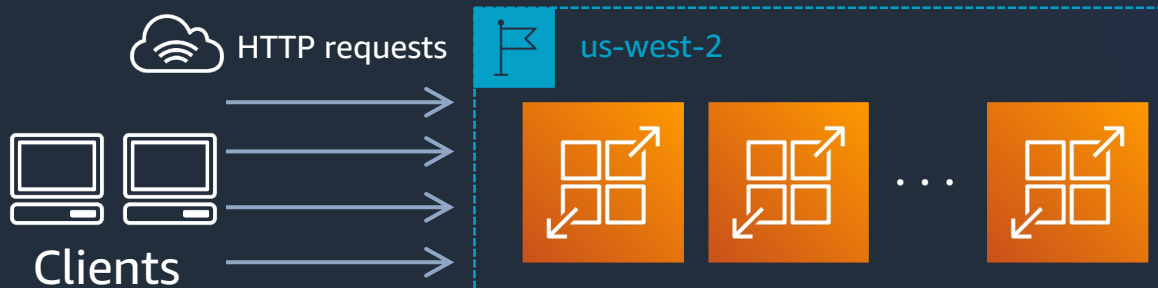
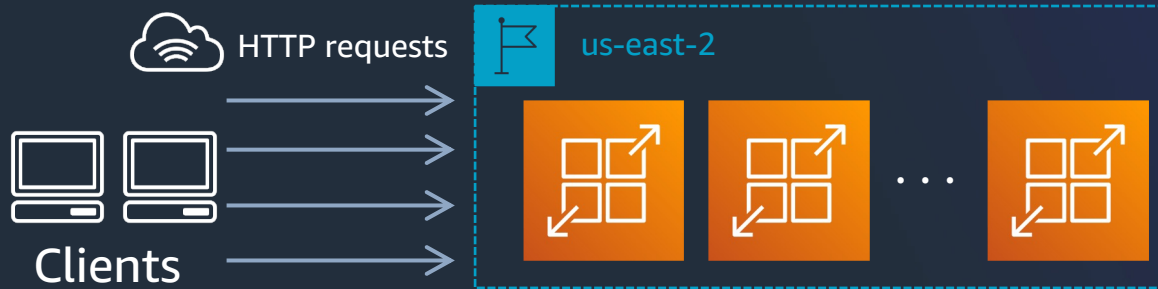
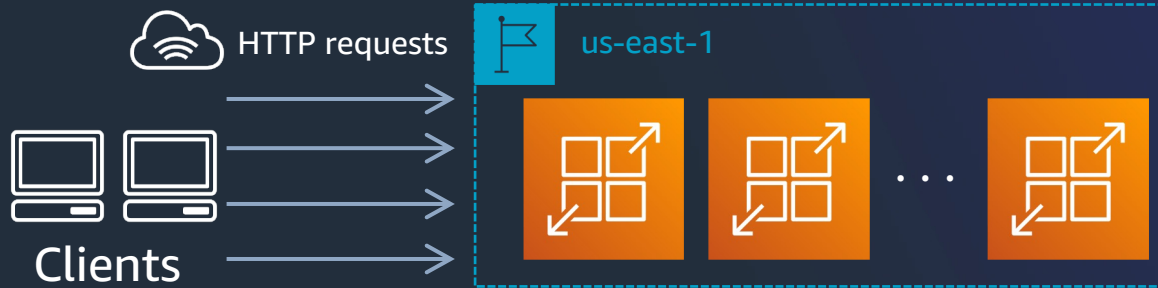


## Managed experience

- App hosting
- Request routing
- Load balancing
- Automatic scaling
- Networking
- Observability
- CI/CD
- Safe deployment (blue-green)
- And more

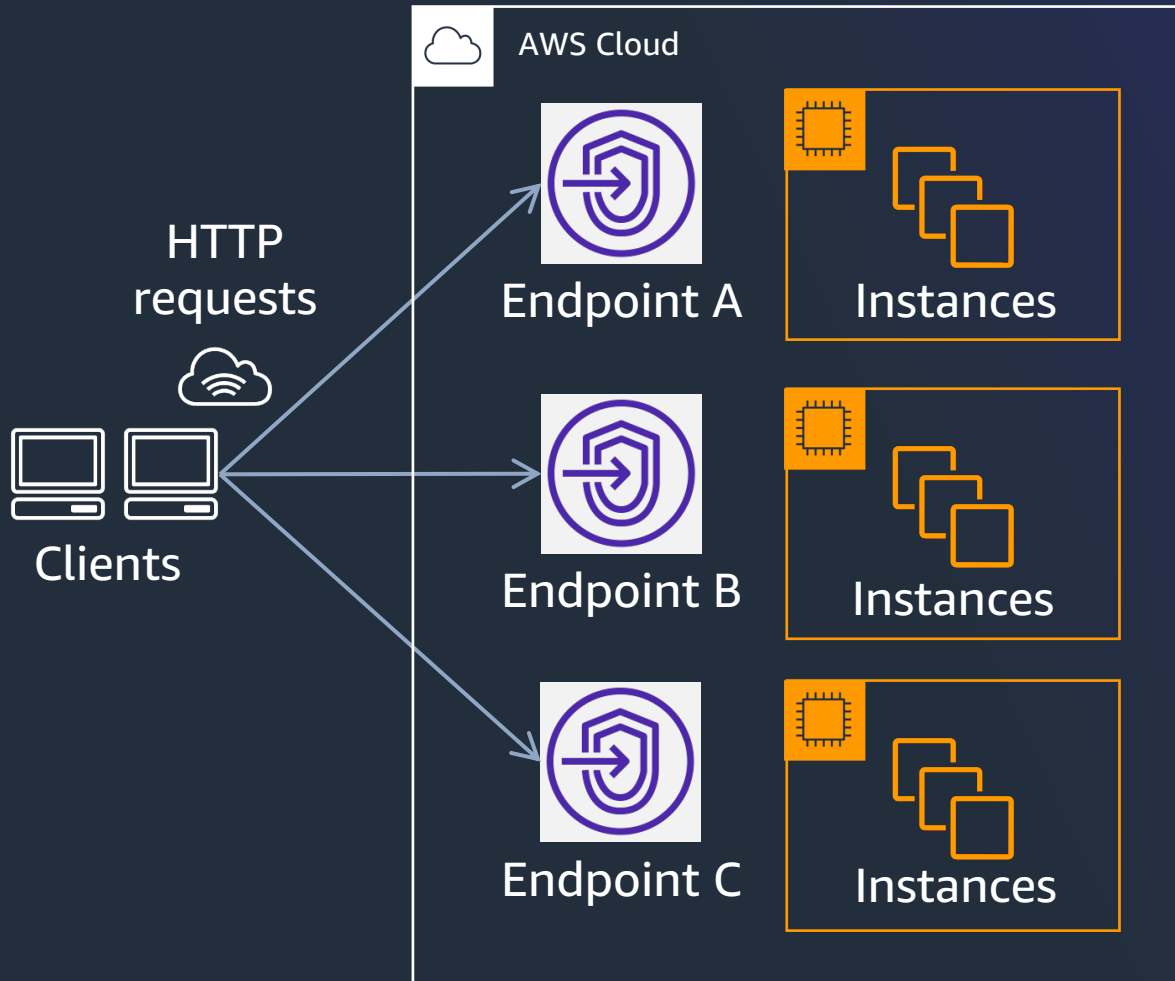


# User point of view: Workloads



- API/web application
- HTTP requests
- Concurrent requests
- Long-time running servers

# Responsibility of request router



To support managed user experience and managed experience on request routing and load balancing, App Runner needs a request router on the backend to help underlying multi-tenant request routing, different AWS Regions, and customer traffic load balancing



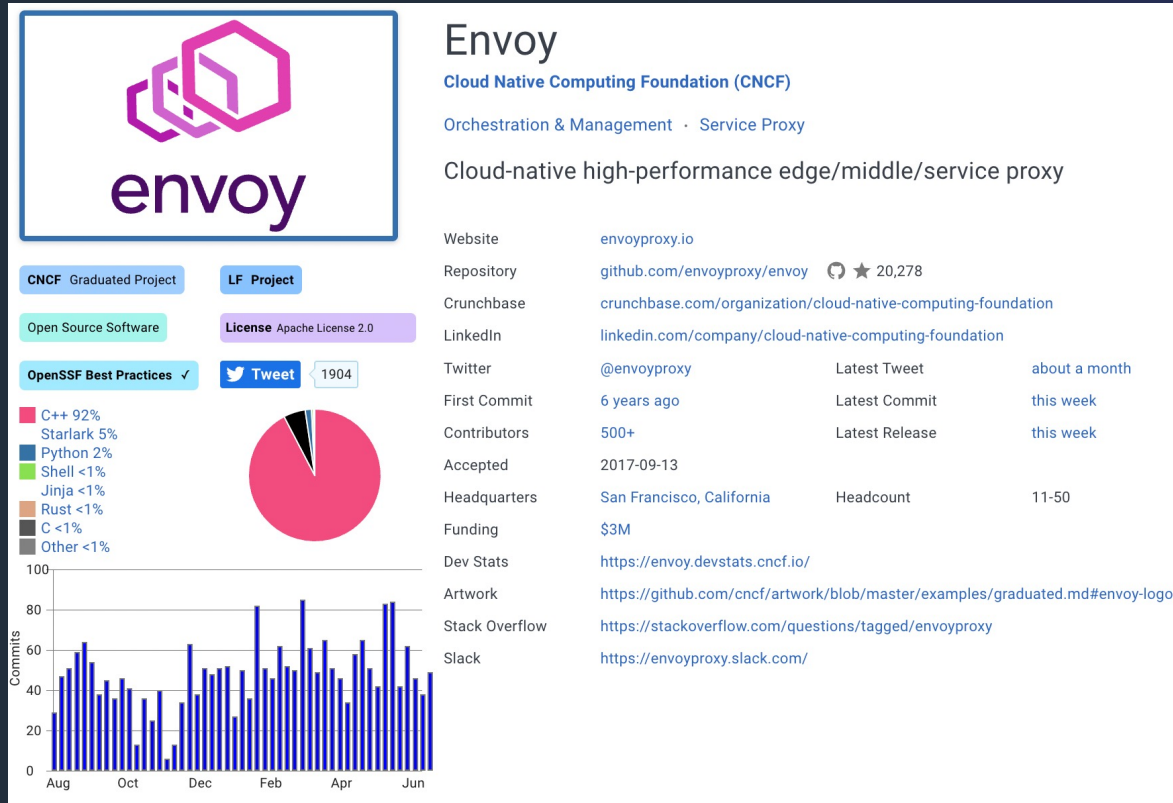
# Product requirement

- Multi-Regions and Region-expansible
- Multi-tenant
- Automatic scaling
- Safe deployment
- Observable

# Technical challenges

- High throughput
- High performance
- Availability
- Reliability
- Extensibility
- Security
- Large number of concurrent connections
- Observability

# Envoy-based request router




<https://www.envoyproxy.io/>

Envoy is an open-source edge and service proxy, designed for cloud-native applications

## Features

- Out-of-process architecture (lightweight and portable)
- HTTP L7 routing support
- gRPC ready
- Best-in-class observability
- Microservice friendly
- Managed open-source version available (AWS AppMesh)

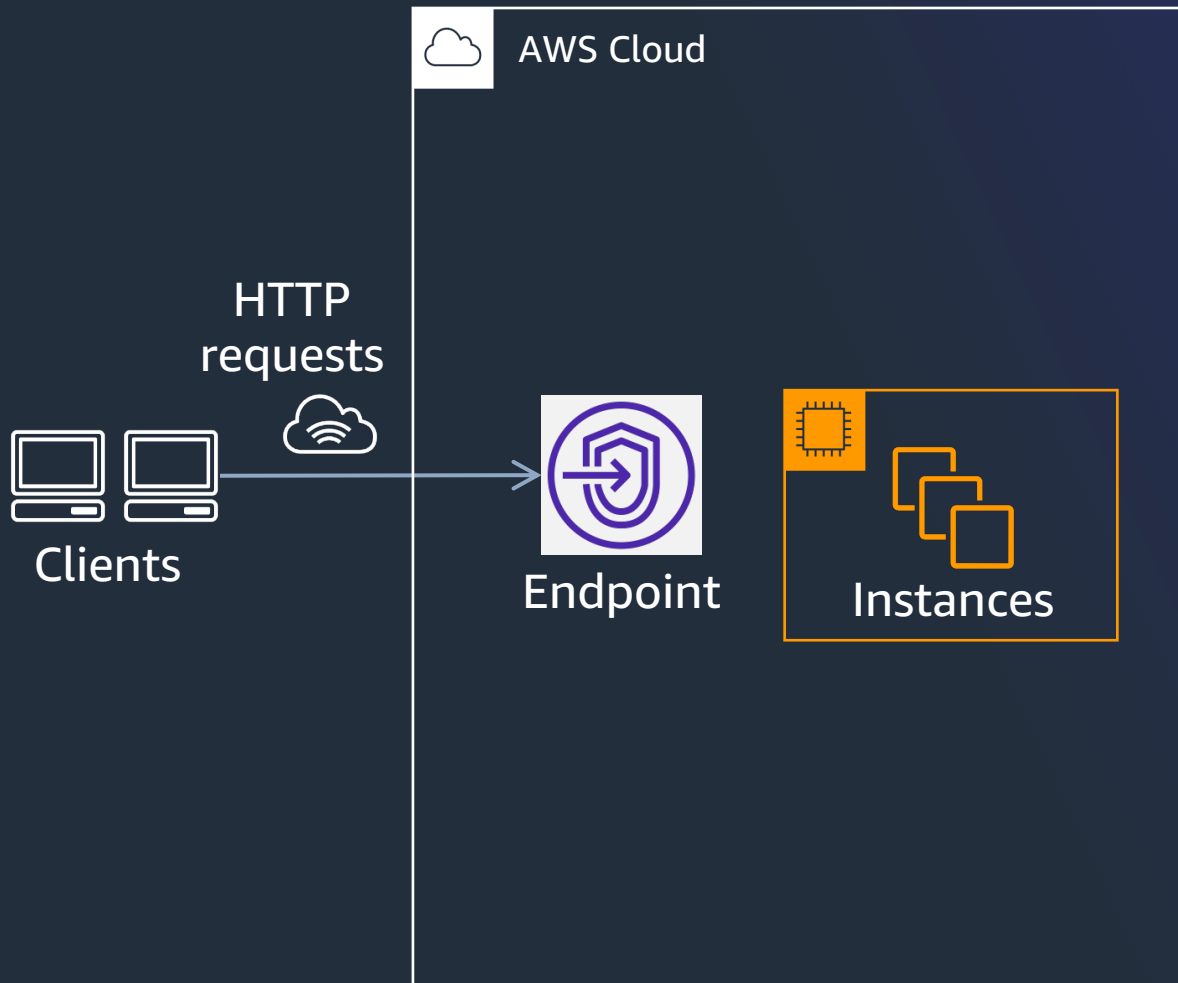
# Envoy-based request router



## Out-of-process architecture

- Envoy works with any programming language
  - Can write applications in Go, Java, C++, or any other language
  - Can bridge the gap between them – its behavior is identical, regardless of the application's programming language or the operating system they're running on
- The out-of-process architecture is beneficial, as it gives consistency across programming languages/applications stacks – you can get an independent lifecycle

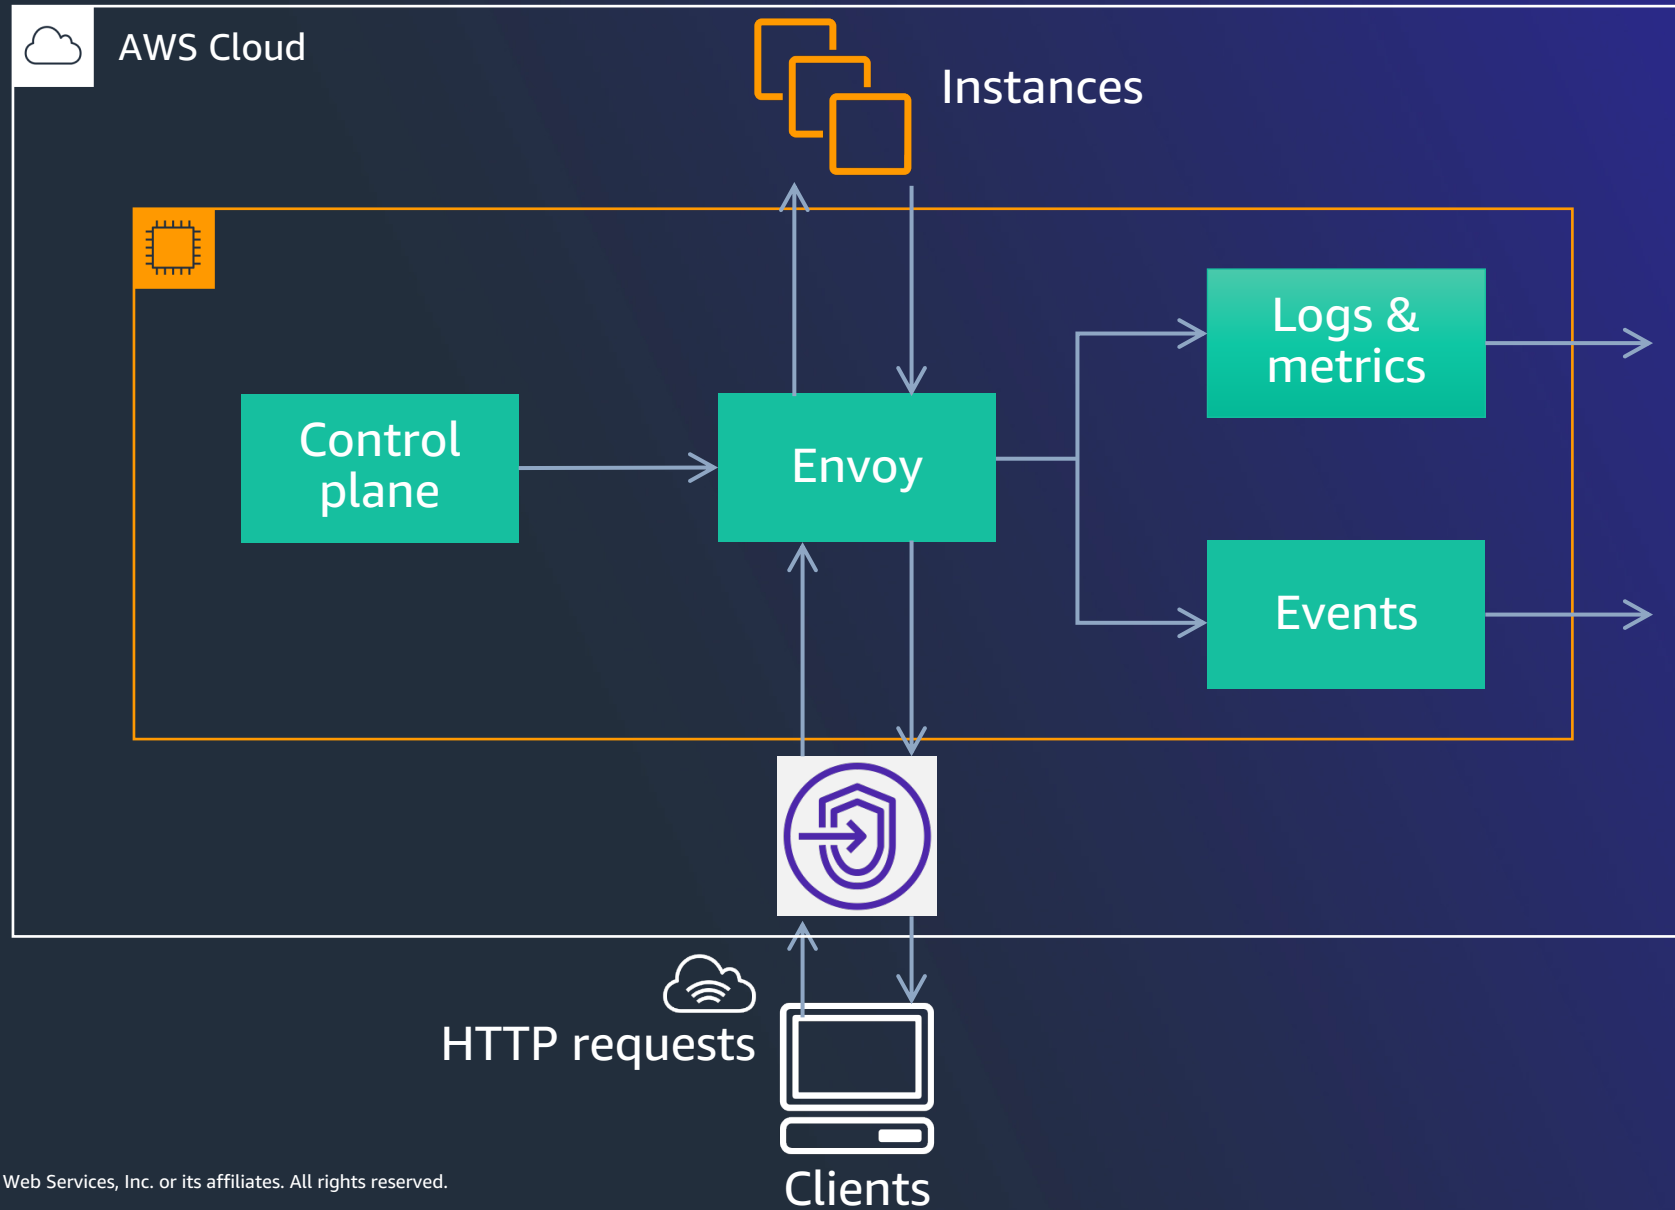
# Building Envoy-based request router



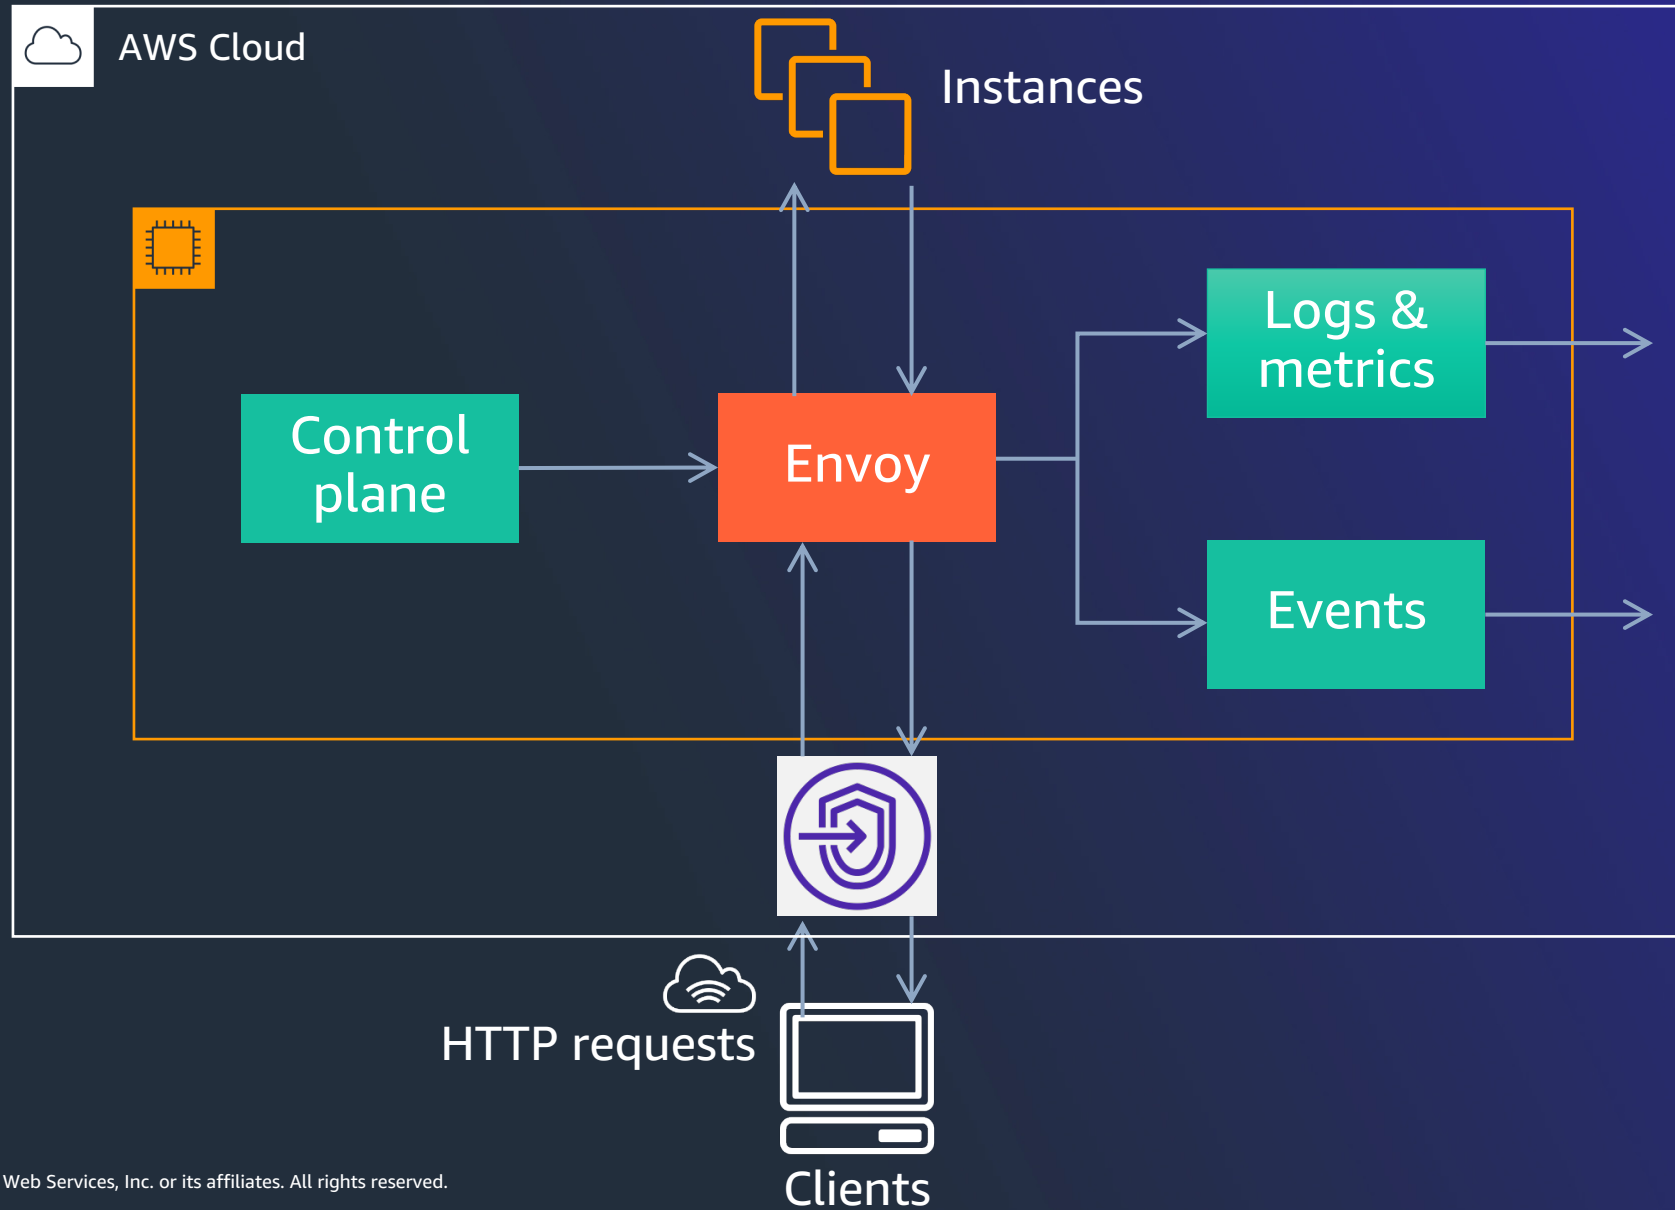
## Building blocks

- Bootstrap configuration – dynamic resources → container Instances

# Architecture




# Architecture






# Building Envoy-based request router



## Building blocks

- Listeners: Named network location (IP address + port); Envoy receives requests through it
- Routes: Route configuration in HCM filter; match the incoming requests (URI, headers, etc.) and define where traffic is sent
- Clusters: Group of upstream hosts that accept the traffic for a route; list of hosts or IP addresses on which the services are listening
- Endpoints: Service application instances


# Building Envoy-based request router



## Listeners

- Listening external HTTPS requests
- Operate on packet's payload


# Building Envoy-based request router



## Routes

- Envoy will be CRUX of LB and L7 request proxying layer
- Maintain mapping of service URL to application instance IPs; route incoming requests to an appropriate endpoint
- Domain map – for example <https://helloworld.us-east-1.awsapprunner.com>


# Building Envoy-based request router



## Clusters

- Envoy cluster → AWS App Runner service
- Each service version is assigned an Envoy cluster
- Configuration updated via Envoy cluster discovery service (CDS)


# Building Envoy-based request router



## Endpoints

- Upstream endpoints updated via Envoy endpoint discovery service (EDS)
- Dynamic and automatic scaling
- Unhealthy endpoints reaper


# Building Envoy-based request router: HTTP connection manager (HCM)



## HTTP connection manager (HCM)

- Network-level filter: Translate raw bytes into HTTP
- Handles
  - Access logging
  - Request ID generation
  - Tracing
  - Header manipulation
  - Retry policy
  - Timeout
  - Traffic weights
  - Route matcher


# Building Envoy-based request router: HTTP routing



- The router uses the information from the incoming request (e.g., host or authority headers) and matches it to an upstream cluster through virtual hosts and routing rules
- HTTP filters use the route configuration (`route_config`) that contains the route table

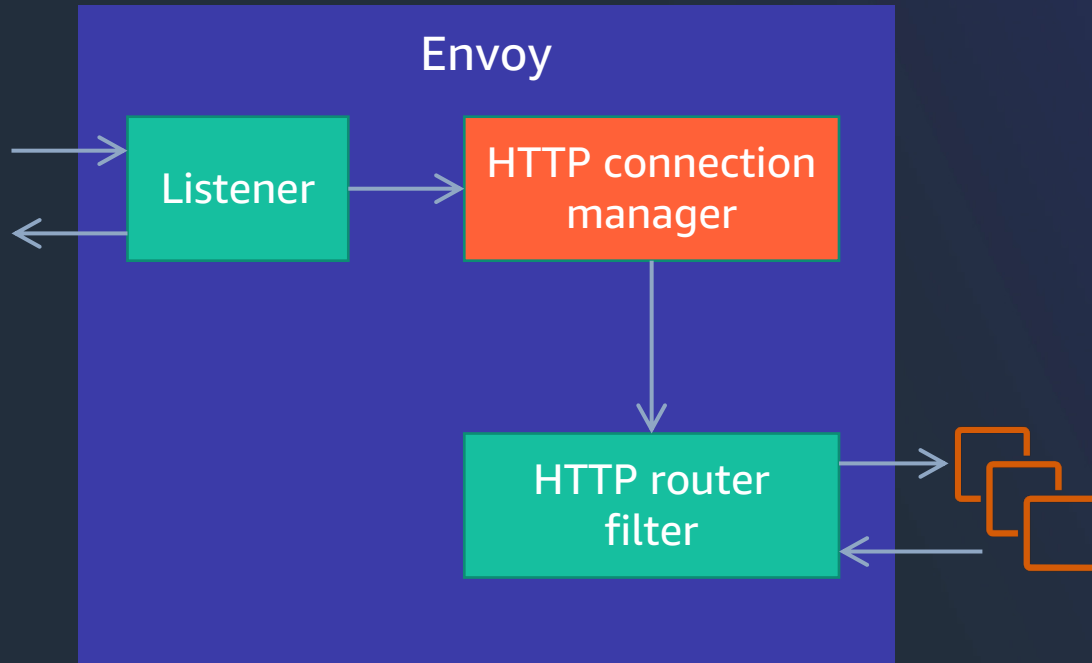


# Building Envoy-based request router: Request matching




- Set of headers can be specified to match request – router checks request headers against all specified headers in the route config
- Some matches include range, present, string, invert match, etc.
- HCM filters also support query parameter matching, TLS context matching, and gRPC route matching for routing

# Building Envoy-based request router: Traffic splitting




- Envoy supports traffic splitting to different routes within the same virtual host
- We can split traffic between two or more upstream clusters
- Traffic split can be on
  - Runtime percentage
  - Weighted clusters

# Building Envoy-based request router: Health check




- Envoy supports two types of health checks: Active vs. passive
- With active health checking, Envoy periodically sends a request to the endpoints to check its status
- With passive health checking, Envoy monitors how the endpoints respond to connections

# Building Envoy-based request router: Load balancing




- Load balancing is a way of distributing traffic between multiple endpoints in a single upstream cluster
- The reason for distributing traffic across numerous endpoints is to make the best use of the available resources (cost effective)

# Building Envoy-based request router: Retries




- Envoy allows setting retry policy at the virtual host and the route level based on conditions
- For example, if there is a 429 status code returned from upstream, then retry will happen seamlessly on upstream of other available instance

# Building Envoy-based request router: Circuit breakers



- Circuit breaker pattern prevents additional failures managing access to failing services
- It allows us to fail quickly and apply back pressure downstream as soon as possible
- We can configure the circuit breaker thresholds for each route priority separately and globally

# Building Envoy-based request router: Timeout




- Envoy supports numerous configurable timeouts that depend on the scenarios you're using the proxy for
- Example timeouts
  - **request\_timeout** specifies the amount of time Envoy waits for the entire request to be received
  - **idle\_timeout** represents when a downstream or upstream connection gets terminated if there are no active streams (default 1 hour)

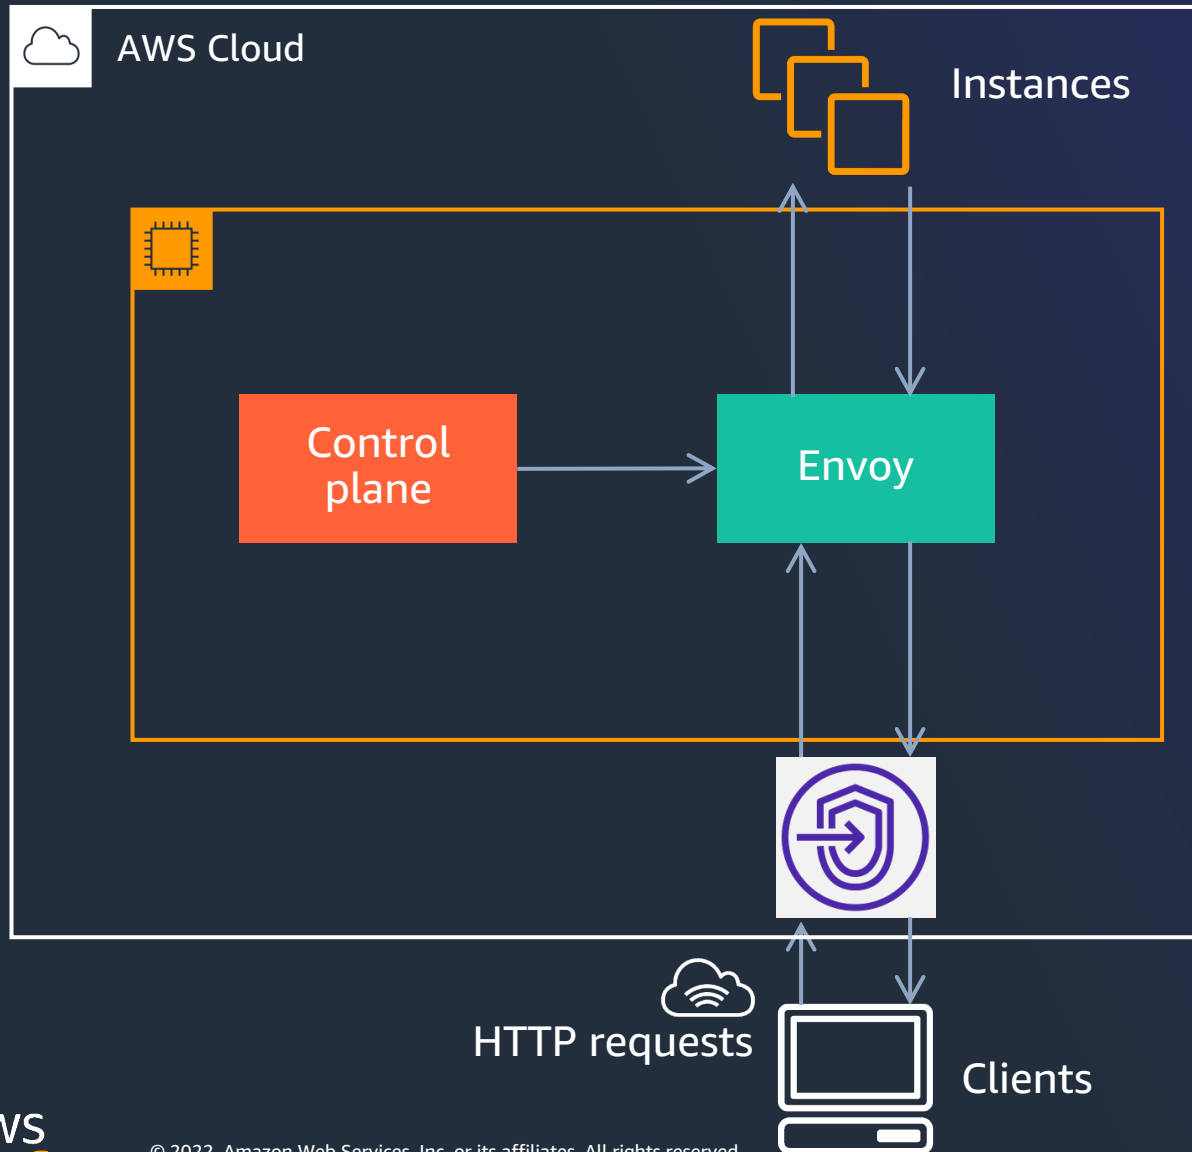


# Building Envoy-based request router: Rate limiter

Envoy supports global and local rate limits to upstream endpoints




# Building Envoy-based request router: Control plane



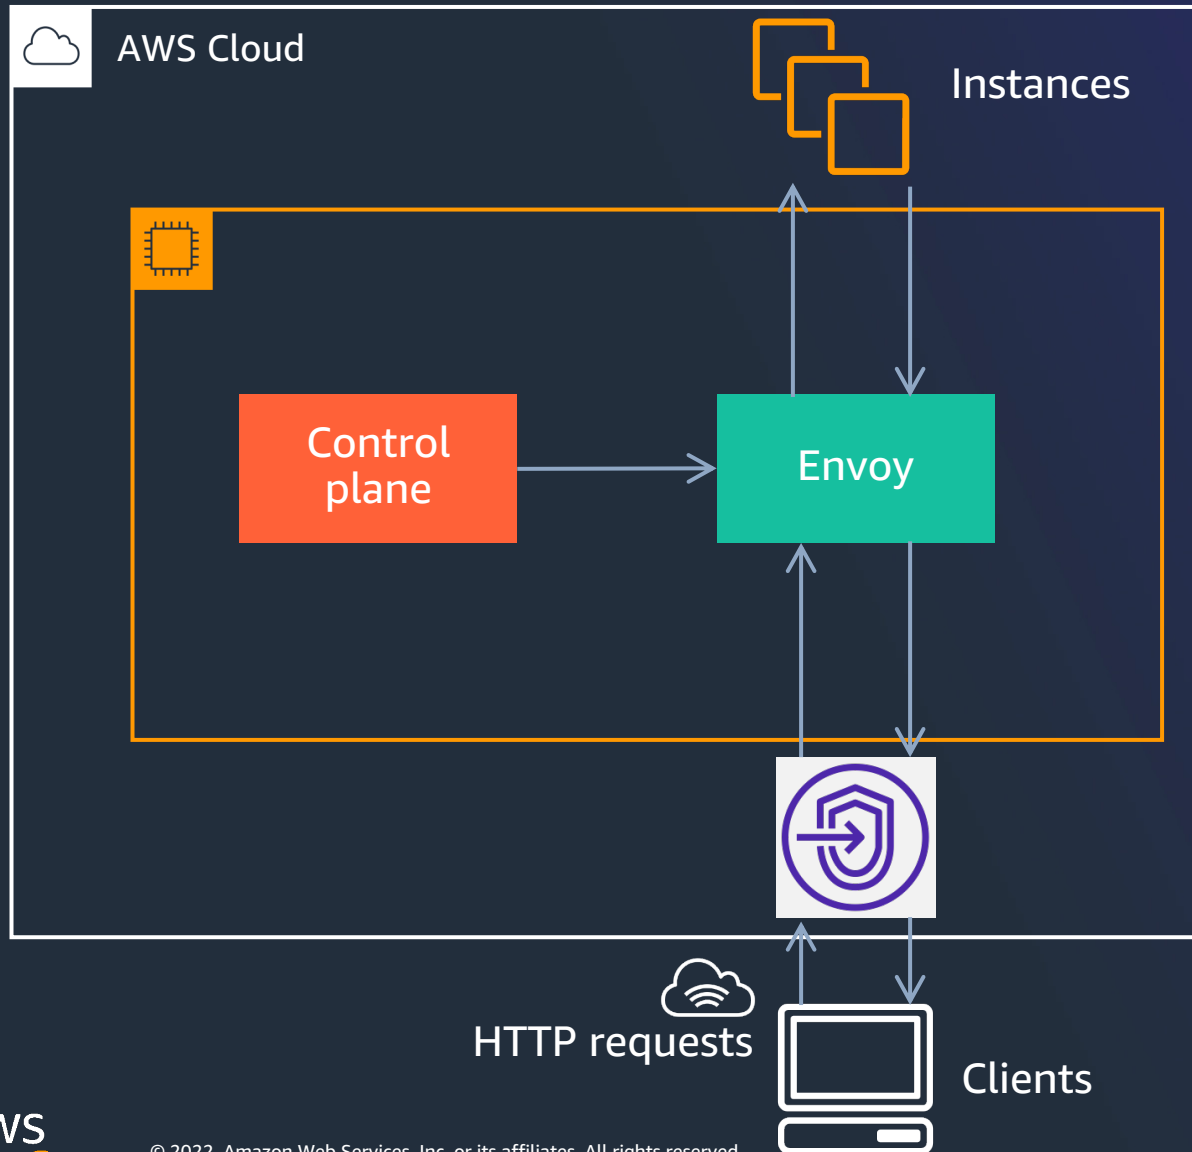
- Custom control plane for Envoy
- Responsible for
  - Service cluster discovery (CDS) and service endpoint discovery (EDS)
  - Route and other configuration
- gRPC server listening on defined port

# Building Envoy-based request router: Envoy runtime



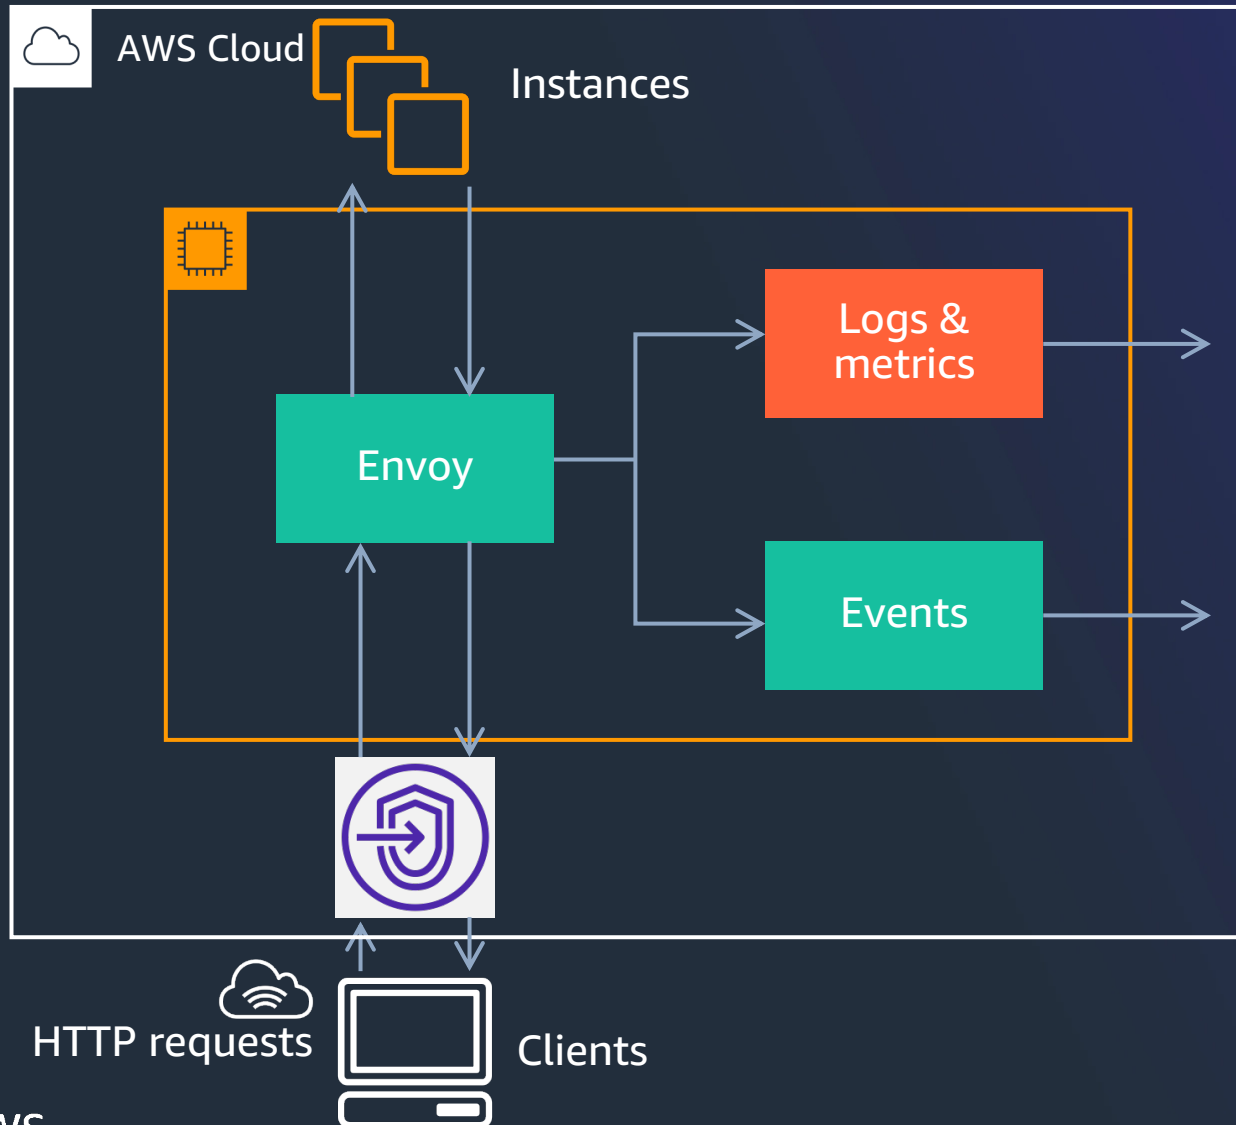
- Envoy needs to be updated with latest service information at runtime
- Envoy discovers dynamic resources by querying management servers using xDS (discovery service APIs)
- Dynamic resources:
  - Clusters
  - Endpoints
- Envoy requests resources via subscriptions by initiating gRPC streams from Envoy client to the management server

# Building Envoy-based request router: Envoy runtime



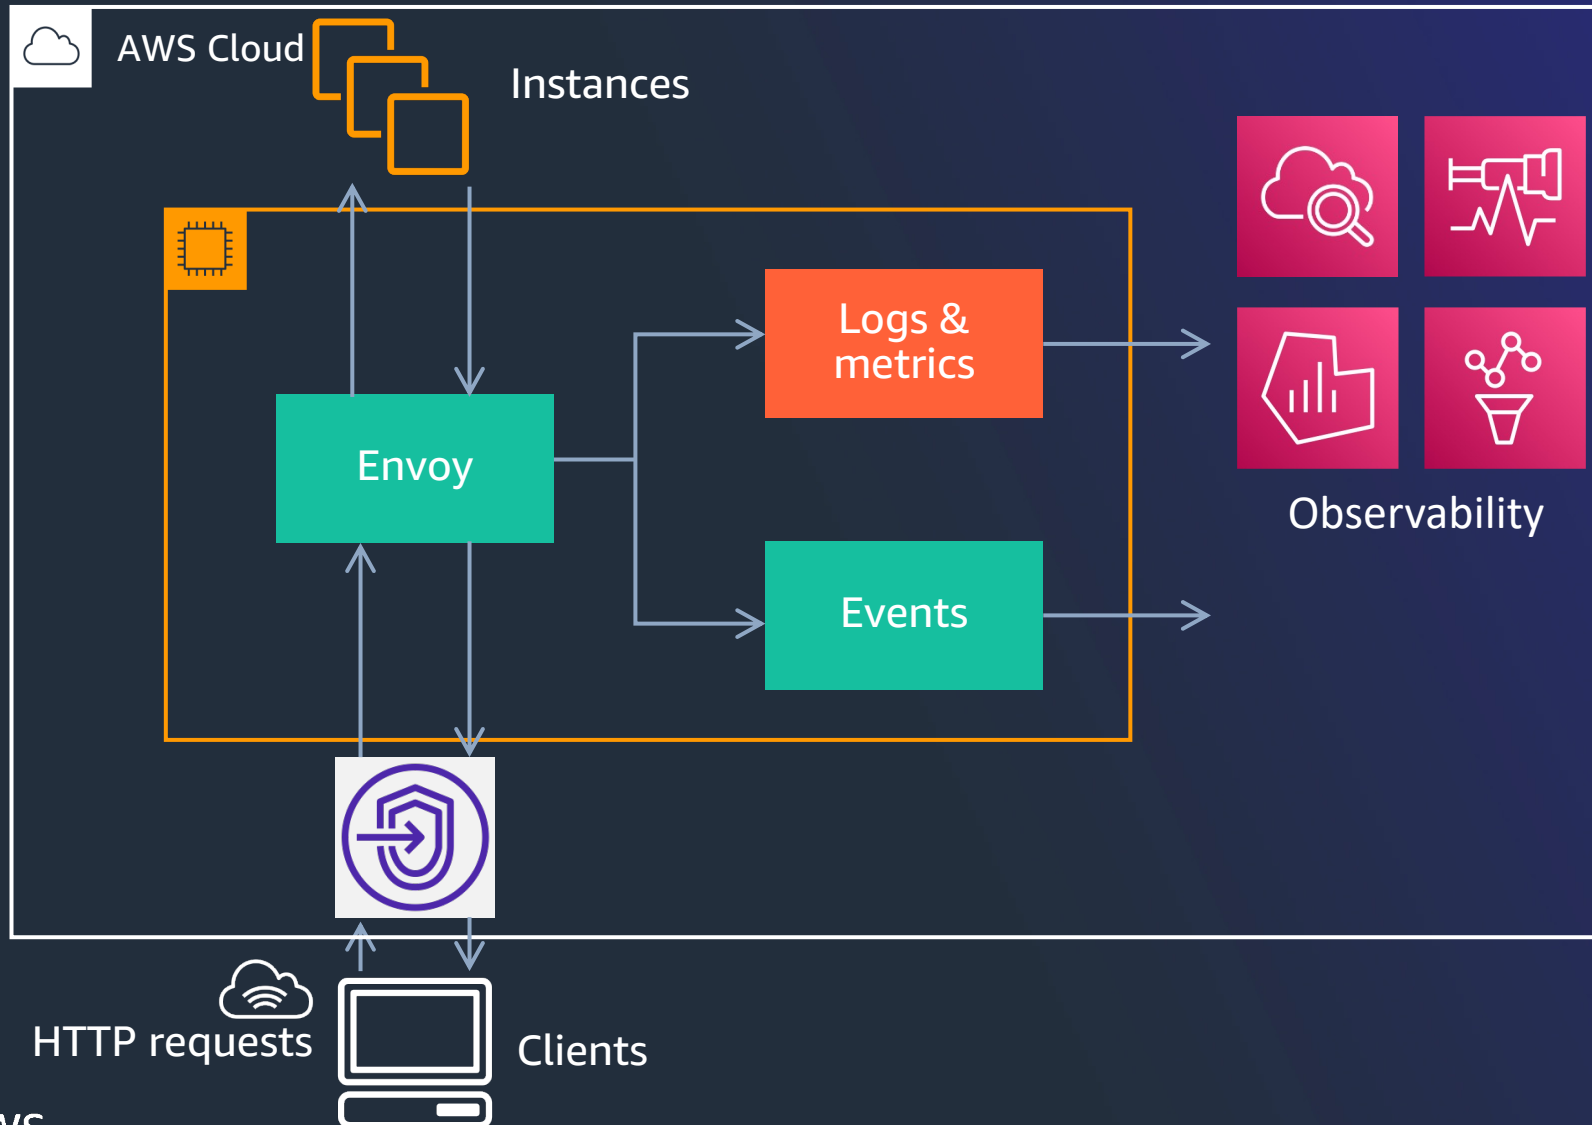
- Clusters configured on Envoy through **CDS xDS API** – [bit.ly/3RZmMq6](https://bit.ly/3RZmMq6)
- Endpoints are configured using **EDS xDS API** – [bit.ly/3S0f8MR](https://bit.ly/3S0f8MR)
- AWS App Runner uses CDS API to dynamically update clusters and EDS to update endpoints per cluster

# Building Envoy-based request router: Logs and metrics



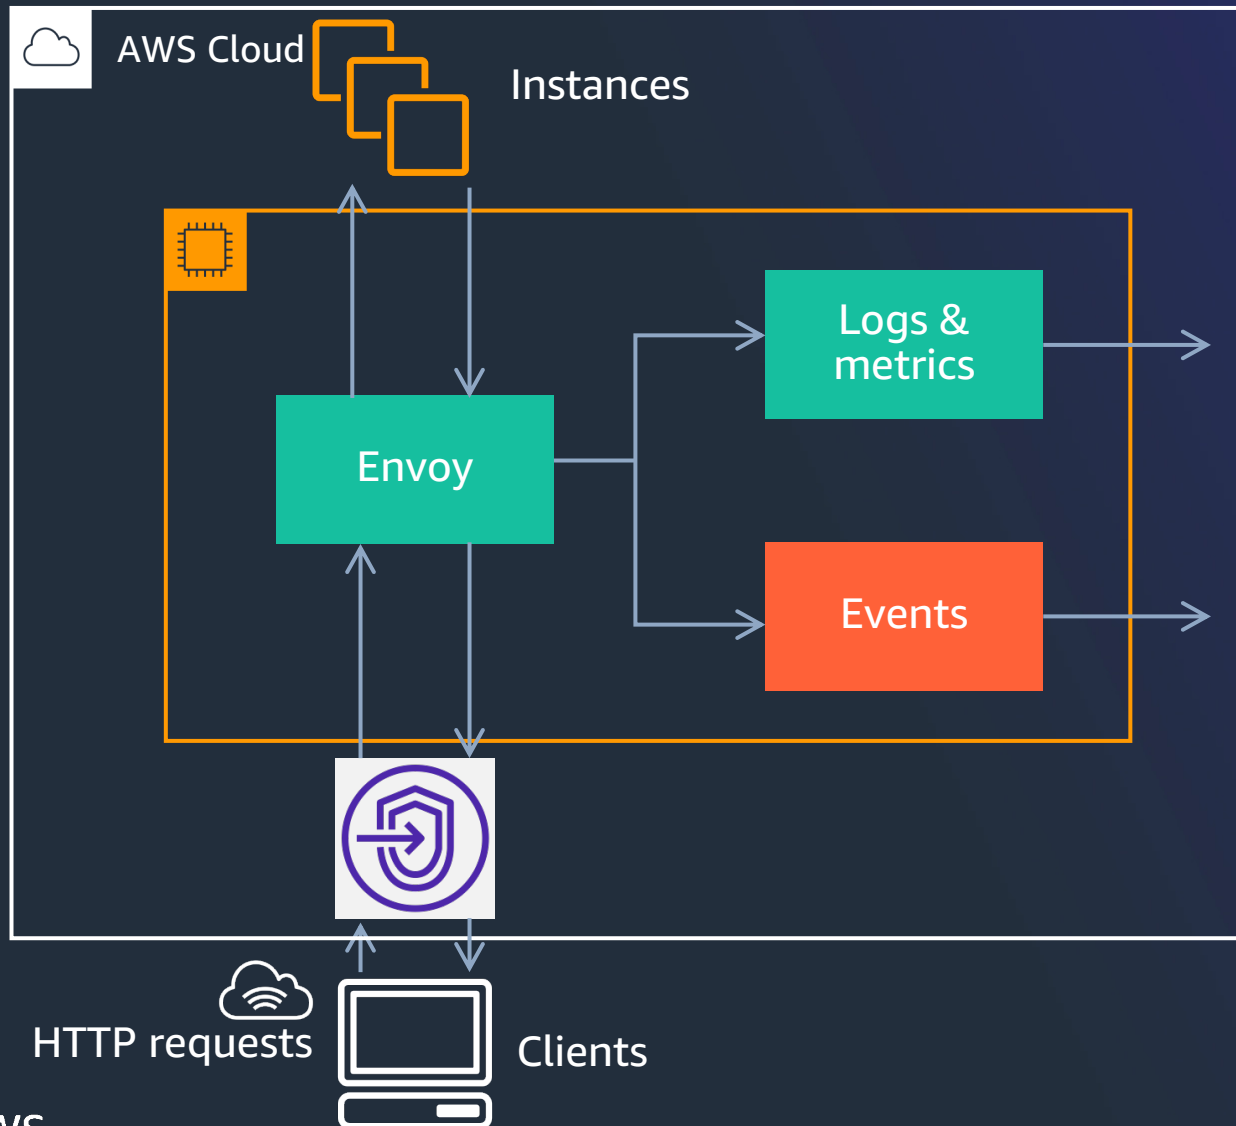
- Responsible for Envoy logs and metrics collection
- Envoy configures exposed metrics locally
- A stats\_sink that receives metrics emitted from Envoy
- Metrics such as 2xx, 4xx, 5xx request count are eventually forwarded to logs and metrics service
- gRPC service listening on defined port

# Building Envoy-based request router: Logs and metrics



- Amazon CloudWatch
- Amazon Managed Service for Prometheus
- AWS Distro for OpenTelemetry
- Amazon Managed Grafana
- AWS X-Ray
- And more

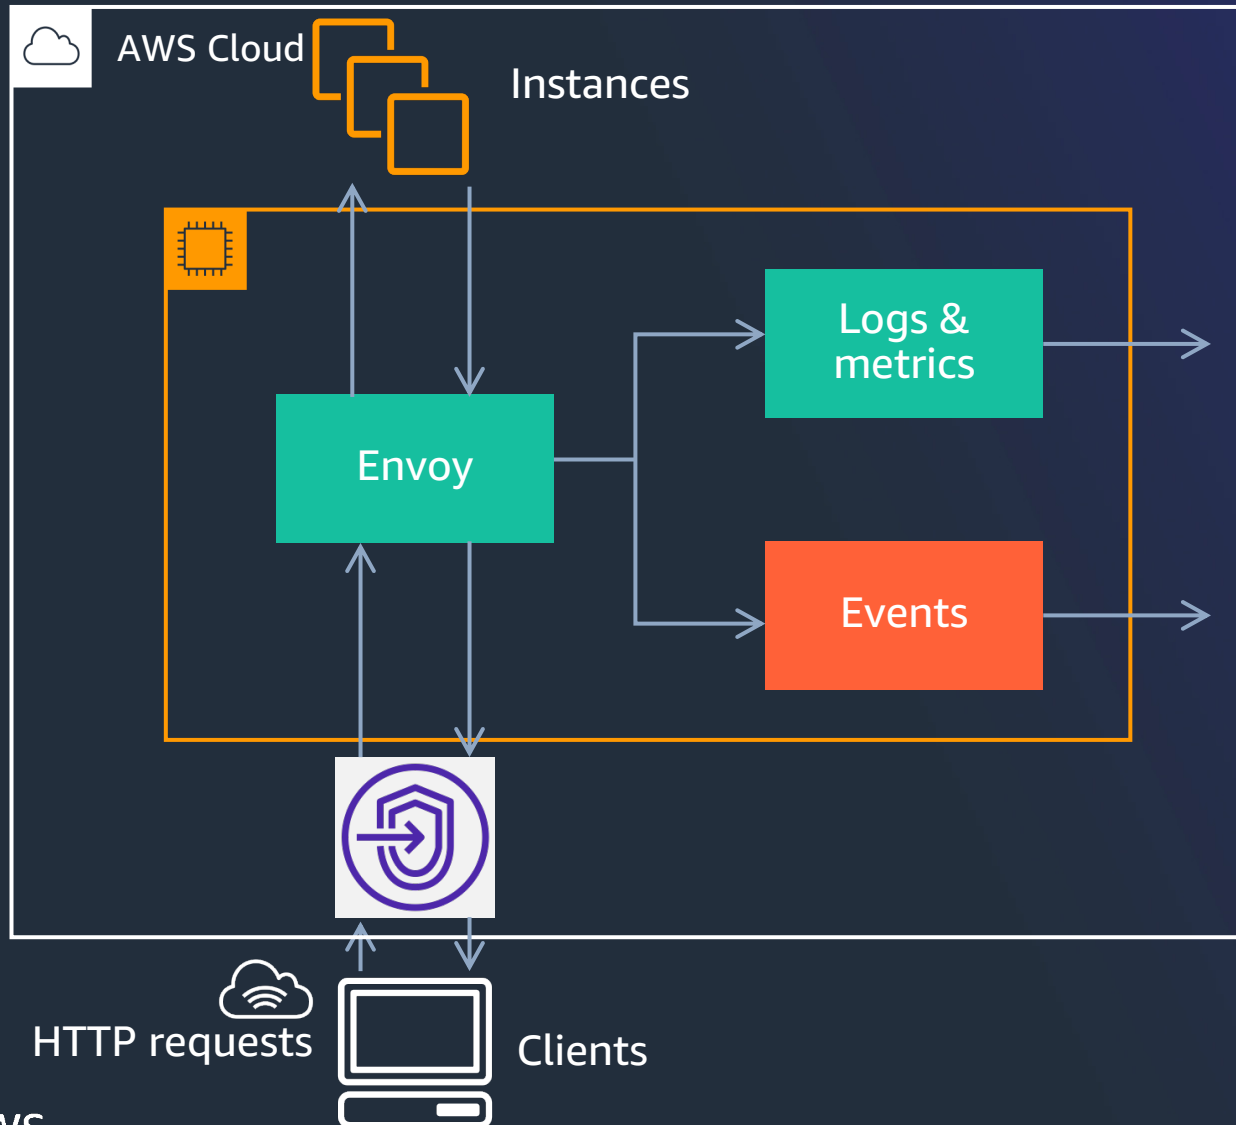
# Building Envoy-based request router: Event streamer



- Responsible for reporting service status and health check results
- gRPC service listening on defined port



# Building Envoy-based request router: Event streamer



Used in supporting service

- Safe deploy (blue-green)
- Automatic scaling
- Load balancing
- Request routing
- ...

# Building Envoy-based request router: Production readiness

- Function readiness
- Capacity management – CPU/memory consumption
- Performance and scalability
- Security (e.g., DDoS)
- Other operational readiness – monitors, alarms, etc.

# Wrap up

Envoy brings value in building request router with many features:

- High throughput and performance request routing and load balancing
- Portability and extensibility
- Reliability and availability
- Observability

# More on roadmap

- GitHub issue 2: Allow private endpoints for App Runner services – [github.com/aws/apprunner-roadmap/issues/2](https://github.com/aws/apprunner-roadmap/issues/2)
- GitHub issue 52: Set additional X-Forwarded-Headers – [github.com/aws/apprunner-roadmap/issues/52](https://github.com/aws/apprunner-roadmap/issues/52)
- GitHub issue 58: AWS WAF support – [github.com/aws/apprunner-roadmap/issues/58](https://github.com/aws/apprunner-roadmap/issues/58)
- GitHub issue 104: Configurable timeout – [github.com/aws/apprunner-roadmap/issues/104](https://github.com/aws/apprunner-roadmap/issues/104)

The screenshot shows the GitHub interface for the repository 'aws/apprunner-roadmap'. The 'Issues' tab is selected, showing 121 open issues. A banner at the top asks if the user wants to contribute to the repository, with a link to the contributing guidelines. Below the banner, there are filters for 'is:issue is:open' and buttons for 'Labels' (9) and 'Milestones' (0). A 'New issue' button is also present. The list of issues includes:

- Support pipenv and poetry for python code based service (#152 opened yesterday by bryantbiggs)
- Support for ca-central-1 Region (Canada-Montreal) (#151 opened 5 days ago by 0sudash9)
- AWS AppRunner start command running multiple time (#150 opened 12 days ago by myanees284)
- #2 #92 (#149 opened 16 days ago by jereytucker)
- Gpu Supports (#148 opened 26 days ago by fightnny)
- Improve speed UpdateService (#147 opened 29 days ago by kamitath)
- CloudWatch embedded metrics (#146 opened on Jul 28 by rschick)
- Support for eu-west-2 (London) (#145 opened on Jul 18 by pauldougan)
- Single Sign-On Application Exposure for App Runner (#144 opened on Jul 5 by jtelleriar)
- Changing environment variables in AppRunner Console overrides HTTP healthcheck config (#143 opened on Jul 4 by VladStarr)

# Thank you!

**Yiming Peng**

 pengyiming

 @pymhq



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



**Please complete  
the session survey**