
KubeCon
CloudNativeCon
North America 2022
BUILDING FOR THE ROAD AHEAD
DETROIT 2022
TiKV
TiKV
HELM

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Decentralized Routing for a Sharded Application on Service Mesh

Pankaj Sikka & Vinay Gonuguntla

Decentralized Routing for a Sharded Application on Service Mesh

KubeCon CloudNativeCon North America 2022
BUILDING FOR THE ROAD AHEAD
DETROIT 2022

Vinay Gonuguntla
Staff Software Engineer
Intuit

Pankaj Sikka
Staff Software Engineer
Intuit

KubeCon CloudNativeCon North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

October 24-28, 2021

Agenda

01 **Service Mesh**
What Servicemesh is and it's use-cases at Intuit

02 **Routing for Sharded Applications**
Sharded applications, routing on mesh

03 **Design & Challenges**
What we built , what we learnt

04 **Demo**
Decentralized routing demo

05 **What's Next?**
Future investment

Who we are - Intuit

End User Award in 2019 and 2022

End user of Cloud Native and mobile open source tech

Created, open-sourced, used, and maintained by Intuit

Intuit Scale

245+
Clusters

16000+
Namespaces

~77,320
Nodes

7000000+
PODS

2000+
Services

SERVICE MESH

Service Mesh

- **Service mesh** - infrastructure layer to facilitate service-to-service communication
- Security, Observability, Routing and Access Control
- Examples - Istio, Linkerd, Consul

Service Mesh

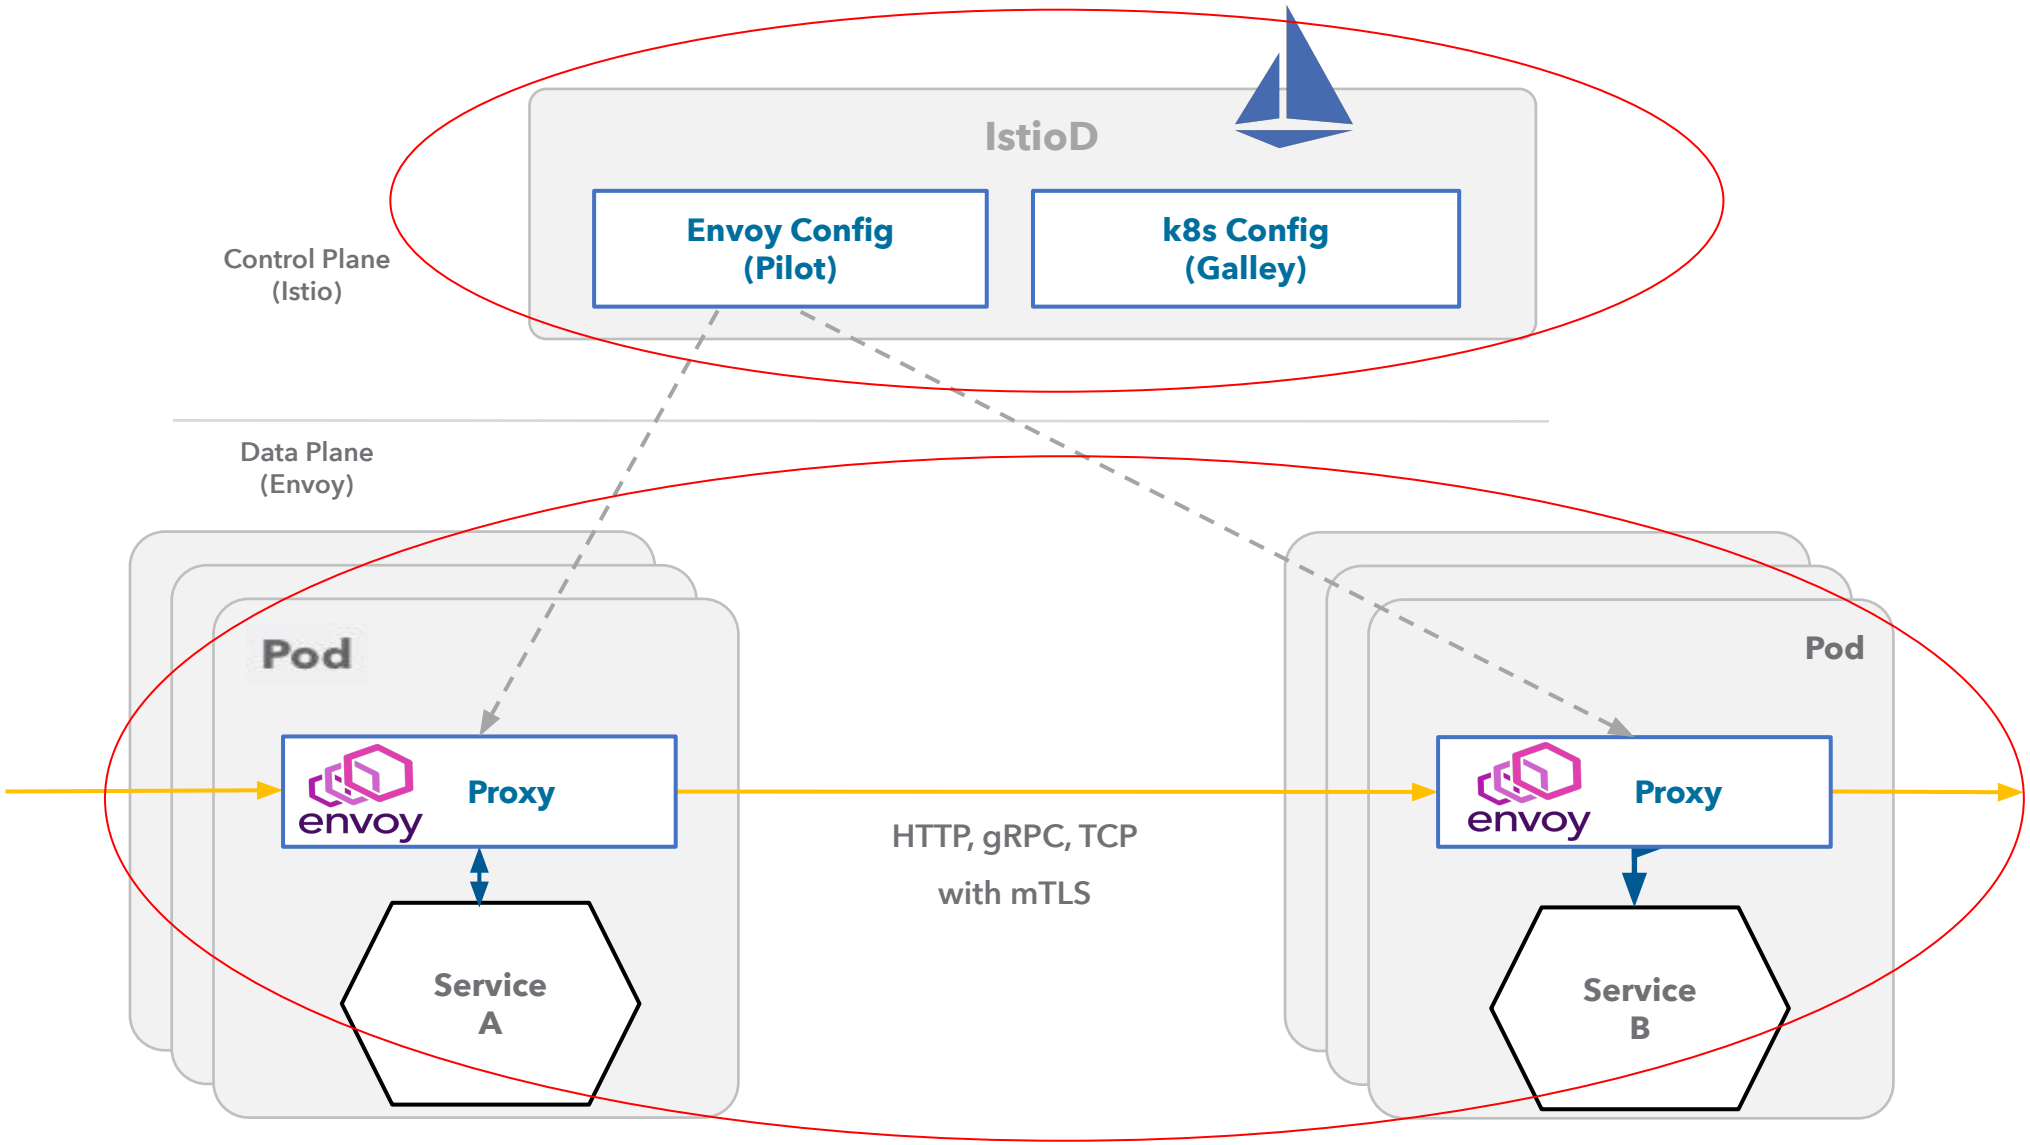

Service Mesh at Intuit

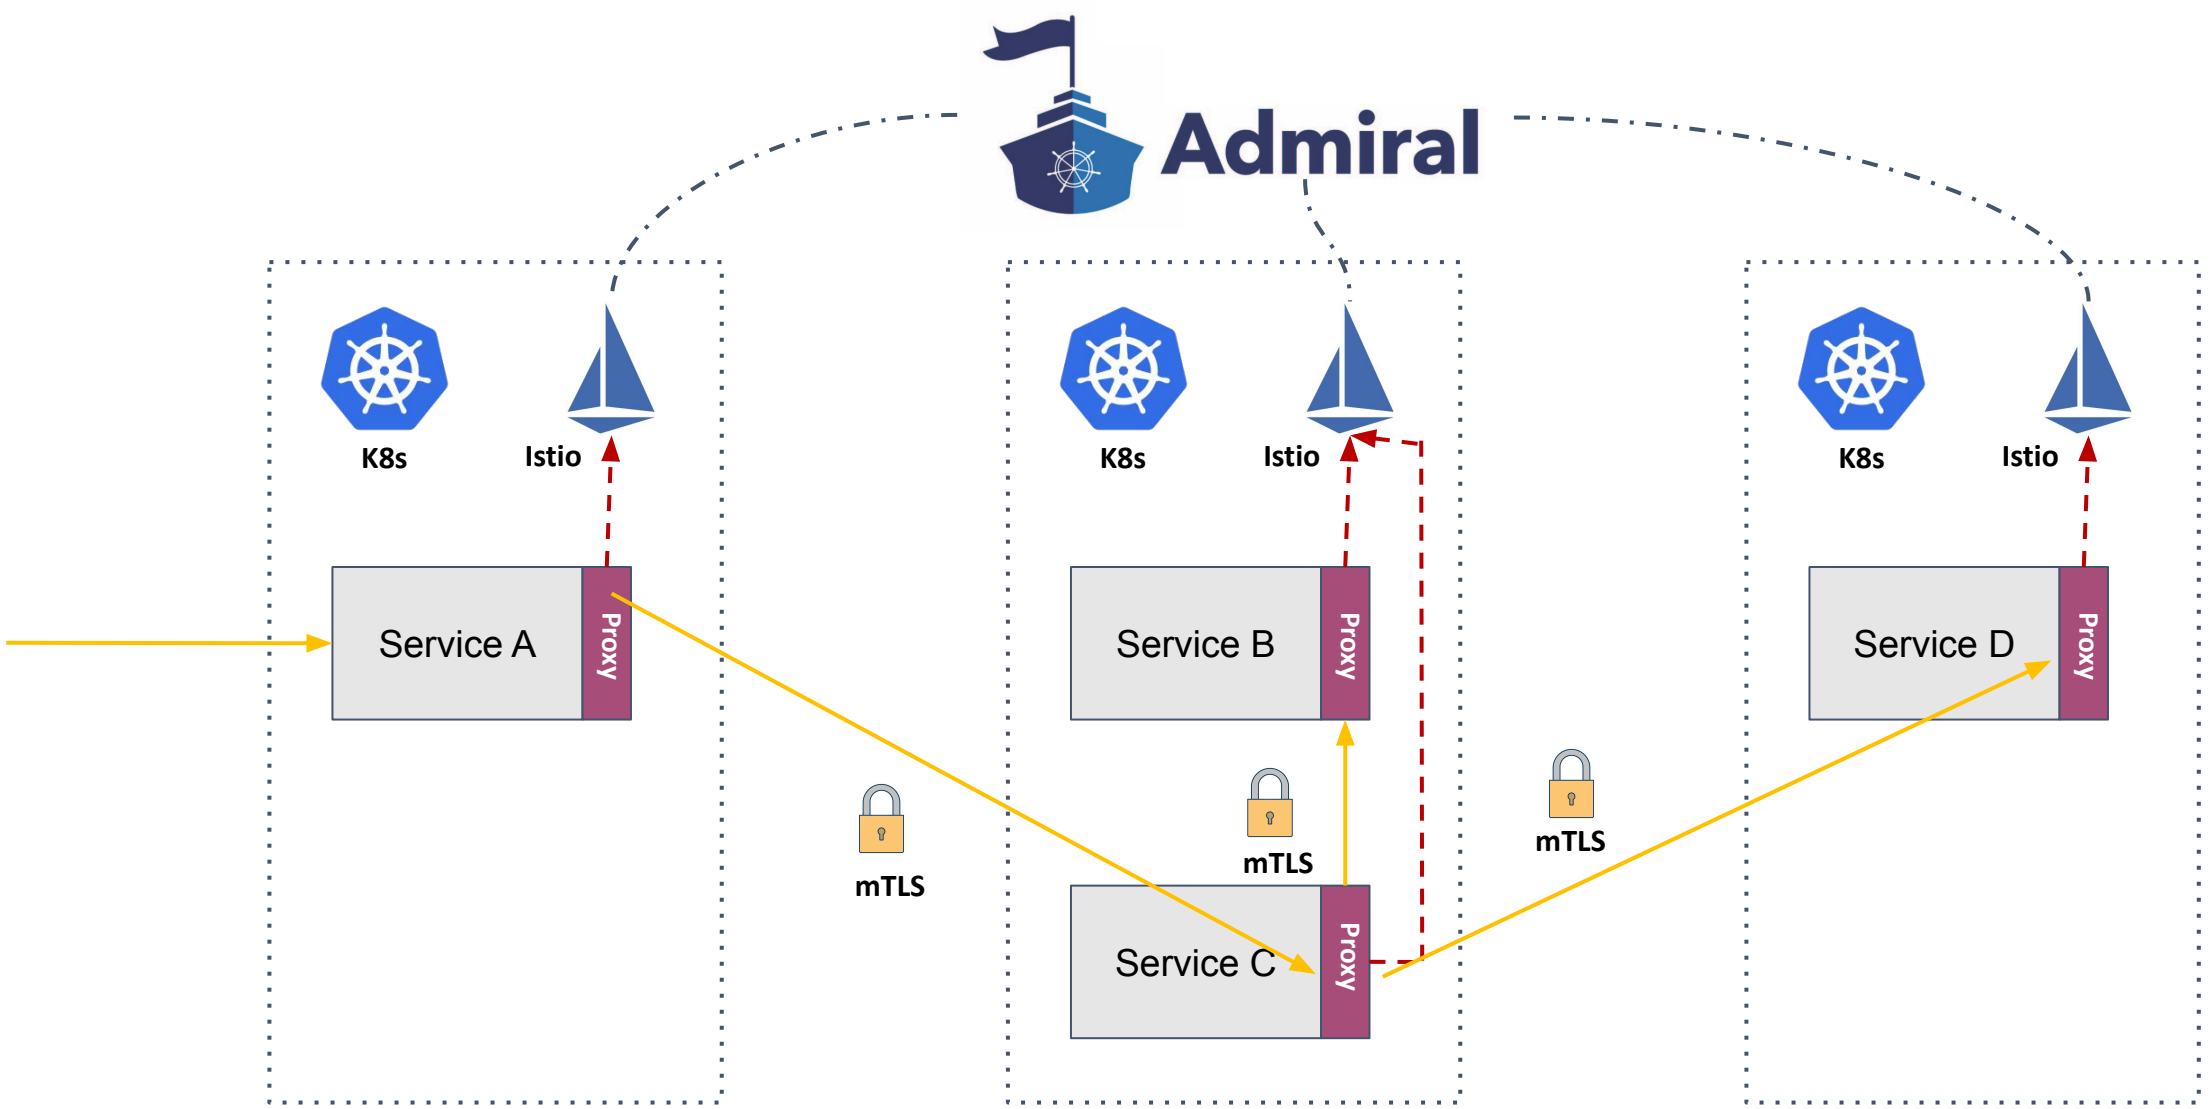

ROUTING FOR SHARDED APPLICATIONS

What is a sharded application ?

Sharded Routing - MESH

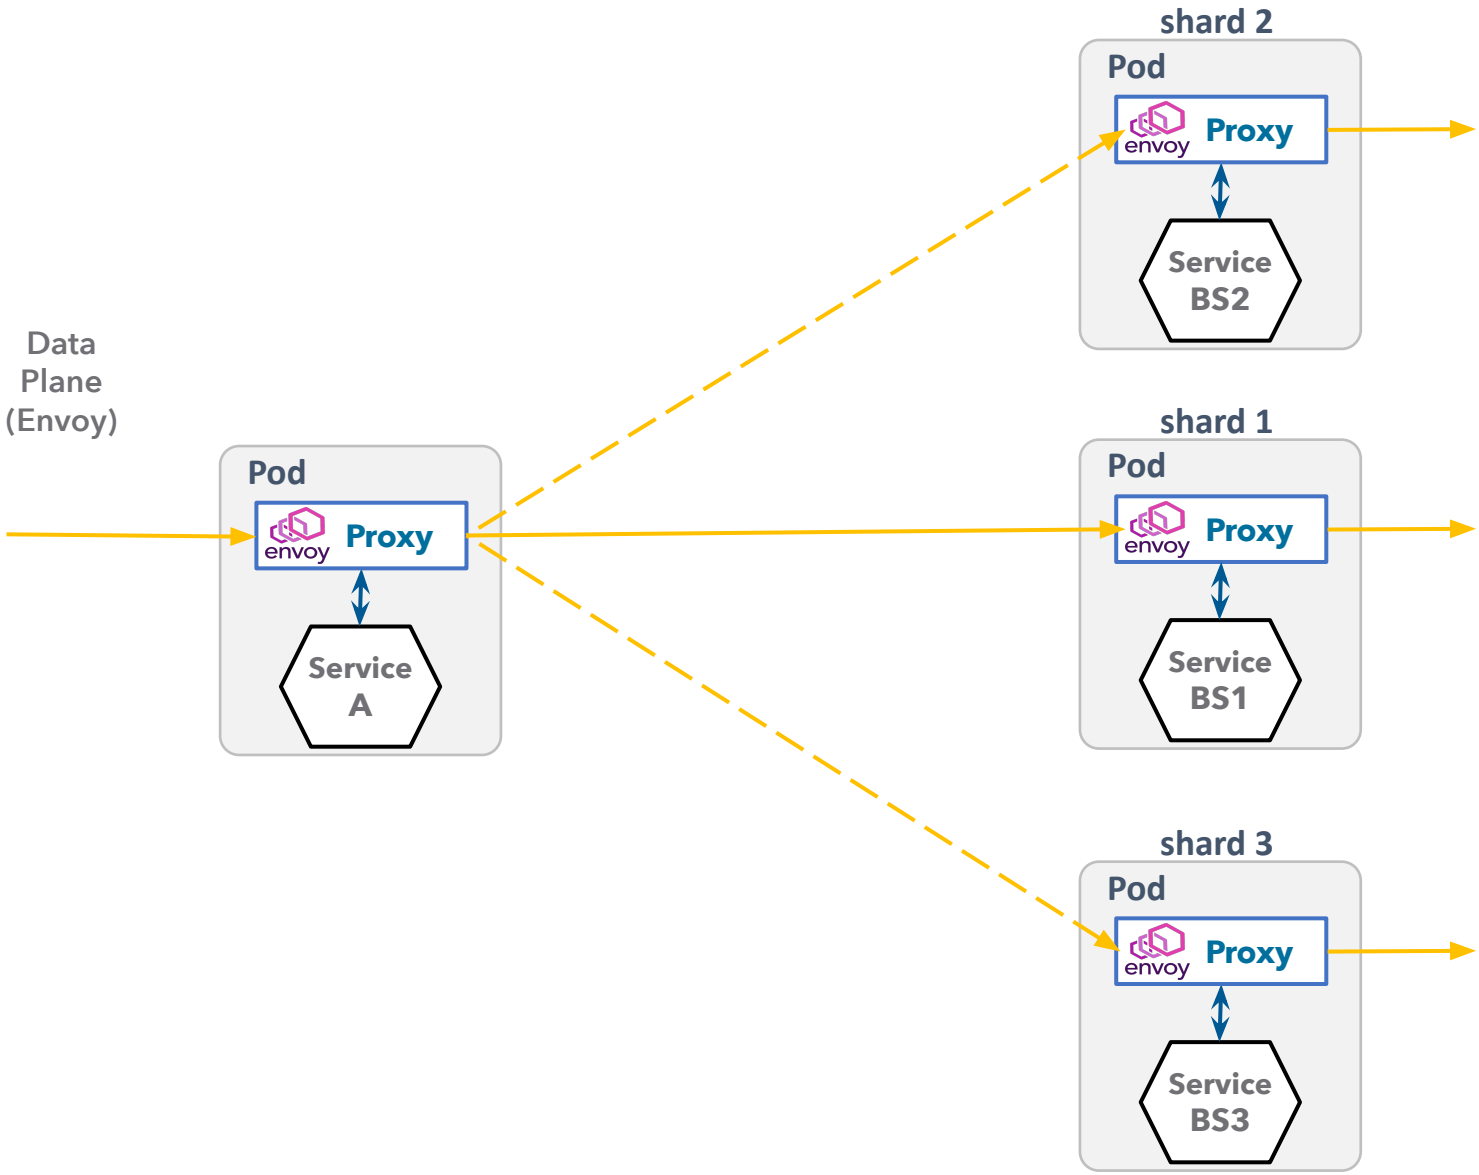

Sharded Routing - Goals

- Client services need to have minimal/no changes
- Scale for millions of users
- Support moving data among shards
- Service owners should control the routing configuration

Sharded Routing - Existing solutions

Istio VirtualService

```yaml
apiVersion: networking.istio.io/v1beta1
kind: VirtualService
metadata:
  name: demo-route
spec:
  hosts:
  - demo.greeting.mesh
  http:
  - name: "demo-1-route"
    match:
    - uri:
        prefix: "/customer/1"
    - uri:
        prefix: "/customer/2"
    route:
    - destination:
        host: shard1.customerservice.mesh
        subset: v1
  - name: "demo-2-route"
    match:
    - uri:
        prefix: "/customer/3"
    route:
    - destination:
        host: shard2.customerservice.mesh
        subset: v1
```

Limitations

- Millions of records, in a single file
- Not scalable due to size constraints
- Near real time update, very challenging
- Retry on moved shards not possible

Sharded Routing - Approach

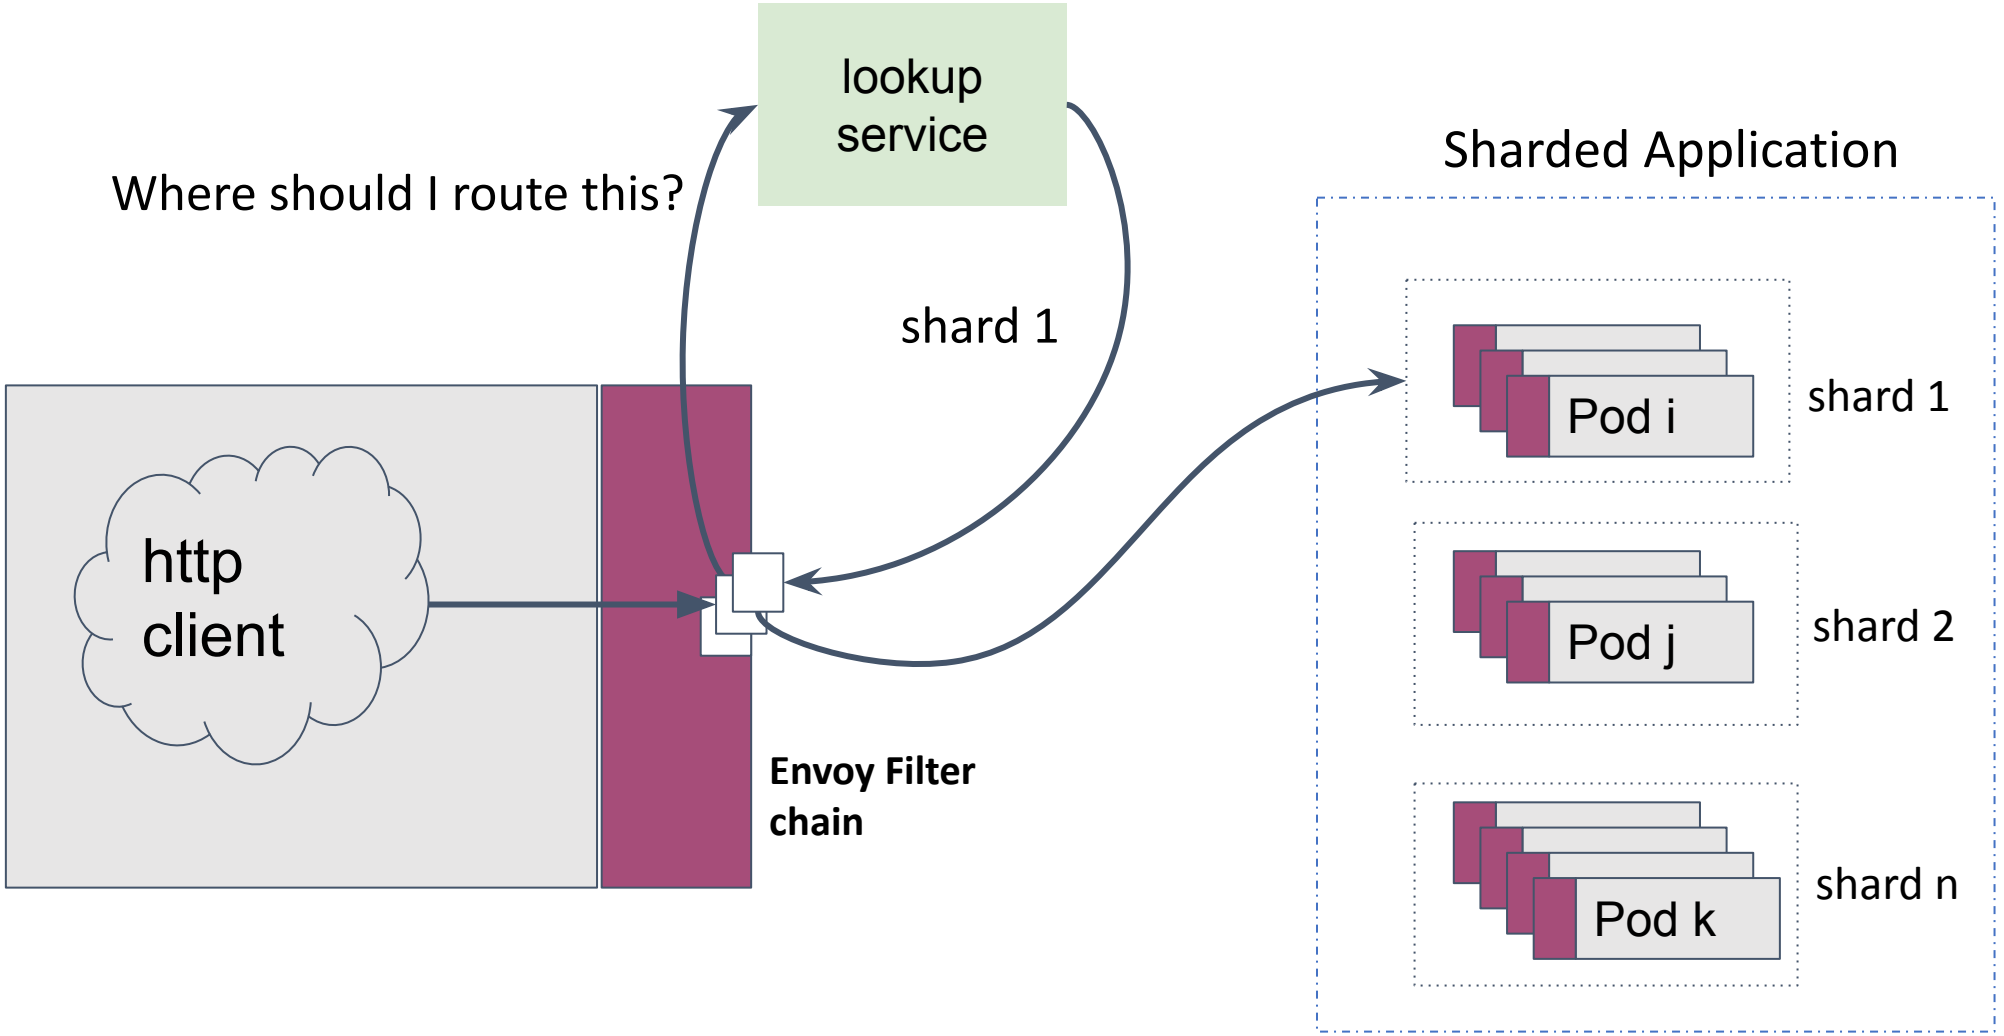

Dynamic Routing in Intuit

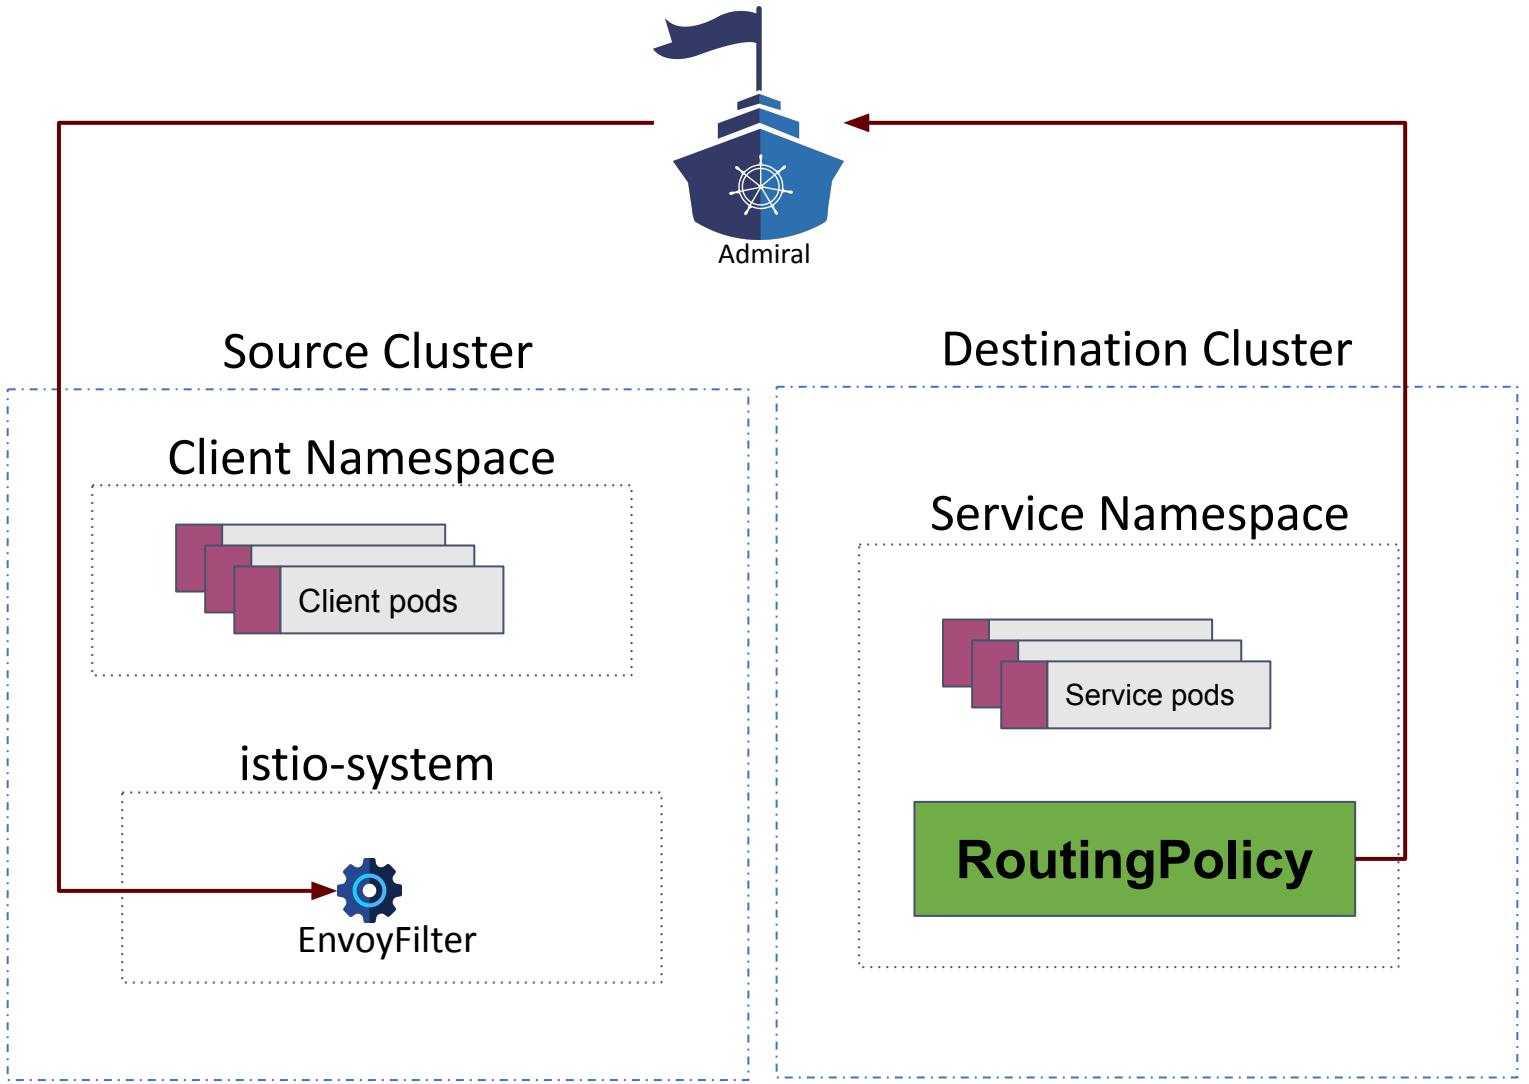

Dynamic Routing in Intuit


```
apiVersion: networking.istio.io/v1alpha3
kind: EnvoyFilter
...
spec:
  configPatches:
  - applyTo: HTTP_FILTER
    match:
      context: SIDECAR_OUTBOUND
...
        config:
             configuration:
               '@type':
type.googleapis.com/google.protobuf.StringValue
       value: |-
         routingAlgorithm: ShardedRouting
         cachePrefix: prefix-v1-a
         cacheTtlSec: "3600"
         lookupUrl: http://demo.lookup.mesh/
         env: stage
         wasmPath:/etc/istio/dynamicrouter.wasm
        hosts: demo.greeting.mesh
         plugin: demo
...
```

```
apiVersion: admiral.io/v1alpha1
kind: RoutingPolicy
metadata:
  annotations:
    admiral.io/env: stage
...
  name: demo-routing-policy
spec:
  plugin: demo
  hosts:
    - demo.greeting.mesh
  config:
    routingAlgorithm: ShardedRouting
    cachePrefix: prefix-v1-a
    cacheTtlSec: "3600"
    lookupUrl: http://demo.lookup.mesh/
    env: stage
    wasmPath: /etc/istio/dynamicrouter.wasm
```

Challenges

1 Workload selector limitations

OR operation in workload selector is not supported

2 TinyGo limitations

The need for external lookup service

3 Pre-built proxy image

Move to dynamically loading the business logic

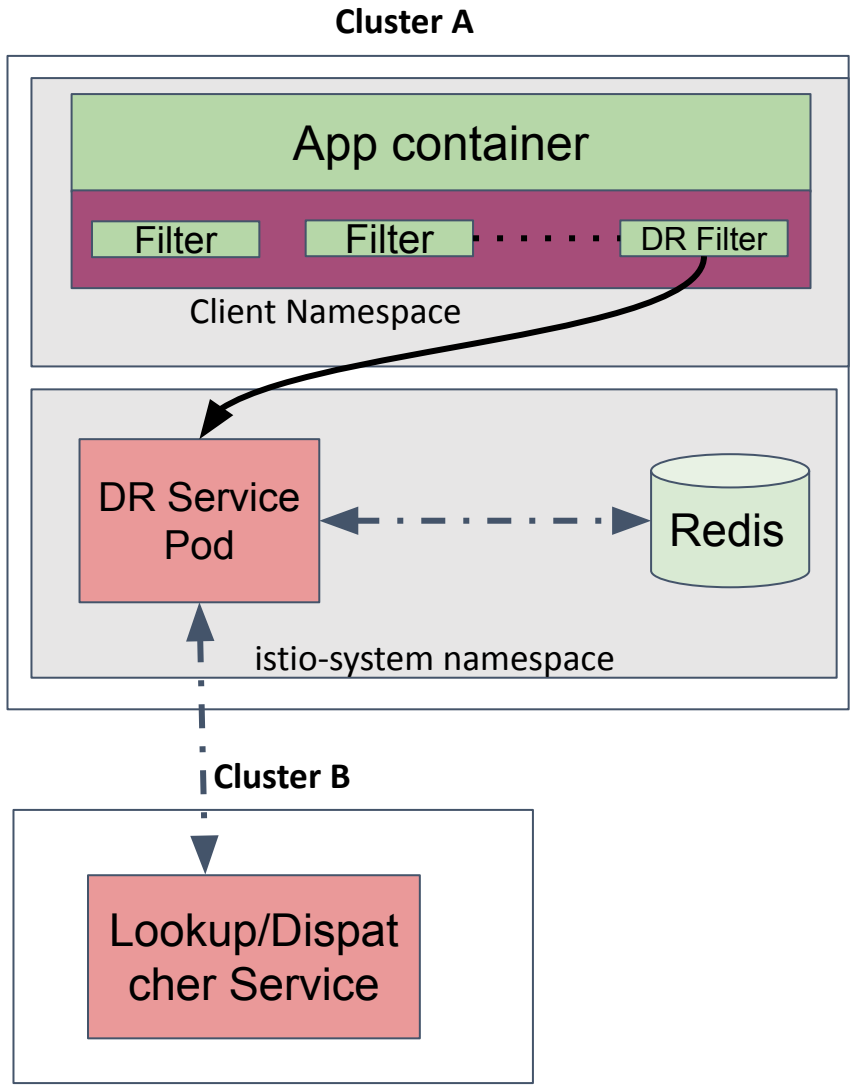

DEMO

Demo Setup

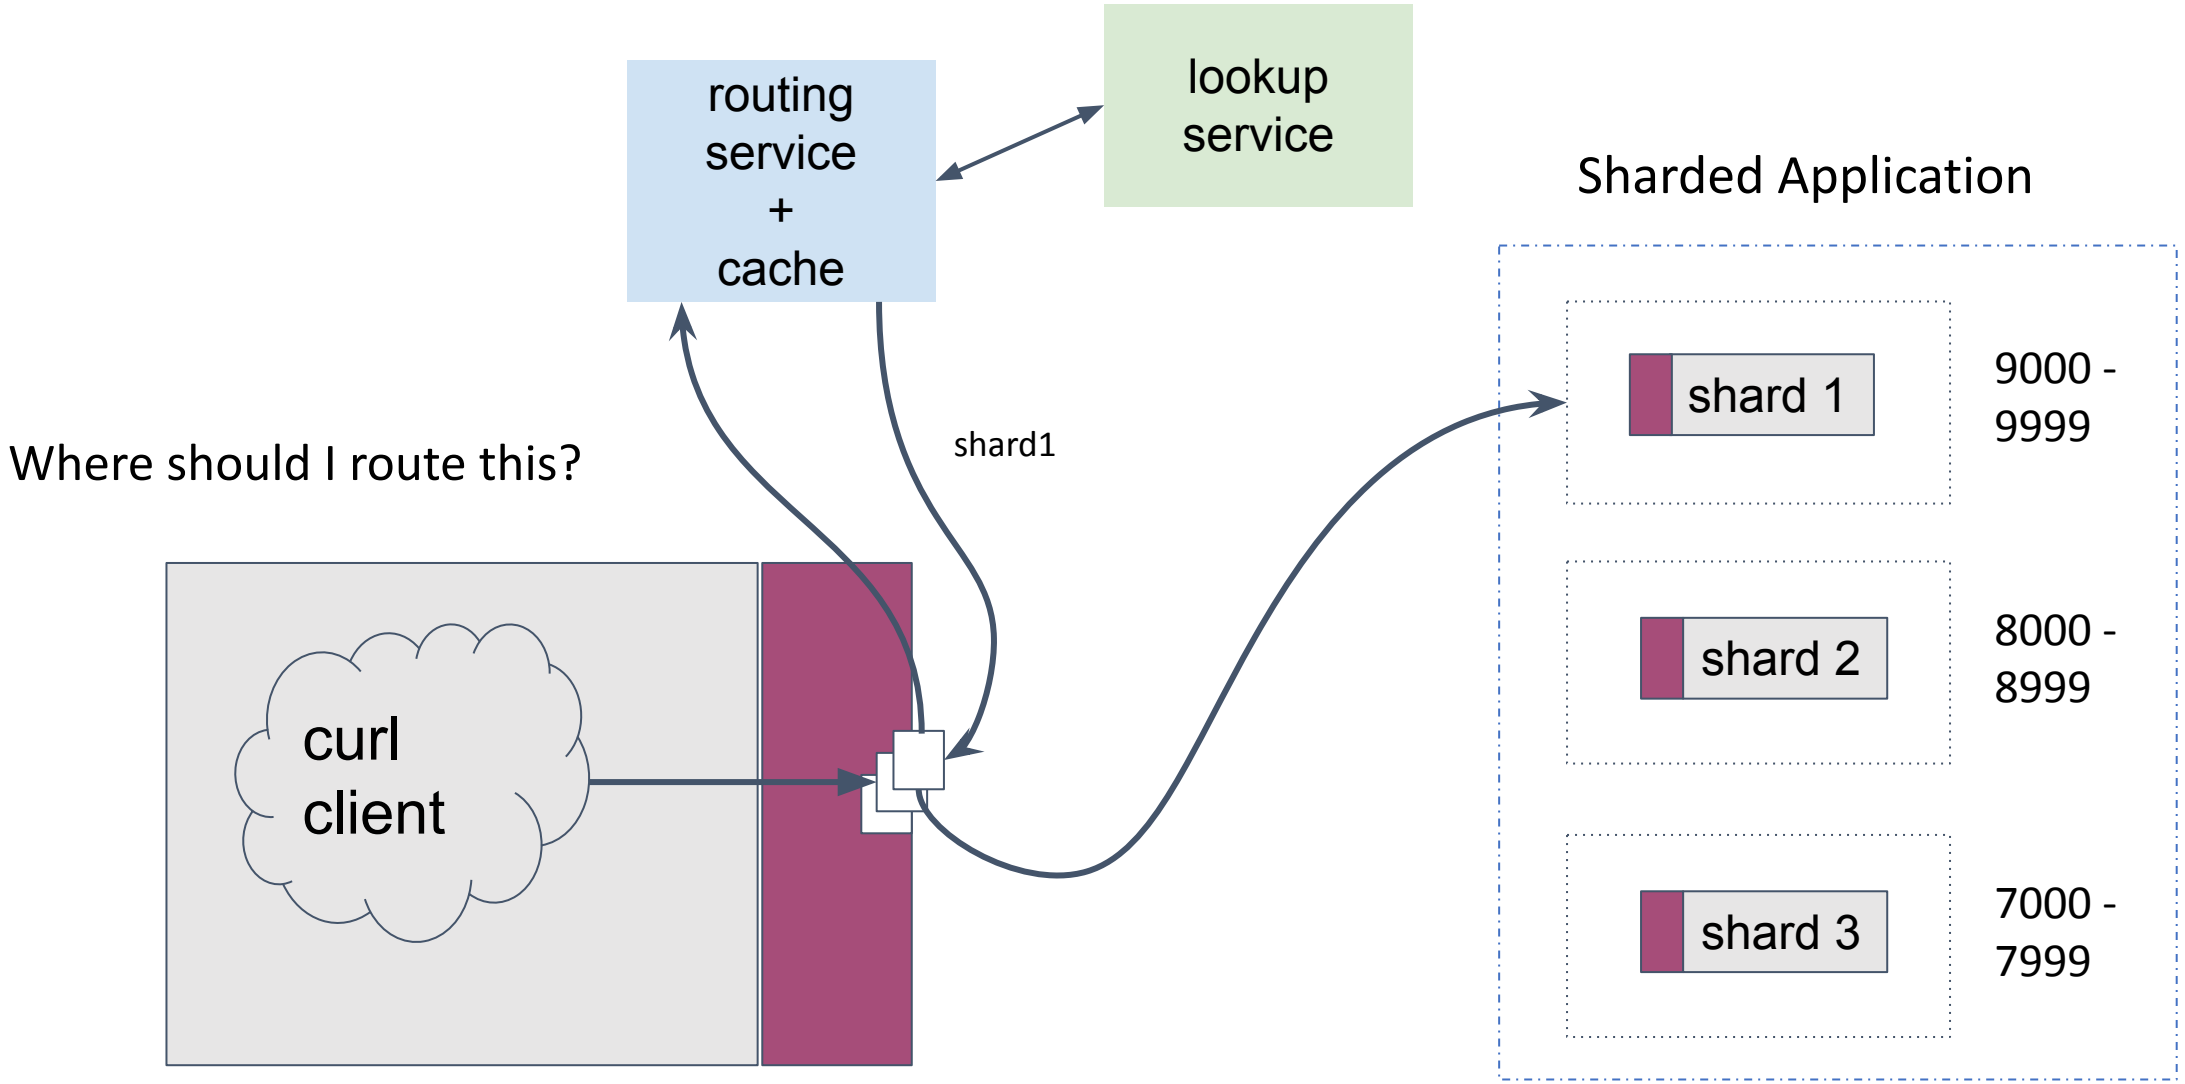
routing service + cache
lookup service
Sharded Application
Where should I route this?
shard1
curl client
shard 1
9000 - 9999
shard 2
8000 - 8999
shard 3
7000 - 7999

Future Work

- Adding rate limiting for services using this approach
- Work with Istio community to allow workload selector to be applied to multiple pods
- Explore using C++ and rust

Q/A

KubeCon | CloudNativeCon
North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Please scan the QR Code above to
leave feedback on this session