

# DoorDash's Journey From StatsD To Prometheus With 10 Million Metrics/Second

Benjamin Raskin  
Emma Wang

# Introductions

# Who are we?

**Ben Raskin**



Solutions Architect



**Emma Wang**



Infrastructure Software Engineer



# Agenda

1. StatsD vs Prometheus
2. Client Side Migration
3. Enablement
4. Learnings/Results

# Challenges with StatsD

# Challenges with StatsD

- **Lack of naming standardization**/limited support for tags
  - Service A: service\_a.page\_views.us-east-1
  - Service B: service\_b.us-east-1.page\_views
- Number of metrics **scales with user traffic**
  - Requires an aggregation tier; otherwise metrics can grow exponentially
- **Sent over UDP** (potential for packet loss)
- **Delta counters**
  - No way to interpolate missing data points
- **Inability to aggregate percentiles**/lack of histograms



# Why Prometheus?

# DoorDash Observability Requirements



## Use Open Source

- Cost efficiency
- Integration
  - Open source data formats and open source query languages.
- No vendor lock-in



## Governance and Control

- Standard Conventions
  - Standards for common tags, metrics naming conventions.
- Cost accounting
  - Ability to breakdown usage by team, service, etc.



## Self-Service

- Productivity Empowerment
  - For newly registered services, the metrics can be automatically discovered and collected.
  - Automate the process of generating basic dashboards and alerts.

# Why DoorDash Chose Prometheus

- Prometheus has emerged as the dominant standard for open source metrics and aligned well with the organization's strategy and requirements.
  - Open Source Standard
  - Query Language (PromQL)
  - Tag-Based Metric Ingestion Format
  - Pull-based system (allows for client-side aggregation)
  - Native support for histograms and cumulative counters
- Strong support for the more fundamental tools used in DoorDash. eg: [Kubernetes](#)

# Client Side Migration

# Metric generation



# Custom Metrics

- Application and service metrics
  - microservices use prometheus native libraries to generate metrics.
  - pay attention to python clients with a multi-process pattern.
- Internal libraries
  - provide the internal library to generate `HTTP/GRPC` metrics automatically.
  - `JVM` prometheus exporter is included in the doordash docker base images to export `JVM` metrics.

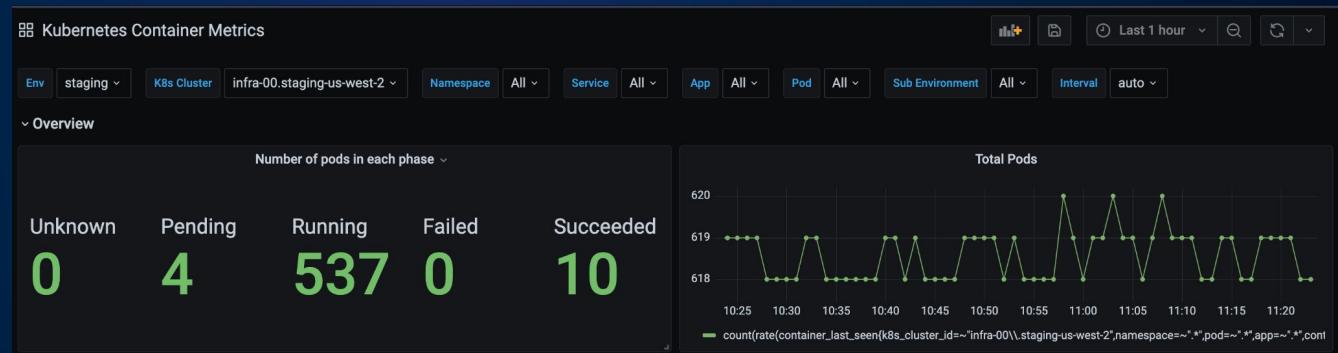
# Some Other Metrics

- Open Source Exporters
  - Wide-range of available exporters maintained by the open-source community
  - AWS Cloudwatch Metrics: official prometheus cloudwatch exporter  
VS YACE cloudwatch exporter
  - StatsD exporter
  - JVM exporter
  - Pgbounce exporter
- Short-Lived Jobs
  - can not be scraped by metrics collectors
  - use the [prom aggregation gateway](#) with a small modification

# Standard Tags For All Metrics

## Common Tags:

- service
- app
- k8s\_cluster
- environment
- sub\_environment
- region
- zone
- cost\_origin



# Service Discovery (thru standardization)

- Service discovery done through Kubernetes annotations, which tell Prometheus what endpoints to scrape.
  - Central Observability Team (COT) creates gold-standard Kubernetes manifest templates, which services use.
- Default scrape options are also defined by the COT.
  - Scrape frequency
  - Standard labels (prev slide)

# Reducing Cardinality & Improving Performance

## Relabel rules

- Rewrite the label set of a target.
- Consolidate common labels and drop unused metrics or labels.

## Rollup rules using M3

- Eliminate labels that are not needed by the service owners.
- For example, the instance or pod id can be removed.

## Mapping rules using M3

- Define the storage policy for each time series.
- Drop unnecessary metrics after rollups are applied.

## Recording rules

- Pre-compute expensive expressions to speed up query time.

During the  
scrape event

Server side  
aggregation

Server side  
aggregation

Post Storage

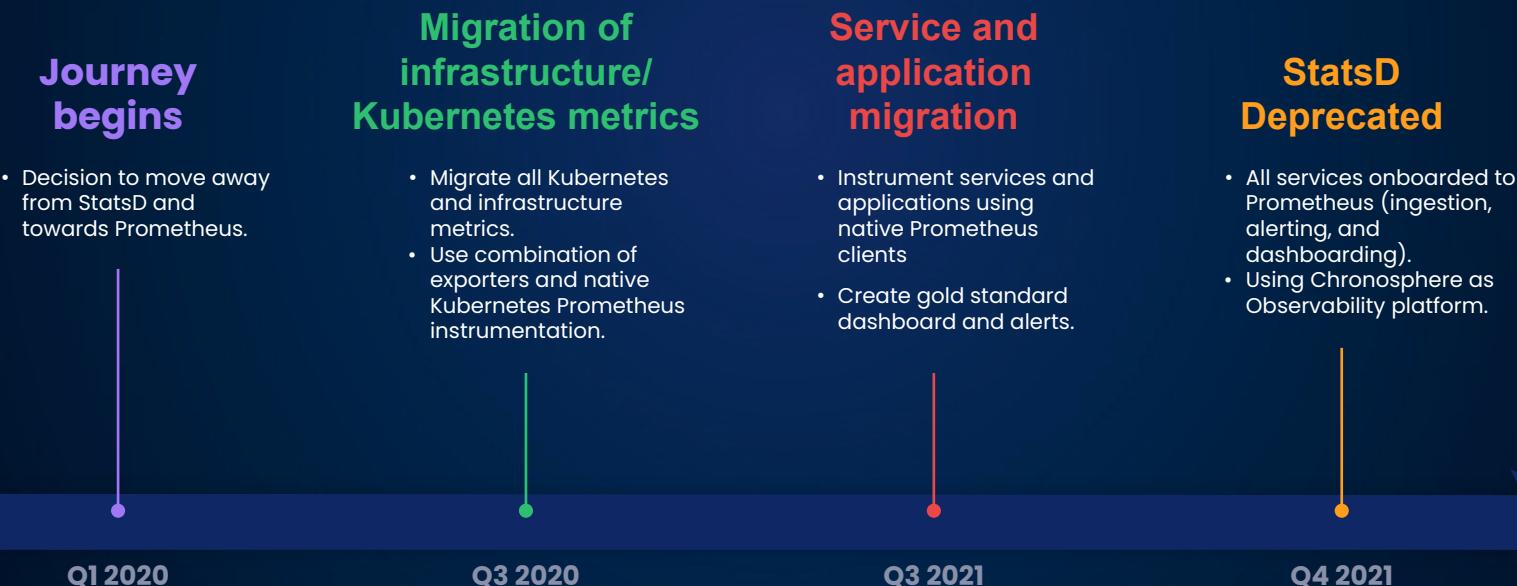
# Enablement

# Learning (and teaching) to fish

- Prometheus data model
  - Best practices!
  - How to name things...
- PromQL
  - Counters (rate, rate, rate!)
  - Histograms vs. timers
- Not everything needs to be built from scratch
  - Open-source and gold standards
- Managing resources through Terraform

# Timeline, Results, and Learnings

# Timeline



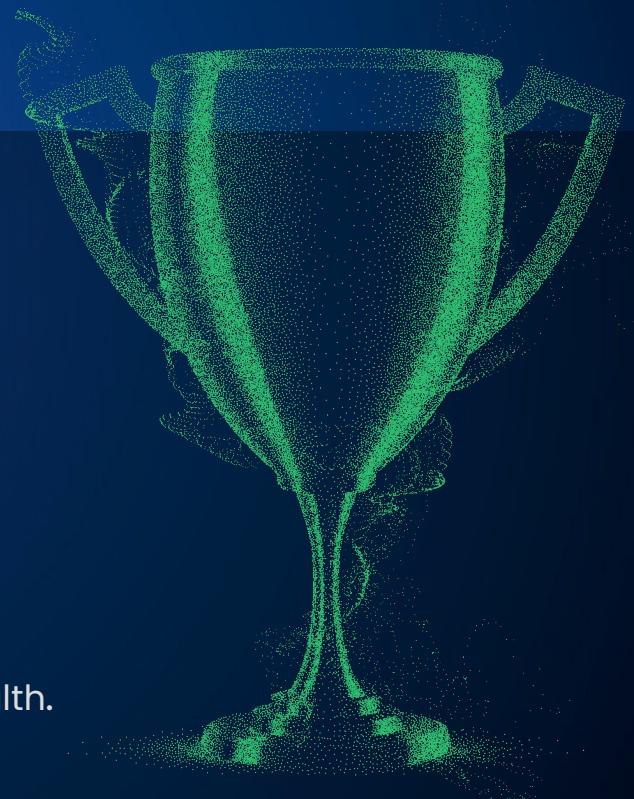
# Achievements and Wins

starting with more than **7000** alerts, **1500** dashboards and metrics for **130** services, this has been a massive collaboration across all teams and engineers.

One year post migration, we now have:

- ~27k alerts
- ~2.2k dashboards
- 15M metrics/s ingested
- 10M metrics/s persisted (after aggregation)

The Prometheus based monitoring system is stable, scalable, and flexible, providing a greatly improved view of DoorDash service health.



# Learnings

Switching from a percentile to a histogram is tough

The instance label is an important concern for overall volume of metrics expected after the move

Prometheus is not friendly with short lifecycle jobs

Automate the monitoring onboarding process for teams



Come by Booth G15 to:

- Win the Ultimate Gaming Setup
- Test your gaming skills at OIIY Legend
- Win daily prizes at 4pm
- Talk to the best in observability

Trusted by the best  
in the industry

Abnormal

Snap Inc.

Robinhood

DOORDASH



Scan the QR code

# Thank You



# Use Codification For Alerts

Choose **Terraform** based on:

- source control
- Locks on state files
- Alert templating

Some caveats:

- State management
- Performance at scale