



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

KEDA -

Real Time And Serverless Scaling In Kubernetes

Zbynek Roubalik, Principal Software Engineer @ Red Hat
Jeff Hollan, Director of Product @ Snowflake

- Jeff Hollan

- Director of Product @ Snowflake
- KEDA founding member / maintainer while at Microsoft (Director of product for serverless / containers)
- <https://twitter.com/jeffhollan>
- <https://linkedin.com/in/jeffhollan>

- Zbyněk Roubalík

- Principal Software Engineer @ Red Hat
- KEDA founding member / maintainer, Knative (TOC), Microsoft MVP
- <https://github.com/zroubalik>
- <https://twitter.com/zroubalik>
- <https://www.linkedin.com/in/zbynek-roubalik/>

Agenda

1. What is KEDA?
2. KEDA project and community
3. Demo!
4. KEDA concepts and architecture
5. Future development
6. Q&A

A tale of two scalings....

Goal: Provide enough pizza for everyone at a KubeCon party 🍕


Strategy #1: Start with 1 pizza, see if it gets eaten, and add another

Strategy #2: Find out how many people accepted the invite.
Bring the amount of pizzas estimated to feed the group

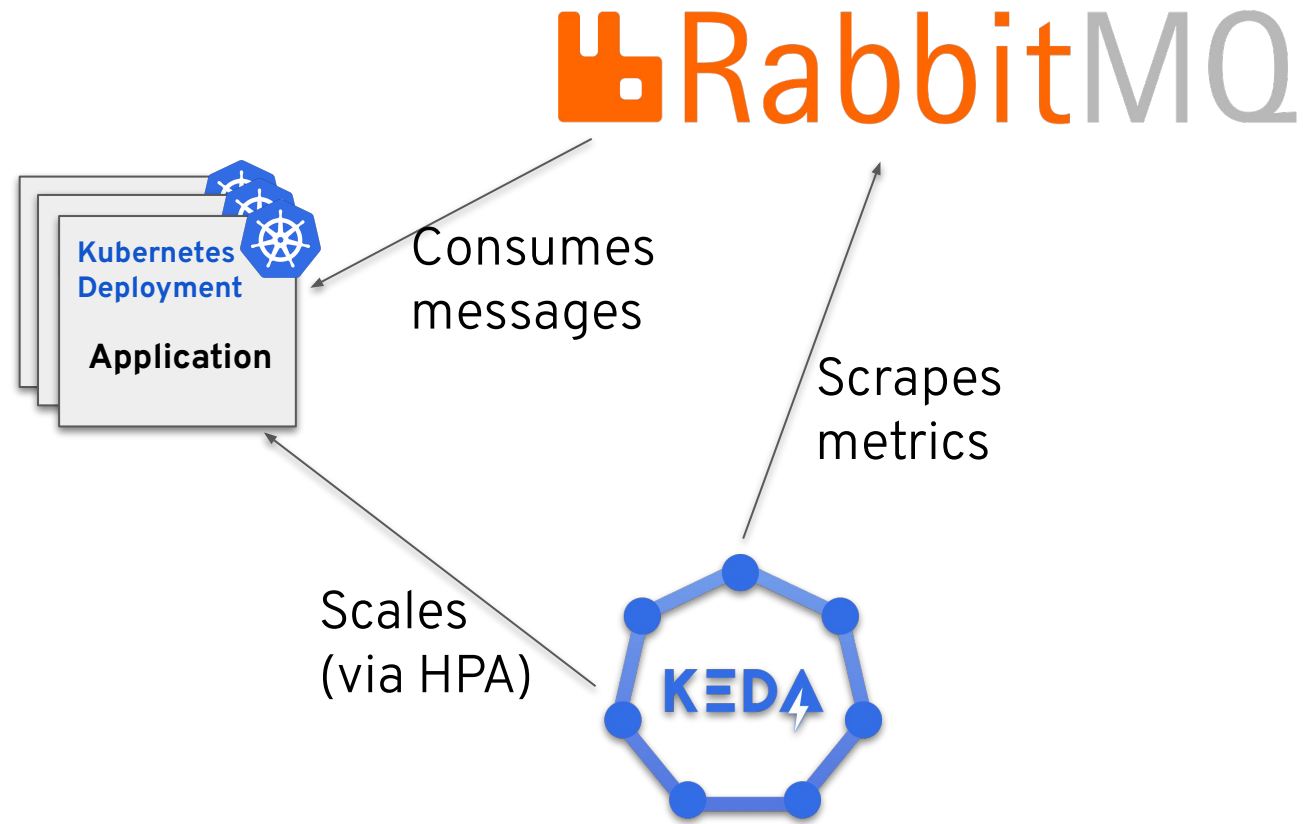


- Project aims to make **K**ubernetes **E**vent **D**riven **A**utoscaling dead simple
- Allows you to scale any deployment resource or job based on **events**, not only on CPU / Memory
- 55+ integrated event sources (Prometheus, RabbitMQ, Kafka, SQS, PostgreSQL,)
- <https://keda.sh>

Hello KEDA Demo

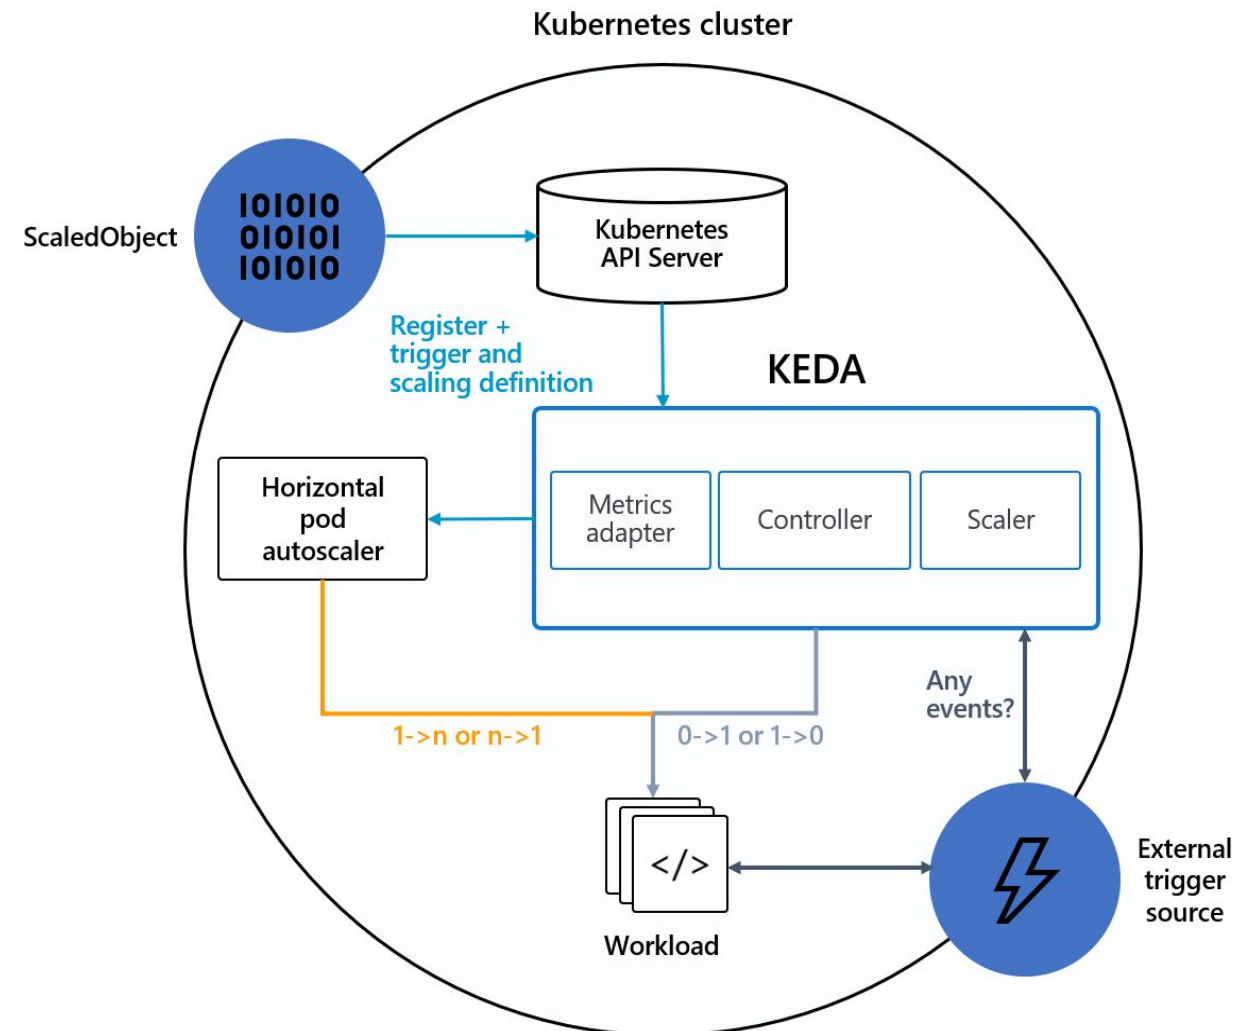


Hello KEDA Demo



Architecture

- KEDA is built on top of Kubernetes
- Use **ScaledObject/ScaledJob** to define scaling metadata
- Manages workloads to scale to 0
- Registers itself as Kubernetes Metric Adapter
- Publishes metrics to the Horizontal Pod Autoscaler (HPA) which makes most scale decisions



ScaledObject

- Can target **Deployment**, **StatefulSet** or **Custom Resource** with **/scale**
- **Multiple scalers** can be defined as triggers for the target workload
- User can specify **HPA related settings** to tweak the scaling behavior

```
apiVersion: keda.sh/v1alpha1
kind: ScaledObject
metadata:
  name: example-so
spec:
  scaleTargetRef:
    name: example-deployment
  minReplicaCount: 0
  maxReplicaCount: 100
  triggers:
  - type: kafka
    metadata:
      bootstrapServers: kafka.svc:9092
      consumerGroup: my-group
      topic: test-topic
      lagThreshold: '5'
```

ScaledJob

- Schedule **Kubernetes Job** based on events
- Useful option to handle **processing long running executions**

```
apiVersion: keda.sh/v1alpha1
kind: ScaledJob
metadata:
  name: example-sj
spec:
  jobTargetRef:
    ... # standard k8s Job definition

  maxReplicaCount: 100
  triggers:
  - type: kafka
    metadata:
      bootstrapServers: kafka.svc:9092
      consumerGroup: my-group
      topic: test-topic
      lagThreshold: '5'
```

KEDA vs Custom Metrics Adapter

```
apiVersion: keda.sh/v1alpha1
kind: ScaledObject
metadata:
  name: example-scaled-object
spec:
  scaleTargetRef:
    name: example-deployment
  minReplicaCount: 0
  maxReplicaCount: 100
  triggers:
  - type: kafka
    metadata:
      bootstrapServers: kafka.svc:9092
      consumerGroup: my-group
      topic: test-topic
      lagThreshold: '5'
```

```
apiVersion: autoscaling/v2beta2
kind: HorizontalPodAutoscaler
metadata:
  name: web-prometheus-adapter
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: web
  minReplicas: 1
  maxReplicas: 50
  metrics:
  - type: Object
    object:
      metric:
        name: http_requests_per_second_per_pod
      describedObject:
        apiVersion: v1
        kind: Service
        name: web-prometheus-adapter-service
    target:
      type: Value
      value: 25
```

```
rules:
  default: false
  custom:
  - seriesQuery: 'http_requests_received_total'
    resources:
      overrides:
        namespace: {resource: "namespace"}
        service: {resource: "service"}
    name:
      matches: ""
      as: "http_requests_per_second_per_pod"
    metricsQuery:
      'max(irate(<<.Series>>{<<.LabelMatchers>>}[1
m])) by (<<.GroupBy>>)'
```

Advanced features

- Ability to specify **Fallback replicas count** - in case of problems
- Users can still tweak HPA settings if they want to (**scaling behavior**)
- Ability to **Pause autoscaling**
- KEDA exposes Prometheus metrics
- Users can extend KEDA implementing **External scalers** via gRPC interface or **Metrics API scalers** via Rest API.
- <https://keda.sh/docs/2.8/concepts/scaling-deployments/#scaledobject-spec>

Authentication

- Re-use trigger authentication across ScaledObject/ScaledJobs with **TriggerAuthentication** (namespaced) or **ClusterTriggerAuthentication**
- Out-of-the-box integration with sources such as:
 - Environment variables (on scale target)
 - Kubernetes secrets
 - Pod Identity (“No secret authentication” - Azure / AWS)
 - HashiCorp Vault
 - Azure Key Vault

What is next?

- Cache metrics values in KEDA Metrics Server
 - Reduced load on the source of metrics (eg. Prometheus server...)
 - Smoother autoscaling (apply AI/ML model to incoming metrics)
- Custom logic for evaluation when there are multiple triggers in ScaledObject
- CloudEvents integration
- OpenTelemetry integration
- Open interface for Predictive autoscaling

What is next?

- Environmental impact
 - CNCF TAG Environmental Sustainability
 - Integrate CO2 emission intensity, power consumption,... data into KEDA scaling decision
 - Carbon aware autoscaling
 - - [POC](#) & [recording](#)

```
apiVersion: keda.sh/v1alpha1
kind: ScaledObject
metadata:
  name: my-scaled-object
spec:
  scaleTargetRef:
    name: my-resource
  minReplicaCount: 1
  maxReplicaCount: 100
  environmentalImpact:
    carbon:
      - carbonIntensity: 90
        allowedMaxReplicaCount: 81
      - carbonIntensity: 100
        allowedMaxReplicaCount: 53
      - carbonIntensity: 110
        allowedMaxReplicaCount: 44
  triggers:
    - type: kafka
      metadata:
        bootstrapServers: kafka.svc:9092
        consumerGroup: my-group
        topic: test-topic
        lagThreshold: '5'
```




Please scan the QR Code above to
leave feedback on this session