



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Like Peas and Carrots:

Argo CD and Crossplane for Infrastructure Management

Jesse Suen, Akuity

Victor Farci, Upbound



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

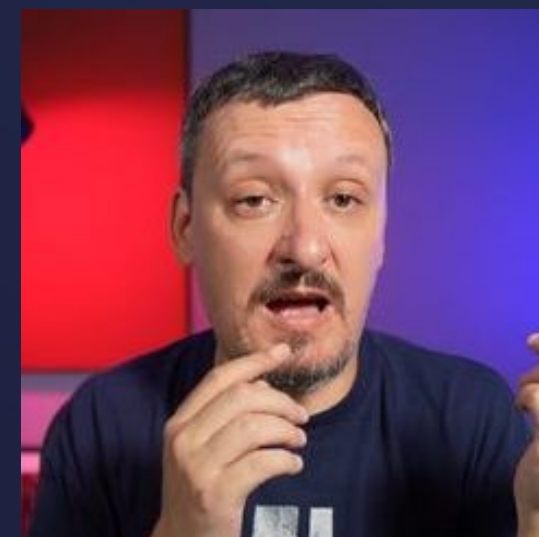
October 24-28, 2021



Jesse Suen

CTO

Akuity



Victor Farcic

Developer Advocate

Upbound

How We Got Here?



Configuration Management



CFEngine



CHEF™



puppet



ANSIBLE

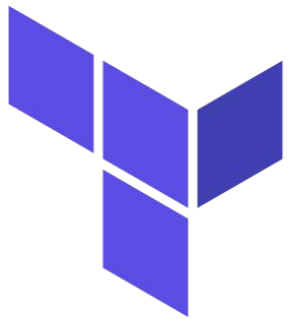
Mutable - Baremetal



Infrastructure as Code (IaC)



AWS CLOUDFORMATION



HashiCorp

Terraform



Pulumi

Immutable - Reactive



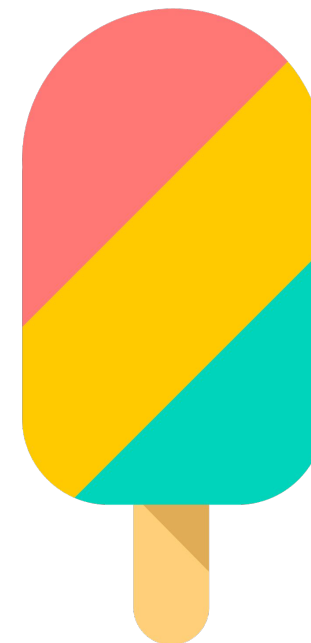
Control Planes - Now

- API
- CRDs
- Continuous drift-detection and reconciliation
- Ecosystem



Custom Resource Definitions (CRDs)

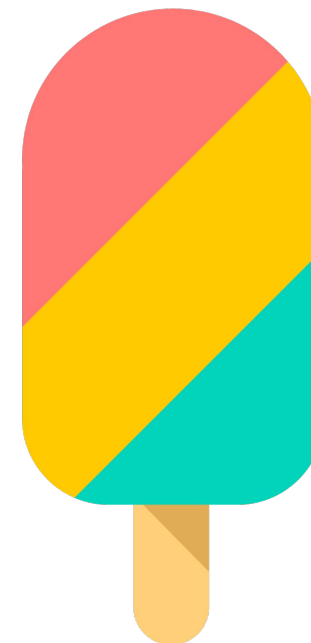
```
apiVersion: ec2.aws.crossplane.io/v1alpha1
kind: Instance
metadata:
  name: dot
  labels:
    app: dot
spec:
  forProvider:
    imageId: ami-052efd3df9dad4825
    region: us-east-1
    instanceType: t2.micro
    subnetIdSelector:
      matchLabels:
        app: dot
```



Crossplane

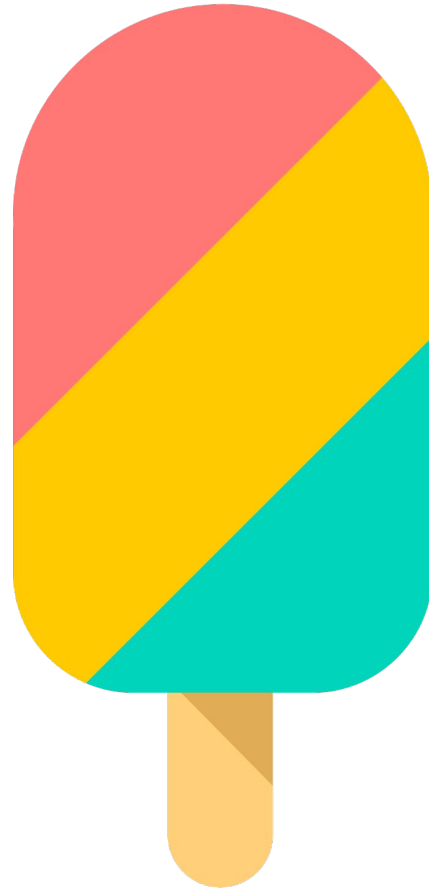
Composite Resources (XRs)

```
apiVersion: devopstoolkitseries.com/v1alpha1
kind: ClusterClaim
metadata:
  name: a-team-eks
spec:
  id: a-team-eks
  compositionSelector:
    matchLabels:
      provider: aws
      cluster: eks
  parameters:
    nodeSize: medium
    minNodeCount: 3
  writeConnectionSecretToRef:
    name: a-team-eks
```



Crossplane

Internal Developer Platforms (IDPs)



Crossplane

Why Crossplane?

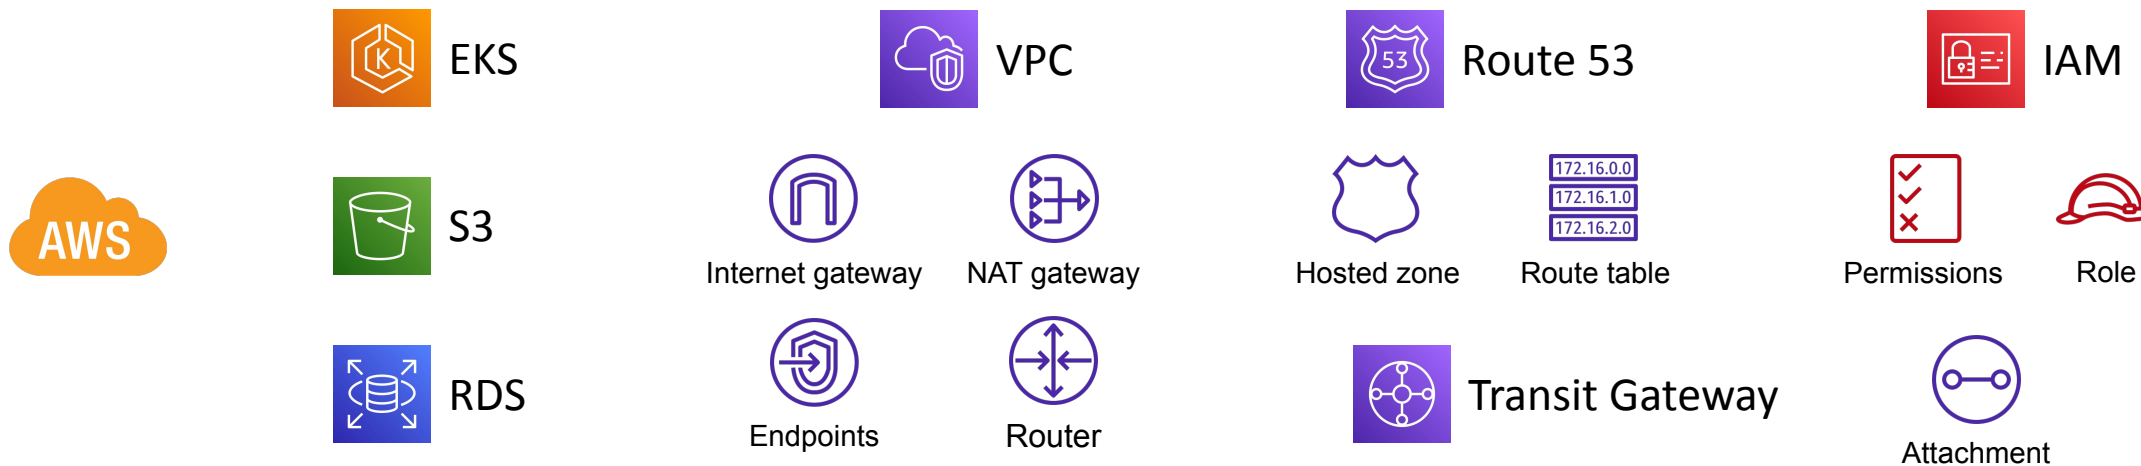
- Manage infrastructure the same way as our apps (Argo CD)
- Coordination between infra and apps (IAM)
- Want standardized, self-service infra with simplified knobs
- Want to treat our clusters and infra as cattle

Akuity: How we use Crossplane



Providers:

- aws-provider
- helm-provider
- kubernetes-provider



Akuity Compositions

EKS

Cluster	
Roles	Policies
OIDCProvider	NodeGroup
KMS Key	KMS Alias
Addons	AkuityAddons

IRSA

Policy
Role
RolePolicyAttachment

RDS

DBCluster	DBInstance
DBSubnetGroup	SecurityGroup
KMS Key	KMS Alias
DBClusterParameterGroup	
DBParameterGroup	

EKSNetwork

VPC	
Gateway	InternetGateway
3 Public Subnets	3 Private Subnets
3 NAT IPs	3 NAT Gateways
Routes	RouteTables

RDSNetwork

VPC	
3 Private Subnets	RouteTables

Akuity Addon Compositions



Karpenter

IRSA Helm Release
InstanceProfile Provisioner



Cert-Manager

IRSA
Namespace ServiceAccount



External Secrets

IRSA Helm Release
ClusterSecretStore



AWS Load Balancer Controller

IRSA Helm Release

AWS Distro Open Telemetry (ADOT)

IRSA
Namespace ServiceAccount



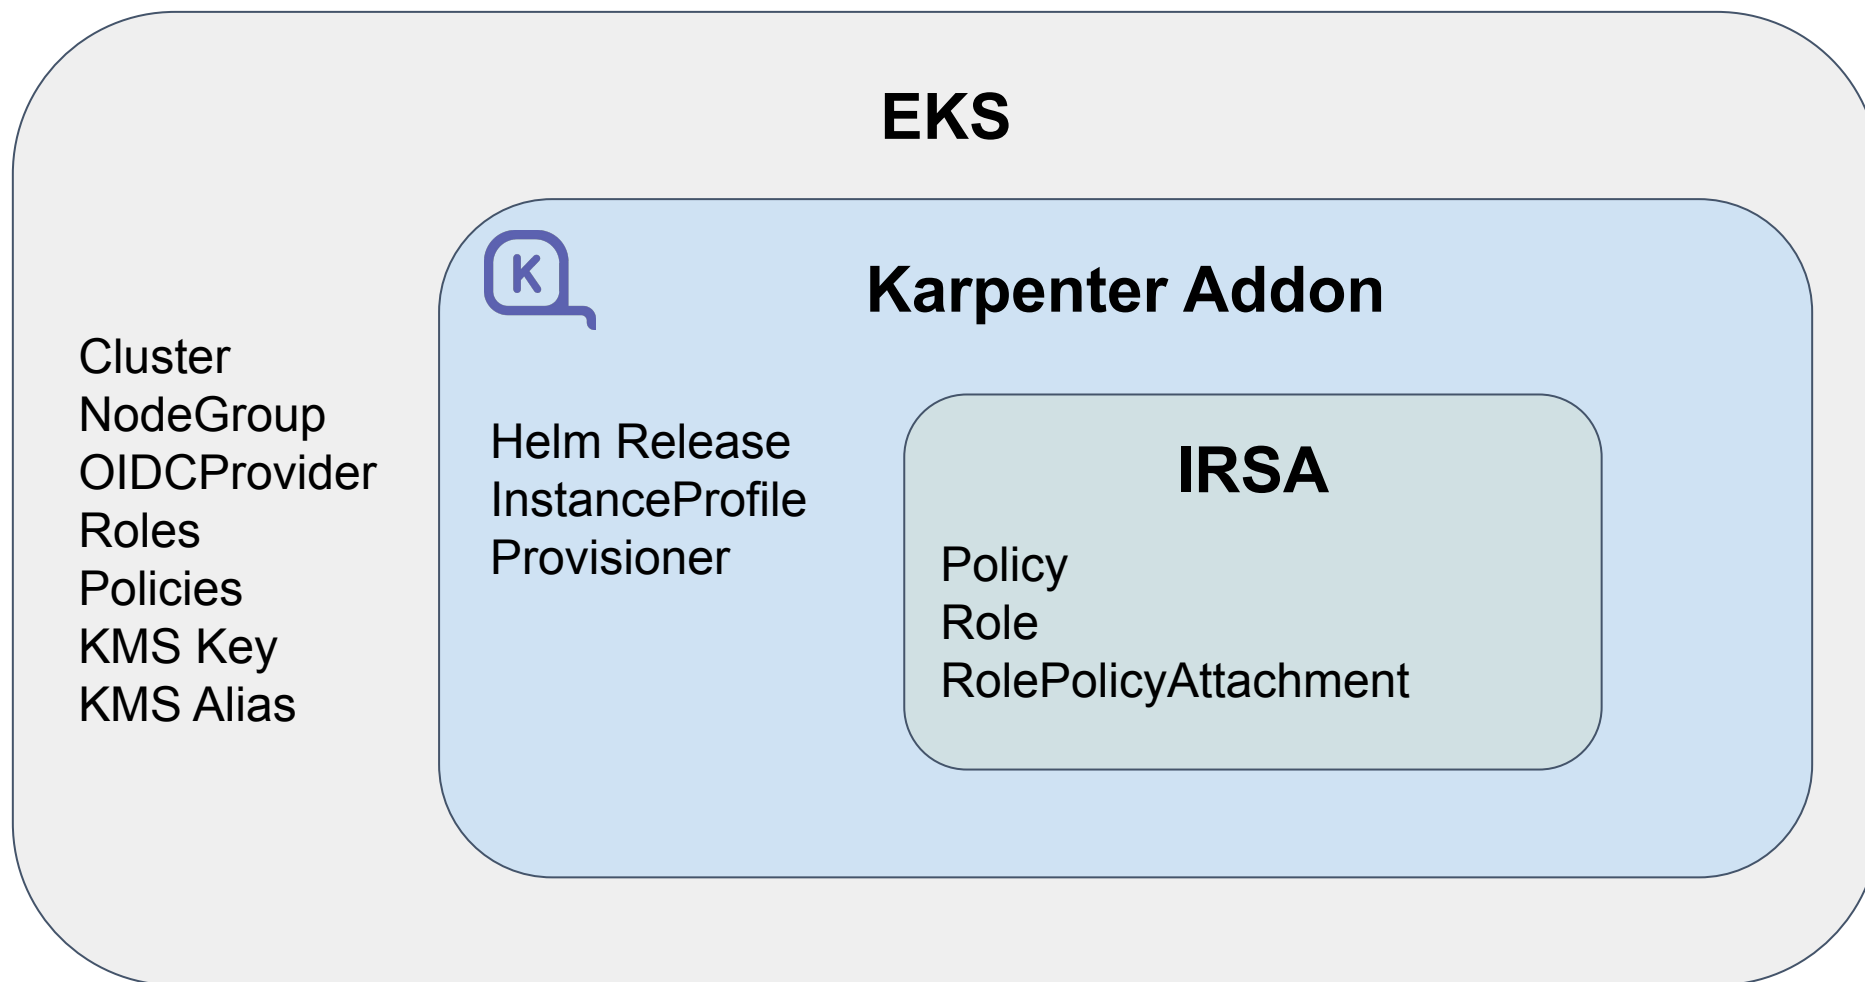
External DNS

IRSA ConfigMap
Namespace ServiceAccount

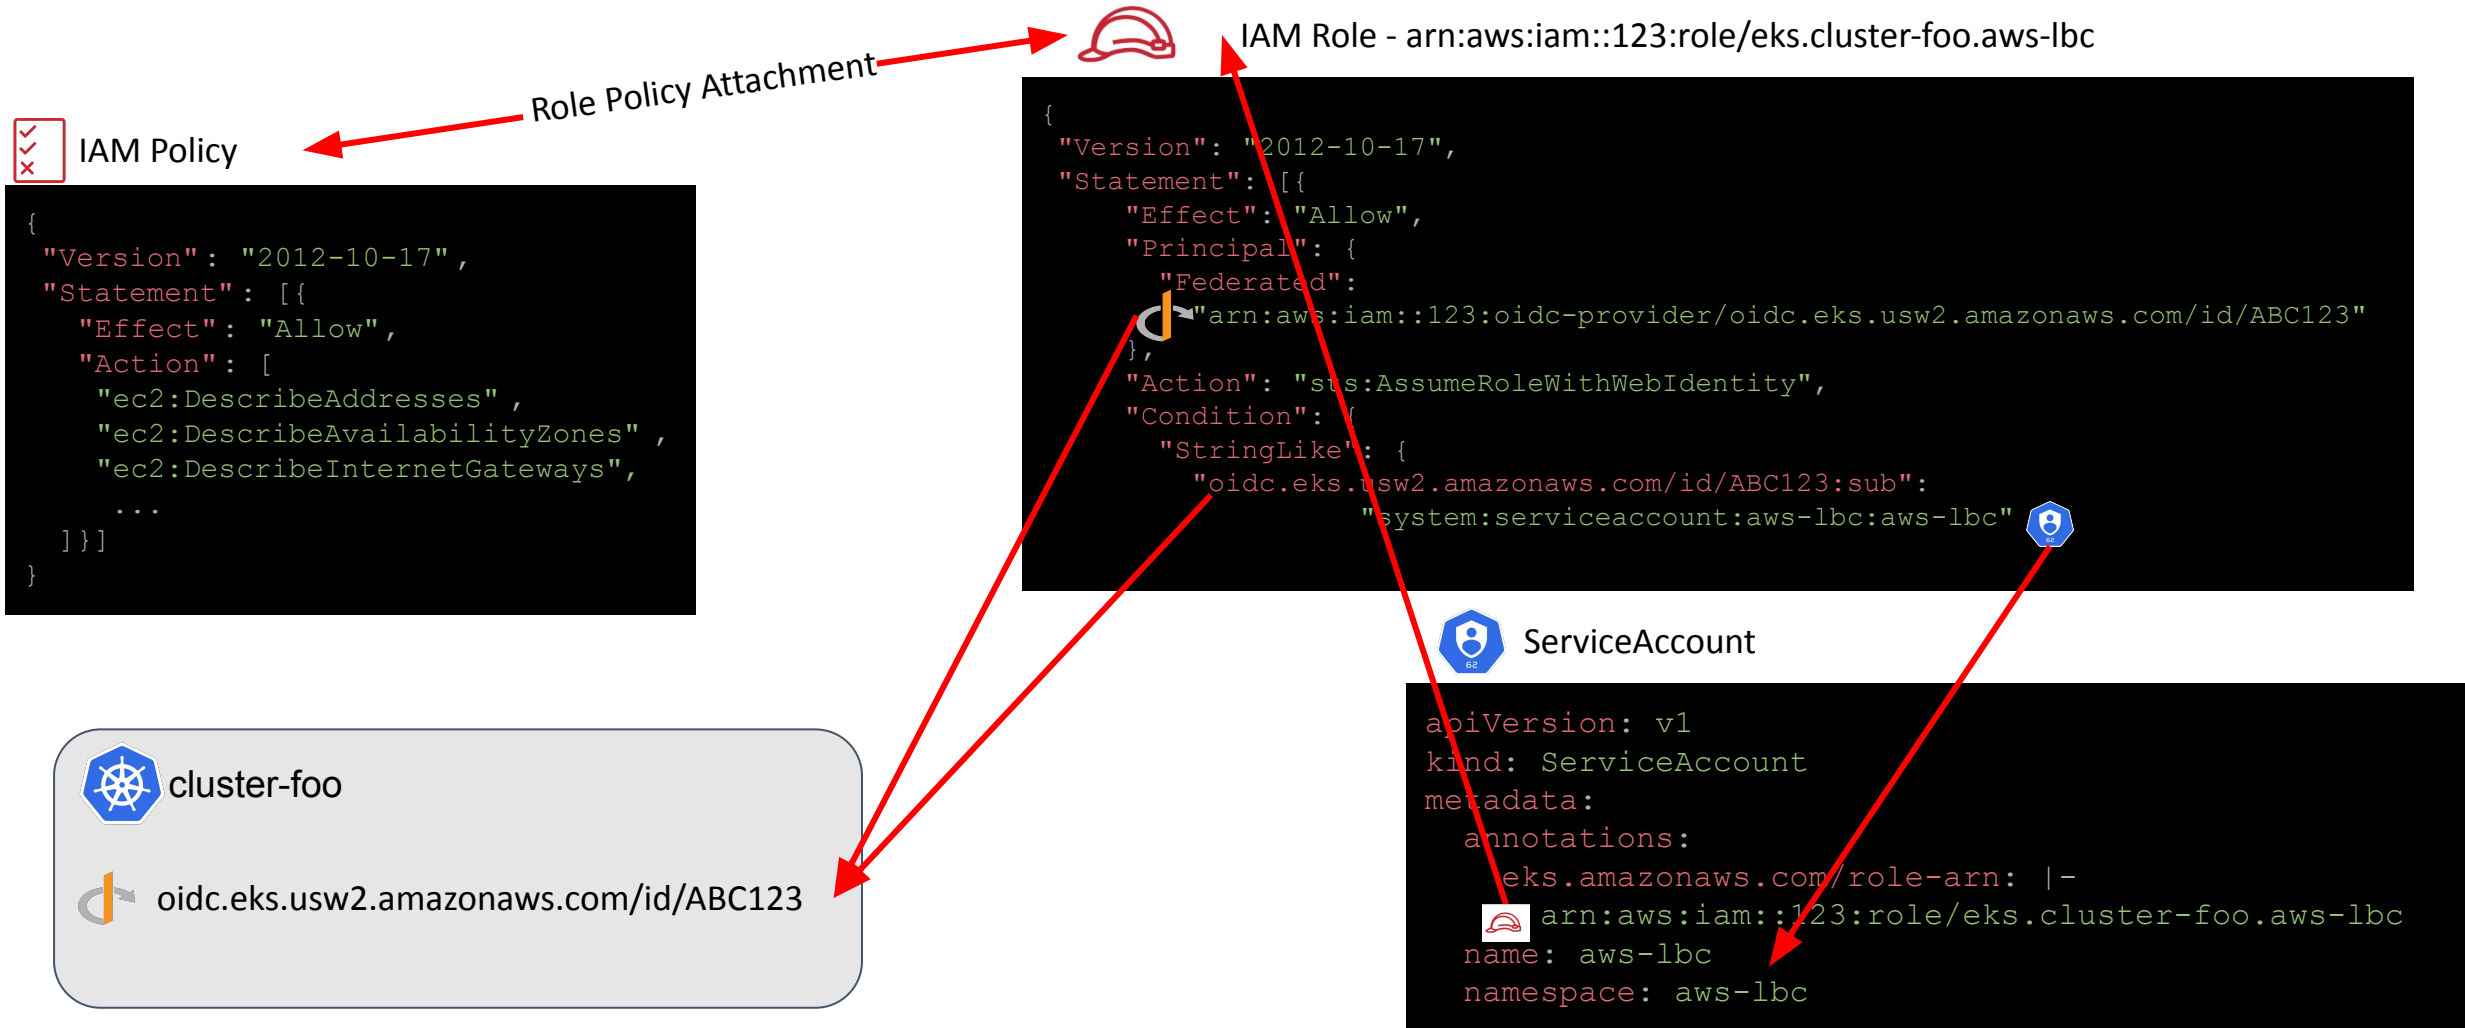
Reduced knobs for simplified experience

```
apiVersion: addons.kuivity.io/v1alpha1
kind: Karpenter
metadata:
  name: mycluster
spec:
  capacityType: spot
  instanceTypes: [m5.large, m5.xlarge, c5.large, c5.xlarge]
  limits:
    cpu: "100"
    ...
```

Akuity EKS Addon Compositions

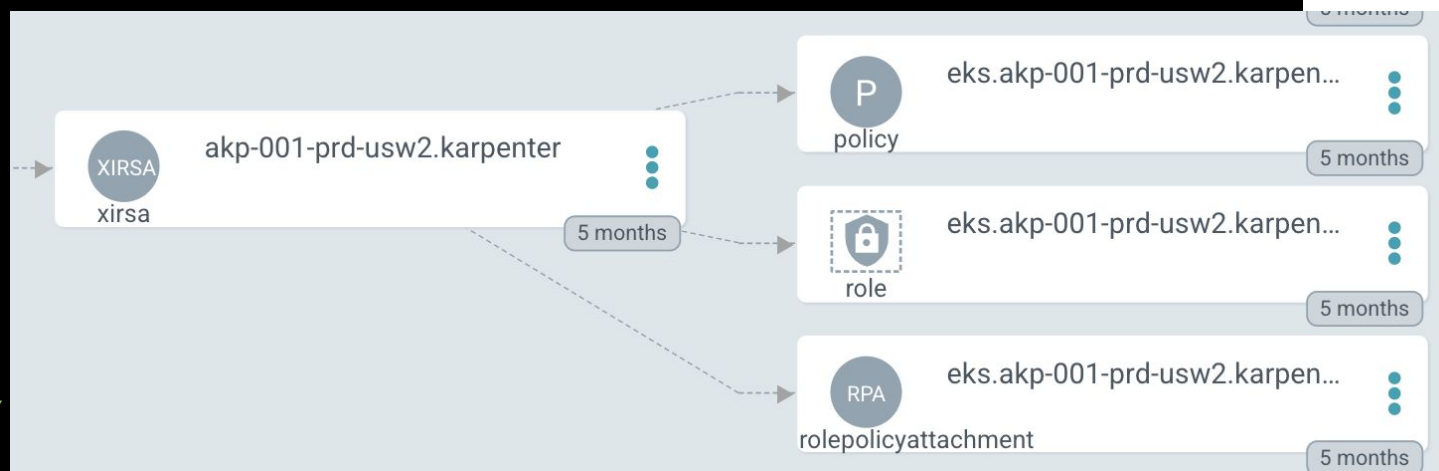


Challenge: Cross referencing between AWS / Kubernetes



EKS - IRSA Composition

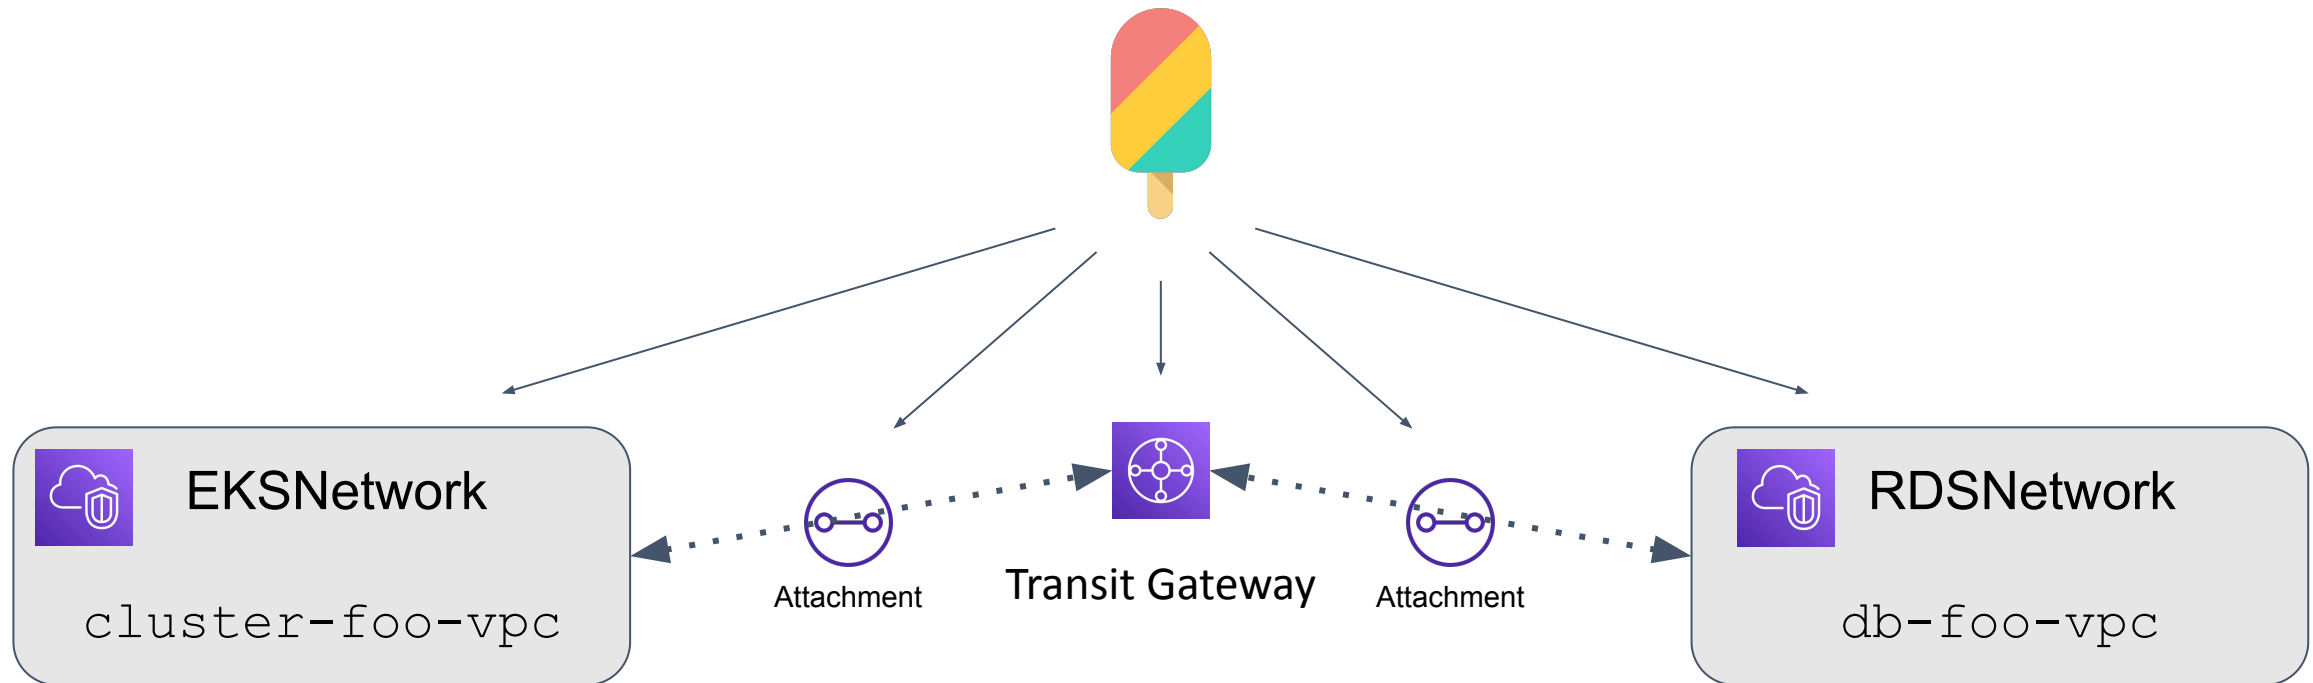
```
apiVersion: aws.akuity.io/v1alpha1
kind: IRSA
metadata:
  name: cluster-foo.aws-lbc
spec:
  serviceName: aws-lbc
  namespace: aws-lbc
  clusterRef:
    endpoint: https://ABCD1234.gr7.usw2.eks.amazonaws.com
    id: cluster-foo
    oidc:
      provider: oidc.eks.usw2.amazonaws.com/id/ABCD1234
      providerARN: arn:aws:iam::123:oidc-provider/oidc.eks.usw2.amazonaws.com/id/ABCD1234
  policyDocument: |
    {
      "Version": "2012-10-17",
      "Statement": [
        {
          "Effect": "Allow",
          "Action": [
            "ec2:DescribeAddresses",
            "ec2:DescribeAvailabilityZones",
            ...
```



Managed Resources

Also deploy managed resources directly (not in a Composition):

- e.g. Route53, Routes, TransitGateways, Bespoke IAM



Crossplane vs. Argo CD for Addons



Pros:

- Cluster works out of the box with no extra steps
- IAM can be coordinated

Cons:

- Cumbersome addons upgrades: require a new composition version
- All clusters updated simultaneously



Pros:





- GitOps-based update process
- Finer control of versions in envs
- Easy to modify ConfigMap/Secrets
- Addons can be made optional

Cons:

- Harder to coordinate with IAM
- Extra step after cluster creation

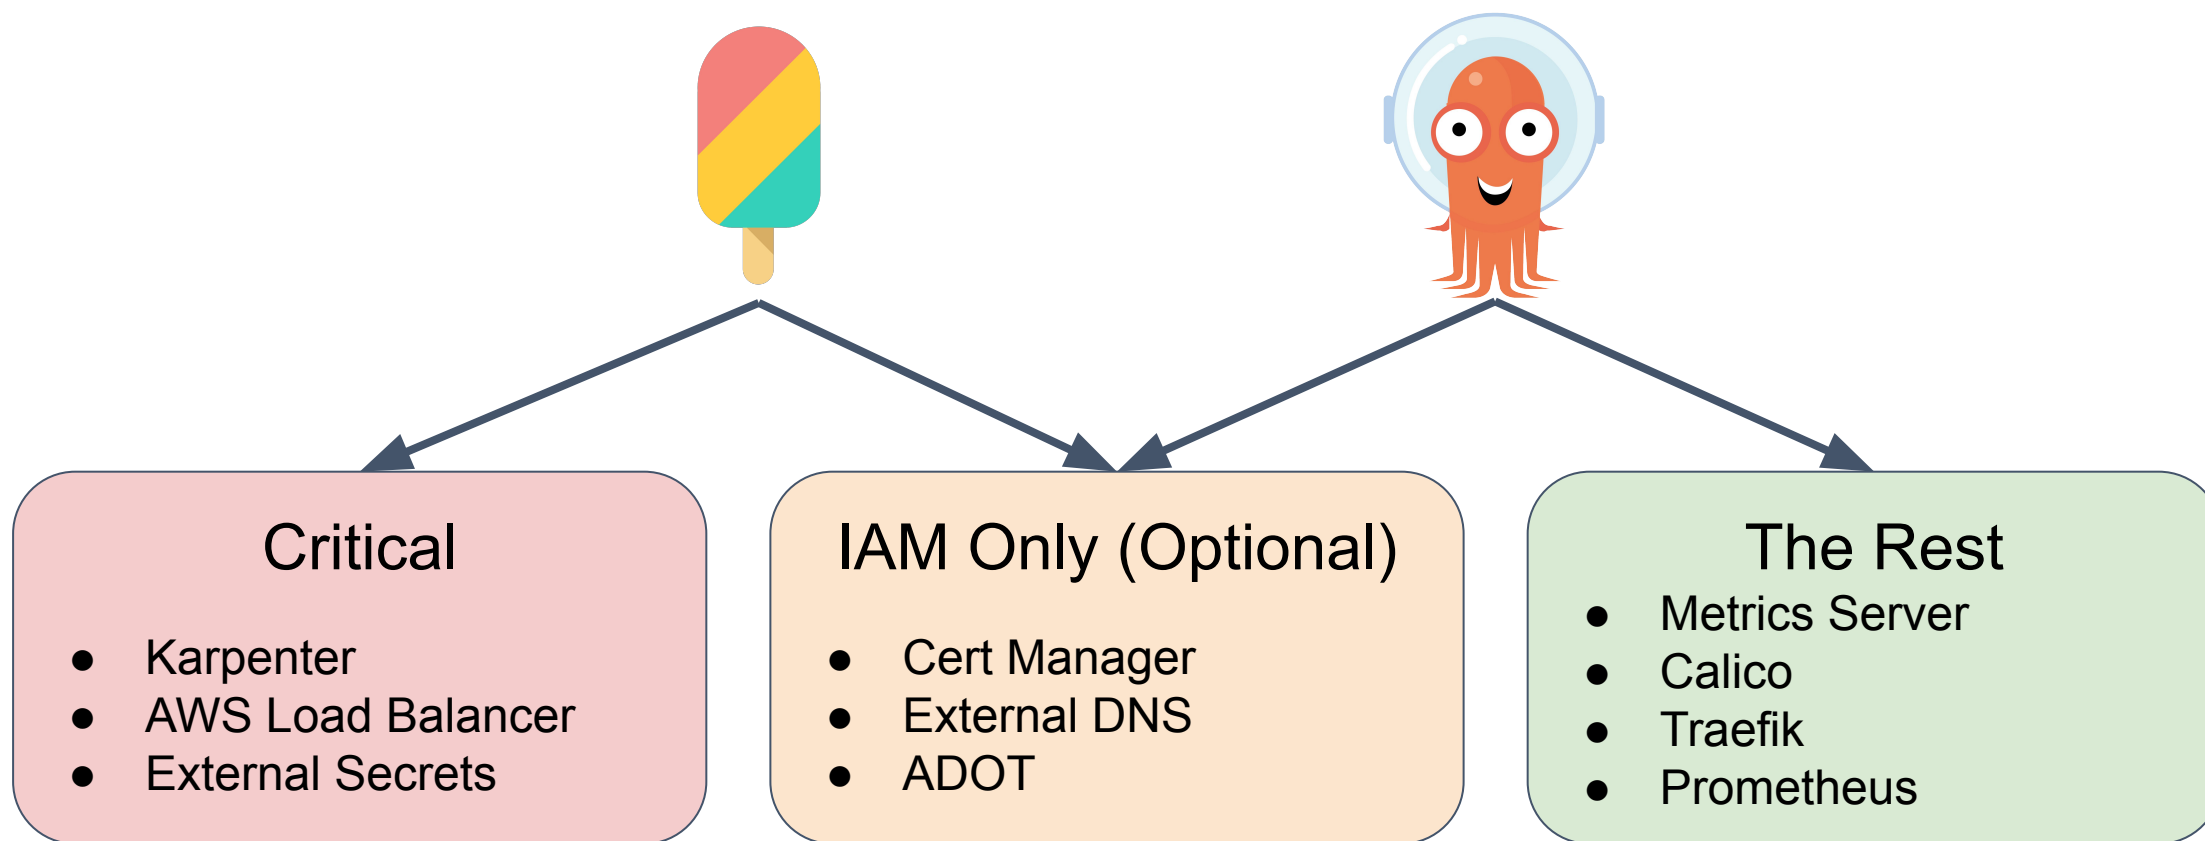
Crossplane vs. Argo CD for Addons

Solution: Use Crossplane for some addons, Argo CD for others

Type	Description	Managed by
Critical	Essential for cluster to function (autoscaling). Fully managed by Crossplane.	
IAM Only	Optional addons requiring AWS resources. IAM managed by Crossplane. Deployment/Config managed by Argo CD.	 
Other	No dependency to AWS. Normal applications, fully managed by Argo CD	

Crossplane vs. Argo CD for Addons

Crossplane vs. Argo CD for Kubernetes Resources



Best Practice: Controlling Blast Radius

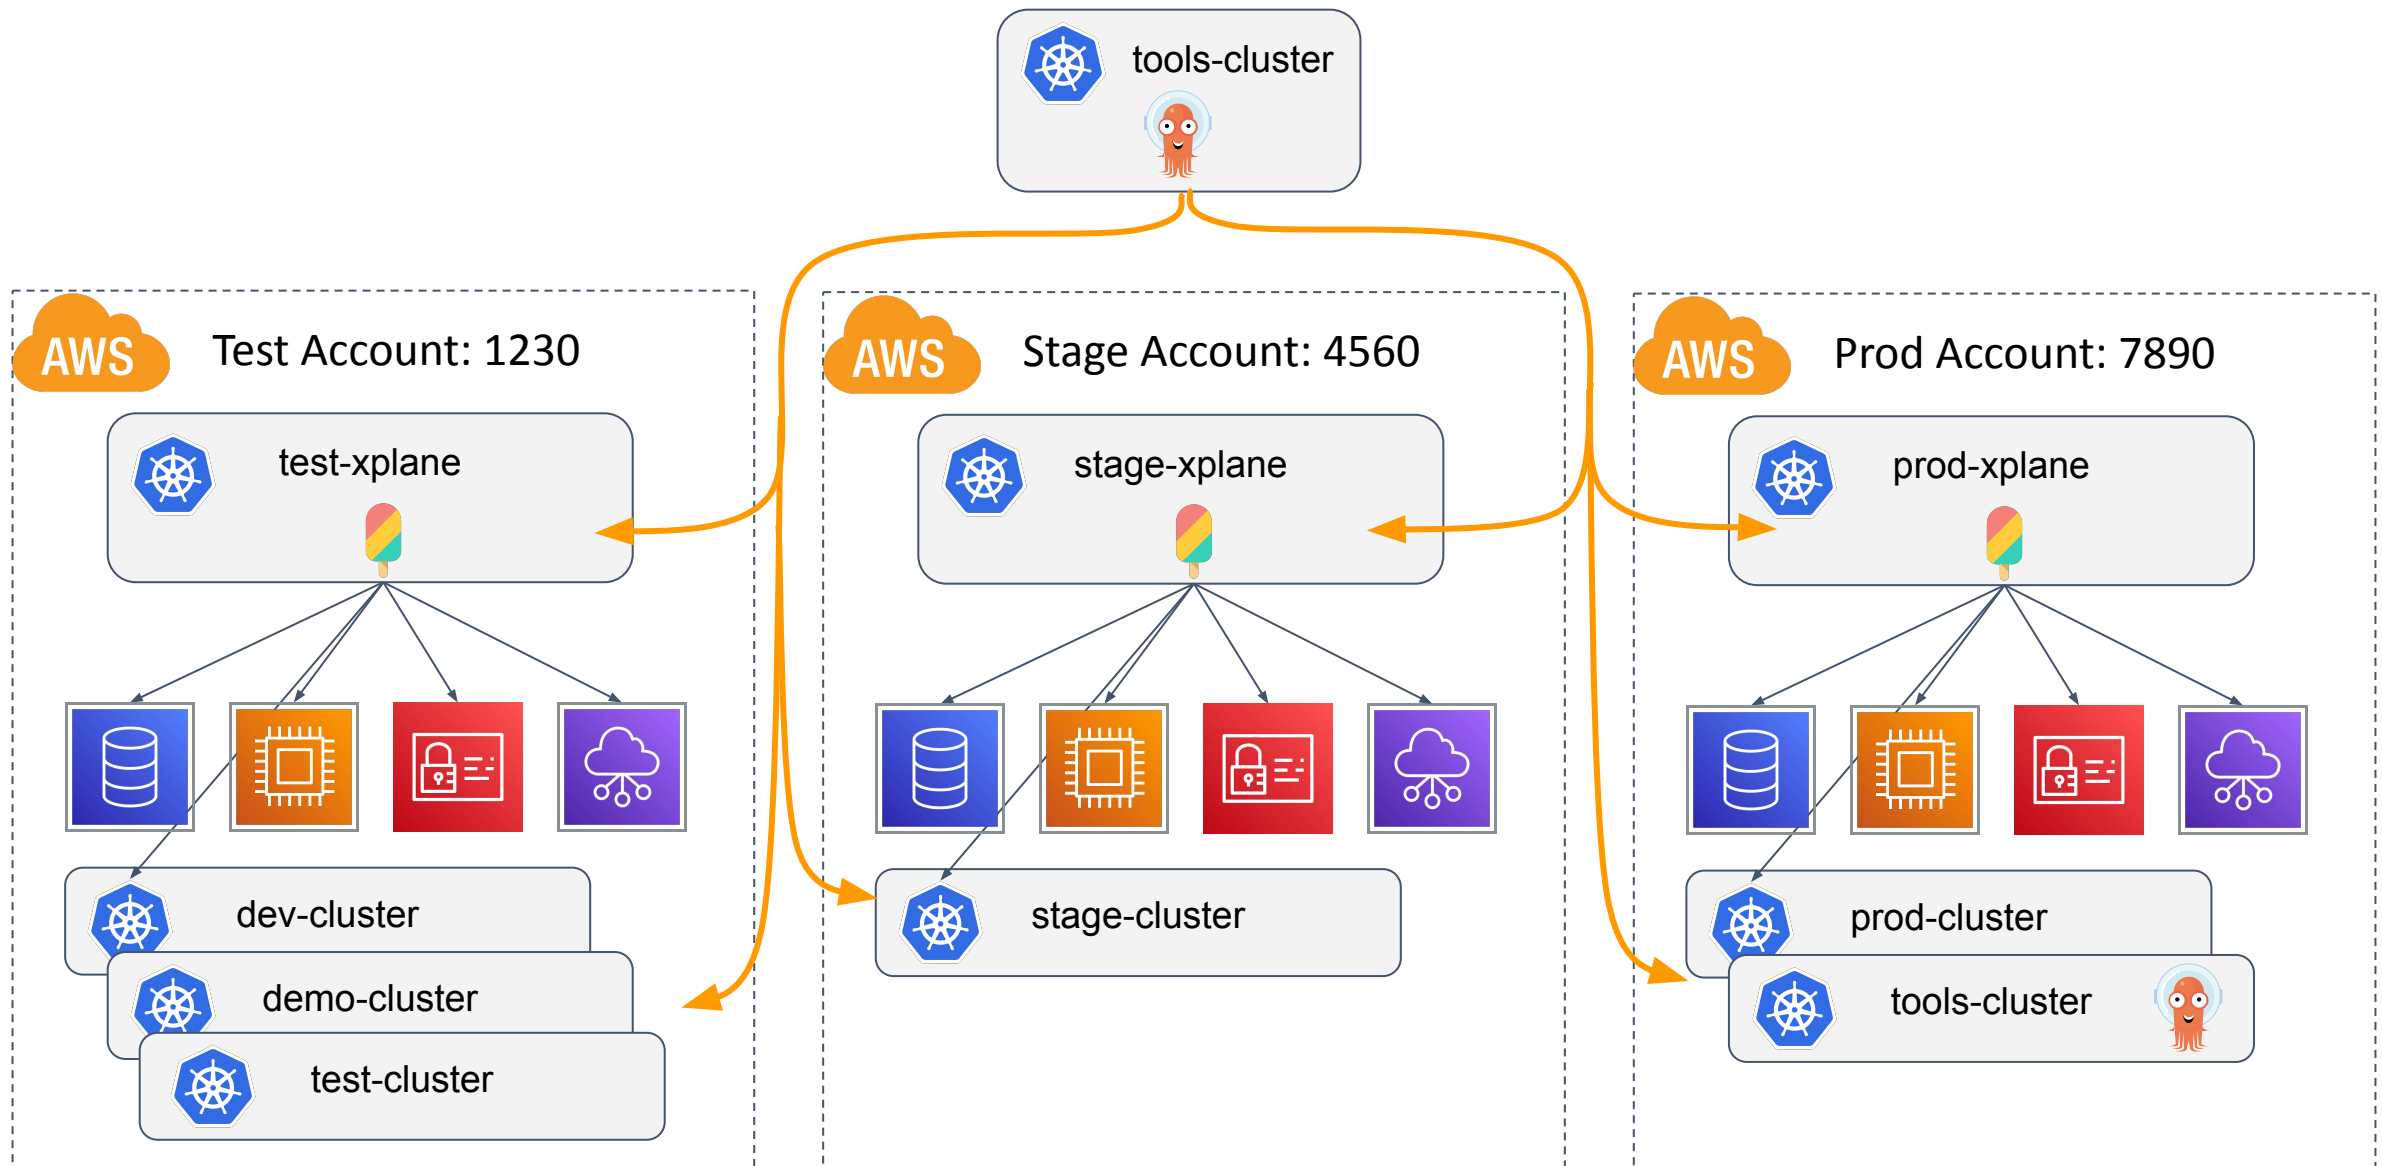
Problem:

- Rolling out new Compositions, Providers versions is risky
- Bugs in your composition affects all instances simultaneously

Solution:

- Run multiple Crossplane instances (per environment)
- Promote Compositions progressively to each environment

Best Practice: Crossplane per Env





Tips & Tricks

Argo CD: Pruning

Problem:

- Argo CD wants to prune the XRs produced by XRCs

Solution:

- Use “annotation” based resource tracking

```
kind: ConfigMap
metadata:
  name: argocd-cm
data:
  application.resourceTrackingMethod: annotation
```

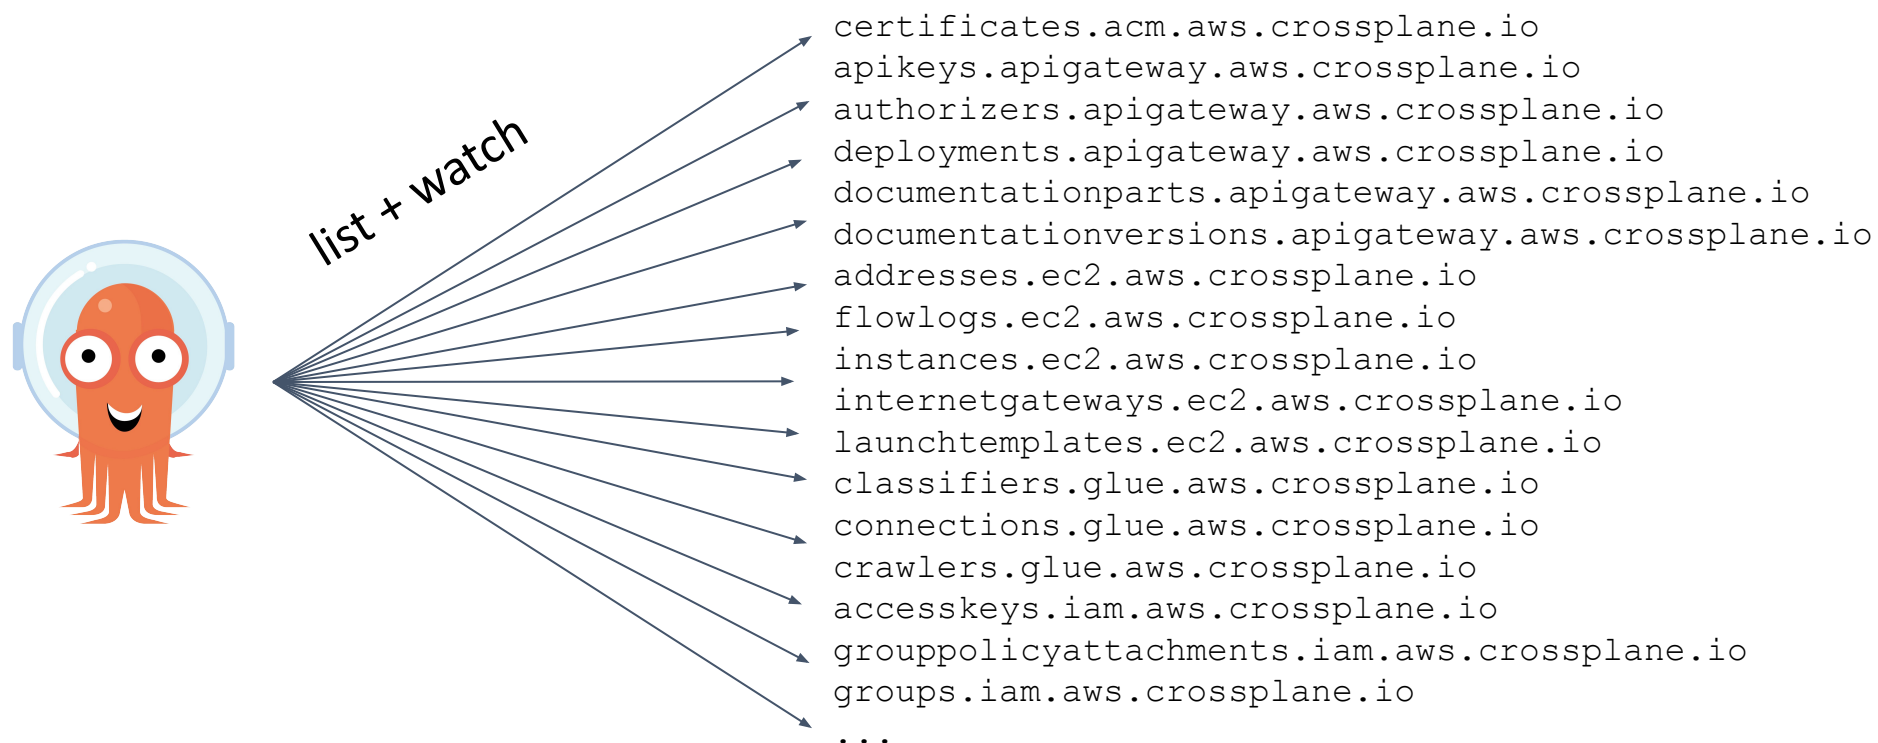
- Requires Argo CD v2.4.9 ([#8683](#))

Problem:

- Crossplane installs hundreds of CRDs
- You're probably not using most of them
- Argo CD list/watches each CRD, even if it's not used
- Increased K8s API and app-controller memory pressure

Argo CD: Resource Exclusions

Hundreds of open connections to K8s API



Argo CD: Resource Exclusions

Reduce connections by excluding unused resources

```
kind: ConfigMap
metadata:
  name: argocd-cm
data:
  resource.exclusions: |
    - apiGroups:
      - apigateway.aws.crossplane.io
      - apigatewayv2.aws.crossplane.io
      - glue.aws.crossplane.io
      - cache.aws.crossplane.io
      - kafka.aws.crossplane.io
      - dax.aws.crossplane.io
      - docdb.aws.crossplane.io
      - dynamodb.aws.crossplane.io
    ...
```



certificates.acm.aws.crossplane.io
~~apikeys.apigateway.aws.crossplane.io~~
~~authorizers.apigateway.aws.crossplane.io~~
~~deployments.apigateway.aws.crossplane.io~~
~~documentationparts.apigateway.aws.crossplane.io~~
~~documentationversions.apigateway.aws.crossplane.io~~
addresses.ec2.aws.crossplane.io
flowlogs.ec2.aws.crossplane.io
instances.ec2.aws.crossplane.io
internetgateways.ec2.aws.crossplane.io
launchtemplates.ec2.aws.crossplane.io
~~classifiers.glue.aws.crossplane.io~~
~~connections.glue.aws.crossplane.io~~
~~crawlers.glue.aws.crossplane.io~~
accesskeys.iam.aws.crossplane.io
grouppolicyattachments.iam.aws.crossplane.io
groups.iam.aws.crossplane.io
...

Argo CD: Resource Exclusions

Declutter Argo CD UI by excluding ProviderConfigUsages
(Crossplane implementation detail)

```
kind: ConfigMap
metadata:
  name: argocd-cm
data:
  resource.exclusions: |
    - kinds:
      - ProviderConfigUsage
    apiGroups:
      - "*"

```


Argo CD: Performance Tuning

Problem:

- Crossplane installs hundreds of CRDs
- API Discovery of those CRDs get throttled ([#3272](#))
- Currently, Argo CD Discovery is not efficient ([#448](#))

(Temporary) Workaround:

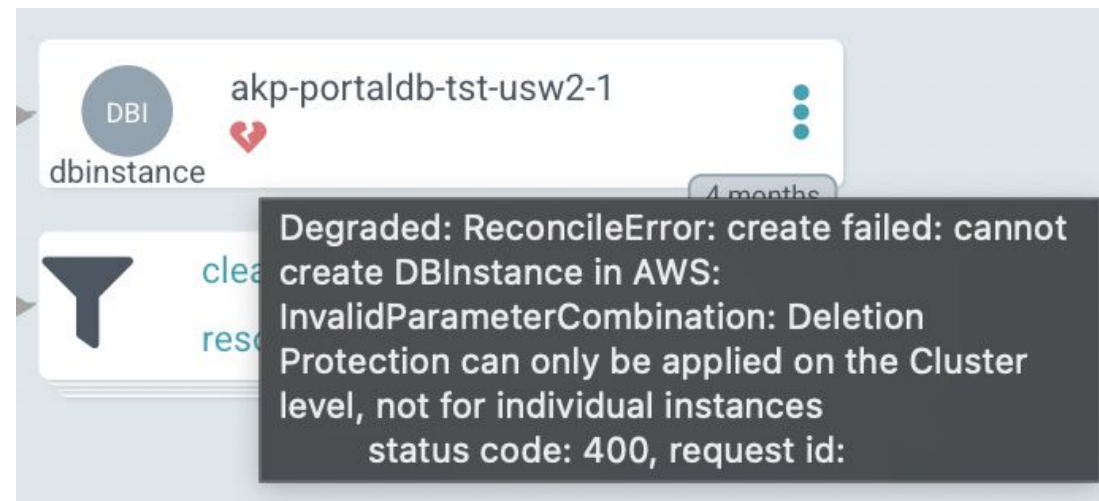
- Increase Kubernetes QPS limit to app-controller

```
env:  
- name: ARGOCD_K8S_CLIENT_QPS  
  value: "300"
```

Argo CD: Health Checks

Tip & Trick:

- Write Argo CD health check for Crossplane resources



```
status:
  conditions:
  - lastTransitionTime: "2022-10-14T23:53:45Z"
    message: "create failed: cannot create DBInstance in AWS: InvalidParameterCombination:
      Deletion Protection can only be applied on the Cluster level, not for individual
      instances\n\tstatus code: 400, request id: "
    reason: ReconcileError
    status: "False"
    type: Synced
```

Argo CD: Health Checks

```
kind: ConfigMap
metadata:
  name: argocd-cm
data:
  resource.customizations.health.rds.aws.crossplane.io_DBInstance: |
    hs = {}
    hs.status = "Healthy"
    synced = true
    for i, condition in ipairs(obj.status.conditions) do
      if condition.type == "Synced" and condition.status ~= "True"
        hs.status = "Degraded"
        synced_message = condition.reason
        if condition.message then
          synced_message = synced_message .. ": " .. condition.message
        end
      end
    end
  end
  return hs
```

Future Improvement: Shared Resource Health Checks ([#4212](#))



Challenges

Adopting existing cloud resources

- Many AWS resources receive random IDs
 - e.g. VPCs, Subnets, SecurityGroups, HostedZones
- Crossplane will recreate them if XR is recreated
- `external-name` annotation helps, but breaks GitOps UX

```
apiVersion: ec2.aws.crossplane.io/v1beta1
kind: VPC
metadata:
  annotations:
    crossplane.io/external-name: vpc-0a1b2c3d4e5f6789
  name: foo-vpc
```

Tip: Use formulated XR names (avoid `generateName`)!

Conditional Resources ([#2712](#))

- In a Composition, I want to create resources depending on the values of input parameters.

Use Cases:

- Enable/Disable Cluster Addons
- Prod vs. Dev Databases

Current Workaround:

- Violate DRY. We resort to code generation to create variations of compositions (RDS, RDSDev)

Cross-Resource Referencing ([#1770](#))

- In a Composition, I would like to reference fields of *another* resource as input to the current resource

Use Cases:

- Enables loosely coupled compositions
- Simplify deploy process

Current Workaround: Staged deploys

1. Deploy two VPCs
2. Extract VPC ID of each
3. Copy & paste VPC ID into TransitGateway Attachments

Cheaper Control Planes

- Currently running dedicated EKS clusters just for Crossplane
- Multiple accounts/environments == expensive!
- Instead of EKS as our crossplane cluster, K3s in EKS



Thank You!



Please scan the QR Code above to
leave feedback on this session