



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

NRI: Extending Containerd and CRI-O With Common Plugins



BUILDING FOR THE ROAD AHEAD

DETROIT 2022



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

October 24-28, 2021



Mike Brown
OSS Software
Engineer
IBM



Krisztian Litkey
Software Engineer
Intel Finland

We'll Cover...

- Some History And Background
- What Is NRI ?
- How Does It Work ?
- What Can It Do ?
- Writing An NRI Plugin ?
- Examples, Use Cases
- Q & A

History And Background

- various extension mechanisms over the years
- OCI hooks
 - custom actions executed at various container lifecycle events
 - runc, cri-o
- runtime wrapper using runtime classes
 - custom actions implemented by (custom) runtime wrapper
 - a bit of a kludge
 - runc/crun/etc..., containerd/cri-o
- NRI
 - custom actions at pod or container creation/stop
 - containerd
- NRI evolution
 - subject of this presentation...

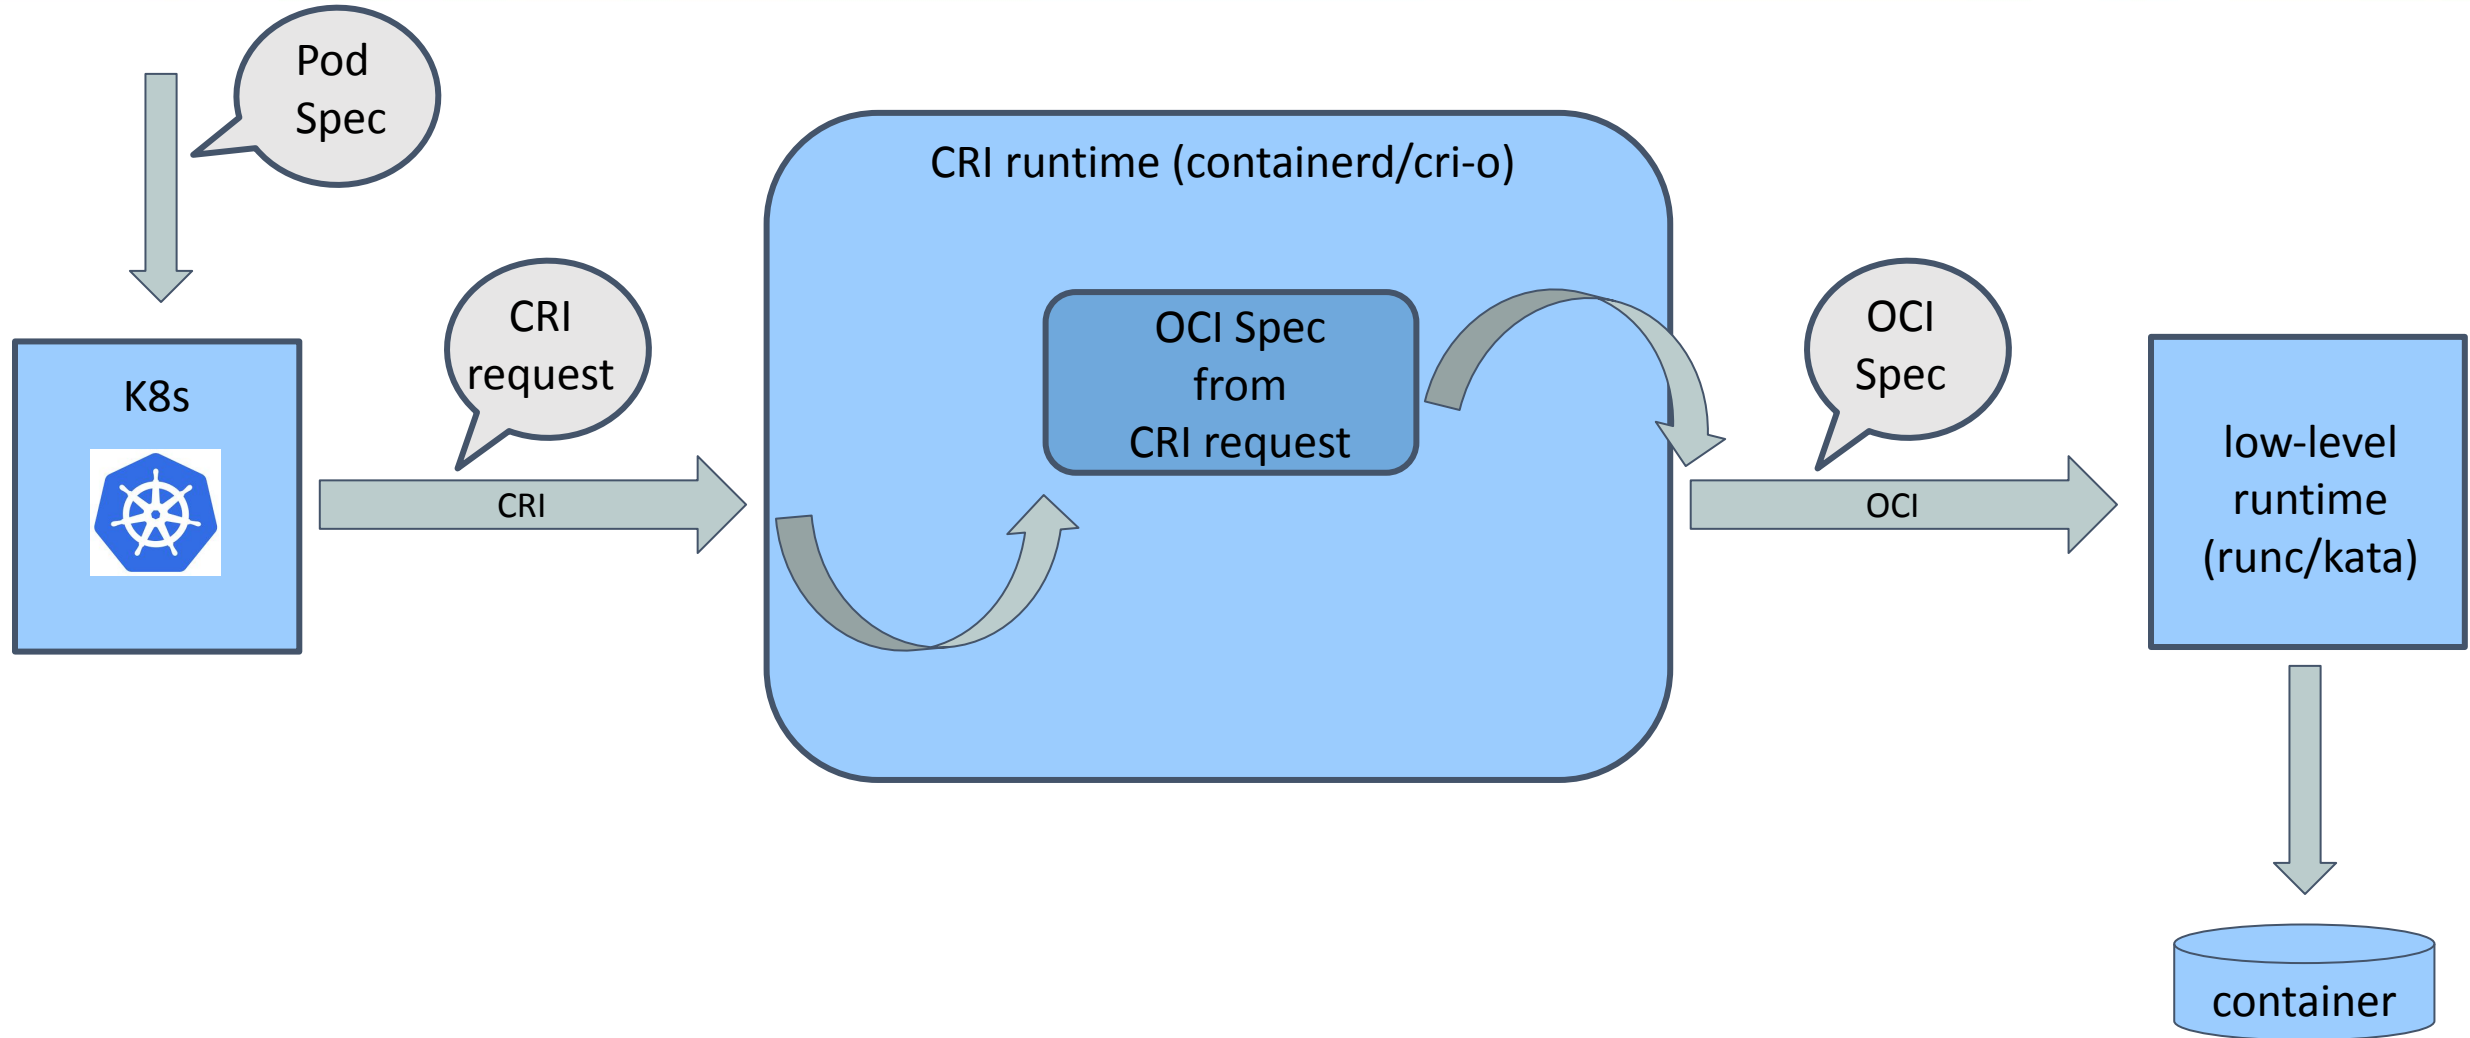
What Is NRI ?

- NRI: Node Resource Interface, a common framework for
 - plugging extensions into OCI-compatible runtimes
 - implementing custom container configuration logic
- ‘common’
 - plugins work identically in supported runtimes
 - support present in commonly used runtimes
- ‘framework’
 - multiple NRI pieces
 - collectively achieve NRI’s goal
- ‘pluggable extensions’
 - alter behavior without runtime modifications
- ‘custom logic’
 - your cluster, your plugin, your rules
 - within the allowed boundaries imposed by NRI

What Is NRI...

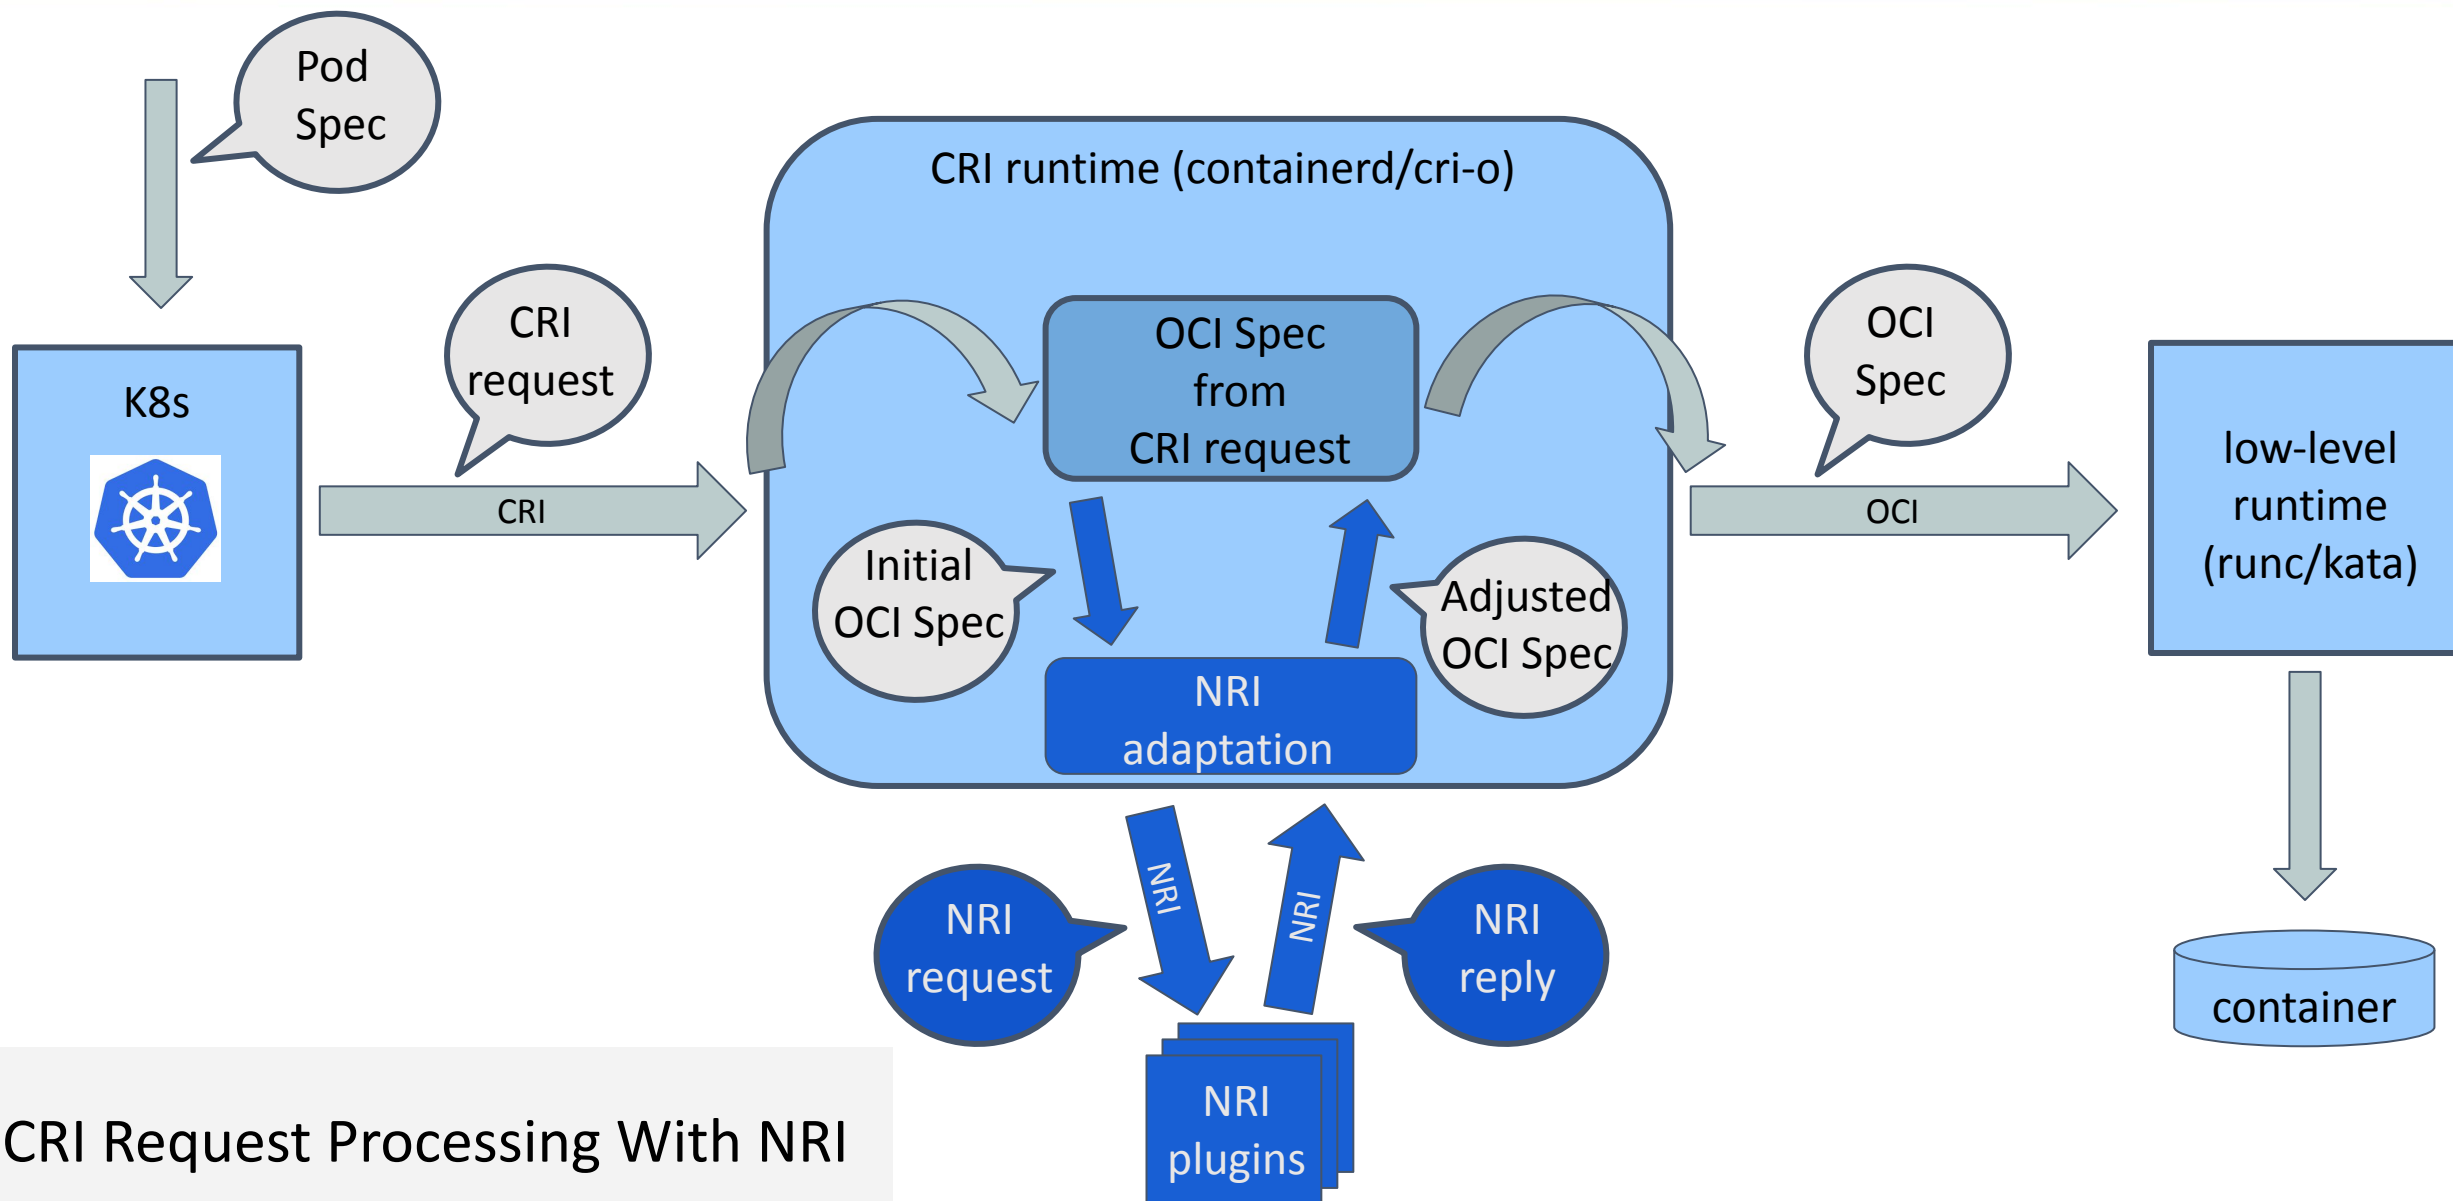
- plumbing to enforce how containers are configured
 - initially, during container creation
 - subsequently, by container updates
 - potentially, in response to some external events
- => your cluster, your plugin, your rules
 - within the allowed boundaries imposed by NRI

How Does NRI Work ?



CRI Request Processing

How Does NRI Work...

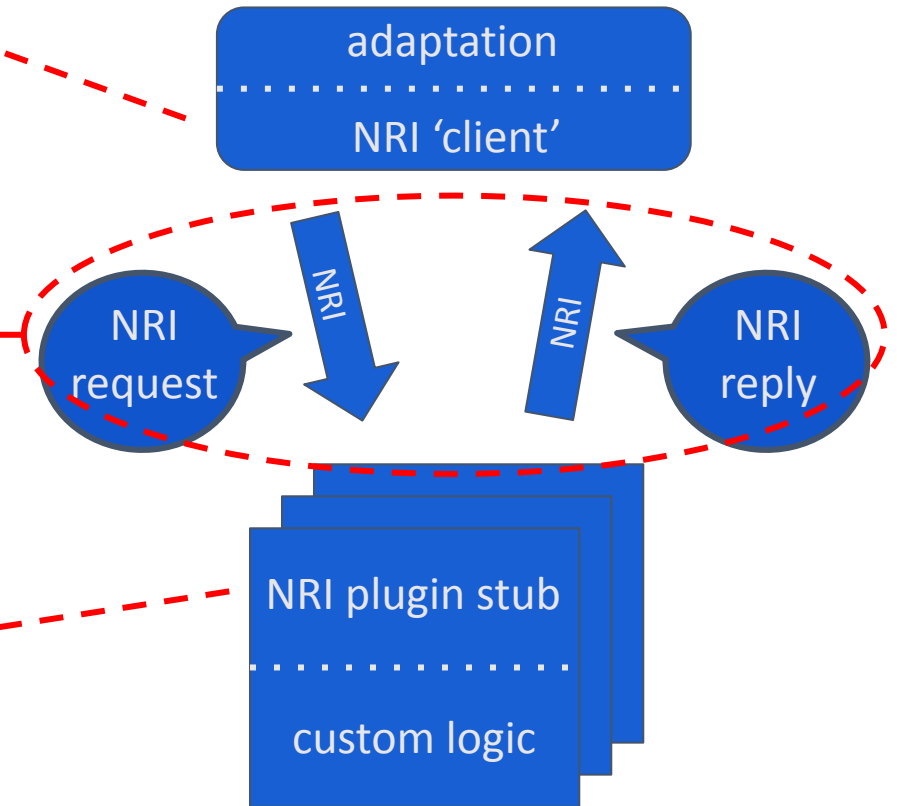


How Does NRI Work...

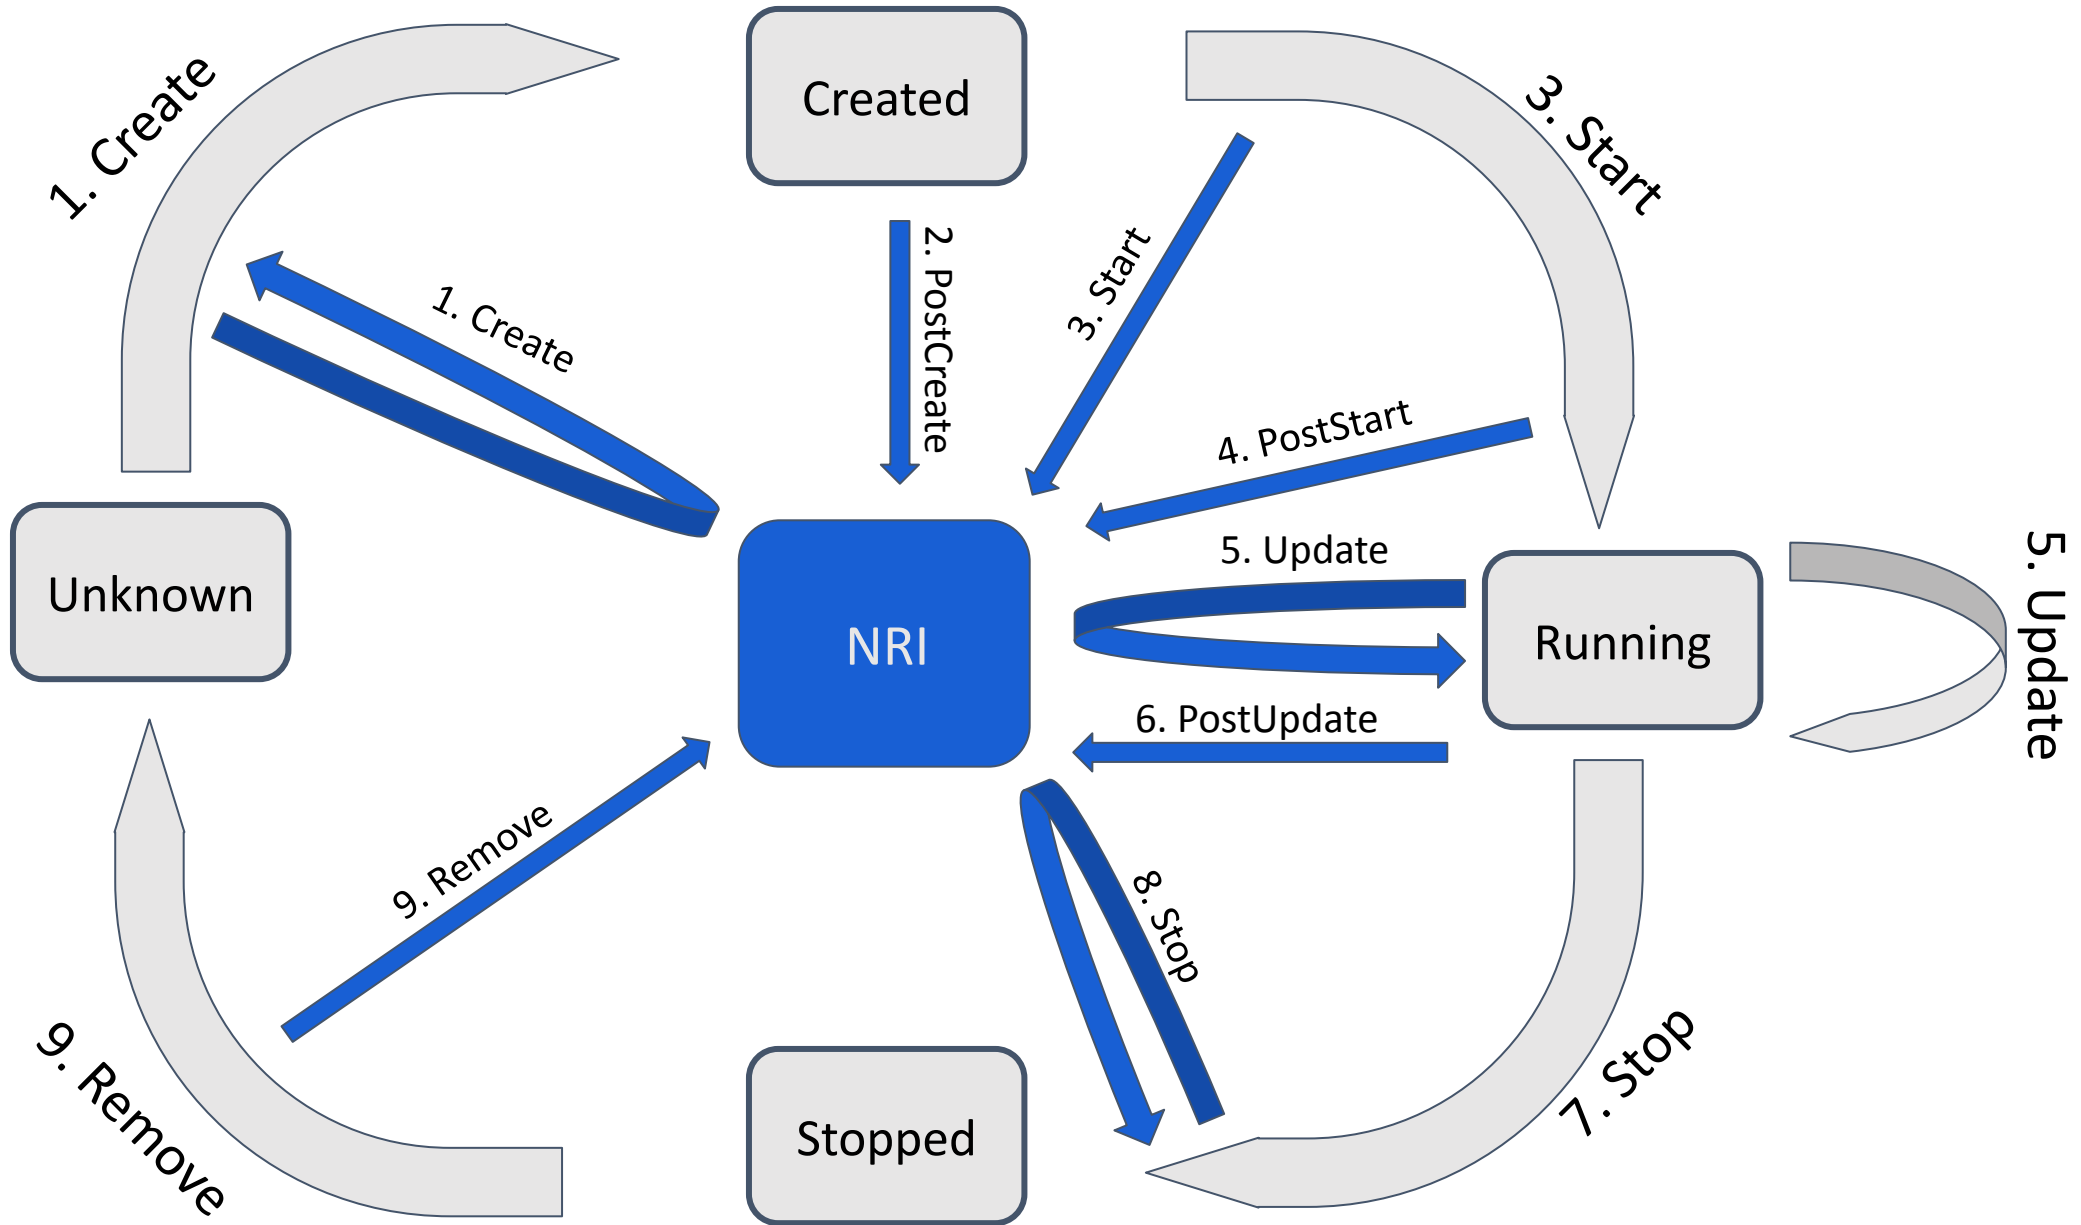
- runtime adaptation client
 - runtime agnostic
 - integrate to NRI / interact with plugins

- NRI protocol
 - protobuf-based 'API'
 - ttRPC bindings
 - execution model
 - full set of events/request
 - data model: event/request details
 - Pods, containers (subset of OCI Spec)
 - NRI container adjustments, updates

- plugin stub
 - low-level details of plugin logic
 - IPC, connection, registration, etc.




NRI Plugin Events



NRI Plugin Input And Output

container creation response:

- container being created
 - annotations, env. vars.
 - mounts, devices
 - OCI hooks
 - 'resources'
 - scheduling params
 - CPU/memory pinning
 - memory, huge page limits
 - LLC cache allocation
 - blkio controller parameters
- existing containers
 - 'resources'



container update response:

- container being updated
 - 'resources'
- other containers
 - 'resources'

container stop response:

- stopped container
 - nothing
- remaining containers
 - 'resources'

container request includes:

- labels, annotations, command, env. vars.
- mounts, devices, OCI hooks
- 'resources'
- runtime handler, cgroup path

What Else NRI Plugins Can Do...

- plugins can also
 - synchronize state on connection establishment
 - make unsolicited changes to containers
 - subscribe and react to Pod lifecycle events
 - control other things, outside the scope of the OCI Spec
- plugins can't
 - make conflicting changes to the same container
 - NRI catches this and treats it as an error
- plugins shouldn't
 - try to manipulate directly anything covered by the OCI Spec
 - even if those things are not covered/controllable via NRI

How To Write An NRI Plugin

- clone the NRI plugin template, fill in the missing details
 - implement data collection
 - if you need to collect pod/container data subscribe for the corresponding events
 - if you need data external to NRI/OCI, implement collection (eg. metrics)
 - implement custom container setup
 - subscribe for container creation requests
 - provide container customization response
 - implement custom container update, if needed
 - subscribe for container update, stop requests
 - provide update for other containers in create response
 - provide update for containers in update, stop responses
- proactive custom container configuration logic
 - collect any data/metrics you need for updating container configuration
 - trigger 'calculation' of new configurations
 - set new configurations using an unsolicited update request
- and... look at our sample plugins for simple examples

Examples And Use Cases

- ‘real world’ plugins
 - annotation-based device injection
 - CDI device injection (obsolete)
 - OCI hook injection
- debugging and development plugins
 - event logger
 - container differ
 - plugin template
- experimental resource assignment plugin
 - CRI Resource Manager with ‘--use-nri-plugin’

Thank You ! / Q & A

Contact info:

- Krisztian Litkey, krisztian.litkey@intel.com, klihub on github
- Mike Brown, brownwm@us.ibm.com, mikebrow on github

NRI Artifacts:

- <https://github.com/containerd/nri/pull/16>
- <https://github.com/containerd/containerd/pull/6019>
- <https://github.com/cri-o/cri-o/pull/5318>
- any questions or comments ?



Please scan the QR Code above to
leave feedback on this session