

One VTOrc to Rule Them All

High Availability In a Distributed Database System

Deepthi Sigireddi, Manan Gupta



Deepthi Sigireddi



Technical Lead

@ATechGirl

Manan Gupta



Maintainer

@GuptaManan100

Vitess Overview

What is Vitess?

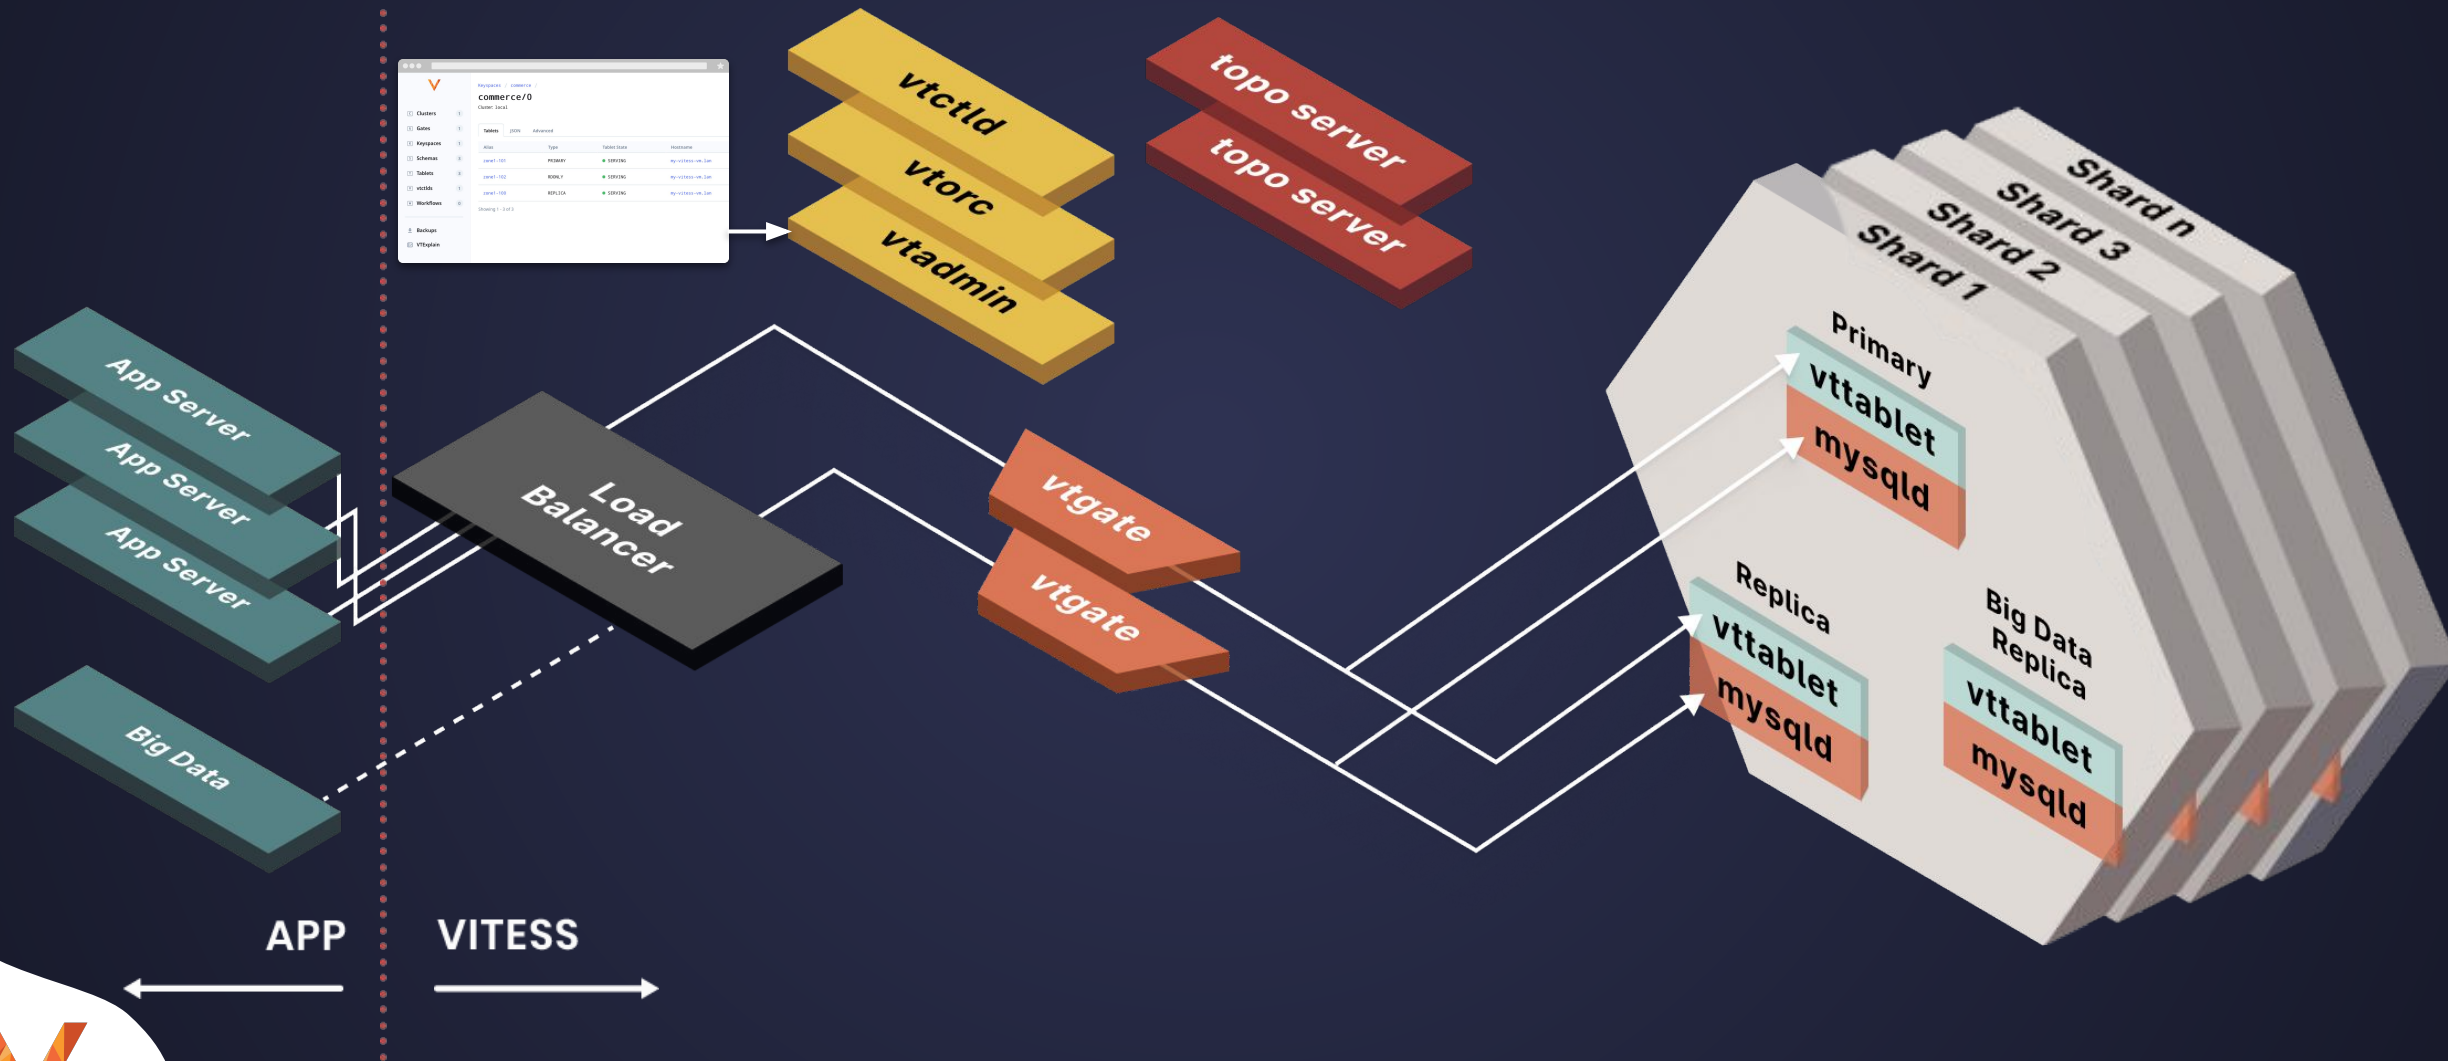
Cloud
Native
Database

Massively
Scalable

Highly
Available

MySQL
5.7/8.0
Compatible

Architecture



VTOrc Overview

Problem Statement

- Resiliency to MySQL failures
- High availability
- Data Durability
- Minimize downtime / recovery time

Before there was VTOrc ...

- openark/orchestrator
- Integrated with Vitess
- enable_semi_sync flag on VTTablet
- Works well enough most of the time
- But not all the time

VTOrc

- VTOrc is now GA (v15)
- In production

Using the Vitess Kubernetes operator and [...] orchestrator (VTOrc) has been very pleasant, as it removed a lot of operational overhead. Vitess seems rock solid so far, and we look forward to seeing what the future will bring for this awesome project.

- Principal Ops Architect, AAA gaming studio

VTOrc

- VTOrc is the agent that detects failures
- Durability through Replication
 - Policies allow trade-offs
- High availability through failover
 - Planned / unplanned leader election



Design Principles


- Engineering approach
- Single leader system
- Fulfill requests while respecting durability policy
- Leader election process
 - Planned versus unplanned
- Forward Progress
- Race conditions

Leader Election

- Revocation
- Election
- Propagation

Planned Leader Election

- Revocation
 - Current leader is asked to step down
- Leader selection
 - A new leader is chosen
- Propagation
 - Completed requests



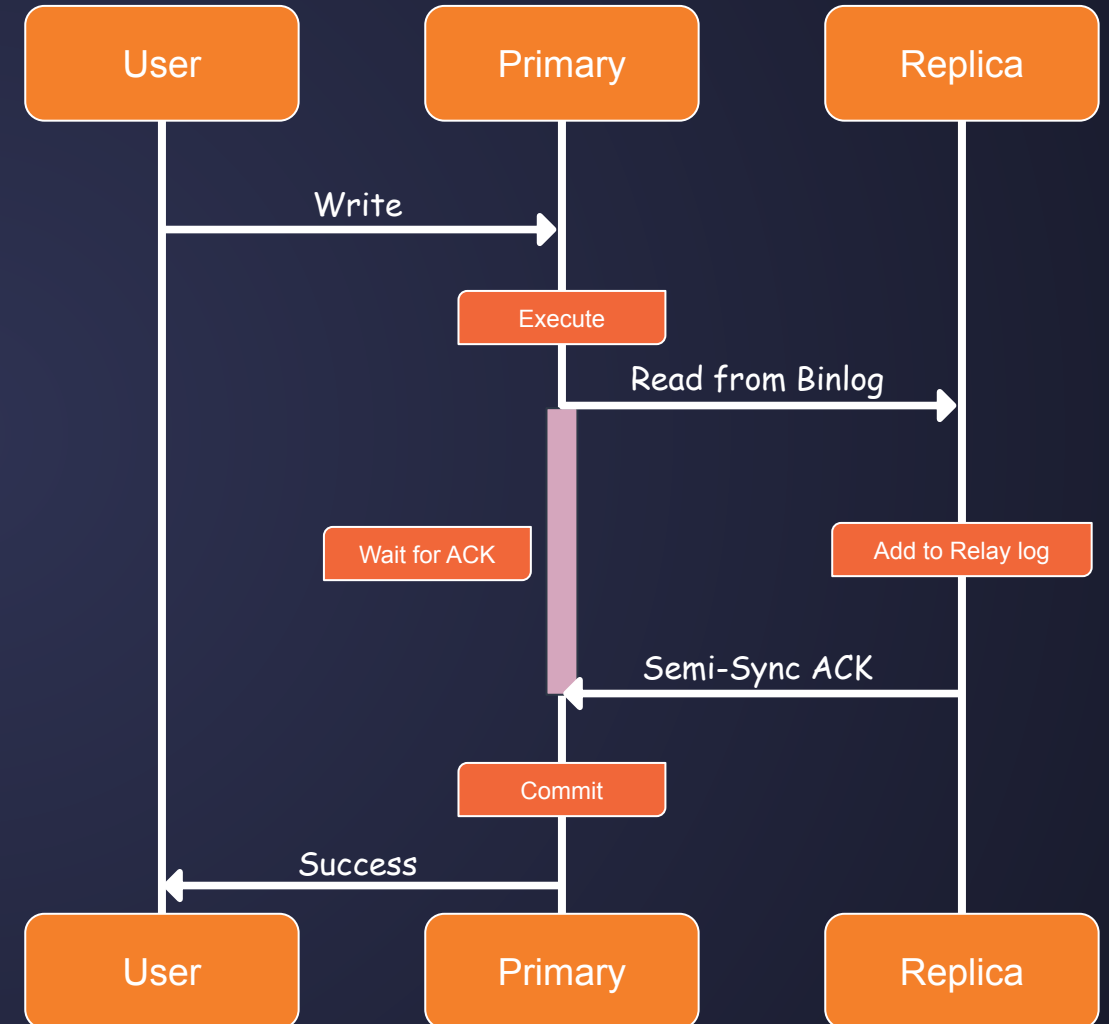
Unplanned Leader Election

- Revocation
 - Reach “m” followers
- Leader selection
 - A new leader is chosen
 - Based on durability policy
- Propagation
 - Completed requests



Durability Policies & Semi-Sync

- Semi-Sync in MySQL
- Durability Policy
 - Who can be the primary?
 - How many semi-sync ACKs required for each primary?
 - Who can send these ACKs?
- Increased Flexibility



Revocation and Quorum

- What is “m”?
- How do we know we have reached sufficient tablets to guarantee safety?
- Intersecting Quorum
- Quorum for accepting transactions
- Quorum for revocation

Revocation and Quorum




Revocation and Quorum



Demo

Semi-Sync Durability

- `semi_sync_ks`
- Durability Policy - semi-sync
 - Any replica can be the primary
 - 1 semi-sync ACK required
 - Any replica can send the ACK




Revocation

- **Quorums for Accepting Transactions**

- [(100, 101), (100, 102), (100, 103)]
- [(101, 100), (101, 102), (101, 103)]
- [(102, 100), (102, 101), (102, 103)]
- [(103, 100), (103, 101), (103, 102)]


- **Quorums for Revocations -**

- [100, 103] ❌
- [100, 102, 103] ✅
- [100, 101, 102, 103] ✅
- [101] ❌




Cross-Cell Durability

- `cross_cell_ks`
- Cell = Failure Domain
- Durability Policy - cross-cell
 - Any replica can be the primary
 - 1 semi-sync ACK required
 - A replica from a different cell can send an ACK



Revocation

- **Quorums for Accepting Transactions**
 - $[(300, 200), (300, 201), (300, 202)]$
 - $[(200, 300)]$
 - $[(201, 300)]$
 - $[(202, 300)]$
- **Quorums for Revocations -**
 - $[200, 201]$ ❌
 - $[300, 202, 201]$ ✅
 - $[300, 201]$ ✅
 - $[300, 200, 201, 202]$ ✅



Custom Durability Policies

- Durability Policy
 - Who can be the primary?
 - How many semi-sync ACKs required for each primary?
 - Who can send these ACKs?
- Increased Flexibility

Custom Durability Policies

- **Quorums for Accepting Transactions**

- [(A, B, C), (A, B, D), (A, C, D)]
- [(B, D), (B, A)]
- [(C, B)]

- **Quorums for Revocations -**

- [A, D] ❌
- [A, B, D] ✅
- [B, C] ✅
- [C, D] ❌

More Failure Scenarios

- Maintains the cluster's desired state
- Primary is Read-Only
- Replica's replication is stopped
- Replica is writable
- Semi-sync settings are incorrect
- Shard has no primary
- Primary is replicating from a different tablet

Q & A

Resources

Vitess: Introduction and New Features
Friday, Oct 28 11:55 am EDT

Blog Post Series

- <https://planetscale.com/blog/blog-series-consensus-algorithms-at-scale-part-1>

VTORC Documentation

- <https://vitess.io/docs/15.0/reference/vtorc/>
- <https://vitess.io/docs/15.0/user-guides/configuration-basic/vtorc/>



<https://sched.co/182J8>

Thank you!