Prometheus Intro, Deep Dive

Goutham Veeramachaneni
Grafana Labs, @putadent

Ganesh Vernekar
Grafana Labs, @_codesome

Prometheus team members

What is Prometheus?

Metrics-based monitoring & alerting stack.

- Instrumentation for applications and systems
- Metrics collection and storage
- Querying, alerting, dashboarding
- For all levels of the stack!

Made for dynamic cloud environments.

History

- Started 2012 at SoundCloud
- Fully publicised in 2015
- Joined CNCF, Prometheus v1.0.0 released in 2016
- Prometheus v2.0.0 released in 2017

Architecture

Architecture

Targets

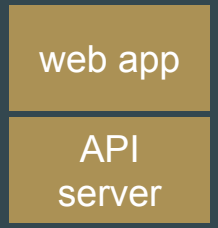

Architecture

Targets

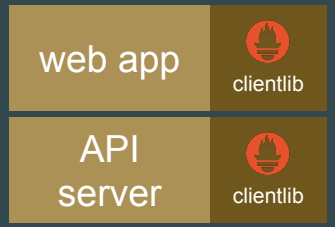

Architecture

Targets

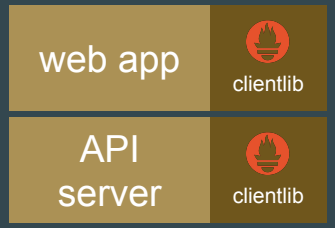

Instrumentation & Exposition

Architecture

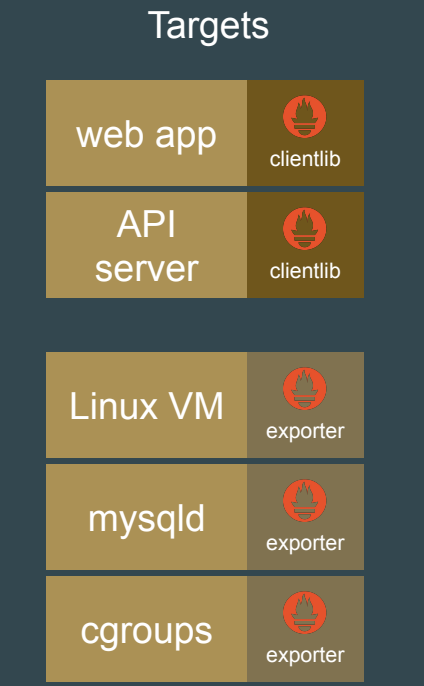

Architecture

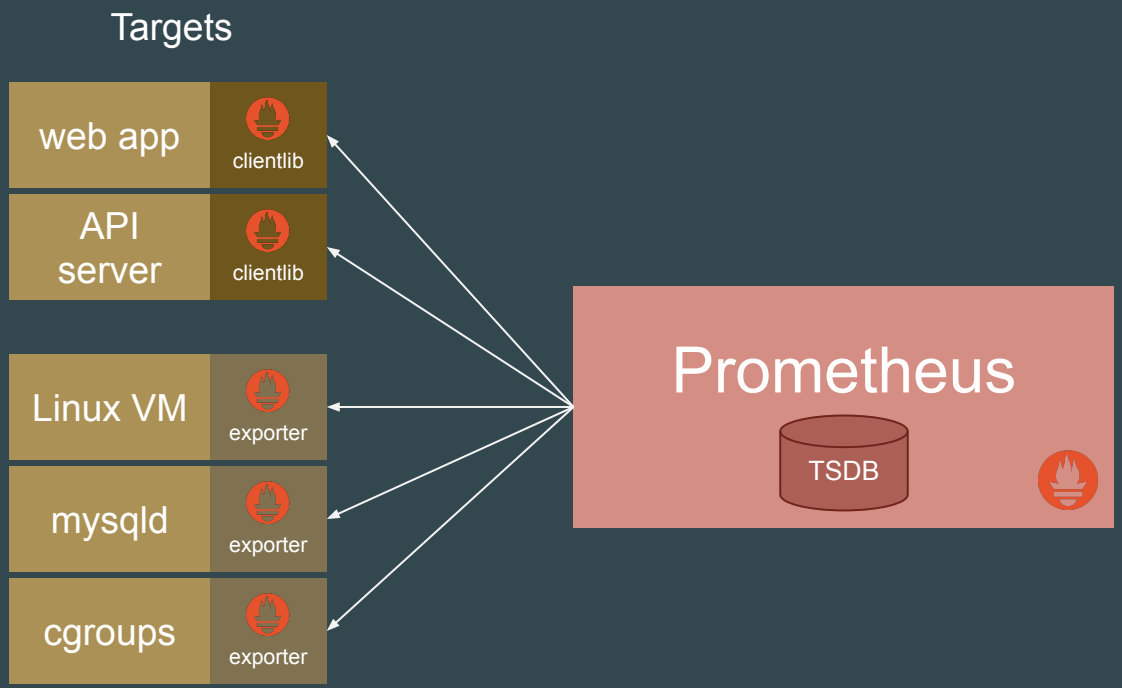

Architecture

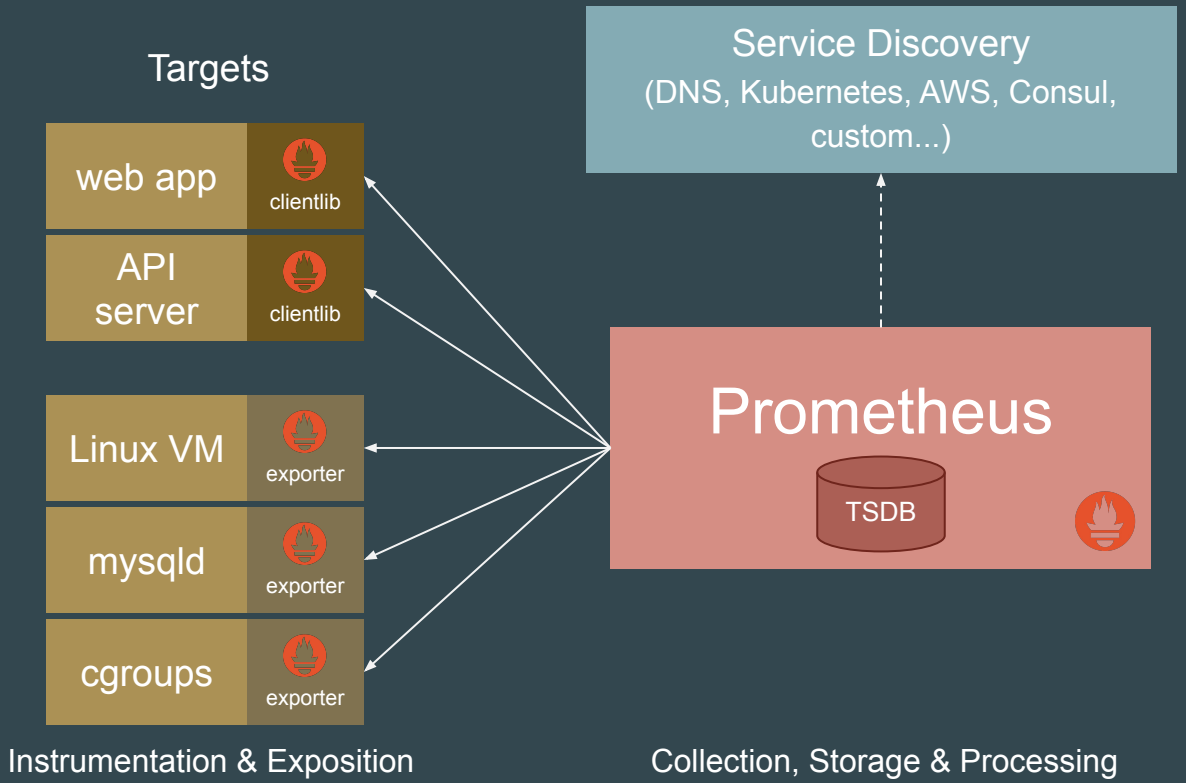

Architecture

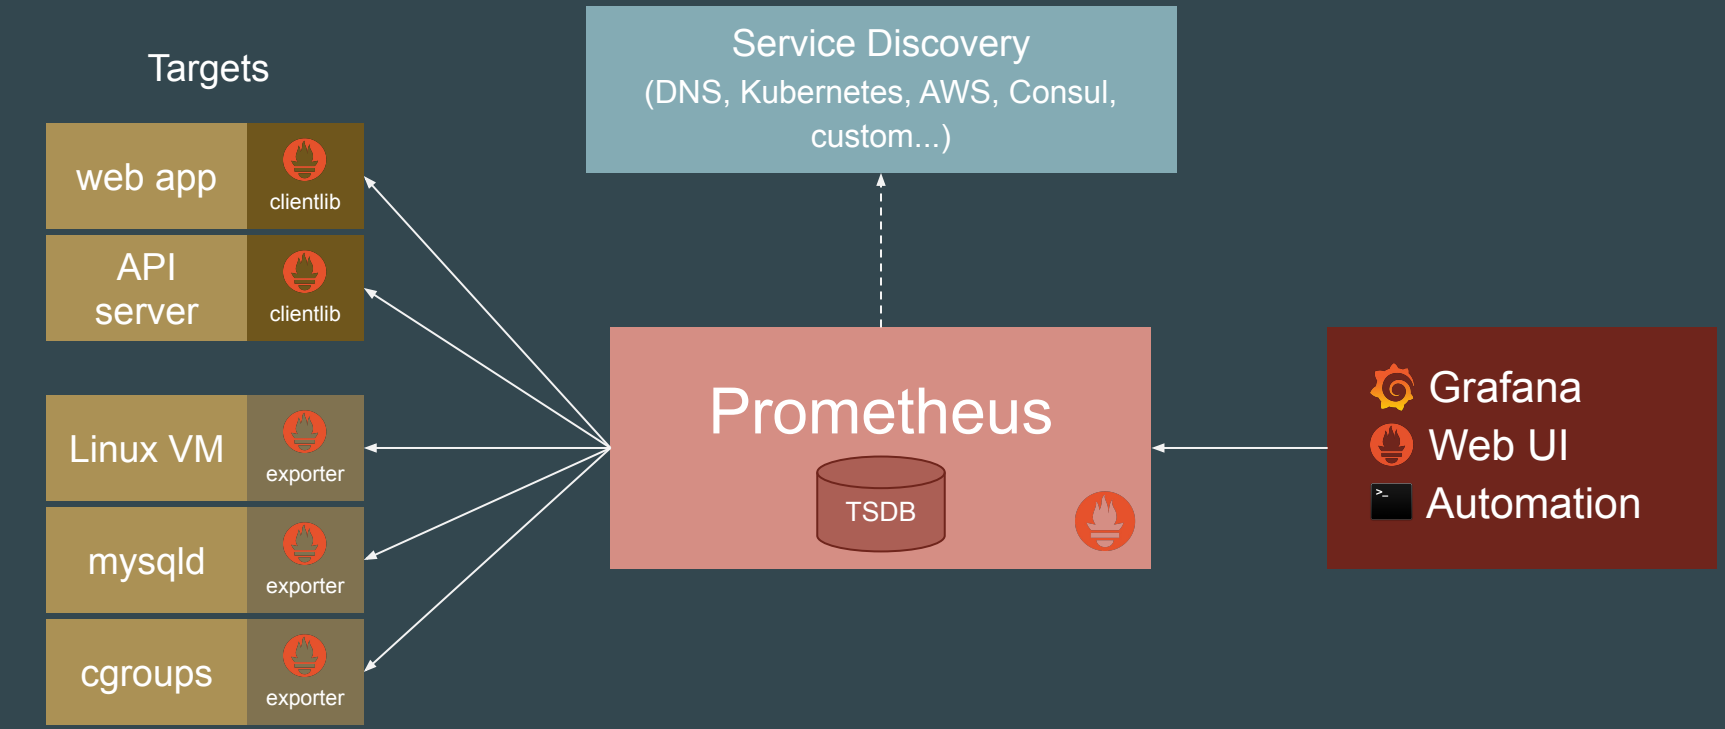

Architecture

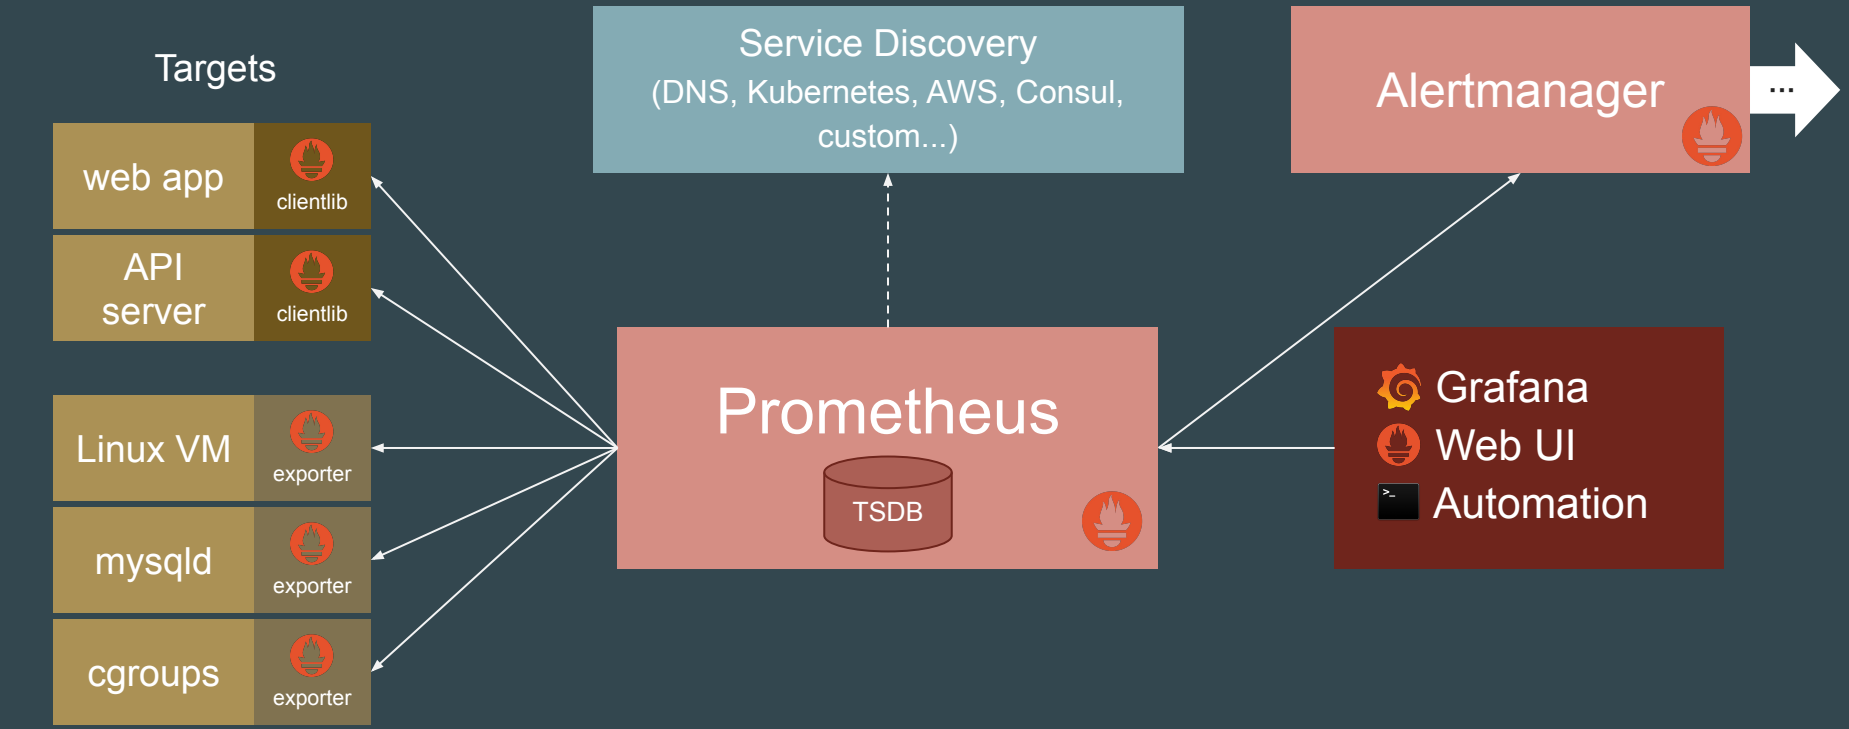

Selling Points

- Dimensional data model
- Powerful query language (PromQL)
- Simple & efficient server
- Service discovery integration

Data Model

What is a time series?

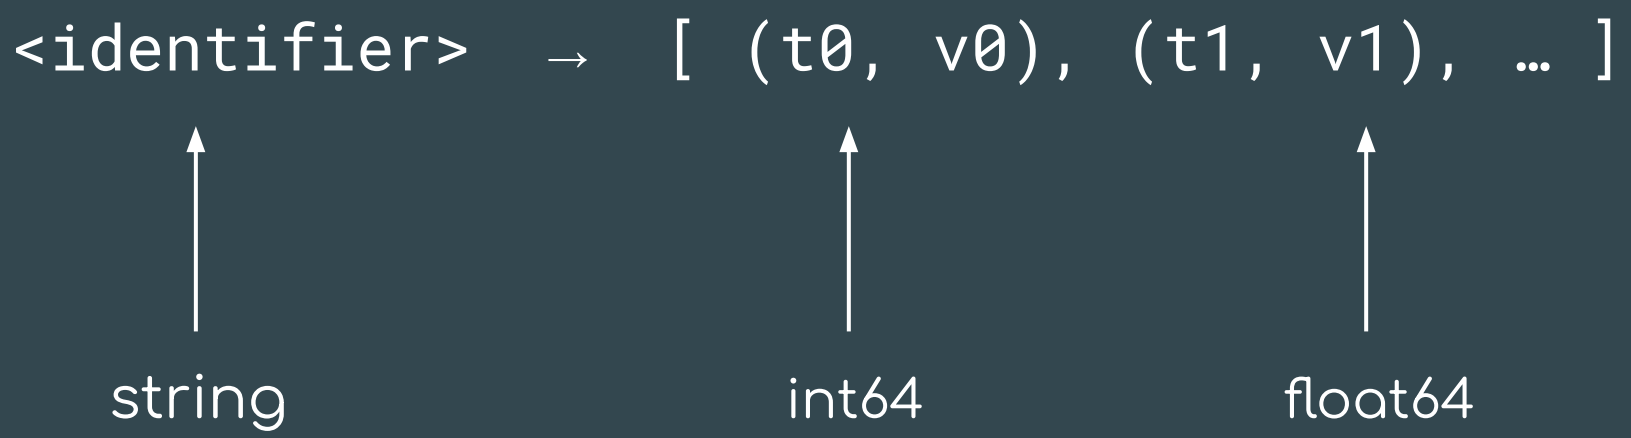

Data Model

What identifies a time series?

```
http_requests_total{job="nginx",instance="1.2.3.4:80",status="200"} 28
```

metric name

labels

- Flexible
- No hierarchy
- Explicit dimensions

Querying

Querying

PromQL

- Functional query language
- Great for time series computations
- Not SQL-style

Querying

All partitions in my entire infrastructure with more than 100GB capacity that are not mounted on root?

```
node_filesystem_bytes_total{mountpoint!="/"} / 1e9 > 100
```

```
{device="sda1", mountpoint="/home", instance="10.0.0.1"}        118.8
{device="sda1", mountpoint="/home", instance="10.0.0.2"}        118.8
{device="sdb1", mountpoint="/data", instance="10.0.0.2"}        451.2
{device="xdvc", mountpoint="/mnt", instance="10.0.0.3"}         320.0
```

Querying

What's the ratio of request errors across all service instances?

```
  sum(rate(http_requests_total{status="500"}[5m]))
/ sum(rate(http_requests_total[5m]))
```

```
{}                                  0.029
```

Querying

What's the ratio of request errors across all service instances?

```
    sum by(path) (rate(http_requests_total{status="500"}[5m]))
  / sum by(path) (rate(http_requests_total[5m]))
```

{path="/status"}	0.0039
{path="/"}	0.0011
{path="/api/v1/topics/:topic"}	0.087
{path="/api/v1/topics}	0.0342

Alerting

generate an alert for each
path with an error rate of >5%

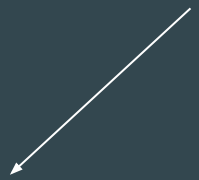

```
alert: Many500Errors
expr: |
  (
      sum by(path) (rate(http_requests_total{status="500"}[5m]))
    /
      sum by(path) (rate(http_requests_total[5m]))
  ) * 100 > 5
for: 5m
labels:
  severity: "critical"
annotations:
  summary: "Many 500 errors for path {{$labels.path}} ({{$value}}%)"
```

Efficiency

Local storage is scalable enough for many orgs:

- 1 million+ samples/s
- Millions of series
- 1-2 bytes per sample

Good for keeping a few weeks or months of data. Some people keep years, with careful backups.

Bridging the gap

Not everything speaks Prometheus – exporters help

- Translate from other metric systems (statsd, CloudWatch, ...)
- Transform system-specific metrics (Linux, MySQL, HAProxy, ...)
- Do it yourself (JSON exporter, Python, Go, ...)

Conclusion

Prometheus helps you make sense of complex dynamic environments via its:

- Dimensional data model
- Powerful query language
- Simplicity + efficiency
- Service discovery integration

What's new

Some of the new features released in the last ~12 months

...only looking at prometheus/prometheus

PromQL

```
temperature_kelvin - avg_over_time(temperature_kelvin[1d] offset -12h)
```

```
  rate(http_requests_total[1m])
and
  topk(5, rate(http_requests_total[1h] @ end()))
```

v2.25 (feature flag) v2.33 (stable)

```
sin(rad(wind_turbine_angle_degrees))
x_value atan2 y_value
```

v2.31

Remote Write Receiver

```
$ prometheus --enable-feature=remote-writer-receiver
```

v2.25 (feature flag) v2.33 (stable)

Agent mode

v2.32

```
$ prometheus --enable-feature=agent
```

https://prometheus.io/blog/2021/11/16/agent/

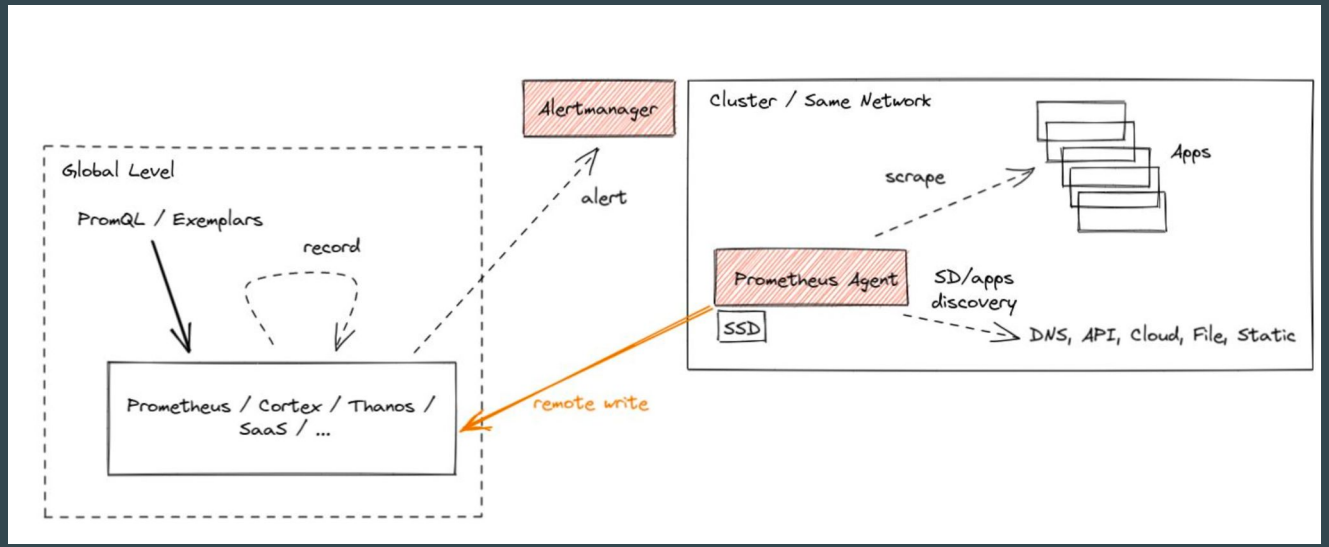

LTS releases

LTS for Long Term Support.

A release of Prometheus supported for 6+ months.

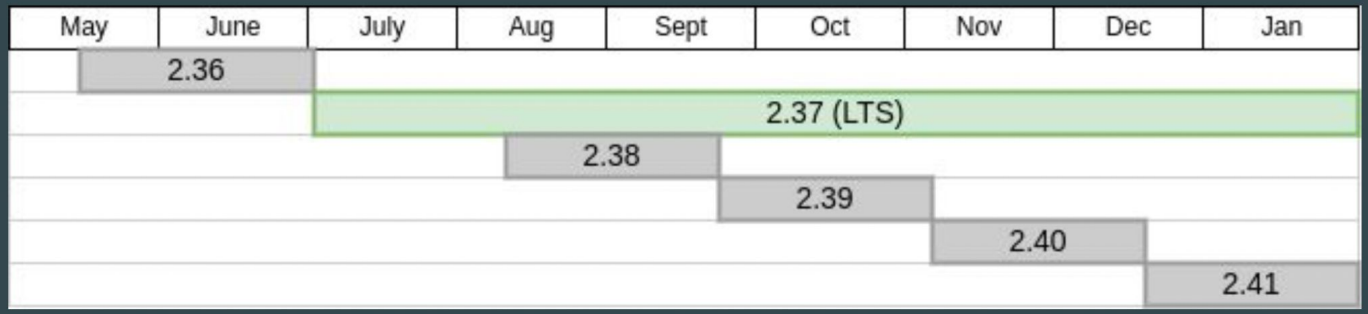

https://prometheus.io/docs/introduction/release-cycle/

Out-of-order ingestion v2.39

```
$ cat config.yml
...
storage:
  tsdb:
    out_of_order_time_window: 2h
...
```

What's Coming

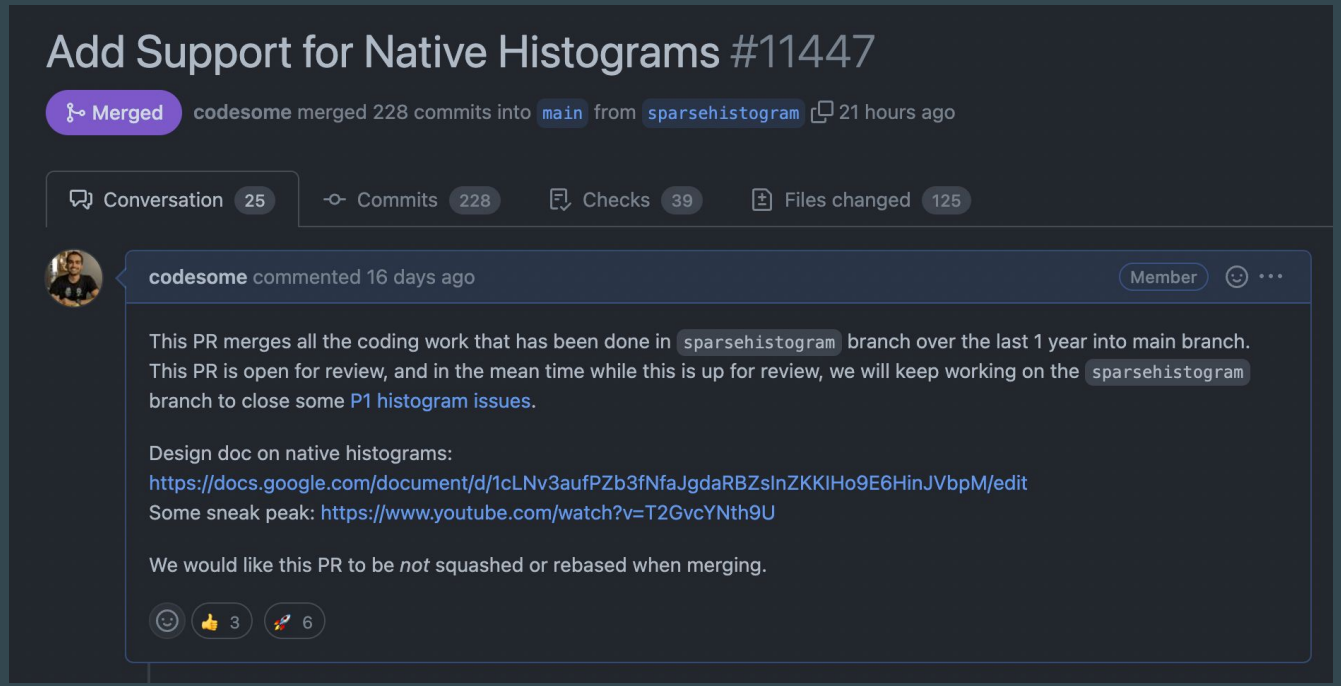

Thanks!