
KubeCon
CloudNativeCon
North America 2022
BUILDING FOR THE ROAD AHEAD
DETROIT 2022
TiKV
TiKV
HELM

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Remote Control Planes With Konnectivity; What, Why And How?

Remote Control Planes With Konnectivity

KubeCon | CloudNativeCon
North America 2022
BUILDING FOR THE ROAD AHEAD
DETROIT 2022

KubeCon | CloudNativeCon
North America 2022
BUILDING FOR THE ROAD AHEAD
DETROIT 2022

October 24-28, 2021

Rastislav Szabo
Senior Sw. Engineer
Kubermatic

Jussi Nummelin
Sr. Principal Engineer
Mirantis

Outline

- What?
  - What do we mean by “Remote Control Planes”
- Why?
  - Example use-cases for remote control planes
- How?
  - Building blocks and concepts to make this happen with a bit of history
  - Real-world integration examples

Acknowledgements

KubeCon CloudNativeCon North America 2022

- We’re not the inventors, merely happy users
- There’s very little documentation on the topic → We wanted to raise awareness
- Kudos to the original “inventors” and maintainers

Local Kubernetes Control Plane

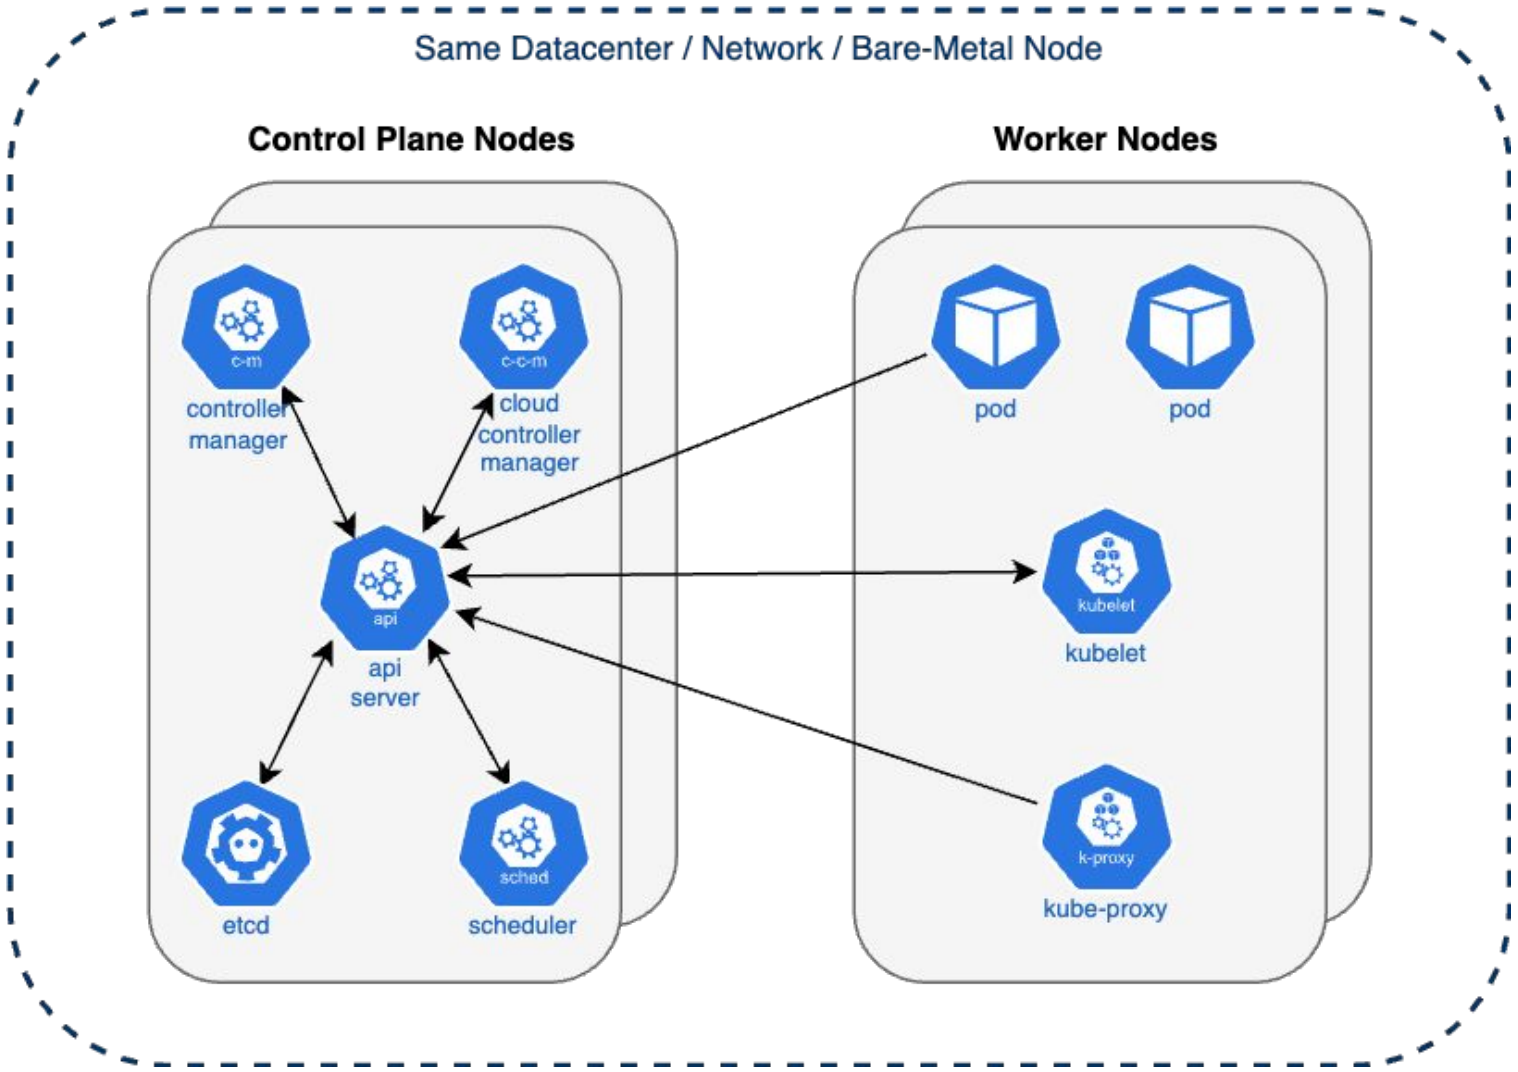

Control plane nodes ⇔ worker nodes connection:

- Bidirectional
- Unrestricted
- Fast
- Reliable
- Secure
- Often node-local / rack-local / same L2 network / same datacenter

Remote Kubernetes Control Plane

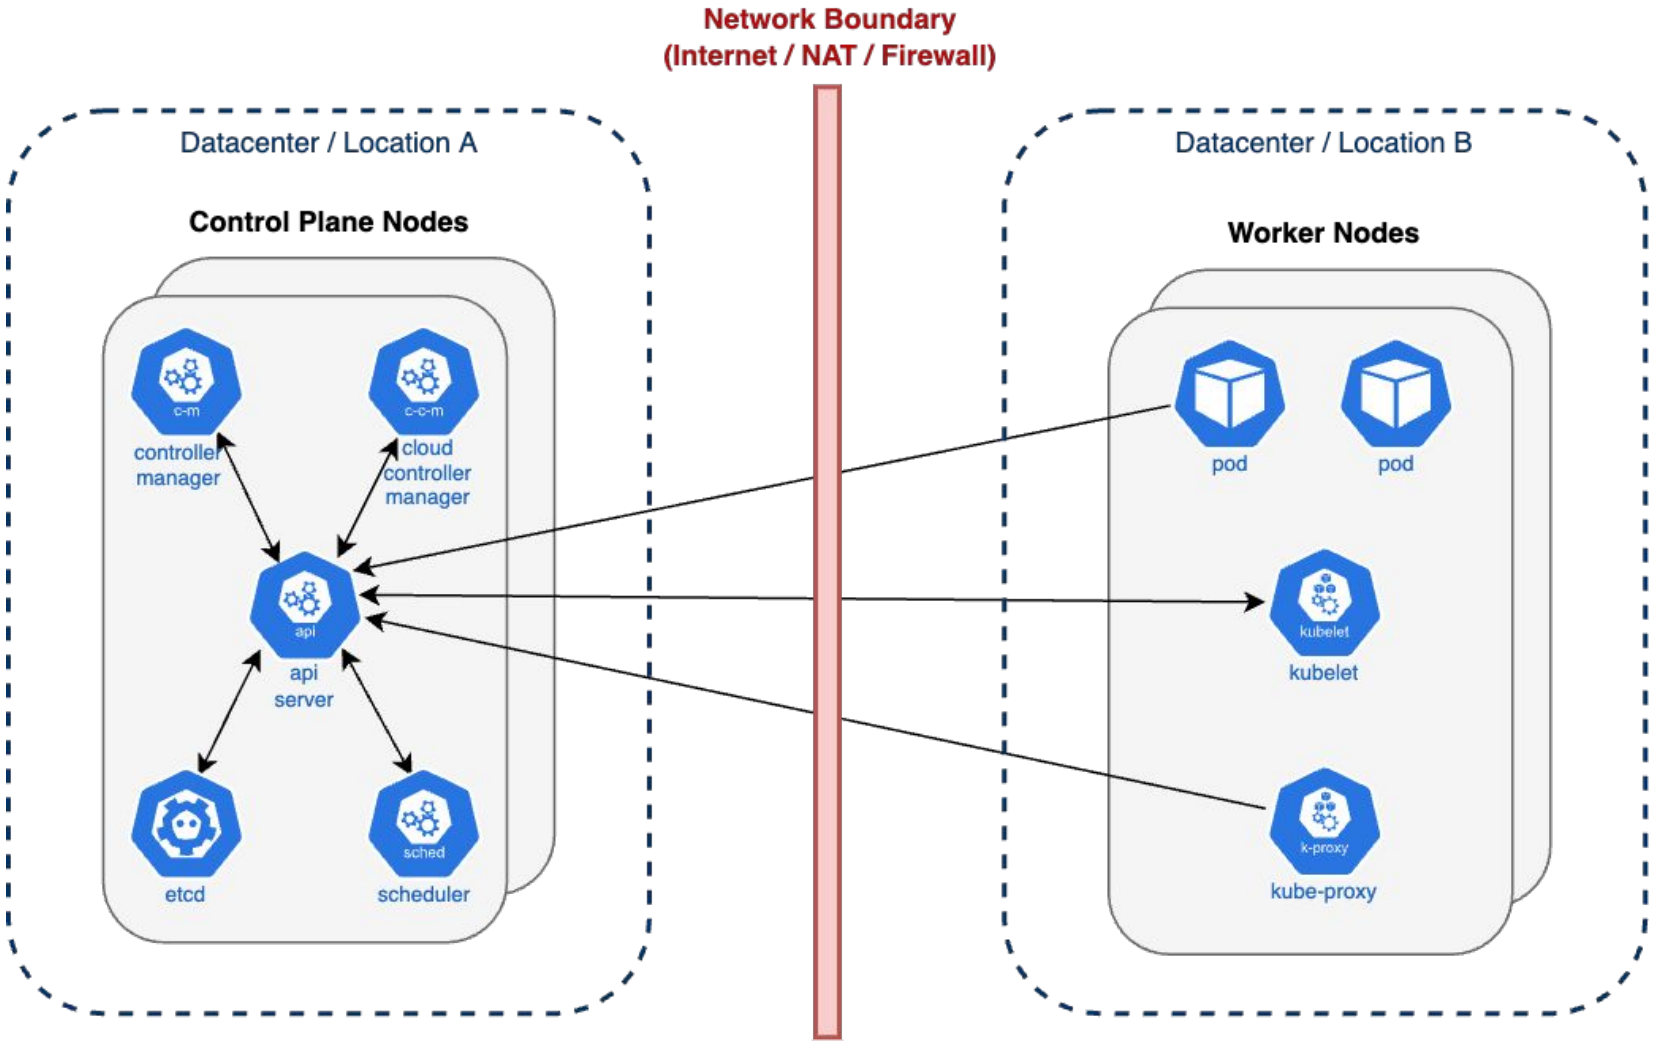

Control plane nodes ⇔ worker nodes connection:

- One-directional (NAT)
- Restricted (Firewall)
- Slower
- Less reliable
- Insecure
- Possibly over public Internet

Remote Control Plane Use Cases / 1

- Trust Segmentation
  - Control Plane in “protected / trusted” network, worker nodes in “untrusted” network
- Human Error Reduction
  - Only specialized personnel can access the control plane nodes / infrastructure
- Kubernetes at Edge
  - Worker nodes (e.g. resource-constrained) at edge, control plane in a datacenter
- Hybrid Cloud
  - Worker nodes on different platform / cloud than the control plane
  - Easy migration of worker-nodes between platforms

Remote Control Plane Use Cases / 2

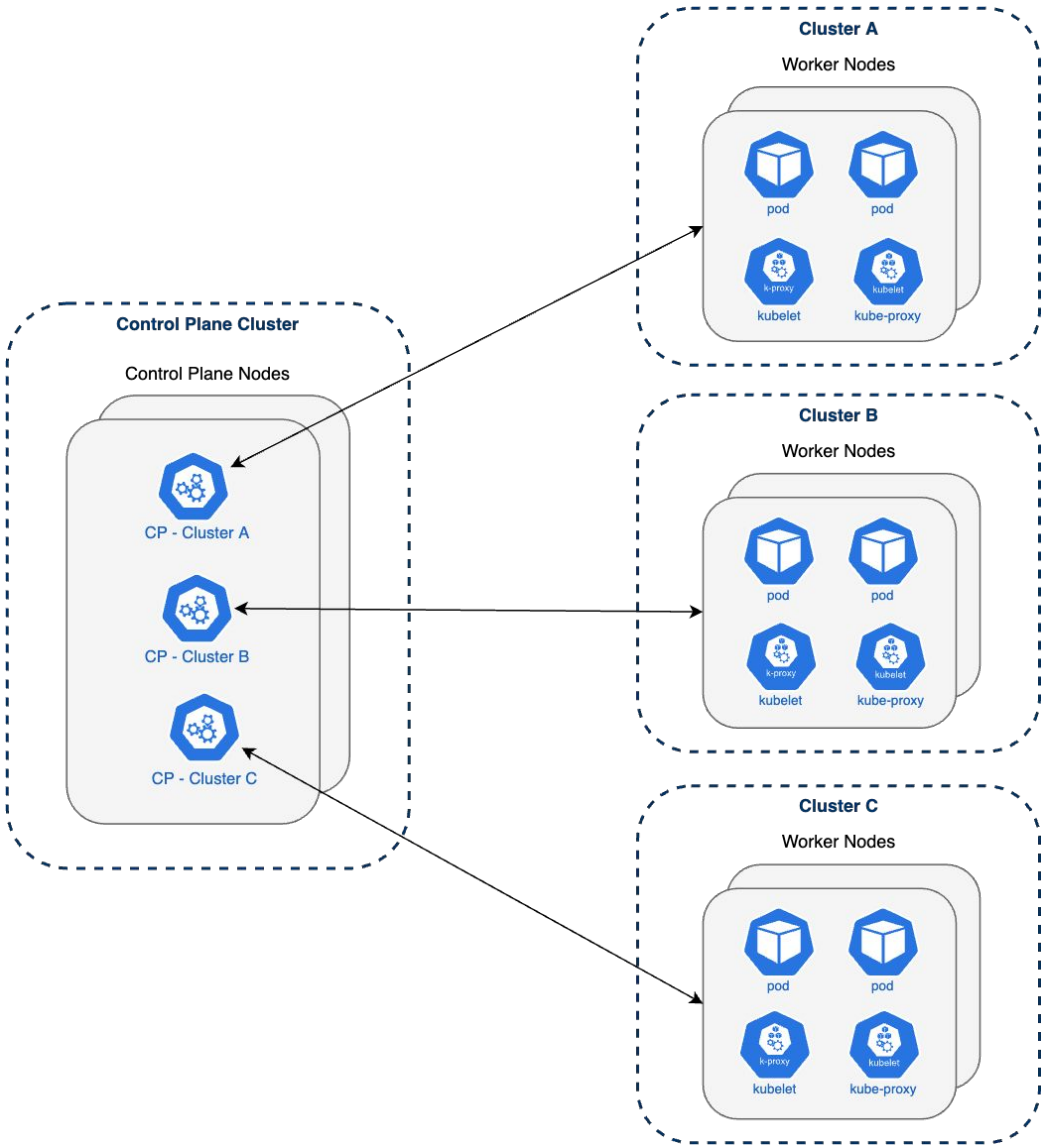

Co-located control-plane for multiple clusters:

- Easy operation for 100s / 1000s clusters
- HA control plane with wise resource usage
- Same control-plane experience across different (hybrid) cloud platforms
- Fast cluster spin-up time - good for temporary / short-lived clusters
- Build your own Kubernetes as a service

History

- Kubernetes supported SSH tunnels in the past, deprecated at v1.9
- Various custom solutions using things like VPN tunnels
- KEP-1281 was born in Spring 2019
- KEP-1281 is the architectural foundation for API server → workerplane communication routing

Concepts - EgressSelector

- *EgressSelector* defines how API talks to external components
- Almost like a typical network proxy
- Types: cluster, etcd, controlplane
  - cluster: pods/logs,pods/exec,svc/proxy,...
  - etcd: obviously KAS -> etcd connections
  - controlplane: admission etc. webhooks

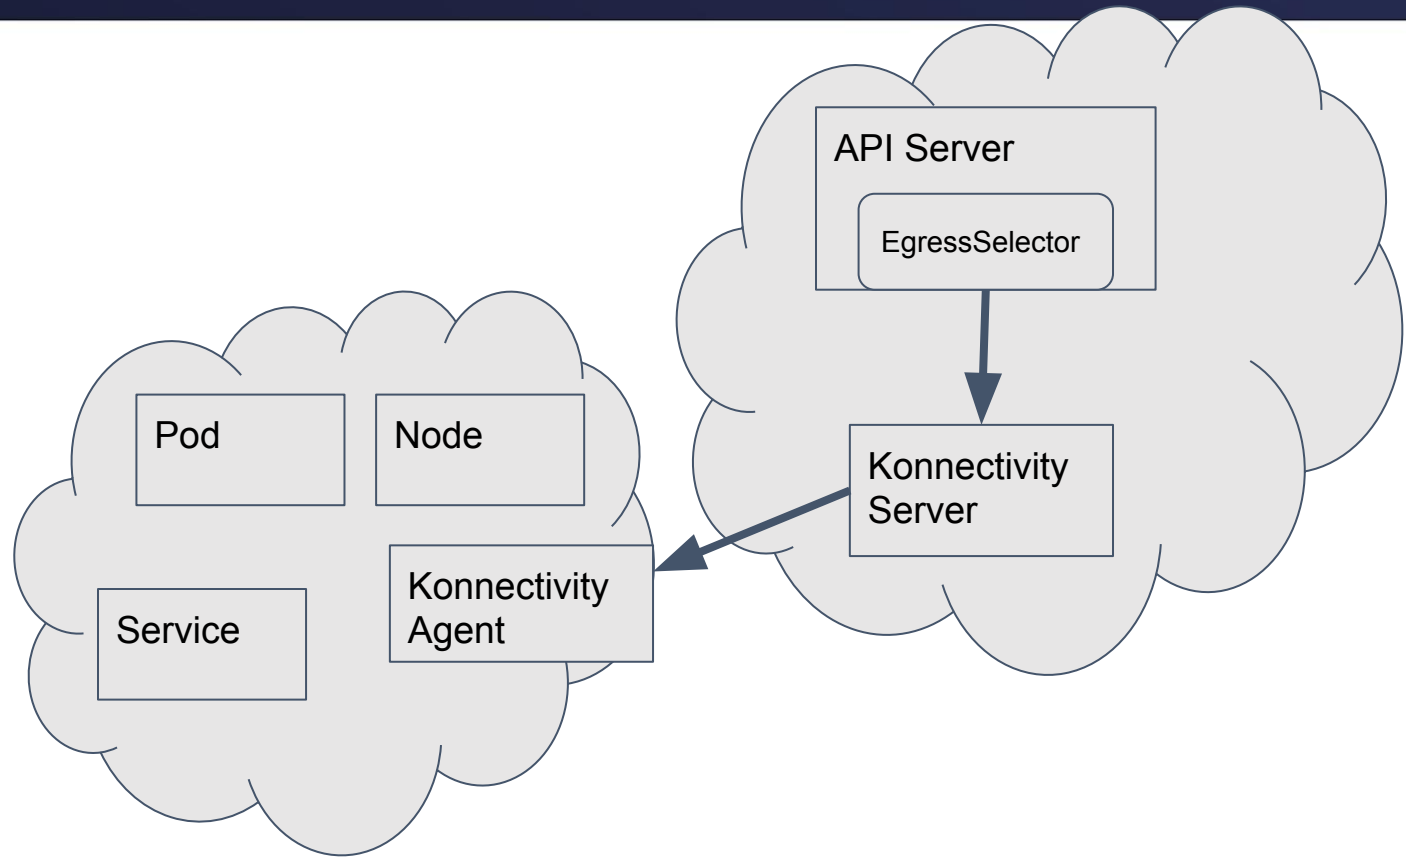


```
apiVersion: apiserver.k8s.io/v1beta1
kind: EgressSelectorConfiguration
egressSelections:
- name: cluster
  connection:
    proxyProtocol: GRPC
    transport:
      uds:
        udsName: /run/k0s/konnectivity-server/konnectivity-server.sock
```

Concepts - Tunneling protocol

- Custom protocol on top of gRPC over TLS
  - Works through NAT/FW/etc.
- Agent opens the bi-directional connection to server
- Much like SSH reverse tunnels

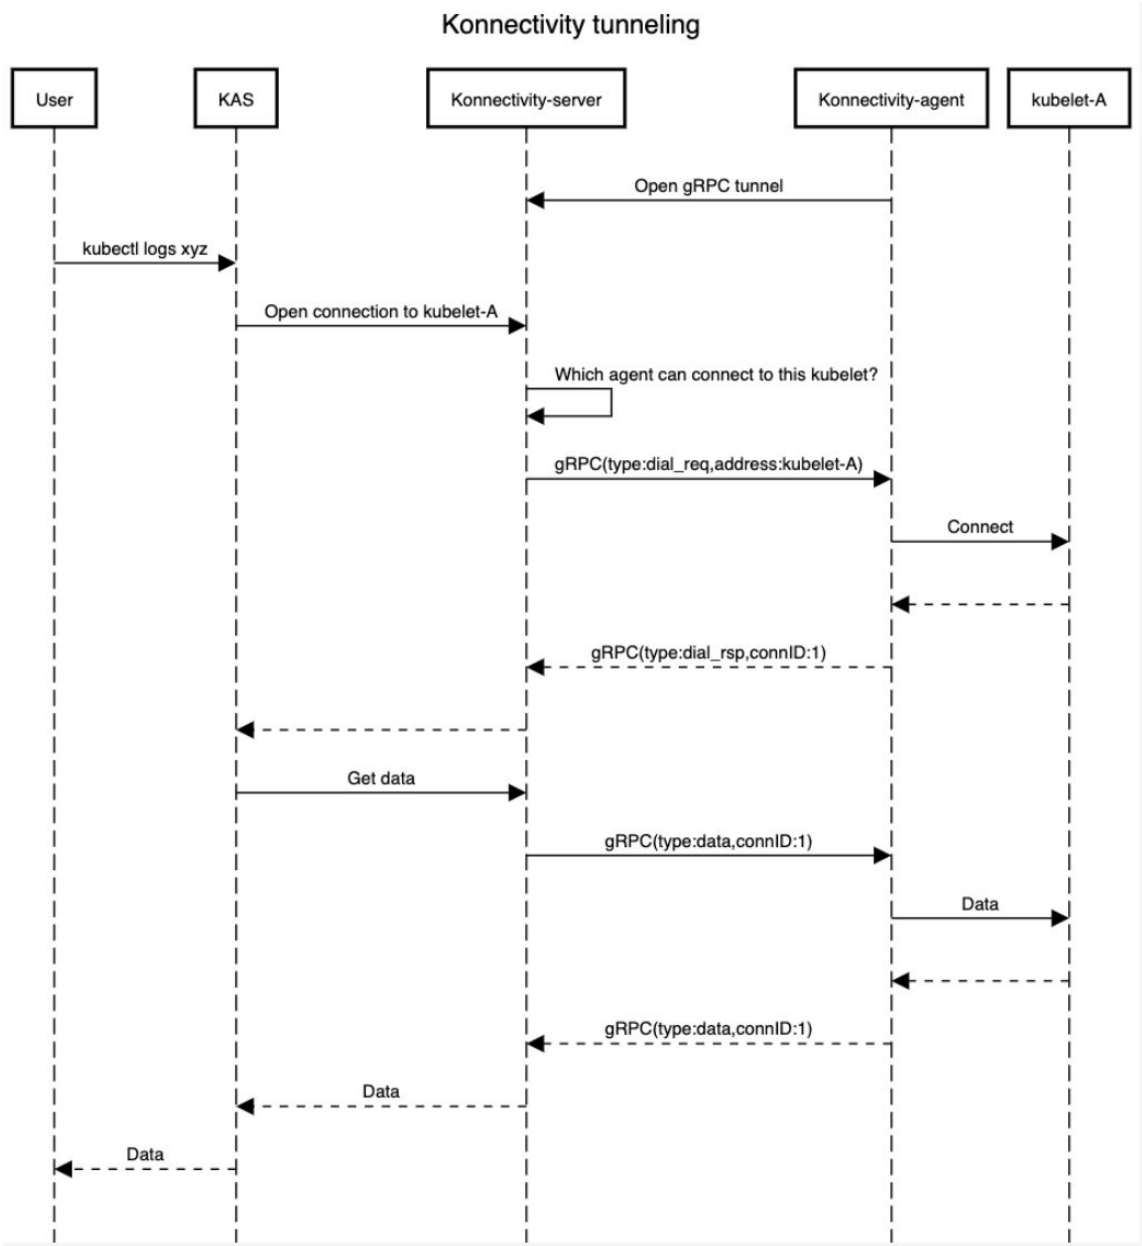

Real-World Integrations - Kubermatic

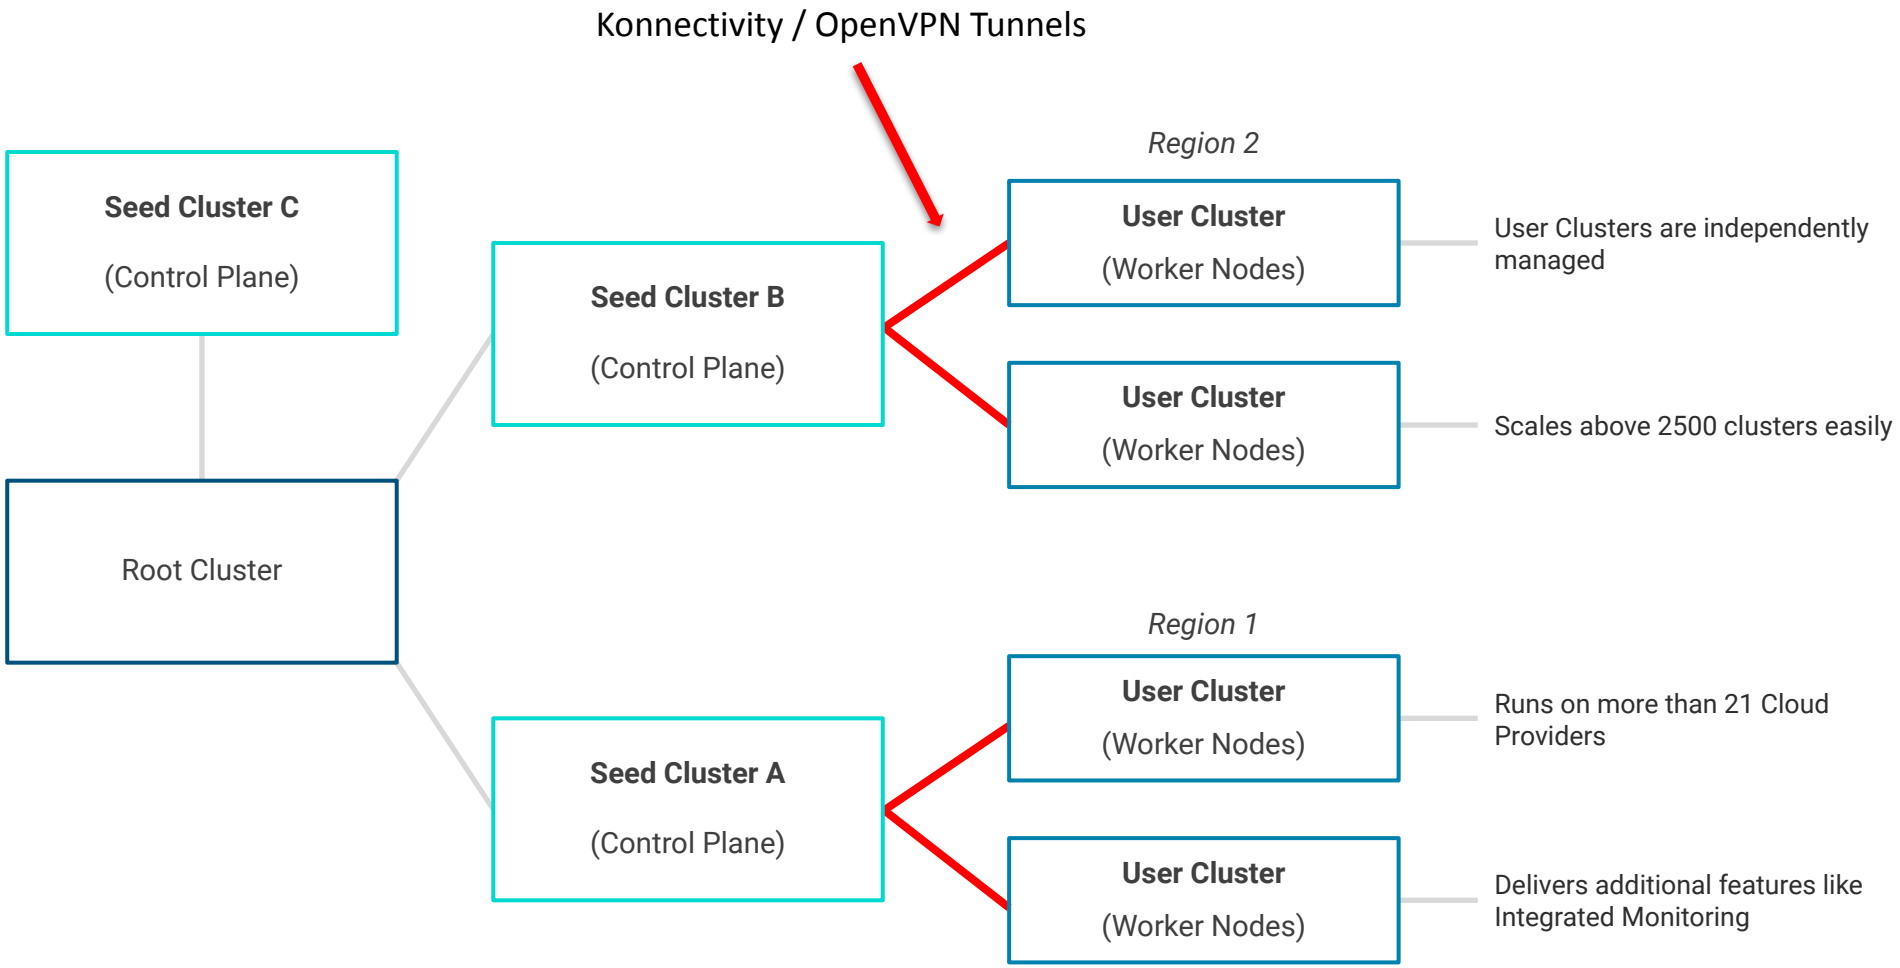

Real-World Integrations - Kubermatic

OpenVPN

Seed Cluster (Control Plane)

35 containers

NAME	READY	STATUS	RESTARTS	AGE
apiserver-65cbb55b6d-5dtlv	3/3	Running	0	2m52s
apiserver-65cbb55b6d-qc7vk	3/3	Running	0	4m25s
controller-manager-84487fdd55-5dlnm	2/2	Running	0	4m25s
dns-resolver-5979674cb-2d5t2	2/2	Running	0	2m52s
dns-resolver-5979674cb-9wn97	2/2	Running	0	4m24s
etcd-0	1/1	Running	0	15m
etcd-1	1/1	Running	0	2m20s
etcd-2	1/1	Running	0	15m
kube-state-metrics-74c7497b4d-jk55x	1/1	Running	0	14m
kubernetes-dashboard-847dd7bd4b-bp6c6	1/1	Running	2 (14m ago)	15m
kubernetes-dashboard-847dd7bd4b-f9dsl	1/1	Running	2 (14m ago)	15m
machine-controller-7fbbf69f88-r9tkn	1/1	Running	1 (14m ago)	15m
machine-controller-webhook-b87768f6-g4nhj	1/1	Running	0	15m
metrics-server-74488f6f47-nlklc	3/3	Running	0	2m52s
metrics-server-74488f6f47-xf2v2	3/3	Running	2 (4m1s ago)	4m24s
openvpn-server-697b5955b8-jx99p	3/3	Running	0	4m25s
operating-system-manager-69f9fdd4c7-7tzcj	1/1	Running	0	15m
prometheus-0	1/1	Running	0	4m30s
scheduler-5d84d49c45-67rmt	2/2	Running	0	4m25s
usercluster-controller-74b6954f5-sxh4g	1/1	Running	0	4m25s
usercluster-webhook-6646995fd4-dzgbc	1/1	Running	0	15m

User Cluster (Worker Nodes)

10 containers

calico-kube-controllers-57fb8785bf-jl226	1/1	Running	0	24m
canal-8mvtr	2/2	Running	0	20m
coredns-58b65bfd4d-d92mp	1/1	Running	0	24m
coredns-58b65bfd4d-w8df2	1/1	Running	0	24m
kube-proxy-6p5tv	1/1	Running	0	20m
node-local-dns-fvtrq	1/1	Running	0	20m
openvpn-client-68bc8f76bd-m8b85	2/2	Running	0	8m59s
user-ssh-keys-agent-s8xvx	1/1	Running	0	20m

Konnectivity

-15 containers

Seed Cluster (Control Plane)

18 containers

NAME	READY	STATUS	RESTARTS	AGE
apiserver-5cc59db7-5m2mr	2/2	Running	0	11m
apiserver-5cc59db7-txvpc	2/2	Running	0	8m6s
controller-manager-65f7b96c8c-lhkm9	1/1	Running	0	11m
etcd-0	1/1	Running	0	11m
etcd-1	1/1	Running	0	11m
etcd-2	1/1	Running	0	11m
kube-state-metrics-74c7497b4d-jk55x	1/1	Running	0	9m49s
kubernetes-dashboard-847dd7bd4b-bp6c6	1/1	Running	2 (10m ago)	11m
kubernetes-dashboard-847dd7bd4b-f9dsl	1/1	Running	2 (10m ago)	11m
machine-controller-7fbbf69f88-r9tkn	1/1	Running	1 (10m ago)	11m
machine-controller-webhook-b87768f6-g4nhj	1/1	Running	0	11m
operating-system-manager-69f9fdd4c7-7tzcj	1/1	Running	0	11m
prometheus-0	1/1	Running	0	9m49s
scheduler-67fb45677-8k9s6	1/1	Running	0	11m
usercluster-controller-65f8596445-278vq	1/1	Running	0	11m
usercluster-webhook-6646995fd4-dzgbc	1/1	Running	0	11m

User Cluster (Worker Nodes)

12 containers

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-57fb8785bf-jl226	1/1	Running	0	26m
canal-8mvtr	2/2	Running	0	22m
coredns-58b65bfd4d-d92mp	1/1	Running	0	25m
coredns-58b65bfd4d-w8df2	1/1	Running	0	25m
konnectivity-agent-77c956d6cd-n8jst	1/1	Running	0	7s
konnectivity-agent-77c956d6cd-pzrqw	1/1	Running	0	7s
kube-proxy-6p5tv	1/1	Running	0	22m
metrics-server-677cf6b8dd-g7bvq	1/1	Running	0	7s
metrics-server-677cf6b8dd-wbwcr	1/1	Running	0	7s
node-local-dns-fvtrq	1/1	Running	0	22m
user-ssh-keys-agent-s8xvx	1/1	Running	0	22m

Real-World Integrations - K0s

- Konnectivity built-in in k0s
- Enables our “controlplane isolation” feature
- Enables some very interesting deployment architectures

Challenges

- Lack of documentation
- HA setup is bit tricky; Especially in dynamic environments
- Debuggability; When things are not working as expected users see bit obscure errors
- Lack of contributors; PRs can take quite a while to land
- Real-world Testing; Resource leaks are known to happen

KubeCon | CloudNativeCon
North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Please scan the QR Code above to
leave feedback on this session