



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

SIG-NETWORK: Intro and Deep Dive

Andrew Stoycos, Red Hat

Bowei Du, Google

Rob Scott, Google

Surya Seetharaman, Red Hat

- Networking APIs
 - Service, Endpoint(Slice), Ingress, Gateway API, DNS - Rob Scott
 - NetworkPolicy, Admin Network Policy, Multi Network - Surya
- Networking Components (Kube-Proxy, KPNG) - Andrew
- Features in development - Bowei
 - Topology
 - Terminating Endpoints
 - NetworkPolicy Status
 - InternalTrafficPolicy

Networking APIs

- Enables grouping Pods together and exposing as a network Service
- Services are assigned IP address(es) they can be reached on
- Requests to those addresses will be routed to one of the associated Pods (endpoints)


```
apiVersion: v1
kind: Service
metadata:
  name: store
spec:
  selector:
    app: store
  ports:
    - name: tcp
      protocol: TCP
      port: 80
      targetPort: 8080
```

- Enables grouping Pods together and exposing as a network Service
- Services are assigned IP address(es) they can be reached on
- Requests to those addresses will be routed to one of the associated Pods (endpoints)


```
apiVersion: v1
kind: Service
metadata:
  name: store
spec:
  selector:
    app: store
  ports:
    - name: tcp
      protocol: TCP
      port: 80
      targetPort: 8080
```

Service Types

- ClusterIP
- NodePort
- LoadBalancer
- ExternalName




Reaching a Service



```
apiVersion: v1
kind: Service
metadata:
  name: store
spec:
  clusterIP: 10.1.2.3
  selector:
    app: store
  ports:
    - name: tcp
      protocol: TCP
      port: 80
      targetPort: tcp
```



Reaching a Service

```
curl store.prod.svc.cluster.local:80
```



```
apiVersion: v1
kind: Service
metadata:
  name: store
  namespace: prod
spec:
  clusterIP: 10.1.2.3
  selector:
    app: store
  ports:
    - name: tcp
      protocol: TCP
      port: 80
      targetPort: tcp
```

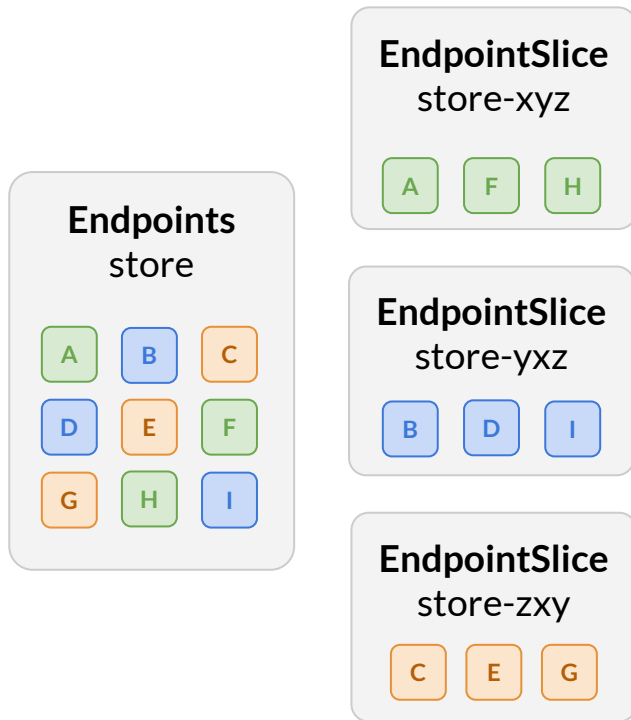
Reaching a Service



```
apiVersion: v1
kind: Service
metadata:
  name: store
  namespace: prod
spec:
  clusterIP: 10.1.2.3
  selector:
    app: store
  ports:
    - name: tcp
      protocol: TCP
      port: 80
      targetPort: tcp
```

Endpoints and EndpointSlices

- Track IPs and Ports for Pods backing a Service
- Endpoints limited to 1000 Pods per Service
- EndpointSlices are sharded Endpoints, much more scalable
- New features enabled by EndpointSlices:
 - Dual Stack
 - Topology
 - Terminating Endpoints



Gateway and Ingress

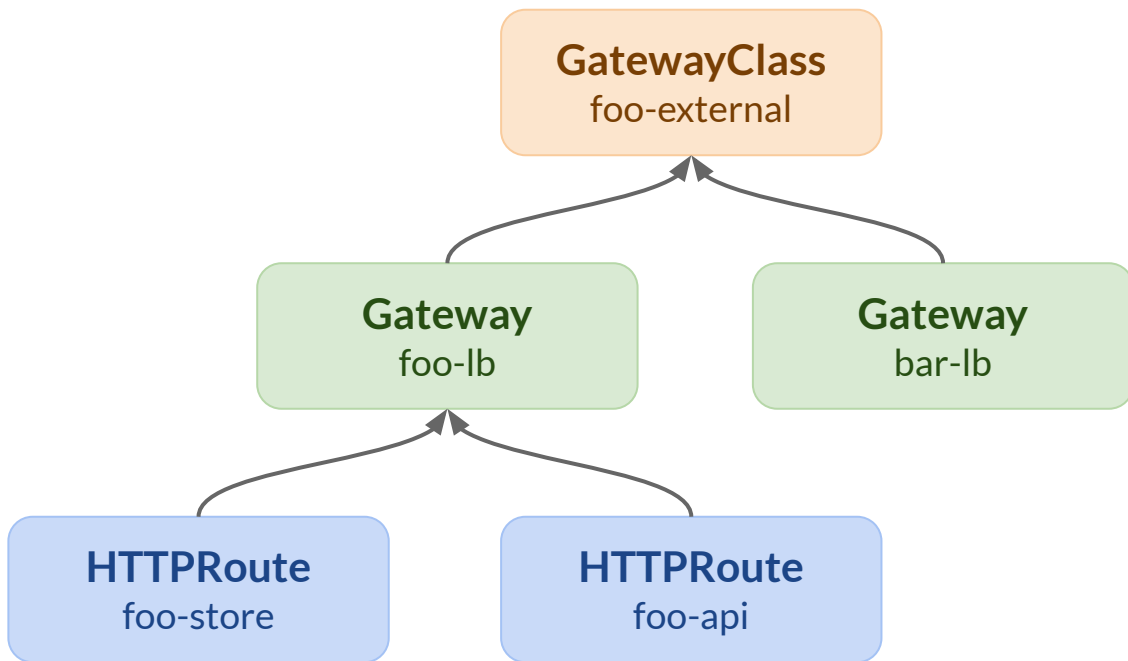
The Ingress API

- Host and Path Matching
- Forward to Service
- TLS Configuration
- Stable for 5 years
- Simple and broadly implementable

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: minimal-ingress
spec:
  ingressClassName: nginx-example
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          service:
            name: test
            port:
              number: 80
```

Gateway API

- Next generation of Kubernetes routing and load balancing APIs
- Designed to be expressive and extensible
- Role oriented resource model
- 15 implementations
- Graduated to beta in July



Simple Path Match

Ingress

```
ingressClassName: nginx
rules:
- http:
    paths:
    - path: /login
      pathType: Prefix
      backend:
        service:
          name: auth-svc
          port:
            number: 8080
```


HTTPRoute

```
parentRefs:
- name: nginx
rules:
- matches:
  - path:
      type: PathPrefix
      value: /login
  backendRefs:
  - name: auth-svc
    port: 8080
```

Get Involved

- Focus Areas: Mesh, L4, L7 -> GA, Conformance
- Weekly community meetings on Mondays
- Weekly GAMMA meetings on Tuesdays
- Contributors from all backgrounds welcome

gateway-api.sigs.k8s.io



#sig-network-gateway-api



kubernetes-sigs/gateway-api


Network Policy

Contributions welcome! Stable for over 5 years.

[network-policy-api](https://kubernetes.io/docs/concepts/services-networking/network-policies/) && <https://kubernetes.io/docs/concepts/services-networking/network-policies/>

Network Policy (NP)

- How can app owners control traffic to/from their workloads?
 - example; backends can get traffic only from frontends, databases can only get traffic from backends etc..



Network Policy (NP)

- How can app owners control traffic to/from their workloads?
 - example; backends can get traffic only from frontends, databases can only get traffic from backends etc..
- An API that let's users define simple ingress/egress rules
- API design is implicit in nature
- Network policy peers
 - pod, namespace, ipBlock

```
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-ingress-to-backend-from-frontend
  namespace: foo
spec:
  podSelector:
    matchLabels:
      app: backend
  policyTypes:
    - Ingress
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            app: frontend
```


Admin Network Policy

Contributions welcome! Under active development!

<https://github.com/kubernetes-sigs/network-policy-api> && <https://network-policy-api.sigs.k8s.io/>

Admin Network Policy (ANP)

- Network Policies were designed for app owners...
- How can admins enforce policies cluster-wide??





Admin Network Policy (ANP)

- Network Policies were designed for app owners...
- How can admins enforce policies cluster-wide??
- Cluster-scoped policy API
 - AdminNetworkPolicy
 - BaselineAdminNetworkPolicy
- API design is explicit in nature

```
apiVersion: policy.networking.k8s.io/v1alpha1
kind: AdminNetworkPolicy
metadata:
  name: cluster-wide-deny-example
spec:
  priority: 1
  subject:
    namespaces:
      matchLabels:
        kubernetes.io/metadata.name: sensitive-ns
  ingress:
    - name: "default-deny-to-sensitive-ns"
      action: "Deny"
      from:
        - namespaces:
            related: "NotSelf"
```


Admin Network Policy (ANP)

- Network Policies were designed for app owners...
- How can admins enforce policies cluster-wide??
- Cluster-scoped policy API
 - AdminNetworkPolicy
 - BaselineAdminNetworkPolicy
- API design is explicit in nature
- v1alpha1 supports east-west traffic
- TBD: [north-south traffic support](#)
- Next Steps: Implementations!



Get Involved

- Focus Areas: Network Policies, Admin Network Policies
- Bi-Weekly community meetings on Monday at 4pm EDT/1pm PDT
- All are Welcome!



#sig-network-policy-api



kubernetes-sigs/network-policy-api



Multi Network

Contributions welcome! Ongoing design! New SIG!

https://docs.google.com/document/d/1ztx9TOQ9Hij9PG9aPv6jyDLhe_FB7haV_yjJlcb-0Y/edit#heading=h.5t6qeinm4208

Multi-Network

- Apps might need to access different isolated networks via performance oriented interfaces like SR-IOV
- Represent more complex networking solutions via an API
- In early stages (design phase)



Get Involved

- Focus Areas: Multi Network ([intro slide deck](#))
- Bi-Weekly community meetings on Wednesday at 8am PST/ 5pm CEST
- Welcome!




#sig-network-multi-network

Networking Components

Kube-Proxy

Kube-Proxy

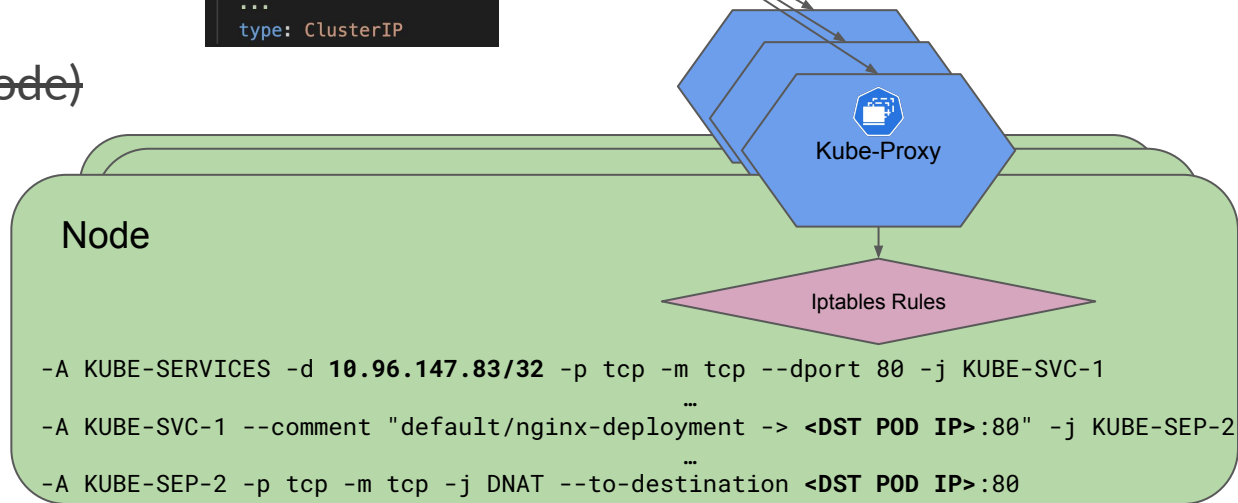
- We talked about some great K8s Networking APIs including **Services** and **Endpoints**
- Kube-Proxy implements the Services API
- Node agent which resolves K8s Networking Objects into datapath rules



Kube-Proxy

- Implemented in core Kubernetes
- Can program datapath rules in three different modes
 - iptables (default)
 - ipvs
 - ~~○ userspace ("legacy" mode)~~

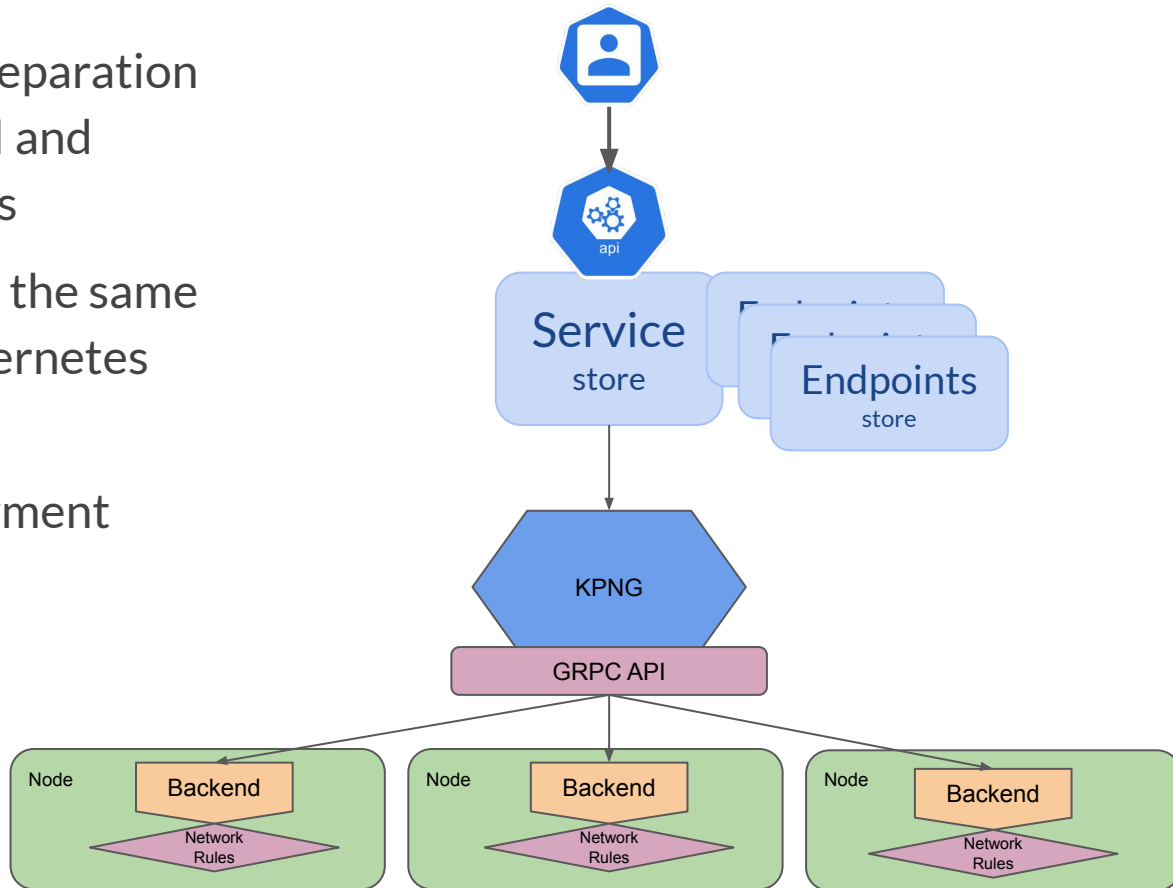
```
apiVersion: v1
kind: Service
metadata:
  name: nginx-deployment
  namespace: default
...
spec:
  clusterIP: 10.96.147.83
  ...
  - port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: nginx
  ...
  type: ClusterIP
```



KPNG

(Kube-Proxy Next Generation)


- Provides a fundamental separation between the Services API and Backend Implementations
- Allows backends to reuse the same bits which watch the kubernetes API
- Extremely flexible deployment model



- Current Backends

- Iptables
- NFT
- Ipv6
- Userspace
- Windows
- Ebpf (POC)

- More to come from Viewers like You!



Networking Components: Get Involved

- Now that you know a bit more....Get involved!
 - [So You Want To Write A Service Proxy...](#)
 - [Kube-Proxy good-first-issues](#)
 - [KPNG good-first-issues](#)



Features in Development

Topology Aware Hints

Pods

Zone A



Zone B




Zone C



EndpointSlices


Original



Allocate proportional endpoints
to each zone

Topology Hints

For: Zone A



For: Zone B



For: Zone C

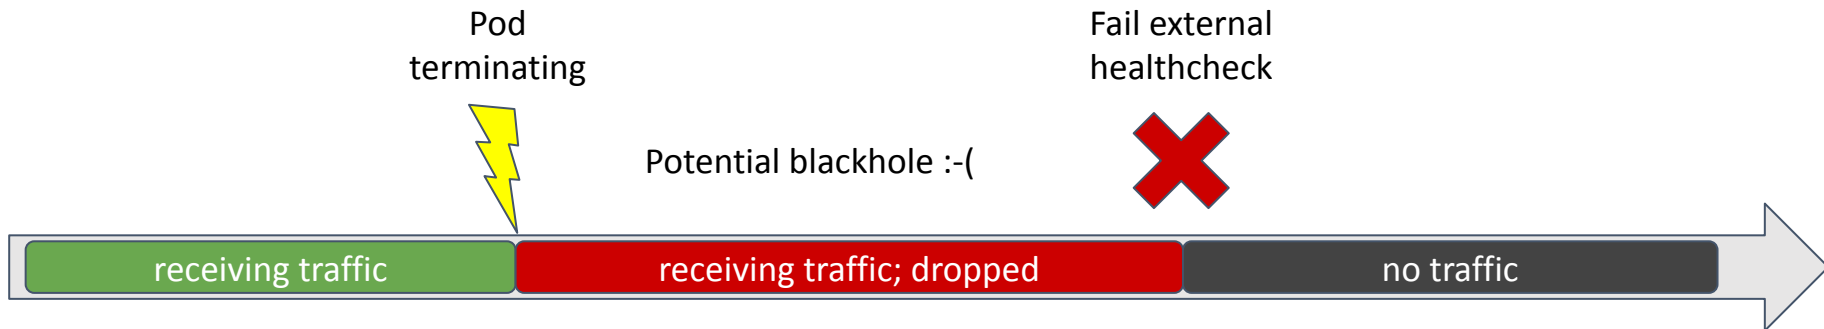


Terminating Endpoints

Tracking terminating endpoints ([KEP](#), [KEP](#))

Gracefully keep sending traffic to an `ExternalTrafficPolicy=local` Pods to include those that are terminating. This avoids a **blackhole** between when the Pods terminate and when the LB notices the Pods are shutting down.

Adds “terminating” state to EndpointSlice.

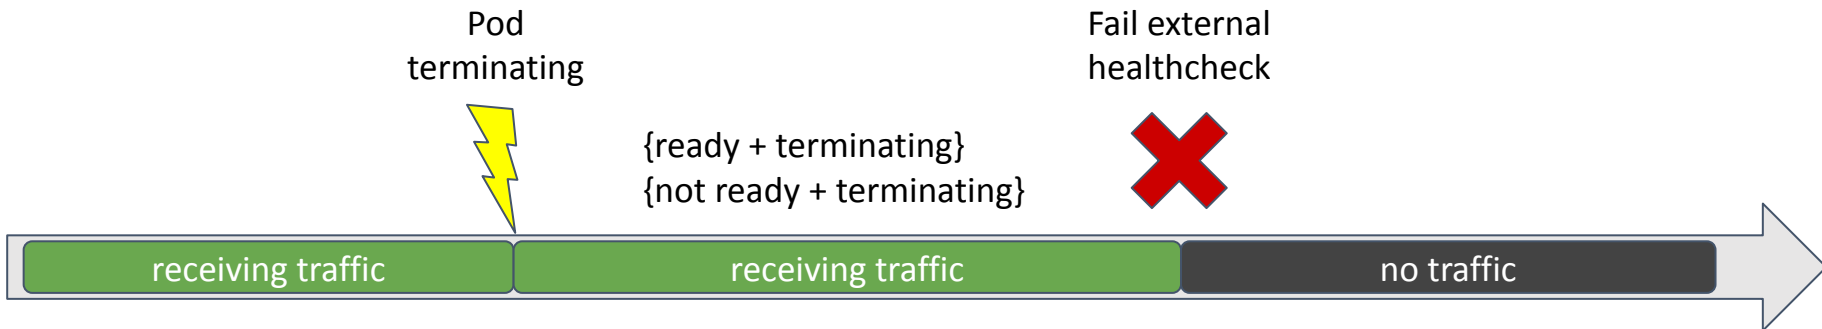


Terminating Endpoints

Tracking terminating endpoints ([KEP](#), [KEP](#))

Gracefully keep sending traffic to an `ExternalTrafficPolicy=local` Pods to include those that are terminating. This avoids a **blackhole** between when the Pods terminate and when the LB notices the Pods are shutting down.

Adds “terminating” state to EndpointSlice.



NetworkPolicy Status

NetworkPolicy now has a Status and Conditions fields.

This allows implementations to signal state about your policy:

- Optional feature support such as portRanges.

```
apiVersion: v1
kind: NetworkPolicy
metadata:
  name: my-policy
spec:
  ...
status:
  conditions:
  - ...
```


InternalTrafficPolicy

Makes your service route only to endpoints local to the Node for traffic sent to the Service ClusterIP.

Typical use case: Service backed by a Daemonset for a logging service with traffic that you never want to traverse outside of the node.

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app.kubernetes.io/name: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
  internalTrafficPolicy: Local
```

Questions?