



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

**DETROIT 2022**

# SIGs Aren't Silos

A Case Study Into Solving Inter-Domain Problems In Kubernetes Development



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

**DETROIT 2022**

**October 24-28, 2021**



**Swetha Repakula**

Software Engineer  
*Google*




**Antonio Ojea**

Software Engineer  
*Google*

# Kubernetes community: Governance



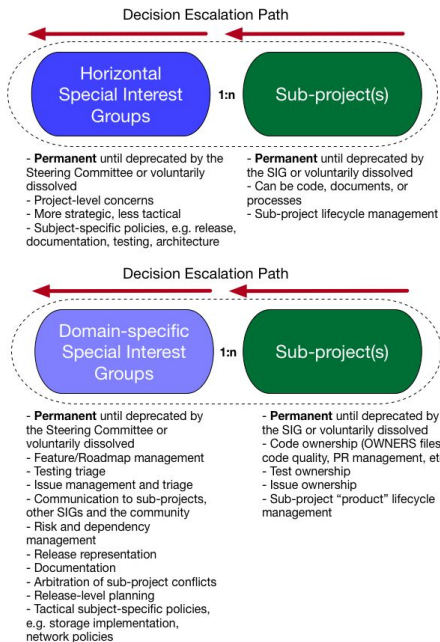
## Kubernetes Community Governance Model



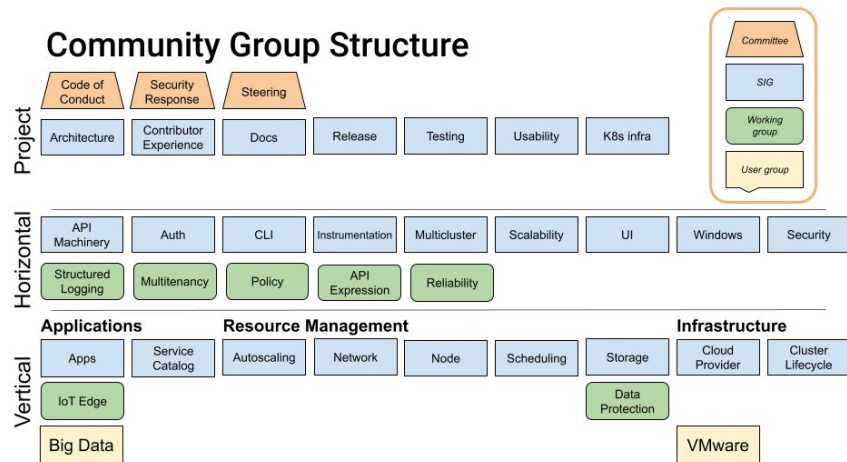
- Project-level governance
- Binding/final arbitration
- Management of sub-structures
- Security process management
- Manage project-level policies such as the Code of Conduct



- Organized to solve a **specific problem** then **dissolve**
- Cross-SIG collaboration around a specific effort
- Does not own code
- Can spawn sub-projects in participating SIGs
- Define and address gaps




## Community Group Structure



# Kubernetes components



K8s cluster





# Kubernetes community: People/Friends



# SIG Network and Node bugs

Pods with failed status IP address reused on new pods, but traffic still going to old pods across namespaces. #109414

New issue

🔒 Closed declangallagher opened this issue on Apr 11 · 11 comments · Fixed by #110255

"Terminated" pod on shutdown node listed in service endpoints. #109718

New issue

🔒 Closed tcolgate opened this issue on Apr 29 · 35 comments

Kubernetes sending traffic to draining nodes #110195

New issue

🔒 Closed Maria-Milusheva opened this issue on May 24 · 10 comments

Kubelet Readiness Probes do not run during pod termination #110309

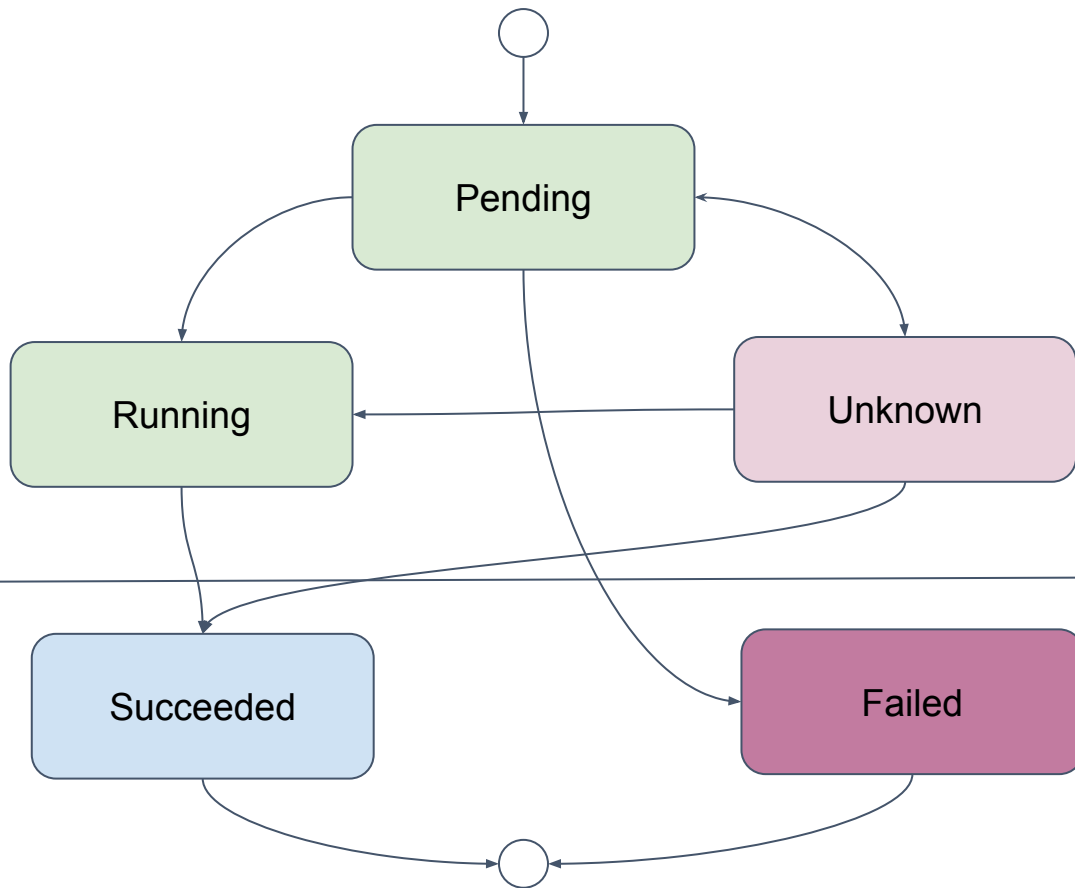
New issue

🔒 Closed rphillips opened this issue on May 31 · 5 comments · Fixed by #110191

# Symptoms

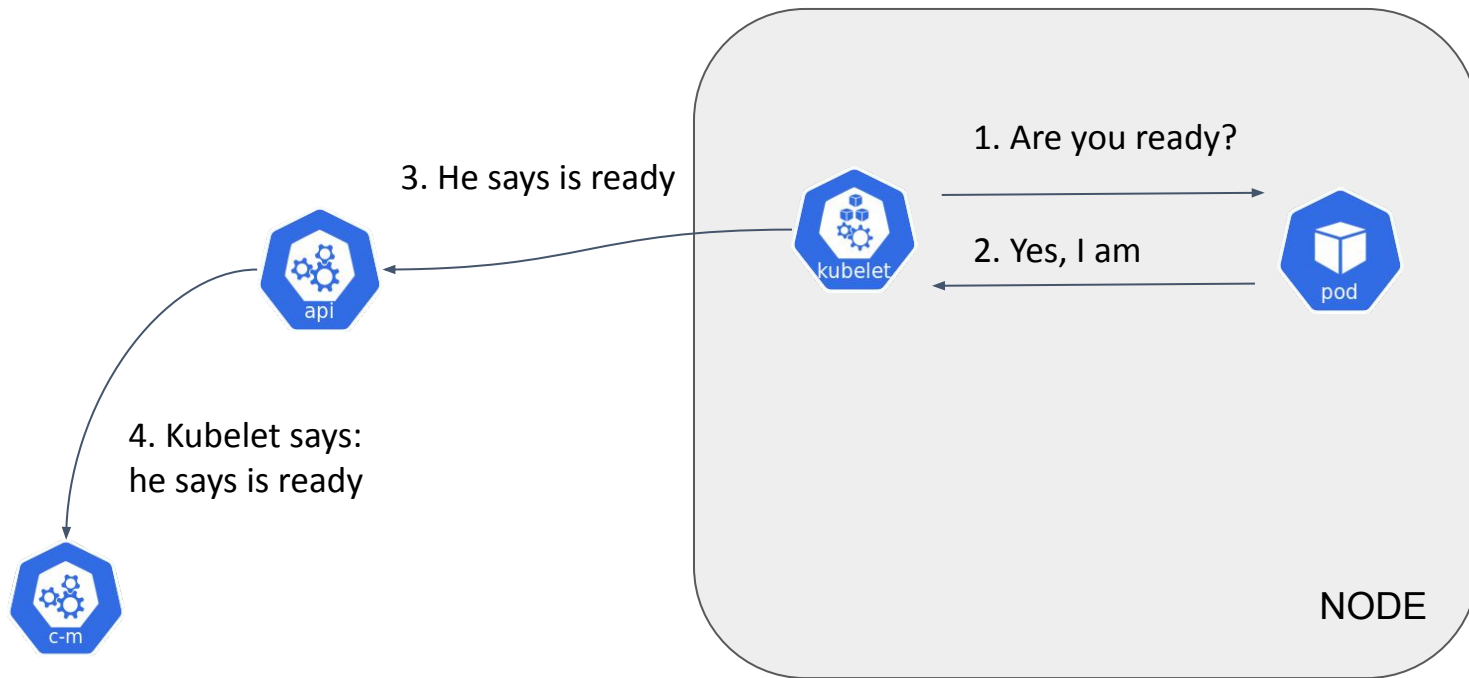
- Traffic was being routed to non-existent pods
- Traffic was being routed to the wrong pods
- IPs in EndpointSlices were not matching the Node Pod CIDR

# Pod Lifecycle: Pod Phases

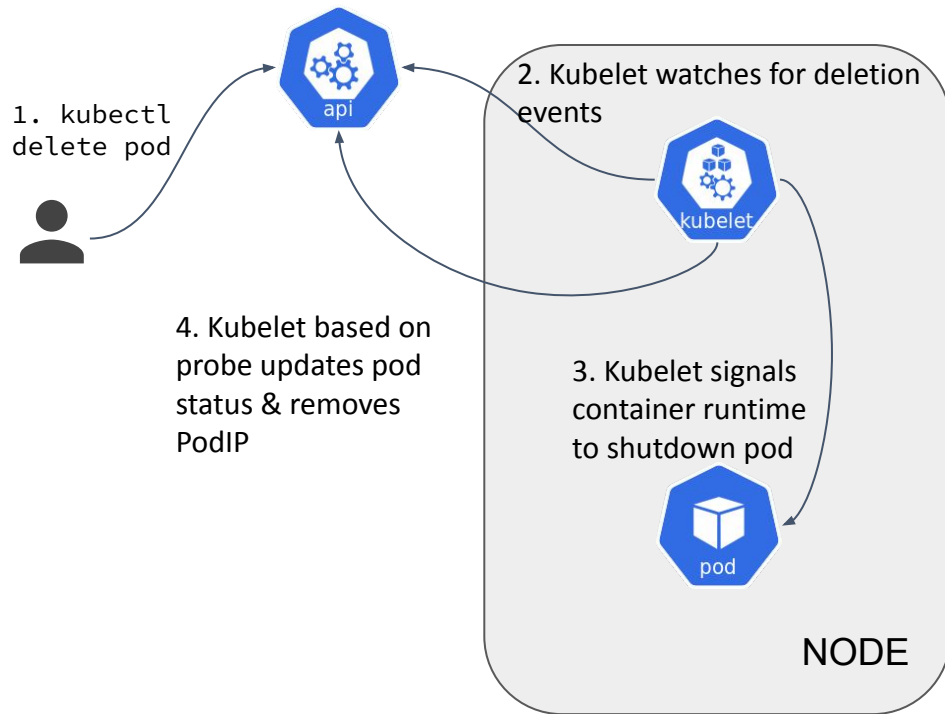




# Pod Lifecycle: Pod Readiness



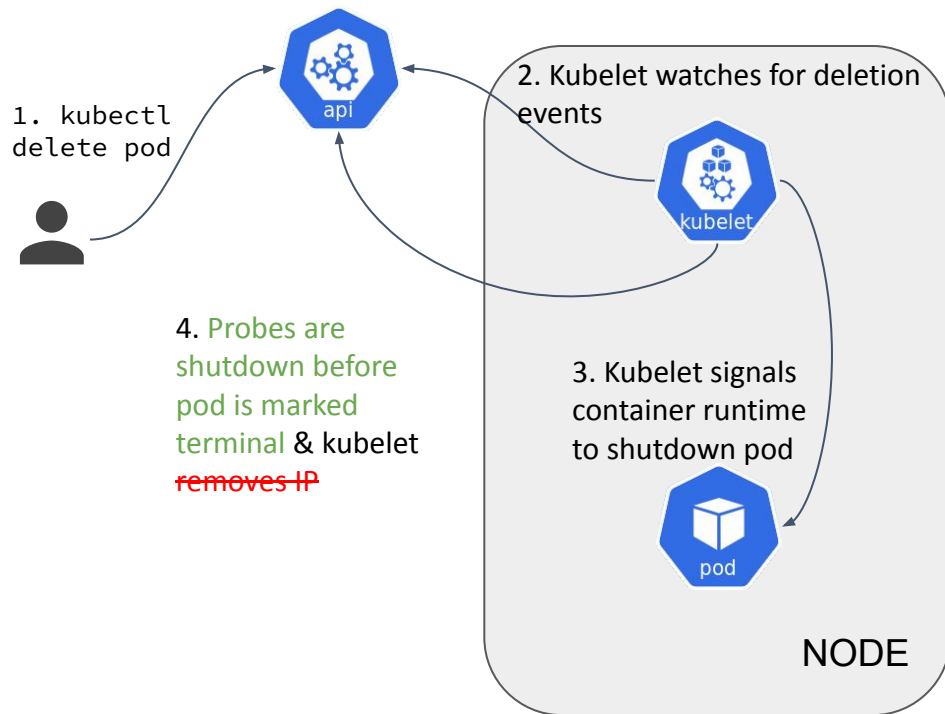
# Pod Shutdown - Before Refactor



When nodes are shutdown, or if pods are evicted, pods will not be deleted until PodGC kicks in (based on threshold).

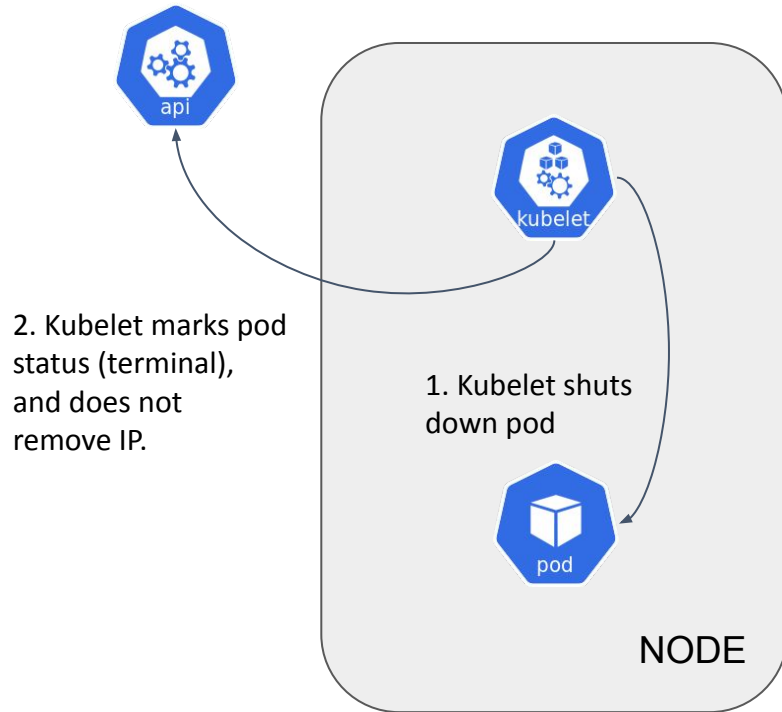
```
$ kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP            NODE
busybox  0/1     Completed 0           15m   10.244.0.2    121-control-plane
busybox3 0/1     StartError 0           3m56s 10.244.0.6    121-control-plane
evictme  0/1     Evicted    0           63s   <none>        121-control-plane
```

# Pod Shutdown - After Refactor



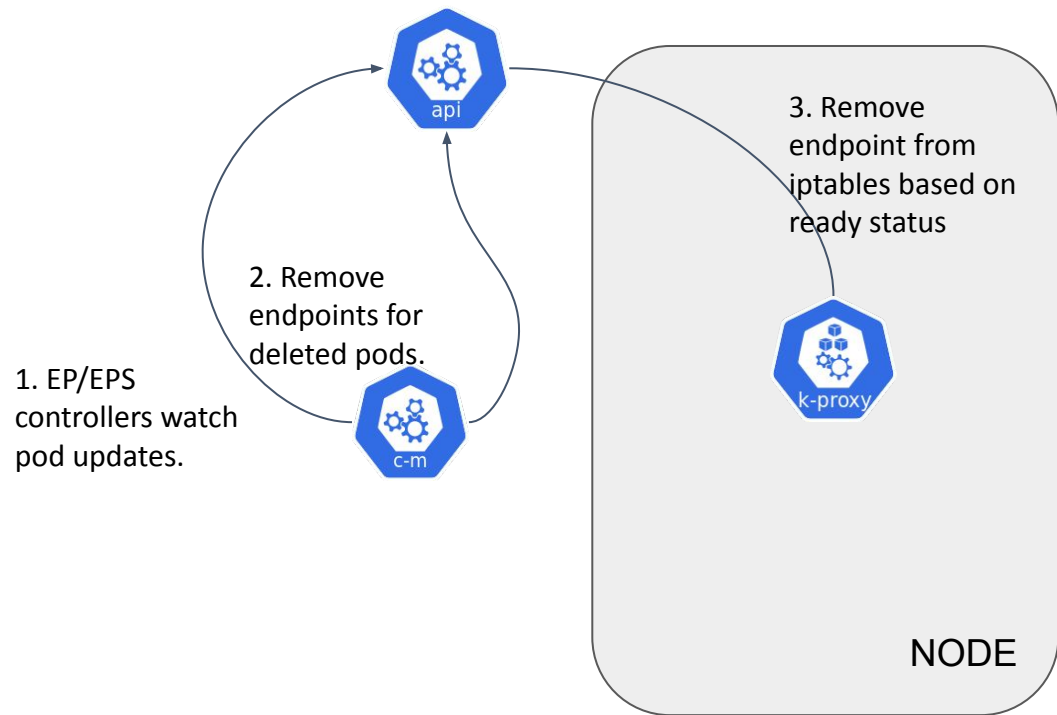
```
$ kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE
busybox   0/1     Completed 0          22m   10.244.0.5    kind-control-plane
busybox3  0/1     StartError 0          5m36s  10.244.0.6    kind-control-plane
evictme   0/1     Evicted    0          63s   10.244.0.7    kind-control-plane
```

# Pod Shutdown - Pod Eviction



When nodes are shutdown, or if pods are evicted, pods will not be deleted until PodGC kicks in (based on threshold).

# Service/Endpoints Removal



# The bug ...

```
448     for _, pod := range pods {
449         if len(pod.Status.PodIP) == 0 {
450             klog.V(5).Infof("Failed to find an IP for pod %s/%s", pod.Namespace, pod.Name)
451             continue
452         }
453         if !tolerateUnreadyEndpoints && pod.DeletionTimestamp != nil {
454             klog.V(5).Infof("Pod is being deleted %s/%s", pod.Namespace, pod.Name)
455             continue
456         }
457     }
```

pkg/controller/endpoint/endpoints\_controller.go (v1.22.0)




# The bug ...

```
448     for _, pod := range pods {
449         if len(pod.Status.PodIP) == 0 {
450             klog.V(5).Infof("Failed to find an IP for pod %s/%s", pod.Namespace, pod.Name)
451             continue
452         }
453         if !tolerateUnreadyEndpoints && pod.DeletionTimestamp != nil {
454             klog.V(5).Infof("Pod is being deleted %s/%s", pod.Namespace, pod.Name)
455             continue
456         }
457     }
```

pkg/controller/endpoint/endpoints\_controller.go (v1.22.0)

# Pod Shutdown - Pod Eviction



When nodes are shutdown, or if pods are evicted, pods will not be deleted until PodGC kicks in (based on threshold).

# The bug ...

```
448     for _, pod := range pods {
449         if len(pod.Status.PodIP) == 0 {
450             klog.V(5).Infof("Failed to find an IP for pod %s/%s", pod.Namespace, pod.Name)
451             continue
452         }
453         if !tolerateUnreadyEndpoints && pod.DeletionTimestamp != nil {
454             klog.V(5).Infof("Pod is being deleted %s/%s", pod.Namespace, pod.Name)
455             continue
456         }
457     }
```

pkg/controller/endpoint/endpoints\_controller.go (v1.22.0)

# The bug ...

```
622 func addEndpointSubset(subsets []v1.EndpointSubset, pod *v1.Pod, epa v1.EndpointAddress,  
623     epp *v1.EndpointPort, tolerateUnreadyEndpoints bool) ([]v1.EndpointSubset, int, int) {  
624     var readyEps int  
625     var notReadyEps int  
626     ports := []v1.EndpointPort{}  
627     if epp != nil {  
628         ports = append(ports, *epp)  
629     }  
630     if tolerateUnreadyEndpoints || podutil.IsPodReady(pod) {  
631         subsets = append(subsets, v1.EndpointSubset{  
632             Addresses: []v1.EndpointAddress{epa},  
633             Ports:     ports,  
634         })  
635         readyEps++  
636     } else if shouldPodBeInEndpoints(pod) {  
637         klog.V(5).Infof("Pod is out of service: %s/%s", pod.Namespace, pod.Name)  
638         subsets = append(subsets, v1.EndpointSubset{  
639             NotReadyAddresses: []v1.EndpointAddress{epa},  
640             Ports:             ports,  
641         })  
642         notReadyEps++  
643     }  
644     return subsets, readyEps, notReadyEps  
645 }
```

pkg/controller/endpoint/endpoints\_controller.go (v1.22.0)

# The bug ...

```
622 func addEndpointSubset(subsets []v1.EndpointSubset, pod *v1.Pod, epa v1.EndpointAddress,
623     epp *v1.EndpointPort, tolerateUnreadyEndpoints bool) ([]v1.EndpointSubset, int, int) {
624     var readyEps int
625     var notReadyEps int
626     ports := []v1.EndpointPort{}
627     if epp != nil {
628         ports = append(ports, *epp)
629     }
630     if tolerateUnreadyEndpoints || podutil.IsPodReady(pod) {
631         subsets = append(subsets, v1.EndpointSubset{
632             Addresses: []v1.EndpointAddress{epa},
633             Ports:     ports,
634         })
635         readyEps++
636     } else if shouldPodBeInEndpoints(pod) {
637         klog.V(5).Infof("Pod is out of service: %s/%s", pod.Namespace, pod.Name)
638         subsets = append(subsets, v1.EndpointSubset{
639             NotReadyAddresses: []v1.EndpointAddress{epa},
640             Ports:             ports,
641         })
642         notReadyEps++
643     }
644     return subsets, readyEps, notReadyEps
645 }
```

pkg/controller/endpoint/endpoints\_controller.go (v1.22.0)

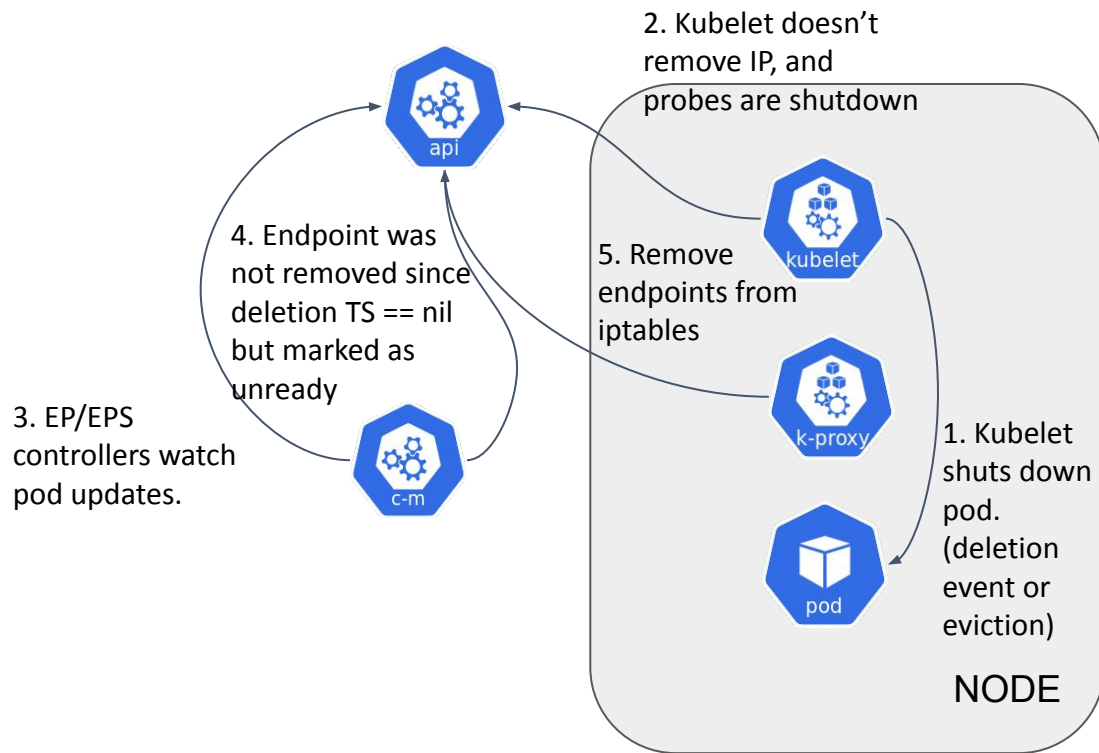
# The bug ...

```
622 func addEndpointSubset(subsets []v1.EndpointSubset, pod *v1.Pod, epa v1.EndpointAddress,
623     epp *v1.EndpointPort, tolerateUnreadyEndpoints bool) ([]v1.EndpointSubset, int, int) {
624     var readyEps int
625     var notReadyEps int
626     ports := []v1.EndpointPort{}
627     if epp != nil {
628         ports = append(ports, *epp)
629     }
630     if tolerateUnreadyEndpoints || podutil.IsPodReady(pod) {
631         subsets = append(subsets, v1.EndpointSubset{
632             Addresses: []v1.EndpointAddress{epa},
633             Ports:     ports,
634         })
635         readyEps++
636     } else if shouldPodBeInEndpoints(pod) {
637         klog.V(5).Infof("Pod is out of service: %s/%s", pod.Namespace, pod.Name)
638         subsets = append(subsets, v1.EndpointSubset{
639             NotReadyAddresses: []v1.EndpointAddress{epa},
640             Ports:             ports,
641         })
642         notReadyEps++
643     }
644     return subsets, readyEps, notReadyEps
645 }

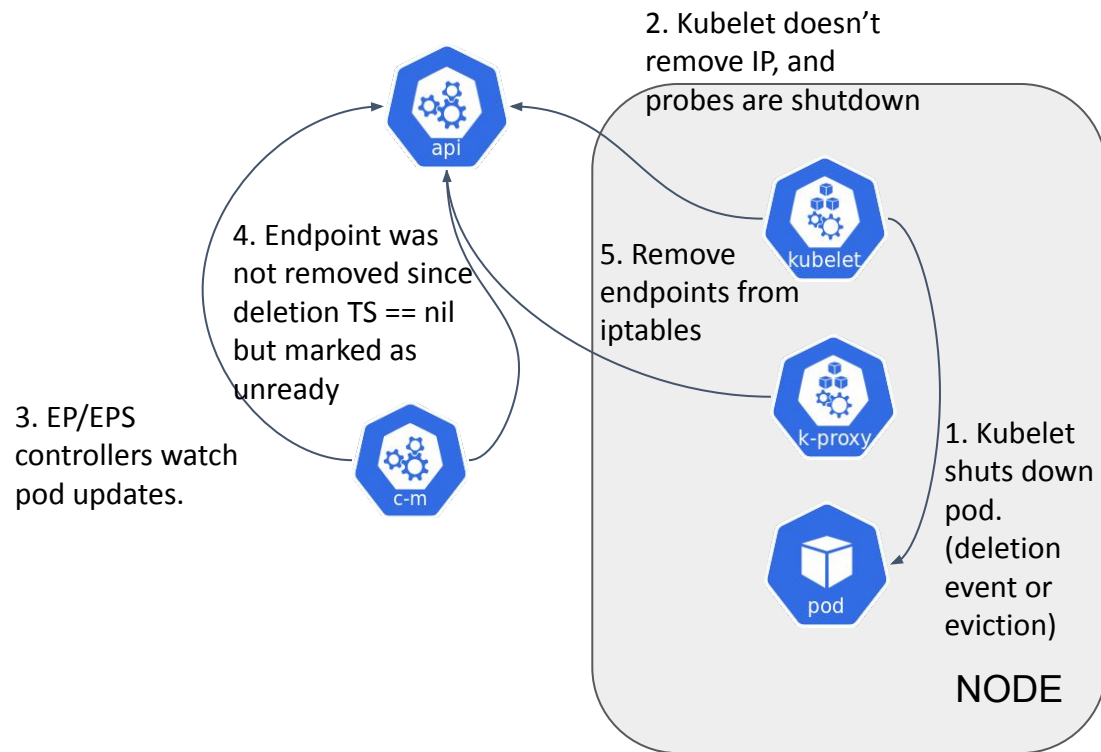
647 func shouldPodBeInEndpoints(pod *v1.Pod) bool {
648     switch pod.Spec.RestartPolicy {
649     case v1.RestartPolicyNever:
650         return pod.Status.Phase != v1.PodFailed && pod.Status.Phase != v1.PodSucceeded
651     case v1.RestartPolicyOnFailure:
652         return pod.Status.Phase != v1.PodSucceeded
653     default:
654         return true
655     }
656 }
```



# The bug ...

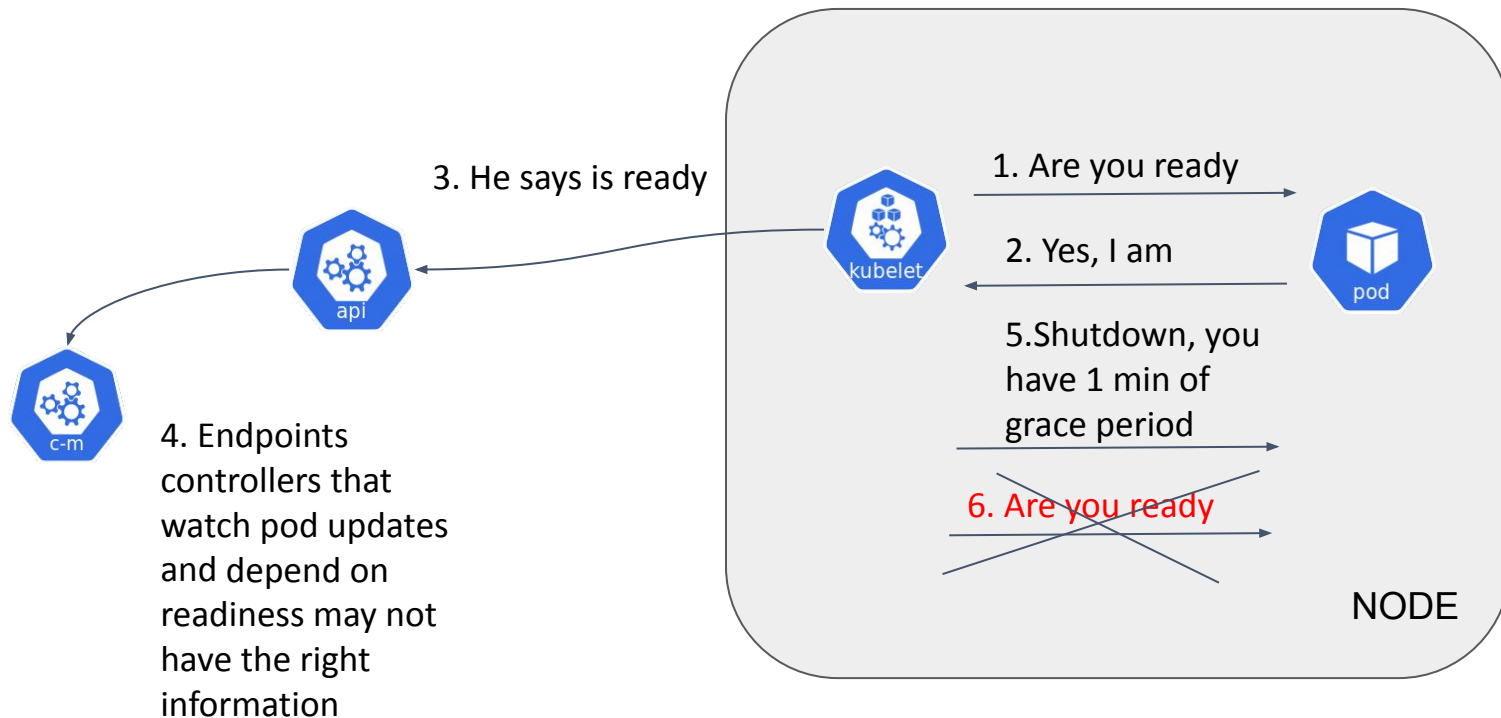


# ... and its consequences

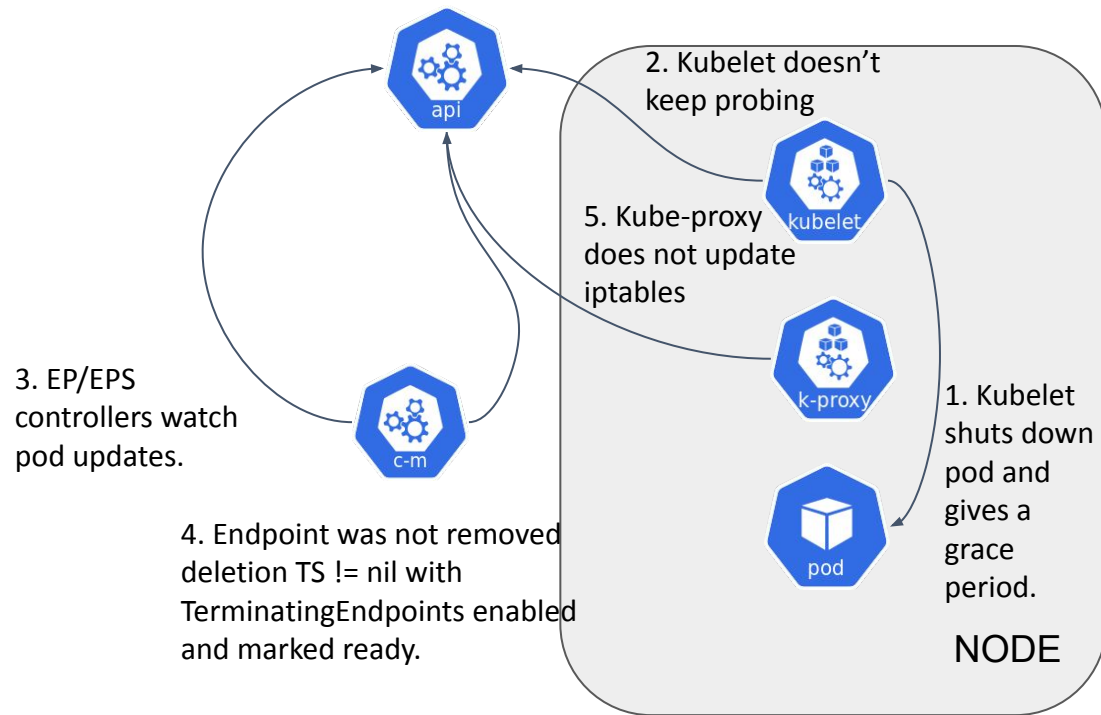


Endpoints and EndpointSlice resources have an endpoint that doesn't point to a current pod but marked as unready, so kube-proxy will removed from iptables.

# The other bug ... TerminatingEndpoints



# ... and its consequences



Traffic could be routed pods that are no longer ready. If pod readiness is not checked during the termination grace period, the pod keeps receiving traffic.

# The bugfix

## kubelet: only shutdown probes for pods that are terminated

This fixes a bug where terminating pods would not run their readiness probes. Terminating pods are found within the possiblyRunningPods map.

[< Prev](#) [Next >](#)

master (#110191)  
v1.26.0-alpha.2 v1.25.0-alpha.1

rphillips committed on Jun 6

commit f25ca15e1c47ebd8036c1fb5c06fe8fe6714807a

4 pkg/kubelet/kubelet\_pods.go

```
@@ -1105,8 +1105,8 @@ func (kl *Kubelet) HandlePodCleanups() error {  
1105 1105     }  
1106 1106  
1107 1107     // Stop probing pods that are not running  
1108 -     klog.V(3).InfoS("Clean up probes for terminating and terminated pods")  
1109 -     kl.probeManager.CleanupPods(runningPods)  
1108 +     klog.V(3).InfoS("Clean up probes for terminated pods")  
1109 +     kl.probeManager.CleanupPods(possiblyRunningPods)
```

# The lessons learned

- All behavior changes have potential to break others
- Networking and Pod/Node lifecycle are tightly coupled
- Behavior Contracts need to be enforced with testing
- “Broken telephone” bugs are hard to solve
- Bugs can become a feature



# Thank You!

Thank you!

