



DETROIT 2022

Windows Operational Readiness

Amim Knabben (VMware)

Xinqi Li (VMware)



Speaker introductions

BUILDING FOR THE ROAD AHEAD

DETROIT 2022



BUILDING FOR THE ROAD AHEAD

DETROIT 2022

October 24-28, 2022



Amim Knabben
Sr. Software Engineer
VMware

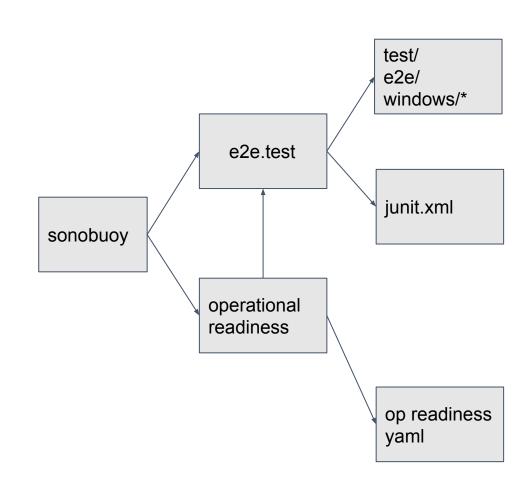


Xinqi Li Software Engineer VMware

Agenda



- Introduction
- KEP implementation
- The sonobuoy plugin
- Project usage
- Wrap up





BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Introduction

Conformance tests on Linux



 A suite of tests is defined as Linux conformance test in Kubernetes e2e with [Conformance] tag.

```
- testname: DNS, cluster

codename: '[sig-network] DNS should provide /etc/hosts entries for the cluster [Conformance]'

description: When a Pod is created, the pod MUST be able to resolve cluster dns

entries such as kubernetes.default via /etc/hosts.

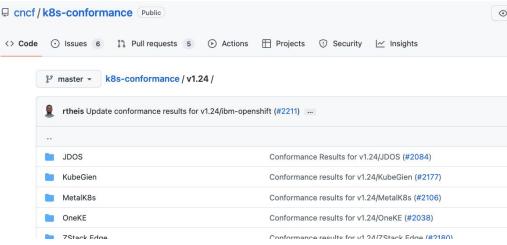
release: v1.14

file: test/e2e/network/dns.go
```

Each provider is encouraged to submit the conformance test result to prove the

cluster behave as expected.

conformance.yaml



No Conformance tests on Windows



- Windows is not required by official conformance testing.
- Is still important to verify that the behavior of the Windows nodes match the ones tested in the conformance suite as much as possible.
- It raises a question, for the enterprise what is considered "as much as possible"?

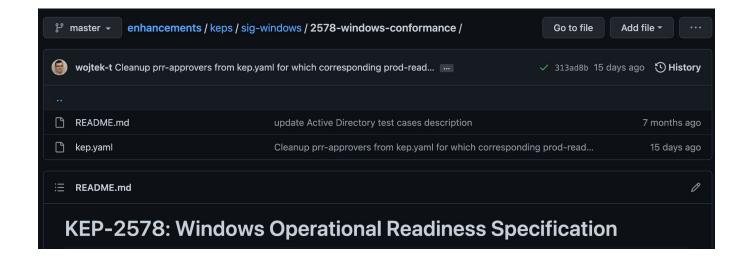
https://github.com/kubernetes/community/blob/master/contributors/devel/sig-architecture/conformance-tests.md#windows--linux-considerations

History of Windows Testing









KEP-2578



- NOT a CNCF official "conformance" standard.
- SIG-Windows definition of the MINIMUM Windows end-to-end TESTS required to pass in order to be considered operational.
- Architecture and definition for a testing tool implementation.

https://github.com/kubernetes/enhancements/tree/master/keps/sig-windows/2578-windows-conformance#kep-2578-windows-operational-readiness-specification

KEP-2578 – KEP categories



Required

- Basic networking (Core.Network)
- Basic storage (Core.Storage)
- Basic scheduling (Core. Scheduling)
- Basic concurrent functionality (Core.Concurrent)

Optional

- Windows HostProcess Operational Readiness (Extend. HostProcess)
- Active Directory (Extend.ActiveDirectory)
- Network Policy (Extend.NetworkPolicy)
- Windows Advanced Networking / Service Proxy (Extend.ServiceProxy)
- Windows worker configuration

KEP-2578 – KEP categories Example



Basic Networking

- Ability to access Linux container IPs by service IP from Windows containers
- Ability to access Windows container IPs by service IP from Linux containers
- Ability to access Linux container IPs by NodePort IP from Windows containers
- Ability to access Windows container IPs by NodePort IP from Linux containers
- Ability to access Linux container IPs by pod IP from Windows containers
- Ability to access Windows container IPs by pod IP from Linux containers
- Ability to schedule multiple containers, with distinct IP addresses, on the same node.
- Ability to delete and recreate services in such a way that load balancing rules for pods are recovered by whatever chosen service proxy is being utilized.
- Ability to delete and recreate pods for StatefulSets which preserve their ability to serve as routed endpoints for services, while also having unchanging IP addresses.
- Ability to access internal pods by internal Kubernetes DNS service endpoints.
- Ability to access external services by Kubernetes DNS services endpoints (for non-airgapped clusters).
- Ability to access internal pods through custom DNS records, injected by the Kubernetes pod Specification.
- Ability to route services from pods from the EndpointSlice API.

Linux Conformance VS. Windows Op-Readiness



	Linux Conformance	Windows Operational Readiness
Who defined?	CNCF	KEP-2578
Where are the test cases live?	Upstream Kubernetes repo	Upstream Kubernetes repo
How to run it?	Sonobuoy	Windows Operational Readiness tool



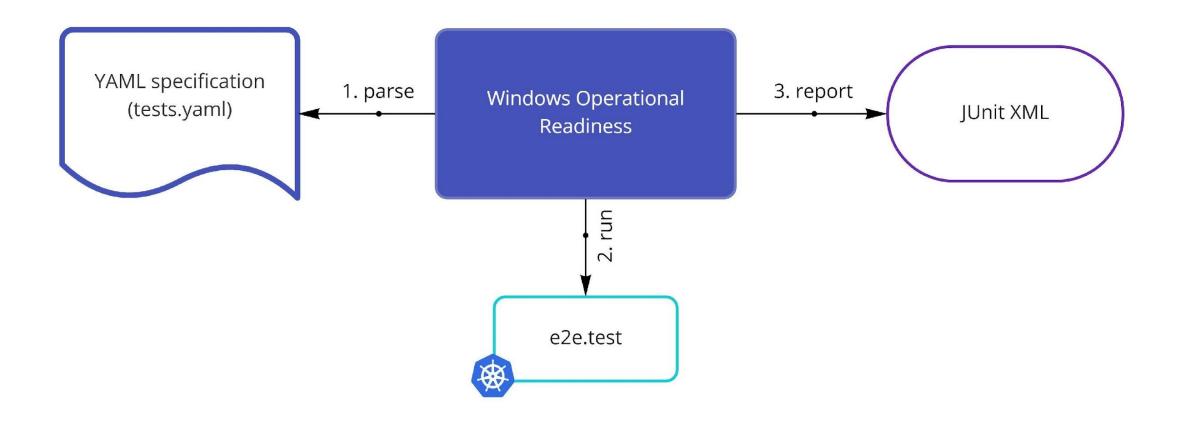
BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Implementation

High level architecture





YAML definition



- Category: Represents the test group.
- Description: A verbose definition of the test.
- Focus: Ginkgo tests to focus in.
- Skip: Ginkgo tests to be skipped.
- Kubernetes Version: the kubernetes version to run this test suite.

```
kubernetesVersions:
    -1.25
    -1.24
    -1.23
 5
    testCases:
      # Network category
      - category: Core.Network
9
        description: Ability to access Windows container IP by pod IP
10
        focus:
        - 'should have stable networking for Linux and Windows pods'
11
        skip:
12
        _ ''
13
      category: Core.Network
14
15
        description: Ability to expose windows pods by creating the service ClusterIP
        focus:
        - 'should be able to up and down services'
18
        skip:
        _ 11
19
      - category: Core.Network
20
        description: Ability to expose windows pods by creating the service NodePort
21
22
        focus:
        - 'should be able to create a functioning NodePort service for Windows'
24
        skip:
        _ !!!
25
26
      - category: Core-Network
```

Choosing one or more categories

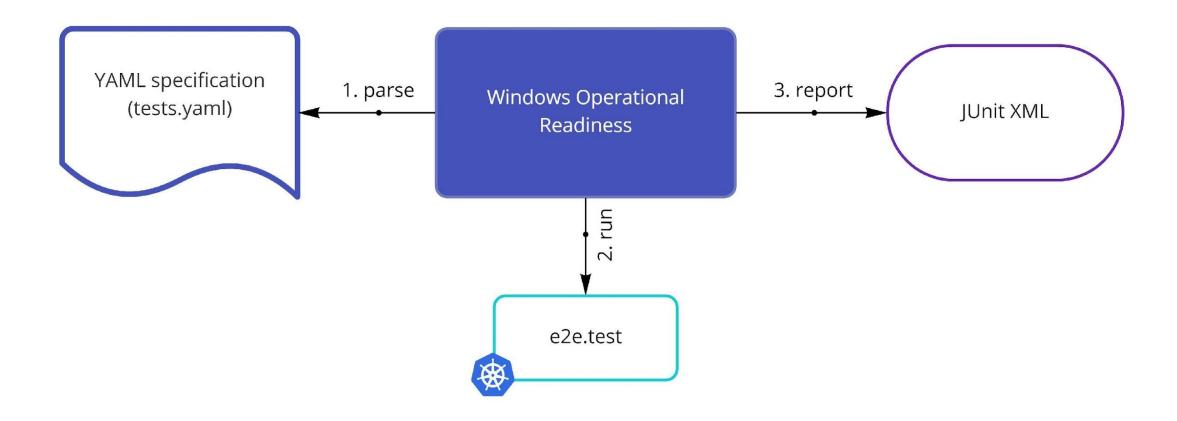


./op-readiness --provider=local --kubeconfig=<path-to-kubeconfig> --category=Core.Network --category=Extend.ActiveDirectory

```
+ > ./op-readiness --category Core.Network --category Extend.ActiveDirectory --provider=local --kubeconfig=${HOME}/.kube/config
2022-10-01T16:51:30-03:00
                                        [Core.Network] 1 / 17 - Running Operational Readiness Test: Ability to access Windows container IP by pod IP
                                INFO
                                        {"msg":"Test Suite starting", "total":1, "completed":0, "skipped":0, "failed":0}
2022-10-01T16:51:31-03:00
                                INFO
                                        Running Suite: Kubernetes e2e suite
2022-10-01T16:51:31-03:00
                                INFO
                                INFO
2022-10-01T16:51:31-03:00
2022-10-01T16:51:31-03:00
                                INFO
                                        Random Seed: 1664653890 - Will randomize all specs
2022-10-01T16:51:31-03:00
                                INFO
                                        Will run 1 of 6965 specs
2022-10-01T16:51:31-03:00
                                INFO
                                        Oct 1 16:51:31.798: INFO: >>> kubeConfig: /home/aknabben/.kube/config
                                INFO
2022-10-01T16:51:31-03:00
                                        Oct 1 16:51:31.811: INFO: Waiting up to 30m0s for all (but 0) nodes to be schedulable
2022-10-01T16:51:31-03:00
                                INFO
                                        Oct 1 16:51:31.839: INFO: Waiting up to 10m0s for all pods (need at least 0) in namespace 'kube-system' to be running and ready
2022-10-01T16:51:31-03:00
                                INFO
                                        Oct 1 16:51:31.863: INFO: 8 / 8 pods in namespace 'kube-system' are running and ready (0 seconds elapsed)
2022-10-01T16:51:31-03:00
                                INFO
                                        Oct 1 16:51:31.863: INFO: expected 2 pod replicas in namespace 'kube-system', 2 are Running and Ready.
2022-10-01T16:51:31-03:00
                                INFO
2022-10-01T16:51:31-03:00
                                INFO
                                        Oct 1 16:51:31.863: INFO: Waiting up to 5m0s for all daemonsets in namespace 'kube-system' to start
                                        Oct 1 16:51:31.867: INFO: 1 / 1 pods ready in namespace 'kube-system' in daemonset 'kindnet' (0 seconds elapsed)
2022-10-01T16:51:31-03:00
                                INFO
                                        Oct 1 16:51:31.867: INFO: 1 / 1 pods ready in namespace 'kube-system' in daemonset 'kube-proxy' (0 seconds elapsed)
2022-10-01T16:51:31-03:00
                                INFO
                                        Oct 1 16:51:31.867: INFO: e2e test version: v1.24.0
2022-10-01T16:51:31-03:00
                                INFO
2022-10-01T16:51:31-03:00
                                        Oct 1 16:51:31.869: INFO: kube-apiserver version: v1.24.0
                                INFO
                                        Oct 1 16:51:31.869: INFO: >>> kubeConfig: /home/aknabben/.kube/config
2022-10-01T16:51:31-03:00
                                INFO
2022-10-01T16:51:31-03:00
                                INFO
                                        Oct 1 16:51:31.875: INFO: Cluster IP family: ipv4
2022-10-01T16:51:31-03:00
                                INFO
2022-10-01T16:51:31-03:00
                                INFO
                                        [sig-windows] Hybrid cluster network for all supported CNIs
2022-10-01T16:51:31-03:00
                                INFO
                                          should have stable networking for Linux and Windows pods
2022-10-01T16:51:31-03:00
                                INFO
2022-10-01T16:51:31-03:00
                                INFO
                                          test/e2e/windows/hybrid_network.go:55
                                        [BeforeEach] [sig-windows] Hybrid cluster network
2022-10-01T16:51:31-03:00
                                INFO
2022 10 01T16.51.21 02.00
                                          test/e2e/windows/framework go:28
```

High level architecture





Calling upstream tests



- The ./e2e.test binary is called internally.
- Windows parameters are set by default.
- Filters and skip are applied per test.

```
cmd := exec.Command(testCtx.E2EBinary, args...)
stdout, err := cmd.StdoutPipe()
if err != nil {
        return err
}
stderr, err := cmd.StderrPipe()
if err != nil {
        return err
}
```

```
// RunTest runs the binary set in the test context with the parameters from flags.
     func (o *OpTestCase) RunTest(testCtx *TestContext, idx int) error {
             args := []string{
44
                     "--provider", testCtx.Provider,
45
                     "--kubeconfig", testCtx.KubeConfig,
46
                     "--report-dir", testCtx.ReportDir,
47
                     "--report-prefix", strconv.Itoa(idx),
                     "--node-os-distro", "windows",
49
                     "--non-blocking-taints", "os,node-role.kubernetes.io/master,node-role
50
                     "--ginkgo.flakeAttempts", "1",
51
52
53
54
             if testCtx.DryRun {
                     args = append(args, "--ginkgo.dryRun")
55
56
57
```

Report generated by JUnit



kubernetes-jenkins

/kubernetes-jenkins/pr-logs/pull/kubernetes-sigs_windows-operational-readiness/

- If **-report-dir** or **ARTIFACTS** env var set.
- Junit reports will be saved and dumped.
- Can be parsed by dashboard.

Name	Size
<u></u>	- .
<u>clusters/</u>	-
<u>services-6857/</u>	-
<u>junit_1001.xml</u>	533
<u>junit_101.xml</u>	538
<u>junit_201.xml</u>	2924825
<u>junit_301.xml</u>	512
<u>junit_401.xml</u>	572
<u>junit_501.xml</u>	299
<u>junit_601.xml</u>	299
<u>junit_701.xml</u>	489
junit 801.xml	498
junit_901.xml	486

2 Ways to run Windows Op-readiness tests



- 1. Run the "windows operational readiness" binary directly
 - \$ go run op-readiness



- 2. Run as a Sonobuoy plugin
 - \$ sonobuoy run

https://github.com/kubernetes-sigs/windows-operational-readiness



BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Sonobuoy plugin

Sonobuoy plugin





- Advantages of running as a Pod having in-cluster access.
- Tooling and support around report extracting, parsing and publishing.
- Sonobuoy is the STANDARD tool for running and submitting Conformance tests.

Running Operational Readiness via Sonobuoy



- Define the op readiness configuration (sonobuoy plugin):
 - Core.Network
 - Extend.NetworkPolicy
- Change binary parameters if needed (spec field)
- A new image is published on master commit, get the

latest version at

gcr.io/k8s-staging-win-op-rdnss/k8s-win-op-rdnss:latest

```
cat sonobuoy-plugin.yaml
sonobuoy-config:
 driver: Job
 plugin-name: op-readiness
 result-format: manual
spec:
 command:
 - /app/run.sh
 args:
 - --e2e-binary
 - /app/e2e.test
 --category
 - Core.Network
 --category
 - Extend.NetworkPolicy
 image: gcr.io/k8s-staging-win-op-rdnss/k8s-win-op-rdnss:latest
 name: plugin
 volumeMounts:
 - mountPath: /tmp/sonobuoy/results
   name: results
```

Running the plugin



```
> make sonobuoy-plugin
sonobuoy delete
INFO[0000] already deleted
                                                         kind=namespace namespace=sonobuoy
INFO[0000] delete request issued
                                                         kind=clusterrolebindings
INFO[0000] delete request issued
                                                         kind=clusterroles
sonobuoy run --sonobuoy-image projects.registry.vmware.com/sonobuoy/sonobuoy:v0.56.9 --plugin sonobuoy-plugin.yaml --wait=0
INFO[0000] create request issued
                                                         name=sonobuoy namespace= resource=namespaces
INFO[0000] create request issued
                                                         name=sonobuoy-serviceaccount namespace=sonobuoy resource=serviceaccounts
INFO[0000] create request issued
                                                         name=sonobuoy-serviceaccount-sonobuoy namespace= resource=clusterrolebindings
                                                         name=sonobuoy-serviceaccount-sonobuoy namespace= resource=clusterroles
INFO[0000] create request issued
INFO[0000] create request issued
                                                         name=sonobuoy-config-cm namespace=sonobuoy resource=configmaps
                                                         name=sonobuoy-plugins-cm namespace=sonobuoy resource=configmaps
INFO[0000] create request issued
INFO[0000] create request issued
                                                         name=sonobuoy namespace=sonobuoy resource=pods
INFO[0000] create request issued
                                                         name=sonobuoy-aggregator namespace=sonobuoy resource=services
```

- Use make sonobuoy-plugin to run in-cluster.
- Wait and run make sonobuoy-results will summarize the failed tests.

```
windows-operational-readiness on | main [$?] via | v1.19 took 6m  
+ ) make sonobuoy-results  
sonobuoy results --mode=report 202210091253_sonobuoy_06e8484b-6411-407d-9365-373d9821b74b.tar.gz  
Plugin: op-readiness  
Status: failed  
Total: 6965  
Passed: 0  
Failed: 1  
Skipped: 6964  

Failed tests: [sig-network] Netpol NetworkPolicy between server and client should deny ingress from pods on other namespaces [Feature:NetworkPolicy]  
Run Details:  
API Server version: v1.24.0  
Node health: 1/1 (100%)  
Pods health: 11/20 (55%)
```



BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Project usage

Windows dev tools



- The easiest way to learn internals of Kubernetes on Windows locally, no cloud required !!!
- Run a Kubernetes Windows cluster locally.
- Use vagrant to create a Hybrid OS cluster.
- All is required is a robust PC and Virtualbox
- Easy to start with make.
- Easy to run the windows operational readiness.

https://sigs.k8s.io/sig-windows-dev-tools

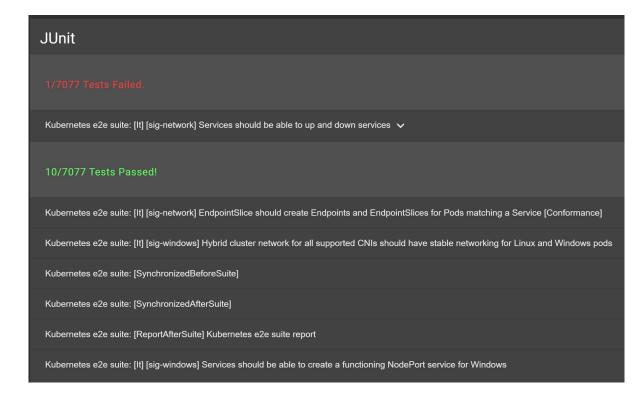


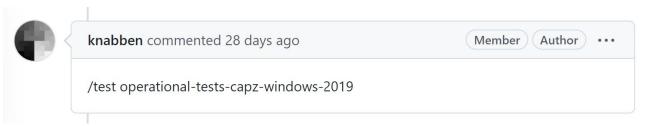


Prow jobs



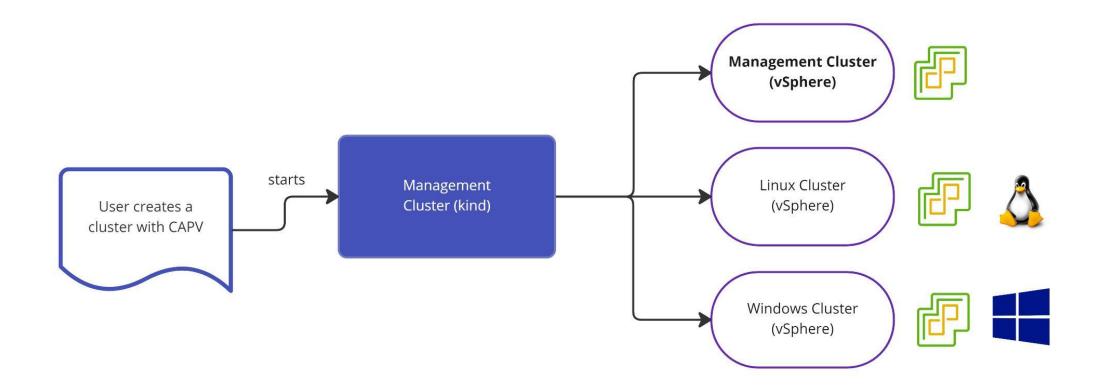
- Added a new test and want to test upstream?
- Run it on ClusterAPI on Azure.
- The project has a presubmit job.
- The job is NOT enabled by default.
- Run manually via Prow bot.





ClusterAPI (vSphere)

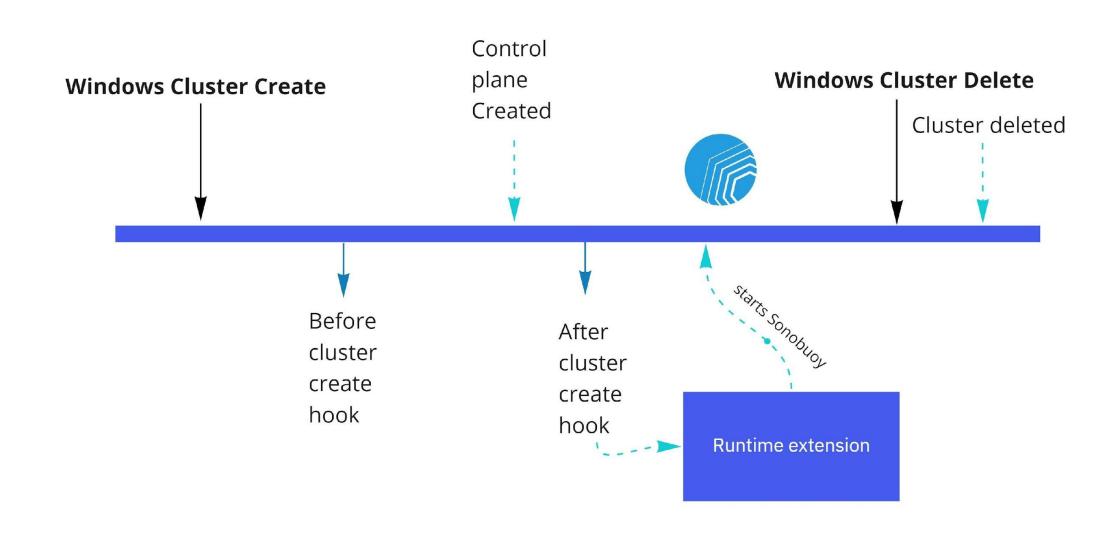




 Kubernetes project to bring declarative, Kubernetes-style APIs to cluster creation, configuration, and management.

ClusterAPI - Runtime Extensions (exp)







BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Wrap up

Special thanks



- SIG-Windows + SIG-Testing + SIG-Network
 - Jay Vyas (VMware) for all mentoring and leadership.
 - Mark Rossetti (Microsoft) and James Sturtevant (Microsoft) for mentoring and guidance on Windows + Kubernetes.
 - Aravindh Puthiyaparambil (RedHat), Claudiu Belu (Cloudbase) and Luther Monson (Spectro Cloud), Douglas Landgraf (RedHat) for working on new Windows features and always ready to help the community.
 - Adelina and Ben the Elder for early discussions on this topic !!!
 - Antonio Ojea for helping with networking e2e tests.

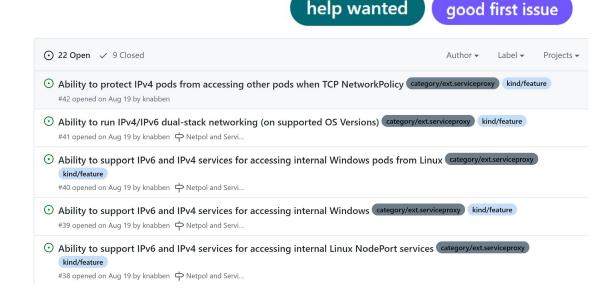
How To Contribute



- Right now ~40% of the tests are implemented
- There are ~20 issues open for the missing tests
- First comers are welcome.
- We are happy to mentor and help onboarding!
- https://sigs.k8s.io/windows-operational-readiness
- Join us on #sig-windows at Kubernetes Slack!

Interested in learning K8s on windows? Try out

https://github.com/kubernetes-sigs/sig-windows-dev-tools





BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Thank you

Amim Knabben (VMware)

Xinqi Li (VMware)



Please scan the QR Code above to leave feedback on this session