



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Tutorial: How To Write a Reconciler Using K8s Controller-Runtime!

*Scott Rigby, Somtochi Onyekwere,
Niki Manoledaki & Soulé Ba, Weaveworks;
Amine Hilaly, Amazon Web Services*

Tutorial: How To Write a Reconciler Using K8s Controller-Runtime!



KubeCon



CloudNativeCon

North America 2022


BUILDING FOR THE ROAD AHEAD

DETROIT 2022

October 24-28, 2021



Scott Rigby
Developer Experience
Engineer
Weaveworks



Somtochi Onyekwere
Developer Experience
Engineer
Weaveworks



Niki Manoledaki
Software Engineer
Weaveworks



Soulé Ba
Consulting Reliability
Engineer
Weaveworks



Amine Hilaly
Software Development
Engineer
Amazon Web Services

Agenda

- Demo - 5m
- Local dev setup - 10m
- Challenge - 5m
- Intro to concepts - 10m
- Step-by-step write the controller! - 40m
- Wrap up + Q&A - 10m

Challenge

The challenge

Goal

As a hypothetical KubeCon speaker
In order to have fun and learn things
I want to be able to manage my KubeCon proposal submissions declaratively

We have a CFP API

- Code and docs [here](#)
- It supports Speaker and Proposal objects, with a one to many relationship

Write a K8s controller that

- Defines CRDs for speaker and proposal
- Reconciles CFP API submissions with our declarative CRDs

Local dev setup

Repository

Please fork & clone this repository:

<https://bit.ly/3TFnryi>



**github.com/scottrigby/
how-to-write-a-reconciler-using-k8s-controller-runtime**

Prerequisites


Repository: <https://bit.ly/3TFnryi>

- Just watch
- GitPod (preferred)
 - Open the link in the repository



- Please run it now
- Vagrantfile
 - Follow instructions in Readme
- DIY everything!
 - Requirements:
 - Go
 - Kind
 - Docker
 - Kustomize

**JUST
WATCH**



GITPOD



VAGRANTFILE



**DIY
EVERYTHING!**



Demo

Demo time!



Intro to concepts

- Reconciliation
- K8s controller-runtime
- Kubebuilder
- Conditions
- Status & observedGeneration

Reconciliation

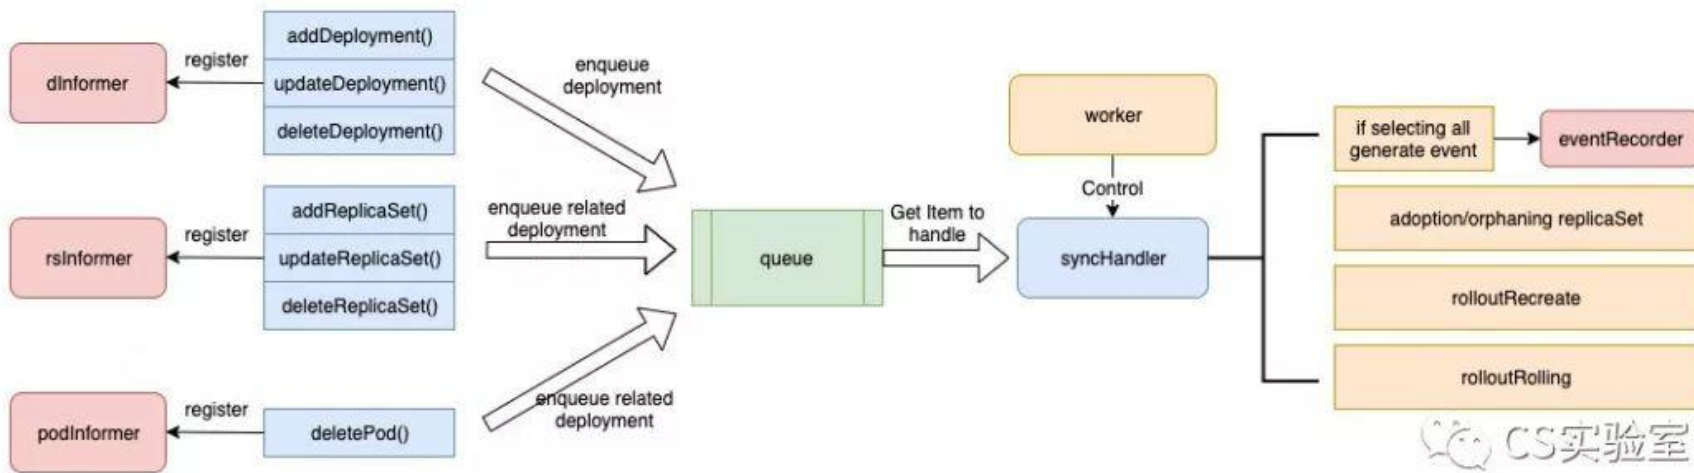
```
for {  
    desired := getDesiredState()  
    current := getCurrentState()  
    makeChanges(desired, current)  
}
```

source: kubernetes/community repo
developer guide on [“Writing Controllers”](#)



K8s controller-runtime

Native controllers like Deployment and ReplicaSet were built using core components like Informers, Listers, work queues etc...



K8s controller-runtime

- Controller-runtime is a set of go libraries for building controllers.
- Product of years of experience building controllers.
- Hides the complexity of building the reconcilers
- Allows you to focus on the logic of the sync function
- Well tested and actively maintained

K8s controller-runtime

- Standardize logging `pkg/log`
- Leader election `pkg/leaderelection`
- Webhooks `pkg/webhooks`
- Ratelimiting `pkg/ratelimiter`
- Better and simpler testing `pkg/envtest`

Kubebuilder is framework for building operators following Kubernetes API best practices

- It builds on controller-runtime and controller-tools
- Also provides code -generation with markers

Commands to scaffold

```

```
kubebuilder init --domain my.domain --repo my.domain/guestbook
kubebuilder create api --group webapp --version v1 --kind Guestbook
```

```

Other frameworks include operator-framework

Kubebuilder book: <https://book.kubebuilder.io/>

- The conditions is a set of types (Ready, PodScheduled...) with a status (True, False or Unknown) that make up the ‘computed state’ of a Resource.
- The set of conditions describes the ‘current’ state.
- Conditions are just a way of communicating changes of state between components, and this state can always be reconstructed by observing the system.
- Condition types should indicate state in the “abnormal-true” polarity.

<https://github.com/kubernetes/community/blob/4c9ef2d/contributors/devel/sig-architecture/api-conventions.md>
<https://github.com/kubernetes-sigs/cli-utils/blob/master/pkg/kstatus/README.md>

Status & observedGeneration

- A /status subresource MUST be provided to enable system components to update statuses of resources they manage.
- If the generation and the observedGeneration of a resource does not match, it means there are changes that the controller has not yet seen, and therefore not acted upon.

<https://github.com/kubernetes/community/blob/4c9ef2d/contributors/devel/sig-architecture/api-conventions.md#spec-and-status>



Step-by-step write the controller!

Step 1

Run

```
$ git checkout tags/s1 -b s1-branch
```

run `make test` to run the test cases

Step 2

Run

```
$ git checkout tags/s2 -b s2-branch
```

run `make test` to run the test cases

Step 3

Run

```
$ git checkout tags/s3 -b s3-branch
```

run `make test` to run the test cases

Step 4

Run

```
$ git checkout tags/s4 -b s4-branch
```

run `make test` to run the test cases

Step 5

Run

```
$ git checkout tags/s5 -b s5-branch
```

run `make test` to run the test cases

Step 6

Run

```
$ git checkout tags/s6 -b s6-branch
```

run `make test` to run the test cases

Step 7

Run

```
$ git checkout tags/s7 -b s7-branch
```


run `make test` to run the test cases

Wrap up + Q&A

What we learned

TO-DO: Add the most salient points here to remind folks they actually learned things, since their brain is probably full by now

- Cloud Native Glossary: glossary.cncf.io
- Kubernetes community developer guide:
<https://github.com/kubernetes/community/blob/master/contributors/devel/sig-api-machinery/controllers.md>
- Kubebuilder book: sigs.k8s.io/kubebuilder
- Kubernetes controller-runtime Project:
github.com/kubernetes-sigs/controller-runtime
- Flux pkg/runtime: pkg.go.dev/github.com/fluxcd/pkg/runtime
- Finalizers to Control Deletion:
kubernetes.io/blog/2021/05/14/using-finalizers-to-control-deletion
- Implementing observedGeneration:
alenkacz.medium.com/kubernetes-operator-best-practices-implementing-observedgeneration-250728868792



Q&A

Session QR Codes will be sent via
email before the event



Please scan the QR Code above
to leave feedback on this session