



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Using eBPF Superpowers to generate Kubernetes Security Policies

Mauricio Vásquez-Bernal & Alban Crequy

Agenda

- K8s security policies
- Generating security policies
- Auditing security policies
- Limitations

K8s Security Policies

- K8s offers different kind of policies to secure the clusters
 - Seccomp
 - Pod Capabilities
 - Network Policies
- Configuring those policies requires knowledge about the application
 - Which system calls should it perform?
 - Which Linux capabilities does it need?
 - Which network endpoints does it communicate with?
- Most of the times the application developer is not the one configuring those

Generating Security Policies

- Observe the application and use this data to generate policies
 1. Observe the application
 - Technologies like eBPF allow to capture the activity of an application
 - Low overhead
 - Great flexibility
 2. Generate a policy
 - Depending on the policy kind this is easy or very difficult

Seccomp

Secure Computing

- Linux kernel mechanism to limit the system calls a process can make
- A seccomp profile defines a list of syscalls and an action for them:
 - KILL: Kills the calling process
 - TRAP: Sends SIGSYS to the calling thread
 - ERRNO: Returns an error without executing the syscall
 - LOG: Logs and execute the syscall
- <https://man7.org/linux/man-pages/man2/seccomp.2.html>

Seccomp Profile

```
{
  "defaultAction": "SCMP_ACT_ERRNO",
  "syscalls": [
    {
      "names": [
        "futex",
        [...],
        "recvfrom",
      ],
      "action": "SCMP_ACT_ALLOW"
    }
  ]
}
```


Seccomp in Kubernetes

- Security profile defined in the pod security context (k8s >= 1.19)
 - Can be defined at pod or container level

spec:

securityContext:

seccompProfile:

type: Localhost

localhostProfile: myprofile.json



File under /var/lib/kubelet/seccomp on the node.

<https://kubernetes.io/docs/tutorials/security/seccomp/>

Security Profiles Operator (SPO)

- Tool to handle security profiles in Kubernetes
- Provides SeccompProfile Custom Resource
 - Saves profile on nodes to be referenced in pod spec

```
apiVersion: security-profiles-operator.x-  
k8s.io/v1beta1  
kind: SeccompProfile  
metadata:  
  namespace: ns  
  name: profile1  
spec:  
  defaultAction: SCMP_ACT_LOG
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: test-pod  
spec:  
  securityContext:  
    seccompProfile:  
      type: Localhost  
      localhostProfile: operator/ns/profile1.json  
  containers:  
    - name: test-container  
      image: nginx
```

Defining Seccomp Profiles

- An allow-list is preferred
 - More secure
 - Must include all system calls needed by the application
- How to know which syscalls are needed my application?
 - Analyze the code
 - Use strace
 - Define a profile with SCMP_ACT_LOG and check the audit log
 - **Use a tool that records the syscalls performed in a container**

Demo

Demo

Auditing Seccomp Policies

- How to know if a given security profile will break my application?
- Use "defaultAction: SCMP_ACT_LOG"
 - Check syslog
 - Problem: Difficult to understand
- Alternative: "seccomp audit" gadget

Demo

Linux Capabilities

Linux Capabilities

- Since Linux 2.2 privileges are divided into units
 - It's possible to give a process only some privileges
- Those privileges are called capabilities:
 - CAP_CHOWN: change file UIDs and GIDs
 - CAP_NET_BIND_SERVICE: bind ports < 1024
 - CAP_SYS_BOOT: call reboot()
 - Many more: <https://man7.org/linux/man-pages/man7/capabilities.7.html>

Linux Capabilities in Kubernetes

- The container runtime grants some capabilities by default
- Others can be added / dropped by using the container securityContext

```
securityContext:  
  capabilities:  
    drop:  
      - all  
    add: ["NET_ADMIN", "SYS_TIME"]
```

Demo

Network Policies

- Kubernetes mechanism to control how a pod can communicate with network “entities”
- Operate at IP or port level (Layer 3 or 4)
- Can be used for inbound (**Ingress**) or outbound (**Egress**) traffic
- <https://kubernetes.io/docs/concepts/services-networking/network-policies/>

The NetworkPolicy Resource

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - ipBlock:
            cidr: 172.17.0.0/16
      ports:
        - protocol: TCP
          port: 6379
  egress:
  ...
```

Pods that the policy
applies to

Layer 3 selectors are supported:

- podSelector
- namespaceSelector
- namespaceSelector and podSelector
- ipBlock

Layer 4 selectors

Defining Network Policies

- Network policies are usually created when defining the architecture of a solution
- In some cases, those are defined after the solution is deployed
- For those, using a policy advisor can be an option

Demo

Auditing Network Policies

- Some tools allows to check when a network policy blocks traffic
- Cilium supports `auditMode: true`
 - <https://docs.cilium.io/en/v1.8/gettingstarted/policy-creation/>
- Other tools can determine the iptables rule that blocked a packet
 - <https://github.com/box/kube-iptables-tailer>

- The application must generate all possible events while observing it
 - If an event is not generated, then it'll be blocked by the policy
- This technique assumes the application is trusted when observing
 - It's not compromised, and no malicious activity is generated while observing it

- [You and Your Security Profiles; Generating Security Policies with the Help of eBPF - John Fastabend & Natalia Reka Ivanko, Isovalent](#)

Thanks



Please scan the QR Code above to
leave feedback on this session