



KubeCon



CloudNativeCon

North America 2022

BUILDING FOR THE ROAD AHEAD

DETROIT 2022

What's New With SIG-Windows

What's New With SIG-Windows



BUILDING FOR THE ROAD AHEAD

DETROIT 2022

October 24-28, 2021



Mark Rossetti
Software Engineer
Microsoft



James Sturtevant
Software Engineer
Microsoft



Amim Knabben
Software Engineer
VMware



BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Agenda

- Updates
- Upcoming Work
- Additional Resources
- Q&A

Updates

Enhancements Updates

- OS field in pod-spec
 - Stable in 1.25
- HostProcess Containers
 - Stable in 1.26
- HostNetwork support for Windows containers
 - Targeting Alpha in v1.26 (Will require containerd vNext)
- NodeService log viewer
 - Targeting Alpha in v1.26
- CRI-based stats support for Windows

Pod OS Field

- Promoted to Stable in v1.25
- Currently NOT used for scheduling
- Allows for enhanced pod validation and error reporting
- No default values!

```
apiVersion: v1
kind: Pod
metadata:
  name: ping
spec:
  os:
    name: windows
  containers:
  - name: ping-test
    image: mcr.microsoft.com/windows/nanoserver:1809
    command:
    - "ping"
    - "-t"
    - "127.0.0.1"
  nodeSelector:
    "kubernetes.io/os": windows
```

Use when scheduling Windows where **restricted** pod security profiles are used!

Windows container accidentally scheduled to linux nodes

Old behavior

```
$ kubectl get pod
NAME    READY   STATUS    RESTARTS   AGE
ping    0/1     ErrImagePull    0          111s
```

```
Events:
  Type     Reason      Age    From      Message
  ----     -
  Normal   Scheduled   7s     default-scheduler   Successfully assigned default/ping to marosset-capz-2022-ep-md-0-94drk
  Normal   Pulling     6s     kubelet    Pulling image "mcr.microsoft.com/windows/nanoserver:1809"
  Warning   Failed      6s     kubelet    Failed to pull image "mcr.microsoft.com/windows/nanoserver:1809": rpc error: code = NotFound desc = failed to pull and unpack image "mcr.microsoft.com/windows/nanoserver:1809": failed to unpack image on snapshotter overlaysfs : no match for platform in manifest sha256:62a8d022600141cd93d7e74cb190de58c9ad273ca238424028af88ad46495ca7: not found
  Warning   Failed      6s     kubelet    Error: ErrImagePull
  Normal   BackOff     4s (x2 over 5s)   kubelet    Back-off pulling image "mcr.microsoft.com/windows/nanoserver:1809"
  Warning   Failed      4s (x2 over 5s)   kubelet    Error: ImagePullBackOff
```

New behavior

```
$ kubectl get pods
NAME    READY   STATUS             RESTARTS   AGE
ping    0/1     PodOSNotSupported    0          3s
```


```
Events:
  Type     Reason             Age    From      Message
  ----     -
  Warning   PodOSNotSupported  48s    kubelet    Failed to admit pod as the OS field doesn't match node OS
  Normal   Scheduled          48s    default-scheduler   Successfully assigned default/ping to marosset-capz-2022-ep-md-0-4b8tw
```

HostProcess Containers Updates

- Going stable in v1.26
- Requires containerd v1.7
- New volume-mount behavior
 - Volume mounts will show in expected locations (not under %CONTAINER_SANDBOX_MOUNT_POINT%)
 - InClusterConfig works!!
- New base image
 - <http://bit.ly/hpc-base-image>
- <https://bit.ly/k8s-hpc-kep>
- More information in <https://sched.co/182Fe>

Windows Node Exporter

- Runs as a DaemonSet
- Can use with Prometheus Operator
 - ServiceMonitor or PodMonitor
- Integrates with Grafana to create dashboards



Windows container networking stack inspector

- Captures Network Traffic across
 - Nodes
 - Pods
- Traces
 - Packets
 - Packets Received/Transmitted/Dropped
 - VFP Counters
 - HNS Counters
 - HNS state
- Uses PacketMon to capture packets
 - Convert to Wireshark format

```
t 2022-10-14 04:31:14.1079323 +0000 UTC:
    00-0D-3A-14-F2-30 > 00-15-5D-65-50-05, ethertype IPv4 (0x0800), length 60: 152
    [..], ack 2793264986, win 501, length 0
Response from aksnpwin000000 StartCapture (10.224.0.33) sent at 2022-10-14 04:31:14.10
04:31:11.985876900 PktGroupId 562949954527102, PktNumber 1, Appearance 1, Direction Tx
    Filter 0, OriginalSize 64, LoggedSize 64
Response from aksnpwin000000 StartCapture (10.224.0.33) sent at 2022-10-14 04:31:14.10
    00-15-5D-65-50-05 > 12-34-56-78-9A-BC, ethertype IPv4 (0x0800), length 64: 10.
    wpad. (22)
Response from aksnpwin000000 StartCapture (10.224.0.33) sent at 2022-10-14 04:31:14.10
04:31:11.993994900 PktGroupId 1303417, PktNumber 1, Appearance 1, Direction Rx , Type
    0, OriginalSize 139, LoggedSize 128
Response from aksnpwin000000 StartCapture (10.224.0.33) sent at 2022-10-14 04:31:14.10
    12-34-56-78-9A-BC > 00-15-5D-65-50-05, ethertype IPv4 (0x0800), length 139: 10.
    Domain 0/1/0 (97)
Response from aksnpwin000000 StartCapture (10.224.0.33) sent at 2022-10-14 04:31:14.10
04:31:11.994318700 PktGroupId 562949954527103, PktNumber 1, Appearance 1, Direction Tx
    Filter 0, OriginalSize 90, LoggedSize 90
Response from aksnpwin000000 StartCapture (10.224.0.33) sent at 2022-10-14 04:31:14.10
    00-15-5D-65-50-05 > 12-34-56-78-9A-BC, ethertype IPv4 (0x0800), length 90: 10.
    wpad.default.svc.cluster.local. (48)
StopCapture successfully ran on node: aksnpwin000000 (IP: 10.224.0.33) at time: 2022-1
Packet capture ended on node: aksnpwin000000 (IP: 10.224.0.33).
PS C:\hpc> .\wcnspect.exe counters
```

Performance and Scale


Kube-Proxy perf improvements

- 2000 Cluster IPS and 5000 endpoints
 - Found issue Syncing proxy rules made a lot of calls to HNS Subsystem
 - Caching locally and improvements in HNS
 - Available in K8s v1.22+
- Thanks to David Schott and Sravanth Bangari!

Winkernel proxier	Windows Server 2019	Windows Server 2022
Without improvements	114 minutes	192 minutes
With improvements	110 minutes	10 minutes
With improvements + OS updates (shipping separately)	22 minutes	1 minute

Perf stuff

Performance Dashboard



- Uses the Windows Exporter
- Runs Daily and measures
- Memory usage
- CPU
- OpenFile handles
- Network
- Integration With Perf-test Framework
- Next up:
- Larger scale tests
- Turning into soak

KEP 2258: Node Log Viewer

#sig-windows

Aravindh Puthiyaparambil
OpenShift Staff Engineer, Red Hat



Aravindh Puthiyaparambil

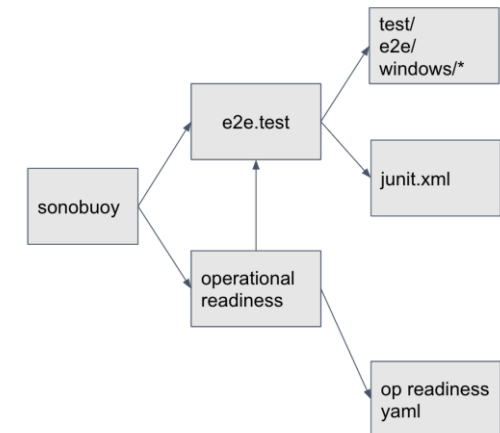


Windows Operational Readiness Spec

Make it easy for an enterprise to compare Windows K8S distributions via:

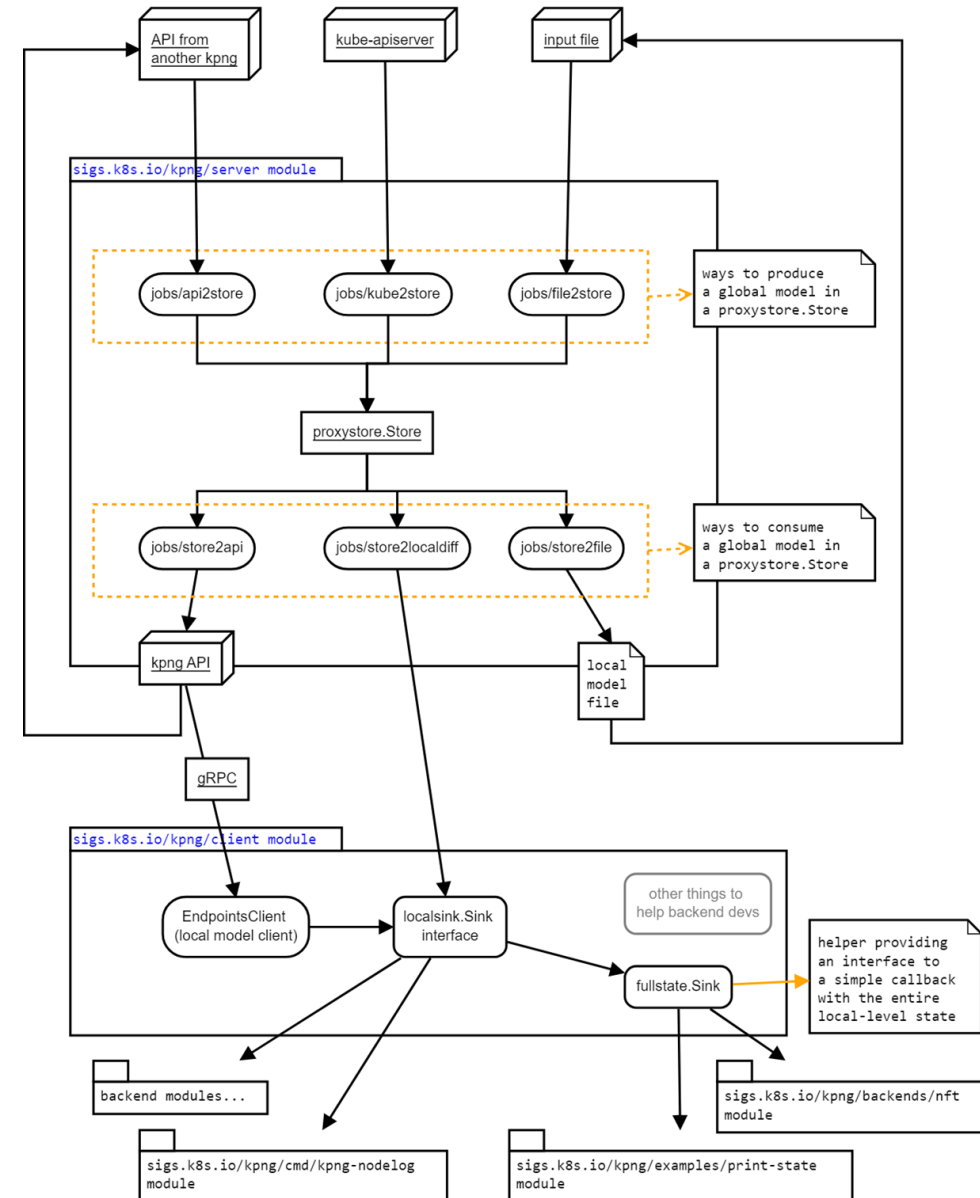
- Learning what **features** are supported on your Cluster.
- Discover the **healthiness** of your Windows cluster.
- Simplify the ginkgo E2E tests to run, without having to **select/grep** from 1000s of K8s e2e strings.

```
1  kubernetesVersions:
2  - 1.25
3  - 1.24
4  - 1.23
5
6  testCases:
7    # Network category
8    - category: Core.Network
9      description: Ability to access Windows container IP by pod IP
10     focus:
11       - 'should have stable networking for Linux and Windows pods'
12     linux_image: .nan
13     skip:
14       - ''
15     windows_image: .nan
16   - category: Core.Network
17     description: Ability to expose windows pods by creating the service ClusterIP
18     focus:
19       - 'should be able to up and down services'
20     linux_image: .nan
21     skip:
22       - ''
23     windows_image: .nan
24   - category: Core.Network
```



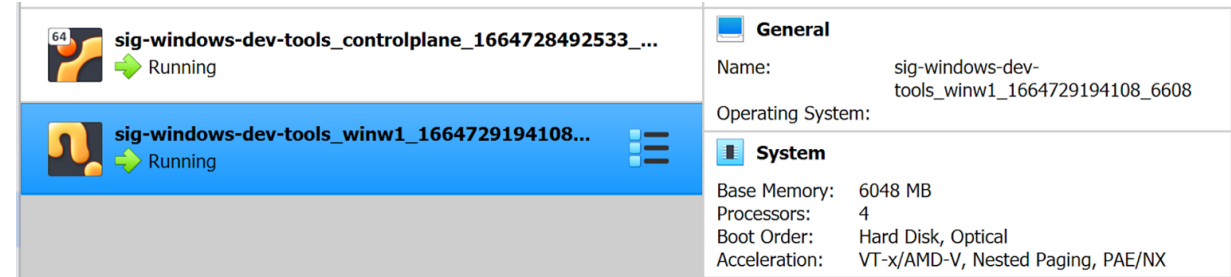
KPNG now supports Windows!

- The Next generation Kube Proxy
- Allows running Windows proxier as a separate process from the APIServer watch
- Windows Kernelspace and Userspace backends on KPNG
- Easy to develop new backends
- Collaborative effort
 - Mark Rossetti
 - Dmitri Mititelu
 - Jay Vyas
 - Douglas Landgraf
- Windows (and Linux) Userspace ARE removed from Kubernetes v1.26 in-tree.



SIG Windows Development Tools

- Easy onboarding for developers working with Windows clusters.
- What's new?:
 - Run on a Windows machine with **WSL2**.
 - Use latest **binaries** OR build Kubernetes from **source** using a hash.
 - Latest CNI plugins:
 - **Antrea** v1.8.0
 - **Calico** v3.23.2



General	
Name:	sig-windows-dev-tools_winw1_1664729194108_6608
Operating System:	

System	
Base Memory:	6048 MB
Processors:	4
Boot Order:	Hard Disk, Optical
Acceleration:	VT-x/AMD-V, Nested Paging, PAE/NX

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

PS C:\Users\vagrant> get-service -name kubelet

Status  Name          DisplayName
-----  ----          -
Running kubelet      kubelet service

PS C:\Users\vagrant> get-service -name kube-proxy

Status  Name          DisplayName
-----  ----          -
Running kube-proxy  kube-proxy service
```

What's Next

Hyper-V Isolated Containers

- Support for starting Hyper-V Isolated Containers added for containerd v1.7
- Tracked with <https://github.com/containerd/containerd/issues/6862>
- Uses runtime classes to tell containerd which mode to run containers in
- Uses Pod annotations to specify UVM memory/CPU sizes

```
apiVersion: v1
kind: Pod
metadata:
  name: wcow-test
  labels:
    app: wcow
  annotations:
    io.microsoft.virtualmachine.computetopology.memory.sizeinmb: "4096"
    io.microsoft.virtualmachine.computetopology.processor.count: "4"
spec:
  replicas: 2
  selector:
    matchLabels:
      app: wcow
  spec:
    runtimeClassName: runhcs-wcow-hypervisor <-----
    containers:
      - name: servercore
        image: mcr.microsoft.com/windows/servercore:1809
        ports:
          - containerPort: 80
            protocol: TCP
```

Additional Resources

Special Thanks and New Contributors

- Fabian Fulga – For adding adding test coverage to multiple components and documentation updates in sig-windows-tools docs
- David Schott – For perf improvements and various fixes in the winkernel proxy and building tools like WCNSpect to help trouble shoot networking issues on Windows
- Andrew Collins – For sharing findings from his perf investigations with the community.
- Aravindh Puthiyaparambil - For helping with bug triage, backlog grooming, docs, and the node service log viewer.
- And all our community members!

How To Contribute

- Visit our [SIG-Windows community page](#)
- Follow our [contributing guide](#)
- Join our [community meetings](#) (12:30pm EST every tuesday)
- Join our SIG-Windows pairing sessions: <https://github.com/kubernetes-sigs/sig-windows-samples/blob/master/PAIRING.md> (right after the main meeting)
- Help us write additional documentation and user stories
- File bugs or review open PRs from our project boards - [Issues](#) - [PRs](#)


Community Resources

Mark Rossetti

Co-Chair
Microsoft
@Mark Rossetti [Slack]
@marosset [GitHub]




#sig-windows



<https://groups.google.com/forum/#!forum/kubernetes-sig-windows>

Jay Vyas

Technical Lead
VMWare
@jayunit100 [Slack / GitHub]



<https://kubernetes.io/docs/setup/production-environment/windows/>



<https://github.com/kubernetes/community/tree/master/sig-windows>

James Sturtevant

Technical Lead
Microsoft
@jsturtevant [Slack / GitHub]

Claudiu Belu

Technical Lead
Cloudbase Solutions
@Claudiu Belu [Slack]
@claudiubelu [Github]



[YouTube Playlist](#)



<https://zoom.us/j/94389601840>
Every Tuesday @ 12.30pm EST

Q & A



Please scan the QR Code above to
leave feedback on this session