



Arbitrary Code & File Execution in R/O FS – Am I Write?



Golan Myers

- Security Consultant @ WithSecure
- Focus on low-level aspect of container security
- Enthusiastic about all things containers
- Containers, containers, containers, containers!!!!!!!



Agenda

- What are RO/FS
- Why are they used
- Kubernetes and RO/FS
- An attacker's perspective
- 3 methods to bypass R/O restrictions
- Remediation and mitigation
- Final thoughts and conclusions




What Is A R/O FS

Why Should We Use It?

- **Generally** 3 main “actions” – Read/Write/Execute
- **R/O FS** - Read/Execute
- **Security-wise** – better control/mgmt of containerised applications

```
apiVersion: v1
kind: Pod
metadata:
  name: pod
spec:
  containers:
  - name: container
    image: alpine
    securityContext:
      readOnlyRootFileSystem: true
```





I'm in, but it's R/O

What do I
do now?

A Threat Actor's Approach To a R/O Pod



- Foothold in the environment
- Enumeration from an internal perspective
- Intermediary

Attack Scenario

- Adversary with a foothold on an application pod with a R/O FS
- Execution of arbitrary code and executables vital to further the attack
- Low Privileged



```
apiVersion: v1
kind: Pod
metadata:
  name: method1-pod
spec:
  containers:
  - name: nginx
    image: nginx:latest
    securityContext:
      readOnlyRootFilesystem: true
      runAsUser: 101
    ports:
    - containerPort: 80
  volumeMounts:
  - mountPath: /var/run
    name: run
  - mountPath: /var/cache/nginx
    name: nginx-cache
  securityContext:
    seccompProfile:
      type: RuntimeDefault
  volumes:
  - name: run
    emptyDir: {}
  - name: nginx-cache
    emptyDir: {}
```

Method #1

- R/O FS
- Nginx image
- Contains Bash
- No standard network tools to retrieve data (wget, curl, etc..)

DEMO

nginx@method1-pod:/\$

newkalier@demokali: ~

newkalier@demokali: ~ 71x15

(newkalier@demokali)-[~]

\$ serve

Serving HTTP on 0.0.0.0 port 8080 (http://0.0.0.0:8080/) ...


newkalier@demokali: ~ 71x15

(newkalier@demokali)-[~]

\$ listen

listening on [any] 9898 ...

What we saw




```
apiVersion: v1
kind: Pod
metadata:
  name: method2-pod
spec:
  containers:
    - args:
      - alpine
      name: alpine
      image: alpine
      command:
        - "sleep"
        - "3600"
      securityContext:
        readOnlyRootFilesystem: true
        runAsUser: 65534
    securityContext:
      seccompProfile:
        type: RuntimeDefault
```

Method #2

- R/O FS
- Alpine image
- Contains sh
- No standard network tools to retrieve data (wget*, curl, etc..)

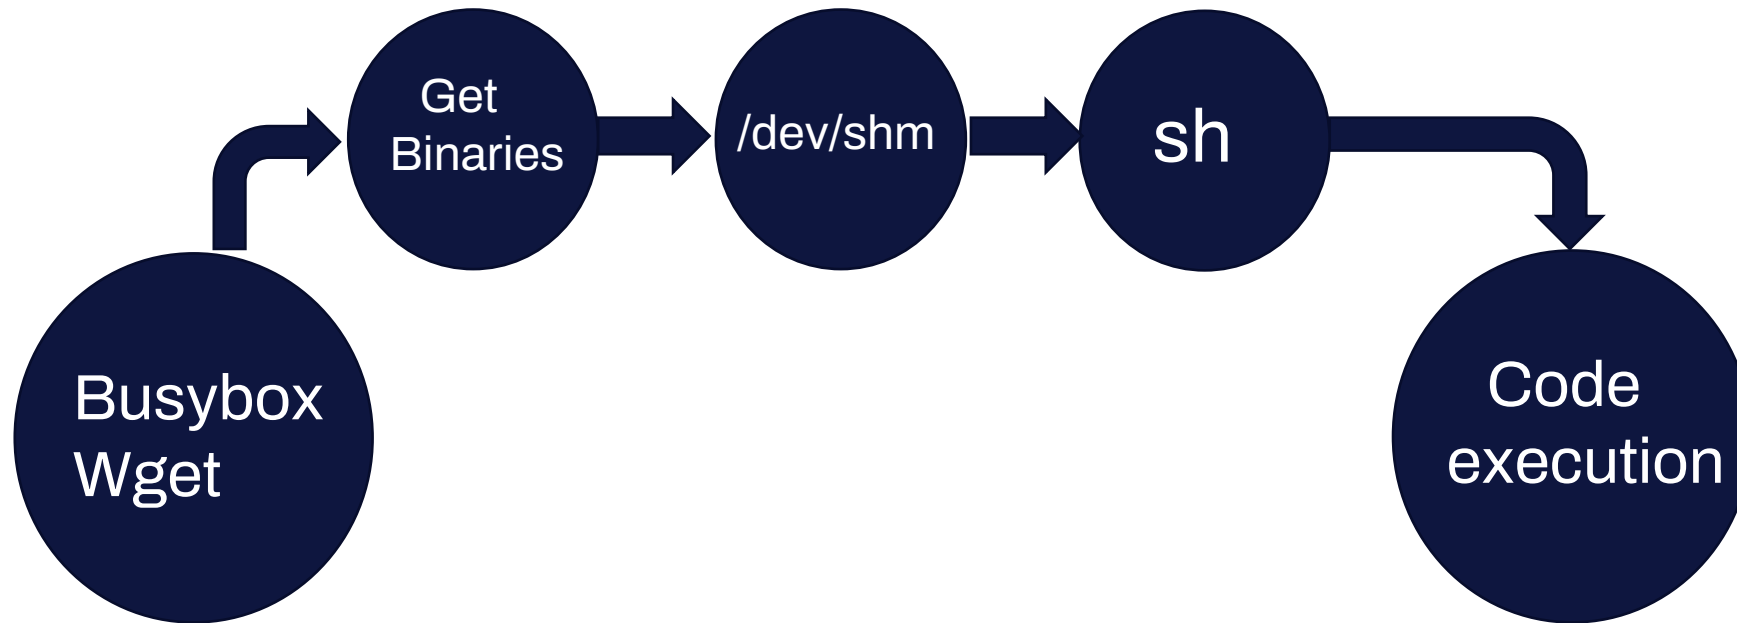
DEMO

PgUp



```
(newkalier@demokali)-[~]  
$ listen  
listening on [any] 9898 ...
```


What we saw



```
apiVersion: v1
kind: Pod
metadata:
  name: method2-pod
spec:
  containers:
  - args:
    - alpine
    name: alpine
    image: alpine
    command:
      - "sleep"
      - "3600"
    securityContext:
      readOnlyRootFilesystem: true
      runAsUser: 65534
  securityContext:
    seccompProfile:
      type: RuntimeDefault
```

Method #3

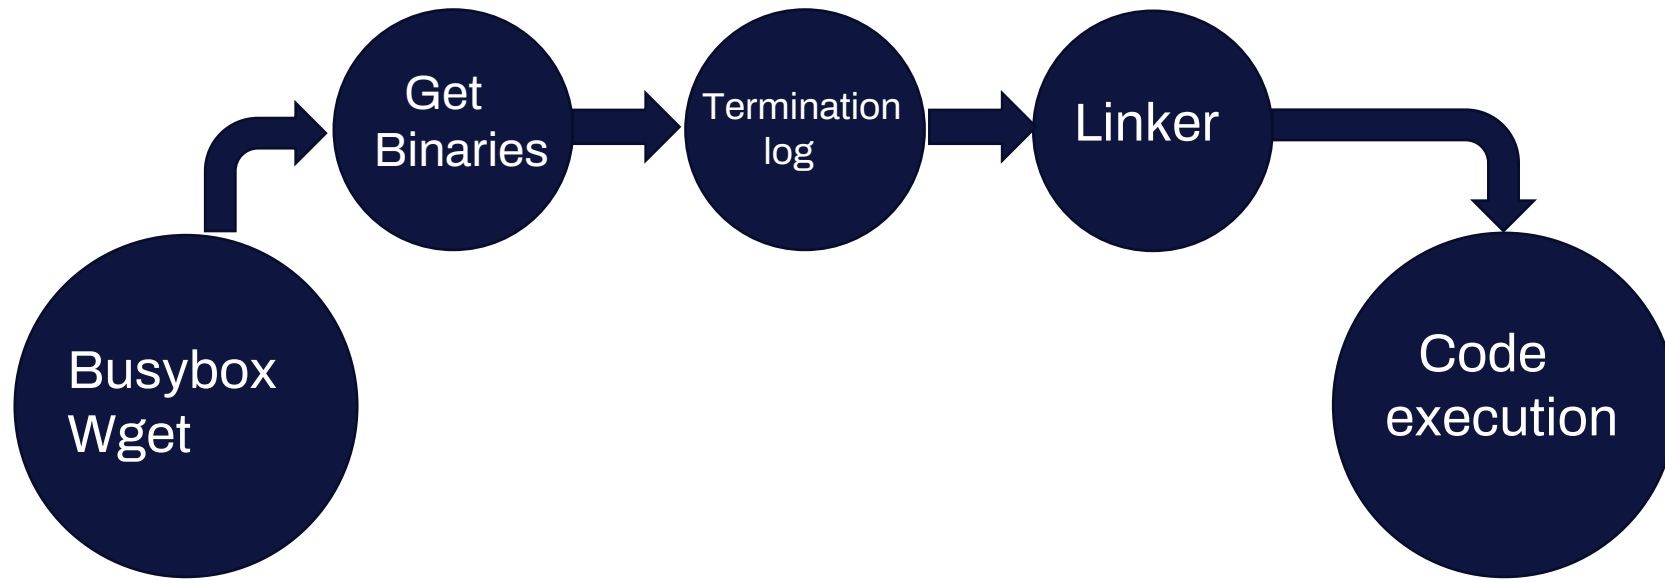
- R/O FS
- Alpine image
- Contains sh
- No standard network tools to retrieve data (wget*, curl, etc..)

DEMO


```
(newkalier@demokali)-[~]
```

```
$
```

What we saw



Detecting File Execution From /dev/shm

- rule: Execution from /dev/shm

desc: This rule detects file execution from the /dev/shm directory, a common tactic for threat actors to stash their readable+writable+(sometimes)executable files.

condition: >

spawned_process and
(proc.exe startswith "/dev/shm/" or
(proc.cwd startswith "/dev/shm/" and proc.exe startswith "./") or
(shell_procs and proc.args startswith "-c /dev/shm") or
(shell_procs and proc.args startswith "-i /dev/shm") or
(shell_procs and proc.args startswith "/dev/shm") or
(proc.args contains "/dev/shm" or proc.cwd startswith
"/dev/shm") or
(proc.cwd startswith "/dev/shm/" and proc.args startswith "./"))
and

not container.image.repository in (falco_privileged_images,
trusted_images)

output: "File execution detected from /dev/shm
(proc.cmdline=%proc.cmdline connection=%fd.name
user.name=%user.name user.loginuid=%user.loginuid
container.id=%container.id evt.type=%evt.type evt.res=%evt.res
proc.pid=%proc.pid proc.cwd=%proc.cwd proc.ppid=%proc.ppid
proc.pcmdline=%proc.pcmdline proc.sid=%proc.sid
proc.exepath=%proc.exepath user.uid=%user.uid
user.loginname=%user.loginname group.gid=%group.gid
group.name=%group.name container.name=%container.name
image=%container.image.repository)"
priority: WARNING

Defending against these attacks

- **Seccomp** – Docker default runtime profile/custom profile etc
- **SELinux** – prevent access to execmem
- **Detection** – unexpected network connections (all methods), execution of files that reside /dev/shm (method 2)
- **Noexec** – specify termination-log location

Rapping Up

- R/O FS for container IS recommended
- Just R/O FS \neq Security



</dev/audience



Golan Myers:

<https://labs.withsecure.com/publications/executing-arbitrary-code-executables-in-read-only-filesystems>

